AWS Certified Developer - Associate

By Cloud Mentor Pro

Updated: 15/02/2026

© Copyright by Cloud Mentor Pro | Confidential

Table of Contents

- Section 1
  - *Getting started with AWS*
  - *IAM & AWS CLI*
  - *EC2 Fundamentals*
- Section 2
  - *EC2 Instance Storage*
  - *High Availability & Scalability: ELB & ASG*
- Section 3
  - *RDS, Aurora & ElastiCache*
  - *Route53*
- *Section 4*
  - *Amazon S3*
  - *Amazon S3 – Advanced*
  - *Amazon S3 – Security*

© Copyright by Cloud Mentor Pro | Confidential

Table of Contents

- Section 5
  - *Amazon VPC*
- Section 6
  - *ECS, ECR & Fargate - Docker in AWS*
- Section 7
  - *CloudFront*
  - *AWS Elastic Beanstalk*
- *Section 8*
  - *AWS CloudFormation*
- Section 9
  - *AWS Integration & Messaging: SQS, SNS & Kinesis*

© Copyright by Cloud Mentor Pro | Confidential

Table of Contents

- Section 10
  - *AWS Monitoring, Troubleshooting & Audit*
- Section 11
  - *AWS Lambda*
- Section 12
  - *AWS DynamoDB*
- Section 13
  - *API Gateway*

© Copyright by Cloud Mentor Pro | Confidential

Table of Contents

- Section 14
  - *AWS Serverless: SAM - Serverless Application Model*
  - *Cloud Development Kit (CDK)*
  - *Cognito: Cognito User Pools, Cognito Identity Pools & Cognito Sync*
  - *Other Serverless: Step Functions & AppSync*
- Section 15
  - *AWS CICD: CodeCommit, CodePipeline, CodeBuild, CodeDeploy*
- Section 16
  - *Advanced Identity*
  - *AWS Security & Encryption: KMS, SSM Parameter Store, IAM & STS*
  - *AWS Other Services*

© Copyright by Cloud Mentor Pro | Confidential

What we'll learn in this course

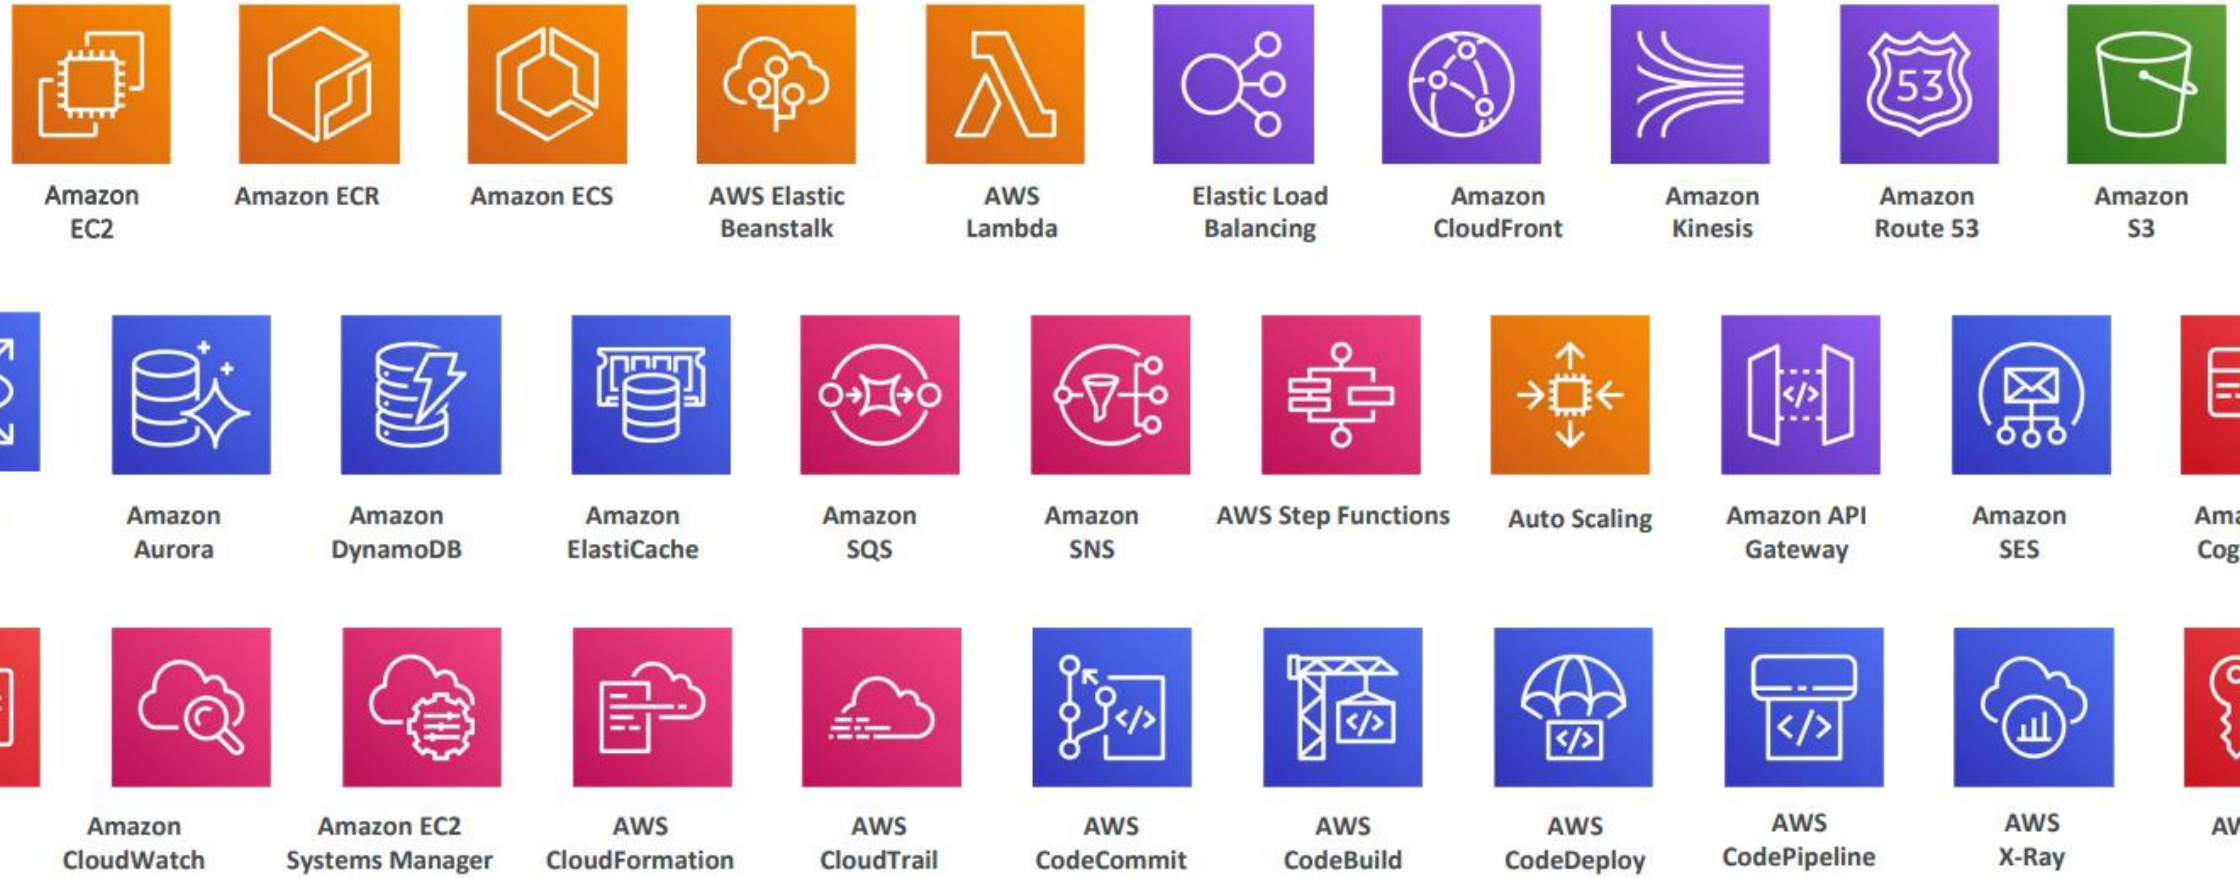

© Copyright by Cloud Mentor Pro | Confidential

Section 1

- *Getting started with AWS*
- *IAM & AWS CLI*
- *EC2 Fundamentals*

© Copyright by Cloud Mentor Pro | Confidential

Getting started with AWS

AWS Global Infrastructure

© Copyright by Cloud Mentor Pro | Confidential

AWS - The leading cloud

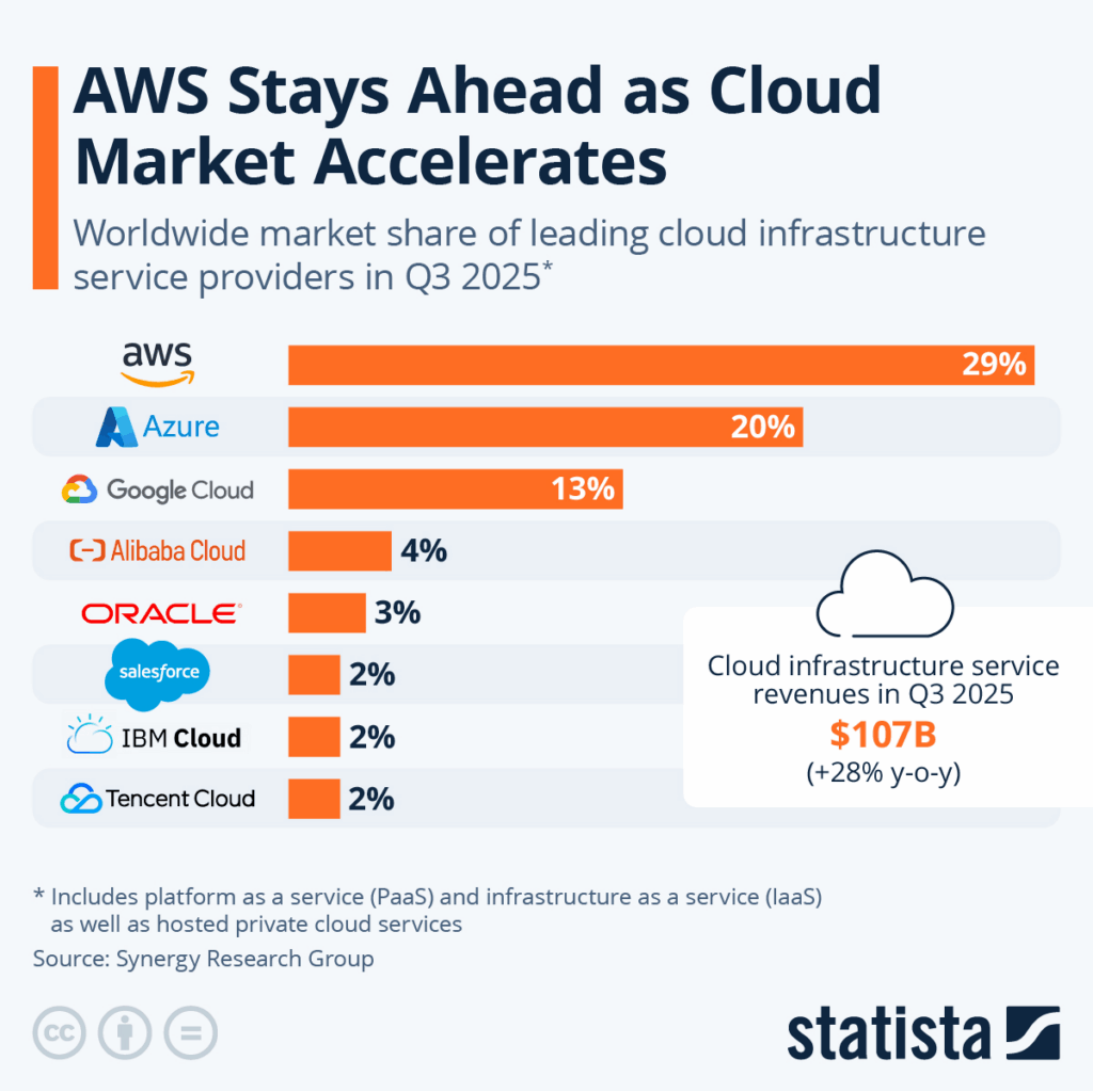

© Copyright by Cloud Mentor Pro | Confidential

AWS Global Infrastructure

- AWS Regions
- AWS Availability Zone
- AWS Data Centers
- AWS Edge Locations/ Points of Presence

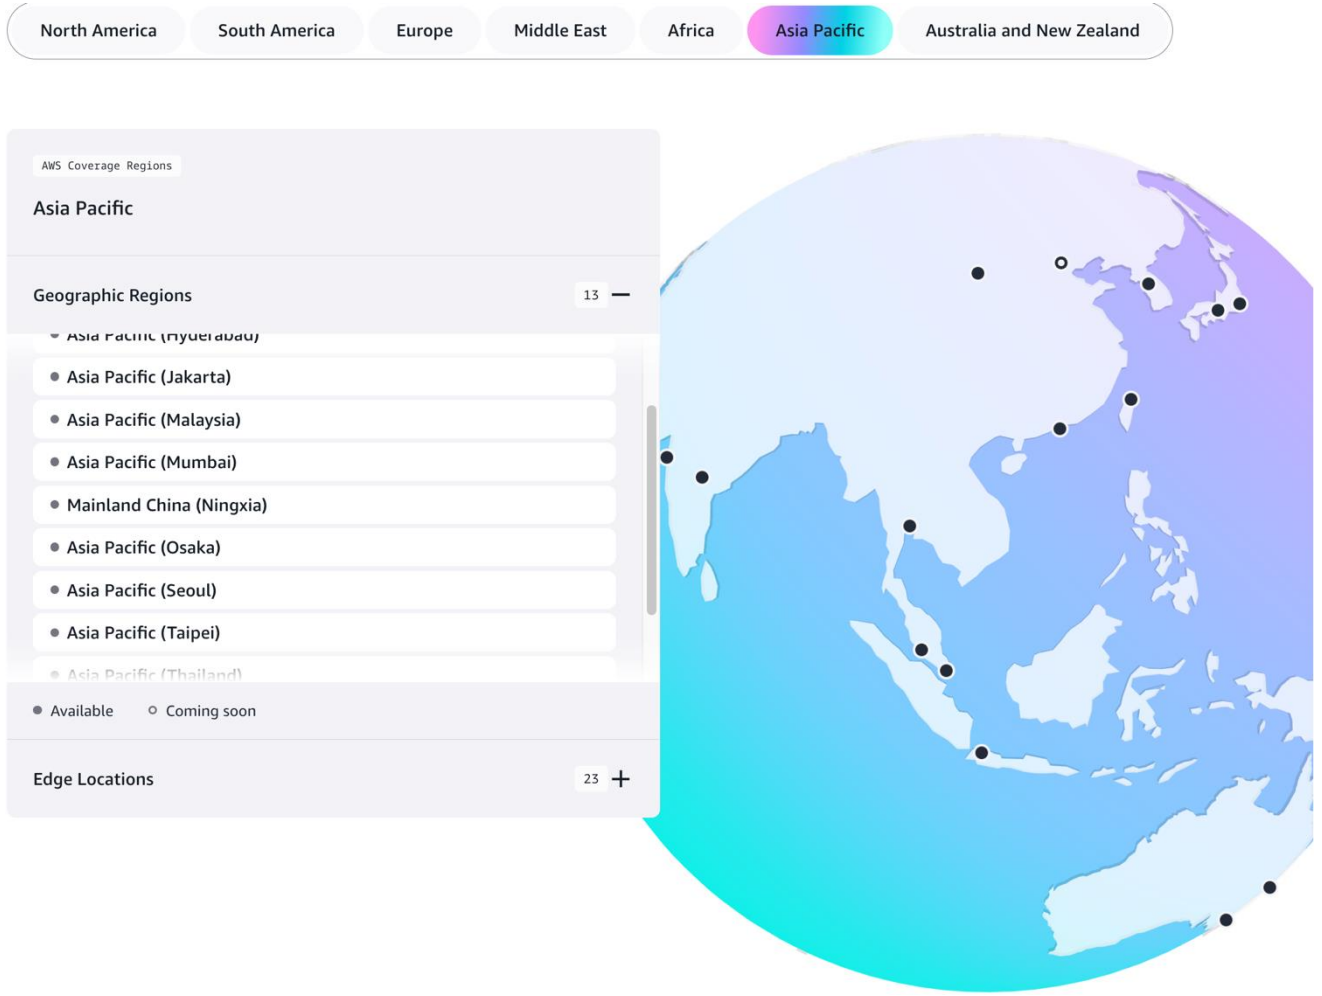

https://aws.amazon.com/about-aws/global-infrastructure/regions_az/

© Copyright by Cloud Mentor Pro | Confidential

AWS Regions

- A region is a cluster of data centers
- Choose an AWS Region
  - Compliance with data governance and legal requirements: data never leaves a region without your explicit permission
  - Proximity to customers: reduced latency
  - Available services within a Region: new services and new features aren't available in every Region
  - Pricing: pricing varies region to region and is transparent in the service pricing page

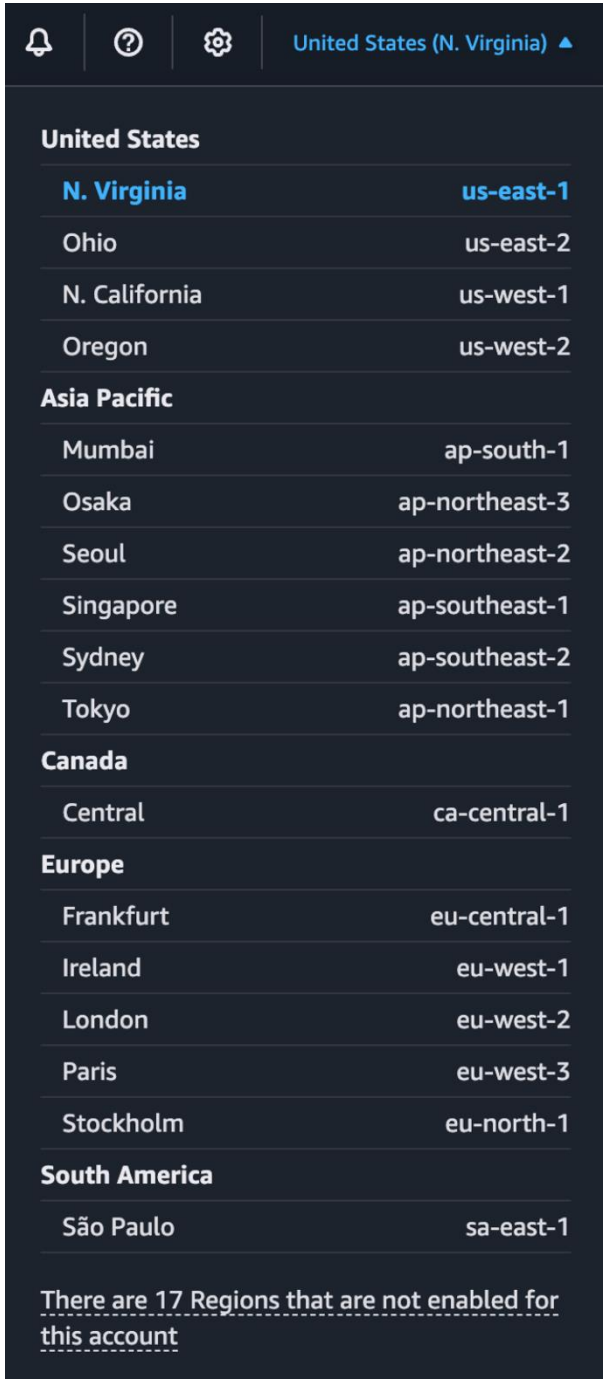

© Copyright by Cloud Mentor Pro | Confidential

AWS Availability Zones

- Each region has many availability zones (usually 3, min is 3, max is 6). Example:
  - ap-southeast-2a
  - ap-southeast-2b
  - ap-southeast-2c
- Each availability zone (AZ) is one or more discrete data centers with redundant power, networking, connectivity
- They're separate from each other, so that they're isolated from disasters
- They're connected with high bandwidth, ultra-low latency networking

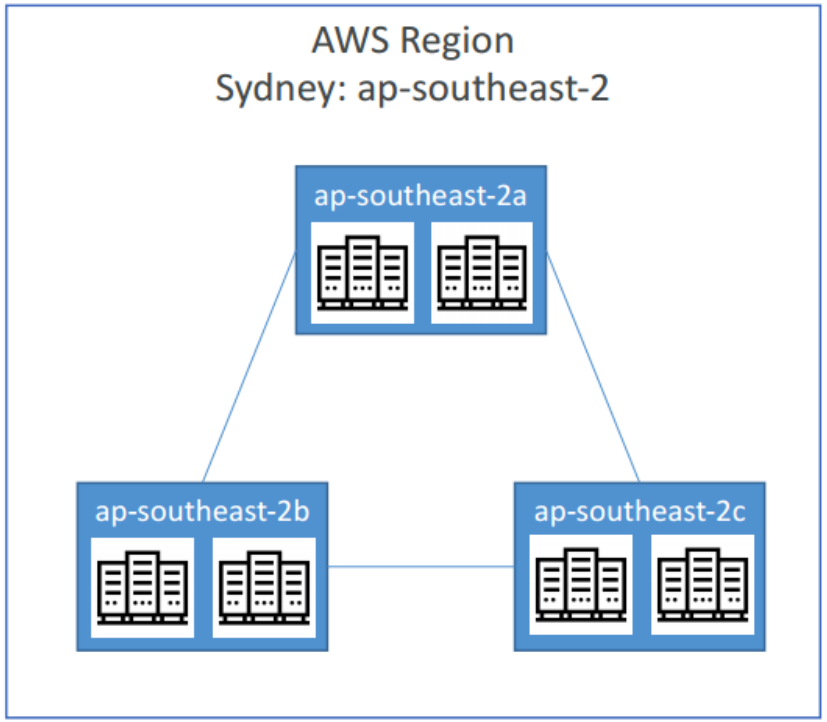

© Copyright by Cloud Mentor Pro | Confidential

AWS Availability Zones

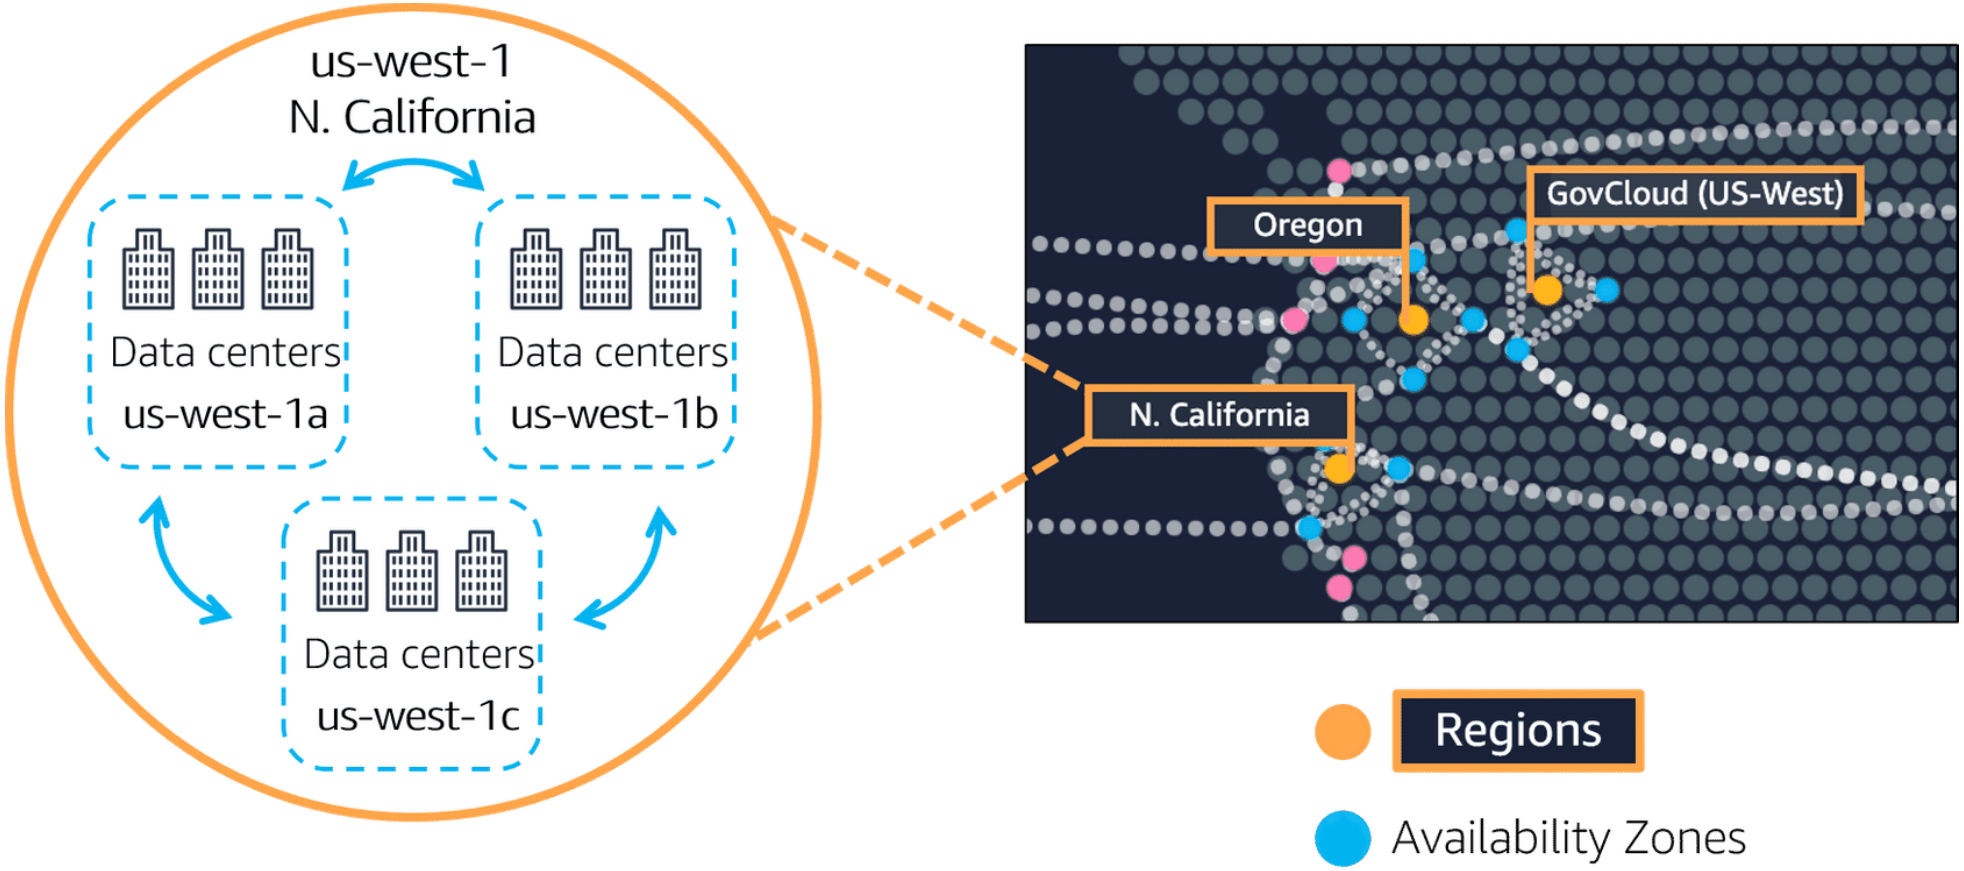

© Copyright by Cloud Mentor Pro | Confidential

AWS Availability Zones

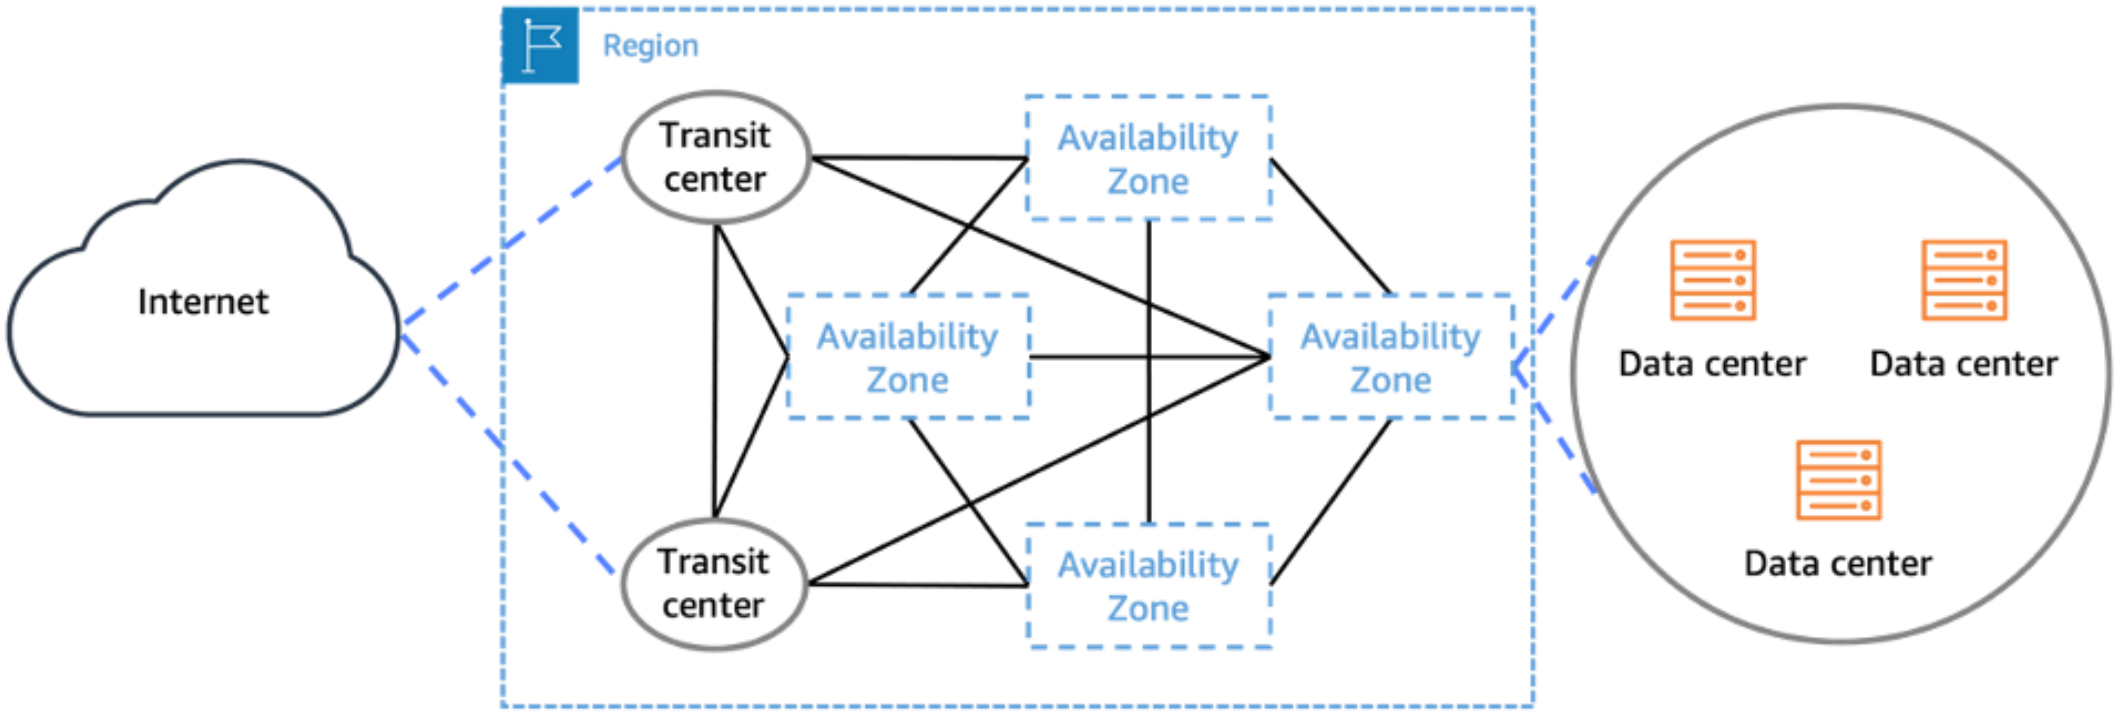

© Copyright by Cloud Mentor Pro | Confidential

AWS Edge Locations (Points of Presence)

- The global edge network currently including 400+ Edge Locations, and 13 Regional Caches in 90+ cities across 48+ countries

- Content is delivered to end users with **lower latency**

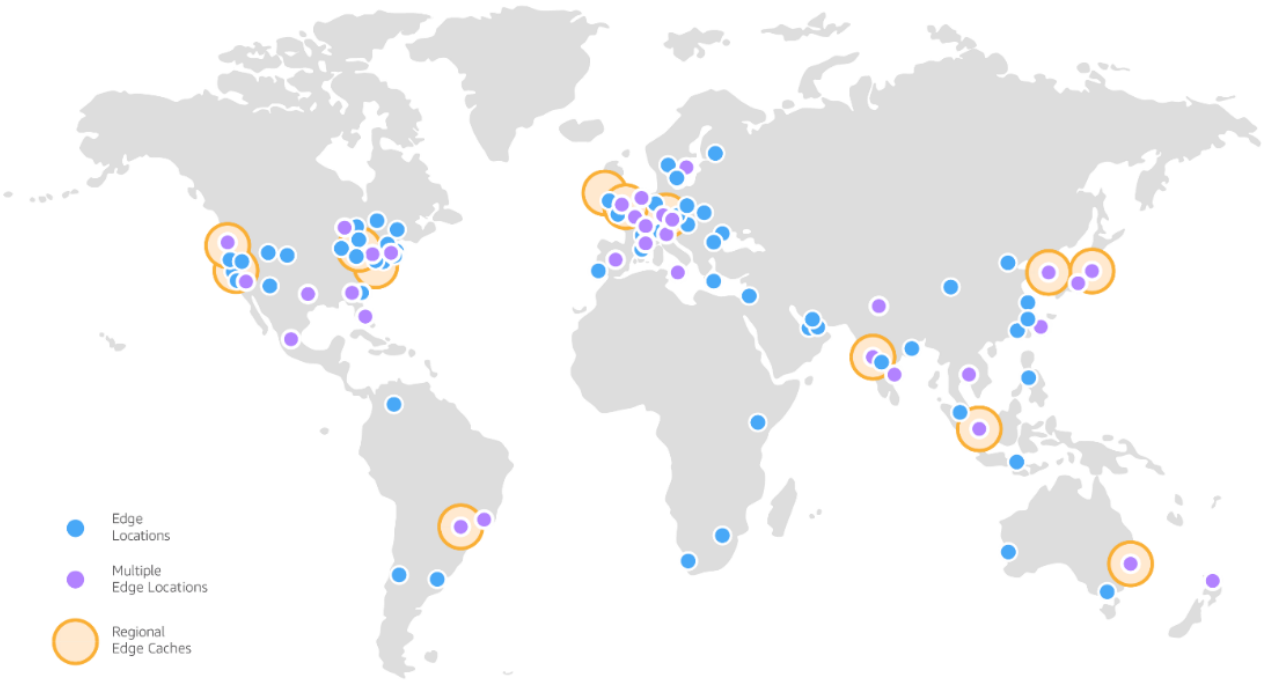

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/whitepapers/latest/aws-fault-isolation-boundaries/points-of-presence.html

AWS Edge Locations (Points of Presence)

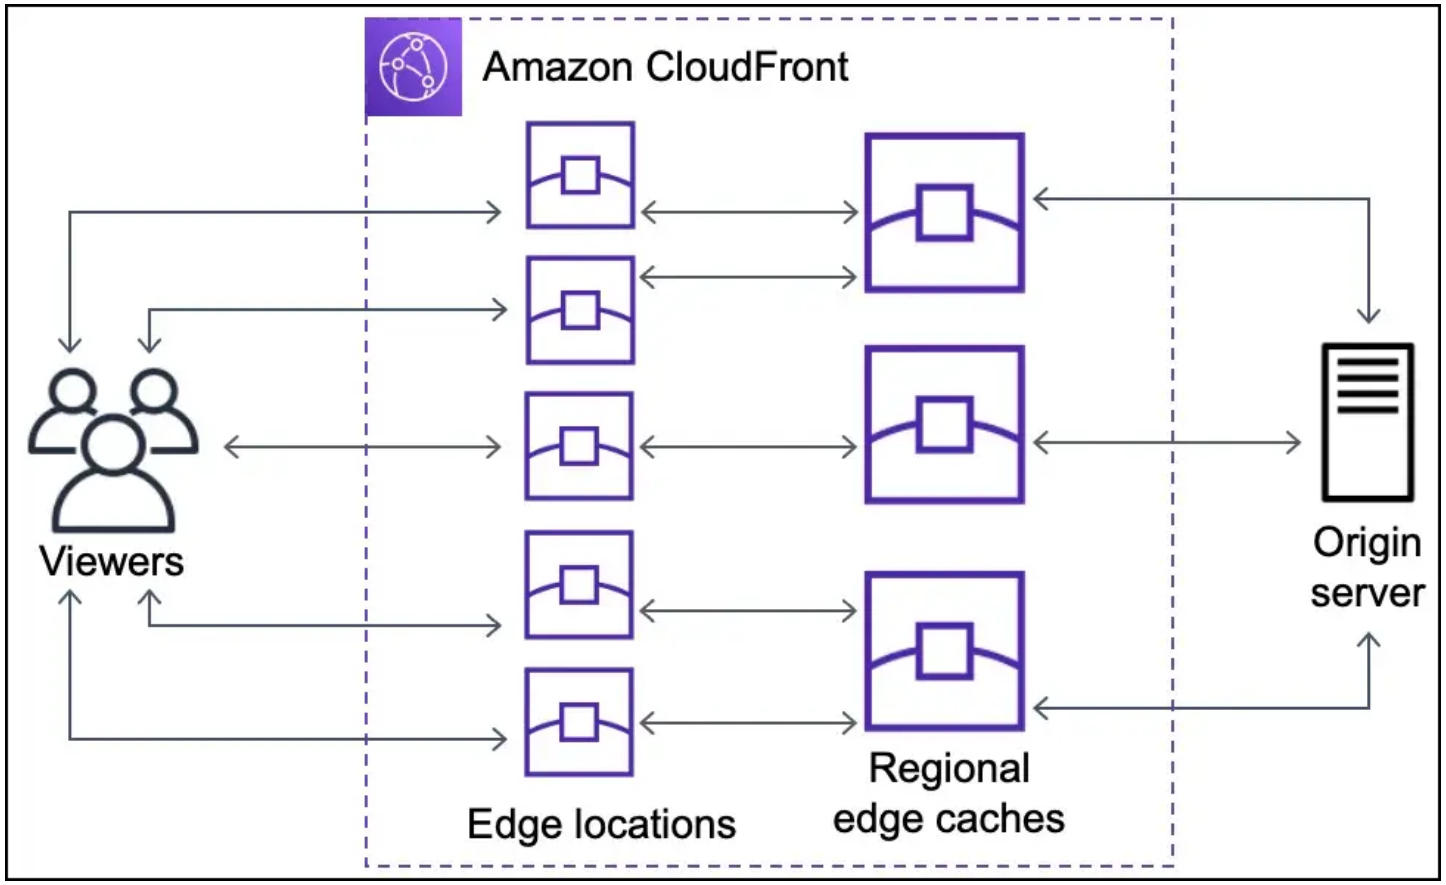

© Copyright by Cloud Mentor Pro | Confidential

IAM & AWS CLI

AWS Identity and Access Management (AWS IAM)

© Copyright by Cloud Mentor Pro | Confidential

IAM: Users & Groups

- IAM = Identity and Access Management, **Global** service
- **Root account** created by default, shouldn't be used or shared
- **Users** are people within your organization, and can be grouped
- **Groups** only contain users, not other groups
- Users don't have to belong to a group, and user can belong to multiple groups

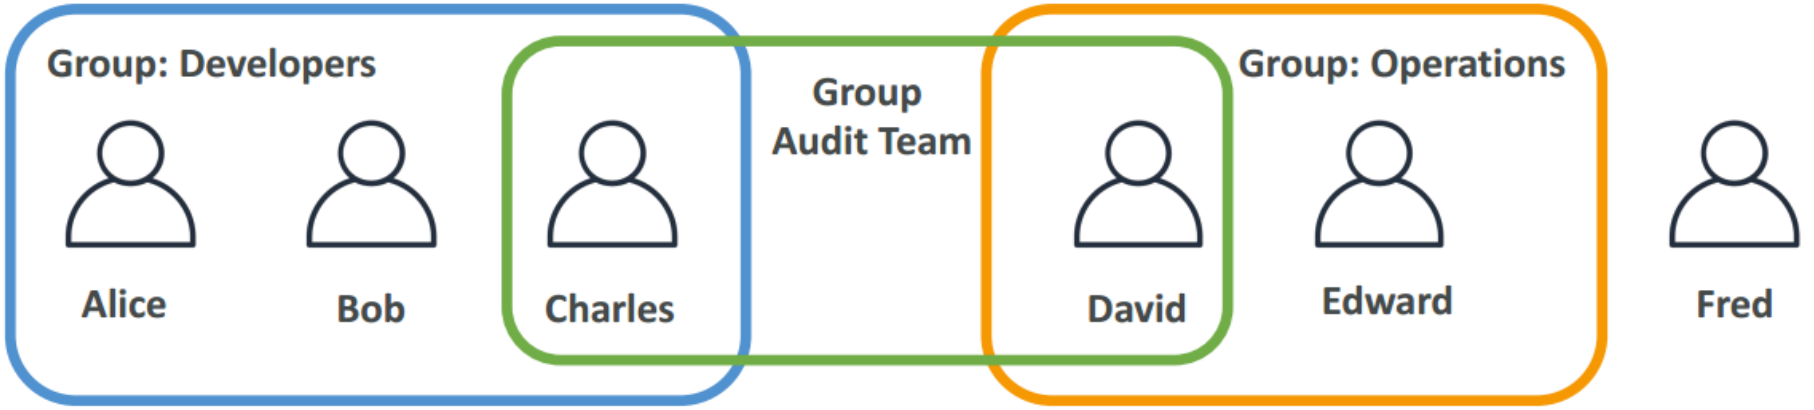

© Copyright by Cloud Mentor Pro | Confidential

IAM: Permissions

- Users or Groups can be assigned JSON documents called policies
- These policies define the permissions of the users
- In AWS you apply the least privilege principle: don't give more permissions than a user needs

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "ec2:Describe*",
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": "elasticloadbalancing:Describe*",
            "Resource": "*"
        },
        {
            "Effect": "Allow",
            "Action": [
                "cloudwatch:ListMetrics",
                "cloudwatch:GetMetricStatistics",
                "cloudwatch:Describe*"
            ],
            "Resource": "*"
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

IAM Policies inheritance

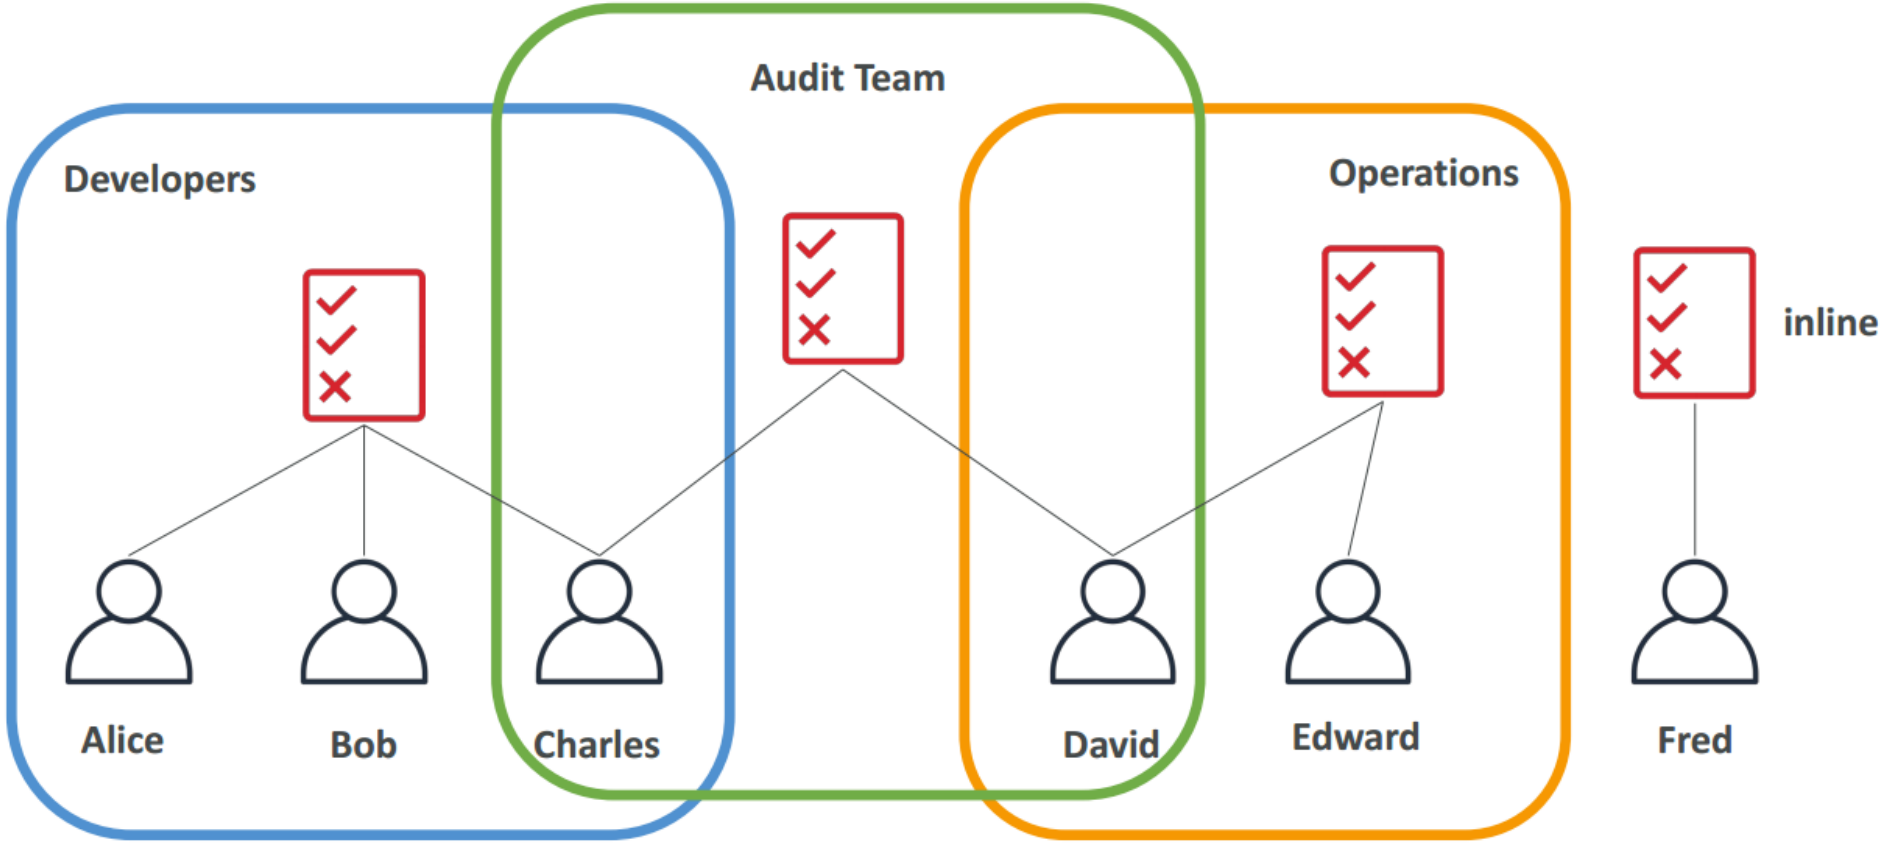

© Copyright by Cloud Mentor Pro | Confidential

IAM Policies Structure

- Consists of
  - **Version**: policy language version
  - **Id**: an identifier for the policy (optional)
  - **Statement**: one or more individual statements (required)
- Statements consists of
  - **Sid**: an identifier for the statement (optional)
  - **Effect**: Allow, Deny
  - **Principal**: account/user/role
  - **Action**: list of actions this policy allows or denies
  - **Resource**: list of resources to which the actions applied to
  - **Condition**: conditions for when this policy is in effect (optional)

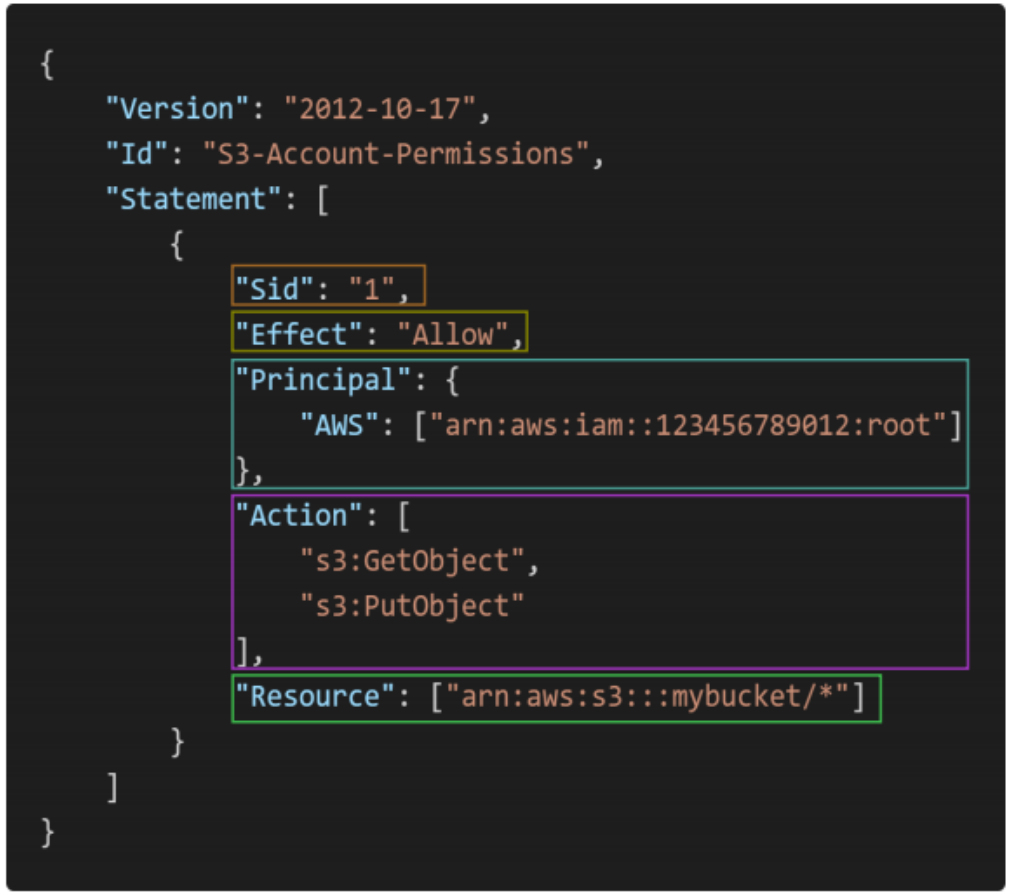

© Copyright by Cloud Mentor Pro | Confidential

IAM – Password Policy

- Strong passwords = higher security
- In AWS, you can setup a password policy:
  - Set a minimum password length
  - Require specific character types
  - Allow all IAM users to change their own passwords
  - Require users to change their password after some time (password expiration)
  - Prevent password re-use

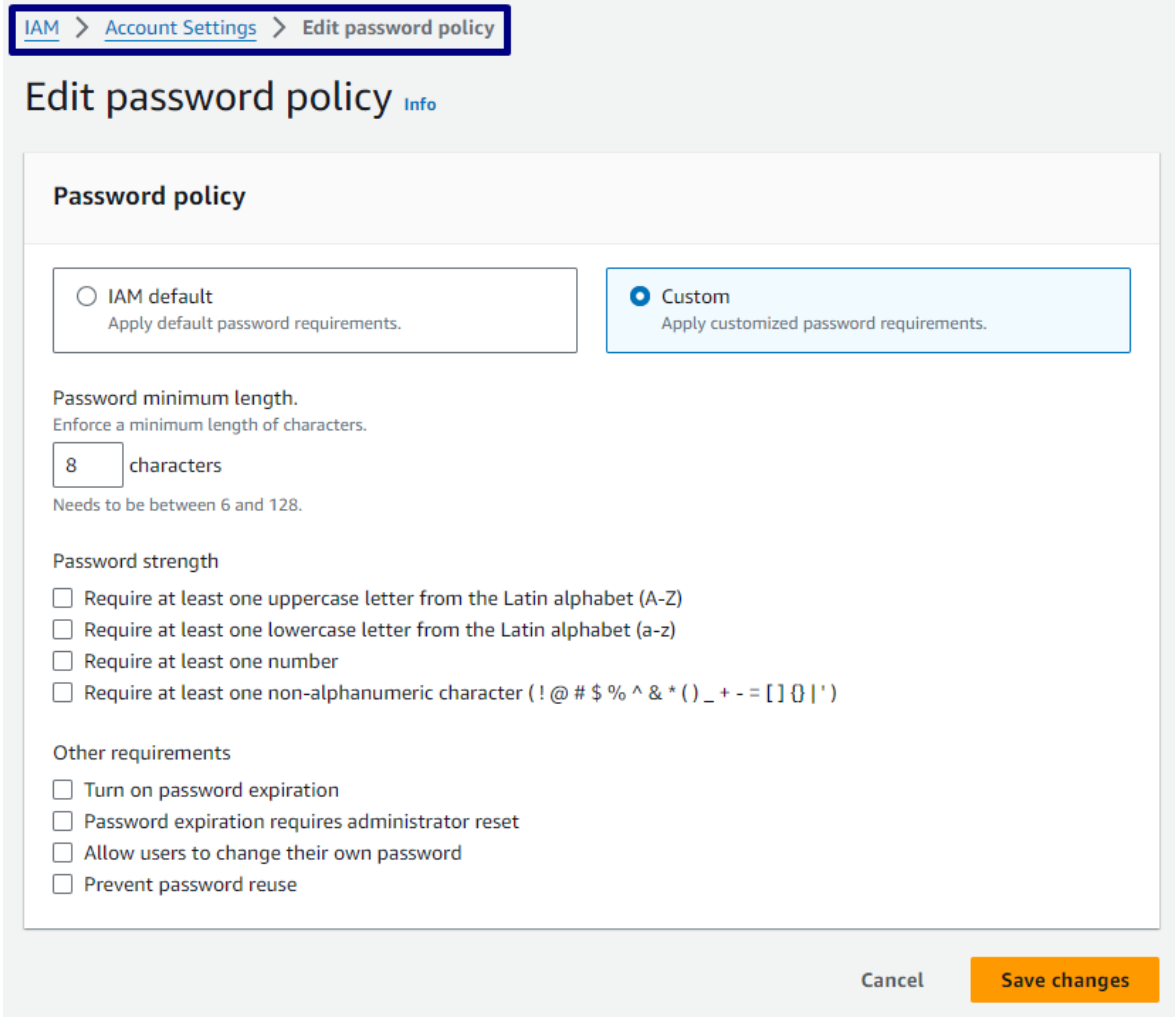

© Copyright by Cloud Mentor Pro | Confidential

Multi Factor Authentication - MFA

- Protect your Root Accounts and IAM users
- MFA = password you know + security device you own

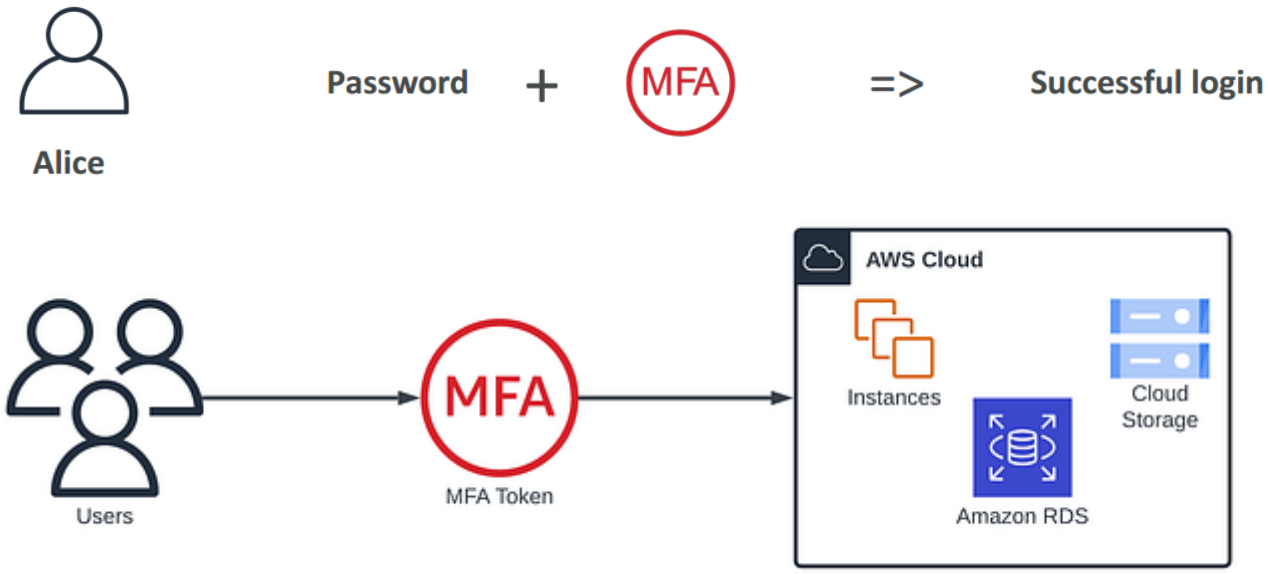

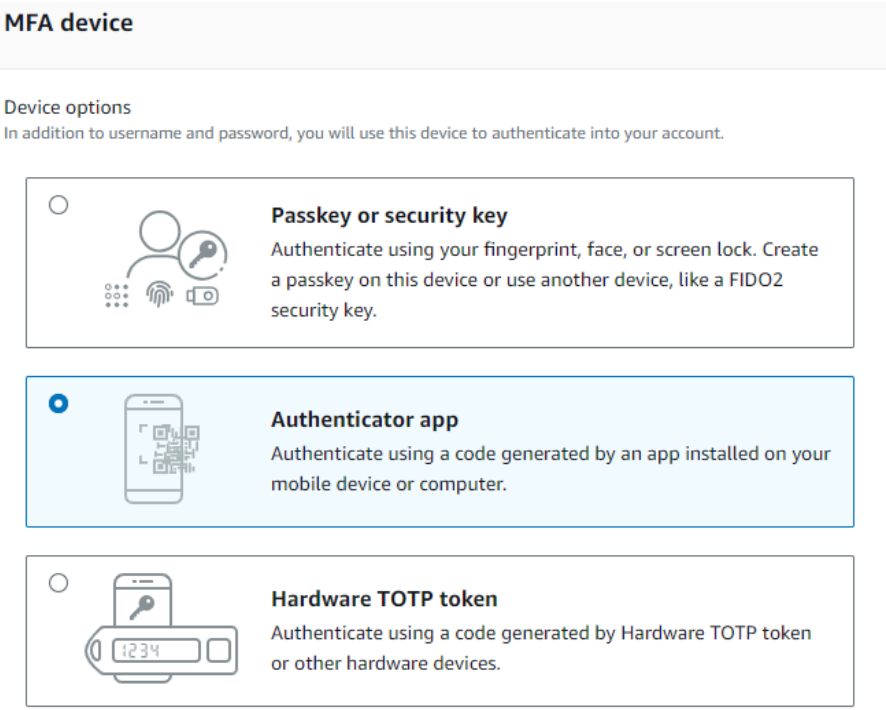

© Copyright by Cloud Mentor Pro | Confidential

MFA devices options in AWS

Virtual MFA device

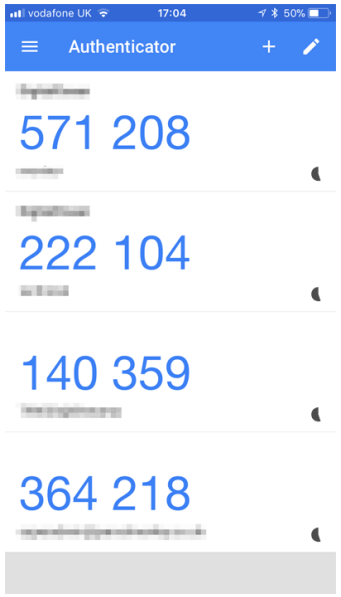

Google Authenticator
(phone only)

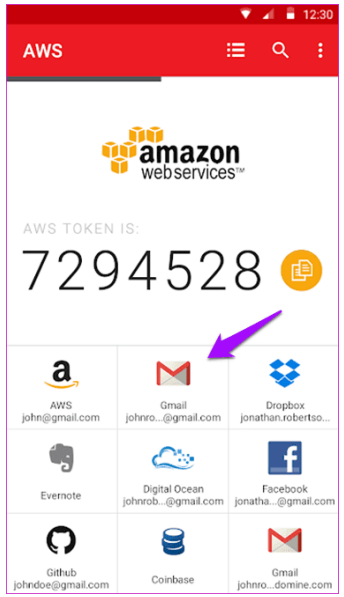

Authy
(phone only)

Universal 2nd Factor (U2F) Security Key

YubiKey by Yubico (3rd party)

© Copyright by Cloud Mentor Pro | Confidential

How can users access AWS ?

- To access AWS, you have three options:
  - AWS Management Console (protected by password + MFA)
  - AWS Command Line Interface (CLI): protected by access keys
  - AWS Software Developer Kit (SDK) - for code: protected by access keys

- Access Keys are generated through the AWS Console
  - Access Keys are secret, just like a password. Don't share them
  - Access Key ID ~= username
  - Secret Access Key ~= password

© Copyright by Cloud Mentor Pro | Confidential

What's the AWS CLI?

- A tool that enables you to interact with AWS services using commands in your command-line shell
- Direct access to the public APIs of AWS services
- It's open-source https://github.com/aws/aws-cli

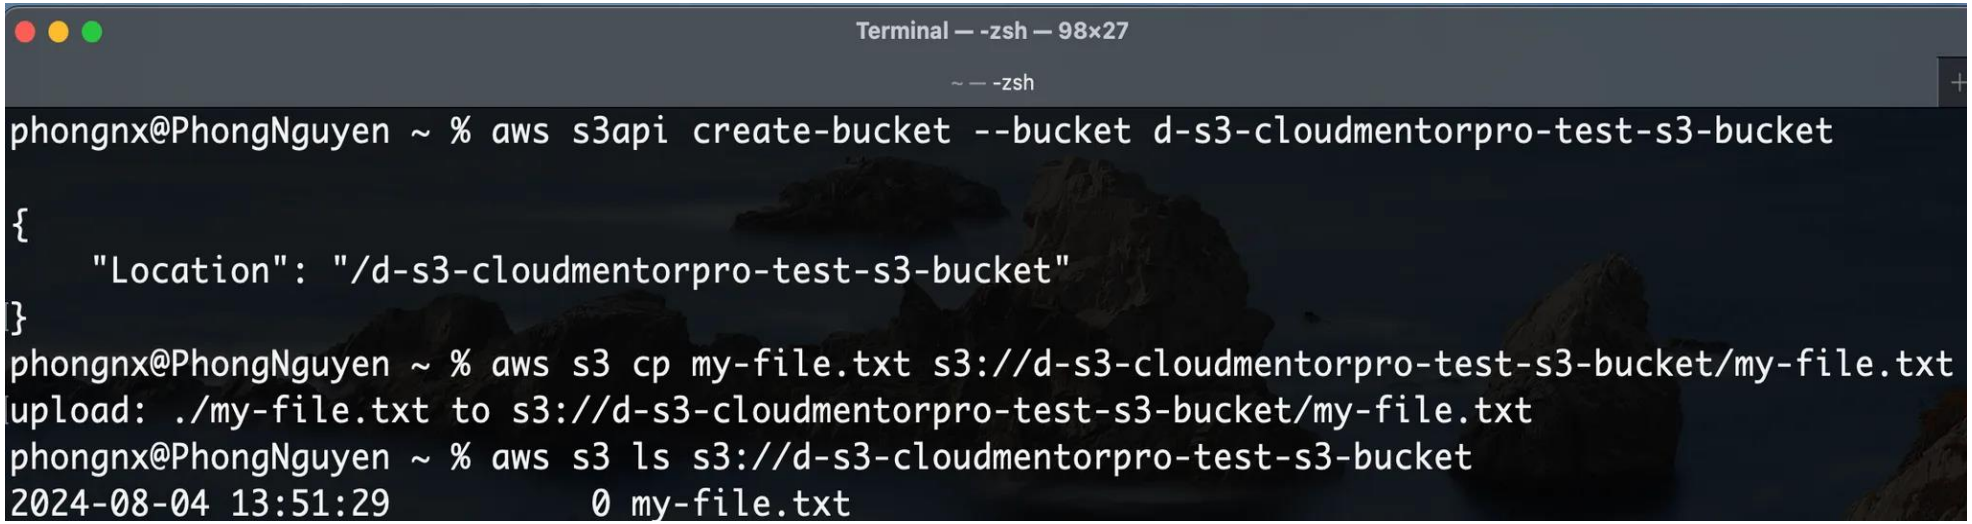

© Copyright by Cloud Mentor Pro | Confidential

AWS SDK

- AWS Software Development Kit (AWS SDK)
- Language-specific APIs (set of libraries)
- Enables you to access and manage AWS services programmatically
- Embedded within your application
- Supports
  - SDKs (JS, Python, PHP, .NET, Ruby, Java, Go, Node.js, , C++)
  - Mobile SDKs (Android, iOS, …)
  - IoT Device SDKs (Embedded C, Arduino, …)
- Example: AWS CLI is built on AWS SDK for Python

```python
import boto3

def lambda_handler(event, context):
    s3 = boto3.client('s3')
    response = s3.list_buckets()
    # your code
```

© Copyright by Cloud Mentor Pro | Confidential

IAM Roles for Services

- Some AWS service will need to perform actions on your behalf
- To do so, we will assign **permissions** to AWS services with **IAM Roles**

- **Common roles:**
  - EC2 Instance Roles
  - Lambda Function Roles
  - Roles for CloudFormation

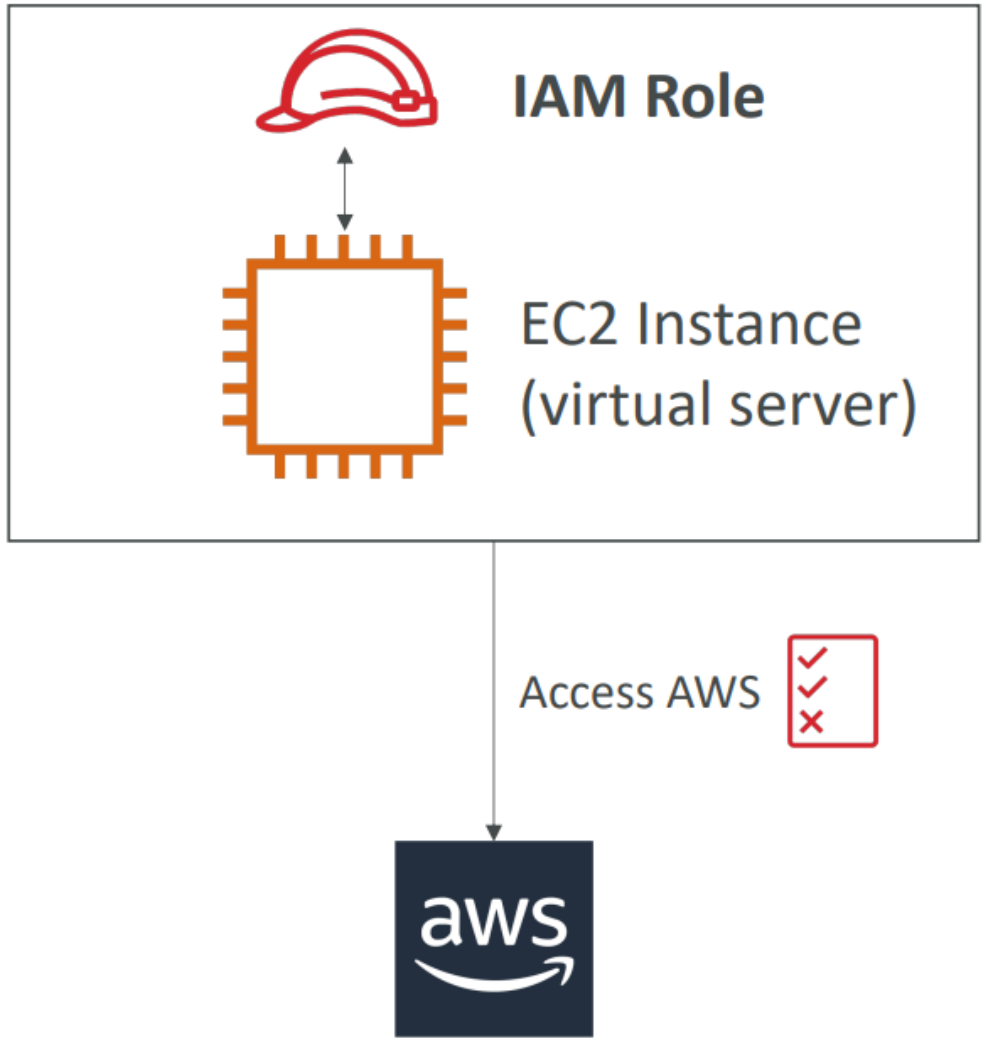

© Copyright by Cloud Mentor Pro | Confidential

IAM Security Tools

- IAM Credentials Report (account-level)
  - a report that lists all your account's users and the status of their various credentials

- IAM Access Advisor (user-level)
  - Access advisor shows the service permissions granted to a user and when those services were last accessed.
  - You can use this information to revise your policies.

© Copyright by Cloud Mentor Pro | Confidential

IAM Guidelines & Best Practices

- Don’t use the root account except for AWS account setup
- One physical user = One AWS user
- Assign users to groups and assign permissions to groups
- Create a strong password policy
- Use and enforce the use of Multi Factor Authentication (MFA)
- Create and use Roles for giving permissions to AWS services
- Use Access Keys for Programmatic Access (CLI / SDK)
- Audit permissions of your account using IAM Credentials Report & IAM Access Advisor
- **Never share IAM users & Access Keys**

© Copyright by Cloud Mentor Pro | Confidential

Shared Responsibility Model for IAM

- Infrastructure (global network security)
- Configuration and vulnerability analysis
- Compliance validation

You

- Users, Groups, Roles, Policies management and monitoring
- Enable MFA on all accounts
- Rotate all your keys often
- Use IAM tools to apply appropriate permissions
- Analyze access patterns & review permissions

© Copyright by Cloud Mentor Pro | Confidential

IAM Section – Summary

- **Users:** mapped to a physical user, has a password for AWS Console
- **Groups:** contains users only
- **Policies:** JSON document that outlines permissions for users or groups
- **Roles:** for EC2 instances or AWS services
- **Security:** MFA + Password Policy
- **AWS CLI:** manage your AWS services using the command-line
- **AWS SDK:** manage your AWS services using a programming language
- **Access Keys:** access AWS using the CLI or SDK
- **Audit:** IAM Credential Reports & IAM Access Advisor

© Copyright by Cloud Mentor Pro | Confidential

EC2 Fundamentals Part 1 (Basic)

Amazon EC2 – Basics

© Copyright by Cloud Mentor Pro | Confidential

Amazon EC2

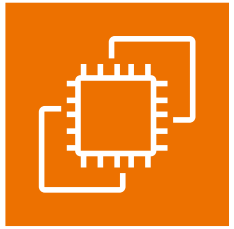

- EC2 is one of the most popular of AWS' offering
- EC2 = Elastic Compute Cloud = Infrastructure as a Service
- It mainly consists in the capability of :
  - Renting virtual machines (EC2)
  - Storing data on virtual drives (EBS)
  - Distributing load across machines (ELB)
  - Scaling the services using an auto-scaling group (ASG)

- Knowing EC2 is fundamental to understand how the Cloud works

© Copyright by Cloud Mentor Pro | Confidential

EC2 - Sizing & configuration options

- Operating System (**OS**): Linux, Windows or Mac OS
- How much compute power & cores (CPU)
- How much random-access memory (RAM)
- How much storage space:
  - Network-attached (**EBS & EFS**)
  - hardware (**EC2 Instance Store**)
- Network card: speed of the card, Public IP address
- Firewall rules: **security group**
- Bootstrap script (configure at first launch): EC2 User Data

© Copyright by Cloud Mentor Pro | Confidential

EC2 - User Data

- It is possible to bootstrap our instances using an EC2 User data script.
- bootstrapping means launching commands when a machine starts
- That script is only run once at the instance first start
- EC2 user data is used to automate boot tasks such as:
  - Installing updates
  - Installing software
  - Downloading common files from the internet
  - Anything you can think of
- The EC2 User Data Script runs with the root user

© Copyright by Cloud Mentor Pro | Confidential

Hands-On: Launching an EC2 Instance running Linux

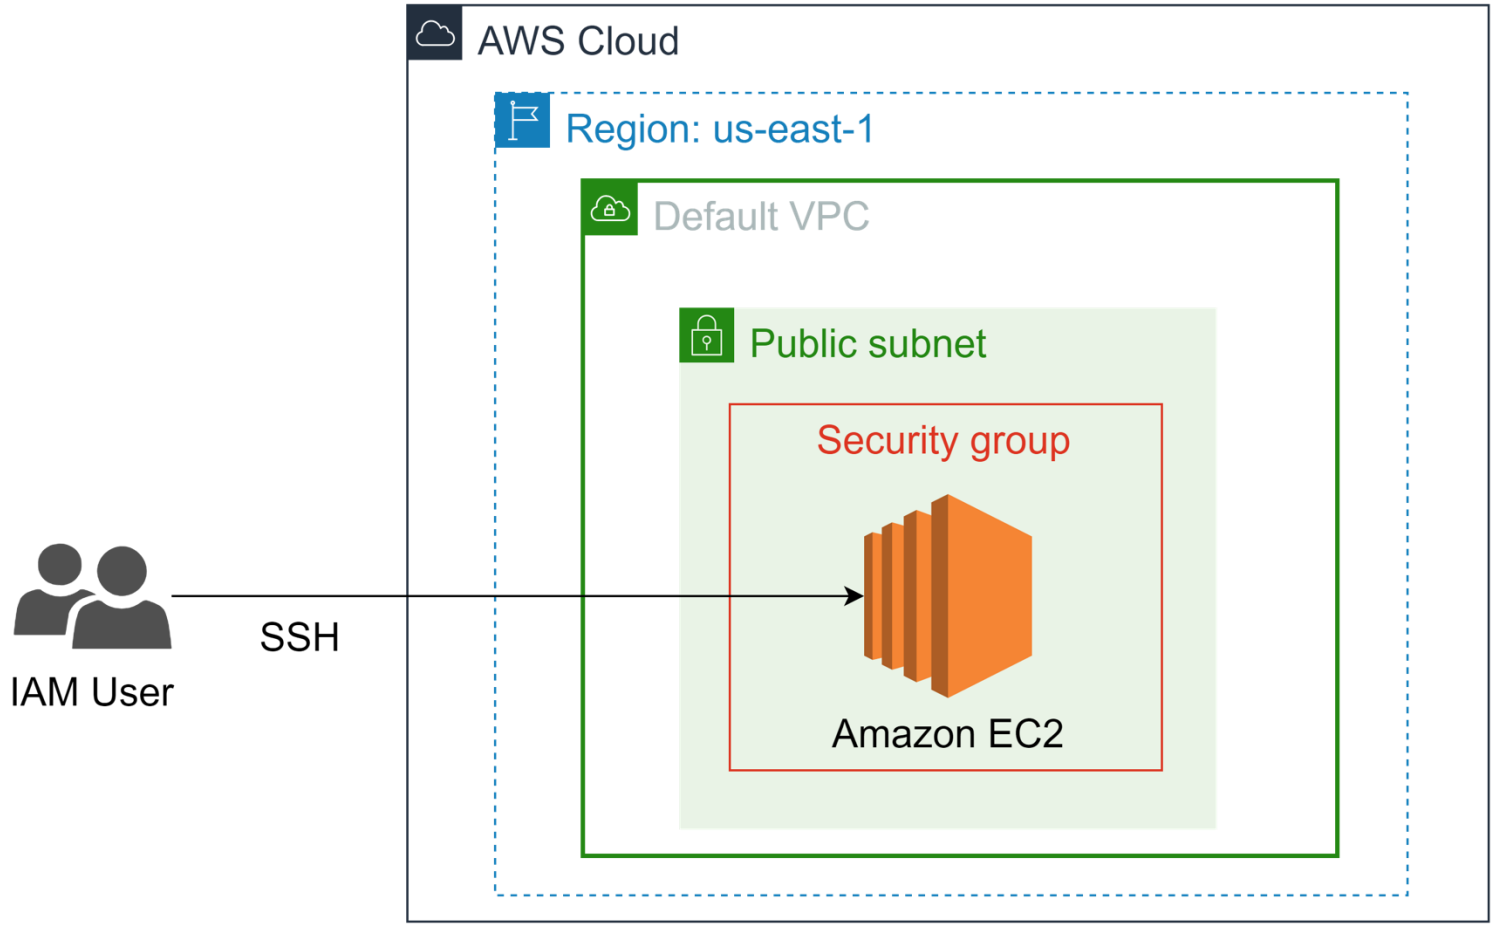

© Copyright by Cloud Mentor Pro | Confidential

EC2 - Instance Types

- You can use different types of EC2 instances that are optimised for different use cases (https://aws.amazon.com/ec2/instance-types/)
- AWS has the following naming convention:

m5.2xlarge

  - m: instance class
  - 5: generation (AWS improves them over time)
  - 2xlarge: size within the instance class

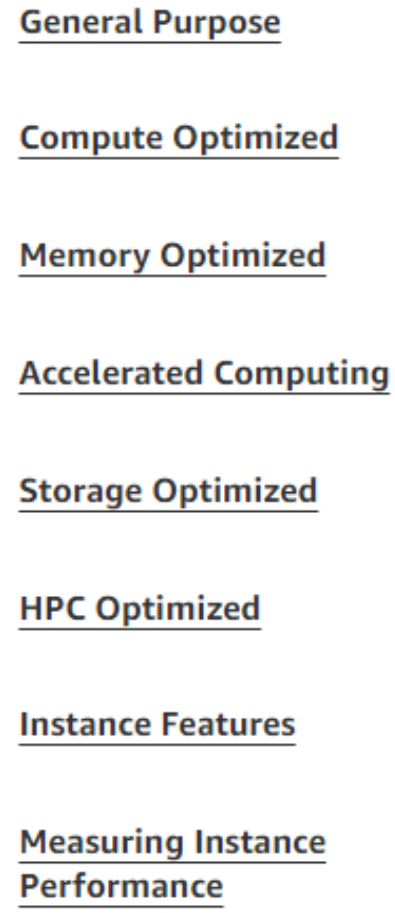

© Copyright by Cloud Mentor Pro | Confidential

EC2 Instance Types - General Purpose

- General purpose instances provide a balance of
  - Compute
  - Memory
  - Networking
- Great for
  - web servers
  - code repositories

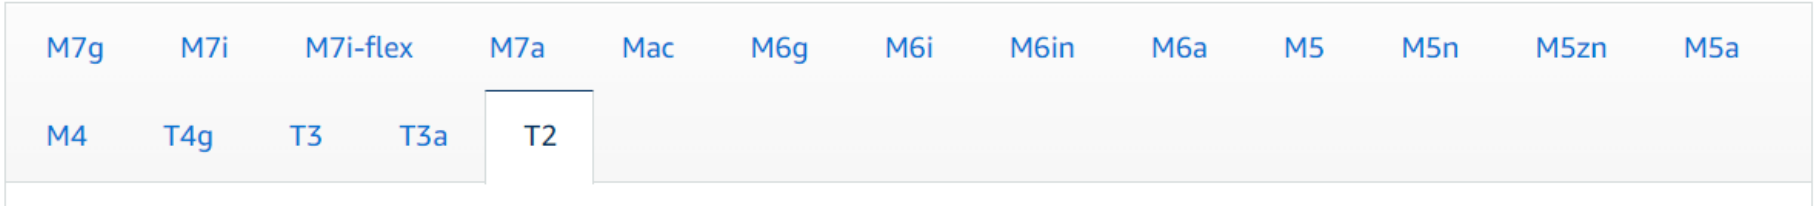

© Copyright by Cloud Mentor Pro | Confidential

https://aws.amazon.com/ec2/instance-types/

EC2 Instance Types - Compute Optimized

- Ideal for compute bound applications that benefit from high performance processors. **Great for:**
  - Compute intensive applications
  - Batch processing workloads
  - Media transcoding
  - High performance web servers, computing (HPC)
  - Scientific modeling, machine learning
  - Dedicated gaming servers and ad server engines

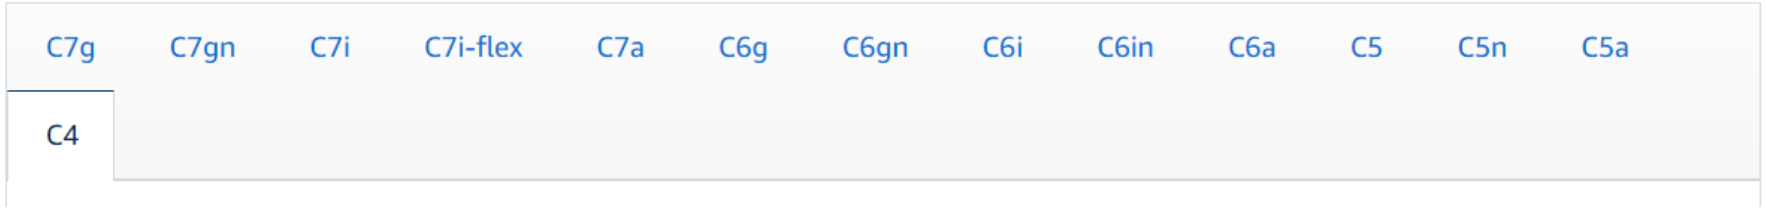

© Copyright by Cloud Mentor Pro | Confidential

https://aws.amazon.com/ec2/instance-types/

EC2 Instance Types - Memory Optimized

- Fast performance for workloads that process large data sets in memory.
- Great for
  - Relational/non-relational databases
  - Distributed web scale cache stores
  - In-memory databases optimized for BI (business intelligence)
  - Applications performing real-time processing of big unstructured data

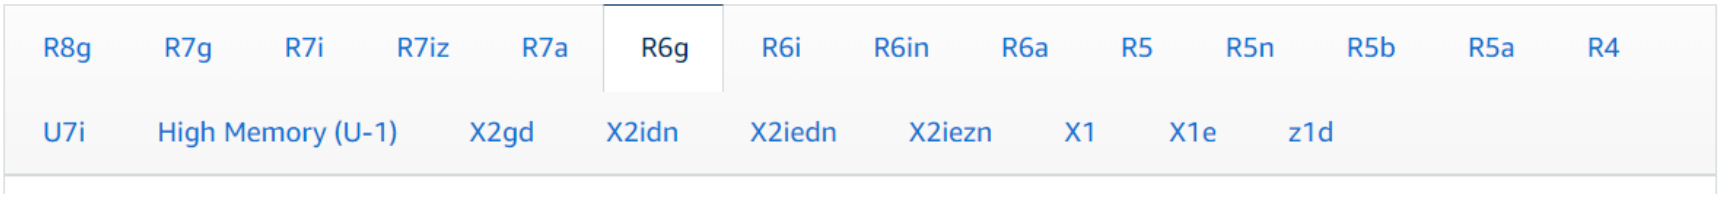

© Copyright by Cloud Mentor Pro | Confidential

https://aws.amazon.com/ec2/instance-types/

EC2 Instance Types - Storage Optimized

- High, sequential read and write access to very large data sets on local storage.
- Great for
  - High frequency online transaction processing (OLTP) systems
  - Relational & NoSQL databases
  - Cache for in-memory databases (for example, Redis)
  - Data warehousing applications
  - Distributed file systems

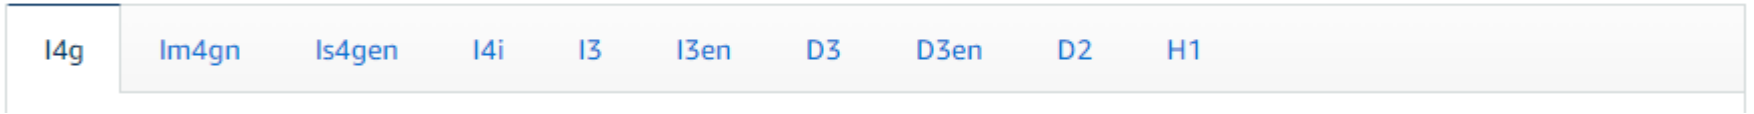

© Copyright by Cloud Mentor Pro | Confidential

https://aws.amazon.com/ec2/instance-types/

EC2 - Security Groups

- Control how traffic is allowed into or out of our EC2(**firewall**)
- Can be attached to multiple instances
- All inbound traffic is blocked by default
- All outbound traffic is authorised by default
- Security groups only contain **allow** rules
- Security groups rules can reference by IP or by security group

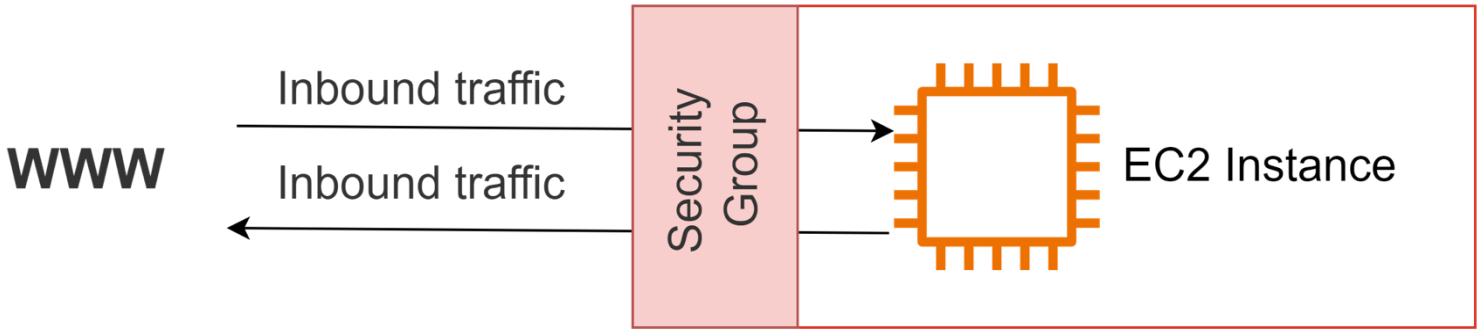

© Copyright by Cloud Mentor Pro | Confidential

EC2 - Security Groups

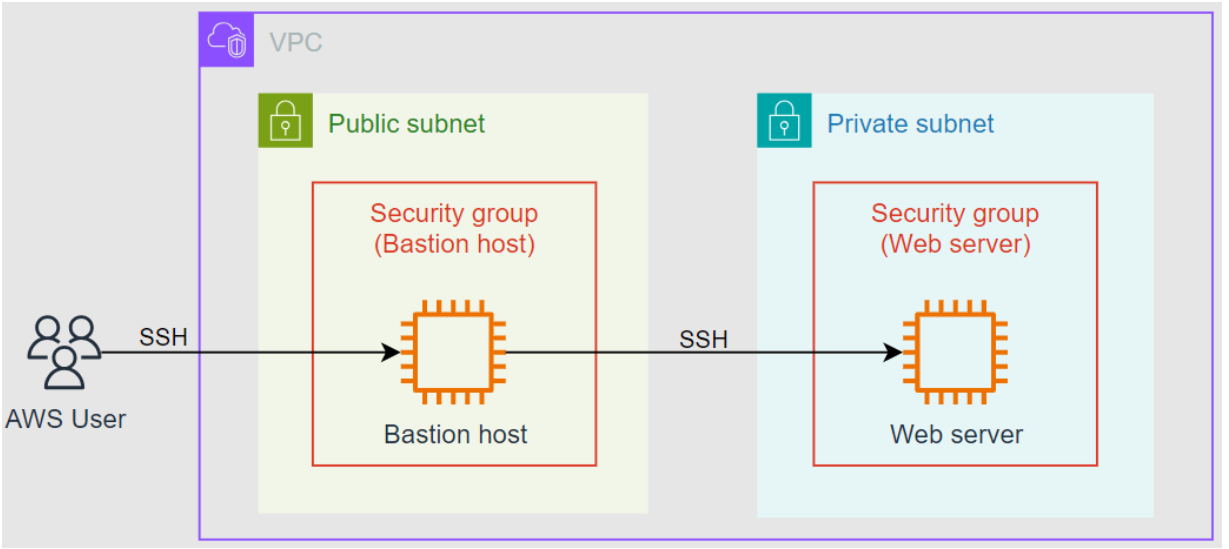

Inbound

Bastion SG

Type	Protocol	Port range	Source
SSH	TCP	22	0.0.0.0/0

Web server SG

Type	Protocol	Port range	Source
SSH	TCP	22	`Chọn SG của bastion host`

Outbound

Type	Protocol	Port range	Source
All traffic	All	All	0.0.0.0/0

Type	Protocol	Port range	Source
All traffic	All	All	0.0.0.0/0

© Copyright by Cloud Mentor Pro | Confidential

Classic Ports to know

- 22 = SSH (Secure Shell) - log into a Linux instance
- 21 = FTP (File Transfer Protocol) – upload files into a file share
- 22 = SFTP (Secure File Transfer Protocol) – upload files using SSH
- 80 = HTTP – access unsecured websites
- 443 = HTTPS – access secured websites
- 3389 = RDP (Remote Desktop Protocol) – log into a Windows instance

© Copyright by Cloud Mentor Pro | Confidential

EC2 - Connect to your EC2 instance

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/connect.html

EC2 Instance Connect

- Connect to your EC2 instance within your browser
- No need to use your key file
- Need to make sure the port 22 is still opened!
- **Works only Amazon Linux or Ubuntu**

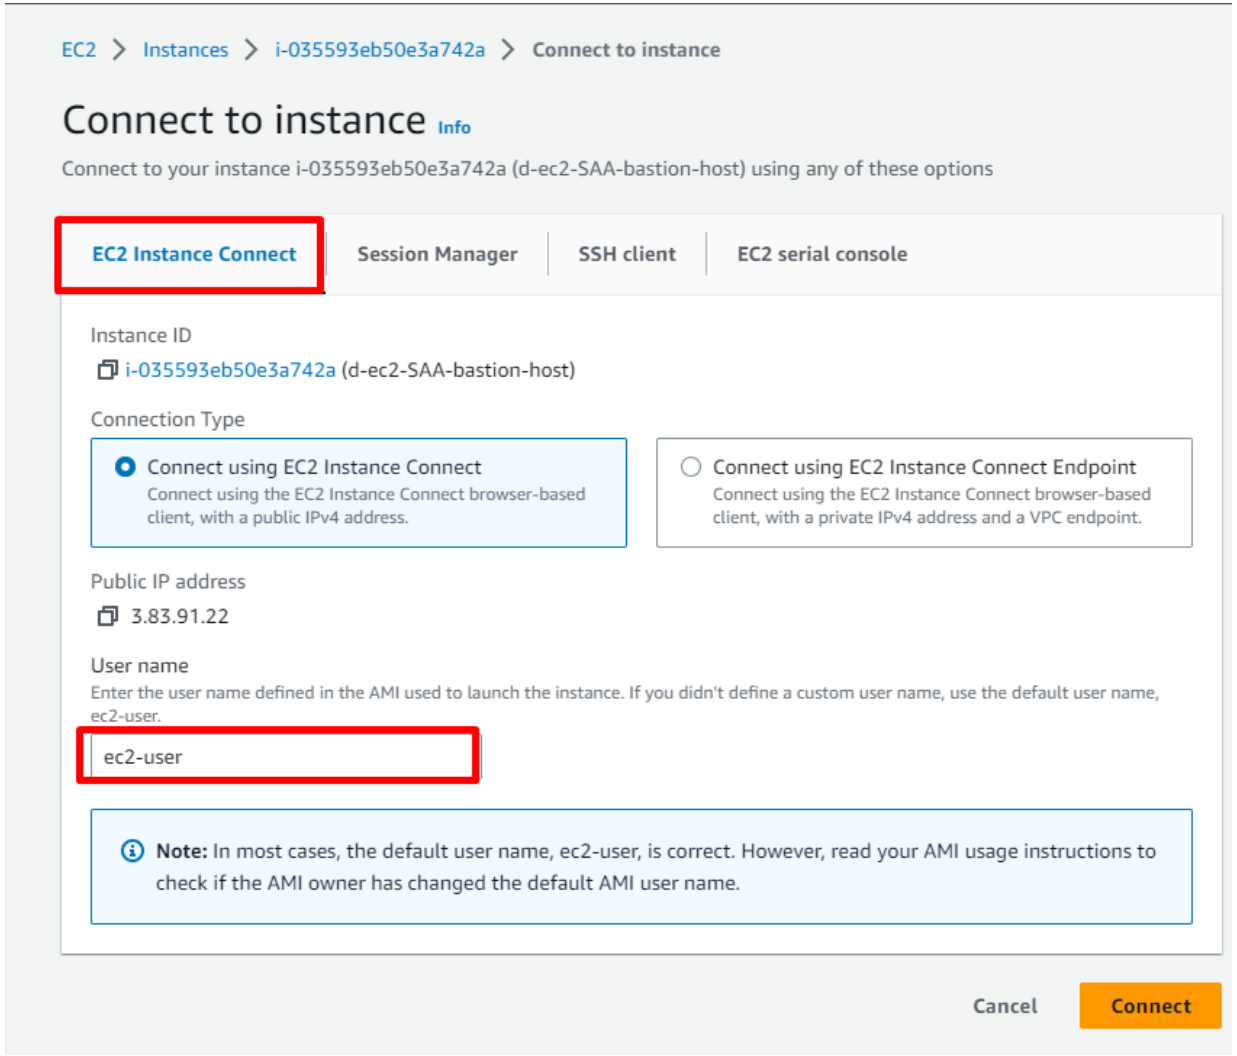

© Copyright by Cloud Mentor Pro | Confidential

EC2 Instances Purchasing Options

- **On-Demand Instances** – short workload, predictable pricing, pay by second
- **Reserved** (1 & 3 years)
  - **Reserved Instances** – long workloads
  - **Convertible Reserved Instances** – long workloads with flexible instances
- **Savings Plans** (1 & 3 years) –commitment to an amount of usage, long workload
- **Spot Instances** – short workloads, cheap, can lose instances (less reliable)
- **Dedicated Hosts** – book an entire physical server, control instance placement
- **Dedicated Instances** – no other customers will share your hardware
- **Capacity Reservations** – reserve capacity in a specific AZ for any duration

© Copyright by Cloud Mentor Pro | Confidential
https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/instance-purchasing-options.html

EC2 On Demand

- Has the highest cost but no upfront payment
- **Short-term** and **un-interrupted** workloads
- Pay for what you use
- No long-term commitment
- **Use case:**
  - Dev/stg/test environment
  - Critical batch job
  - System that runs only in business hour (ex: 8am – 17 pm)
  - System that scales frequently

© Copyright by Cloud Mentor Pro | Confidential

EC2 Reserved Instances

- Up to 72% discount compared to On-demand
- **Reserved Instances** – reserve instance configuration (Type, Region, Tenancy, OS)
- **Reservation Period** – 1 year (+discount) or 3 years (+++discount)
- **Payment Options** – No Upfront (+), Partial Upfront (++), All Upfront (+++)
- **Reserved Instance's Scope** – Regional or Zonal (reserve capacity in an AZ)
- Recommended for steady-state usage applications (think database)
- **Convertible Reserved Instances** – flexible instances
  - Can change the type, family, OS,..
  - Up to 66% discount

© Copyright by Cloud Mentor Pro | Confidential

EC2 Reserved Instances

- Use case:
  - System that runs 24/7 for a long time
  - Workload does not change for a long time

© Copyright by Cloud Mentor Pro | Confidential

EC2 Savings Plans

- Commit to a consistent **amount of usage, in USD per hour ($10/hour for 1 or 3 years)**
- Up to 72% discount compared to On-demand
- Locked to a specific instance family & AWS region (e.g., M5 in us-east-1)
- Flexible across:
  - Instance Size (e.g., m5.xlarge, m5.2xlarge)
  - OS (e.g., Linux, Windows)
  - Tenancy (Host, Dedicated, Default)

© Copyright by Cloud Mentor Pro | Confidential

EC2 Savings Plans

- Use case:
  - Same as for RI, system that runs 24/7 for a long time
  - Workload need to change INSTANCE TYPE in a long run

© Copyright by Cloud Mentor Pro | Confidential

EC2 Spot Instances

- Discount of up to 90% compared to On-demand
- If the current spot price > your max price you can choose to stop or terminate your instance within a **2 minutes** grace period.
- The MOST cost-efficient instances in AWS
- **Spot Block:** "block" spot instance during **1 to 6 hours** without interruptions
- **Useful for workloads that can be interrupted**
  - Batch jobs
  - Data analysis
  - Workloads with a flexible start and end time

© Copyright by Cloud Mentor Pro | Confidential

EC2 Dedicated Hosts

- A physical server with EC2 instance capacity fully dedicated to your use
- Use your existing server- bound software licenses (per-socket, per-core, pe—VM software licenses)

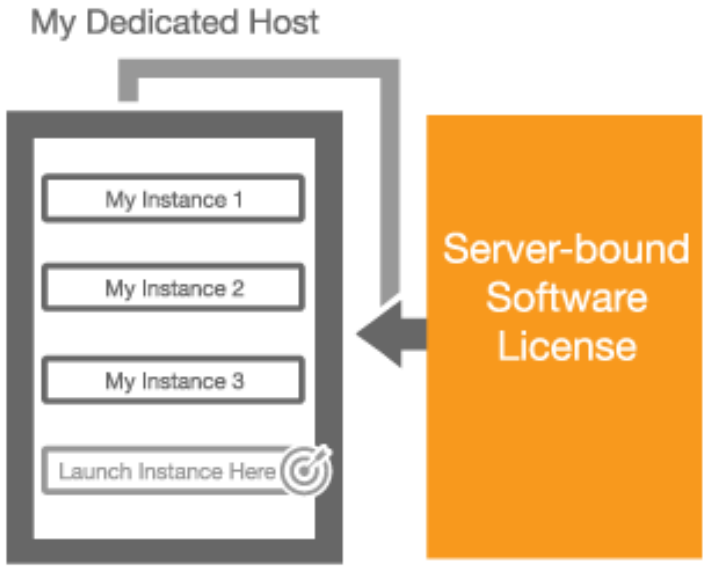

© Copyright by Cloud Mentor Pro | Confidential

EC2 Dedicated Instances

- Instances run on hardware that's dedicated to you
- May share hardware with other instances in same account
- No control over instance placement (can move hardware after Stop / Start)
- **Use case:**
  - System that requires strict policy which needs to run in an isolated environment (Ex: Bank system, Hospital system)

© Copyright by Cloud Mentor Pro | Confidential

Understanding AWS Tenancy

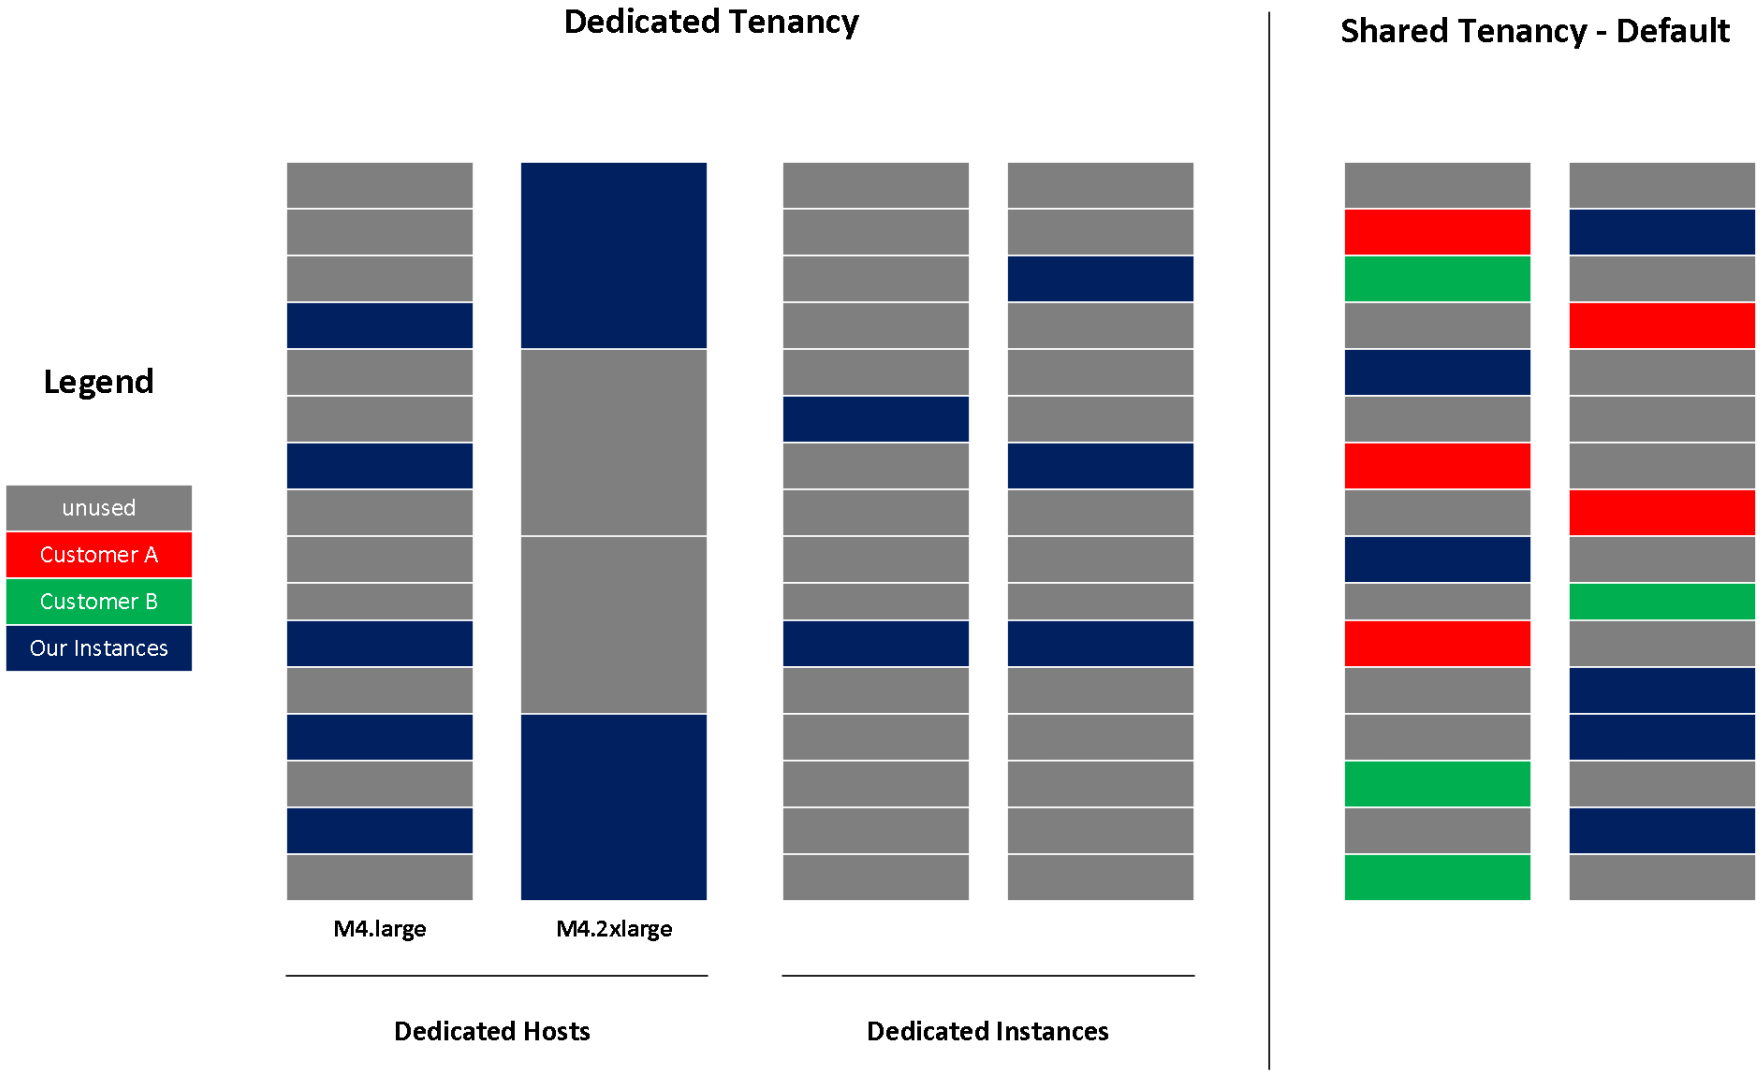

© Copyright by Cloud Mentor Pro | Confidential

EC2 Capacity Reservations

- Reserve On-Demand instances capacity in a specific AZ for any duration
- You always have access to EC2 capacity when you need it
- No time commitment (create/cancel anytime), no billing discounts
- You're charged at On-Demand rate whether you run instances or not
- Suitable for short-term, uninterrupted workloads that needs to be in a specific AZ

© Copyright by Cloud Mentor Pro | Confidential

Which purchasing option is right for me?

- On demand: coming and staying in resort whenever we like, we pay the full price
- Reserved: like planning ahead and if we plan to stay for a long time, we may get a good discount.
- Savings Plans: pay a certain amount per hour for certain period and stay in any room type (e.g., King, Suite, Sea View, …)
- Spot instances: the hotel allows people to bid for the empty rooms and the highest bidder keeps the rooms. You can get kicked out at any time
- Dedicated Hosts: We book an entire building of the resort
- Capacity Reservations: you book a room for a period with full price even you don't stay in it

© Copyright by Cloud Mentor Pro | Confidential

AWS charges for IPv4 addresses

- Starting February 1st 2024, there's a charge for all Public IPv4 created in your account
- $0.005 per hour of Public IPv4 (~$3.6 per month)
- For new accounts in AWS, you have a free tier for the EC2 service: 750 hours of Public IPv4 per month for the first 12 months
- For all other services there is no free tier

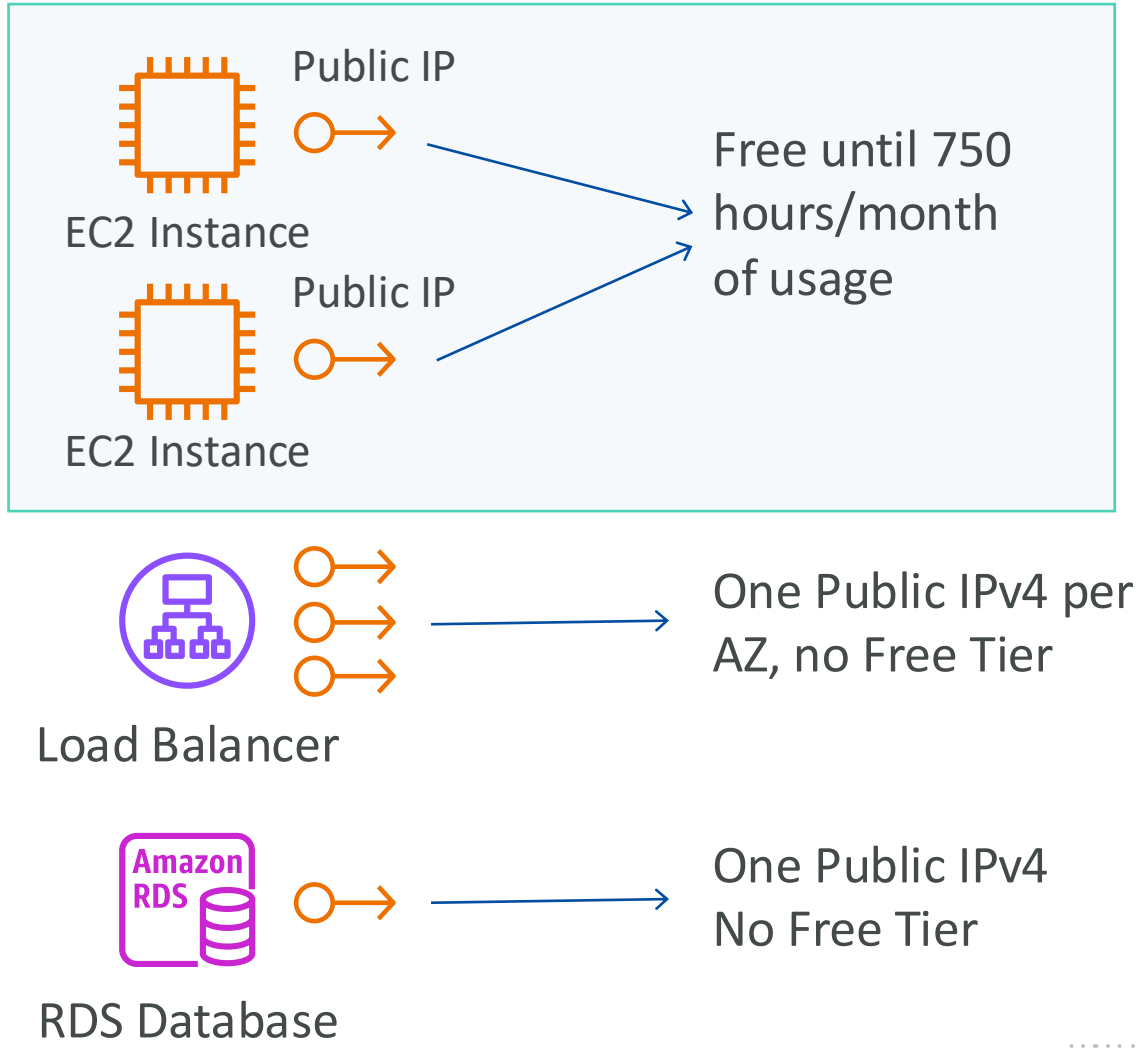

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- *EC2 Basic*

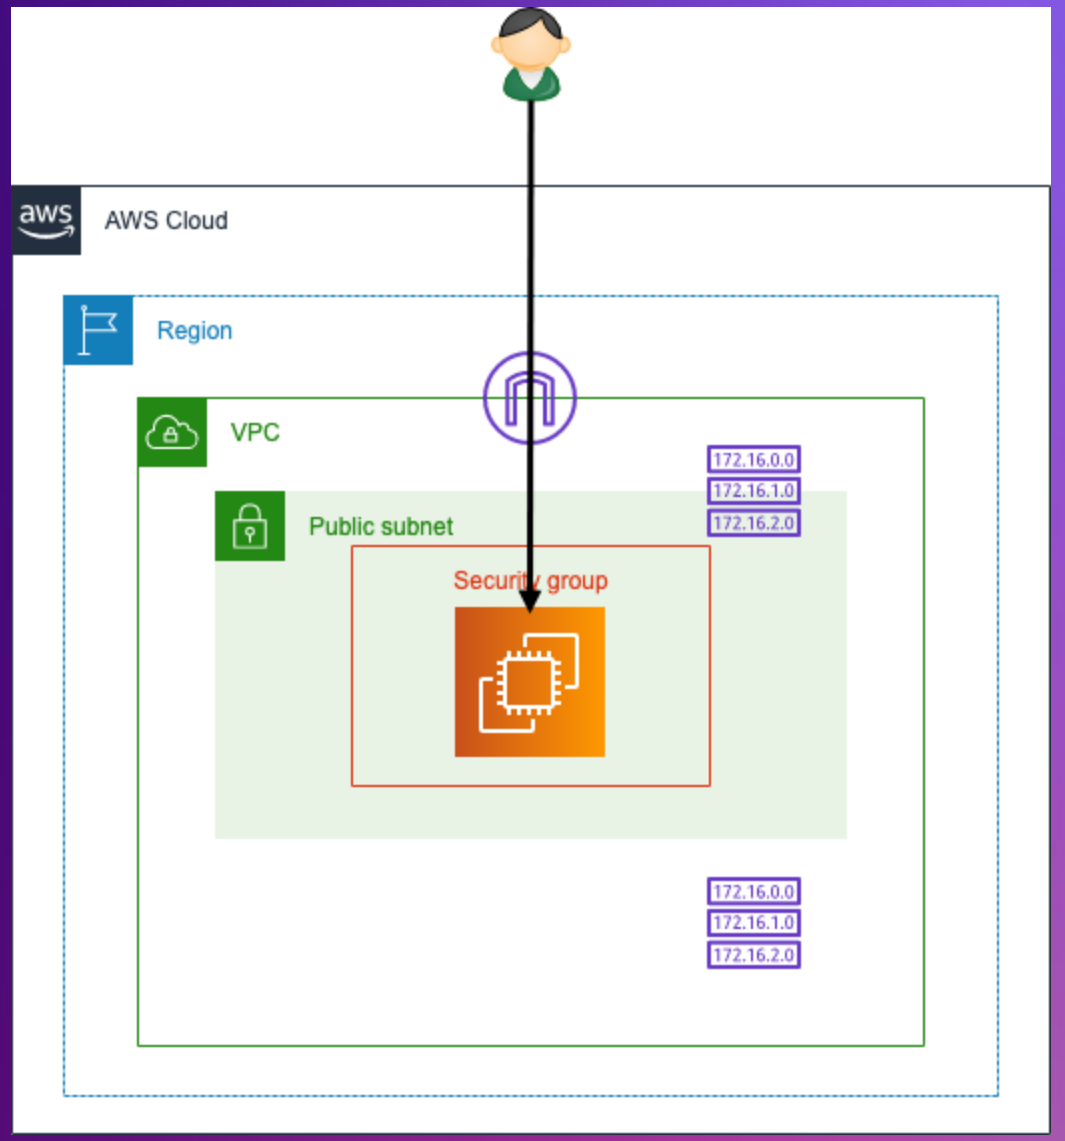

© Copyright by Cloud Mentor Pro | Confidential

Section 2

- *EC2 Instance Storage*
- *High Availability & Scalability: ELB & ASG*

© Copyright by Cloud Mentor Pro | Confidential

EC2 Instance Storage

EBS Volume, EC2 Instance Store, Elastic File System(EFS)

© Copyright by Cloud Mentor Pro | Confidential

What's an EBS Volume?

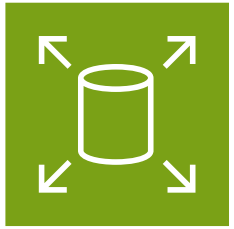

- An **EBS (Elastic Block Store) Volume** is a **network drive** you can attach to your instances while they run
- It allows your instances to **persist data**, even after their termination
- They can only be mounted to one instance at a time (except: io1/io2)
- It's locked to an Availability Zone (AZ)
- Analogy: Think of them as a "network USB stick"

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-volumes.html

What's an EBS Volume?

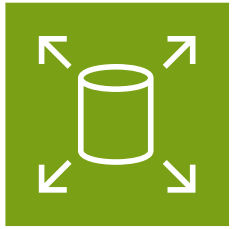

- By default, when you create an EC2 Instance, it will automatically attach an EBS Volume, this is called **Root EBS Volume**.
- This volume contains OS and system files.
- **You can add additional volumes if you want to**. See example in the next slide.

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-volumes.html

EBS Volume - Example

- Example: You have an EC2 Windows instance, and you want to add additional volume for D and E Drive

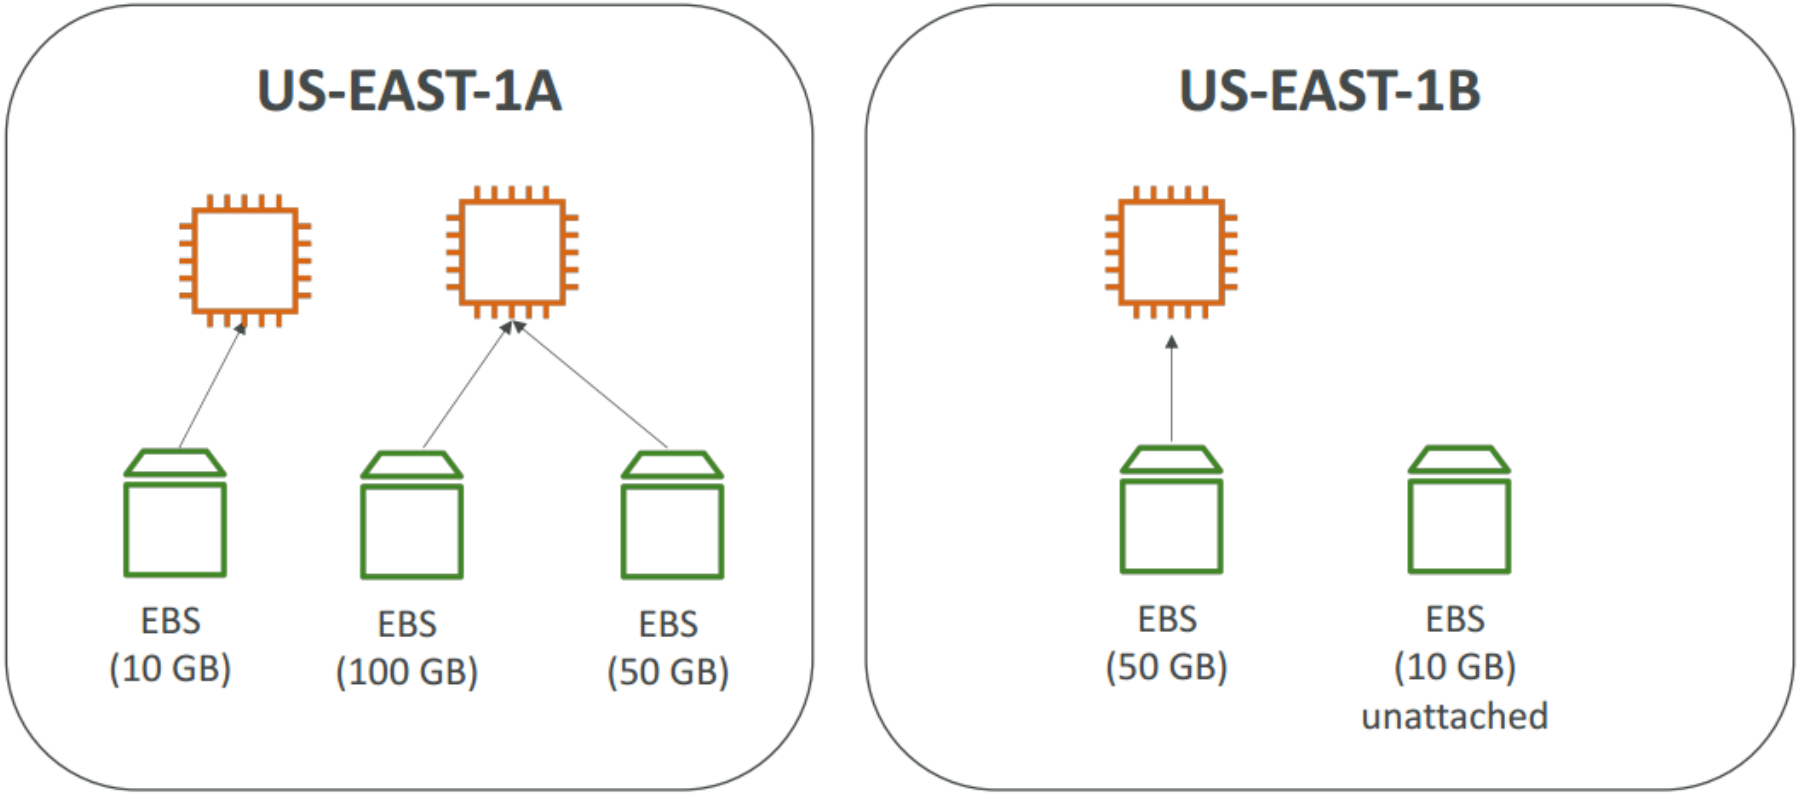

© Copyright by Cloud Mentor Pro | Confidential

EBS Volume - Example

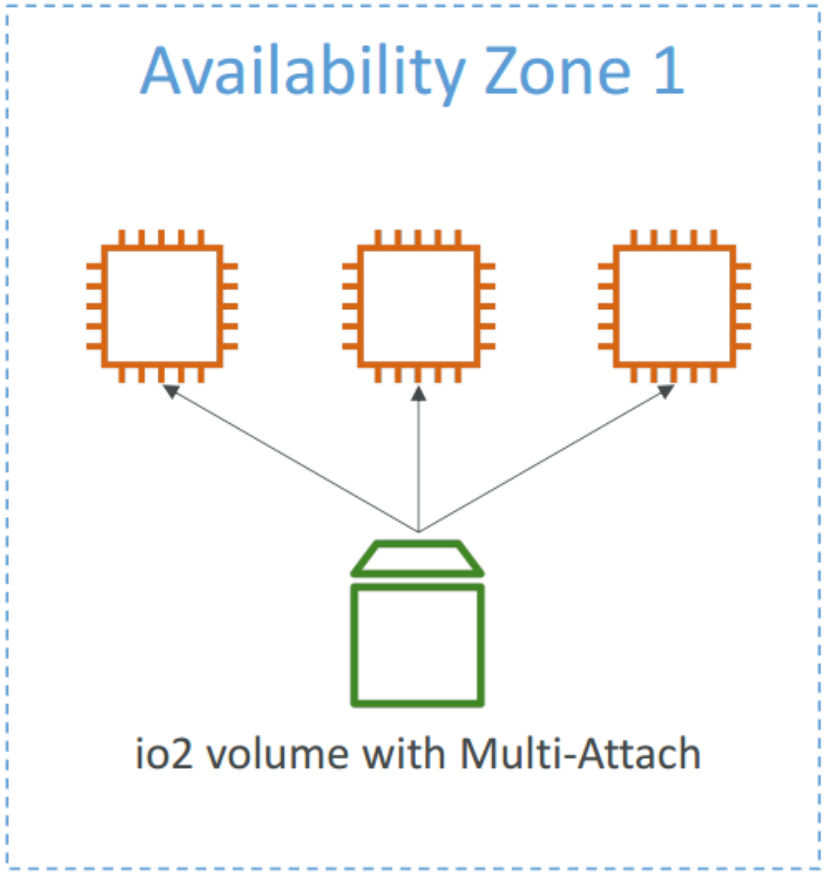

EBS **Multi-Attach (io1/io2 family only) -**
Up to 16 EC2 Instances at a time

© Copyright by Cloud Mentor Pro | Confidential

EBS – Delete on Termination attribute

- By default, the **root EBS volume** is deleted (attribute **enabled**)
- By default, **Non-root EBS volume** is not deleted (attribute **disabled**)
- **Use case**: preserve root volume when instance is terminated

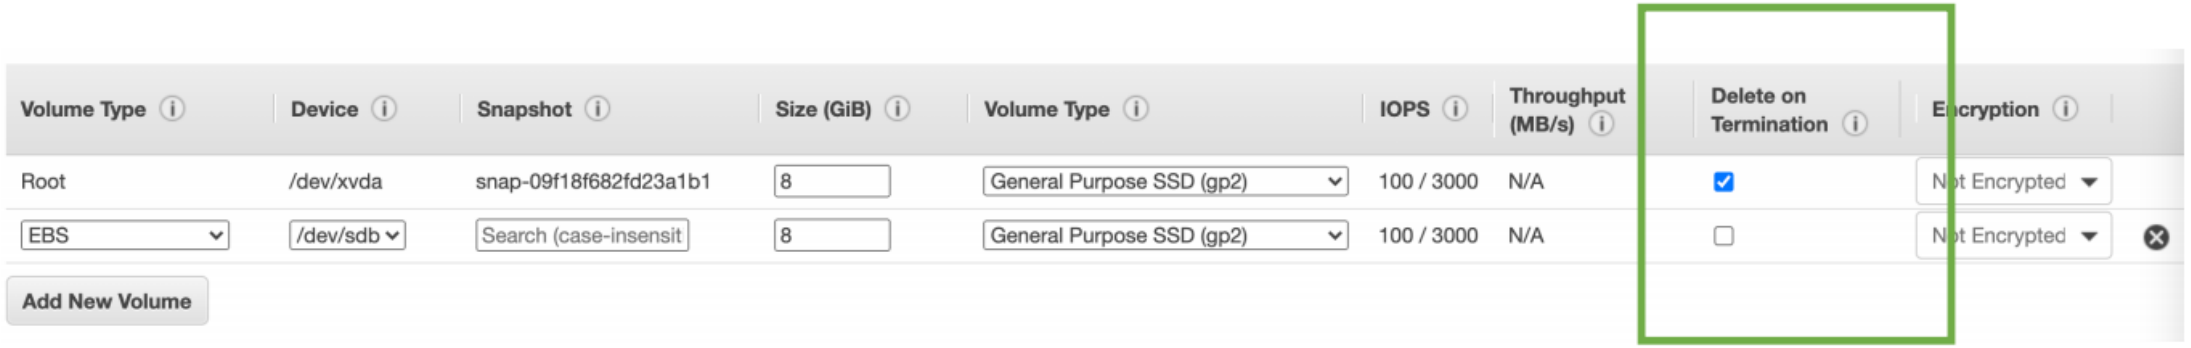

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/preserving-volumes-on-termination.html

© Copyright by Cloud Mentor Pro | Confidential

EBS Snapshots

- Make a **backup snapshot** of your EBS volume at a point in time
- Not necessary to detach volume to do snapshot, but recommended
- Can **copy snapshots across AZ or Region**

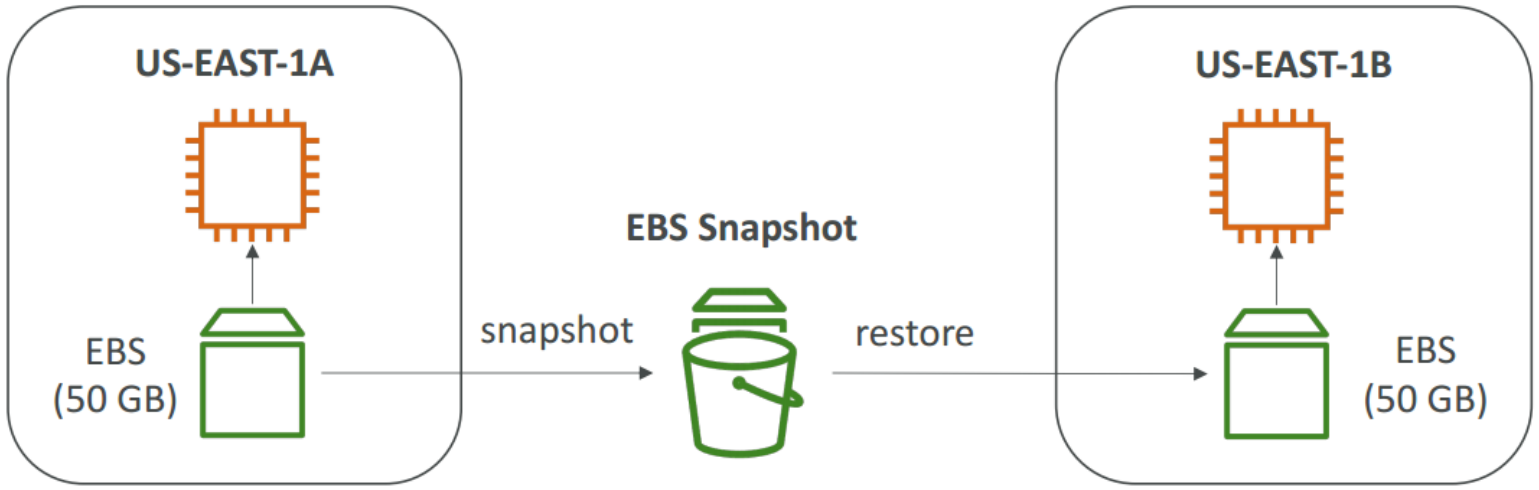

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-snapshots.html

AMI Overview

- AMI = Amazon Machine Image
- AMI are a customization of an EC2 instance
  - You add your own software, configuration, operating system, monitoring…
  - Faster boot / configuration time because all your software is pre-packaged
- AMI are built for a **specific region** (and can be copied across regions)
- You can launch EC2 instances from:
  - **A Public AMI:** AWS provided
  - **Your own AMI:** you make and maintain them yourself
  - **An AWS Marketplace AMI:** an AMI someone else made (and potentially sells)

© Copyright by Cloud Mentor Pro | Confidential

AMI Process (from an EC2 instance)

- Start an EC2 instance and customize it
- Stop the instance (for data integrity)
- Build an AMI – this will also create EBS snapshots
- Launch instances from other AMIs

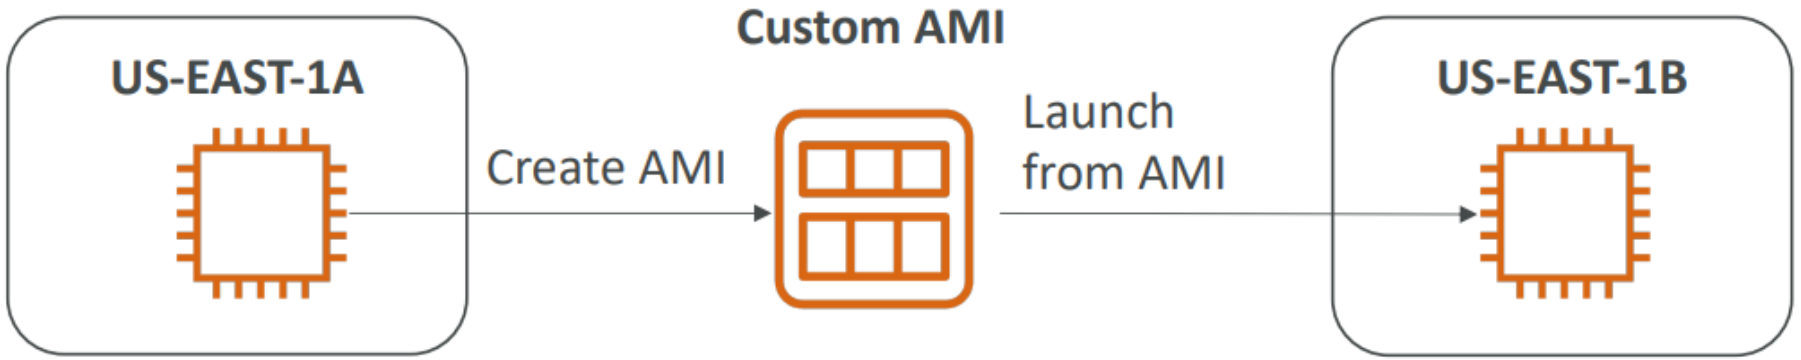

© Copyright by Cloud Mentor Pro | Confidential

EBS Snapshots vs AMI

- EBS Snapshot is an copy of a particular volume only.
- You cannot launch an EC2 instance from an EBS Snapshot.
- Use case:
  - Backup data for a particular volume, for example: D Drive, E Drive, etc…

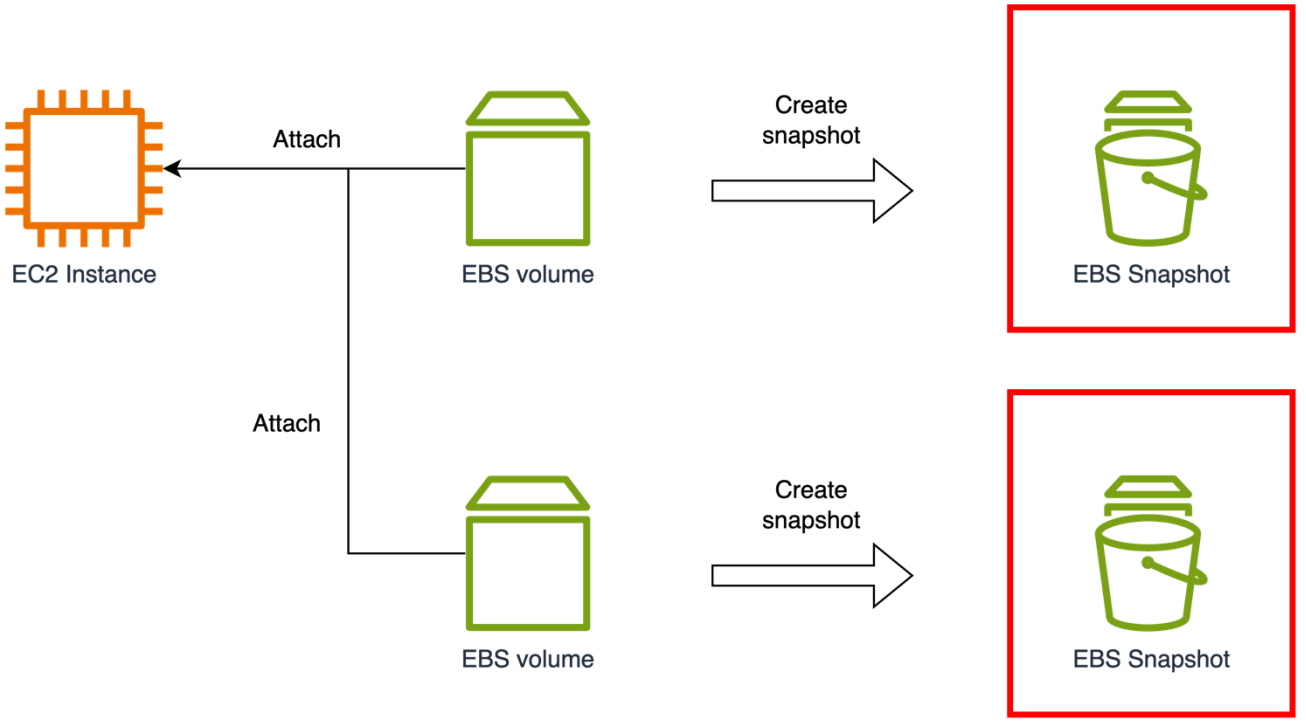

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-snapshots.html

EBS Snapshots vs AMI

- AMI is a backup of your **whole EC2 instance, including EBS snapshot.**
- **You can launch an EC2 instance from an AMI.**
- **Use case:**
  - Backup your ec2 instance
  - When your system scales, you need multiple ec2 to handle the load

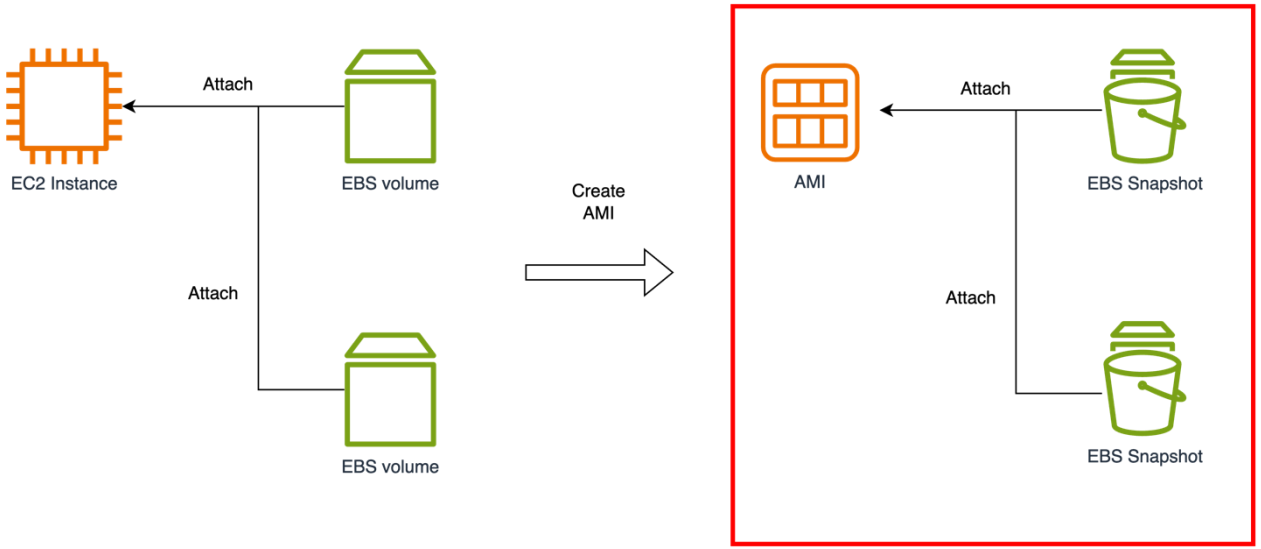

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-snapshots.html

EBS Snapshots Features

- EBS Snapshot Archive
  - Move a Snapshot to an ”archive tier” that is 75% cheaper
  - Takes within 24 to 72 hours for restoring the archive
- Recycle Bin for EBS Snapshots
  - Setup rules to retain deleted snapshots so you can recover them after an accidental deletion
  - Specify retention (from 1 day to 1 year)
- Fast Snapshot Restore (FSR)
  - Force full initialization of snapshot to have no latency on the first use ($$$)

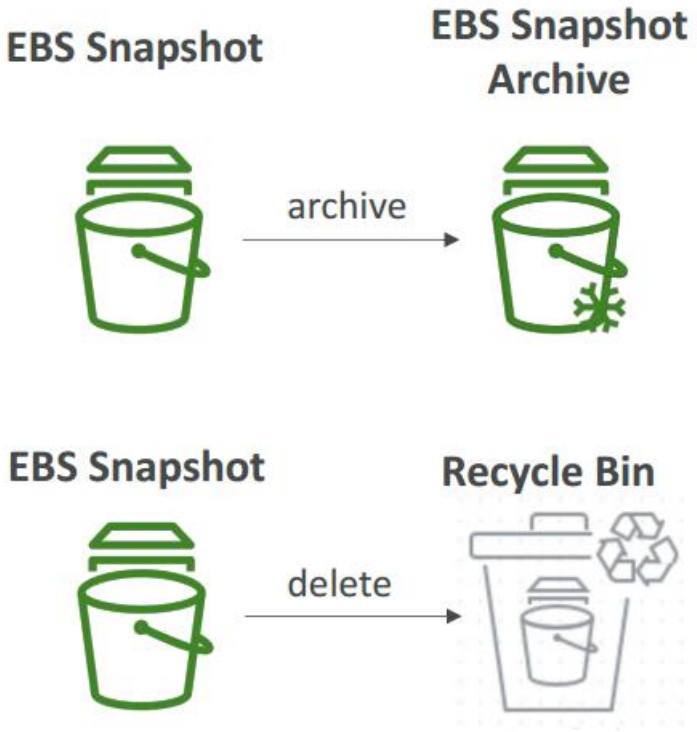

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-snapshots.html

EBS Volume Types

- EBS Volumes come in 6 types
  - gp2 / gp3 (SSD)
  - io1 / io2 Block Express (SSD
  - st1 (HDD):
  - sc1 (HDD):
- EBS Volumes are characterized in Size | Throughput | IOPS (I/O Ops Per Sec)
- **Only gp2/gp3 and io1/io2 Block Express can be used as boot volumes**

© Copyright by Cloud Mentor Pro | Confidential

General Purpose SSD

- General purpose SSD volume that balances price and performance
- Cost effective storage, low-latency

- **gp3:**
  - Baseline of 3,000 IOPS and throughput of 125 MiB/s
  - Can increase IOPS up to **16,000** and throughput up to 1000 MiB/s independently
- **gp2:**
  - Small gp2 volumes can burst IOPS to 3,000
  - Size of the volume and IOPS are linked, max IOPS is **16,000**
  - 3 IOPS per GB, means at 5,334 GB we are at the max IOPS

© Copyright by Cloud Mentor Pro | Confidential

Provisioned IOPS (PIOPS) SSD

- **Highest-performance** SSD volume for **mission-critical low-latency or high-throughput workloads**
- Or applications that need more than 16,000 IOPS
- Great for **databases workloads**
- io1 (4 GiB - 16 TiB):
  - Max PIOPS: 64,000 for Nitro EC2 instances & 32,000 for other
  - Can increase PIOPS independently from storage size
- io2 Block Express (4 GiB – 64 TiB):
  - Sub-millisecond latency
  - Max PIOPS: 256,000 with an IOPS:GiB ratio of 1,000:1
- **Supports EBS Multi-attach**

© Copyright by Cloud Mentor Pro | Confidential

Hard Disk Drives (HDD)

- Throughput Optimized HDD (st1)
  - Big Data, Data Warehouses, Log Processing
  - Max throughput 500 MiB/s – max IOPS 500

- Cold HDD (sc1):
  - For data that is infrequently accessed
  - Scenarios where lowest cost is important
  - Max throughput 250 MiB/s – max IOPS 250

© Copyright by Cloud Mentor Pro | Confidential

EBS - Volume Types Summary - SSD

Solid state drive (SSD) volumes				
	General Purpose SSD volumes		Provisioned IOPS SSD volumes	
Volume type	**gp3**	**gp2**	**io2 Block Express 3**	**io1**
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)		99.999% durability (0.001% annual failure rate)	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)
Use cases	Transactional workloads Virtual desktops Medium-sized, single-instance databases Low-latency interactive applications Boot volumes Development and test environments		Sub-millisecond latency Sustained IOPS performance **More than 64,000 IOPS** or 1,000 MiB/s of throughput	Sustained IOPS performance or **more than 16,000 IOPS** I/O-intensive database workloads
Volume size	1 GiB - 16 TiB		4 GiB - 64 TiB	4 GiB - 16 TiB
Max IOPS per volume	16,000 (64 KiB I/O)	16,000 (16 KiB I/O)	256,000 (16 KiB I/O)	64,000 (16 KiB I/O)
Max throughput per volume	1,000 MiB/s	250 MiB/s	4,000 MiB/s	1,000 MiB/s
Amazon EBS Multi-attach	Not supported		Supported	
NVMe reservations	Not supported		Supported	Not supported
Boot volume	Supported			

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/ebs-volume-types.html

EBS - Volume Types Summary - HDD

Hard disk drive (HDD) volumes		
	Throughput Optimized HDD volumes	**Cold HDD volumes**
Volume type	**st1**	**sc1**
Durability	99.8% - 99.9% durability (0.1% - 0.2% annual failure rate)	
Use cases	**Big data** **Data warehouses** **Log processing**	Throughput-oriented storage for data that is **infrequently accessed** Scenarios where the **lowest storage cost** is important
Volume size	125 GiB - 16 TiB	
Max IOPS per volume (1 MiB I/C	500	250
Max throughput per volume	500 MiB/s	250 MiB/s
Amazon EBS Multi-attach	Not supported	
Boot volume	Not supported	

© Copyright by Cloud Mentor Pro | Confidential

EBS Multi-Attach – io1/io2 family

- Attach the same EBS volume to multiple EC2 instances in the **same AZ**
- Use case:
  - Achieve **higher application availability** in clustered Linux applications (ex: Teradata)
  - Applications must manage concurrent write operations
- **Up to 16 EC2 Instances at a time**

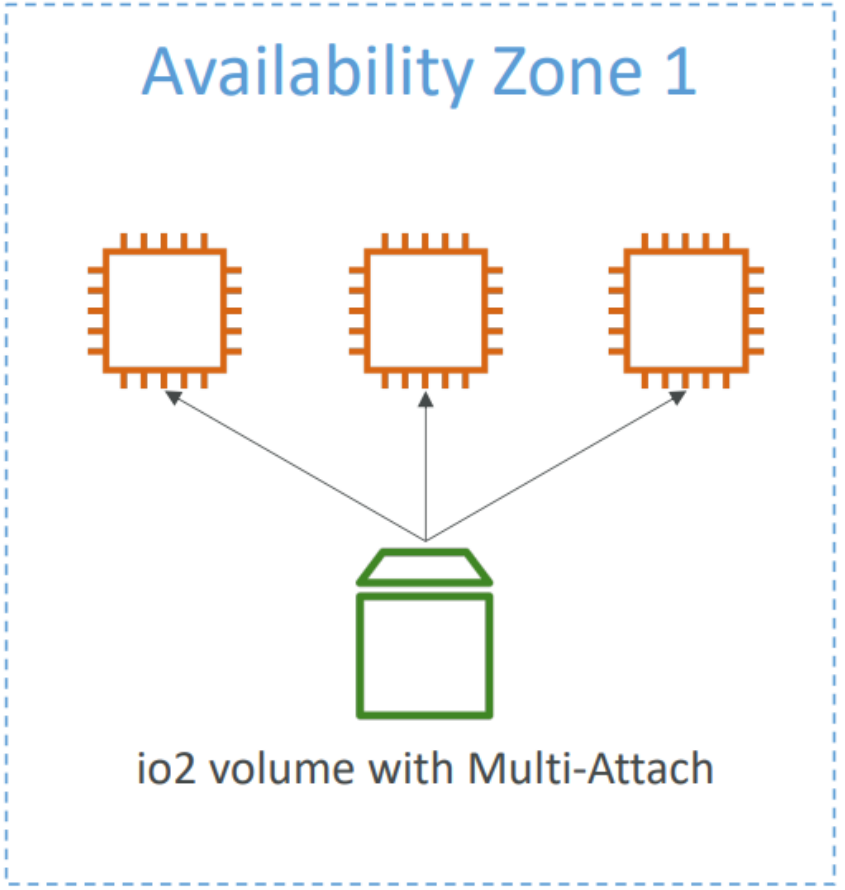

© Copyright by Cloud Mentor Pro | Confidential

EBS Encryption

- When you create an encrypted EBS volume, you get the following:
  - Data at rest is encrypted inside the volume
  - All the data in flight moving between the instance and the volume is encrypted
  - All snapshots are encrypted
  - All volumes created from the snapshot
- Encryption has a minimal impact on latency
- EBS Encryption leverages keys from KMS (AES-256)
- Copying an unencrypted snapshot allows encryption
- Snapshots of encrypted volumes are encrypted

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/ebs/latest/userguide/encryption-examples.html

Encryption: encrypt an unencrypted EBS volume

- Create an EBS snapshot of the volume
- Encrypt the EBS snapshot (using copy)
- Create new ebs volume from the snapshot (the volume will also be encrypted)
- Now you can attach the encrypted volume to the original instance

© Copyright by Cloud Mentor Pro | Confidential

EC2 Instance Store

- EC2 Instance Store is **high-performance** hardware disk
- Better I/O performance
- Risk of data loss when EC2 are stopped (**ephemeral**)
- Good for **buffer / cache / scratch data / temporary** content
- Some instance types **do not support** instance store volumes

Very high IOPS

Instance Size	100% Random Read IOPS	Write IOPS
i3.large *	100,125	35,000
i3.xlarge *	206,250	70,000
i3.2xlarge	412,500	180,000
i3.4xlarge	825,000	360,000
i3.8xlarge	1.65 million	720,000
i3.16xlarge	3.3 million	1.4 million
i3.metal	3.3 million	1.4 million
i3en.large *	42,500	32,500
i3en.xlarge *	85,000	65,000
i3en.2xlarge *	170,000	130,000
i3en.3xlarge	250,000	200,000
i3en.6xlarge	500,000	400,000
i3en.12xlarge	1 million	800,000
i3en.24xlarge	2 million	1.6 million
i3en.metal	2 million	1.6 million

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html

Amazon EFS – Elastic File System

- Managed **NFS** (network file system)
- Can be **mounted on many EC2**
- EFS works with EC2 instances in **Multi-AZ**
- Highly available, scalable, expensive (3x gp2)
- Use cases:
  - **Content management**, **web serving, data sharing**
  - Uses **NFSv4.1 protocol**
  - Compatible with **Linux** based AMI (not Windows)
  - **POSIX file system** (~Linux)

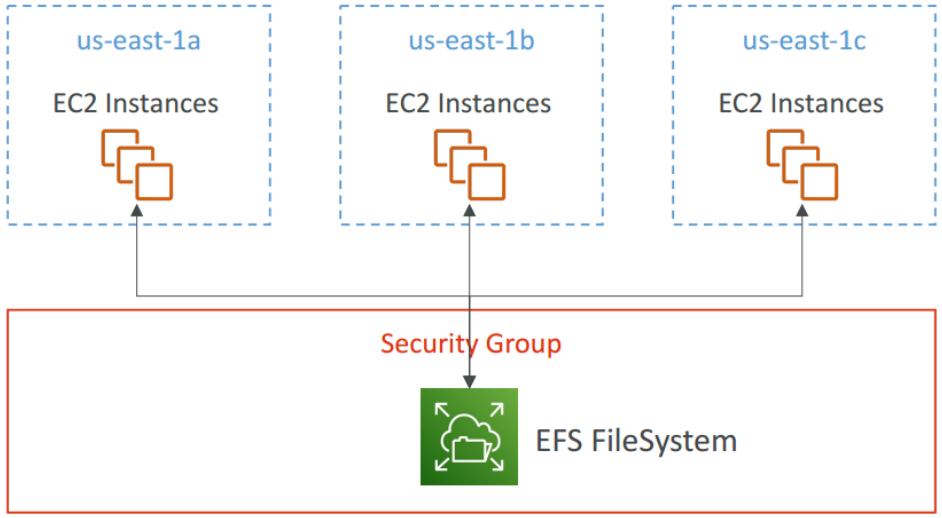

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/efs/latest/ug/whatisefs.html

Amazon EFS – Storage Classes

Storage class	Designed for	First byte read latency	Durability	Minimum billing charge per file	Minimum storage duration
EFS Standard	**Active data** requiring fast sub-millisecond latency performance	Sub-millisecond	99.999999999%	Not applicable	Not applicable
EFS IA Infrequent Access	**Inactive data** that is accessed only a **few times each quarter**	Tens of milliseconds		128 KiB	Not applicable
EFS Archive	**Inactive data** that is accessed a **few times each year or less**	Tens of milliseconds		128 KiB	90 days

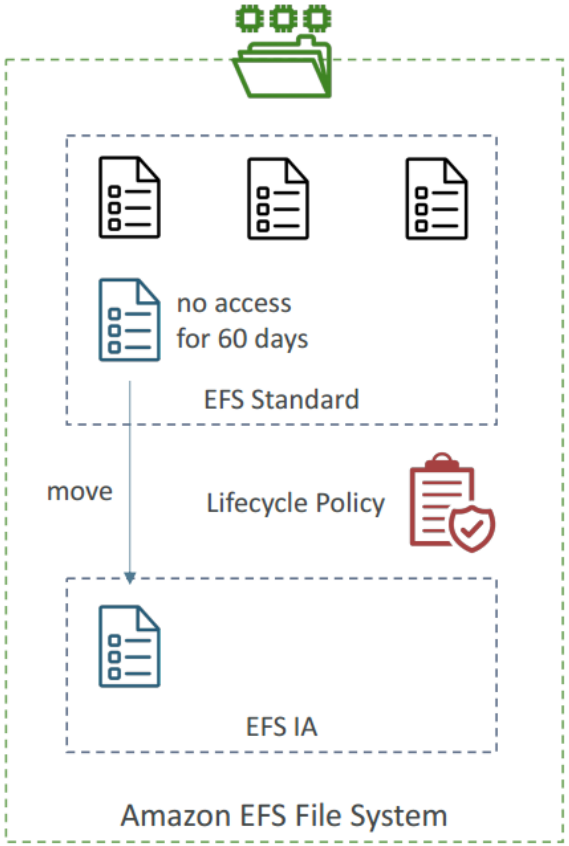

© Copyright by Cloud Mentor Pro | Confidential

https://aws.amazon.com/efs/storage-classes/

EBS vs EFS vs Instance store

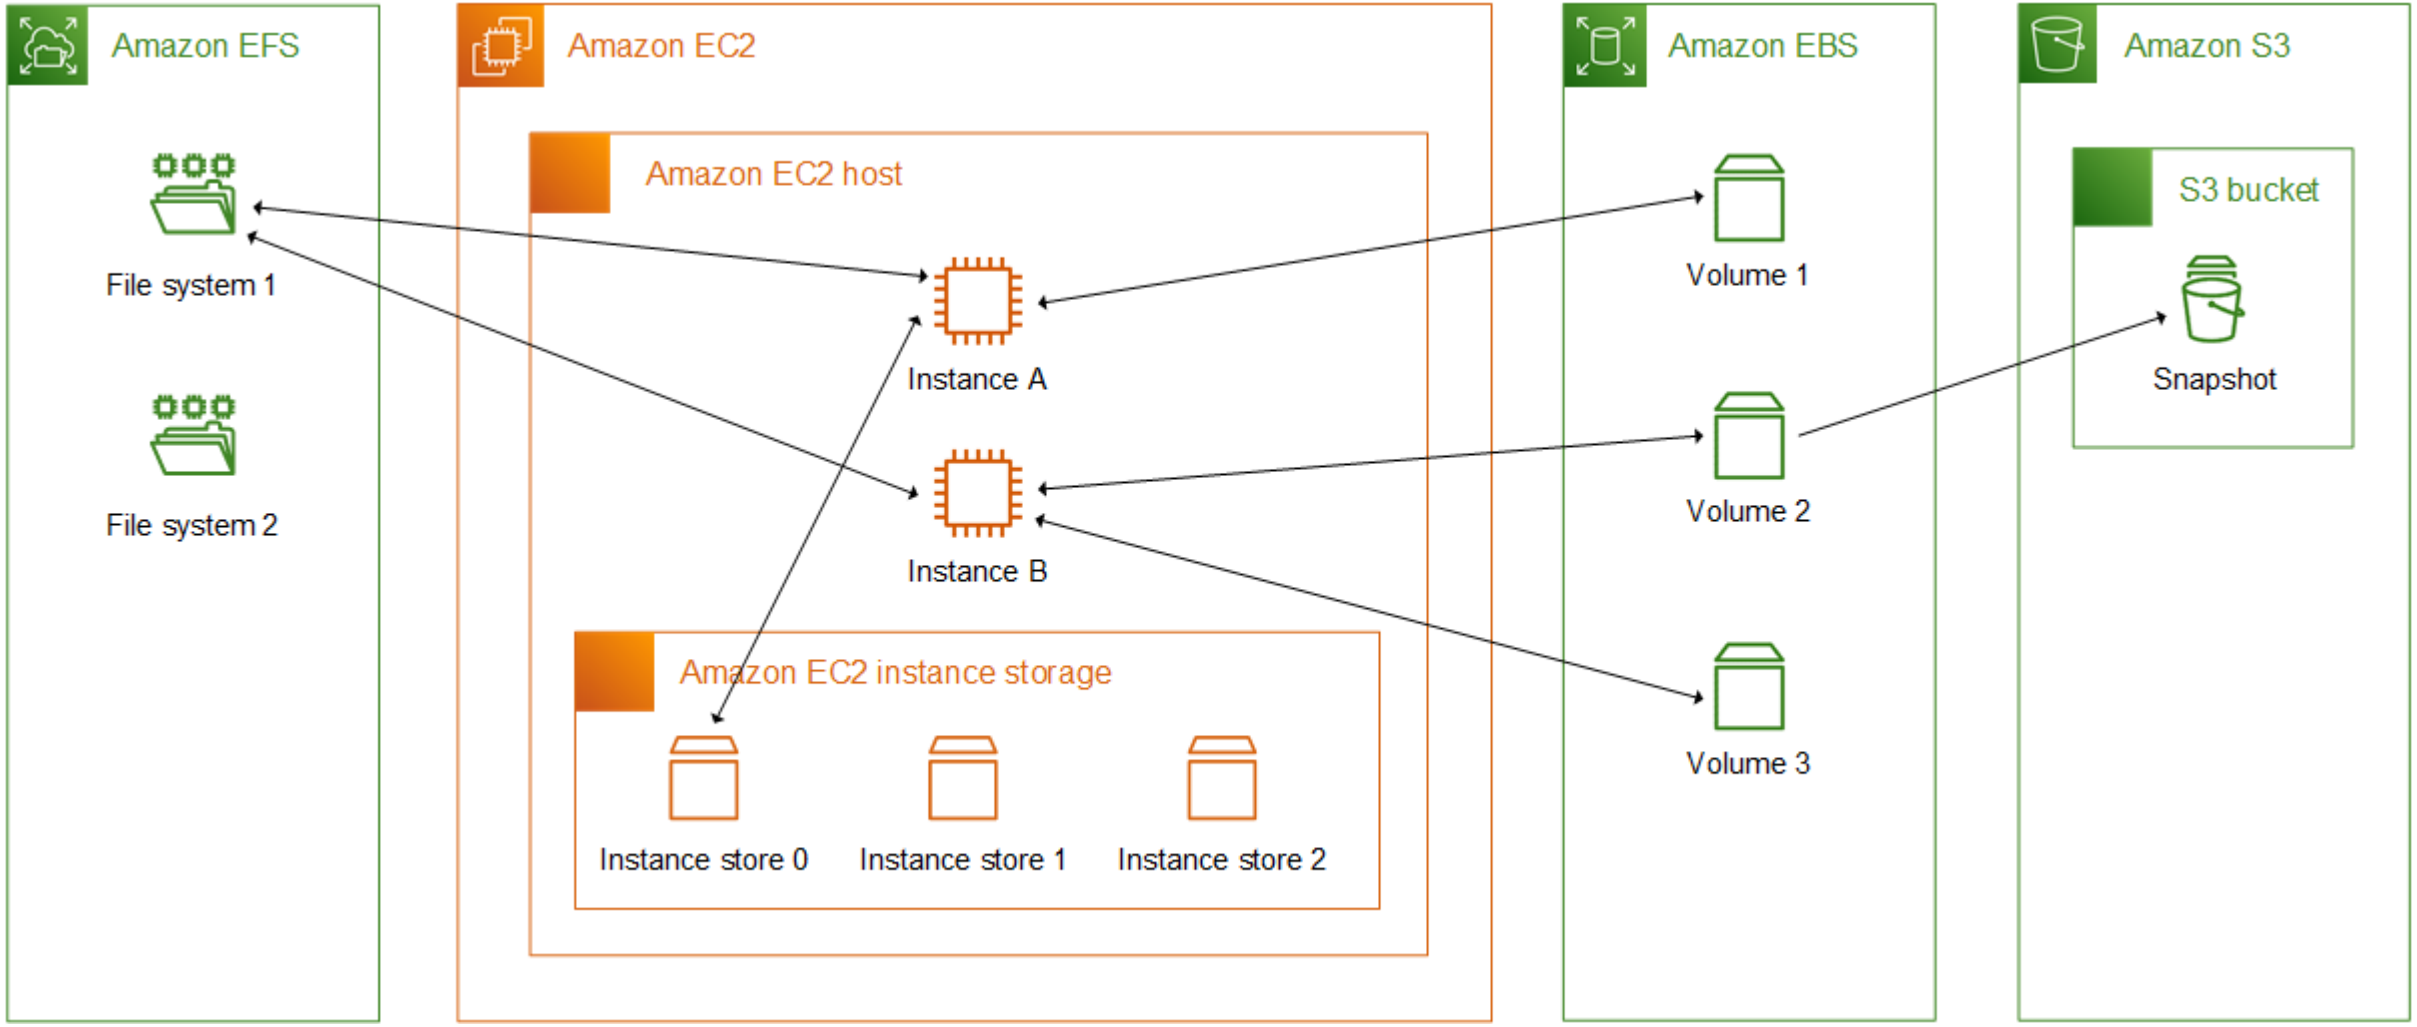

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Storage.html

High Availability & Scalability

ELB & ASG

© Copyright by Cloud Mentor Pro | Confidential

Scalability & High Availability

- Scalability means that an application / system can handle greater loads by adapting.
- There are two kinds of scalability:
  - Vertical Scalability
  - Horizontal Scalability (= elasticity)

- **Scalability is linked but different to High Availability**

© Copyright by Cloud Mentor Pro | Confidential

Vertical Scalability

- Vertically scalability (= **scale up / down**) means increasing the size of the instance
- Vertical scalability is very common for non distributed systems, such as a **database**.
- RDS, ElastiCache are services that can scale vertically.

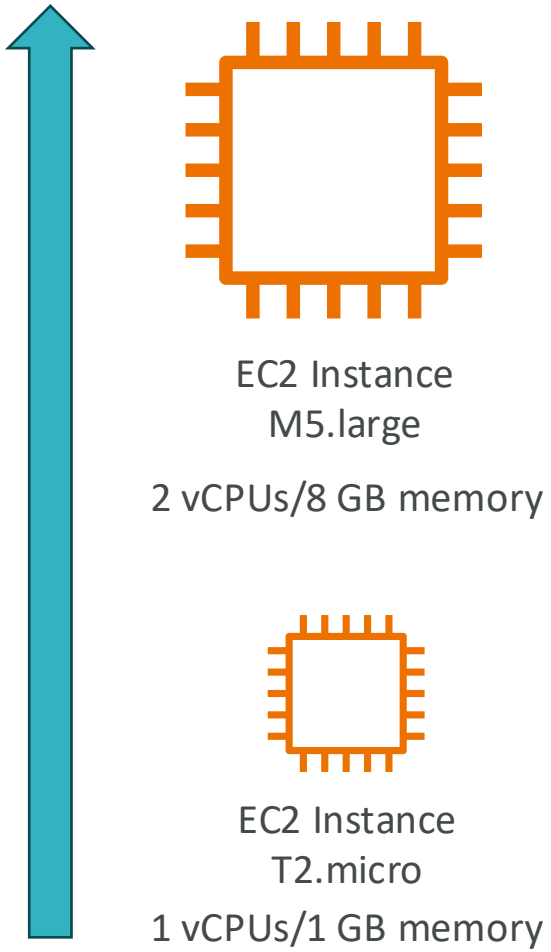

© Copyright by Cloud Mentor Pro | Confidential

Horizontal Scalability

- Horizontal Scalability (= scale out / in) means increasing the number of instances / systems for your application
- Horizontal scaling implies **distributed systems**.
- This is very common for **web applications / modern applications**

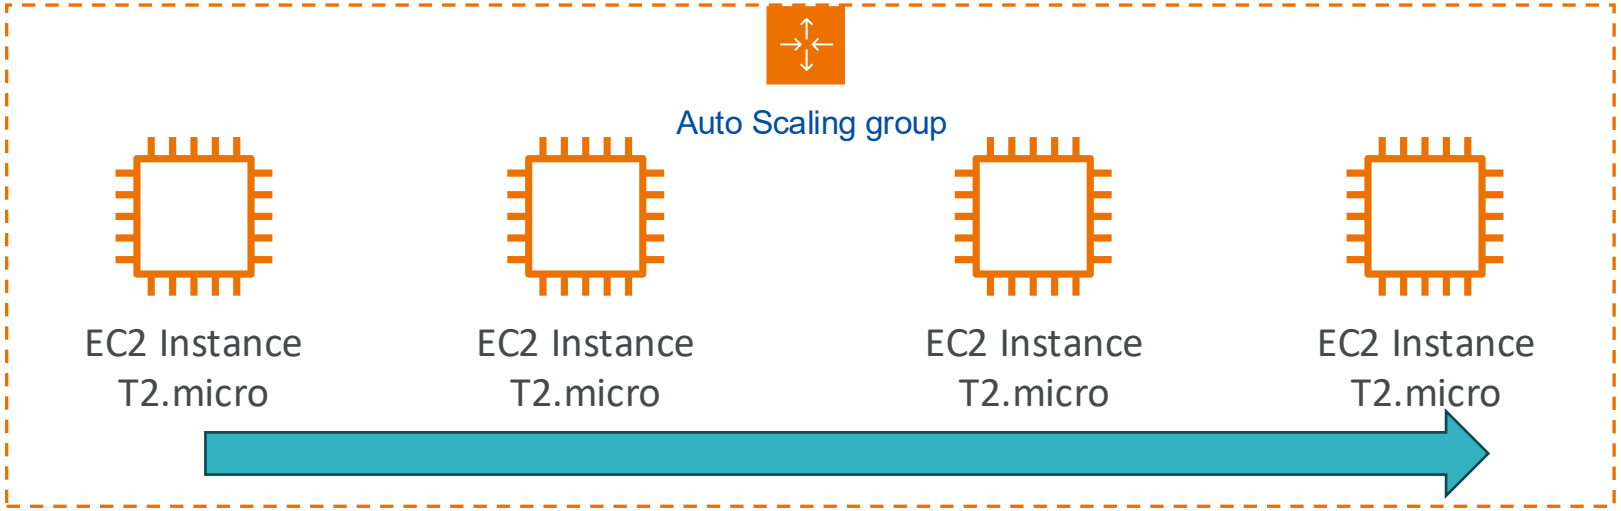

© Copyright by Cloud Mentor Pro | Confidential

High Availability

- High Availability usually goes hand in hand with horizontal scaling
- High availability means running your application / system in at least 2 data centers (== Availability Zones)
- The goal of high availability is to survive a data center loss
- The high availability can be passive (for RDS Multi AZ for example)
- The high availability can be active (for horizontal scaling)

© Copyright by Cloud Mentor Pro | Confidential

High Availability & Scalability For EC2

- **Vertical Scaling:** Increase instance size (= scale up / down)
  - From: t2.nano - 0.5G of RAM, 1 vCPU
  - To: u-12tb1.metal – 12.3 TB of RAM, 448 vCPUs
- **Horizontal Scaling:** Increase number of instances (= scale out / in)
  - Auto Scaling Group
  - Load Balancer
- **High Availability:** Run instances for the same application across multi AZ
  - Auto Scaling Group multi AZ
  - Load Balancer multi AZ

© Copyright by Cloud Mentor Pro | Confidential

High Availability & Scalability For EC2

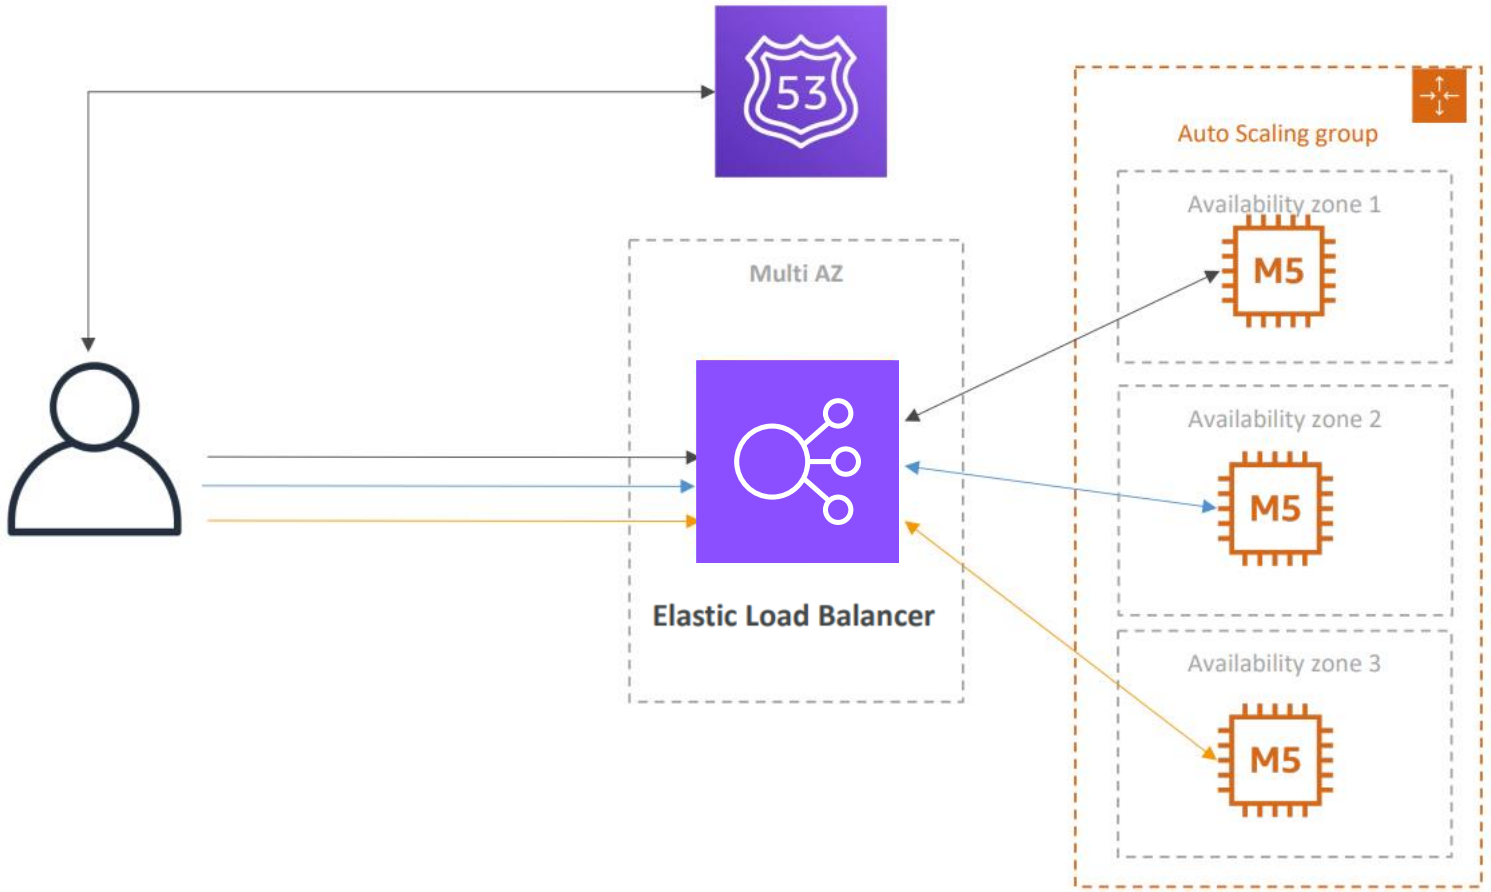

© Copyright by Cloud Mentor Pro | Confidential

What is load balancing?

- Load Balances are servers that forward traffic to multiple servers (e.g., EC2 instances) downstream

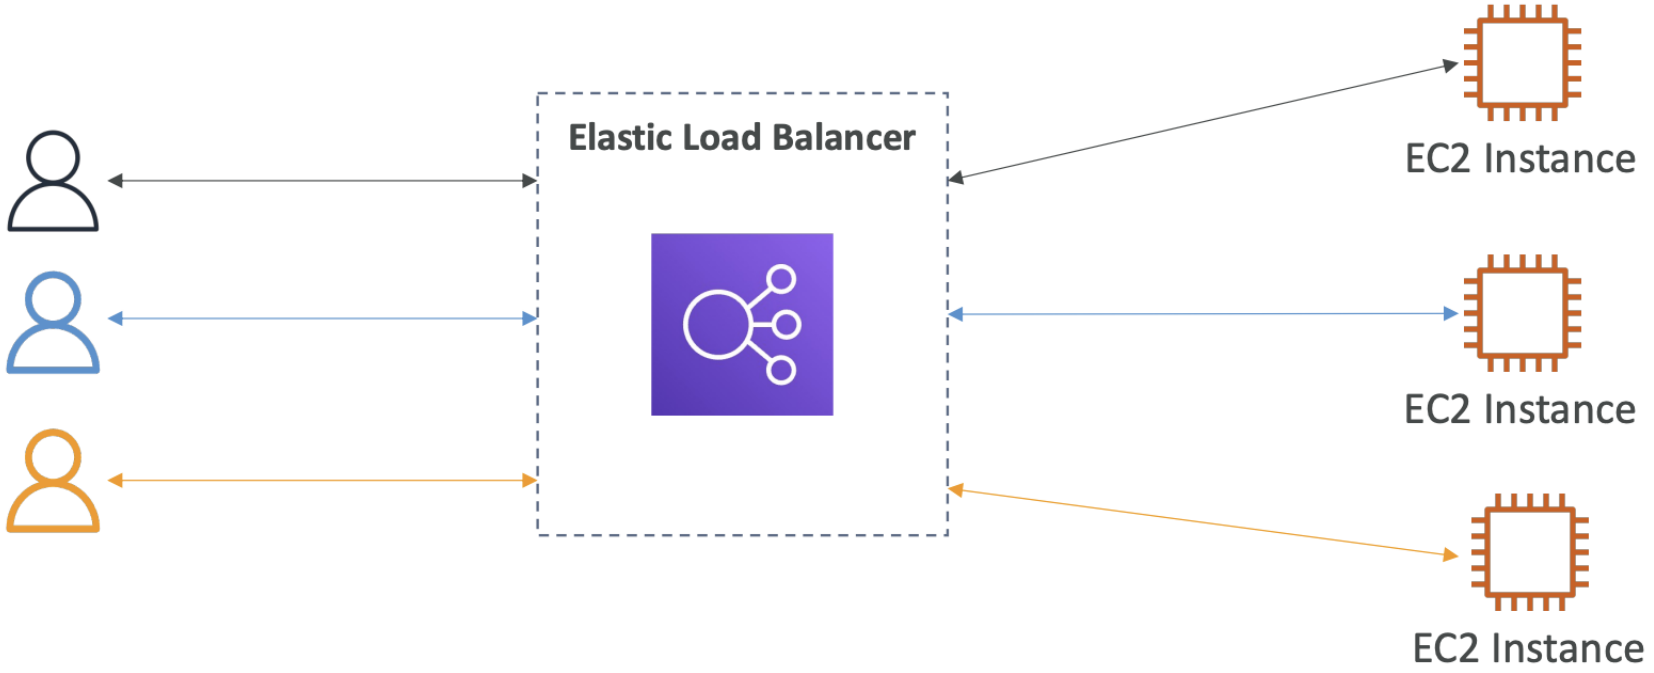

© Copyright by Cloud Mentor Pro | Confidential

Why use a load balancer?

- Spread load across multiple downstream instances
- Expose a single point of access (DNS) to your application
- Seamlessly handle failures of downstream instances
- Do regular health checks to your instances
- Provide SSL termination (HTTPS) for your websites
- Enforce stickiness with cookies
- High availability across zones
- Separate public traffic from private traffic

© Copyright by Cloud Mentor Pro | Confidential

Why use an Elastic Load Balancer?

- An Elastic Load Balancer is a managed load balancer
  - AWS guarantees that it will be working
  - AWS takes care of upgrades, maintenance, high availability
  - AWS provides only a few configuration knobs

- It is integrated with many AWS offerings / services
  - EC2, EC2 Auto Scaling Groups, Amazon ECS
  - AWS Certificate Manager (ACM), CloudWatch
  - Route 53, AWS WAF, AWS Global Accelerator

© Copyright by Cloud Mentor Pro | Confidential

Health Checks

- Health Checks are crucial for Load Balancers
- They enable the load balancer to know if instances it forwards traffic to are available to reply to requests
- The health check is done on a port and a route (**/health** is common)
- If the response is not 200 (OK), then the instance is unhealthy

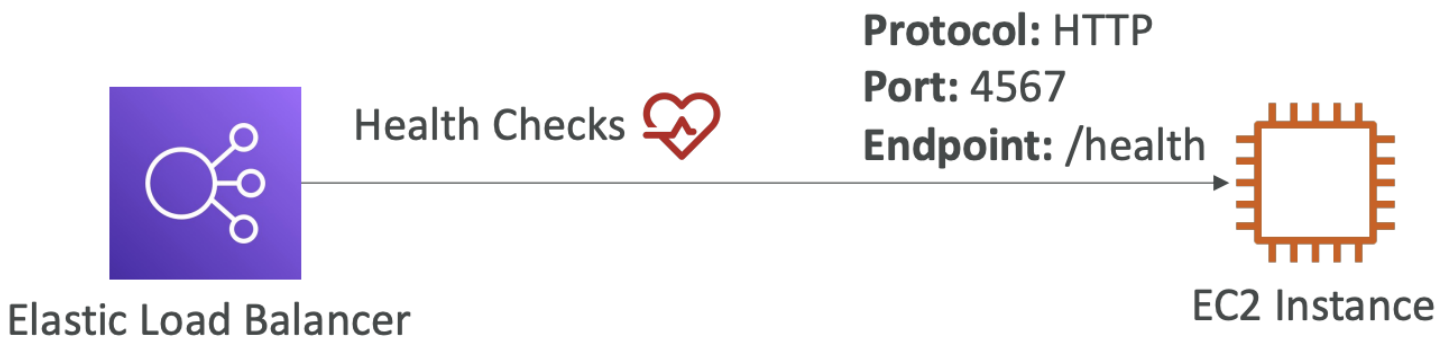

© Copyright by Cloud Mentor Pro | Confidential

Load Balancer Security Groups

Load Balancer Security Group:

Type	Protocol	Port Range	Source	Description
HTTP	TCP	80	0.0.0.0/0	Allow HTTP from an...
HTTPS	TCP	443	0.0.0.0/0	Allow HTTPS from a...

Application Security Group: Allow traffic only from Load Balancer

Type	Protocol	Port Range	Source	Description
HTTP	TCP	80	sg-054b5ff5ea02f2b6e (load-b	Allow Traffic only...

© Copyright by Cloud Mentor Pro | Confidential

Types of load balancer on AWS

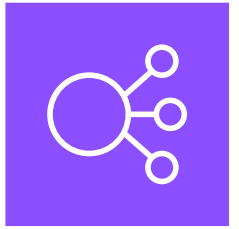

Feature	Application Load Balancer	Network Load Balancer	Gateway Load Balancer	Classic Load Balancer
Load Balancer type	Layer 7	Layer 4	Layer 3 Gateway + Layer 4 Load Balancing	Layer 4/7
Target type	IP, Instance, Lambda	IP, Instance, Application Load Balancer	IP, Instance	
Protocol listeners	HTTP, HTTPS, gRPC	TCP, UDP, TLS	IP	TCP, SSL/TLS, HTTP, HTTPS

© Copyright by Cloud Mentor Pro | Confidential

https://aws.amazon.com/elasticloadbalancing/features/#Product_comparisons

Application Load Balancer (v2)

- Application load balancers is Layer 7 (HTTP) of OSI model (**support HTTP/HTTPS**)

- Load balancing to multiple HTTP applications across machines (target groups)
- Load balancing to multiple applications on the same machine (ex: containers)
- Support for HTTP/2 and WebSocket
- Support redirects (from HTTP to HTTPS for example)

© Copyright by Cloud Mentor Pro | Confidential

Application Load Balancer (v2)

- Routing tables to different target groups:
  - Routing based on **path** in URL (example.com/users & example.com/posts)
  - Routing based on **hostname** in URL (one.example.com & other.example.com)
  - Routing based on **Query String, Headers** (example.com/users?id=123&order=false)

- ALB are a great fit for **micro services & container-based application** (example: Docker & Amazon ECS)
- Has a port mapping feature to redirect to a **dynamic port in ECS**

© Copyright by Cloud Mentor Pro | Confidential

Application Load Balancer (v2) Target Groups

- EC2 instances (can be managed by an Auto Scaling Group) – HTTP
- ECS tasks (managed by ECS itself) – HTTP
- Lambda functions – HTTP request is translated into a JSON event
- IP Addresses – must be private IPs

- ALB can route to multiple target groups
- Health checks are at the target group level

© Copyright by Cloud Mentor Pro | Confidential
https://docs.aws.amazon.com/elasticloadbalancing/latest/application/load-balancer-target-groups.html

Application Load Balancer (v2)
HTTP Based Traffic

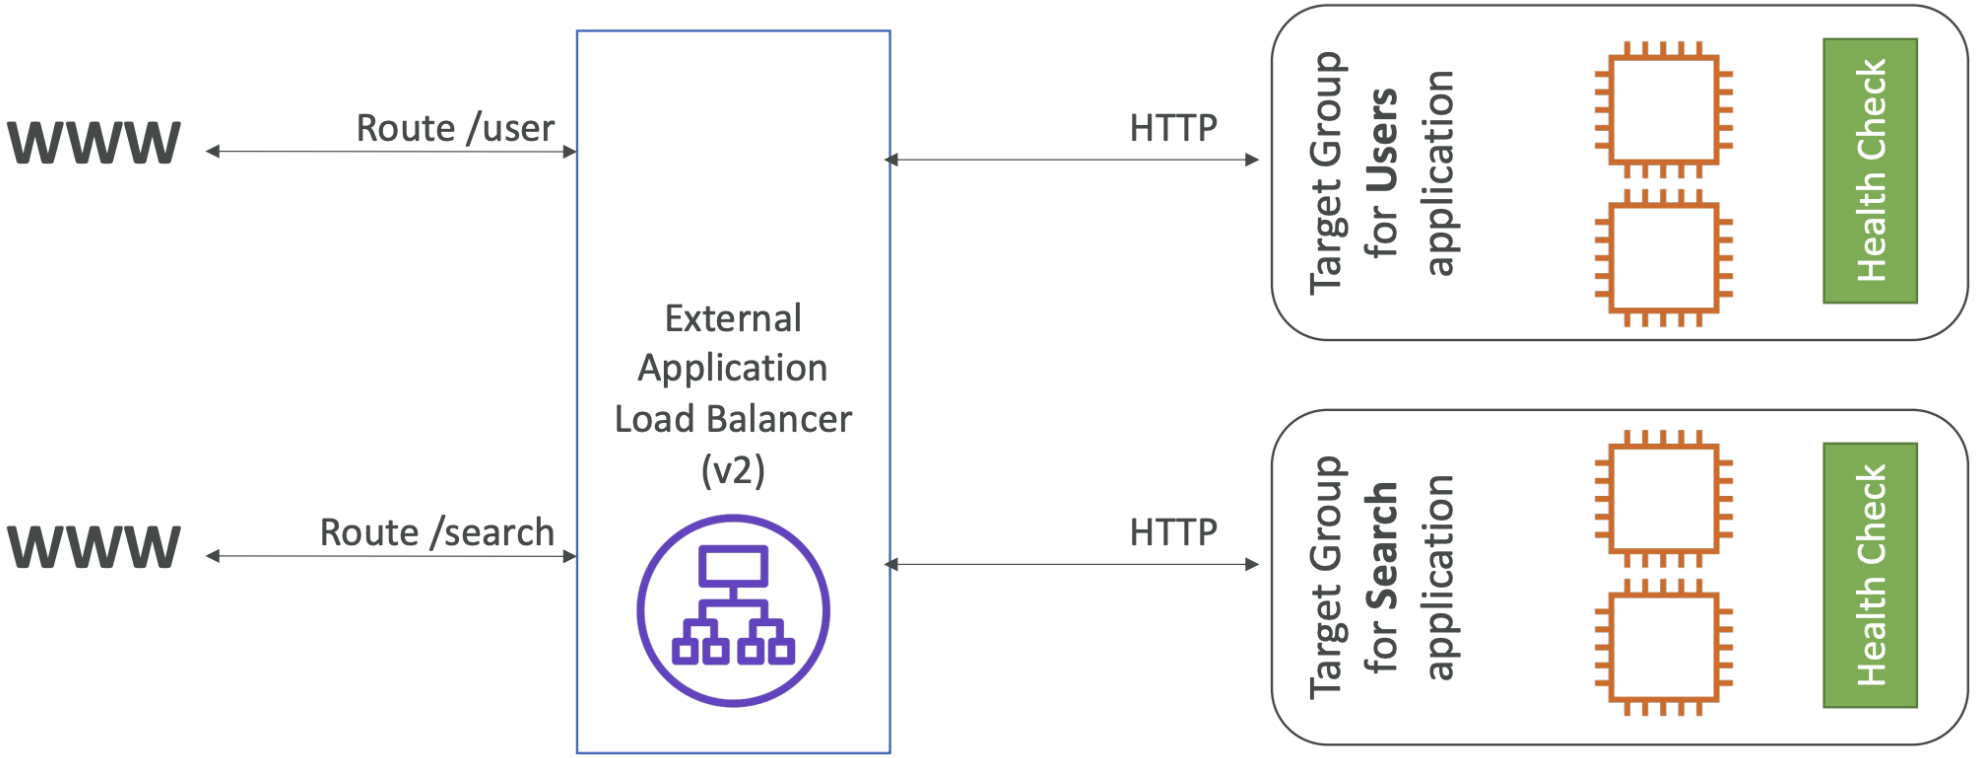

© Copyright by Cloud Mentor Pro | Confidential

Application Load Balancer (v2) Query Strings/Parameters Routing

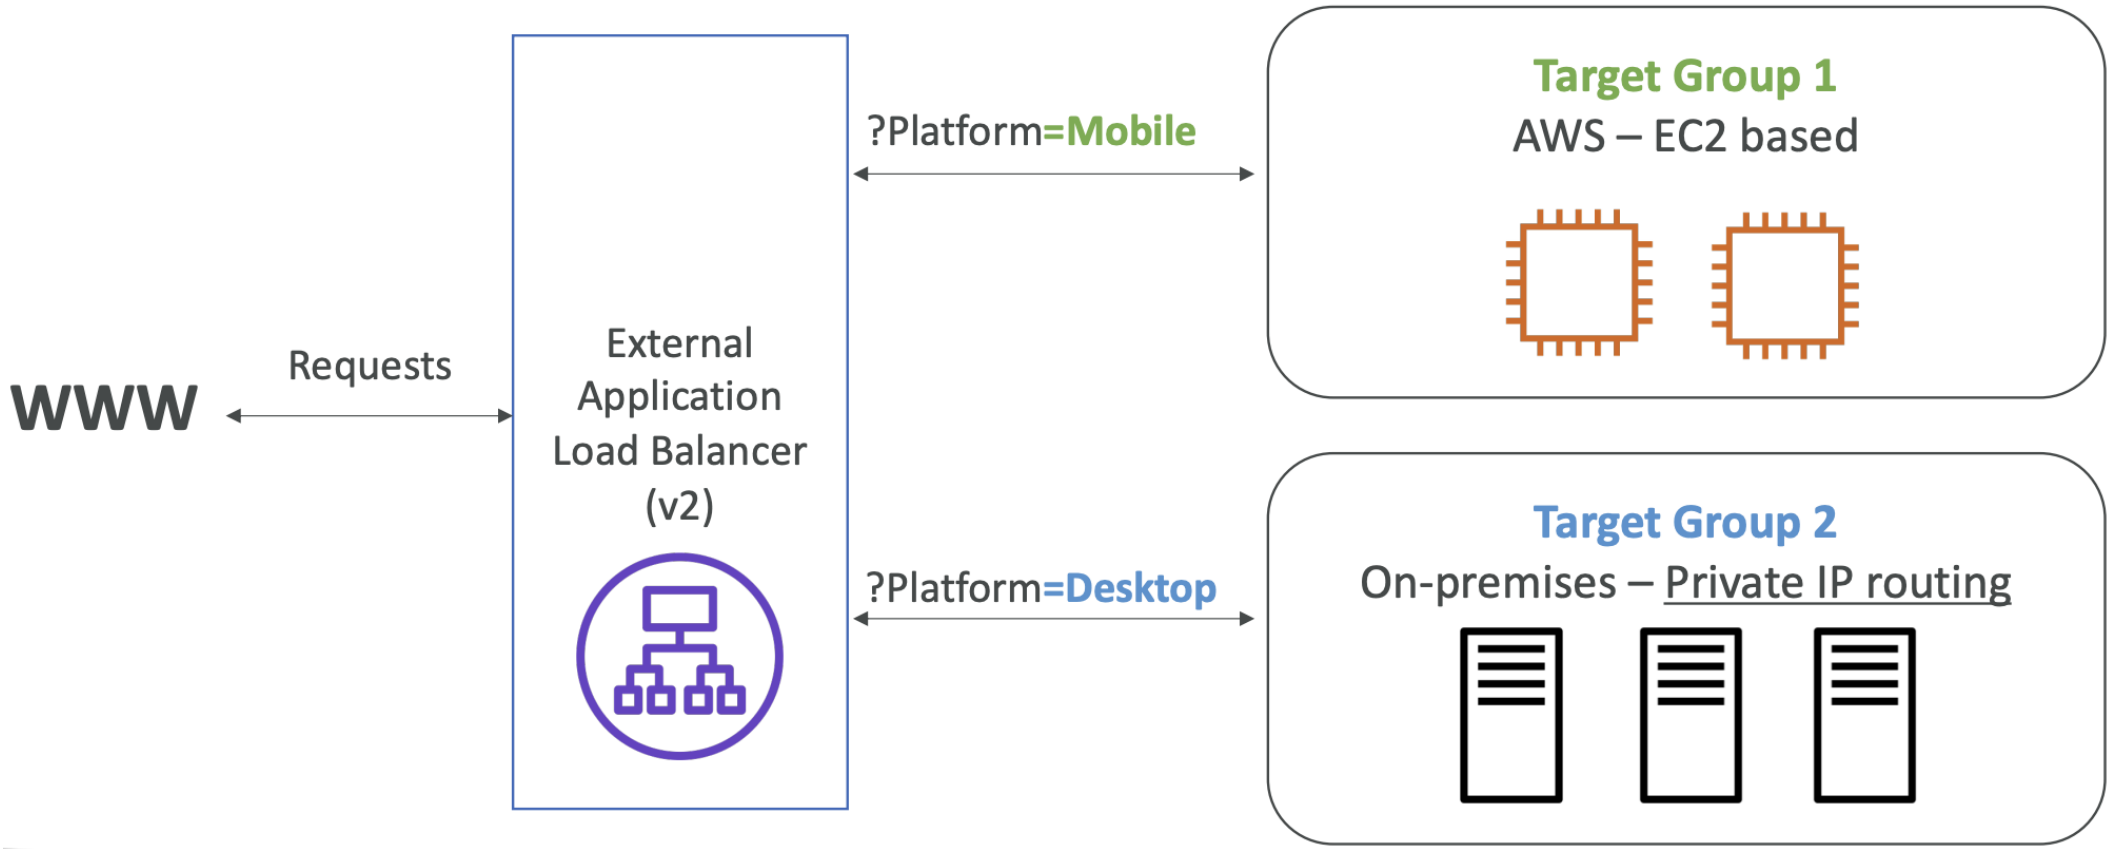

© Copyright by Cloud Mentor Pro | Confidential

Application Load Balancer (v2) Good to Know

- Fixed hostname (XXX.region.elb.amazonaws.com)
- The application servers don’t see the IP of the client directly
  - The true IP of the client is inserted in the header **X-Forwarded-For**

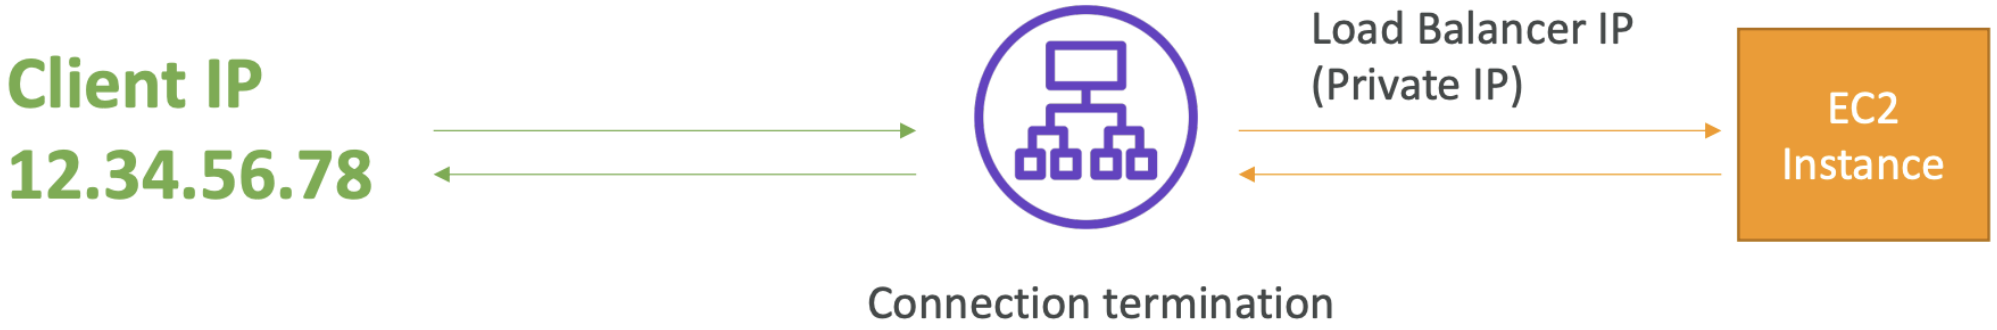

© Copyright by Cloud Mentor Pro | Confidential

Network Load Balancer (v2)

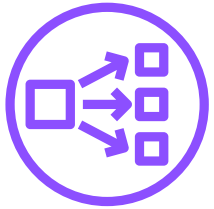

- Network load balancers (Layer 4) allow to:
  - Forward TCP & UDP traffic to your instances
  - Handle millions of request per seconds
  - **Less latency ~100 ms** (vs 400 ms for ALB)

- **NLB has one static IP per AZ, and supports assigning Elastic IP** (helpful for whitelisting specific IP)

- NLB are used for extreme performance, TCP or UDP traffic
- Not included in the AWS free tier

© Copyright by Cloud Mentor Pro | Confidential

Network Load Balancer –Target Groups

- EC2 instances
- IP Addresses – must be private IPs
- Application Load Balancer
- Health Checks support the TCP, HTTP and HTTPS Protocols

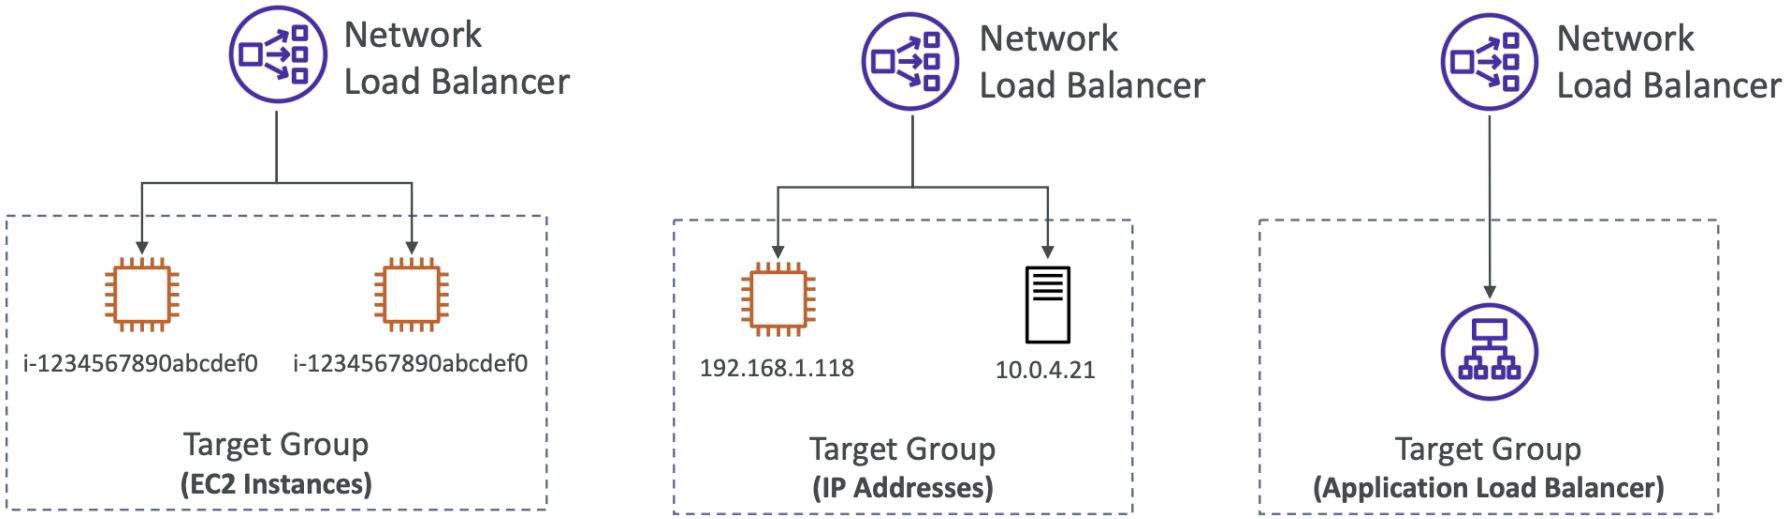

© Copyright by Cloud Mentor Pro | Confidential

Network Load Balancer (v2)
TCP (Layer 4) Based Traffic

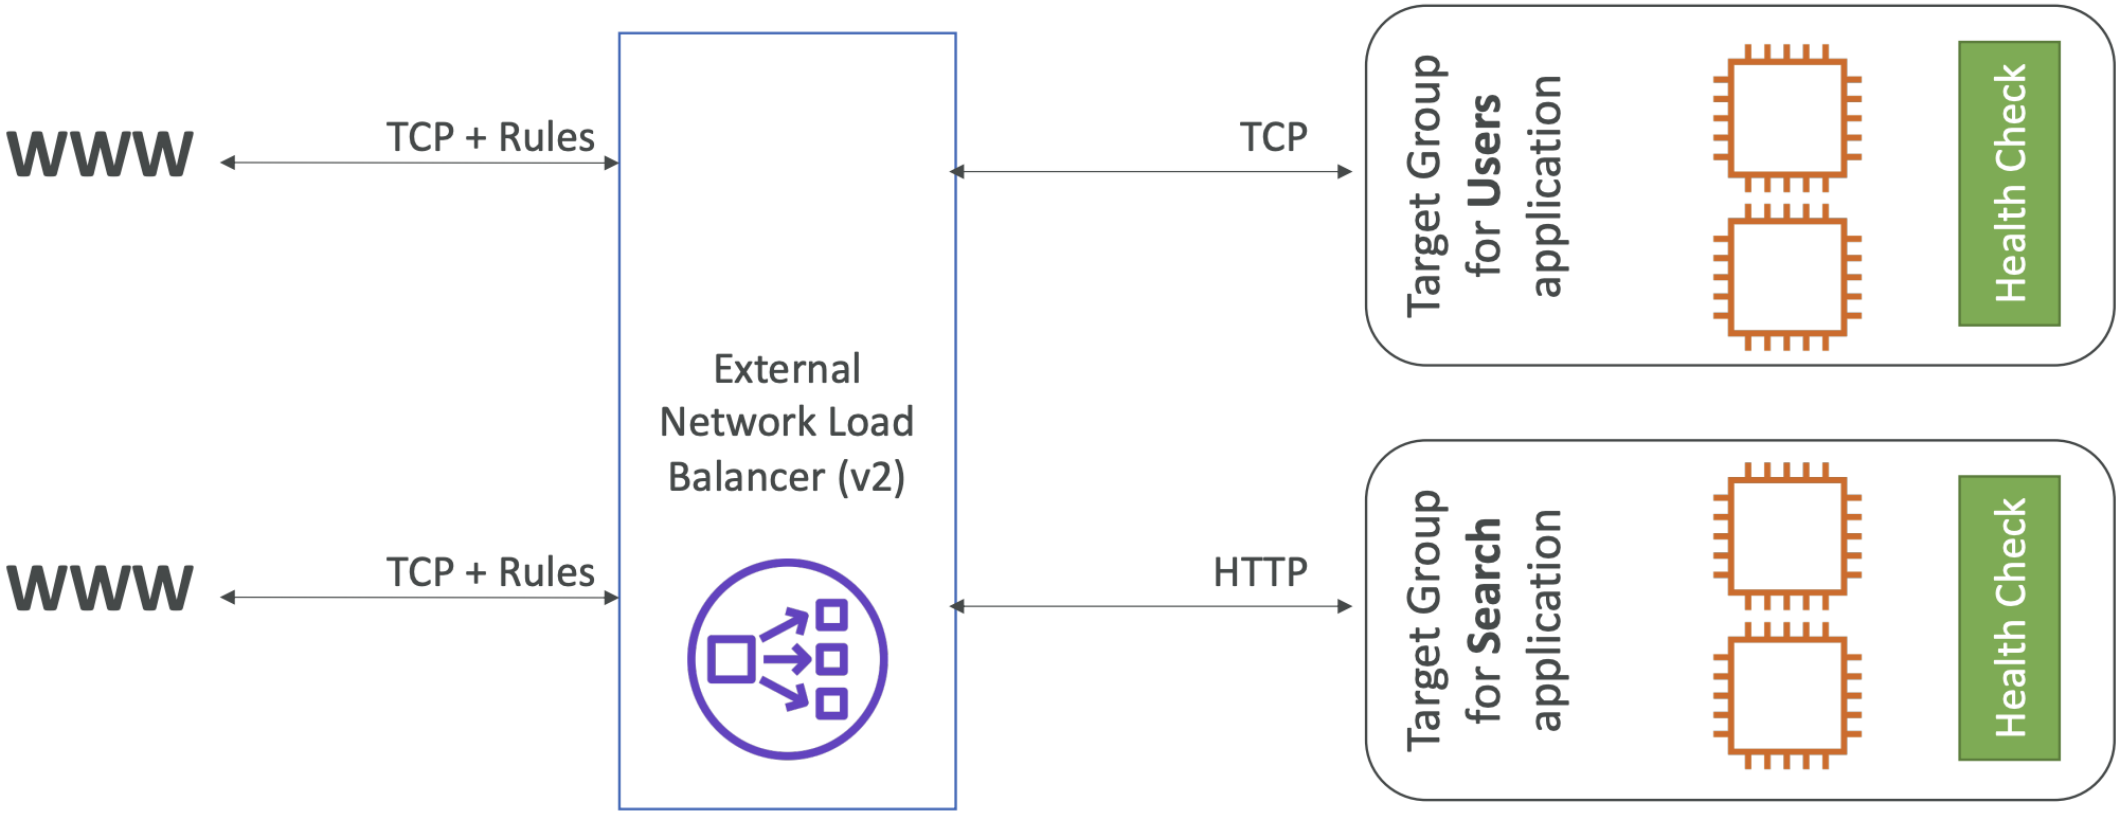

© Copyright by Cloud Mentor Pro | Confidential

Gateway Load Balancer

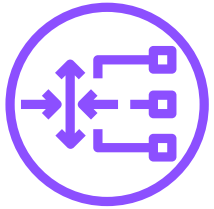

- Deploy, scale, and manage a fleet of 3rd party network virtual appliances in AWS
- Example: Firewalls, Intrusion Detection and Prevention Systems, Deep Packet Inspection Systems, payload manipulation, …

- Operates at Layer 3 (Network Layer) – IP Packets
- Combines the following functions:
  - **Transparent Network Gateway** – single entry/exit for all traffic
  - **Load Balancer** – distributes traffic to your virtual appliances
- Uses the GENEVE protocol on port 6081

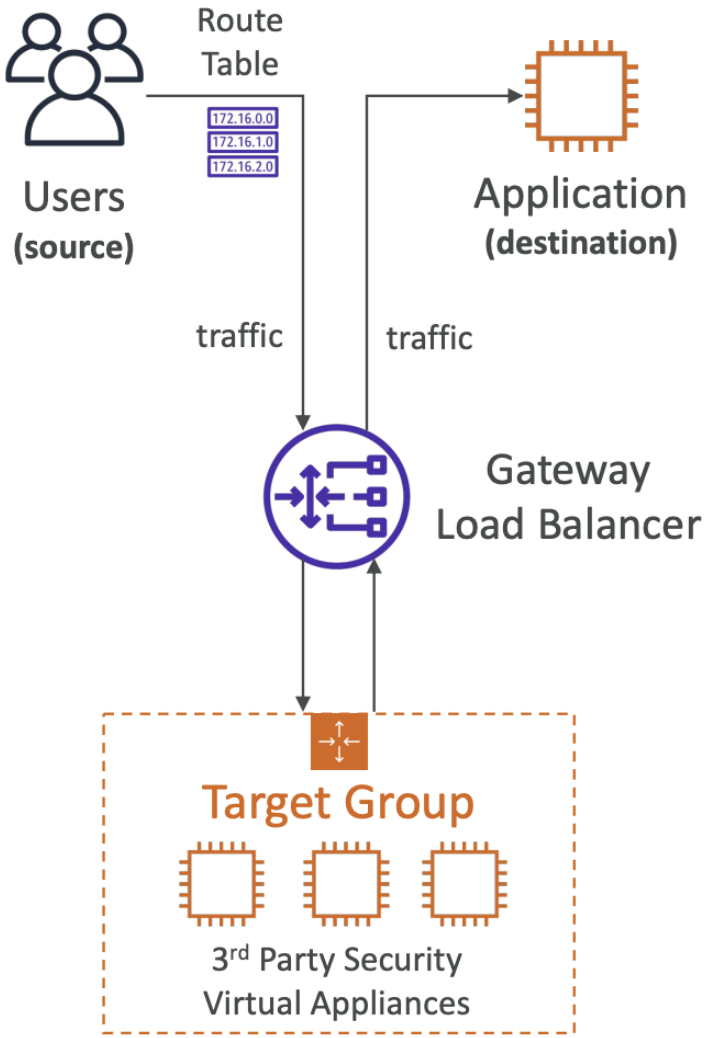

© Copyright by Cloud Mentor Pro | Confidential

Gateway Load Balancer –Target Groups

- EC2 instances
- IP Addresses – must be private IPs

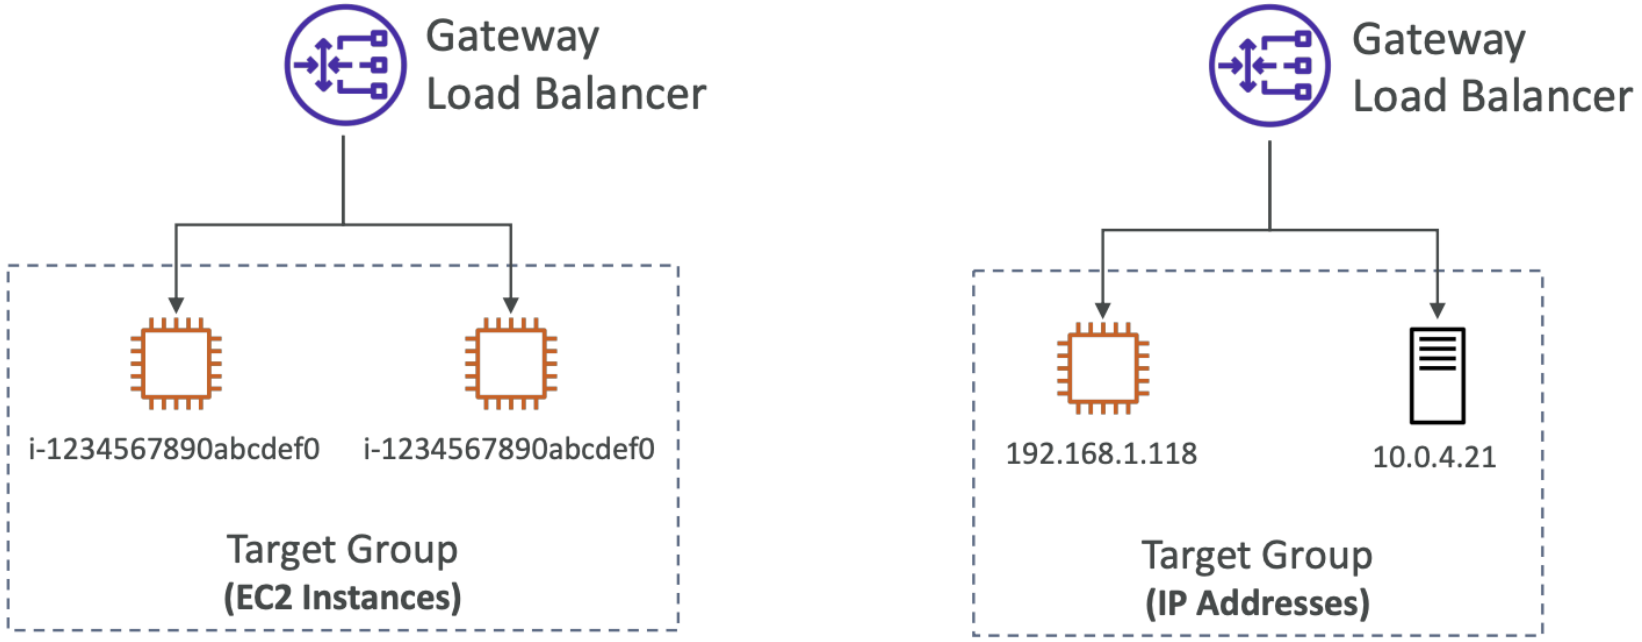

© Copyright by Cloud Mentor Pro | Confidential

Sticky Sessions (Session Affinity)

- It is possible to implement stickiness so that the same client is always redirected to the same instance behind a load balancer
- This works for **Classic Load Balancer, Application Load Balancer, and Network Load Balancer**
- For both CLB & ALB, the "cookie" used for stickiness has an expiration date you control
- **Use case:** make sure the user doesn't lose his session data
- Enabling stickiness may bring imbalance to the load over the backend EC2 instances

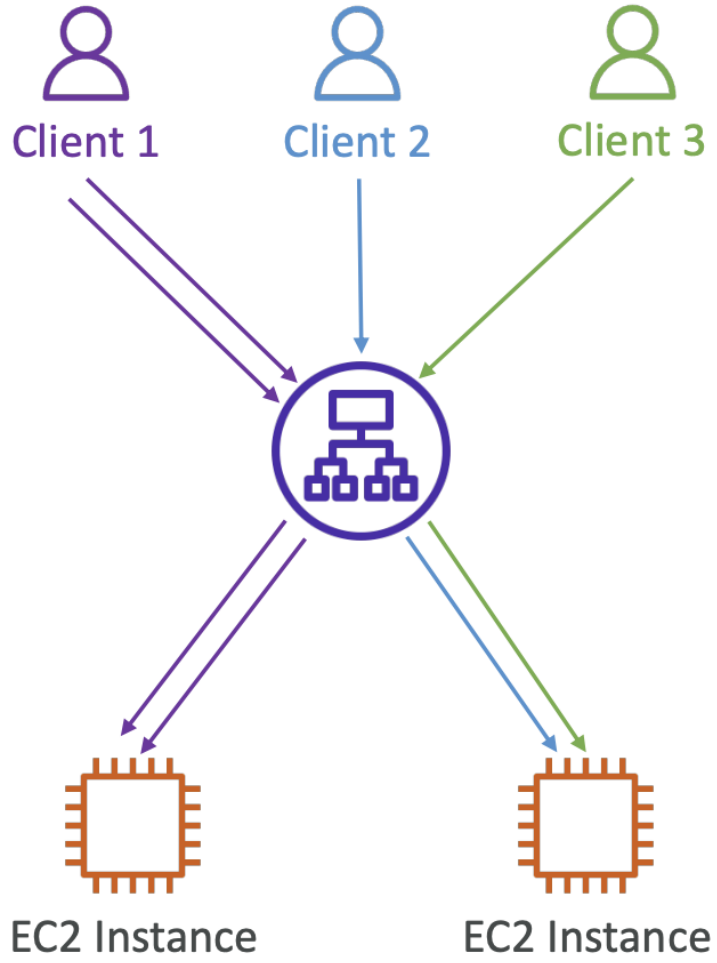

© Copyright by Cloud Mentor Pro | Confidential

Cross-Zone Load Balancing

With Cross Zone Load Balancing:

each load balancer instance distributes evenly across all registered instances in all AZ

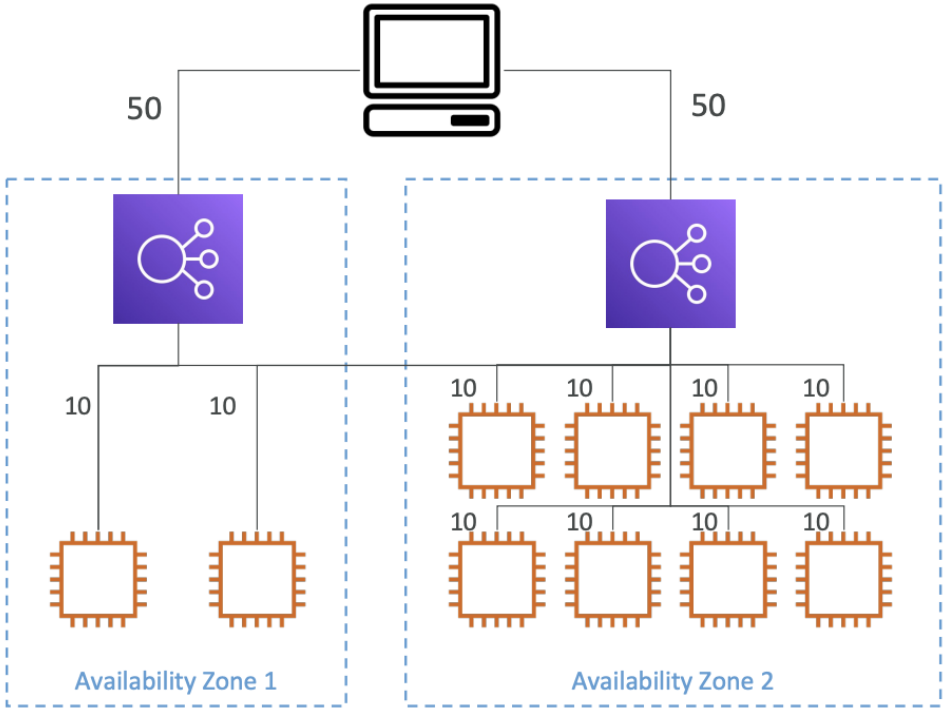

Without Cross Zone Load Balancing:

Requests are distributed in the instances of the node of the Elastic Load Balancer

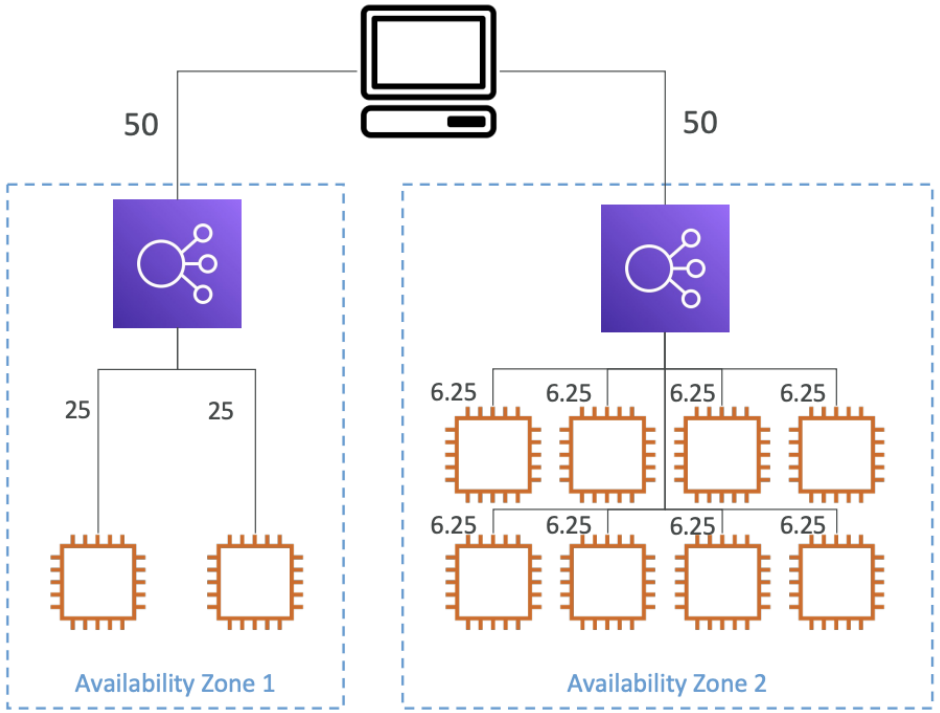

© Copyright by Cloud Mentor Pro | Confidential

Cross -Zone Load Balancing

- Application Load Balancer
  - Enabled by default (can be disabled at the Target Group level)
  - No charges for inter AZ data

- Network Load Balancer & Gateway Load Balancer
  - Disabled by default
  - You pay charges ($) for inter AZ data if enabled

© Copyright by Cloud Mentor Pro | Confidential

SSL/TLS - Basics

- An SSL Certificate allows traffic between your clients and your load balancer to be encrypted in transit (in-flight encryption)

- **SSL** refers to Secure Sockets Layer, used to encrypt connections
- **TLS** refers to Transport Layer Security, which is a newer version
- Nowadays, **TLS certificates are mainly used**, but people still refer as SSL

- Public SSL certificates are issued by Certificate Authorities (CA)
- Comodo, Symantec, GoDaddy, GlobalSign, Digicert, Letsencrypt, etc…

- SSL certificates have an expiration date (you set) and must be renewed

© Copyright by Cloud Mentor Pro | Confidential

Load Balancer - SSL Certificates

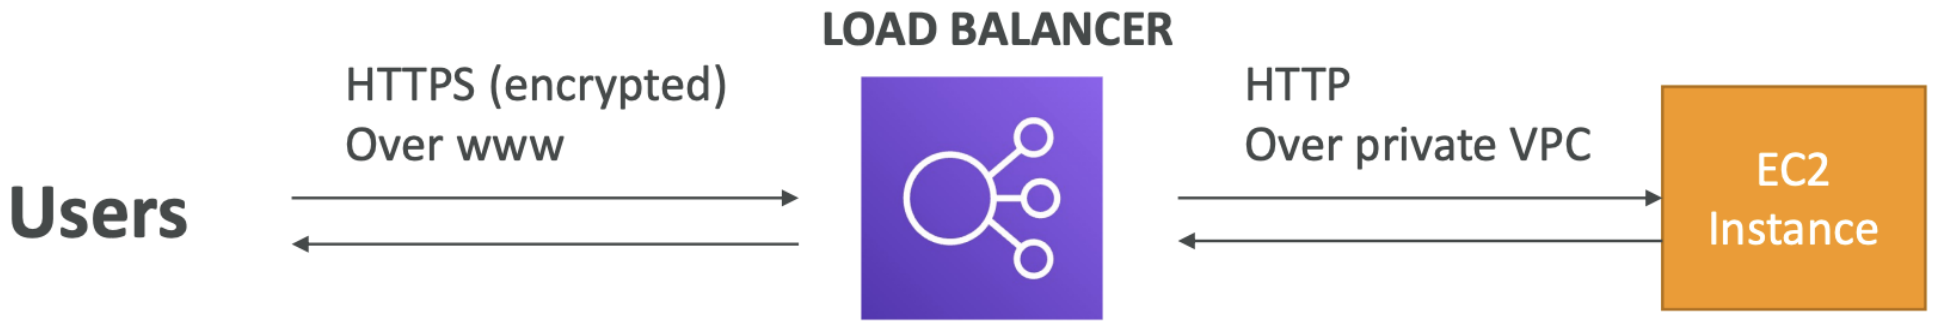

- The load balancer uses an X.509 certificate (SSL/TLS server certificate)
- You can manage certificates using ACM (AWS Certificate Manager)
- You can create upload your own certificates alternatively
- HTTPS listener:
  - You must specify a default certificate
  - You can add an optional list of certs to support multiple domains
  - **Clients can use SNI (Server Name Indication) to specify the hostname they reach**

© Copyright by Cloud Mentor Pro | Confidential

SSL – Server Name Indication (SNI)

- SNI solves the problem of loading **multiple SSL certificates onto one web server** (to serve multiple websites)
- **Note:**
  - Only works for ALB & NLB (newer generation), CloudFront
  - Does not work for CLB (older gen)

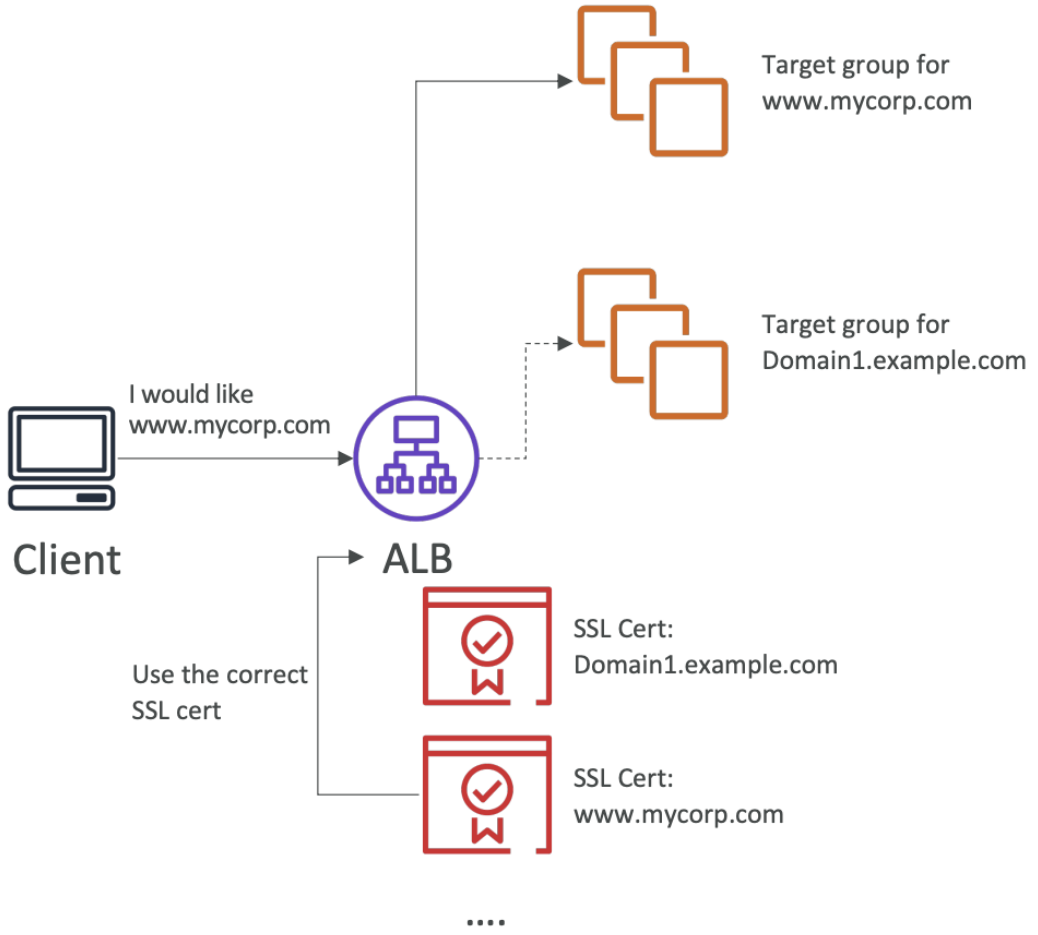

© Copyright by Cloud Mentor Pro | Confidential

Connection Draining

- **Feature naming**
  - Connection Draining – for CLB
  - Deregistration Delay – for ALB & NLB

- Time to complete "in-flight requests" while the instance is de-registering or unhealthy
- Stops sending new requests to the EC2 instance which is de-registering
- Between 1 to 3600 seconds (default: 300 seconds)
- Can be disabled (set value to 0)
- Set to a low value if your requests are short

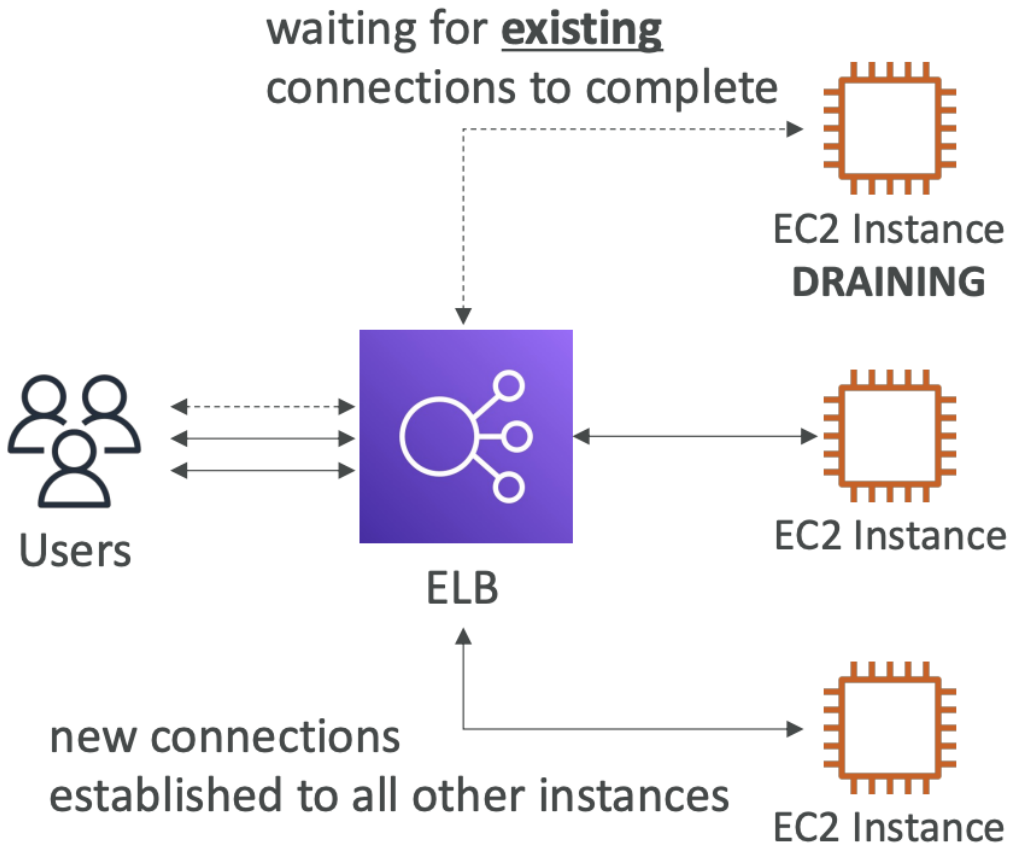

© Copyright by Cloud Mentor Pro | Confidential

What's an Auto Scaling Group?

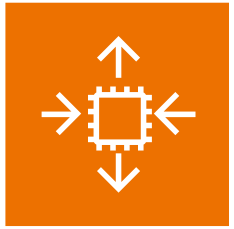

- In real-life, the load on your websites and application can change

- The goal of an Auto Scaling Group (ASG) is to:
  - Scale out (add EC2 instances) to match an increased load
  - Scale in (remove EC2 instances) to match a decreased load
  - Ensure we have a minimum and a maximum number of EC2 instances running
  - **Automatically register** new instances to a load balancer
  - **Re-create** an EC2 instance in case a previous one is terminated (ex: if unhealthy)

- ASG are free (you only pay for the underlying EC2 instances)

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling Group in AWS

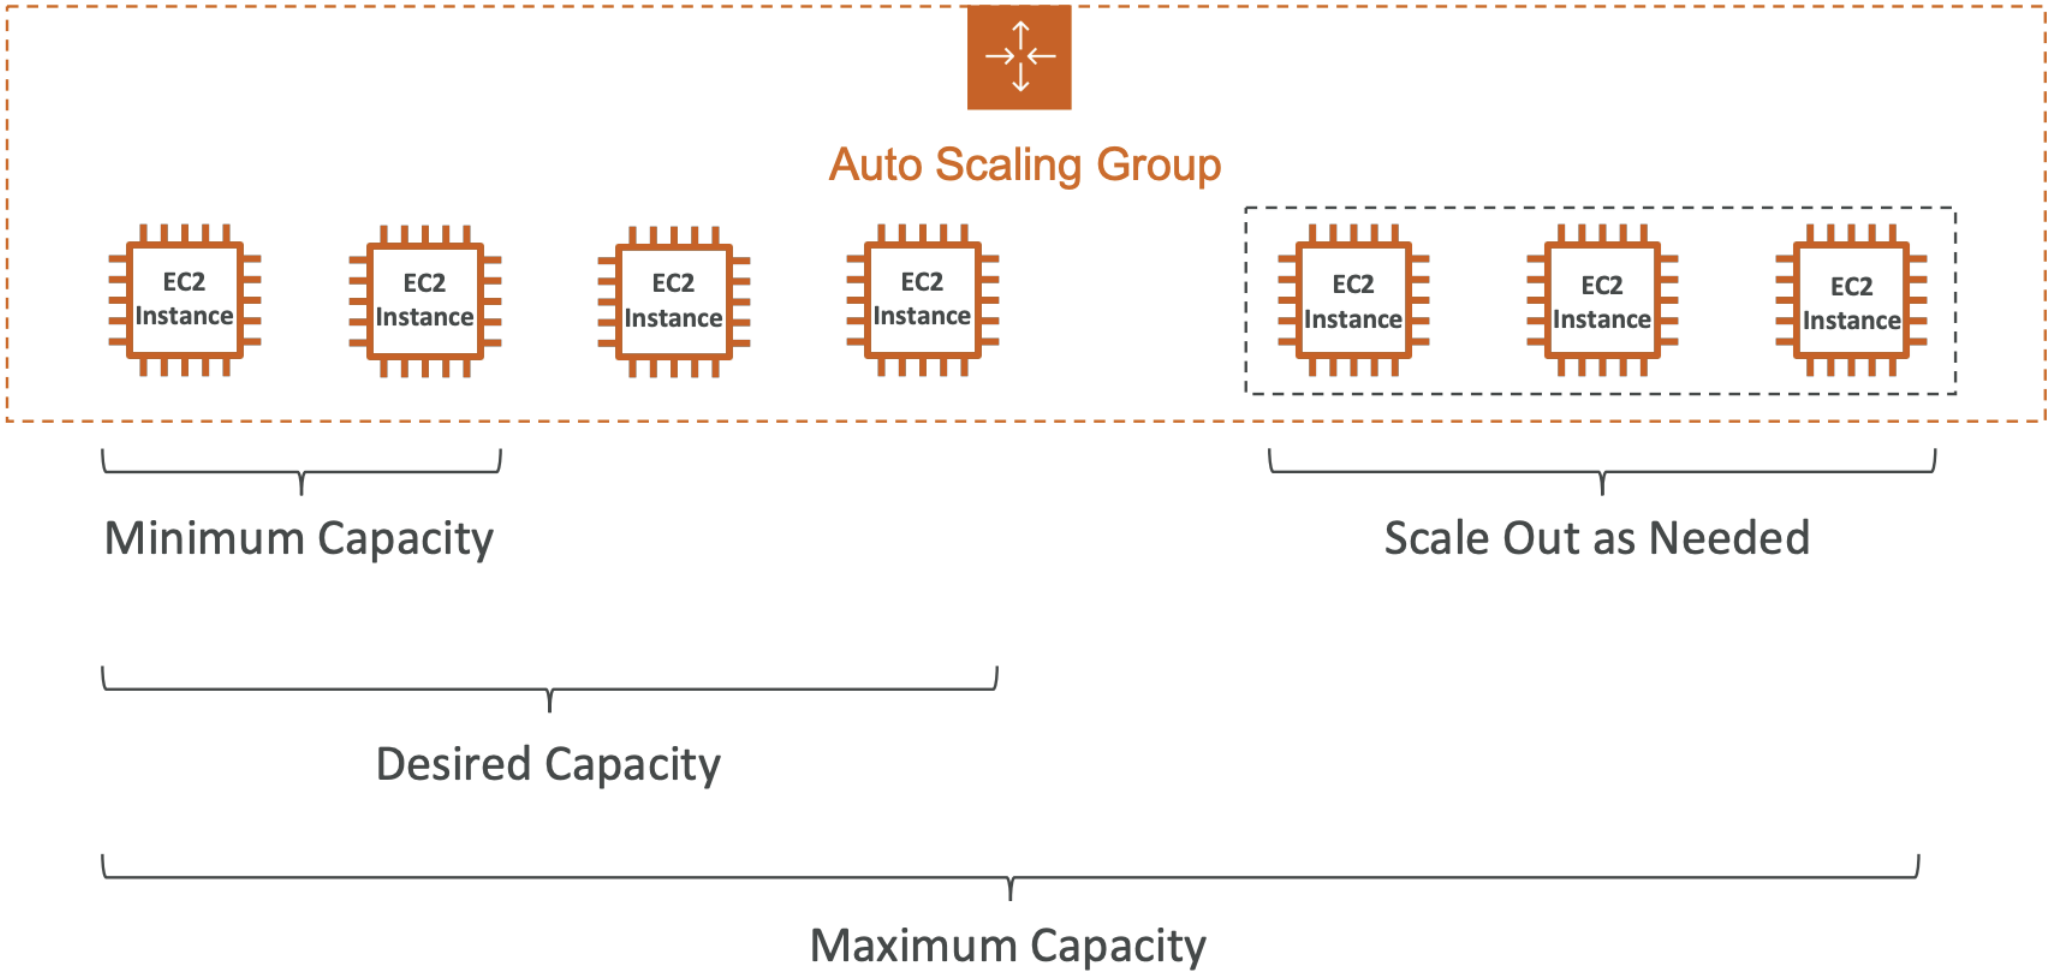

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling Group in AWS With Load Balancer

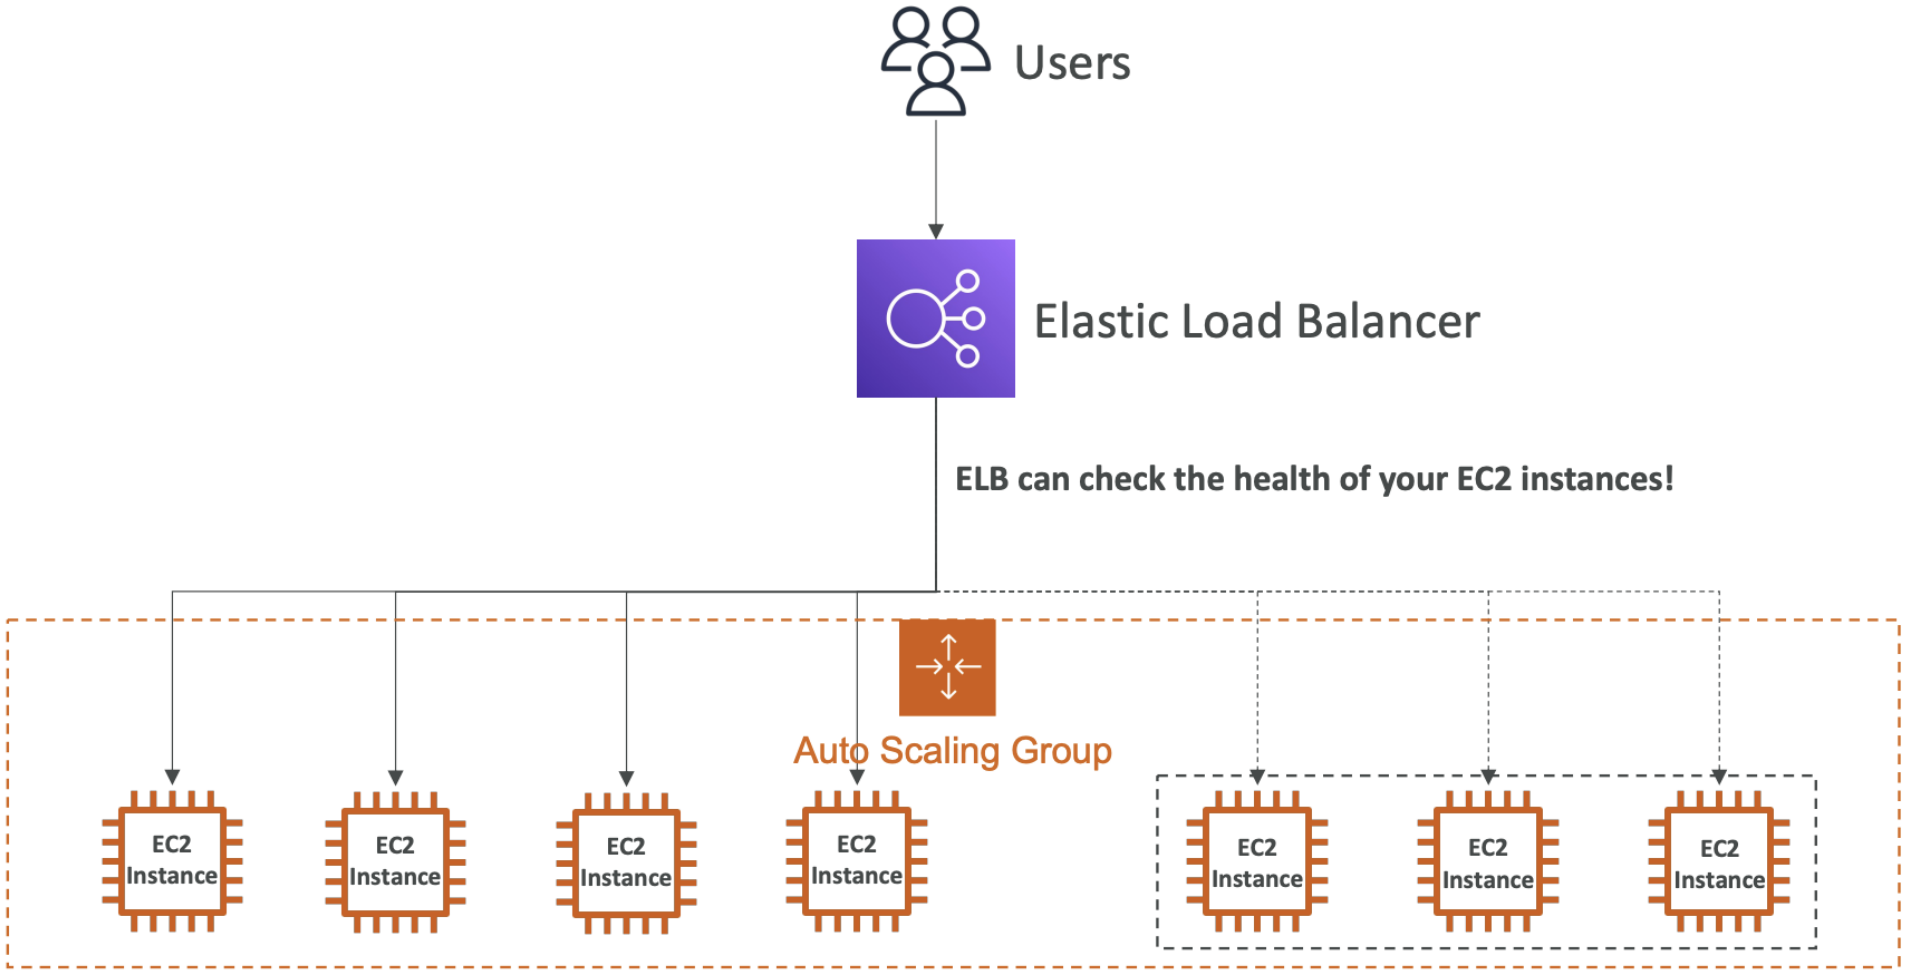

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling Group Attributes

- A Launch Template (older “Launch Configurations” are deprecated)
  - AMI + Instance Type
  - EC2 User Data
  - EBS Volumes
  - Security Groups
  - SSH Key Pair
  - IAM Roles for your EC2 Instances
  - Network + Subnets Information
  - Load Balancer Information
- Min Size / Max Size / Initial Capacity
- Scaling Policies

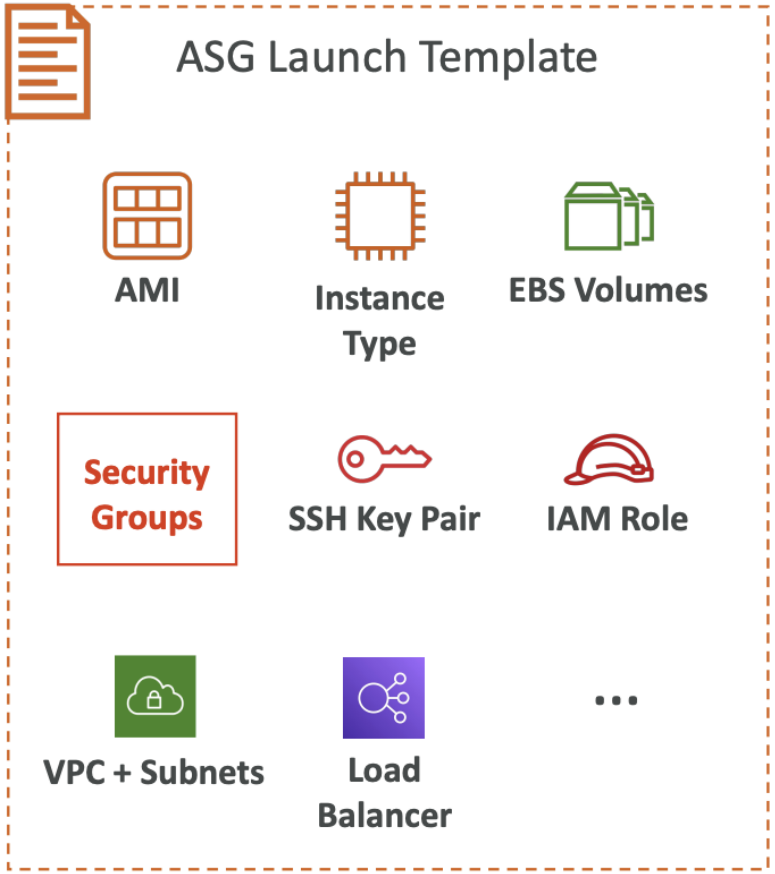

https://docs.aws.amazon.com/autoscaling/ec2/userguide/launch-templates.html

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling - CloudWatch Alarms & Scaling

- It is possible to scale an ASG based on CloudWatch alarms
- An alarm monitors a metric (such as **Average CPU**, or a **custom metric**)
- Metrics such as Average CPU are computed for the overall ASG instances
- Based on the alarm:
  - We can create scale-out policies (increase the number of instances)
  - We can create scale-in policies (decrease the number of instances)

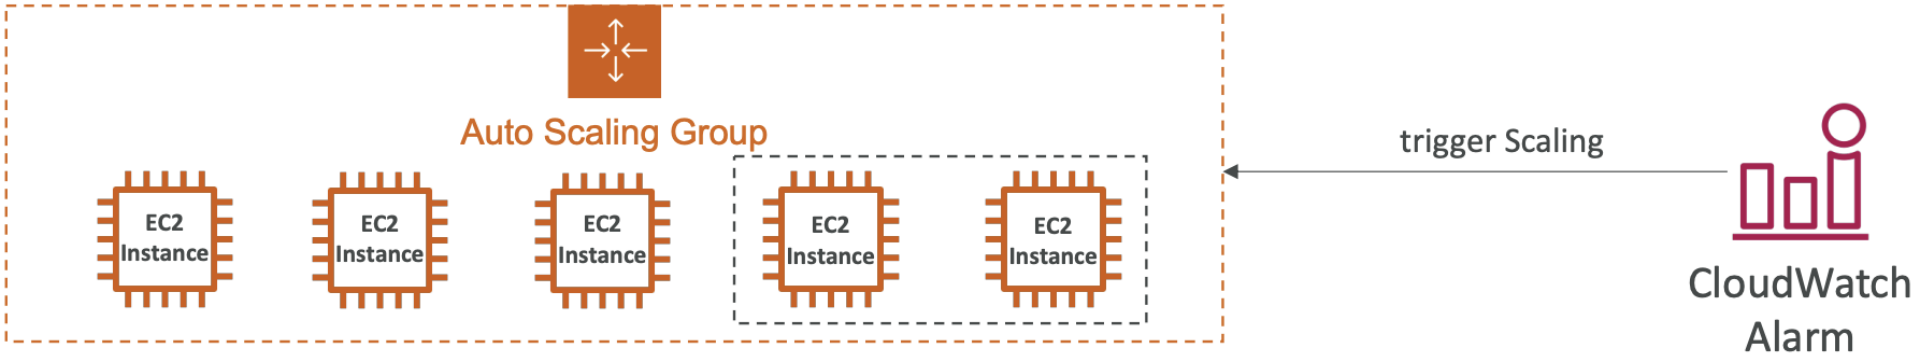

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling Groups – Scaling Policies

- Dynamic Scaling
  - Target Tracking Scaling (strongly recommend)
    - Scale a resource based on a **target value for a specific CloudWatch metric**
    - Example: I want the average ASG CPU to stay at around 40%
  - Simple / Step Scaling
    - When a CloudWatch alarm is triggered (example CPU > 70%), then add 2 units
    - When a CloudWatch alarm is triggered (example CPU < 30%), then remove 1
- Scheduled Scaling
  - Scale based on known usage patterns
  - Example: increase the min capacity to 10 at 5 pm on Fridays

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling Groups – Scaling Policies

- Predictive scaling:
  - Uses **machine learning** to predict capacity requirements based on historical data from CloudWatch

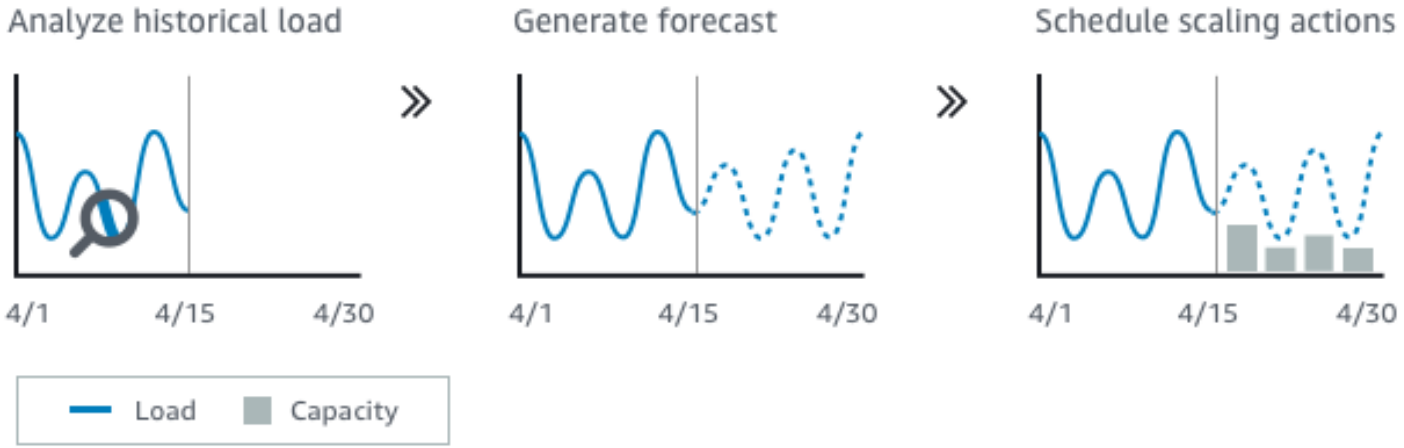

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/autoscaling/plans/userguide/how-it-works.html

Good metrics to scale on

- **CPUUtilization**: Average CPU utilization across your instances
- **RequestCountPerTarget**: to make sure the number of requests per EC2 instances is stable
- **Average Network In / Out** (if you're application is network bound)
- **Any custom metric** (that you push using CloudWatch)

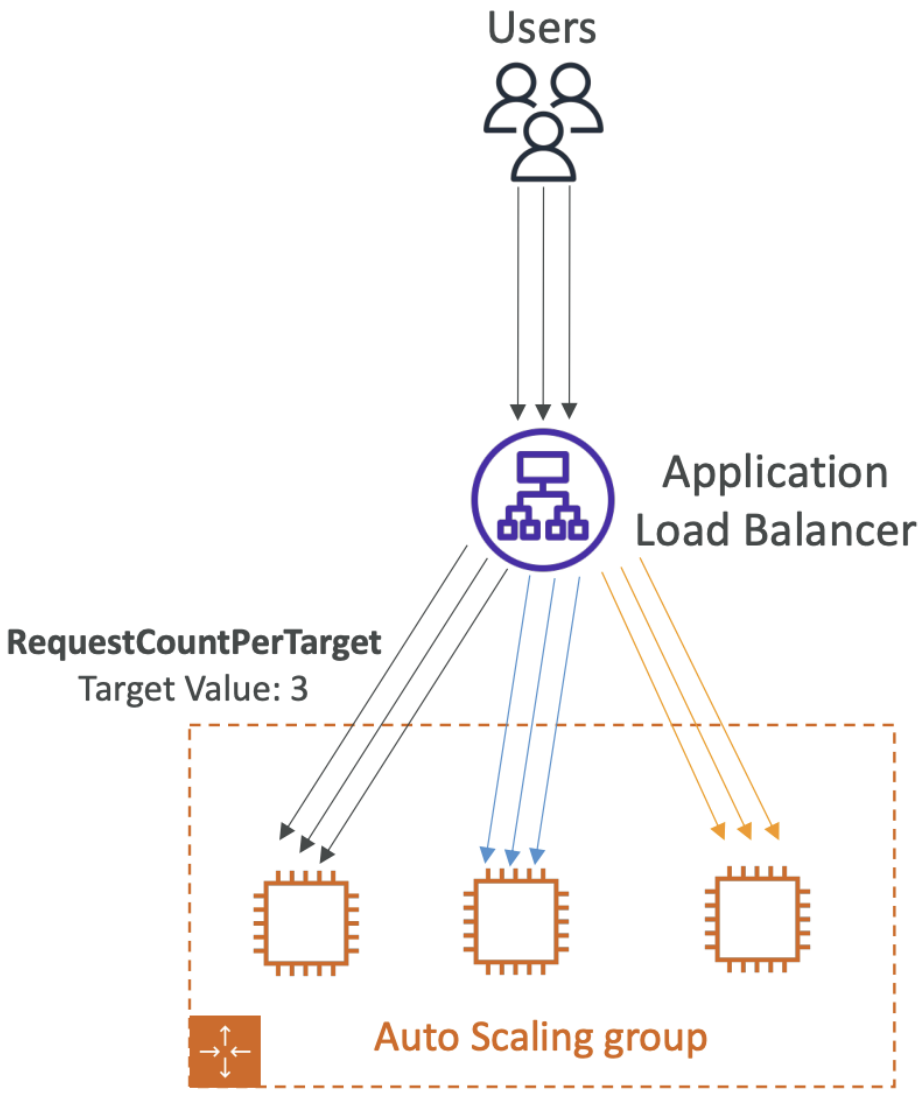

© Copyright by Cloud Mentor Pro | Confidential

Auto Scaling Groups - Scaling Cooldowns

- After a scaling activity happens, you are in the **cooldown period (default 300 seconds)**
- During the cooldown period, the ASG will not launch or terminate additional instances (to allow for metrics to stabilize)

- Advice: Use a ready-to-use AMI to reduce configuration time in order to be serving request fasters and reduce the cooldown period

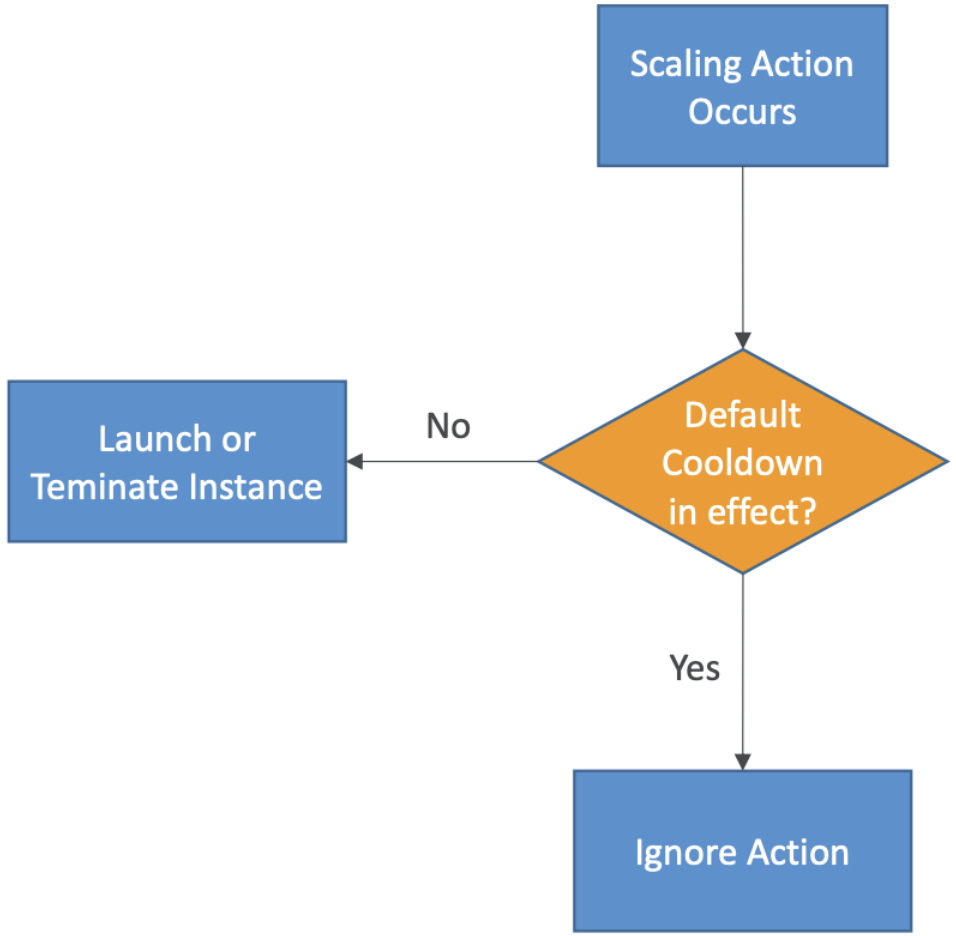

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/autoscaling/ec2/userguide/ec2-auto-scaling-scaling-cooldowns.html

Auto Scaling – Instance Refresh

- Goal: update launch template and then re-creating all EC2 instances
- For this we can use the native feature of Instance Refresh
- Setting of minimum healthy percentage
- Specify warm-up time (how long until the instance is ready to use)

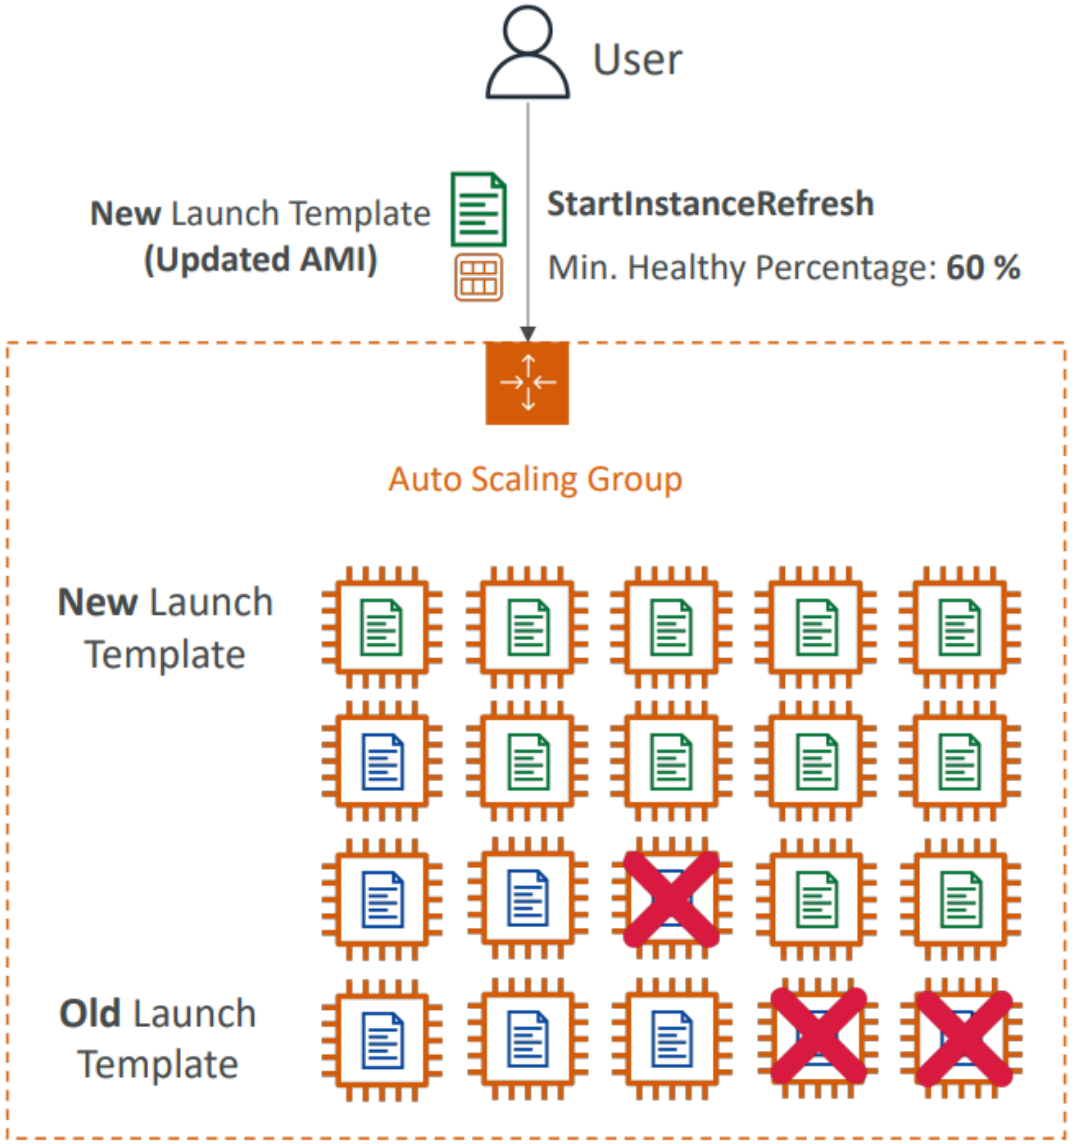

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- *Application Load Balancer*
- *Application Load Balancer and EC2 Auto Scaling Group*

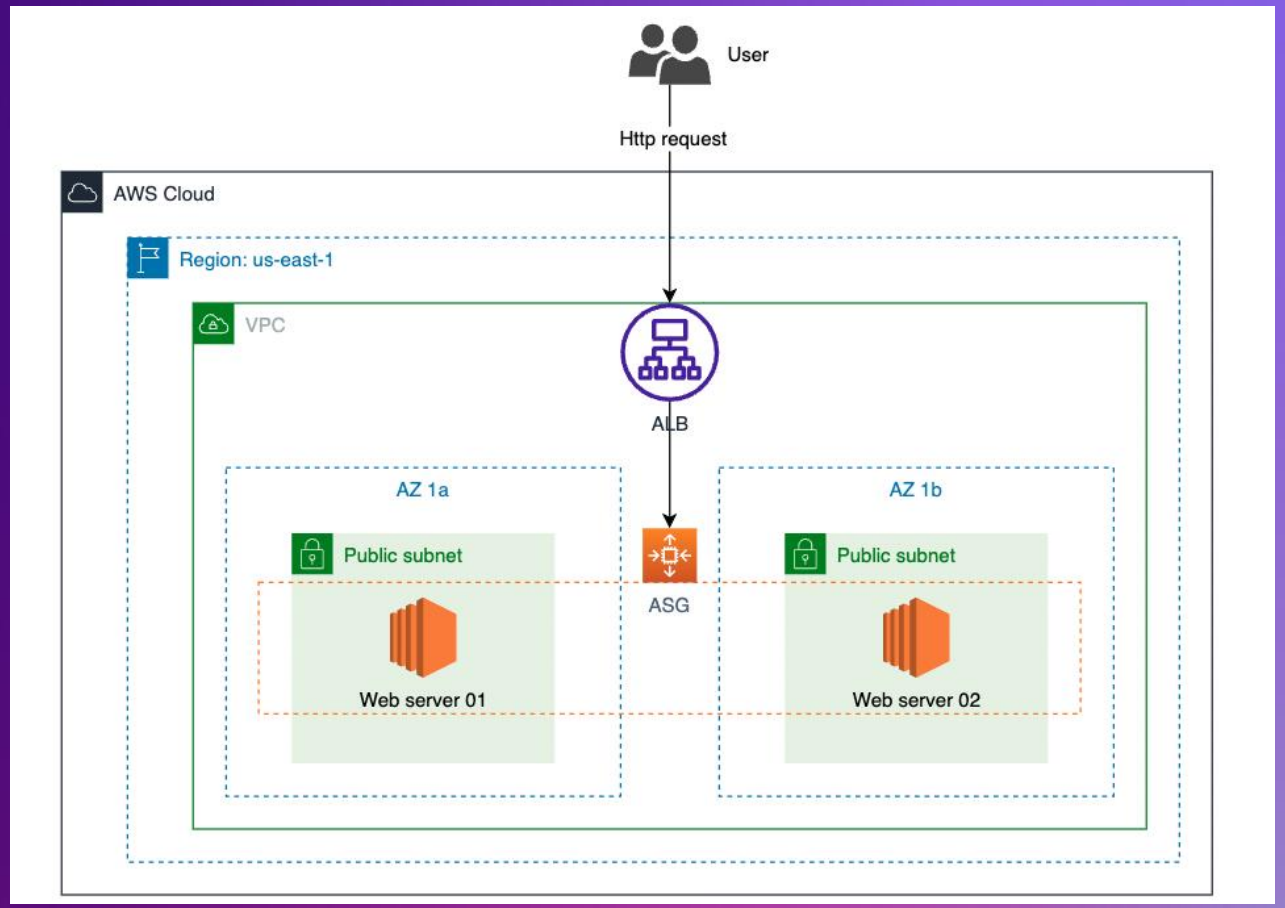

© Copyright by Cloud Mentor Pro | Confidential

Section 3

- *RDS, Aurora & ElastiCache*
- *Route53*

© Copyright by Cloud Mentor Pro | Confidential

RDS, Aurora & ElastiCache

© Copyright by Cloud Mentor Pro | Confidential

Amazon RDS Overview

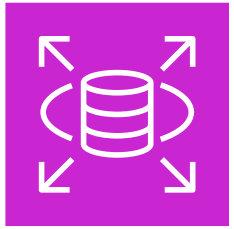

- RDS stands for Relational Database Service
- It’s a managed DB service for DB use SQL as a query language
- It allows you to create databases in the cloud that are managed by AWS
  - Postgres
  - MySQL
  - MariaDB
  - Oracle
  - Microsoft SQL Server
  - IBM DB2
  - Aurora (AWS Proprietary database)

© Copyright by Cloud Mentor Pro | Confidential

Advantage over using RDS versus deploying DB on EC2

- RDS is a managed service:
  - Automated provisioning, OS patching
  - Continuous backups and restore to specific timestamp (Point in Time Restore)!
  - Monitoring dashboards
  - Read replicas for improved read performance
  - Multi AZ setup for DR (Disaster Recovery)
  - Maintenance windows for upgrades
  - Scaling capability (vertical and horizontal)
  - Storage backed by EBS (gp2 or io1)

© Copyright by Cloud Mentor Pro | Confidential

RDS – Storage Auto Scaling

- Helps you increase storage on your RDS DB instance dynamically
- When RDS detects you are running out of free database storage, it scales automatically
- Avoid manually scaling your database storage
- You have to set **Maximum Storage Threshold** (maximum limit for DB storage)
- Automatically modify storage if:
  - Free storage is less than 10% of allocated storage
  - Low-storage lasts at least 5 minutes
  - 6 hours have passed since last modification
- Useful for applications with unpredictable workloads
- Suppor ts all RDS database engines

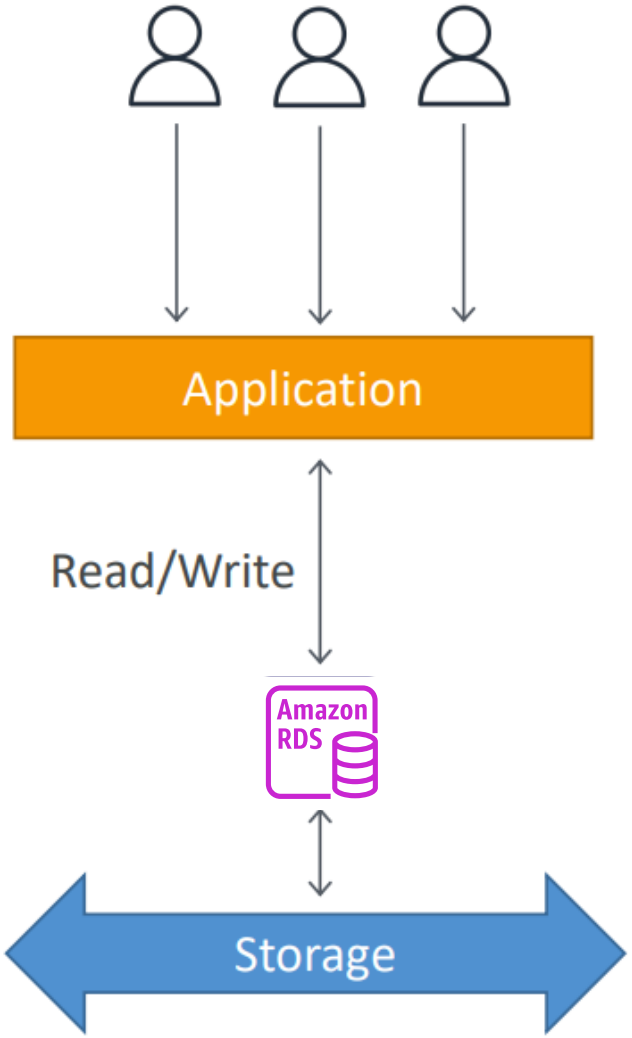

© Copyright by Cloud Mentor Pro | Confidential

RDS Read Replicas for read scalability

- Up to 15 Read Replicas
- Within AZ, Cross AZ or Cross Region
- Replication is **ASYNC**, so reads are eventually consistent
- Replicas can be promoted to their own DB
- Applications must update the connection string to leverage read replicas

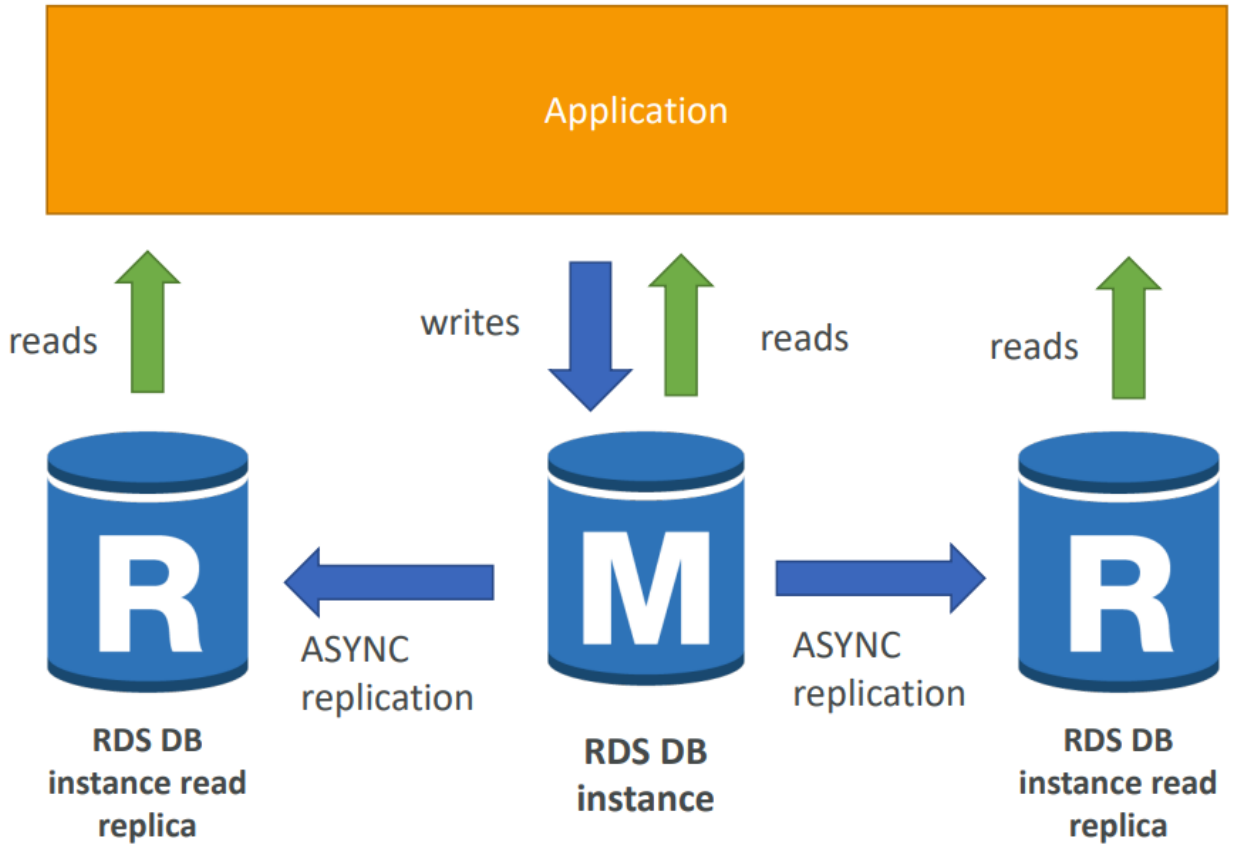

© Copyright by Cloud Mentor Pro | Confidential

RDS Read Replicas – Use Cases

- You have a production database that is taking on normal load
- You want to run a reporting application to run some analytics
- You create a Read Replica to run the new workload there
- The production application is unaffected
- Read replicas are used for SELECT (=read) only kind of statements (not INSERT, UPDATE, DELETE)

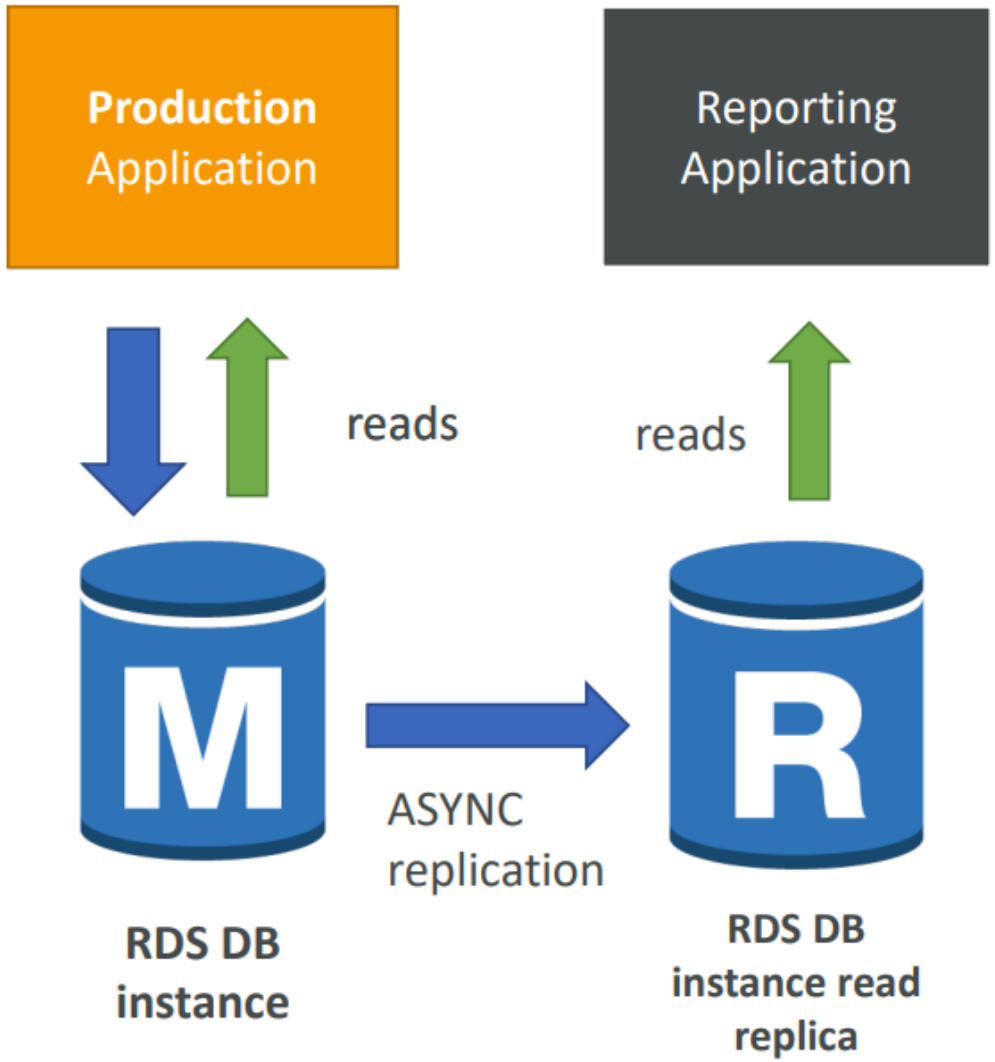

© Copyright by Cloud Mentor Pro | Confidential

RDS Read Replicas – Network Cost

- In AWS there's a network cost when data goes from one AZ to another
- **For RDS Read Replicas within the same region, you don't pay that fee**

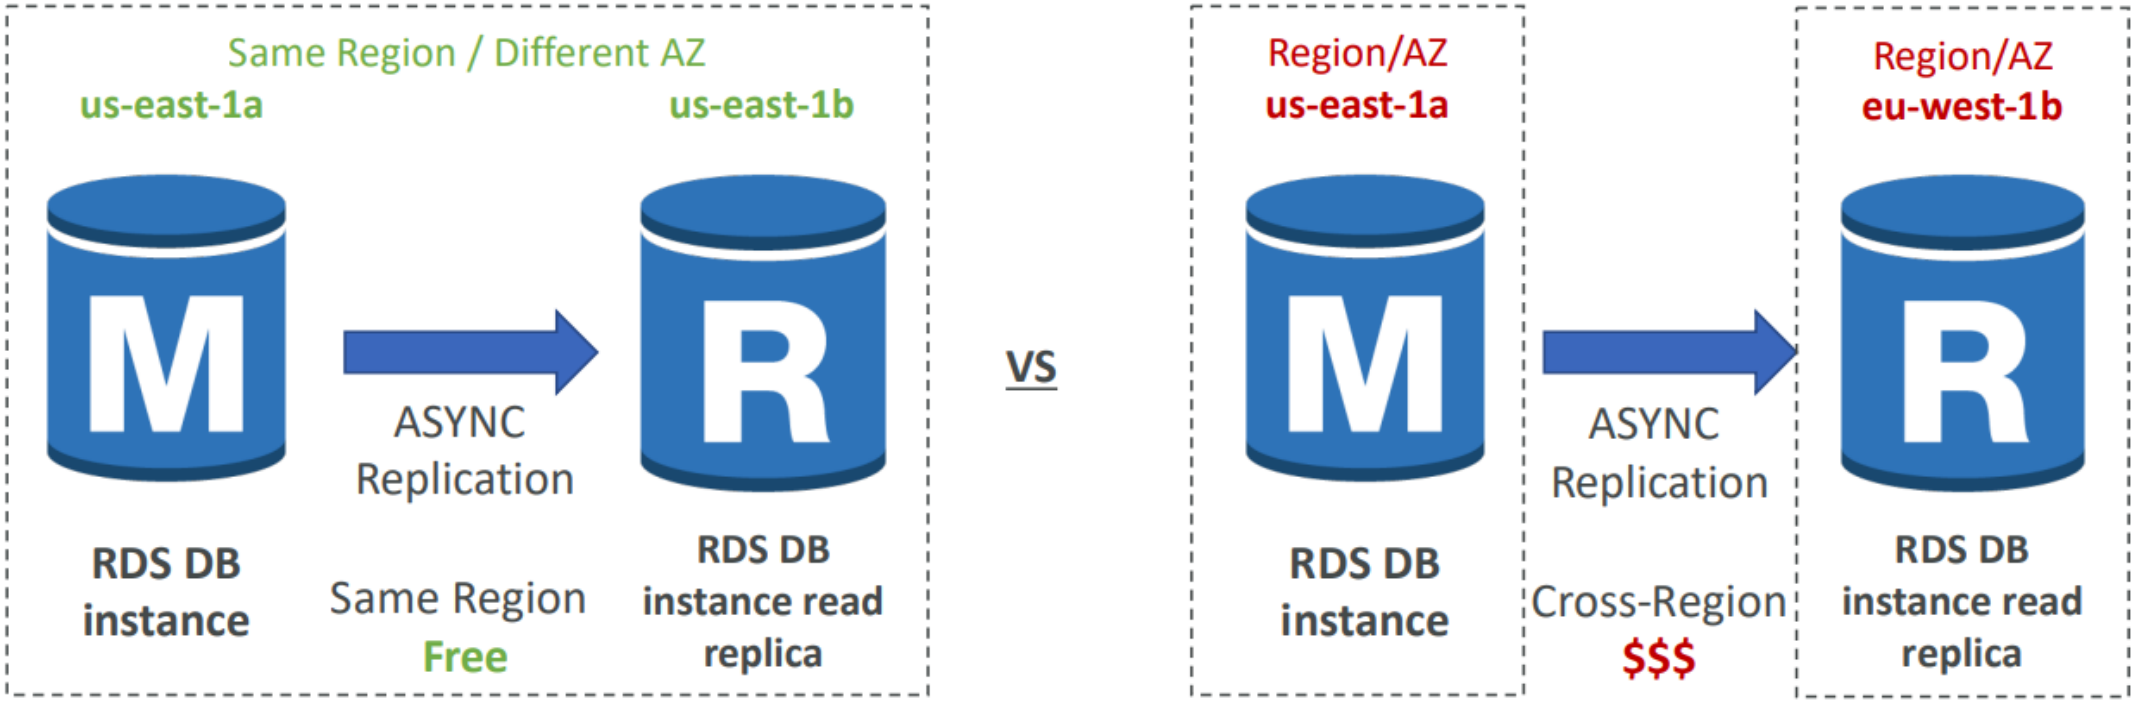

© Copyright by Cloud Mentor Pro | Confidential

RDS Multi AZ (Disaster Recovery)

- **SYNC** replication
- One DNS name – automatic app failover to standby
- Increase **availability**
- Failover in case of loss of AZ, loss of network, instance or storage failure
- No manual intervention in apps
- Not used for scaling

- Note: The Read Replicas be setup as Multi AZ for Disaster Recovery (DR)

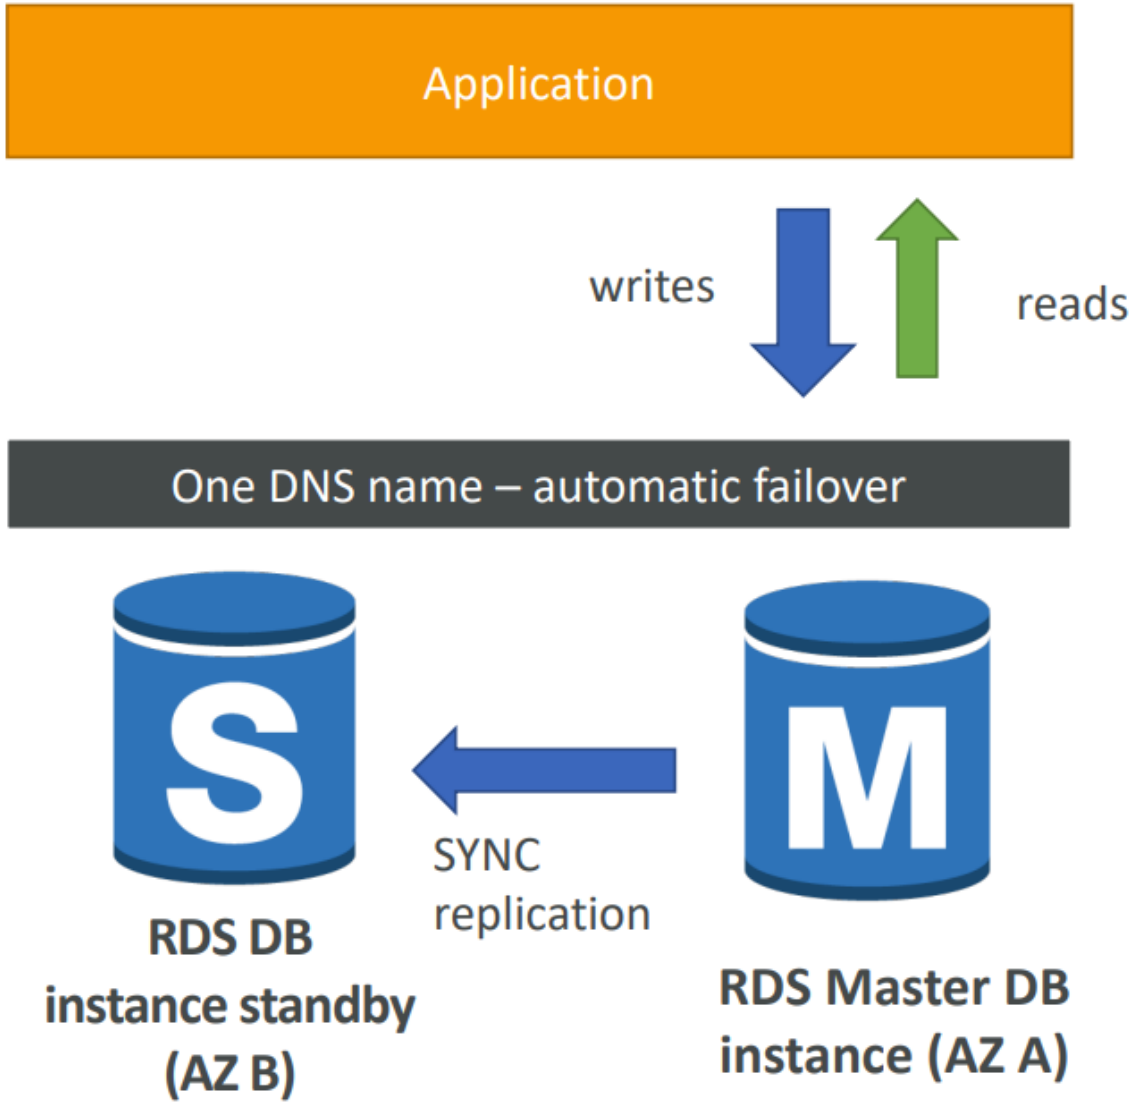

© Copyright by Cloud Mentor Pro | Confidential

RDS – From Single-AZ to Multi-AZ

- Zero downtime operation (no need to stop the DB)
- Just click on "modify" for the database
- The following happens internally:
  - A snapshot is taken
  - A new DB is restored from the snapshot in a new AZ
  - Synchronization is established between the two databases

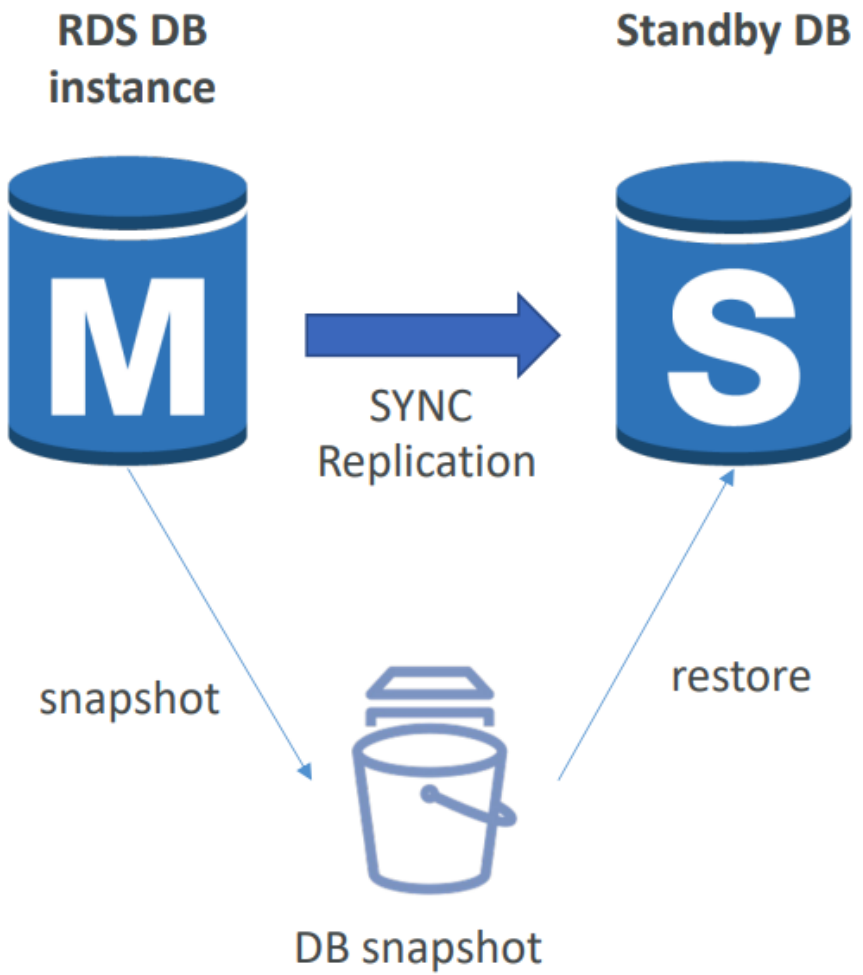

© Copyright by Cloud Mentor Pro | Confidential

Amazon Aurora

- Aurora is a proprietary technology from AWS (**not open sourced**)
- **Postgres** and **MySQL** are both supported as Aurora DB
- Aurora is "AWS cloud optimized" and claims 5x performance improvement over MySQL on RDS, over 3x the performance of Postgres on RDS
- Aurora storage automatically grows in increments of 10GB, up to 128 TB.
- Aurora can have up to **15 replicas** and the replication process is faster than MySQL (sub 10 ms replica lag)
- Failover in Aurora is instantaneous. **It's HA (High Availability) native.**
- Aurora costs more than RDS (**20% more**) – but is more efficient

© Copyright by Cloud Mentor Pro | Confidential

Aurora High Availability and Read Scaling

- 6 copies of your data across 3 AZ:
  - 4 copies out of 6 needed for writes
  - 3 copies out of 6 need for reads
  - Self healing with peer-to-peer replication
  - Storage is striped across 100s of volumes
- One Aurora Instance takes writes (master)
- Automated failover for master in less than 30 seconds
- Master + up to 15 Aurora Read Replicas serve reads
- **Support for Cross Region Replication**

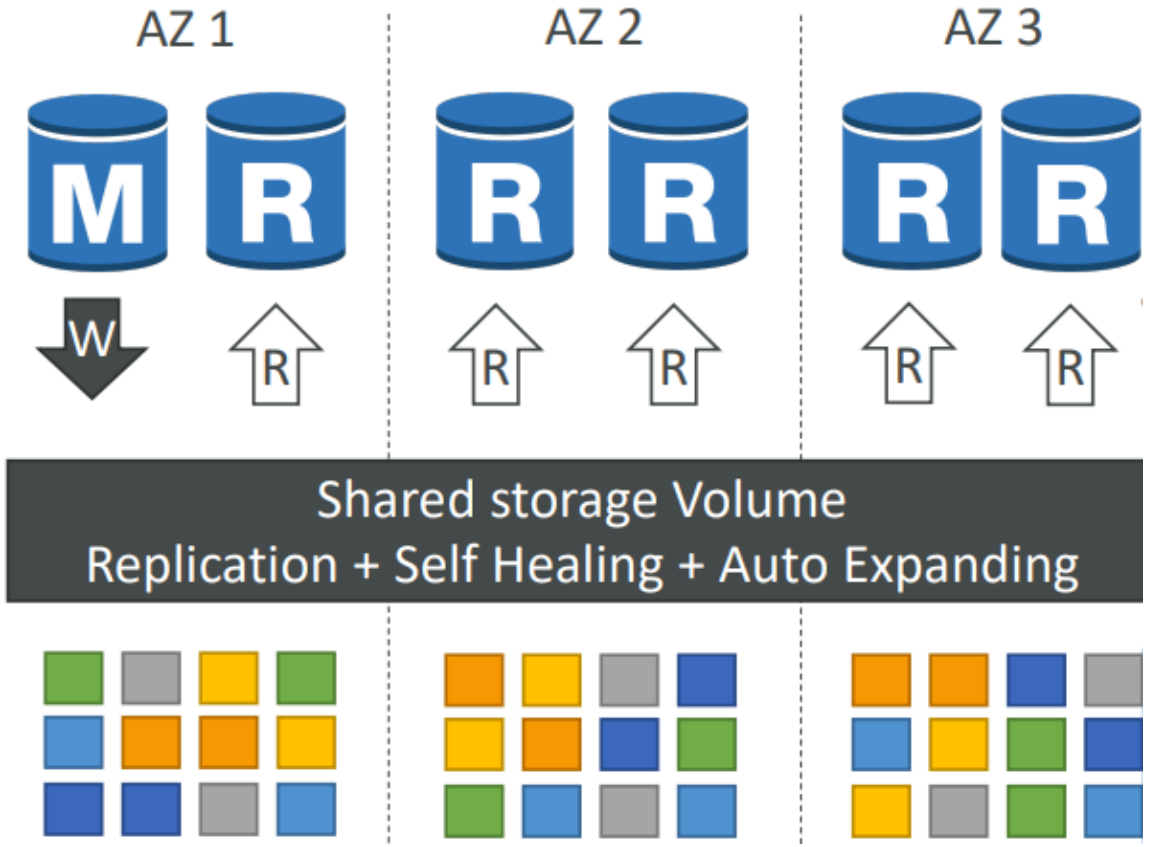

© Copyright by Cloud Mentor Pro | Confidential

Aurora DB Cluster

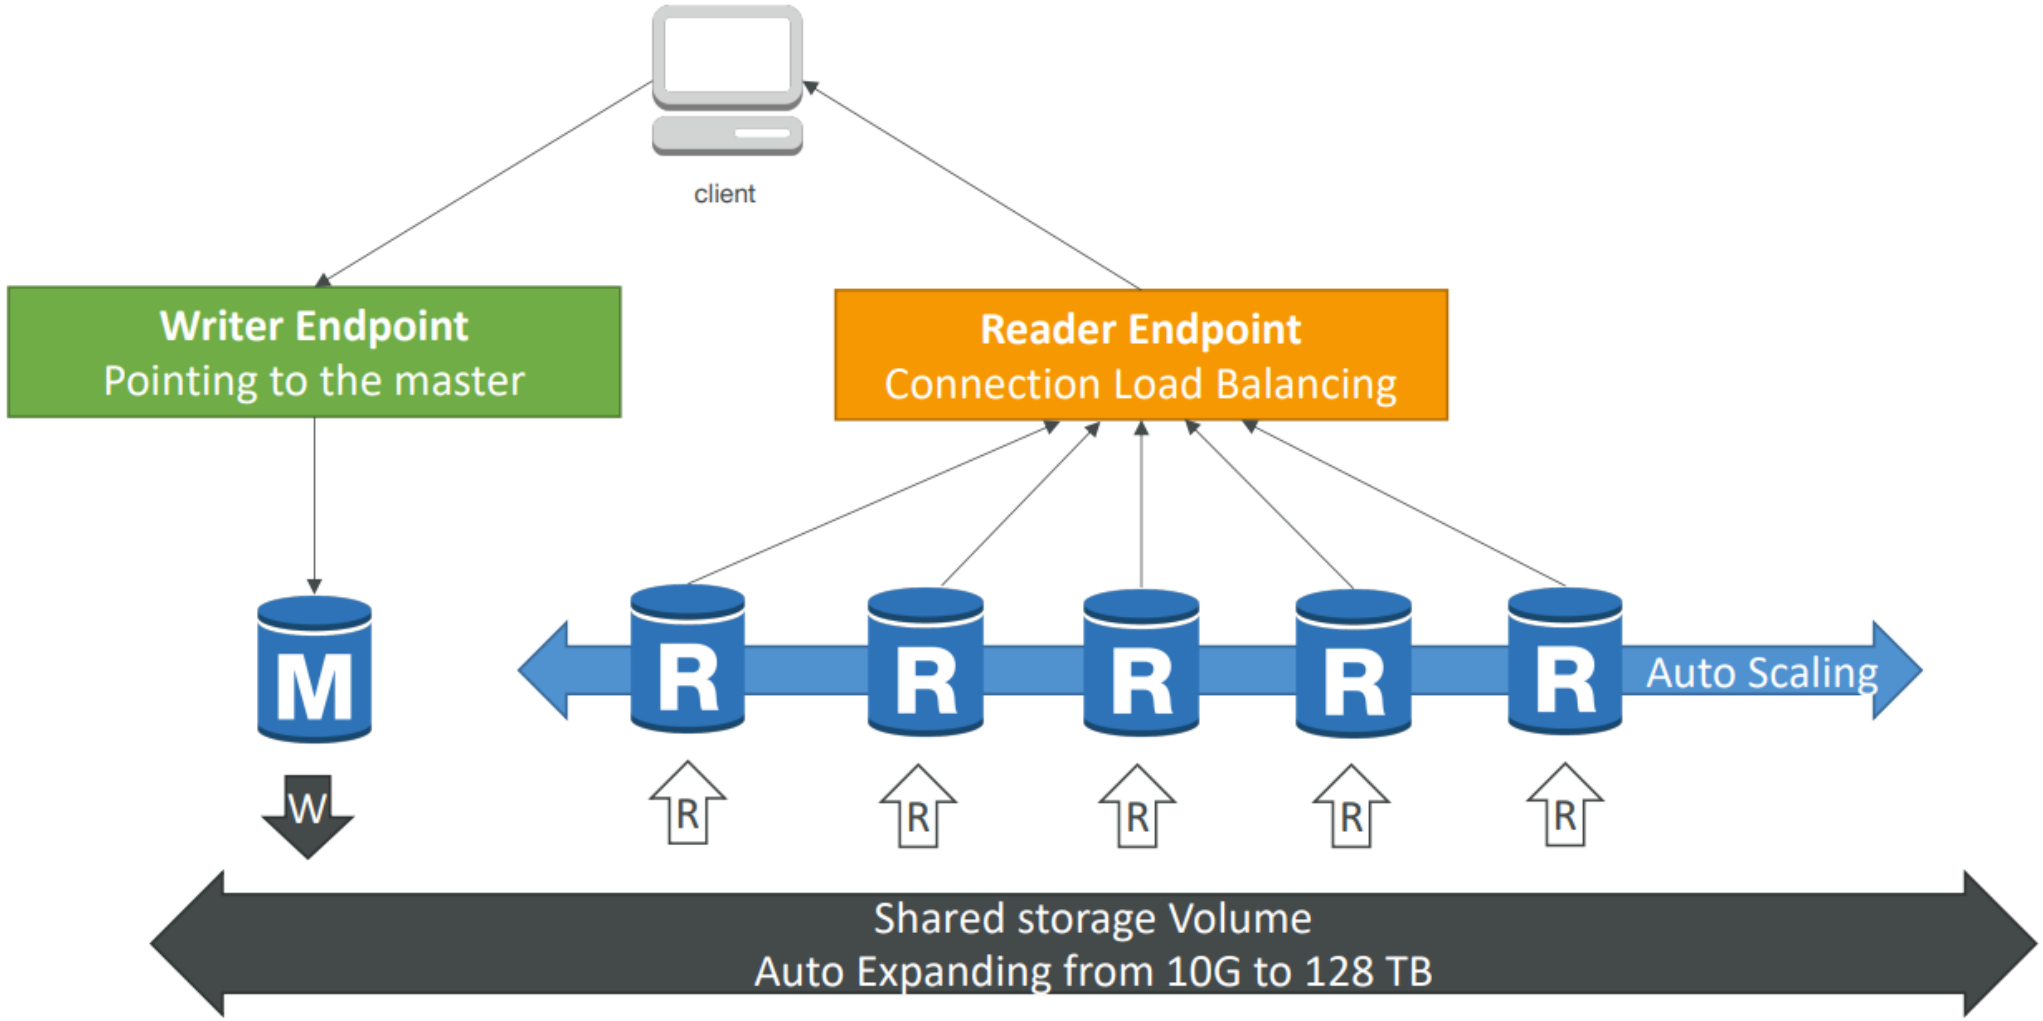

© Copyright by Cloud Mentor Pro | Confidential

Features of Aurora

- Automatic fail-over
- Backup and Recovery
- Isolation and security
- Industry compliance
- Push-button scaling
- Automated Patching with Zero Downtime
- Advanced Monitoring
- Routine Maintenance
- Backtrack: restore data at any point of time without using backups

© Copyright by Cloud Mentor Pro | Confidential

RDS & Aurora Security

- **At-rest encryption:**
  - Database master & replicas encryption using AWS KMS – must be defined as launch time
  - If the master is not encrypted, the read replicas cannot be encrypted
  - To encrypt an un-encrypted database, go through a DB snapshot & restore as encrypted
- **In-flight encryption:** TLS-ready by default, use the AWS TLS root certificates client-side
- **IAM Authentication:** IAM roles to connect to your database (instead of username/pw)
- **Security Groups:** Control Network access to your RDS / Aurora DB
- **No SSH available** except on RDS Custom
- **Audit Logs can be enabled** and sent to CloudWatch Logs for longer retention

© Copyright by Cloud Mentor Pro | Confidential

Amazon RDS Proxy

- Fully managed database proxy for RDS
- Allows apps to pool and share DB connections established with the database
- **Improving database efficiency by reducing the stress on database resources (e.g., CPU, RAM) and minimize open connections (and timeouts)**
- Serverless, autoscaling, highly available (multi-AZ)
- **Reduced RDS & Aurora failover time by up 66%**
- Supports RDS (MySQL, PostgreSQL, MariaDB, MS SQL Server) and Aurora (MySQL, PostgreSQL)
- No code changes required for most apps
- **Enforce IAM Authentication for DB, and securely store credentials in AWS Secrets Manager**
- **RDS Proxy is never publicly accessible (must be accessed from VPC)**

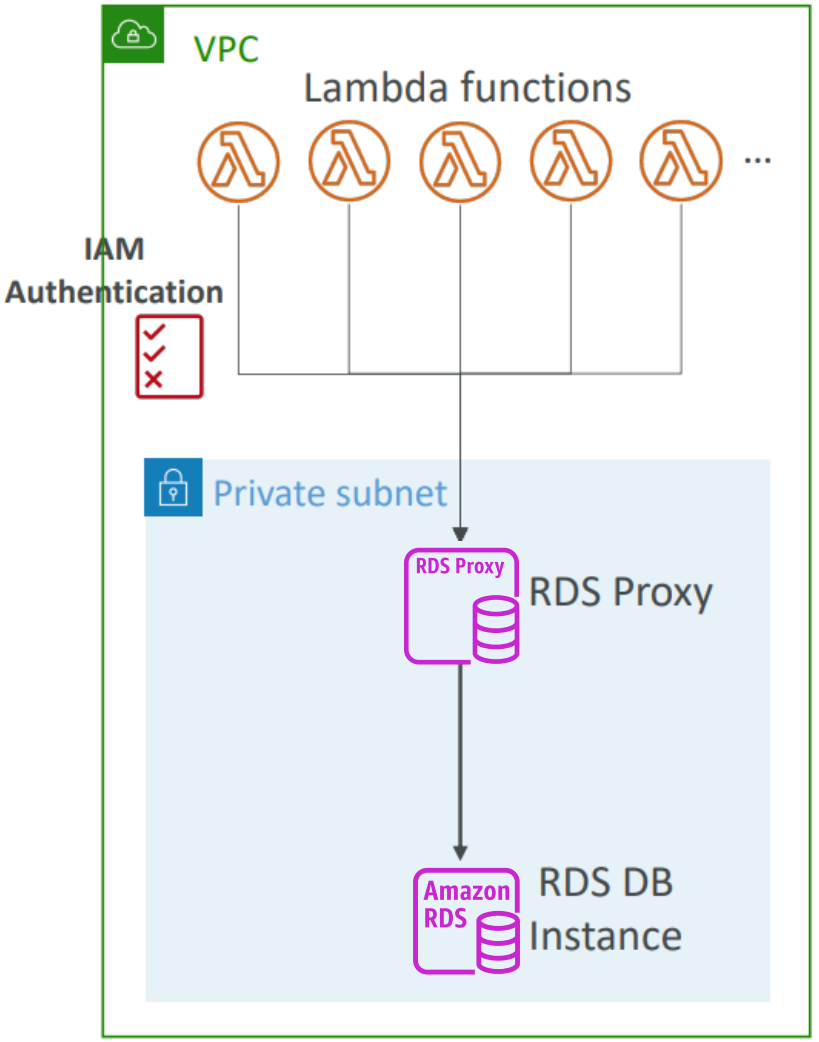

© Copyright by Cloud Mentor Pro | Confidential

Amazon ElastiCache Overview

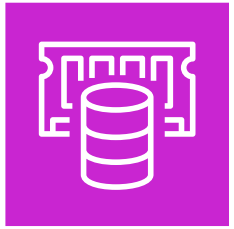

- The same way RDS is to get managed Relational Databases…
- ElastiCache is to get managed Redis or Memcached
- Caches are in-memory databases with really high performance, low latency
- Helps reduce load off of databases for read intensive workloads
- Helps make your application stateless
- AWS takes care of OS maintenance / patching, optimizations, setup, configuration, monitoring, failure recovery and backups
- Using ElastiCache involves heavy application code changes

© Copyright by Cloud Mentor Pro | Confidential

ElastiCache Solution Architecture - DB Cache

- Applications queries ElastiCache, if not available, get from RDS and store in ElastiCache.
- Helps relieve load in RDS
- Cache must have an invalidation strategy to make sure only the most current data is used in there.

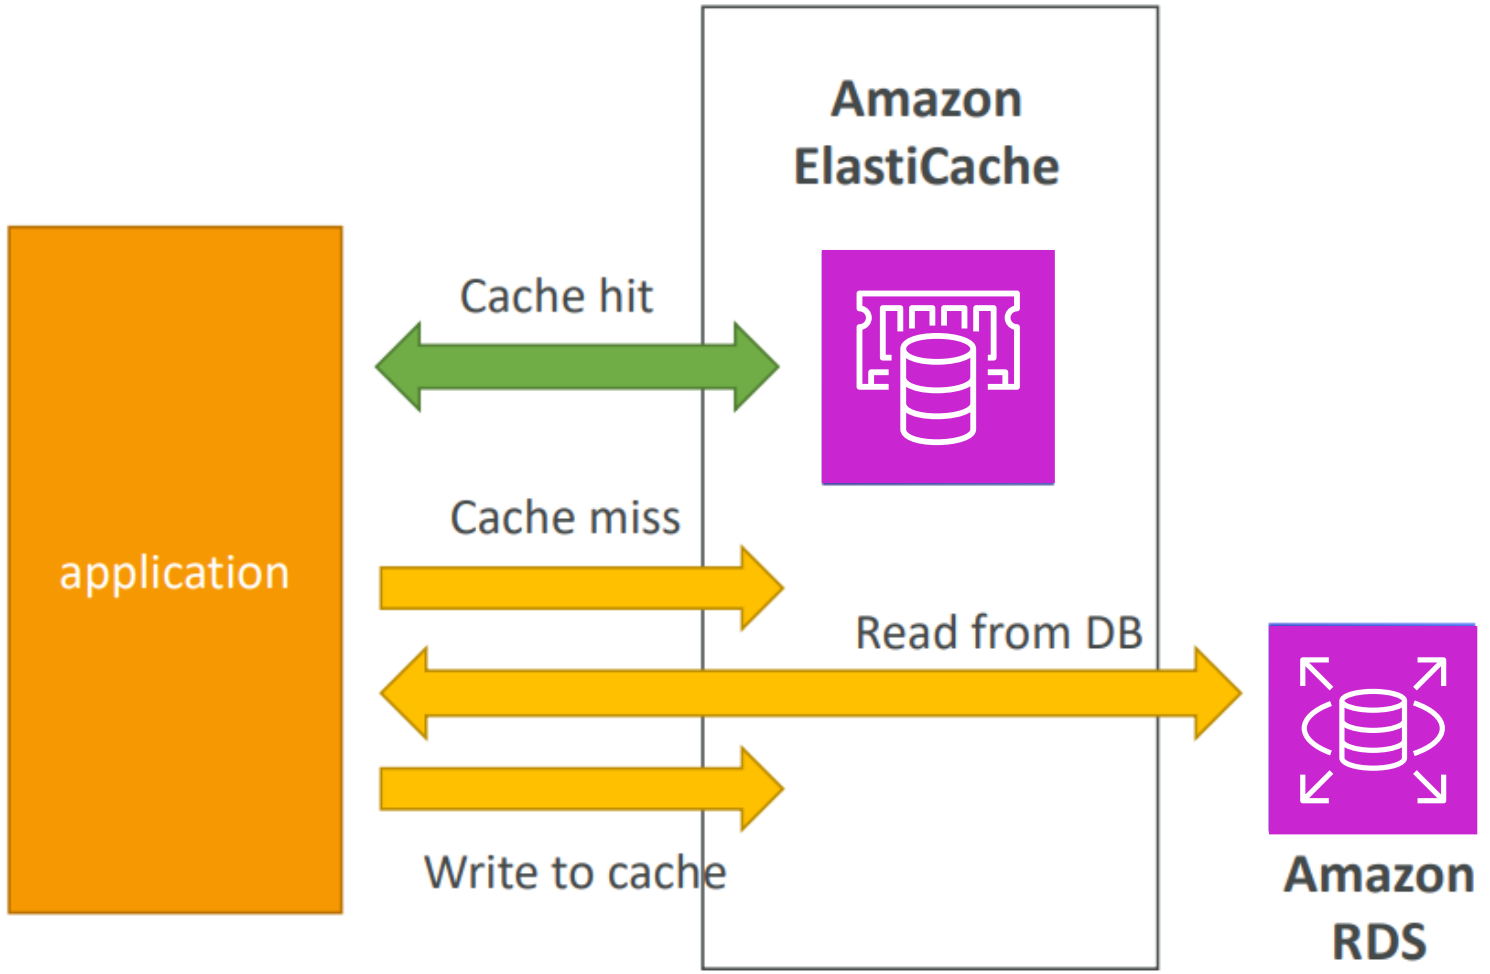

© Copyright by Cloud Mentor Pro | Confidential

ElastiCache
Solution Architecture – User Session Store

- User logs into any of the application
- The application writes the session data into ElastiCache
- The user hits another instance of our application
- The instance retrieves the data and the user is already logged in

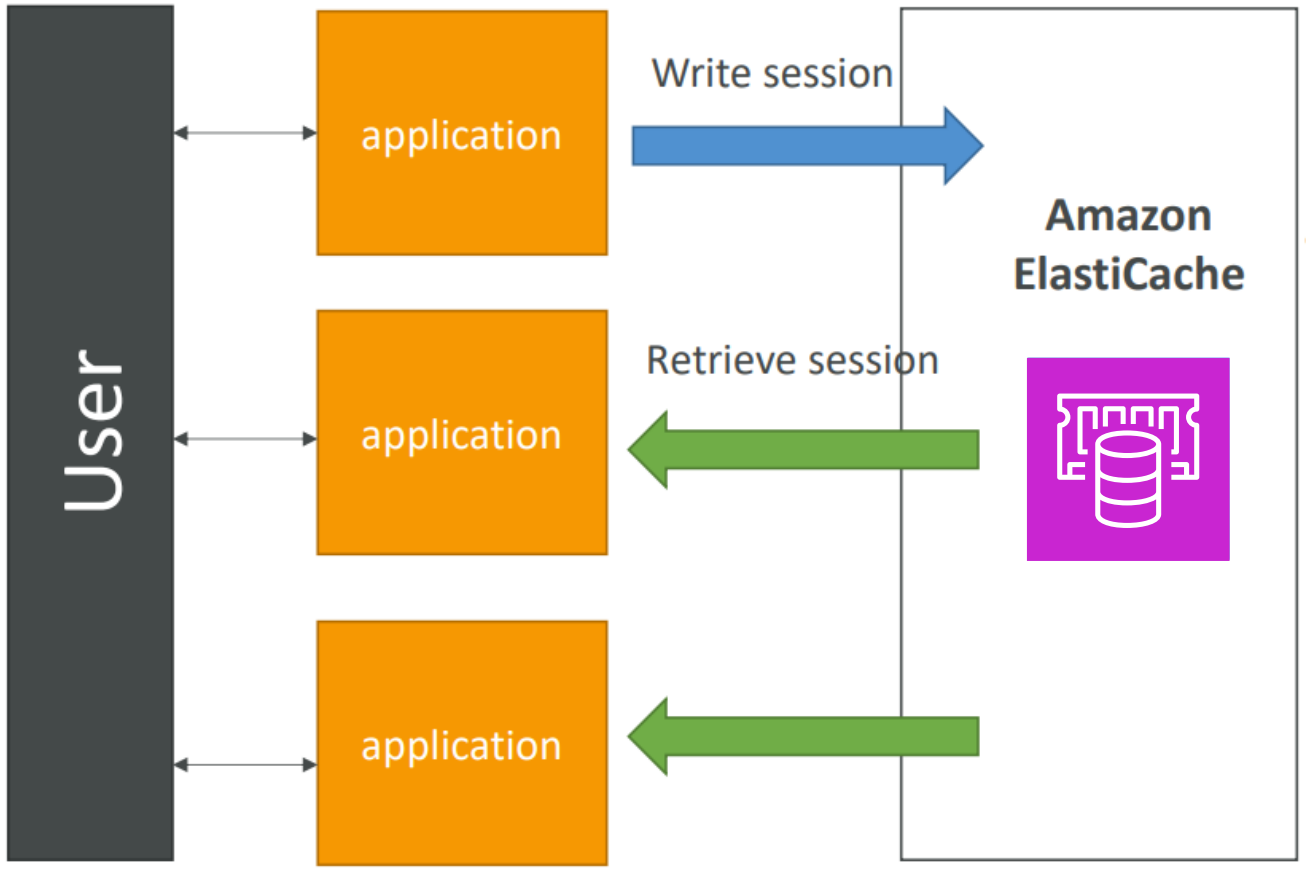

© Copyright by Cloud Mentor Pro | Confidential

ElastiCache – Redis vs Memcached

REDIS

- **Multi AZ** with Auto-Failover
- **Read Replicas** to scale reads and have **high availability**
- Data Durability using AOF persistence
- **Backup and restore features**
- Supports Sets and Sorted Sets

MEMCACHED

- Multi-node for partitioning of data (sharding)
- **No high availability (replication)**
- **Non persistent**
- **No backup and restore**
- Multi-threaded architecture

© Copyright by Cloud Mentor Pro | Confidential

Caching Implementation Considerations

- Read more at: https://aws.amazon.com/caching/implementation-considerations/

- **Is it safe to cache data?** Data may be out of date, eventually consistent
- **Is caching effective for that data?**
  - Pattern: data changing slowly, few keys are frequently needed
  - Anti patterns: data changing rapidly, all large key space frequently needed
- **Is data structured well for caching?**
  - example: key value caching, or caching of aggregations results
- Which caching design pattern is the most appropriate?

© Copyright by Cloud Mentor Pro | Confidential

Lazy Loading / Cache-Aside / Lazy Population

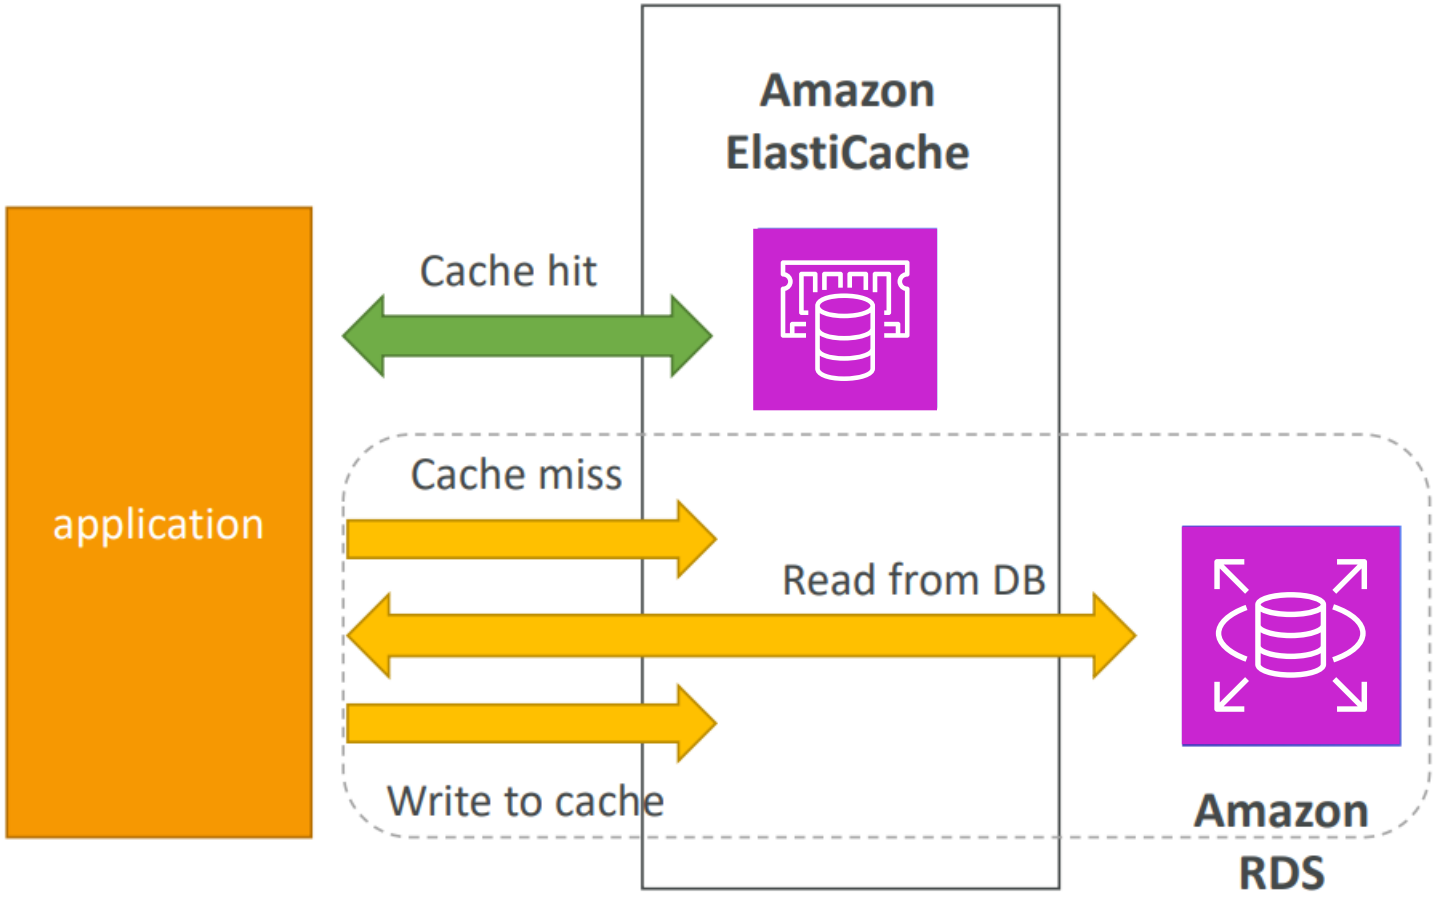

- Pros
  - Only requested data is cached (the cache isn't filled up with unused data)
  - Node failures are not fatal (just increased latency to warm the cache)
- Cons
  - Cache miss penalty that results in 3 round trips, noticeable delay for that request
  - Stale data: data can be updated in the database and outdated in the cache

© Copyright by Cloud Mentor Pro | Confidential

Lazy Loading / Cache-Aside / Lazy Population Python Pseudocode

```python
# Python

def get_user(user_id):
    # Check the cache
    record = cache.get(user_id)

    if record is None:
        # Run a DB query
        record = db.query("select * from users where id = ?", user_id)
        # Populate the cache
        cache.set(user_id, record)
        return record
    else:
        return record

# App code
user = get_user(17)
```

© Copyright by Cloud Mentor Pro | Confidential

Write Through – Add or Update cache when database is updated

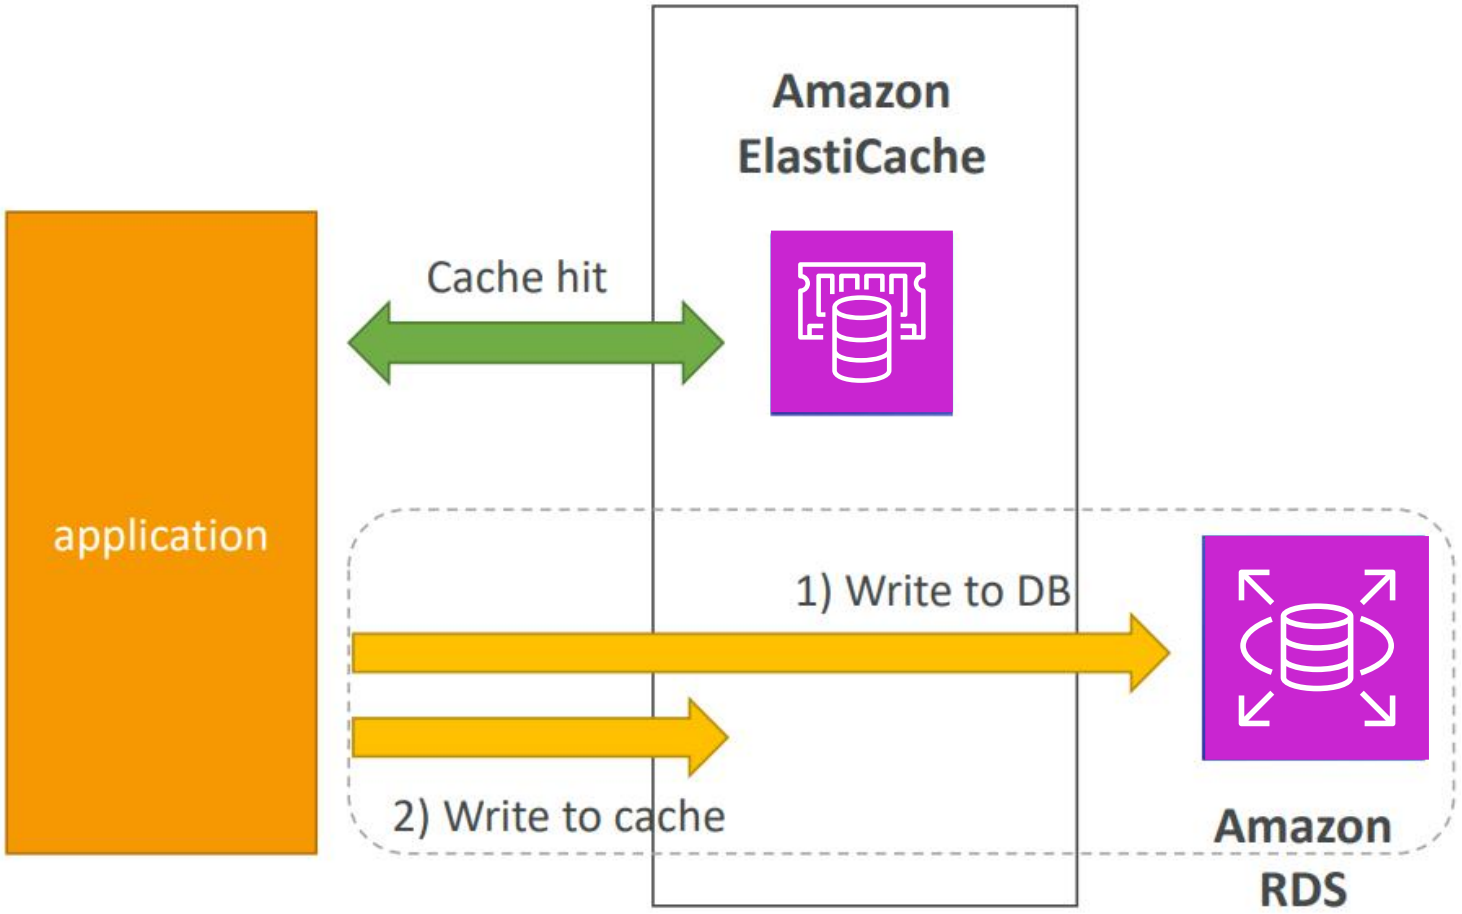

- Pros:
  - Data in cache is never stale, reads are quick
  - Write penalty vs Read penalty (each write requires 2 calls)
- Cons:
  - Missing Data until it is added / updated in the DB. Mitigation is to implement Lazy Loading strategy as well
  - Cache churn – a lot of the data will never be read

© Copyright by Cloud Mentor Pro | Confidential

Write –Through
Python Pseudocode

```python
# Python

def save_user(user_id, values):

    # Save to DB

    record = db.query("update users ... where id = ?", user_id, values)

    # Push into cache

    cache.set(user_id, record)

    return record

# App code

user = save_user(17, {"name": "Nate Dogg"})
```

© Copyright by Cloud Mentor Pro | Confidential

Cache Evictions and Time-to-live (TTL)

- Cache eviction can occur in three ways:
  - You delete the item explicitly in the cache
  - Item is evicted because the memory is full and it's not recently used (LRU)
  - You set an item **time-to-live (or TTL)**
- TTL are helpful for any kind of data:
  - Leaderboards
  - Comments
  - Activity streams
- TTL can range from few seconds to hours or days
- If too many evictions happen due to memory, you should scale up or out

© Copyright by Cloud Mentor Pro | Confidential

Final words of wisdom

- Lazy Loading / Cache aside is easy to implement and works for many situations as a foundation, especially on the read side
- Write-through is usually combined with Lazy Loading as targeted for the queries or workloads that benefit from this optimization
- Setting a TTL is usually not a bad idea, except when you're using Write-through. Set it to a sensible value for your application
- Only cache the data that makes sense (user profiles, blogs, etc…)

© Copyright by Cloud Mentor Pro | Confidential

Amazon Route 53

DNS, Records, Hosted Zones, Routing Policies, Health Checks

© Copyright by Cloud Mentor Pro | Confidential

What is DNS?

- Domain Name System which translates the human friendly hostnames into the machine IP addresses
- www.google.com => 172.217.18.36
- DNS is the backbone of the Internet
- DNS uses hierarchical naming structure

.com

example.com

www.example.com

api.example.com

© Copyright by Cloud Mentor Pro | Confidential

DNS Terminologies

- **Domain Registrar:** Amazon Route 53, GoDaddy, …
- **DNS Records:** A, AAAA, CNAME, NS, …
- **Zone File:** contains DNS records
- **Name Server:** resolves DNS queries (Authoritative or Non-Authoritative)
- **Top Level Domain (TLD):** .com, .us, .in, .gov, .org, …
- **Second Level Domain (SLD):** amazon.com, google.com, …

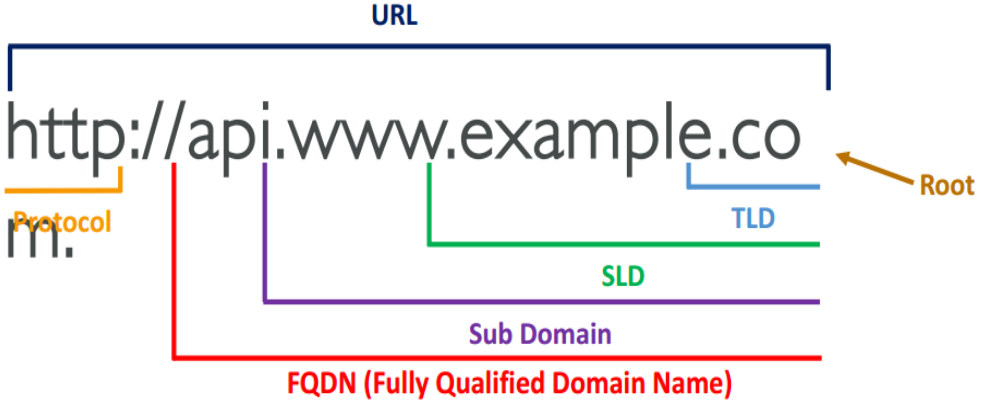

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/welcome-dns-service.html

How DNS Works

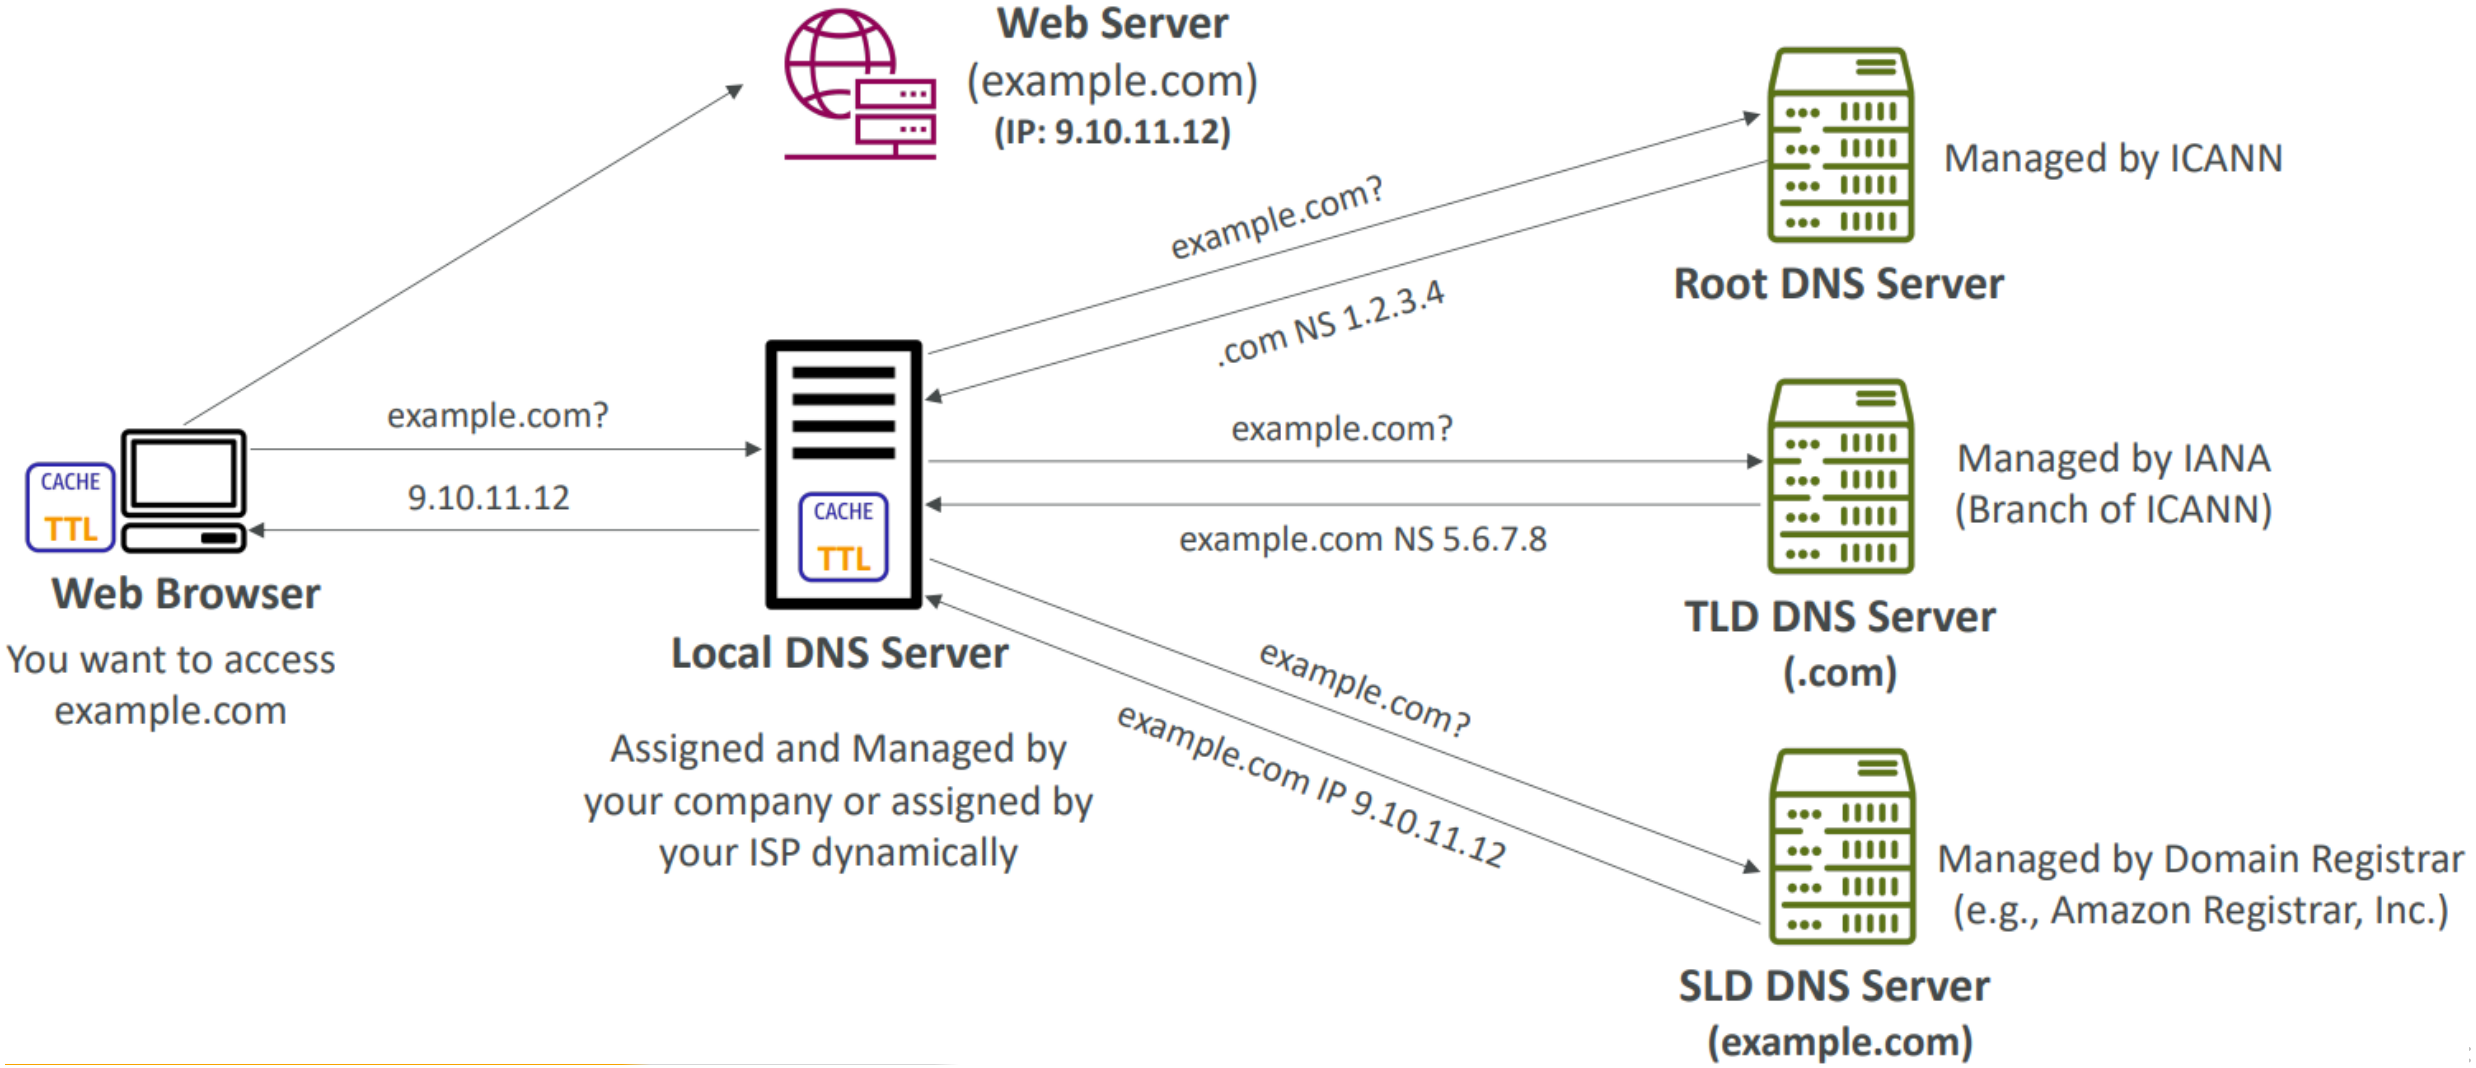

© Copyright by Cloud Mentor Pro | Confidential

Amazon Route 53

- A highly available, scalable, fully managed and Authoritative DNS
  - Authoritative = the customer (you) can update the DNS records
- Route 53 is also a **Domain Registrar**
- Ability to **check the health** of your resources
- The only AWS service which provides 100% availability SLA

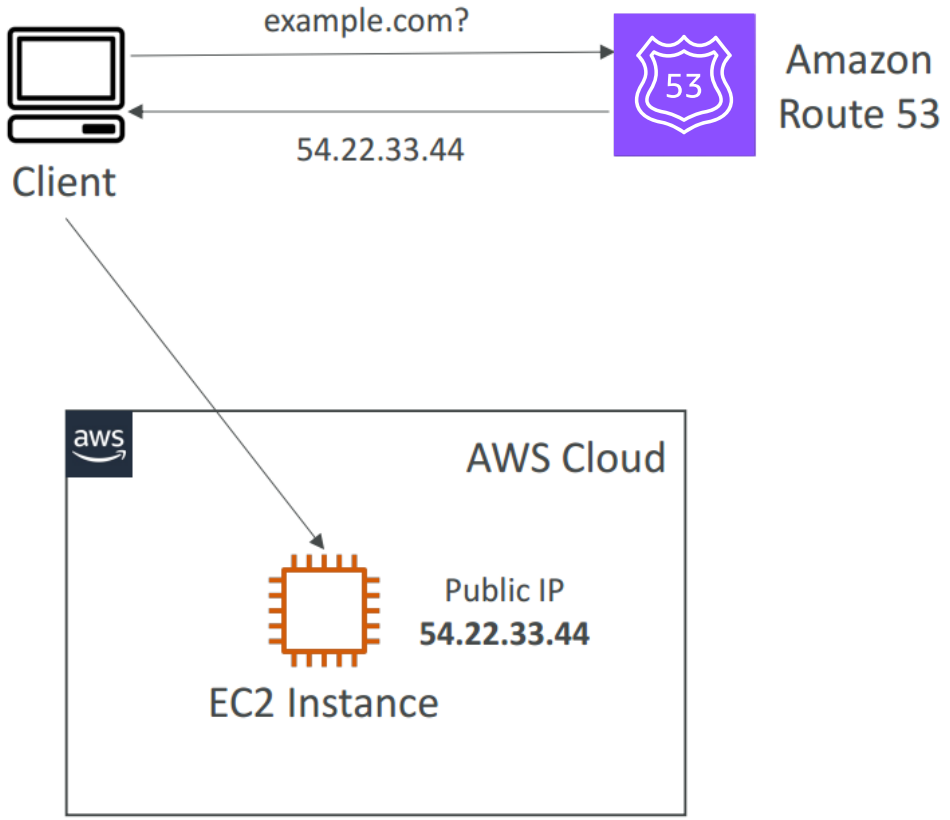

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/welcome-dns-service.html

Route 53 – Records

- How you want to route traffic for a domain
- Each record contains:
  - **Domain/subdomain Name** – e.g., example.com
  - **Record Type** – e.g., A or AAAA
  - **Value** – e.g., 12.34.56.78
  - **Routing Policy** – how Route 53 responds to queries
  - **TTL** – amount of time the record cached at DNS Resolvers
- Route 53 supports the following DNS record types:
- (must know) A / AAAA / CNAME / NS
- (advanced) CAA / DS / MX / NAPTR / PTR / SOA / TXT / SPF / SRV

© Copyright by Cloud Mentor Pro | Confidential

Route 53 – Record Types

- **A** – maps a hostname to IPv4
- **AAAA** – maps a hostname to IPv6
- **CNAME** – maps a hostname to another hostname
  - The target is a domain name which must have an A or AAAA record
  - Can't create a CNAME record for the top node of a DNS namespace (Zone Apex)
  - Example: you can't create for example.com, but you can create for www.example.com
- **NS** – Name Servers for the Hosted Zone
  - Control how traffic is routed for a domain
  - Define which organization manages your domain

© Copyright by Cloud Mentor Pro | Confidential

Route 53 – Hosted Zones

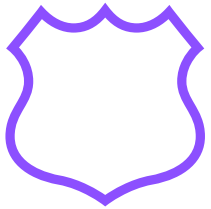

- A container for records that define how to route traffic to a domain and its subdomains

- **Public Hosted Zones** – contains records that specify how to route traffic on the Internet (public domain names) application1.mypublicdomain.com
- Private Hosted Zones – contain records that specify how you route traffic within one or more VPCs (private domain names) application1.company.internal

- You pay $0.50 per month per hosted zone

© Copyright by Cloud Mentor Pro | Confidential

Route 53 – Public vs. Private Hosted Zones

Public Hosted

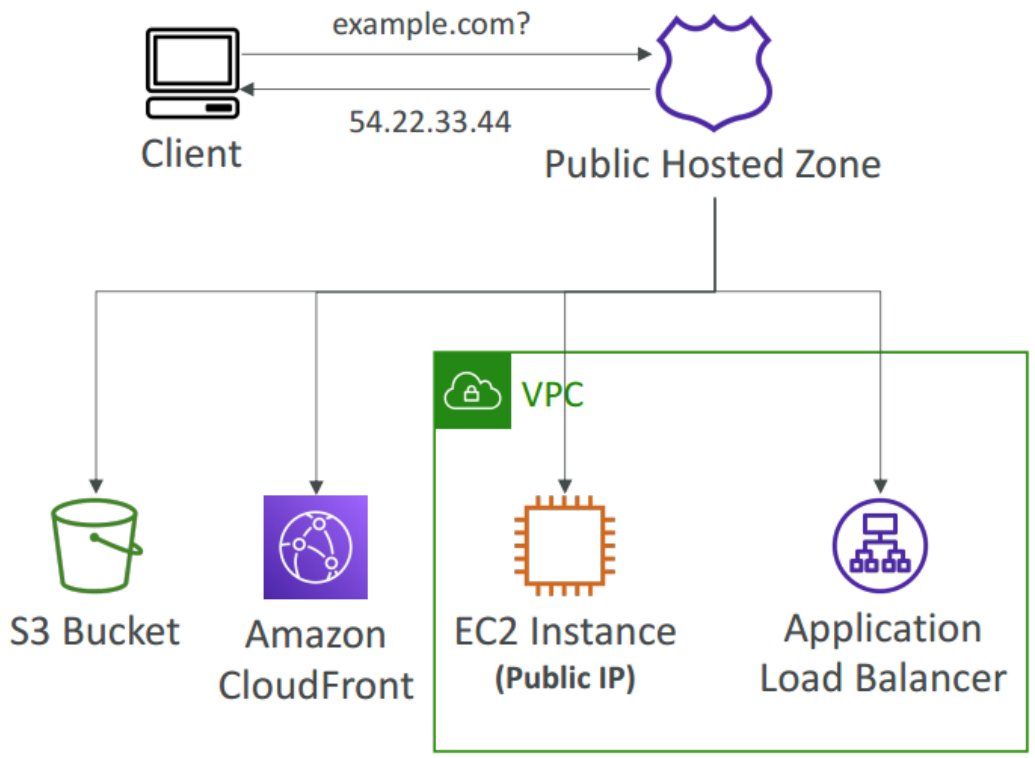

Private Hosted Zones

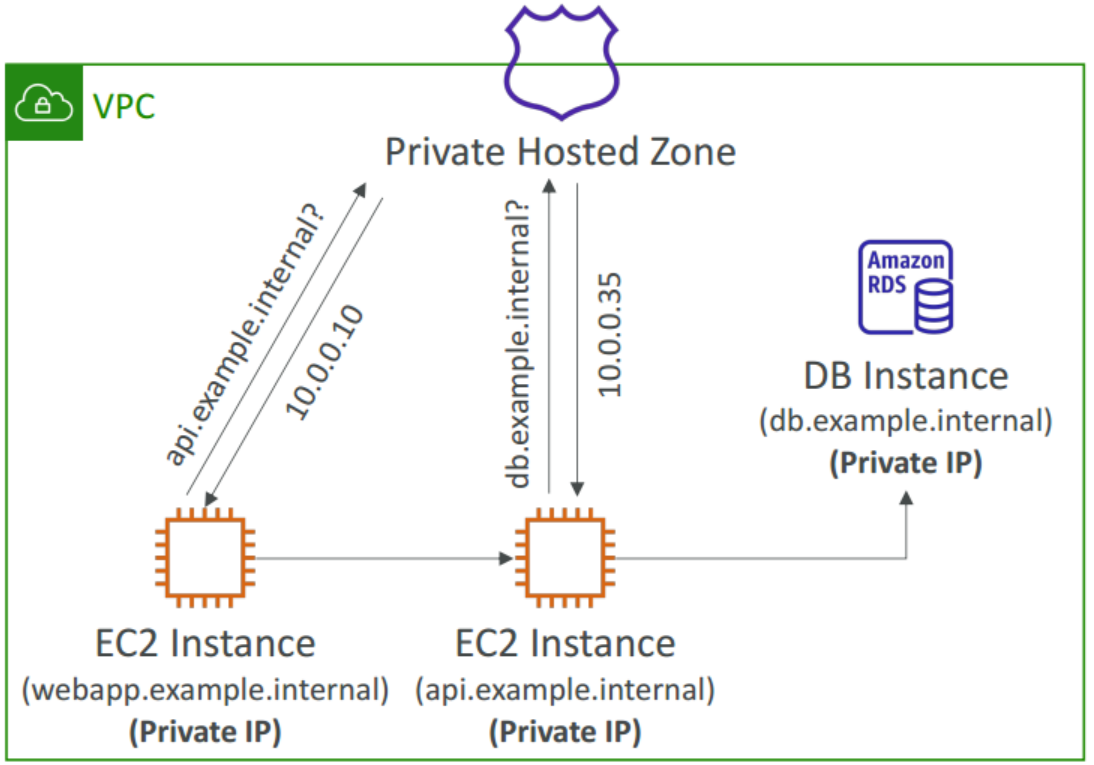

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/hosted-zones-working-with.html

Route 53 – Records TTL (Time To Live)

- High TTL – e.g., 24 hr
  - Less traffic on Route 53
  - Possibly outdated records
- Low TTL – e.g., 60 sec.
  - More traffic on Route 53 ($$)
  - Records are outdated for less time
  - Easy to change records
- Except for Alias records, TTL is mandatory for each DNS record

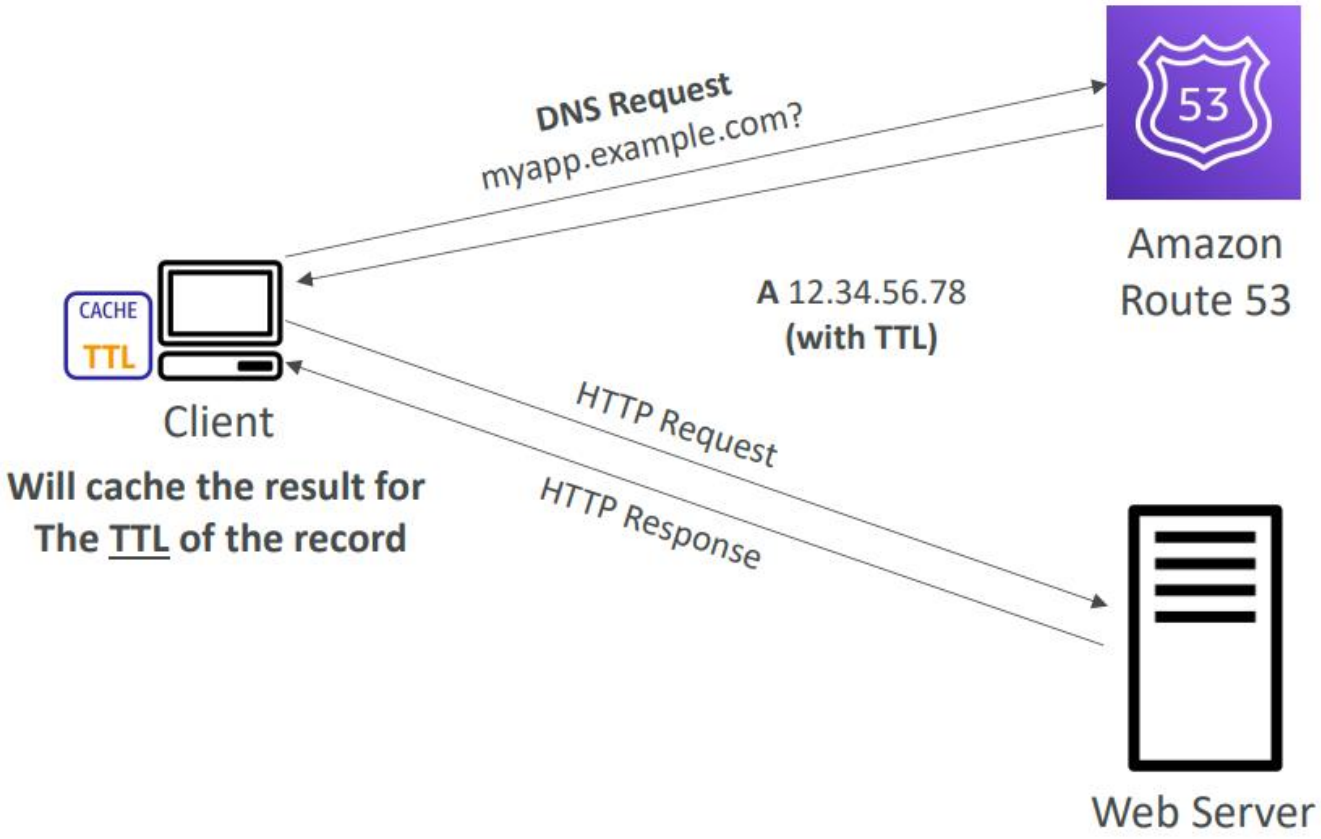

© Copyright by Cloud Mentor Pro | Confidential

CNAME vs Alias

- AWS Resources (Load Balancer, CloudFront...) expose an AWS hostname:
  - lb1-1234.us-east-2.elb.amazonaws.com and you want to route your domain myapp.mydomain.com to that host name
- CNAME:
  - Points a hostname to any other hostname. (app.mydomain.com => blabla.anything.com)
  - ONLY FOR NON ROOT DOMAIN (aka. something.mydomain.com)
- Alias:
  - Points a hostname to an AWS Resource (app.mydomain.com => blabla.amazonaws.com)
  - Works for ROOT DOMAIN and NON ROOT DOMAIN (aka mydomain.com)
  - Free of charge
  - Native health check

© Copyright by Cloud Mentor Pro | Confidential

Route 53 – Alias Records

- Maps a **hostname to an AWS resource**
- Automatically recognizes changes in the resource's IP addresses
- Unlike CNAME, it can be **used for the top node** e.g.: example.com (**ROOT DOMAIN**)
- Alias Record is always of type A/AAAA for AWS resources (IPv4 / IPv6)
- **You can't set the TTL**

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/resource-record-sets-values-alias.html

Route 53 – Routing Policies

- Define how Route 53 responds to DNS queries
- Don’t get confused by the word “Routing”
  - It’s not the same as Load balancer routing which routes the traffic
  - DNS does not route any traffic, it only responds to the DNS queries
- Route 53 Supports the following Routing Policies
  - Simple
  - Weighted
  - Failover
  - Latency based
  - Geolocation
  - Multi-Value Answer
  - Geoproximity (using Route 53 Traffic Flow feature)

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy.html

Routing Policies – Simple

- Typically, route traffic to a **single resource**
- Can specify **multiple values in the same record**
- If multiple values are returned, a random one is chosen by the client
- When Alias enabled, specify only one AWS resource
- **No health checks**

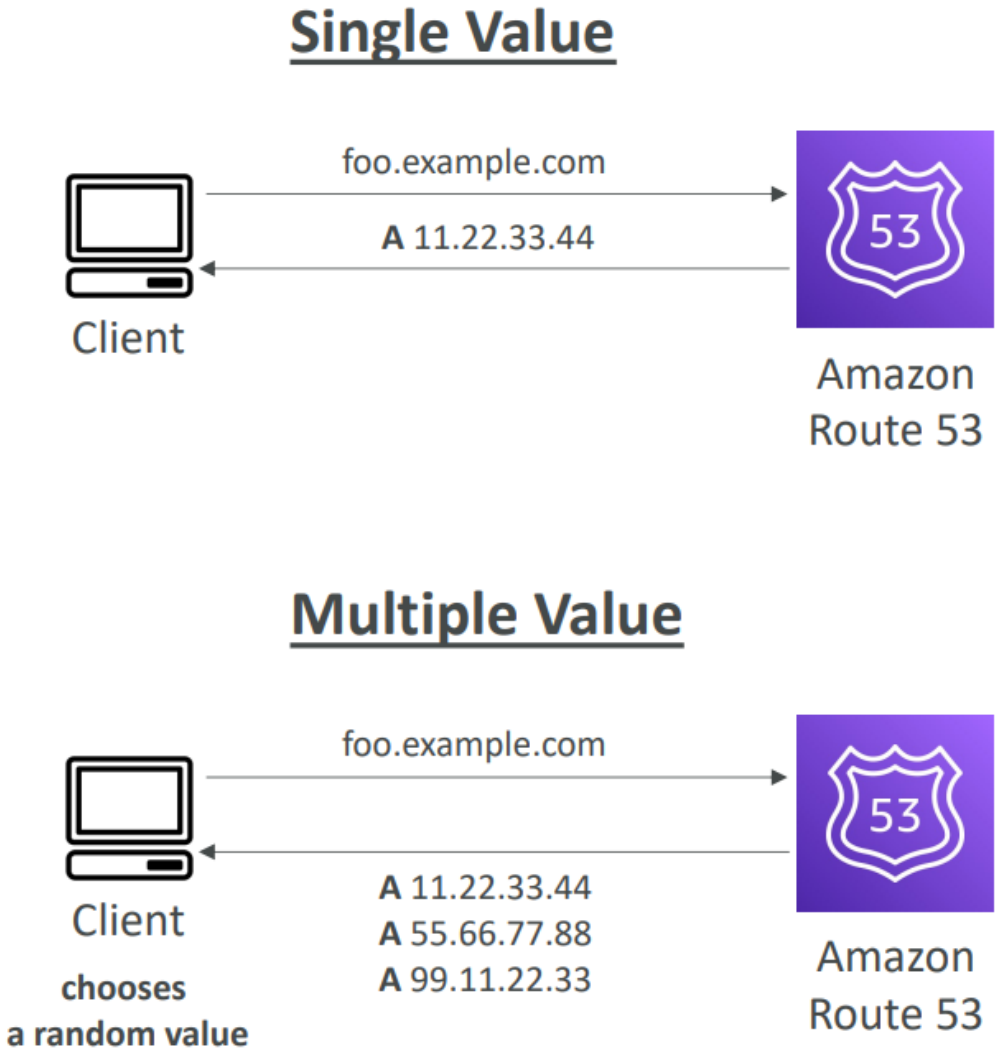

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-simple.html

Routing Policies – Weighted

- Control the **% of the requests** that go to each specific resource
- DNS records must have the same name and type
- Can be associated with **Health Checks**
- Use cases: load balancing between regions, testing new application versions…
- **Assign a weight of 0 to a record to stop sending traffic to a resource**
- **If all records have weight of 0, then all records will be returned equally**

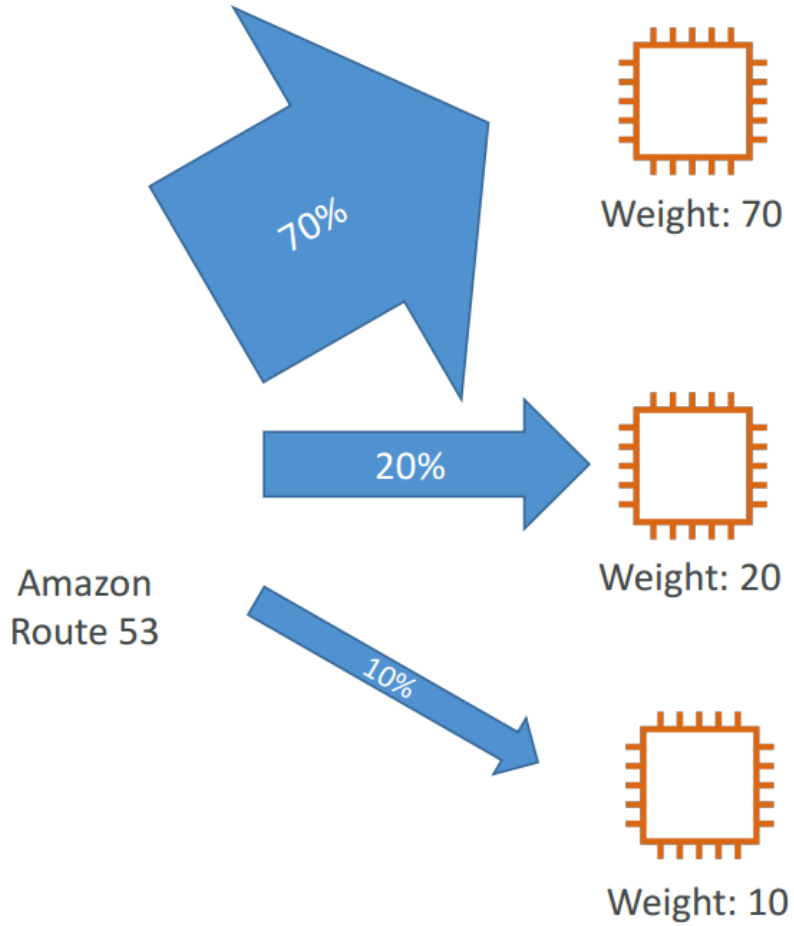

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-weighted.html

Routing Policies – Latency-based

- Redirect to the resource that has the **least latency** close to us
- Super helpful when **latency for users is a priority**
- **Latency is based on traffic between users and AWS Regions**
- Associated with **Health Checks** (has a failover capability)

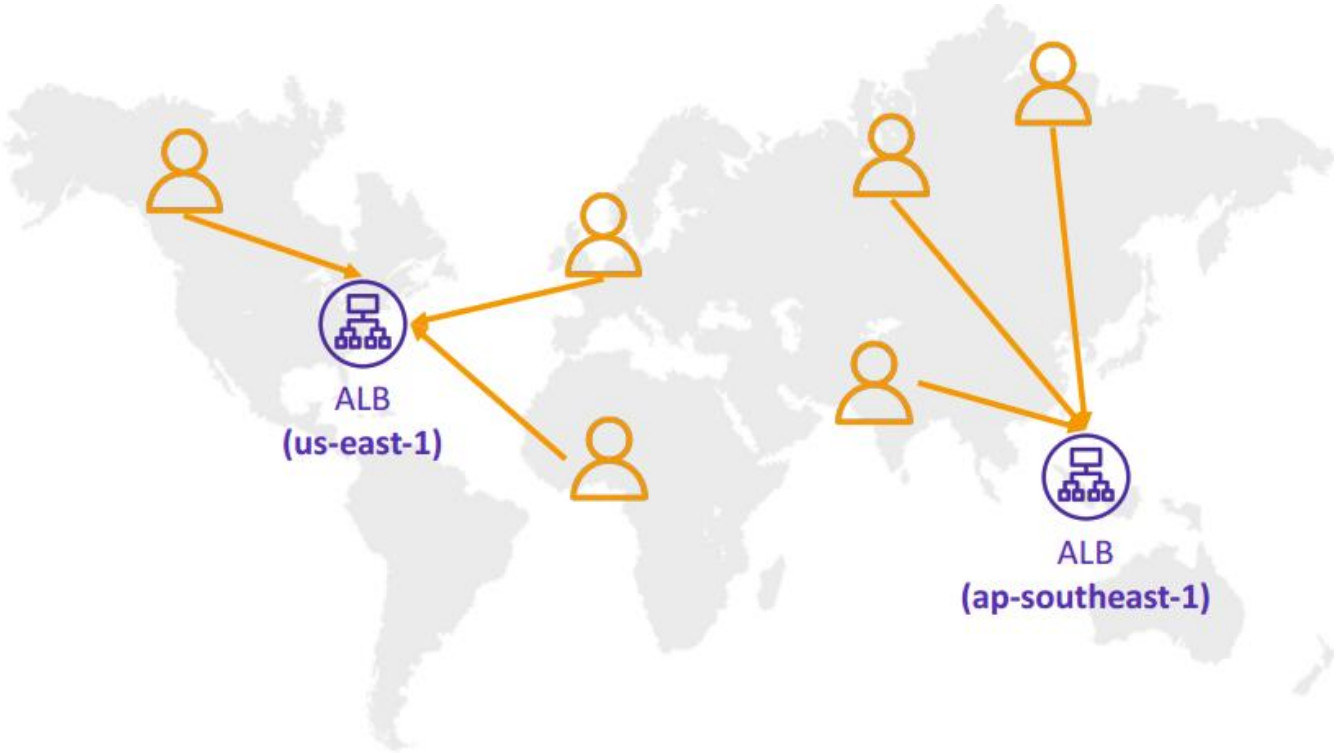

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-latency.html

Route 53 – Health Checks

- HTTP Health Checks are only for **public resources**
- Health Check => **Automated DNS Failover**:
  1. Health checks that monitor an **endpoint** (application, server, other AWS resource)
  2. Health checks that monitor **other health checks** (**Calculated Health Checks**)
  3. Health checks that monitor **CloudWatch Alarms** (full control) – e.g., throttles of DynamoDB, alarms on RDS, custom metrics, … (helpful for private resources)

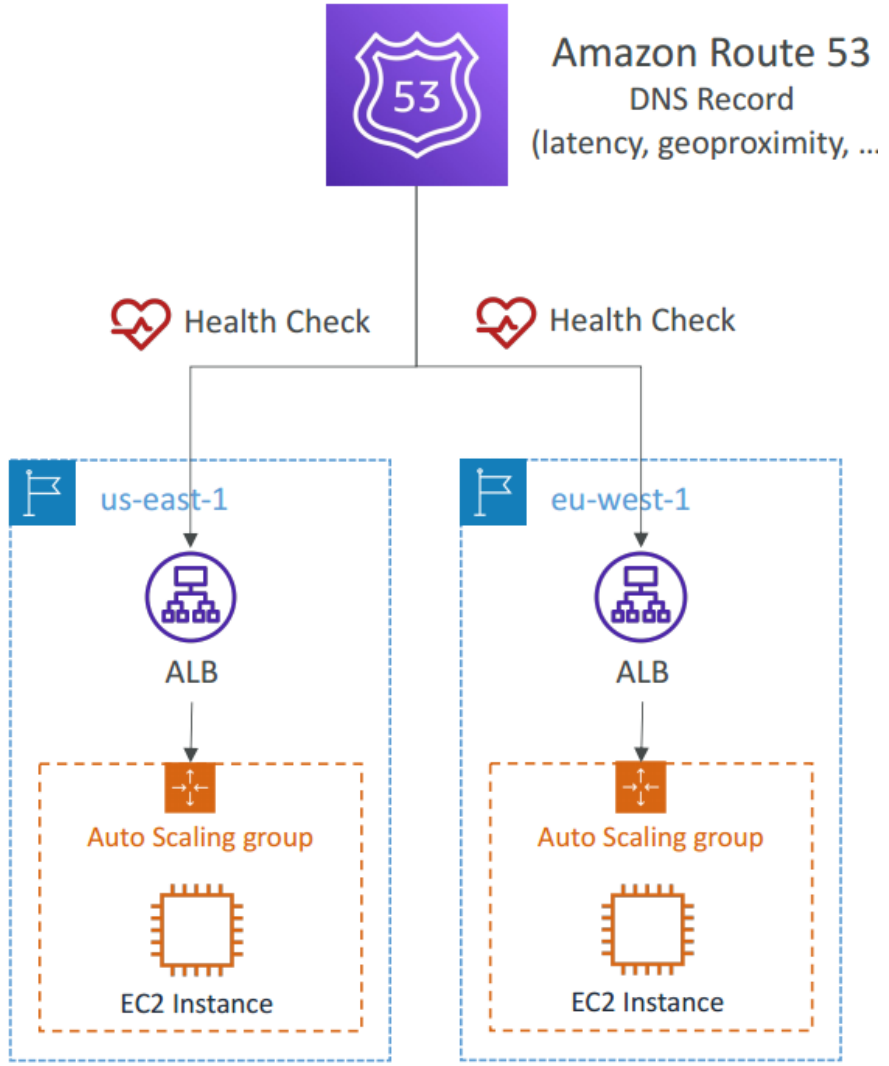

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/health-checks-types.html

Health Checks - Monitor CloudWatch Alarms

- Route 53 health checkers are outside the VPC
- They **can't access private endpoints** (private VPC or on-premises resource)

- You can create a **CloudWatch Metric** and associate a **CloudWatch Alarm,** then create a **Health Check that checks the alarm itself**

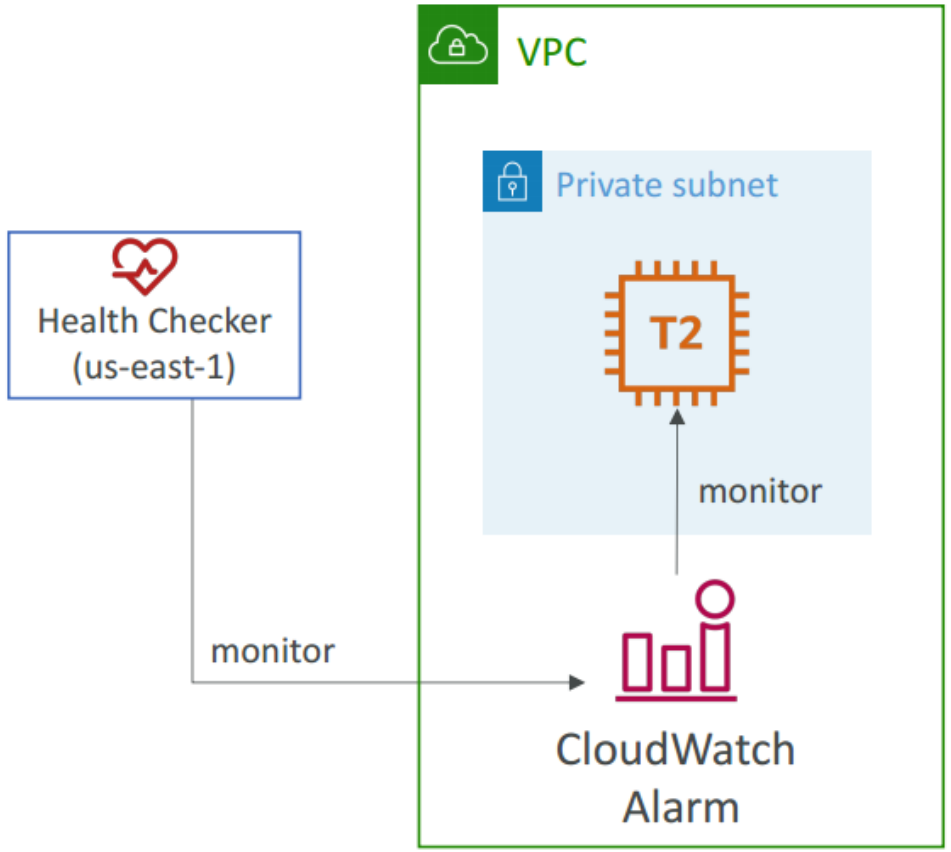

© Copyright by Cloud Mentor Pro | Confidential

Routing Policies – Failover (Active-Passive)

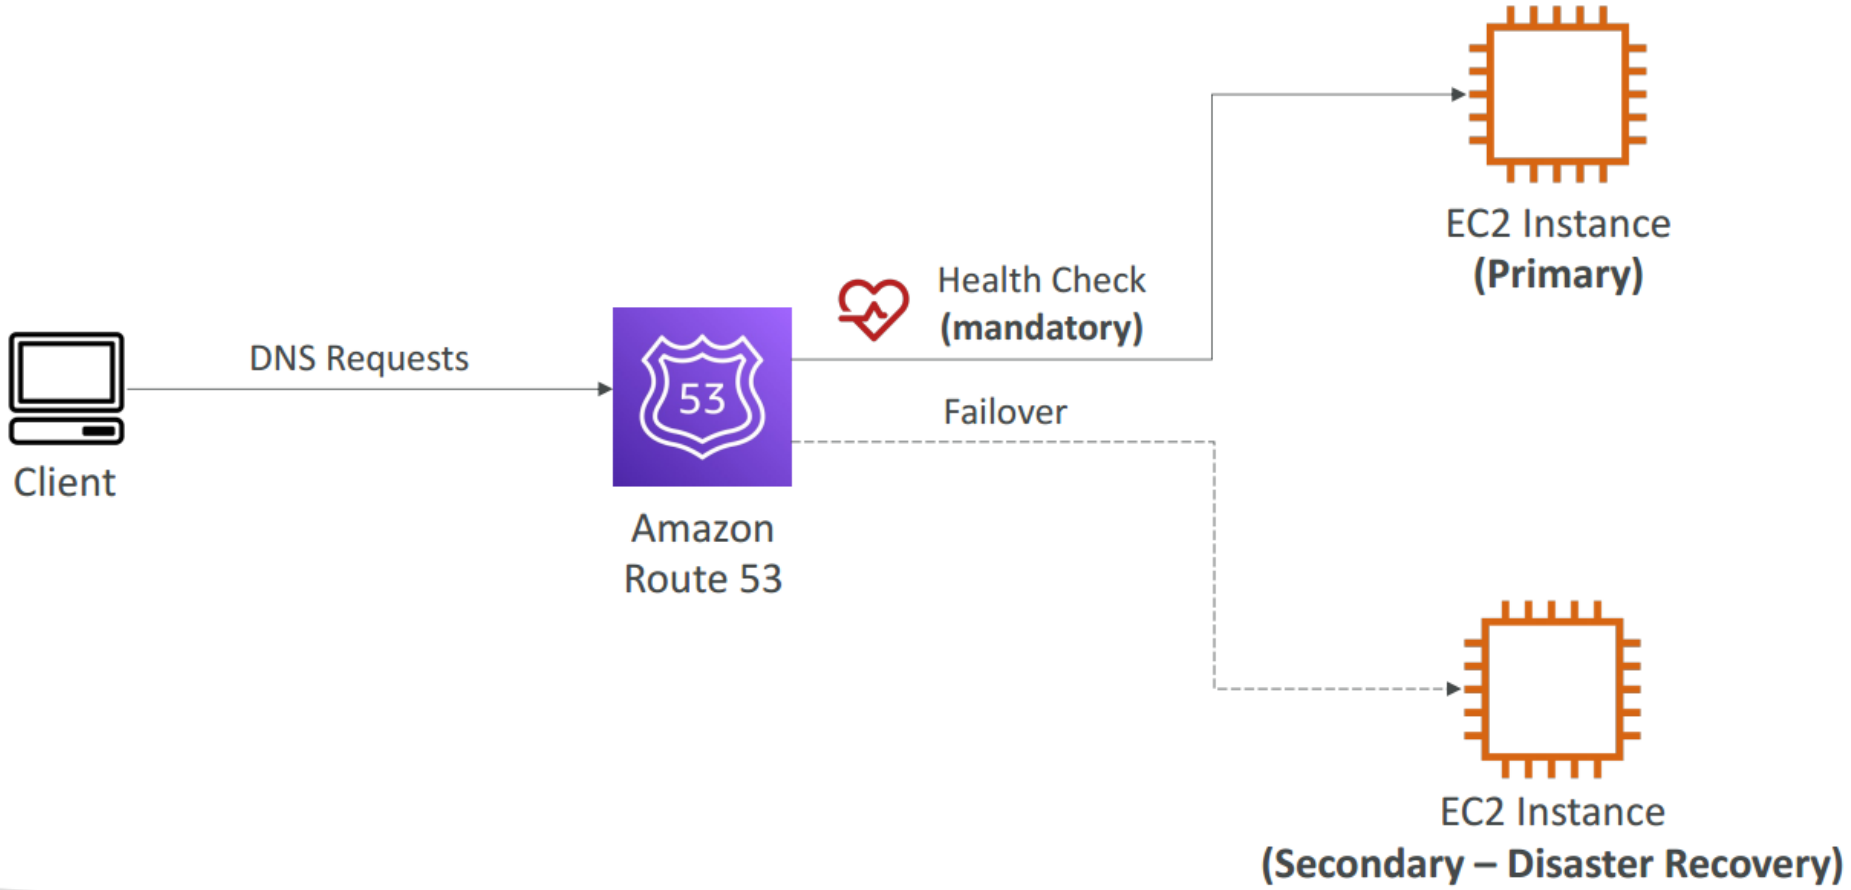

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-failover.html

Routing Policies – Geolocation

- Different from Latency-based!
- This routing is based on **user location**
- Specify location by Country or by US State
- Should create a “Default” record (in case there’s no match on location)
- Use cases: **website localization, restrict content distribution, load balancing,** …
- Can be associated with **Health Checks**

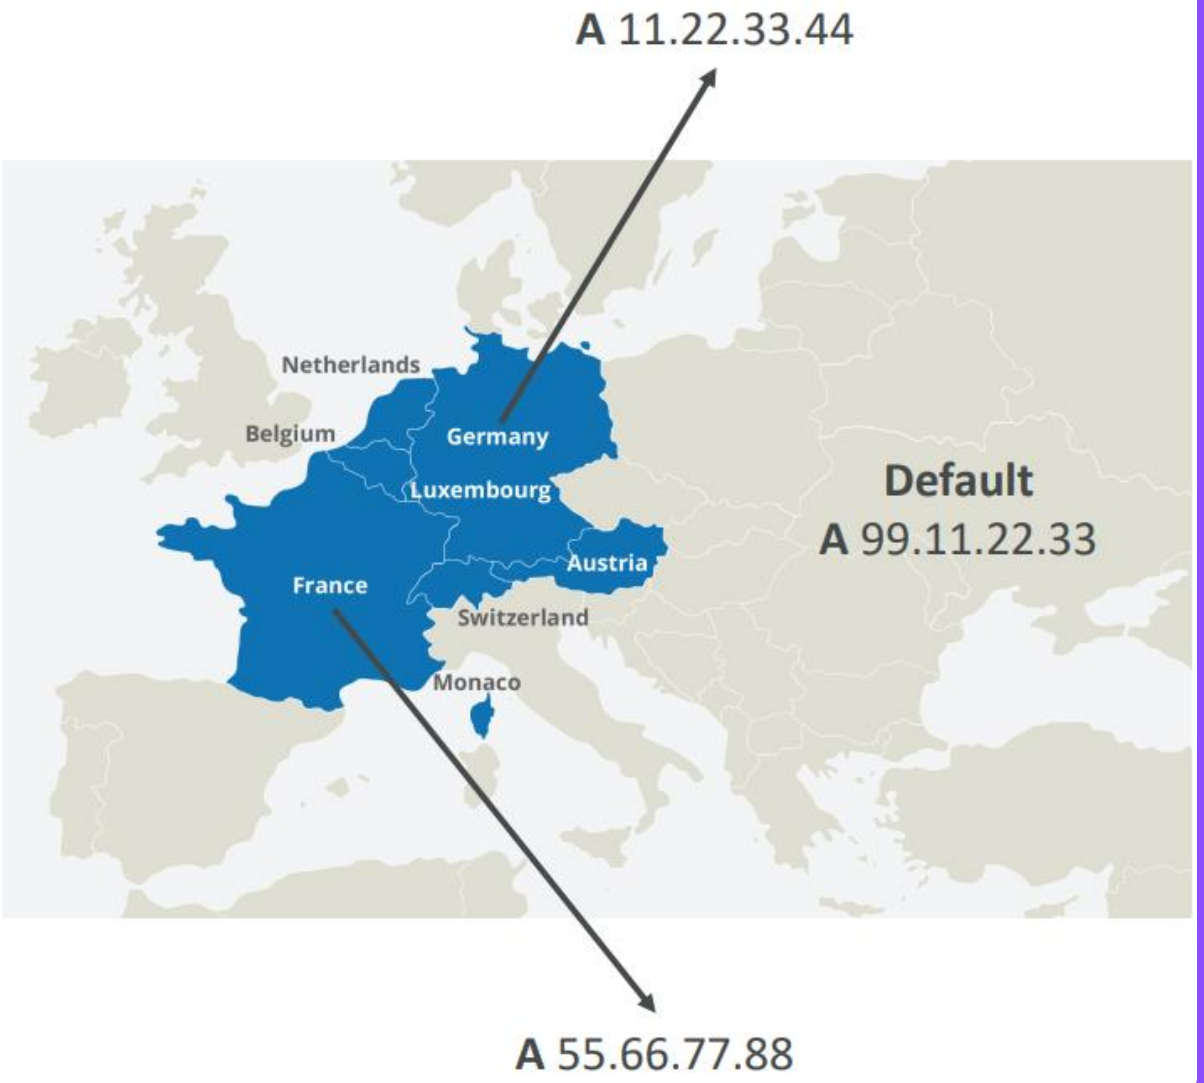

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geo.html

Routing Policies – Geoproximity

- Route traffic to your resources based on the geographic location of users and resources
- Ability to **shift more traffic to resources** based on the defined **bias**
- To change the size of the geographic region, specify bias values:
  - To expand (1 to 99) – more traffic to the resource
  - To shrink (-1 to -99) – less traffic to the resource
- Resources can be:
  - AWS resources (specify AWS region)
  - Non-AWS resources (specify Latitude and Longitude)
- You must use Route 53 **Traffic Flow** to use this feature

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-geoproximity.html

Routing Policies – Geoproximity

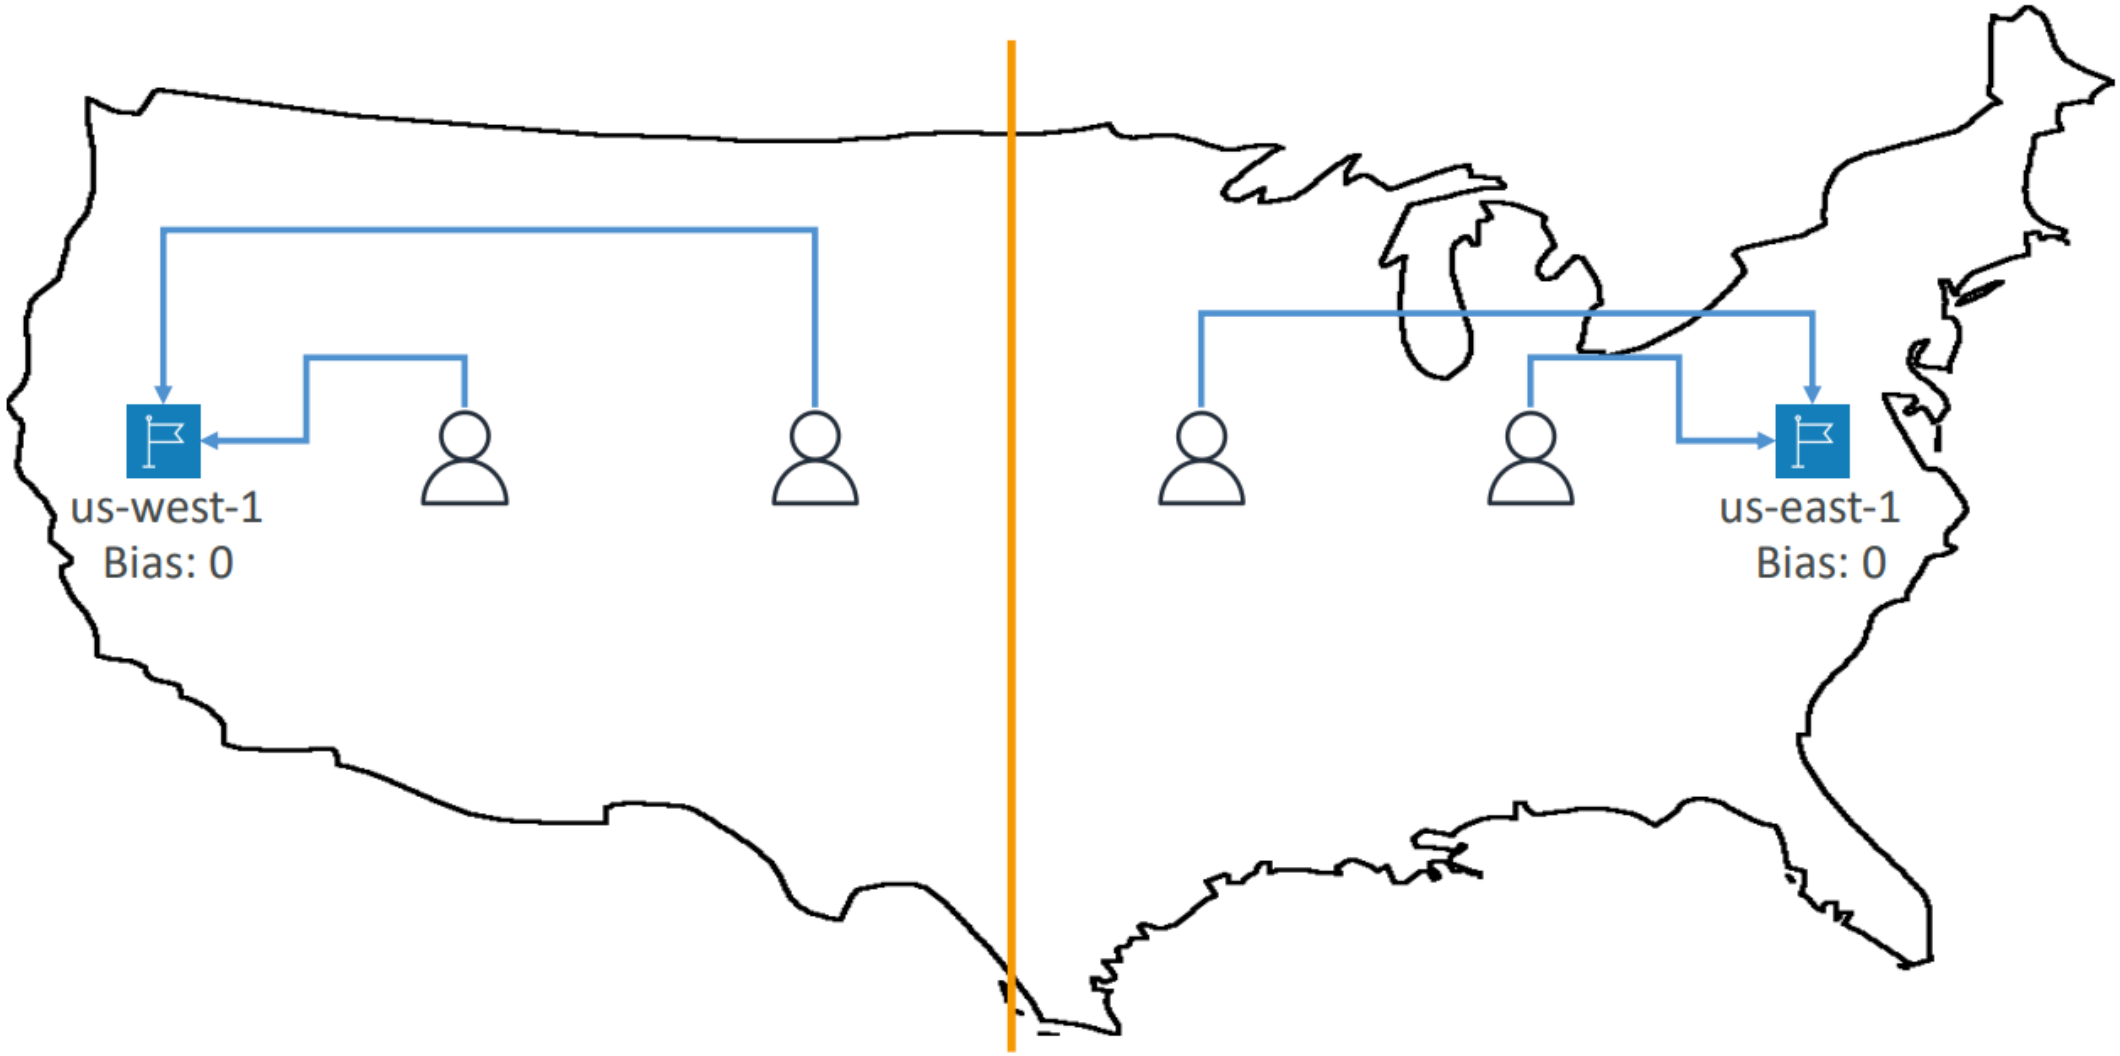

© Copyright by Cloud Mentor Pro | Confidential

Routing Policies – Geoproximity

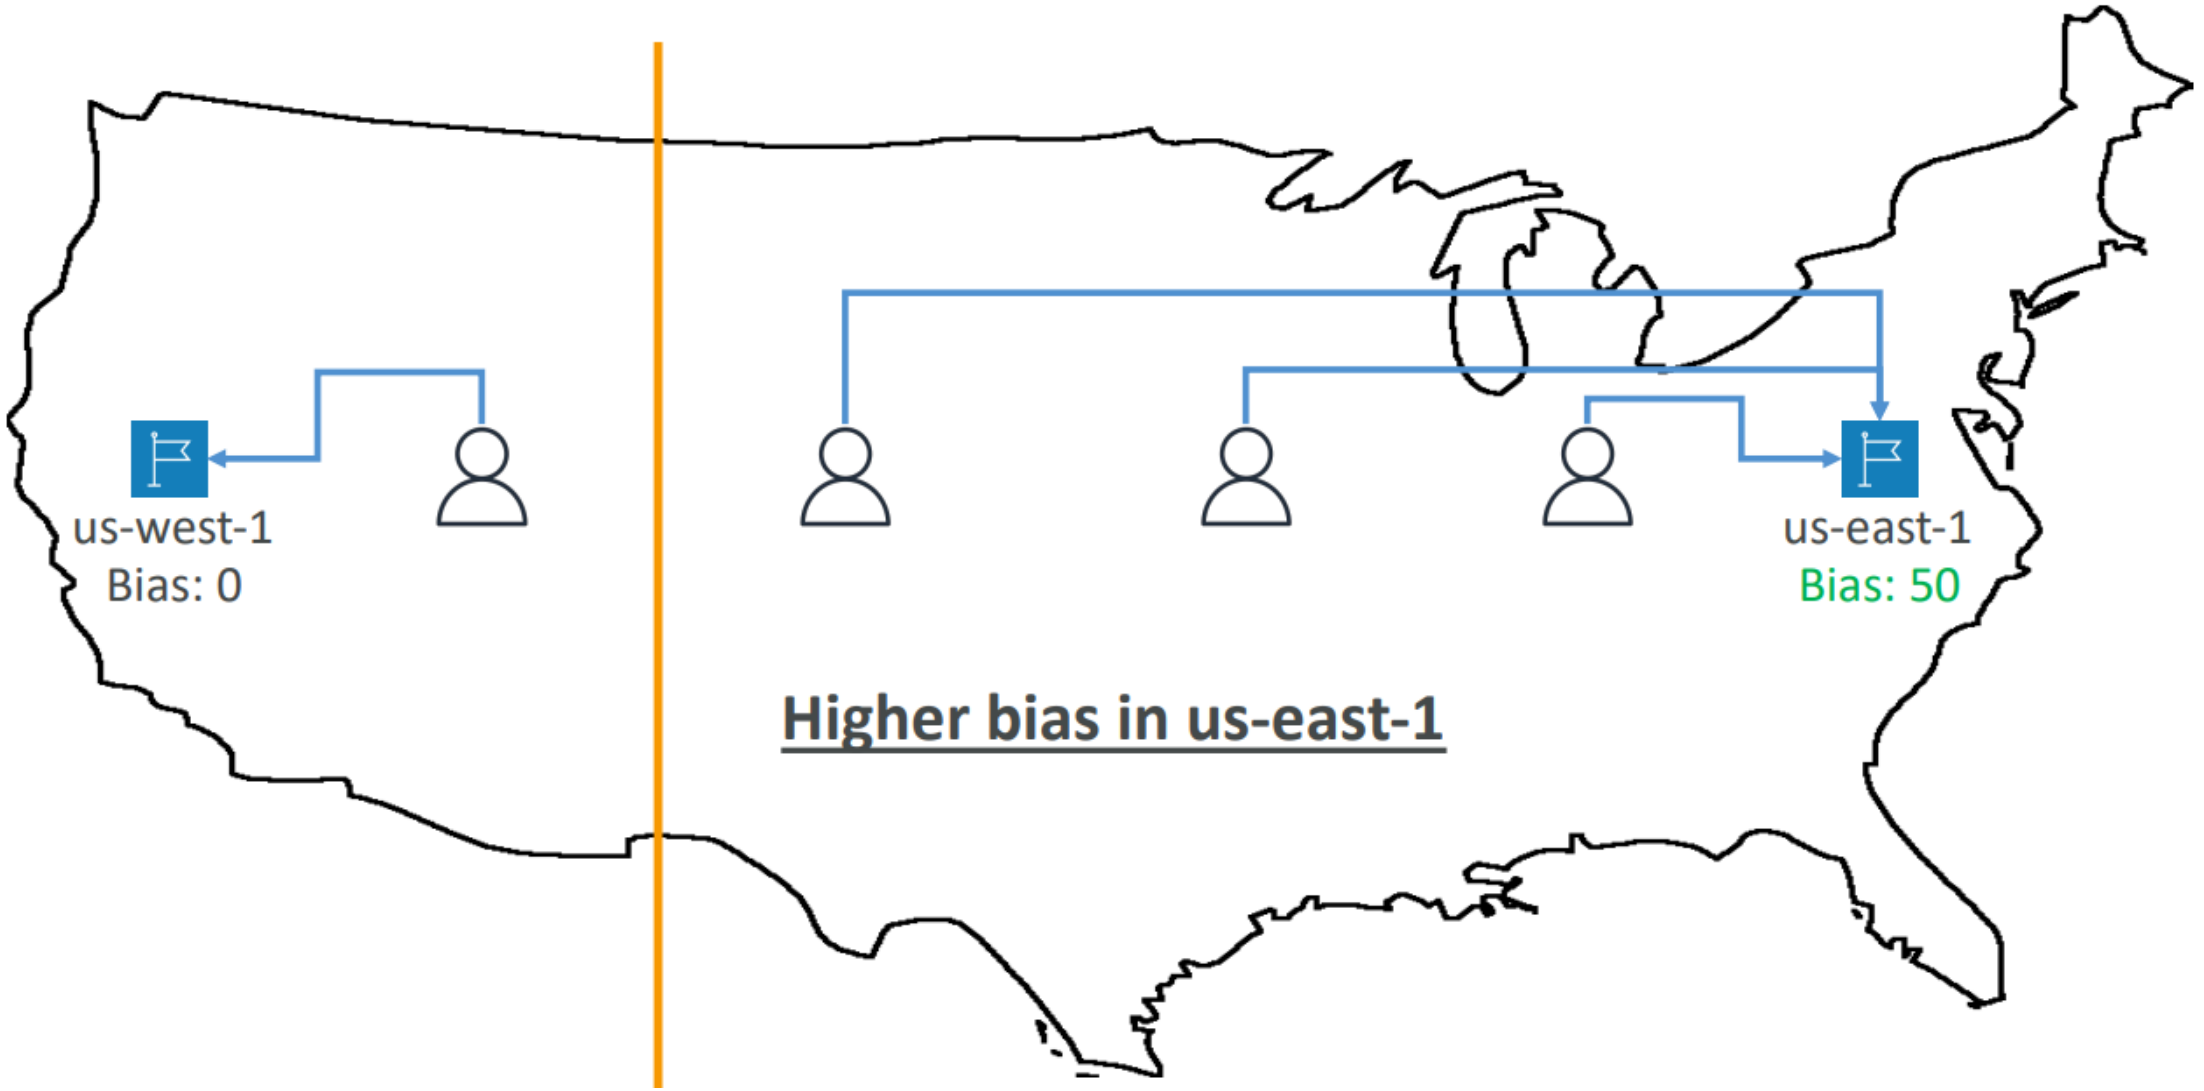

© Copyright by Cloud Mentor Pro | Confidential

Route 53 –Traffic flow

- Simplify the process of creating and maintaining records in large and complex configurations
- Visual editor to manage complex routing decision trees
- Configurations can be saved as **Traffic Flow Policy**
  - Can be applied to different Route 53 Hosted Zones (different domain names)
  - Supports versioning

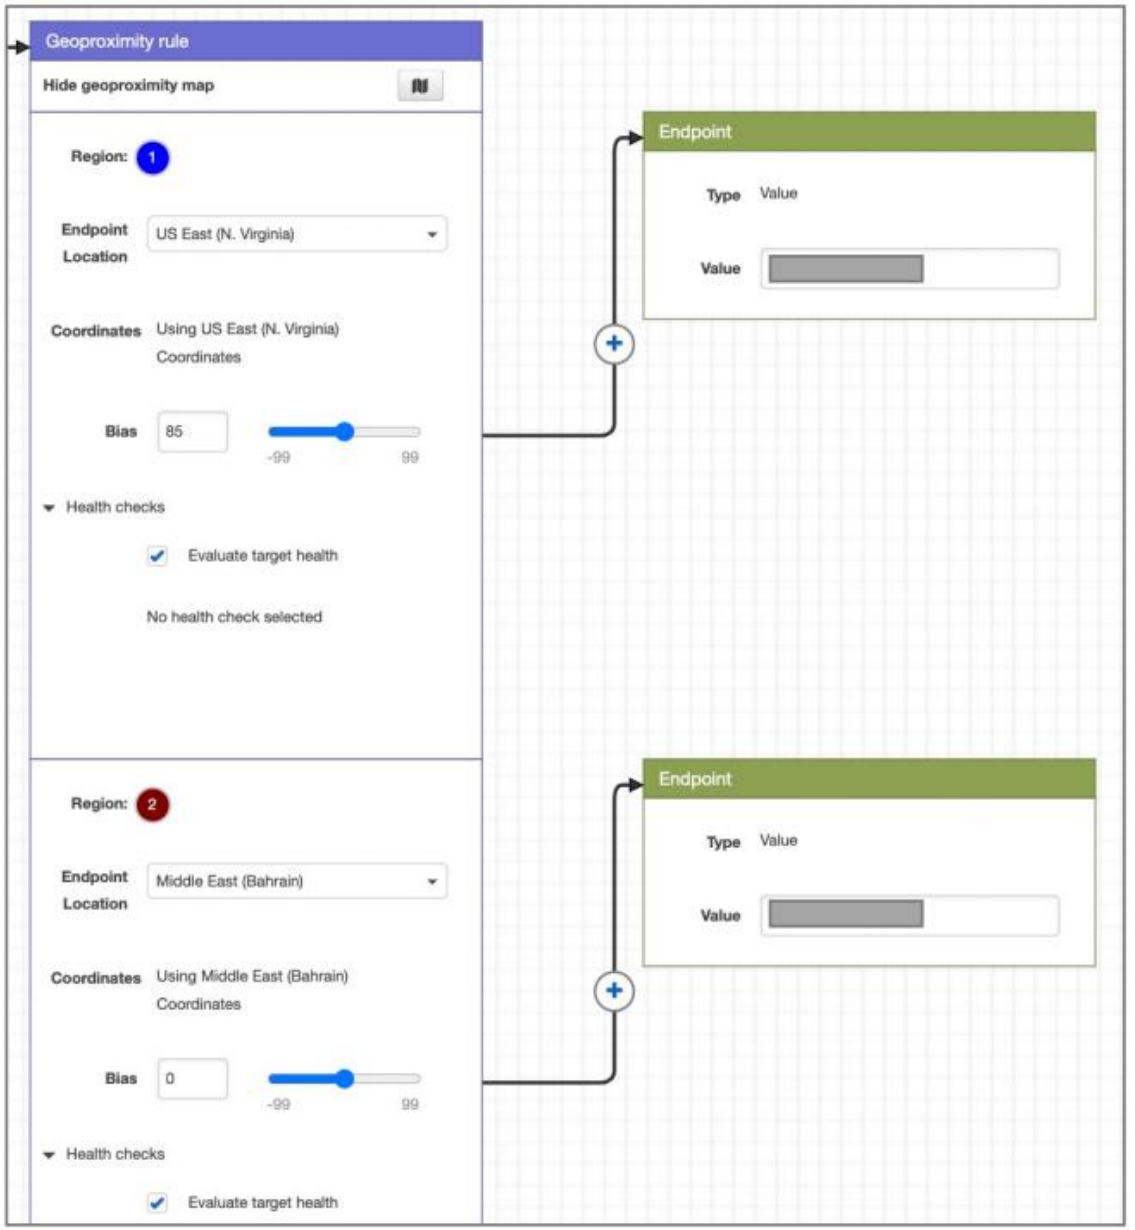

© Copyright by Cloud Mentor Pro | Confidential

Routing Policies – IP-based Routing

- **Routing is based on clients' IP addresses**
- You **provide a list of CIDRs for your clients** and the corresponding endpoints/locations (user-IP-to-endpoint mappings)
- Use cases: **Optimize performance, reduce network costs**…
- Example: route end users from a particular ISP to a specific endpoint

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-ipbased.html

Routing Policies – Multi-Value

- Use when routing traffic to **multiple resources**
- Up to 8 **healthy** records are returned for each Multi-Value query
- It's very similar to simple routing, but with two differences:
  - Routing traffic to **multiple resources**
  - You can have **health checks**

In the following example, there are multiple A Records, each with different values:

Name	Type	Value	TTL	Set ID	Health Check
www.example.com	A Record	192.0.2.2	60	Web1	A
www.example.com	A Record	198.51.100.2	60	Web2	B
www.example.com	A Record	203.0.113.2	60	Web3	C

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/routing-policy-multivalue.html
https://repost.aws/knowledge-center/multivalue-versus-simple-policies

Domain Registar vs. DNS Service

- You buy or register your domain name with a Domain Registrar typically by paying annual charges (e.g., GoDaddy, Amazon Registrar Inc., …)
- The Domain Registrar usually provides you with a DNS service to manage your DNS records
- But you can use another DNS service to manage your DNS records
- Example: purchase the domain from GoDaddy and use Route 53 to manage your DNS records

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/Route53/latest/DeveloperGuide/registrar.html

GoDaddy as Registrar & Route 53 as DNS Service

GoDaddy

Records

We can't display your DNS information because your nameservers aren't managed by us.

Nameservers

Using custom nameservers **Change**

Nameserver
ns-1083.awsdns-07.org
ns-932.awsdns-52.net
ns-1911.awsdns-46.co.uk
ns-481.awsdns-60.com

Amazon Route 53

Public Hosted Zone
stephanetheteacher.com

Hosted zone details — Edit hosted zone

Hosted zone ID Z30IUJCCWPKZUV	Type Public hosted zone	Name servers ns-252.awsdns-31.com ns-1468.awsdns-55.org ns-633.awsdns-15.net ns-1800.awsdns-33.co.uk
Description HostedZone created by Route53 Registrar	Record count 22	
Query log		

© Copyright by Cloud Mentor Pro | Confidential

3rd Party Registrar with Amazon Route 53

If you buy your domain on a 3rd party registrar, you can still use Route 53 as the DNS Service provider

1. Create a Hosted Zone in Route 53 2.
2. Update NS Records on 3rd party website to use Route 53 Name Servers

- **Domain Registrar != DNS Service**
- But every Domain Registrar usually comes with some DNS features

© Copyright by Cloud Mentor Pro | Confidential

3rd Party Registrar with Amazon Route 53

Before

- Domain Management: GoDaddy
- DNS Management : GoDaddy

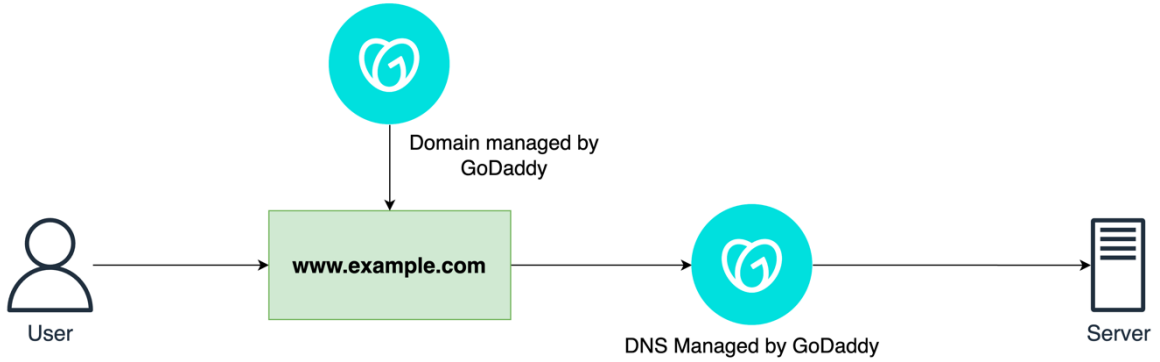

After

- Domain Management : GoDaddy
- DNS Management : Route53

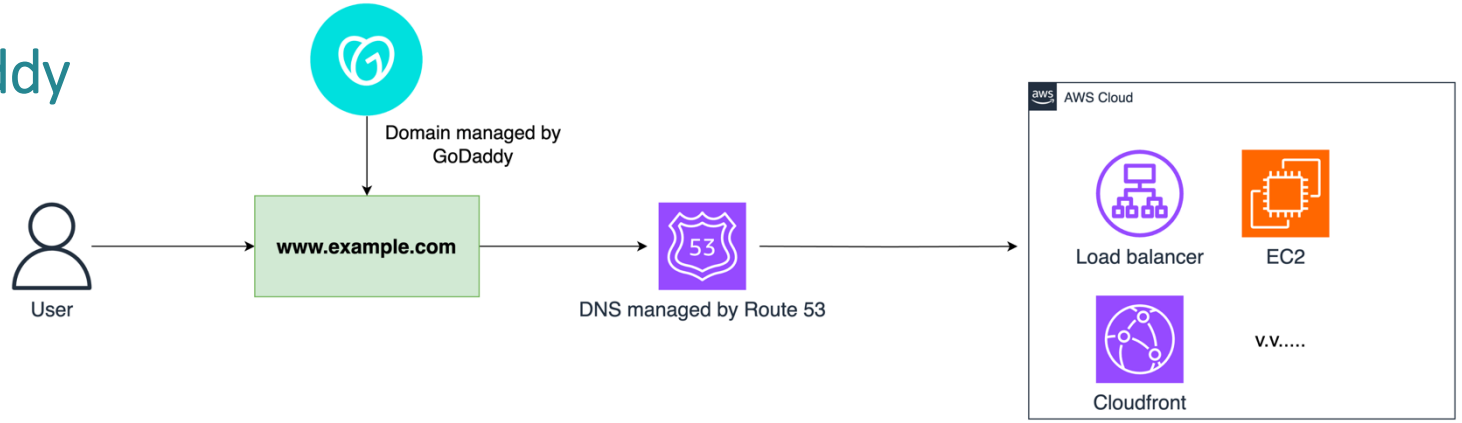

© Copyright by Cloud Mentor Pro | Confidential

https://blog.cloudmentor.pro/blog/aws/huong-dan-chuyen-doi-dns-tu-ben-thu-3-sang-route-53

Exercise

- *Create Amazon RDS for MySQL*

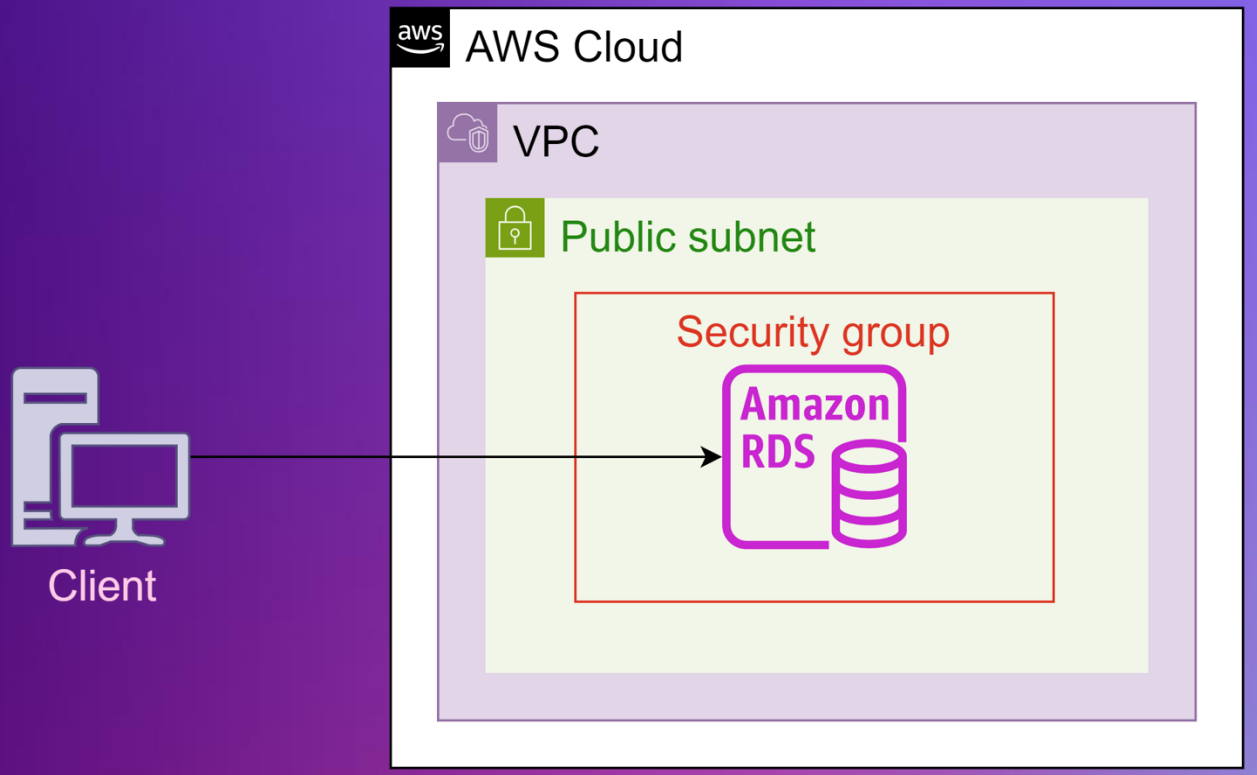

© Copyright by Cloud Mentor Pro | Confidential

Section 4

- *Amazon S3*
- *Amazon S3 – Advanced*
- *Amazon S3 – Security*

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3

© Copyright by Cloud Mentor Pro | Confidential

Section introduction

- Amazon S3 is one of the main building blocks of AWS
- It's advertised as "infinitely scaling" storage
- Benefits
  - **Scalability**
  - S3 provides the most **durable storage** in the cloud and industry leading availability.
  - S3 is **secure, private, and encrypted** by default
  - Multiple storage classes with the **best price performance** for any workload and automated data lifecycle management
  - Resiliency, flexibility, latency, and throughput, to ensure **never limits performance**

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Use cases

- Backup and storage
- Disaster Recovery
- Archive
- Hybrid Cloud storage
- Application hosting
- Media hosting
- Data lakes & big data analytics
- Software delivery
- Static website

Nasdaq stores 7 years of data into S3 Glacier

Sysco runs analytics on its data and gain business insights

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 - Buckets

- Amazon S3 allows people to store objects (files) in “buckets” (directories)
- Buckets must have a globally unique name **(across all regions all accounts)**
- Buckets are defined at the region level
- S3 looks like a global service but buckets are **created in a region**
- Naming convention
  - No uppercase, No underscore
  - 3-63 characters long
  - Not an IP
  - Must start with lowercase letter or number
  - Must NOT start with the prefix xn–
  - Must NOT end with the suffix -**s3alias**

S3 Bucket

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 - Objects

- Objects (files) have a Key
- The key is the FULL path:
  - s3://my-bucket/my_file.txt
  - s3://my-bucket/my_folder1/another_folder/my_file.txt
- The key is composed of prefix + object name
  - s3://my-bucket/my_folder1/another_folder/my_file.txt
- Max Object Size is **5TB** (5000GB)
- **Metadata**：A set of name-value pairs, information regarding the object
- **Tags** (Unicode key / value pair – up to 10) – useful for security / lifecycle
- **Version ID** (if versioning is enabled)

S3 Bucket with objects

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Security

- User-Based
  - **IAM Policies** –allowed for a specific user from IAM
- Resource-Based
  - **Bucket Policies** – bucket wide rules from the S3 console - allows cross account
  - **Object Access Control List (ACL)** – finer grain (can be disabled)
  - **Bucket Access Control List (ACL)** – less common (can be disabled)
- **Note**: an IAM principal can access an S3 object if
  - The user IAM permissions ALLOW it OR the resource policy ALLOWS it
  - AND there's no explicit DENY
- **Encryption**: encrypt objects in Amazon S3 using encryption keys

© Copyright by Cloud Mentor Pro | Confidential
https://docs.aws.amazon.com/AmazonS3/latest/userguide/security.html

S3 Bucket Policies

- JSON based policies
  - **Effect**: Allow / Deny
  - **Principal**: The account or user to apply the policy to
  - **Actions**: Set of API to Allow or Deny
  - **Resources**: buckets and objects

- Use S3 bucket for policy to:
  - Grant **public access** to the bucket
  - Force objects to be **encrypted** at upload
  - Grant access to another account (**Cross Account**)

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Sid": "PublicRead",
            "Effect": "Allow",
            "Principal": "*",
            "Action": [
                "s3:GetObject"
            ],
            "Resource": [
                "arn:aws:s3:::examplebucket/*"
            ]
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

Example: Public Access - Use Bucket Policy

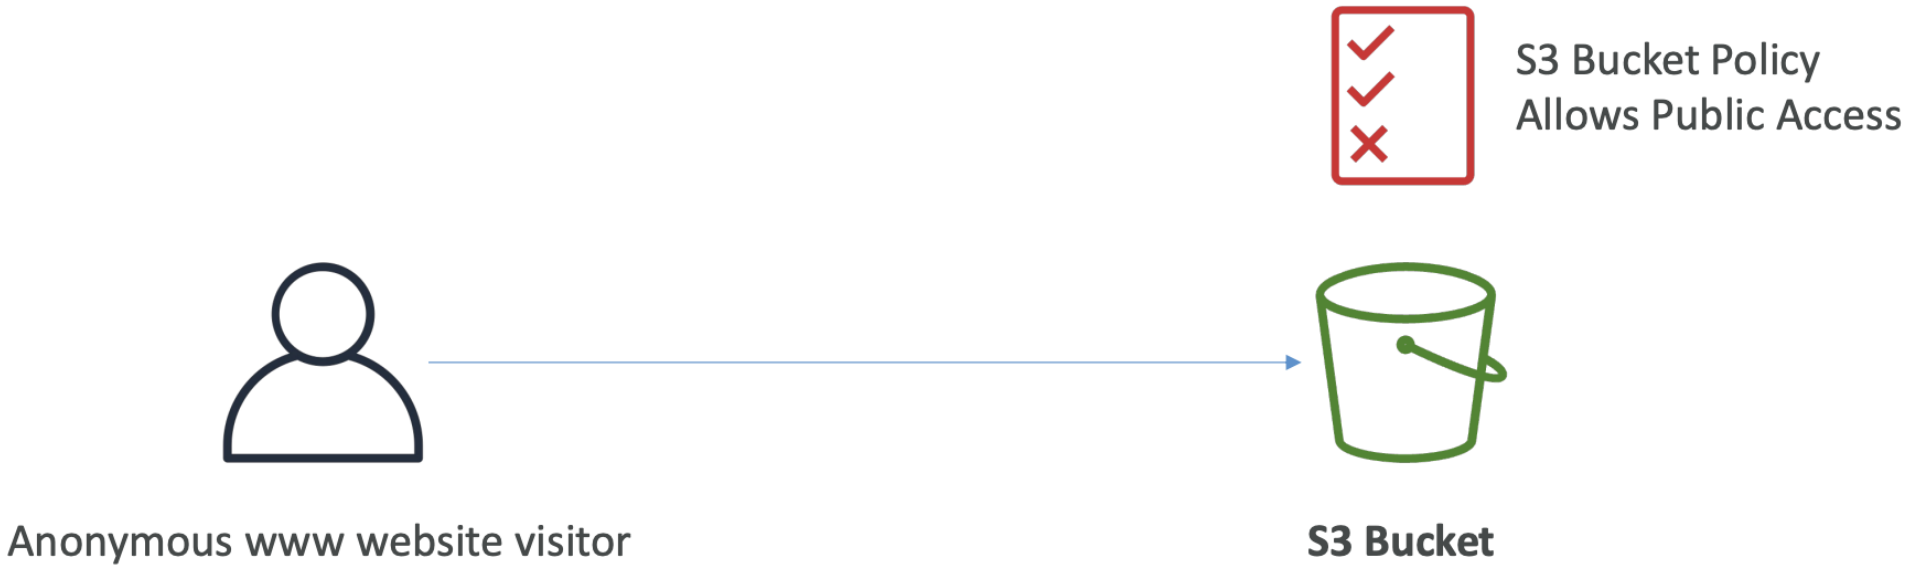

© Copyright by Cloud Mentor Pro | Confidential

Example: User Access to S3 – IAM permissions

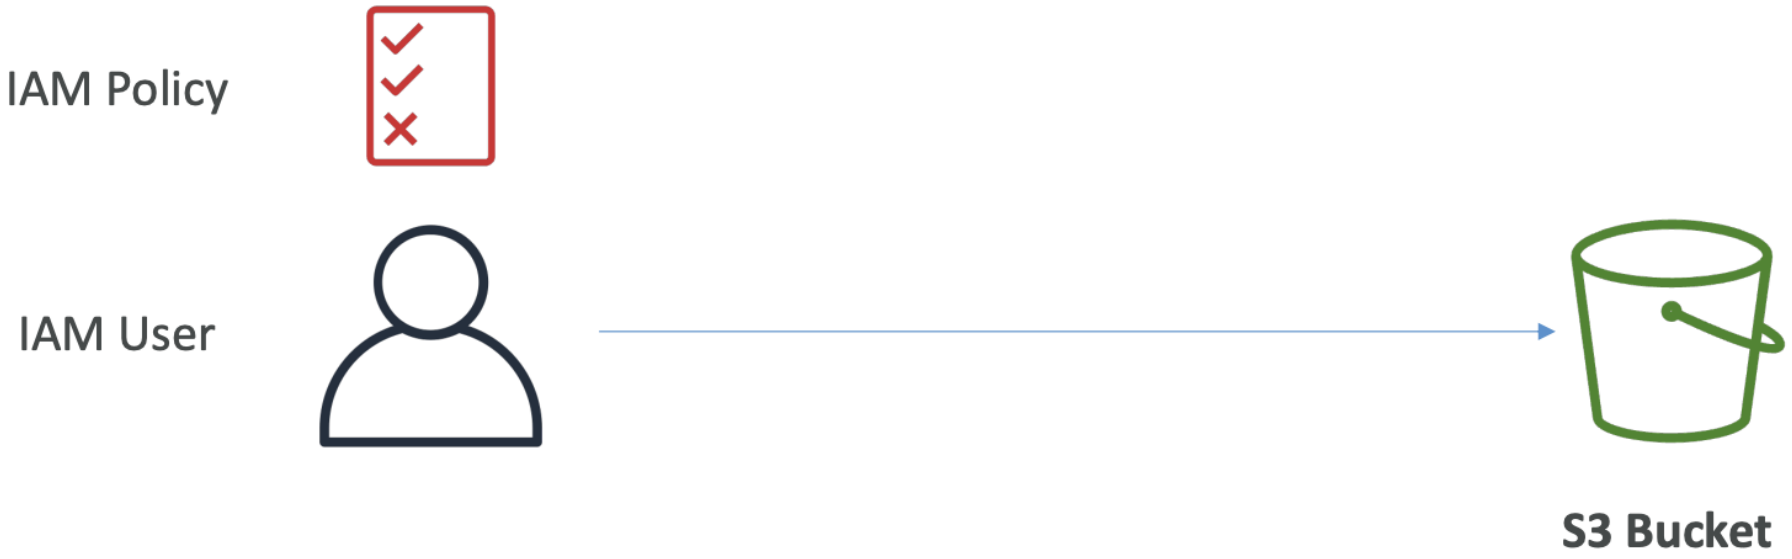

© Copyright by Cloud Mentor Pro | Confidential

Example: EC2 instance access - Use IAM Roles

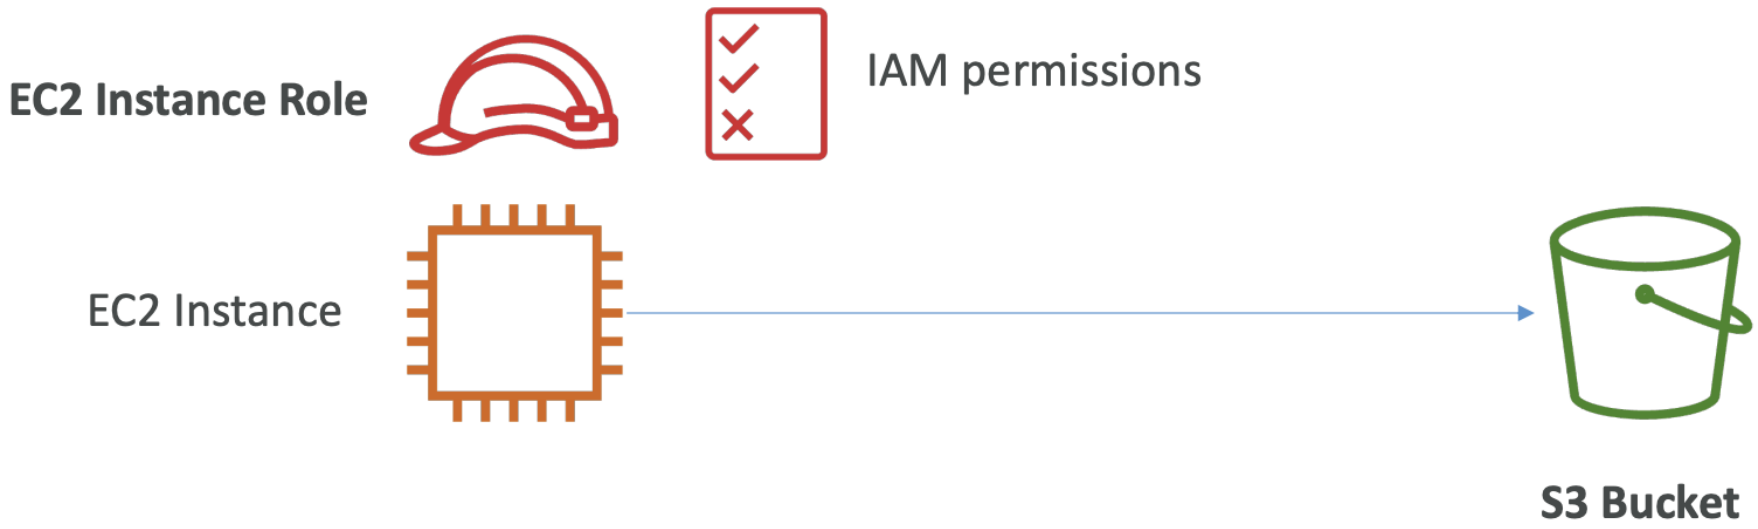

© Copyright by Cloud Mentor Pro | Confidential

Advanced: Cross-Account Access – Use Bucket Policy

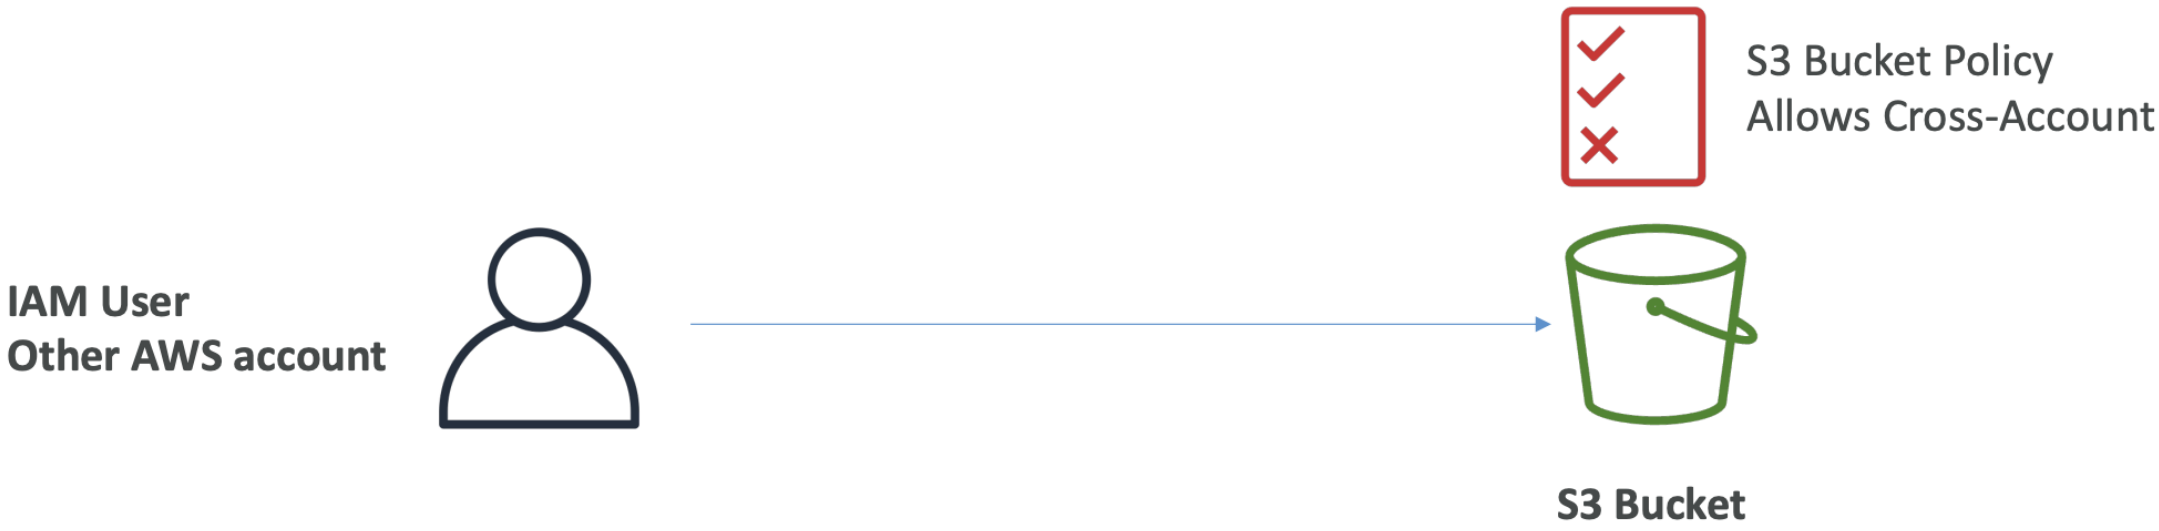

© Copyright by Cloud Mentor Pro | Confidential

Bucket settings for Block Public Access

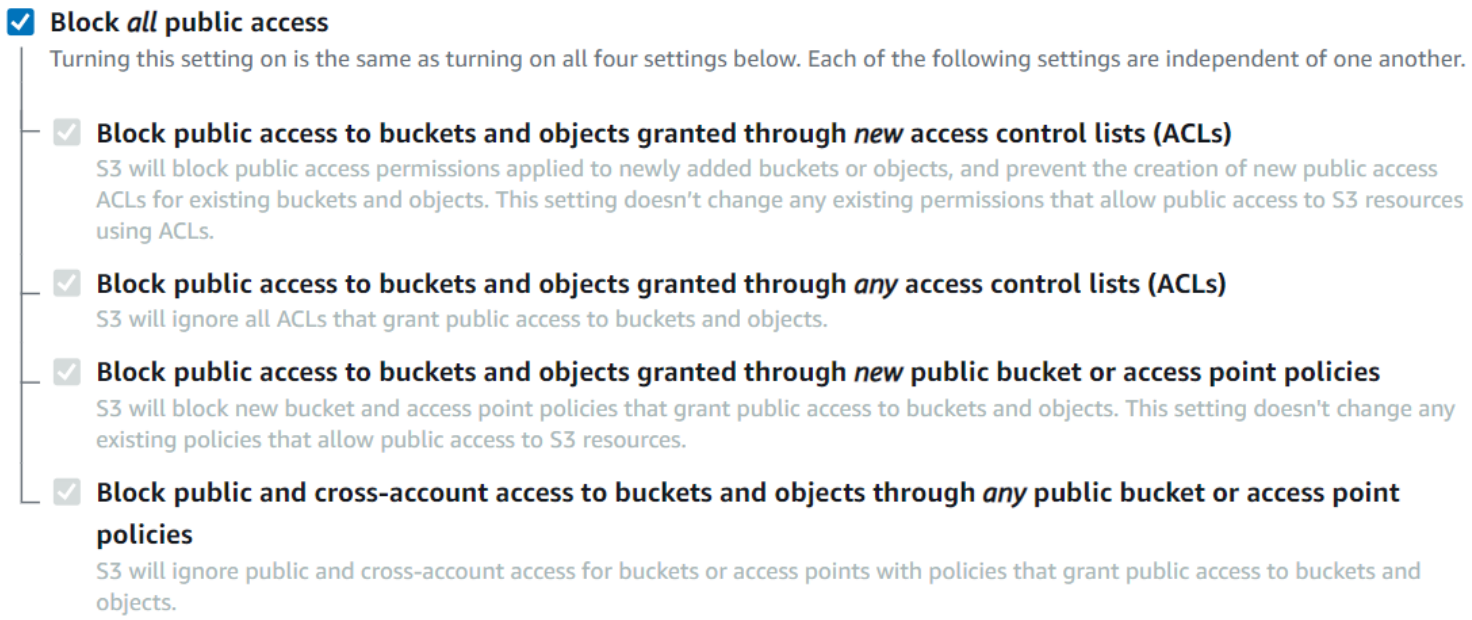

- These settings were created to prevent company data leaks
- If you know your bucket should never be public, leave these on
- Can be set at the account level

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Static Website Hosting

- S3 can host static websites and have them accessible on the Internet
- If you get a **403 Forbidden** error, make sure the bucket policy allows public reads!
  - Block all Public Access: Off (disabled)
  - Bucket policy allows public reads!

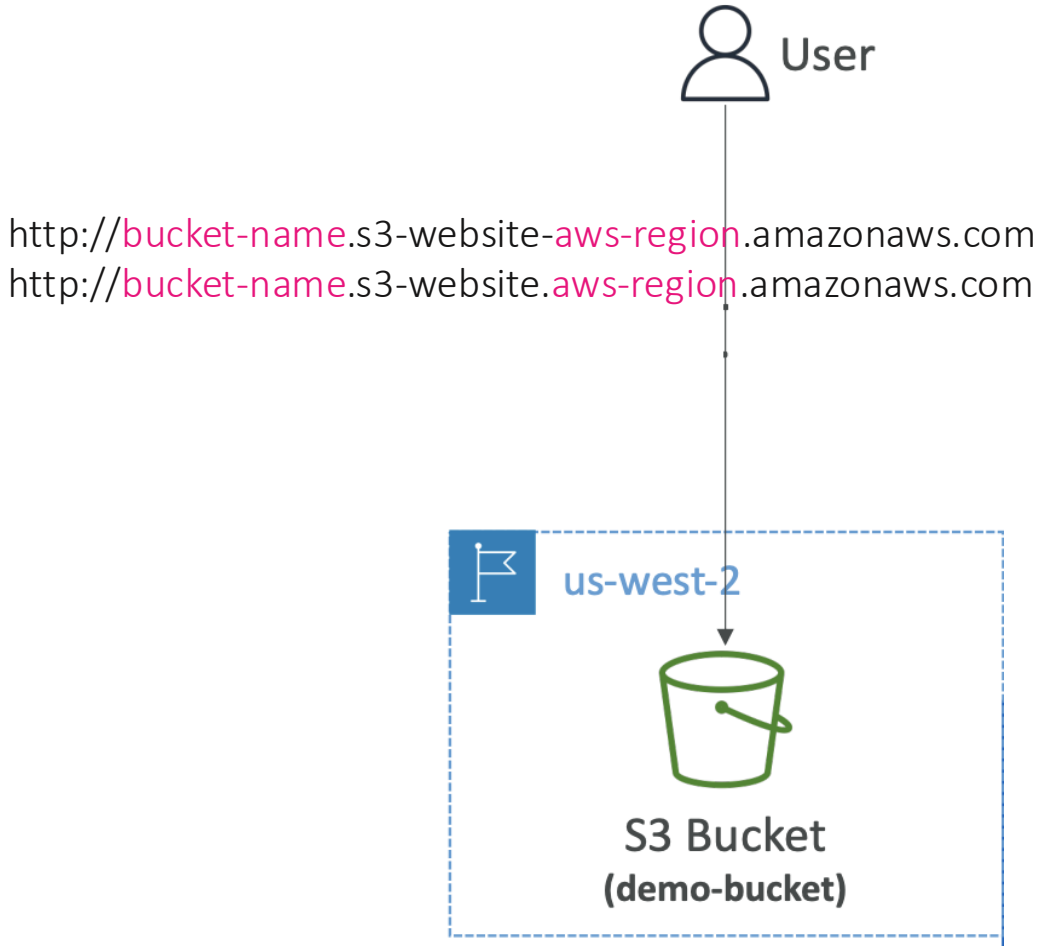

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/AmazonS3/latest/userguide/EnableWebsiteHosting.html

Amazon S3 -Versioning

- You can version your files in Amazon S3
- It is enabled at the **bucket level**
- Same key overwrite will change the "version": 1, 2, 3....
- It is best practice to version your buckets
  - Preserve, retrieve, and restore every **version of every object**
  - Help you **recover objects from accidental deletion** or **overwrite**
- Notes:
  - Any file that is not versioned prior to enabling versioning will have version "null"
  - Suspending versioning does not delete the previous versions

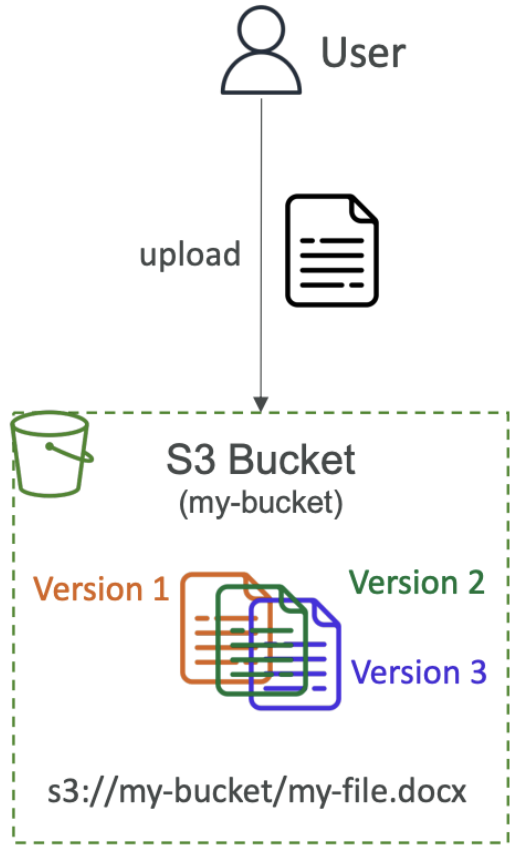

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Replication (CRR & SRR)

- Must enable Versioning in source and destination buckets
- Cross-Region Replication (CRR)
- Same-Region Replication (SRR)
- Buckets can be in different AWS accounts
- Copying is asynchronous
- Must give proper IAM permissions to S3

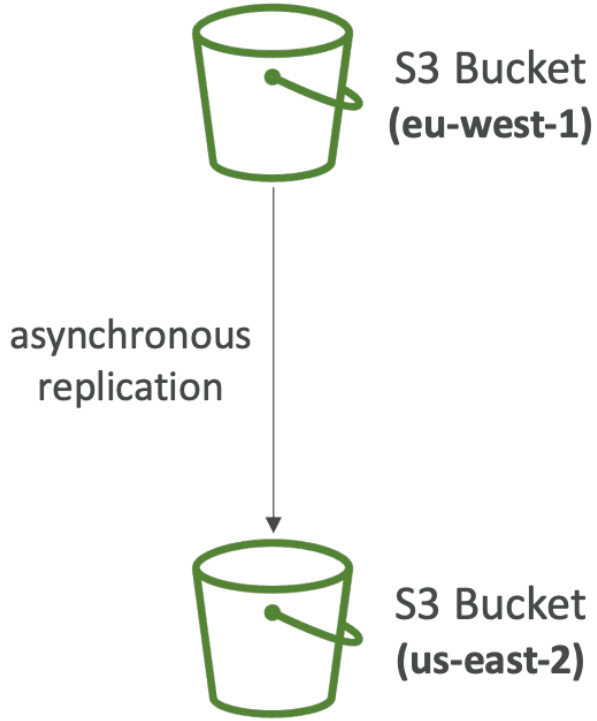

- Use cases:
  - CRR – compliance, lower latency access, replication across accounts
  - SRR – log aggregation, live replication between production and test accounts

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Replication (Notes)

- After you enable Replication, only new objects are replicated
- **Optionally**, you can replicate existing objects using **S3 Batch Replication**
  - Replicates existing objects and objects that failed replication

- For DELETE operations
  - **Can replicate delete markers** from source to target (optional setting)
  - Deletions with a version ID are not replicated (to avoid malicious deletes)

- **There is no "chaining" of replication**
  - If bucket 1 has replication into bucket 2, which has replication into bucket 3
  - Then objects created in bucket 1 are not replicated to bucket 3

© Copyright by Cloud Mentor Pro | Confidential

S3 Storage Classes

- Amazon S3 Standard - General Purpose
- Amazon S3 Standard-Infrequent Access (IA)
- Amazon S3 One Zone-Infrequent Access
- Amazon S3 Glacier Instant Retrieval
- Amazon S3 Glacier Flexible Retrieval
- Amazon S3 Glacier Deep Archive
- Amazon S3 Intelligent Tiering

- Can move between classes manually or using S3 Lifecycle configurations

© Copyright by Cloud Mentor Pro | Confidential

S3 Durability and Availability

- Durability:
  - High durability (99.999999999%, 11 9's) of objects across multiple AZ
  - If you store 10,000,000 objects with Amazon S3, you can on average expect to incur a loss of a single object once every 10,000 years
  - Same for all storage classes

- Availability:
  - Measures how readily available a service is
  - Varies depending on storage class
  - Example: S3 standard has 99.99% availability = not available 53 minutes a year

© Copyright by Cloud Mentor Pro | Confidential

S3 Standard – General Purpose

- 99.99% Availability
- Used for frequently accessed data
- Low latency and high throughput
- Sustain 2 concurrent facility failures

- Use Cases: Big Data analytics, mobile & gaming applications, content distribution…

© Copyright by Cloud Mentor Pro | Confidential

S3 Storage Classes – Infrequent Access

- For data that is less frequently accessed, but requires rapid access when needed
- Lower cost than S3 Standard

- **Amazon S3 Standard-Infrequent Access (S3 Standard-IA)**
  - 99.9% Availability
  - Use cases: Disaster Recovery, backups

- **Amazon S3 One Zone-Infrequent Access (S3 One Zone-IA)**
  - High durability (99.999999999%) in a single AZ; data lost when AZ is destroyed
  - 99.5% Availability
  - Use Cases: Storing secondary backup copies of on-premises data, or data you can recreate

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Glacier Storage Classes

- Low-cost object storage meant for archiving / backup
- Pricing: price for storage + object retrieval cost

- **Amazon S3 Glacier Instant Retrieval**
  - Millisecond retrieval, great for data accessed once a quarter
  - Minimum storage duration of 90 days
- **Amazon S3 Glacier Flexible Retrieval (formerly Amazon S3 Glacier):**
  - Expedited (1 to 5 minutes), Standard (3 to 5 hours), Bulk (5 to 12 hours) – free
  - Minimum storage duration of 90 days
- **Amazon S3 Glacier Deep Archive – for long term storage:**
  - Standard (12 hours), Bulk (48 hours)
  - Minimum storage duration of 180 days

© Copyright by Cloud Mentor Pro | Confidential

S3 Intelligent-Tiering

- Small monthly monitoring and auto-tiering fee
- Moves objects automatically between Access Tiers based on usage
- There are no retrieval charges in S3 Intelligent-Tiering

- *Frequent Access tier (automatic):* default tier
- *Infrequent Access tier (automatic):* objects not accessed for 30 days
- *Archive Instant Access tier (automatic):* objects not accessed for 90 days
- *Archive Access tier (optional):* configurable from 90 days to 700+ days
- *Deep Archive Access tier (optional):* config. from 180 days to 700+ days

© Copyright by Cloud Mentor Pro | Confidential

S3 - Storage Classes

Storage class	Designed for	AZs	Min storage duration	Min billable object size	Other considerations
S3 Standard	**Frequently accessed** data (more than once a month) with millisecond access	>= 3	None	None	**Big Data analytics, mobile & gaming applications, content distribution…**
S3 Standard-IA	Long-lived, **infrequently accessed** data (once a month) with millisecond access	>= 3	30 days	128 KB	Per-GB retrieval fees apply. **Disaster Recovery, backups**
S3 Intelligent-Tiering	Data with unknown, changing, or **unpredictable** access patterns	>= 3	None	None	Monitoring and automation fees per object apply. No retrieval fees. **Moves objects automatically between Access Tiers based on usage**
S3 One Zone-IA	**Recreatable**, infrequently accessed data (once a month) with millisecond access	**1**	30 days	128 KB	Per-GB retrieval fees apply. Not resilient to the loss of the Availability Zone. **Backup copies of on-premises data**
S3 Glacier Instant Retrieval	Long-lived, **archive data accessed once a quarter with millisecond access**	>= 3	90 days	128 KB	Pricing: price for storage + object retrieval cost **Low-cost object storage for archiving / backup**
S3 Glacier Flexible Retrieval	Long-lived archive data accessed once a year with **retrieval times:** **.Expedited (1 – 5 mins)** **.Standard (3 – 5 hours)** **.Bulk (5 – 12 hours)**	>= 3	90 days	NA*	
S3 Glacier Deep Archive	Long-lived archive data accessed less than once a year **retrieval times:** **.Standard (12 hours)** **.Bulk (48 hours)**	>= 3	180 days	NA**	

https://docs.aws.amazon.com/AmazonS3/latest/userguide/storage-class-intro.html

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Advanced

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Moving between Storage Classes

- You can transition objects between storage classes

- For infrequently accessed object, move them to **Standard IA**
- For archive objects that you don't need fast access to, move them to **Glacier or Glacier Deep Archive**

- Moving objects can be automated using a **Lifecycle Rules**

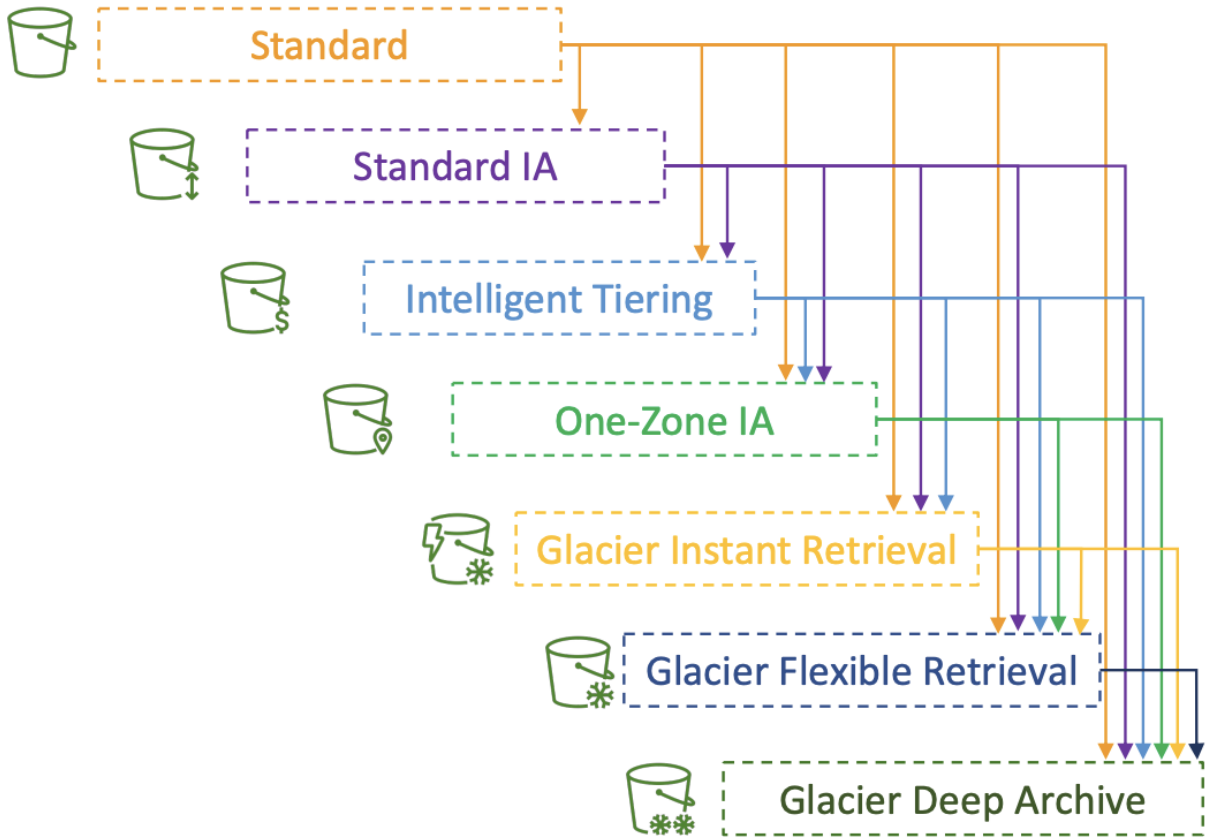

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Lifecycle Rules

- **Transition Actions** – configure objects to transition to another storage class
  - Move objects to Standard IA class 60 days after creation
  - Move to Glacier for archiving after 6 months

- **Expiration actions** – configure objects to expire (delete) after some time
  - Access log files can be set to delete after a 365 days
  - **Can be used to delete old versions of files (if versioning is enabled)**
  - Can be used to delete incomplete Multi-Part uploads

- Rules can be created for a certain prefix (example: *s3://mybucket/mp3/**)
- Rules can be created for certain objects Tags (example: *Department: Finance*)

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Lifecycle Rules (Scenario 1)

- Your application on EC2 creates images thumbnails after profile photos are uploaded to Amazon S3. These thumbnails can be easily recreated, and only need to be kept for 60 days. The source images should be able to be immediately retrieved for these 60 days, and afterwards, the user can wait up to 6 hours. How would you design this?

- S3 source images can be on **Standard**, with a lifecycle configuration to transition them to **Glacier** after 60 days
- S3 thumbnails can be on **One-Zone IA**, with a lifecycle configuration to expire them (delete them) after 60 days

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Lifecycle Rules (Scenario 2)

- A rule in your company states that you should be able to recover your deleted S3 objects immediately for 30 days, although this may happen rarely. After this time, and for up to 365 days, deleted objects should be recoverable within 48 hours.

- **Enable S3 Versioning** in order to have object versions, so that "deleted objects" are in fact hidden by a "delete marker" and can be recovered
- Transition the "noncurrent versions" of the object to **Standard IA**
- Transition afterwards the "noncurrent versions" to **Glacier Deep Archive**

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Analytics – Storage Class Analysis

- Help you decide when to transition objects to the right storage class
- Recommendations for Standard and Standard IA
- Does NOT work for One-Zone IA or Glacier
- Report is updated daily
- 24 to 48 hours to start seeing data analysis
- Good first step to put together Lifecycle Rules

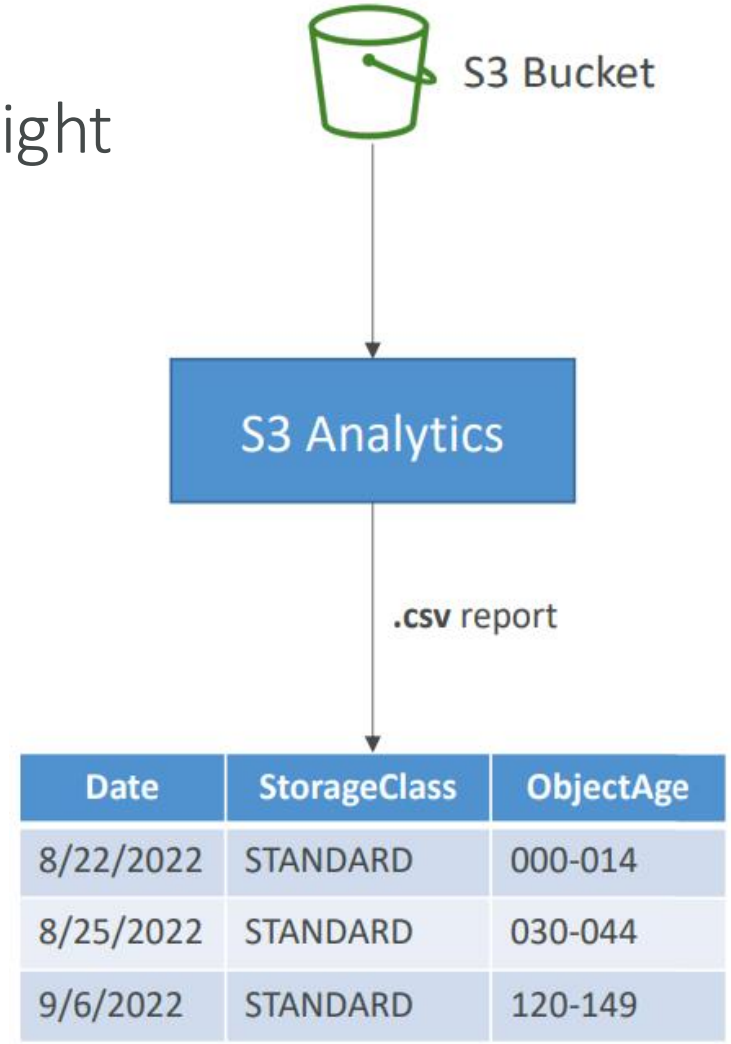

Date	StorageClass	ObjectAge
8/22/2022	STANDARD	000-014
8/25/2022	STANDARD	030-044
9/6/2022	STANDARD	120-149

© Copyright by Cloud Mentor Pro | Confidential

S3 Event Notifications

- S3:ObjectCreated, S3:ObjectRemoved, S3:ObjectRestore, S3:Replication…
- Object name filtering possible (*.jpg)
- Use case: generate thumbnails of images uploaded to S3
- **Can create as many "S3 events" as desired**

- S3 event notifications typically deliver events in seconds but can sometimes take a minute or longer

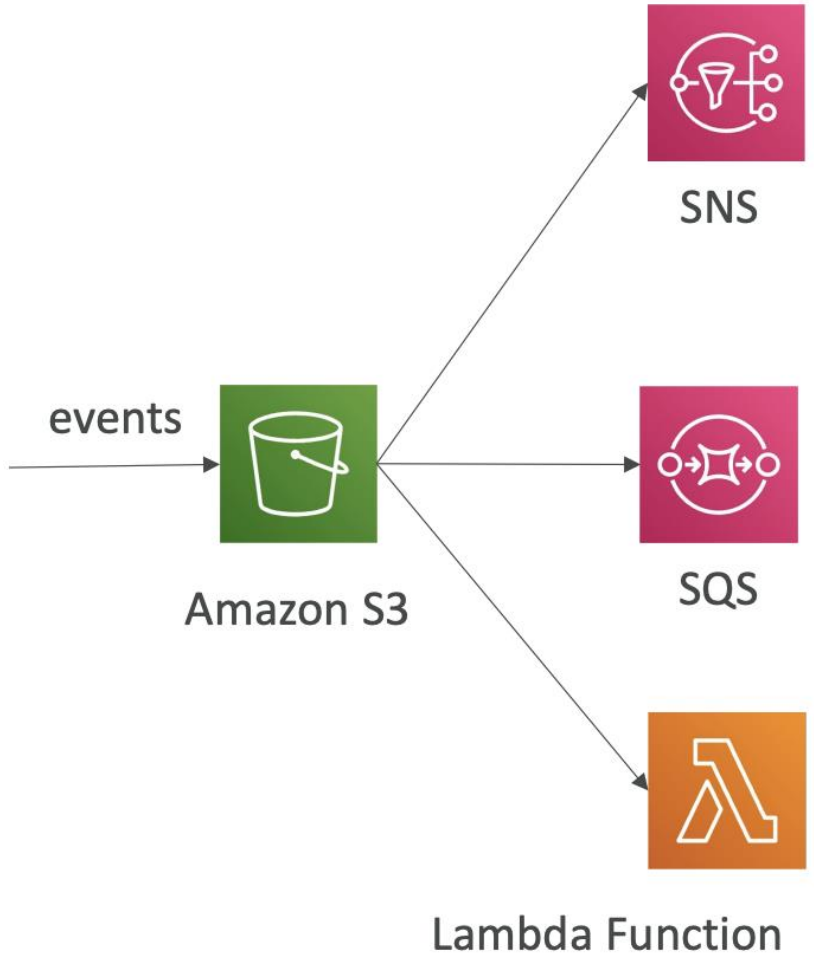

© Copyright by Cloud Mentor Pro | Confidential

S3 Event Notifications – IAM Permissions

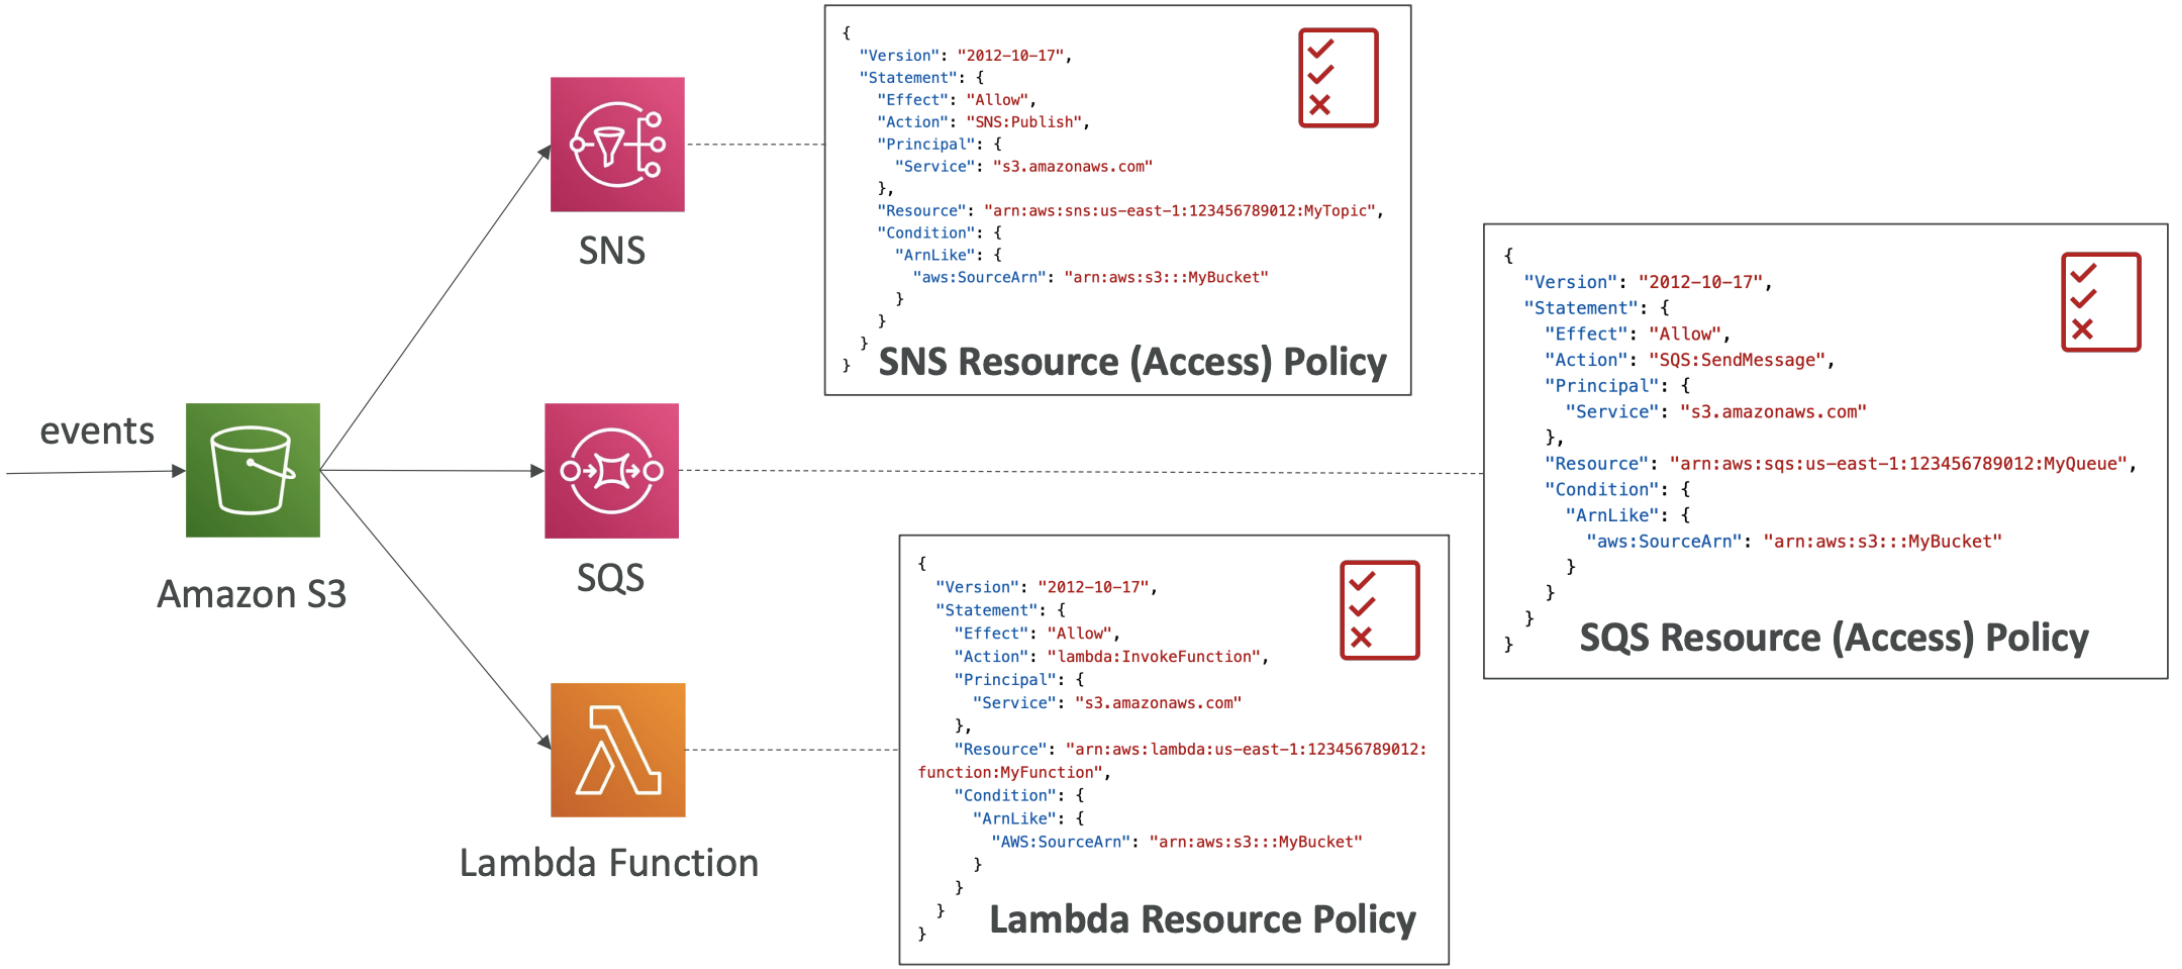

© Copyright by Cloud Mentor Pro | Confidential

S3 Event Notifications with Amazon EventBridge

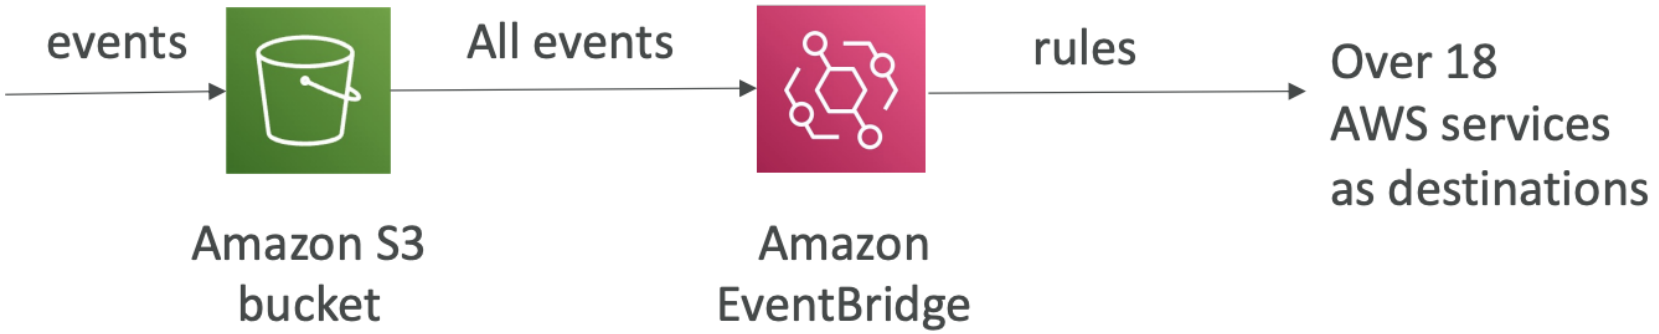

- **Advanced filtering** options with JSON rules (metadata, object size, name...)
- **Multiple Destinations** – ex Step Functions, Kinesis Streams / Firehose…
- **EventBridge Capabilities** – Archive, Replay Events, Reliable delivery

© Copyright by Cloud Mentor Pro | Confidential

S3 – Baseline Performance

- Amazon S3 automatically scales to high request rates, latency 100-200 ms
- Your application can achieve at least **3,500 PUT/COPY/POST/DELETE or 5,500 GET/HEAD requests per second per prefix in a bucket**.
- There are no limits to the number of prefixes in a bucket.
- Example (object path => prefix):
  - bucket/folder1/sub1/file => /folder1/sub1/
  - bucket/folder1/sub2/file => /folder1/sub2/
  - bucket/1/file => /1/
  - bucket/2/file . => /2/
- If you spread reads across all four prefixes evenly, you can achieve 22,000 requests per second for GET and HEAD

© Copyright by Cloud Mentor Pro | Confidential

S3 Performance

- Multi-Part upload:
  - recommended for files > 100MB, must use for files > 5GB
  - Can help parallelize uploads (speed up transfers)

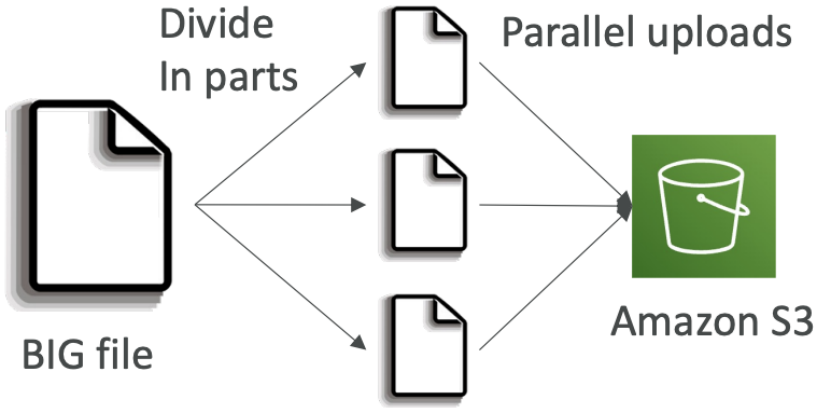

- S3 Transfer Acceleration
  - Increase transfer speed by transferring file to an AWS edge location which will forward the data to the S3 bucket in the target region
  - Compatible with multi-part upload

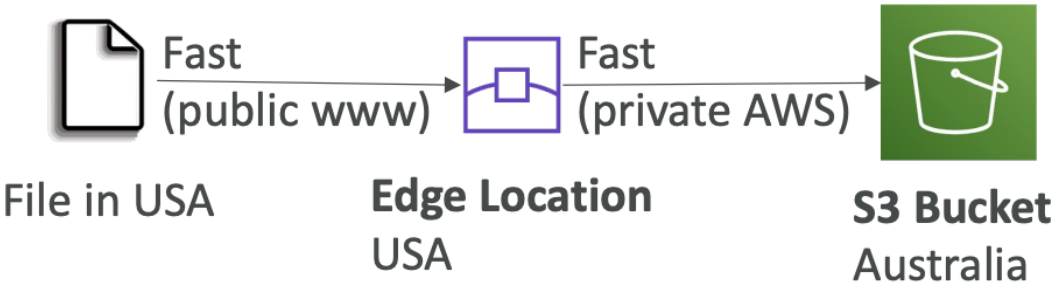

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Security

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Object Encryption

- You can encrypt objects in S3 buckets using one of 4 methods

- Server-Side Encryption (SSE)
  - Server-Side Encryption with Amazon S3-Managed Keys (SSE-S3) – Enabled by Default
    - Encrypts S3 objects using keys handled, managed, and owned by AWS
  - Server-Side Encryption with KMS Keys stored in AWS KMS (SSE-KMS)
    - Leverage AWS Key Management Service (AWS KMS) to manage encryption keys
  - Server-Side Encryption with Customer-Provided Keys (SSE-C)
    - When you want to manage your own encryption keys

- Client-Side Encryption

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Encryption – SSE -S3

- Encryption using keys handled, managed, and owned by AWS
- Object is encrypted server-side
- Encryption type is **AES-256**
- Must set header **"x-amz-server-side-encryption": "AES256"**
- **Enabled by default for new buckets & new objects**

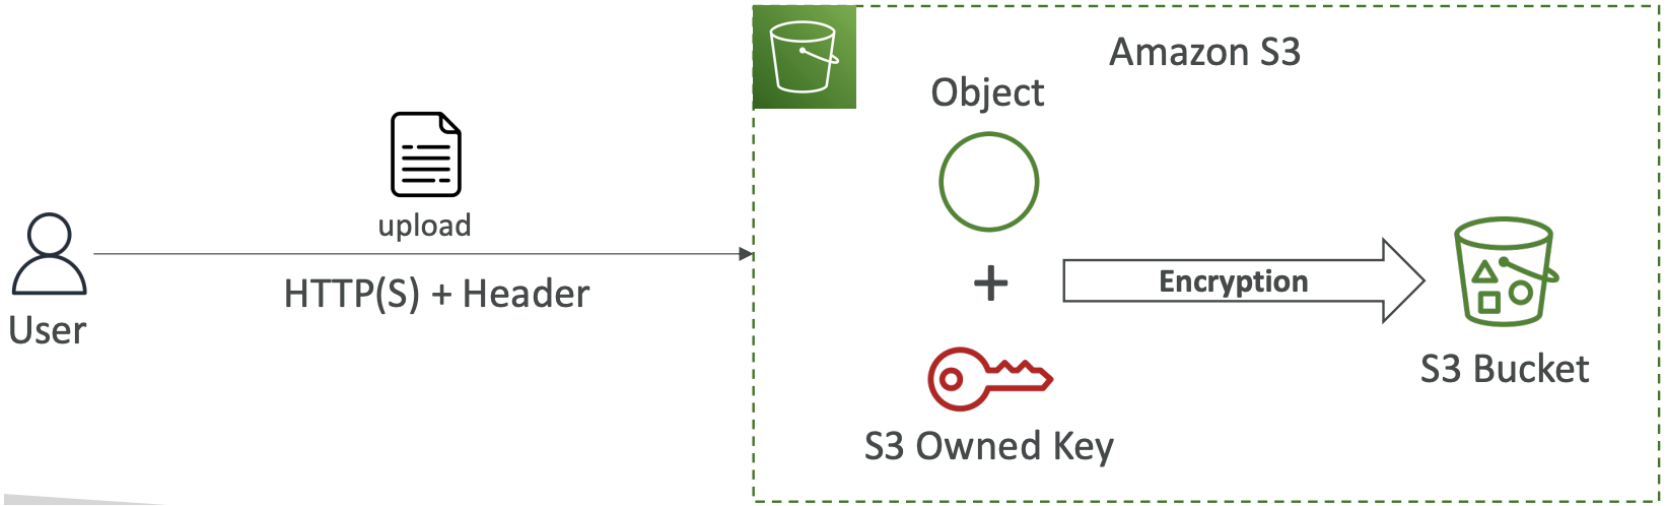

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Encryption – SSE-KMS

- Encryption using keys handled and managed by AWS KMS (Key Management Service)
- KMS advantages: user control + audit key usage using CloudTrail
- Object is encrypted server side
- Must set header **"x-amz-server-side-encryption": "aws:kms"**

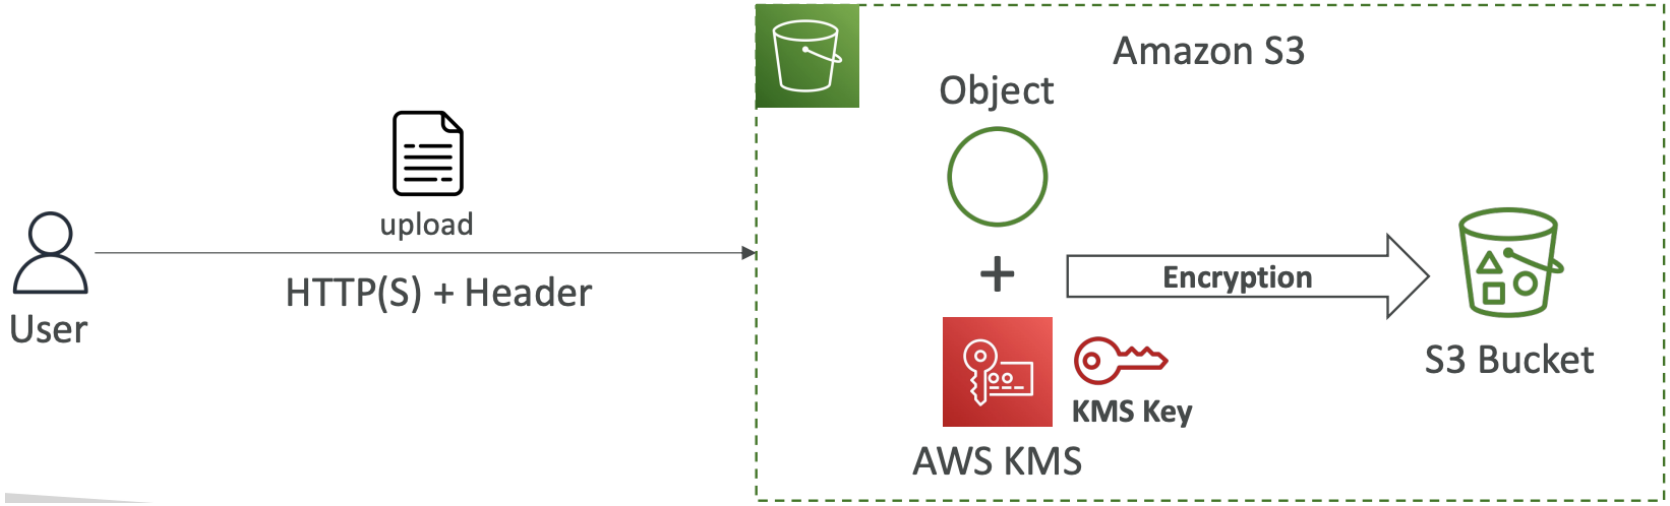

© Copyright by Cloud Mentor Pro | Confidential

SSE-KMS Limitation

- If you use SSE-KMS, you may be impacted by KMS
- When you upload, it calls GenerateDataKey KMS API
- When you download, it calls Decrypt KMS API
- Count towards the KMS quota per second (5500, 10000, 30000 req/s based on region)
- You can request a quota increase using the Service Quotas Console

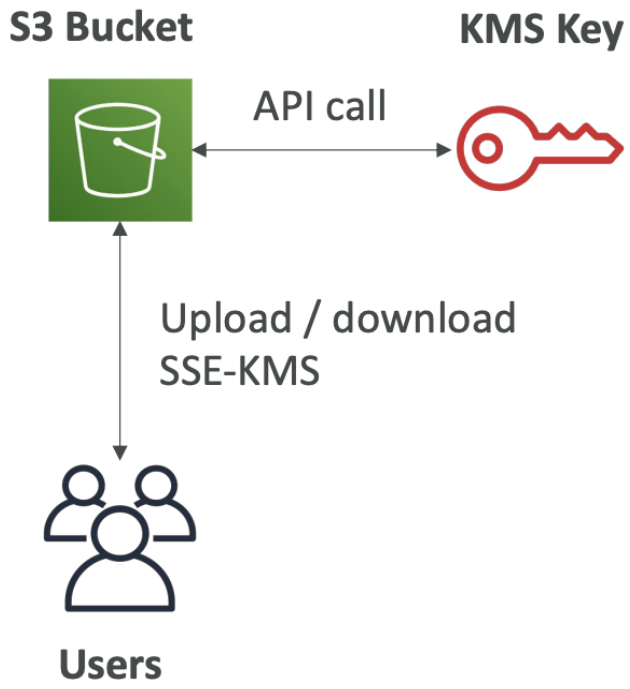

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Encryption – SSE-C

- Server-Side Encryption using keys fully managed by the customer outside of AWS
- Amazon S3 does **NOT** store the encryption key you provide
- **HTTPS must be used**
- Encryption key must provided in HTTP headers, for every HTTP request made

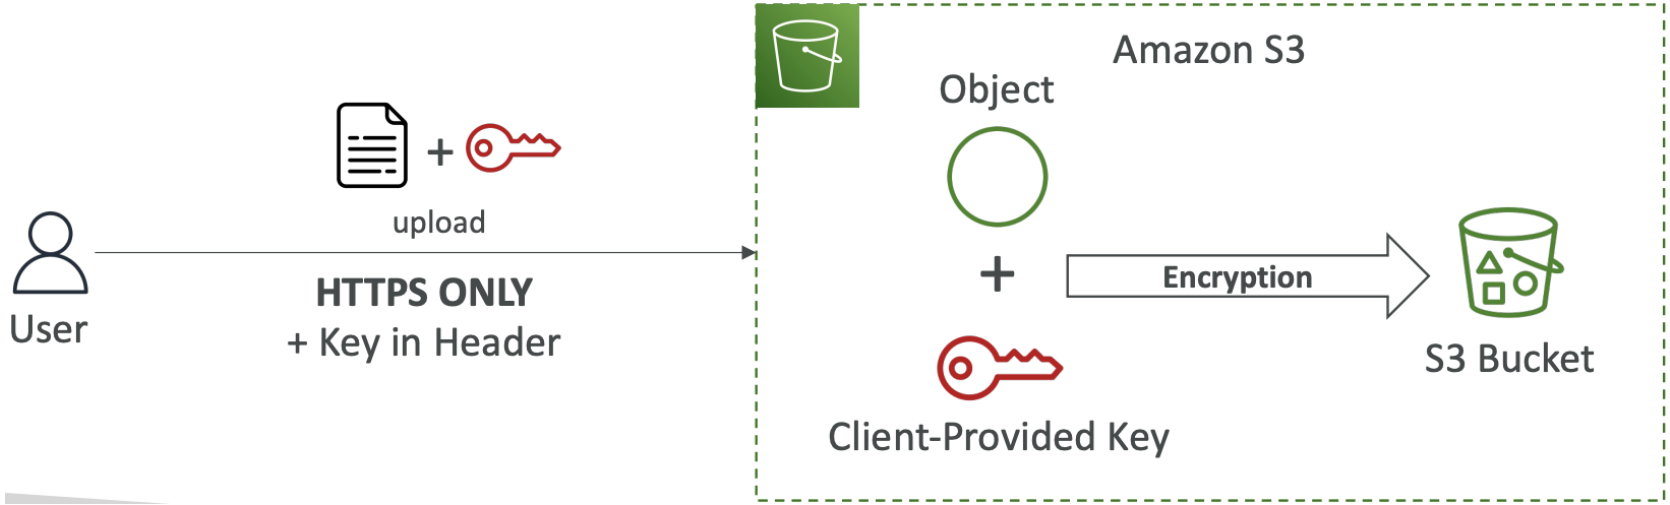

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 Encryption – Client-Side Encryption

- Use client libraries such as **Amazon S3 Client-Side Encryption Library**
- Clients must encrypt data themselves before sending to Amazon S3
- Clients must decrypt data themselves when retrieving from Amazon S3
- Customer fully manages the keys and encryption cycle

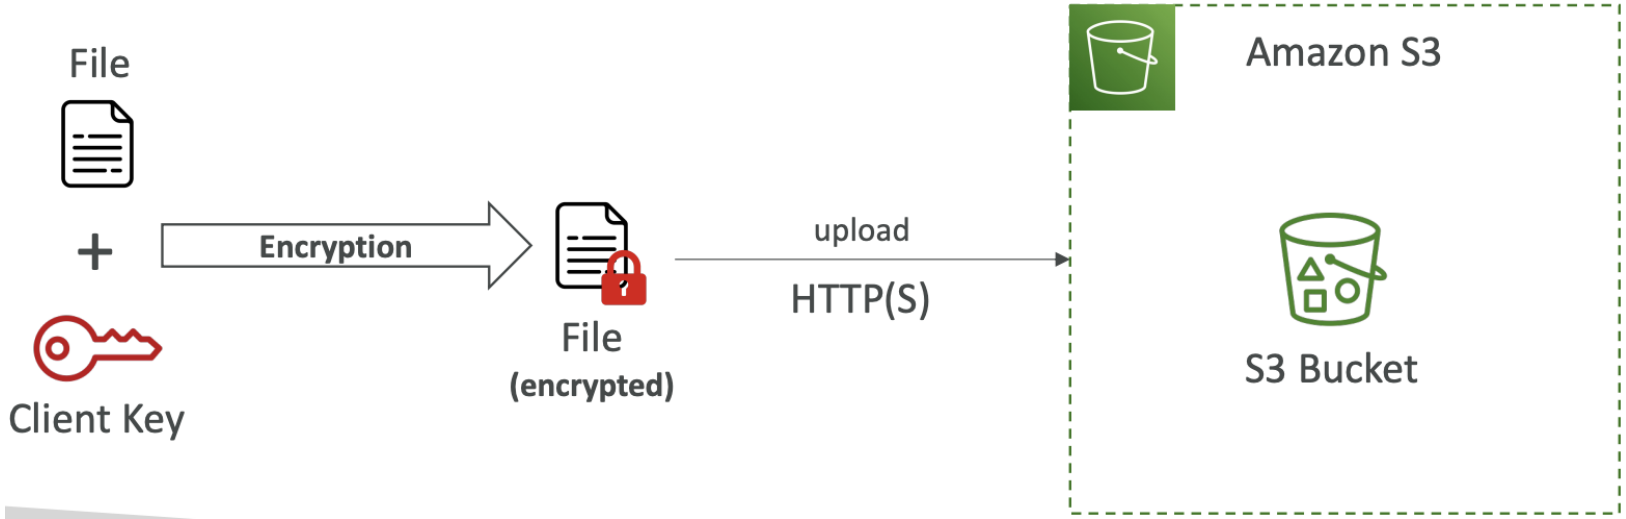

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Encryption in transit (SSL/TLS)

- Encryption in flight is also called SSL/TLS

- Amazon S3 exposes two endpoints:
  - **HTTP Endpoint** – non encrypted
  - **HTTPS Endpoint** – encryption in flight

- **HTTPS is recommended**
- **HTTPS is mandatory for SSE-C**
- Most clients would use the HTTPS endpoint by default

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Force Encryption in Transit aws:SecureTransport

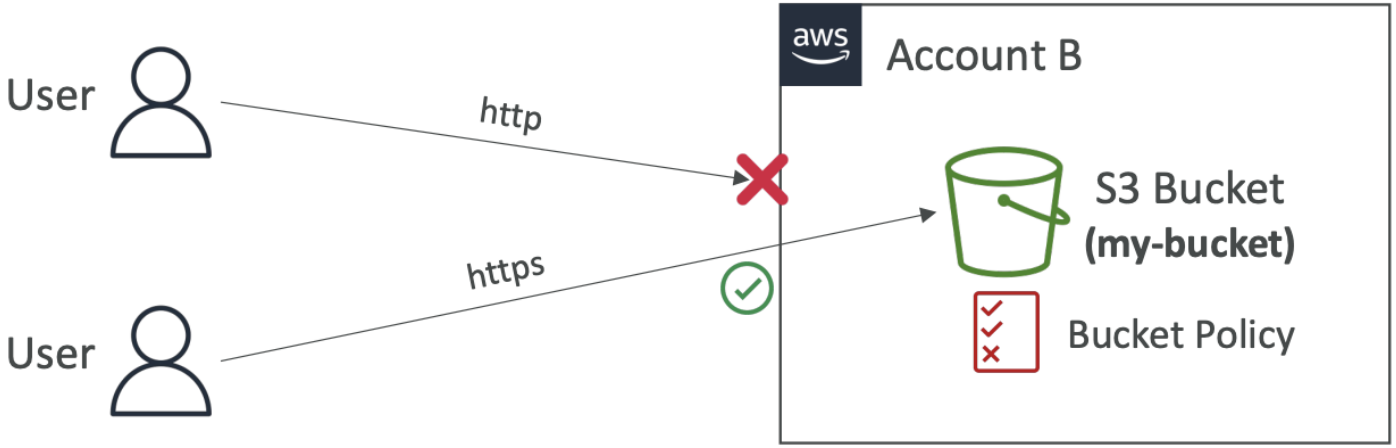


```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Deny",
            "Principal": "*",
            "Action": "s3:GetObject",
            "Resource": "arn:aws:s3:::my-bucket/*",
            "Condition": {
                "Bool": {
                    "aws:SecureTransport": "false"
                }
            }
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Default Encryption vs. Bucket Policies

- SSE-S3 encryption is automatically applied to new objects stored in S3 bucket
- Optionally, you can "force encryption" using a bucket policy and refuse any API call to PUT an S3 object without encryption headers (SSE-KMS or SSE-C)

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Deny",
            "Action": "s3:PutObject",
            "Principal": "*",
            "Resource": "arn:aws:s3:::my-bucket/*",
            "Condition": {
                "StringNotEquals": {
                    "s3:x-amz-server-side-encryption": "aws:kms"
                }
            }
        }
    ]
}
```

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Deny",
            "Action": "s3:PutObject",
            "Principal": "*",
            "Resource": "arn:aws:s3:::my-bucket/*",
            "Condition": {
                "Null": {
                    "s3:x-amz-server-side-encryption-customer-algorithm": "true"
                }
            }
        }
    ]
}
```

- Note: Bucket Policies are evaluated before "Default Encryption"

© Copyright by Cloud Mentor Pro | Confidential

What is CORS?

- Cross-Origin Resource Sharing (CORS)
- **Origin = scheme (protocol) + host (domain) + port**
  - example: https://www.example.com (implied port is 443 for HTTPS, 80 for HTTP)
- **Web Browser** based mechanism to allow requests to other origins while visiting the main origin
- Same origin: http://example.com/app1 & http://example.com/app2
- Different origins: http://www.example.com & http://other.example.com
- The requests won't be fulfilled unless the other origin allows for the requests, using **CORS Headers** (example: **Access-Control-Allow-Origin**)

© Copyright by Cloud Mentor Pro | Confidential

What is CORS?

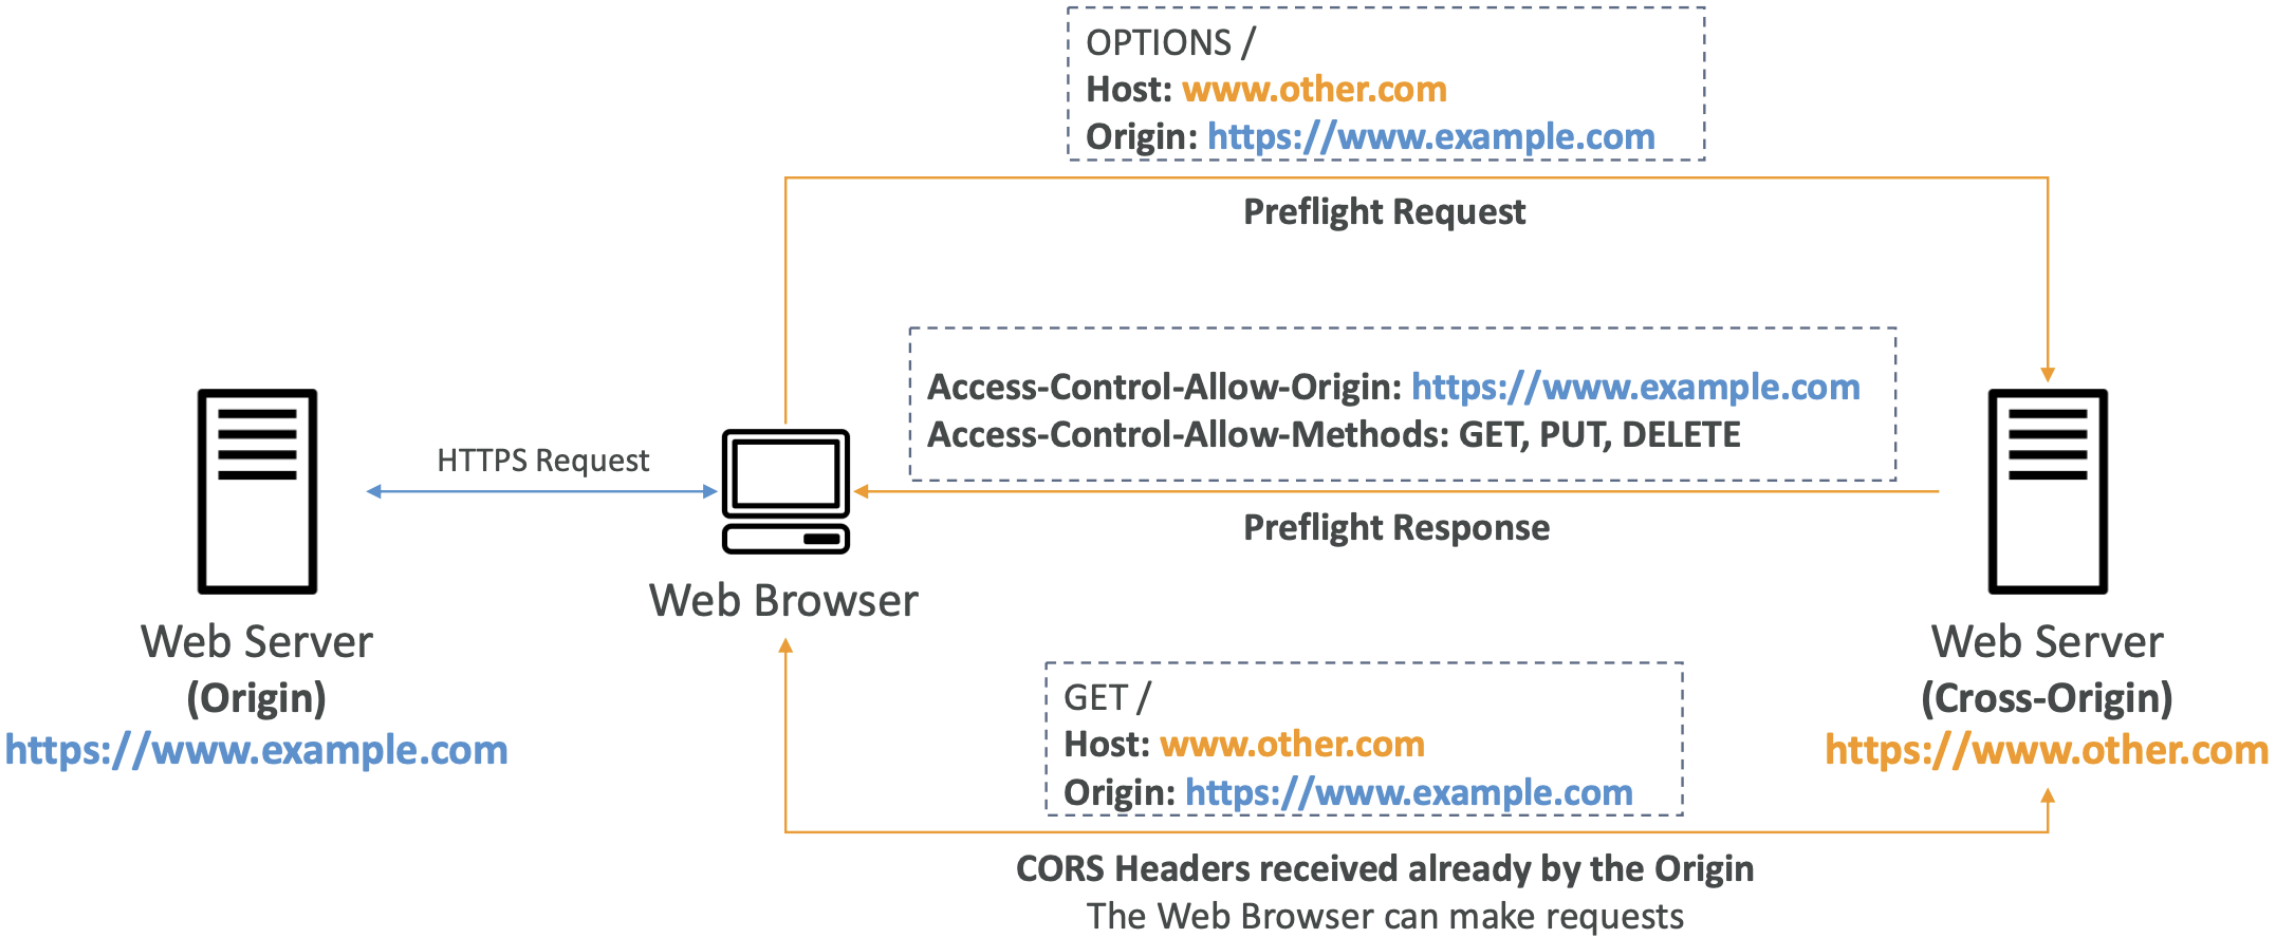

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – CORS

- If a client makes a cross-origin request on our S3 bucket, we need to enable the correct CORS headers
- It's a popular exam question
- You can allow for a specific origin or for * (all origins)

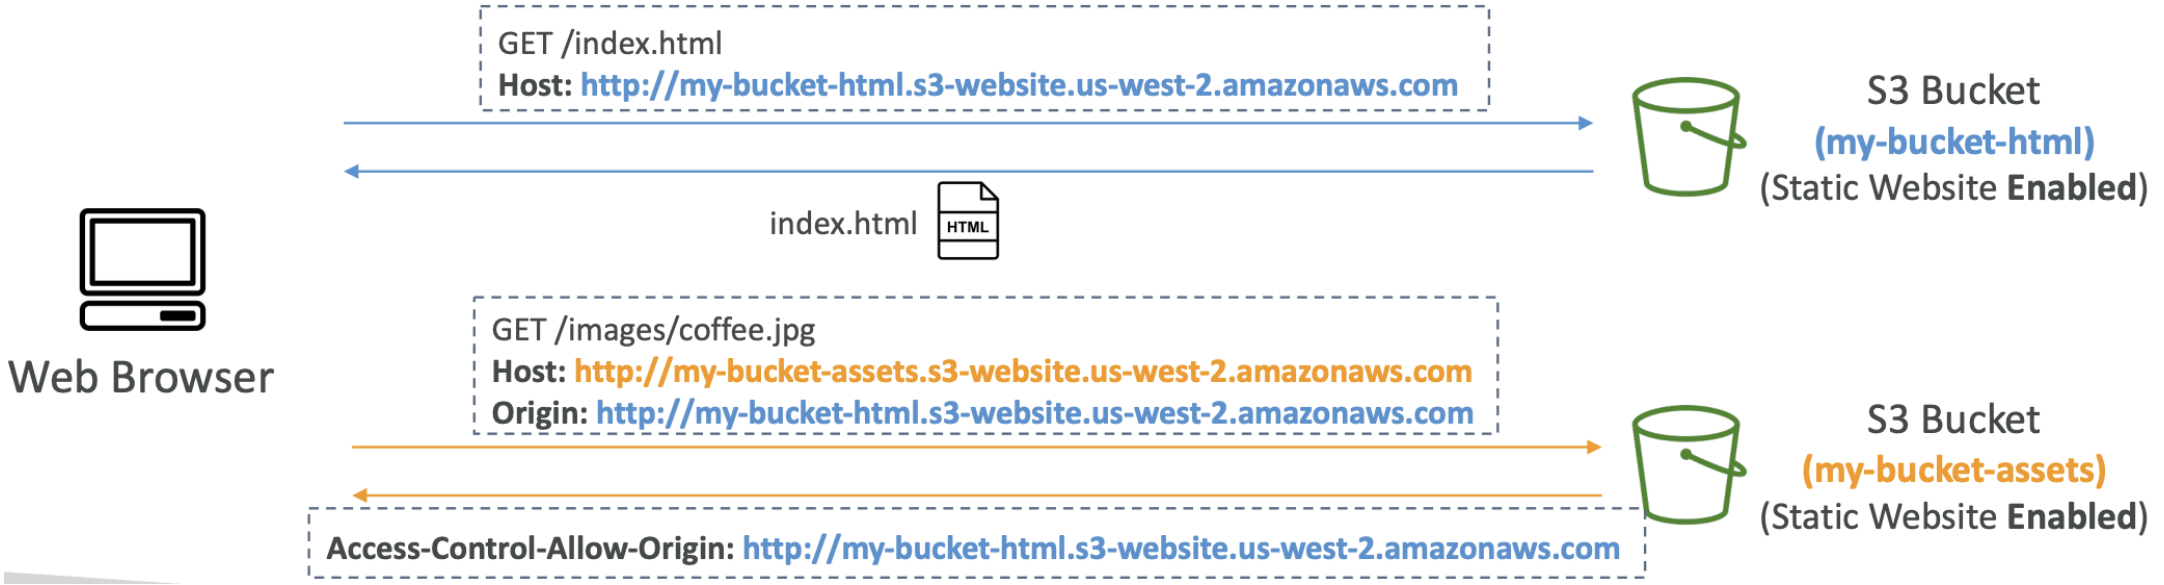

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – MFA Delete

- **MFA (Multi-Factor Authentication)** – force users to generate a code on a device (usually a mobile phone or hardware) before doing important operations on S3
- MFA will be required to:
  - Permanently delete an object version
  - Suspend Versioning on the bucket
- MFA won't be required to:
  - Enable Versioning
  - List deleted versions
- To use MFA Delete, **Versioning must be enabled** on the bucket
- Only the bucket owner (root account) can enable/disable MFA Delete

© Copyright by Cloud Mentor Pro | Confidential

S3 Access Logs

- For audit purpose, you may want to log all access to S3 buckets
- Any request made to S3, from any account, authorized or denied, will be logged into another S3 bucket
- That data can be analyzed using data analysis tools…
- The target logging bucket must be in the same AWS region

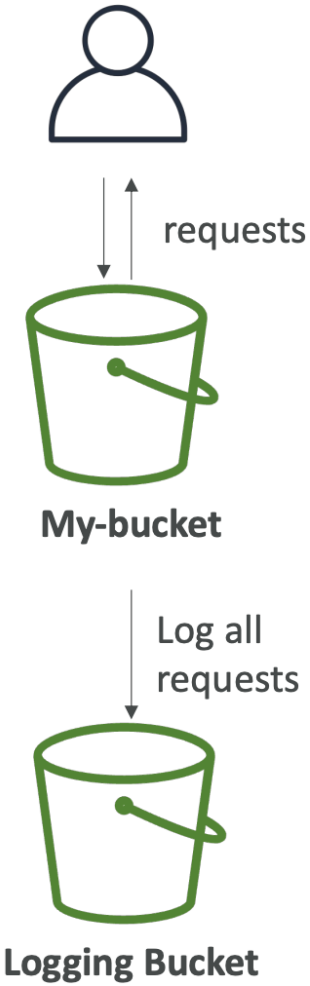

- The log format is at: https://docs.aws.amazon.com/AmazonS3/latest/dev/LogFormat.html

© Copyright by Cloud Mentor Pro | Confidential

S3 Access Logs: Warning

- Do not set your logging bucket to be the monitored bucket
- It will create a logging loop, and **your bucket will grow exponentially**

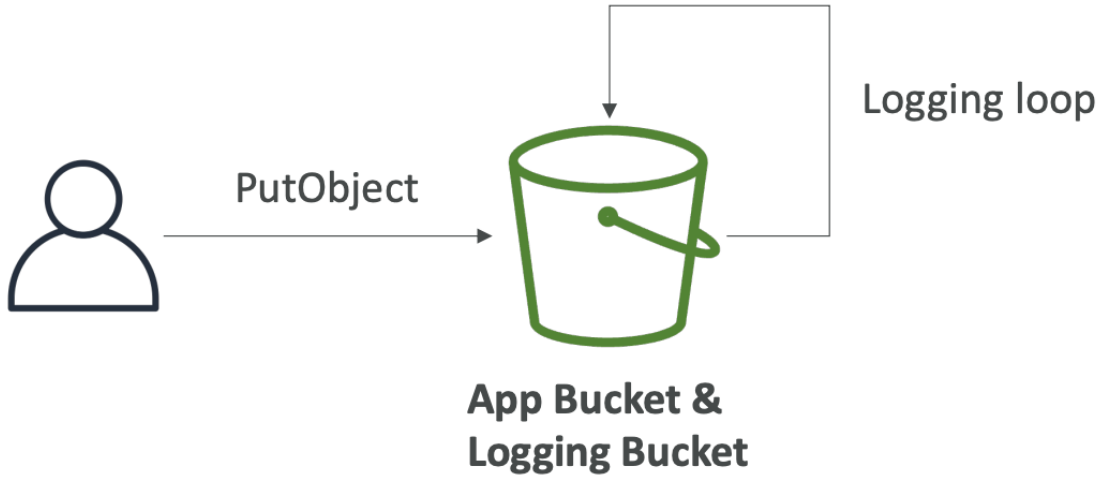

© Copyright by Cloud Mentor Pro | Confidential

Amazon S3 – Pre-Signed URLs

- Generate pre-signed URLs using the **S3 Console, AWS CLI or SDK**
- **URL Expiration**
  - **S3 Console** – 1 min up to 720 mins (12 hours)
  - **AWS CLI** – configure expiration with --expires-in parameter in seconds (default 3600 secs, max. 604800 secs ~ 168 hours)
- Users given a pre-signed URL inherit the permissions of the user that generated the URL for GET / PUT
- Examples:
  - Allow only logged-in users to download a premium video from your S3 bucket
  - Allow an ever-changing list of users to download files by generating URLs dynamically
  - Allow temporarily a user to upload a file to a precise location in your S3 bucket

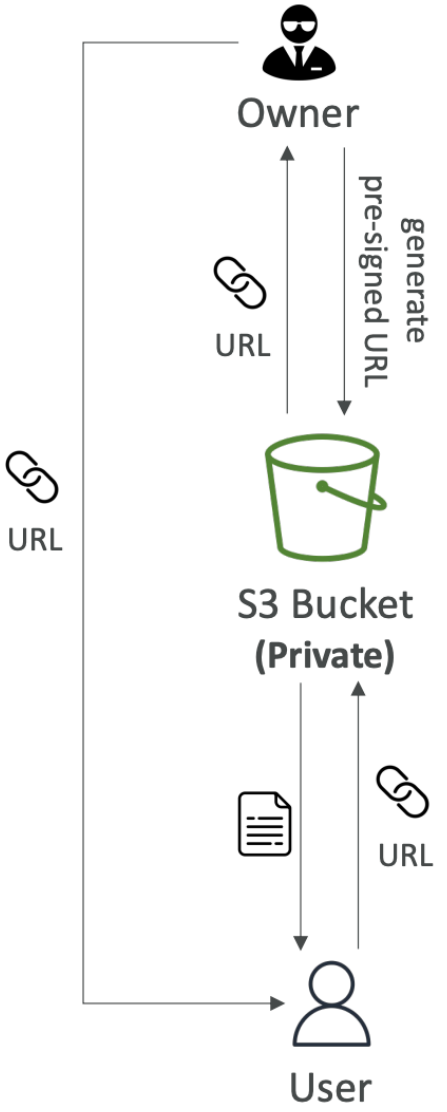

© Copyright by Cloud Mentor Pro | Confidential

S3 Object Lock (versioning must be enabled)

- Adopt a WORM (Write Once Read Many) model
- Block an object version deletion for a specified amount of time
- **Retention mode - Compliance**:
  - Object versions can't be overwritten or deleted by any user, including the root user
  - Objects retention modes can't be changed, and retention periods can't be shortened
- **Retention mode - Governance**:
  - Most users can't overwrite or delete an object version or alter its lock settings
  - Some users have special permissions to change the retention or delete the object
- **Retention Period**: protect the object for a fixed period, it can be extended
- **Legal Hold**:
  - protect the object indefinitely, independent from retention period
  - can be freely placed and removed using the *s3:PutObjectLegalHold* IAM permission

© Copyright by Cloud Mentor Pro | Confidential

S3 – Access Points

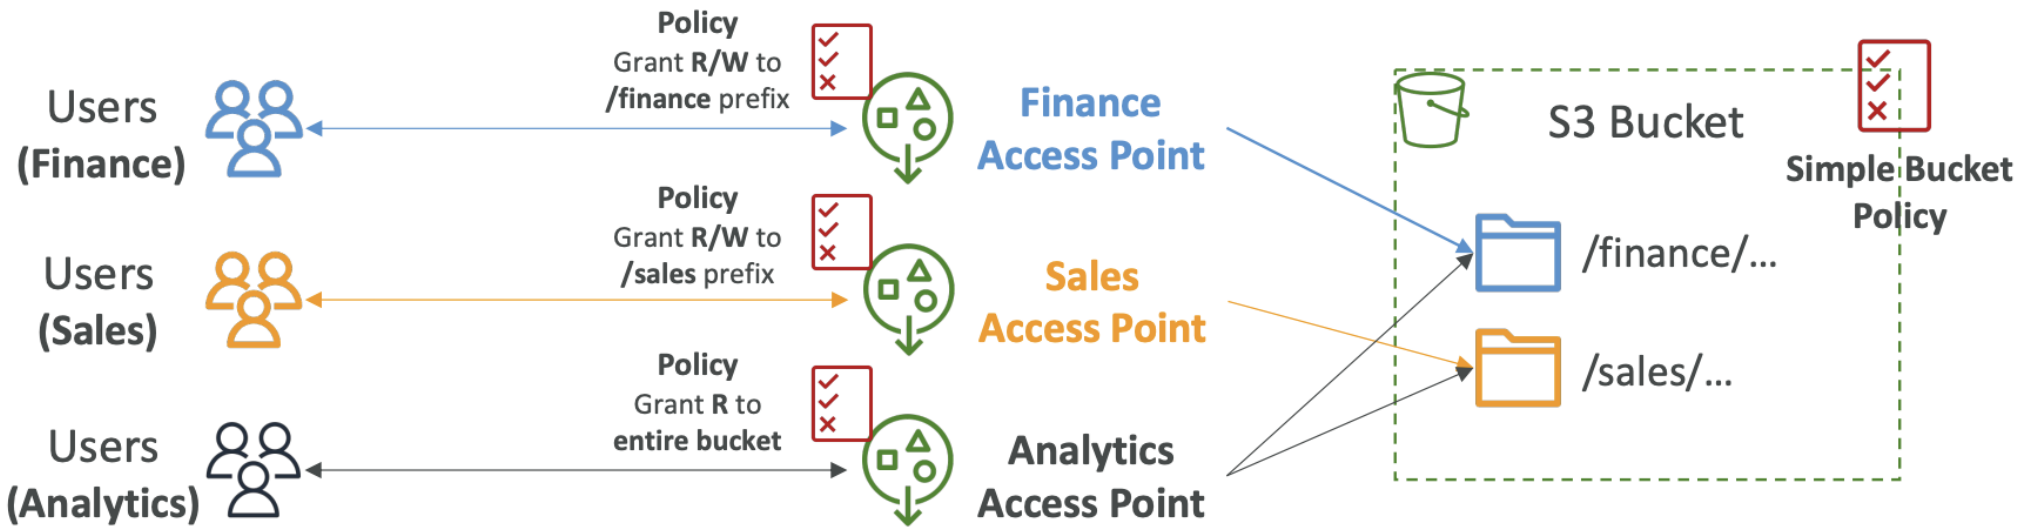

- Access Points simplify security management for S3 Buckets
- Each Access Point has:
  - its own DNS name (Internet Origin or VPC Origin)
  - an access point policy (similar to bucket policy) – manage security at scale

© Copyright by Cloud Mentor Pro | Confidential

S3 – Access Points – VPC Origin

- We can define the access point to be accessible only from within the VPC
- You must create a VPC Endpoint to access the Access Point (Gateway or Interface Endpoint)
- The VPC Endpoint Policy must allow access to the target bucket and Access Point

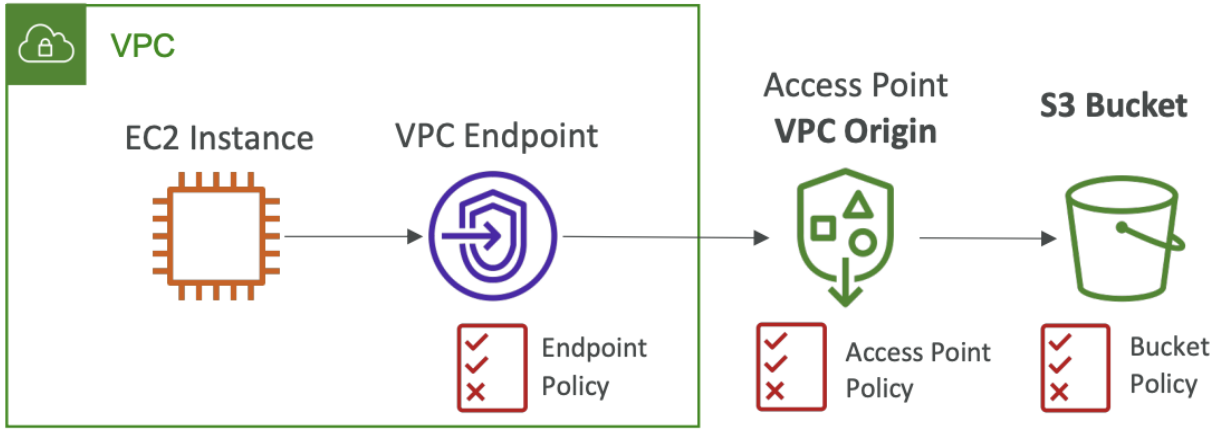


```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Principal": "*",
            "Action": [
                "s3:GetObject"
            ],
            "Effect": "Allow",
            "Resource": [
                "arn:aws:s3:::awsexamplebucket1/*",
                "arn:aws:s3:us-west-2:123456789012:accesspoint/example-vpc-ap/object/*"
            ]
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

S3 Object Lambda

- Use AWS Lambda Functions to change the object before it is retrieved by the caller application

- Use Cases:
  - Redacting personally identifiable information for analytics
  - Converting across data formats, such as converting XML to JSON.
  - Resizing and watermarking

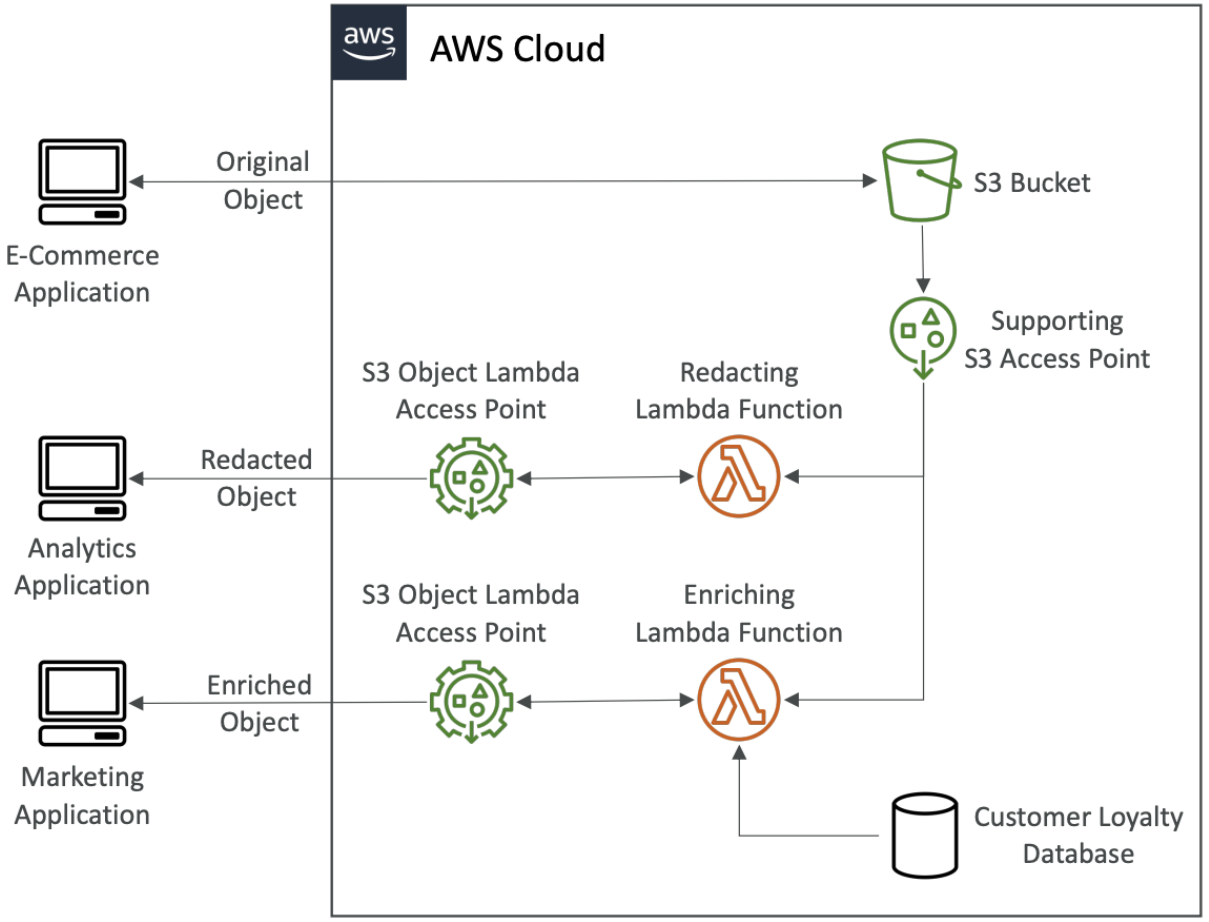

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- *How to Create a static website using Amazon S3*

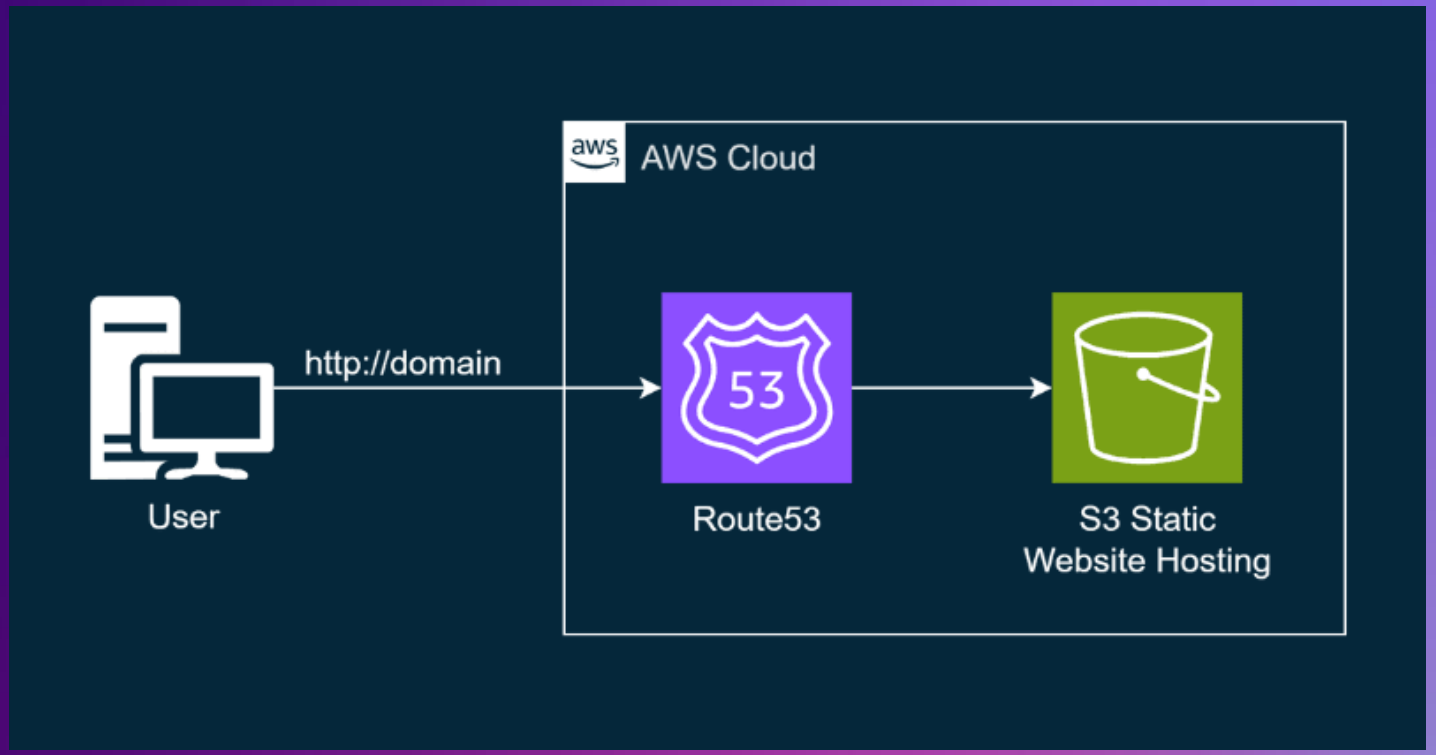

© Copyright by Cloud Mentor Pro | Confidential

Section 5

- *Networking - Amazon VPC*

© Copyright by Cloud Mentor Pro | Confidential

Amazon VPC

© Copyright by Cloud Mentor Pro | Confidential

VPC Components Diagram

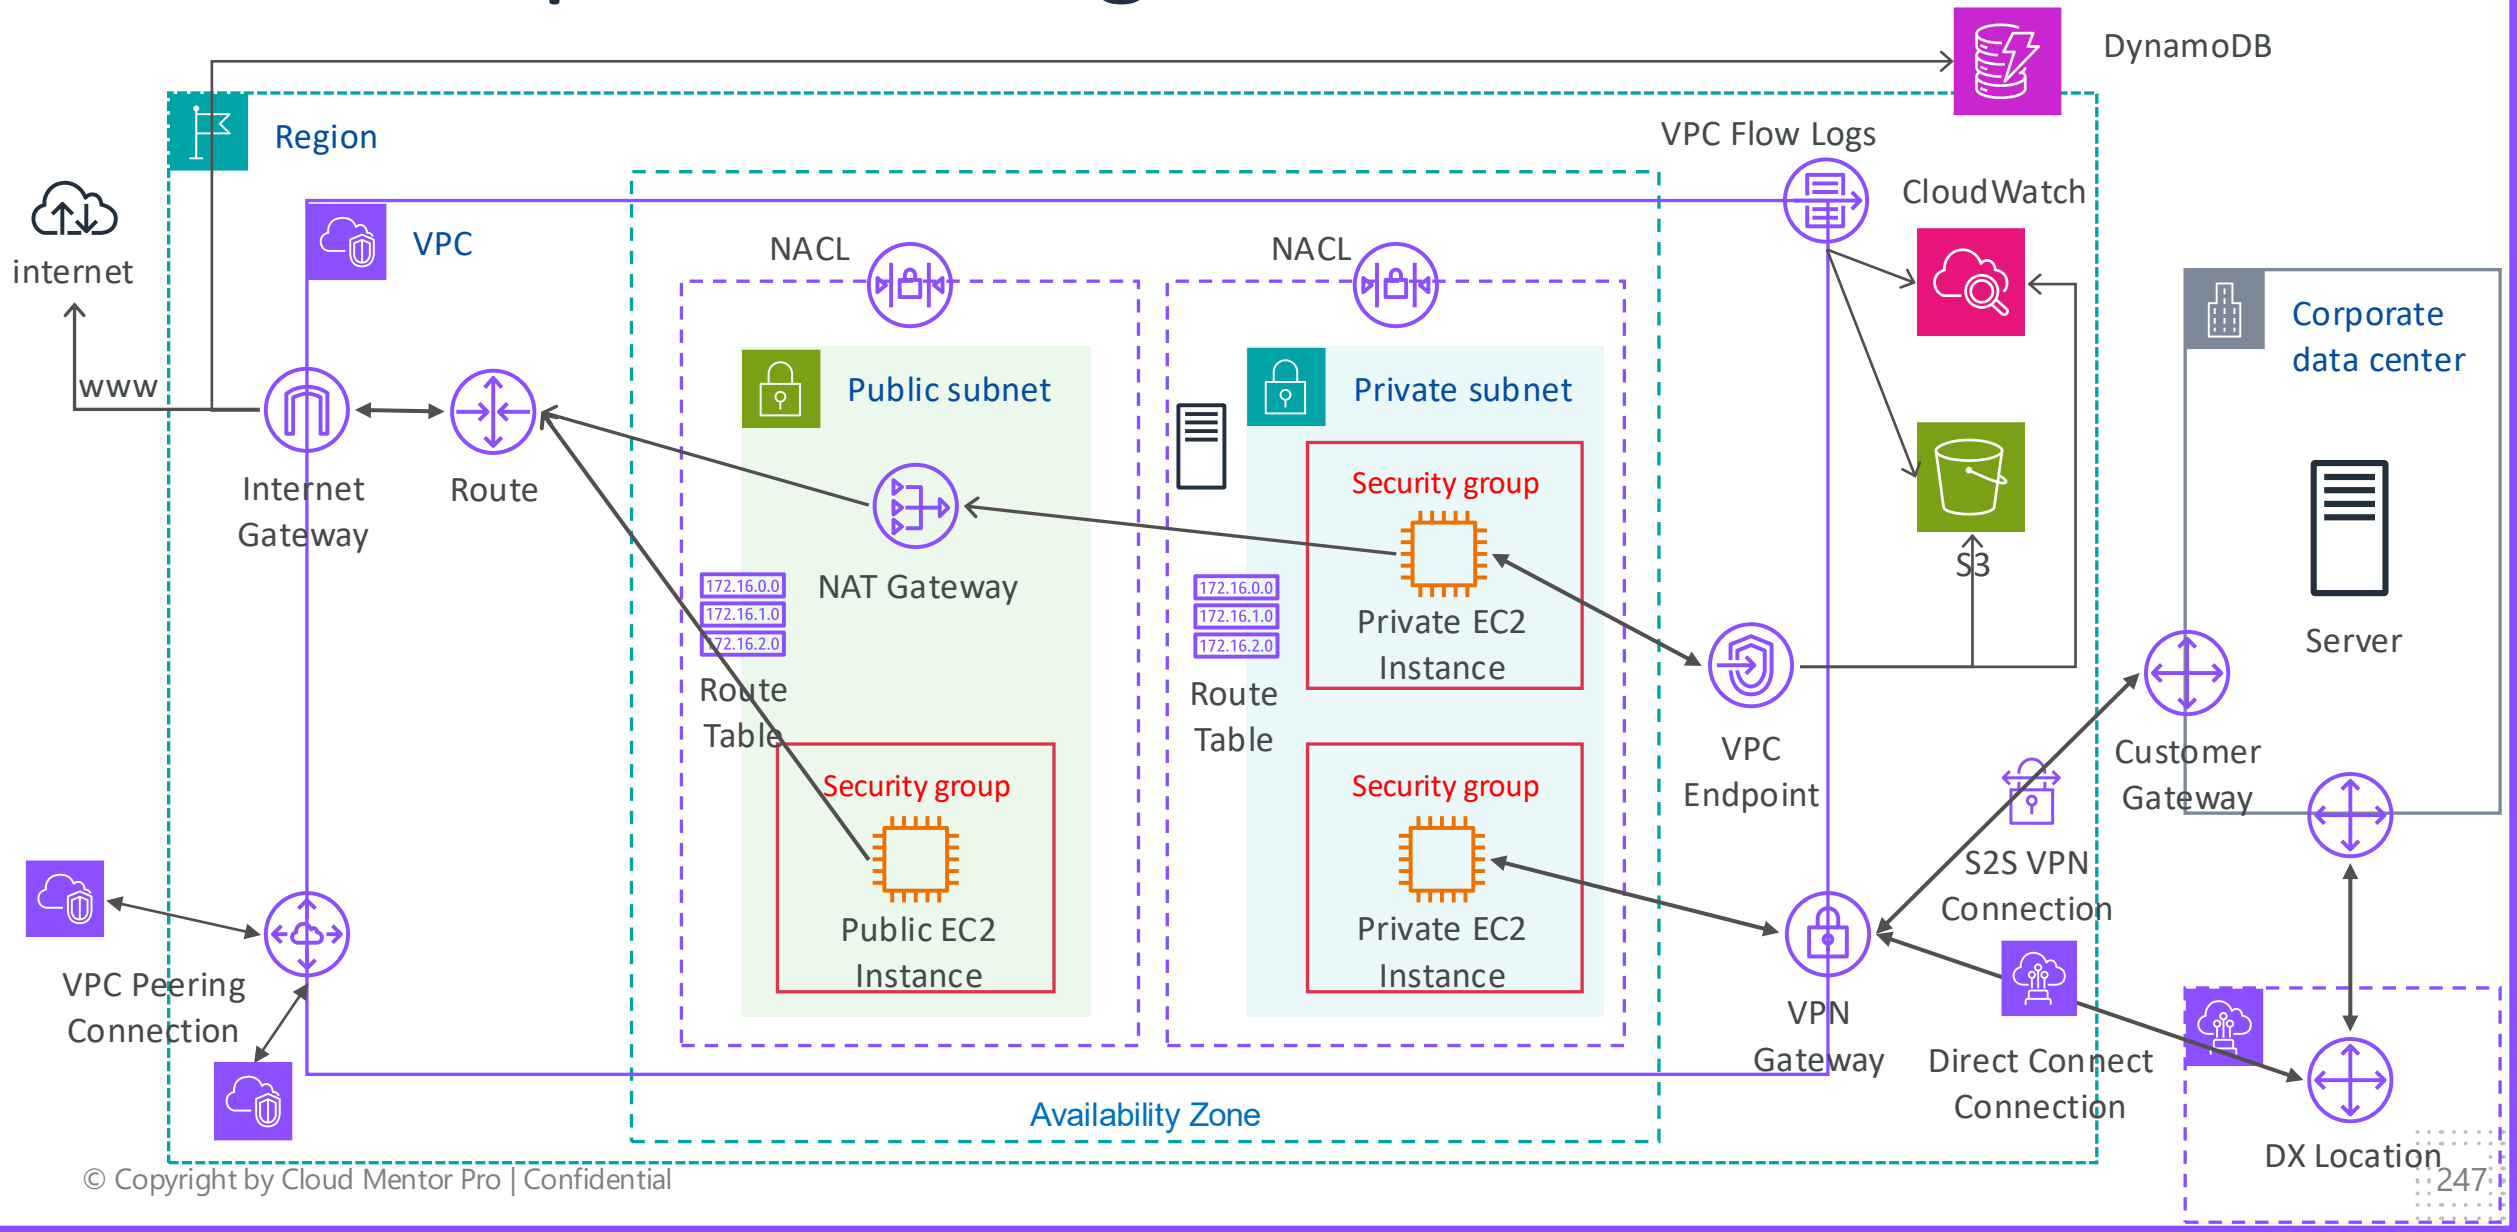

© Copyright by Cloud Mentor Pro | Confidential

Understanding CIDR – IPv4

Classless Inter-Domain Routing – IP Range

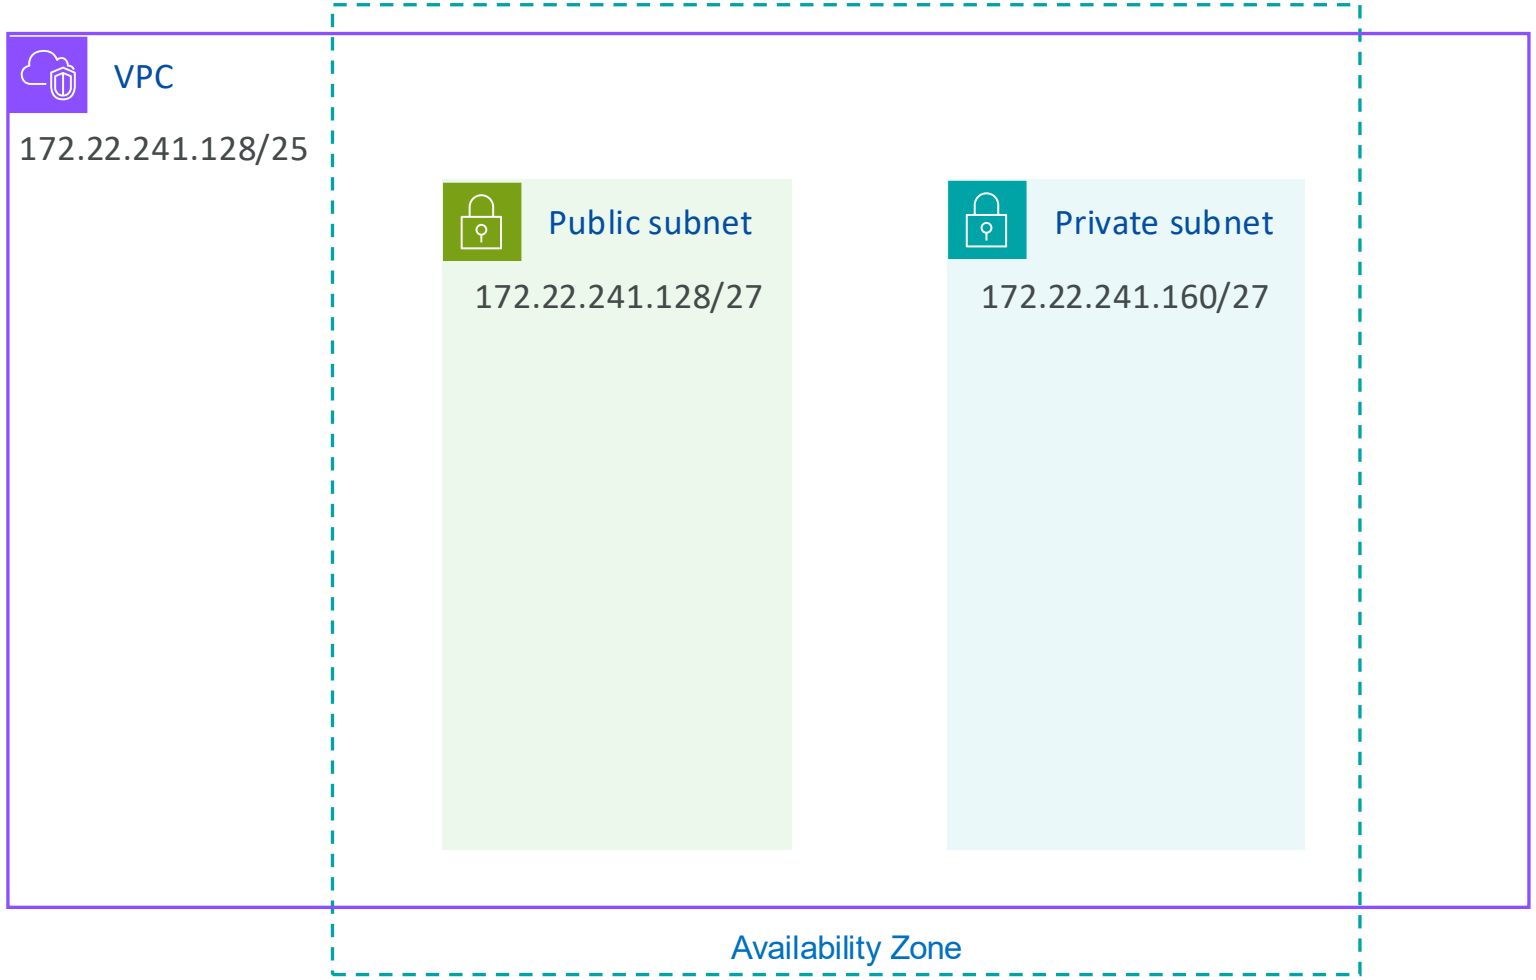

© Copyright by Cloud Mentor Pro | Confidential

Understanding CIDR – IPv4

CIDR: 172.22.241.128/25

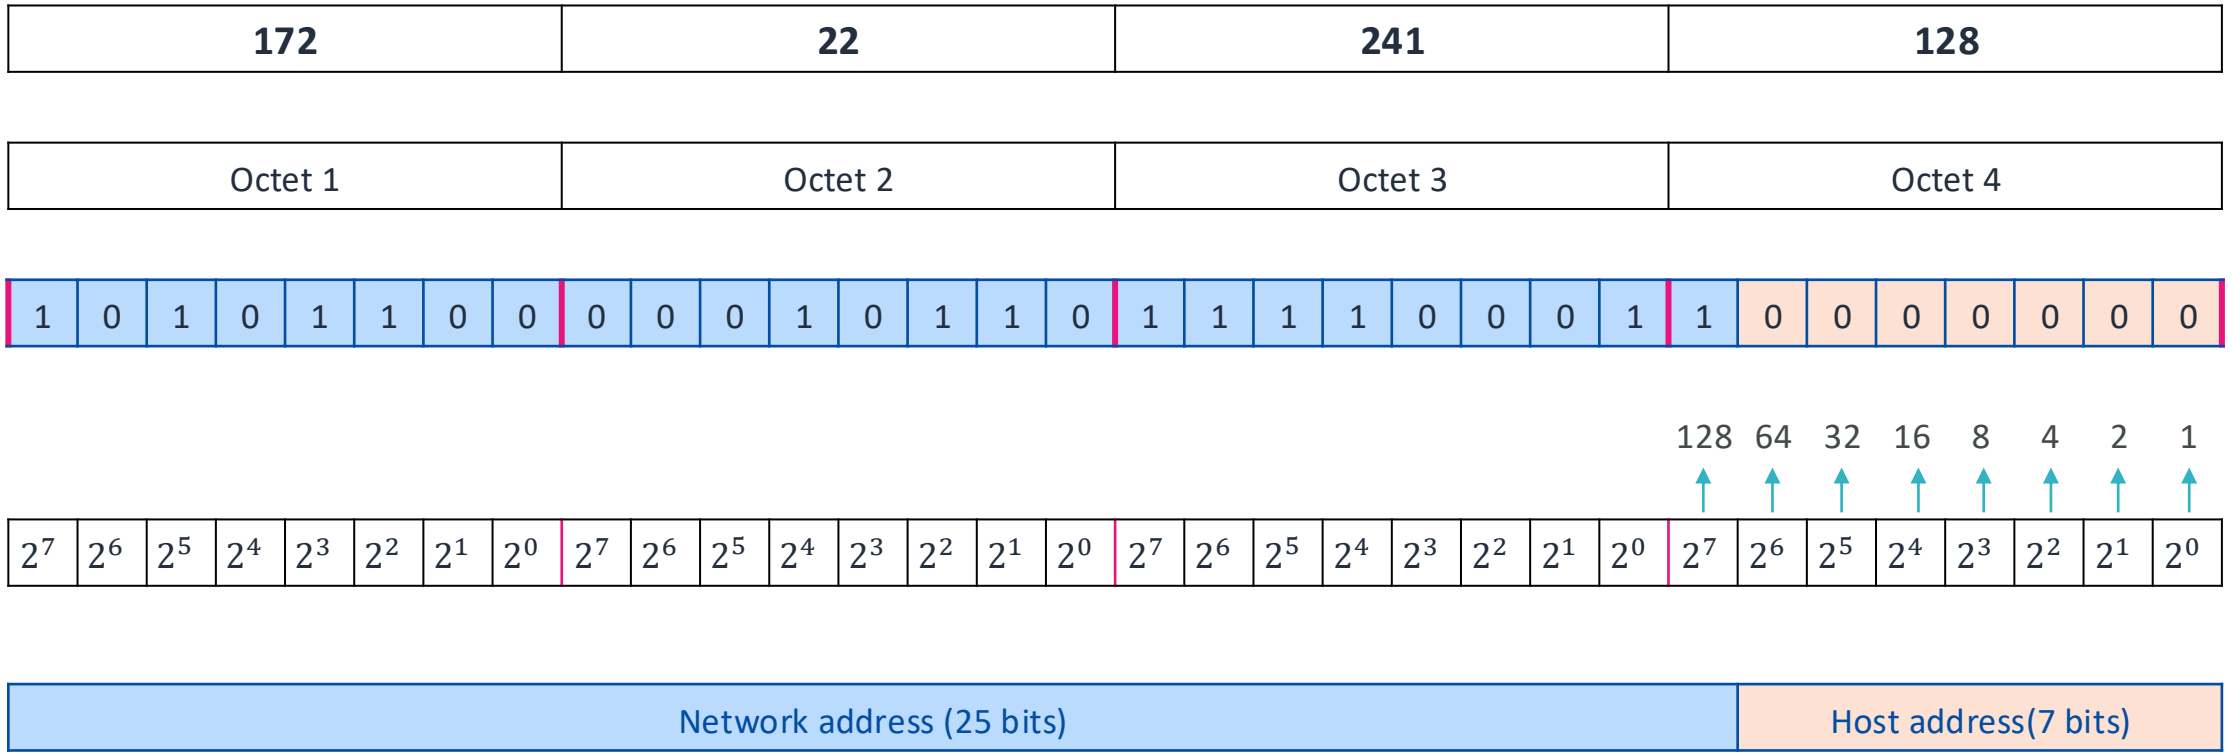

© Copyright by Cloud Mentor Pro | Confidential

Understanding CIDR – IPv4

Number of host address

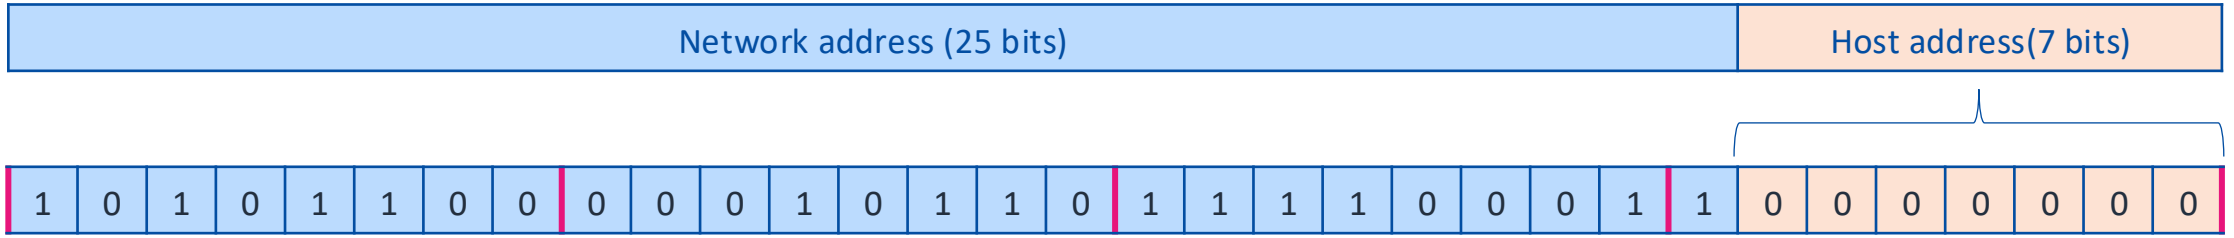

Host address: 172.22.241.128 ~ 172.22.241.255 (128 IPs)

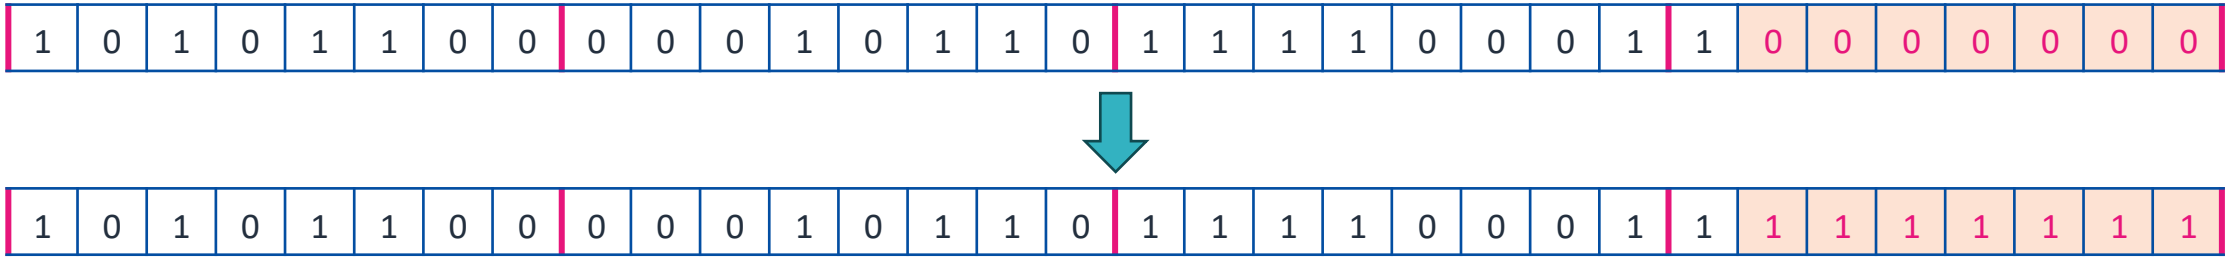

© Copyright by Cloud Mentor Pro | Confidential

Understanding CIDR – Subnet Mask

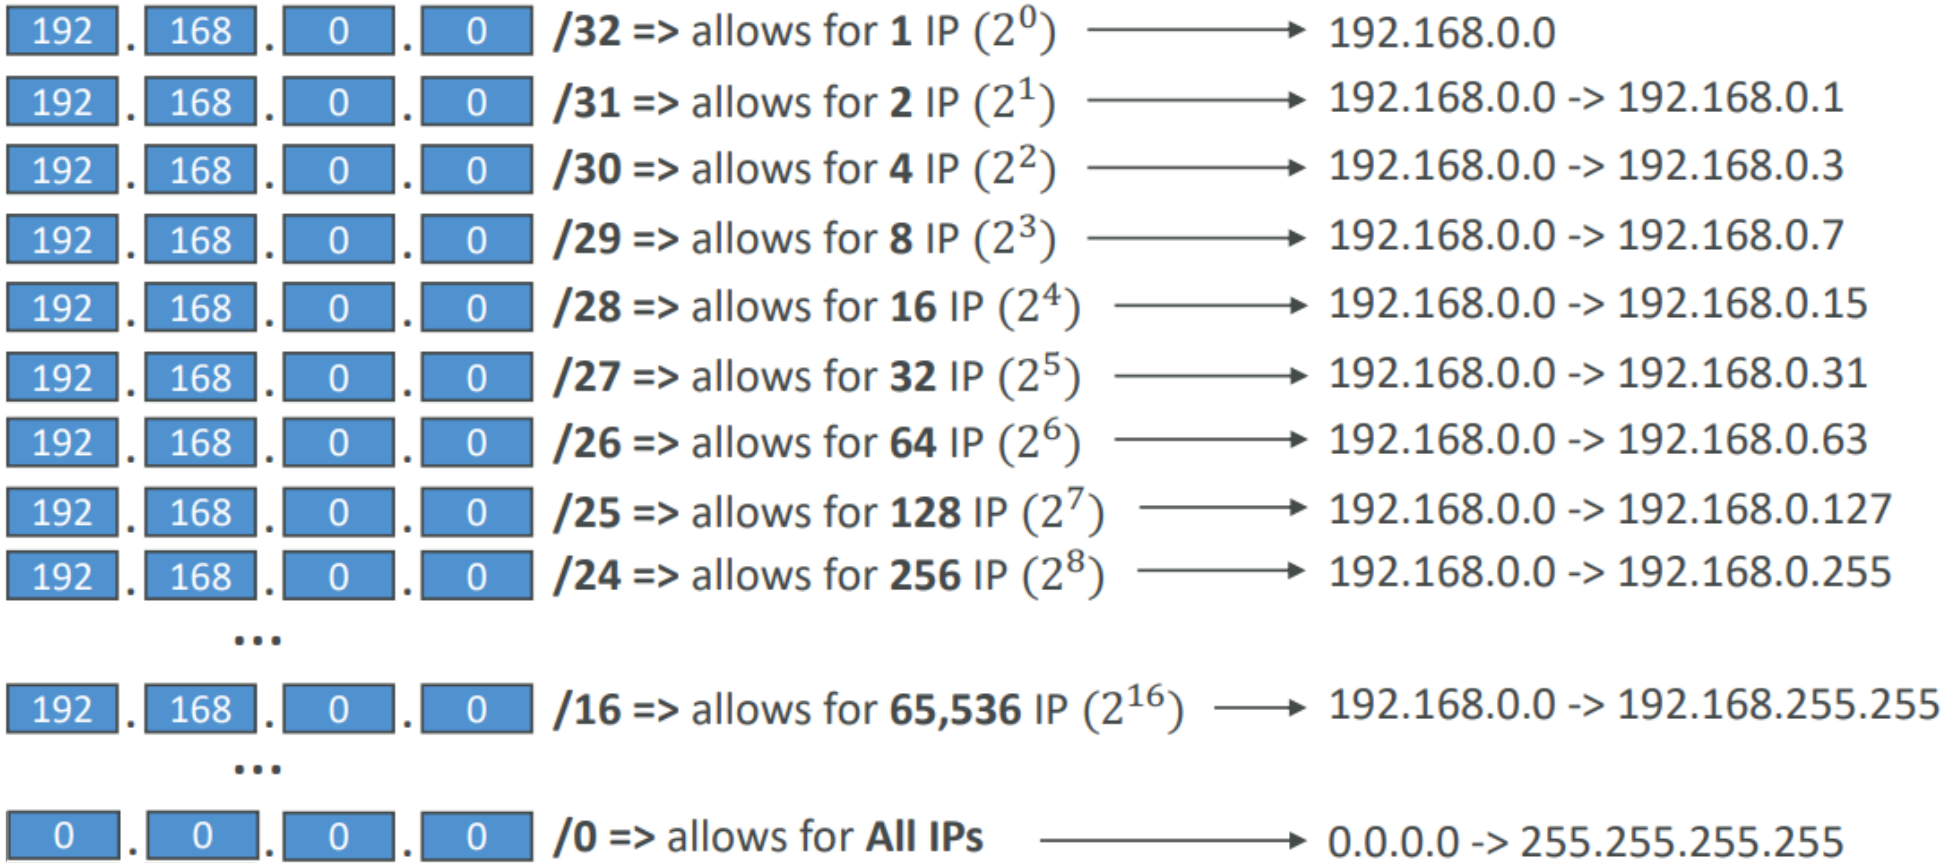

© Copyright by Cloud Mentor Pro | Confidential

Split 4 subnets CIDR: 172.22.241.128/25

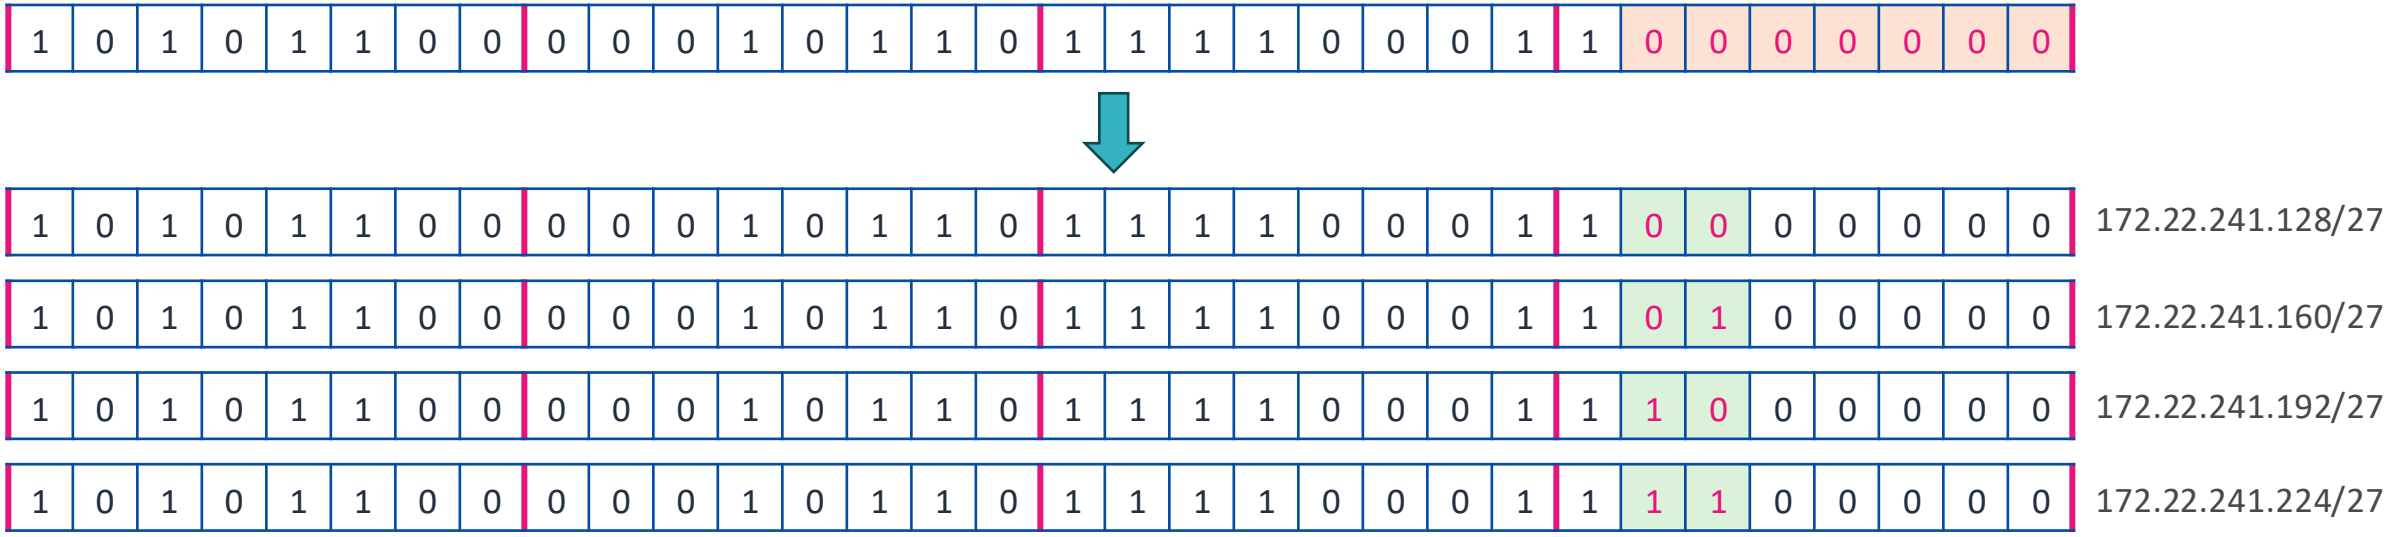

172.22.241.128/27 → 172.22.241.128 → 172.22.241.159

172.22.241.160/27 → 172.22.241.160 → 172.22.241.191

172.22.241.192/27 → 172.22.241.192 → 172.22.241.223

172.22.241.224/27 → 172.22.241.224 → 172.22.241.255

© Copyright by Cloud Mentor Pro | Confidential

https://visualsubnetcalc.com/

Default VPC Walkthrough

- All new AWS accounts have a default VPC
- New EC2 instances are launched into the default VPC if no subnet is specified
- Default VPC has Internet connectivity and all EC2 instances inside it have public IPv4 addresses

© Copyright by Cloud Mentor Pro | Confidential

VPC in AWS – IPv4

- VPC = Virtual Private Cloud
- You can have multiple VPCs in an AWS region (max. 5 per region – soft limit)
- Max. CIDR per VPC is 5, for each CIDR:
  - Min. size is /28 (16 IP addresses)
  - Max. size is /16 (65536 IP addresses)
- We recommend that you specify a CIDR block from the private IPv4 address ranges as specified in RFC 1918

RFC 1918 range	Example CIDR block
10.0.0.0 - 10.255.255.255 (10/8 prefix)	10.0.0.0/16
172.16.0.0 - 172.31.255.255 (172.16/12 prefix)	172.31.0.0/16
192.168.0.0 - 192.168.255.255 (192.168/16 prefix)	192.168.0.0/20

- Your VPC CIDR should NOT overlap with your other networks (e.g.,corporate)

© Copyright by Cloud Mentor Pro | Confidential

State of Hands-on

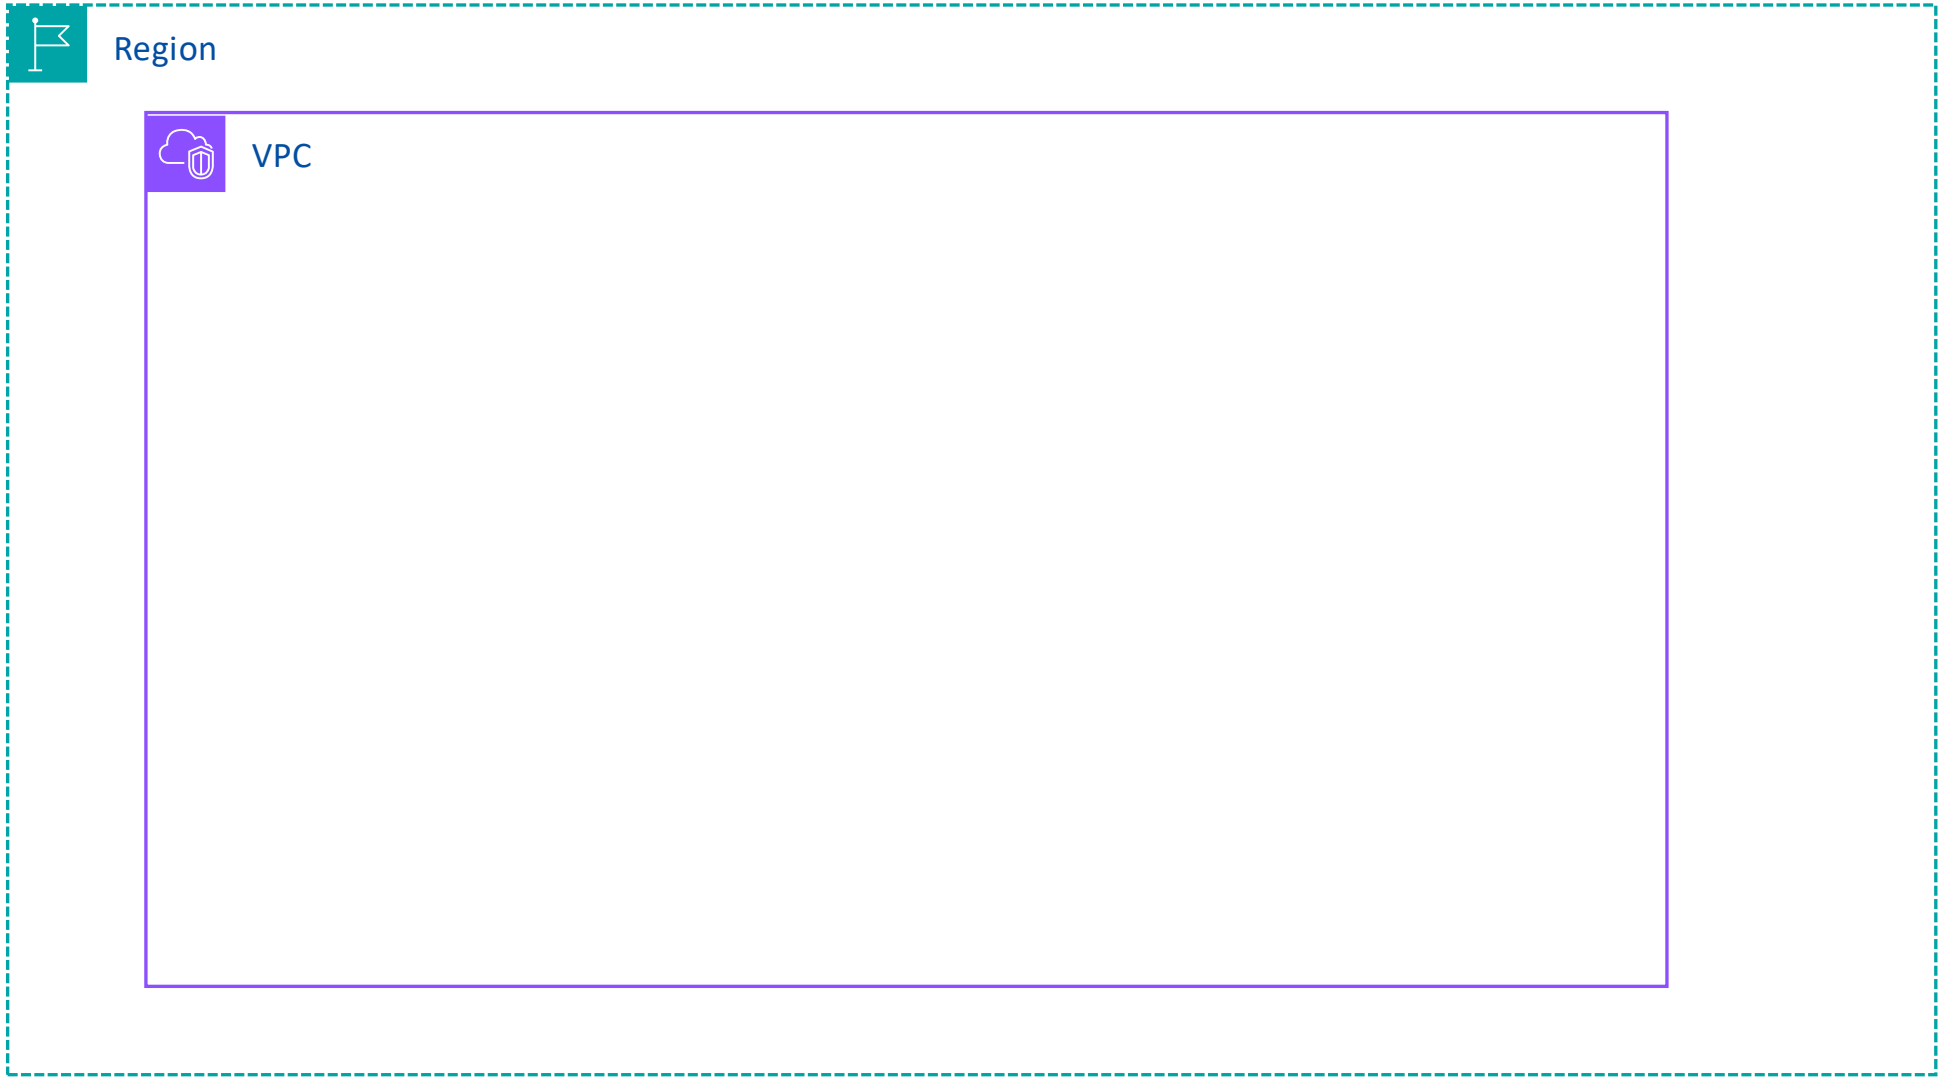

© Copyright by Cloud Mentor Pro | Confidential

Adding Subnets

© Copyright by Cloud Mentor Pro | Confidential

VPC – Subnet (IPv4)

- AWS reserves **5 IP addresses (first 4 & last 1)** in each subnet
- These 5 IP addresses are not available for use and can't be assigned to an EC2 instance
- Example: if CIDR block 10.0.0.0/24, then reserved IP addresses are:
  - **10.0.0.0** – Network Address
  - **10.0.0.1** – reserved by AWS for the VPC router
  - **10.0.0.2** – reserved by AWS for mapping to Amazon-provided DNS
  - **10.0.0.3** – reserved by AWS for future use
  - **10.0.0.255** – Network Broadcast Address. AWS does not support broadcast in a VPC, therefore the address is reserved
- **Exam Tip**, if you need 29 IP addresses for EC2 instances:
  - You can't choose a subnet of size /27 (32 IP addresses, 32 – 5 = 27 < 29)
  - You need to choose a subnet of size /26 (64 IP addresses, 64 – 5 = 59 > 29)

© Copyright by Cloud Mentor Pro | Confidential

Adding Internet Gateway

© Copyright by Cloud Mentor Pro | Confidential

Internet Gateway (IGW)

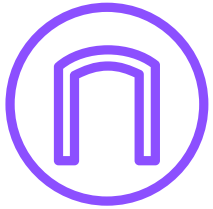

- Allows resources (e.g., EC2 instances) in a VPC connect to the Internet
- Must be created separately from a VPC
- One VPC can only be attached to one IGW and vice versa

- Internet Gateways on their own do not allow Internet access…
- Route tables must also be edited!

© Copyright by Cloud Mentor Pro | Confidential

Editing Route Tables

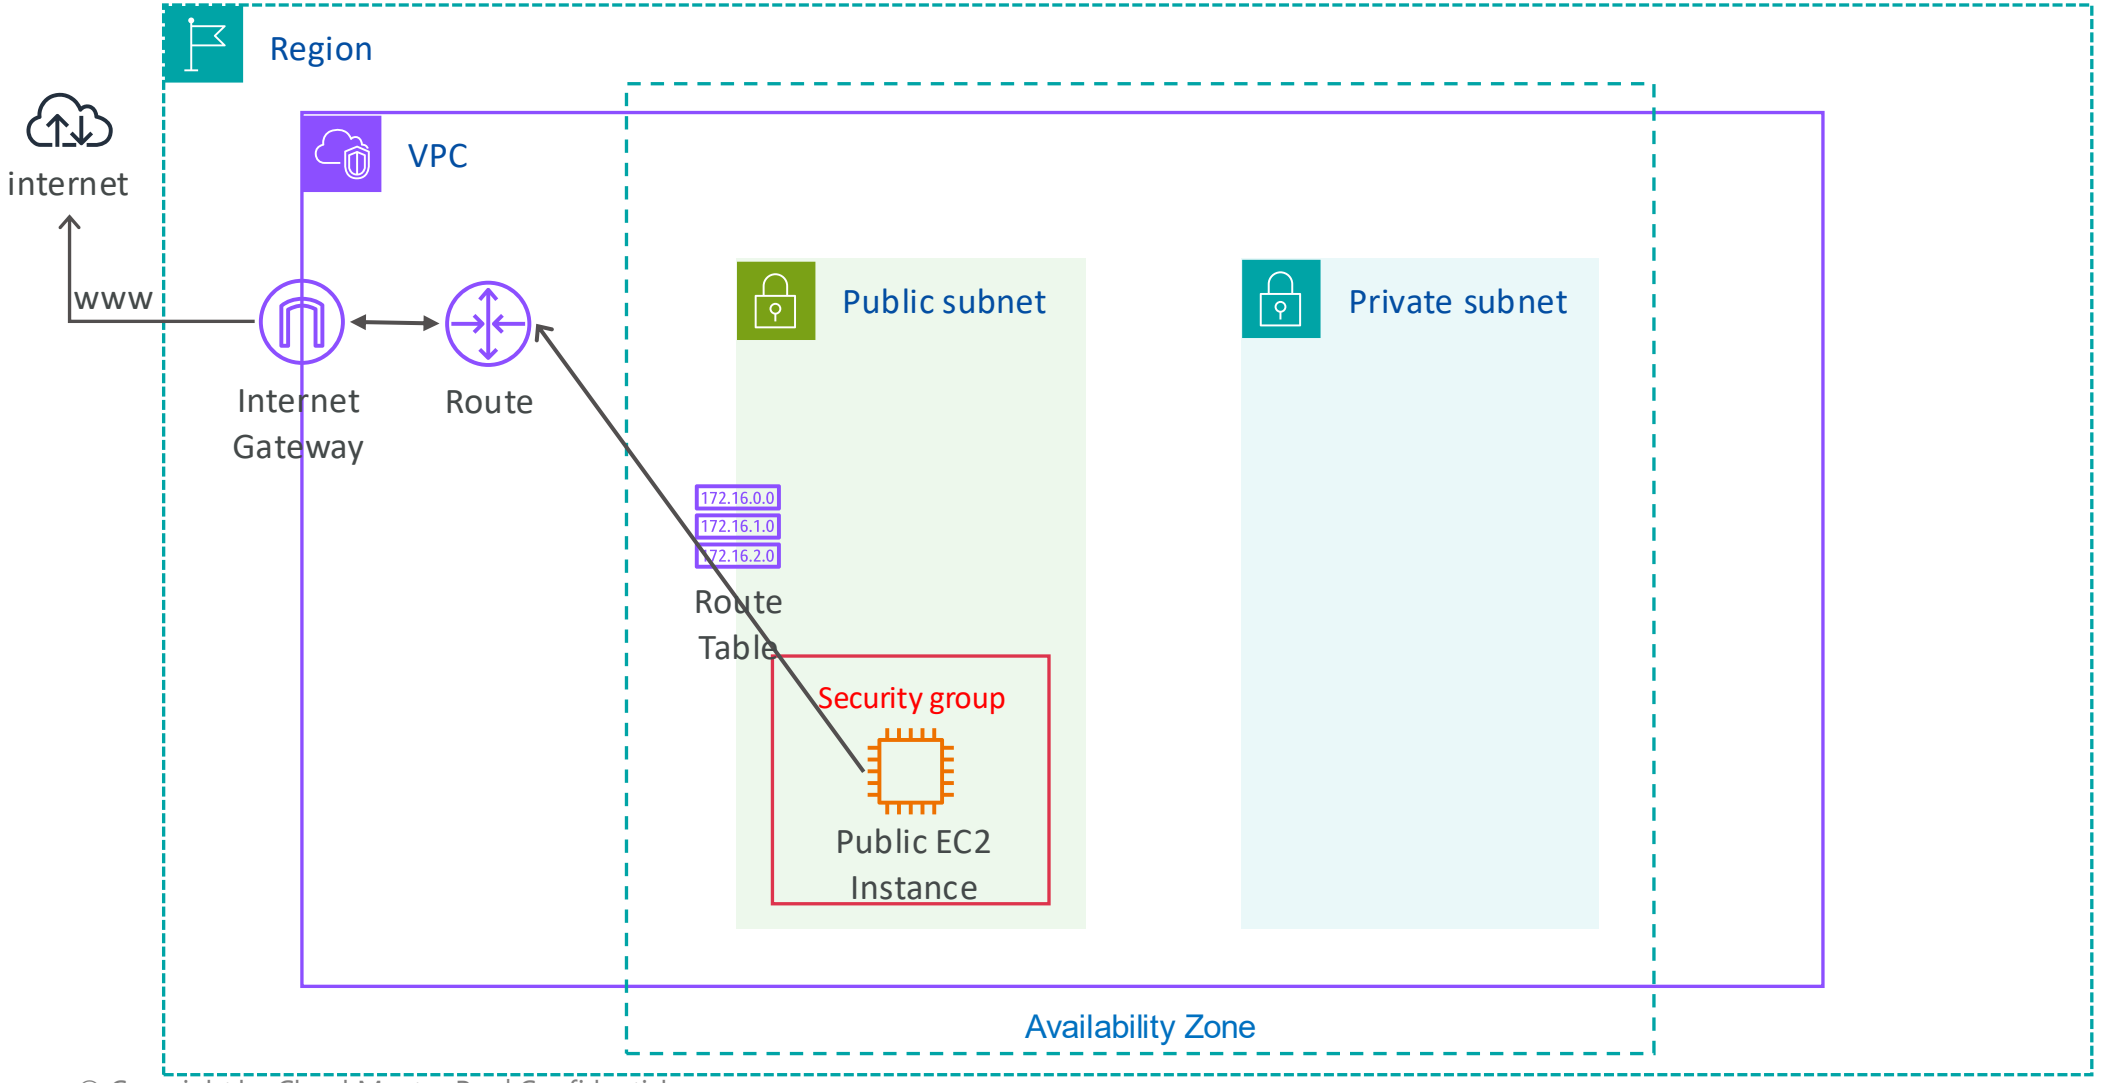

© Copyright by Cloud Mentor Pro | Confidential

Bastion Hosts

- We can use a Bastion Host to SSH into our private EC2 instances
- The bastion is in the public subnet which is then connected to all other private subnets
- Bastion Host security group must allow inbound from the internet on port 22 from restricted CIDR, for example the public CIDR of your corporation
- Security Group of the EC2 Instances must allow the Security Group of the Bastion Host, or the private IP of the Bastion host

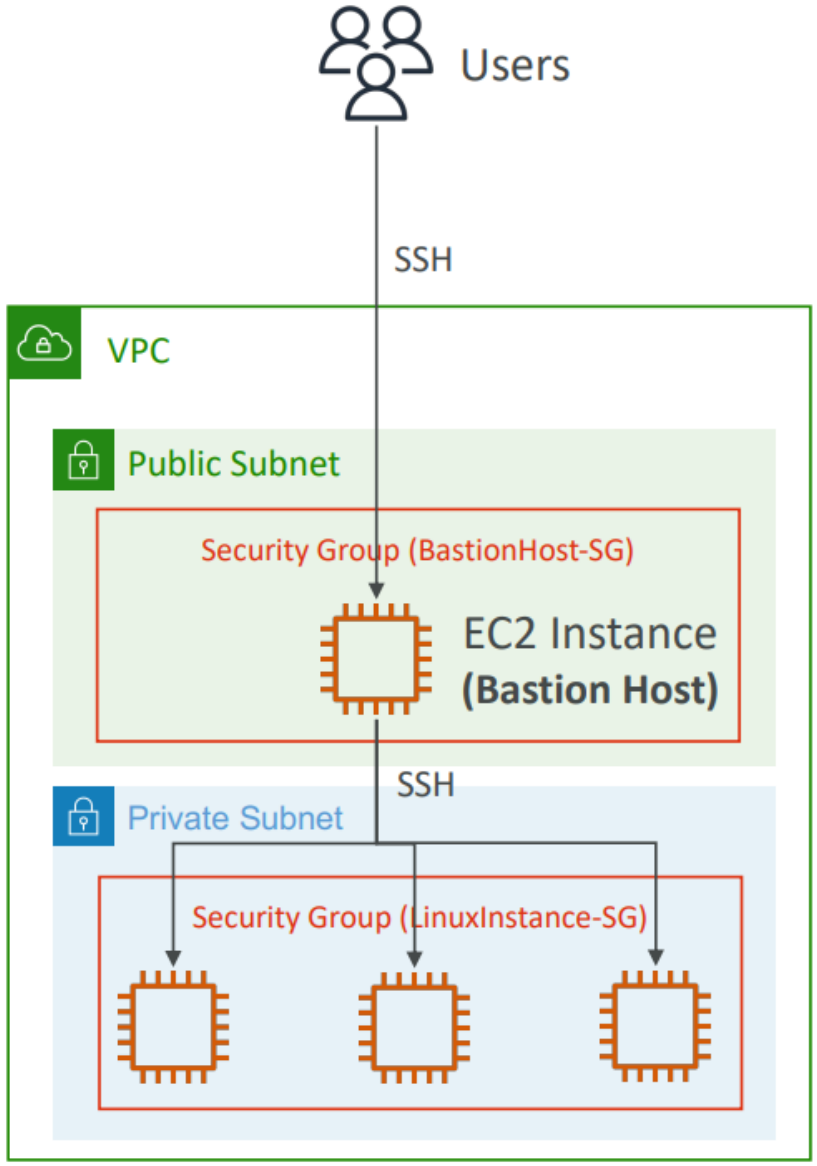

© Copyright by Cloud Mentor Pro | Confidential

NAT Gateway

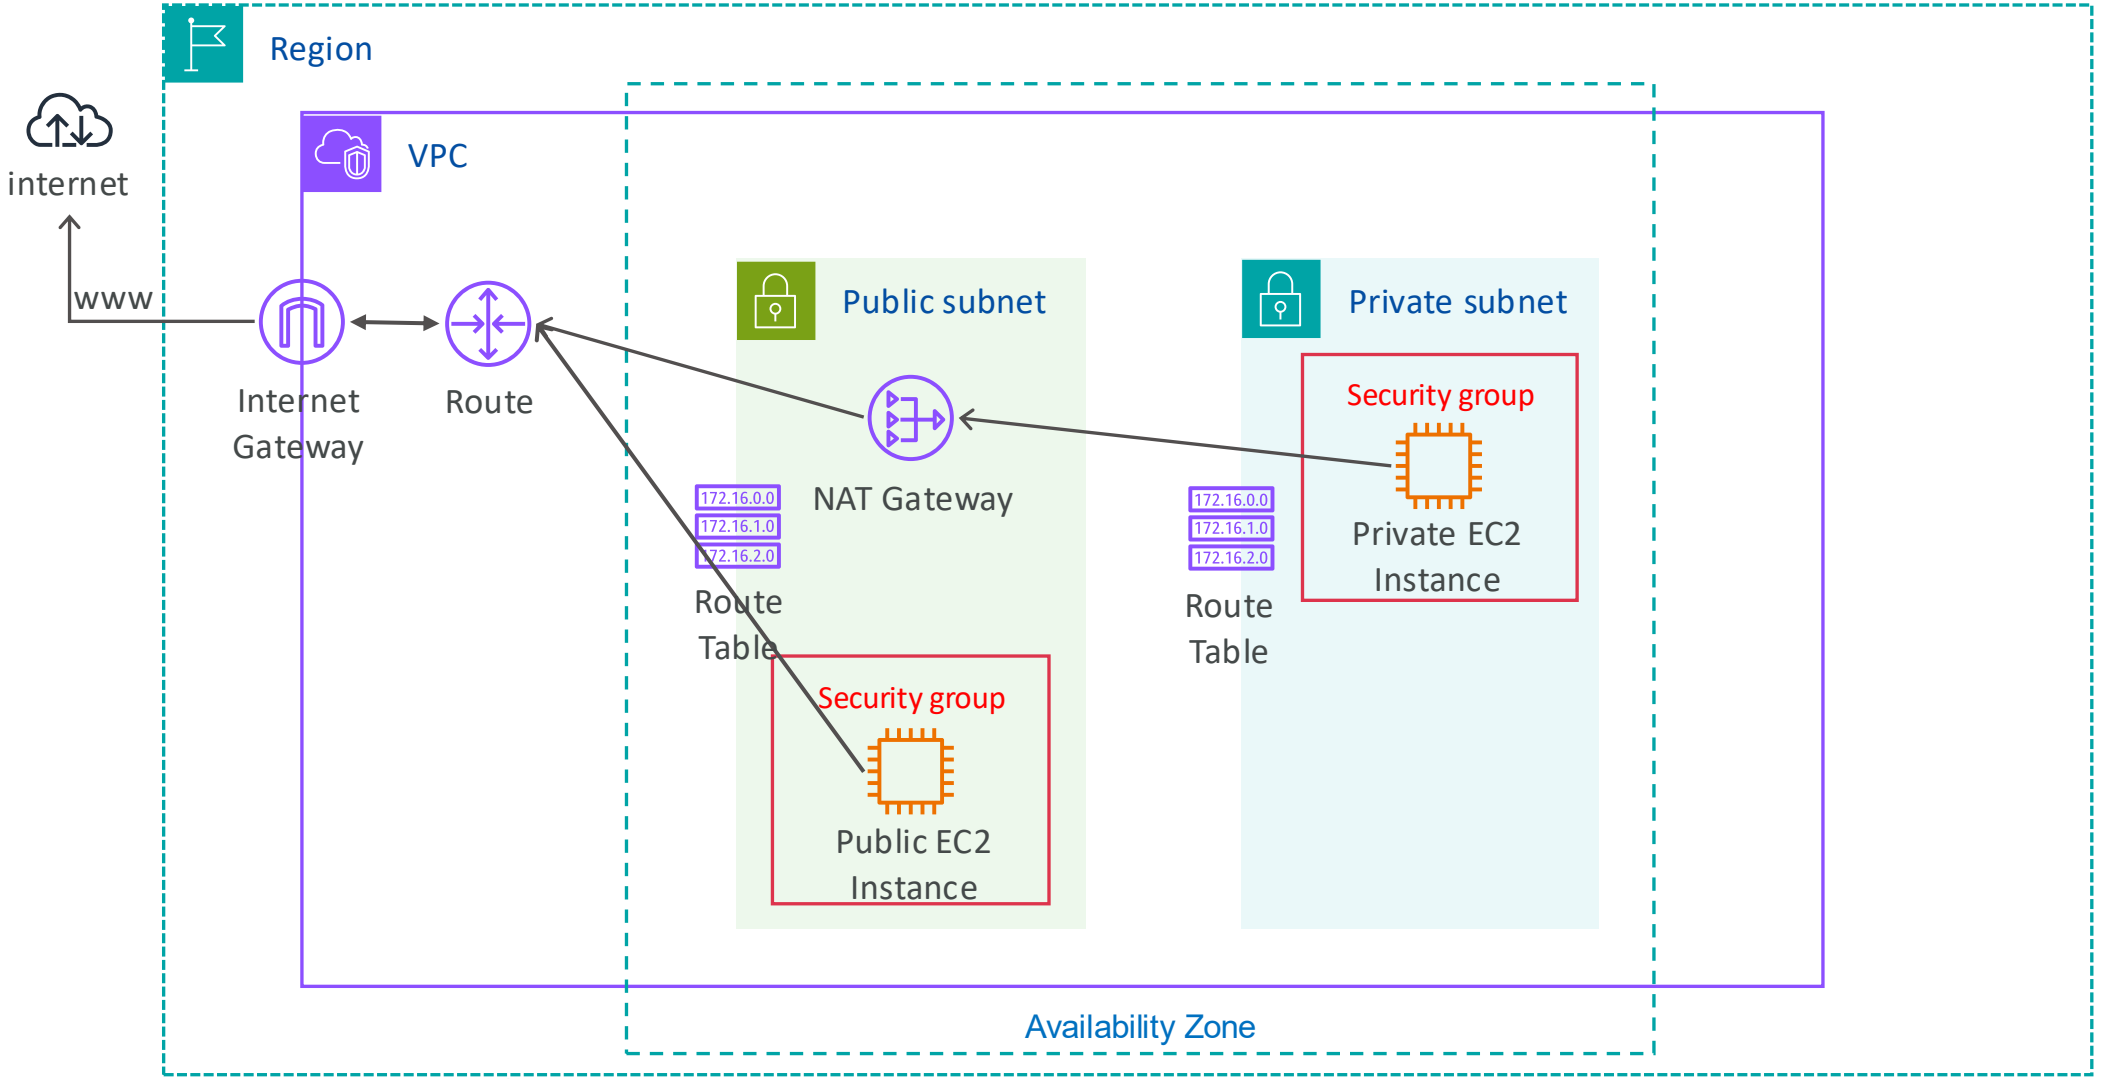

© Copyright by Cloud Mentor Pro | Confidential

NAT Gateway

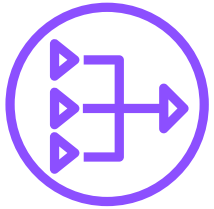

- AWS-managed NAT, higher bandwidth, high availability, no administration
- Pay per hour for usage and bandwidth
- NATGW is created in a specific Availability Zone, uses an Elastic IP
- Requires an IGW (Private Subnet => NATGW => IGW)
- 5 Gbps of bandwidth with automatic scaling up to 100 Gbps
- No Security Groups to manage / required

© Copyright by Cloud Mentor Pro | Confidential

Network Access Control List (NACL)

- NACL are like a firewall which control traffic from and to **subnets**
- **One NACL per subnet**, new subnets are assigned the **Default NACL**
- Can have ALLOW and DENY rules
- Rules only include IP addresses
- NACL are a great way of blocking a specific IP address at the subnet level

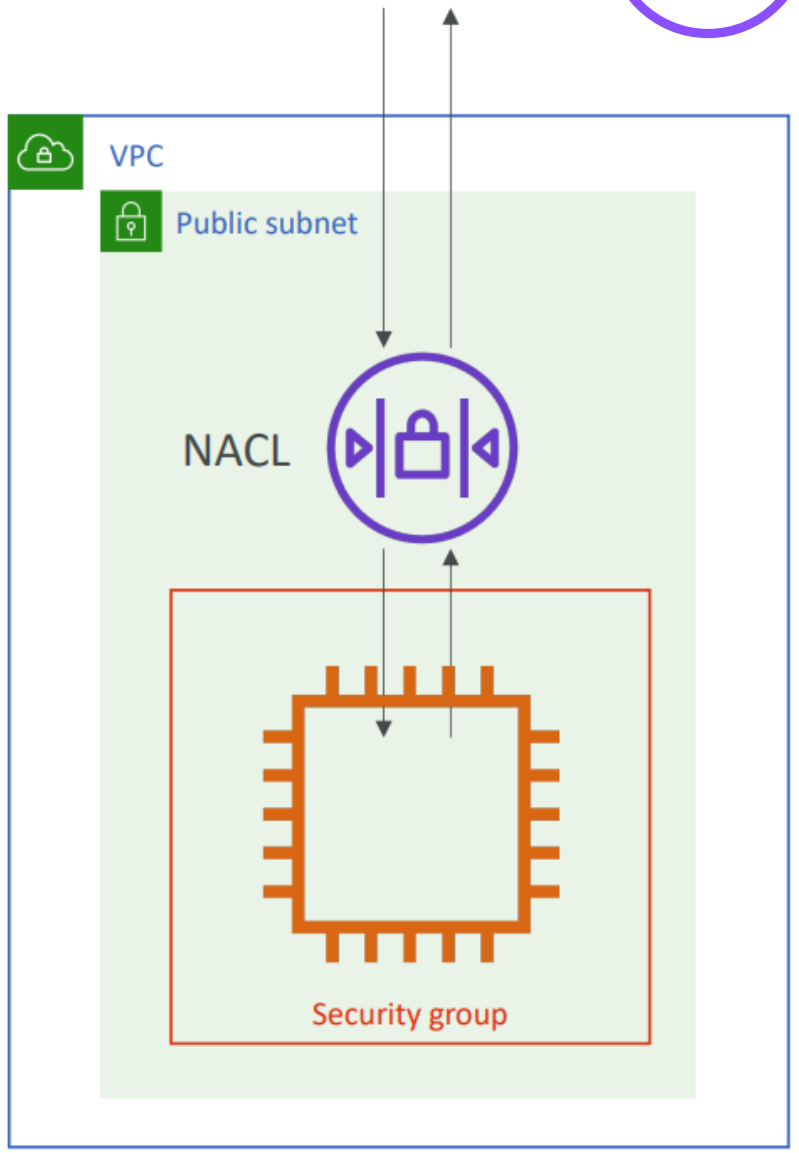

© Copyright by Cloud Mentor Pro | Confidential

NACLs

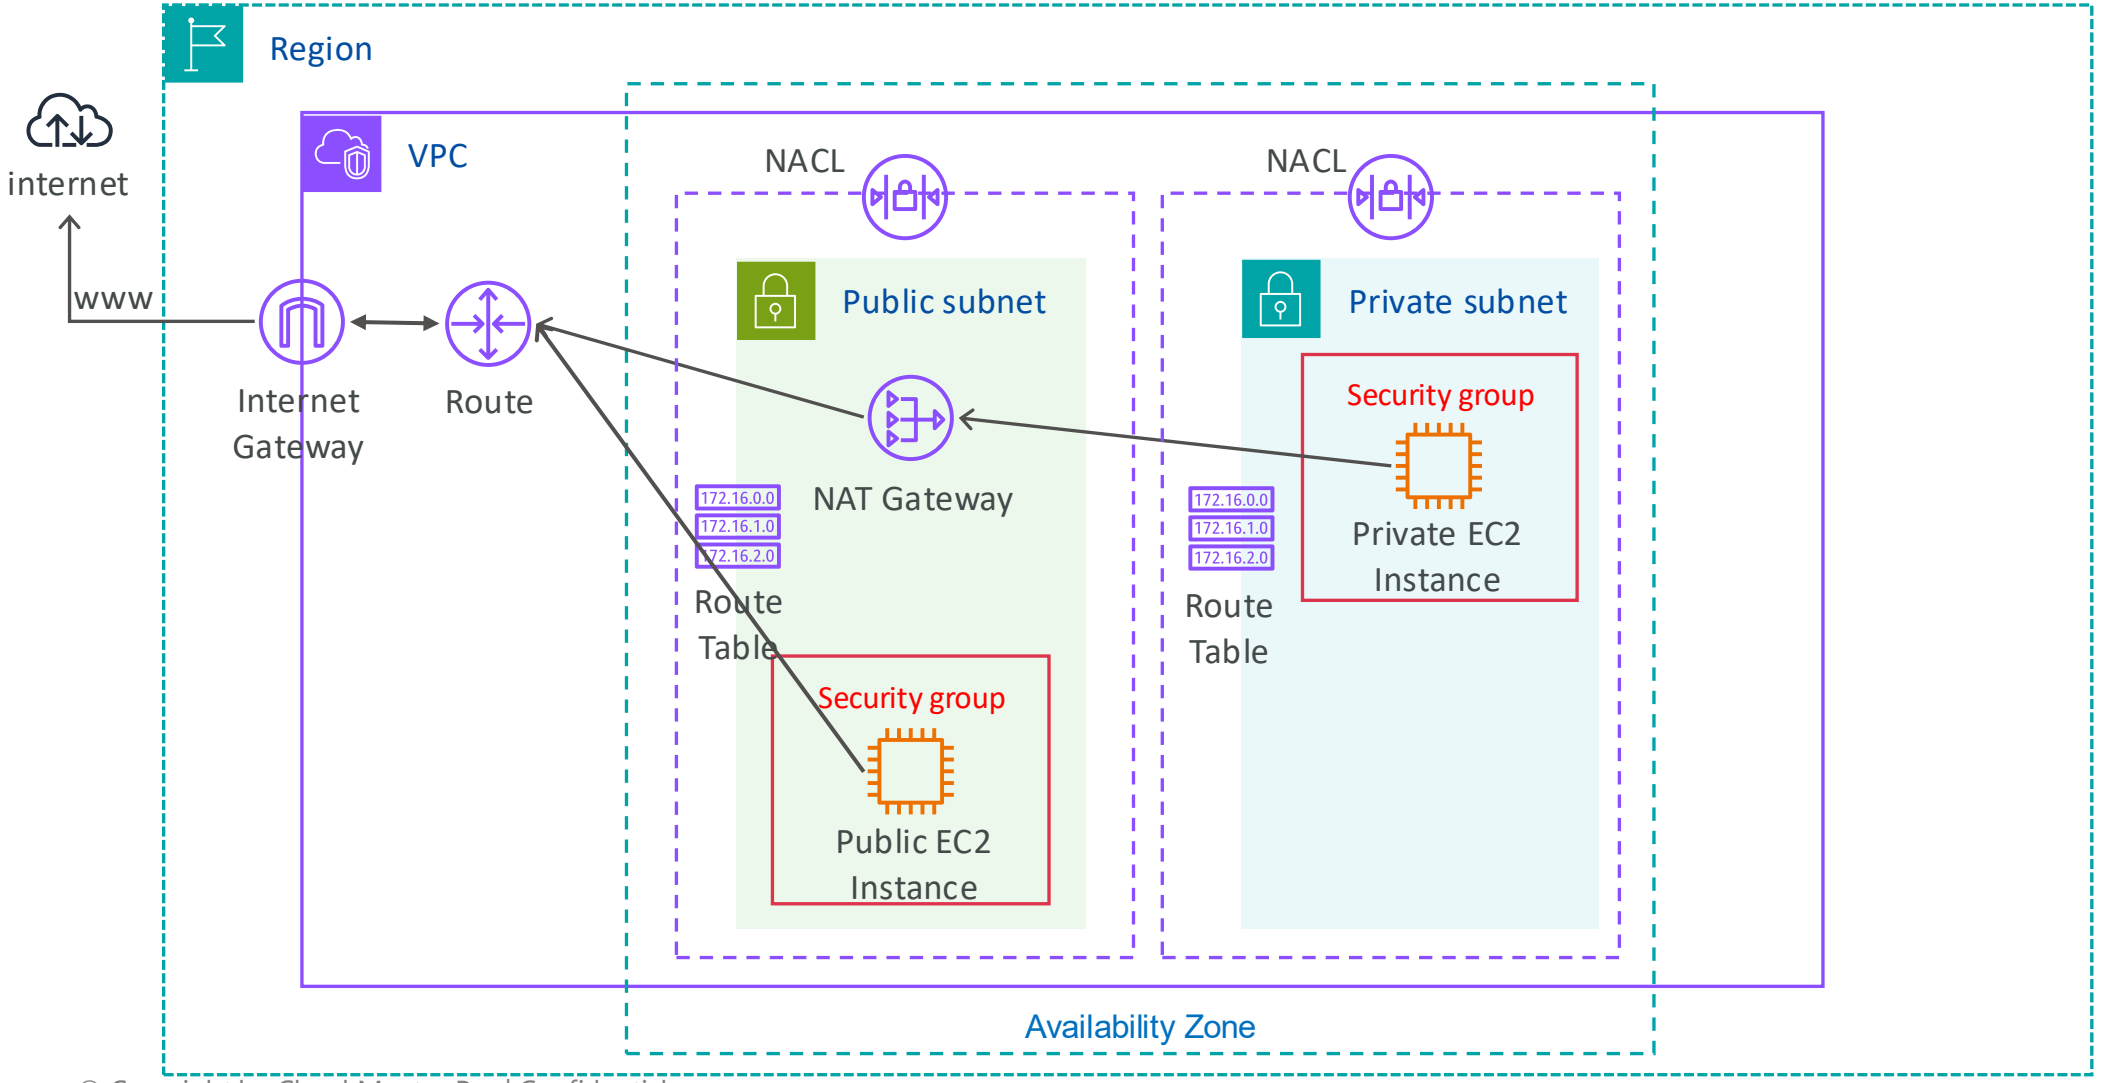

© Copyright by Cloud Mentor Pro | Confidential

Default NACL

- Accepts everything inbound/outbound with the subnets it's associated with
- Do NOT modify the Default NACL, instead create custom NACLs

Default NACL for a VPC that supports IPv4

Inbound Rules

Rule #	Type	Protocol	Port Range	Source	Allow/Deny
100	All IPv4 Traffic	All	All	0.0.0.0/0	ALLOW
*	All IPv4 Traffic	All	All	0.0.0.0/0	DENY

Outbound Rules

Rule #	Type	Protocol	Port Range	Destination	Allow/Deny
100	All IPv4 Traffic	All	All	0.0.0.0/0	ALLOW
*	All IPv4 Traffic	All	All	0.0.0.0/0	DENY

© Copyright by Cloud Mentor Pro | Confidential

Security Group vs. NACLs

Security Group	NACL
Operates at the instance level	Operates at the subnet level
Supports allow rules only	Supports allow rules and deny rules
Stateful: return traffic is automatically allowed, regardless of any rules	**Stateless:** return traffic must be explicitly allowed by rules (think of ephemeral ports)
All rules are evaluated before deciding whether to allow traffic	Rules are evaluated in order (lowest to highest) when deciding whether to allow traffic, first match wins
Applies to an EC2 instance when specified by someone	Automatically applies to all EC2 instances in the subnet that it's associated with

You usually only need a security group to establish connectivity to your AWS resources. In most cases, a NACL isn't necessary, it's only required when your system has strict security requirements.

© Copyright by Cloud Mentor Pro | Confidential

https://docs.aws.amazon.com/vpc/latest/userguide/vpc-network-acls.html

VPC Peering

- Privately connect two VPCs using AWS'network
- Make them behave as if they were in the same network
- Must not have overlapping CIDRs
- VPC Peering connection is NOT transitive (must be established for each VPC that need to communicate with one another)
- **You must update route tables in each VPC's subnets to ensure EC2 instances can communicate with each other**

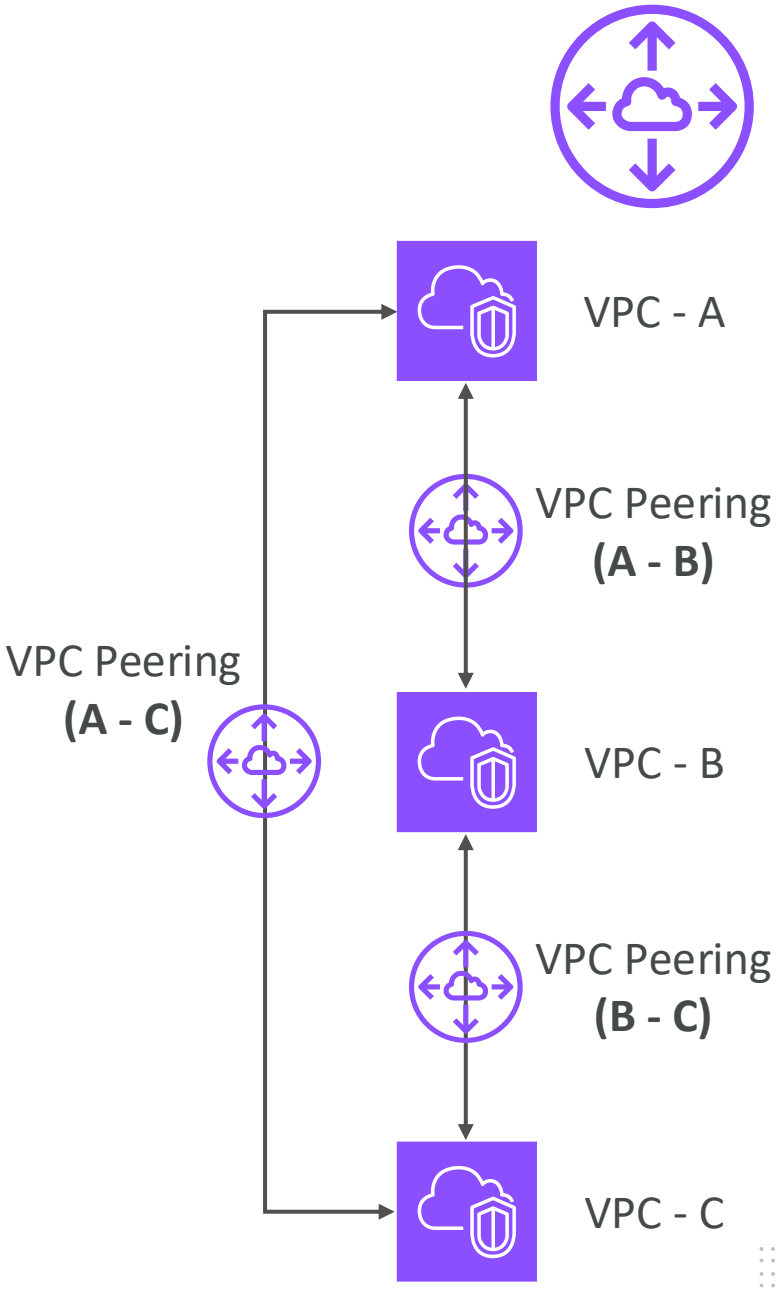

© Copyright by Cloud Mentor Pro | Confidential

VPC Peering

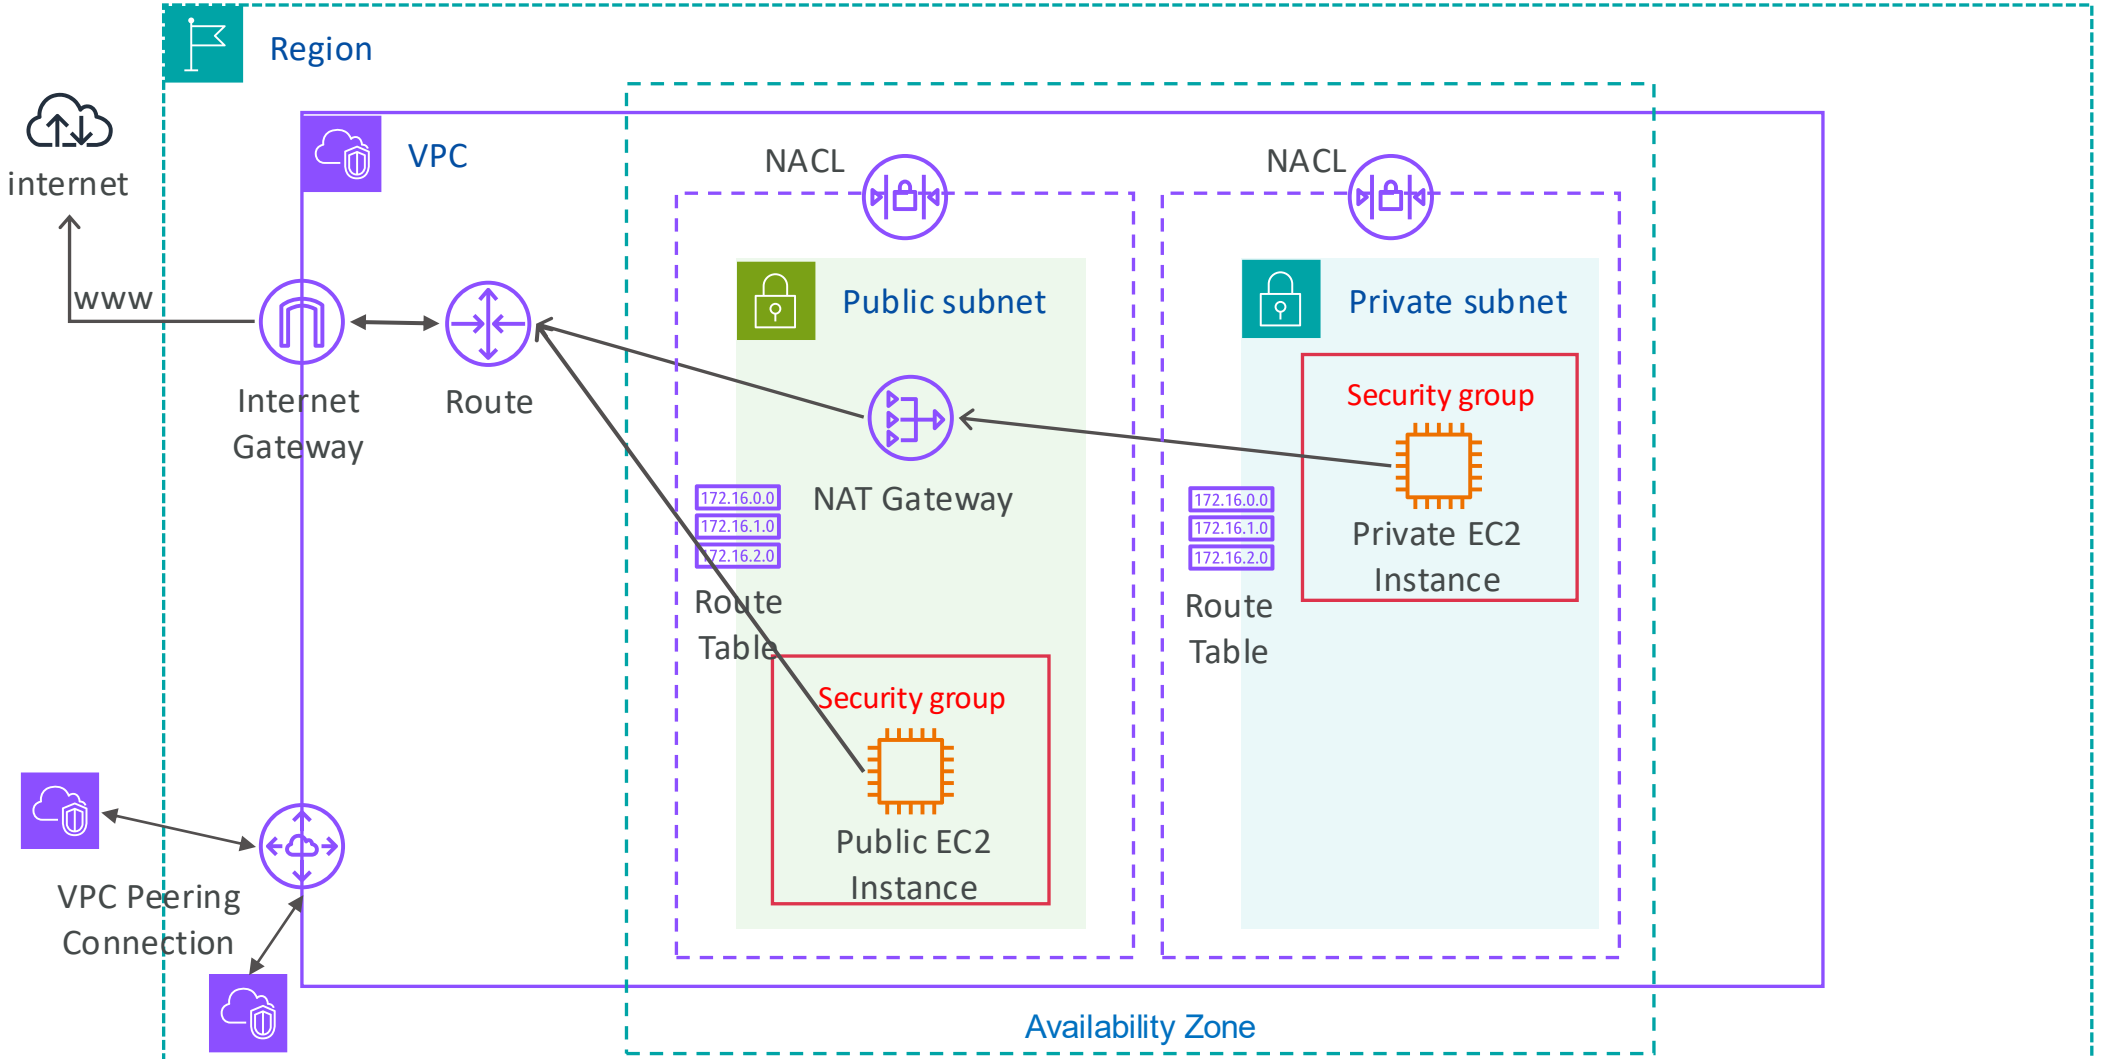

© Copyright by Cloud Mentor Pro | Confidential

VPC Endpoints

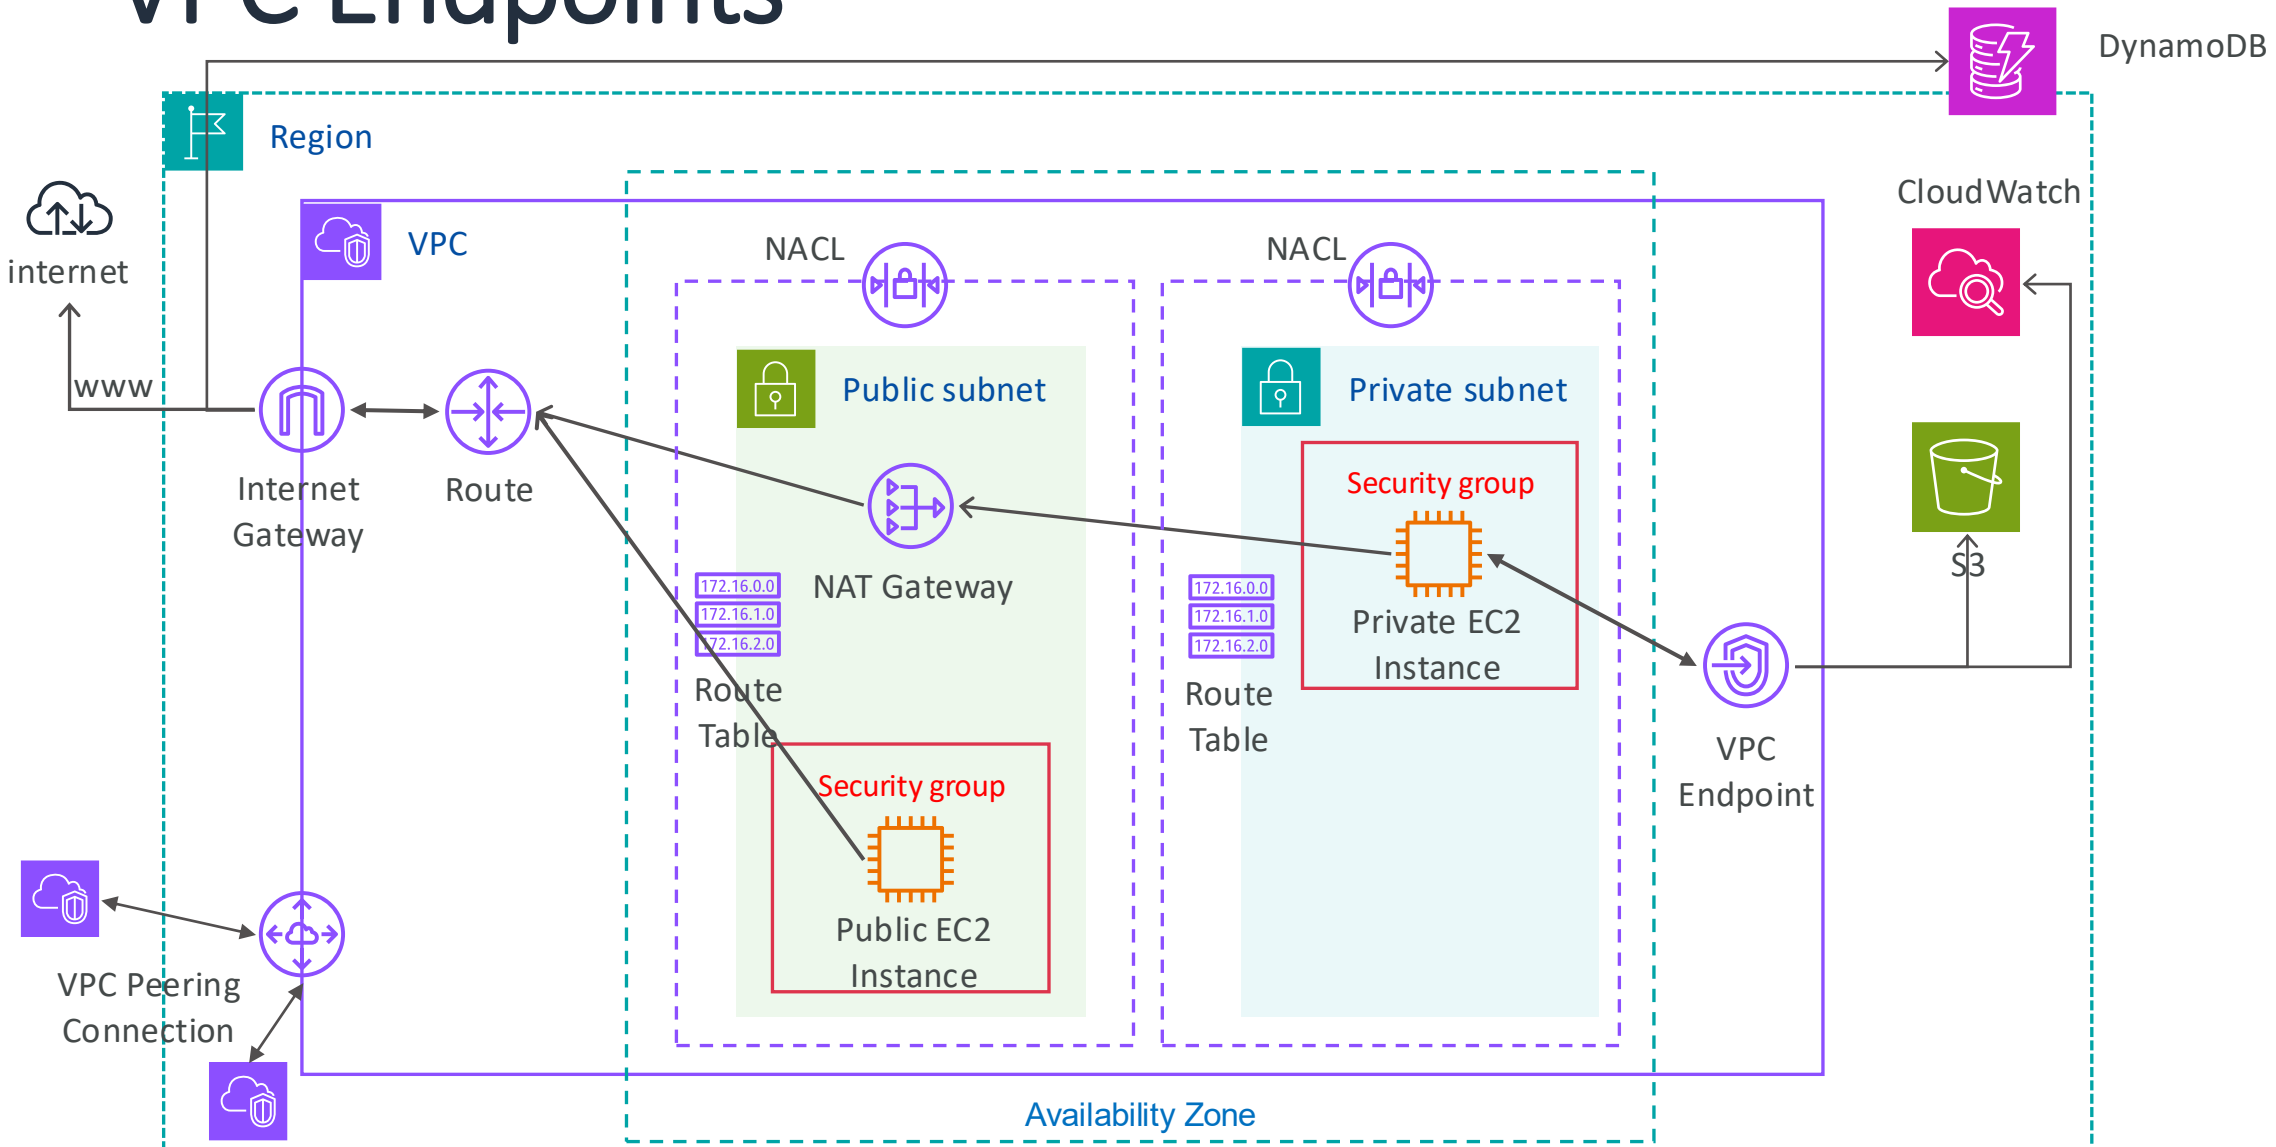

© Copyright by Cloud Mentor Pro | Confidential

VPC Endpoints

- Every AWS service is publicly exposed (public URL)
- VPC Endpoints (powered by AWS PrivateLink) allows you to connect to AWS services using a private network instead of using the public Internet
- They're redundant and scale horizontally
- They remove the need of IGW, NATGW, … to access AWS Services
- In case of issues:
  - Check DNS Setting Resolution in your VPC
  - Check Route Tables

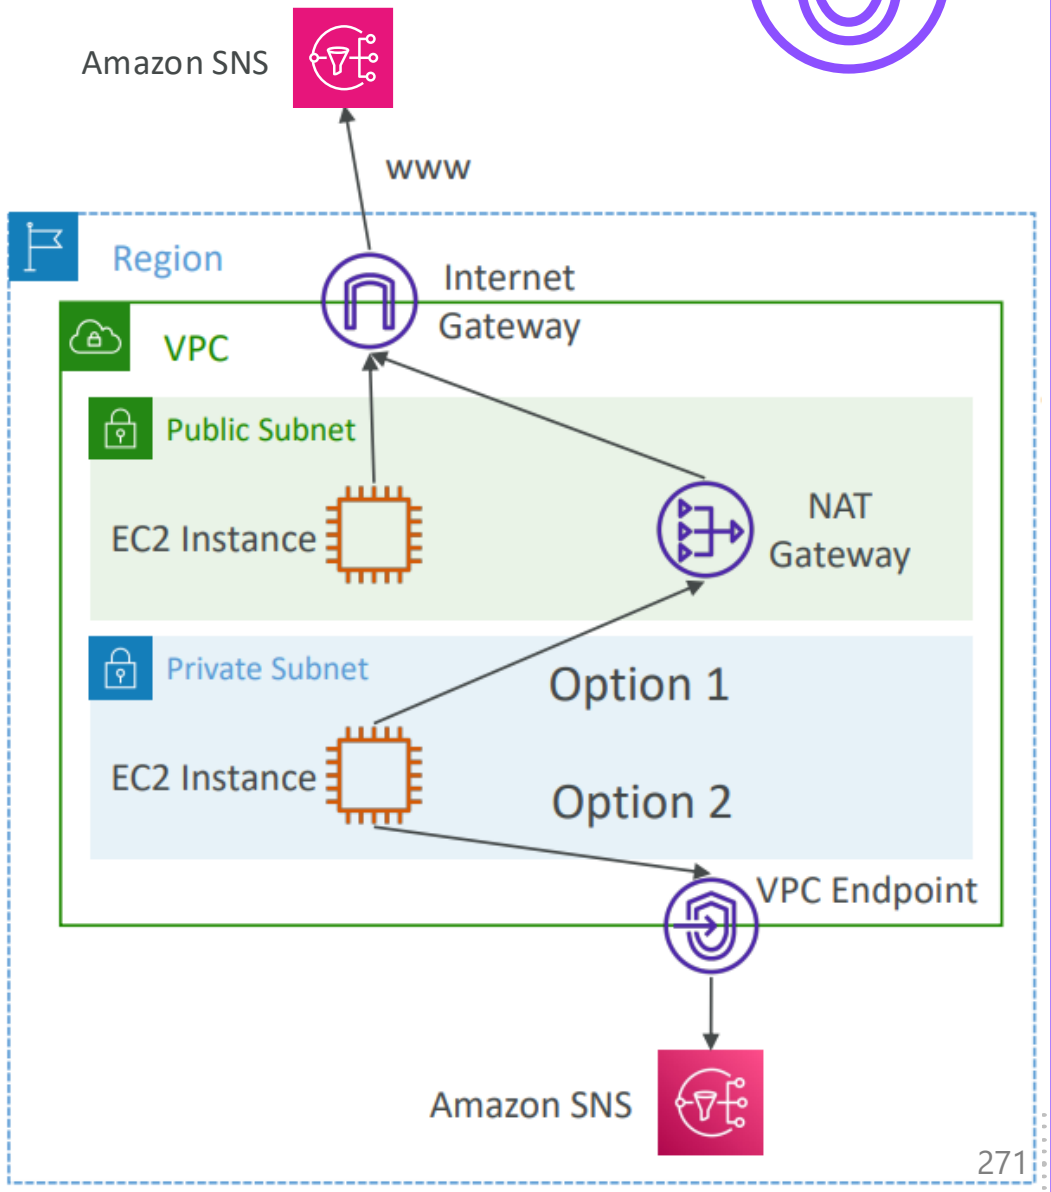

© Copyright by Cloud Mentor Pro | Confidential

Types of Endpoints

- Interface Endpoints (powered by PrivateLink)
  - Provisions an ENI (private IP address) as an entry point (must attach a Security Group)
  - Supports most AWS services
  - $ per hour + $ per GB of data processed
- Gateway Endpoints
  - Provisions a gateway and must be used as a target in a route table (does not use security groups)
  - Supports both S3 and DynamoDB
  - Free

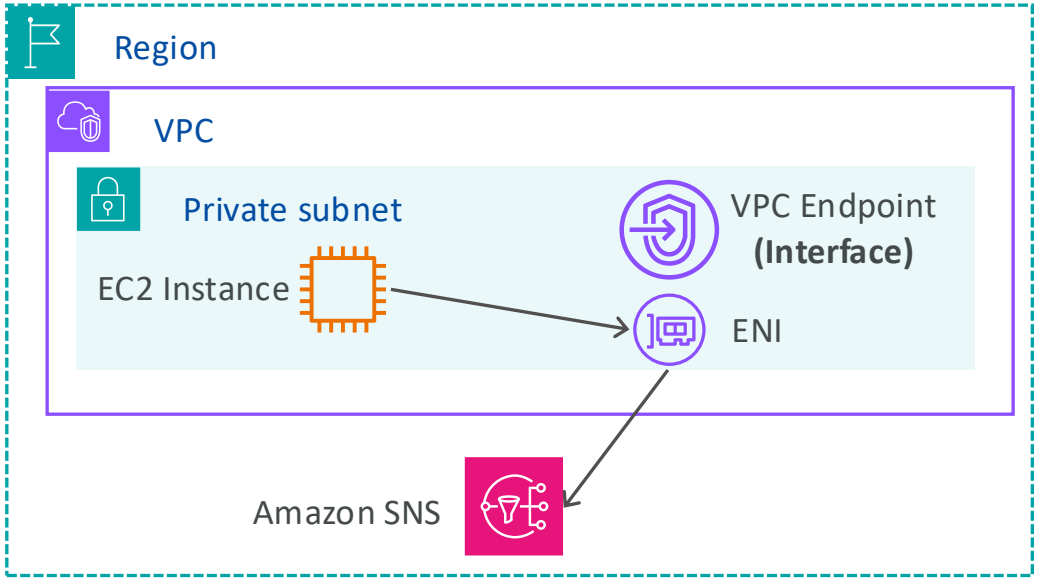

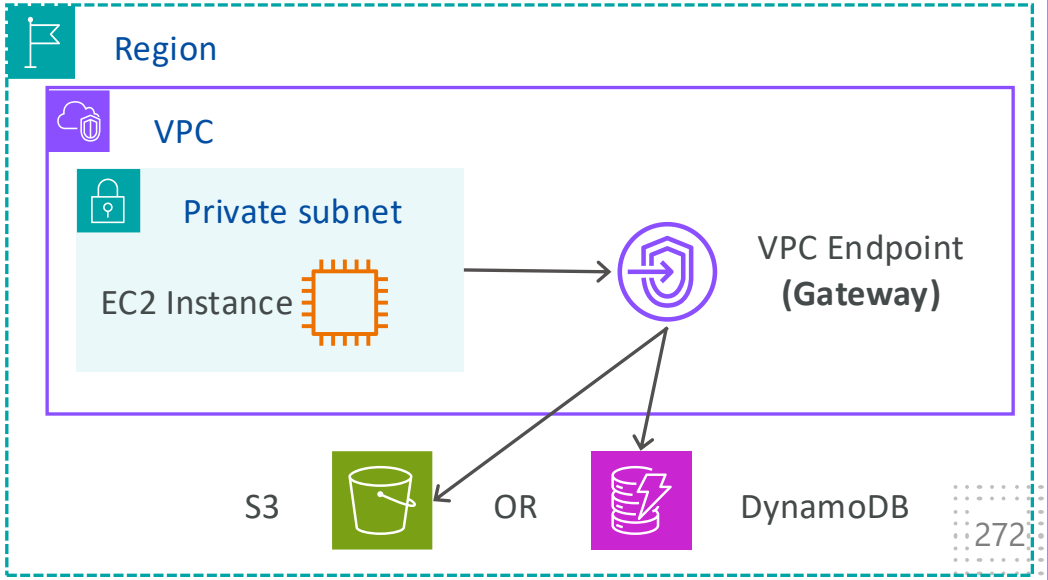

© Copyright by Cloud Mentor Pro | Confidential

VPC Flow Logs

- Capture information about IP traffic going into your interfaces:
  - VPC Flow Logs
  - Subnet Flow Logs
  - Elastic Network Interface (ENI) Flow Logs
- Helps to monitor & troubleshoot connectivity issues
- Flow logs data can go to S3, CloudWatch Logs, and Kinesis Data Firehose
- Captures network information from AWS managed interfaces too: ELB,
- RDS, ElastiCache, Redshift, WorkSpaces, NATGW, Transit Gateway…

© Copyright by Cloud Mentor Pro | Confidential

VPC Flow Logs

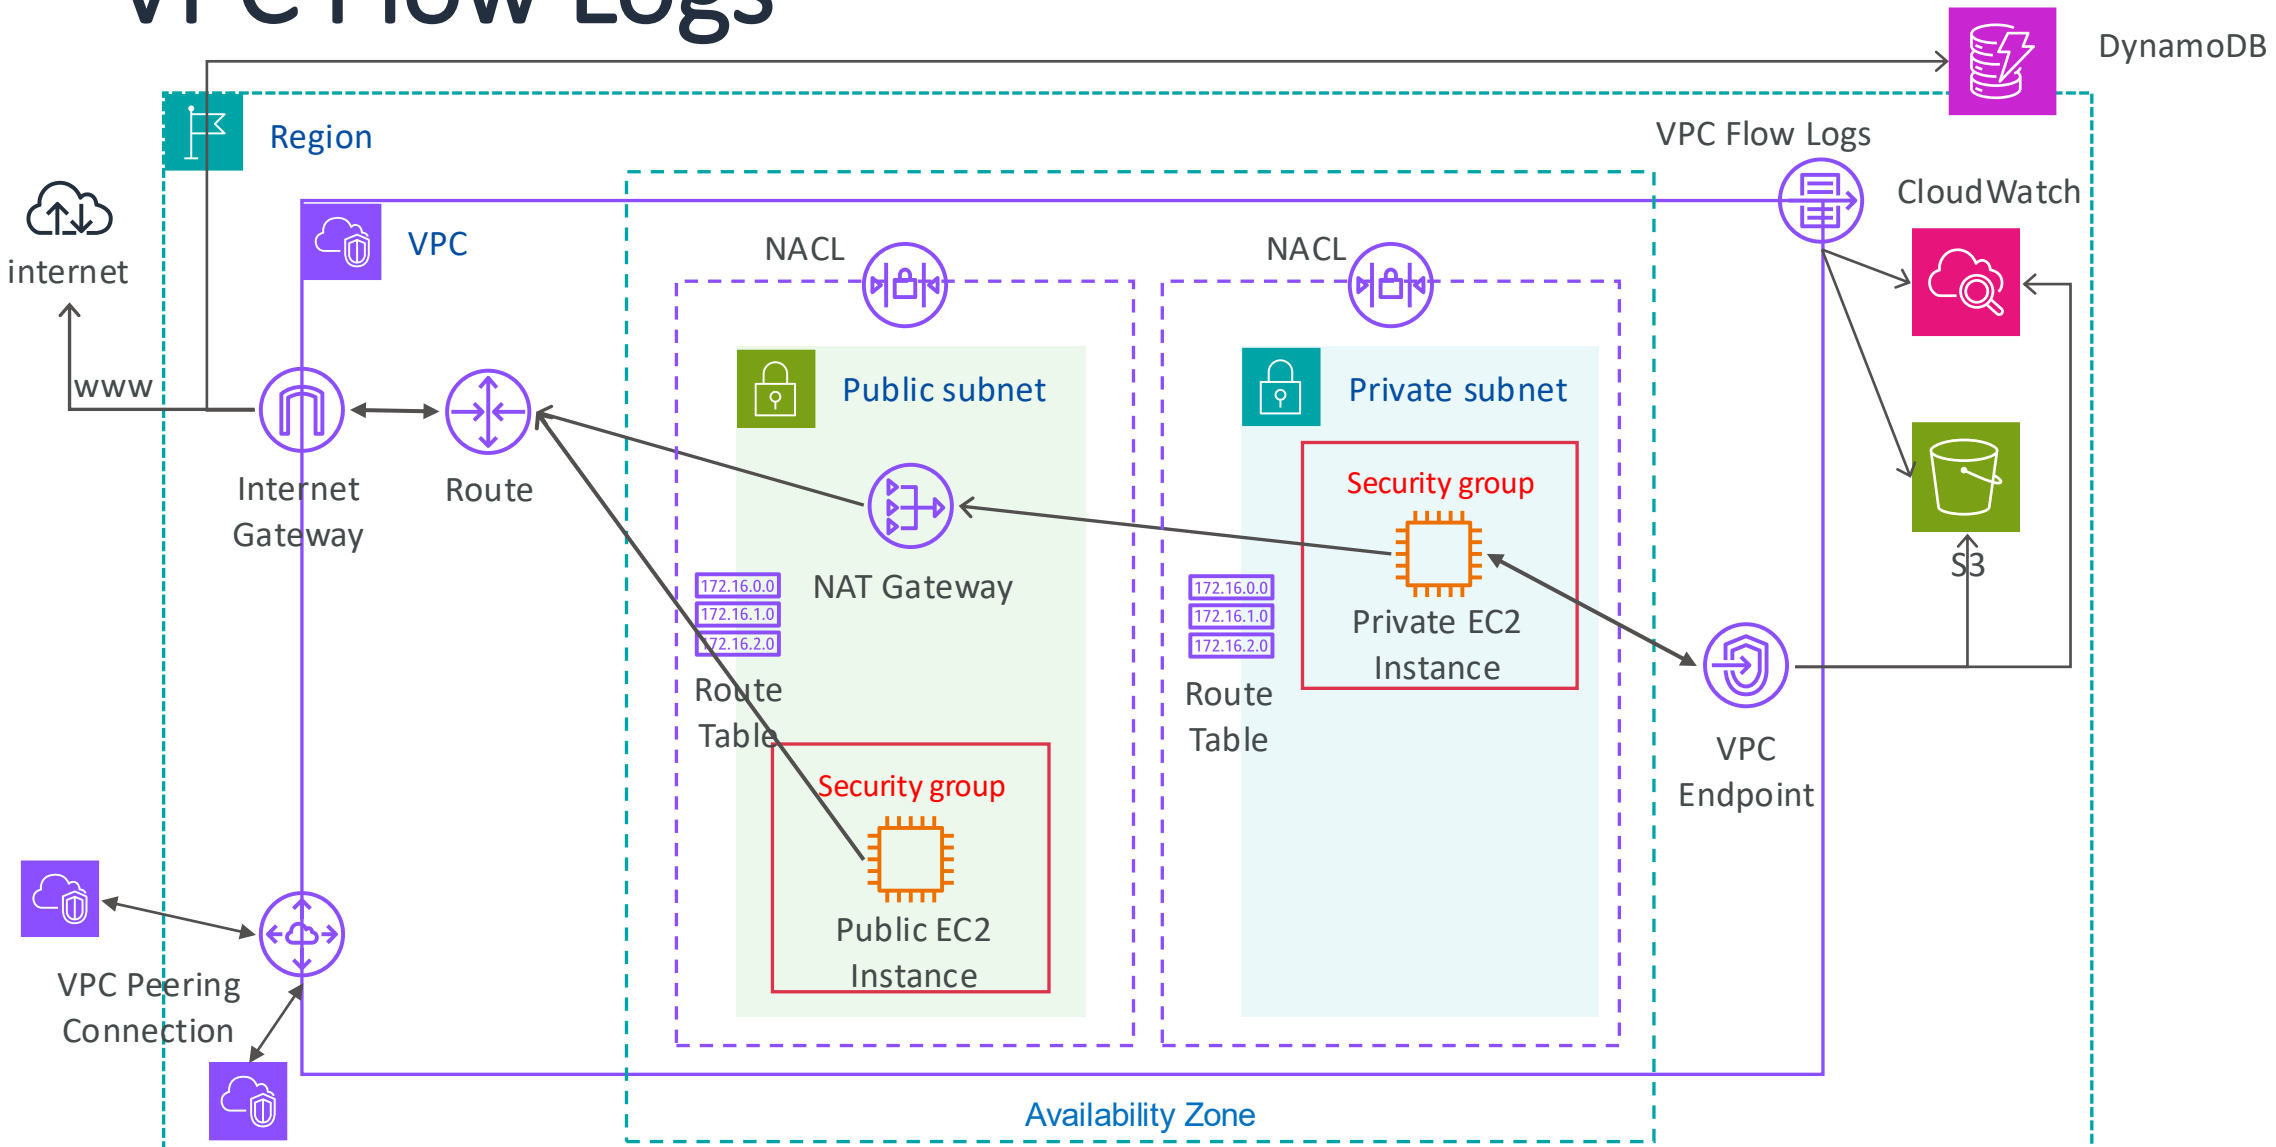

© Copyright by Cloud Mentor Pro | Confidential

VPC Flow Logs Syntax

- **srcaddr & dstaddr** – help identify problematic IP
- **srcport & dstport** – help identity problematic ports
- **Action** – success or failure of the request due to Security Group / NACL
- Can be used for analytics on usage patterns, or malicious behavior
- **Query VPC flow logs using Athena on S3 or CloudWatch Logs Insights**
- Flow Logs examples: https://docs.aws.amazon.com/vpc/latest/userguide/flow-logs-records-examples.html

© Copyright by Cloud Mentor Pro | Confidential

VPC Flow Logs – Architectures

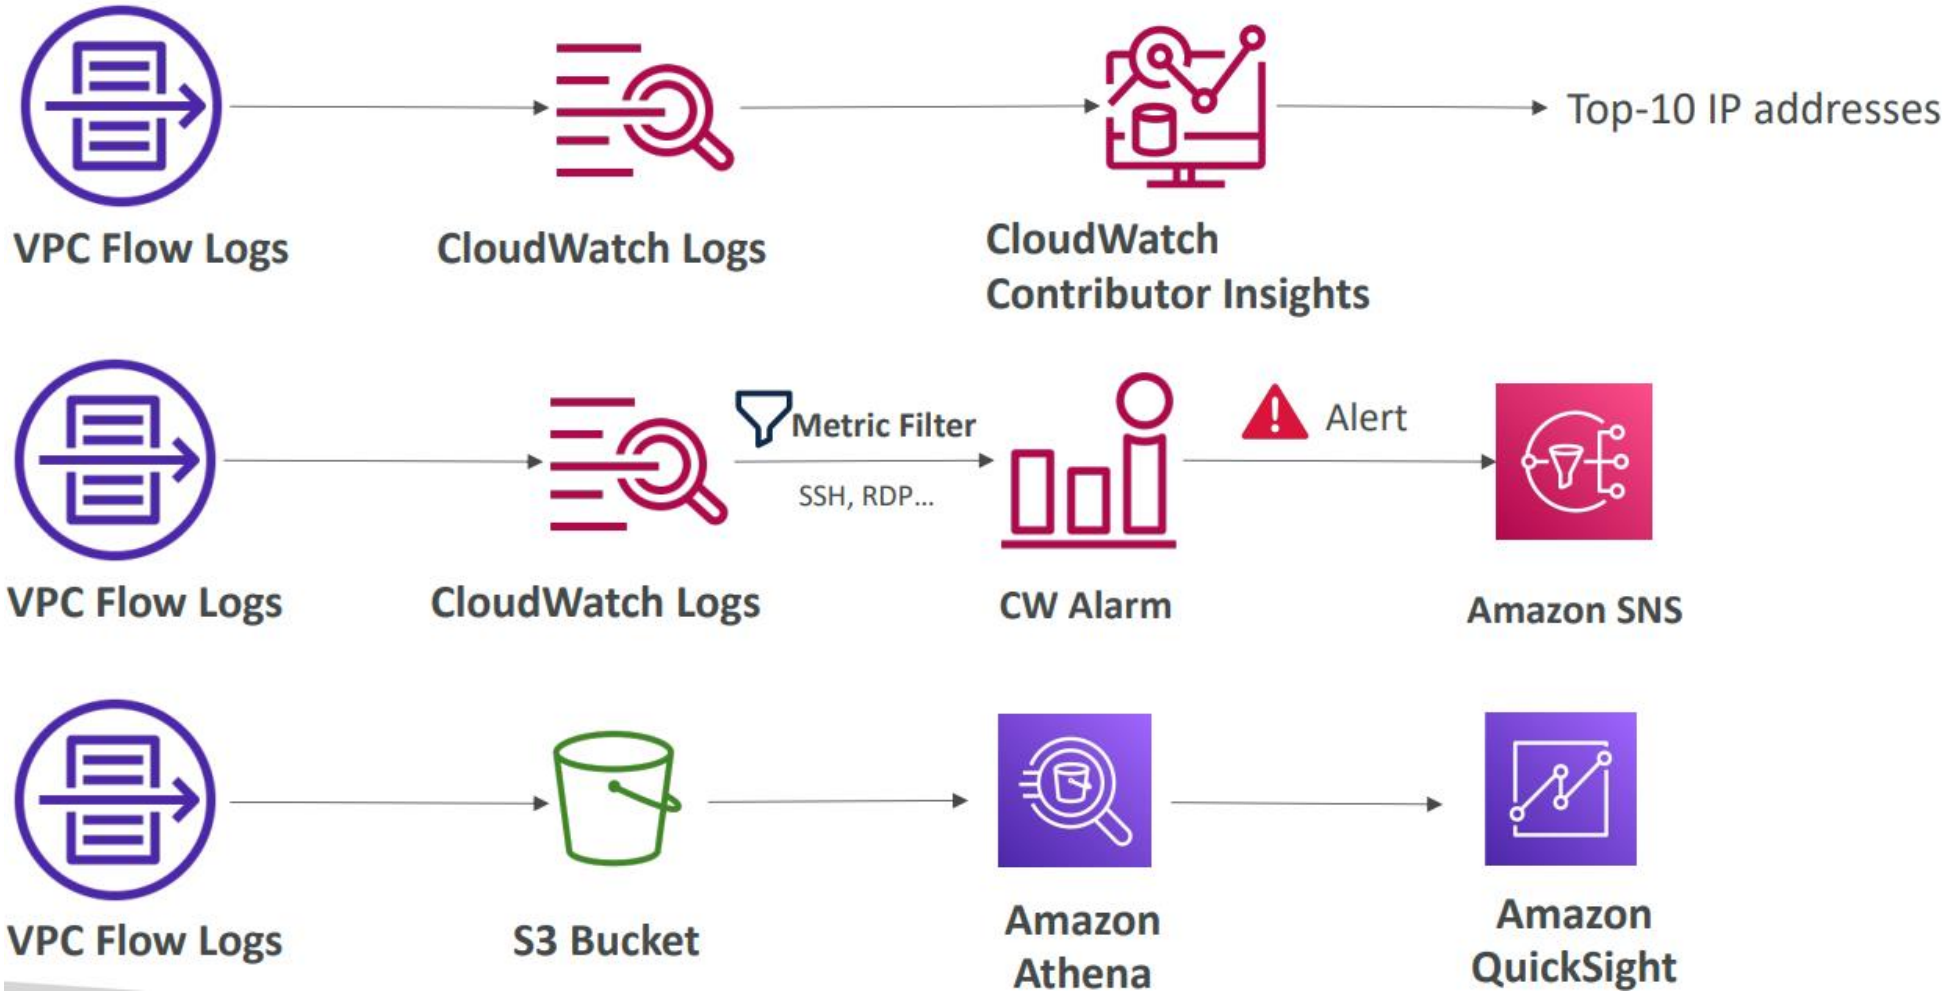

© Copyright by Cloud Mentor Pro | Confidential

AWS Site-to-Site VPN

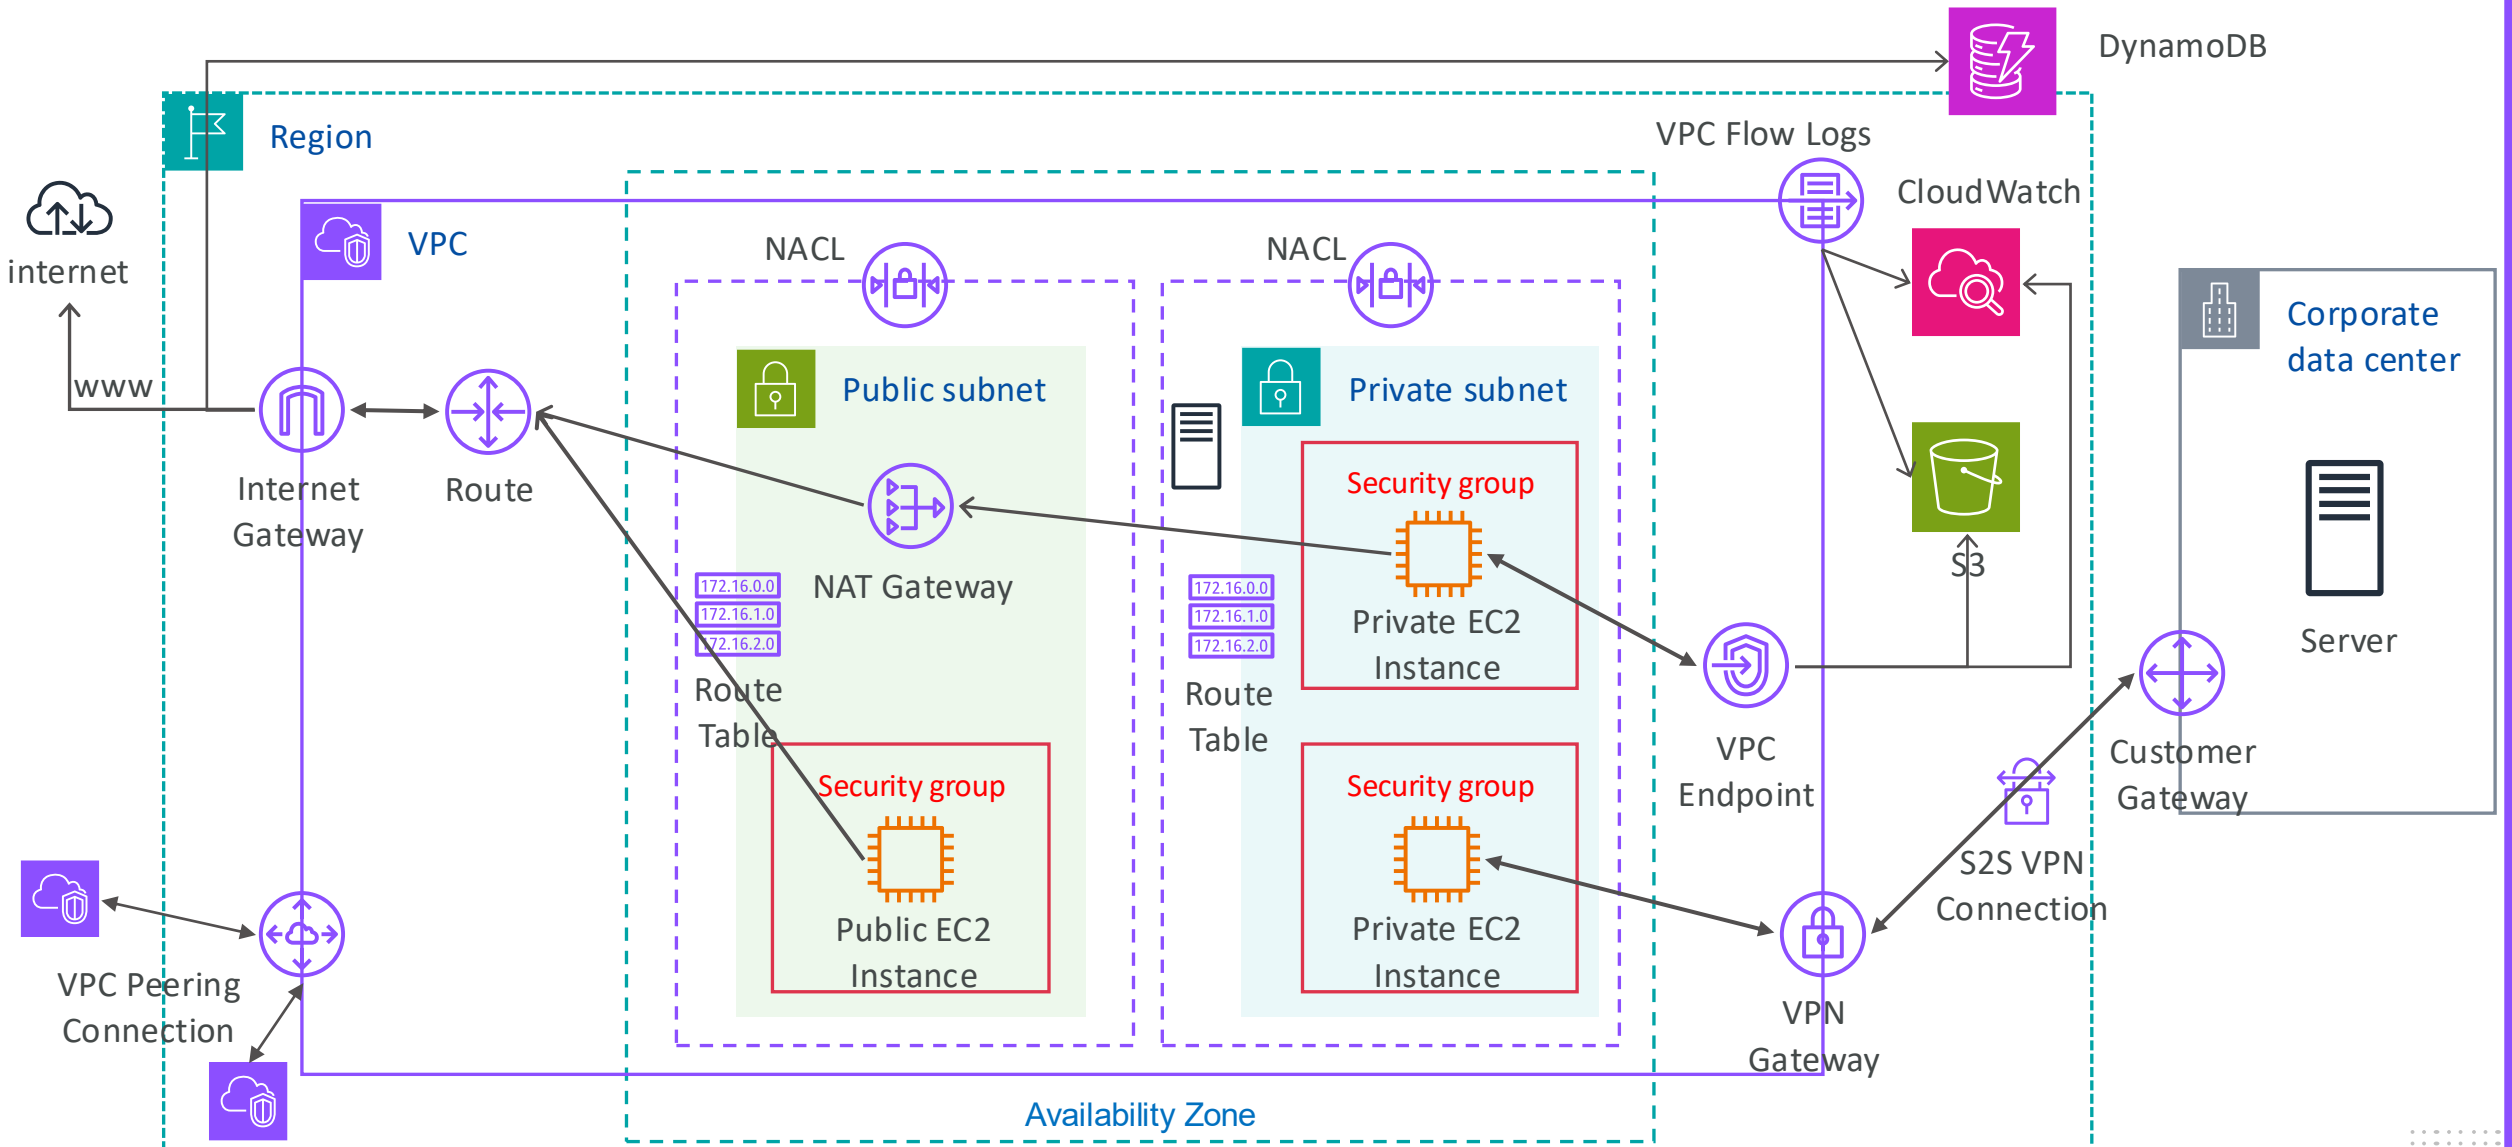

© Copyright by Cloud Mentor Pro | Confidential

Site to Site VPN & Direct Connect

- Site to Site VPN
  - Connect an on-premises VPN to AWS
  - The connection is automatically encrypted
  - Goes over the public internet
- Direct Connect (DX)
  - Establish a physical connection between on-premises and AWS
  - The connection is private, secure and fast
  - Goes over a private network
  - Takes at least a month to establish

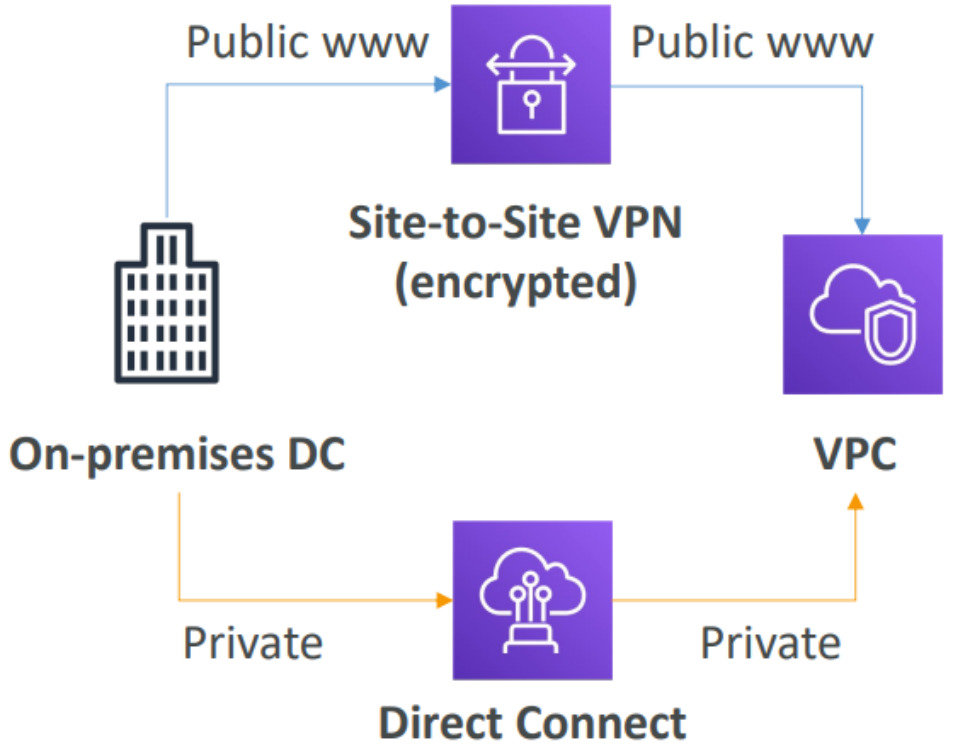

© Copyright by Cloud Mentor Pro | Confidential

Site-to-Site VPN connection as a backup

- In case Direct Connect fails, you can set up a backup Direct Connect connection (expensive), or a Site-to-Site VPN connection

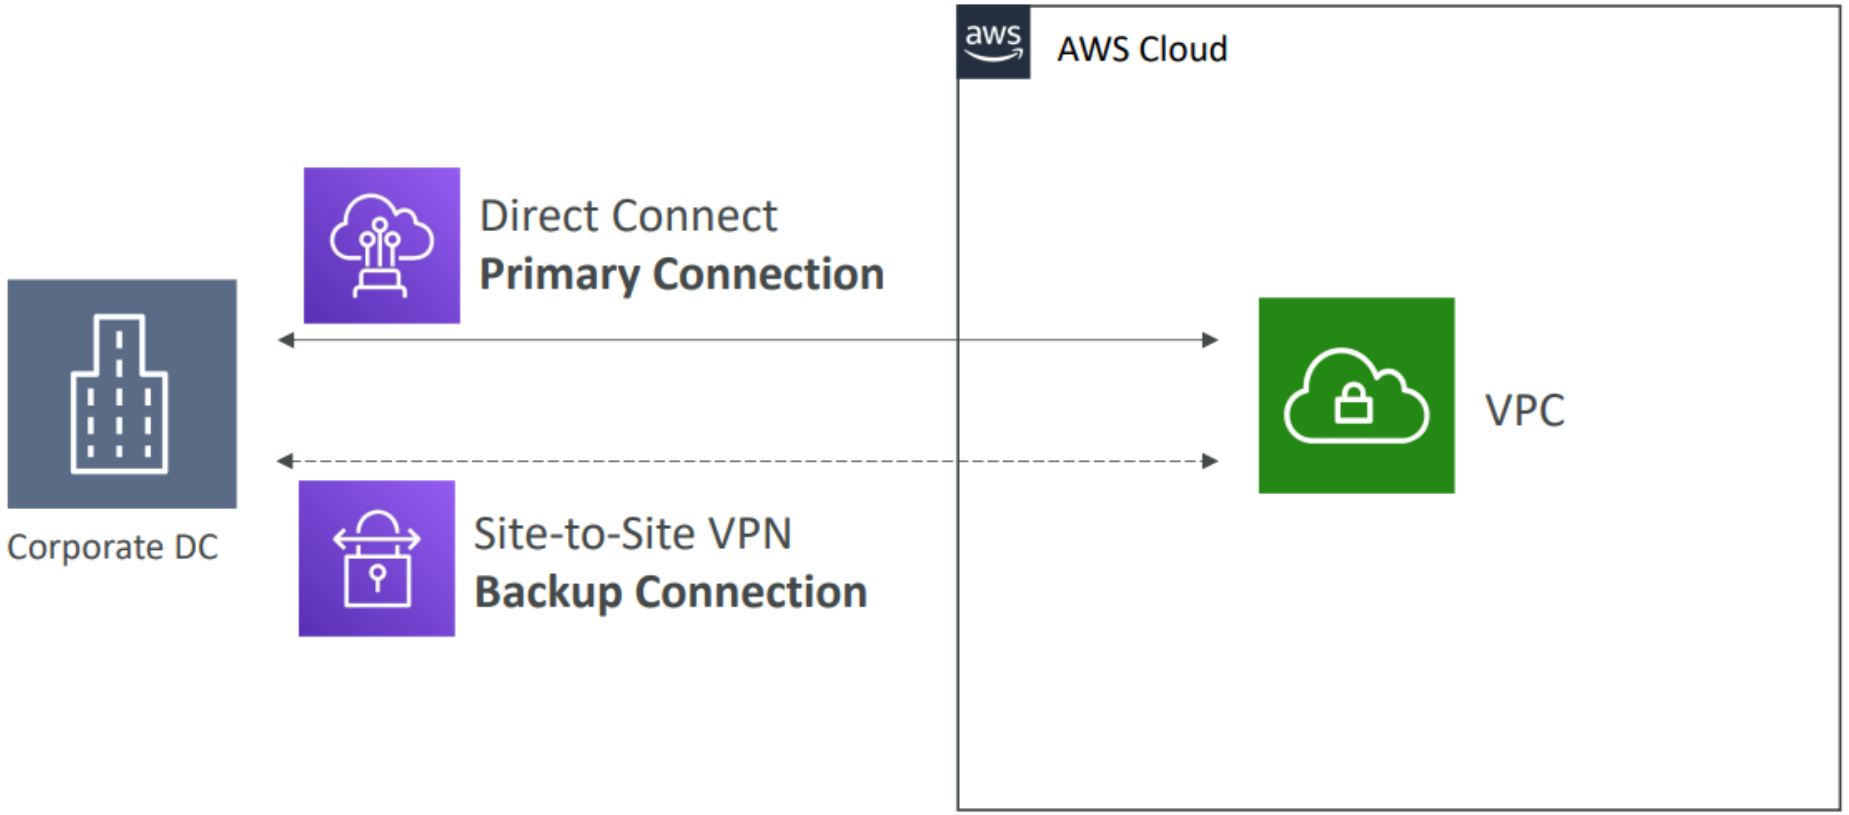

© Copyright by Cloud Mentor Pro | Confidential

VPC Section Summary (1/2)

- **CIDR** – IP Range
- **VPC** – Virtual Private Cloud
- **Subnets** – tied to an AZ, we define a CIDR
- **Internet Gateway** – at the VPC level, provide IPv4 & IPv6 Internet Access
- **Route Tables** – must be edited to add routes from subnets to the IGW, VPC Peering Connections, VPC Endpoints, …
- **Bastion Host** – public EC2 instance to SSH into, that has SSH connectivity to EC2 instances in private subnets
- **NAT Gateway** – managed by AWS, provides scalable Internet access to private EC2 instances, when the target is an IPv4 address

© Copyright by Cloud Mentor Pro | Confidential

VPC Section Summary (2/2)

- **NACL** – stateless, subnet rules for inbound and outbound, don't forget Ephemeral Ports
- **Security Groups** – stateful, operate at the EC2 instance level
- **VPC Peering** – connect two VPCs with non overlapping CIDR, non-transitive
- **VPC Endpoints** – provide private access to AWS Services (S3, DynamoDB, CloudFormation, SSM) within a VPC
- **VPC Flow Logs** – can be setup at the VPC / Subnet / ENI Level, for ACCEPT and REJECT traffic, helps identifying attacks, analyze using Athena or CloudWatch Logs Insights
- **Site-to-Site VPN** – VPN over public internet between on-premises DC and AWS
- **Direct Connect** – establish a direct private connection from On-premise to AWS

© Copyright by Cloud Mentor Pro | Confidential

S3 Data Transfer Pricing

https://blog.cloudmentor.pro/blog/aws/common-data-transfer-cost

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- *Networking - VPC*

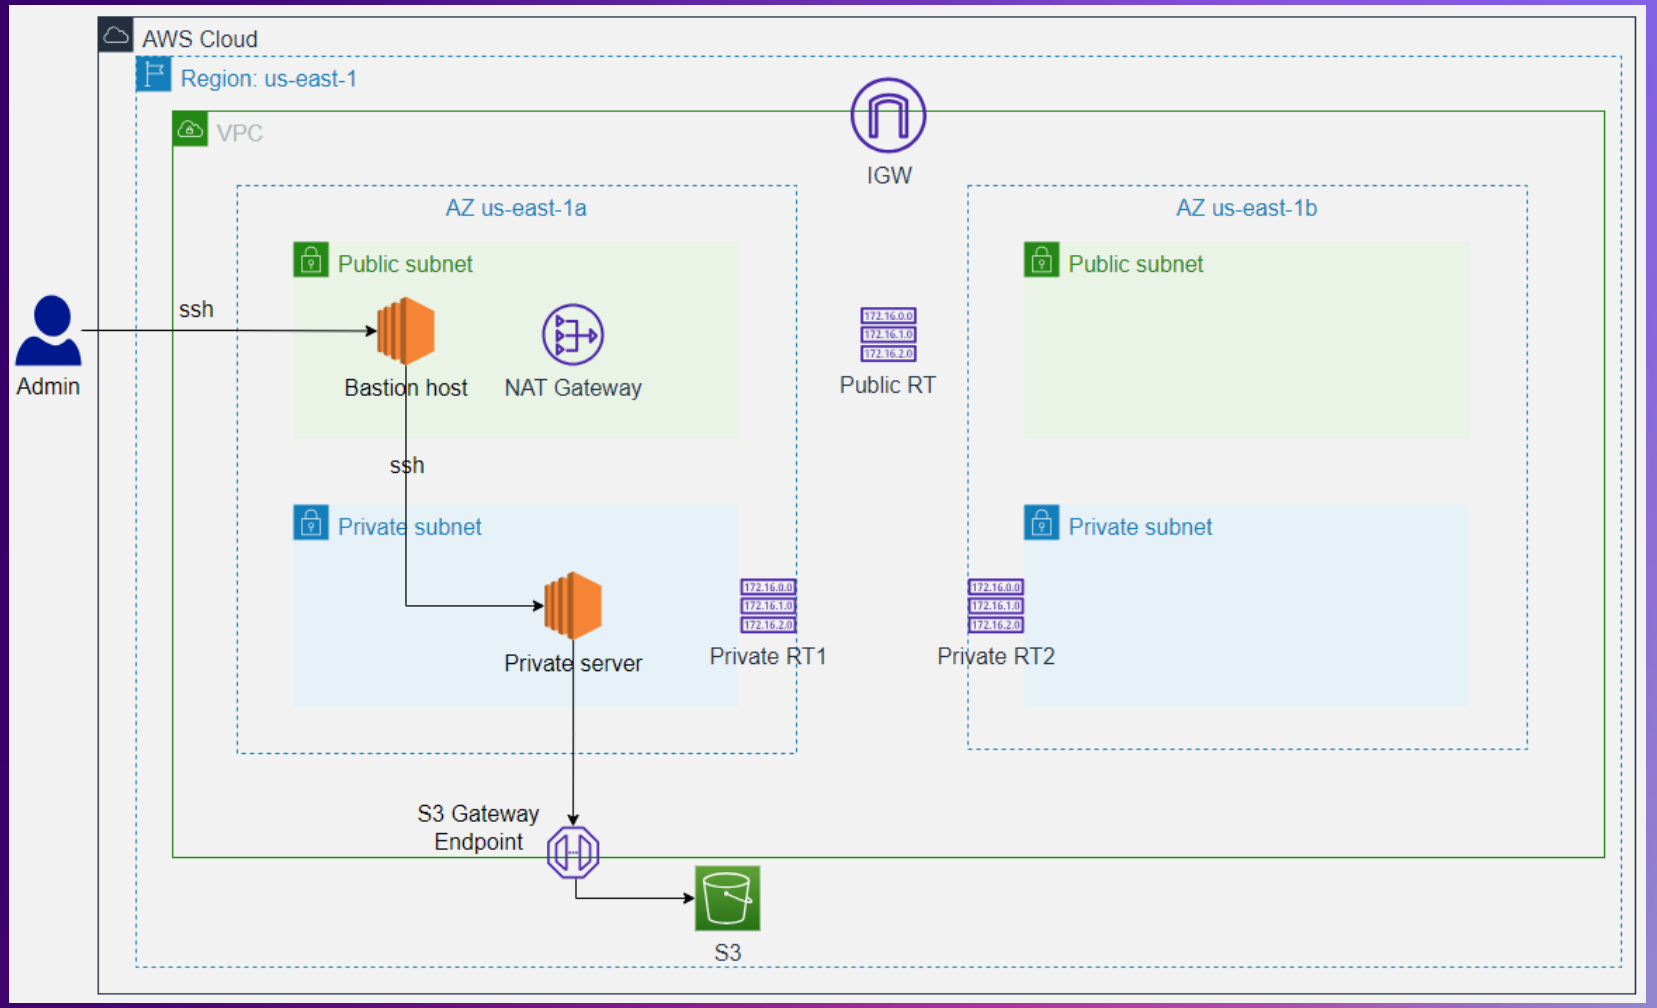

© Copyright by Cloud Mentor Pro | Confidential

Section 6

- *ECS, ECR & Fargate - Docker in AWS*

© Copyright by Cloud Mentor Pro | Confidential

Containers on AWS

ECS, Fargate, ECR & EKS

© Copyright by Cloud Mentor Pro | Confidential

What is Docker?

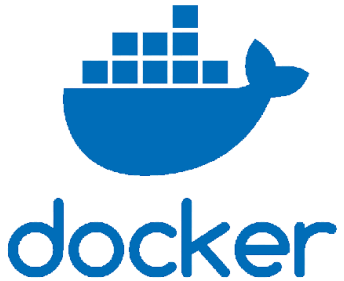

- Docker is a software development platform to deploy apps
- Apps are packaged in **containers** that can be run on any OS
- Apps run the same, regardless of where they're run
  - Any machine
  - No compatibility issues
  - Predictable behavior
  - Less work
  - Easier to maintain and deploy
  - Works with any language, any OS, any technology
- Use cases: microservices architecture, lift-and-shift apps from on-premises to the AWS cloud, …

© Copyright by Cloud Mentor Pro | Confidential

Docker on an OS

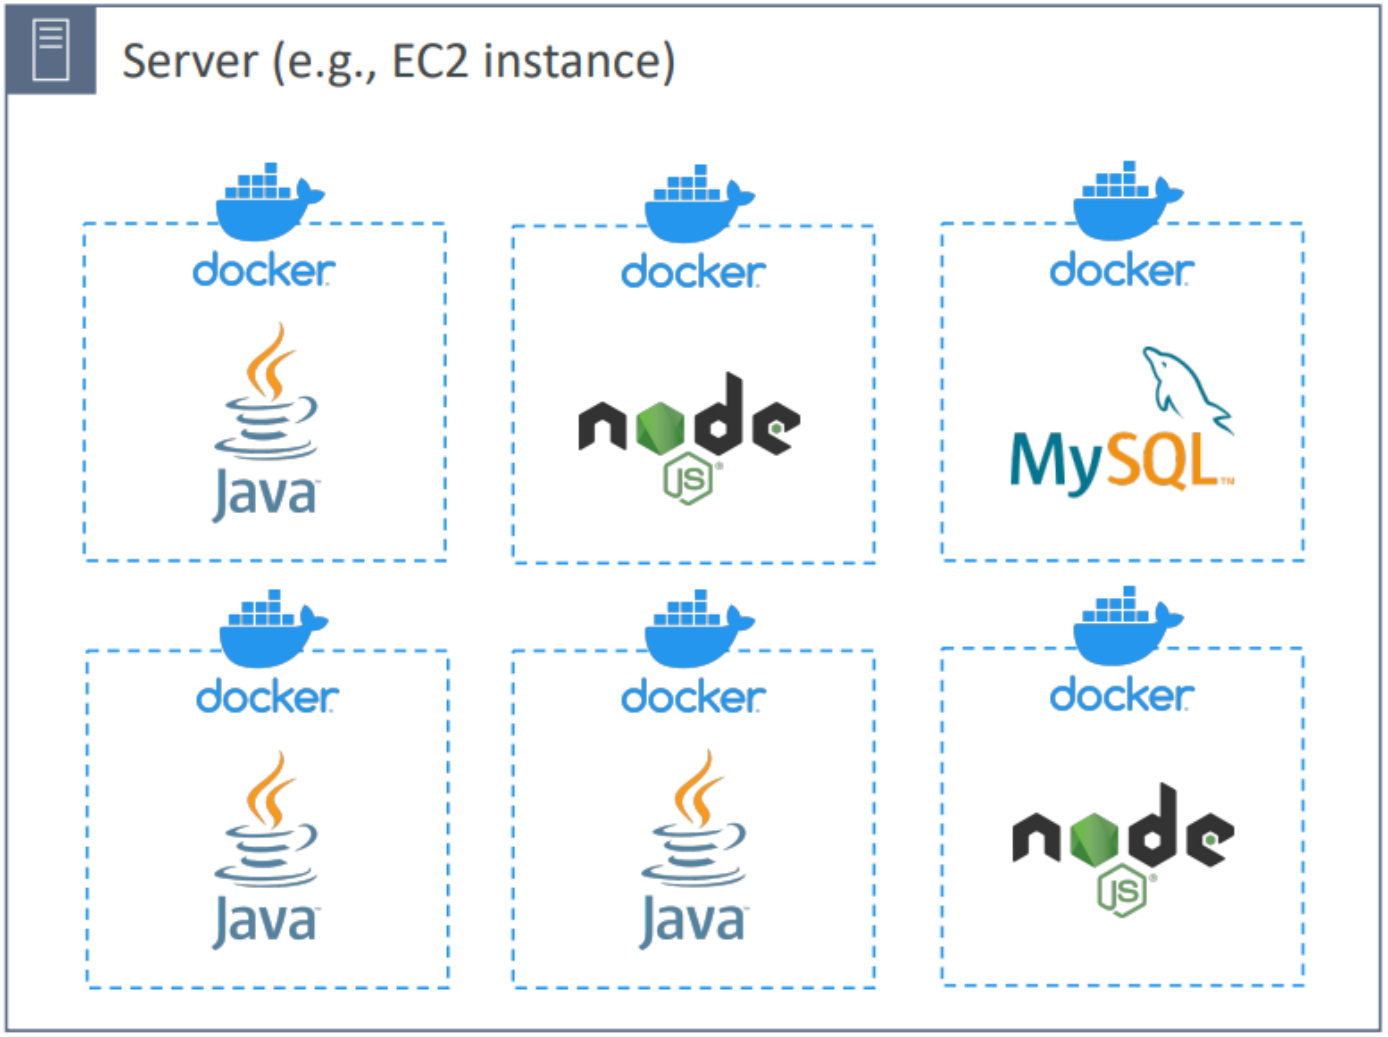

© Copyright by Cloud Mentor Pro | Confidential

Where are Docker images stored?

- Docker images are stored in Docker Repositories

- **Docker Hub** (https://hub.docker.com)
  - Public repository
  - Find base images for many technologies or OS (e.g., Ubuntu, MySQL, …)

- **Amazon ECR (Amazon Elastic Container Registry)**
  - Private repository
  - Public repository (**Amazon ECR Public Gallery** https://gallery.ecr.aws)

© Copyright by Cloud Mentor Pro | Confidential

Docker vs. Virtual Machines

- Docker is “sort of“ a virtualization technology, but not exactly
- Resources are shared with the host => many containers on one server

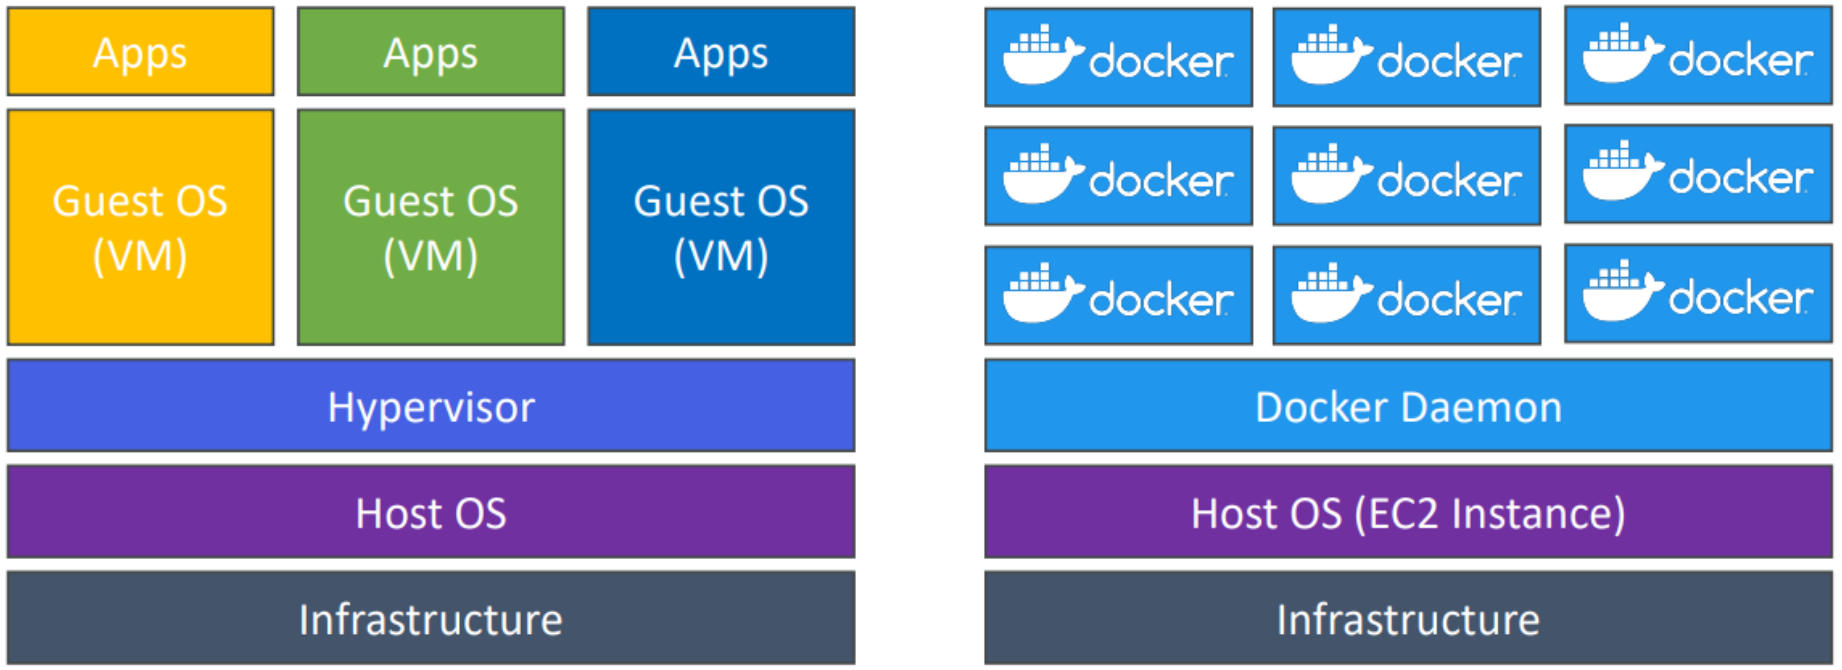

© Copyright by Cloud Mentor Pro | Confidential

Getting Started with Docker

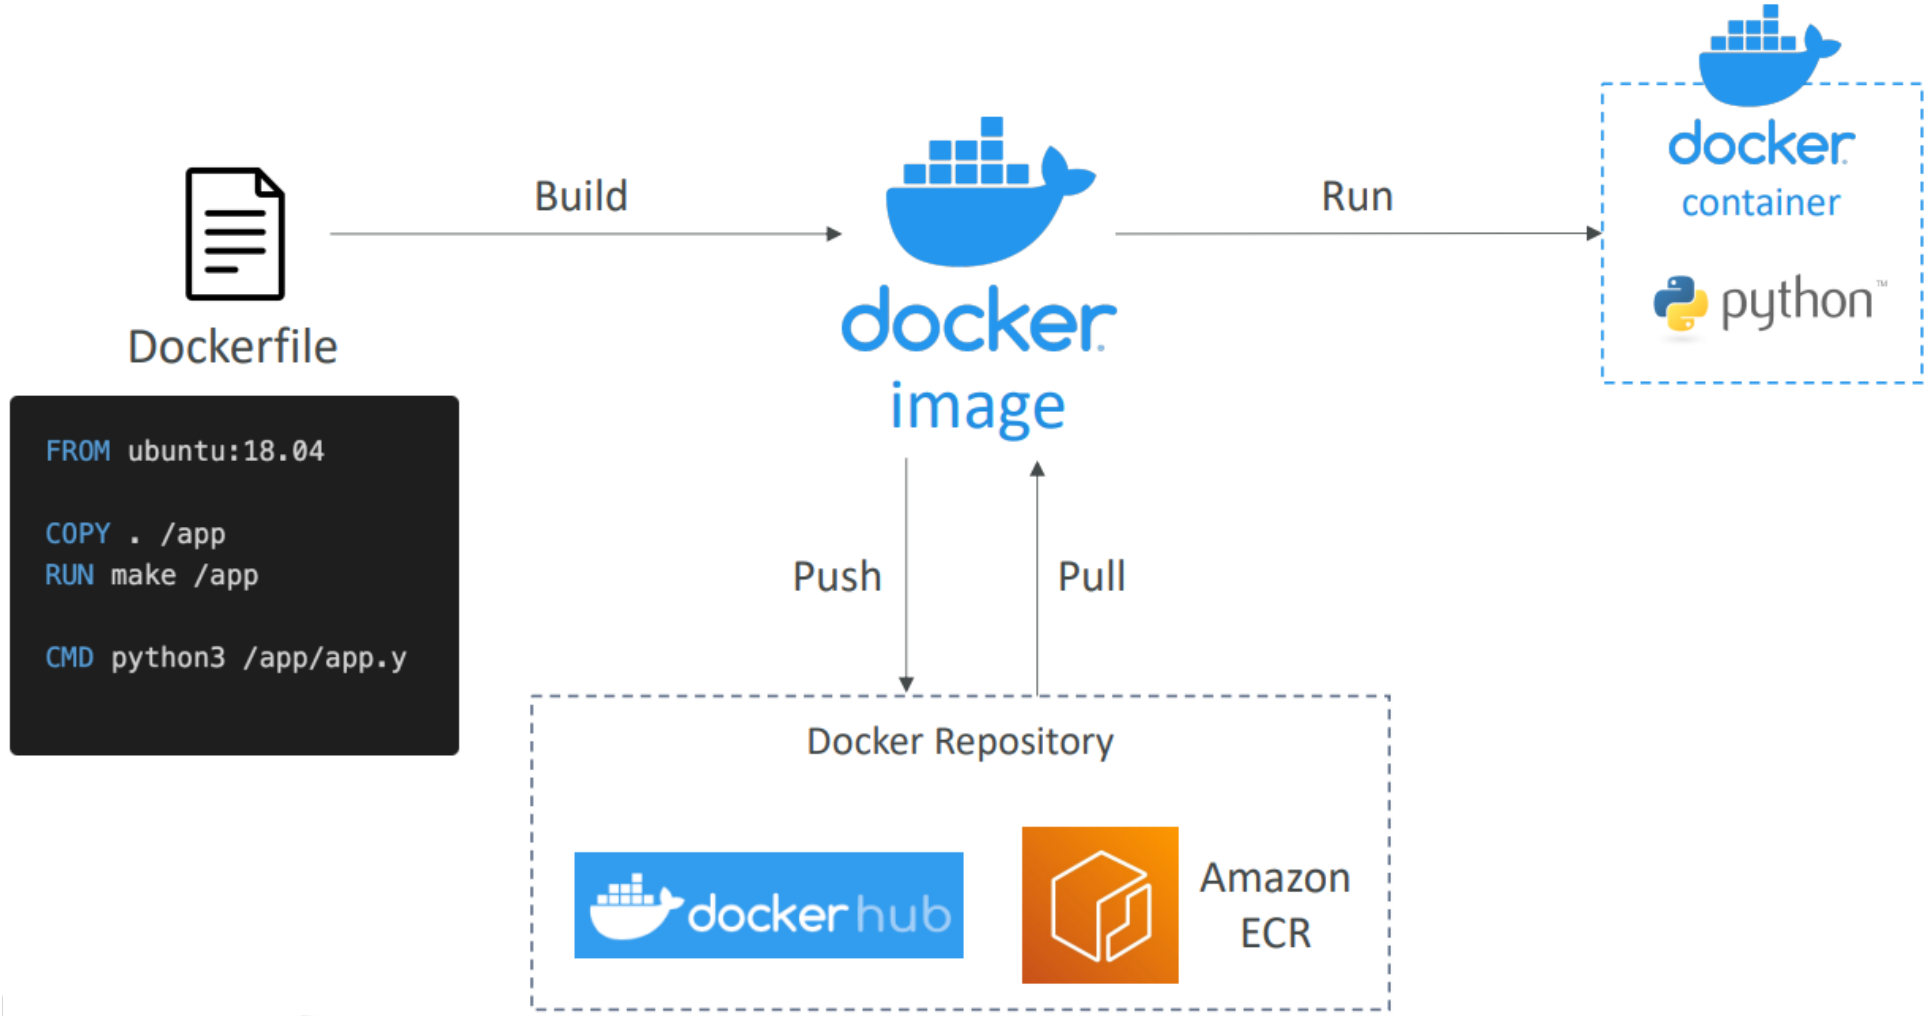

© Copyright by Cloud Mentor Pro | Confidential

Docker Containers Management on AWS

- Amazon Elastic Container Service (Amazon ECS)
  - Amazon's own container platform
- Amazon Elastic Kubernetes Service (Amazon EKS)
  - Amazon's managed Kubernetes (open source)

- AWS Fargate
  - Amazon's own Serverless container platform
  - Works with ECS and with EKS
- Amazon ECR:
  - Store container images

Amazon ECS

Amazon EKS

AWS Fargate

Amazon ECR

© Copyright by Cloud Mentor Pro | Confidential

ECS Cluster

- A Cluster is a group of services and tasks.
- You MUST create a cluster first!

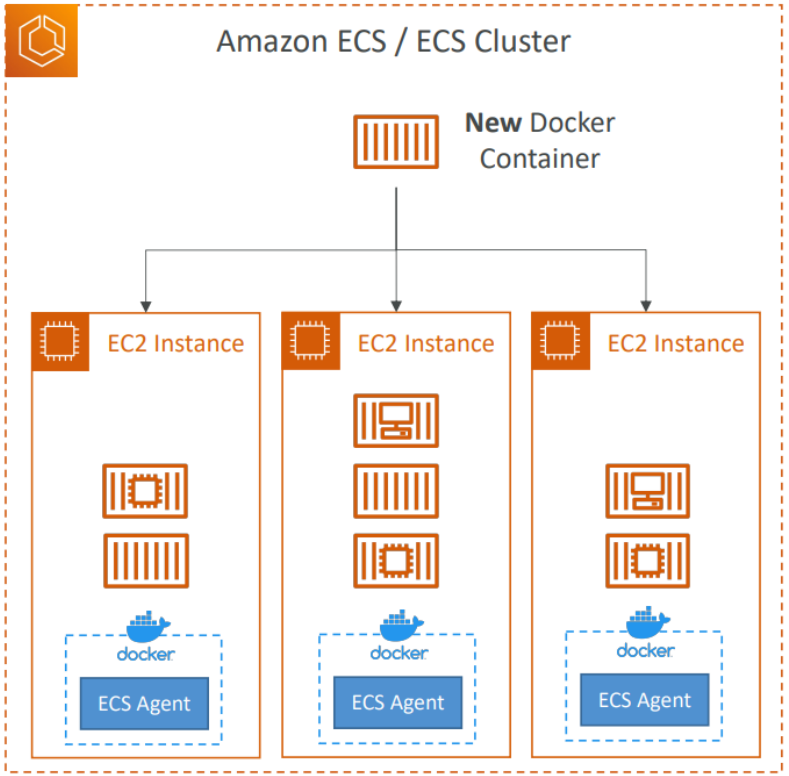

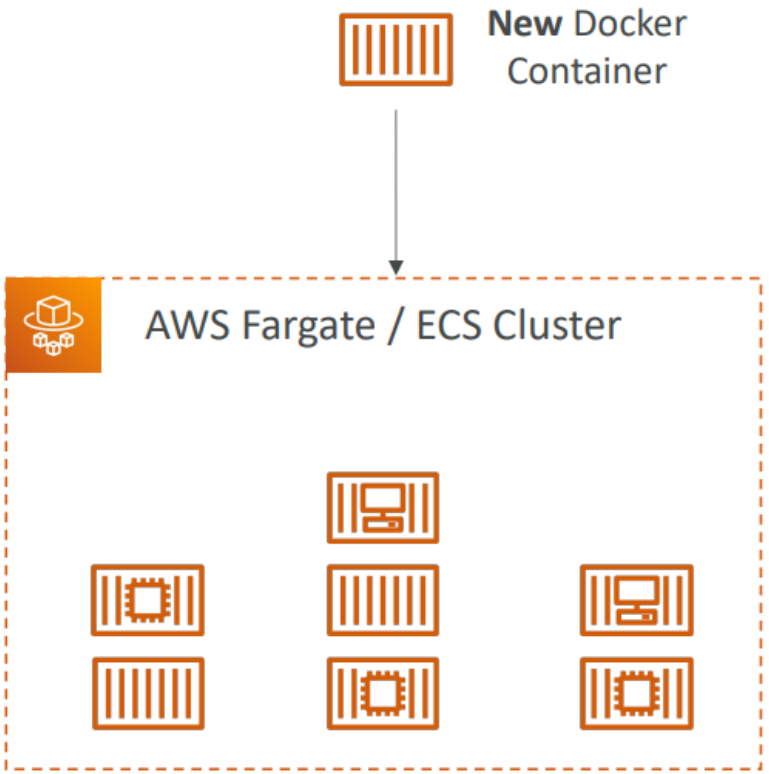

© Copyright by Cloud Mentor Pro | Confidential

ECS Task

- Basic unit of work in ECS
- One or more containers per Task
- Running container with settings specified by Task Definition
- Sidecar container

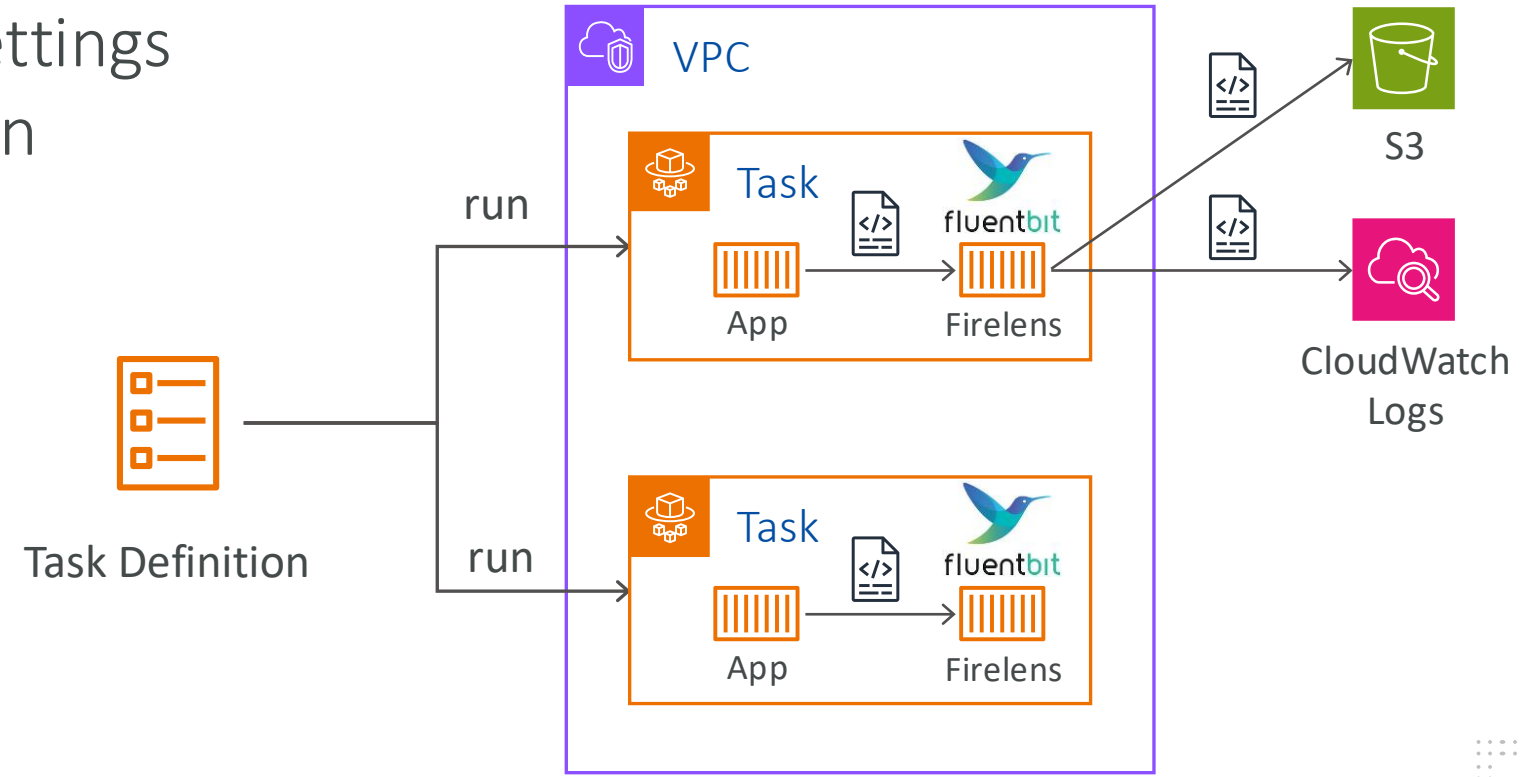

© Copyright by Cloud Mentor Pro | Confidential

ECS Task in comparison to EC2

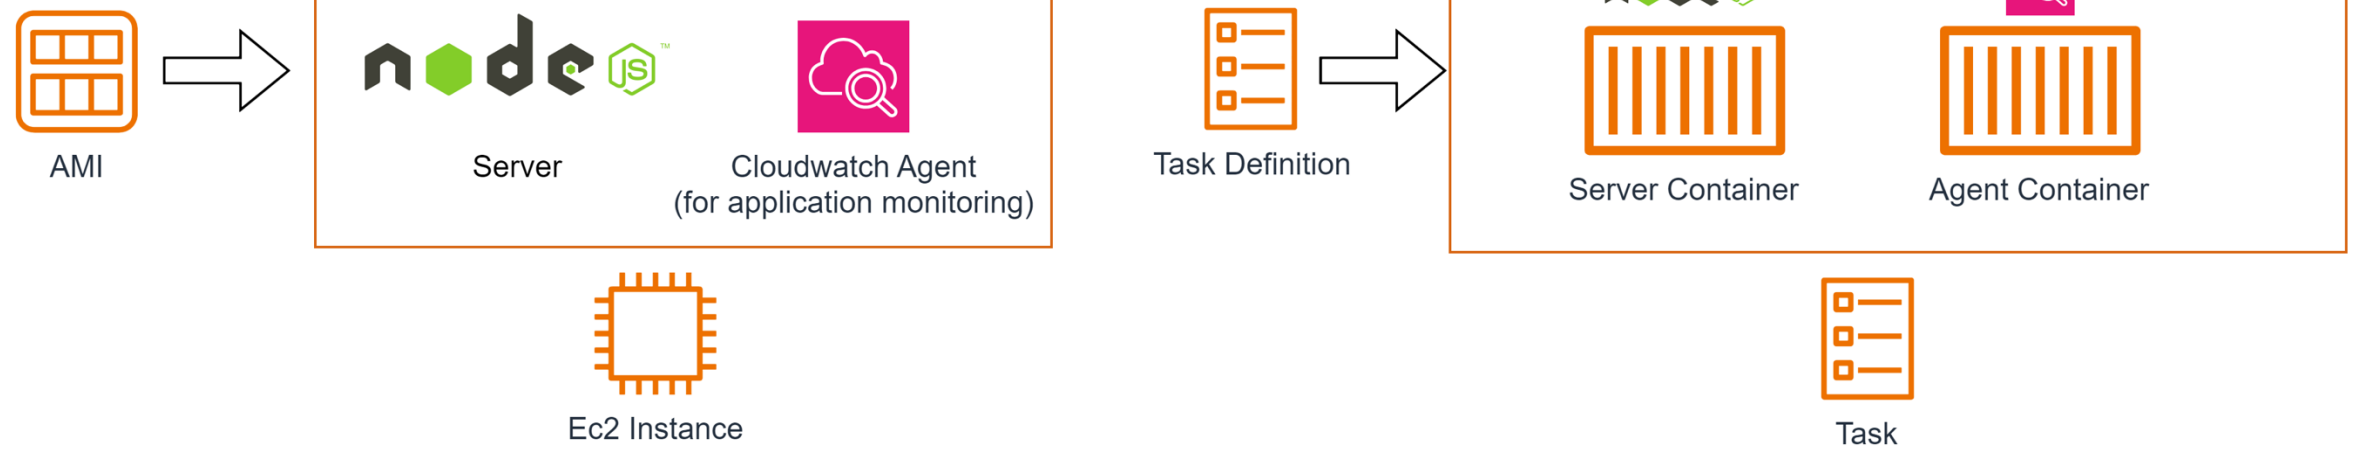

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS –Task Definitions

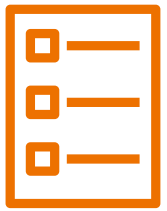

- Task definitions are metadata in JSON form to tell ECS how to run a Docker container
- It contains crucial information, such as:
  - Image Name
  - Port Binding for Container and Host
  - Memory and CPU required
  - Environment variables
  - Networking information
  - IAM Role
  - Logging configuration (ex CloudWatch)
- Up to 10 containers in a Task Definition

```
{
  "family": "webserver",
  "containerDefinitions": [
    {
      "name": "nginx",
      "image": "nginx:latest",
      "memoryReservation": "300",
      "cpu": "256",
      "essential": true,
      "portMappings": [
        {
          "hostPort": "80",
          "containerPort": "80",
          "protocol": "tcp"
        }
      ],
      "logConfiguration": {
        "logDriver": "awslogs",
        "options": {
          "awslogs-group": "/ecs/webserver",
          "awslogs-region": "us-east-1",
          "awslogs-stream-prefix": "ecs"
        }
      }
    }
  ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS
One IAM Role per Task Definition

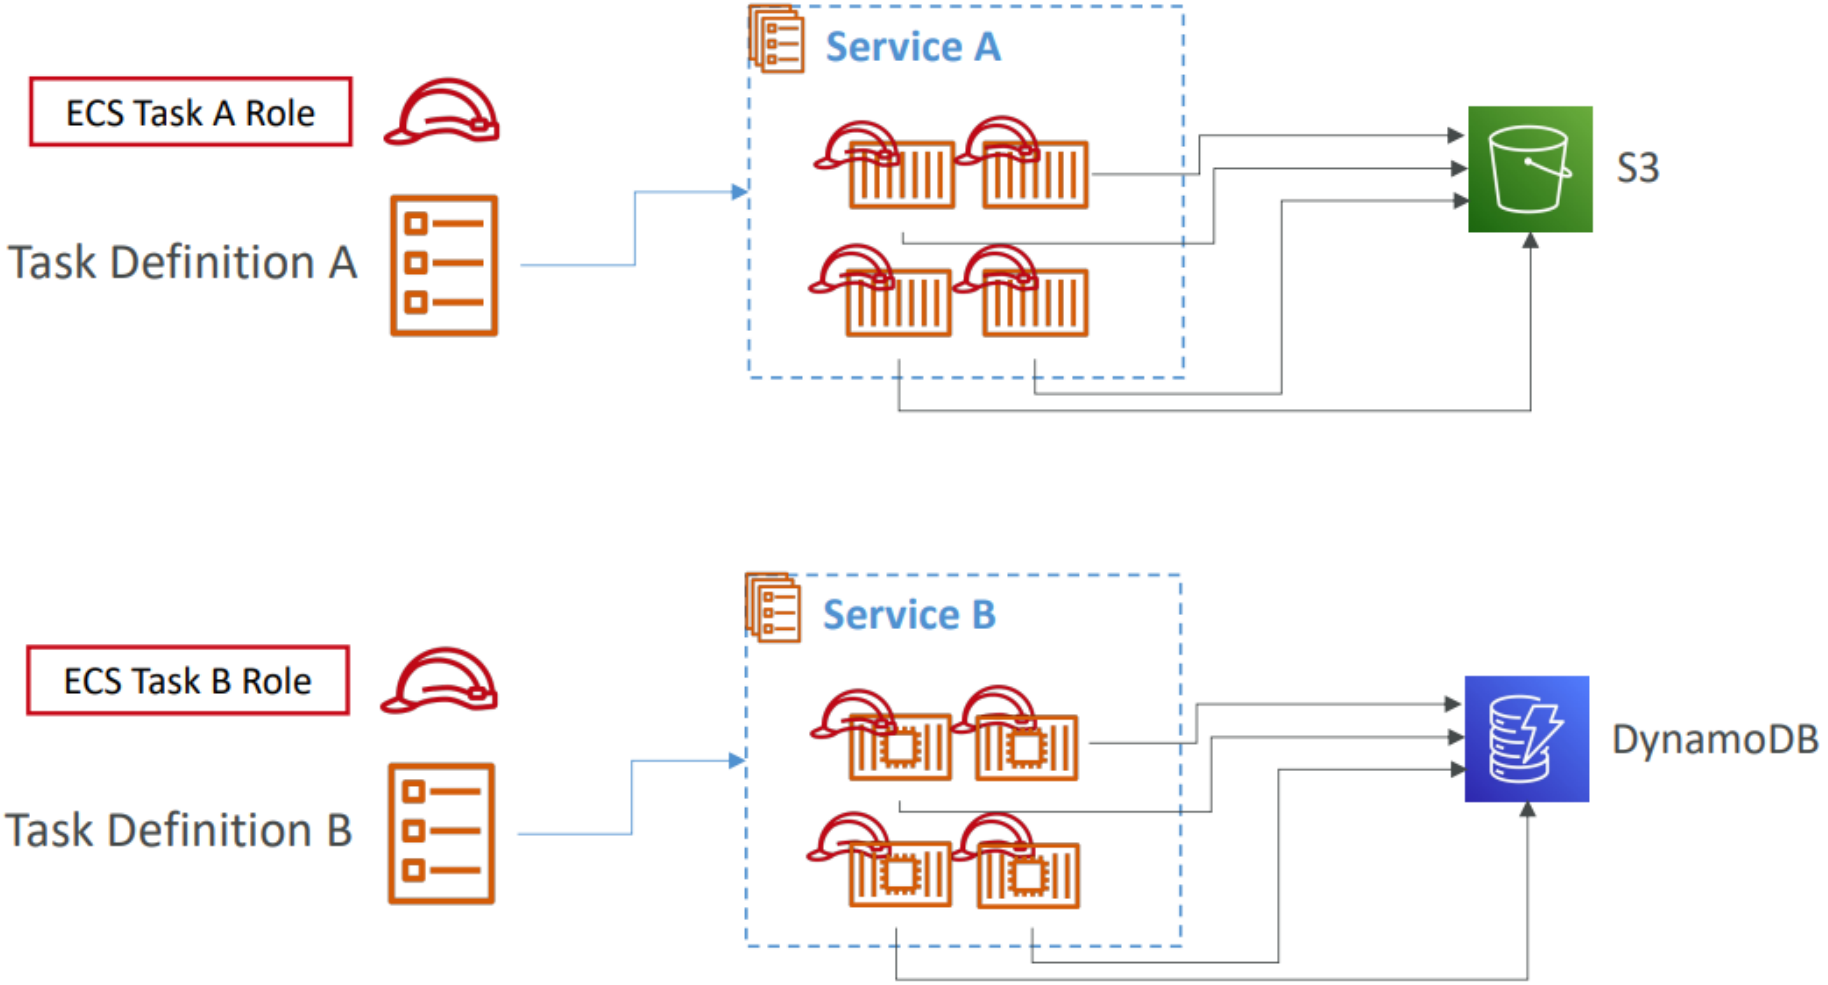

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS - EC2 Launch Type

- ECS tasks are launched in your EC2 instances
- You must provision & maintain the infrastructure (the EC2 instances
- Each EC2 Instance must run the ECS Agent to register in the ECS Cluster
- AWS takes care of starting / stopping containers

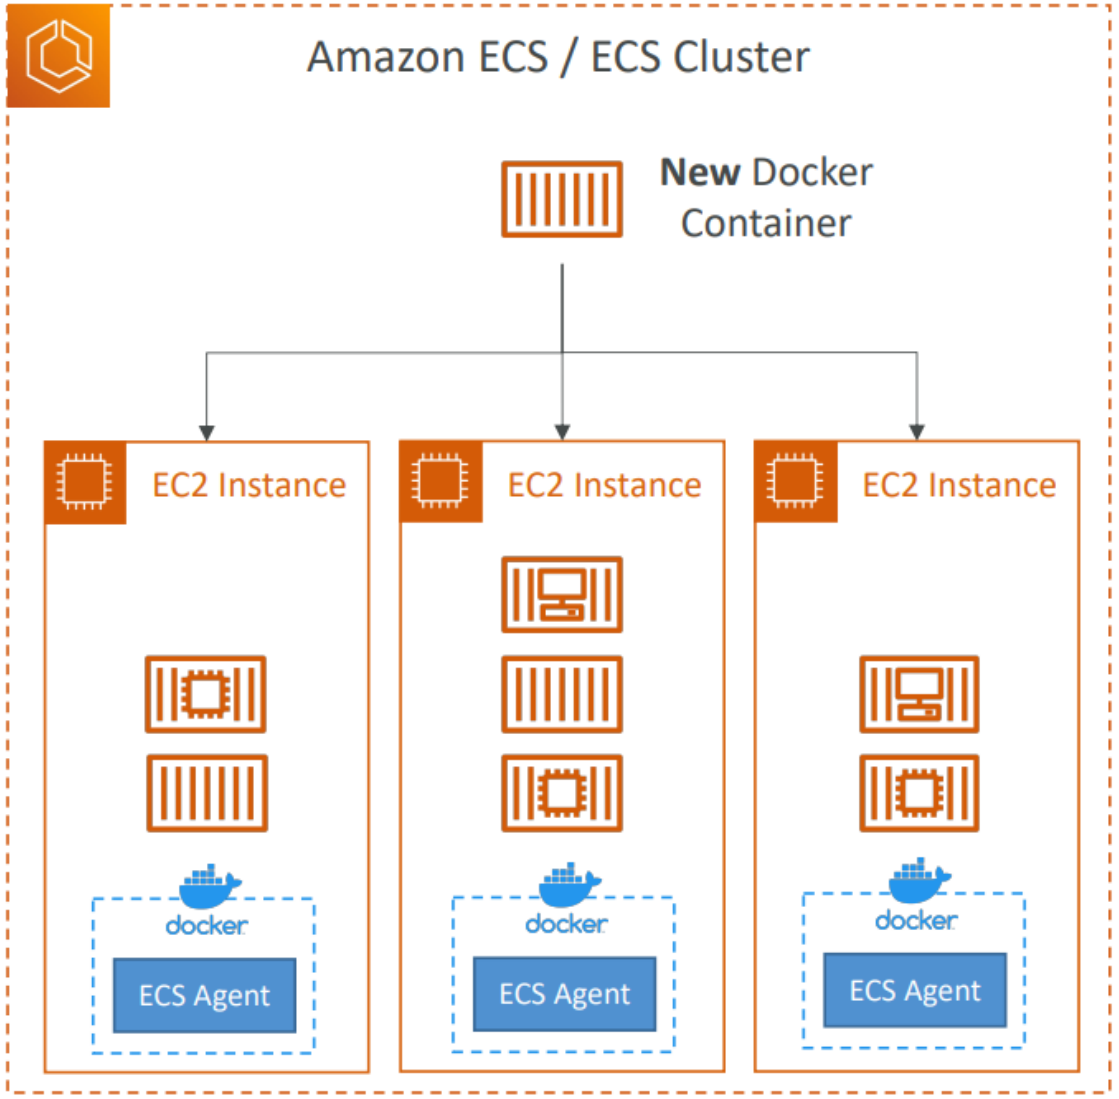

© Copyright by Cloud Mentor Pro | Confidential

ECS Agent

- ECS agent is installed on each of EC2(Container) Instances
- ECS agent handles incoming requests for container deployment
- ECS agent handles the lifecycle of container
- **ECS agent is installed by default when using ECS AMI**

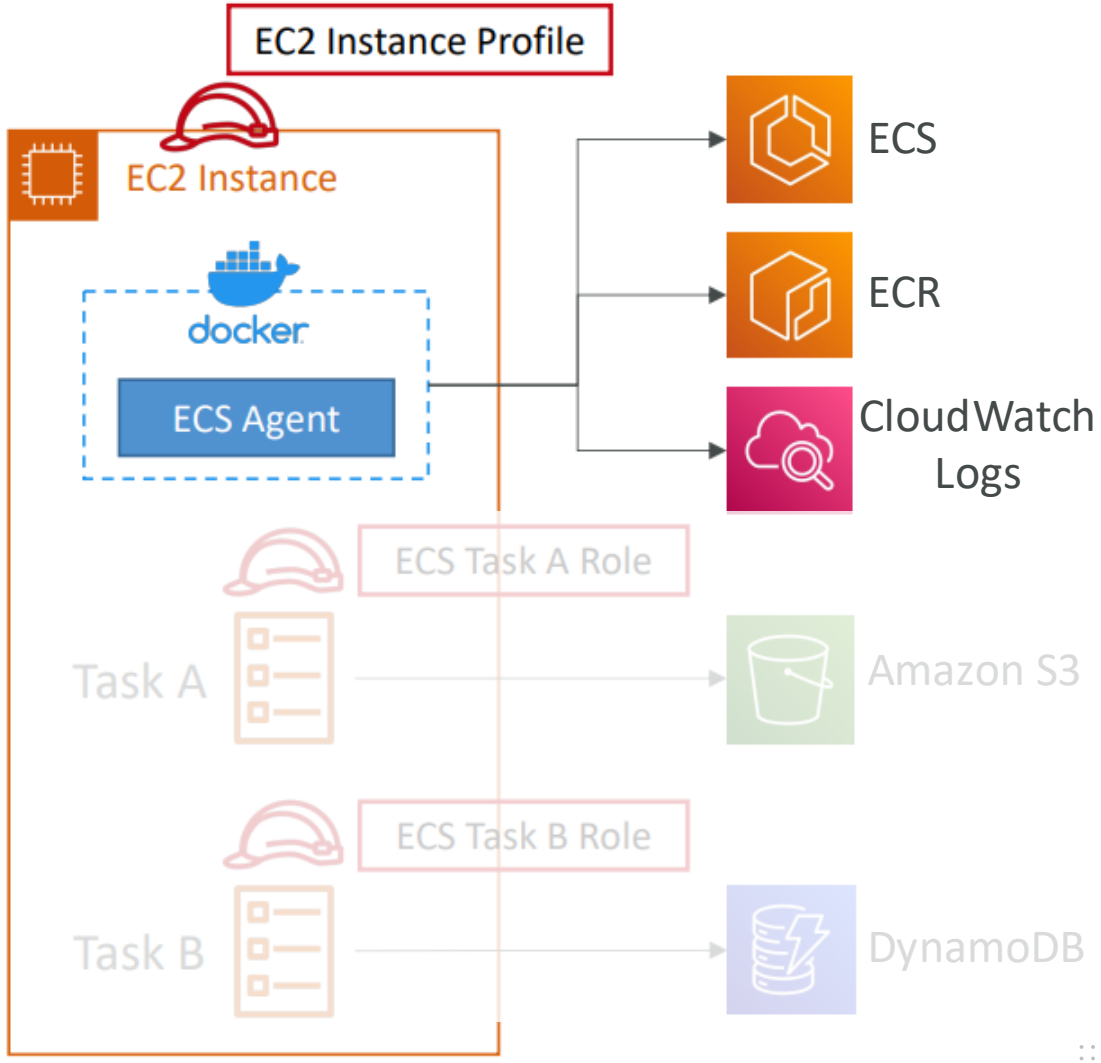

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – IAM Roles for ECS

- EC2 Instance Profile (EC2 Launch Type only):
  - Used by the ECS agent
  - Makes API calls to ECS service
  - Send container logs to CloudWatch Logs
  - Pull Docker image from ECR
  - Reference sensitive data in Secrets Manager or SSM Parameter Store
- ECS Task Role:
  - Allows each task to have a specific role
  - Use different roles for the different ECS Services you run
  - Task Role is defined in the task definition

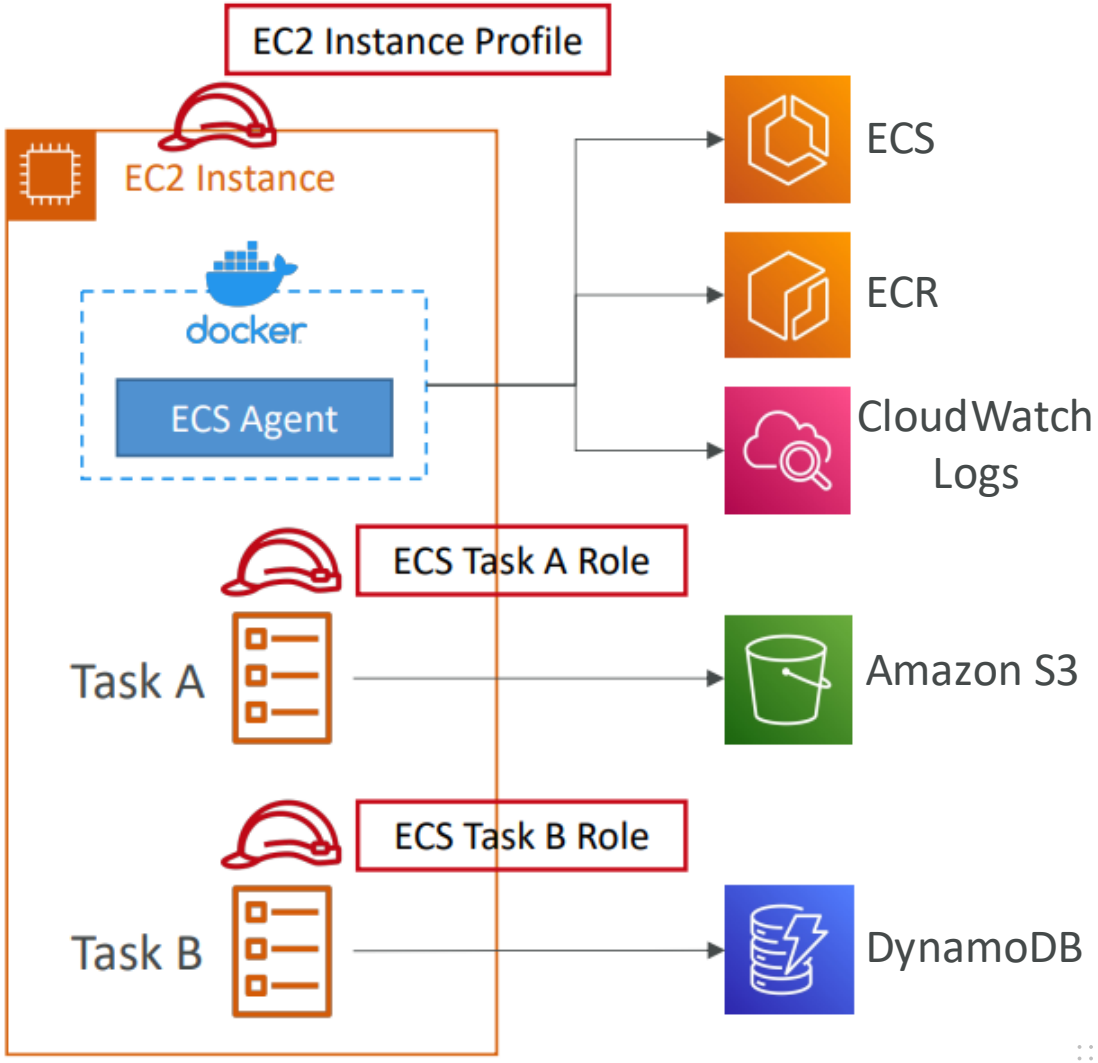

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Fargate Launch Type

- ECS tasks are launched to Fargate (Managed by AWS)
- You do not provision the infrastructure (no EC2 instances to manage)
- It's all Serverless!
- You just create task definitions
- AWS just runs ECS Tasks for you based on the CPU / RAM you need
- To scale, just increase the number of tasks. Simple - no more EC2 instances

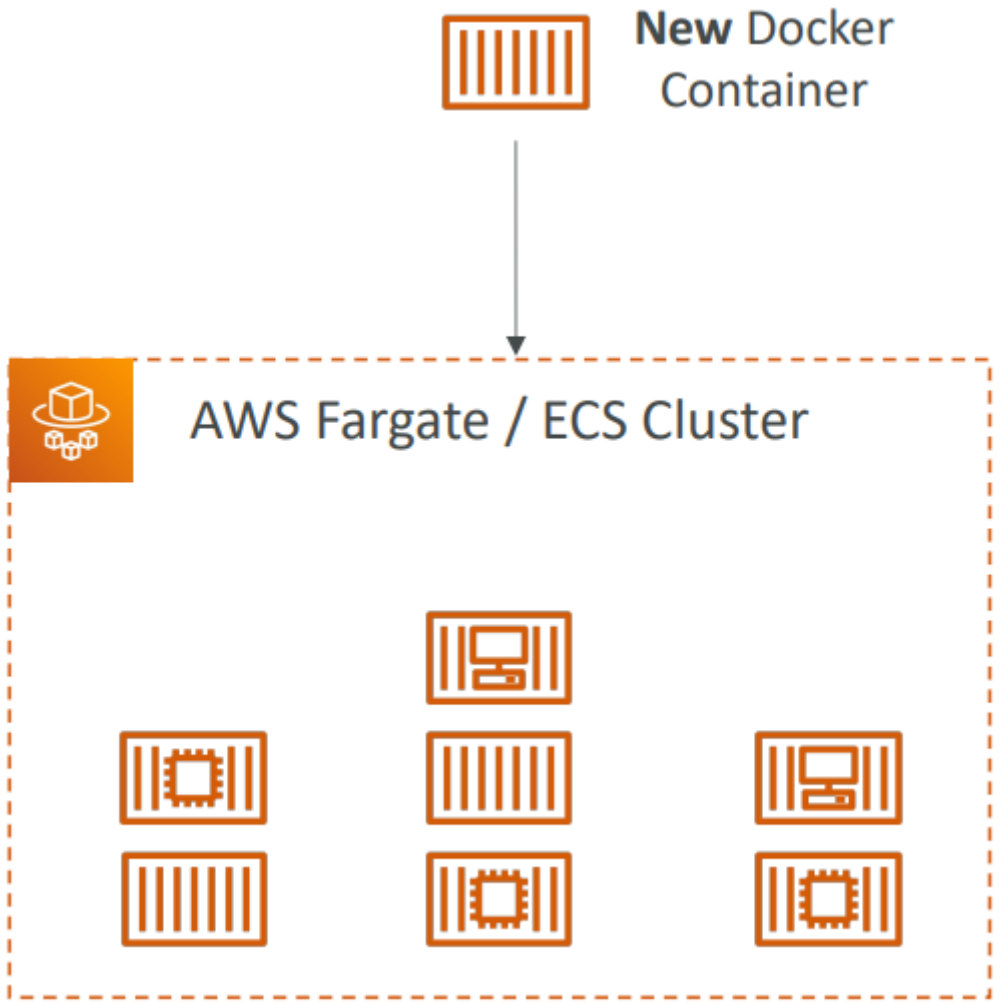

© Copyright by Cloud Mentor Pro | Confidential

ECS Service

- Service supervise task. Its job is keep task running
- Launch instance to maintain scheduling strategy
- Expose tasks to outside world
- Direct network traffic to the correct host and port

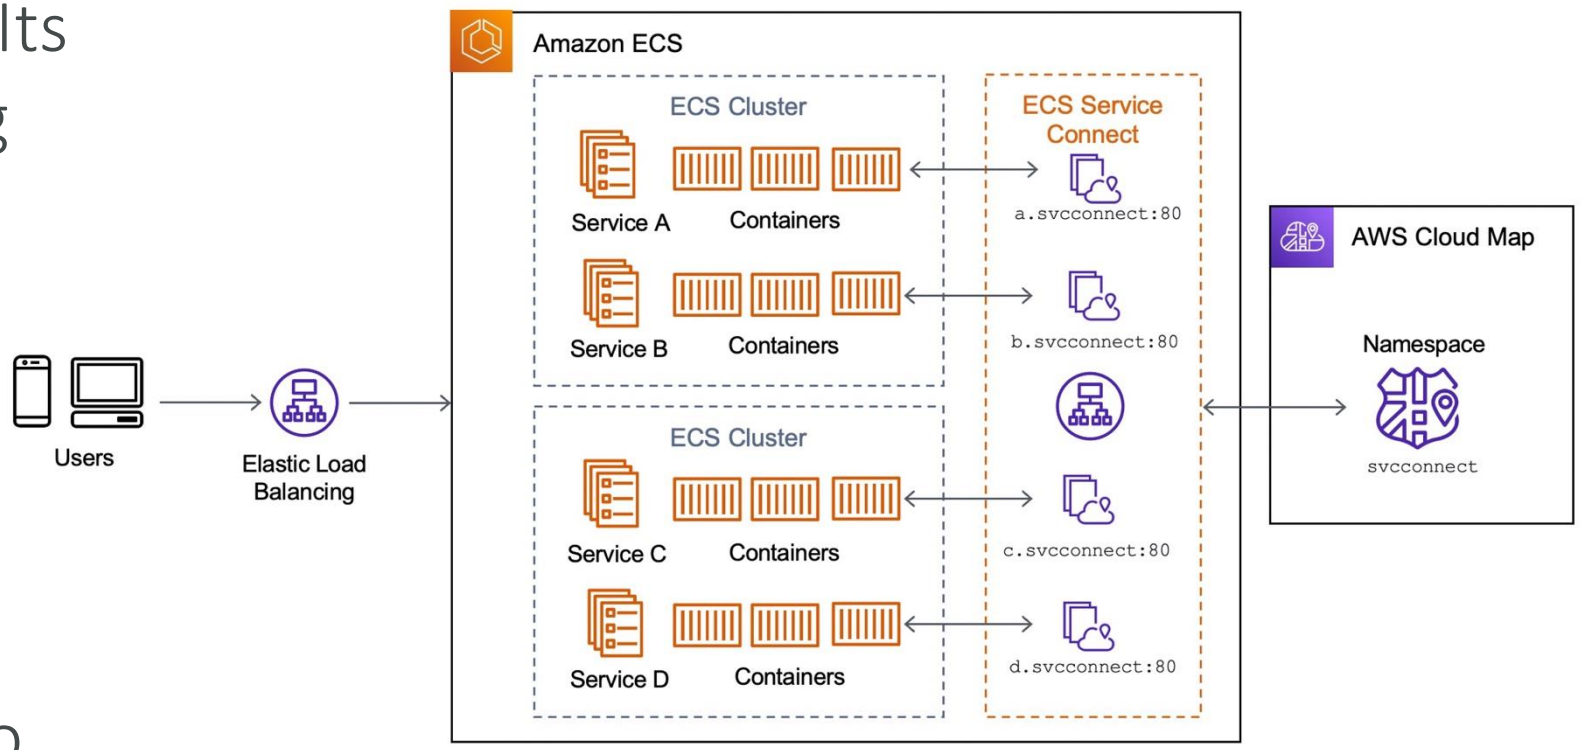

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Load Balancer Integrations

- Application Load Balancer supported and works for most use cases
- Network Load Balancer recommended only for high throughput / high performance use cases, or to pair it with AWS Private Link
- Classic Load Balancer supported but not recommended (no advanced features – no Fargate)

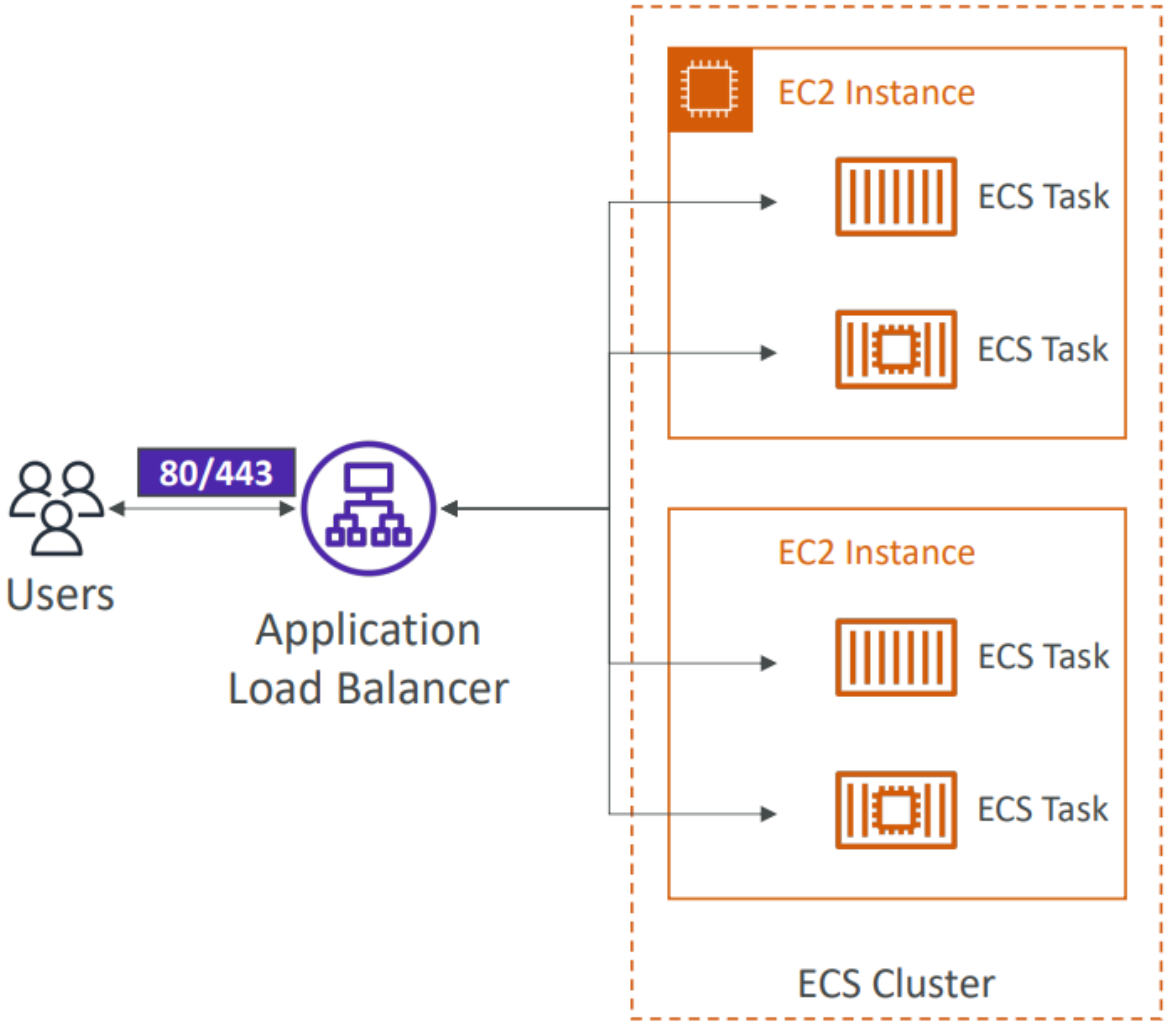

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Data Volumes (EFS)

- Mount EFS file systems onto ECS tasks
- Works for both **EC2** and **Fargate** launch types
- Tasks running in any AZ will share the same data in the EFS file system
- **Fargate + EFS = Serverless**

- Use cases: persistent multi-AZ shared storage for your containers
- Note:
  - Amazon S3 cannot be mounted as a file system

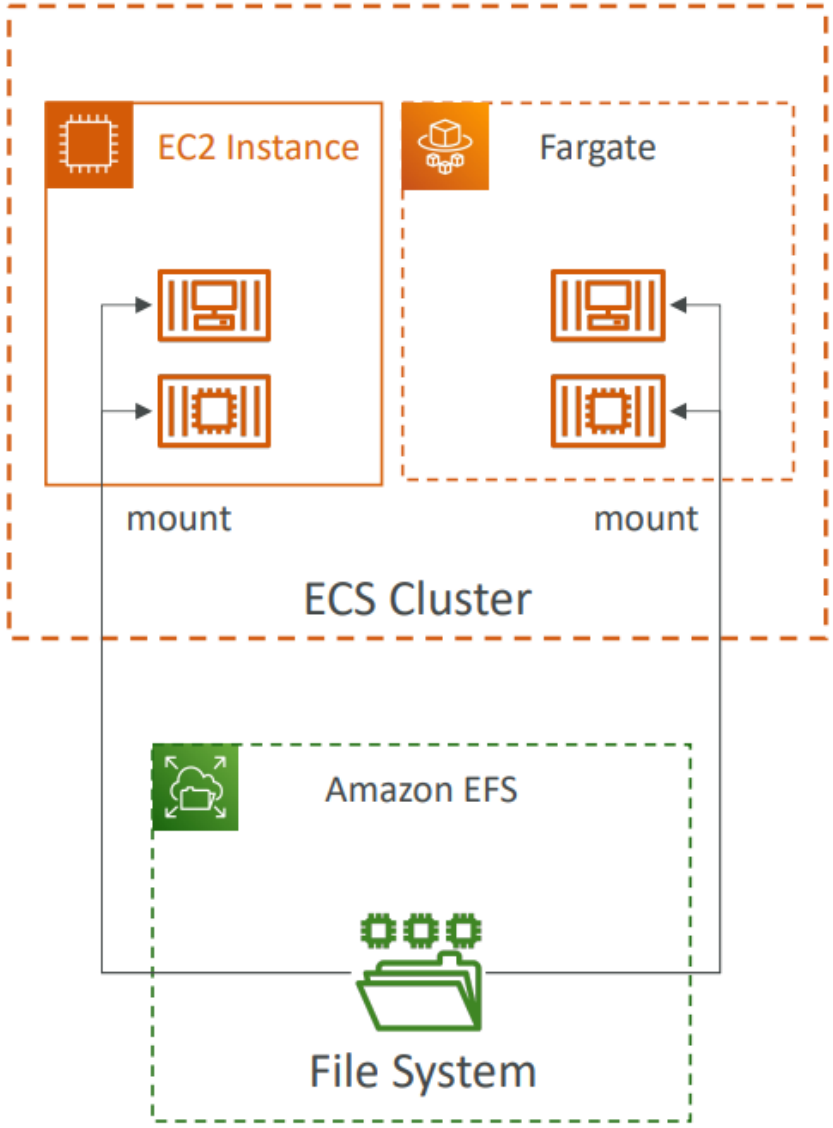

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Data Volumes (Bind Mounts)

- Share data between multiple containers in the same Task Definition (Work for both **EC2** and **Fargate**)
- EC2 Tasks – using EC2 instance storage
  - Data are tied to the lifecycle of the EC2 instance
- Fargate Tasks – using ephemeral storage
  - Data are tied to the container(s) using them
  - 20 GiB – 200 GiB (default 20 GiB)
- Use cases:
  - Share ephemeral data between multiple containers
  - “Sidecar” container pattern, where the “sidecar” container used to send metrics/logs to other destinations (separation of conerns)

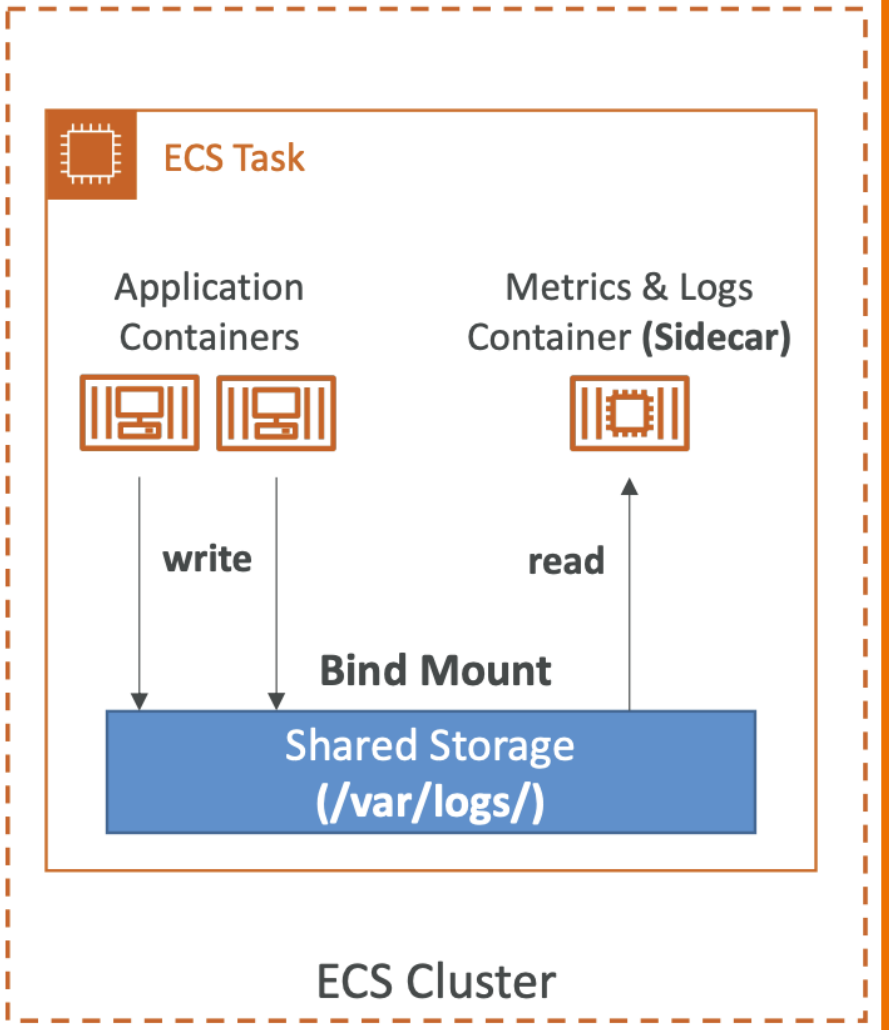

© Copyright by Cloud Mentor Pro | Confidential

ECS Service Auto Scaling

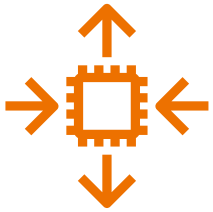

- Automatically increase/decrease the desired number of ECS tasks
- Amazon ECS Auto Scaling uses **AWS Application Auto Scaling**
  - ECS Service Average CPU Utilization
  - ECS Service Average Memory Utilization - Scale on RAM
  - ALB Request Count Per Target – metric coming from the ALB

- **Target Tracking** – scale based on target value for a specific CloudWatch metric
- **Step Scaling** – scale based on a specified CloudWatch Alarm
- **Scheduled Scaling** – scale based on a specified date/time (predictable changes)

- ECS Service Auto Scaling (task level) ≠ EC2 Auto Scaling (EC2 instance level)
- Fargate Auto Scaling is much easier to setup (because **Serverless**)

© Copyright by Cloud Mentor Pro | Confidential

EC2 Launch Type – Auto Scaling EC2 Instances

- Accommodate ECS Service Scaling by adding underlying EC2 Instances

- **Auto Scaling Group Scaling**
  - Scale your ASG based on CPU Utilization
  - Add EC2 instances over time

- **ECS Cluster Capacity Provider**
  - Used to automatically provision and scale the infrastructure for your ECS Tasks
  - Capacity Provider paired with an Auto Scaling Group
  - Add EC2 Instances when you're missing capacity (CPU, RAM…)

© Copyright by Cloud Mentor Pro | Confidential

ECS Scaling – Service CPU Usage Example

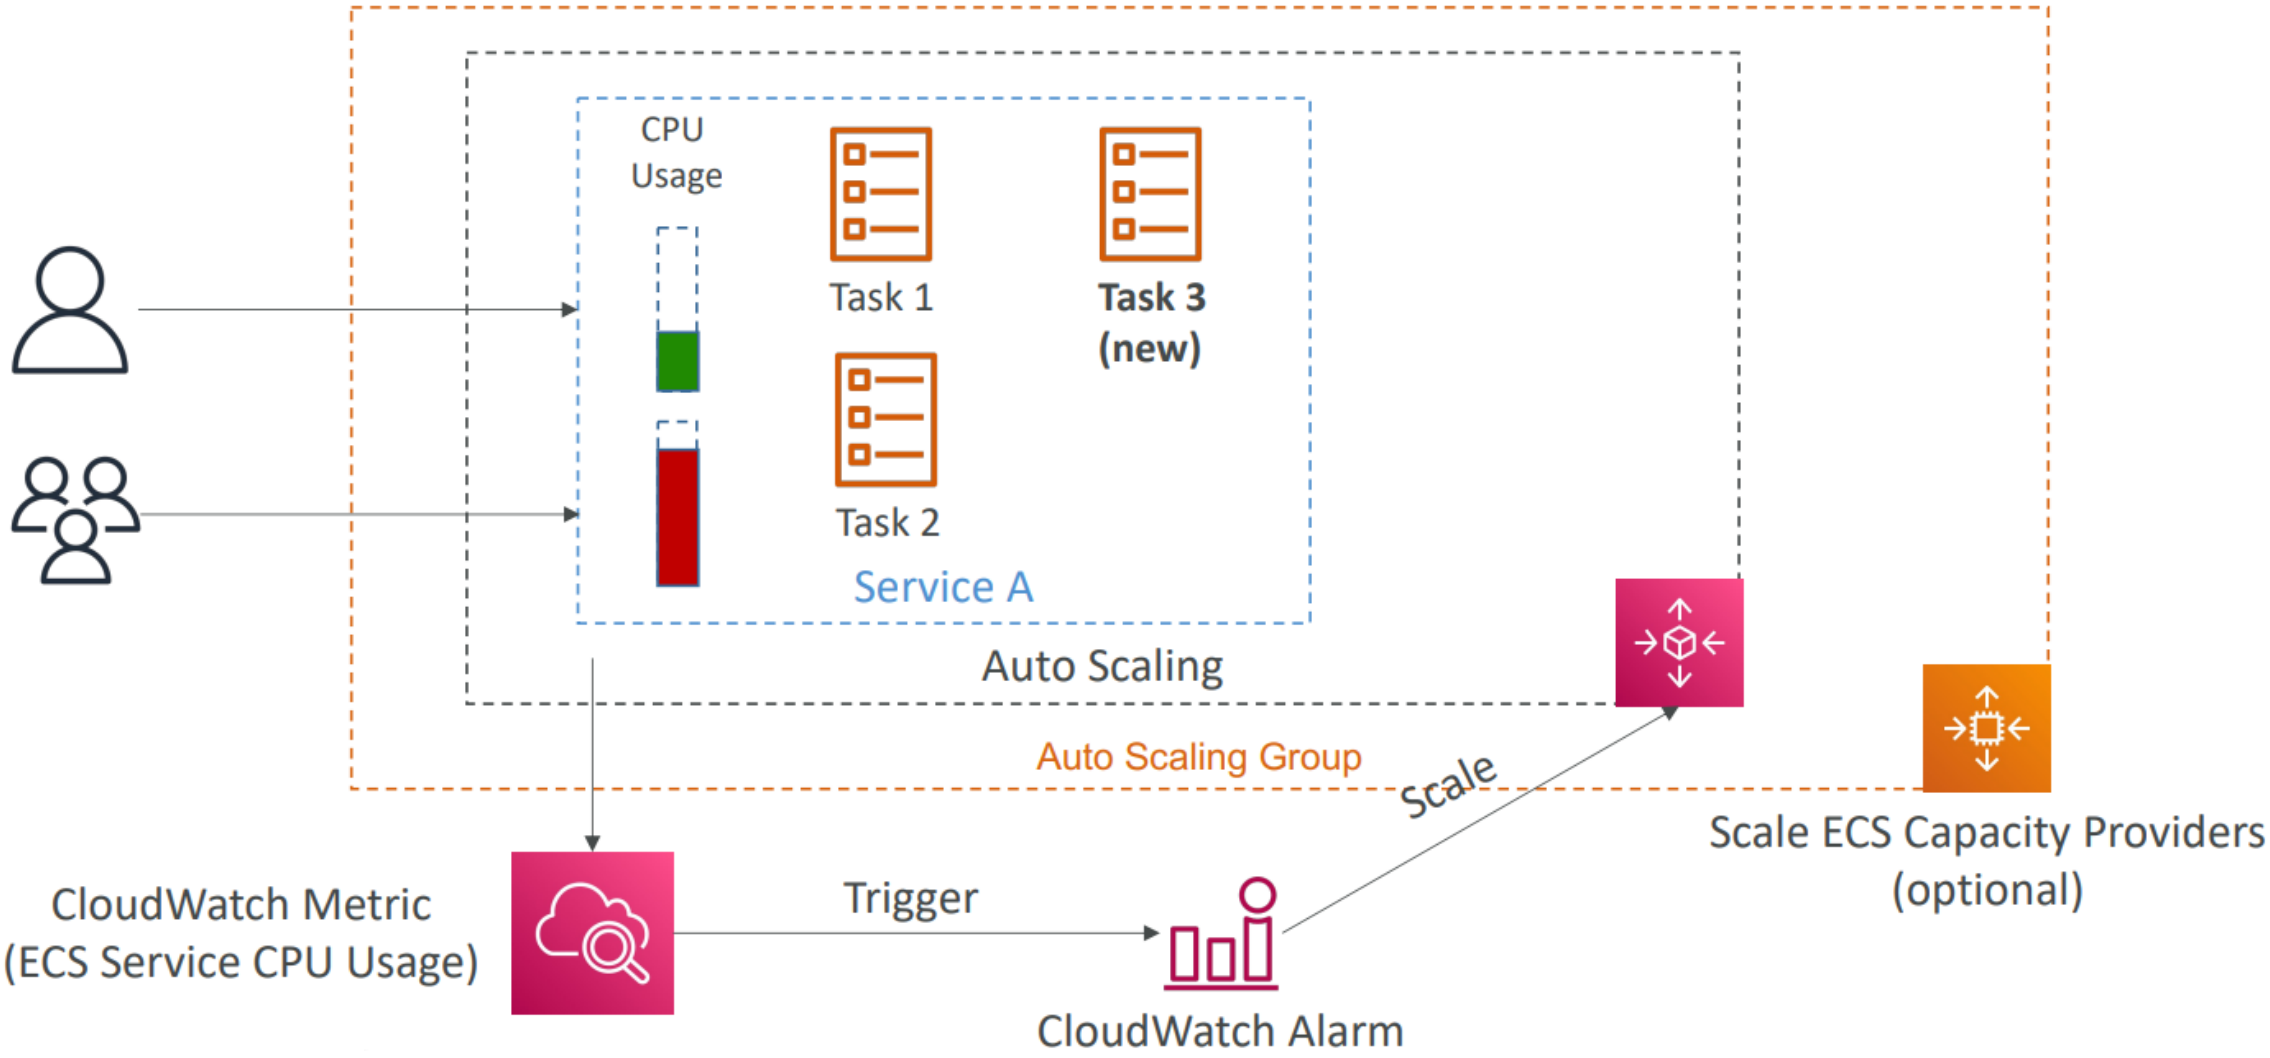

© Copyright by Cloud Mentor Pro | Confidential

ECS Rolling Updates

- When updating from v1 to v2, we can control how many tasks can be started and stopped, and in which order

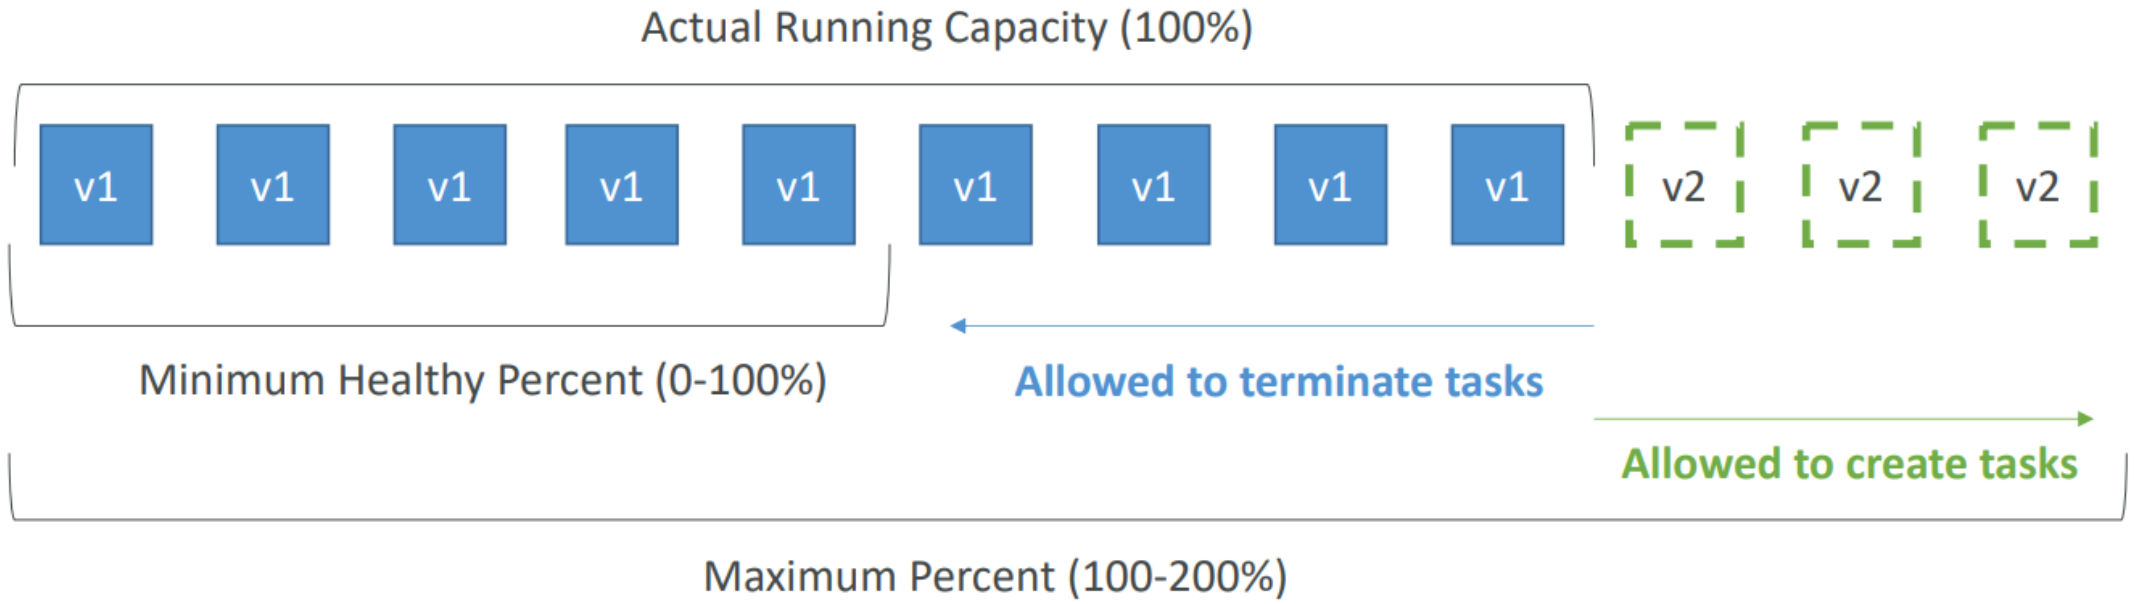

© Copyright by Cloud Mentor Pro | Confidential

ECS Rolling Update – Min 50%, Max 100%

- Starting number of tasks: 4

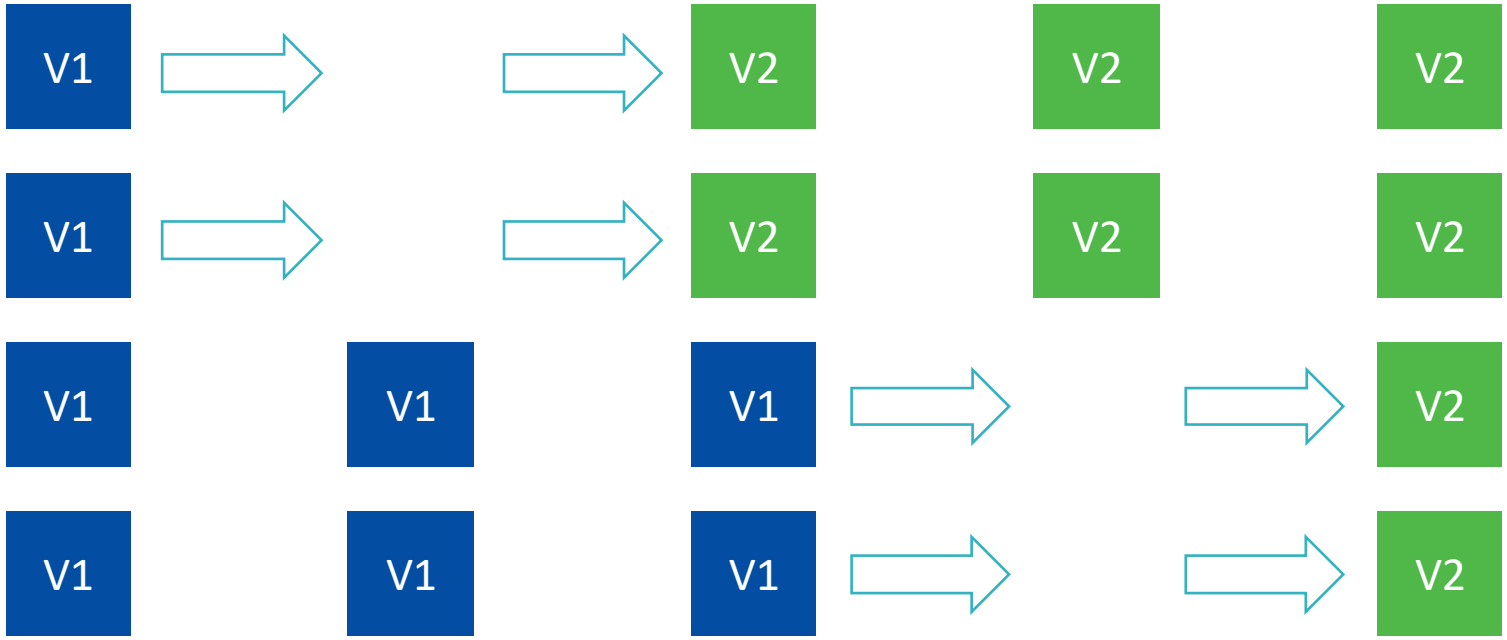

© Copyright by Cloud Mentor Pro | Confidential

ECS Rolling Update – Min 100%, Max 150%

- Starting number of tasks: 4

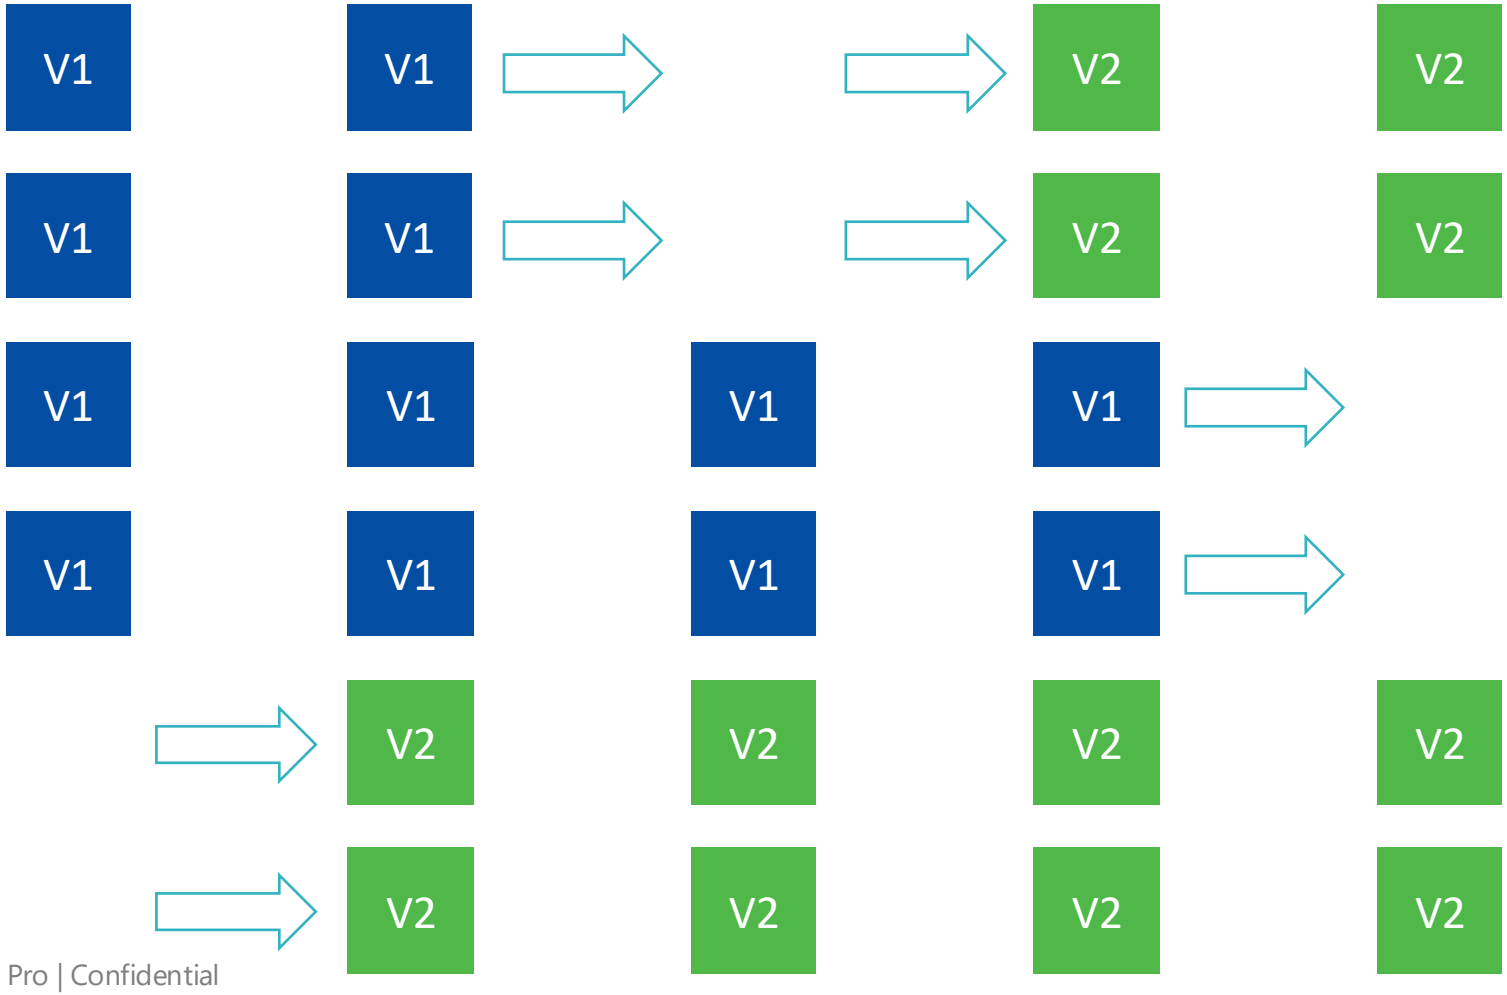

© Copyright by Cloud Mentor Pro | Confidential

ECS tasks invoked by Event Bridge

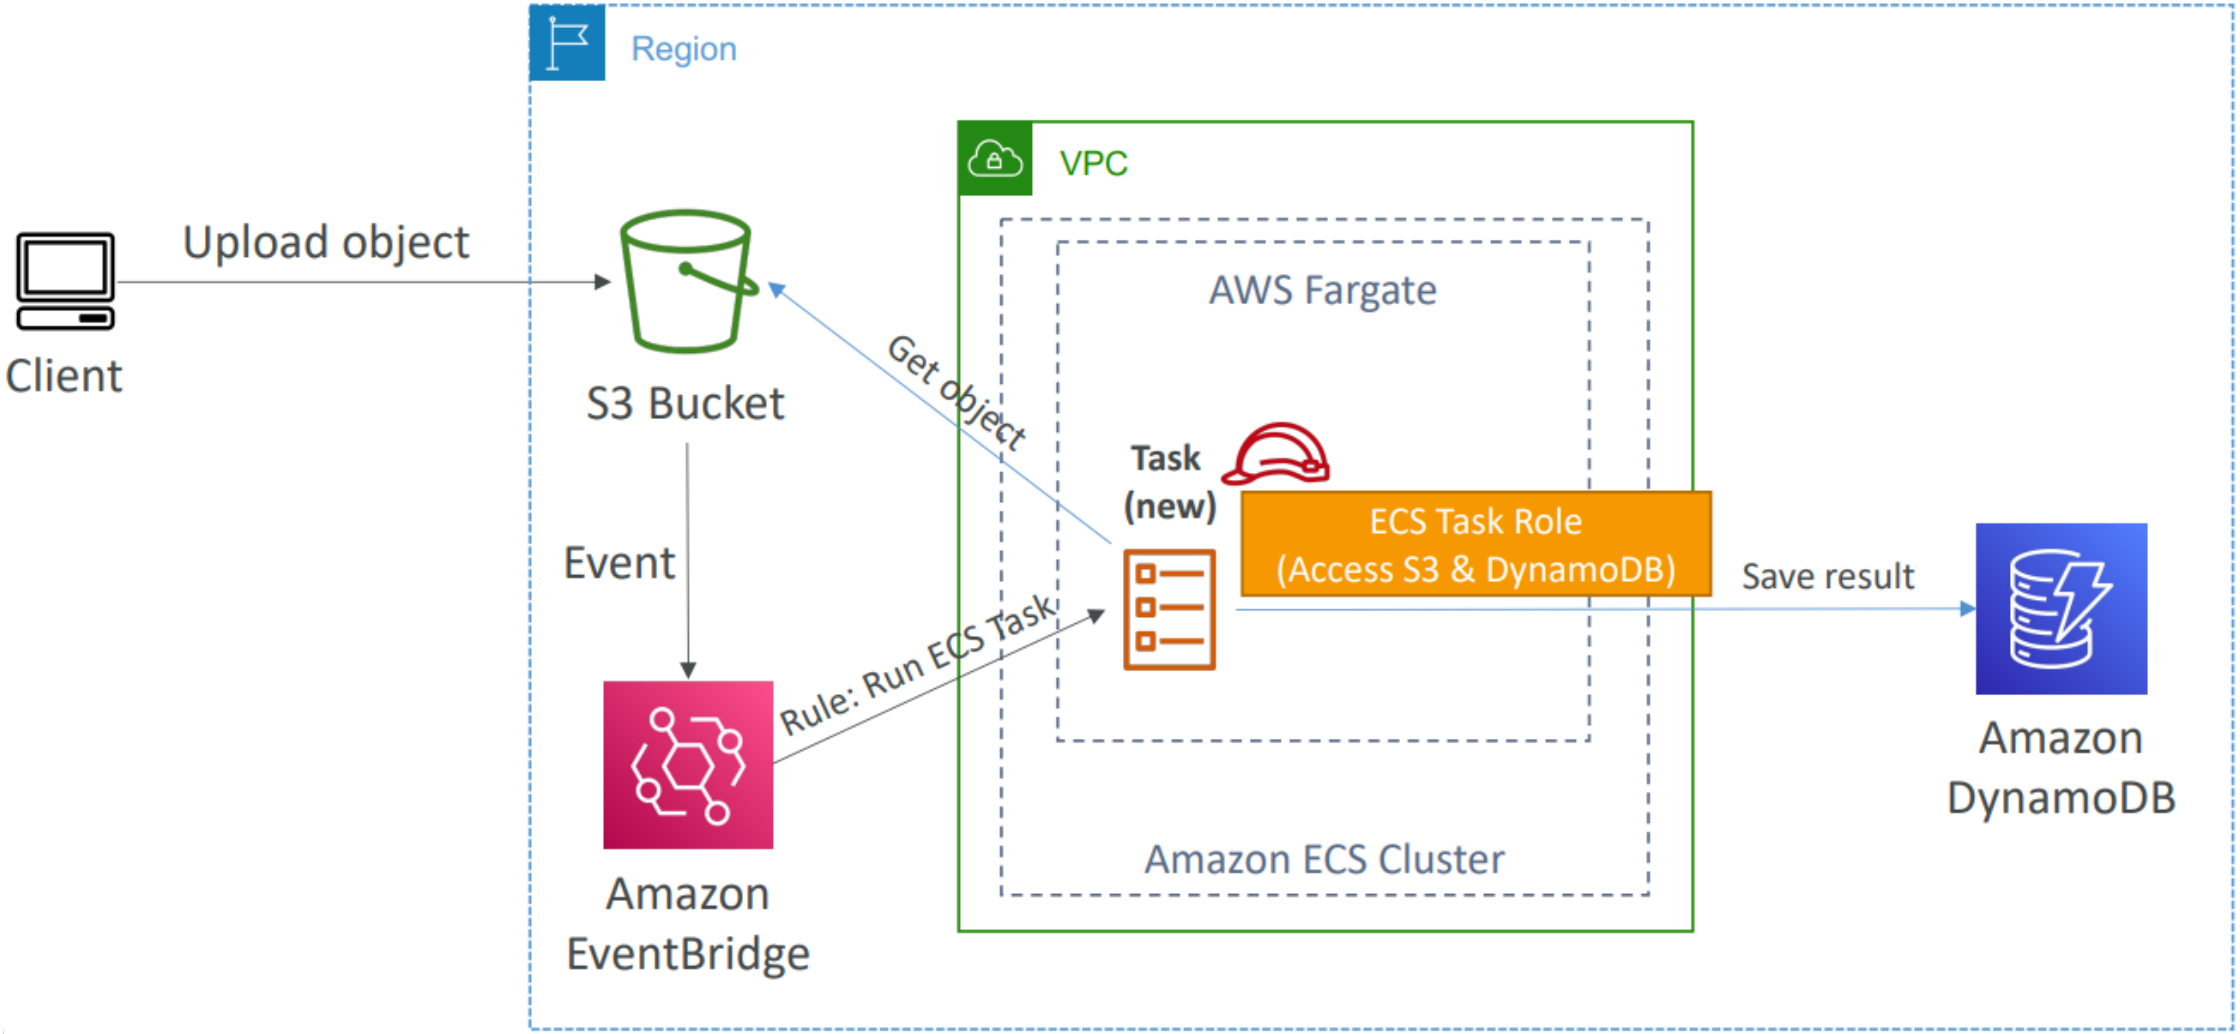

© Copyright by Cloud Mentor Pro | Confidential

ECS tasks invoked by Event Bridge Schedule

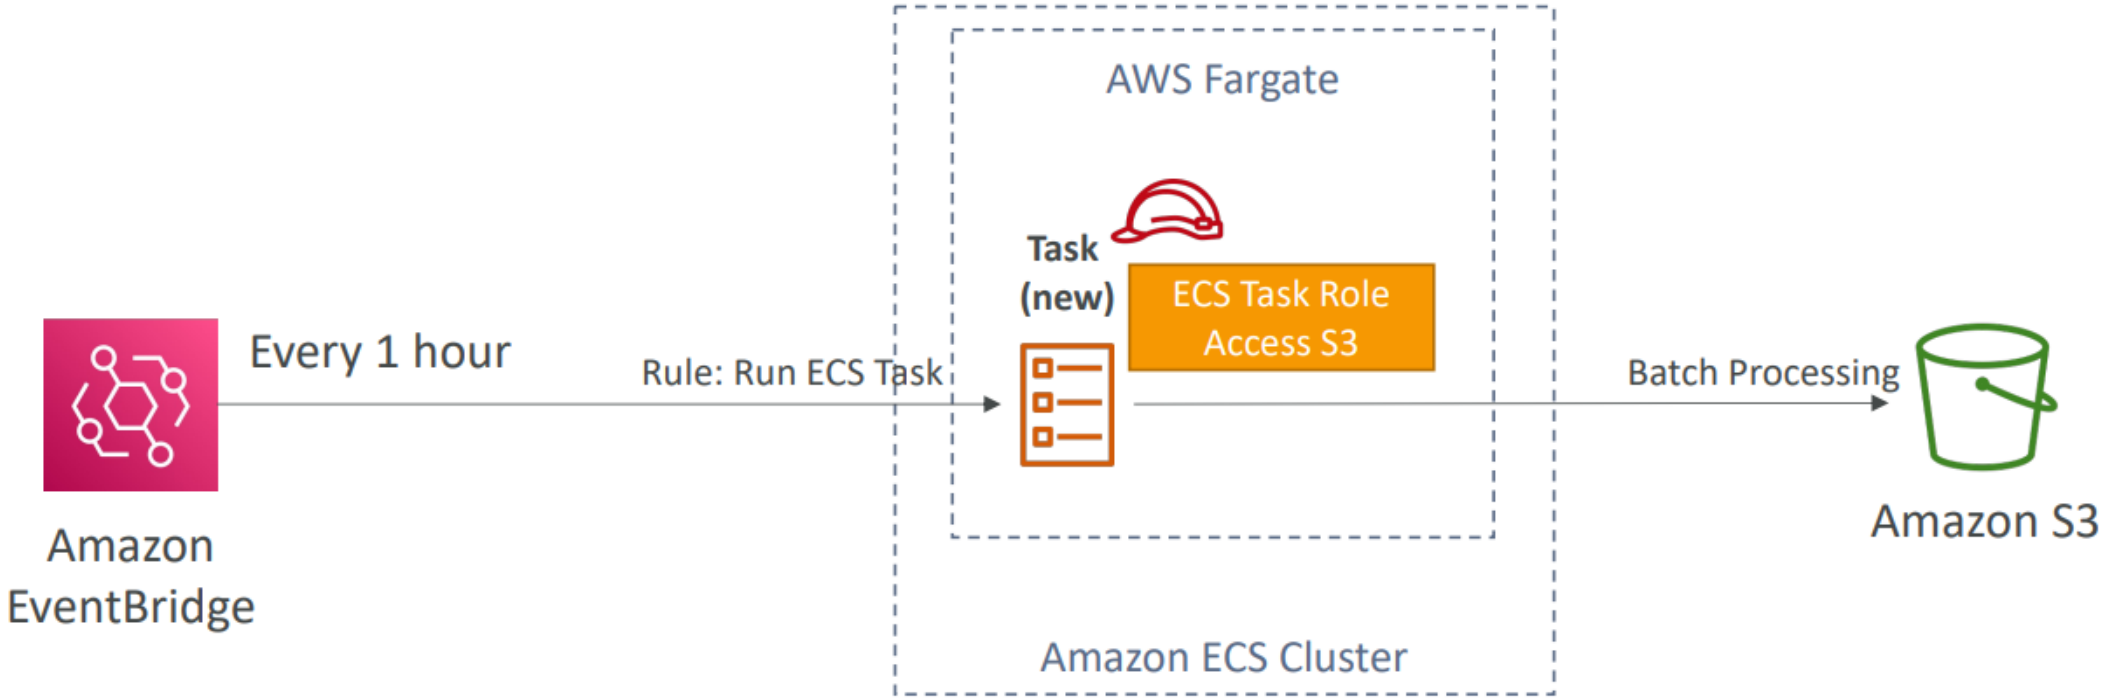

© Copyright by Cloud Mentor Pro | Confidential

ECS – SQS Queue Example

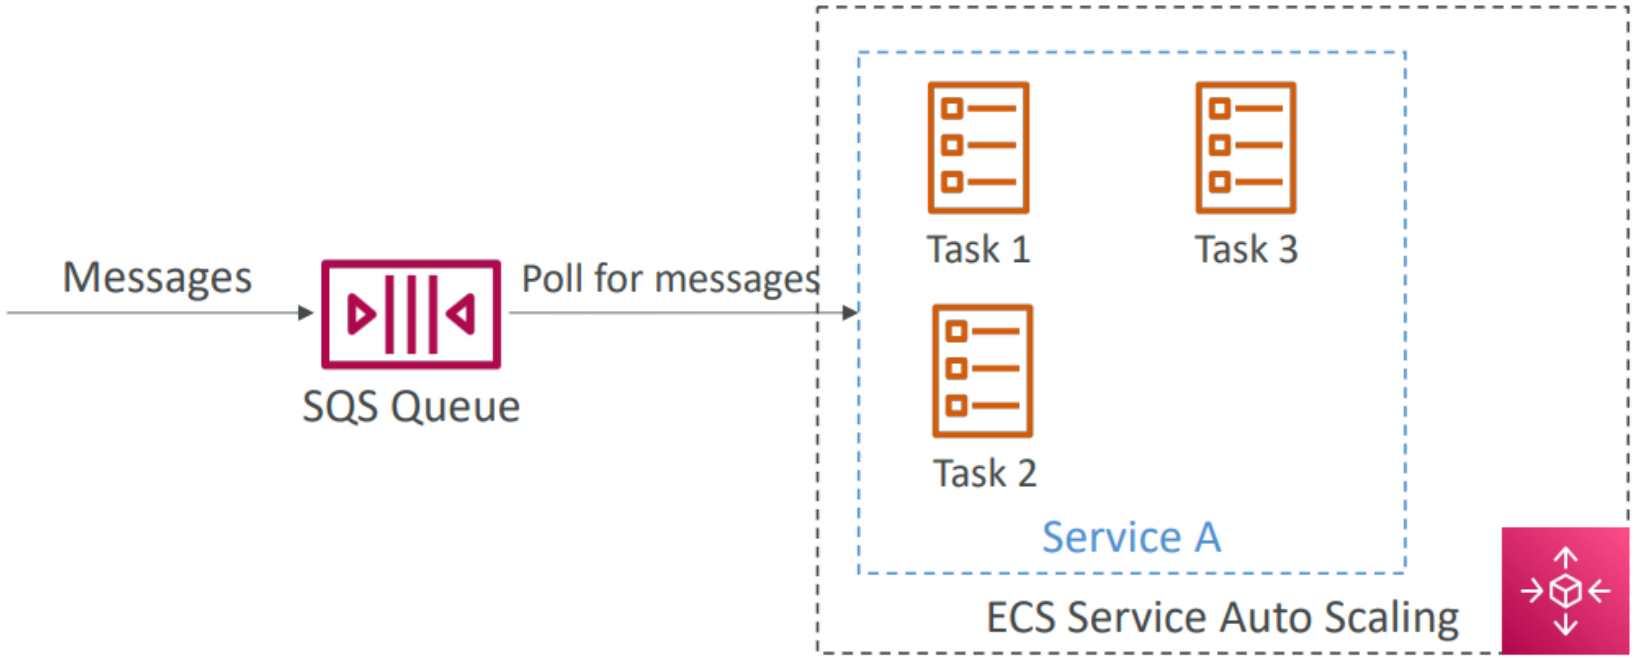

© Copyright by Cloud Mentor Pro | Confidential

ECS – Intercept Stopped Tasks using EventBridge

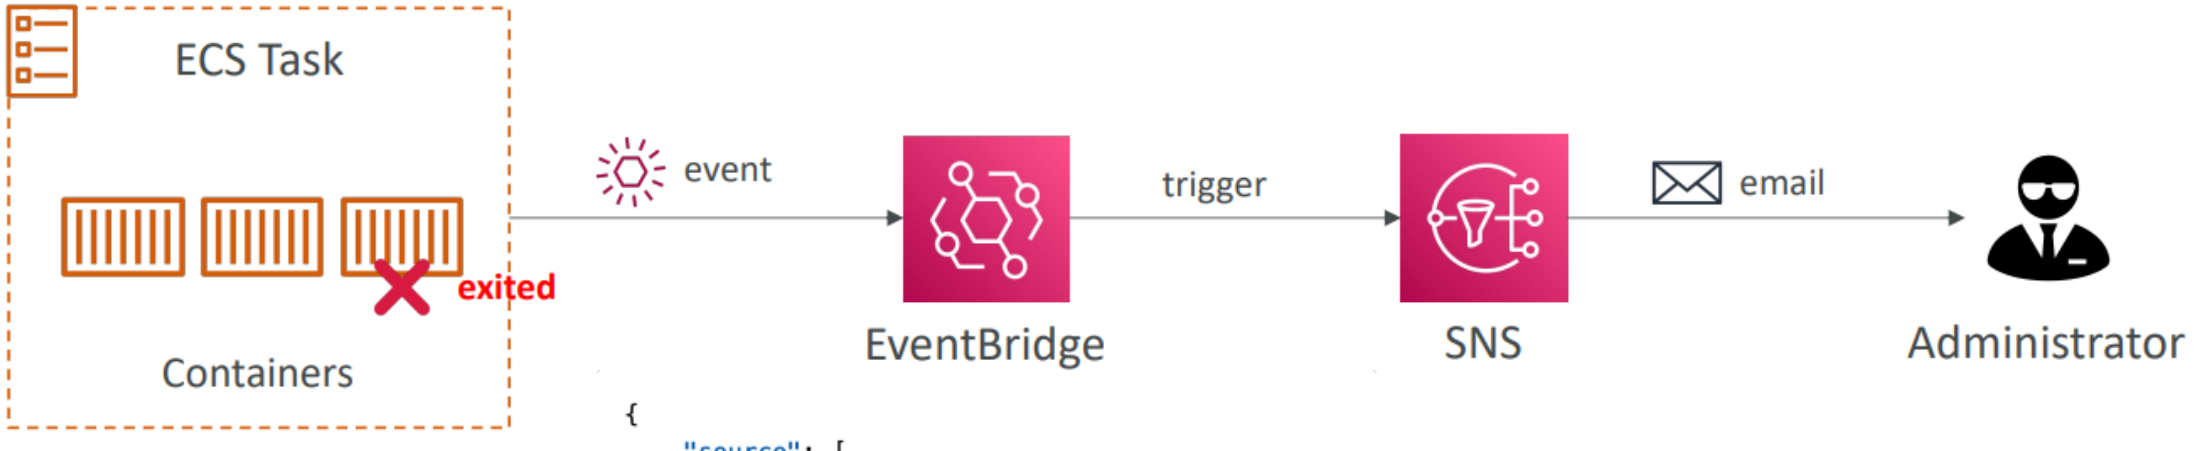


```
{
    "source": [
        "aws.ecs"
    ],
    "detail-type": [
        "ECS Task State Change"
    ],
    "detail": {
        "lastStatus": [
            "STOPPED"
        ],
        "stoppedReason": [
            "Essential container in task exited"
        ]
    }
}
```

Event Pattern

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Load Balancing (EC2 Launch Type)

- We get a Dynamic Host Port Mapping if you define only the container port in the task definition
- The ALB finds the right port on your EC2 Instances
- You must allow on the EC2 instance's Security Group any port from the ALB's Security Group

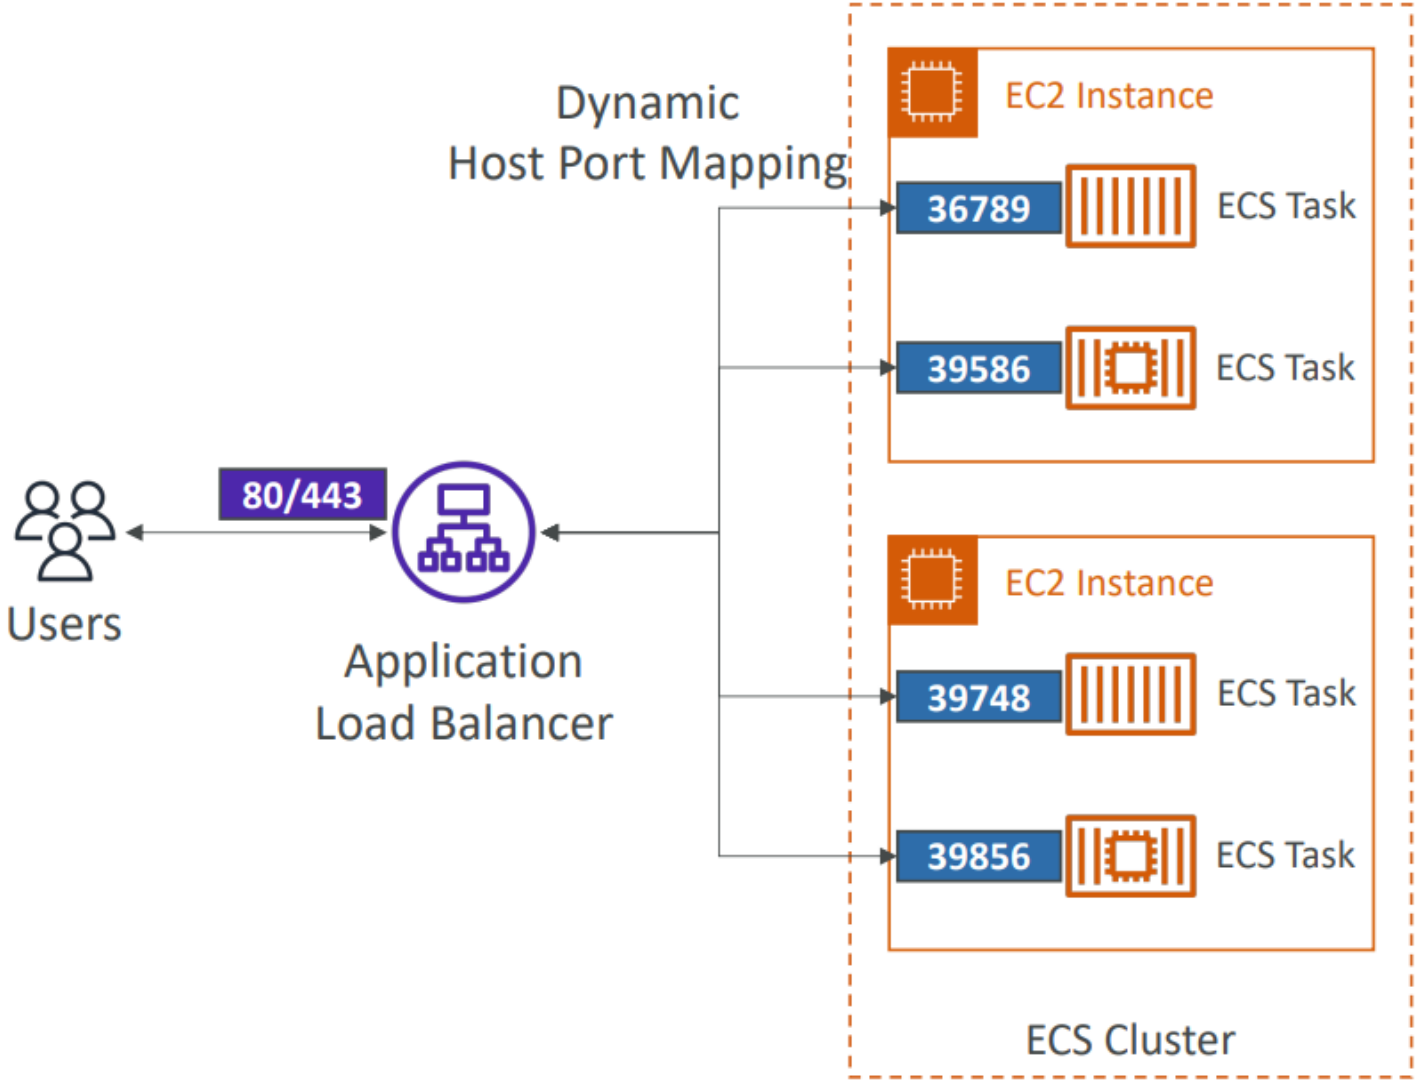

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Load Balancing (Fargate)

- Each task has a unique private IP
- Only define the container port (host port is not applicable)
- Example
  - ECS ENI Security Group
    - Allow port 80 from the ALB
  - ALB Security Group
  - Allow port 80/443 from web

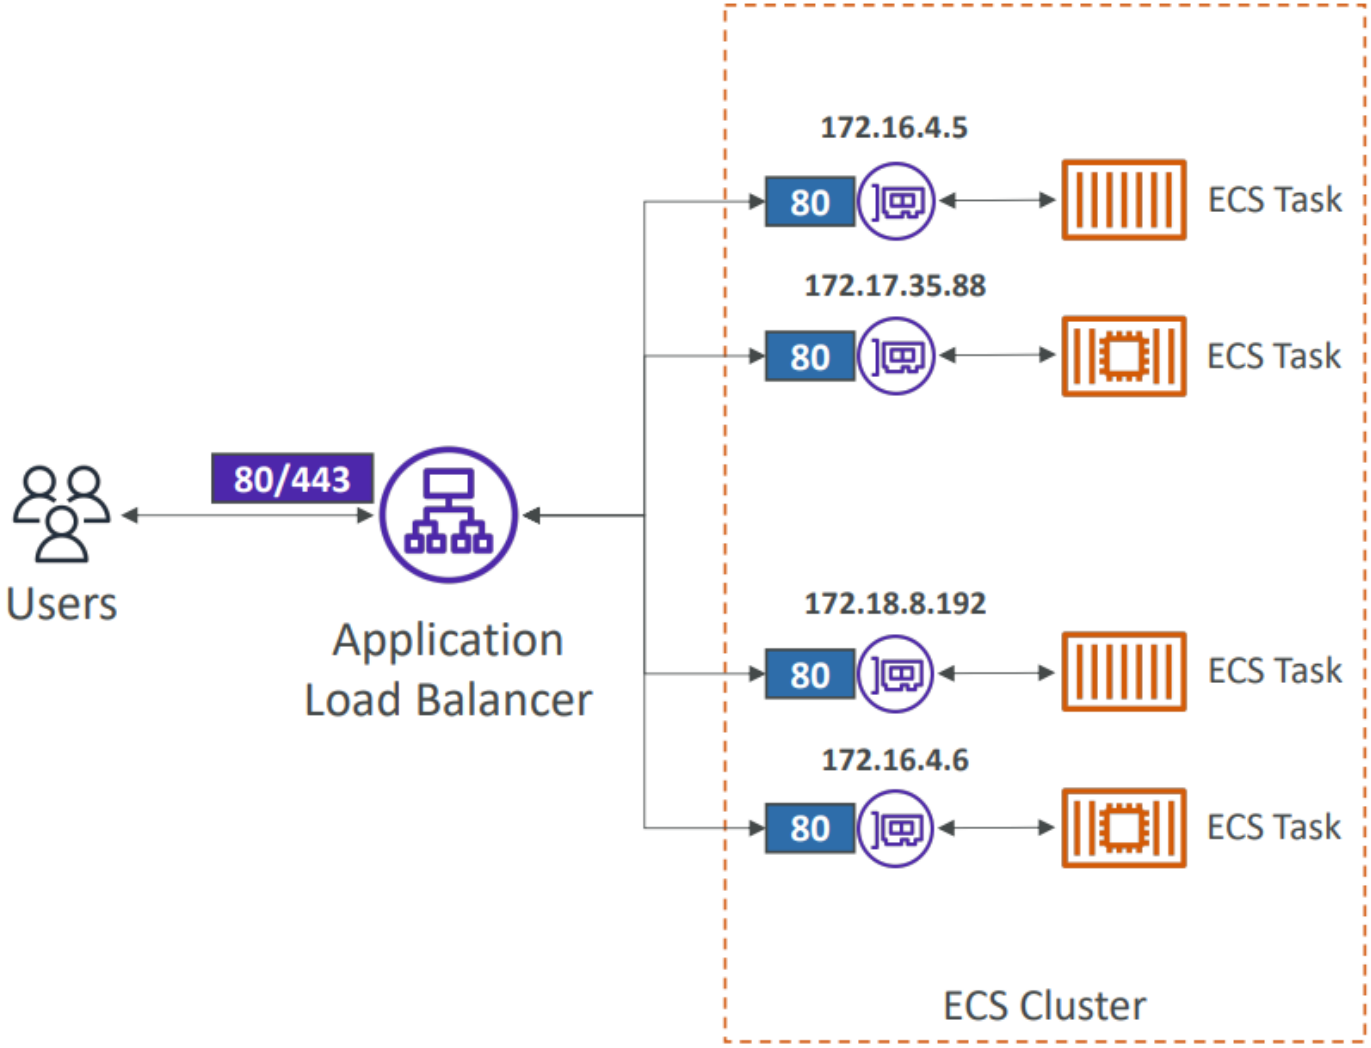

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECS – Environment Variables

- Environment Variable
  - Hardcoded – e.g., URLs
  - SSM Parameter Store – sensitive variables (e.g., API keys, shared configs)
  - Secrets Manager – sensitive variables (e.g., DB passwords)
- Environment Files (bulk) – Amazon S3

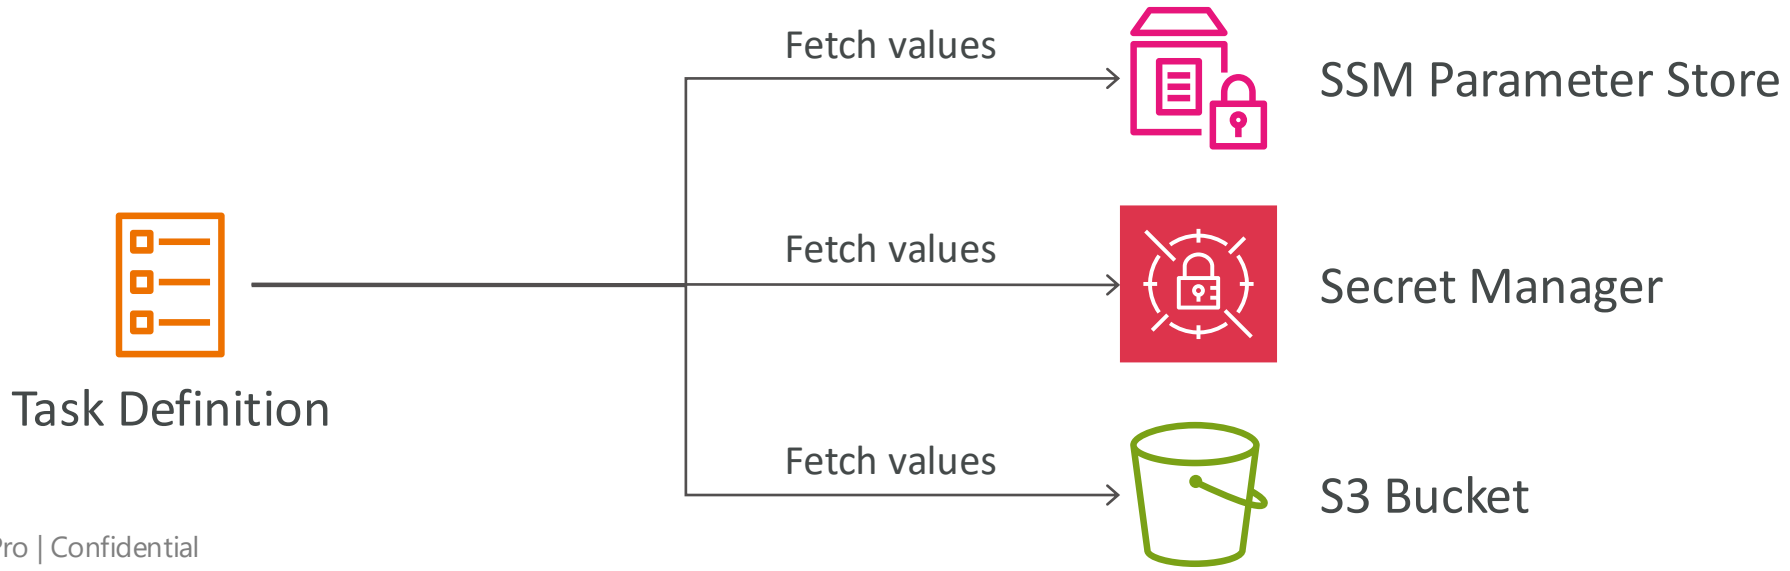

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECR

- ECR = Elastic Container Registry
- Store and manage Docker images on AWS
- **Private** and **Public** repository **(Amazon ECR Public Gallery** https://gallery.ecr.aws**)**
- Fully integrated with ECS, backed by Amazon S3
- Access is controlled through IAM (permission errors => policy)
- Supports image vulnerability scanning, versioning, image tags, image lifecycle, …

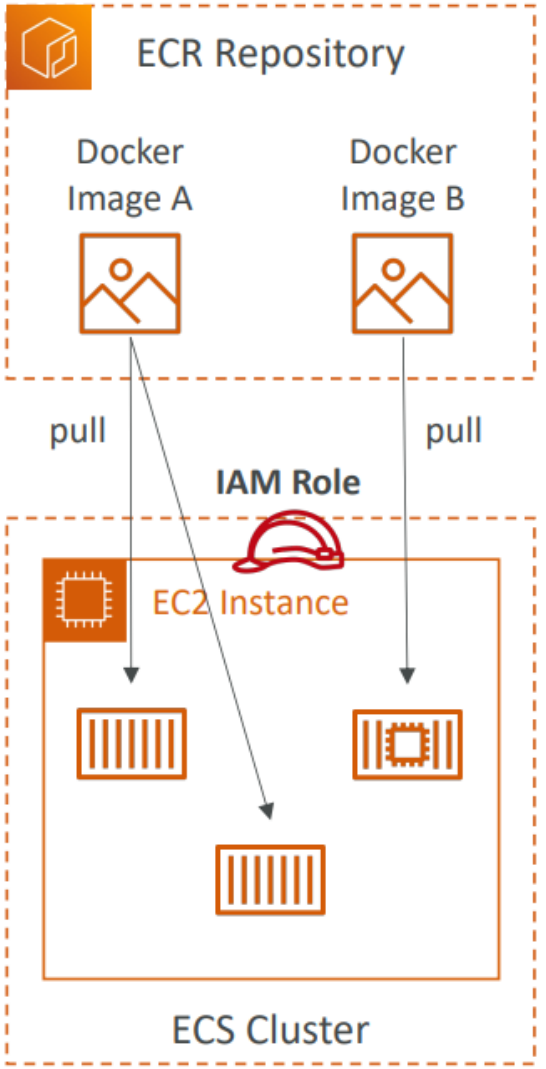

© Copyright by Cloud Mentor Pro | Confidential

Amazon ECR – Using AWS CLI

- Login Command
  - AWS CLI v2

```
aws ecr get-login-password --region region | docker login --username AWS
--password-stdin aws_account_id.dkr.ecr.region.amazonaws.com
```

- Docker Commands
  - Push

```
docker push aws_account_id.dkr.ecr.region.amazonaws.com/demo:latest
```

  - Pull

```
docker pull aws_account_id.dkr.ecr.region.amazonaws.com/demo:latest
```

- In case an EC2 instance (or you) can't pull a Docker image, check IAM permissions

© Copyright by Cloud Mentor Pro | Confidential

AWS Copilot

- CLI tool to build, release, and operate production-ready containerized apps
- Run your apps on **AppRunner, ECS**, and **Fargate**
- Helps you focus on building apps rather than setting up infrastructure
- Provisions all required infrastructure for containerized apps (ECS, VPC, ELB, ECR…)
- Automated deployments with one command using CodePipeline
- Deploy to multiple environments
- Troubleshooting, logs, health status…

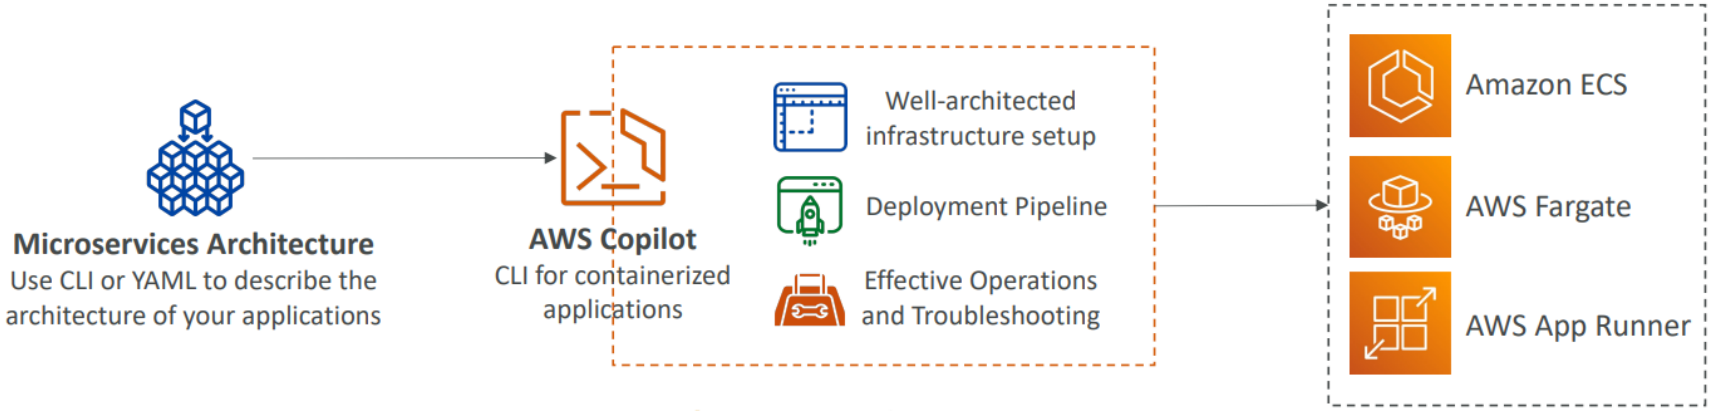

© Copyright by Cloud Mentor Pro | Confidential

Amazon EKS Overview

- Amazon EKS = Amazon Elastic Kubernetes Service
- It is a way to launch **managed Kubernetes clusters on AWS**
- Kubernetes is an **open-source system** for automatic deployment, scaling and management of containerized (usually Docker) application
- It's an alternative to ECS, similar goal but different API
- EKS supports EC2 if you want to deploy worker nodes or **Fargate** to deploy serverless containers
- **Use case:** if your company is already using Kubernetes on-premises or in another cloud, and wants to migrate to AWS using Kubernetes
- Kubernetes is cloud-agnostic (can be used in any cloud – Azure, GCP…)
- For multiple regions, deploy one EKS cluster per region
- Collect logs and metrics using CloudWatch Container Insights

© Copyright by Cloud Mentor Pro | Confidential

Amazon EKS - Diagram

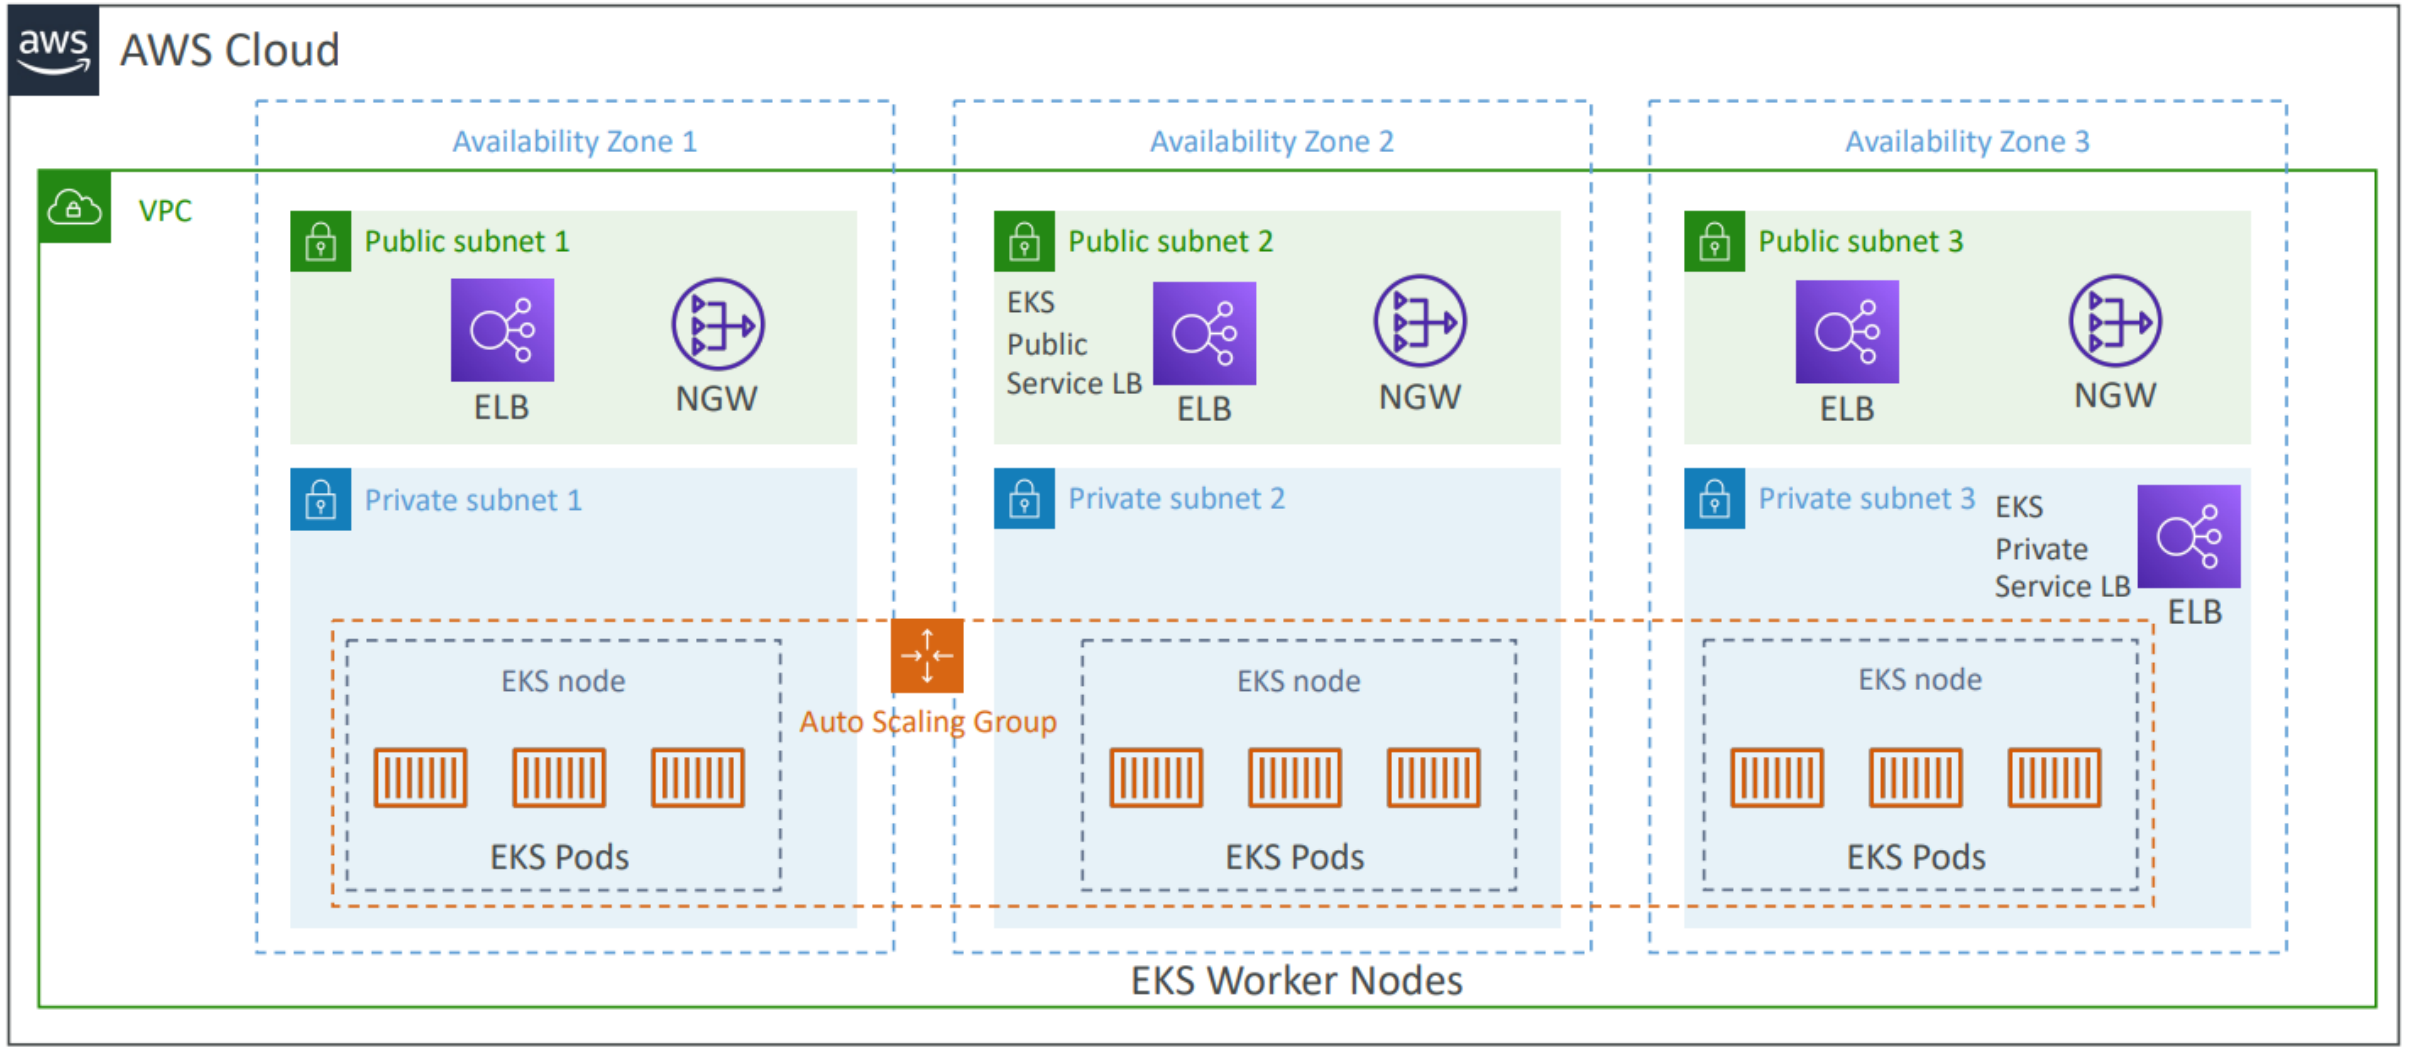

© Copyright by Cloud Mentor Pro | Confidential

Amazon EKS – Node Types

- Managed Node Groups
  - Creates and manages Nodes (EC2 instances) for you
  - Nodes are part of an ASG managed by EKS
  - Supports On-Demand or Spot Instances

- Self-Managed Nodes
  - Nodes created by you and registered to the EKS cluster and managed by an ASG
  - You can use prebuilt AMI - Amazon EKS Optimized AMI
  - Supports On-Demand or Spot Instances

- AWS Fargate
  - No maintenance required; no nodes managed

© Copyright by Cloud Mentor Pro | Confidential

Amazon EKS – Data Volumes

- Need to specify **Storage Class** manifest on your EKS cluster
- Leverages a **Container Storage Interface (CSI)** compliant driver

- **Support for…**
- Amazon EBS
- Amazon EFS (works with Fargate)
- Amazon FSx for Lustre
- Amazon FSx for NetApp ONTAP

© Copyright by Cloud Mentor Pro | Confidential

DEMO

- *How to Run a Application on AWS ECS: A Step-by-Step Guide*

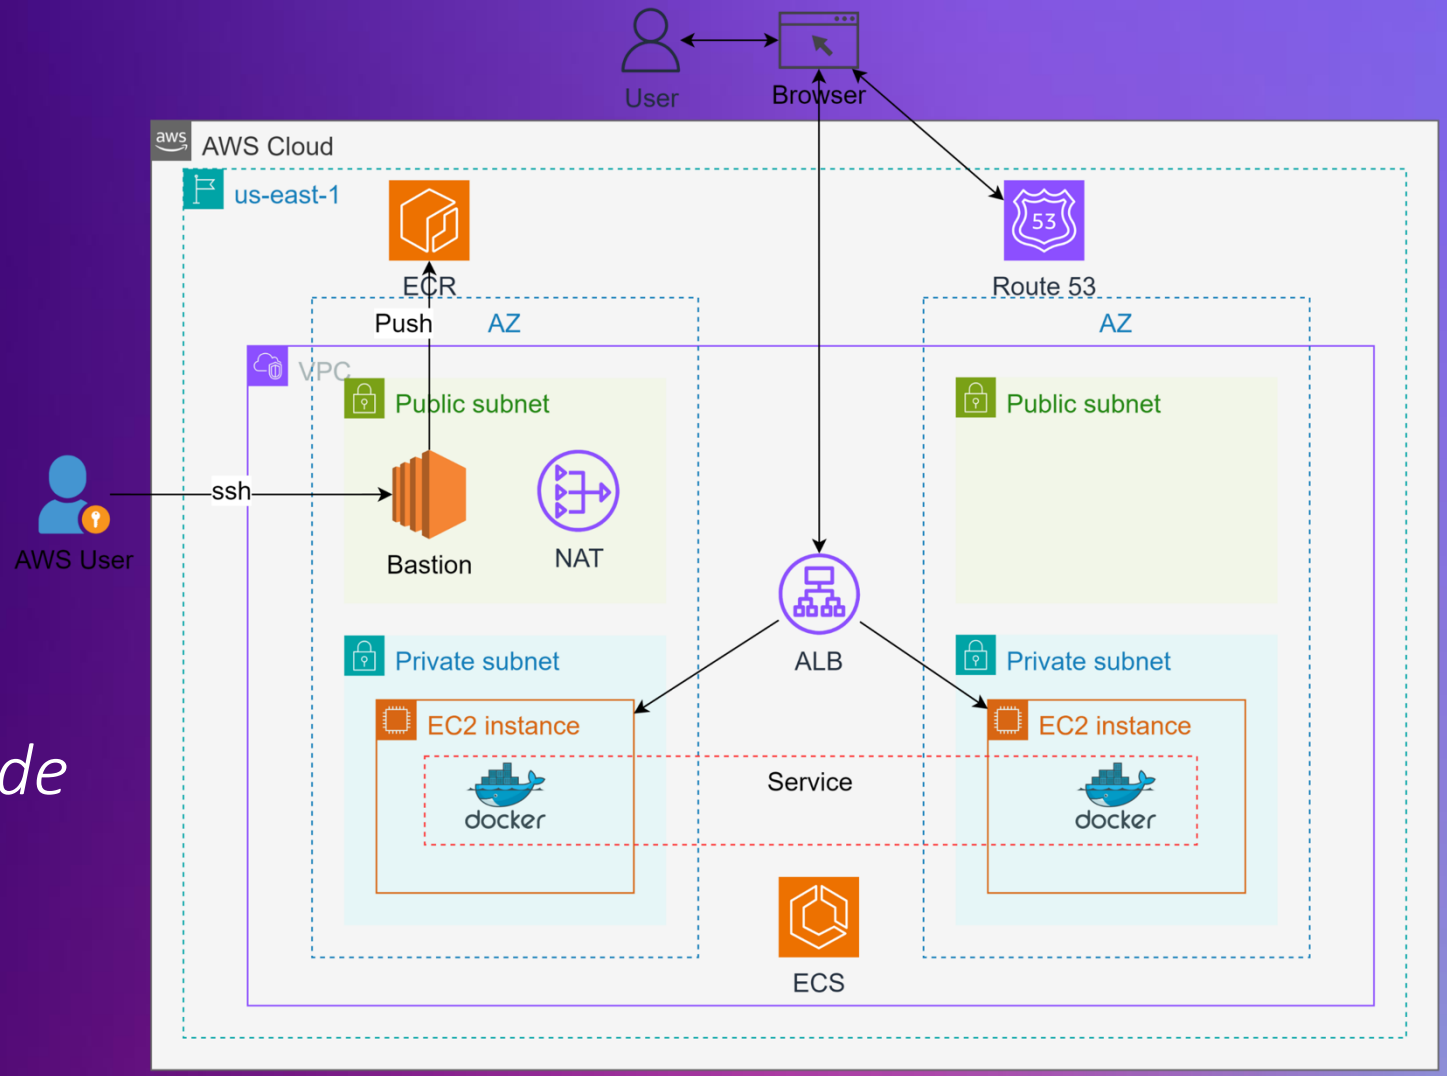

© Copyright by Cloud Mentor Pro | Confidential

Architecture Diagram

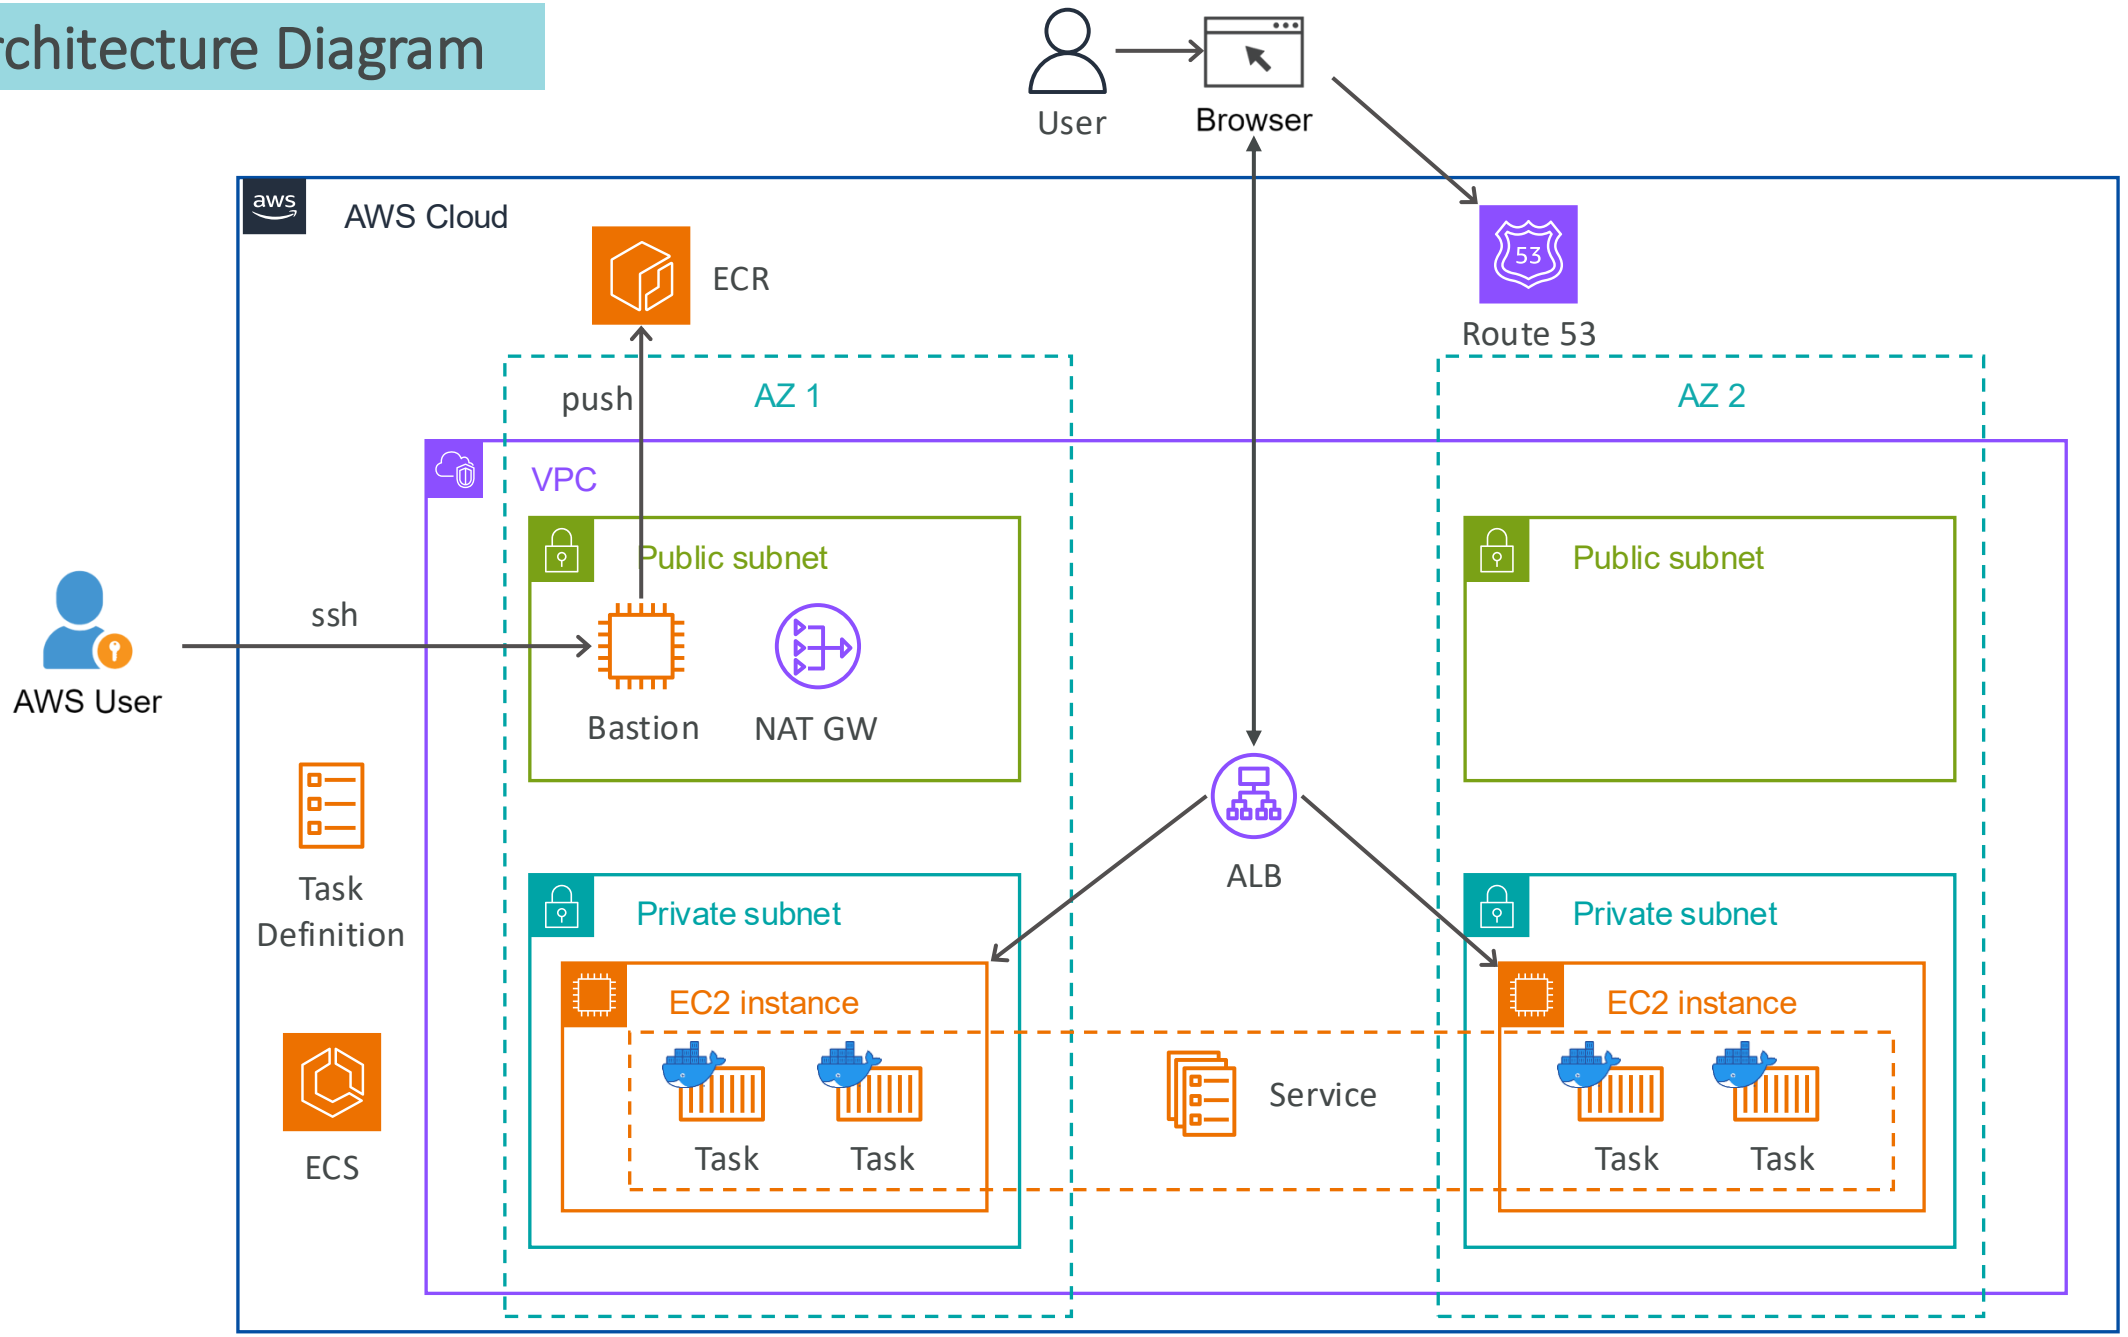

© Copyright by Cloud Mentor Pro | Confidential

1. Create VPC

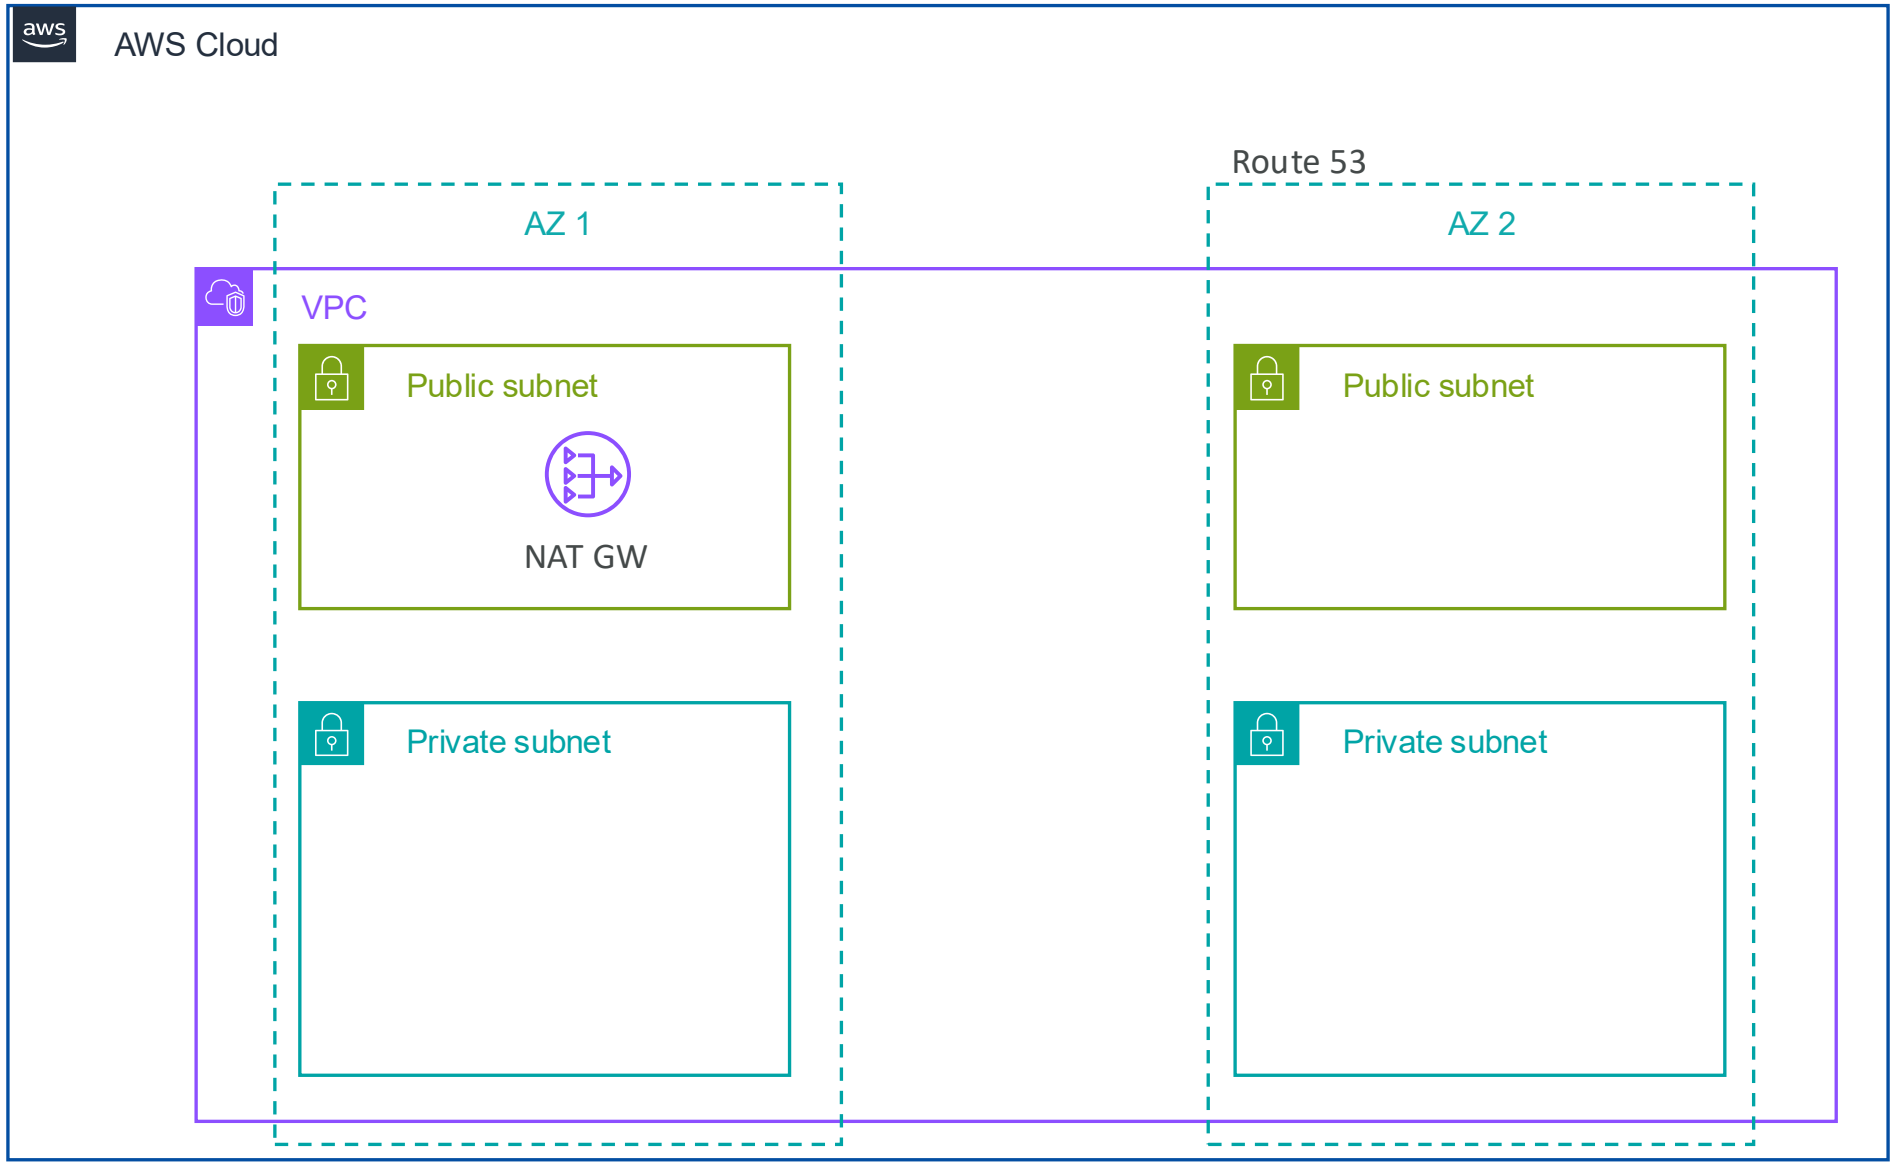

© Copyright by Cloud Mentor Pro | Confidential

2. Create Bastion host

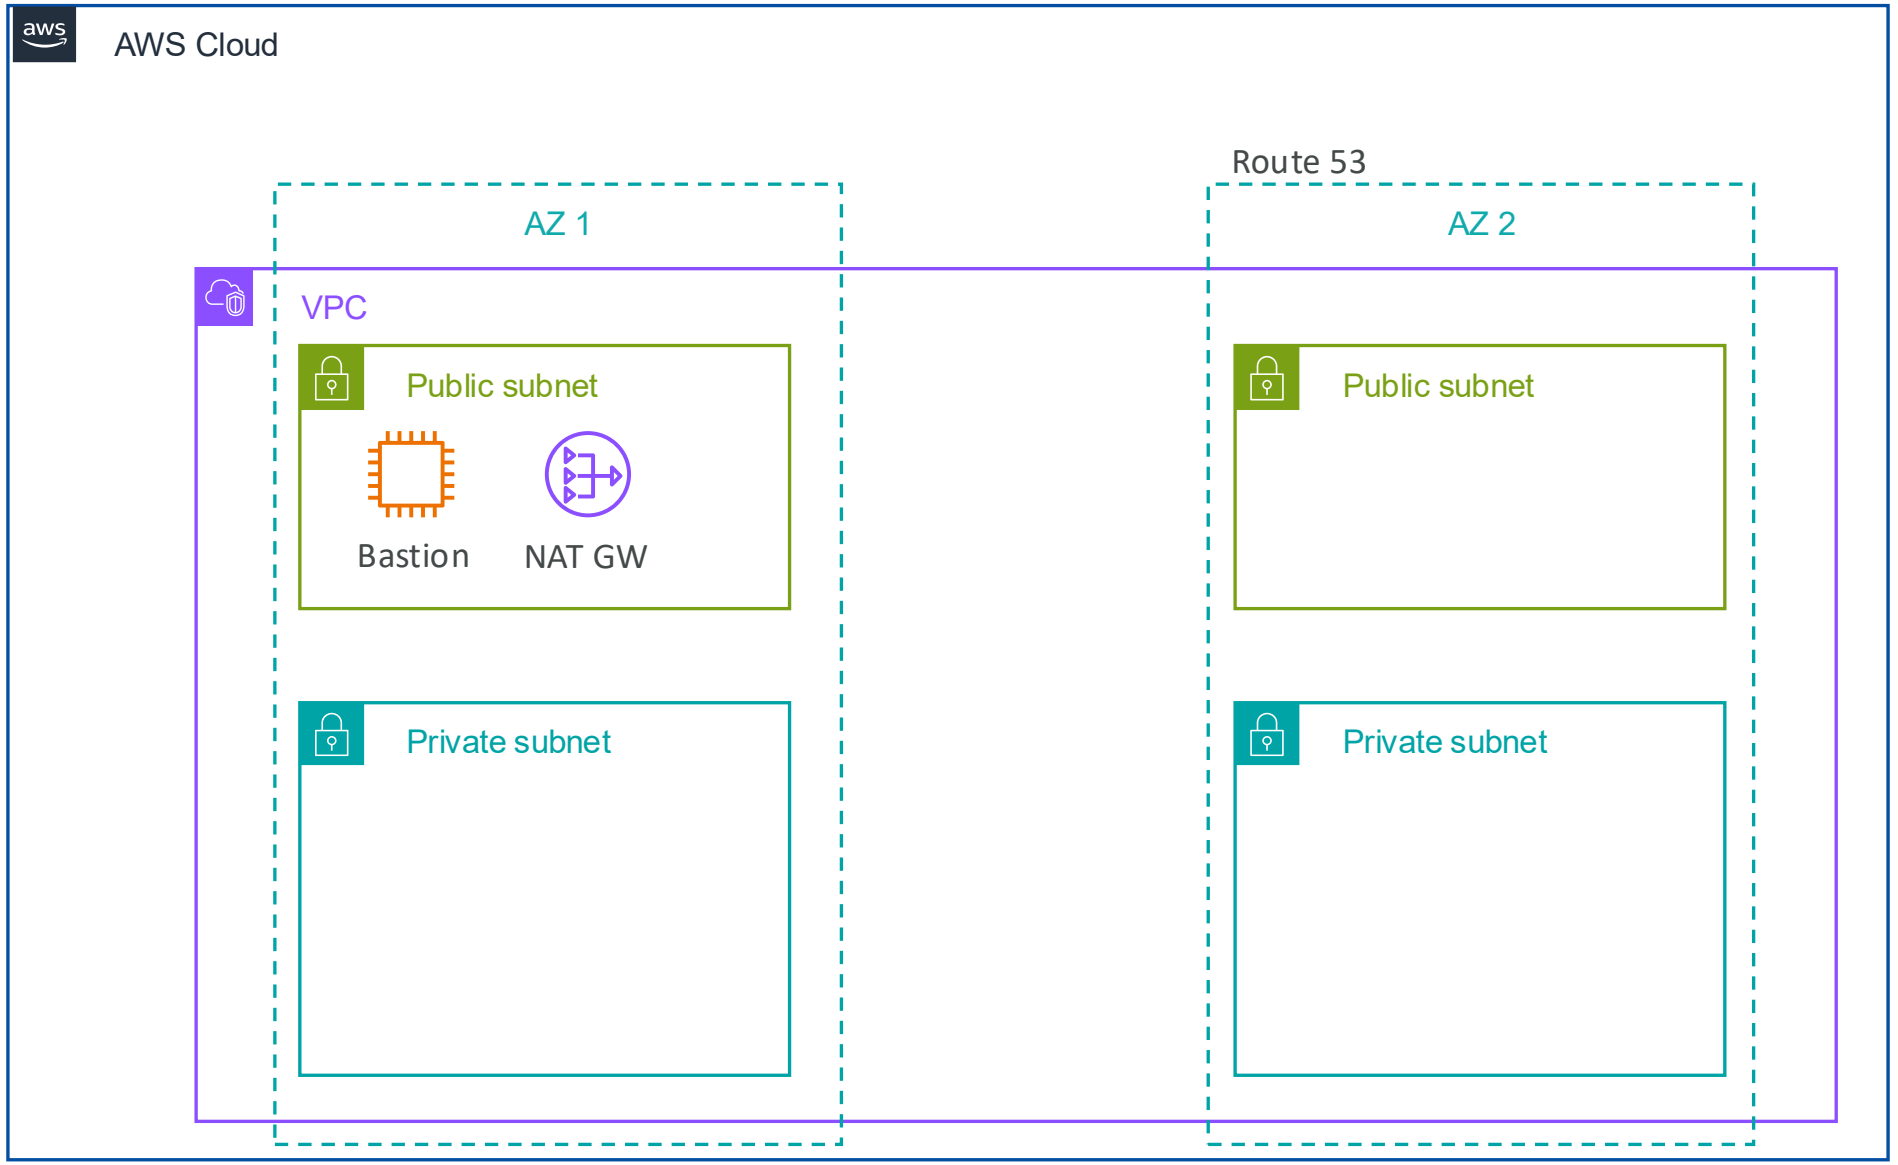

© Copyright by Cloud Mentor Pro | Confidential

3. Create ECR and Build Docker image

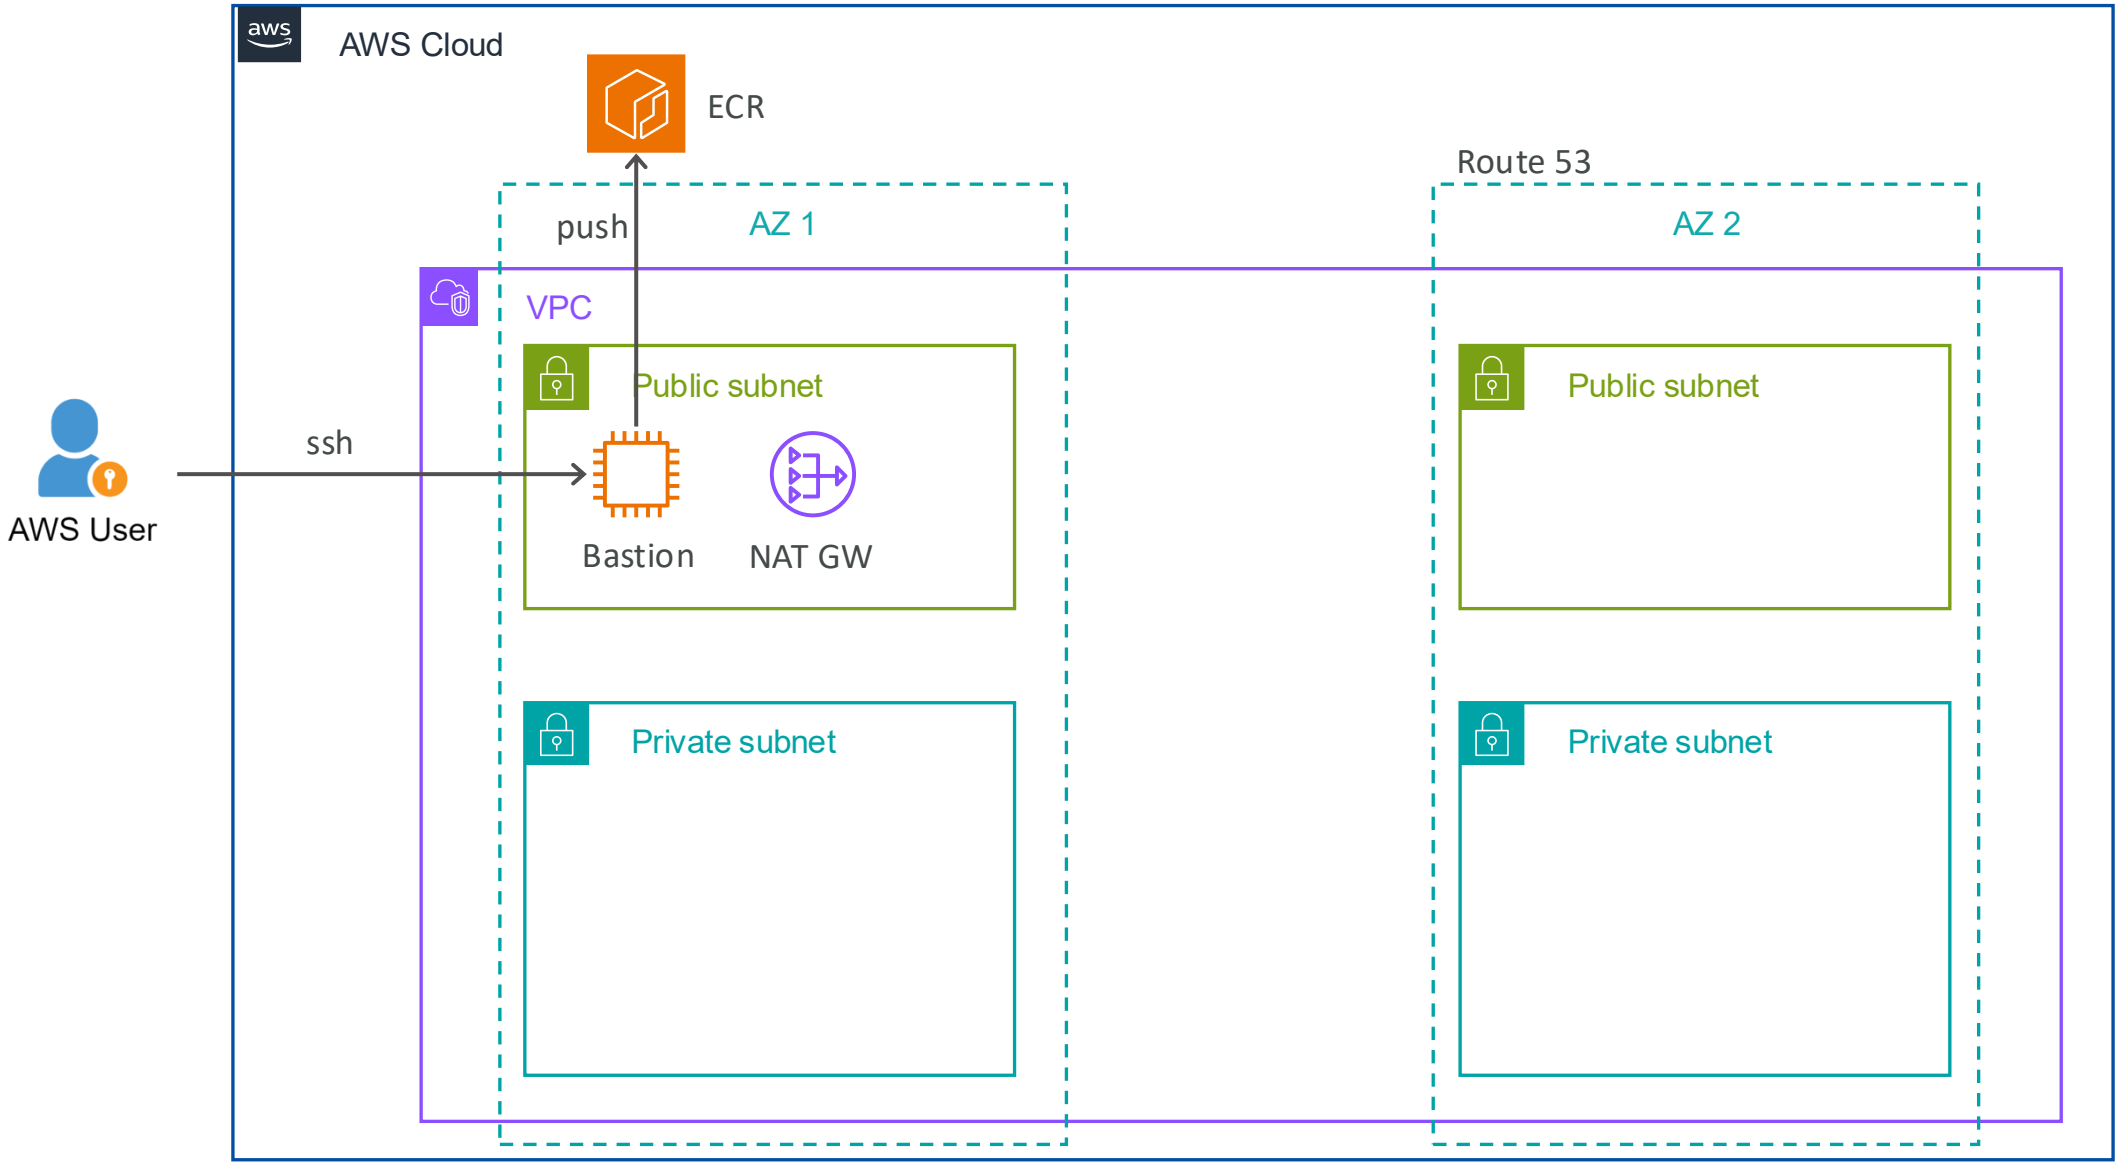

© Copyright by Cloud Mentor Pro | Confidential

4. Create ECS Cluster

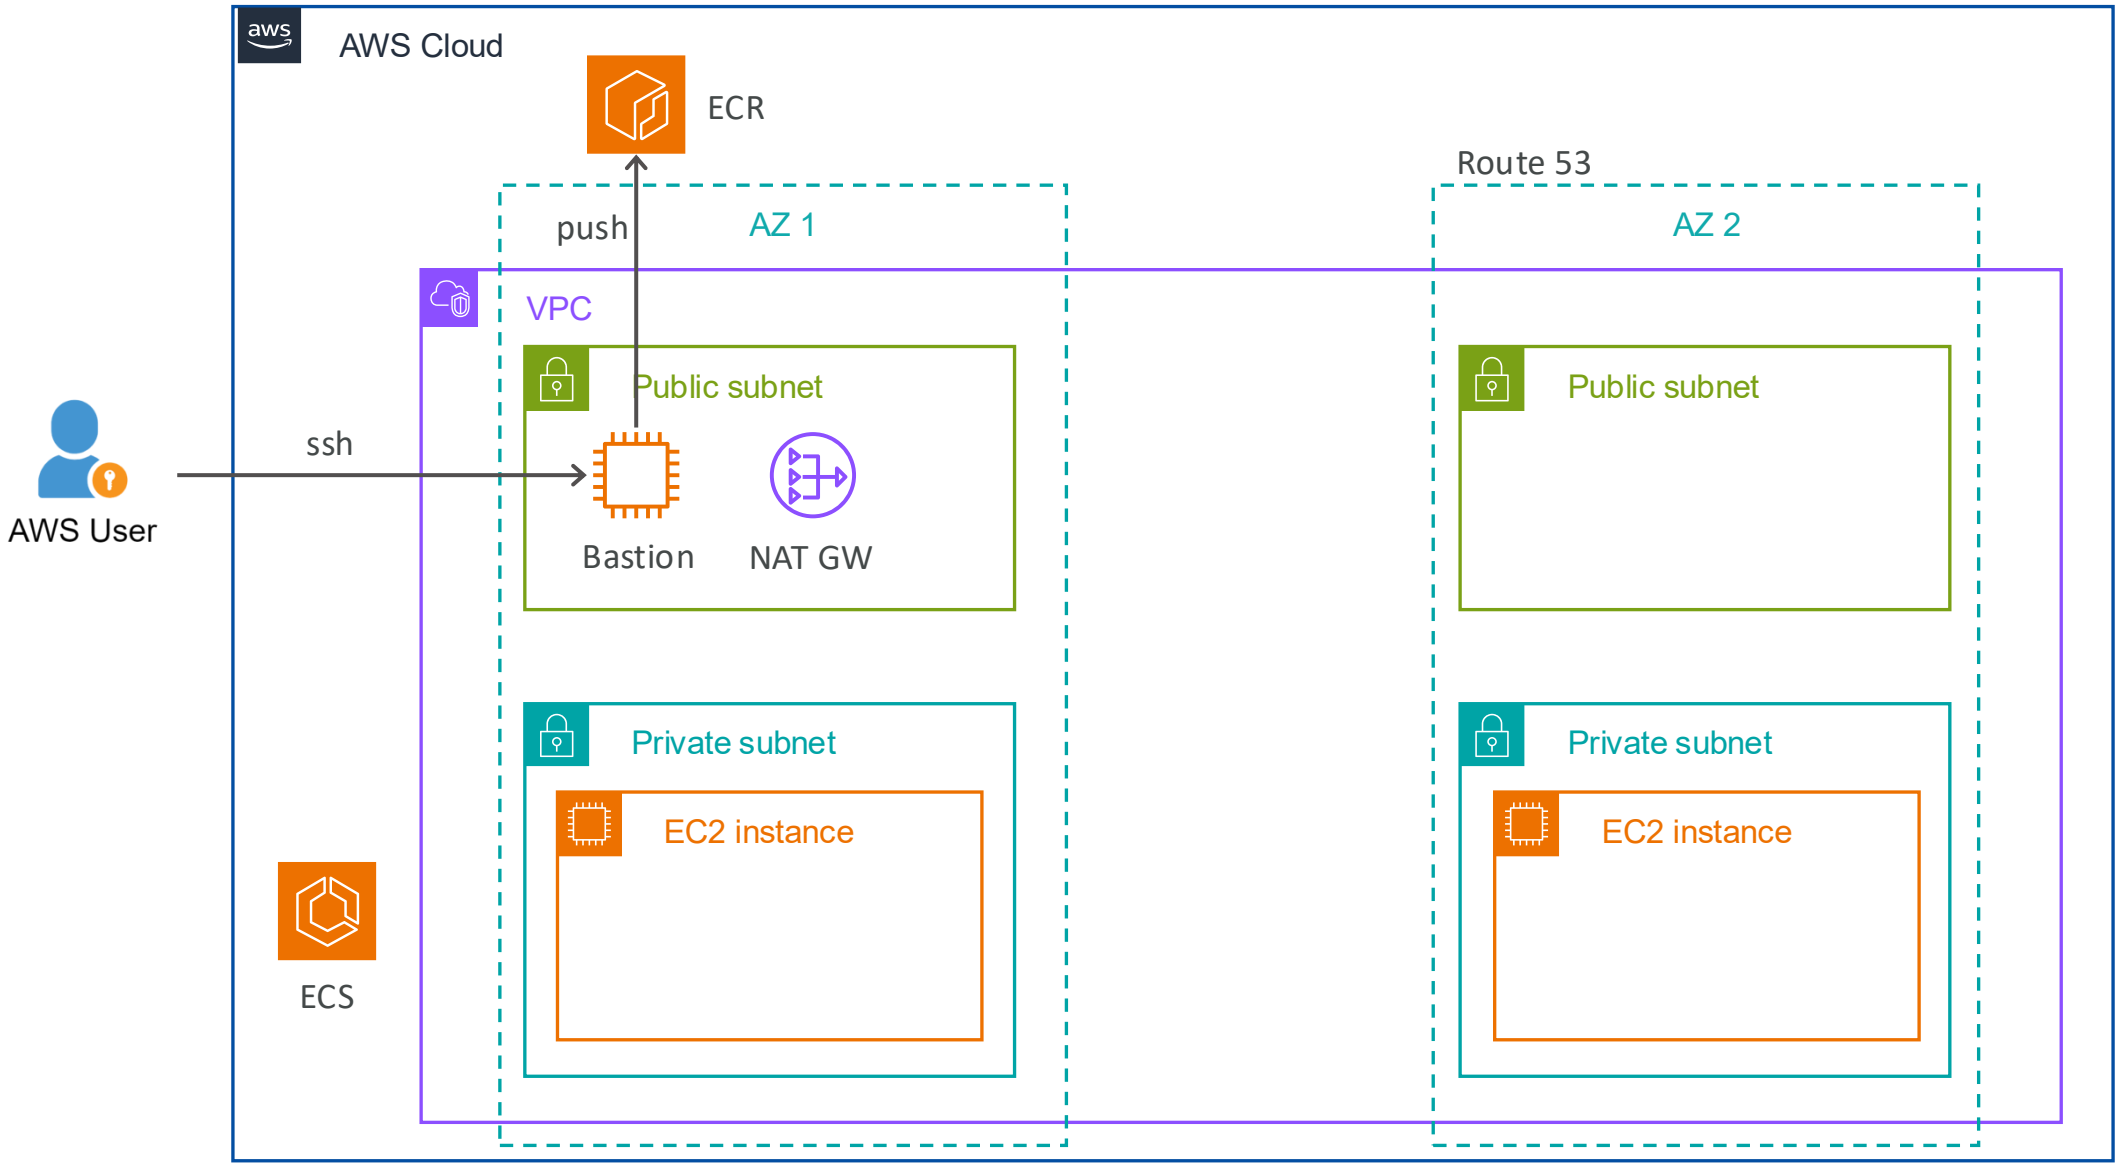

© Copyright by Cloud Mentor Pro | Confidential

5. Create Task definition

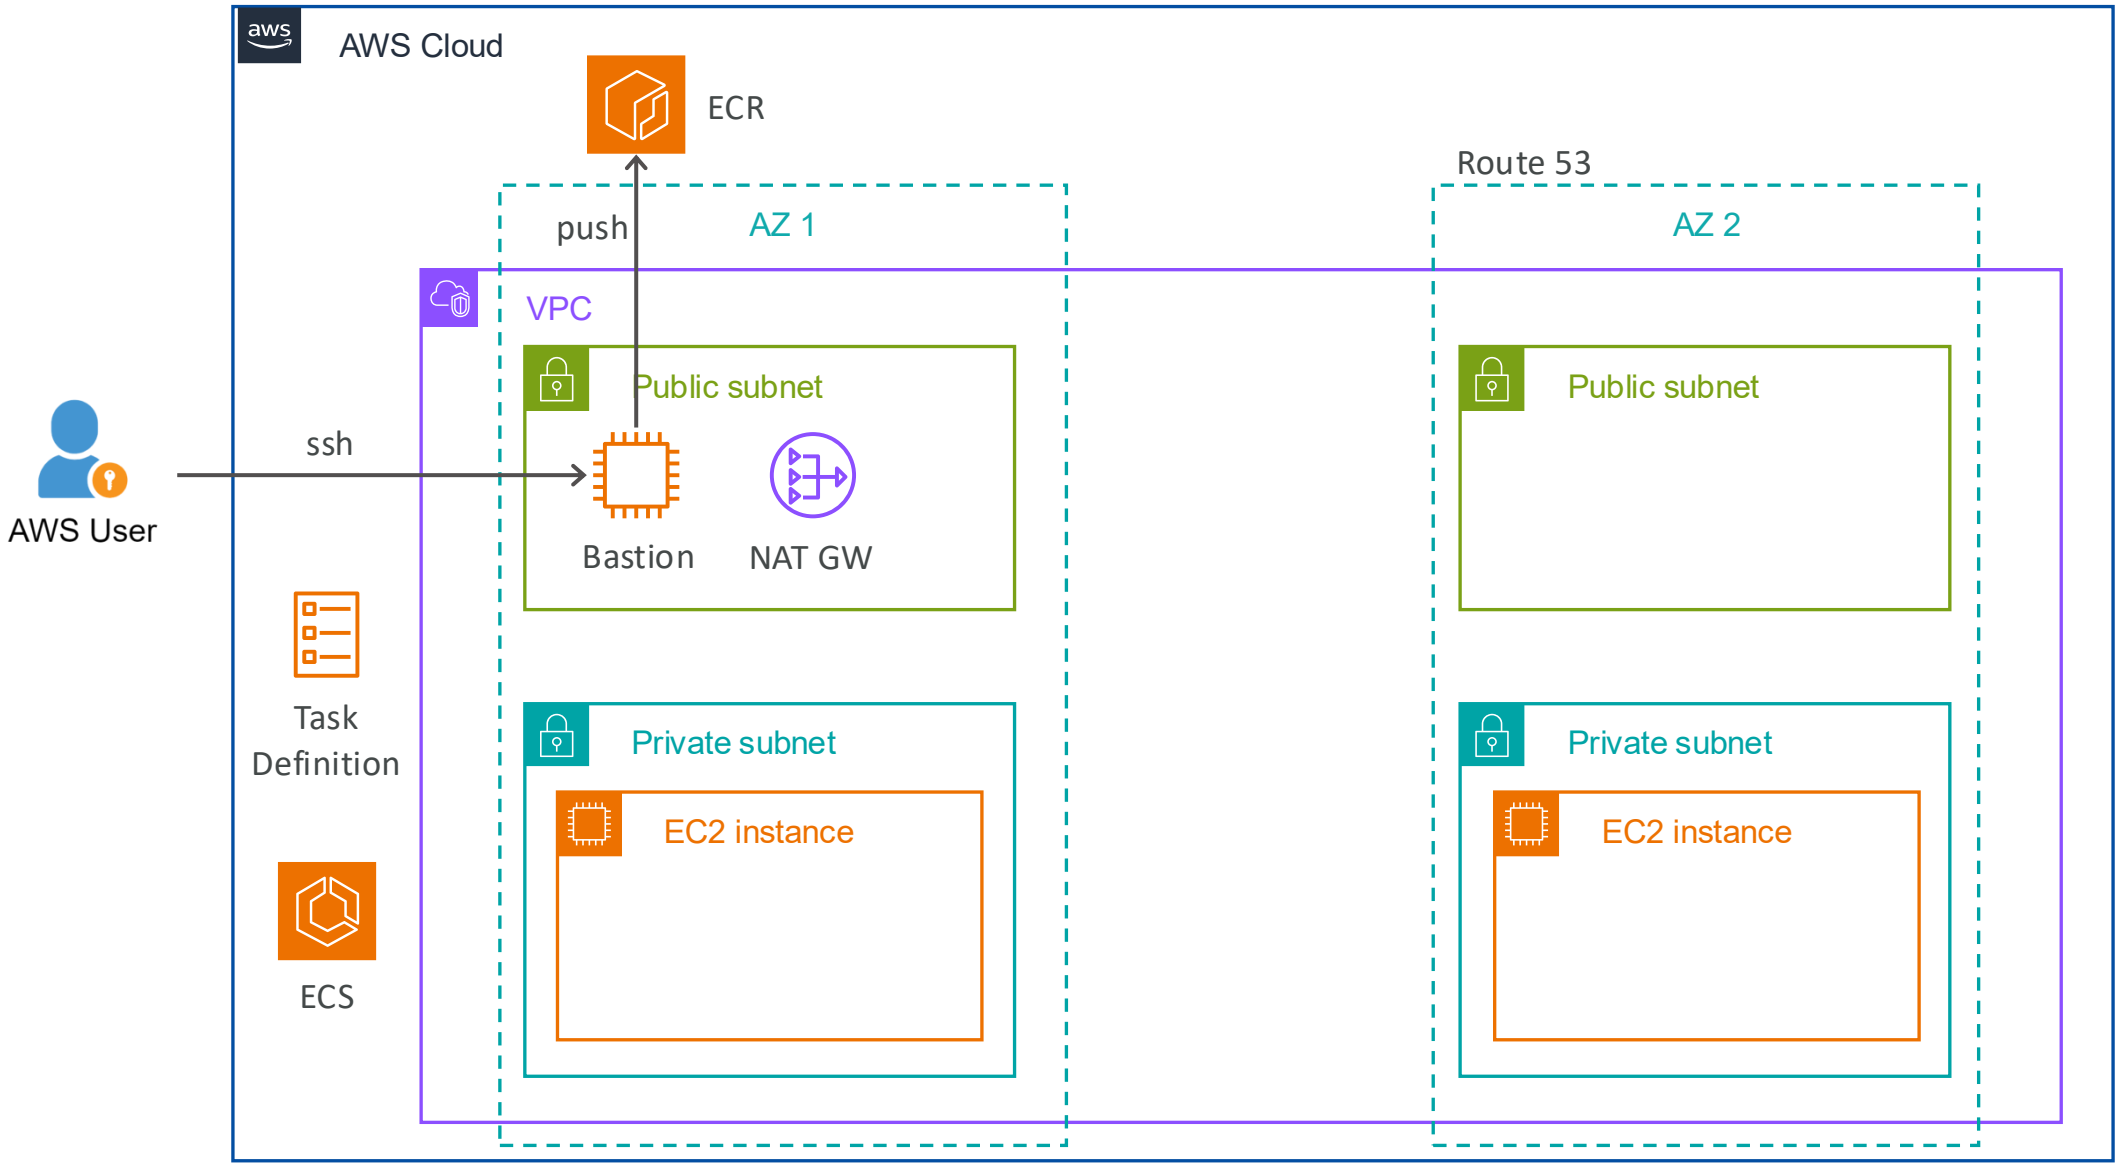

© Copyright by Cloud Mentor Pro | Confidential

6-7. Create TG and ALB

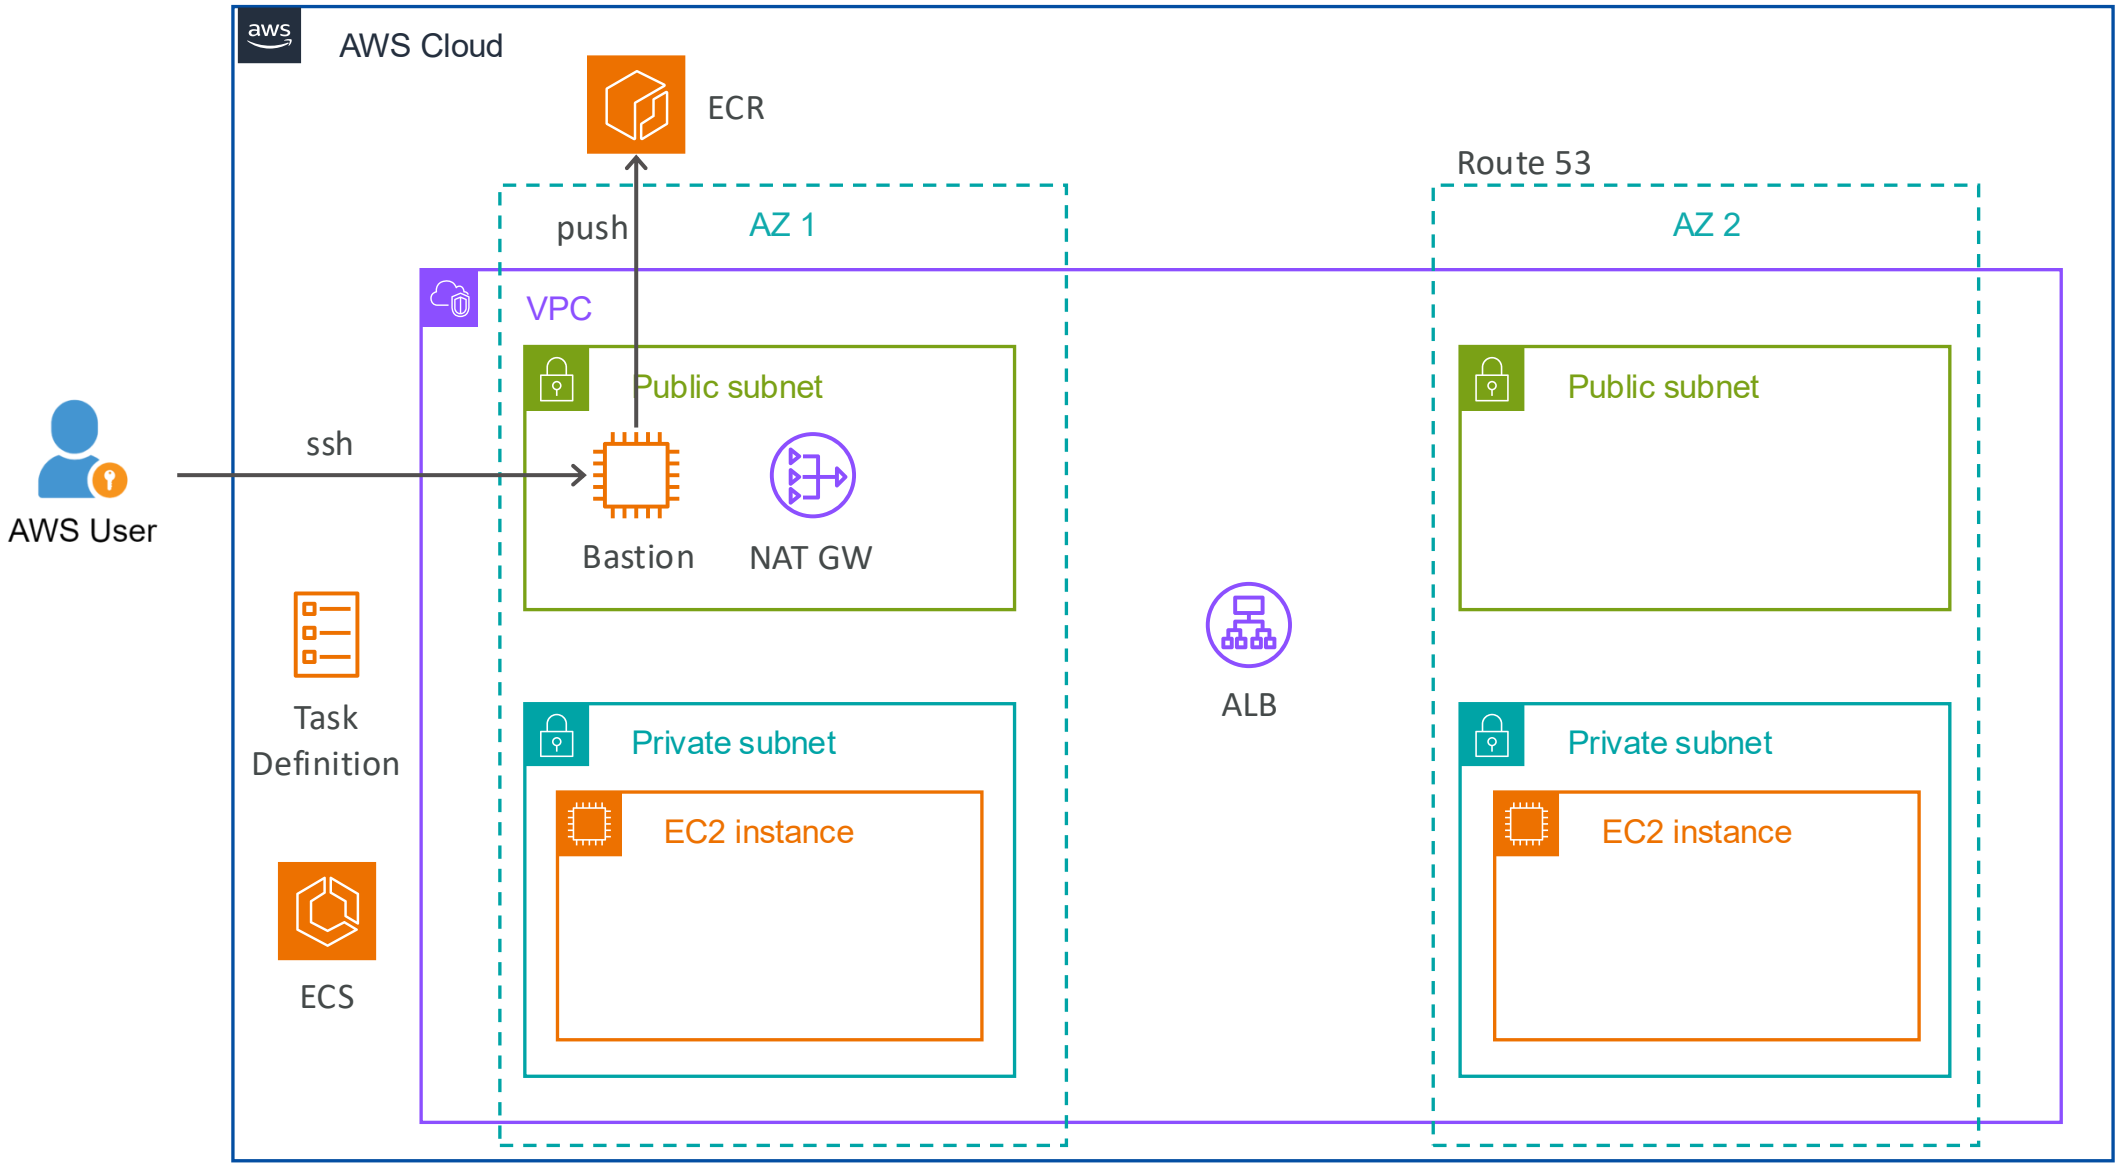

© Copyright by Cloud Mentor Pro | Confidential

8. Create ECS Service

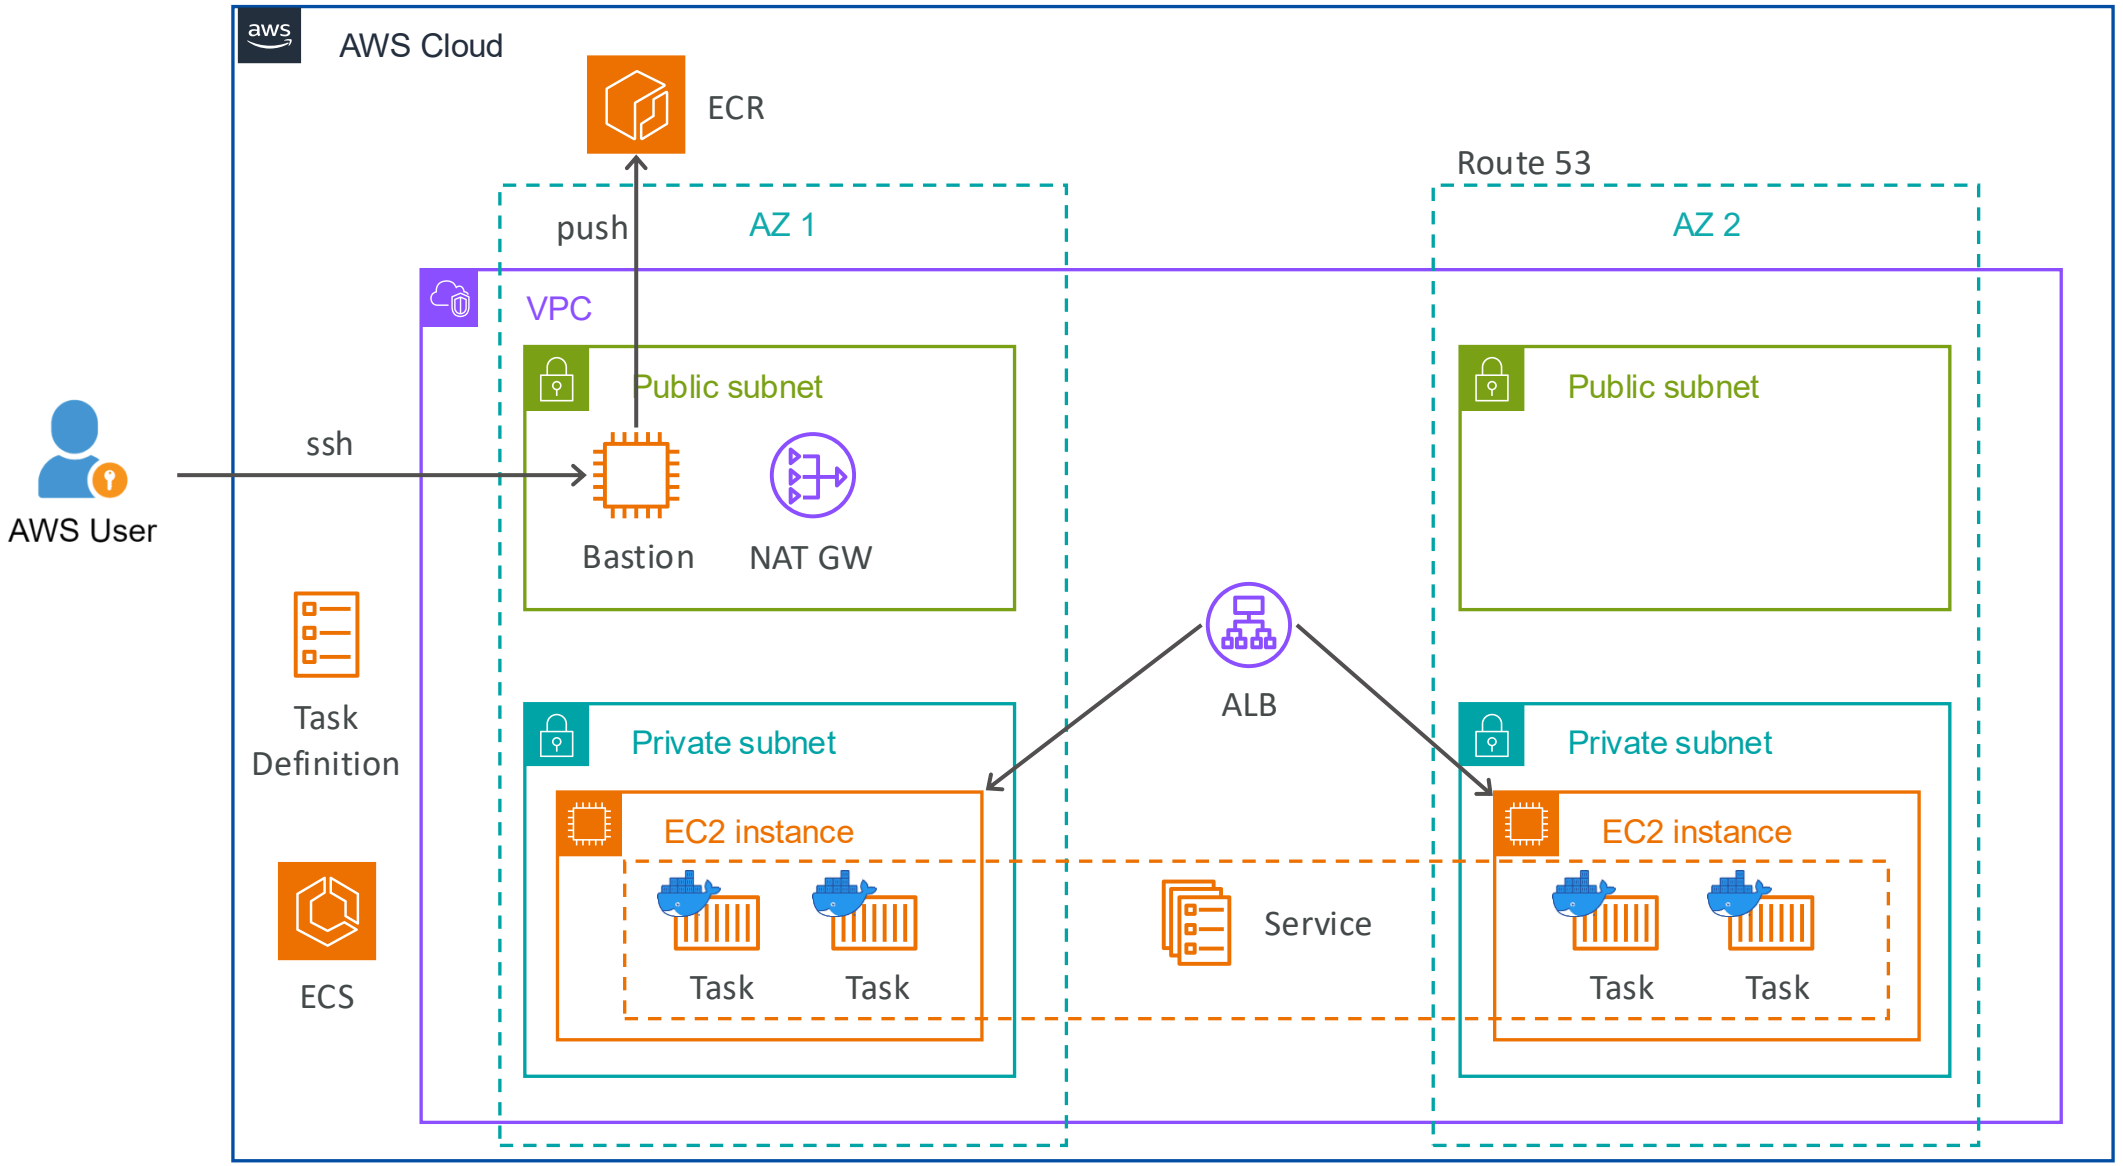

© Copyright by Cloud Mentor Pro | Confidential

9. Test app

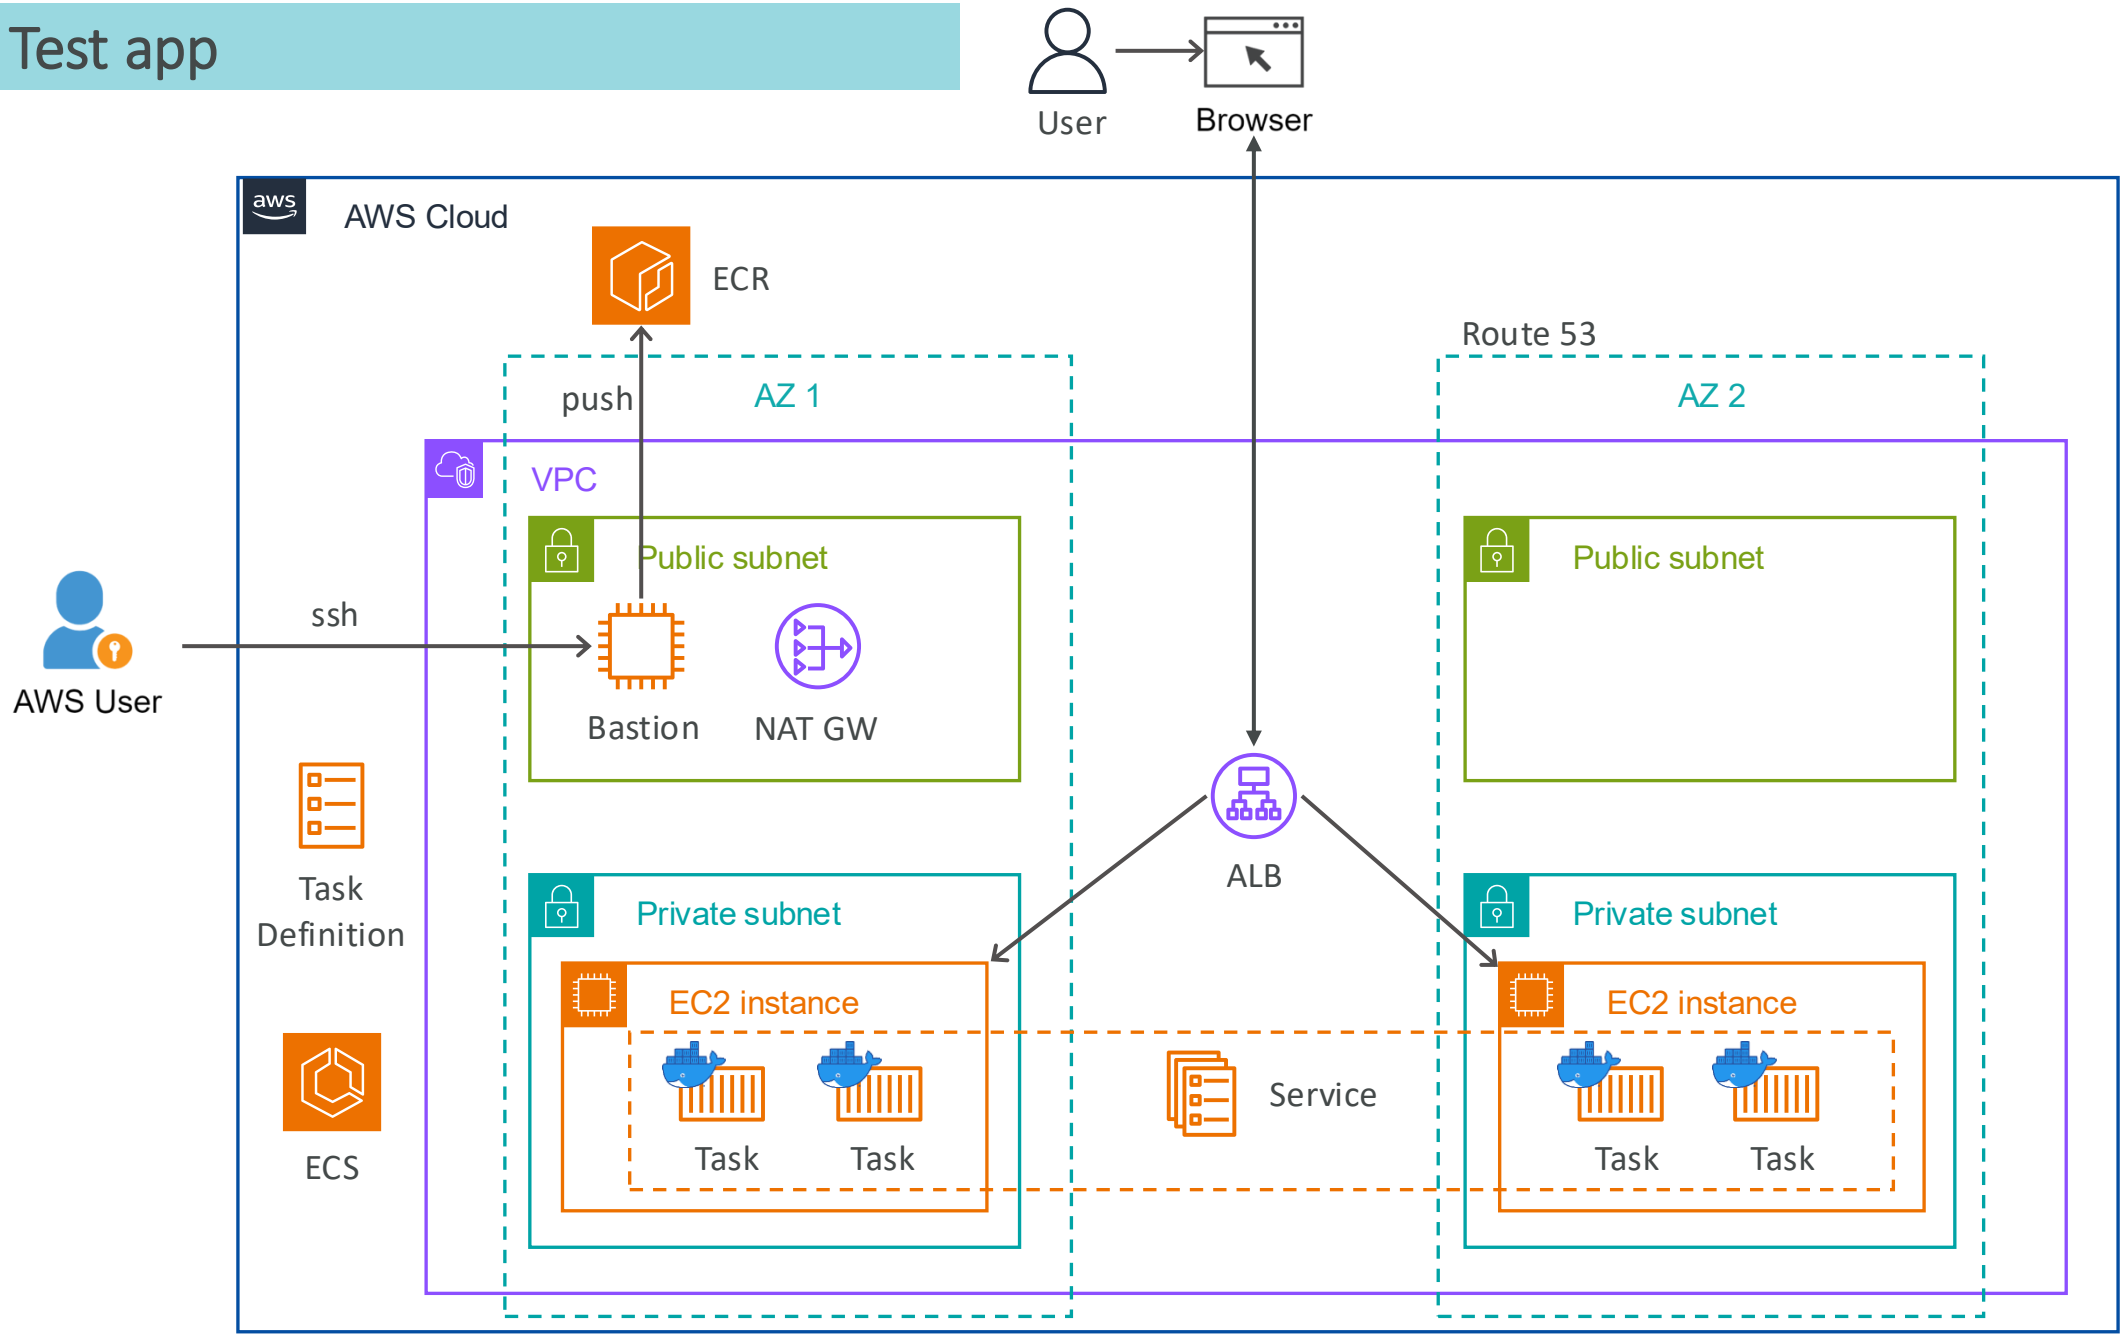

© Copyright by Cloud Mentor Pro | Confidential

10. Setting Route 53

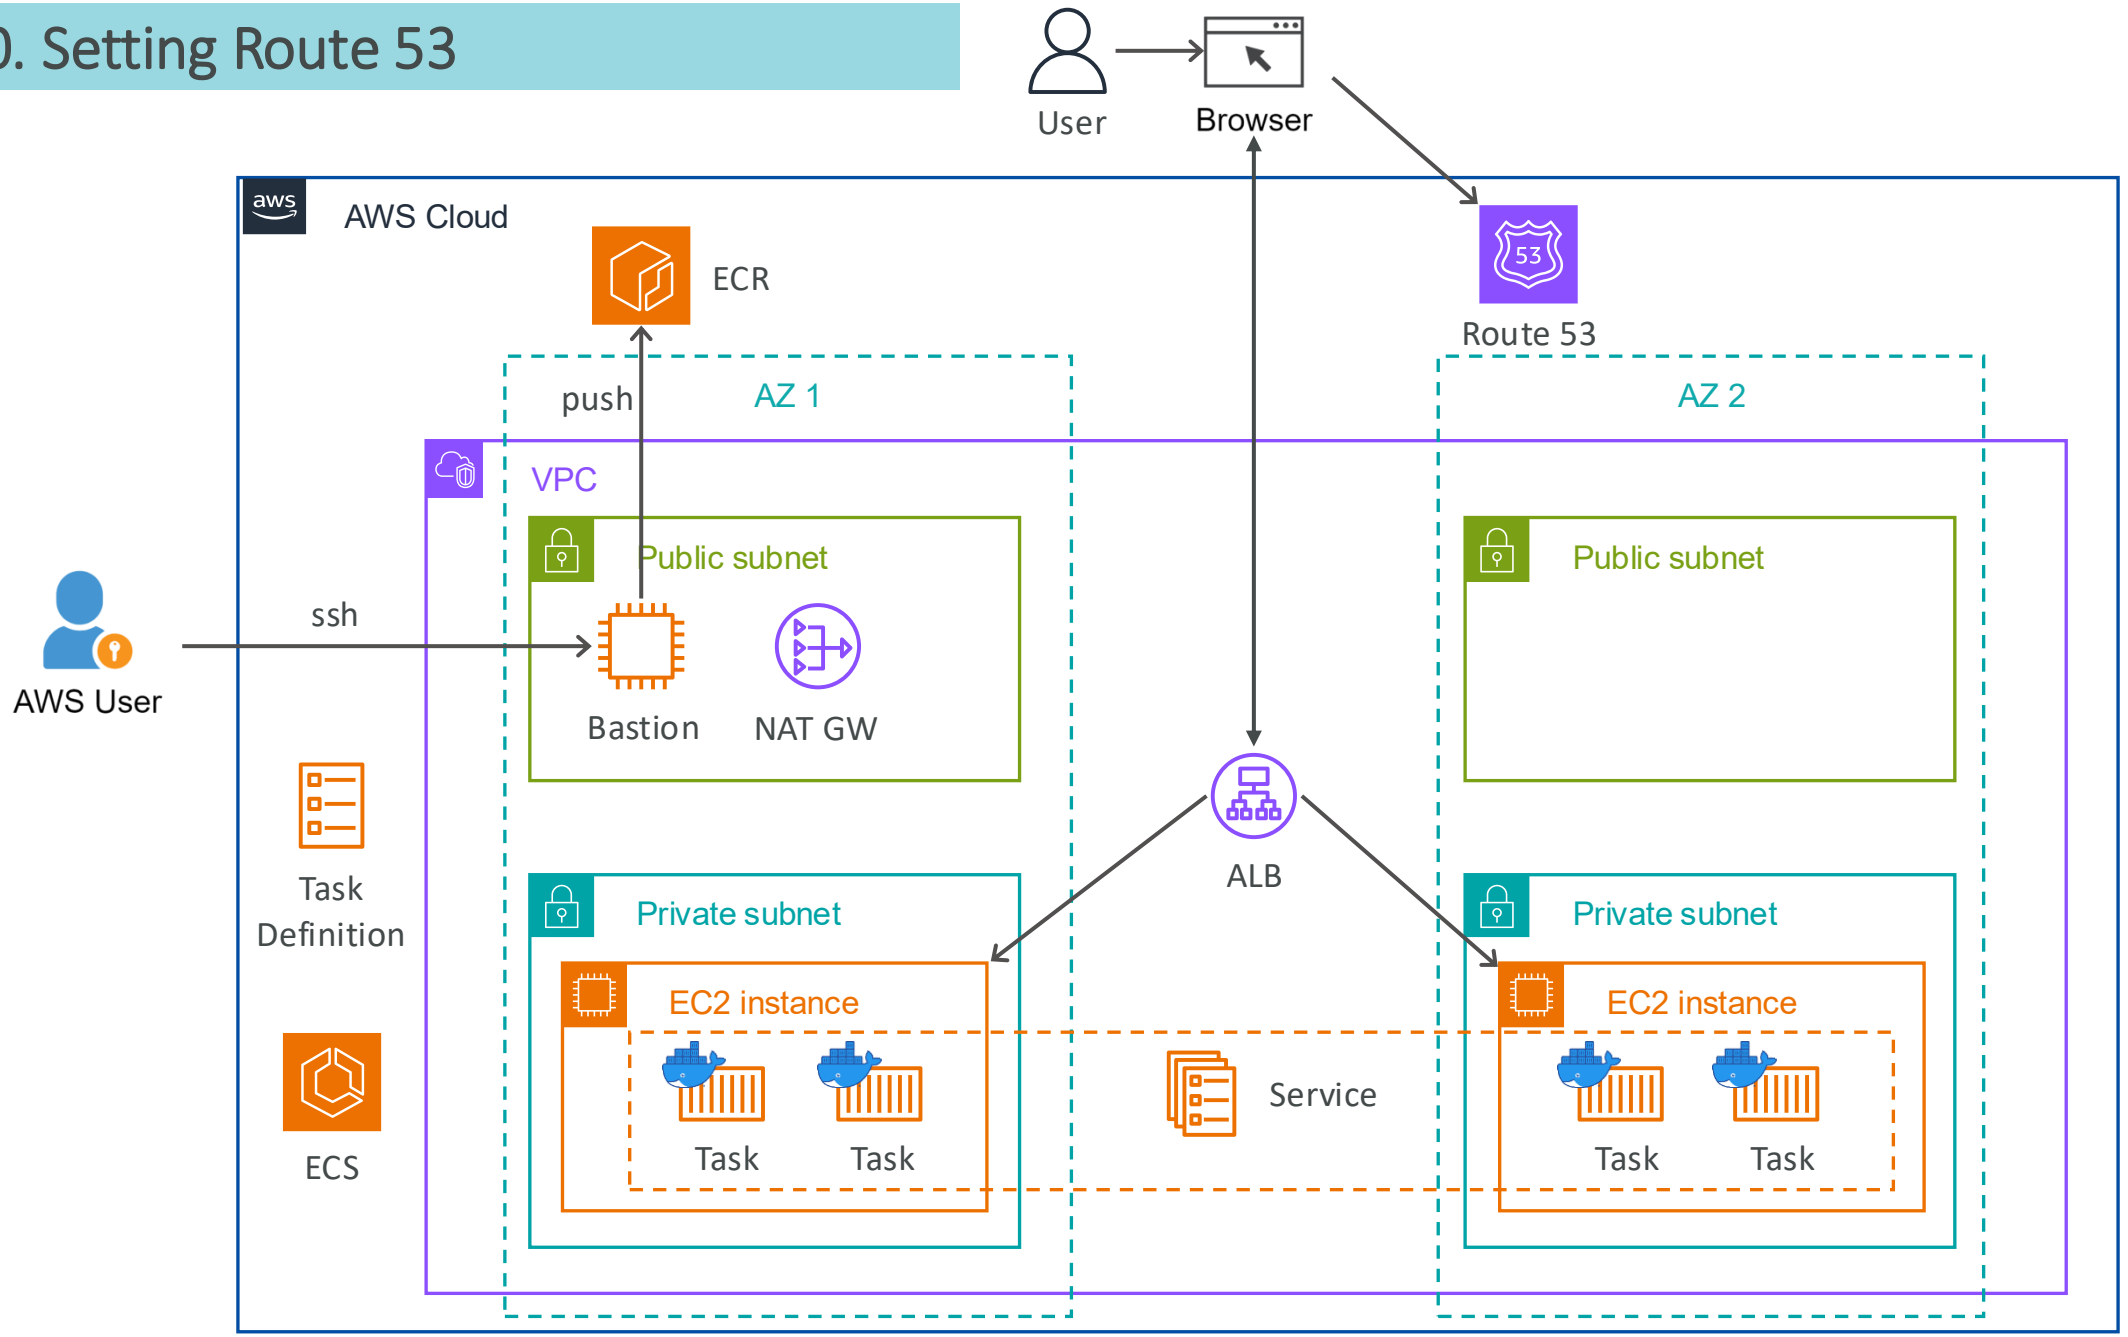

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- *How to Run a Application on AWS ECS: A Step-by-Step Guide*

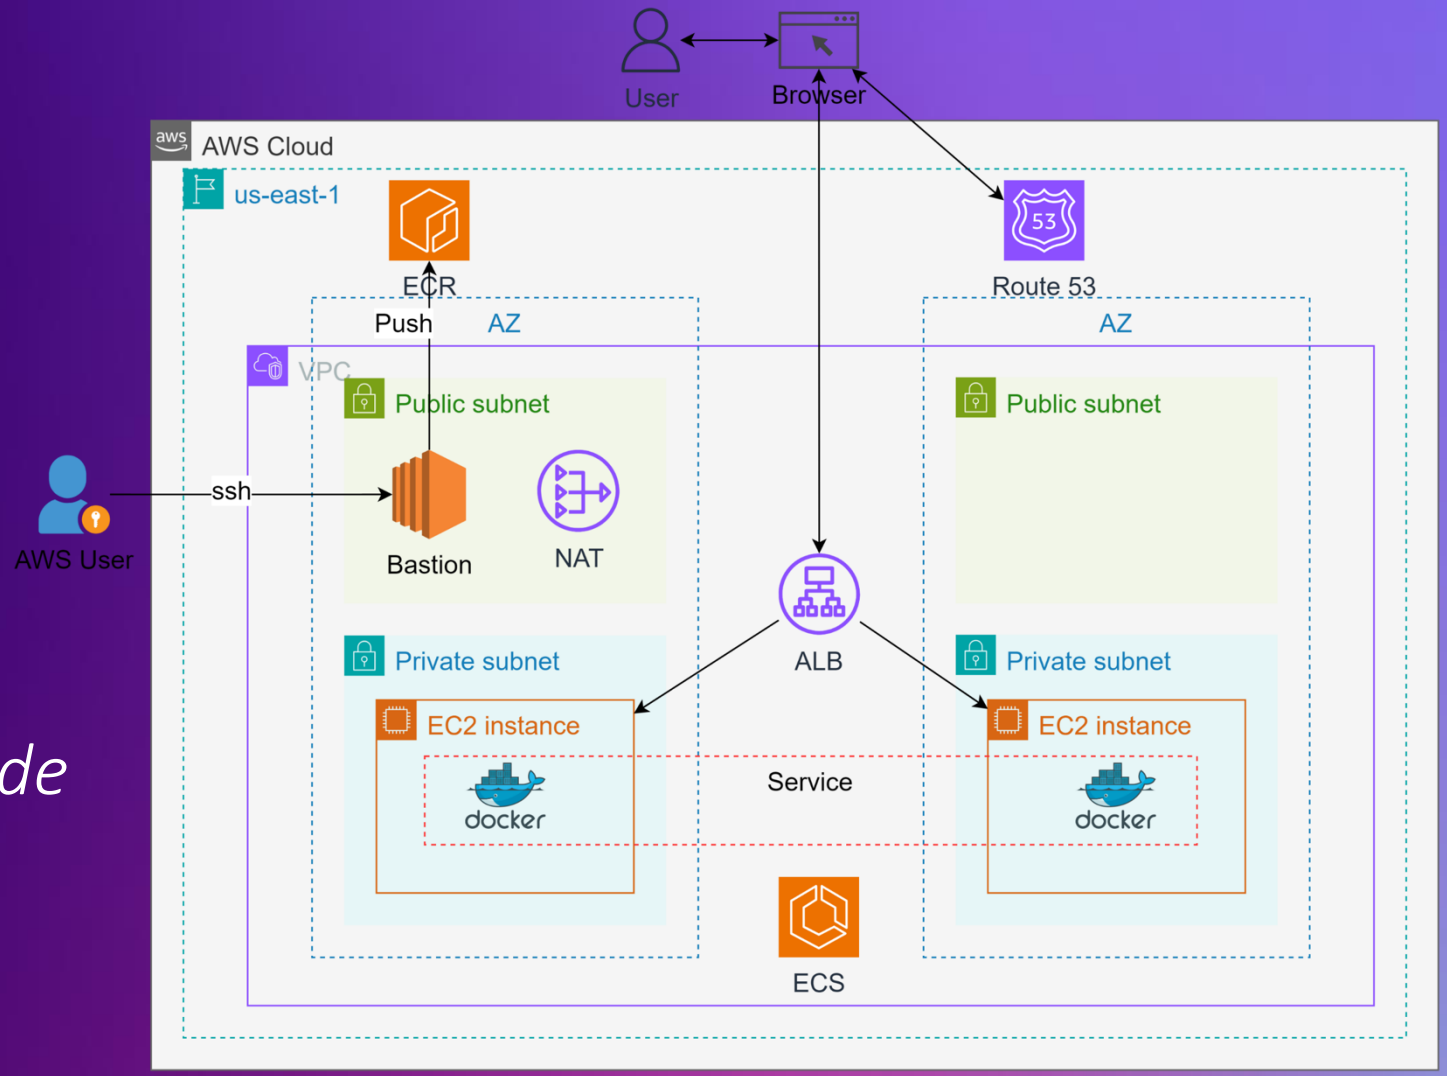

© Copyright by Cloud Mentor Pro | Confidential

Section 7

- *CloudFront*
- *AWS Elastic Beanstalk*

© Copyright by Cloud Mentor Pro | Confidential

Amazon CloudFront

© Copyright by Cloud Mentor Pro | Confidential

Amazon CloudFront

- Content Delivery Network (CDN)
- **Improves read performance, content is cached at the edge**
- Improves users experience
- 600+ Point of Presence globally (edge locations)
- **DDoS protection (because worldwide), integration with Shield, AWS Web Application Firewall**

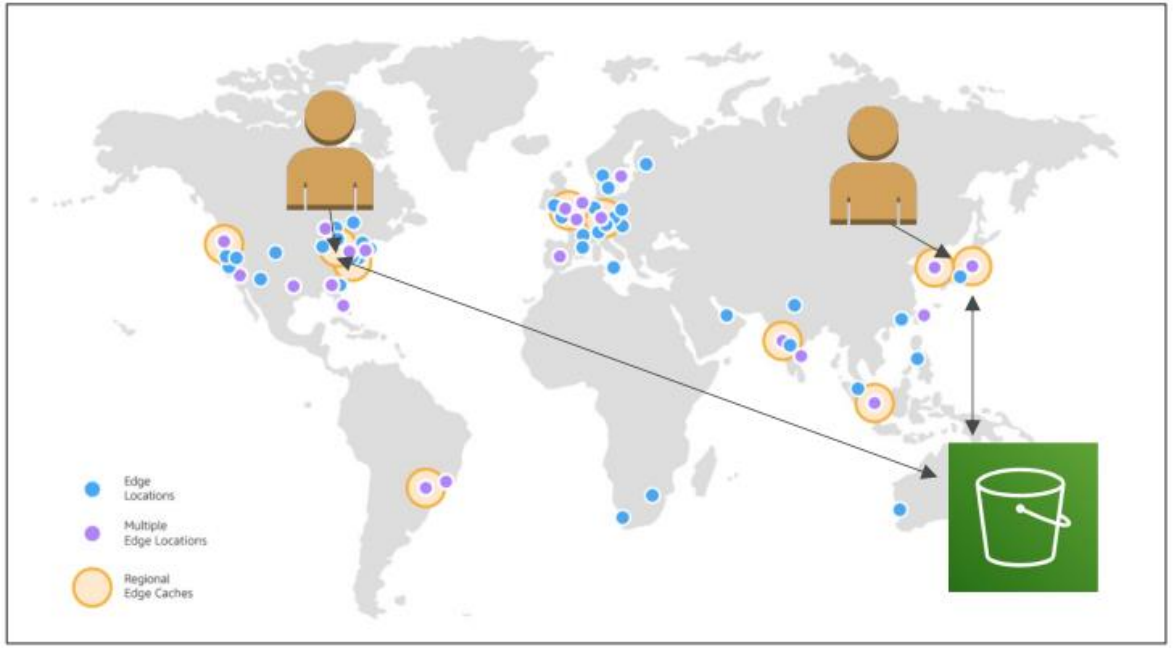

Source: https://aws.amazon.com/cloudfront/features/?nc=sn&loc=2

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Origins

- S3 bucket
  - For distributing files and caching them at the edge
  - Enhanced security with CloudFront **Origin Access Control (OAC)**
  - OAC is replacing Origin Access Identity (OAI)
  - CloudFront can be used as an ingress (to upload files to S3)
- Custom Origin (HTTP)
  - Application Load Balancer
  - EC2 instance
  - S3 website (must first enable the bucket as a static S3 website)
  - Any HTTP backend you want

© Copyright by Cloud Mentor Pro | Confidential

CloudFront at a high level

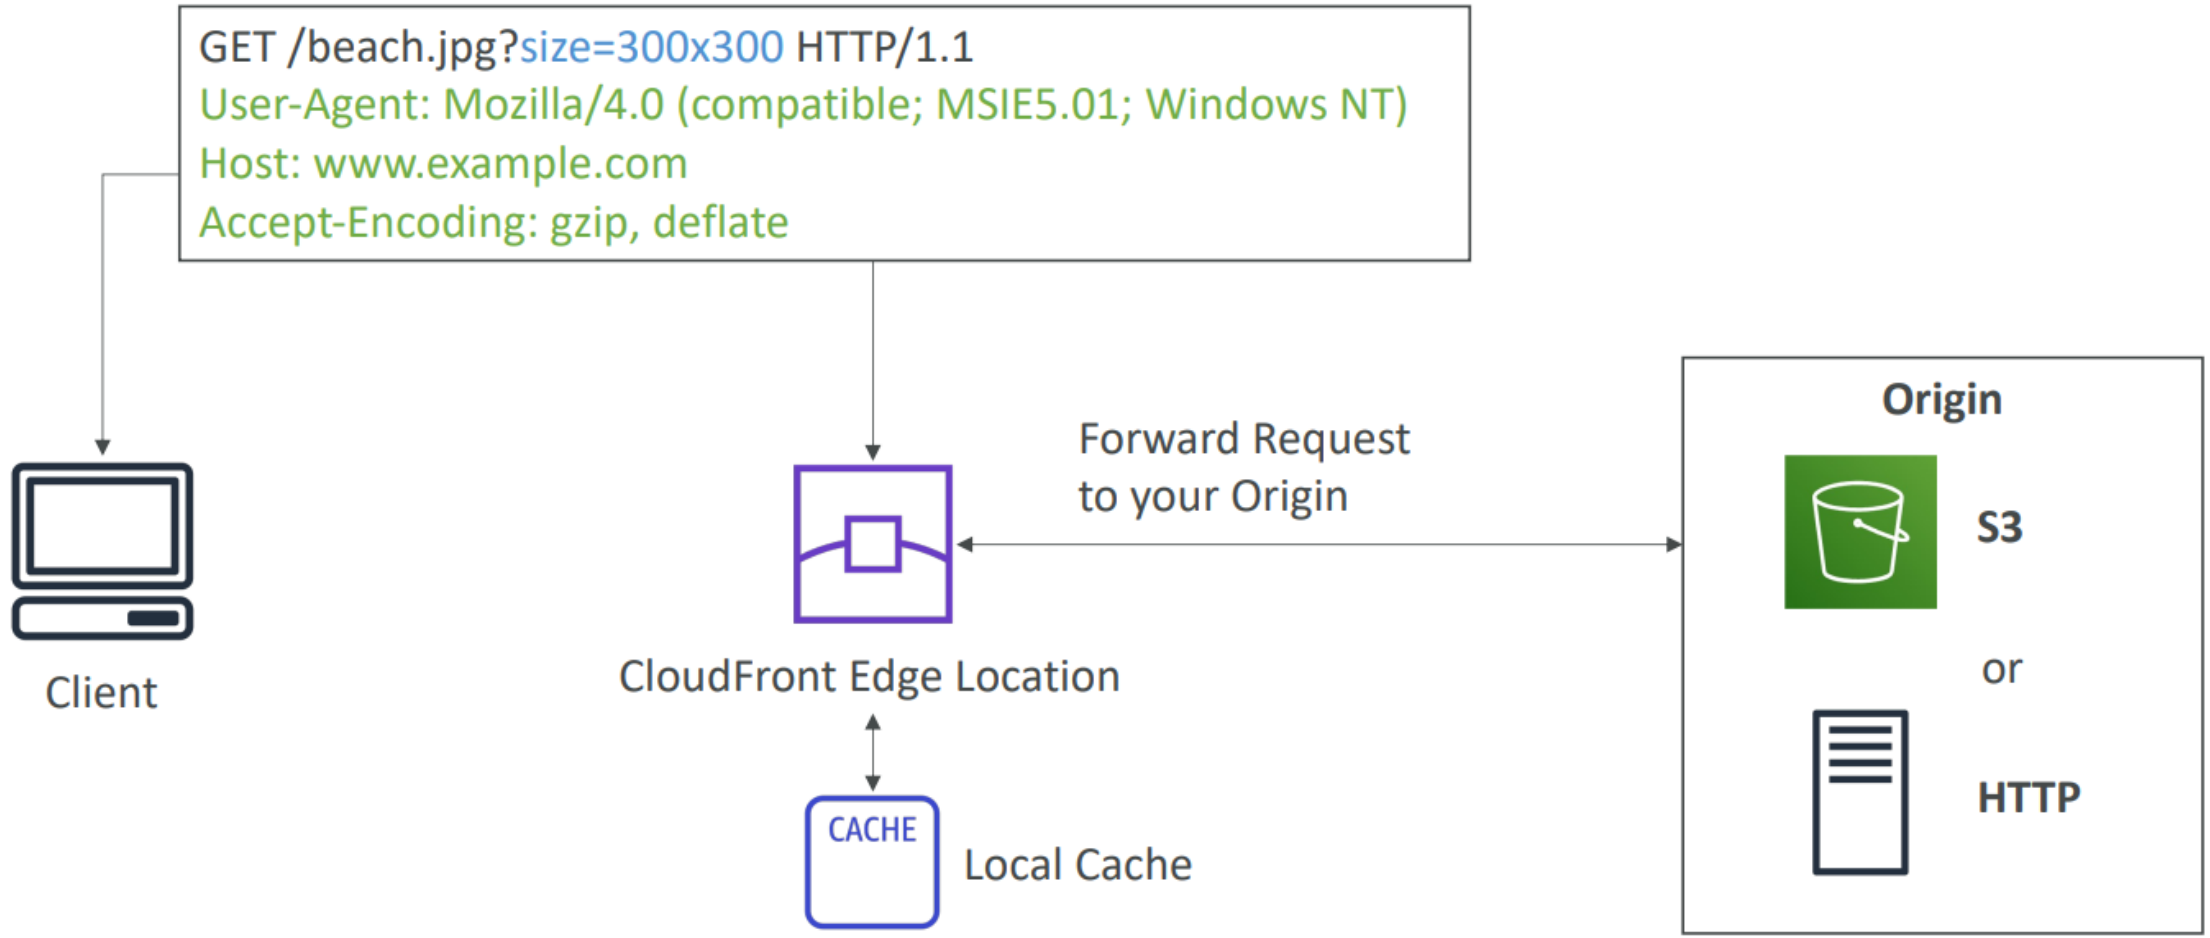

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – S3 as an Origin

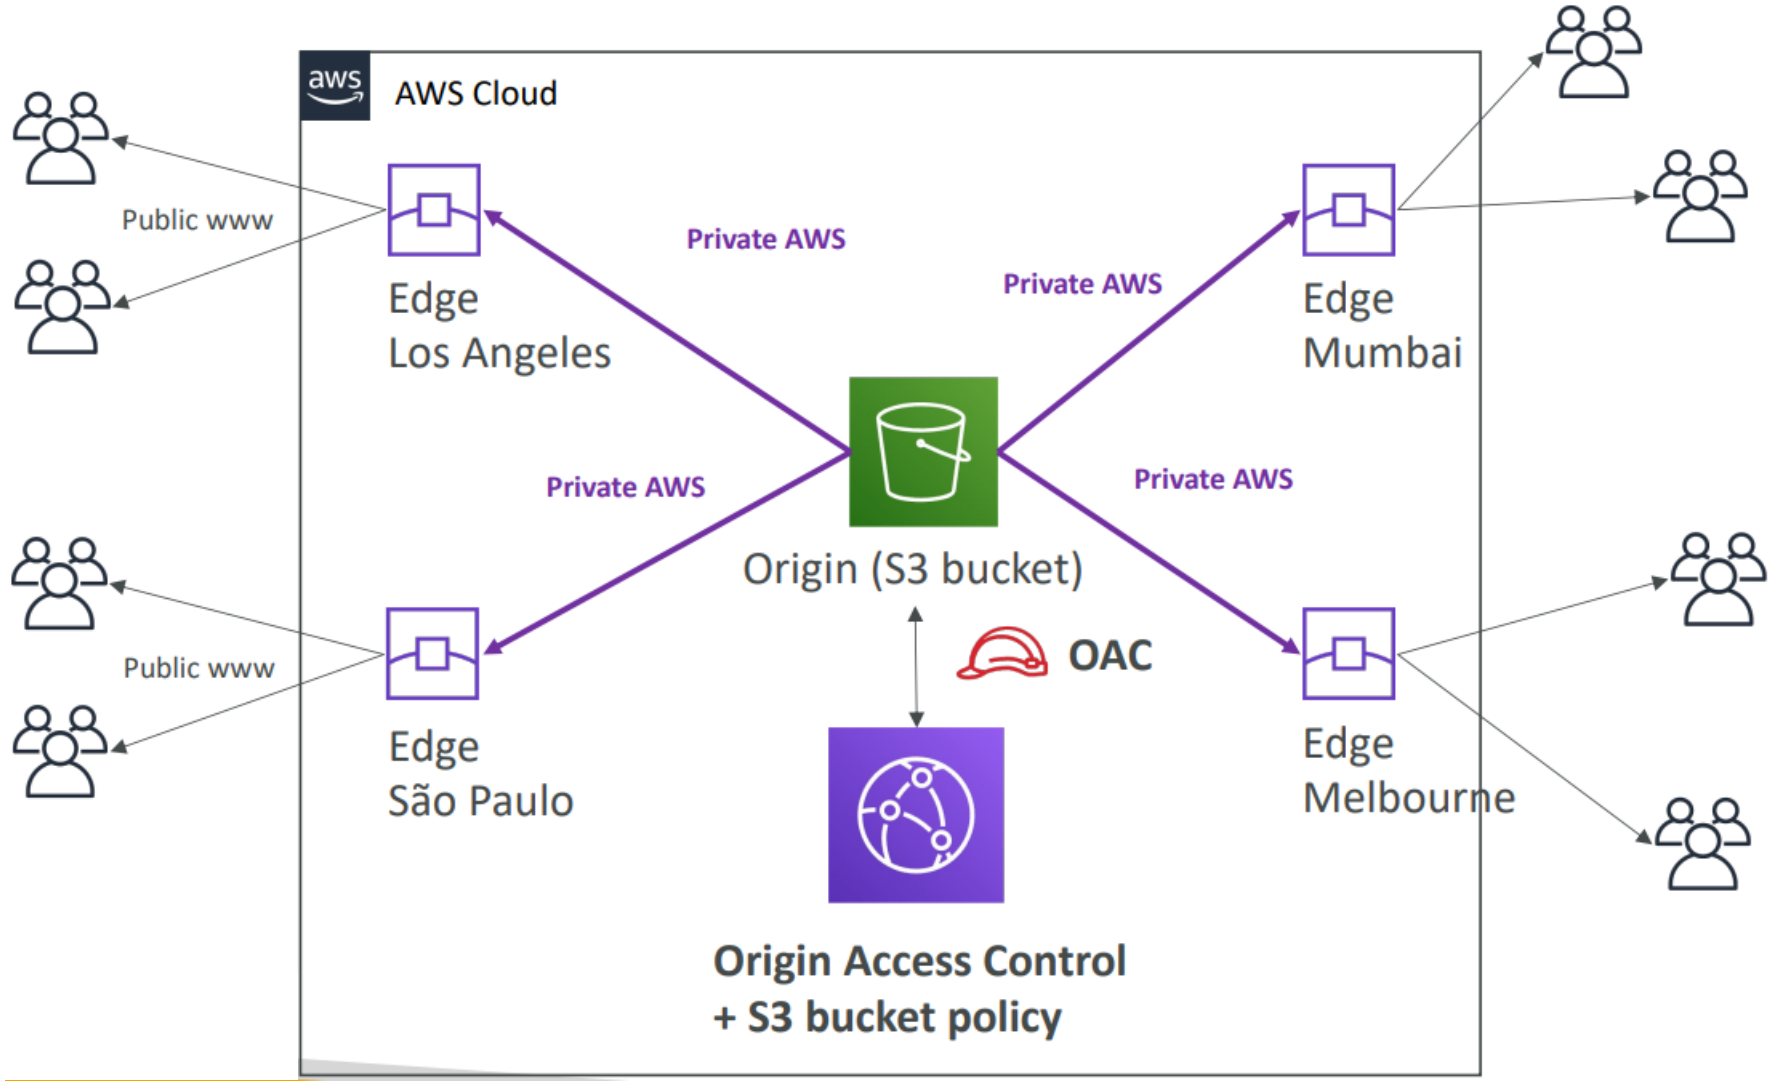

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Origin Groups

- To increase high-availability and do failover
- Origin Group: one primary and one secondary origin
- If the primary origin fails, the second one is used

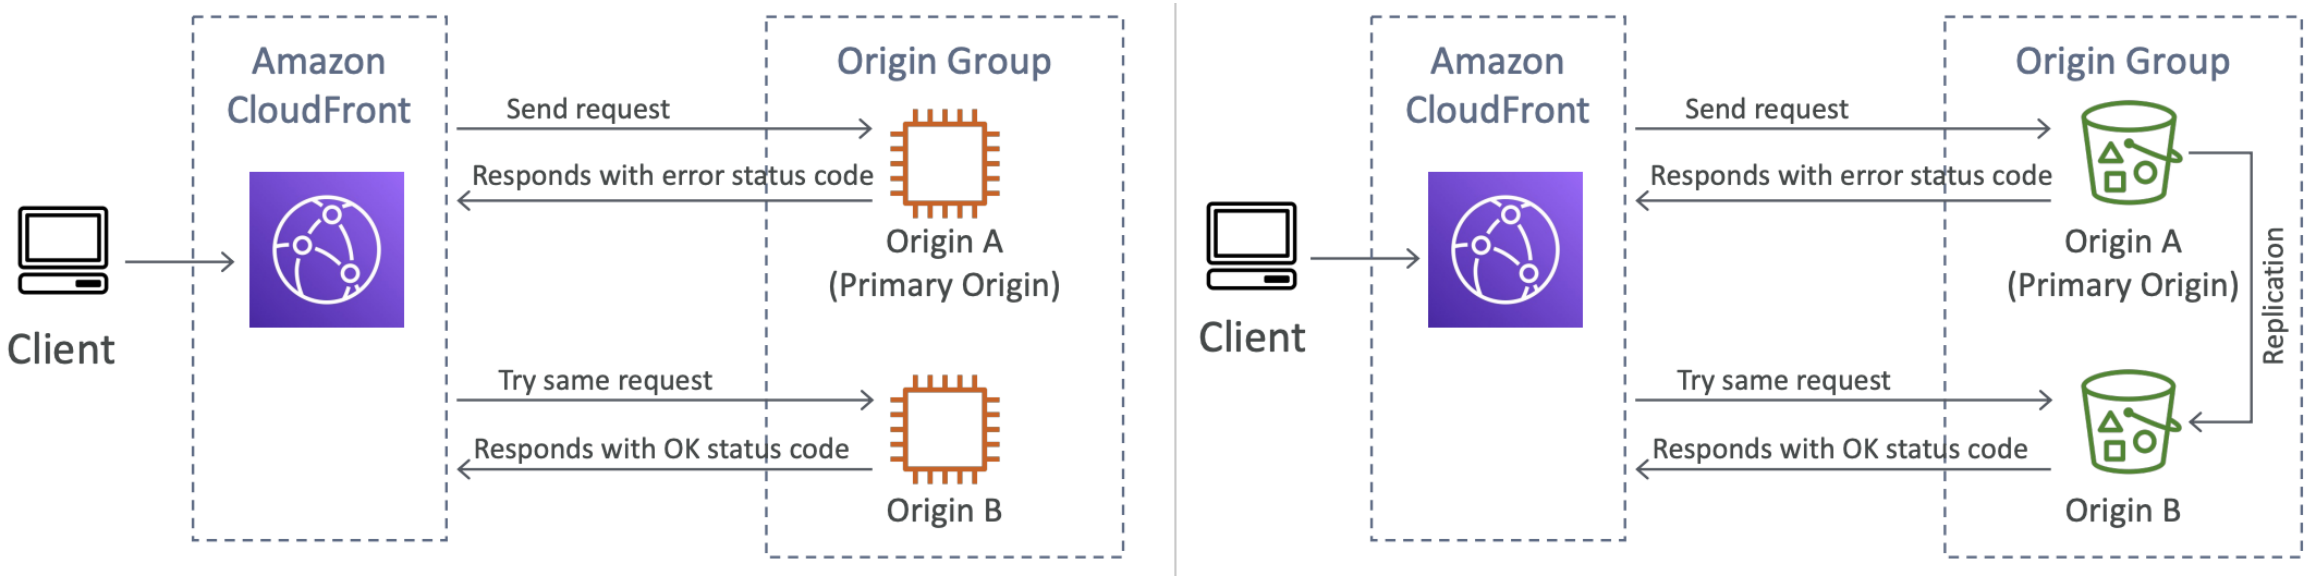

© Copyright by Cloud Mentor Pro | Confidential

CloudFront vs S3 Cross Region Replication

- CloudFront:
  - Global Edge network
  - Files are cached for a TTL (maybe a day)
  - **Great for static content that must be available everywhere**

- S3 Cross Region Replication:
  - Must be setup for each region you want replication to happen
  - Files are updated in near real-time
  - Read only
  - Great for dynamic content that needs to be available at low-latency in few regions

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Caching

- The cache lives at each CloudFront Edge Location
- CloudFront identifies each object in the cache using the Cache Key (see next slide)
- You want to maximize the Cache Hit ratio to minimize requests to the origin
- You can invalidate part of the cache using the CreateInvalidation API

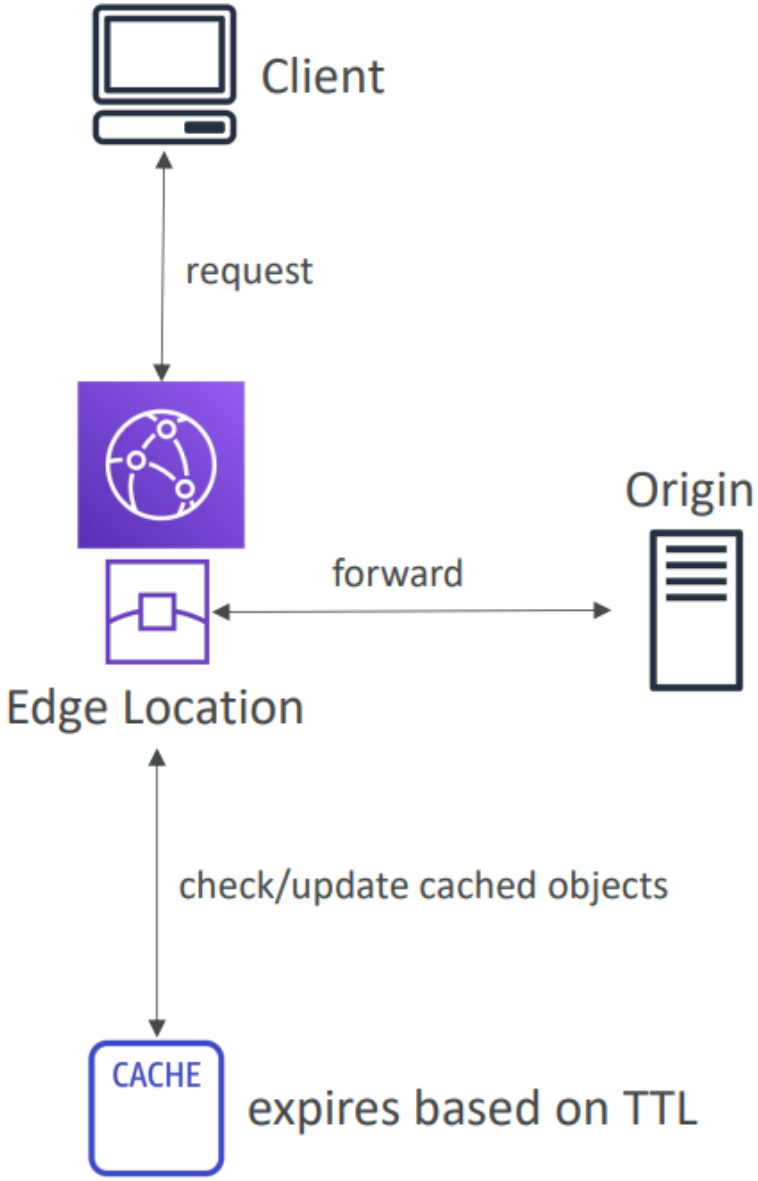

© Copyright by Cloud Mentor Pro | Confidential

What is CloudFront Cache Key?

- A unique identifier for every object in the cache
- By default, consists of hostname + resource portion of the URL
- If you have an application that serves up content that varies based on user, device, language, location…
- You can add other elements (HTTP headers, cookies, query strings) to the Cache Key using CloudFront Cache Policies

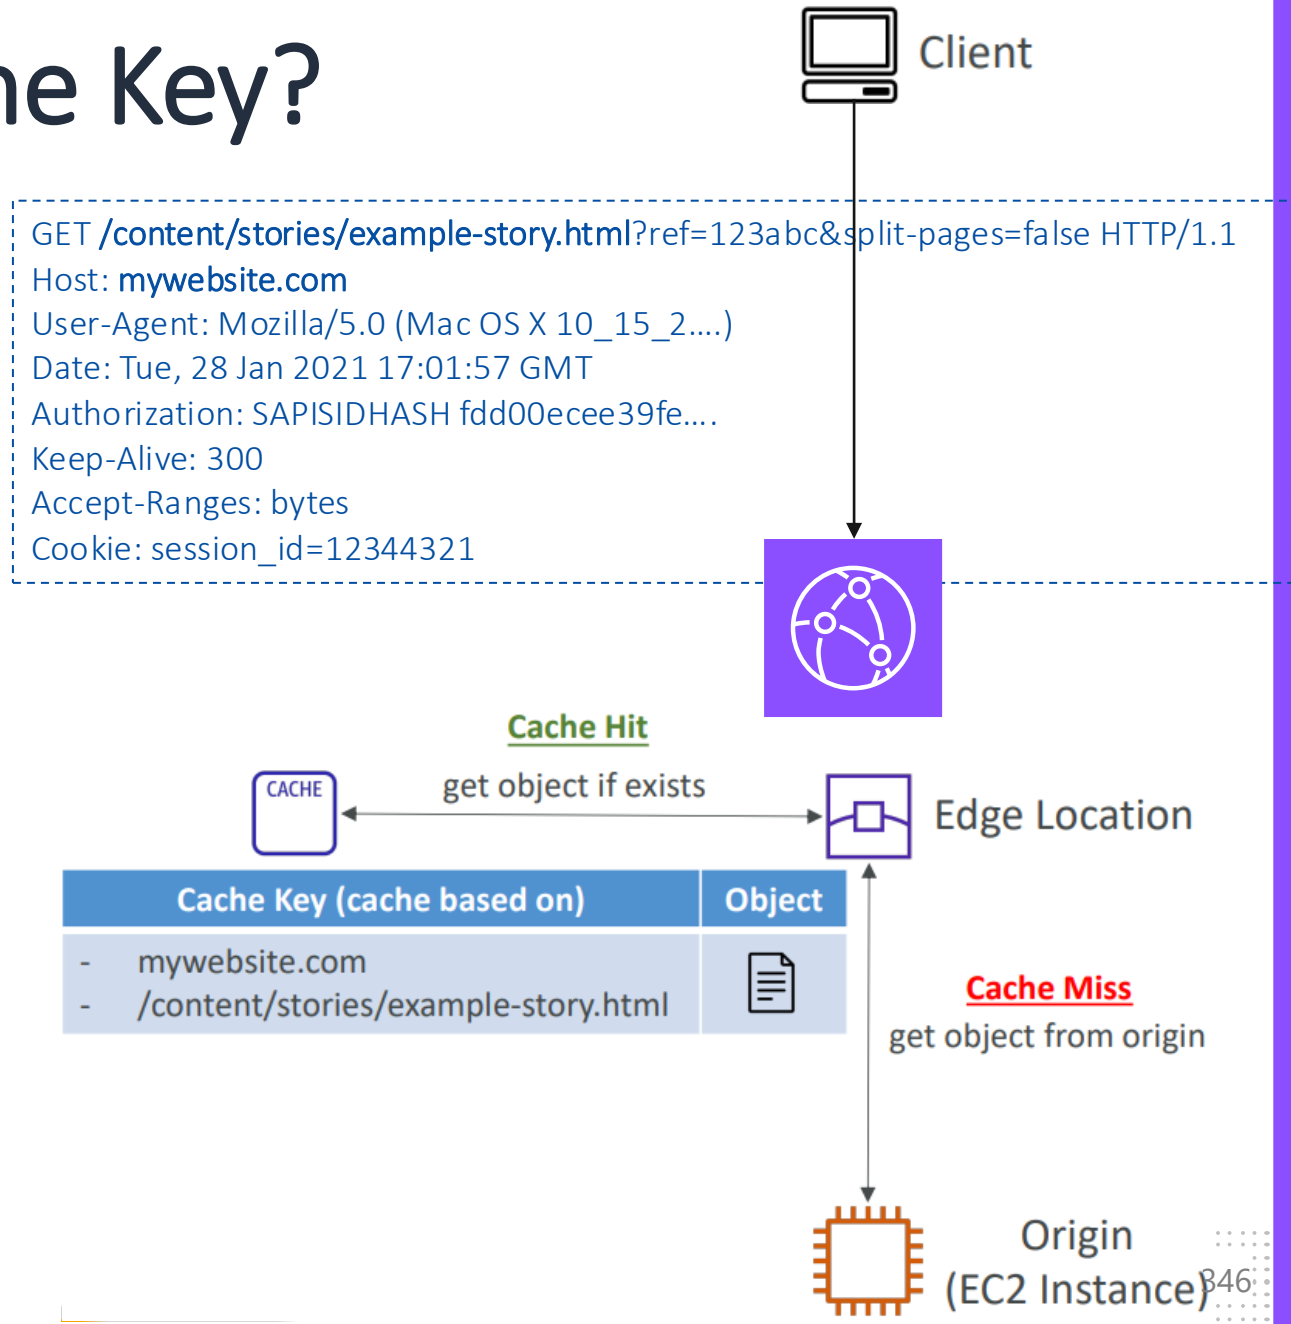

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Cache Invalidations

- In case you update the back-end origin, CloudFront doesn't know about it and will only get the refreshed content after the TTL has expired
- However, you can force an entire or partial cache refresh (thus bypassing the TTL) by performing a CloudFront Invalidation
- You can invalidate all files (*) or a special path (/images/*

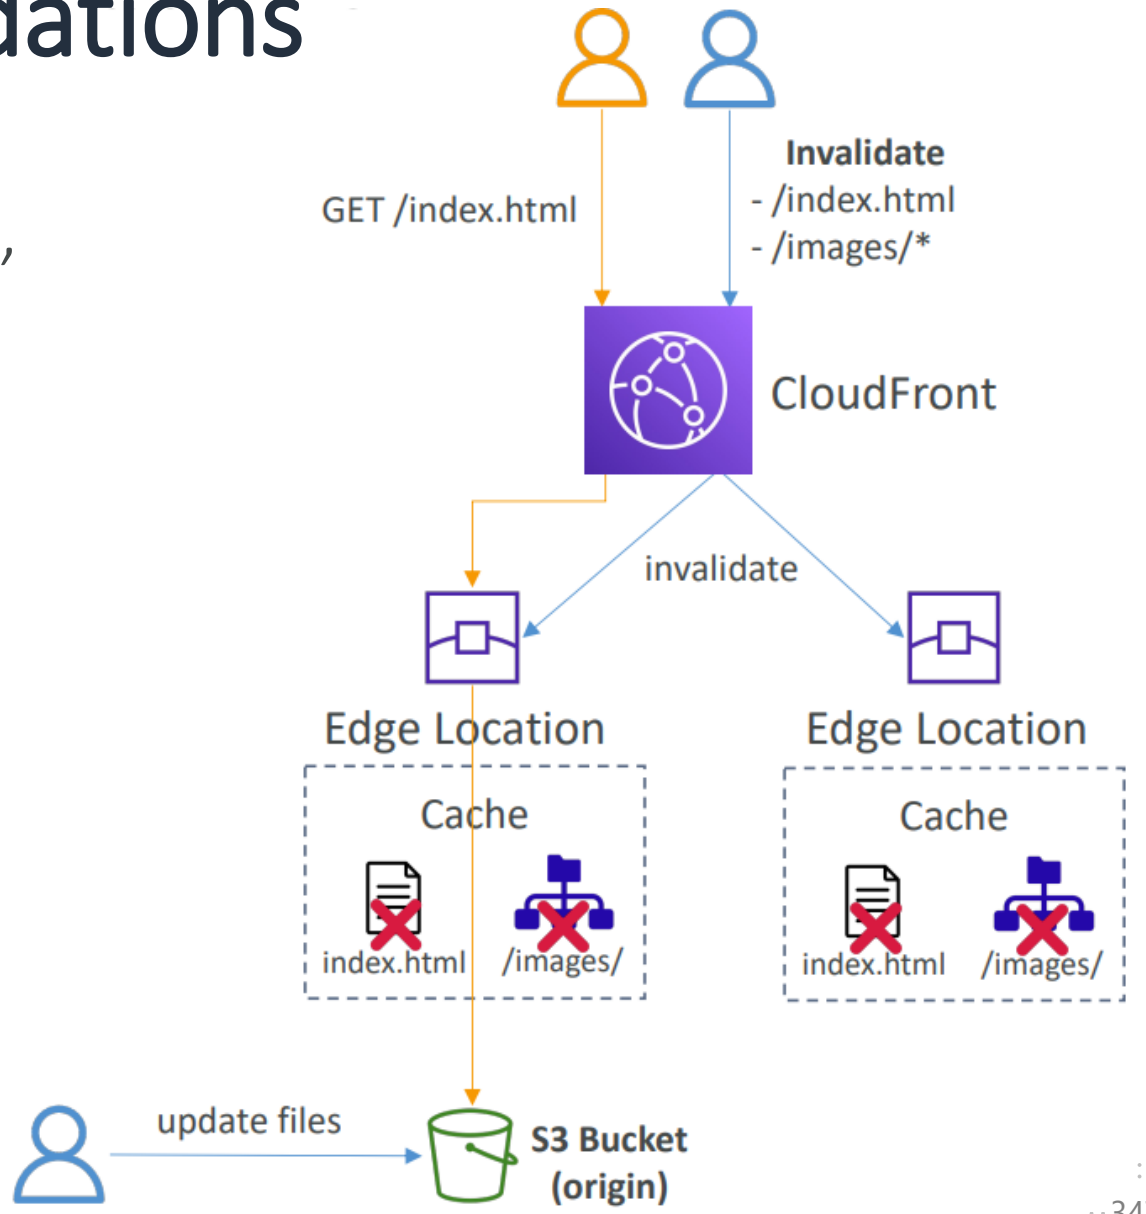

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Cache Behaviors

- Configure different settings for a given URL path pattern
- Example: one specific cache behavior to images/*.jpg files on your origin web server
- Route to different kind of origins/origin groups based on the content type or path pattern
  - /images/*
  - /api/*
  - /* (default cache behavior)
- When adding additional Cache Behaviors, the Default Cache Behavior is always the last to be processed **and is always /***

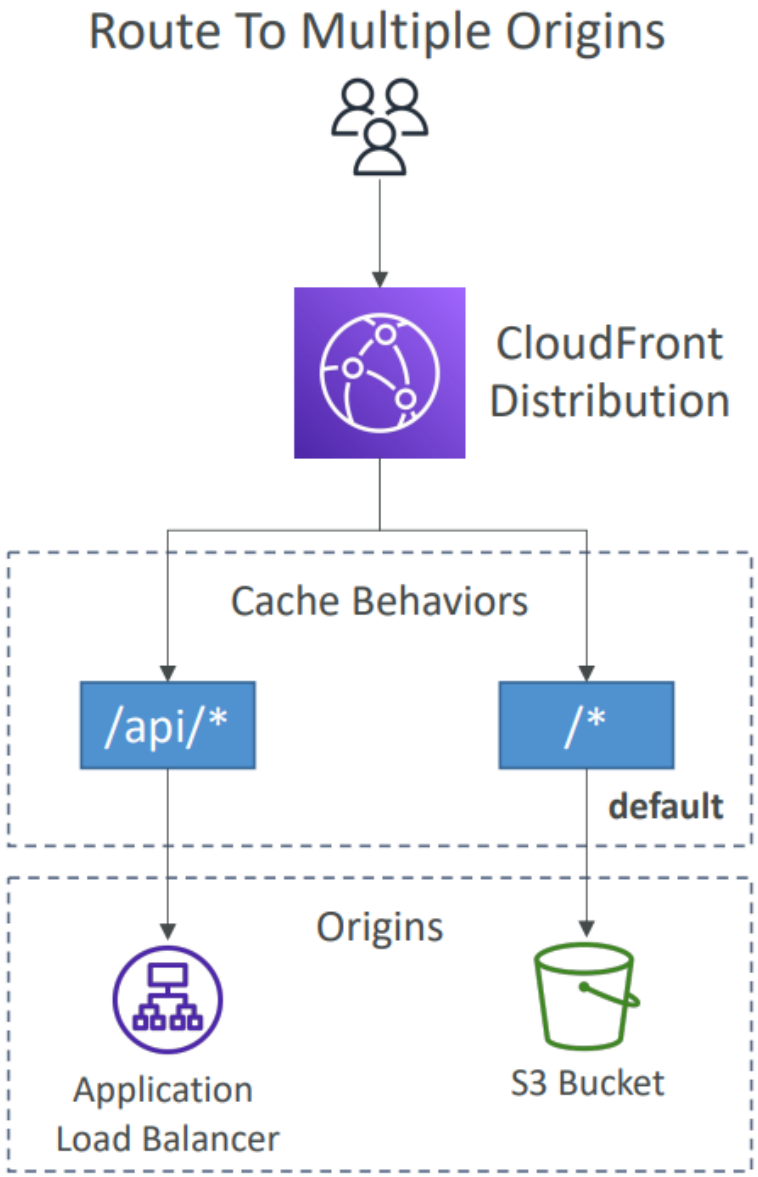

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – ALB or EC2 as an origin Using VPC Origins

- Allows you to deliver content from your applications hosted in your VPC private subnets (no need to expose them on the Internet)
- Deliver traffic to private:
  - Application Load Balancer
  - Network Load Balancer
  - EC2 Instances

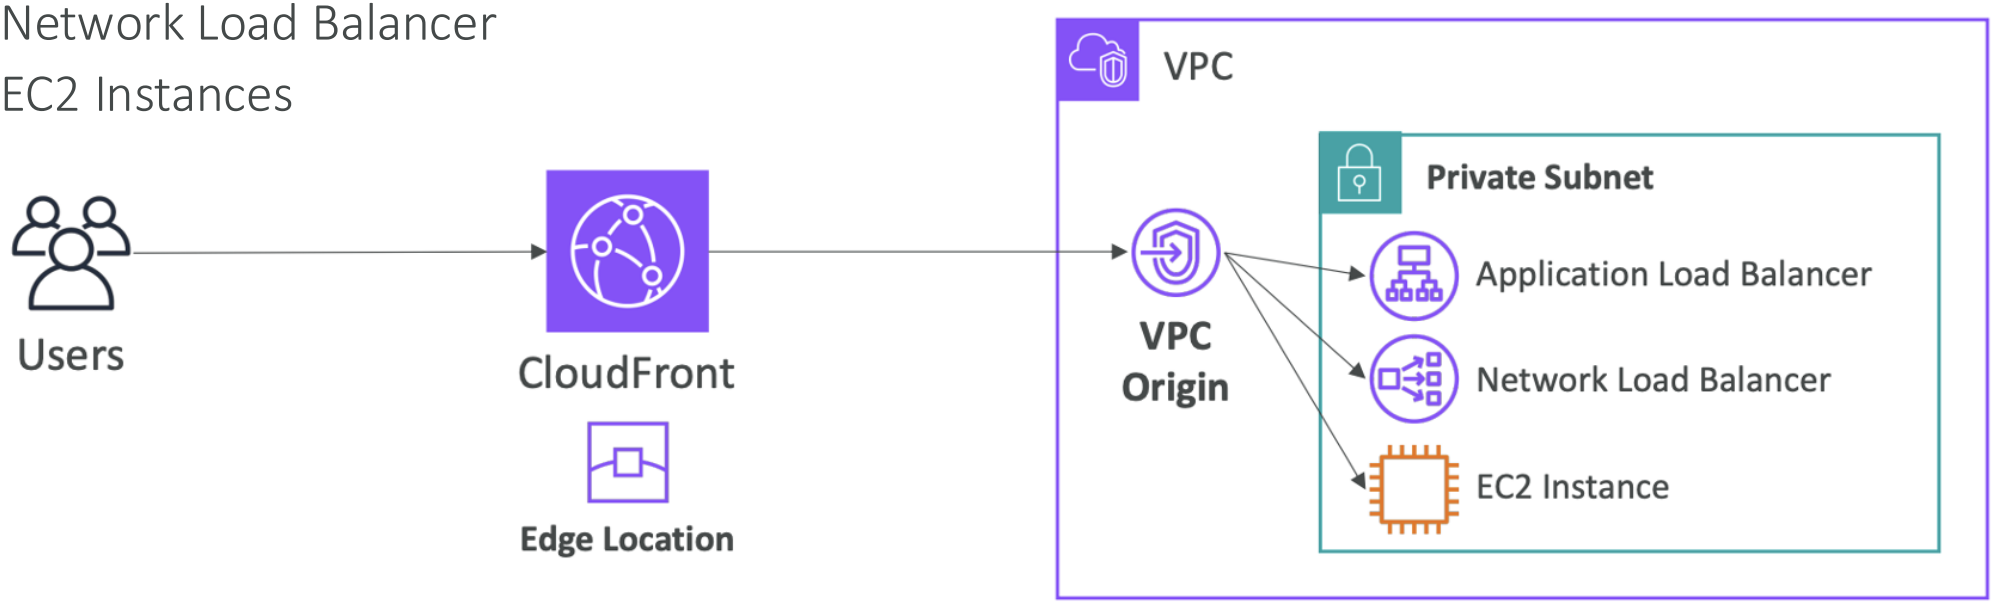

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – ALB or EC2 as an origin

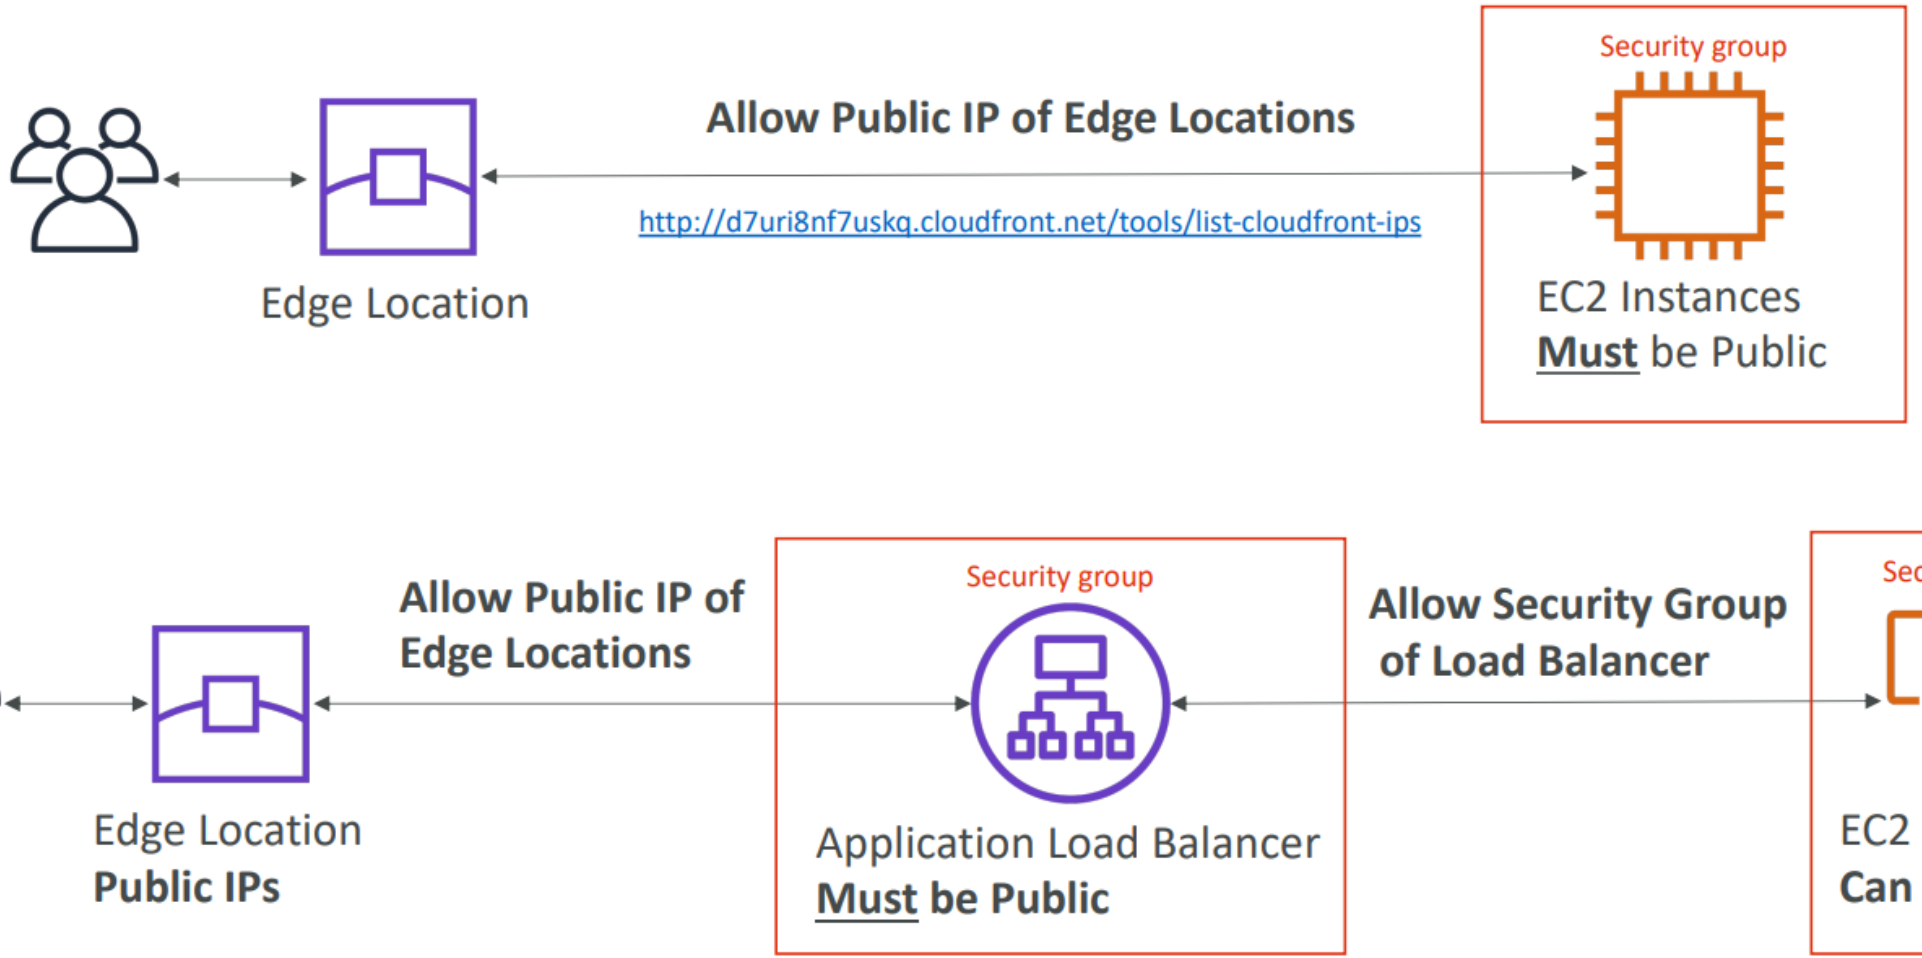

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Geo Restriction

- You can restrict who can access your distribution
  - **Allowlist**: Allow your users to access your content only if they're in one of the countries on a list of approved countries.
  - **Blocklist**: Prevent your users from accessing your content if they're in one of the countries on a list of banned countries.

- The “country” is determined using a 3rd party Geo-IP database
- Use case: Copyright Laws to control access to content

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Signed URL / Signed Cookies

- You want to distribute paid shared content to premium users
- We can use CloudFront Signed URL / Cookie. We attach a policy with:
  - Includes URL expiration
  - Includes IP ranges to access the data from
  - Trusted signers (which AWS accounts can create signed URLs)
- How long should the URL be valid for?
  - Shared content (movie, music): make it short (a few minutes)
  - Private content (private to the user): you can make it last for years
- Signed URL = access to individual files (one signed URL per file)
- Signed Cookies = access to multiple files (one signed cookie for many files)

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Signed URL Diagram

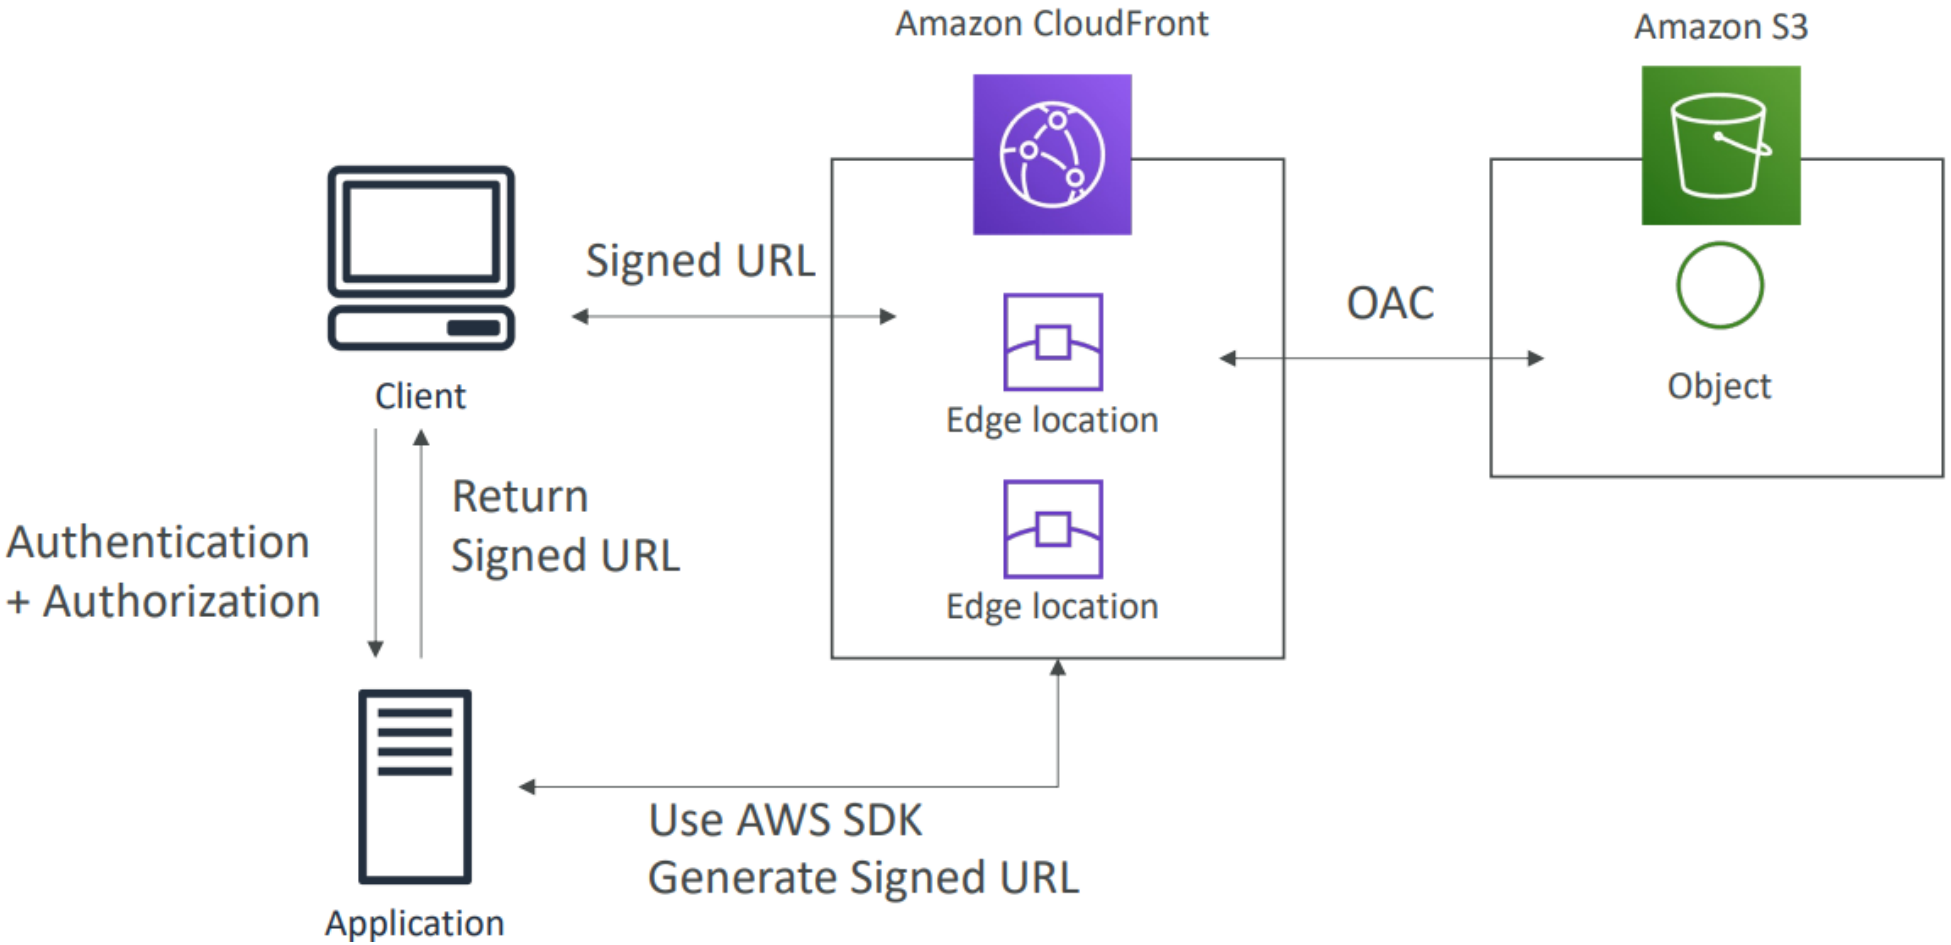

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Field Level Encryption

- Protect user sensitive information through application stack
- Adds an additional layer of security along with HTTPS
- Sensitive information encrypted at the edge close to user
- Uses asymmetric encryption
- Usage:
  - Specify set of fields in POST requests that you want to be encrypted (up to 10 fields)
  - Specify the public key to encrypt them

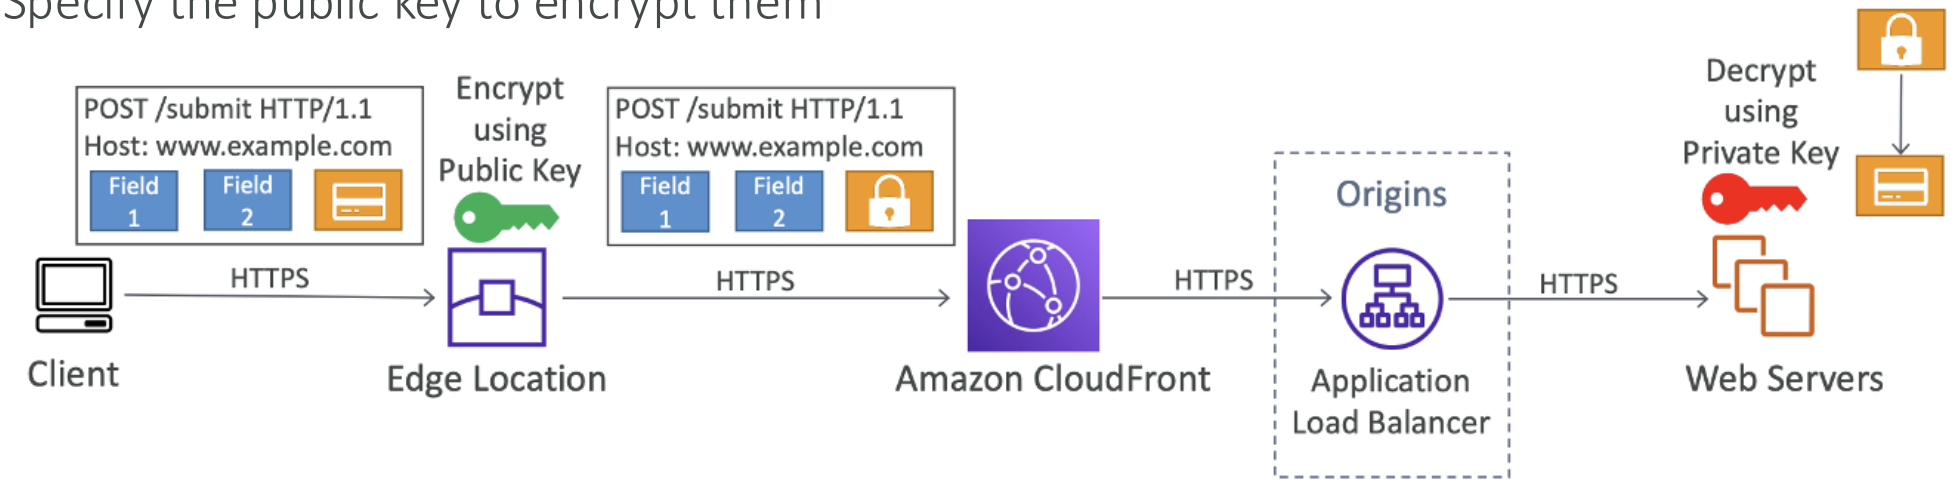

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Real Time Logs

- Get real-time requests received by CloudFront sent to Kinesis Data Streams
- Monitor, analyze, and take actions based on content delivery performance
- Allows you to choose:
  - Sampling Rate – percentage of requests for which you want to receive
  - Specific fields and specific Cache Behaviors (path patterns)

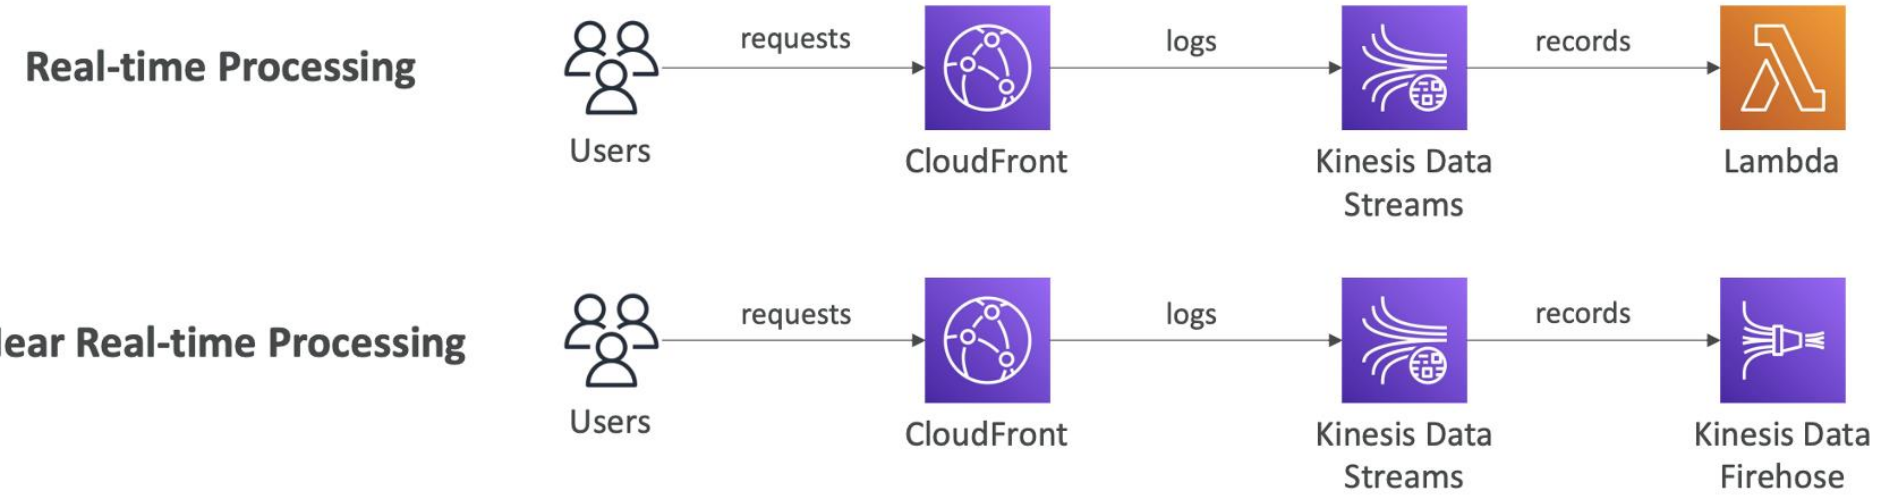

© Copyright by Cloud Mentor Pro | Confidential

CloudFront - Pricing

- CloudFront Edge locations are all around the world
- The cost of data out per edge location varies

Per Month	United States, Mexico, & Canada	Europe & Israel	South Africa, Kenya, & Middle East	South America	Japan	Australia & New Zealand	Hong Kong, Philippines, Singapore, South Korea, Taiwan, & Thailand	India
First 10TB	$0.085	$0.085	$0.110	$0.110	$0.114	$0.114	$0.140	$0.170
Next 40TB	$0.080	$0.080	$0.105	$0.105	$0.089	$0.098	$0.135	$0.130
Next 100TB	$0.060	$0.060	$0.090	$0.090	$0.086	$0.094	$0.120	$0.110
Next 350TB	$0.040	$0.040	$0.080	$0.080	$0.084	$0.092	$0.100	$0.100
Next 524TB	$0.030	$0.030	$0.060	$0.060	$0.080	$0.090	$0.080	$0.100
Next 4PB	$0.025	$0.025	$0.050	$0.050	$0.070	$0.085	$0.070	$0.100
Over 5PB	$0.020	$0.020	$0.040	$0.040	$0.060	$0.080	$0.060	$0.100

© Copyright by Cloud Mentor Pro | Confidential

CloudFront – Price Classes

- You can reduce the number of edge locations for cost reduction
- Three price classes:
  1. Price Class All: all regions – best performance
  2. Price Class 200: most regions, but excludes the most expensive regions
  3. Price Class 100: only the least expensive regions

Edge Locations Included Within	United States, Mexico, & Canada	Europe & Israel	South Africa, Kenya, & Middle East	South America	Japan	Australia & New Zealand	Hong Kong, Philippines, Singapore, South Korea, Taiwan, & Thailand	India
Price Class All	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Price Class 200	Yes	Yes	Yes	x	Yes	x	Yes	Yes
Price Class 100	Yes	Yes	x	x	x	x	x	x

© Copyright by Cloud Mentor Pro | Confidential

CloudFront - Price Class

Prices Class 100

Prices Class 200

Prices Class All

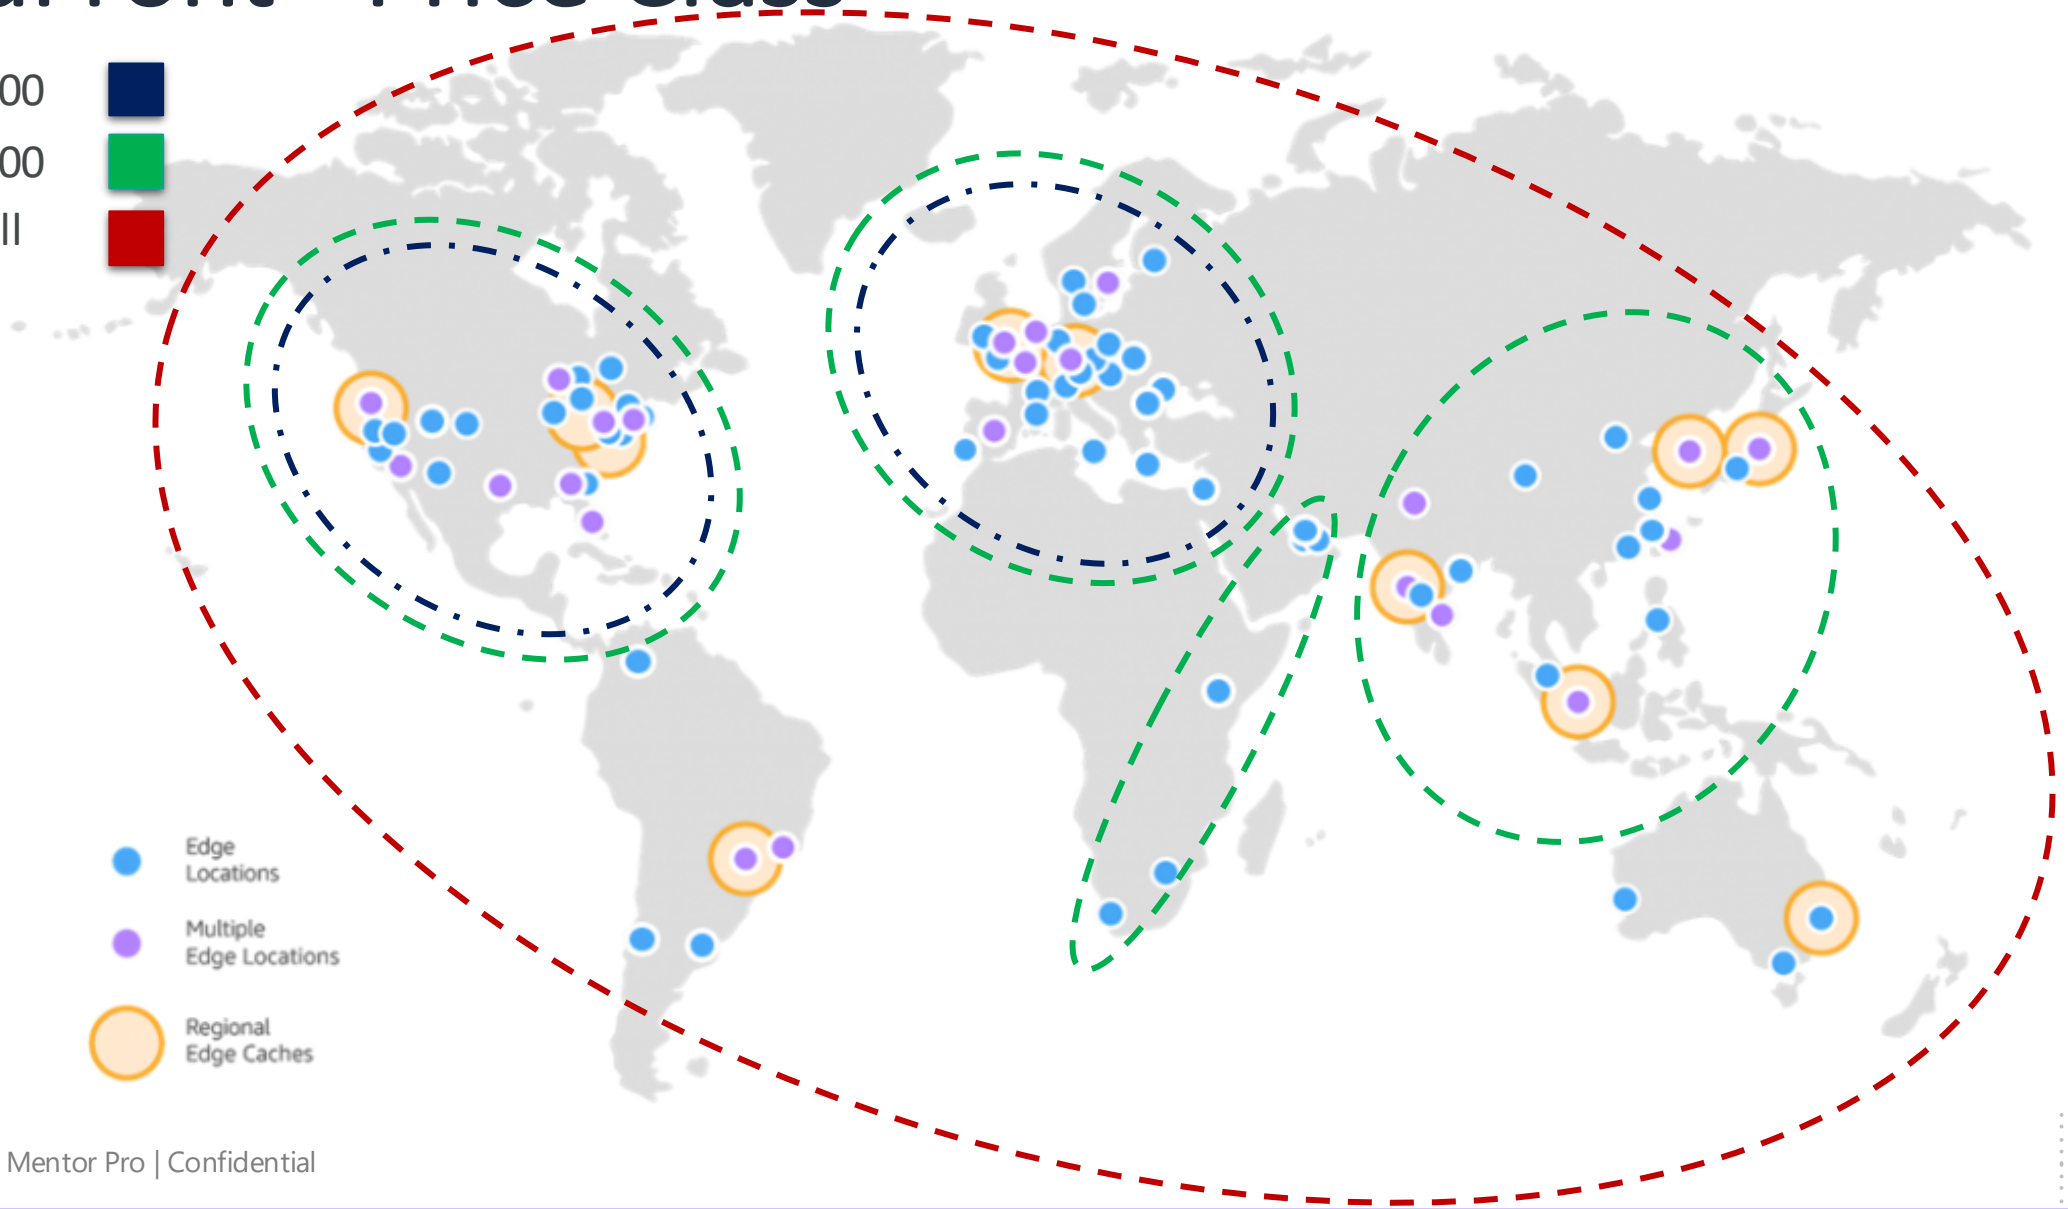

© Copyright by Cloud Mentor Pro | Confidential

AWS Elastic Beanstalk

Deploying applications in AWS safely and predictably

© Copyright by Cloud Mentor Pro | Confidential

Typical architecture: Web App 3-tier

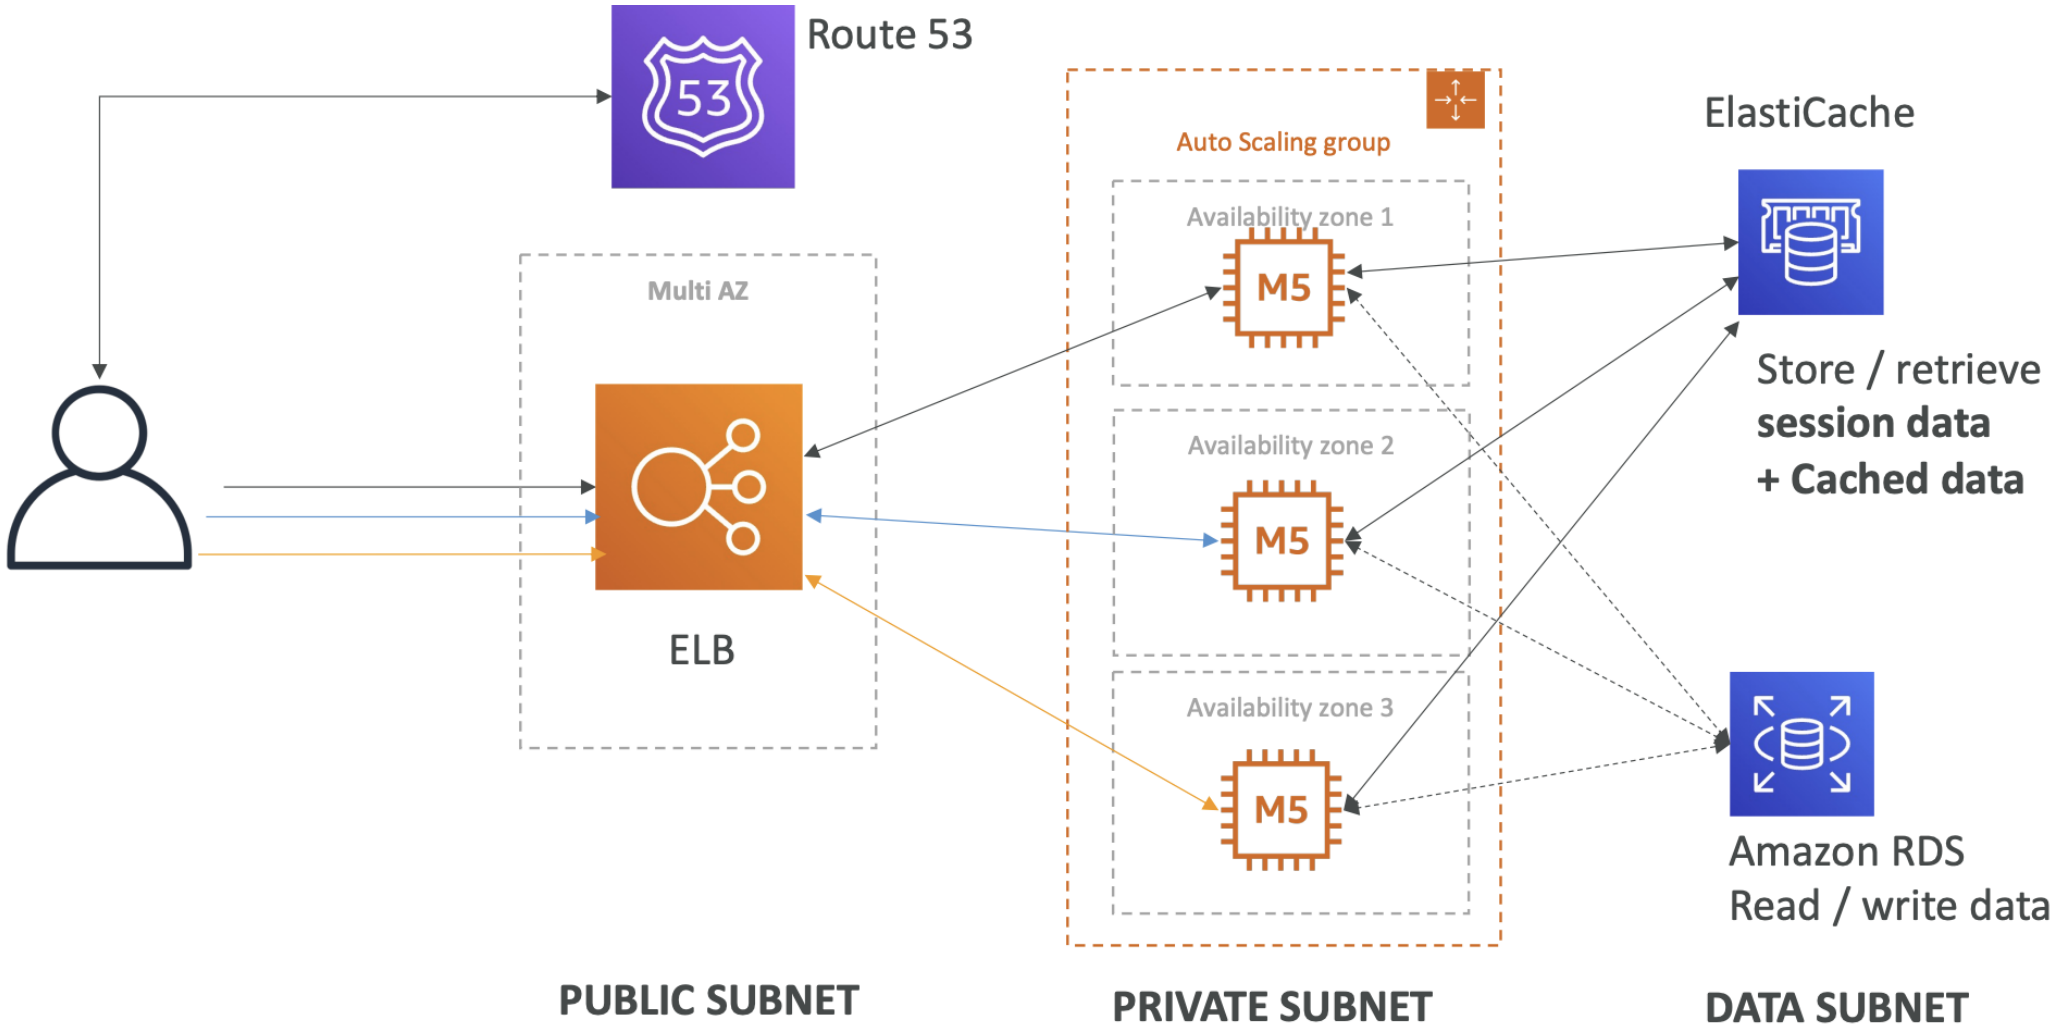

© Copyright by Cloud Mentor Pro | Confidential

Developer problems on AWS

- Managing infrastructure
- Deploying Code
- Configuring all the databases, load balancers, etc
- Scaling concerns

- Most web apps have the same architecture (ALB + ASG)
- All the developers want is for their code to run!
- Possibly, consistently across different applications and environments

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk – Overview

- Elastic Beanstalk is a developer centric view of deploying an application on AWS
- It uses all the component's we've seen before: EC2, ASG, ELB, RDS, …
- Managed service
  - Automatically handles capacity provisioning, load balancing, scaling, application health monitoring, instance configuration, …
  - Just the application code is the responsibility of the developer
- We still have full control over the configuration
- Beanstalk is free but you pay for the underlying instances

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk – Components

- **Application**: collection of Elastic Beanstalk components (environments, versions, configurations, …)
- **Application Version**: an iteration of your application code
- **Environment**
  - Collection of AWS resources running an application version (only one application version at a time)
  - **Tiers**: Web Server Environment Tier & Worker Environment Tier
  - You can create multiple environments (dev, test, prod, …)

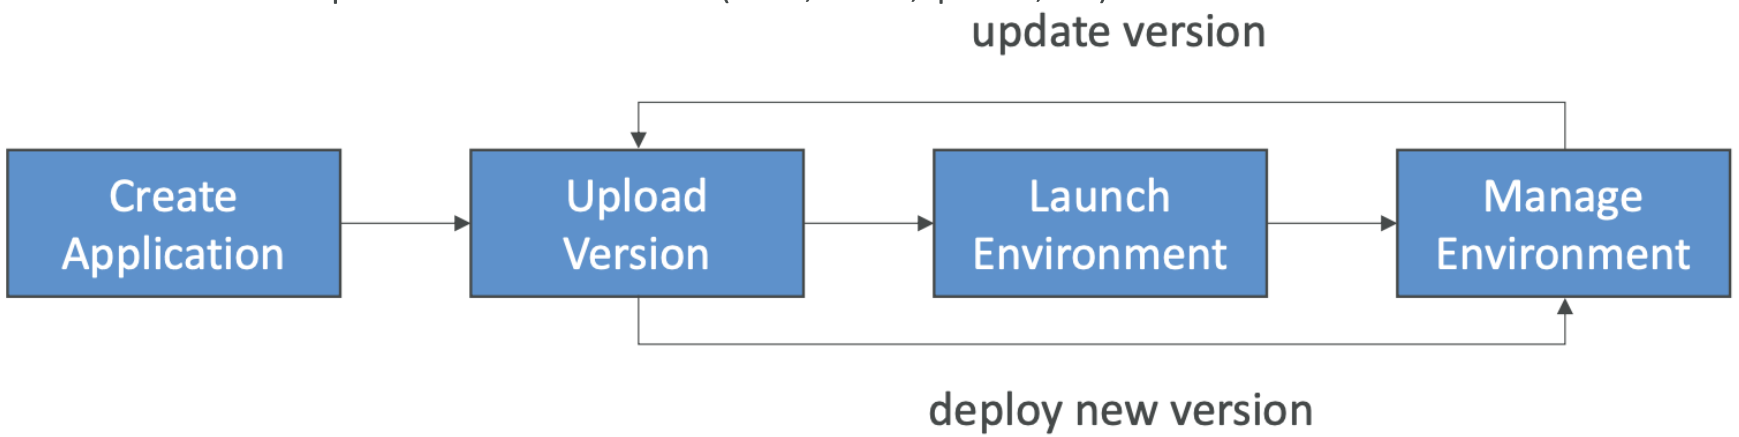

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk – Supported Platforms

- Go
- Java SE
- Java with Tomcat
- .NET Core on Linux
- .NET on Windows Server
- Node.js
- PHP
- Python
- Ruby
- Packer Builder
- Single Container Docker
- Multi-container Docker
- Preconfigured Docker

© Copyright by Cloud Mentor Pro | Confidential

Web Server Tier vs. Worker Tier

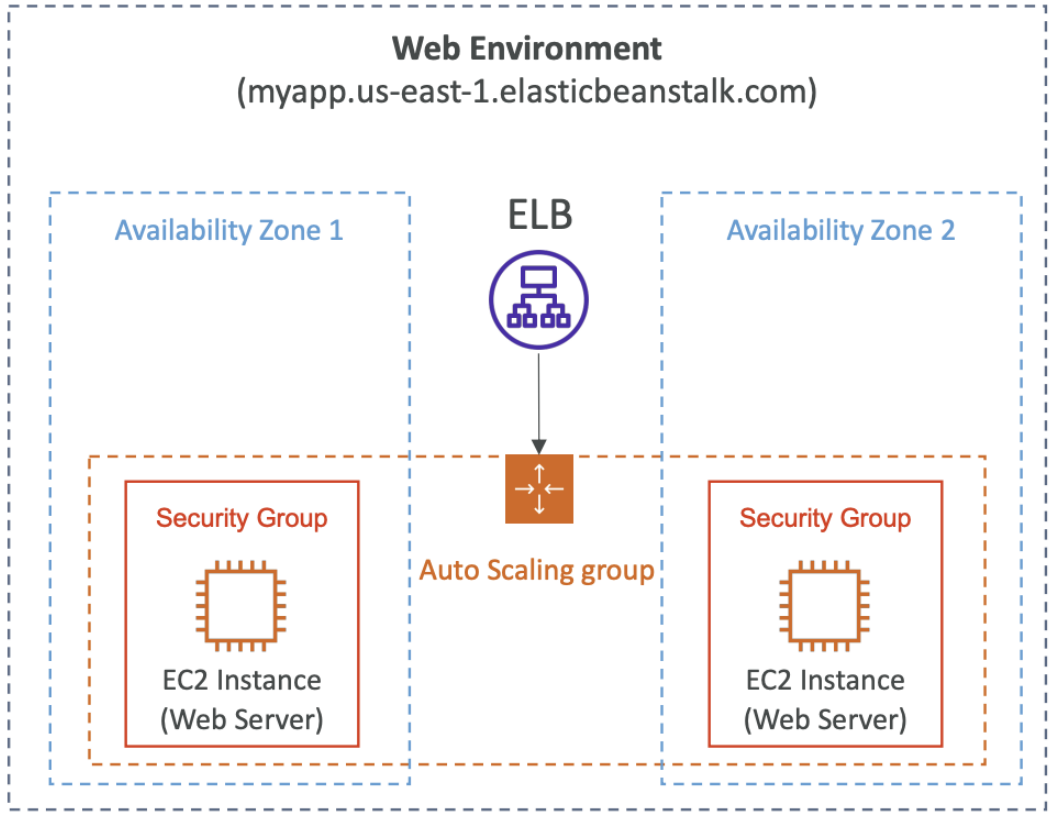

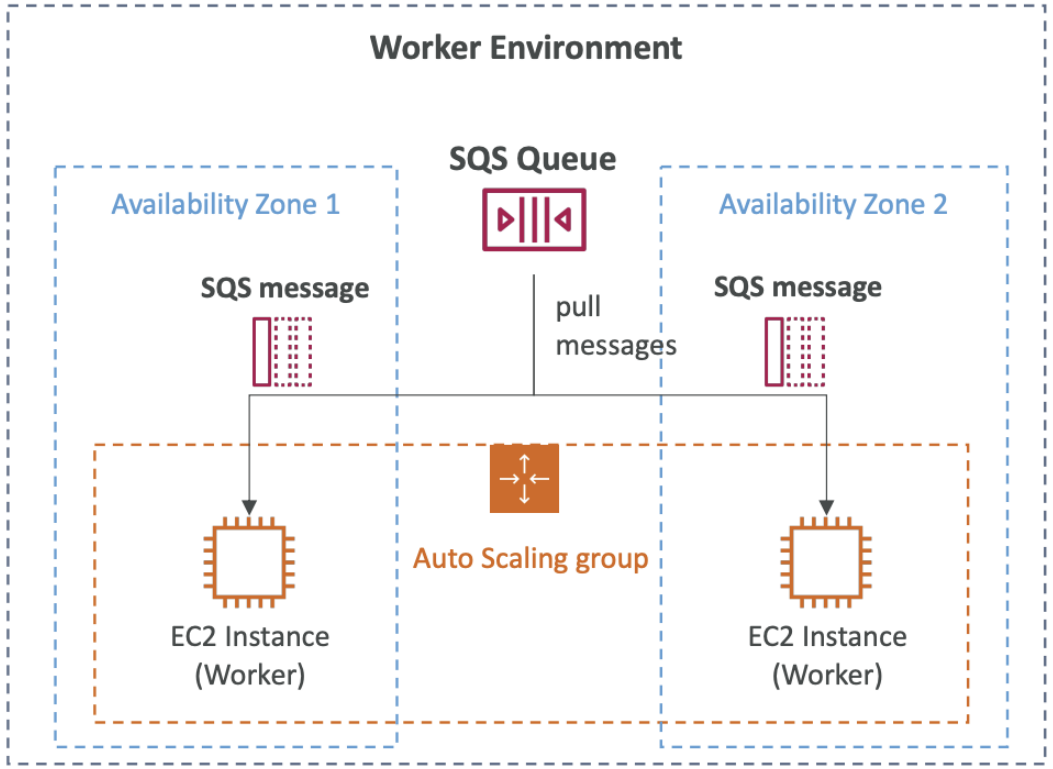

- Scale based on the number of SQS messages
- Can push messages to SQS queue from another Web Server Tier

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Modes

Single Instance
Great for dev

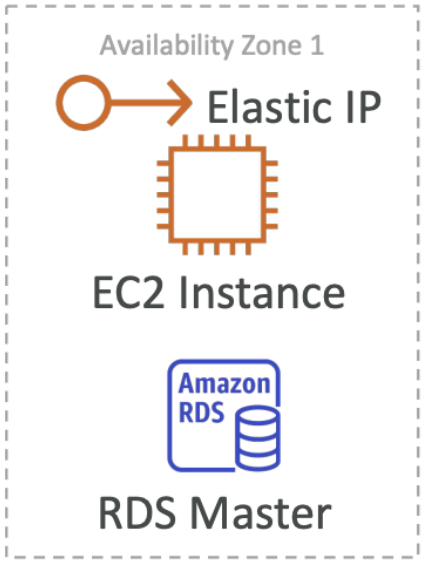

High Availability with Load Balancer
Great for prod

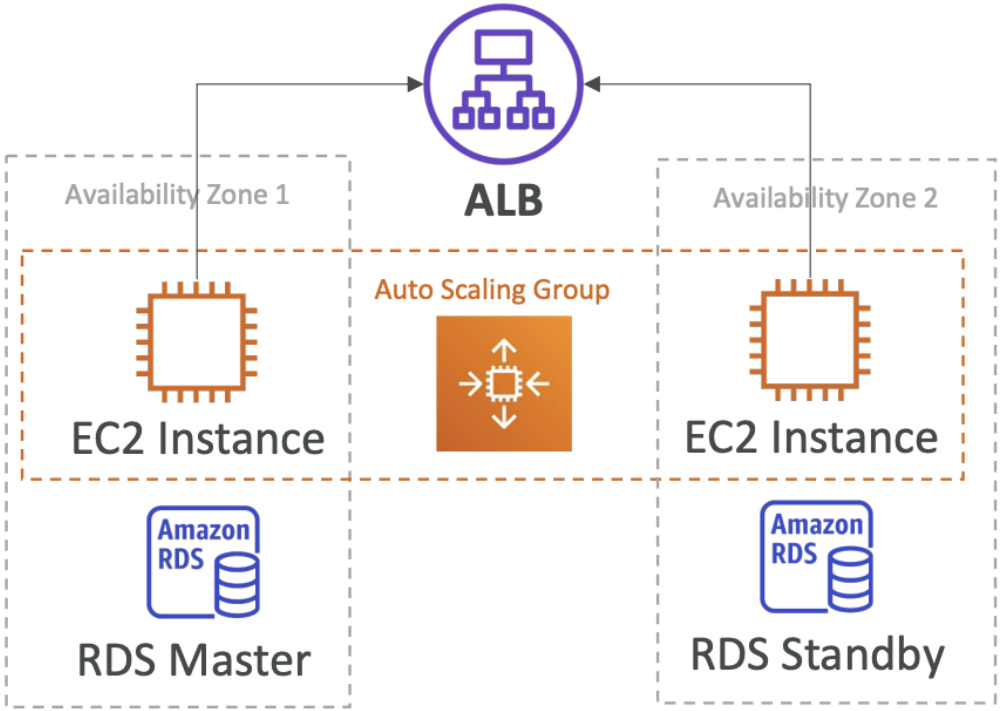

© Copyright by Cloud Mentor Pro | Confidential

Beanstalk Deployment Options for Updates

- All at once (deploy all in one go) – fastest, but instances aren't available to serve traffic for a bit (downtime)
- Rolling: update a few instances at a time (bucket), and then move onto the next bucket once the first bucket is healthy
- Rolling with additional batches: like rolling, but spins up new instances to move the batch (so that the old application is still available)
- Immutable: spins up new instances in a new ASG, deploys version to these instances, and then swaps all the instances when everything is healthy
- Blue Green: create a new environment and switch over when ready
- Traffic Splitting: canary testing – send a small % of traffic to new deployment

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment
All at once

- Fastest deployment
- Application has downtime
- Great for quick iterations in development environment
- No additional cost

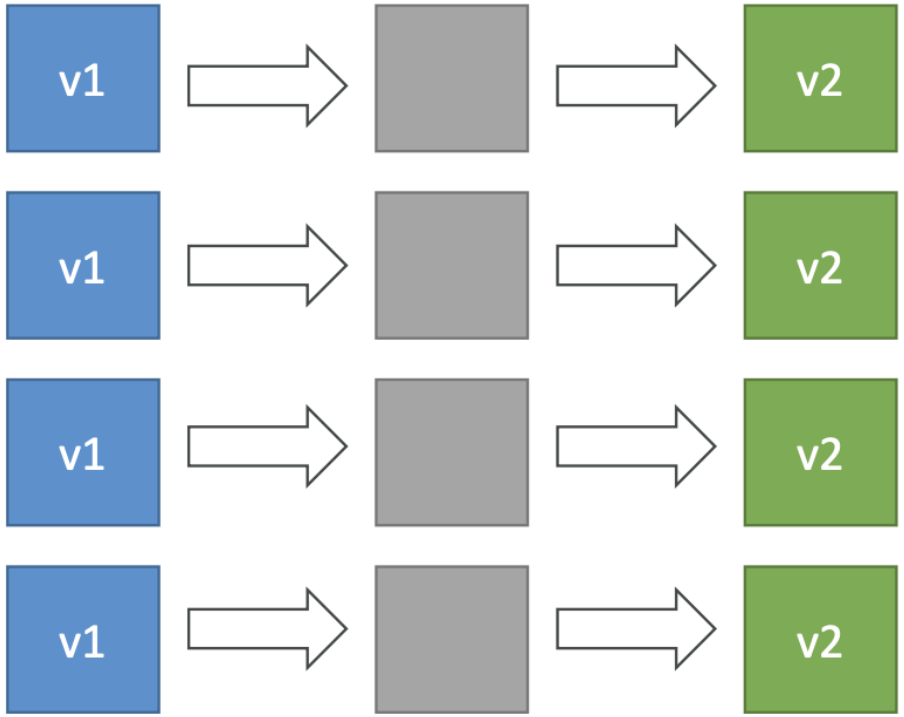

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Rolling

- Application is running below capacity
- Can set the bucket size
- Application is running both versions simultaneously
- No additional cost
- Long deployment

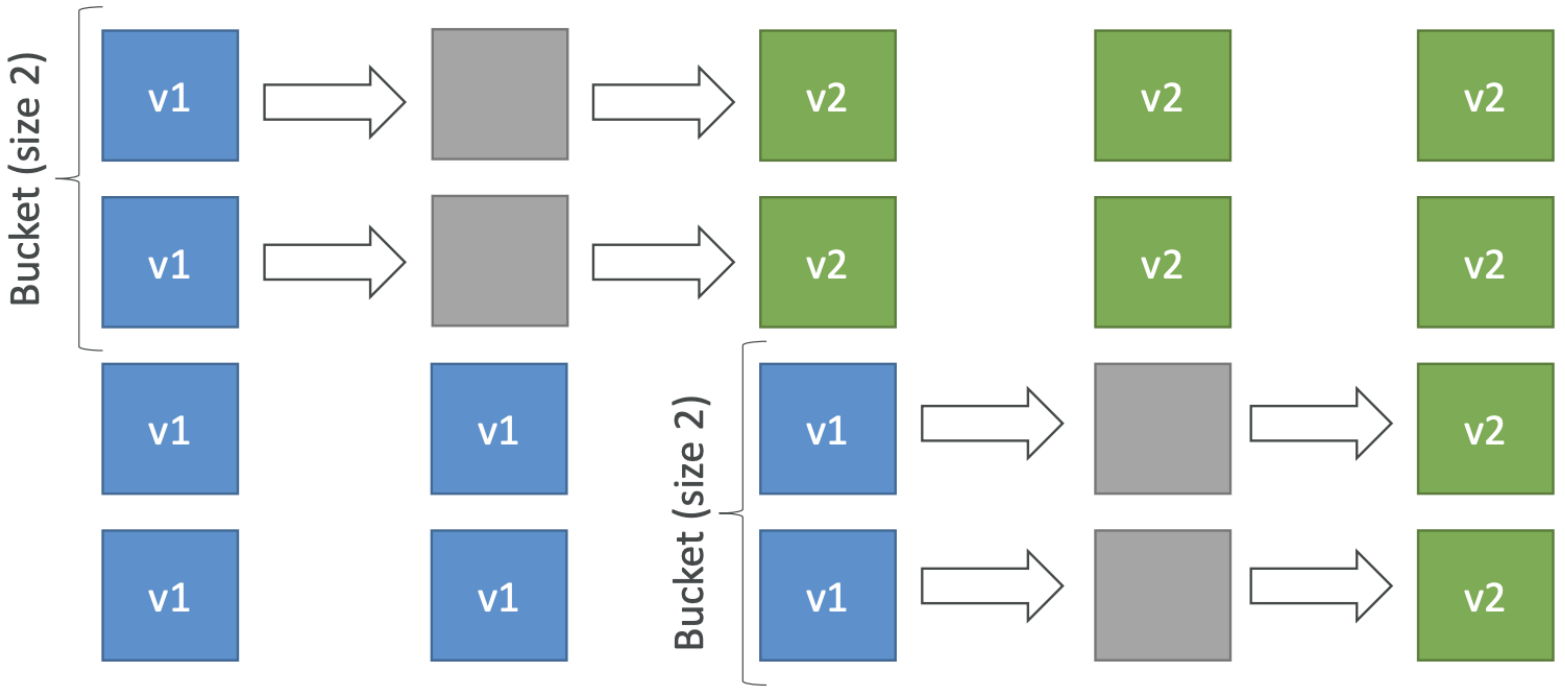

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Rolling with additional batches

- Application is running at capacity
- Can set the bucket size
- Application is running both versions simultaneously
- Small additional cost
- Additional batch is removed at the end of the deployment
- Longer deployment
- Good for prod

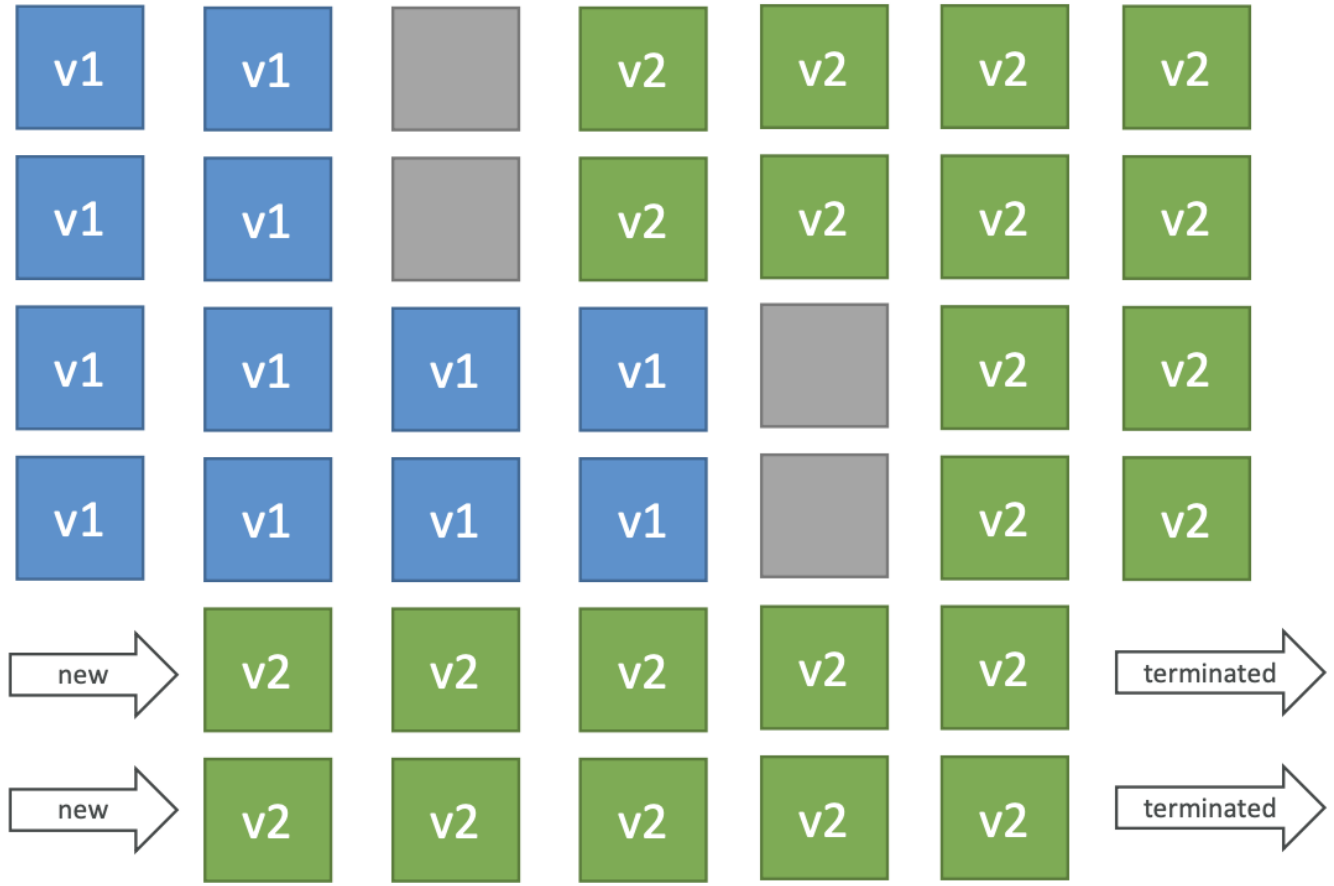

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Immutable

- Zero downtime
- New Code is deployed to new instances on a temporary ASG
- High cost, double capacity
- Longest deployment
- Quick rollback in case of failures (just terminate new ASG)
- Great for prod

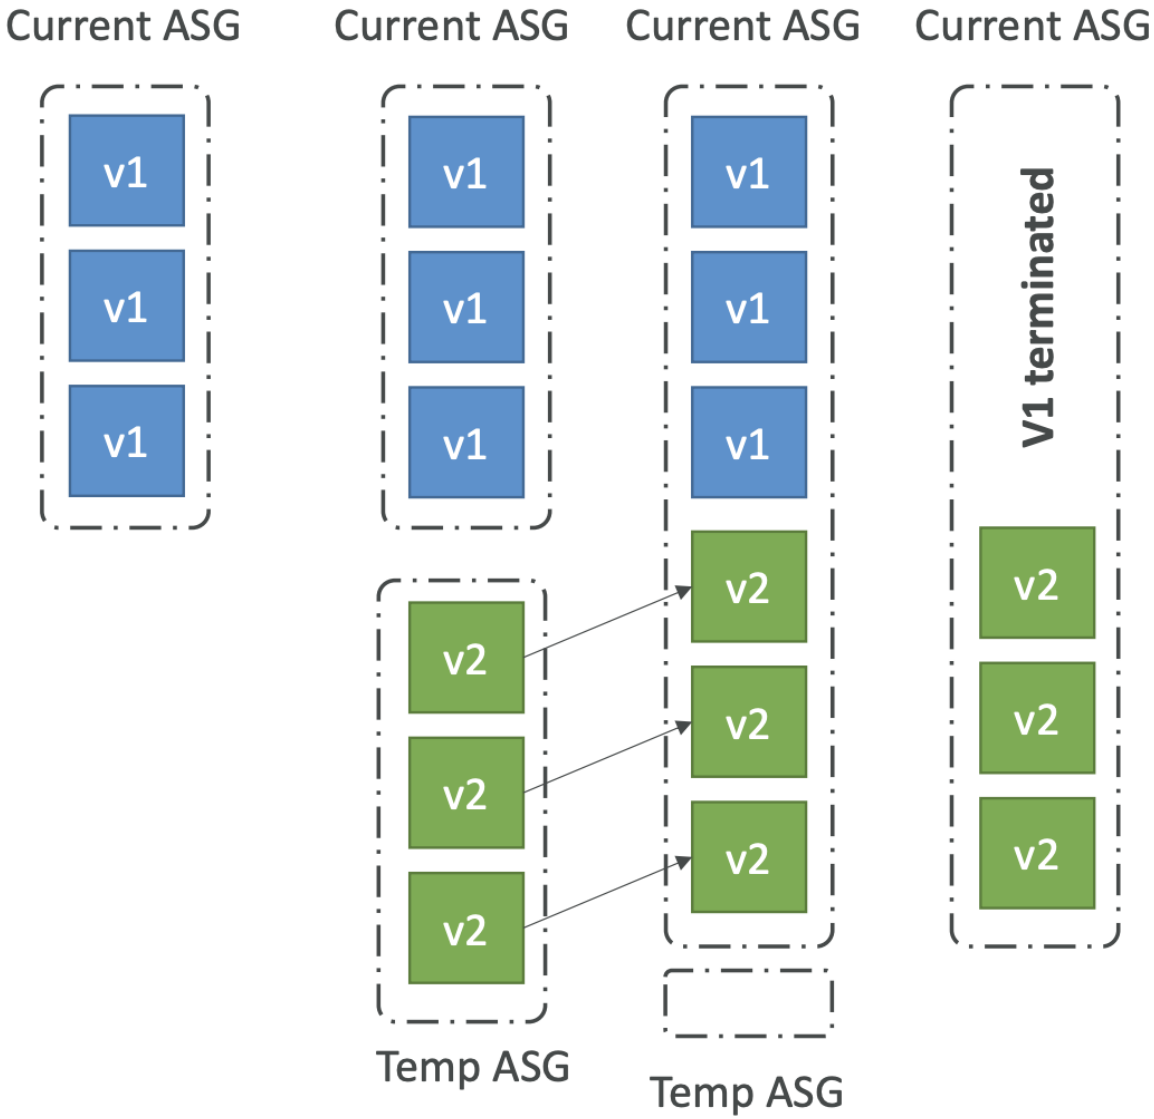

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Blue / Green

- Not a “direct feature” of Elastic Beanstalk
- Zero downtime and release facility
- Create a new “stage” environment and deploy v2 there
- The new environment (green) can be validated independently and roll back if issues
- Route 53 can be setup using weighted policies to redirect a little bit of traffic to the stage environment
- Using Beanstalk, “swap URLs” when done with the environment test

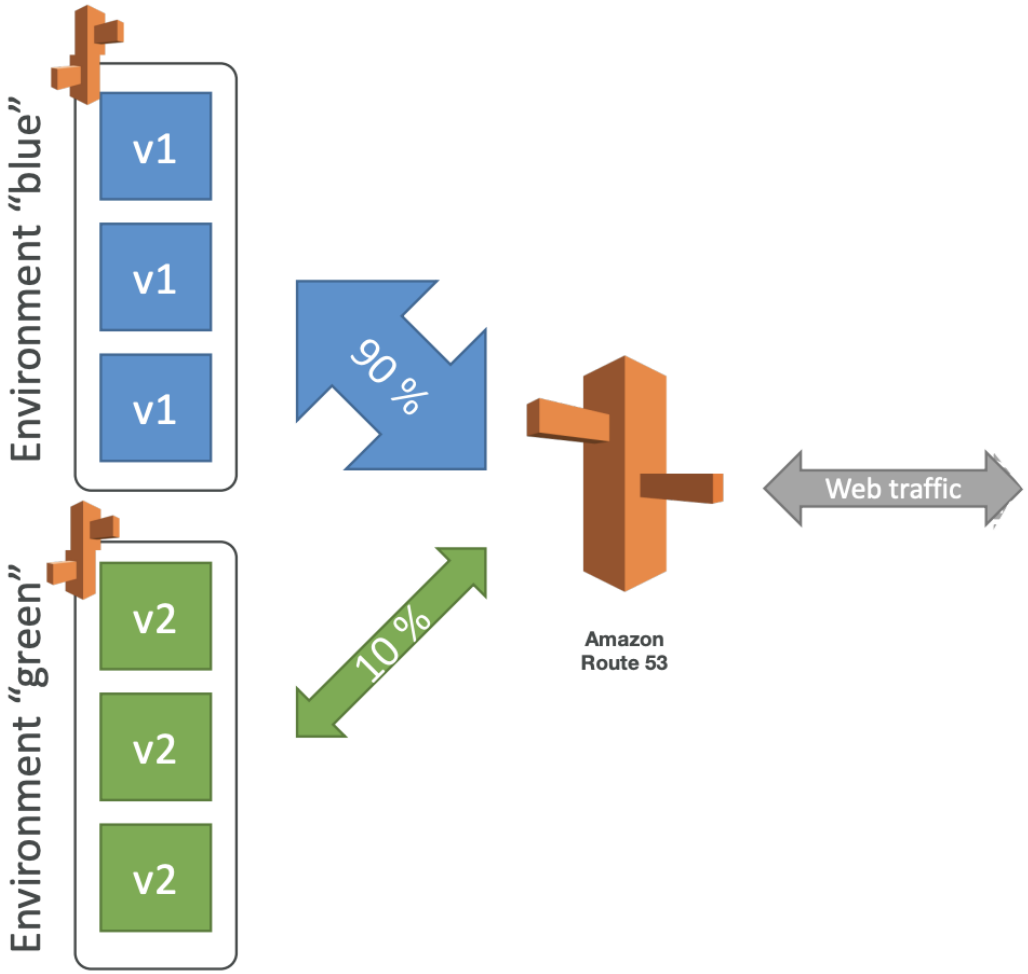

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk -Traffic Splitting

- Canary Testing
- New application version is deployed to a temporary ASG with the same capacity
- A small % of traffic is sent to the temporary ASG for a configurable amount of time
- Deployment health is monitored
- If there's a deployment failure, this triggers an automated rollback (very quick)
- Zero downtime
- New instances are migrated from the temporary to the original ASG
- Old application version is then terminated

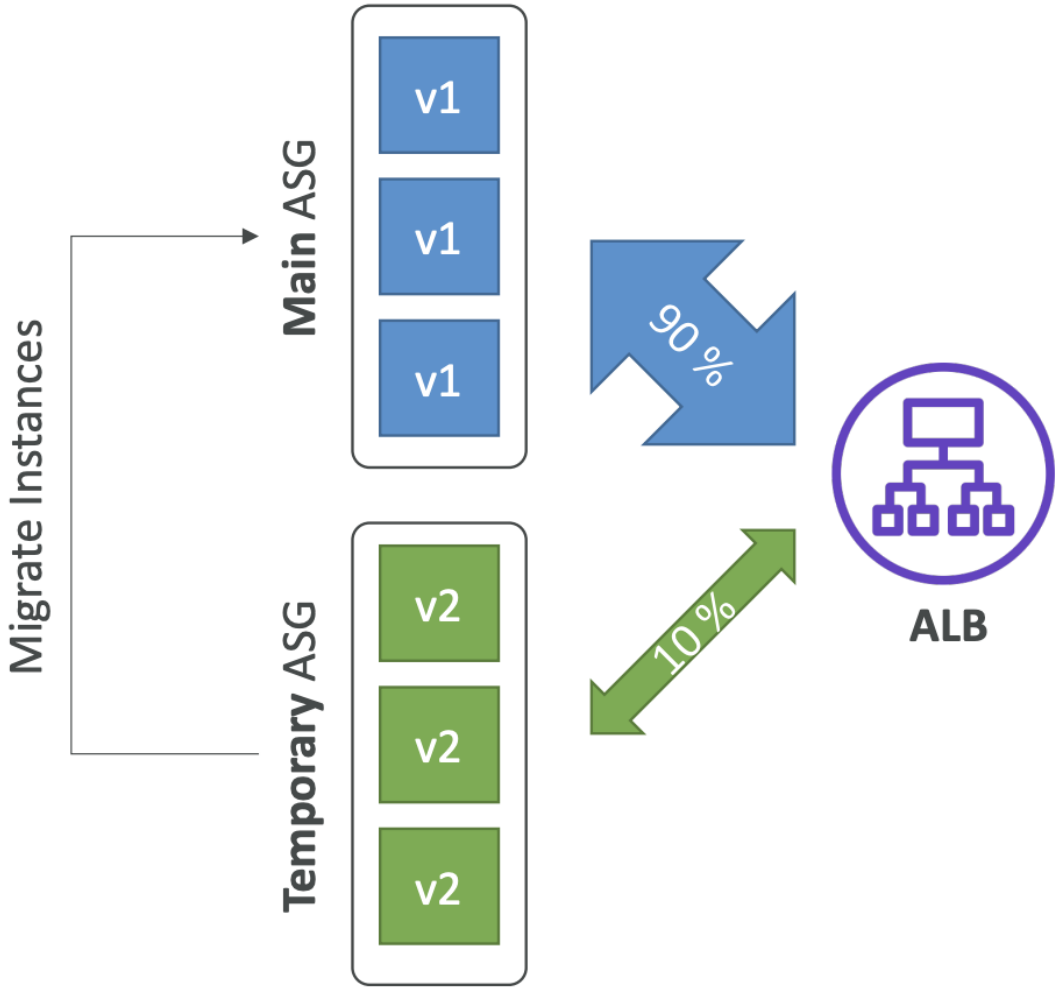

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Summary from AWS Doc

- https://docs.aws.amazon.com/elasticbeanstalk/latest/dg/using-features.deploy-existing-version.html

Deployment methods

Method	Impact of failed deployment	Deploy time	Zero downtime	No DNS change	Rollback process	Code deployed to
All at once	Downtime	◷	⊗ No	✓ Yes	Manual redeploy	Existing instances
Rolling	Single batch out of service; any successful batches before failure running new application version	◷ ◷ †	✓ Yes	✓ Yes	Manual redeploy	Existing instances
Rolling with an additional batch	Minimal if first batch fails; otherwise, similar to **Rolling**	◷ ◷ ◷ †	✓ Yes	✓ Yes	Manual redeploy	New and existing instances
Immutable	Minimal	◷ ◷ ◷ ◷	✓ Yes	✓ Yes	Terminate new instances	New instances
Traffic splitting	Percentage of client traffic routed to new version temporarily impacted	◷ ◷ ◷ ◷ ††	✓ Yes	✓ Yes	Reroute traffic and terminate new instances	New instances
Blue/green	Minimal	◷ ◷ ◷ ◷	✓ Yes	⊗ No	Swap URL	New instances

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk CLI

- We can install an additional CLI called the "EB cli" which makes working with Beanstalk from the CLI easier
- Basic commands are:
  - eb create
  - eb status
  - eb health
  - eb events
  - eb logs
  - eb open
  - eb deploy
  - eb config
  - eb terminate
- It's helpful for your automated deployment pipelines!

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Deployment Process

- Describe dependencies

  (requirements.txt for Python, package.json for Node.js)
- Package code as zip, and describe dependencies
  - Python: requirements.txt
  - Node.js: package.json
- **Console**: upload zip file (creates new app version), and then deploy
- **CLI**: create new app version using CLI (uploads zip), and then deploy

- Elastic Beanstalk will deploy the zip file on each EC2 instance, resolve dependencies and start the application

© Copyright by Cloud Mentor Pro | Confidential

Beanstalk Lifecycle Policy

- Elastic Beanstalk can store at most 1000 application versions
- If you don’t remove old versions, you won’t be able to deploy anymore
- To phase out old application versions, use a lifecycle policy
  - Based on time (old versions are removed)
  - Based on space (when you have too many versions)
- Versions that are currently used won’t be deleted
- Option not to delete the source bundle in S3 to prevent data loss

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Extensions

- A zip file containing our code must be deployed to Elastic Beanstalk
- All the parameters set in the UI can be configured with code using files
- Requirements:
  - in the .ebextensions/ directory in the root of source code
  - YAML / JSON format
  - **.config** extensions (example: logging.config)
  - Able to modify some default settings using: option_settings
  - Ability to add resources such as RDS, ElastiCache, DynamoDB, etc…
- Resources managed by .ebextensions get deleted if the environment goes away

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Under the Hood

- Under the hood, Elastic Beanstalk relies on CloudFormation
- CloudFormation is used to provision other AWS services (we'll see later)

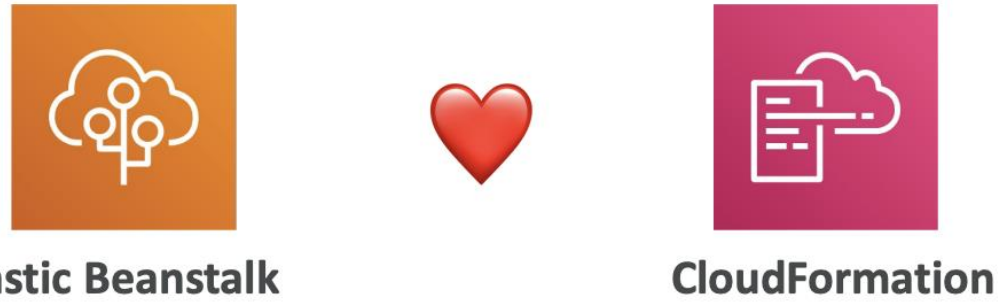

- Use case: you can define CloudFormation resources in your **.ebextensions** to provision ElastiCache, an S3 bucket, anything you want!

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Cloning

- Clone an environment with the exact same configuration
- Useful for deploying a “test” version of your application

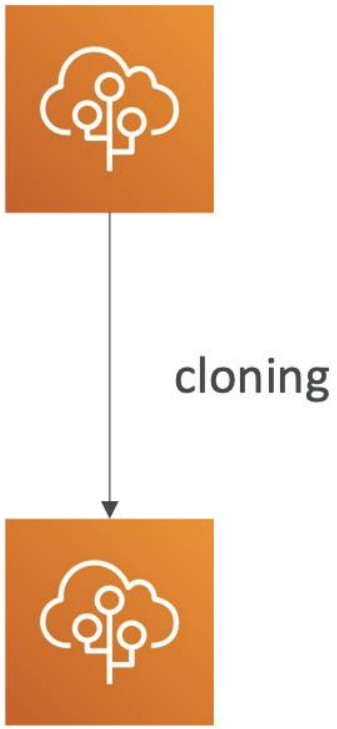

- All resources and configuration are preserved:
  - Load Balancer type and configuration
  - RDS database type (but the data is not preserved)
  - Environment variables

- After cloning an environment, you can change settings

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Migration: Load Balancer

- After creating an Elastic Beanstalk environment, you cannot change the Elastic Load Balancer type (only the configuration)
- To migrate:
  1. create a new environment with the same configuration except LB (can't clone)
  2. deploy your application onto the new environment
  3. perform a CNAME swap or Route 53 update

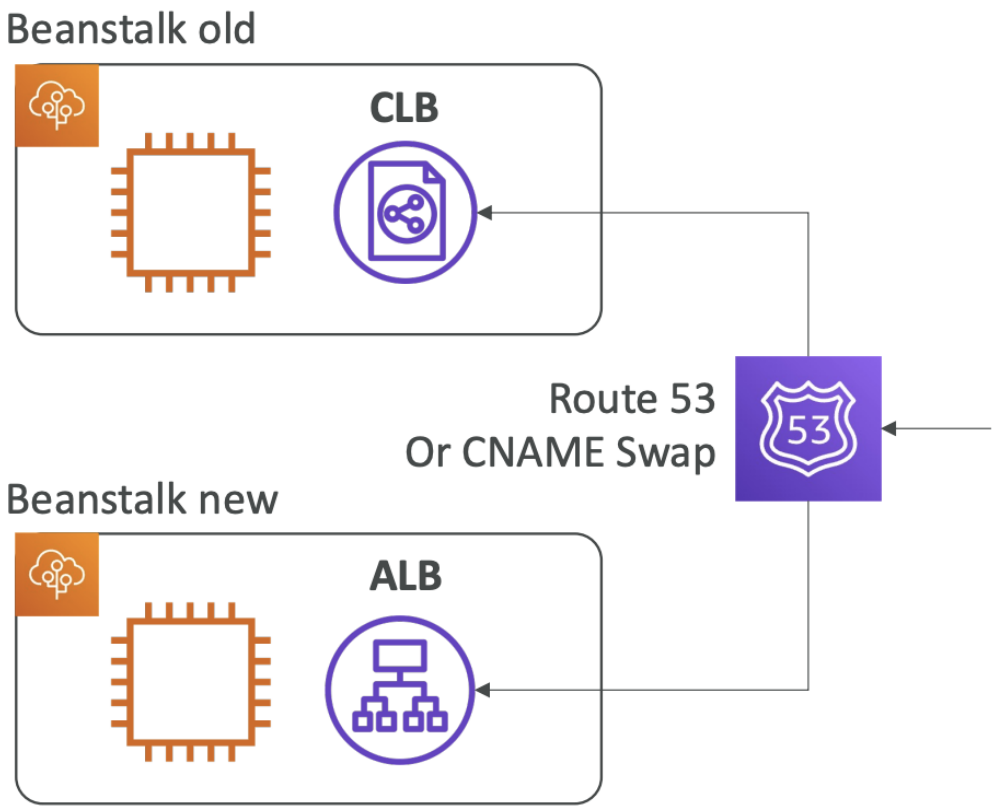

© Copyright by Cloud Mentor Pro | Confidential

RDS with Elastic Beanstalk

- RDS can be provisioned with Beanstalk, which is great for dev / test
- This is not great for prod as the database lifecycle is tied to the Beanstalk environment lifecycle
- The best for prod is to separately create an RDS database and provide our EB application with the connection string

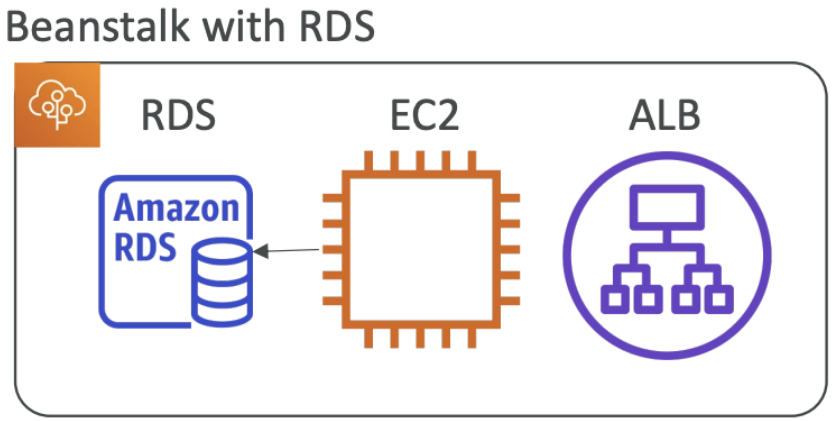

© Copyright by Cloud Mentor Pro | Confidential

Elastic Beanstalk Migration: Decouple RDS

1. Create a snapshot of RDS DB (as a safeguard)
2. Go to the RDS console and protect the RDS database from deletion
3. Create a new Elastic Beanstalk environment, without RDS, point your application to existing RDS
4. perform a CNAME swap (blue/green) or Route 53 update, confirm working
5. Terminate the old environment (RDS won't be deleted)
6. Delete CloudFormation stack (in DELETE_FAILED state)

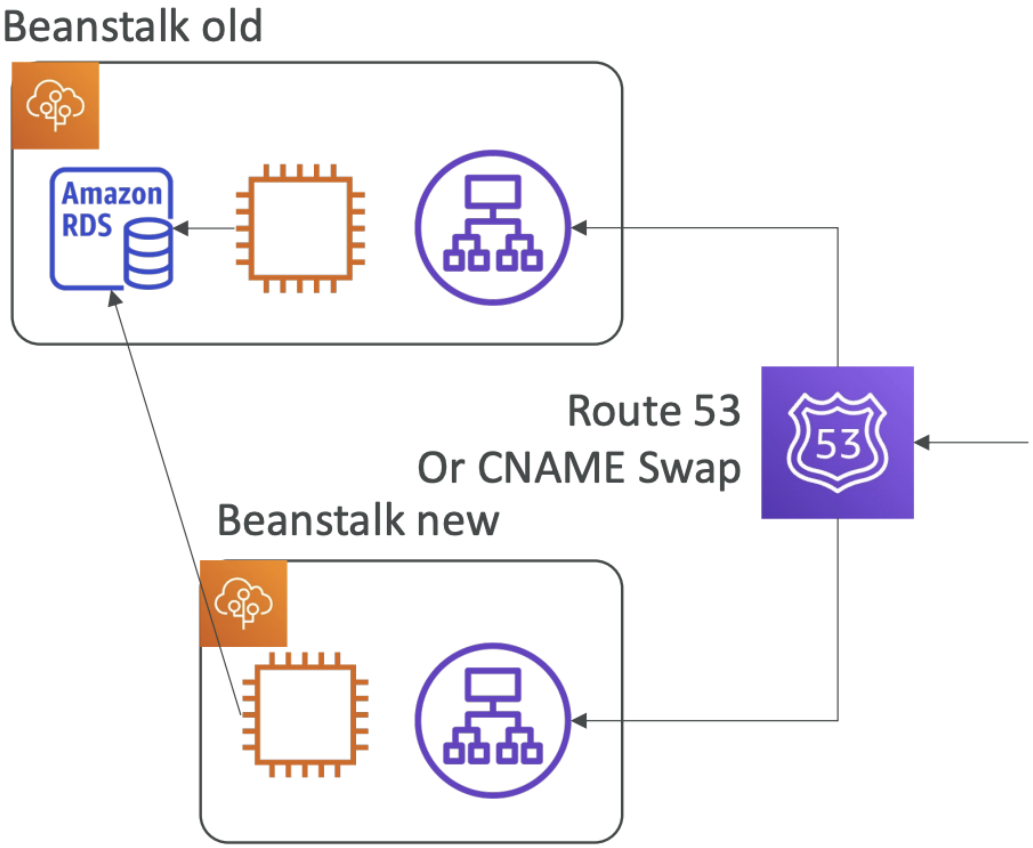

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- *Blue/Green Deployments with Elastic Beanstalk*

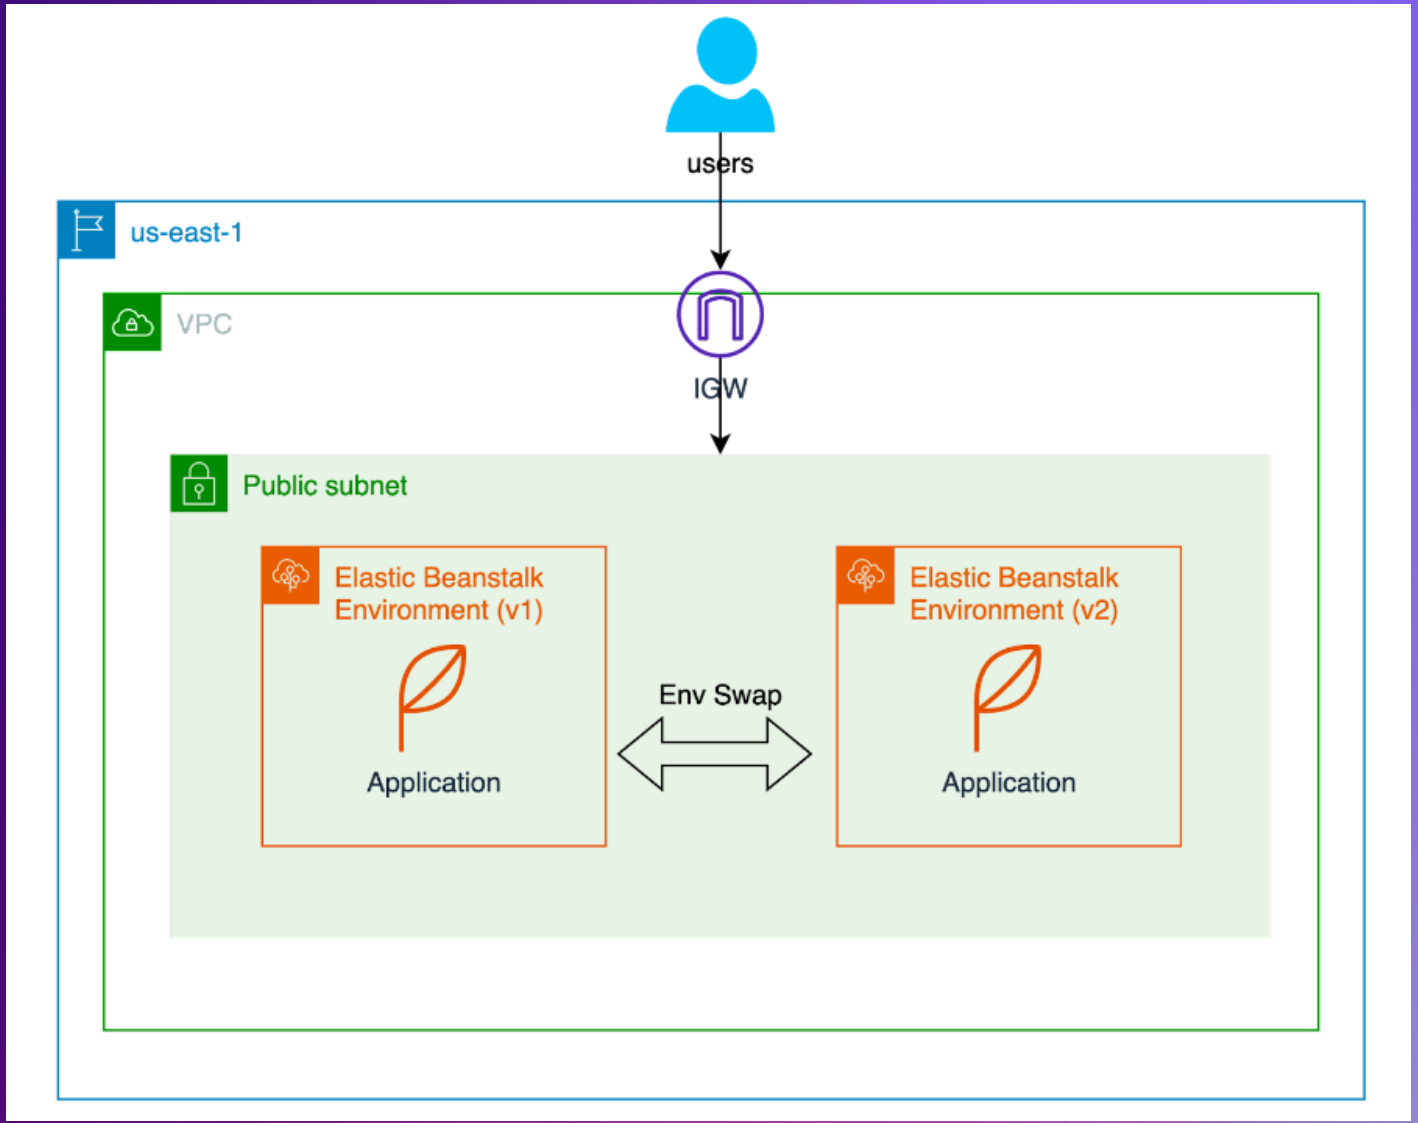

© Copyright by Cloud Mentor Pro | Confidential

Section 8

- AWS CloudFormation

© Copyright by Cloud Mentor Pro | Confidential

AWS CloudFormation

Managing your infrastructure as code

© Copyright by Cloud Mentor Pro | Confidential

AWS CloudFormation

- CloudFormation is a declarative way of outlining your AWS Infrastructure, for any resources (most of them are supported)
- For example, within a CloudFormation template, you say:
  - I want a security group
  - I want two EC2 instances using this security group
  - I want two Elastic IPs for these EC2 instances
  - I want an S3 bucket
  - I want a load balancer (ELB) in front of these EC2 instances
- Then CloudFormation creates those for you, in the **right order**, with the **exact configuration** that you specify

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation –Template Example

```yaml
AWSTemplateFormatVersion: '2010-09-09'
Description: 'AWS CloudFormation Sample Template LAMP_Multi_AZ: C
Parameters: …
Mappings: …
Resources:
  ApplicationLoadBalancer:
    Type: AWS::ElasticLoadBalancingV2::LoadBalancer
    Properties:
      Subnets: !Ref Subnets
  ALBListener:
    Type: AWS::ElasticLoadBalancingV2::Listener
    Properties:
      DefaultActions:
        - Type: forward
          TargetGroupArn: !Ref ALBTargetGroup
      LoadBalancerArn: !Ref ApplicationLoadBalancer
      Port: '80'
      Protocol: HTTP
  ALBTargetGroup:
    Type: AWS::ElasticLoadBalancingV2::TargetGroup
    Properties:
      HealthCheckIntervalSeconds: 10
      HealthCheckTimeoutSeconds: 5
      HealthyThresholdCount: 2
      Port: 80
      Protocol: HTTP
      UnhealthyThresholdCount: 5
      VpcId: !Ref VpcId
      TargetGroupAttributes:
        - Key: stickiness.enabled
          Value: 'true'
```

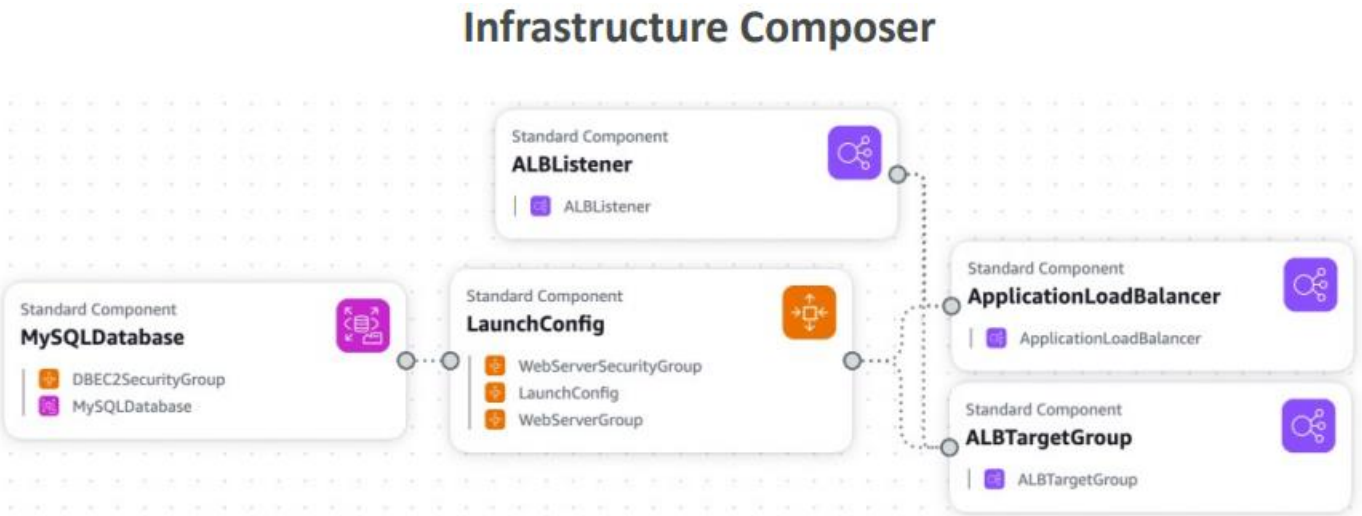

© Copyright by Cloud Mentor Pro | Confidential

Benefits of AWS CloudFormation (1/2)

- Infrastructure as code
  - No resources are manually created, which is excellent for control
  - The code can be version controlled for example using Git
  - Changes to the infrastructure are reviewed through code

- Cost
  - Each resources within the stack is tagged with an identifier so you can easily see how much a stack costs you
  - You can estimate the costs of your resources using the CloudFormation template
  - Savings strategy: In Dev, you could automation deletion of templates at 5 PM and recreated at 8 AM, safely

© Copyright by Cloud Mentor Pro | Confidential

Benefits of AWS CloudFormation (2/2)

- Productivity
  - Ability to destroy and re-create an infrastructure on the cloud on the fly
  - Automated generation of Diagram for your templates!
  - Declarative programming (no need to figure out ordering and orchestration)

- Separation of concern: create many stacks for many apps, and many layers. Ex:
  - VPC stacks
  - Network stacks
  - App stacks

- Don’t re-invent the wheel
  - Leverage existing templates on the web!
  - Leverage the documentation

© Copyright by Cloud Mentor Pro | Confidential

How CloudFormation Works

- Templates must be uploaded in S3 and then referenced in CloudFormation
- To update a template, we can't edit previous ones. We have to re-upload a new version of the template to AWS
- Stacks are identified by a name
- Deleting a stack deletes every single artifact that was created by CloudFormation.

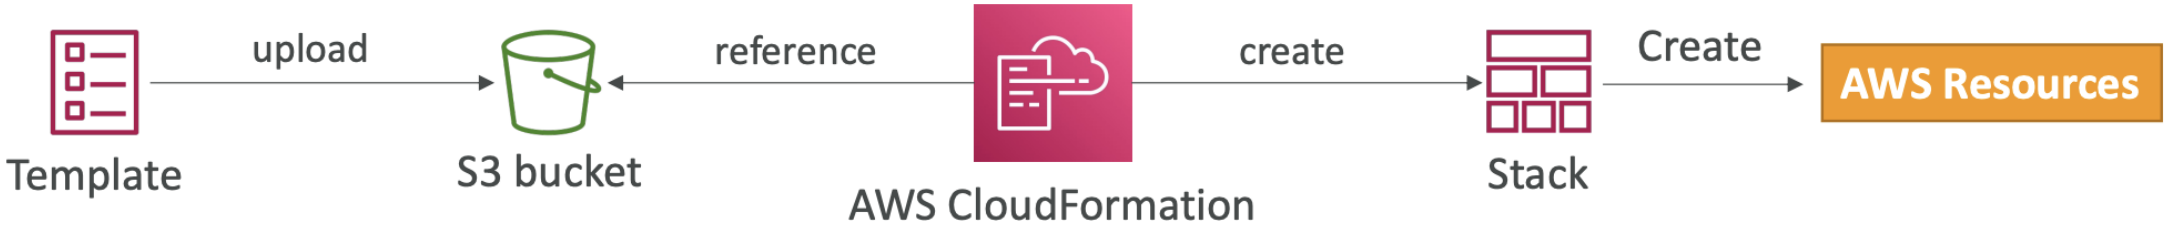

© Copyright by Cloud Mentor Pro | Confidential

Deploying CloudFormation Templates

- Manual way
  - Editing templates in Application Composer or code editor
  - Using the console to input parameters, etc…
  - We'll mostly do this way in the course for learning purposes
- Automated way
  - Editing templates in a YAML file
  - Using the AWS CLI (Command Line Interface) to deploy the templates, or using a Continuous Delivery (CD) tool
  - Recommended way when you fully want to automate your flow

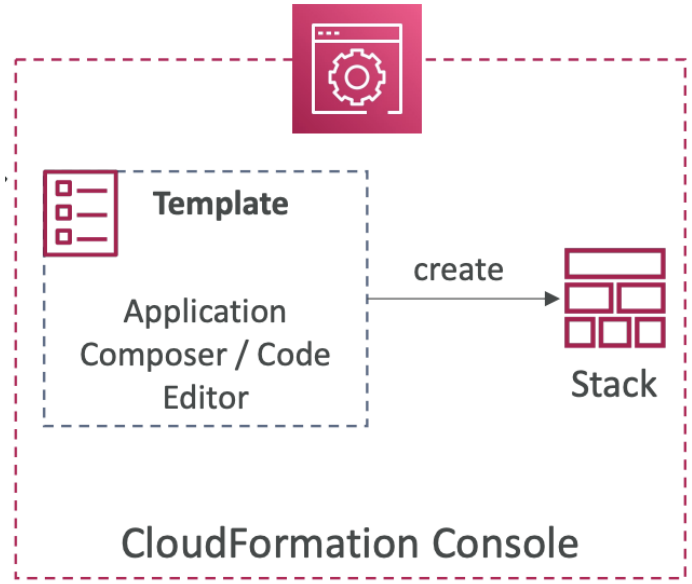

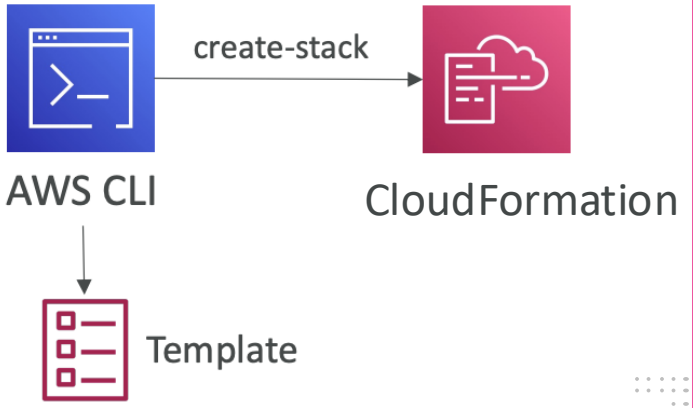

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Building Blocks

- Template's Components
  - **AWSTemplateFormatVersion** – identifies the capabilities of the template "2010-09-09"
  - **Description** – comments about the template
  - **Resources (MANDATORY) – your AWS resources declared in the template**
  - **Parameters** – the dynamic inputs for your template
  - **Mappings** – the static variables for your template
  - **Outputs** – references to what has been created
  - **Conditionals** – list of conditions to perform resource creation

- Template's Helpers
  - References
  - Functions

© Copyright by Cloud Mentor Pro | Confidential

Introductory Example

- We're going to create a simple EC2 instance
- And we're going to add security group to it
- For now, forget about the code syntax
- We'll look at the structure of the files later

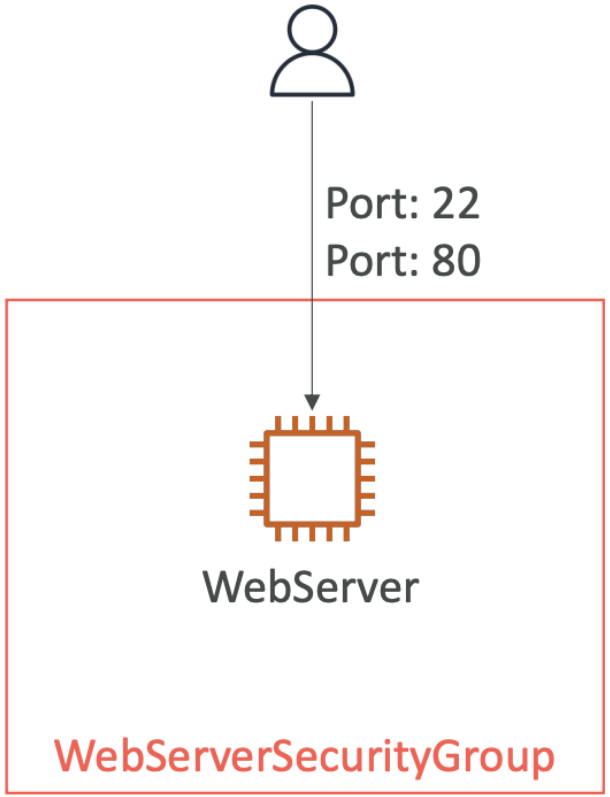

- We'll see how in no-time, we are able to get started with CloudFormation!

© Copyright by Cloud Mentor Pro | Confidential

https://github.com/phongaws/aws-dva-code-2023/blob/main/cloudformation/Exercise/1.simple.ec2.yaml

YAML Crash Course

```yaml
invoice: 34843
date: 2001-01-23
bill-to:
  given: Chris
  family: Dumars
  address:
    lines: |
      458 Walkman Dr.
      Suite #292
    city: Royal Oak
    state: MI
    postal: 48046
products:
  - sku: BL394D
    quantity: 4
    description: Basketball
    price: 450.00
  - sku: BL4438H
    quantity: 1
    description: Super Hoop
    price: 2392.00
```

- YAML and JSON are the languages you can use for CloudFormation
- JSON is horrible for CF
- YAML is great in so many ways
- Let's learn a bit about it!

- Key value Pairs
- Nested objects
- Support Arrays
- Multi line strings
- Can include comments

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Resources

- Resources are the core of your CloudFormation template **(MANDATORY)**
- They represent the different AWS Components that will be created and configured
- Resources are declared and can reference each other
- AWS figures out creation, updates and deletes of resources for us
- There are over 700 types of resources (!)
- Resource types identifiers are of the form:

service-provider::*service-name*::*data-type-name*

© Copyright by Cloud Mentor Pro | Confidential

How do I find Resources documentation?

- I can't teach you all the 700+ resources, but I can teach you how to learn how to use them
- All the resources can be found here: https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-template-resource-type-ref.html
- Then, we just read the docs ☺
- Example here (for an EC2 instance): https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-resource-ec2-instance.html

© Copyright by Cloud Mentor Pro | Confidential

Analysis of CloudFormation Template

- Going back to the example of the introductory lecture, let's learn why it was written this way.
- Relevant documentation can be found here:
  - https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-resource-ec2-instance.html
  - https://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-resource-ec2-securitygroup.html
  - http://docs.aws.amazon.com/AWSCloudFormation/latest/UserGuide/aws-resource-ec2-eip.html

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Resources FAQ

- Can I create a dynamic number of resources?
  - Yes, you can by using CloudFormation Macros and Transform
  - It is not in the scope of this course

- Is every AWS Service supported?
  - Almost. Only a select few niches are not there yet
  - You can work around that using CloudFormation Custom Resources

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Parameters

- Parameters are a way to provide inputs to your AWS CloudFormation template

- They're important to know about if:
  - You want to reuse your templates across the company
  - Some inputs can not be determined ahead of time

- Parameters are extremely powerful, controlled, and can prevent errors from happening in your templates, thanks to types

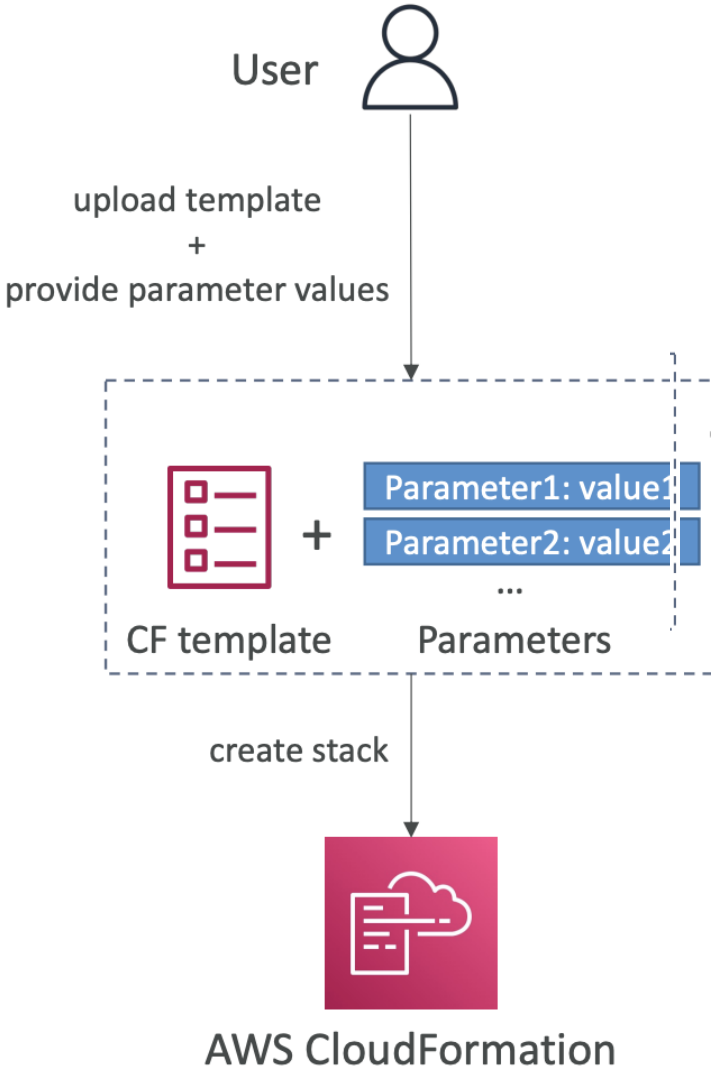

© Copyright by Cloud Mentor Pro | Confidential

When should you use a Parameter?

```
Parameters:
  SecurityGroupDescription:
    Description: Security Group Description
    Type: String
```

- Ask yourself this:
  - Is this CloudFormation resource configuration likely to change in the future?
  - If so, make it a parameter

- You won't have to re-upload a template to change its content ☺

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Parameters Settings

- Parameters can be controlled by all these settings:
  - Type:
    - **String**
    - **Number**
    - **CommaDelimitedList**
    - **List<Number>**
    - **AWS-Specific Parameter** (to help catch invalid values – match against existing values in the AWS account)
    - **List<AWS-Specific Parameter>**
    - **SSM Parameter** (get parameter value from SSM Parameter store)
  - **Description**
  - **ConstraintDescription** (String)
  - **Min/MaxLength**
  - **Min/MaxValue**
  - **Default**
  - **AllowedValues** (array)
  - **AllowedPattern** (regex)
  - **NoEcho** (Boolean)

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Parameters Example

AllowedValues

```yaml
Parameters:
  InstanceType:
    Description: Choose an EC2 instance type
    Type: String
    AllowedValues:
      - t2.micro
      - t2.small
      - t2.medium
    Default: t2.micro

Resources:
  MyEC2Instance:
    Type: AWS::EC2::Instance
    Properties:
      InstanceType: !Ref InstanceType
      ImageId: ami-0c02fb55956c7d316
```

NoEcho

```yaml
Parameters:
  DBPassword:
    Description: The database admin password
    Type: String
    NoEcho: true

Resources:
  MyDBInstance:
    Type: AWS::RDS::DBInstance
    Properties:
      DBInstanceClass: db.t2.micro
      AllocatedStorage: 20
      Engine: mysql
      MasterUsername: admin
      MasterUserPassword: !Ref DBPassword
      DBInstanceIdentifier: mydbinstance
```

© Copyright by Cloud Mentor Pro | Confidential

How to Reference a Parameter?

```
Resources:
  DBSubnet1:
    Type: AWS::EC2::Subnet
    Properties:
      VpcId: !Ref MyVPC
```

- The Fn::Ref function can be leveraged to reference parameters
- Parameters can be used anywhere in a template
- The shorthand for this in YAML is !Ref
- The function can also reference other elements within the template

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Pseudo Parameters

- AWS offers us Pseudo Parameters in any CloudFormation template
- These can be used at any time and are enabled by default
- Important pseudo parameters:

Reference Value	Example Returned Value
AWS::AccountId	123456789012
AWS::Region	us-east-1
AWS::StackId	arn:aws:cloudformation:us-east-1:123456789012:stack/MyStack/1c2fa620-982a-11e3-aff7-50e2416294e0
AWS::StackName	MyStack
AWS::NotificationARNs	[arn:aws:sns:us-east-1:123456789012:MyTopic]
AWS::NoValue	Doesn't return a value

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Mappings

- Mappings are fixed variables within your CloudFormation template
- They're very handy to differentiate between different environments (dev vs prod), regions (AWS regions), AMI types…
- All the values are hardcoded within the template

```
Mappings:
  Mapping01:
    Key01:
      Name: Value01
    Key02:
      Name: Value02
    Key03:
      Name: Value03
```

```
RegionMap:
  us-east-1:
    HVM64: ami-0ff8a91507f77f867
    HVMG2: ami-0a584ac55a7631c0c
  us-west-1:
    HVM64: ami-0bdb828fd58c52235
    HVMG2: ami-066ee5fd4a9ef77f1
  eu-west-1:
    HVM64: ami-047bb4163c506cd98
    HVMG2: ami-0a7c483d527806435
```

© Copyright by Cloud Mentor Pro | Confidential

Accessing Mapping Values (Fn::FindInMap)

- We use Fn::FindInMap to return a named value from a specific key
- !FindInMap [MapName, TopLevelKey, SecondLevelKey]

```
Mappings:
  RegionMap:
    us-east-1:
      HVM64: ami-0ff8a91507f77f867
      HVMG2: ami-0a584ac55a7631c0c
    us-west-1:
      HVM64: ami-0bdb828fd58c52235
      HVMG2: ami-066ee5fd4a9ef77f1

Resources:
  MyEC2Instance:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: !FindInMap [RegionMap, !Ref "AWS::Region", HVM64]
      InstanceType: t2.micro
```

Mappings work great for AMIs Because AMIs are region-specific!

© Copyright by Cloud Mentor Pro | Confidential

When would you use Mappings vs. Parameters?

- Mappings are great when you know in advance all the values that can be taken and that they can be deduced from variables such as
  - Region
  - Availability Zone
  - AWS Account
  - Environment (dev vs prod)
  - etc…
- They allow safer control over the template

- Use parameters when the values are really user specific

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Outputs

- The Outputs section declares ***optional*** outputs values that we can import into other stacks (if you export them first)!
- You can also view the outputs in the AWS Console or in using the AWS CLI
- They're very useful for example if you define a network CloudFormation, and output the variables such as VPC ID and your Subnet IDs
- It's the best way to perform some collaboration cross stack, as you let expert handle their own part of the stack

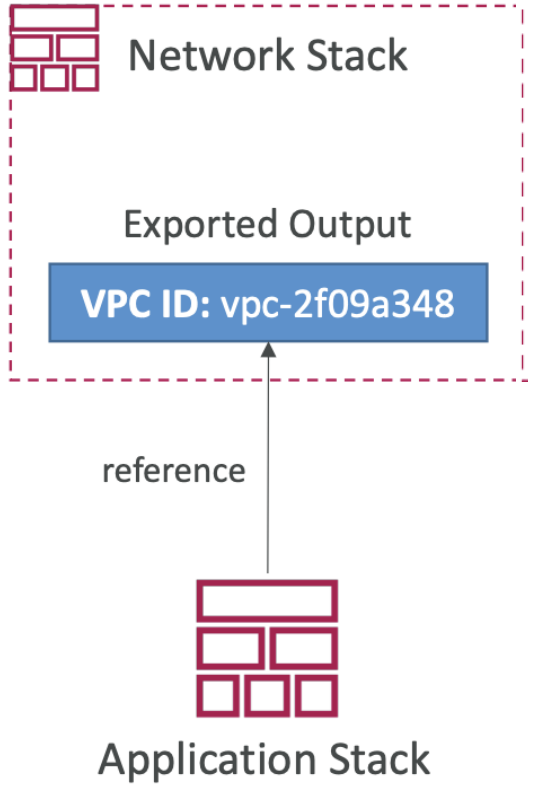

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Outputs

- Creating a SSH Security Group as part of one template
- We create an output that references that security group

```
Outputs:
  StackSSHSecurityGroup:
    Description: The SSH Security Group for our Company
    Value: !Ref MyCompanyWideSSHSecurityGroup
    Export:
      Name: SSHSecurityGroup
```

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Outputs Cross-Stack Reference

- We then create a second template that leverages that security group
- For this, we use the Fn::ImportValue function
- You can't delete the underlying stack until all the references are deleted

```
Resources:
  MySecureInstance:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: ami-0742b4e673072066f
      InstanceType: t2.micro
      AvailabilityZone: us-east-1a
      SecurityGroups:
        - !ImportValue SSHSecurityGroup
```

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Conditions

- Conditions are used to control the creation of resources or outputs based on a condition
- Conditions can be whatever you want them to be, but common ones are:
  - Environment (dev / test / prod)
  - AWS Region
  - Any parameter value
- Each condition can reference another condition, parameter value or mapping

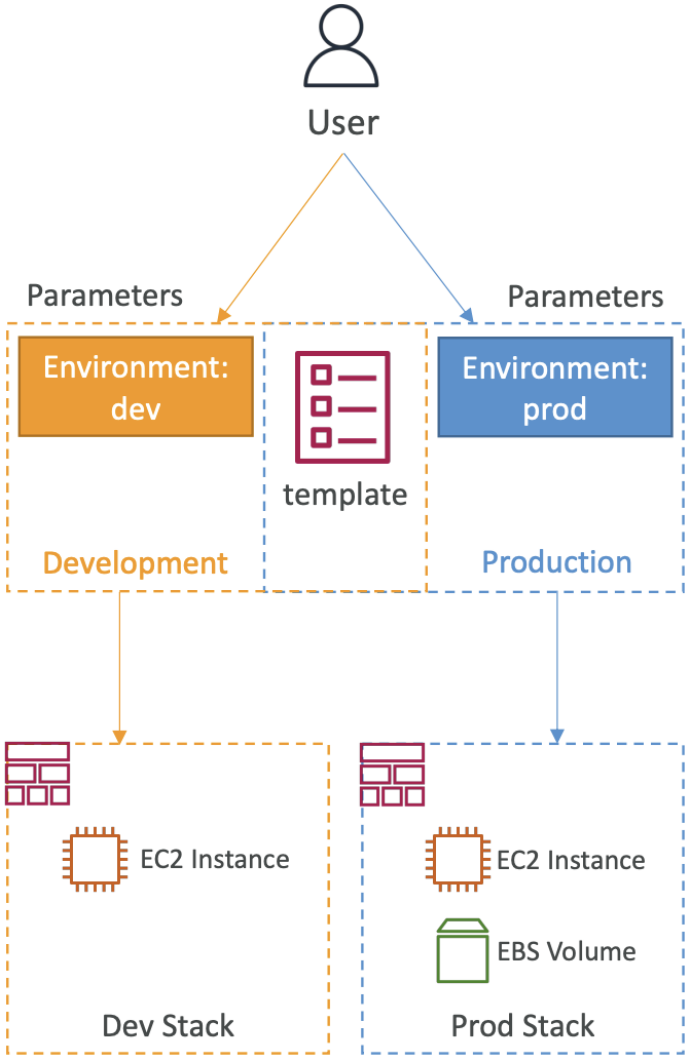

© Copyright by Cloud Mentor Pro | Confidential

How to define a Condition

```
Conditions:
  CreateProdResources: !Equals [ !Ref EnvType, prod ]
```

- The logical ID is for you to choose. It's how you name condition
- The intrinsic function (logical) can be any of the following:
  - Fn::And
  - Fn::Equals
  - Fn::If
  - Fn::Not
  - Fn::Or

© Copyright by Cloud Mentor Pro | Confidential

How to use a Condition

- Conditions can be applied to resources / outputs / etc...

```
Resources:
  MountPoint:
    Type: AWS::EC2::VolumeAttachment
    Condition: CreateProdResources
```

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Intrinsic Functions

Pink = must know

- Ref
- Fn::GetAtt
- Fn::FindInMap
- Fn::ImportValue
- Fn::Join
- Fn::Sub
- Fn::ForEach
- Fn::ToJsonString
- Fn::Base64
- Fn::Cidr
- Fn::GetAZs
- Fn::Select
- Fn::Split
- Fn::Transform
- Fn::Length
- Condition Functions (Fn::If, Fn::Not, Fn::Equals, etc…)

© Copyright by Cloud Mentor Pro | Confidential

Intrinsic Functions – Fn::Ref

- The Fn::Ref function can be leveraged to reference
  - **Parameters** – returns the value of the parameter
  - **Resources** – returns the physical ID of the underlying resource (e.g., EC2 ID)
- The shorthand for this in YAML is !Ref

```
Resources:
  DBSubnet1:
    Type: AWS::EC2::Subnet
    Properties:
      VpcId: !Ref MyVPC
```

© Copyright by Cloud Mentor Pro | Confidential

Intrinsic Functions – Fn::GetAtt

- Attributes are attached to any resources you create
- To know the attributes of your resources, the best place to look at is the documentation
- Example: the AZ of an EC2 instance!

```
Resources:
  EC2Instance:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: ami-0742b4e673072066f
      InstanceType: t2.micro
```

```
EBSVolume:
  Type: AWS::EC2::Volume
  Condition: CreateProdResources
  Properties:
    Size: 100
    AvailabilityZone: !GetAtt EC2Instance.AvailabilityZone
```

© Copyright by Cloud Mentor Pro | Confidential

Intrinsic Functions – Fn::FindInMap

- We use **Fn::FindInMap** to return a named value from a specific key
- !FindInMap [MapName, TopLevelKey, SecondLevelKey]

```
Mappings:
  RegionMap:
    us-east-1:
      HVM64: ami-0ff8a91507f77f867
      HVMG2: ami-0a584ac55a7631c0c
    us-west-1:
      HVM64: ami-0bdb828fd58c52235
      HVMG2: ami-066ee5fd4a9ef77f1

Resources:
  MyEC2Instance:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: !FindInMap [RegionMap, !Ref "AWS::Region", HVM64]
      InstanceType: t2.micro
```

© Copyright by Cloud Mentor Pro | Confidential

Intrinsic Functions – Fn::ImportValue

- Import values that are exported in other stacks
- For this, we use the Fn::ImportValue function

```
Resources:
  MySecureInstance:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: ami-0742b4e673072066f
      InstanceType: t2.micro
      AvailabilityZone: us-east-1a
      SecurityGroups:
        - !ImportValue SSHSecurityGroup
```

© Copyright by Cloud Mentor Pro | Confidential

Intrinsic Functions – Fn::Base64

- Convert **String** to it's **Base64** representation

```
!Base64 "ValueToEncode"
```

- Example: pass encoded data to EC2 Instance's **UserData** property

```
Resources:
  WebServer:
    Type: AWS::EC2::Instance
    Properties:
      ...
      UserData:
        Fn::Base64: |
          #!/bin/bash
          dnf update -y
          dnf install -y httpd
```

© Copyright by Cloud Mentor Pro | Confidential

Intrinsic Functions – Condition Functions

```
Conditions:
  CreateProdResources: !Equals [ !Ref EnvType, prod ]
```

- The logical ID is for you to choose. It's how you name condition
- The intrinsic function (logical) can be any of the following:
  - Fn::And
  - Fn::Equals
  - Fn::If
  - Fn::Not
  - Fn::Or

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Rollbacks

- Stack Creation Fails:
  - Default: everything rolls back (gets deleted). We can look at the log
  - Option to disable rollback and troubleshoot what happened

- Stack Update Fails:
  - The stack automatically rolls back to the previous known working state
  - Ability to see in the log what happened and error messages

- Rollback Failure? Fix resources manually then issue **ContinueUpdateRollback API** from Console
  - Or from the CLI using continue-update-rollback API call

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Service Role

- IAM role that allows CloudFormation to create/update/delete stack resources on your behalf
- Give ability to users to create/update/delete the stack resources even if they don't have permissions to work with the resources in the stack
- Use cases:
  - You want to achieve the least privilege principle
  - But you don't want to give the user all the required permissions to create the stack resources
- User must have **iam:PassRole** permissions

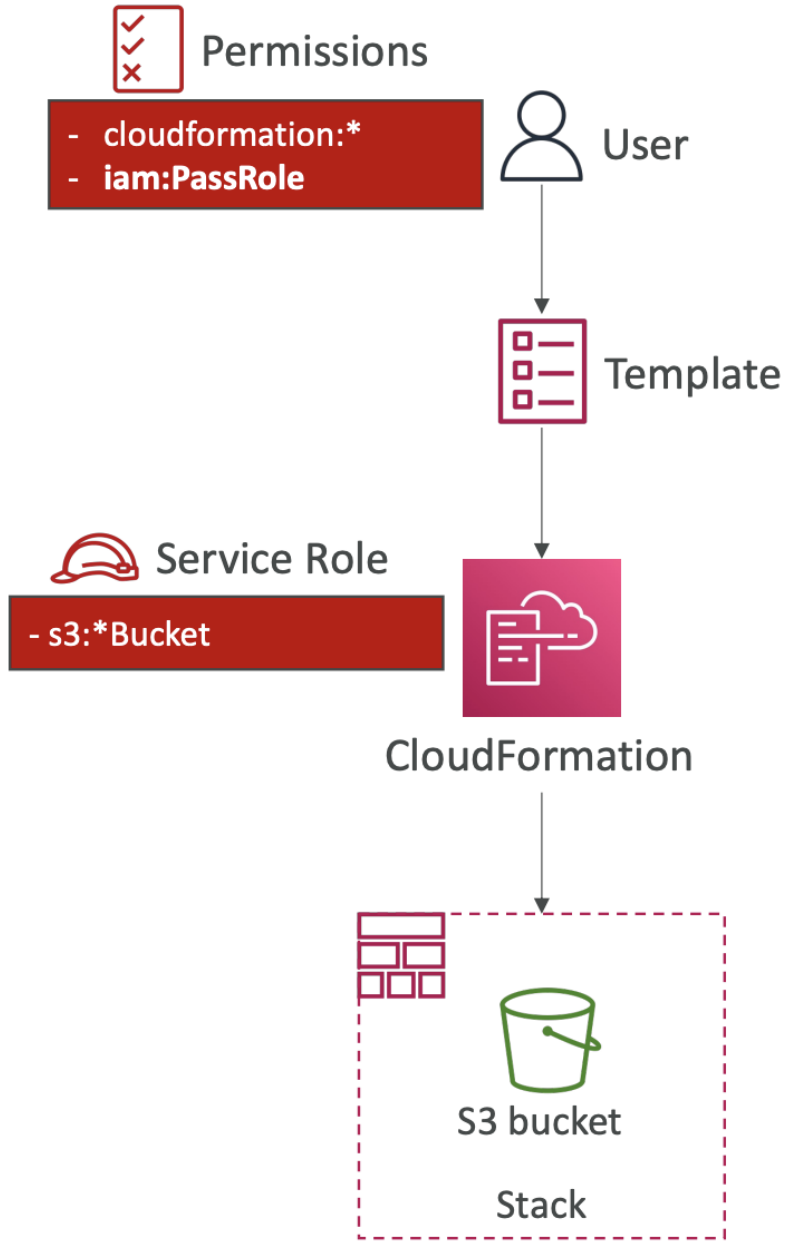

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – DeletionPolicy Delete

- **DeletionPolicy:**
  - Control what happens when the CloudFormation template is deleted or when a resource is removed from a CloudFormation template
  - Extra safety measure to preserve and backup resources
- Default **DeletionPolicy=Delete**
  - ⚠ **Delete** won't work on an S3 bucket if the bucket is not empty

```yaml
Resources:
  MyEC2Instance:
    Type: AWS::EC2::Instance
    Properties:
      ImageId: ami-12345678
      InstanceType: t2.micro
      KeyName: my-key-pair
      SecurityGroupIds:
        - sg-12345678
    DeletionPolicy: Delete
```

```yaml
Resources:
  MyS3Bucket:
    Type: AWS::S3::Bucket
    DeletionPolicy: Delete
```

⚠

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – DeletionPolicy Retain

- DeletionPolicy=Retain:
  - Specify on resources to preserve in case of CloudFormation deletes
- Works with any resources

```yaml
Resources:
  MyDynamoDBTable:
    Type: AWS::DynamoDB::Table
    Properties:
      TableName: MyTable
      AttributeDefinitions:
        - AttributeName: ID
          AttributeType: S
      KeySchema:
        - AttributeName: ID
          KeyType: HASH
      ProvisionedThroughput:
        ReadCapacityUnits: 5
        WriteCapacityUnits: 5
    DeletionPolicy: Retain
```

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – DeletionPolicy Snapshot

- **DeletionPolicy=Snapshot**
- Create one final snapshot before deleting the resource
- Examples of supported resources:
  - EBS Volume, ElastiCache Cluster, ElastiCache ReplicationGroup
  - RDS DBInstance, RDS DBCluster, Redshift Cluster, Neptune DBCluster, DocumentDB DBCluster

```yaml
Resources:
  MyDBInstance:
    Type: AWS::RDS::DBInstance
    Properties:
      DBInstanceClass: db.t2.micro
      AllocatedStorage: 20
      Engine: mysql
      MasterUsername: admin
      MasterUserPassword: "ExamplePassword"
    DeletionPolicy: Snapshot
```

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Stack Policies

- During a CloudFormation Stack update, all update actions are allowed on all resources (default)

- **A Stack Policy is a JSON document that defines the update actions that are allowed on specific resources during Stack updates**
- Protect resources from unintentional updates
- When you set a Stack Policy, all resources in the Stack are protected by default
- Specify an explicit ALLOW for the resources you want to be allowed to be updated

```
{
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "Update:*",
            "Principal": "*",
            "Resource": "*"
        },
        {
            "Effect": "Deny",
            "Action": "Update:*",
            "Principal": "*",
            "Resource": "LogicalResourceId/ProductionDatabase"
        }
    ]
}
```

Allow updates on all resources except the ProductionDatabase

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation –Termination Protection

- To prevent accidental deletes of CloudFormation Stacks, use **TerminationProtection**

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – Custom Resources

- Used to
  - define resources not yet supported by CloudFormation
  - define custom provisioning logic for resources can that be outside of CloudFormation (on-premises resources, 3rd party resources…)
  - have custom scripts run during create / update / delete through Lambda functions (running a Lambda function to empty an S3 bucket before being deleted)
- Defined in the template using
  AWS::CloudFormation::CustomResource or
  Custom::MyCustomResourceTypeName (recommended)
- Backed by a Lambda function (most common) or an SNS topic

© Copyright by Cloud Mentor Pro | Confidential

How to define a Custom Resource?

- **ServiceToken** specifies where CloudFormation sends requests to, such as Lambda ARN or SNS ARN (required & must be in the same region)
- Input data parameters (optional)

```yaml
Resources:
  MyCustomResourceUsingLambda:
    Type: Custom::MyLambdaResource
    Properties:
      ServiceToken: arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME
      # Input values (optional)
      ExampleProperty: "ExampleValue"
```

© Copyright by Cloud Mentor Pro | Confidential

Use Case – Delete content from an S3 bucket

- **You can't delete a non-empty S3 bucket**
- To delete a non-empty S3 bucket, you must first delete all the objects inside it
- We can use a custom resource to empty an S3 bucket before it gets deleted by CloudFormation

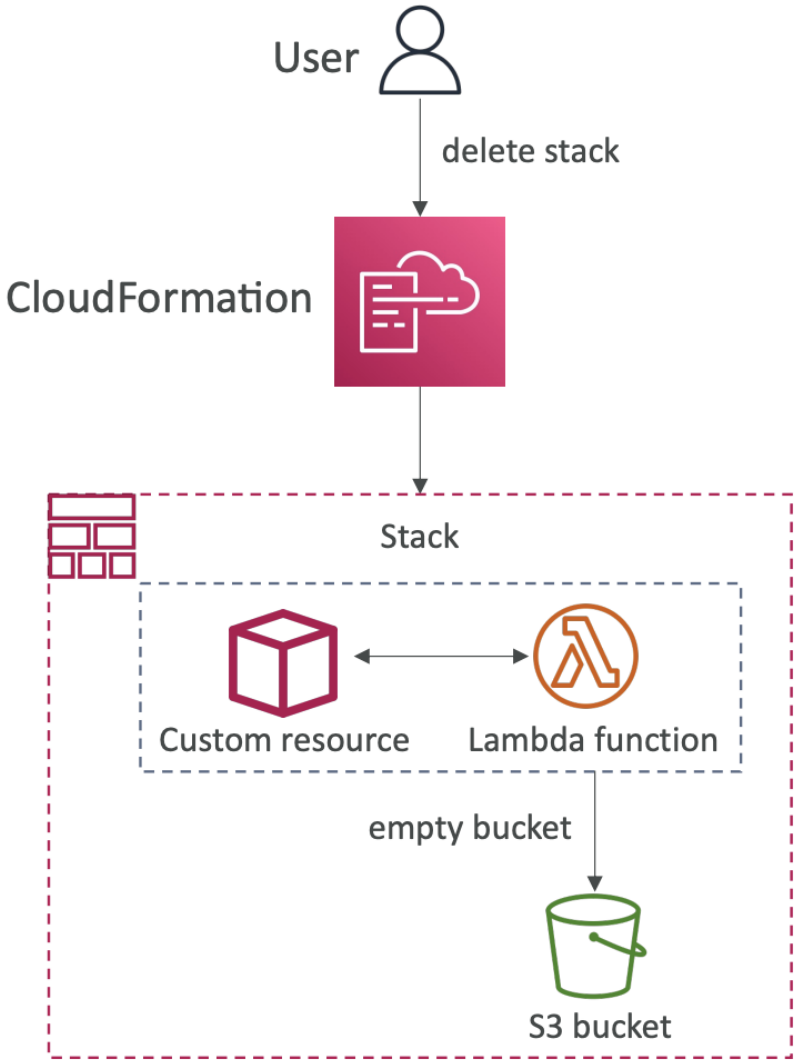

© Copyright by Cloud Mentor Pro | Confidential

CloudFormation – StackSets

- Create, update, or delete stacks across **multiple accounts and regions** with a single operation/template
- Target accounts to create, update, delete stack instances from StackSets
- When you update a stack set, *all* associated stack instances are updated throughout all accounts and regions
- **Can be applied into all accounts of an AWS Organization**
- Only Administrator account (or Delegated Administrator) can create StackSets

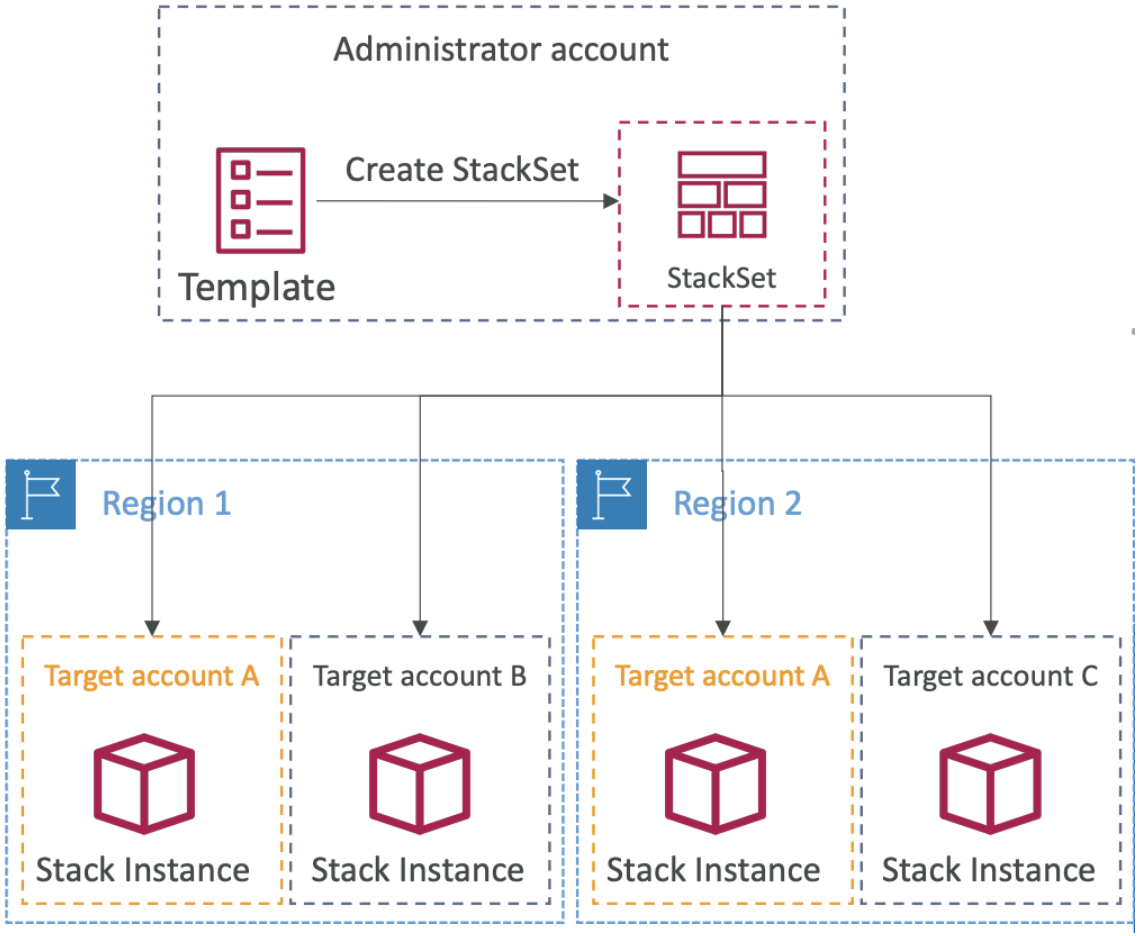

© Copyright by Cloud Mentor Pro | Confidential

Exercise

Create CloudFormation Stack with VPC, Subnet, ALB, EC2, etc...

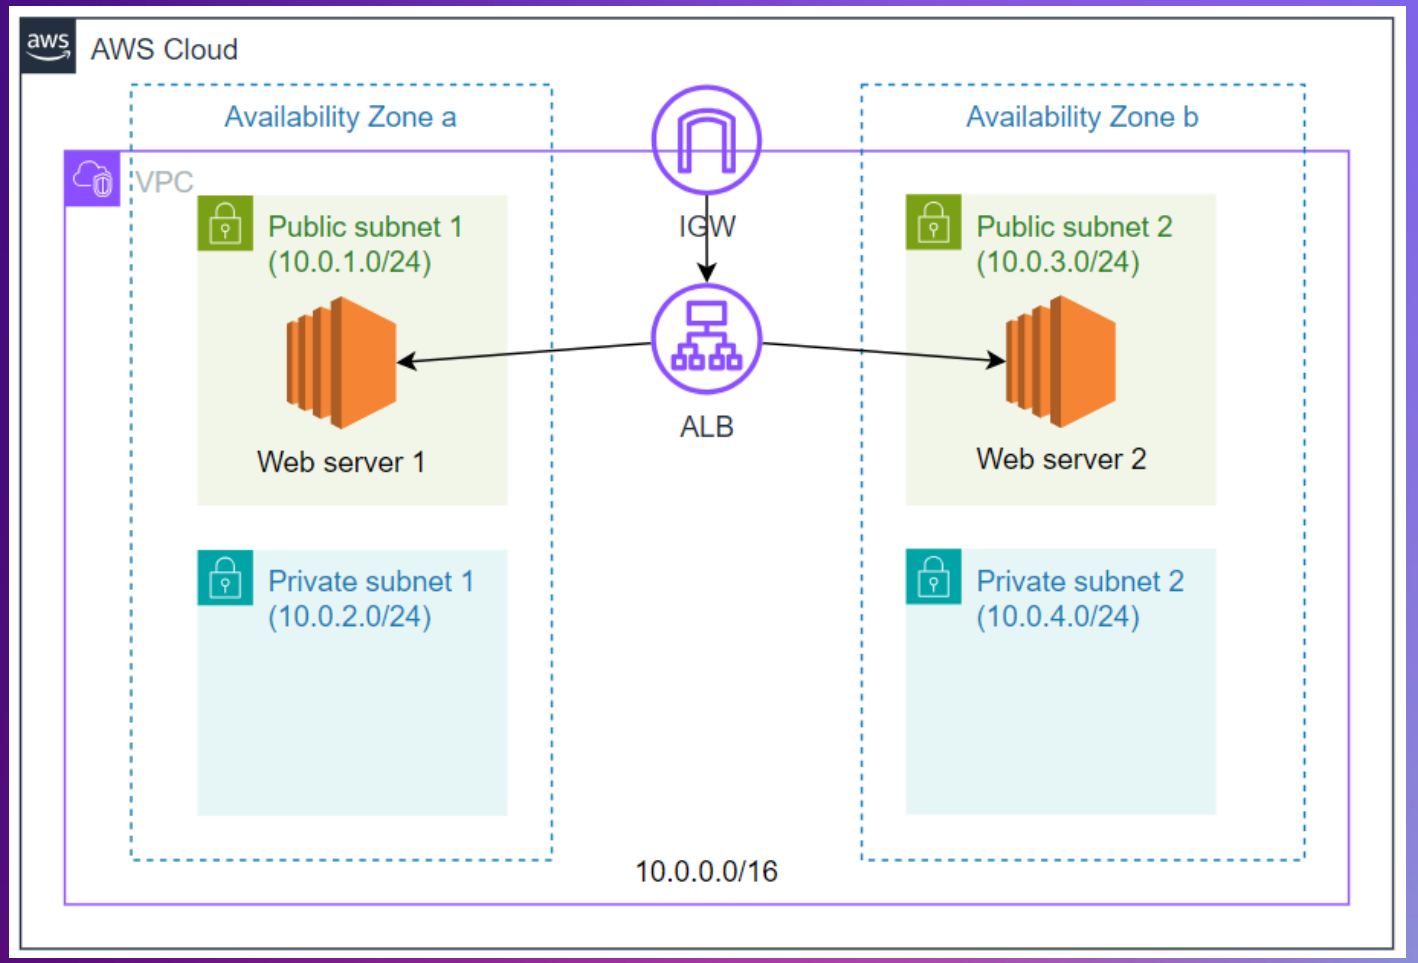

© Copyright by Cloud Mentor Pro | Confidential

Section 9

- AWS Integration & Messaging: SQS, SNS & Kinesis

© Copyright by Cloud Mentor Pro | Confidential

AWS Integration & Messaging

SQS, SNS & Kinesis

© Copyright by Cloud Mentor Pro | Confidential

Section Introduction

- When we start deploying multiple applications, they will inevitably need to communicate with one another
- There are two patterns of application communication

1) Synchronous communications (application to application)

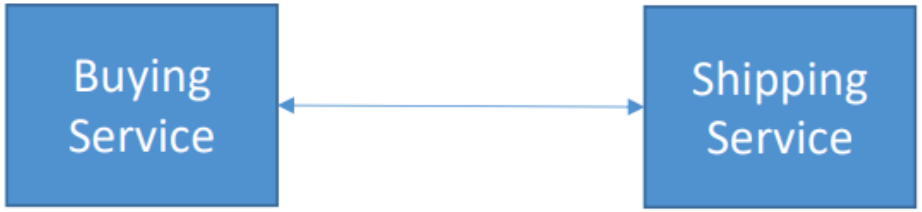

2) Asynchronous / Event based (application to queue to application)

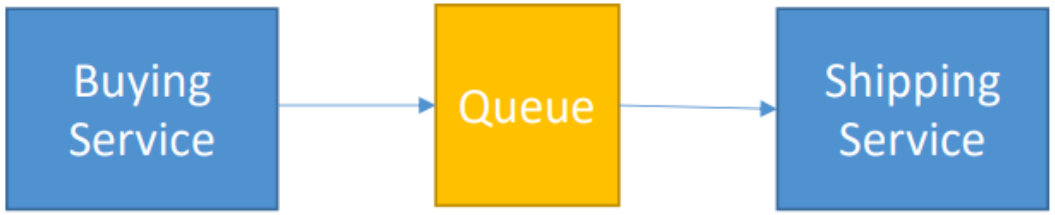

© Copyright by Cloud Mentor Pro | Confidential

Section Introduction

- Synchronous between applications can be problematic if there are sudden spikes of traffic
- What if you need to suddenly encode 1000 videos but usually it's 10?

- In that case, it's better to **decouple** your applications,
  - using SQS: queue model
  - using SNS: pub/sub model
  - using Kinesis: real-time streaming model
- These services can scale independently from our application!

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS - What's a queue?

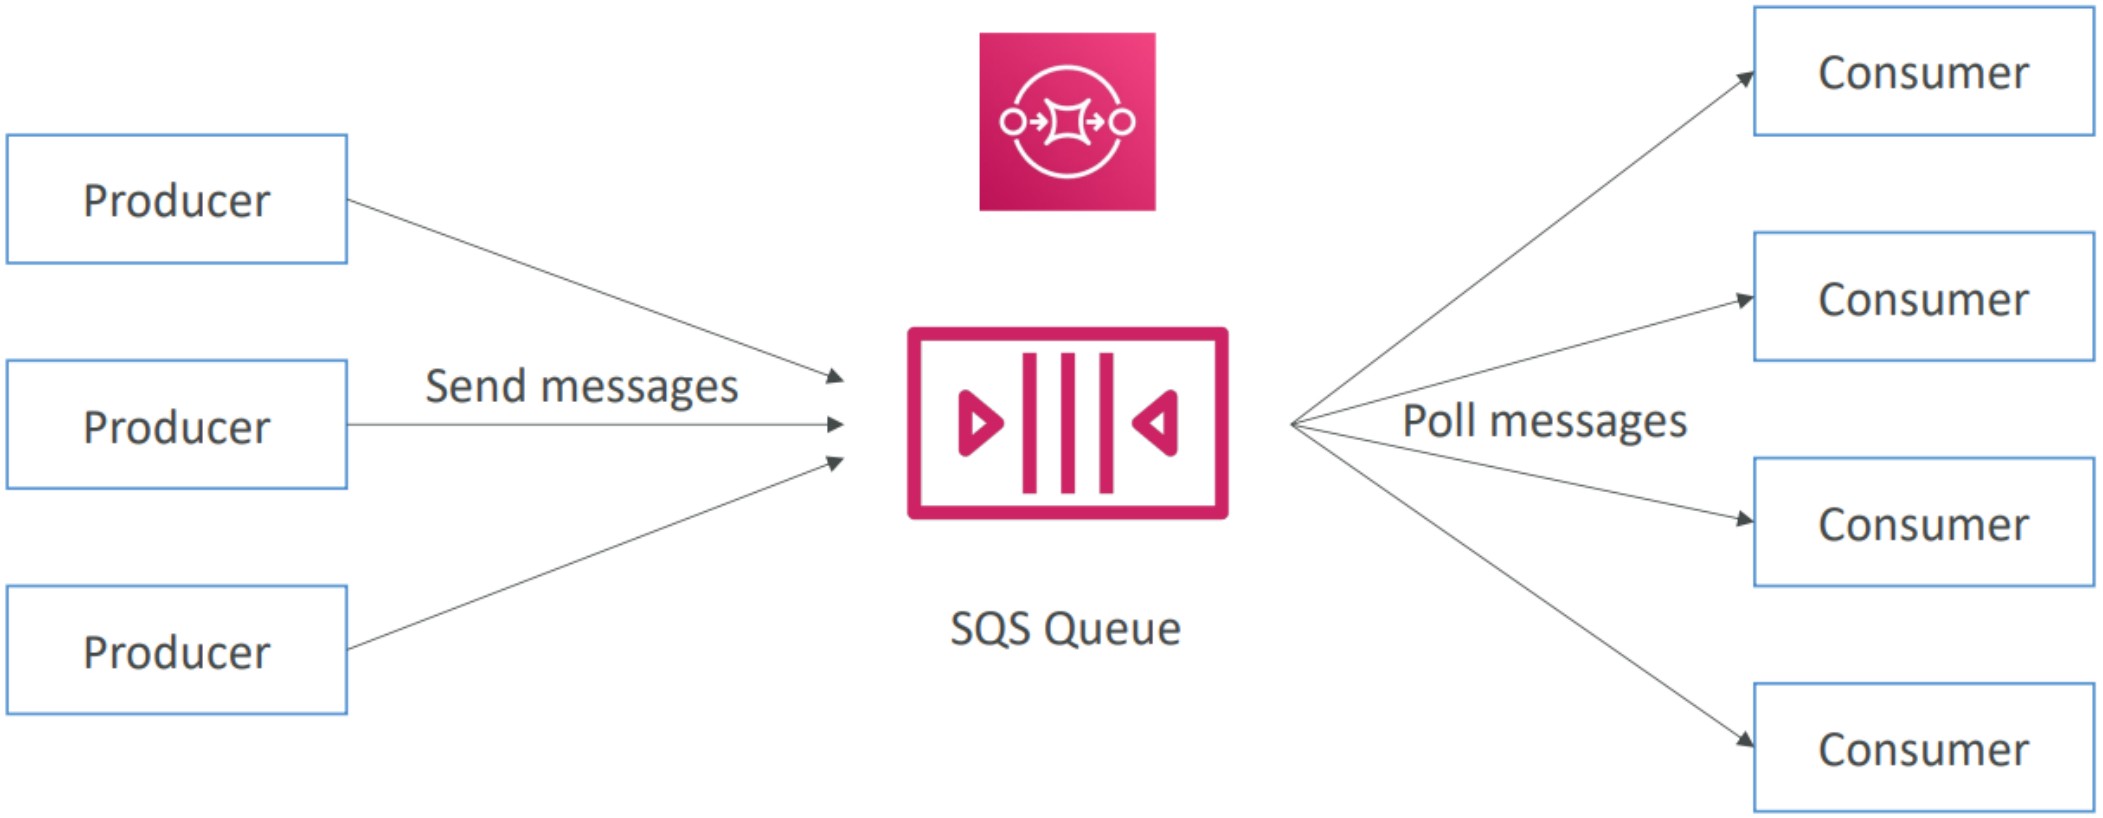

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS – Standard Queue

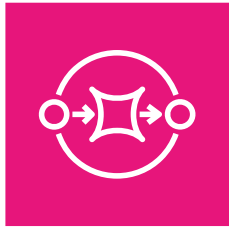

- Oldest offering (over 10 years old)
- Fully managed service, used to decouple applications

- Attributes:
  - Unlimited throughput, unlimited number of messages in queue
  - Default retention of messages: 4 days, maximum of 14 days
  - Low latency (<10 ms on publish and receive)
  - Limitation of 256KB per message sent

- Can have duplicate messages (at least once delivery, occasionally)
- Can have out of order messages (best effort ordering)

© Copyright by Cloud Mentor Pro | Confidential

SQS – Producing Messages

- Produced to SQS using the SDK (SendMessage API)
- The message is persisted in SQS until a consumer deletes it
- Message retention: default 4 days, up to 14 days
- Example: send an order to be processed
  - Order id
  - Customer id
  - Any attributes you want
- SQS standard: unlimited throughput

© Copyright by Cloud Mentor Pro | Confidential

SQS – Consuming Messages

- Consumers (running on EC2 instances, servers, or AWS Lambda)…
- Poll SQS for messages (receive up to 10 messages at a time)
- Process the messages (example: insert the message into an RDS database)
- Delete the messages using the DeleteMessage API

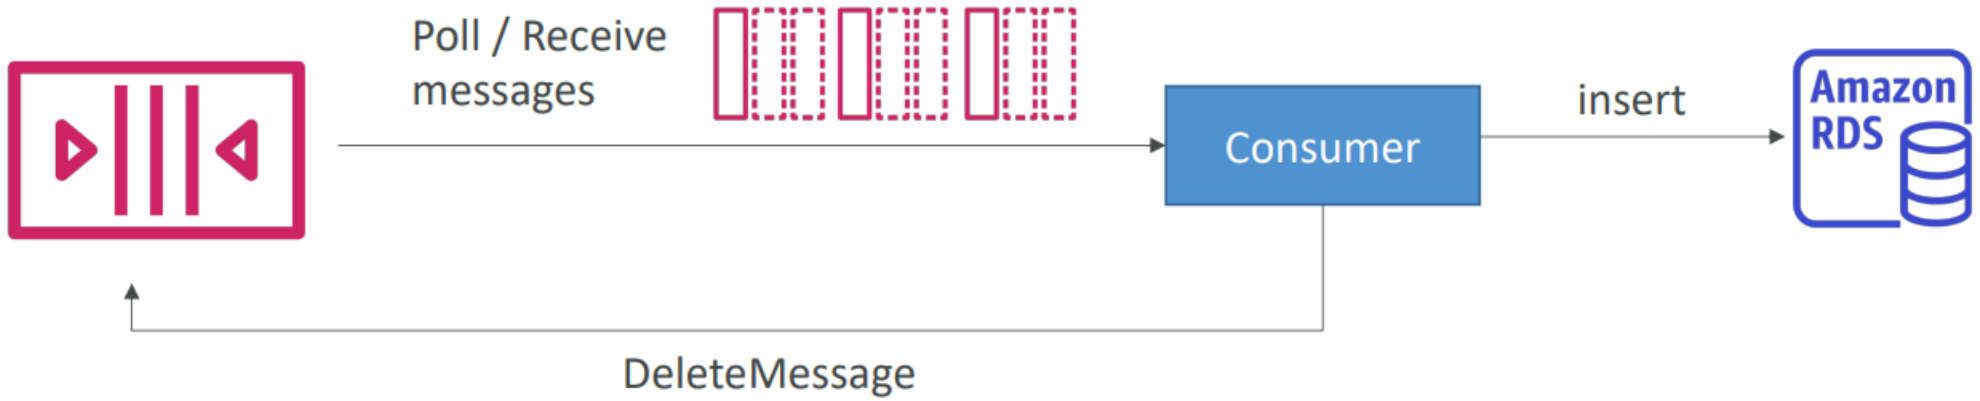

© Copyright by Cloud Mentor Pro | Confidential

SQS – Multiple EC2 Instances Consumers

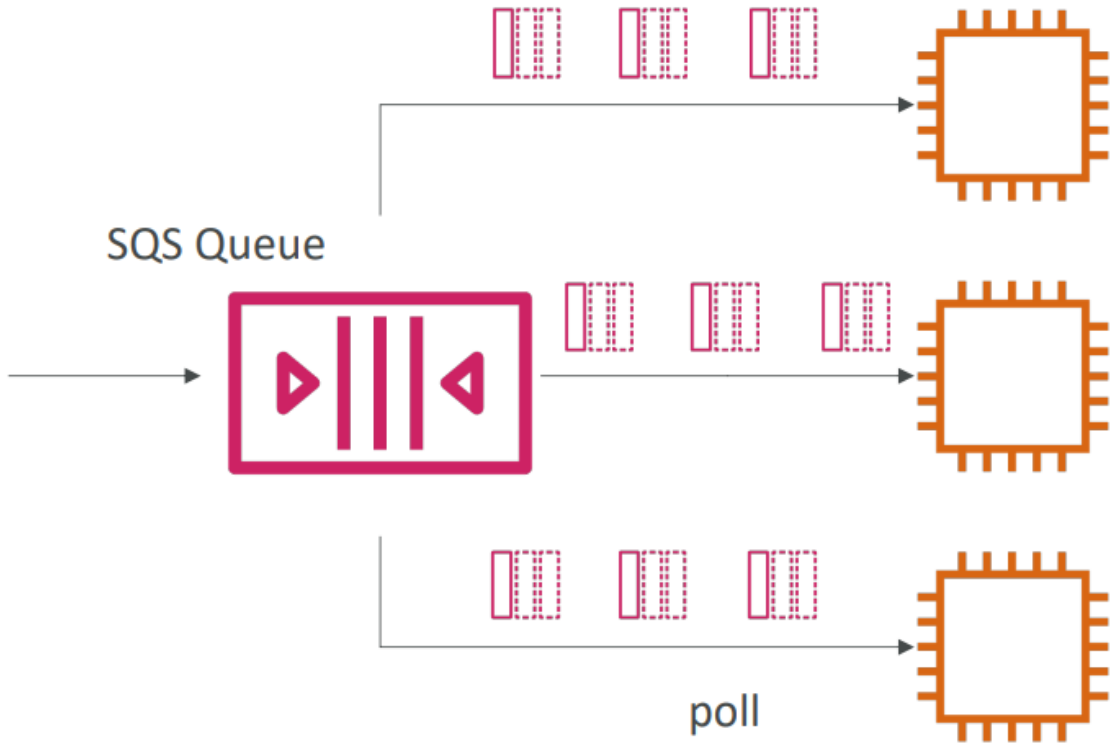

- Consumers receive and process messages in parallel
- At least once delivery
- Best-effort message ordering
- Consumers delete messages after processing them
- We can scale consumers horizontally to improve throughput of processing

© Copyright by Cloud Mentor Pro | Confidential

SQS with Auto Scaling Group (ASG)

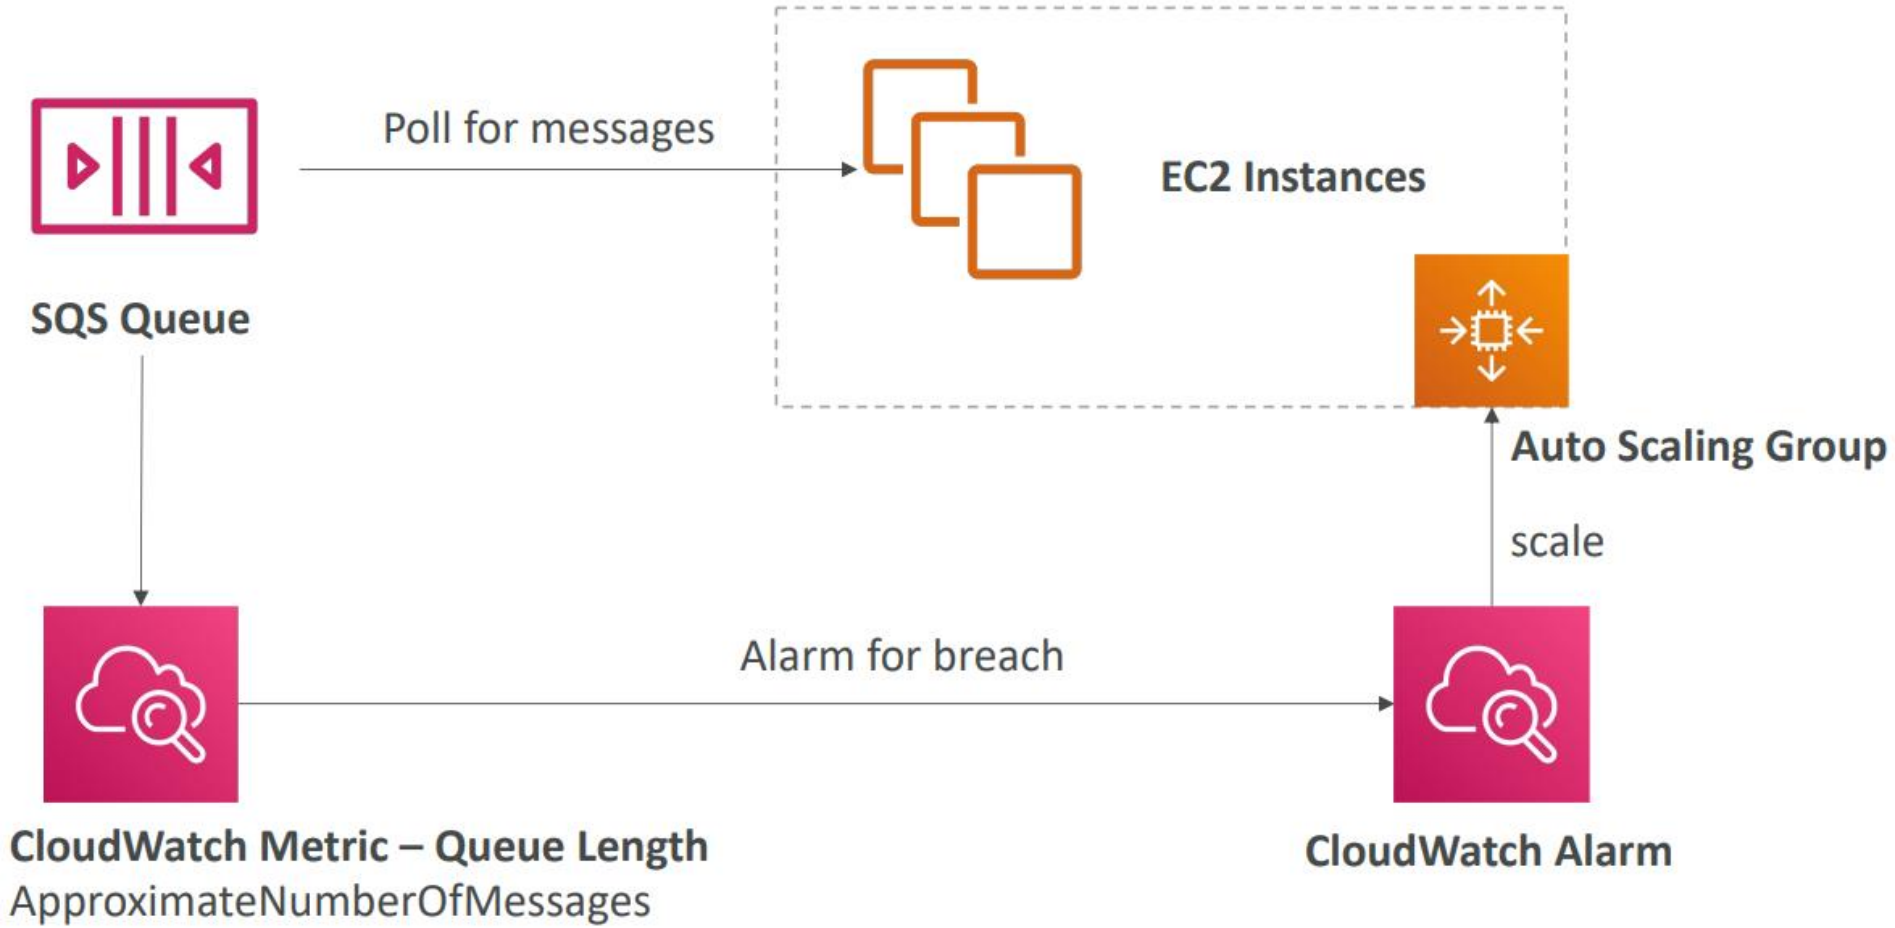

© Copyright by Cloud Mentor Pro | Confidential

SQS to decouple between application tiers

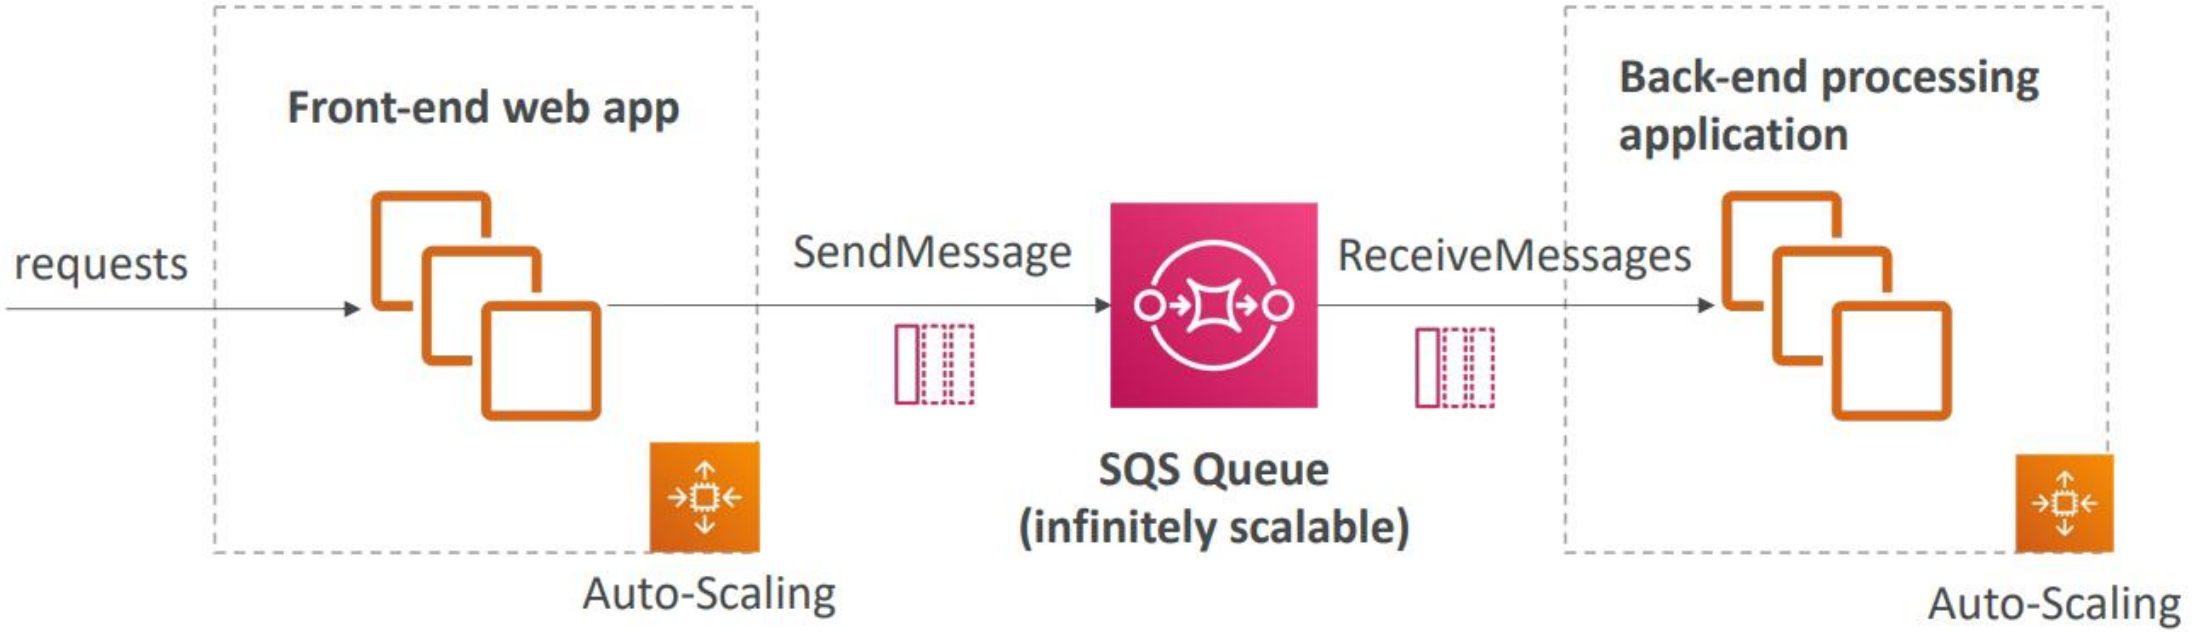

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS - Security

- **Encryption**:
  - In-flight encryption using HTTPS API
  - At-rest encryption using KMS keys
  - Client-side encryption if the client wants to perform encryption/decryption itself

- **Access Controls:** IAM policies to regulate access to the SQS API

- **SQS Access Policies** (similar to S3 bucket policies)
  - Useful for cross-account access to SQS queues
  - Useful for allowing other services (SNS, S3…) to write to an SQS queue

© Copyright by Cloud Mentor Pro | Confidential

SQS Queue Access Policy

Cross Account Access

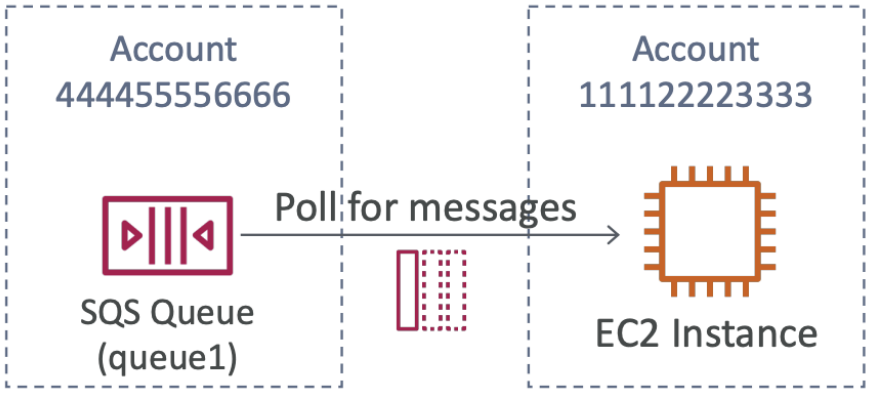


```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Principal": { "AWS": [ "111122223333" ] },
    "Action": [ "sqs:ReceiveMessage" ],
    "Resource": "arn:aws:sqs:us-east-1:444455556666:queue1"
  }]
}
```

Publish S3 Event Notifications To SQS Queue

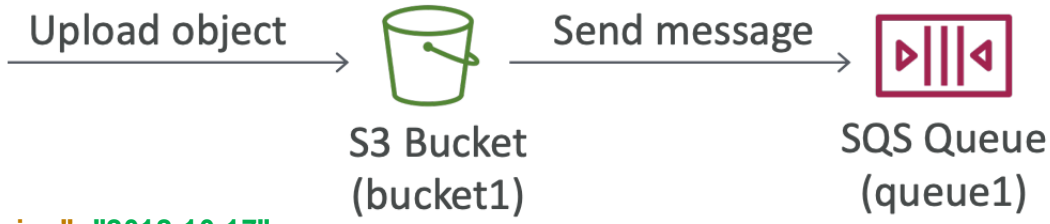


```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Principal": { "AWS": "*" },
    "Action": [ "sqs:SendMessage" ],
    "Resource": "arn:aws:sqs:us-east-1:444455556666:queue1",
    "Condition": {
      "ArnLike": { "aws:SourceArn": "arn:aws:s3:*:*:bucket1" },
      "StringEquals": { "aws:SourceAccount": "<bucket1_owner_account_id>" }
    }
  }]
}
```

© Copyright by Cloud Mentor Pro | Confidential

SQS – Message Visibility Timeout

- After a message is polled by a consumer, it becomes **invisible** to other consumers
- By default, the "message visibility timeout" is **30 seconds**
- That means the message has 30 seconds to be processed
- After the message visibility timeout is over, the message is "visible" in SQS

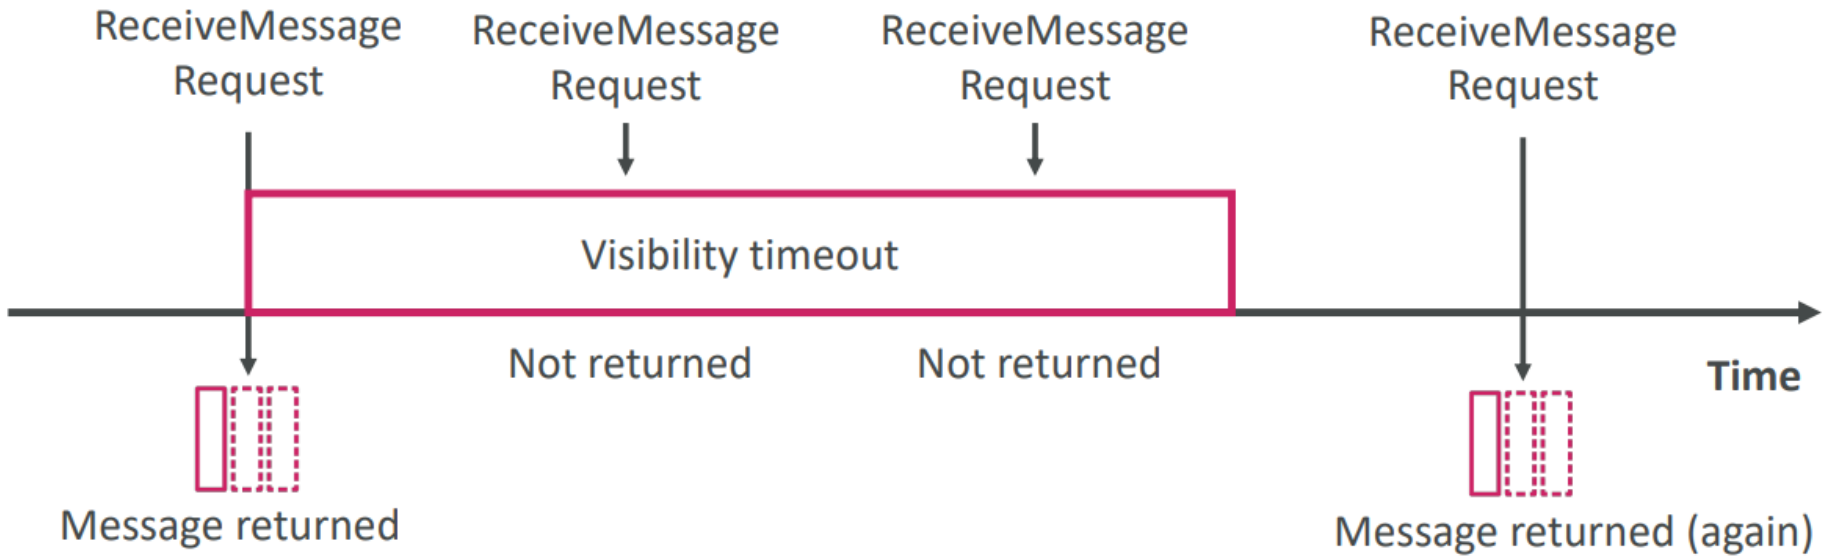

© Copyright by Cloud Mentor Pro | Confidential

SQS – Message Visibility Timeout

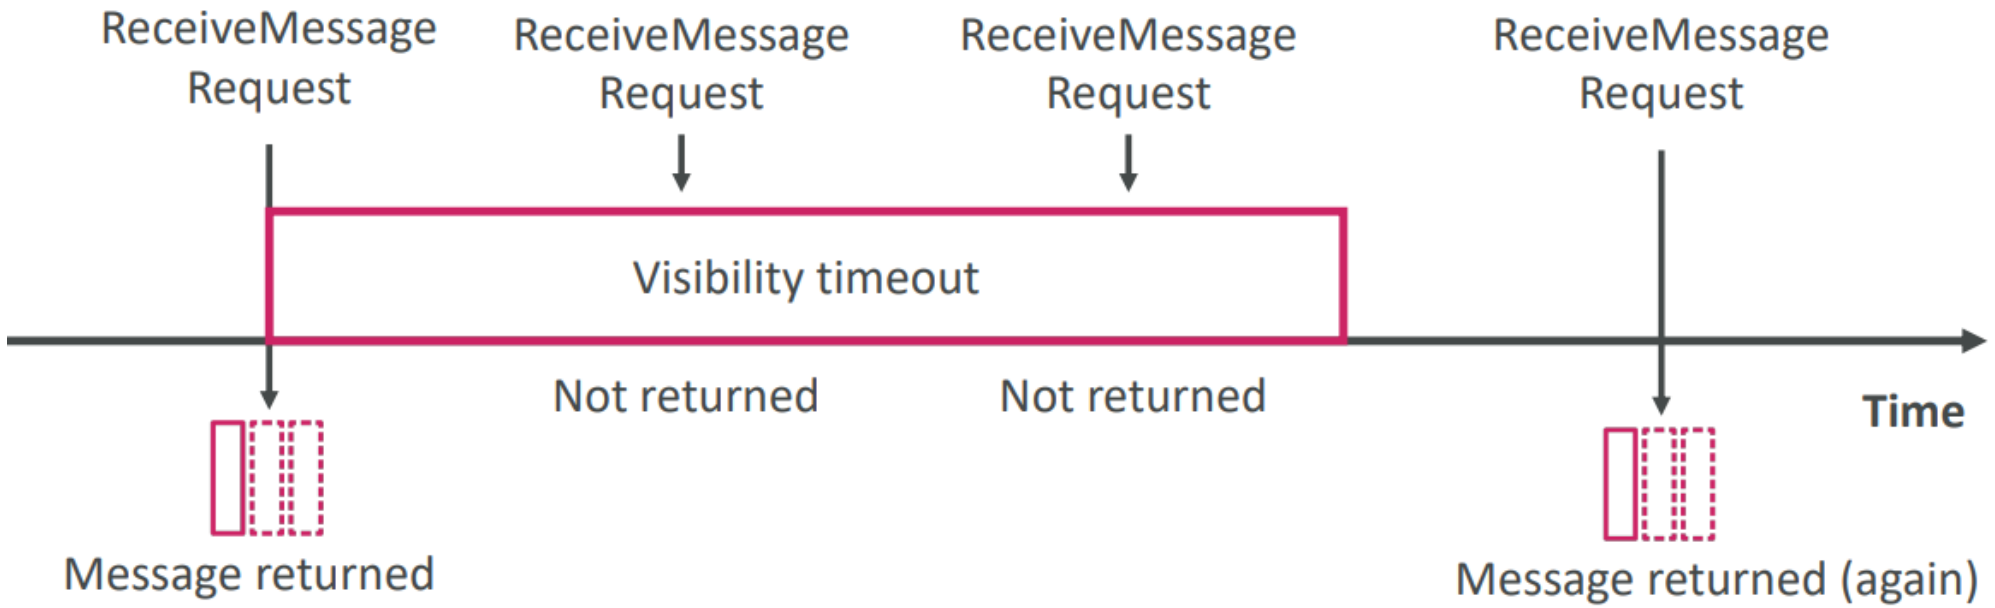

- If a message is not processed within the visibility timeout, it will be processed **twice**
- A consumer could call the **ChangeMessageVisibility** API to get more time
- If visibility timeout is high (hours), and consumer crashes, re-processing will take time
- If visibility timeout is too low (seconds), we may get duplicates

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS – Dead Letter Queue (DLQ)

- If a consumer fails to process a message within the Visibility Timeout… the message goes back to the queue!
- We can set a threshold of how many times a message can go back to the queue
- After the MaximumReceives threshold is exceeded, the message goes into a Dead Letter Queue (DLQ)
- Useful for debugging!
- DLQ of a FIFO queue must also be a FIFO queue
- DLQ of a Standard queue must also be a Standard queue
- Make sure to process the messages in the DLQ before they expire:
  - Good to set a retention of 14 days in the DLQ

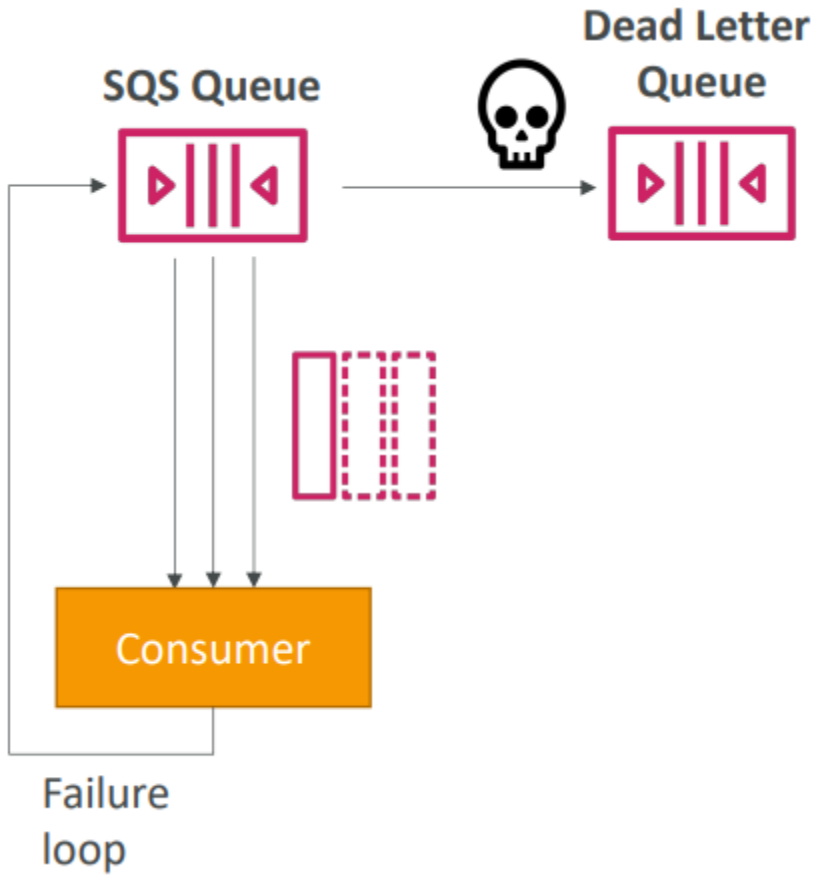

© Copyright by Cloud Mentor Pro | Confidential

SQS DLQ – Redrive to Source

- Feature to help consume messages in the DLQ to understand what is wrong with them
- When our code is fixed, we can redrive the messages from the DLQ back into the source queue (or any other queue) in batches without writing custom code

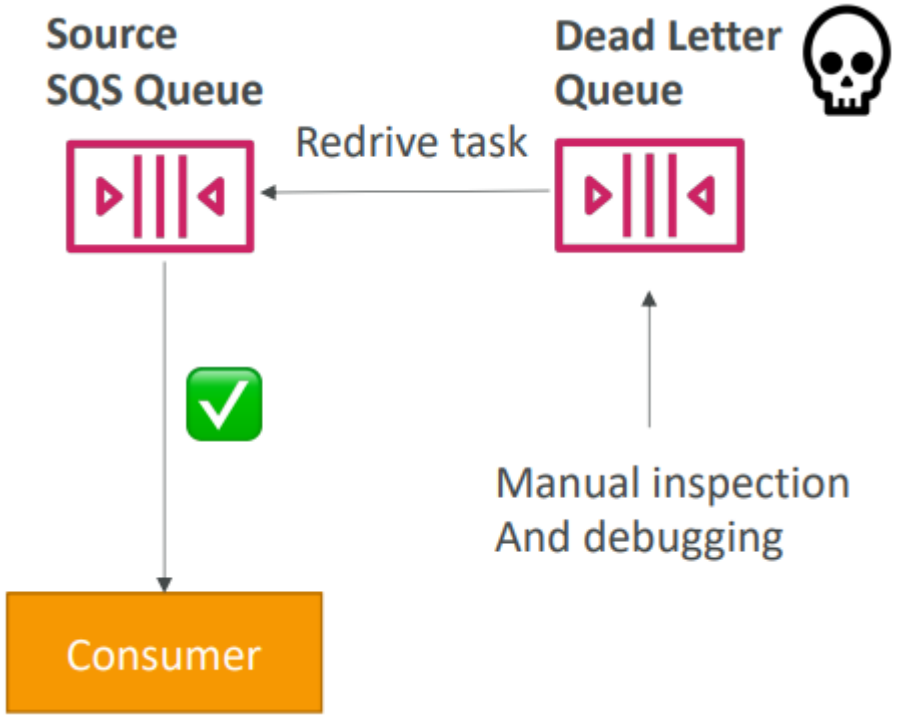

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS – Delay Queue

- Delay a message (consumers don’t see it immediately) up to 15 minutes
- Default is 0 seconds (message is available right away)
- Can set a default at queue level
- Can override the default on send using the DelaySeconds parameter

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS - Long Polling

- When a consumer requests messages from the queue, it can optionally "wait" for messages to arrive if there are none in the queue
- This is called Long Polling
- **LongPolling decreases the number of API calls made to SQS while increasing the efficiency and reducing latency of your application**
- The wait time can be between 1 sec to 20 sec (20 sec preferable)
- Long Polling is preferable to Short Polling
- Long polling can be enabled at the queue level or at the API level using WaitTimeSeconds

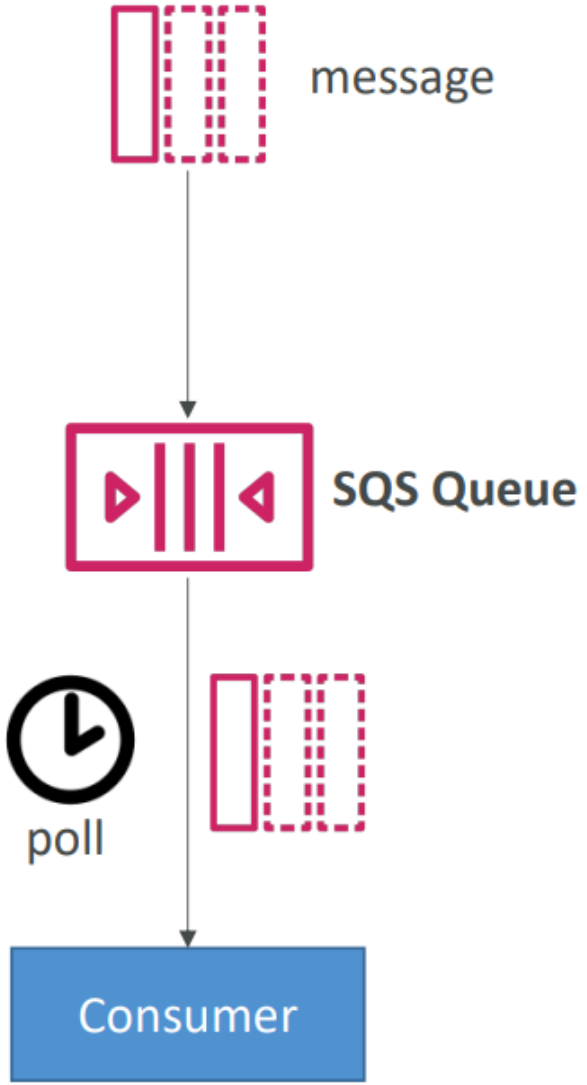

© Copyright by Cloud Mentor Pro | Confidential

SQS Extended Client

- Message size limit is 256KB, how to send large messages, e.g. 1GB?
- Using the SQS Extended Client (Java Library)

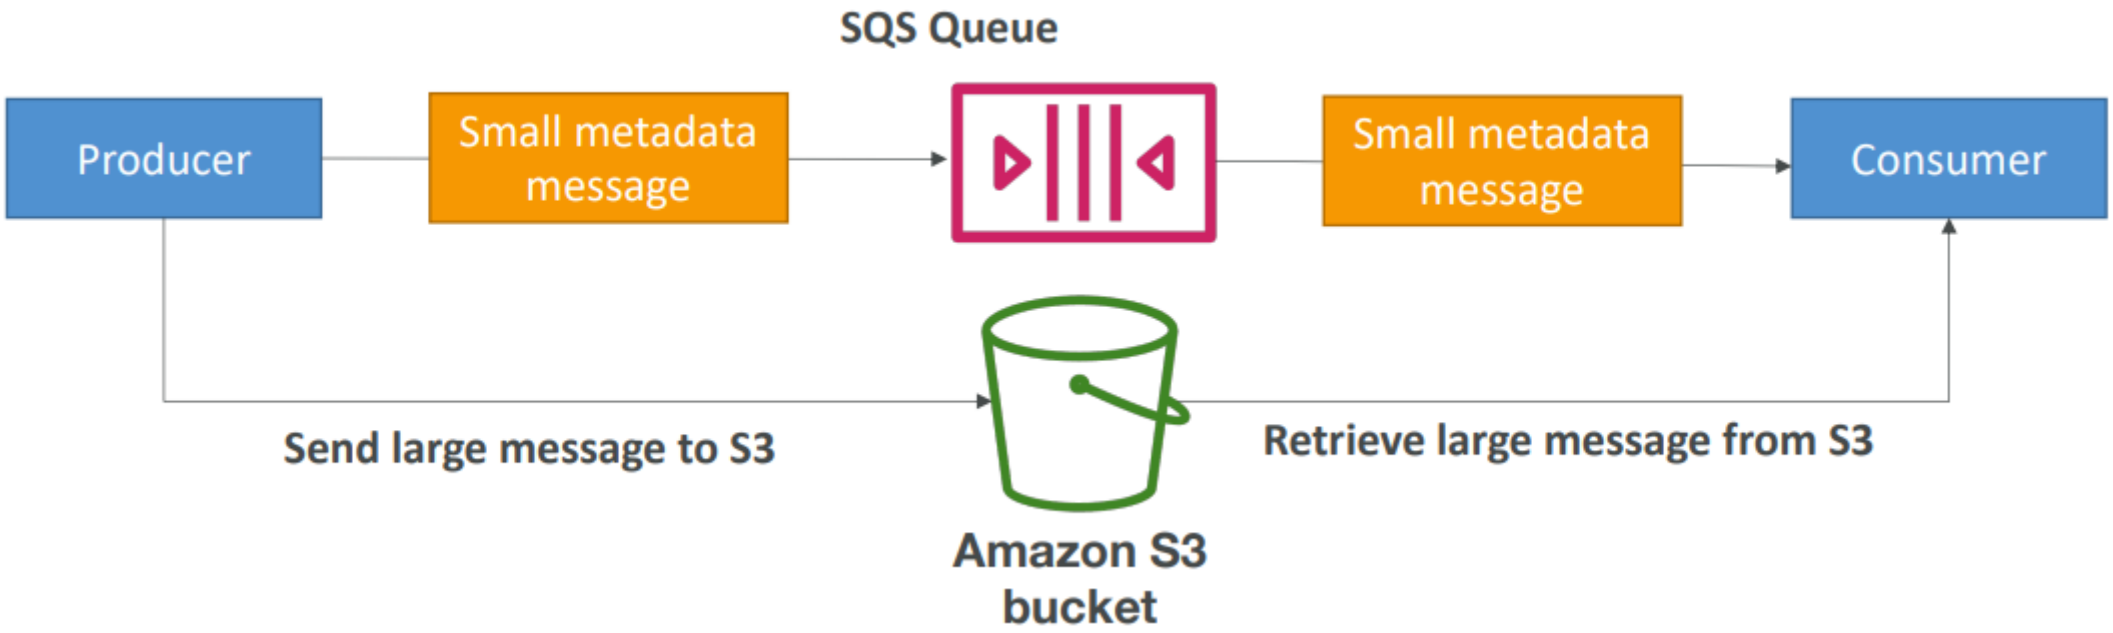

© Copyright by Cloud Mentor Pro | Confidential

SQS – Must know API

- **CreateQueue (MessageRetentionPeriod), DeleteQueue**
- **PurgeQueue**: delete all the messages in queue
- **SendMessage (DelaySeconds), ReceiveMessage, DeleteMessage**
- **MaxNumberOfMessages**: default 1, max 10 (for ReceiveMessage API)
- **ReceiveMessageWaitTimeSeconds**: Long Polling
- **ChangeMessageVisibility**: change the message timeout

- Batch APIs for **SendMessage, DeleteMessage, ChangeMessageVisibility** helps decrease your costs

© Copyright by Cloud Mentor Pro | Confidential

Amazon SQS – FIFO Queue

- FIFO = First In First Out (ordering of messages in the queue)

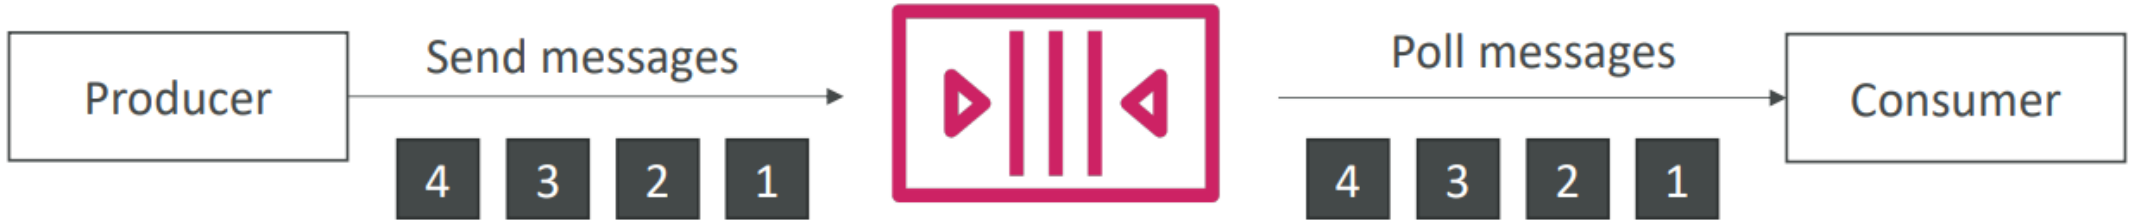

- Limited throughput: 300 msg/s without batching, 3000 msg/s with batching
- Exactly-once send capability (by removing duplicates)
- Messages are processed in order by the consumer

© Copyright by Cloud Mentor Pro | Confidential

SQS FIFO – Deduplication

- De-duplication interval is 5 minutes
- Two de-duplication methods:
  - Content-based deduplication: will do a SHA-256 hash of the message body
  - Explicitly provide a Message Deduplication ID

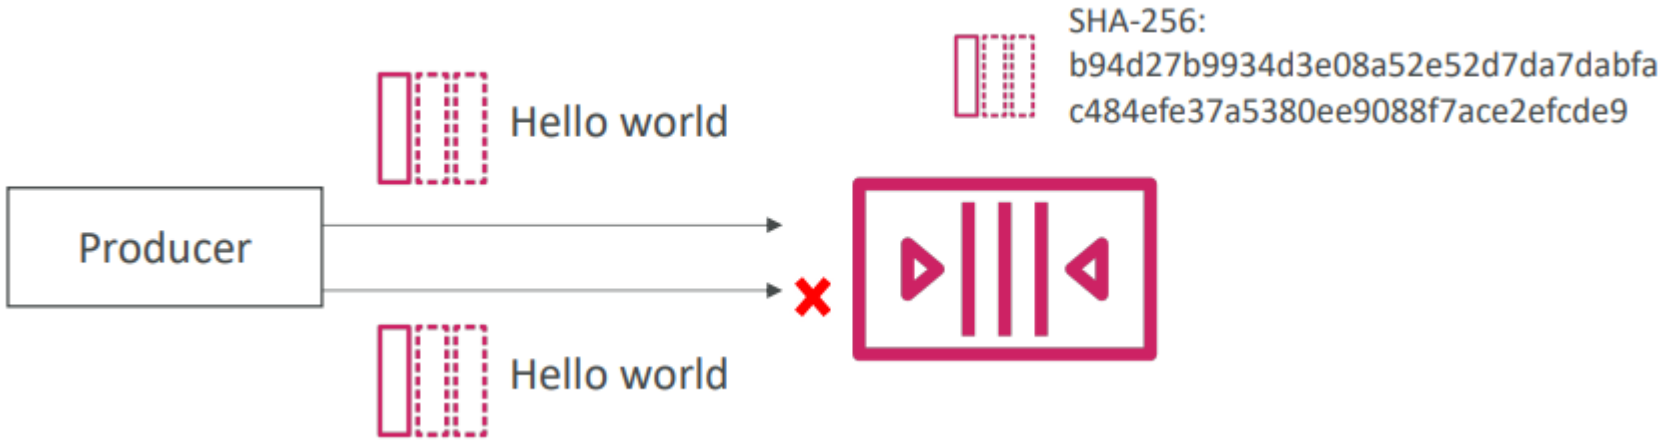

© Copyright by Cloud Mentor Pro | Confidential

SQS FIFO – Message Grouping

- If you specify the same value of MessageGroupID in an SQS FIFO queue, you can only have one consumer, and all the messages are in order
- To get ordering at the level of a subset of messages, specify different values for MessageGroupID
  - Messages that share a common Message Group ID will be in order within the group
  - Each Group ID can have a different consumer (parallel processing!)
  - Ordering across groups is not guaranteed

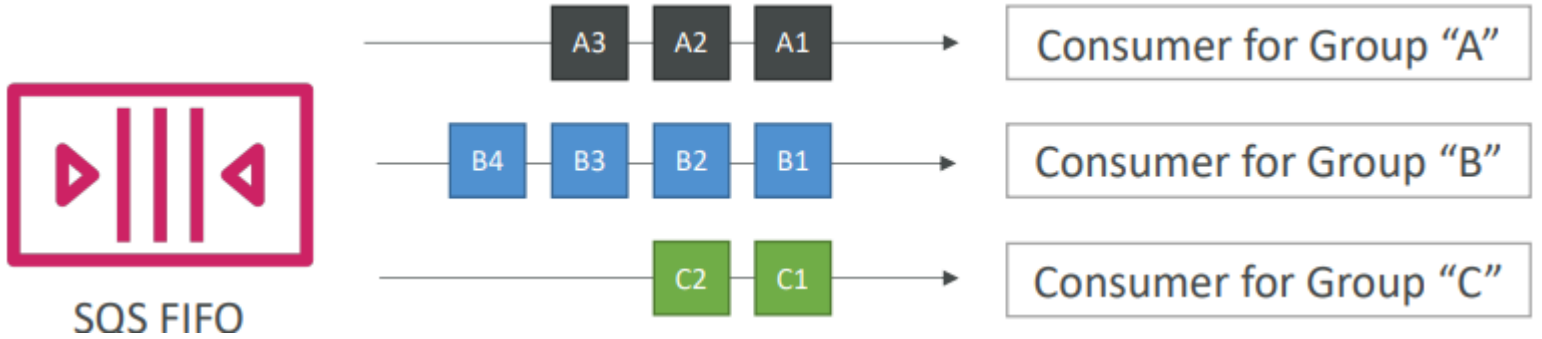

© Copyright by Cloud Mentor Pro | Confidential

Amazon SNS

- What if you want to send one message to many receivers?

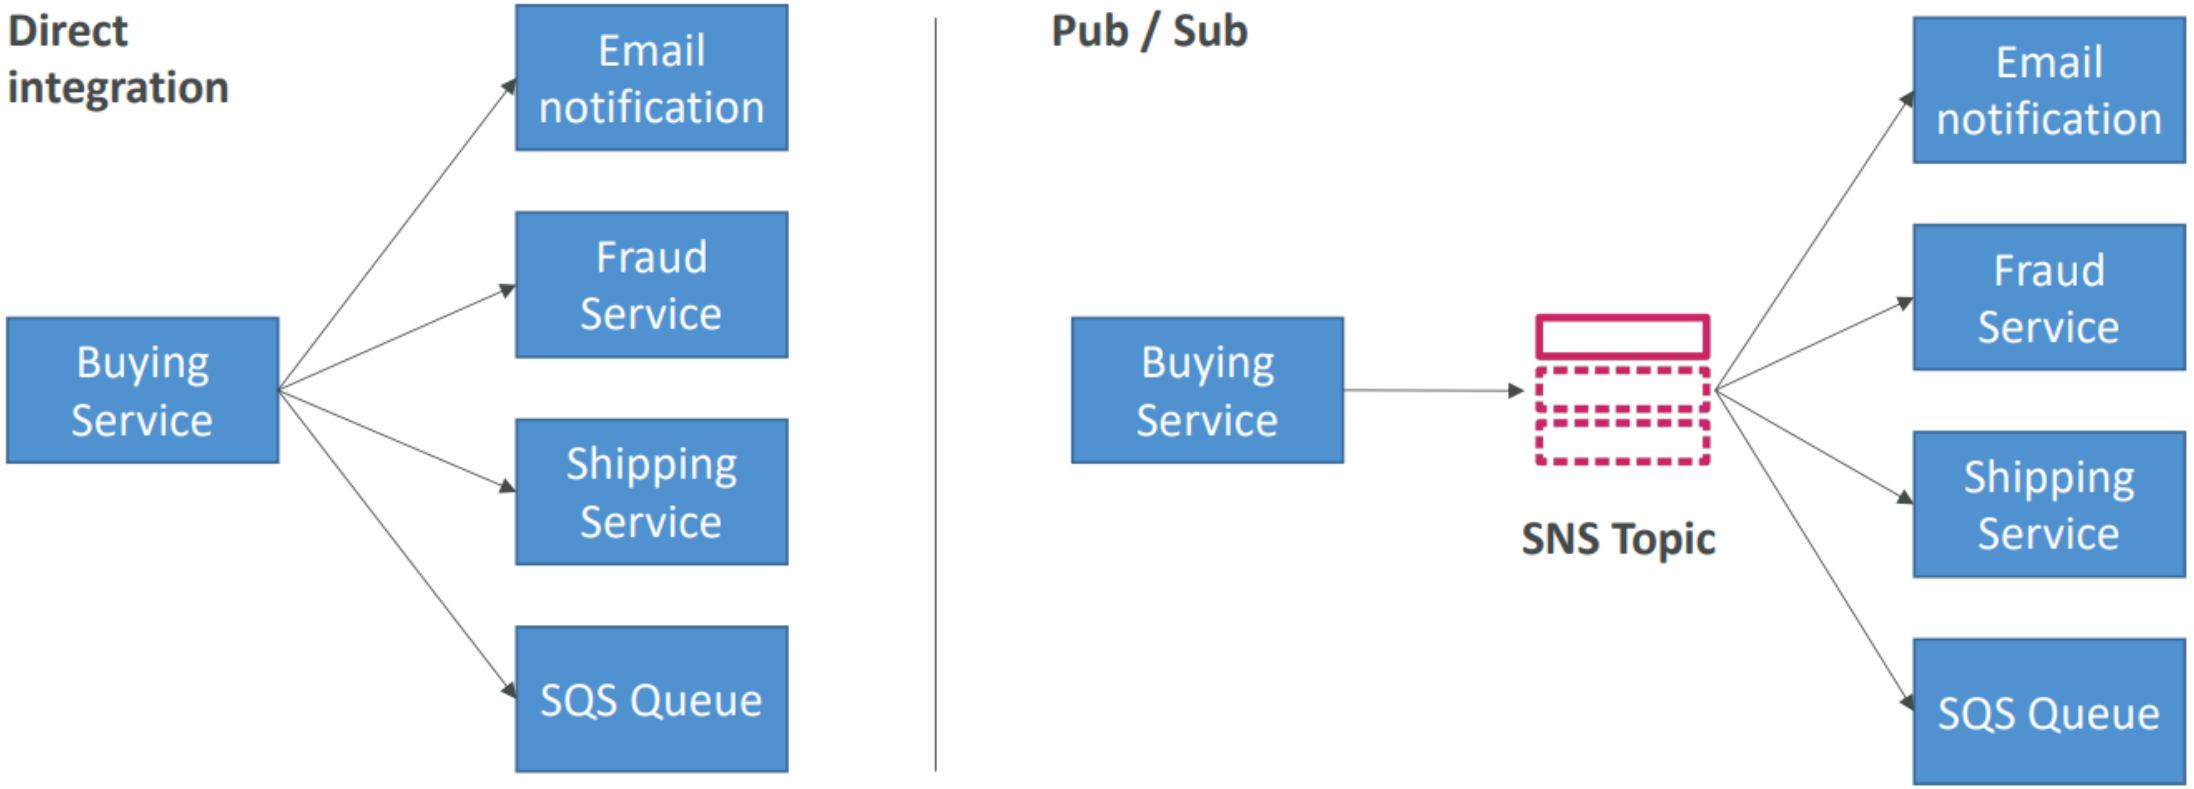

© Copyright by Cloud Mentor Pro | Confidential

Amazon SNS

- The “event producer” only sends message to one SNS topic
- As many “event receivers” (subscriptions) as we want to listen to the SNS topic notifications
- Each subscriber to the topic will get all the messages (note: new feature to filter messages)
- Up to 12,500,000 subscriptions per topic
- 100,000 topics limit

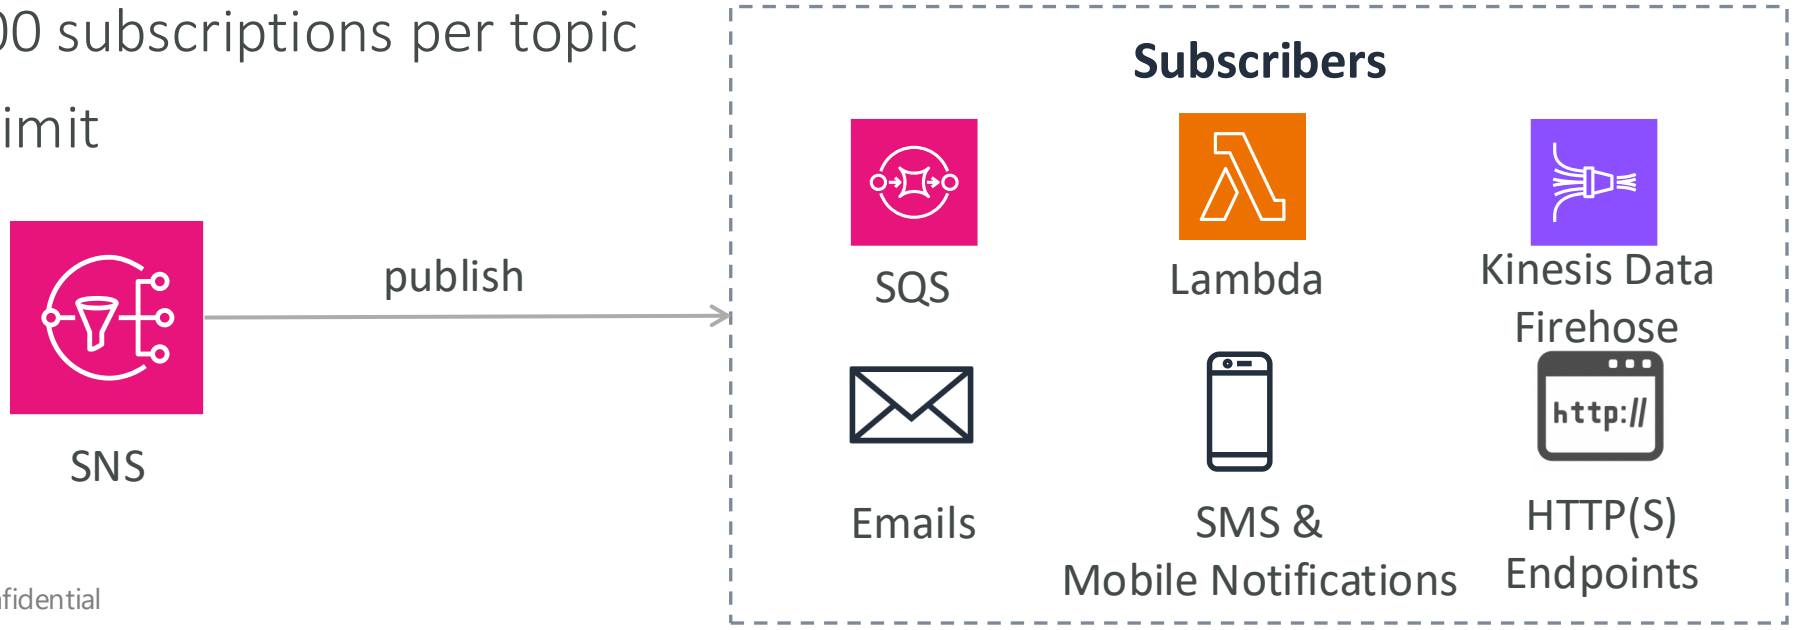

© Copyright by Cloud Mentor Pro | Confidential

SNS integrates with a lot of AWS services

- Many AWS services can send data directly to SNS for notifications

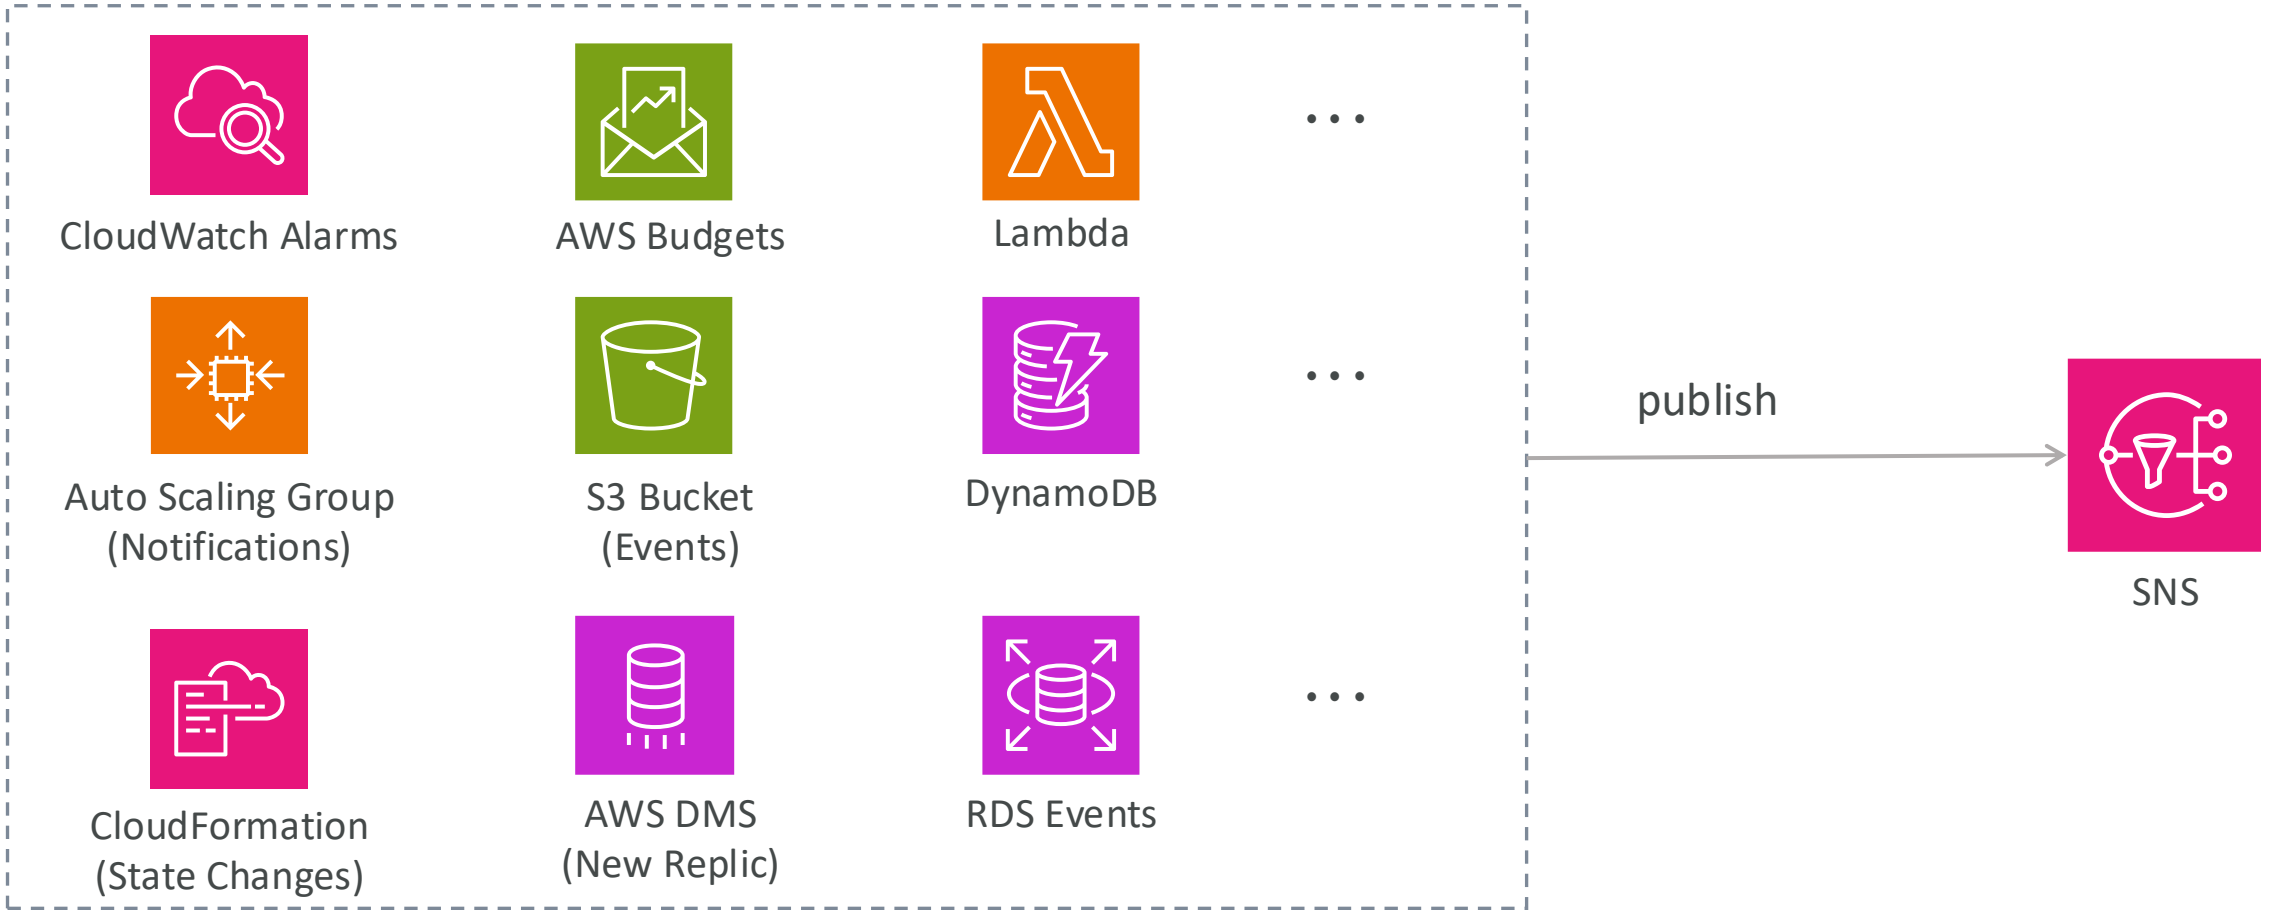

© Copyright by Cloud Mentor Pro | Confidential

Amazon SNS – How to publish

- Topic Publish (using the SDK)
  - Create a topic
  - Create a subscription (or many)
  - Publish to the topic

- Direct Publish (for mobile apps SDK)
  - Create a platform application
  - Create a platform endpoint
  - Publish to the platform endpoint
  - Works with Google GCM, Apple APNS, Amazon ADM…

© Copyright by Cloud Mentor Pro | Confidential

Amazon SNS – Security

- **Encryption**:
  - In-flight encryption using HTTPS API
  - At-rest encryption using KMS keys
  - Client-side encryption if the client wants to perform encryption/decryption itself

- **Access Controls**: IAM policies to regulate access to the SNS API

- **SNS Access Policies** (similar to S3 bucket policies)
  - Useful for cross-account access to SNS topics
  - Useful for allowing other services (S3…) to write to an SNS topic

© Copyright by Cloud Mentor Pro | Confidential

SNS + SQS: Fan Out

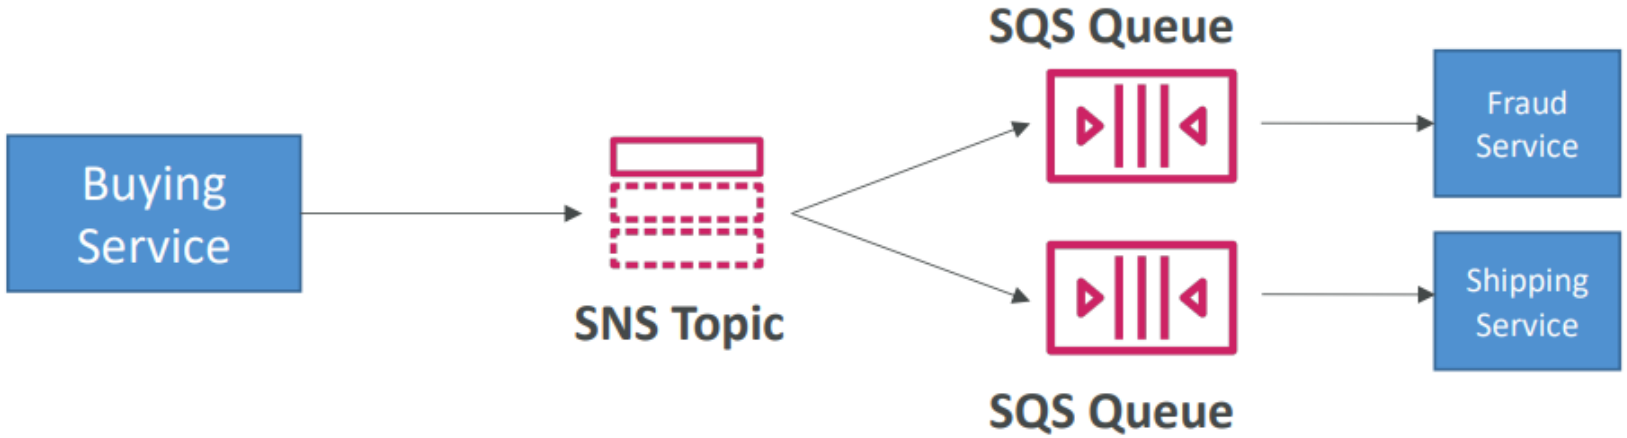

- Push once in SNS, receive in all SQS queues that are subscribers
- Fully decoupled, no data loss
- SQS allows for: data persistence, delayed processing and retries of work
- Ability to add more SQS subscribers over time
- Make sure your SQS queue access policy allows for SNS to write
- Cross-Region Delivery: works with SQS Queues in other regions

© Copyright by Cloud Mentor Pro | Confidential

Application: S3 Events to multiple queues

- For the same combination of: event type (e.g. object create) and prefix (e.g. images/) you can only have one S3 Event rule
- If you want to send the same S3 event to many SQS queues, use fan-out

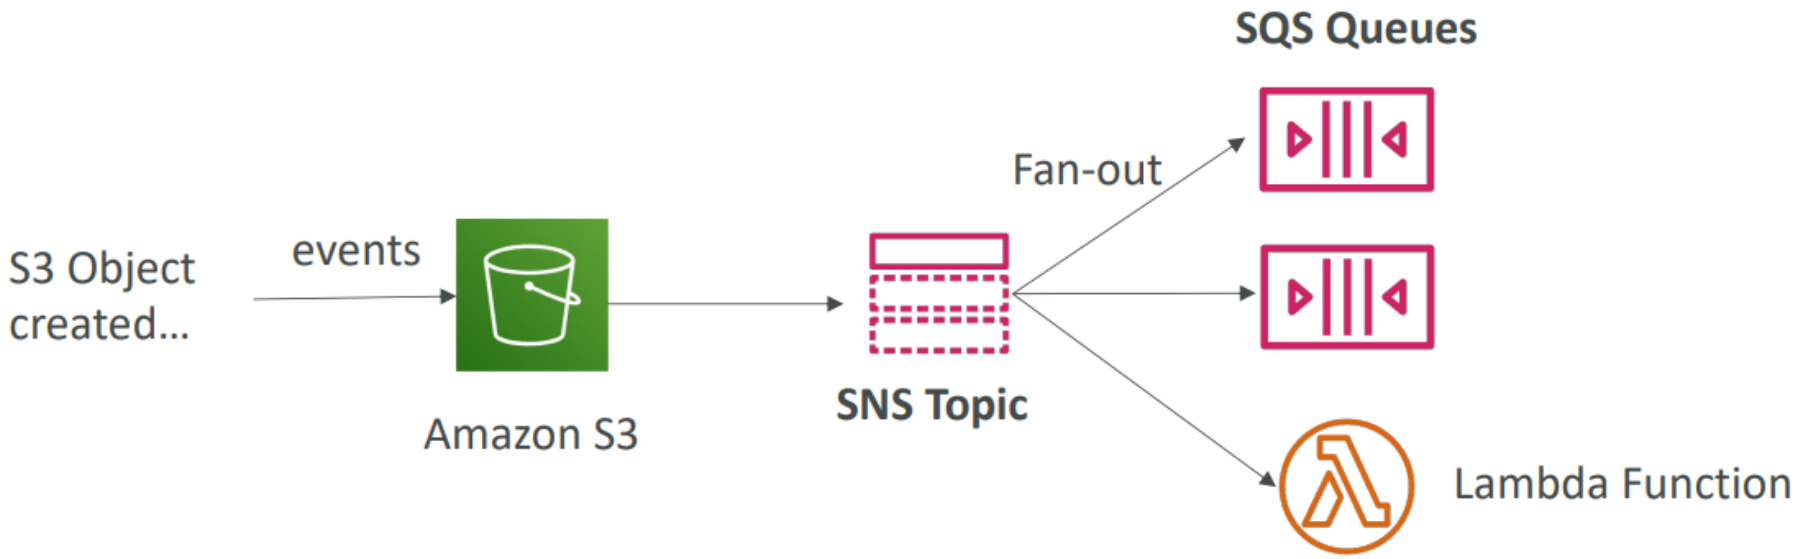

© Copyright by Cloud Mentor Pro | Confidential

Application: SNS to Amazon S3 through Kinesis Data Firehose

- SNS can send to Kinesis and therefore we can have the following solutions architecture:

© Copyright by Cloud Mentor Pro | Confidential

Amazon SNS – FIFO Topic

- FIFO = First In First Out (ordering of messages in the topic)

- Similar features as SQS FIFO:
  - **Ordering** by Message Group ID (all messages in the same group are ordered)
  - **Deduplication** using a Deduplication ID or Content Based Deduplication
- **Can have SQS Standard and FIFO queues as subscribers**
- Limited throughput (same throughput as SQS FIFO)

© Copyright by Cloud Mentor Pro | Confidential

SNS FIFO + SQS FIFO: Fan Out

- In case you need fan out + ordering + deduplication

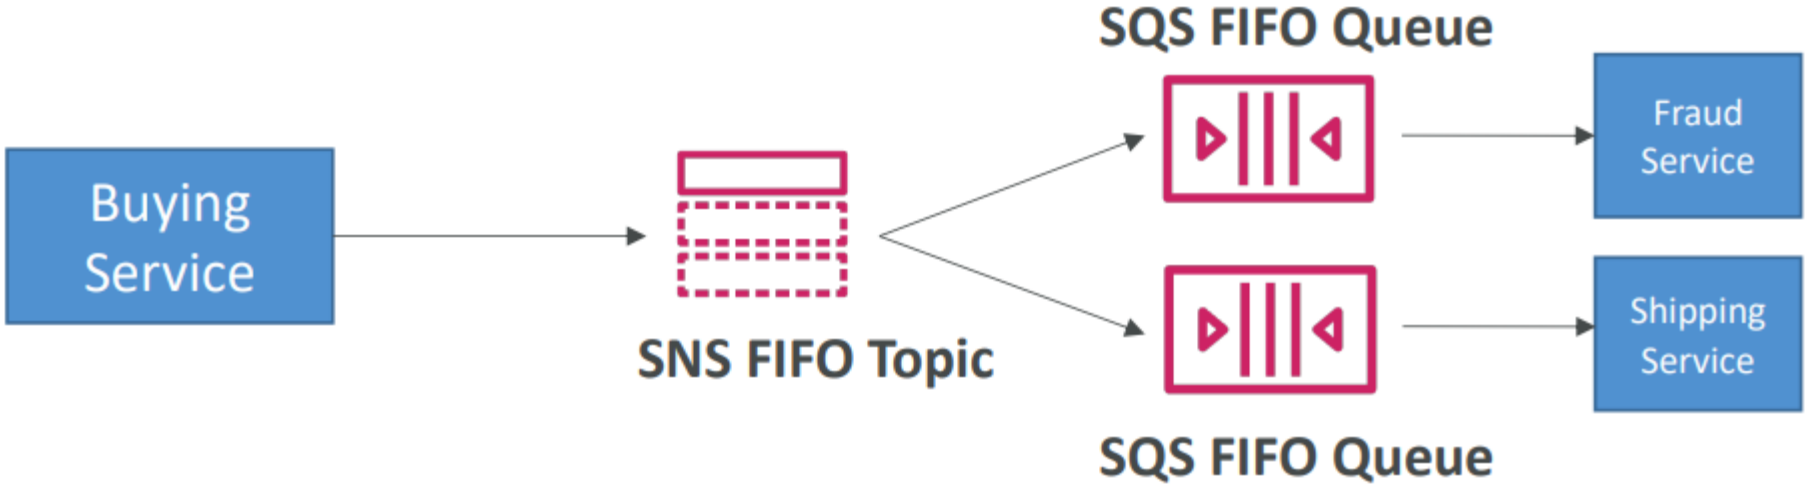

© Copyright by Cloud Mentor Pro | Confidential

SNS – Message Filtering

- JSON policy used to filter messages sent to SNS topic's subscriptions
- If a subscription doesn't have a filter policy, it receives every message

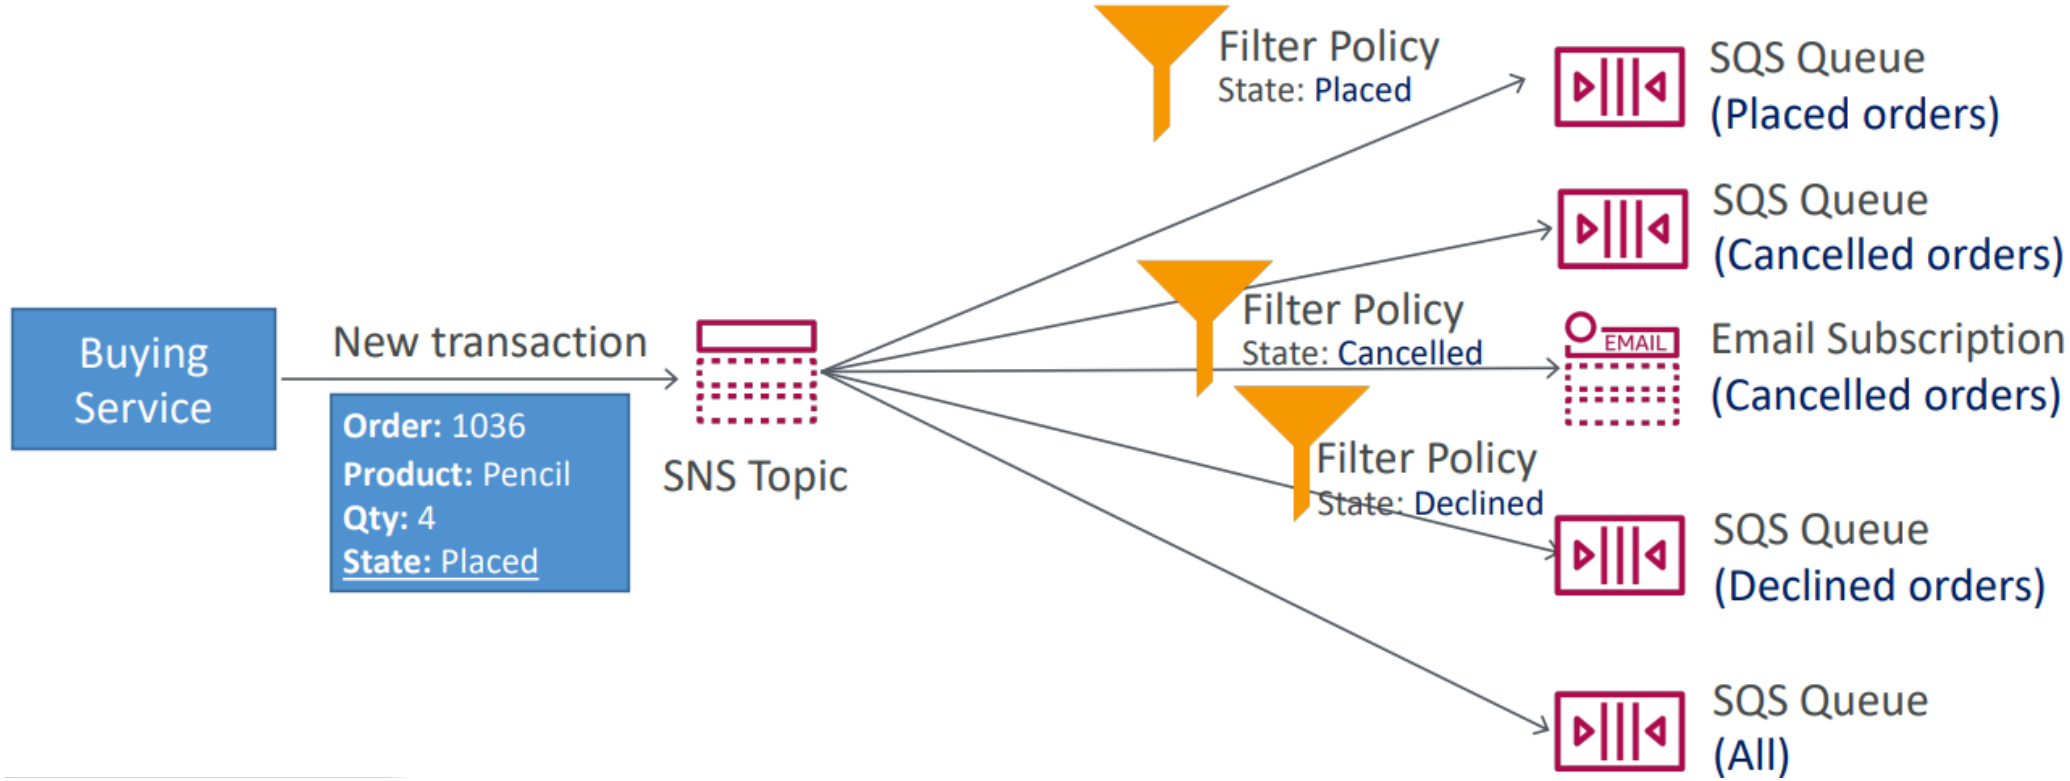

© Copyright by Cloud Mentor Pro | Confidential

Kinesis Data Streams

- Collect and store streaming data in real-time

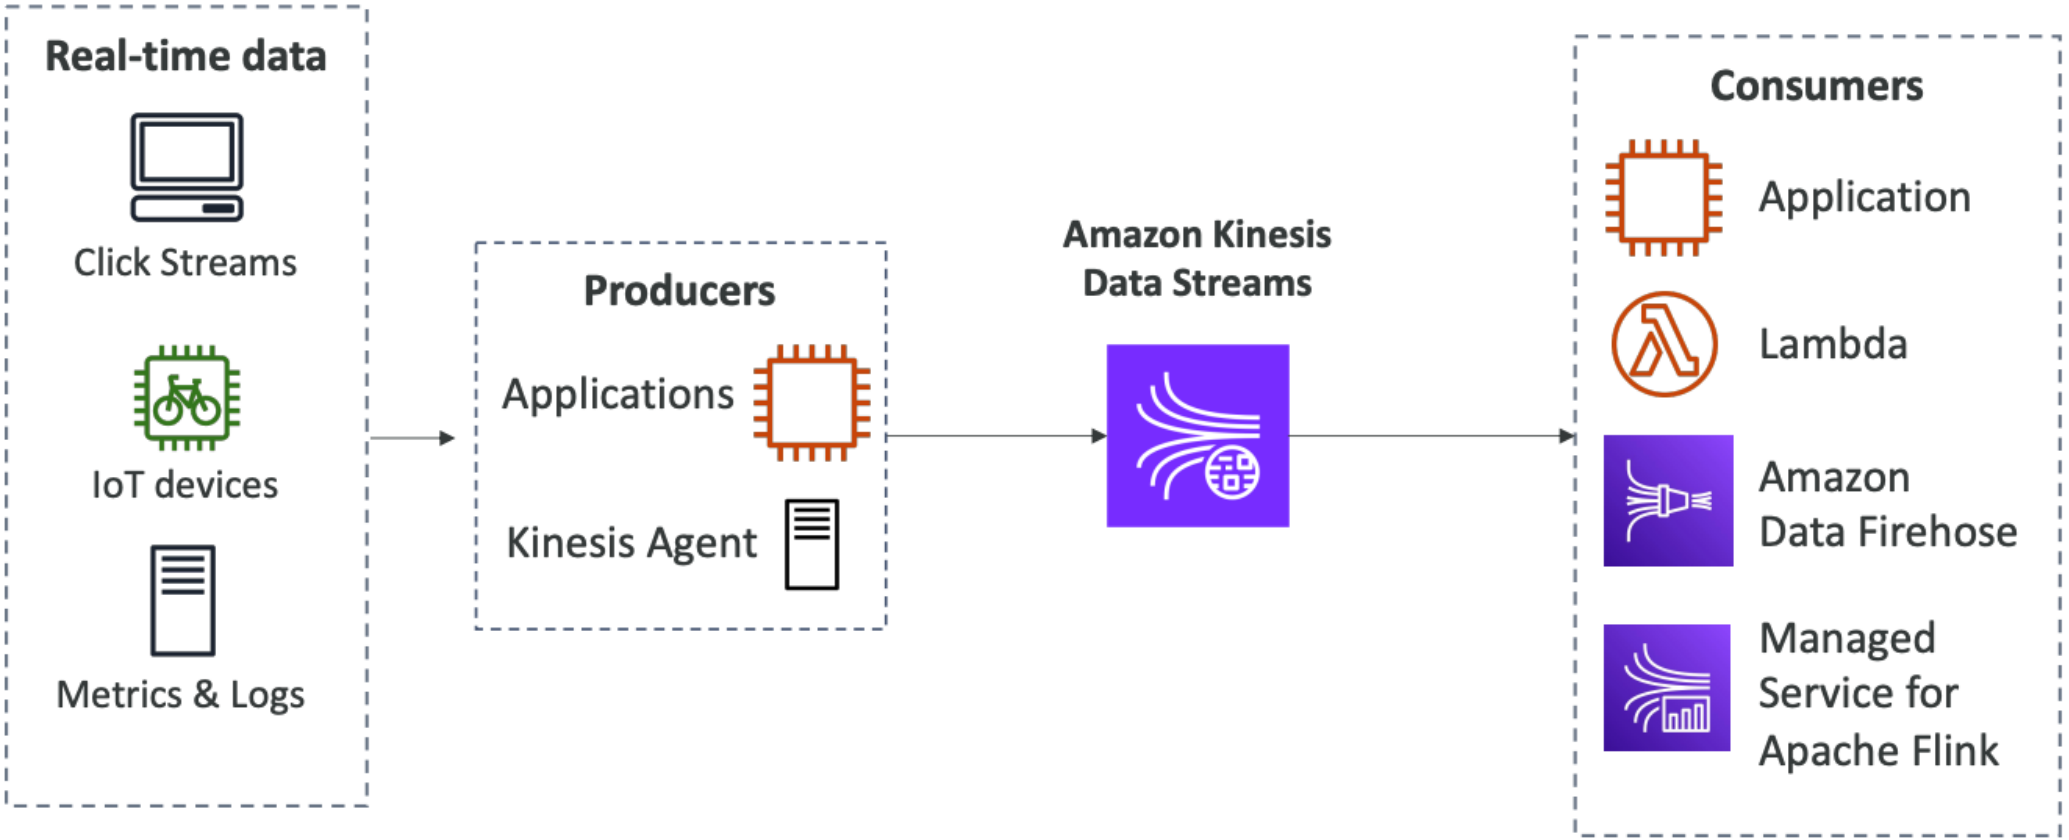

© Copyright by Cloud Mentor Pro | Confidential

Kinesis Data Streams

- Retention between 1 day to 365 days
- Ability to reprocess (replay) data data by consumers
- Data can’t be deleted from Kinesis (until it expires)
- Data up to 1MB (typical use case is lot of “small” real-time data)
- Data ordering guarantee for data with the same “Partition ID”
- At-rest KMS encryption, in-flight HTTPS encryption
- Kinesis Producer Library (KPL) to write an optimized producer application
- Kinesis Client Library (KCL) to write an optimized consumer application

© Copyright by Cloud Mentor Pro | Confidential

Kinesis Data Streams – Capacity Modes

- **Provisioned mode:**
  - Choose number of shards
  - Each shard gets 1MB/s in (or 1000 records per second)
  - Each shard gets 2MB/s out
  - Scale manually to increase or decrease the number of shards
  - You pay per shard provisioned per hour
- **On-demand mode:**
  - No need to provision or manage the capacity
  - Default capacity provisioned (4 MB/s in or 4000 records per second)
  - Scales automatically based on observed throughput peak during the last 30 days
  - Pay per stream per hour & data in/out per GB

© Copyright by Cloud Mentor Pro | Confidential

Amazon Data Firehose

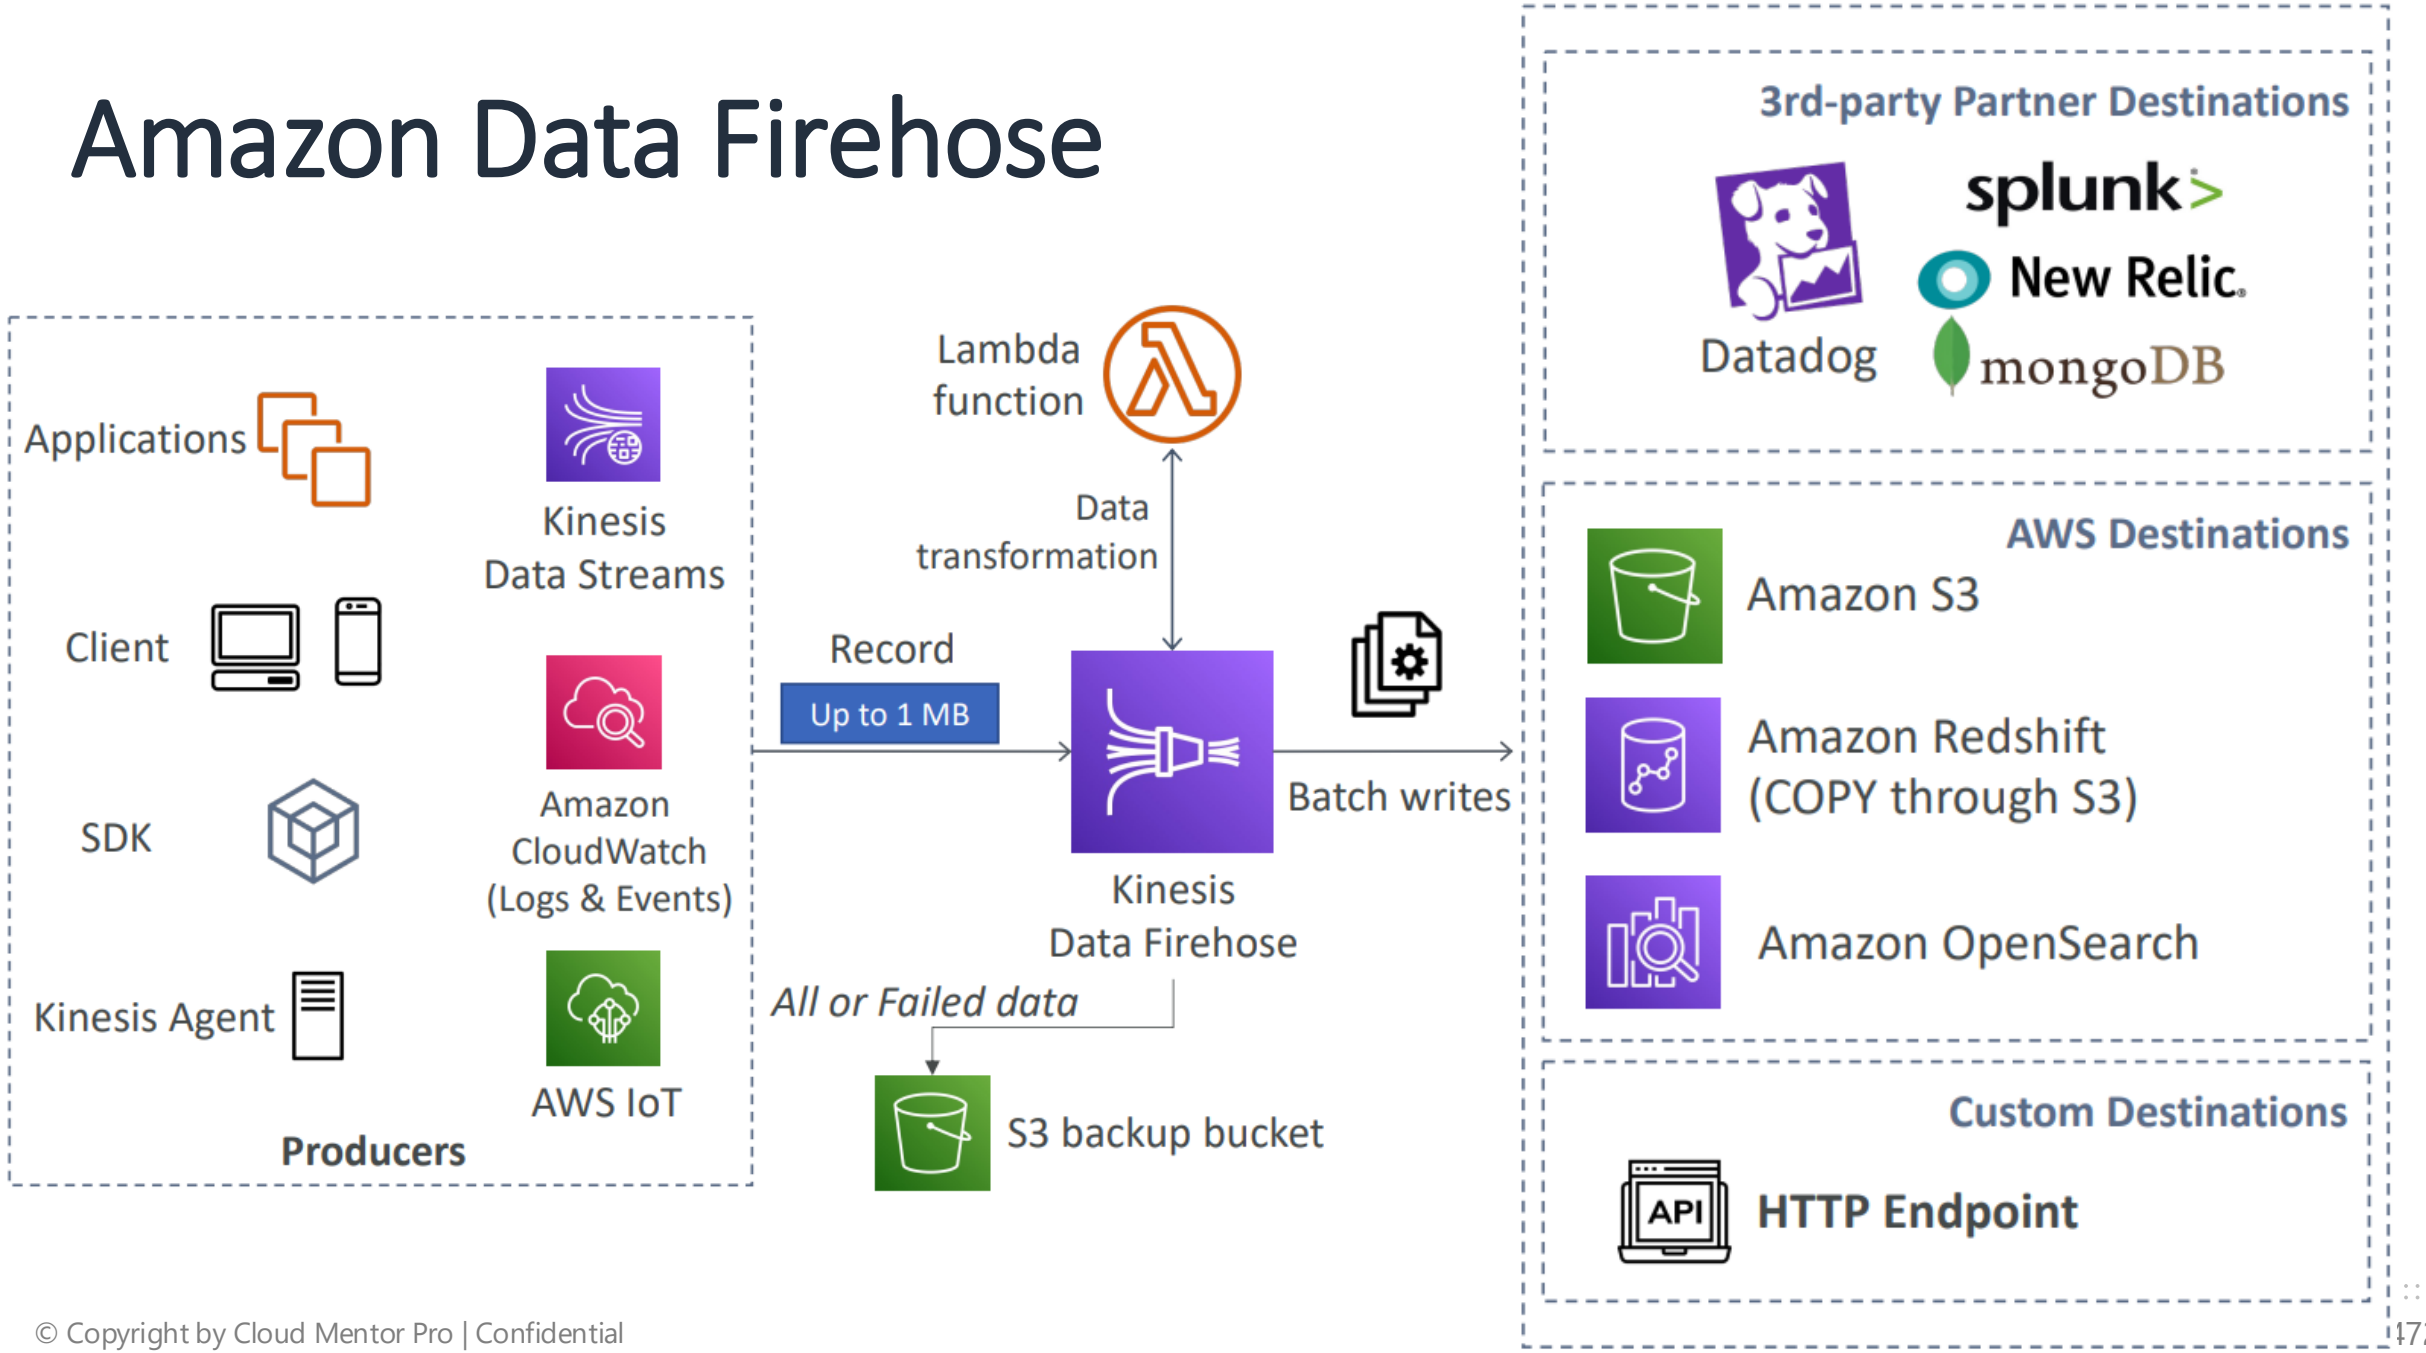

© Copyright by Cloud Mentor Pro | Confidential

Amazon Data Firehose

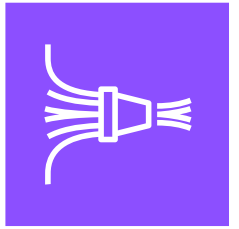

- *Note: used to be called "Kinesis Data Firehose"*
- Fully Managed Service
  - Amazon Redshift / Amazon S3 / Amazon OpenSearch Service
  - 3rd party: Splunk / MongoDB / Datadog / NewRelic / …
  - Custom HTTP
- Automatic scaling, serverless, pay for what you use
- Near Real-Time with buffering capability based on size / time
- Supports CSV, JSON, Parquet, Avro, Raw Text, Binary data
- Conversions to Parquet / ORC, compressions with gzip / snappy
- Custom data transformations using AWS Lambda (ex: CSV to JSON)

© Copyright by Cloud Mentor Pro | Confidential

Kinesis Data Streams vs Amazon Data Firehose

Kinesis Data Streams

- Streaming data collection
- Producer & Consumer code
- Real-time
- Provisioned / On-Demand mode
- Data storage up to 365 days
- Replay Capability

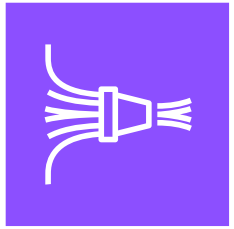

Amazon Data Firehose

- Load streaming data into S3 / Redshift / OpenSearch / 3rd party / custom HTTP
- Fully managed
- Near real-time
- Automatic scaling
- No data storage
- Doesn’t support replay capability

© Copyright by Cloud Mentor Pro | Confidential

Amazon Managed Service for Apache Flink

- *Previously named: Kinesis Data Analytics for Apache Flink*
- Flink (Java, Scala or SQL) is a framework for processing data streams

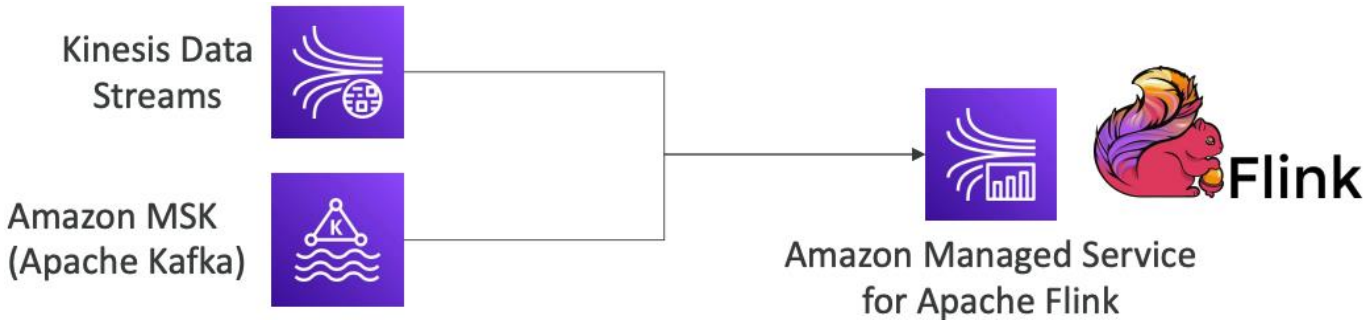

- Run any Apache Flink application on a managed cluster on AWS
  - Provisioned compute resources, parallel computation, automatic scaling
  - Application backups (implemented as checkpoints and snapshots)
  - Use any Apache Flink programming features to transform data
  - Important: Flink does not read from Amazon Data Firehose

© Copyright by Cloud Mentor Pro | Confidential

SQS vs SNS vs Kinesis

SQS:

- Consumer "pull data"
- Data is deleted after being consumed
- Can have as many workers (consumers) as we want
- No need to provision throughput
- Ordering guarantees only on FIFO queues
- Individual message delay capability

SNS:

- Push data to many subscribers
- Up to 12,500,000 subscribers
- Data is not persisted (lost if not delivered)
- Pub/Sub
- Up to 100,000 topics
- No need to provision throughput
- Integrates with SQS for fan-out architecture pattern
- FIFO capability for SQS FIFO

Kinesis:

- Standard: pull data
  - 2 MB per shard
- Enhanced-fan out: push data
  - 2 MB per shard per consumer
- Possibility to replay data
- Meant for real-time big data, analytics and ETL
- Ordering at the shard level
- Data expires after X days
- Provisioned mode or on-demand capacity mode

© Copyright by Cloud Mentor Pro | Confidential

Exercise

Creating and Subscribing to SNS Topics, Adding SNS event for S3 bucket

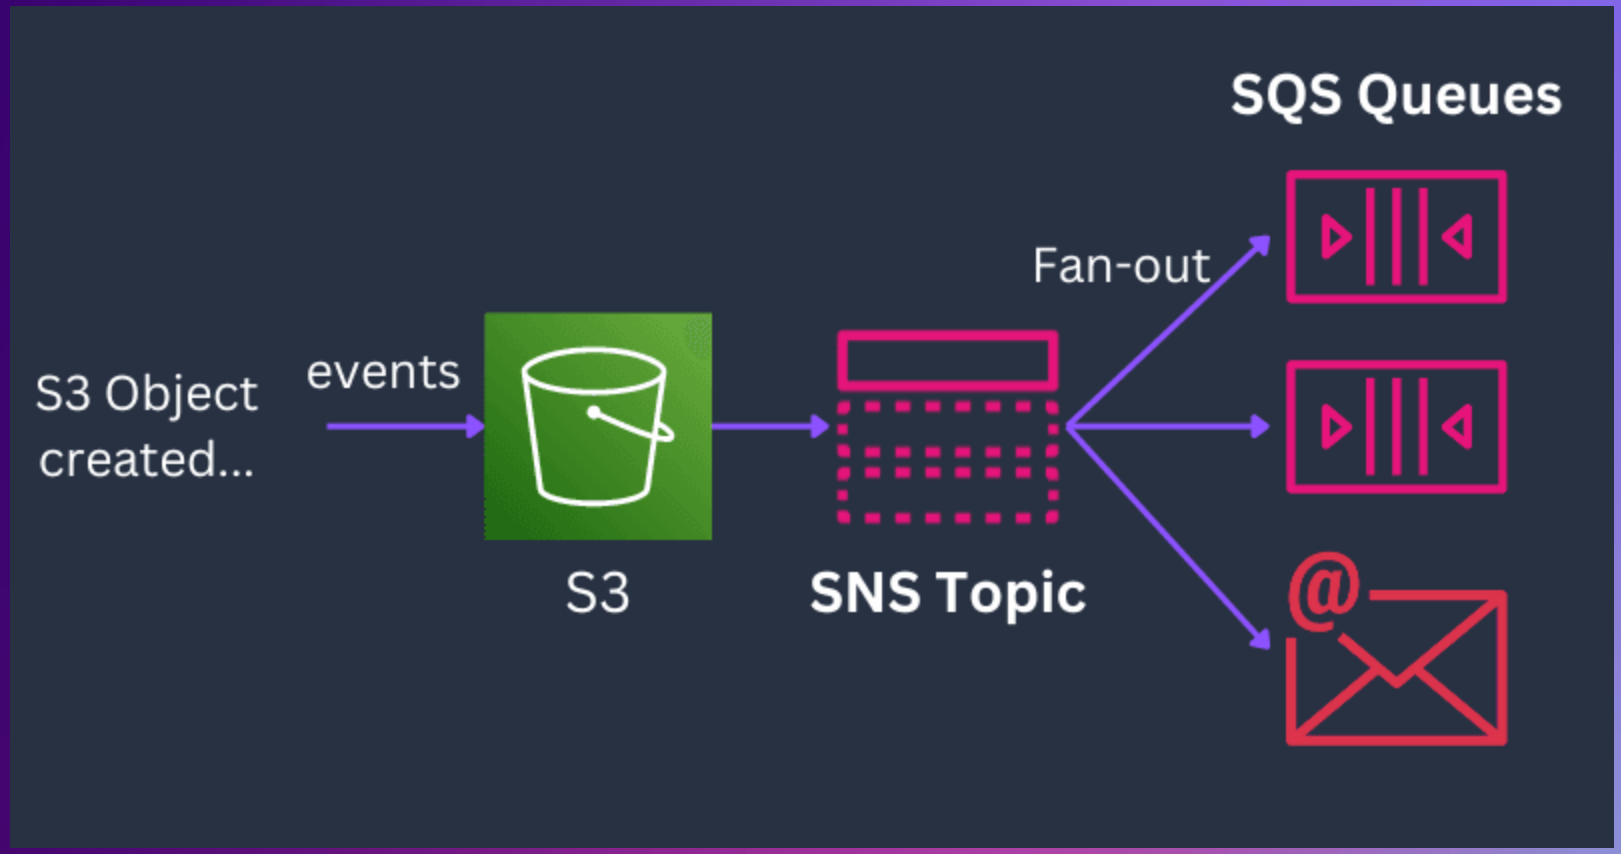

© Copyright by Cloud Mentor Pro | Confidential

Section 10

- AWS Monitoring, Troubleshooting & Audit

© Copyright by Cloud Mentor Pro | Confidential

AWS Monitoring, Troubleshooting & Audit

CloudWatch, X-Ray, Cloudtrail

© Copyright by Cloud Mentor Pro | Confidential

Monitoring in AWS

- AWS CloudWatch:
  - Metrics: Collect and track key metrics
  - Logs: Collect, monitor, analyze and store log files
  - Events: Send notifications when certain events happen in your AWS
  - Alarms: React in real-time to metrics / events
- AWS X-Ray:
  - Troubleshooting application performance and errors
  - Distributed tracing of microservices
- AWS CloudTrail:
  - Internal monitoring of API calls being made
  - Audit changes to AWS Resources by your users

© Copyright by Cloud Mentor Pro | Confidential

Amazon CloudWatch Metrics

- CloudWatch provides metrics for every services in AWS
- Metric is a variable to monitor (CPUUtilization, NetworkIn…)
- Metrics belong to namespaces. This namespaces should be equivalent to service name (ex: EC2, RDS, etc…)
- Dimension is an attribute of a metric (instance id, environment, etc…).
- Up to 30 dimensions per metric
- Metrics have timestamps
- Can create CloudWatch dashboards of metrics

© Copyright by Cloud Mentor Pro | Confidential

Amazon CloudWatch Metrics

- Amazon CloudWatch supports two types of metrics
  - Standard (supported by EC2 and most other services)
  - Custom (typically used for EC2 and on-premises instances)
- With EC2 standard metrics, you can monitor
  - Overall CPU utilization (usually for scaling)
  - Disk Performance (IOPS, Read/Write bytes per second)
  - And much more… but you CAN'T monitor **memory usage (RAM)** and **disk used percentage**
- With EC2 custom metrics, you can monitor:
  - **Memory usage & Disk used percentage**
  - **Application metric** (running status, number of active user, error count, etc…)
  - Anything else defined in the configuration file

© Copyright by Cloud Mentor Pro | Confidential

Amazon CloudWatch Metrics

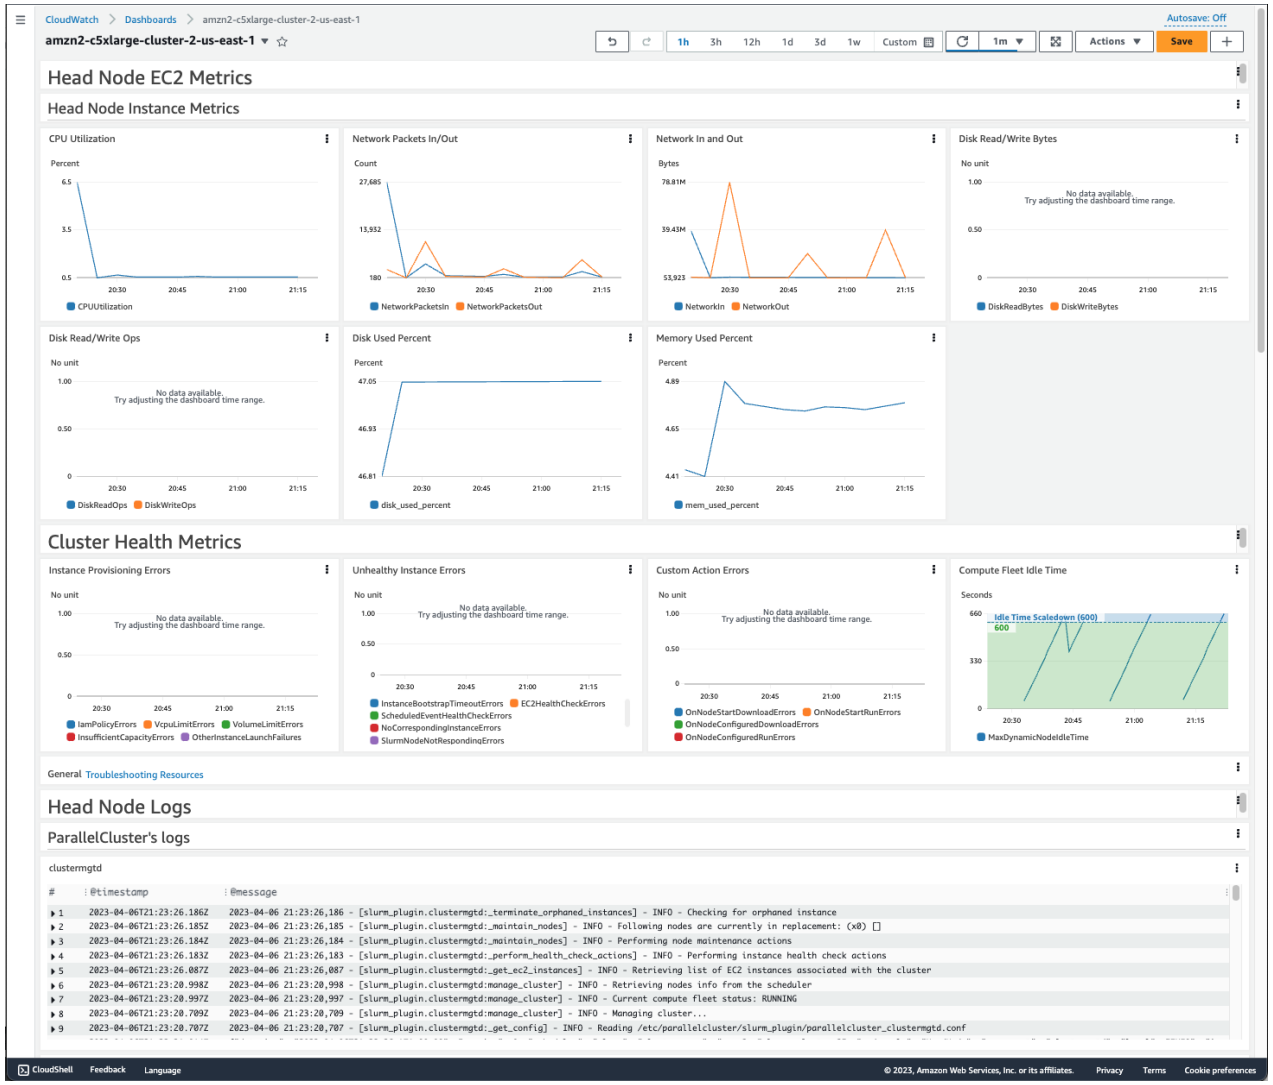

Cloudwatch Dashboard example

© Copyright by Cloud Mentor Pro | Confidential

EC2 Detailed monitoring

- EC2 export STANDARD metrics every 5 minutes by default.
- With detailed monitoring, export frequency becomes 1 minute!
- Use detailed monitoring if you want to scale faster for your ASG!

- The AWS Free Tier allows us to have 10 detailed monitoring metrics

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Custom Metrics

- You can define your own metric and send it to CloudWatch
- Two ways to create CUSTOM metrics
  - PutMetricData API (using AWS SDK, CLI or calling AWS API directly)
  - Cloudwatch Agent (recommended way for EC2 instance)

- You can add custom dimensions to metrics, such as
  - Environment name
  - System name

- Metric resolution (StorageResolution APl parameter - two possible value):
  - Standard: 1 minute (60 seconds)
  - High Resolution: 1/5/10/30 second(s) - Higher cost

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs

- A service that allows you to collect, monitor, and store log files from AWS resources and on-premise systems.
- It helps track application and system performance, troubleshoot issues, and retain logs for auditing and compliance.

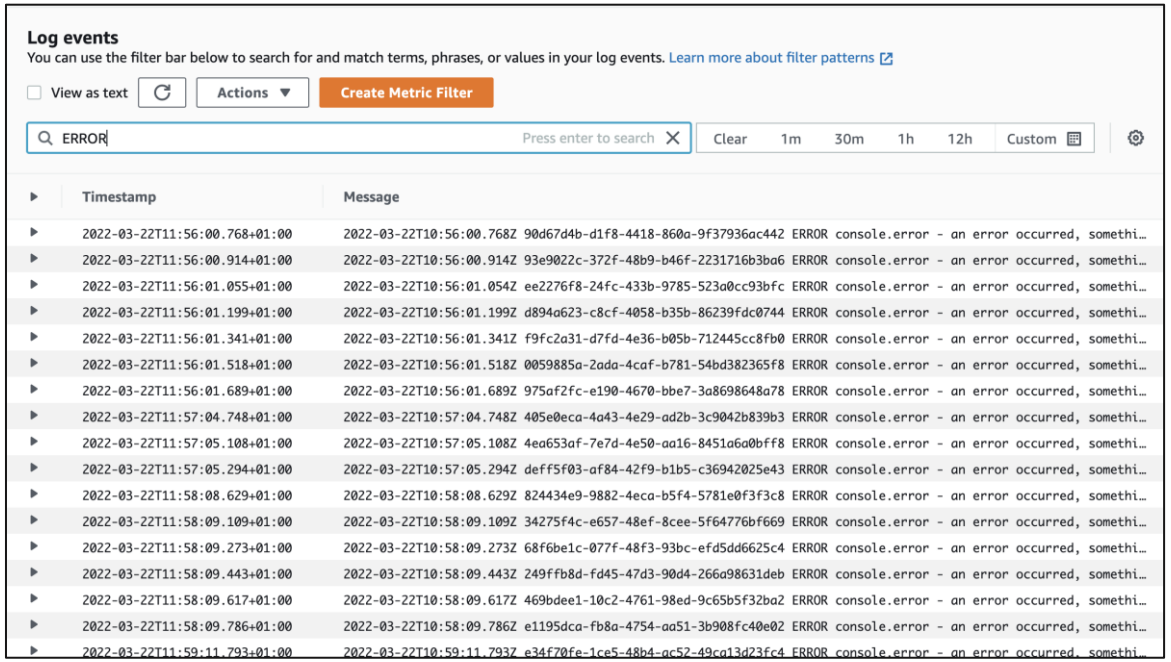

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs - Concepts

- **Log groups**: container for log data. It contains all the related logs from a particular service or application
- **Log stream**: individual sequences of log events inside a log group.
- Can define log expiration policies (never expire, 1 day to 10 years…)
- **CloudWatch Logs can send logs to:**
  - Amazon S3 (exports)
  - Kinesis Data Streams
  - Kinesis Data Firehose
  - AWS Lambda
  - OpenSearch
- Logs are encrypted by default
- Can setup KMS-based encryption with your own keys

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs - Sources

- SDK, CloudWatch Logs Agent, CloudWatch Unified Agent
- Elastic Beanstalk: collection of logs from application
- ECS: collection from containers
- AWS Lambda: collection from function logs
- VPC Flow Logs: VPC specific logs
- API Gateway
- CloudTrail based on filter
- Route53: Log DNS queries

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs Insights

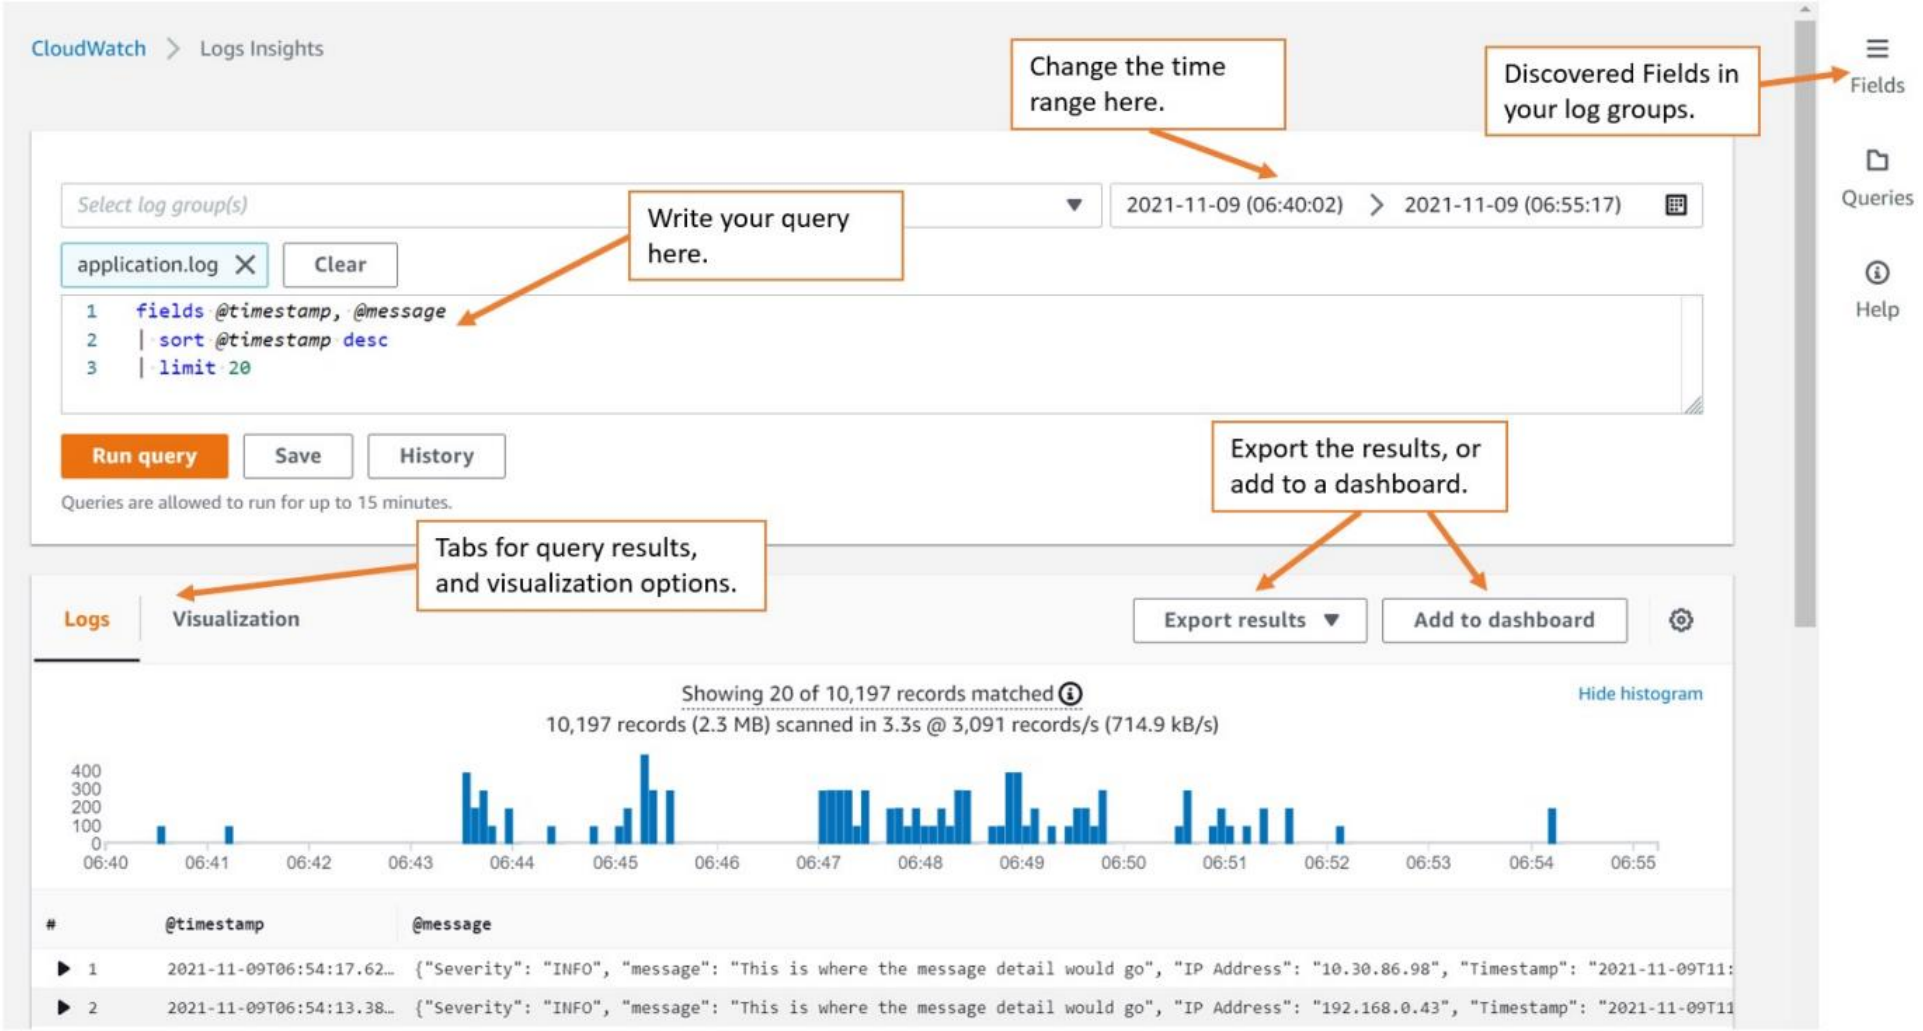

© Copyright by Cloud Mentor Pro | Confidential

https://mng.workshop.aws/operations-2022/detect/cwlogs.html

CloudWatch Logs Insights

- Search and analyze log data stored in CloudWatch Logs
- Example: find a specific IP inside a log, count occurrences of “ERROR” in your logs…
- Provides a purpose -built query language
  - Automatically discovers fields from AWS services and JSON log events
  - Fetch desired event fields, filter based on conditions, calculate aggregate statistics, sort events, limit number of events…
  - Can save queries and add them to CloudWatch Dashboards
- Can query multiple Log Groups in different AWS accounts
- It’s a query engine, not a real-time engine

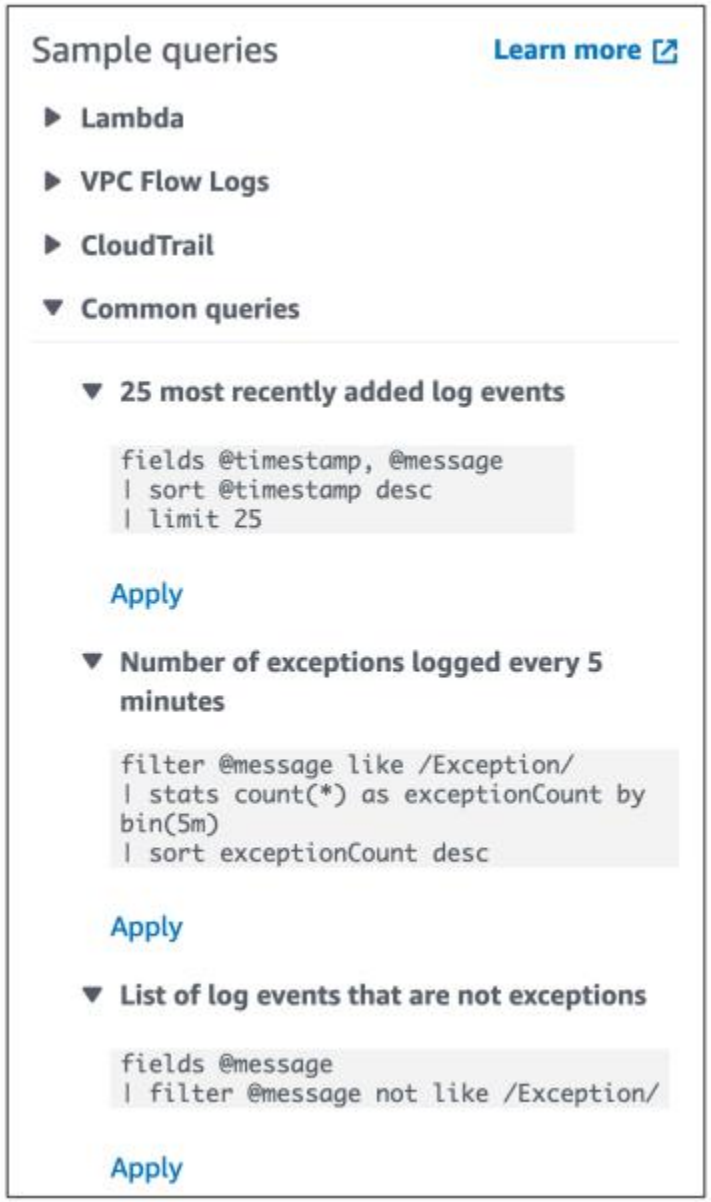

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs – S3 Export

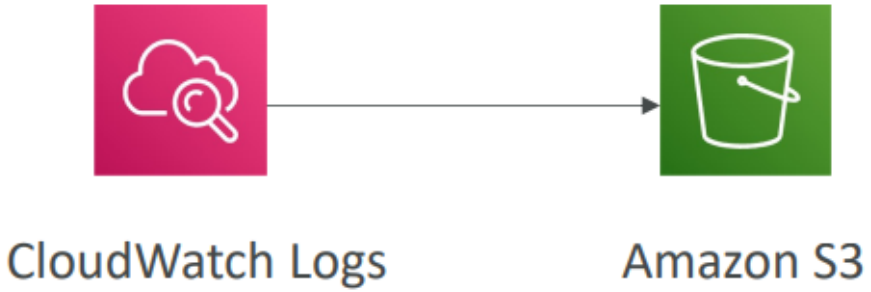

- Log data can take **up to 12 hours** to become available for export
- The API call is **CreateExportTask**

- Not near-real time or real-time… use Logs Subscriptions instead

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs Subscriptions

- Get a real-time log events from CloudWatch Logs for processing and analysis
- Send to Kinesis Data Streams, Kinesis Data Firehose, or Lambda
- Subscription Filter – filter which logs are events delivered to your destination

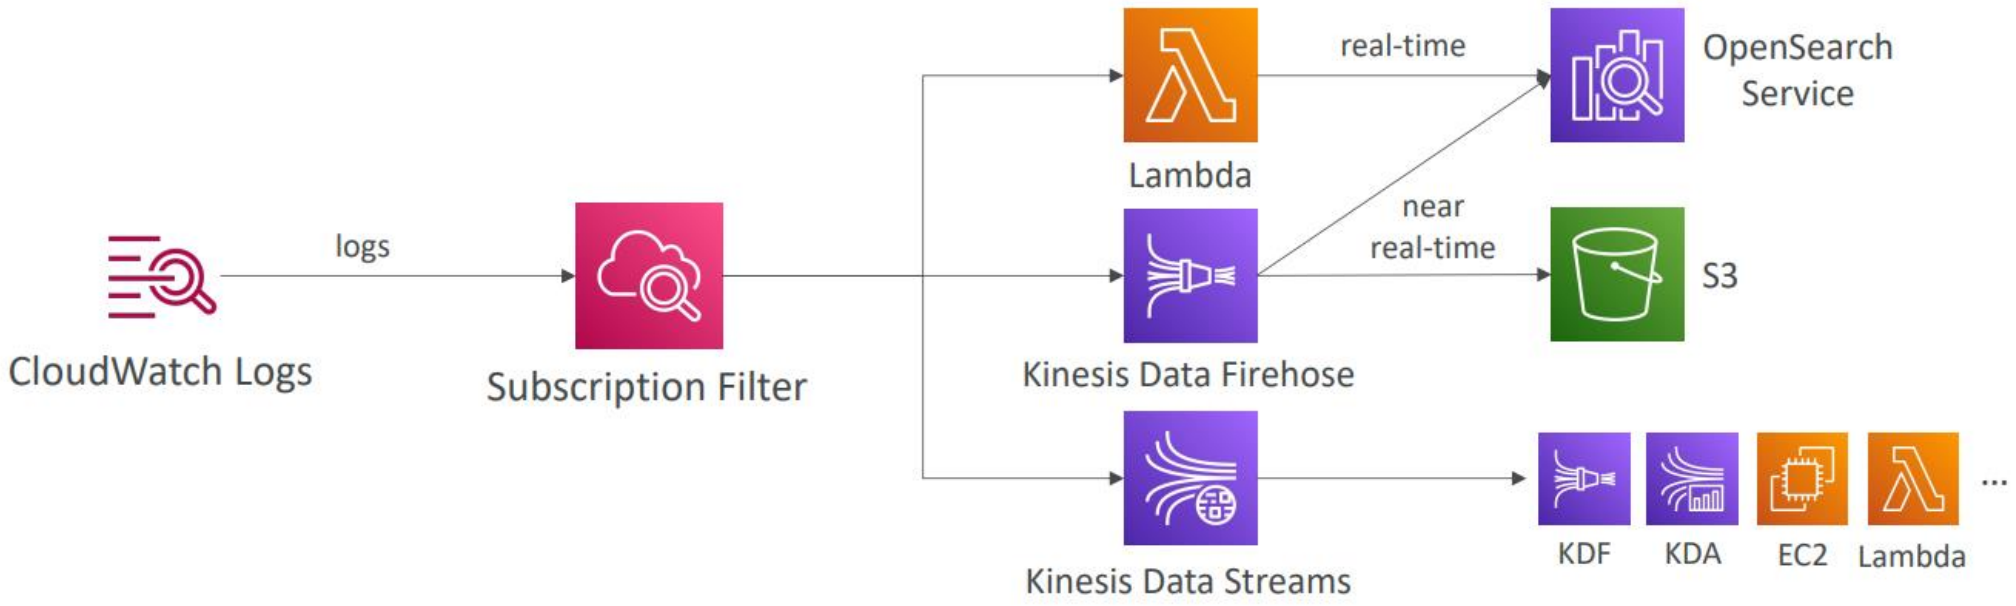

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs Aggregation Multi-Account & Multi Region

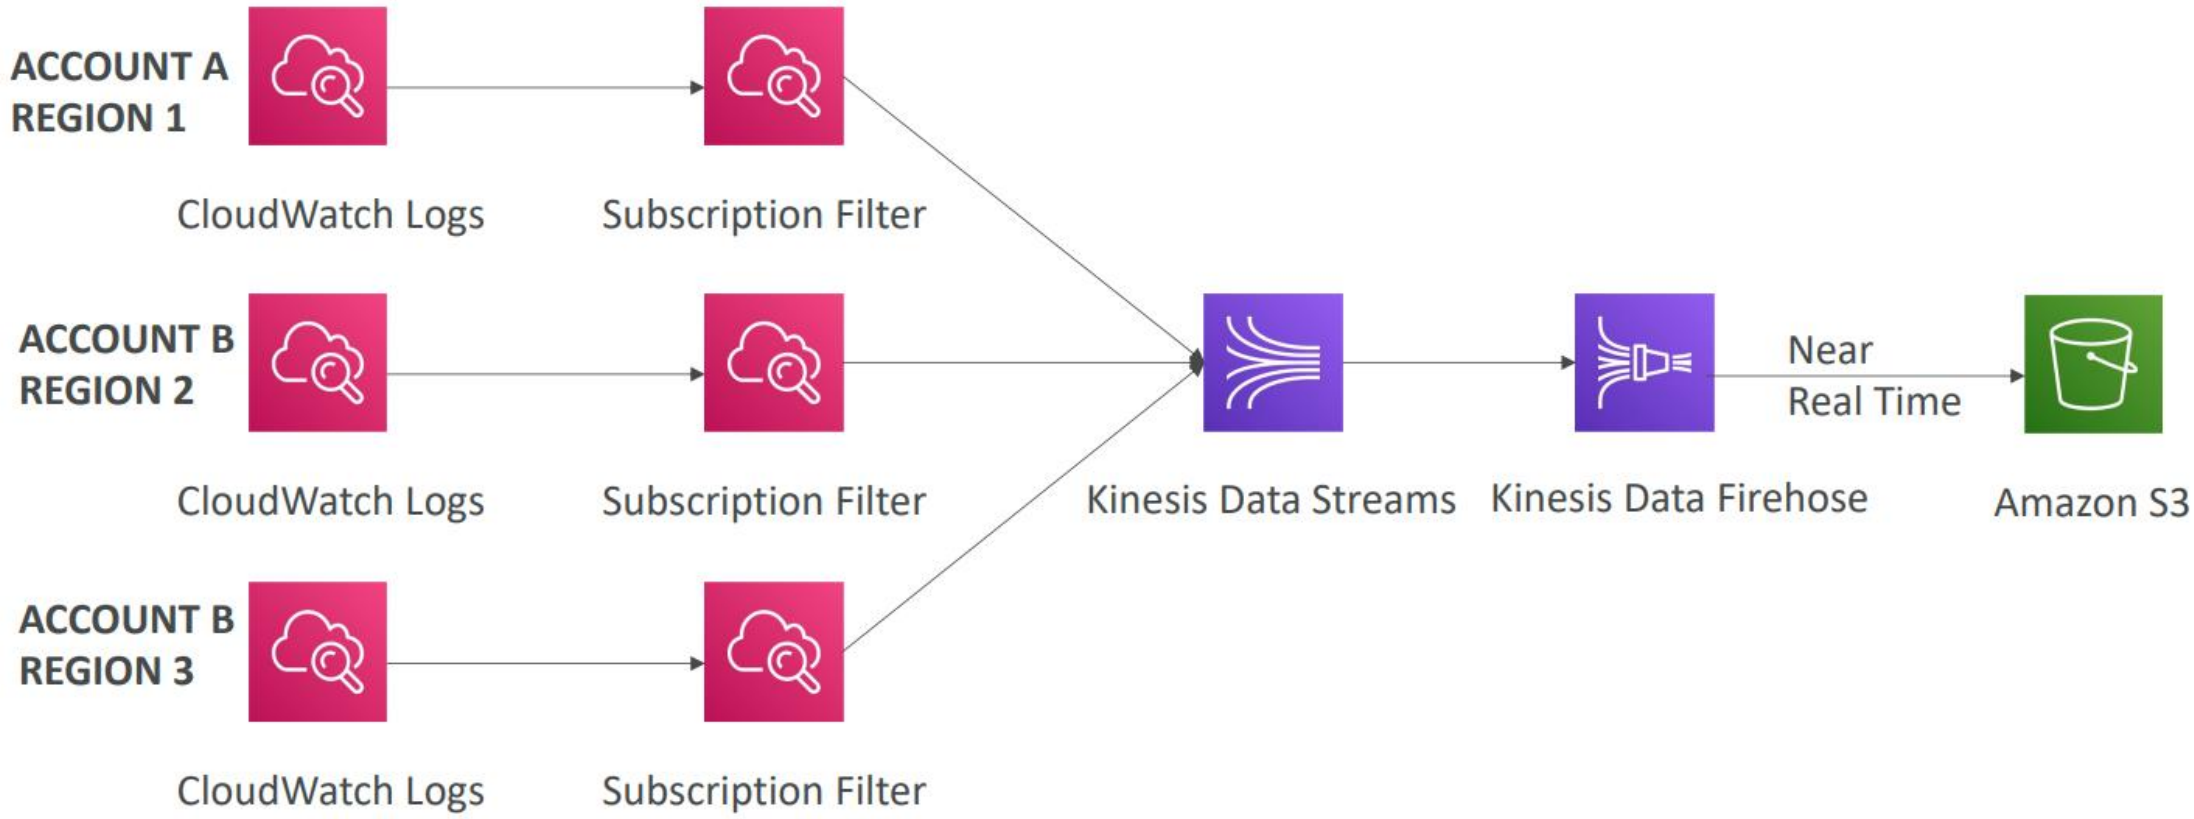

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs Subscriptions

- **Cross-Account Subscription** – send log events to resources in a different AWS account (KDS, KDF)

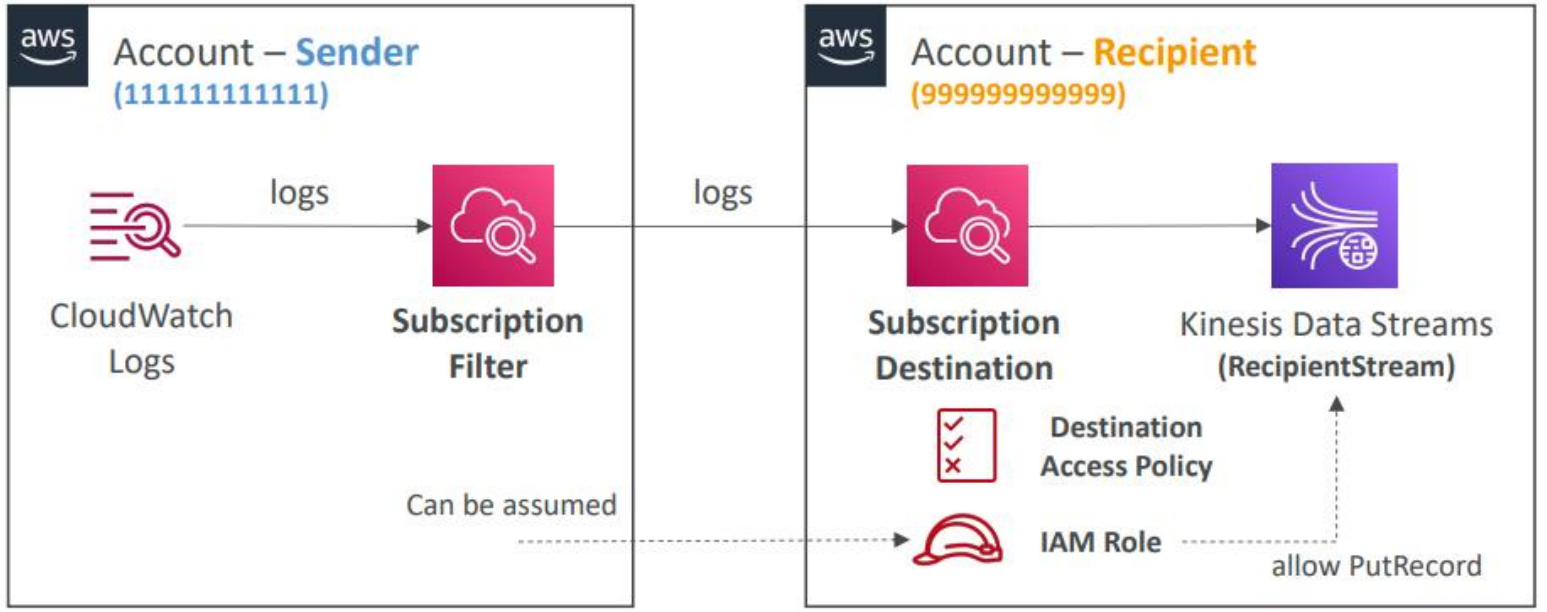

IAM Role (Cross-Account)

```
{
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "kinesis:PutRecord",
      "Resource": "arn:aws:kinesis:us-east-1:
999999999999:stream/RecipientStream"
    }
  ]
}
```

Destination Access Policy

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
        "AWS": "111111111111"
      },
      "Action": "logs:PutSubscriptionFilter",
      "Resource": "arn:aws:logs:us-east-1:999999999999:
destination:testDestination"
    }
  ]
}
```

©

CloudWatch Logs for EC2

- By default, no logs from your EC2 machine will go to CloudWatch
- You need to run a CloudWatch agent on EC2 to push the log files you want
- Make sure IAM permissions are correct
- The CloudWatch agent can be setup on-premises too

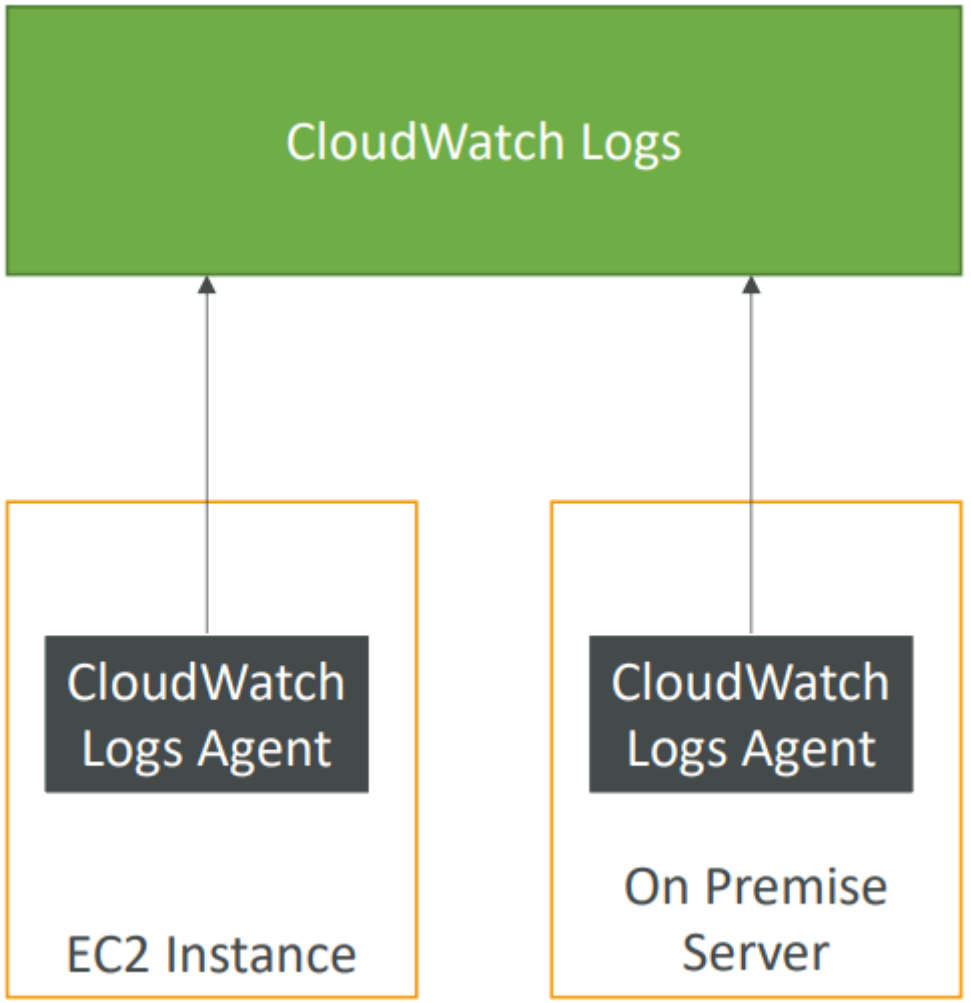

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Agent (Unified Agent)

- A tool that collects & sends CUSTOM METRICS and LOGS from your EC2 or on-premise instance to CloudWatch.
- It helps you monitor both EC2 and on-premises systems more effectively
- Requirement to use CloudWatch Agent:
  - Must be installed and setup correctly
  - IAM Role / Access key with minimum policy CloudWatchAgentServerPolicy
  - Connectivity to CloudWatch (via internet or VPC Endpoints)
  - A valid configuration file (defined metrics and logs to be collected)

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Agent – Metrics

- Collected directly on your Linux server / EC2 instance

- **CPU** (active, guest, idle, system, user, steal)
- **Disk metrics** (free, used, total), Disk IO (writes, reads, bytes, iops)
- **RAM** (free, inactive, used, total, cached)
- **Netstat** (number of TCP and UDP connections, net packets, bytes)
- **Processes** (total, dead, bloqued, idle, running, sleep)
- **Swap Space** (free, used, used %)

- Reminder: out-of-the box metrics for EC2 – disk, CPU, network (high level)

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Logs Metric Filter

- CloudWatch Logs can use filter expressions
  - For example, find a specific IP inside of a log
  - Or count occurrences of"ERROR" in your logs
  - Metric filters can be used to trigger alarms
- **Filters do not retroactively filter data. Filters only publish the metric data points for events that happen after the filter was created.**
- Ability to specify up to 3 Dimensions for the Metric Filter (optional)

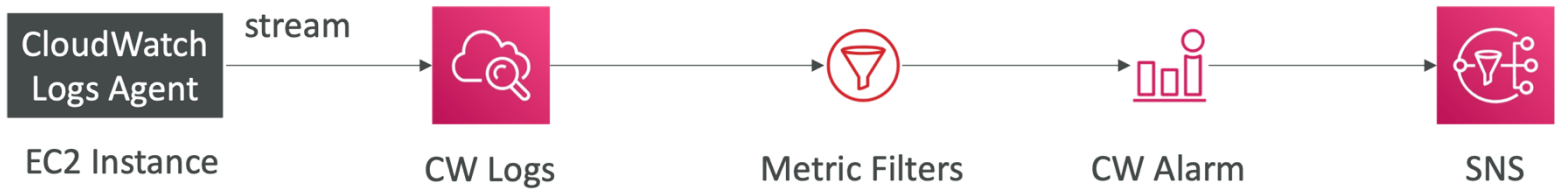

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Alarms

- Alarms are used to trigger notifications for any metric
- Various options (sampling, %, max, min, etc…)
- Alarm States:
  - OK
  - INSUFFICIENT_DATA
  - ALARM
- Period:
  - Length of time in seconds to evaluate the metric
  - High resolution custom metrics: 10 sec, 30 sec or multiples of 60 sec

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Alarm Targets

- Stop, Terminate, Reboot, or Recover an EC2 Instance
- Trigger Auto Scaling Action
- Send notification to SNS (from which you can do pretty much anything)

Amazon EC2

EC2 Auto Scaling

Amazon SNS

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Alarms – Composite Alarms

- CloudWatch Alarms are on a single metric
- Composite Alarms are monitoring the states of multiple other alarms
- AND and OR conditions
- Helpful to reduce “alarm noise” by creating complex composite alarms

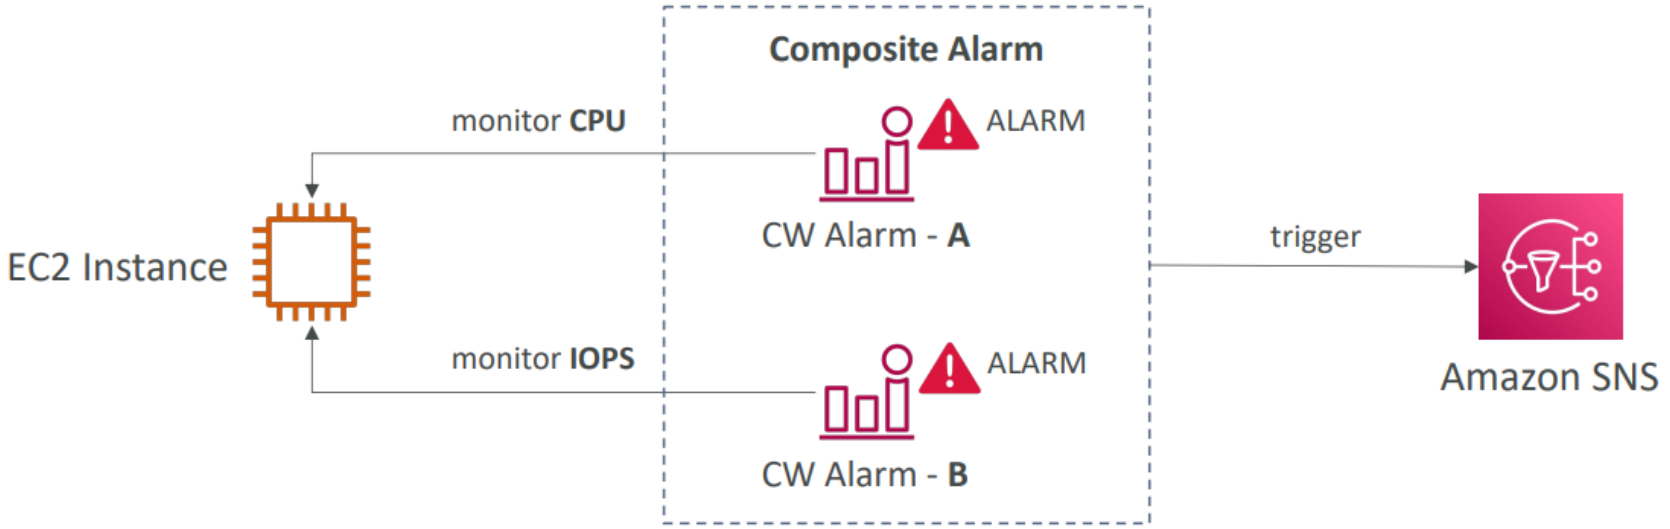

© Copyright by Cloud Mentor Pro | Confidential

EC2 Instance Recovery

- CloudWatch alarm has **Recovery action** which helps restoring your ec2
- <u>Status Check:</u>
  - Instance status = check the EC2 VM
  - System status = check the underlying hardware

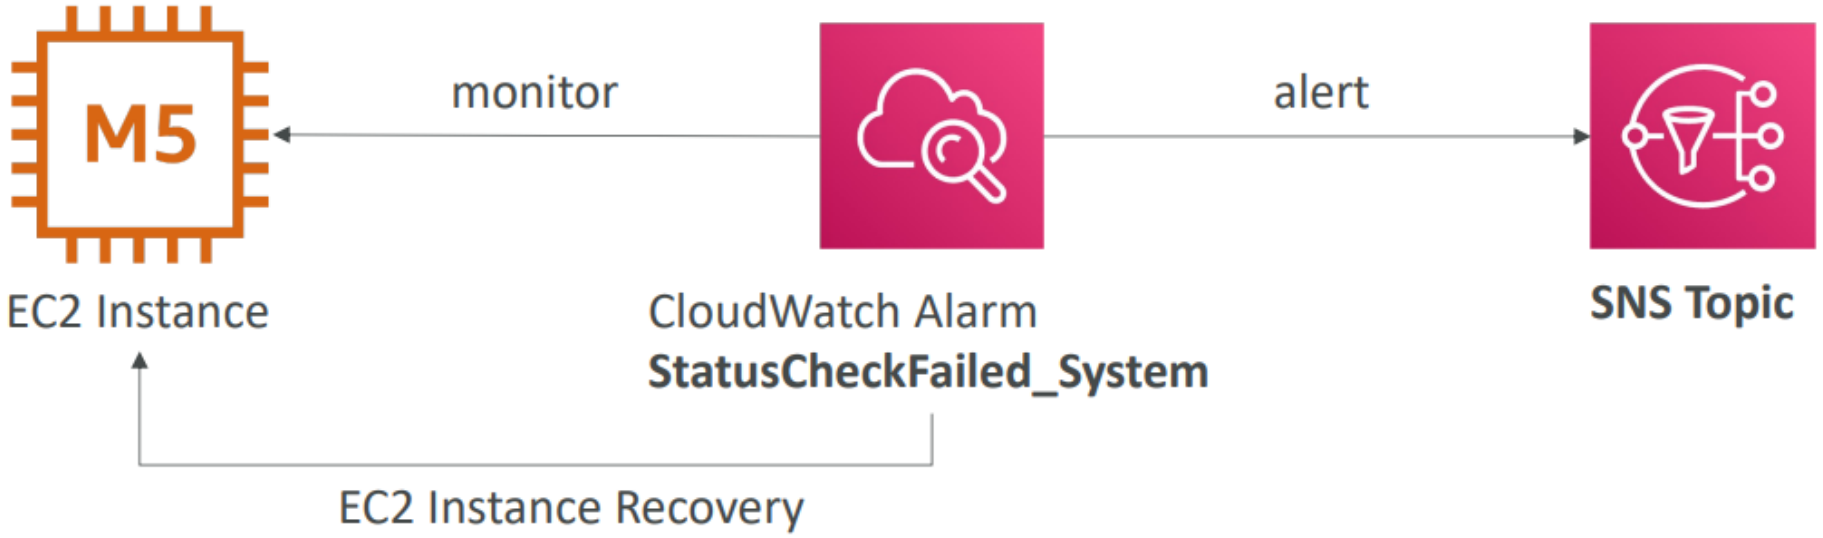

- **Recovery**: Same Private, Public, Elastic IP, metadata, placement group

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Alarm: good to know

- Alarms can be created based on CloudWatch Logs Metrics Filters

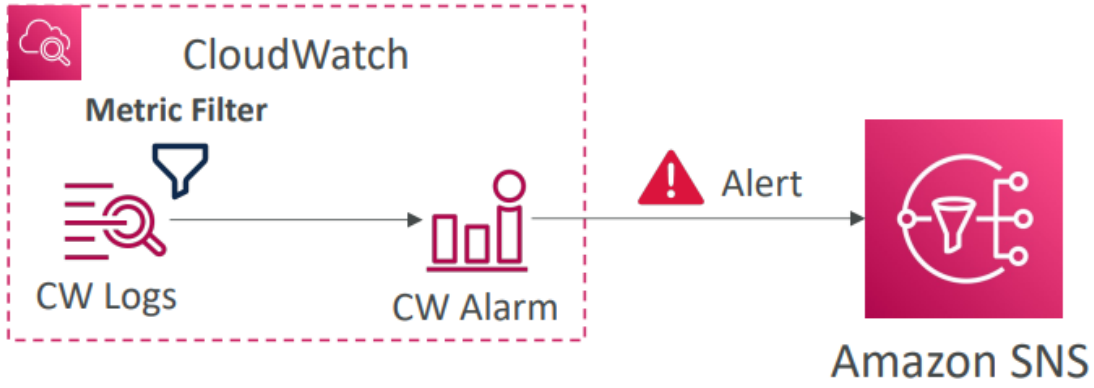

- To test alarms and notifications, set the alarm state to Alarm using CLI

aws cloudwatch set-alarm-state --alarm-name "myalarm" --state-value ALARM --state-reason "testing purposes"

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Synthetics Canary

- Configurable script that monitor your APIs, URLs, Websites, ...
- Reproduce what your customers do programmatically to find issues before customers are impacted
- Checks the availability and latency of your endpoints and can store load time data and screenshots of the UI
- Integration with CloudWatch Alarms
- Scripts written in Node.js or Python
- Programmatic access to a headless Google Chrome browser
- Can run once or on a regular schedule

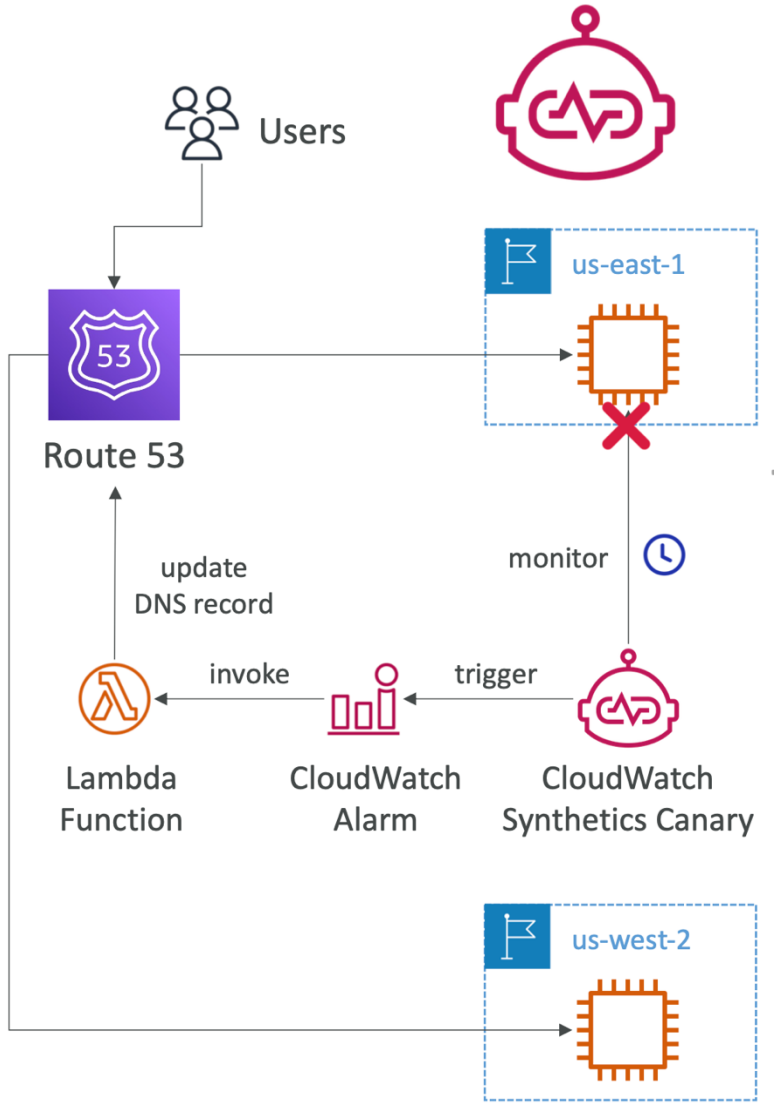

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge (formerly CloudWatch Events)

- Schedule: Cron jobs (scheduled scripts)

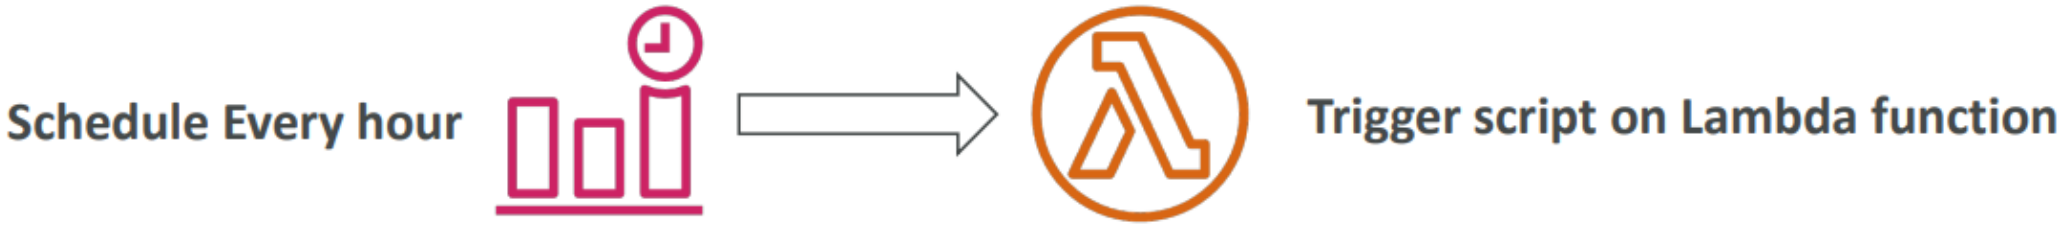

- Event Pattern: Event rules to react to a service doing something

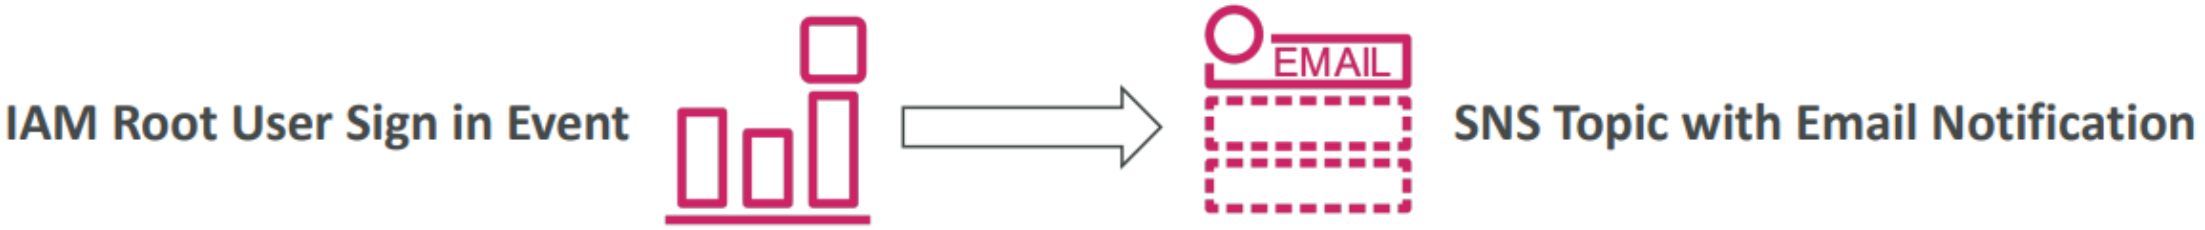

- Trigger Lambda functions, send SQS/SNS messages…

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge Rules

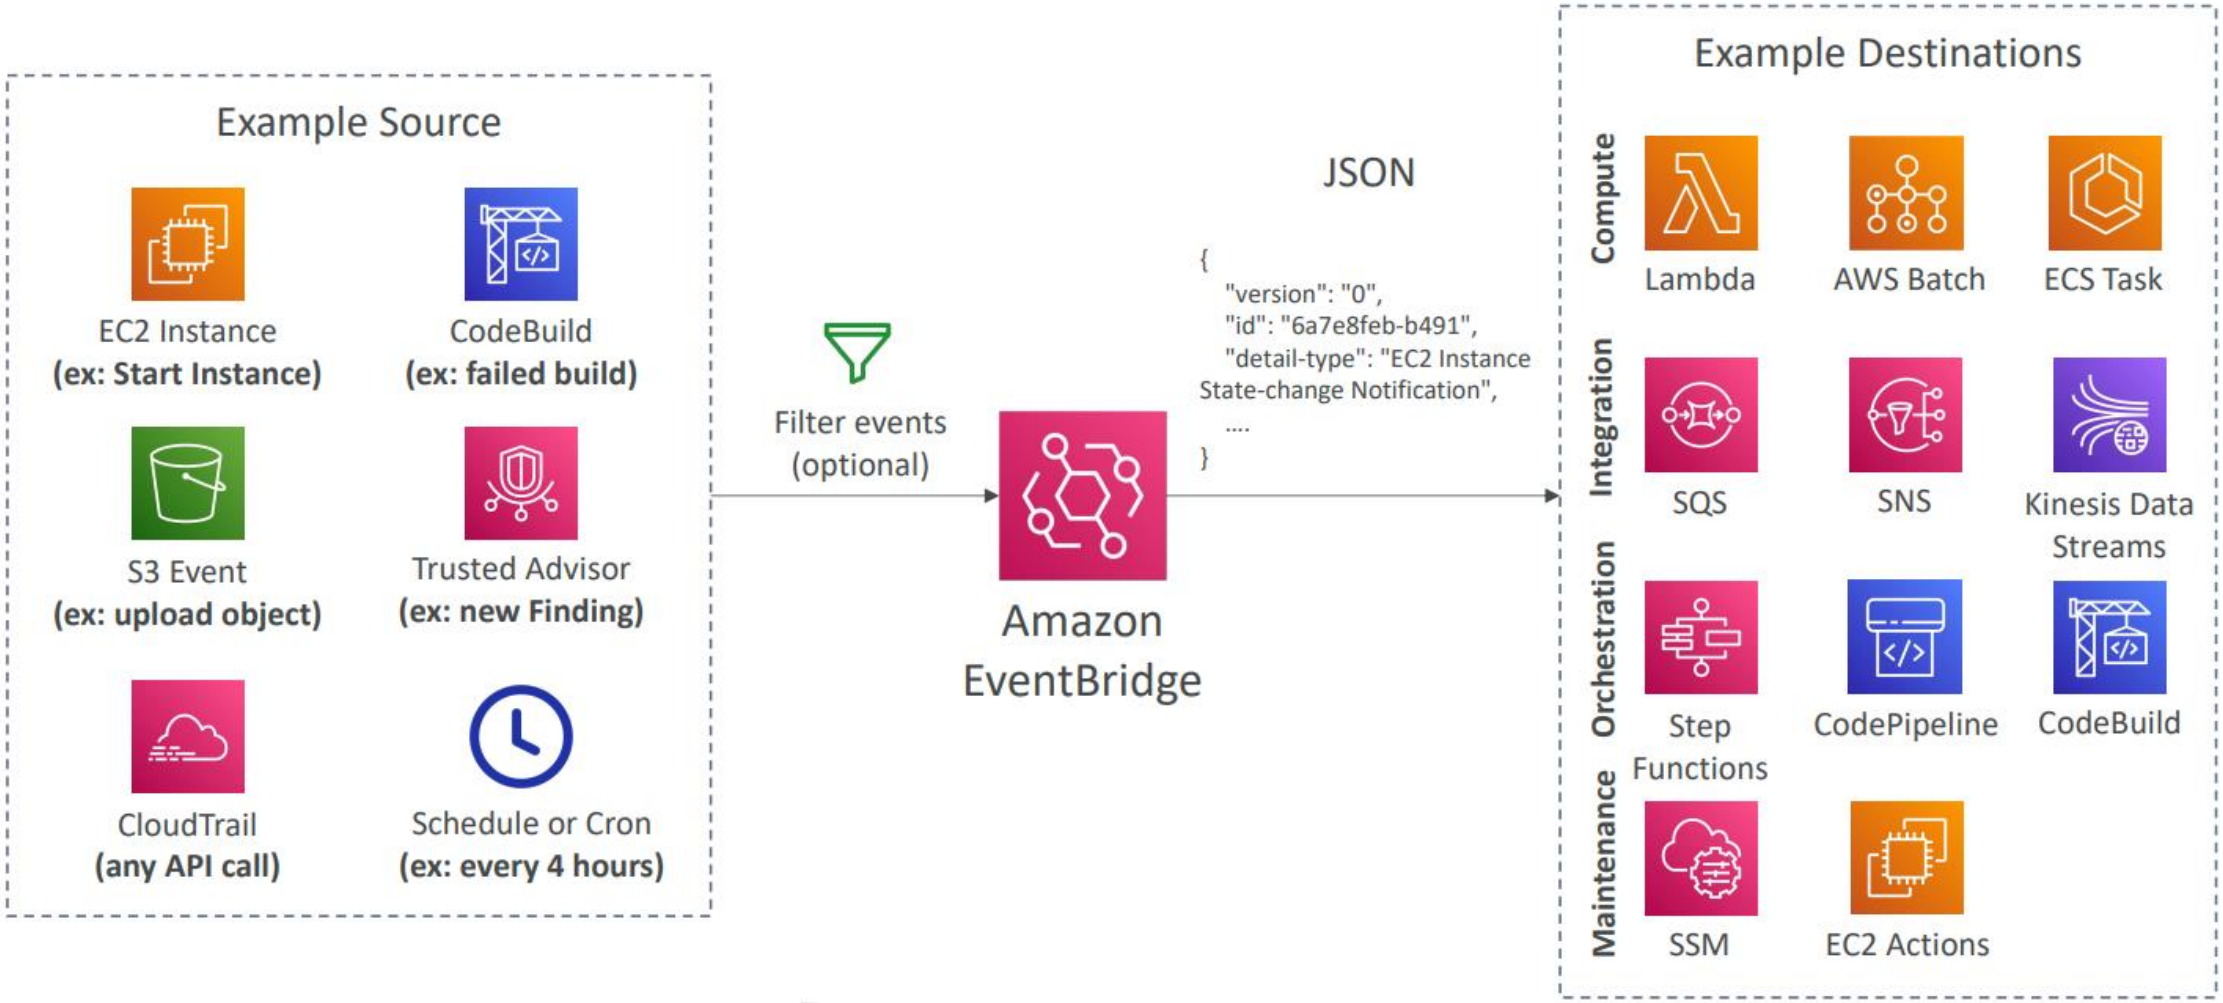

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge

- Event buses can be accessed by other AWS accounts using Resource-based Policies

- You can **archive events** (all/filter) sent to an event bus (indefinitely or set period)
- Ability to **replay archived events**

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge – Schema Registry

- EventBridge can analyze the events in your bus and infer the **schema**
- The **Schema Registry** allows you to generate code for your application, that will know in advance how data is structured in the event bus
- Schema can be versioned

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge – Resource-based Policy

- Manage permissions for a specific Event Bus
- Example: allow/deny events from another AWS account or AWS region
- Use case: aggregate all events from your AWS Organization in a single AWS account or AWS region

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "events:PutEvents",
            "Principal": { "AWS": "111122223333" },
            "Resource": "arn:aws:events:us-east-1:123456789012:
event-bus/central-event-bus"
        }
    ]
}
```

Allow **events** from another AWS account

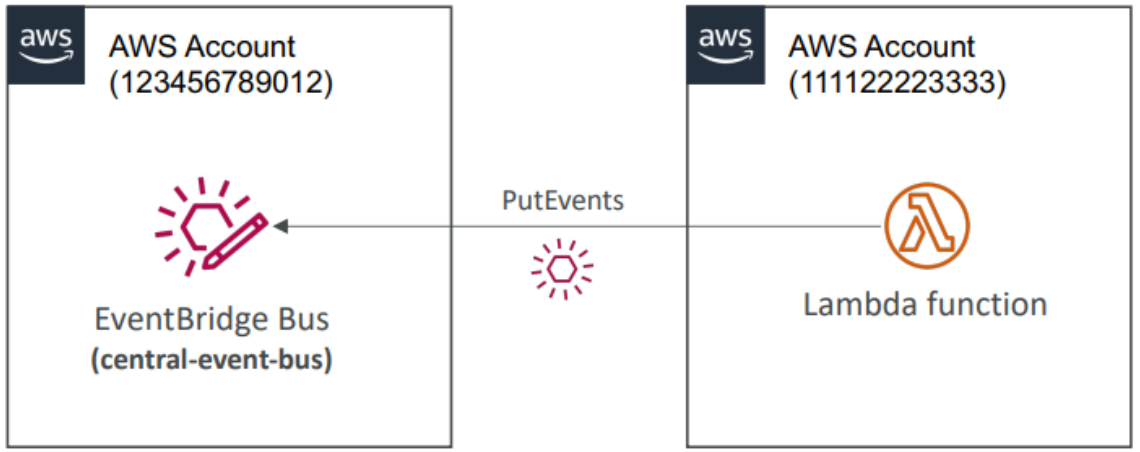

© Copyright by Cloud Mentor Pro | Confidential

AWS X-Ray

- Trace & analyze requests traveling in through your application

- Visualize request flow across different services, identify bottlenecks, troubleshoot performance issues in distributed application

© Copyright by Cloud Mentor Pro | Confidential

AWS X-Ray - Visual analysis of our applications

© Copyright by Cloud Mentor Pro | Confidential

AWS X-Ray advantages

- Troubleshooting performance (bottlenecks)
- Understand dependencies in a microservice architecture
- Pinpoint service issues
- Review request behavior
- Find errors and exceptions
- Are we meeting time SLA?
- Where I am throttled?
- Identify users that are impacted

© Copyright by Cloud Mentor Pro | Confidential

X-Ray compatibility

- AWS Lambda
- Elastic Beanstalk
- ECS
- ELB
- API Gateway
- EC2 Instances or any application server (even on premise)

© Copyright by Cloud Mentor Pro | Confidential

AWS X-Ray Leverages Tracing

- Tracing is an end to end way to following a "request"
- Each component dealing with the request adds its own "trace"
- Tracing is made of segments (+ sub segments)
- Annotations can be added to traces to provide extra-information
- Ability to trace:
  - Every request
  - Sample request (as a % for example or a rate per minute)
- X-Ray Security:
  - IAM for authorization
  - KMS for encryption at rest

© Copyright by Cloud Mentor Pro | Confidential

AWS X-Ray - How to enable it?

- 1) Your code (lava, Python, Go, Node.is, NET) must import the AWS X-Ray SDK
  - Very little code modification needed
  - The application SDK will then capture:
    - Calls to AWS services
    - HTTP / HTTPS requests
    - Database Calls (MySQL, PostgreSQL, DynamoDB)
    - Queue calls (SQS)

- 2) Install the X-Ray daemon or enable X-Ray AWS Integration
  - X-Ray daemon works as a low level UDP packet interceptor (Linux / Windows / Mac...)
  - AWS Lambda / other AWS services already run the X-Ray
  - Each application must have the IAM rights to write data to X-Ray

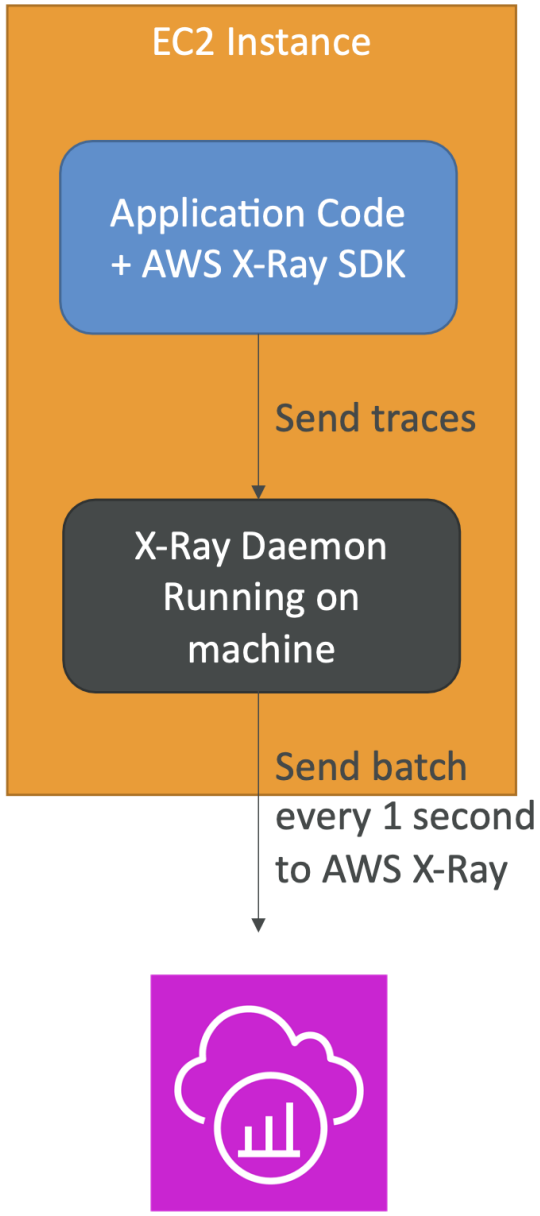

© Copyright by Cloud Mentor Pro | Confidential

The X-Ray magic

- X-Ray service collects data from all the different services
- Service map is computed from all the segments and traces
- X-Ray is graphical, so even non technical people can help troubleshoot

© Copyright by Cloud Mentor Pro | Confidential

AWS X-Ray Troubleshooting

- If X-Ray is not working on EC2
  - Ensure the EC2 IAM Role has the proper permissions
  - Ensure the EC2 instance is running the X-Ray Daemon

- To enable on AWS Lambda:
  - Ensure it has an IAM execution role with proper policy (AWSX-Ray WriteOnlyAccess)
  - Ensure that X-Ray is imported in the code
  - Enable **Lambda X-Ray Active Tracing**

© Copyright by Cloud Mentor Pro | Confidential

X-Ray Instrumentation in your code

- **Instrumentation** means the measure of product's performance, diagnose errors, and to write trace information.
- To instrument your application code, you use the **X-Ray SDK**
- Many SDK require only configuration changes
- You can modify your application code to customize and annotation the data that the SDK sends to X- Ray, using **interceptors, filters, handlers, middleware…**

Example for Node.js & Express

```
var app = express();

var AWSXRay = require('aws-xray-sdk');
app.use(AWSXRay.express.openSegment('MyApp'));

app.get('/', function (req, res) {
  res.render('index');
});

app.use(AWSXRay.express.closeSegment());
```

© Copyright by Cloud Mentor Pro | Confidential

X-Ray Concepts

- Segments: each application / service will send them
- Subsegments: if you need more details in your segment
- Trace: segments collected together to form an end-to-end trace
- Sampling: decrease the amount of requests sent to X-Ray, reduce cost
- Annotations: Key Value pairs used to index traces and use with filters
- Metadata: Key Value pairs, not indexed, not used for searching

- The X-Ray daemon / agent has a config to send traces cross account:
  - Make sure the IAM permissions are correct - the agent will assume the role
  - This allows to have a central account for all your application tracing

© Copyright by Cloud Mentor Pro | Confidential

X-Ray Sampling Rules

- With sampling rules, you control the amount of data that you record
- You can modify sampling rules without changing your code

- By default, the X-Ray SDK records the first request each second, and five percent of any additional requests.
- One request per second is the **reservoir**, which ensures that at least one trace is recorded each second as long the service is serving requests.
- Five percent is the **rate** at which additional requests beyond the reservoir size are sampled.

© Copyright by Cloud Mentor Pro | Confidential

X-Ray Custom Sampling Rules

- You can create your own rules with the reservoir and rate

Example Higher minimum rate for POSTs

- **Rule name** – `POST minimum`
- **Priority** – `100`
- **Reservoir** – `10`
- **Rate** – `0.10`
- **Service name** – `*`
- **Service type** – `*`
- **Host** – `*`
- **HTTP method** – `POST`
- **URL path** – `*`
- **Resource ARN** – `*`

Example Debugging rule to trace all requests for a problematic route

A high-priority rule applied temporarily for debugging.

- **Rule name** – `DEBUG - history updates`
- **Priority** – `1`
- **Reservoir** – `1`
- **Rate** – `1`
- **Service name** – `Scorekeep`
- **Service type** – `*`
- **Host** – `*`
- **HTTP method** – `PUT`
- **URL path** – `/history/*`
- **Resource ARN** – `*`

© Copyright by Cloud Mentor Pro | Confidential

X-Ray Write APIs (used by the X-Ray daemon)

```
"Effect": "Allow",
"Action": [
    "xray:PutTraceSegments",
    "xray:PutTelemetryRecords",
    "xray:GetSamplingRules",
    "xray:GetSamplingTargets",
    "xray:GetSamplingStatisticSummaries"
],
"Resource": [
    "*"
]
```

arn:aws:iam::aws:policy/AWSXrayWriteOnlyAccess

- **PutTraceSegments**: Uploads segment documents to AWS X-Ray
- **PutTelemetryRecords**: Used by the AWS X-Ray daemon to upload telemetry.
  - SegmentsReceivedCount,
  - SegmentsRejectedCounts,
  - BackendConnectionErrors...
- **GetSamplingRules**: Retrieve all sampling rules (to know what/when to send)
- **GetSamplingTargets** & **GetSamplingStatisticSummaries**: advanced
- The X-Ray daemon needs to have an IAM policy authorizing the correct APl calls to function correctly

© Copyright by Cloud Mentor Pro | Confidential

X-Ray Read APIs — continued

```
"Effect": "Allow",
"Action": [
    "xray:GetSamplingRules",
    "xray:GetSamplingTargets",
    "xray:GetSamplingStatisticSummaries",
    "xray:BatchGetTraces",
    "xray:GetServiceGraph",
    "xray:GetTraceGraph",
    "xray:GetTraceSummaries",
    "xray:GetGroups",
    "xray:GetGroup",
    "xray:GetTimeSeriesServiceStatistics"
],
"Resource": [
    "*"
]
```

- **GetServiceGraph**: main graph
- **BatchGetTraces**: Retrieves a list of traces specified by ID. Each trace is a collection of segment documents that originates from a single request.
- **GetTraceSummaries**: Retrieves IDs and annotations for traces available for a specified time frame using an optional filter. To get the full traces, pass the trace IDs to BatchGetTraces.
- **GetTraceGraph**: Retrieves a service graph for one or more specific trace

© Copyright by Cloud Mentor Pro | Confidential

X-Ray with Elastic Beanstalk

- AWS Elastic Beanstalk platforms include the X-Ray daemon
- You can run the daemon by setting an option in the Elastic Beanstalk console or with a configuration file (in.ebextensions/xray-daemon.config)

```
option_settings:
  aws:elasticbeanstalk:xray:
    XRayEnabled: true
```

- Make sure to give your instance profile the correct IAM permissions so that the X-Ray daemon can function correctly
- Then make sure your application code is instrumented with the X-Ray SDK
- Note: The X-Ray daemon is not provided for Multicontainer Docker

© Copyright by Cloud Mentor Pro | Confidential

ECS + X-Ray integration options

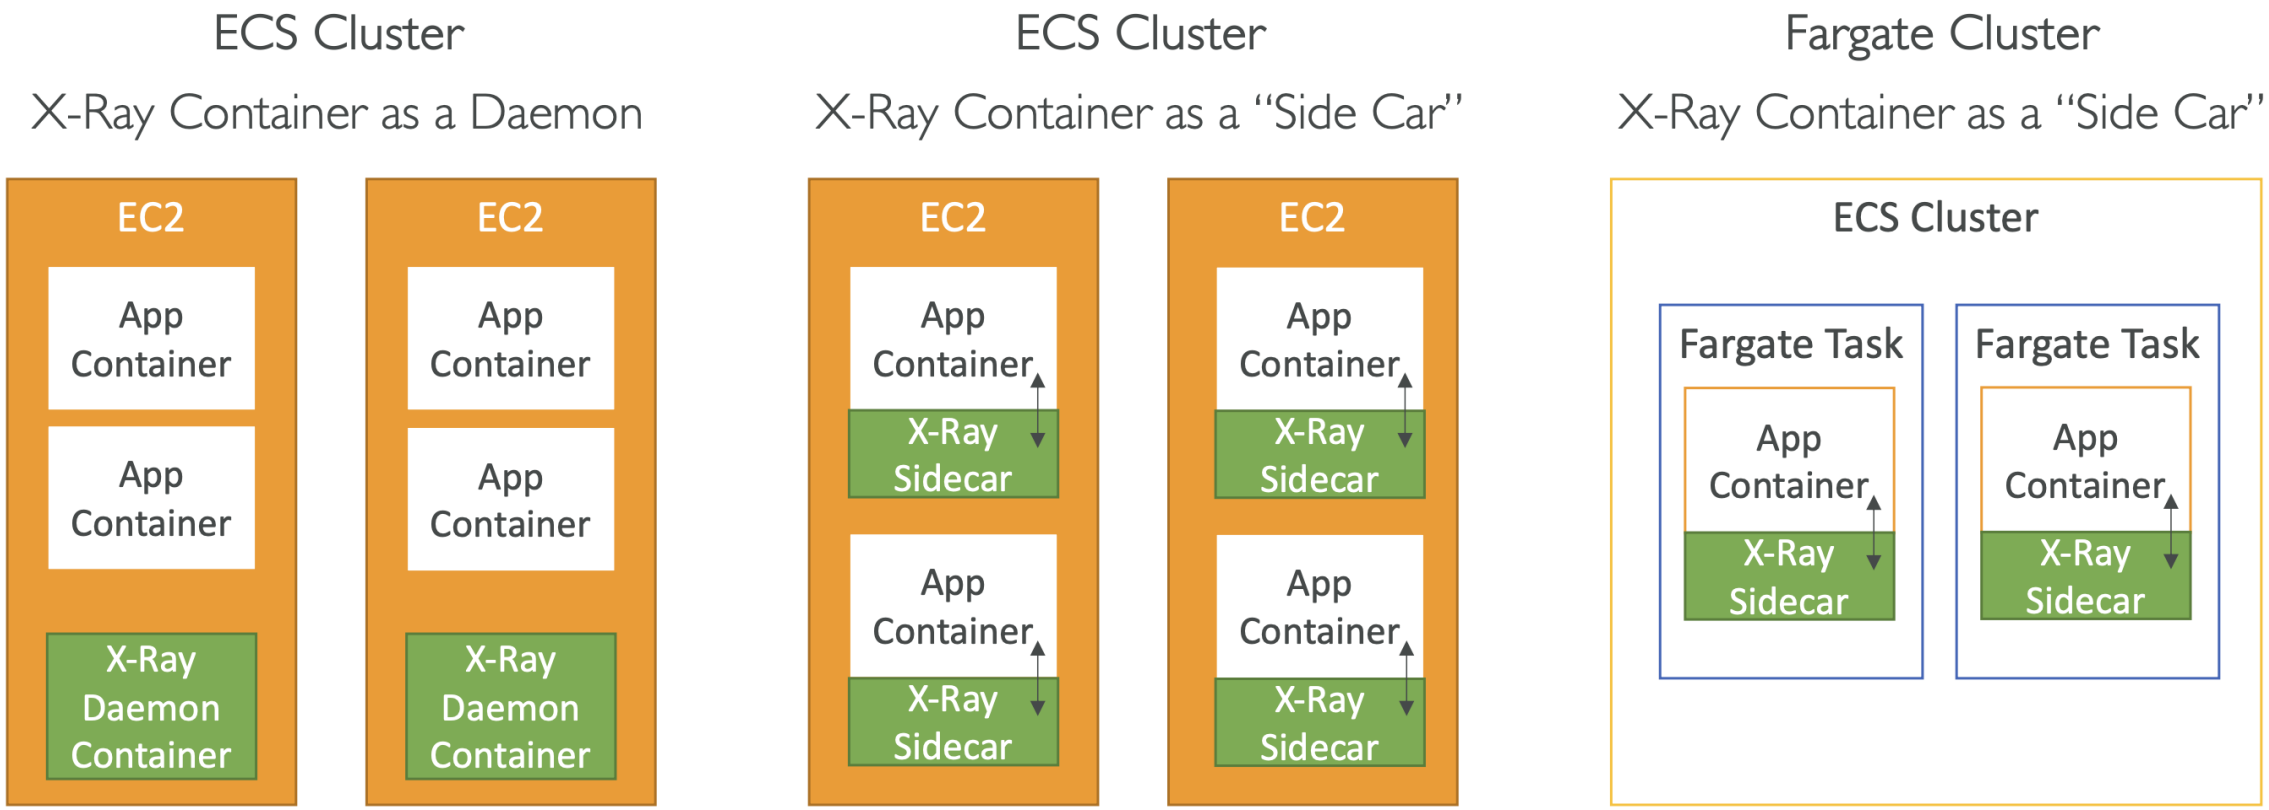

© Copyright by Cloud Mentor Pro | Confidential

ECS + X-Ray: Example Task Definition

```
{
  "name": "xray-daemon",
  "image": "123456789012.dkr.ecr.us-east-2.amazonaws.com/xray-daemon",
  "cpu": 32,
  "memoryReservation": 256,
  "portMappings" : [
      {
          "hostPort": 0,
          "containerPort": 2000,
          "protocol": "udp"
      }
   ],
},
{
  "name": "scorekeep-api",
  "image": "123456789012.dkr.ecr.us-east-2.amazonaws.com/scorekeep-api",
  "cpu": 192,
  "memoryReservation": 512,
  "environment": [
      { "name" : "AWS_REGION", "value" : "us-east-2" },
      { "name" : "NOTIFICATION_TOPIC", "value" : "arn:aws:sns:us-east-2:123456789012:scorekeep-notifications" },
      { "name" : "AWS_XRAY_DAEMON_ADDRESS", "value" : "xray-daemon:2000" }
  ],
  "portMappings" : [
      {
          "hostPort": 5000,
          "containerPort": 5000
      }
  ],
  "links": [
    "xray-daemon"
  ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

AWS Distro for Open Telemetry

- Secure, production-ready AWS-supported distribution of the open-source project Open Telemetry project
- Provides a single set of APIs, libraries, agents, and collector services
- Collects distributed traces and metrics from your apps
- Collects metadata from your AWS resources and services
- **Auto-instrumentation Agents** to collect traces without changing your code
- Send traces and metrics to multiple AWS services and partner solutions
  - X-Ray, CloudWatch, Prometheus...
- Instrument your apps running on AWS (e.g., EC2, ECS, EKS, Fargate, Lambda) as well as on-premises
- **Migrate from X-Ray to AWS Distro for Temeletry if you want to standardize with open-source APIs from Telemetry or send traces to multiple destinations simultaneously**

© Copyright by Cloud Mentor Pro | Confidential

AWS Distro for Open Telemetry

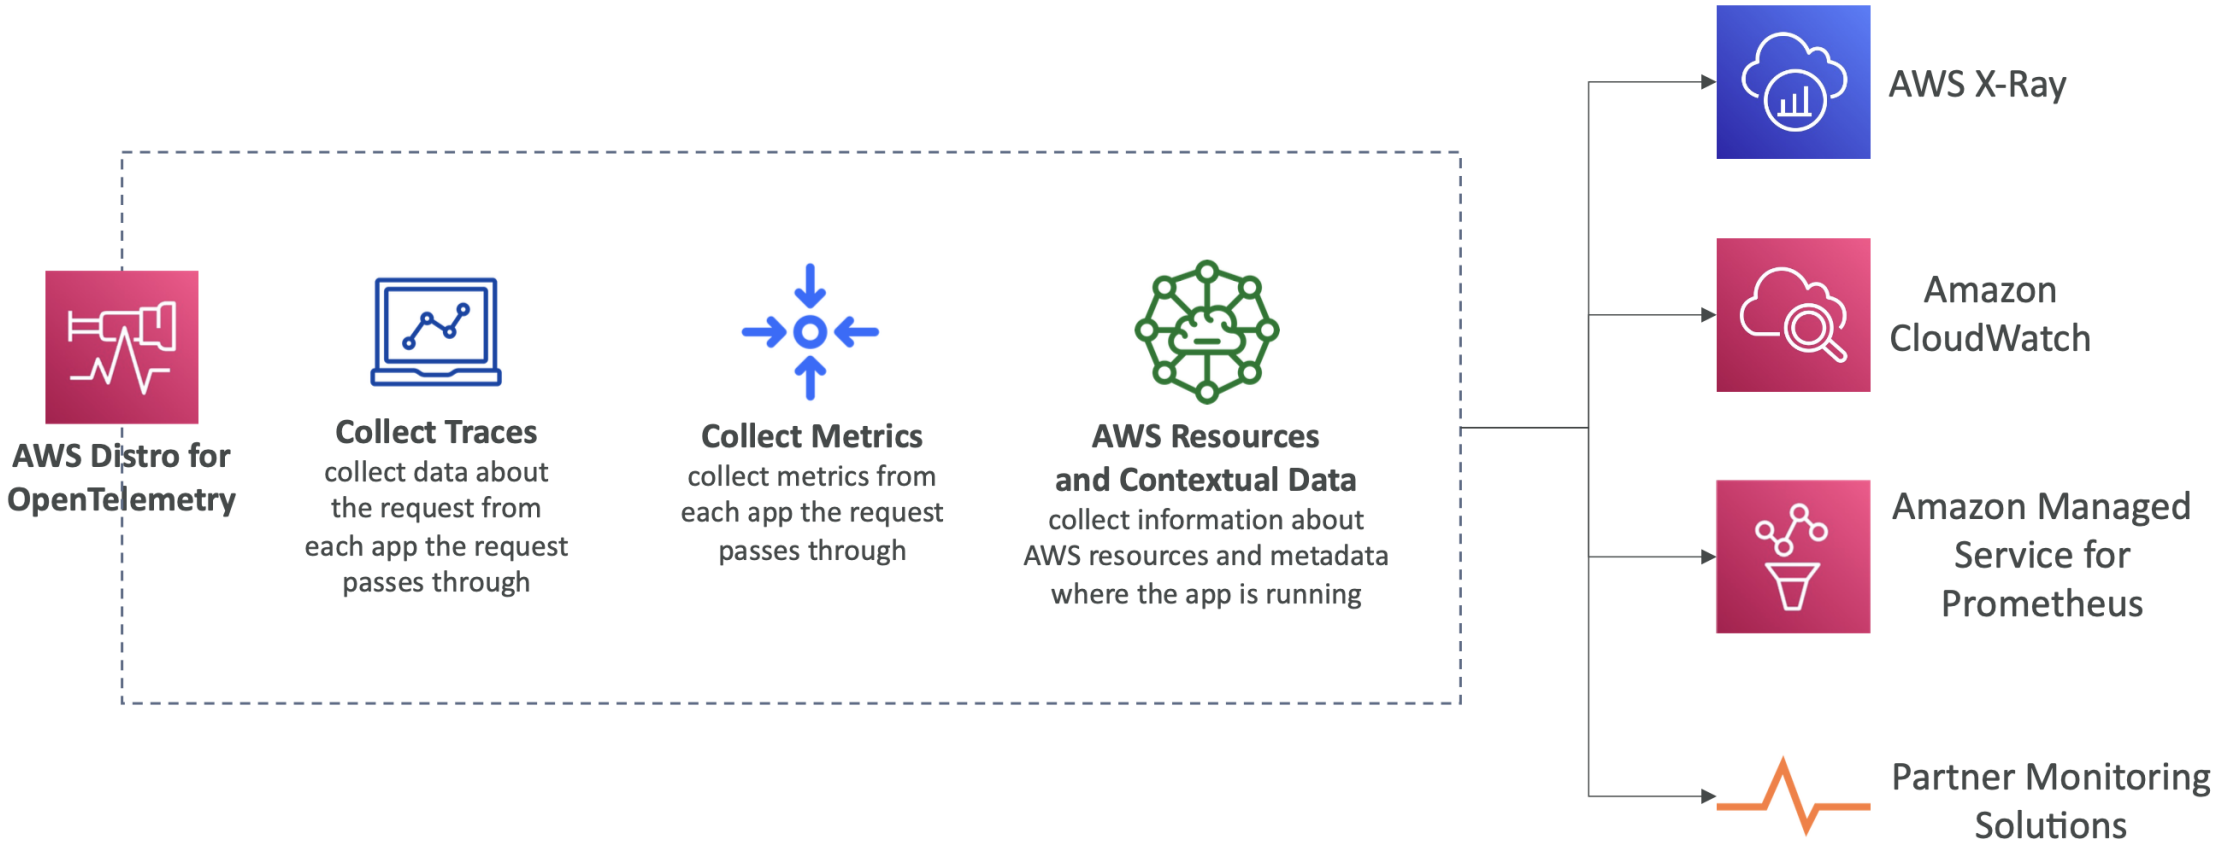

© Copyright by Cloud Mentor Pro | Confidential

AWS CloudTrail

- **Provides governance, compliance and audit** for your AWS Account
- CloudTrail is enabled by default!
- Get **an history of events / API calls made within your AWS Account** by:
  - Console
  - SDK
  - CLI
  - AWS Services
- Can put logs from CloudTrail into CloudWatch Logs or S3
- **A trail can be applied to All Regions (default) or a single Region**.
- If a resource is deleted in AWS, investigate CloudTrail first!

© Copyright by Cloud Mentor Pro | Confidential

CloudTrail Diagram

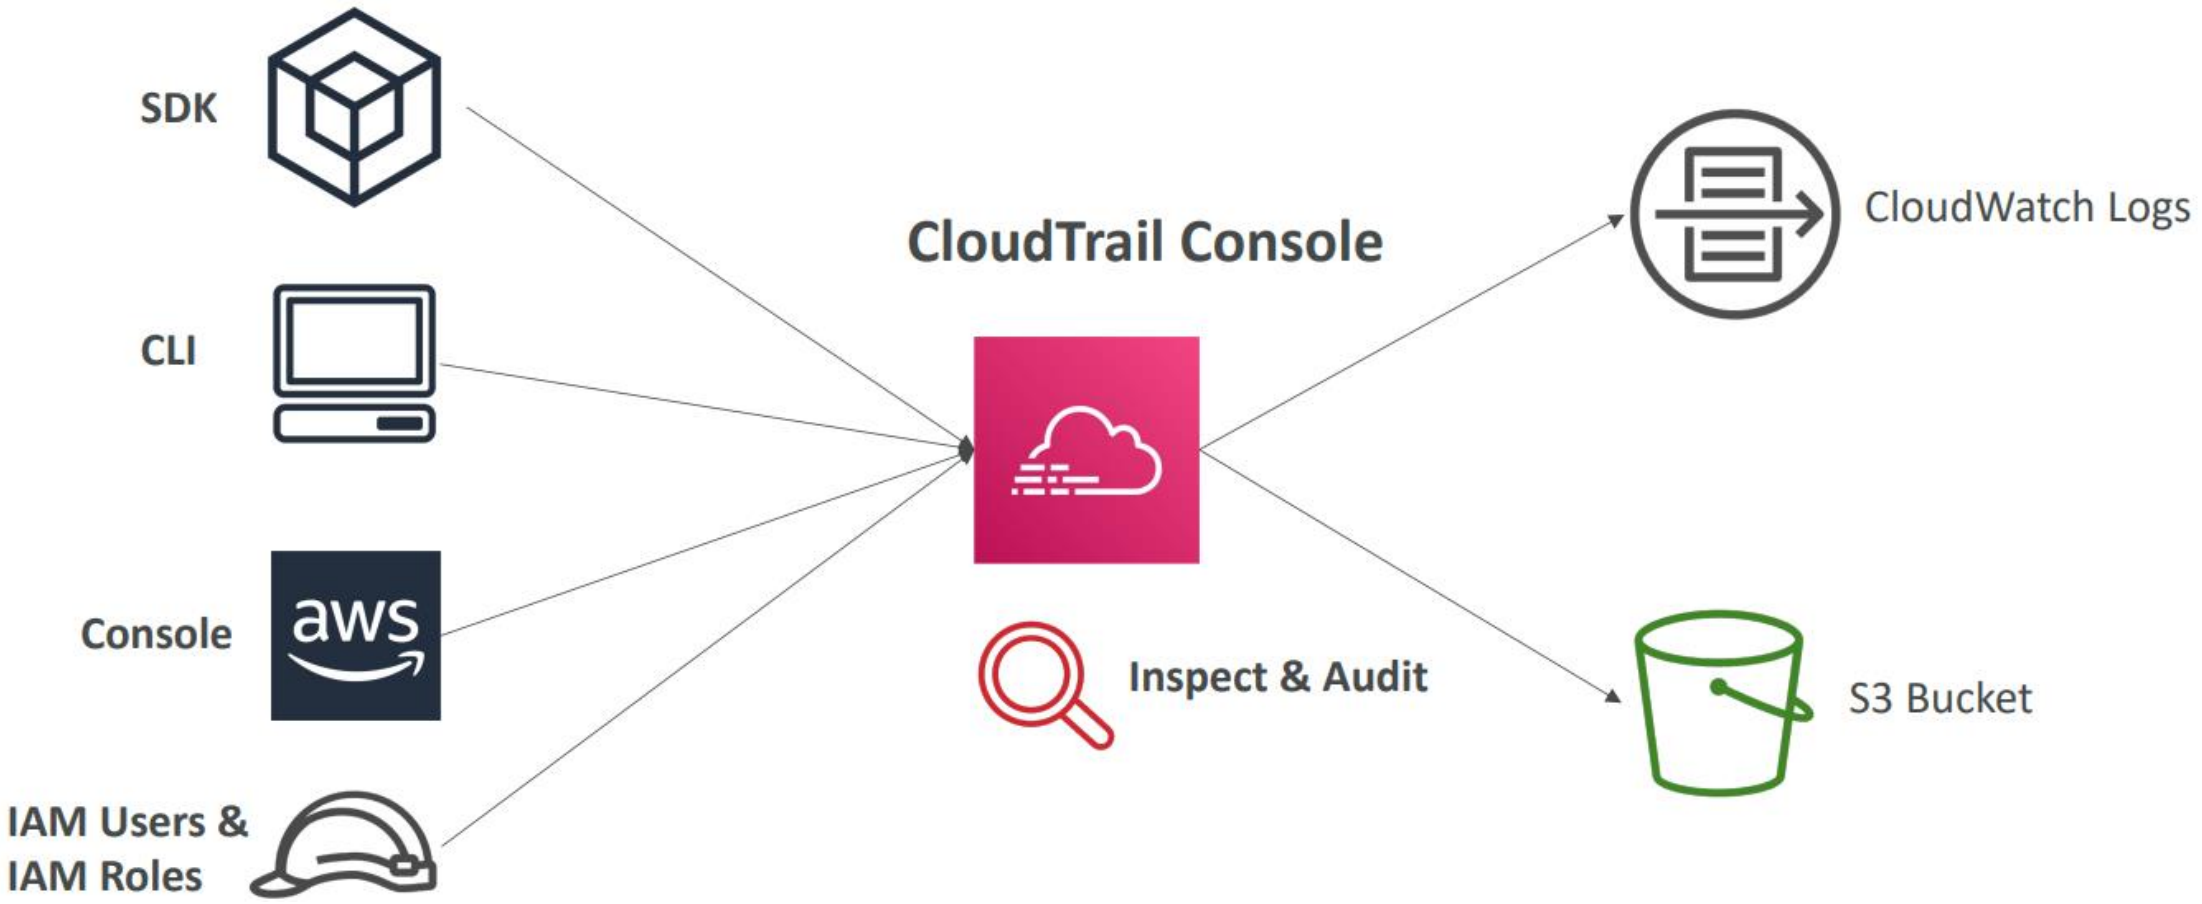

© Copyright by Cloud Mentor Pro | Confidential

CloudTrail Events

- Management Events:
  - Operations that are performed on resources in your AWS account
  - Examples:
    - Configuring security (IAM **AttachRolePolicy**)
    - Configuring rules for routing data (Amazon EC2 **CreateSubnet**)
    - Setting up logging (AWS CloudTrail **CreateTrail**)
  - **By default, trails are configured to log management events.**
  - Can separate **Read Events** (that don't modify resources) from **Write Events** (that may modify resources)
- Data Events:
  - By default, data events are not logged (because high volume operations)
  - Amazon S3 object-level activity (ex: GetObject, DeleteObject, PutObject): can separate Read and Write Events
  - AWS Lambda function execution activity (the Invoke API)
- CloudTrail Insights Events:
  - See next slide

© Copyright by Cloud Mentor Pro | Confidential

CloudTrail Insights

- Enable **CloudTrail Insights to detect unusual activity** in your account:
  - inaccurate resource provisioning
  - hitting service limits
  - Bursts of AWS IAM actions
  - Gaps in periodic maintenance activity
- CloudTrail Insights analyzes normal management events to create a baseline
- And then **continuously analyzes write events to detect unusual patterns**
  - Anomalies appear in the CloudTrail console
  - Event is sent to Amazon S3
  - An EventBridge event is generated (for automation needs)

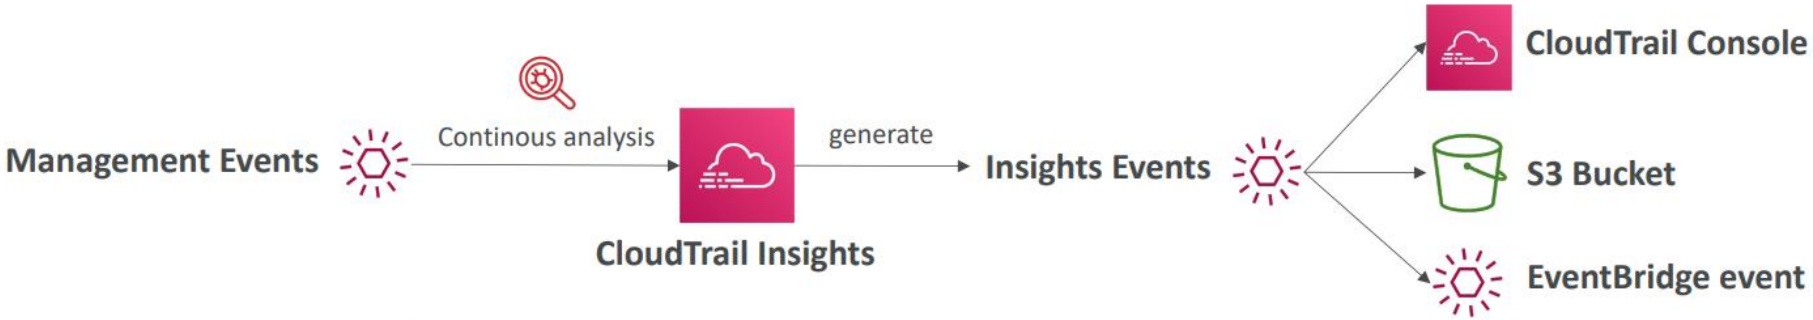

© Copyright by Cloud Mentor Pro | Confidential

CloudTrail Events Retention

- Events are stored for 90 days in CloudTrail
- To keep events beyond this period, log them to S3 and use Athena

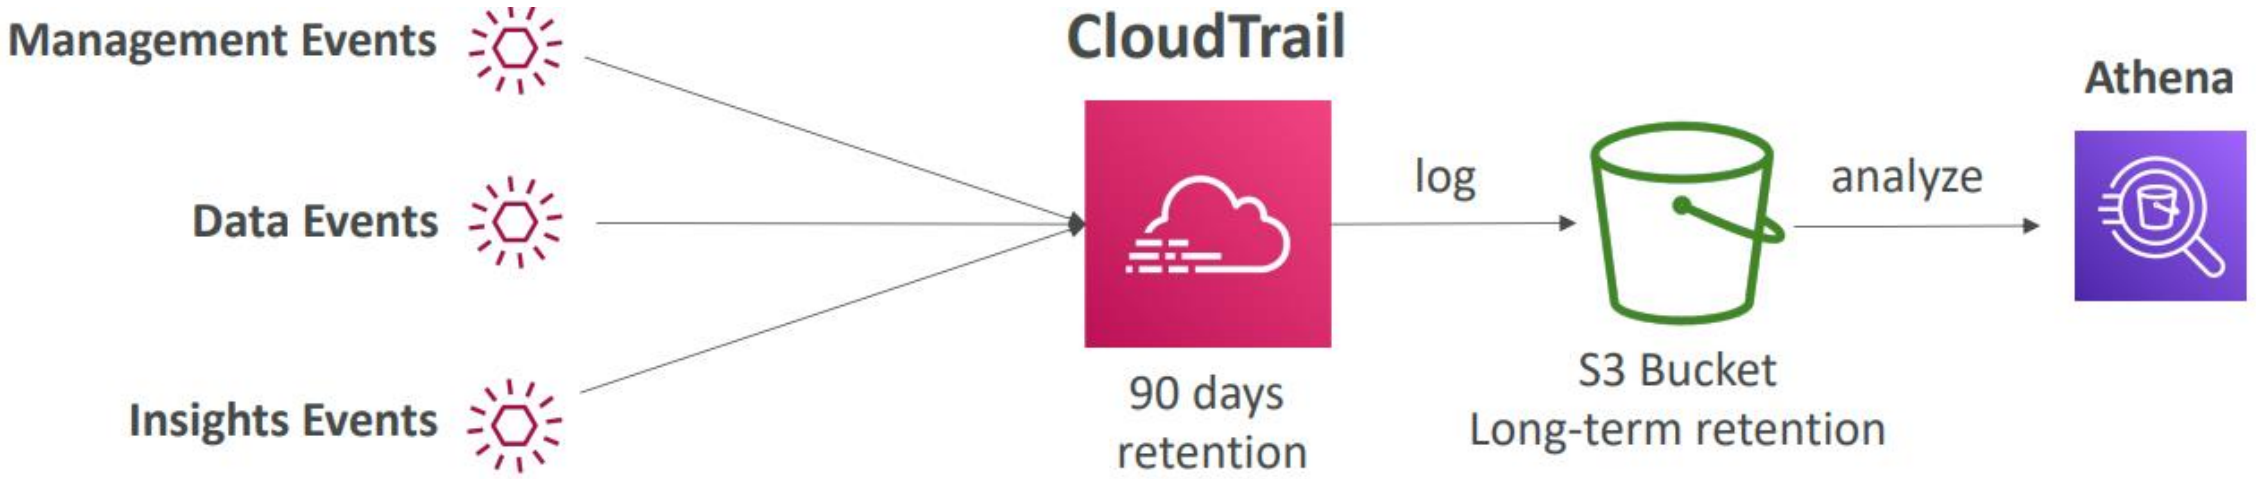

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge – Intercept API Calls

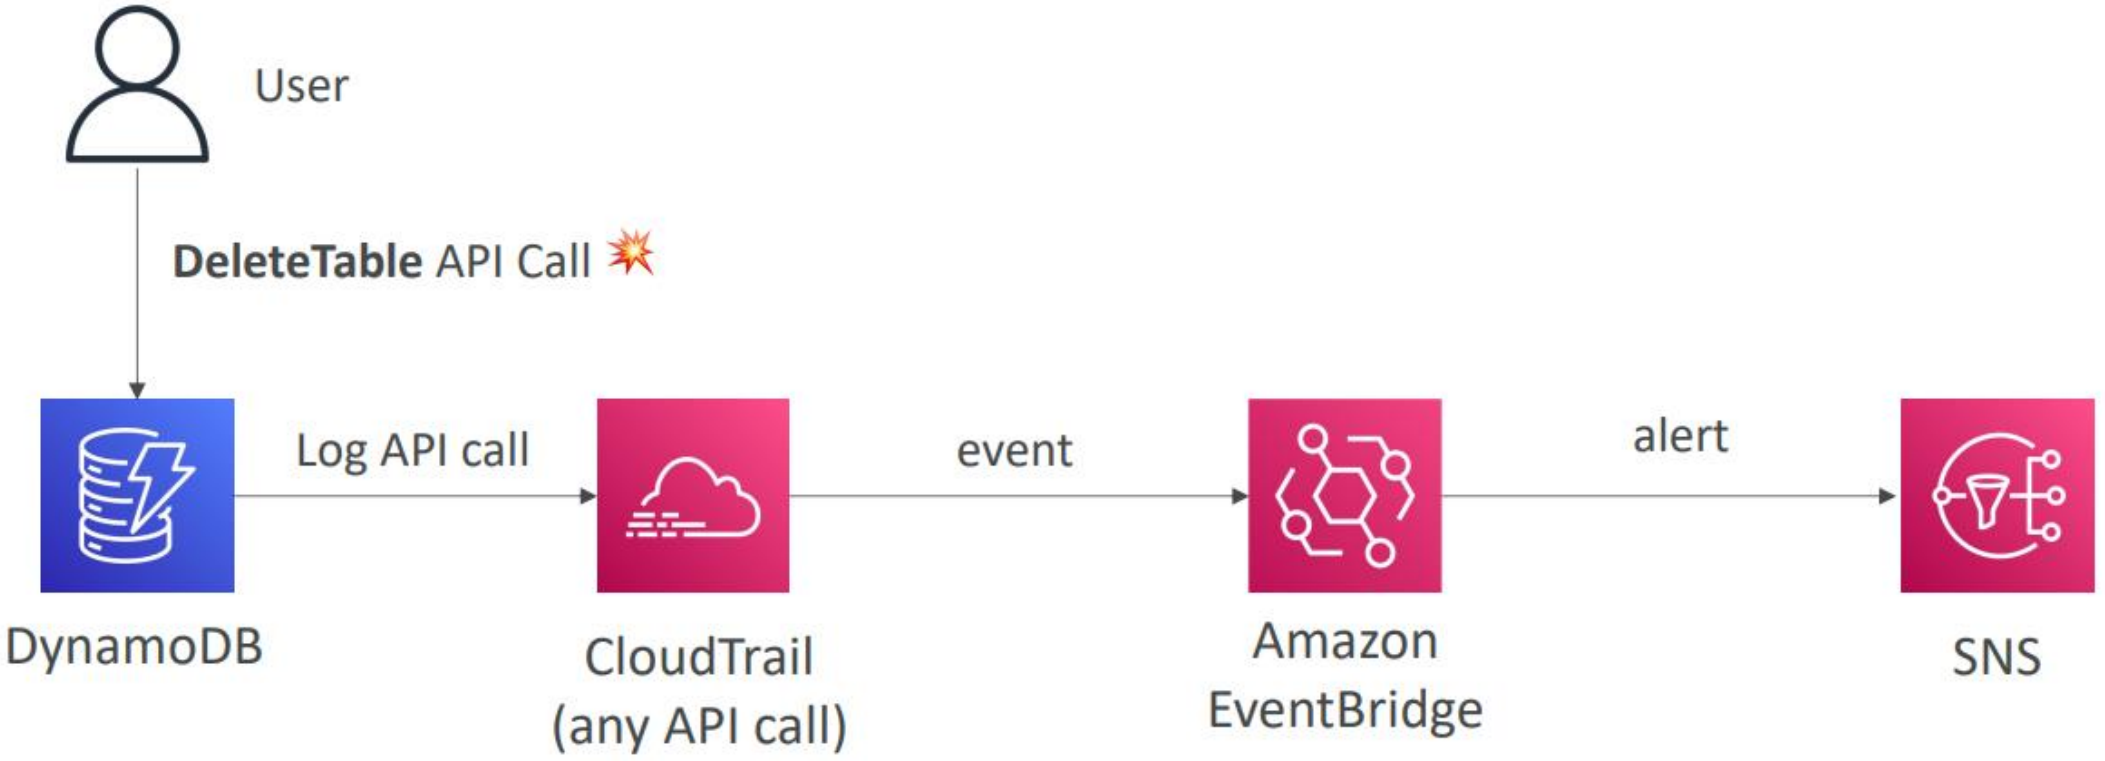

© Copyright by Cloud Mentor Pro | Confidential

Amazon EventBridge + CloudTrail

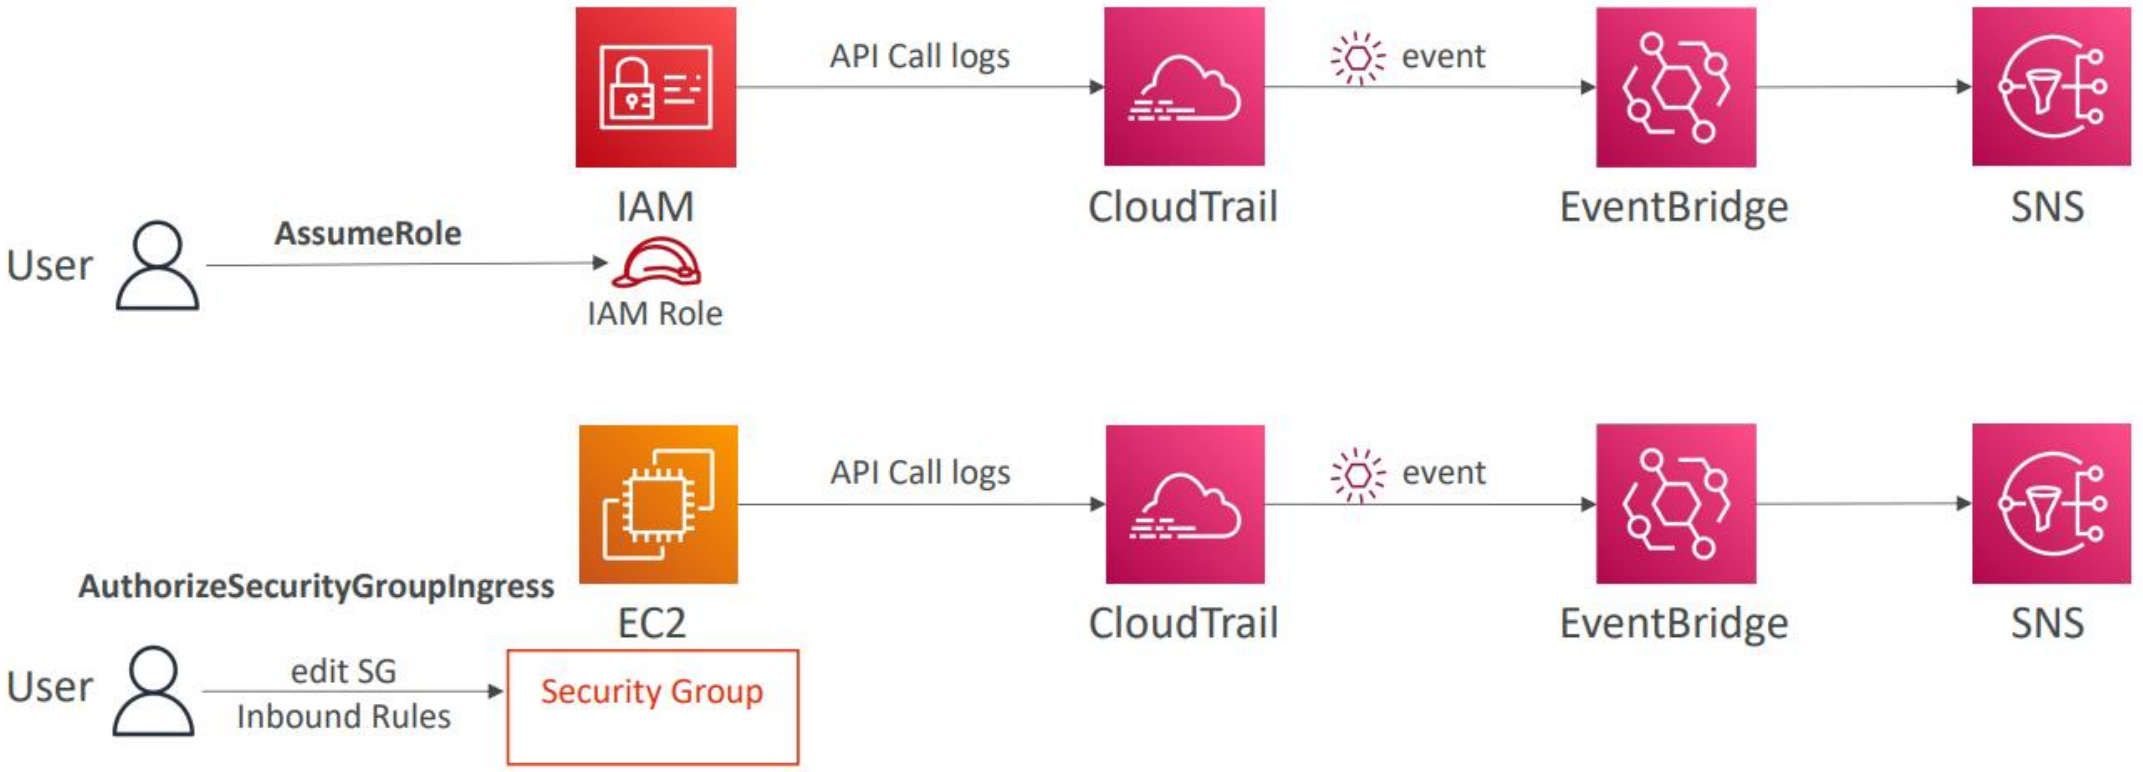

© Copyright by Cloud Mentor Pro | Confidential

Cloud Trail vs CloudWatch vs X-Ray

- CloudTrail
  - Audit API calls made by users / services / AWS console
  - Useful to detect unauthorized calls or root cause of changes
- CloudWatch
  - Cloud Watch Metrics over time for monitoring
  - Cloud Watch Logs for storing application log
  - CloudWatch Alarms to send notifications in case of unexpected metrics
- X-Ray
  - Automated Trace Analysis & Central Service Map Visualization
  - Latency, Errors and Fault analysis
  - Request tracking across distributed systems

© Copyright by Cloud Mentor Pro | Confidential

Exercise

- Monitor CPU Utilization of EC2 via CloudWatch Alarms

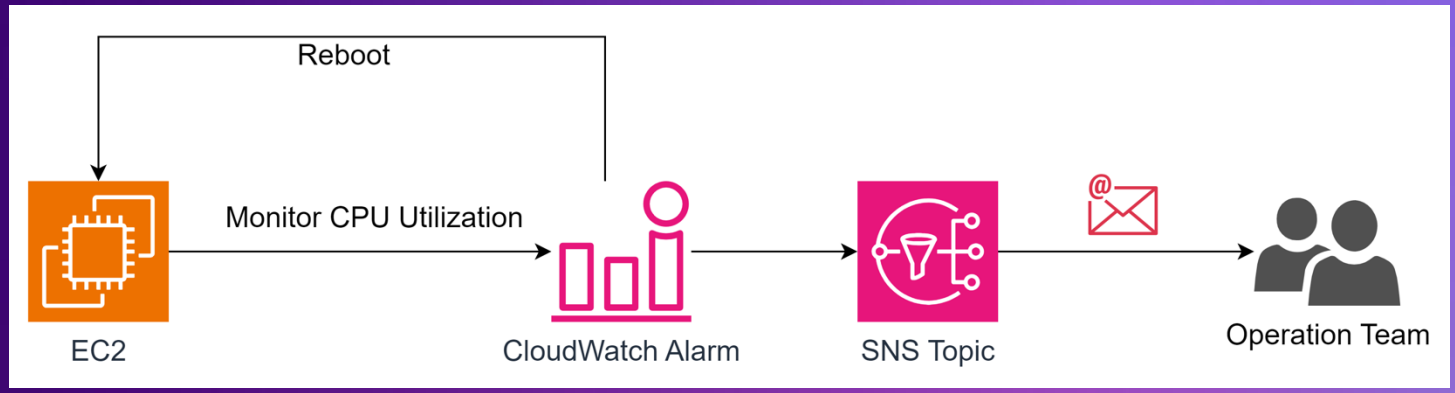

- Create an EventBridge Rule to get notified on EC2 Instance state change

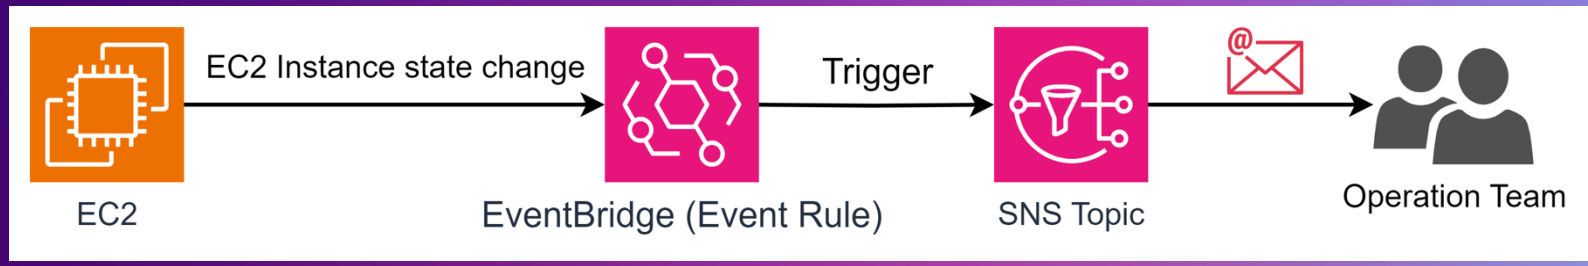

© Copyright by Cloud Mentor Pro | Confidential

Section 11

- *AWS Lambda*

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda

It’s a serverless world

© Copyright by Cloud Mentor Pro | Confidential

What's serverless?

- Serverless is a new paradigm in which the developers don't have to manage servers anymore…
- They just deploy code
- They just deploy… functions !
- Initially… Serverless == FaaS (Function as a Service)
- Serverless was pioneered by AWS Lambda but now also includes anything that's managed: "databases, messaging, storage, etc."
- **Serverless does not mean there are no servers…**

  it means you just don't manage / provision / see them

© Copyright by Cloud Mentor Pro | Confidential

Serverless in AWS

- AWS Lambda
- DynamoDB
- AWS Cognito
- AWS API Gateway
- Amazon S3
- AWS SNS & SQS
- AWS Kinesis Data Firehose
- Aurora Serverless
- Step Functions
- Fargate

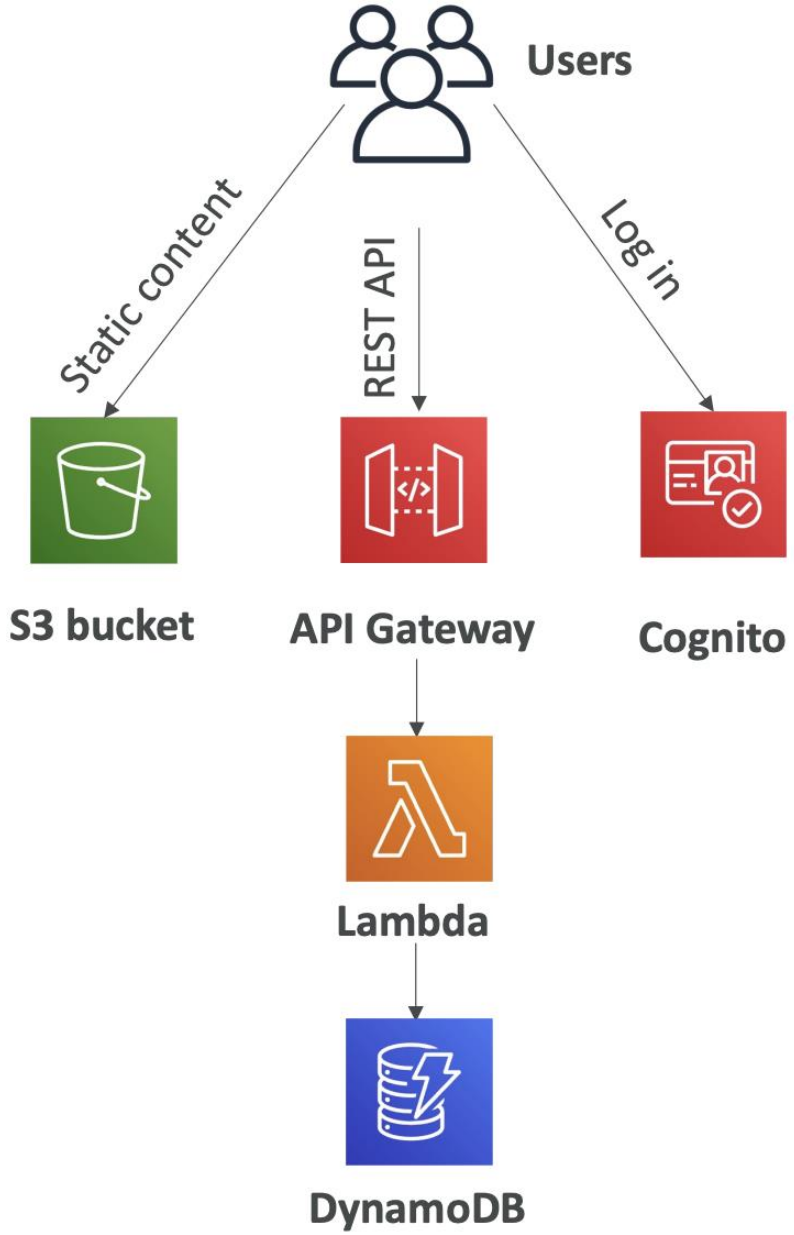

© Copyright by Cloud Mentor Pro | Confidential

Why AWS Lambda

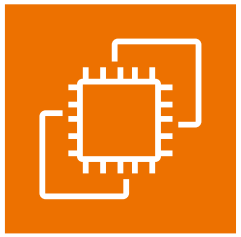

Amazon EC2

- Virtual Servers in the Cloud
- Limited by RAM and CPU
- Continuously running
- Scaling means intervention to add / remove servers

Amazon Lambda

- Virtual functions – no servers to manage!
- Limited by time - short executions
- Run on-demand
- Scaling is automated!

© Copyright by Cloud Mentor Pro | Confidential

Benefits of AWS Lambda

- Easy Pricing:
  - Pay per request and compute time
  - Free tier of 1,000,000 AWS Lambda requests and 400,000 GBs of compute time

- Integrated with the whole AWS suite of services
- Integrated with many programming languages
- Easy monitoring through AWS CloudWatch
- Easy to get more resources per functions (up to 10GB of RAM!)
- Increasing RAM will also improve CPU and network!

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda language support

- Node.js (JavaScript)
- Python
- Java
- C# (.NET Core) / Powershell
- Ruby
- Custom Runtime API (community supported, example Rust or Golang)

- Lambda Container Image
  - The container image must implement the Lambda Runtime API
  - ECS / Fargate is preferred for running arbitrary Docker images

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda Integrations Main ones

API Gateway

Kinesis

DynamoDB

S3

CloudFront

CloudWatch Events
EventBridge

CloudWatch Logs

SNS

SQS

Cognito

© Copyright by Cloud Mentor Pro | Confidential

Example: Serverless Thumbnail creation

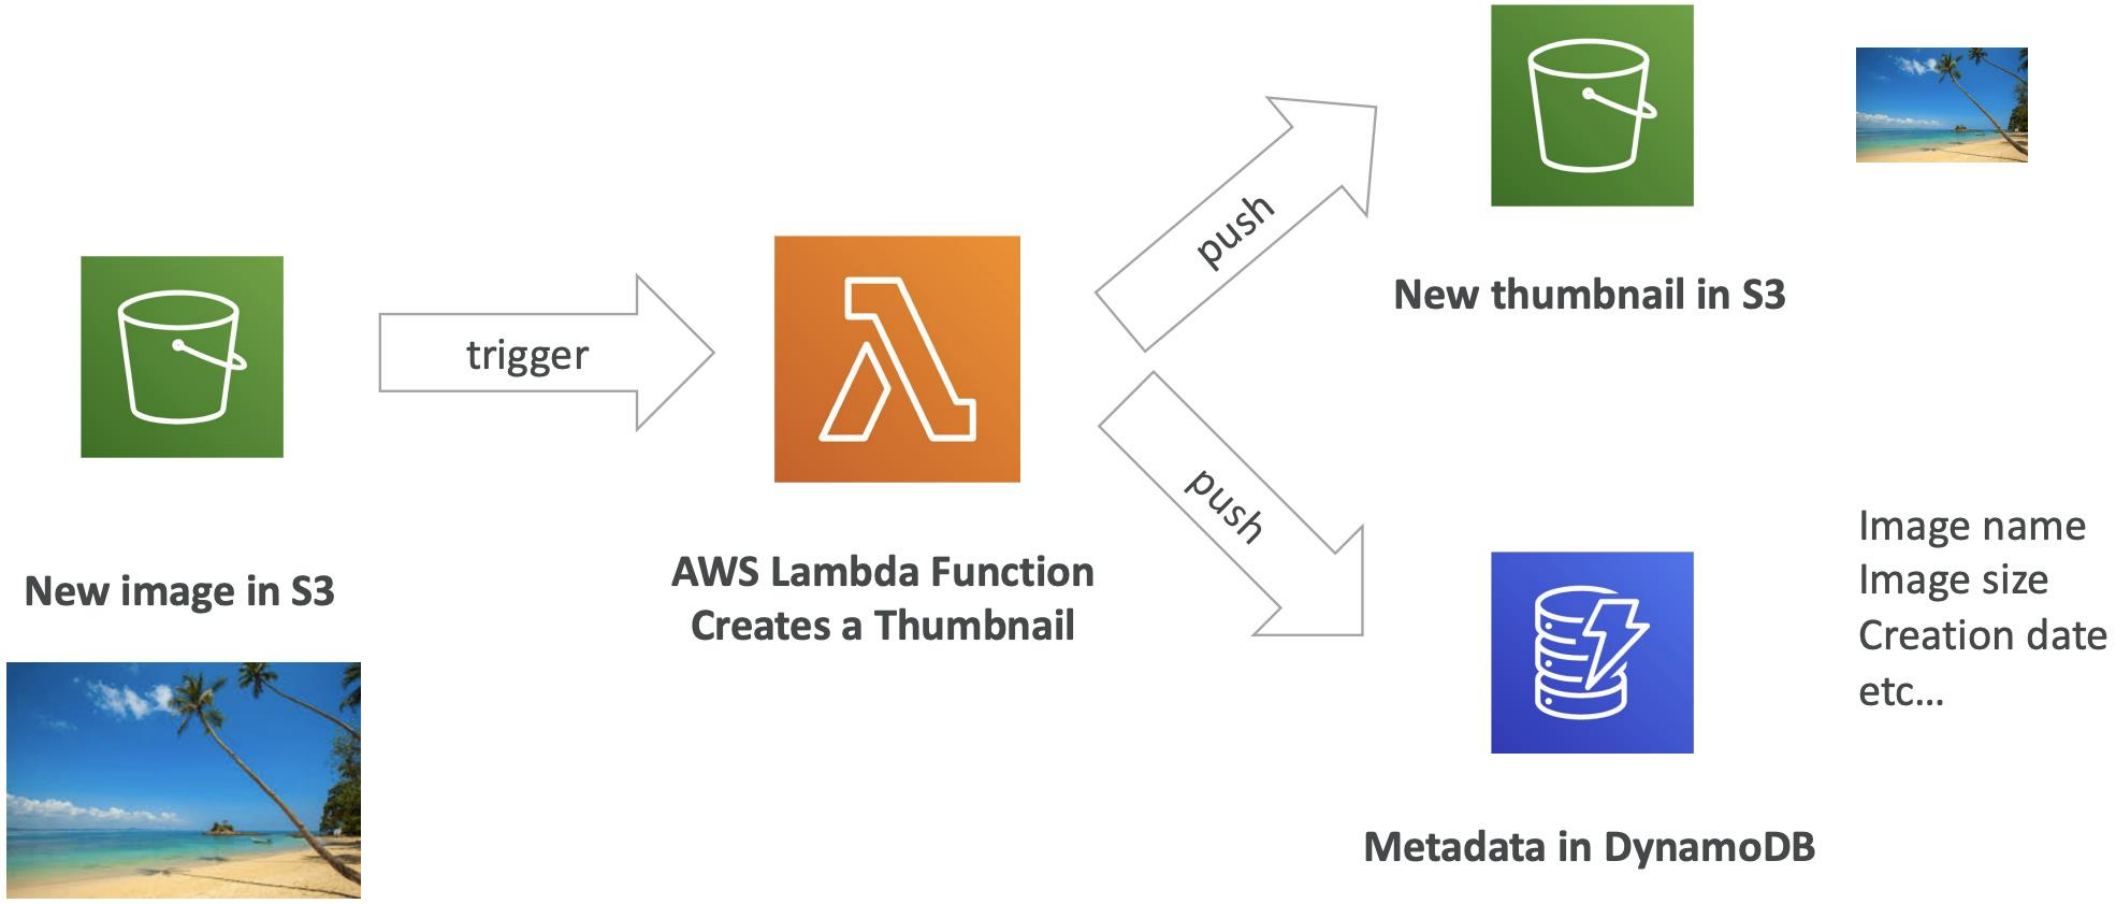

© Copyright by Cloud Mentor Pro | Confidential

Example: Serverless CRON Job

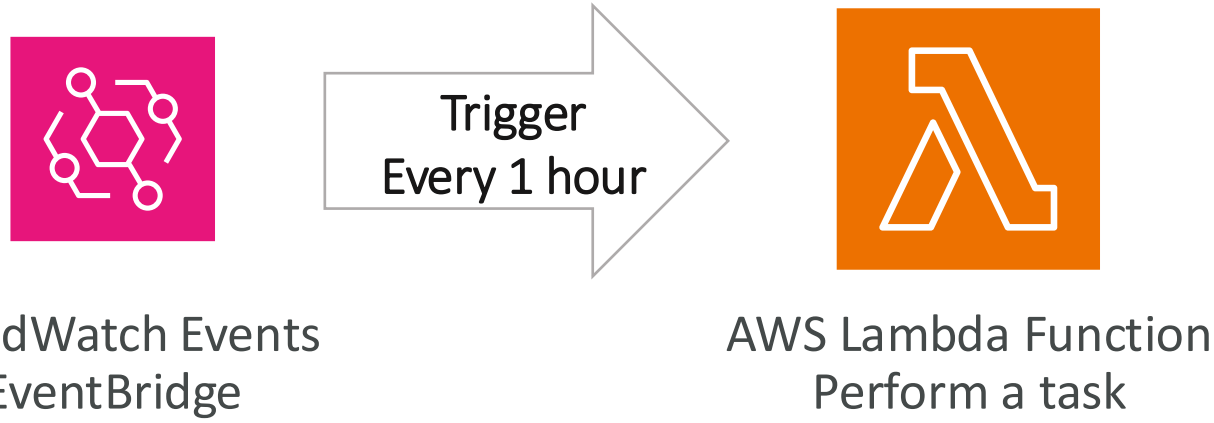

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda Pricing: example

- You can find overall pricing information here: https://aws.amazon.com/lambda/pricing/
- Pay per calls :
  - First 1,000,000 requests are free
  - $0.20 per 1 million requests thereafter ($0.0000002 per request)
- Pay per duration: (in increment of 1 ms)
  - 400,000 GB -seconds of compute time per month for FREE
  - == 400,000 seconds if function is 1GB RAM
  - == 3,200,000 seconds if function is 128 MB RAM
  - After that $1.00 for 600,000 GB-seconds
- It is usually very cheap to run AWS Lambda so it's very popular

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Synchronous Invocations

- Synchronous: CLI, SDK, API Gateway, Application Load Balancer
  - Results is returned right away
  - Error handling must happen client side (retries, exponential backoff, etc…)

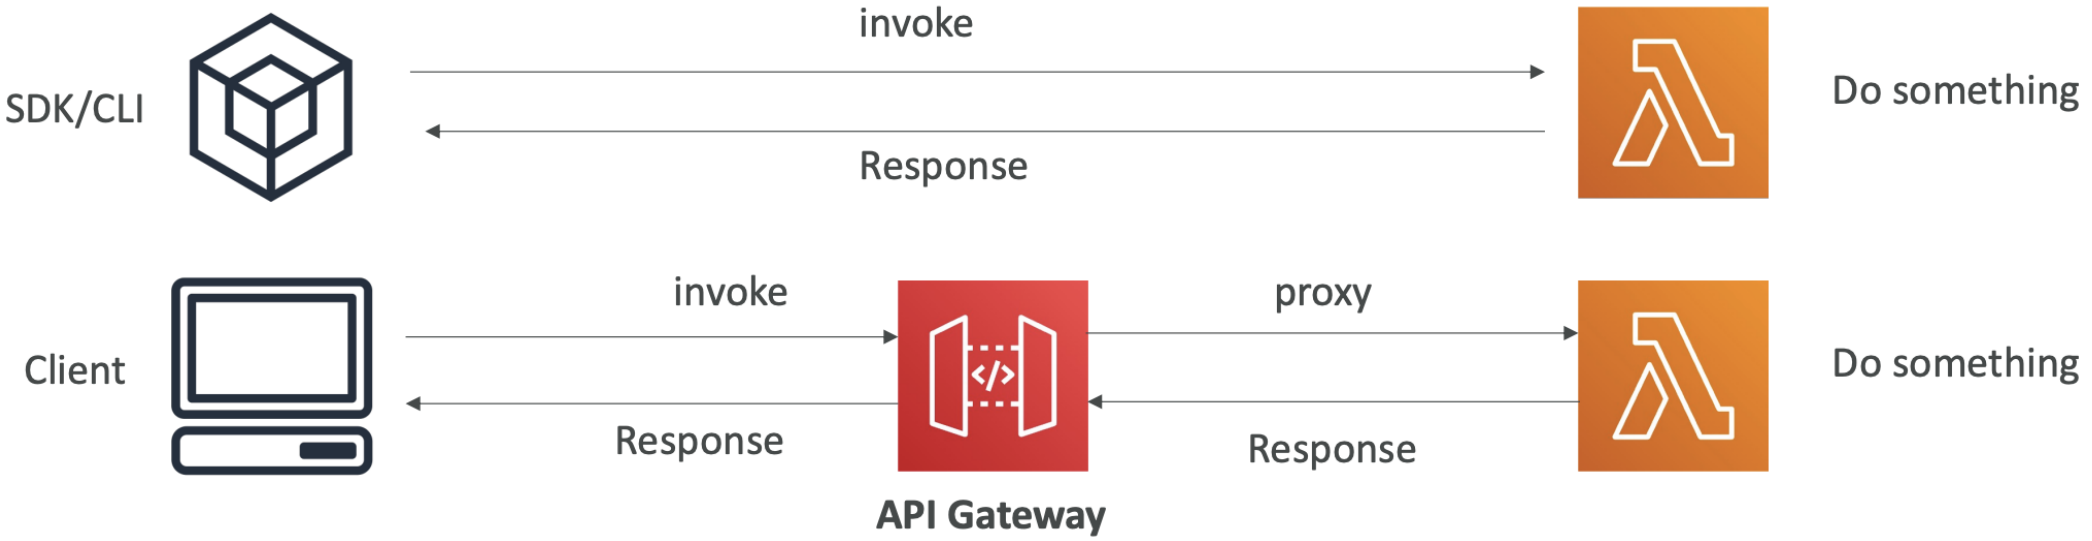

© Copyright by Cloud Mentor Pro | Confidential

Lambda - Synchronous Invocations - Services

- User Invoked:
  - Elastic Load Balancing (Application Load Balancer)
  - Amazon API Gateway
  - Amazon CloudFront (Lambda@Edge)
  - Amazon S3 Batch
- Service Invoked:
  - Amazon Cognito
  - AWS Step Functions
- Other Services:
  - Amazon Lex
  - Amazon Alexa
  - Amazon Kinesis Data Firehose

© Copyright by Cloud Mentor Pro | Confidential

Lambda Integration with ALB

- To expose a Lambda function as an HTTP(S) endpoint…
- You can use the Application Load Balancer (or an API Gateway)
- The Lambda function must be registered in a target group

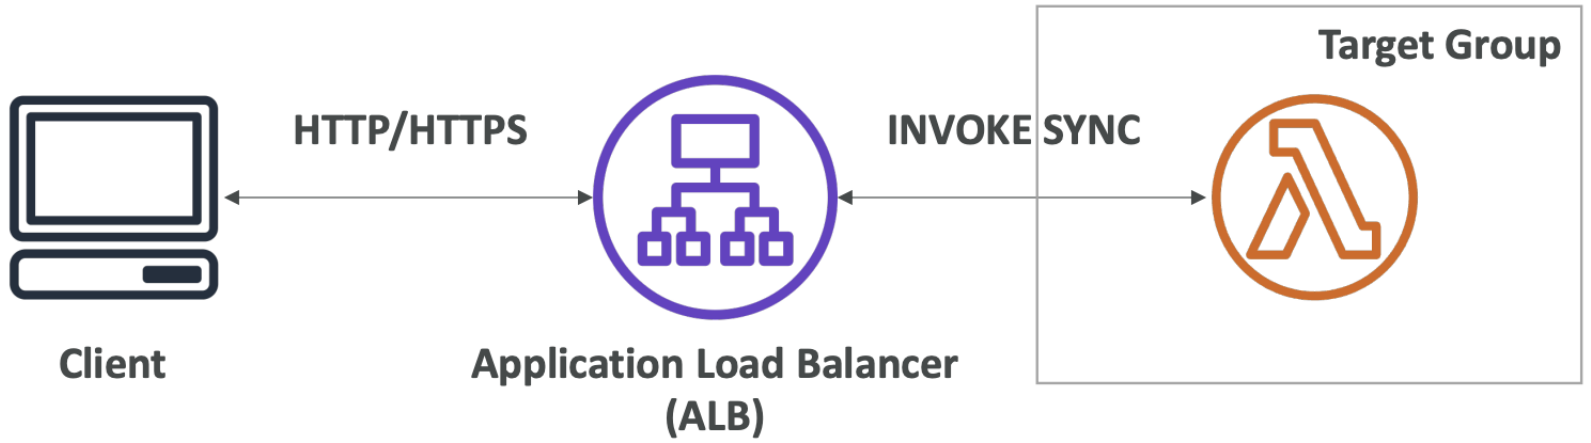

© Copyright by Cloud Mentor Pro | Confidential

ALB to Lambda: HTTP to JSON

Request Payload for Lambda Function

```
{
    "requestContext": {
        "elb": {
            "targetGroupArn": "arn:aws:elasticloadbalancing:us-east-2:12
            49e9d65c45c6791a"
        }
    },
    "httpMethod": "GET",
    "path": "/lambda",
    "queryStringParameters": {
        "query": "1234ABCD"
    },
    "headers": {
        "connection": "keep-alive",
        "host": "lambda-alb-123578498.us-east-2.elb.amazonaws.com",
        "user-agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWe
        Safari/537.36",
        "x-amzn-trace-id": "Root=1-5c536348-3d683b8b04734faae651f476",
        "x-forwarded-for": "72.12.164.125",
        "x-forwarded-port": "80",
        "x-forwarded-proto": "http",
    },
    "body": "",
    "isBase64Encoded": false
}
```

ELB information

HTTP Method & Path

Query String Parameters as Key/Value pairs

Headers as Key/Value pairs

Body (for POST, PUT...) & isBase64Encoded

© Copyright by Cloud Mentor Pro | Confidential

Lambda to ALB conversions: JSON to HTTP

Response from the Lambda Function

```
{
    "statusCode": 200,
    "statusDescription": "200 OK",
    "headers": {
        "Content-Type": "text/html; charset=utf-8"
    },
    "body": "<h1>Hello world!</h1>",
    "isBase64Encoded": false
}
```

Status Code & Description

Headers as Key/Value pairs

Body & isBase64Encoded

© Copyright by Cloud Mentor Pro | Confidential

ALB Multi-Header Values

- ALB can support multi header values (ALB setting)
- When you enable multi-value headers, **HTTP headers and query string parameters** that are sent with **multiple values** are shown **as arrays** within the AWS Lambda event and response objects.

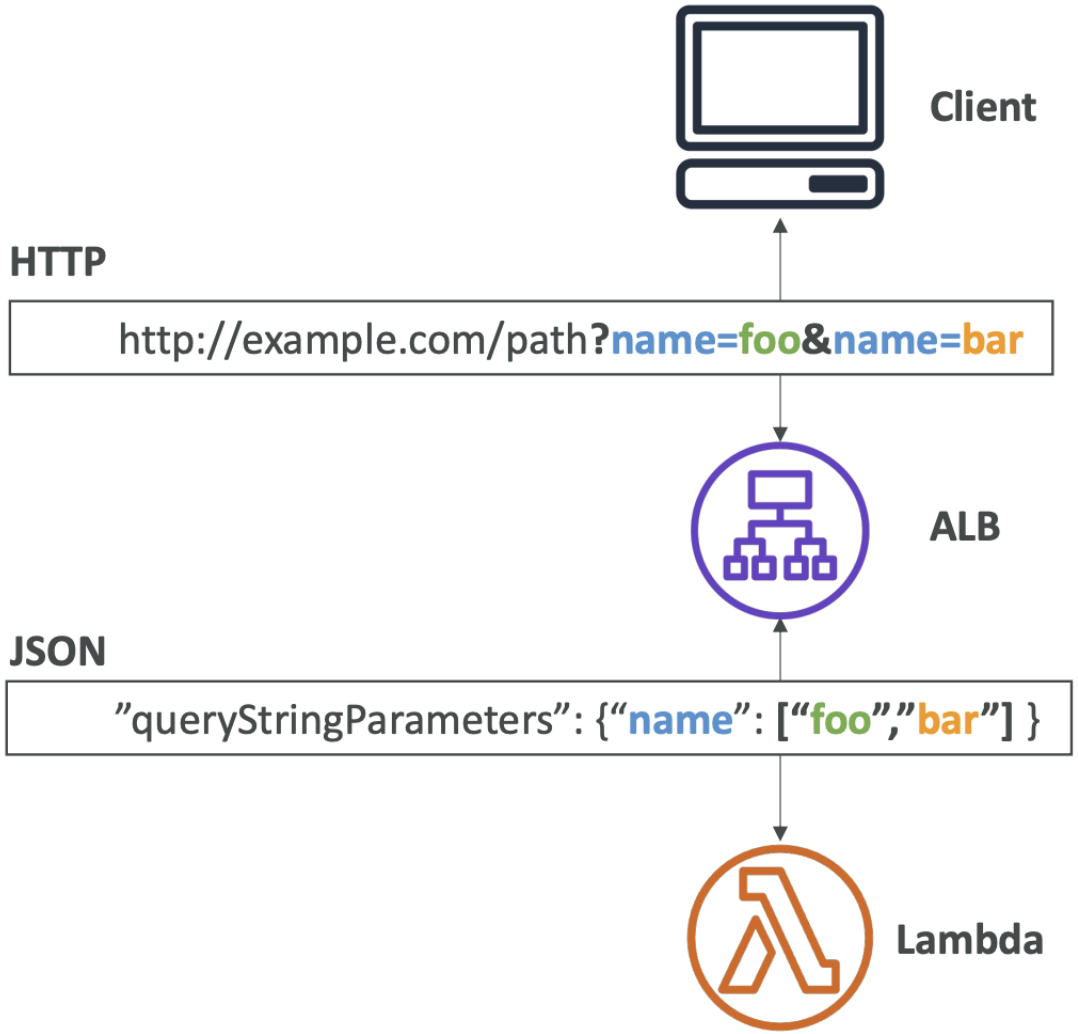

© Copyright by Cloud Mentor Pro | Confidential

ALB + Lambda – Permissions

```
{
    "Statement": {
        "Effect": "Allow",
        "Principal": {
            "Service": "elasticloadbalancing.amazonaws.com"
        },
        "Action": "lambda:InvokeFunction",
        "Resource": "arn:aws:lambda:us-west-2:123456789012:function:alb-function"
    }
}
```

Lambda Resource Policy

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Asynchronous Invocations

- S3, SNS, CloudWatch Events…
- The events are placed in an **Event Queue**
- Lambda attempts to retry on errors
  - 3 tries total
  - 1 minute wait after 1st , then 2 minutes wait
- Make sure the processing is **idempotent** (in case of retries)
- If the function is retried, you will see **duplicate logs entries in CloudWatch Logs**
- Can define a DLQ (dead-letter queue) – **SNS or SQS** – for failed processing (need correct IAM permissions)
- Asynchronous invocations allow you to speed up the processing if you don't need to wait for the result (ex: you need 1000 files processed)

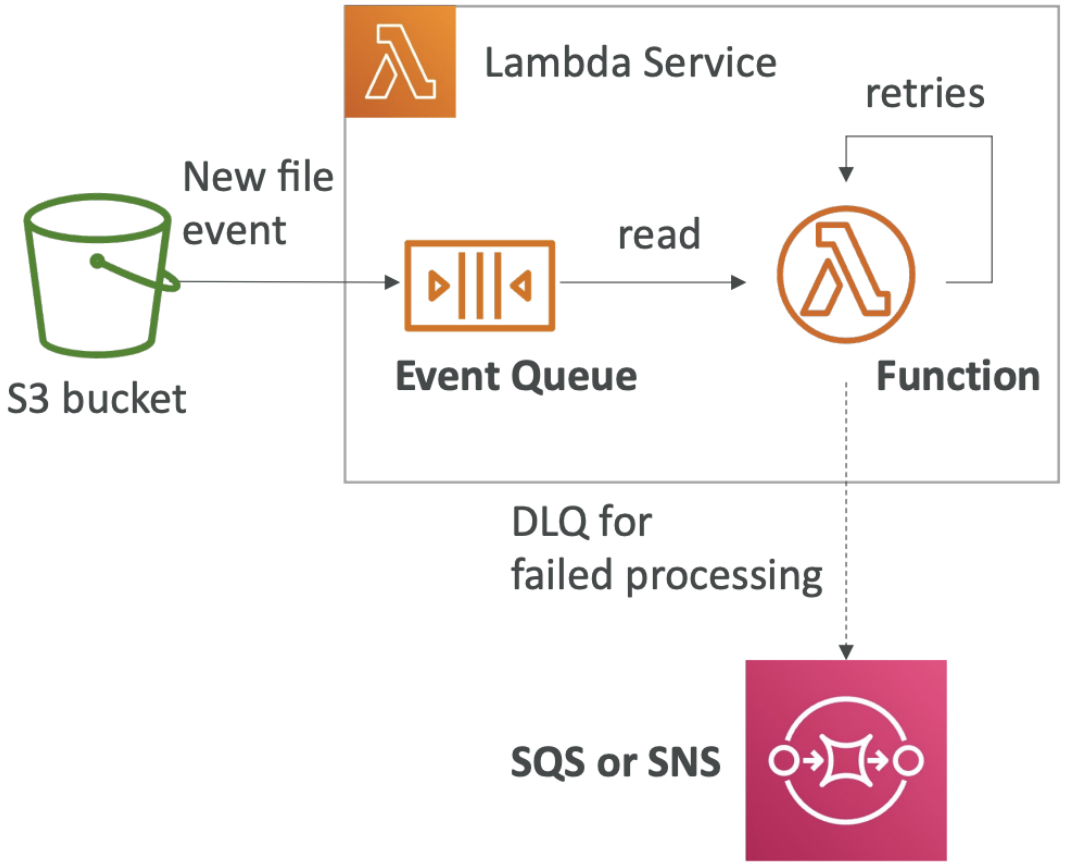

© Copyright by Cloud Mentor Pro | Confidential

Lambda - Asynchronous Invocations - Services

- Amazon Simple Storage Service (S3)
- Amazon Simple Notification Service (SNS)
- Amazon CloudWatch Events / EventBridge
- AWS CodeCommit (CodeCommitTrigger: new branch, new tag, new push)
- AWS CodePipeline (invoke a Lambda function during the pipeline, Lambda must callback)

----- other -----

- Amazon CloudWatch Logs (log processing)
- Amazon Simple Email Service
- AWS CloudFormation
- AWS Config
- AWS IoT
- AWS IoT Events

© Copyright by Cloud Mentor Pro | Confidential

CloudWatch Events / EventBridge

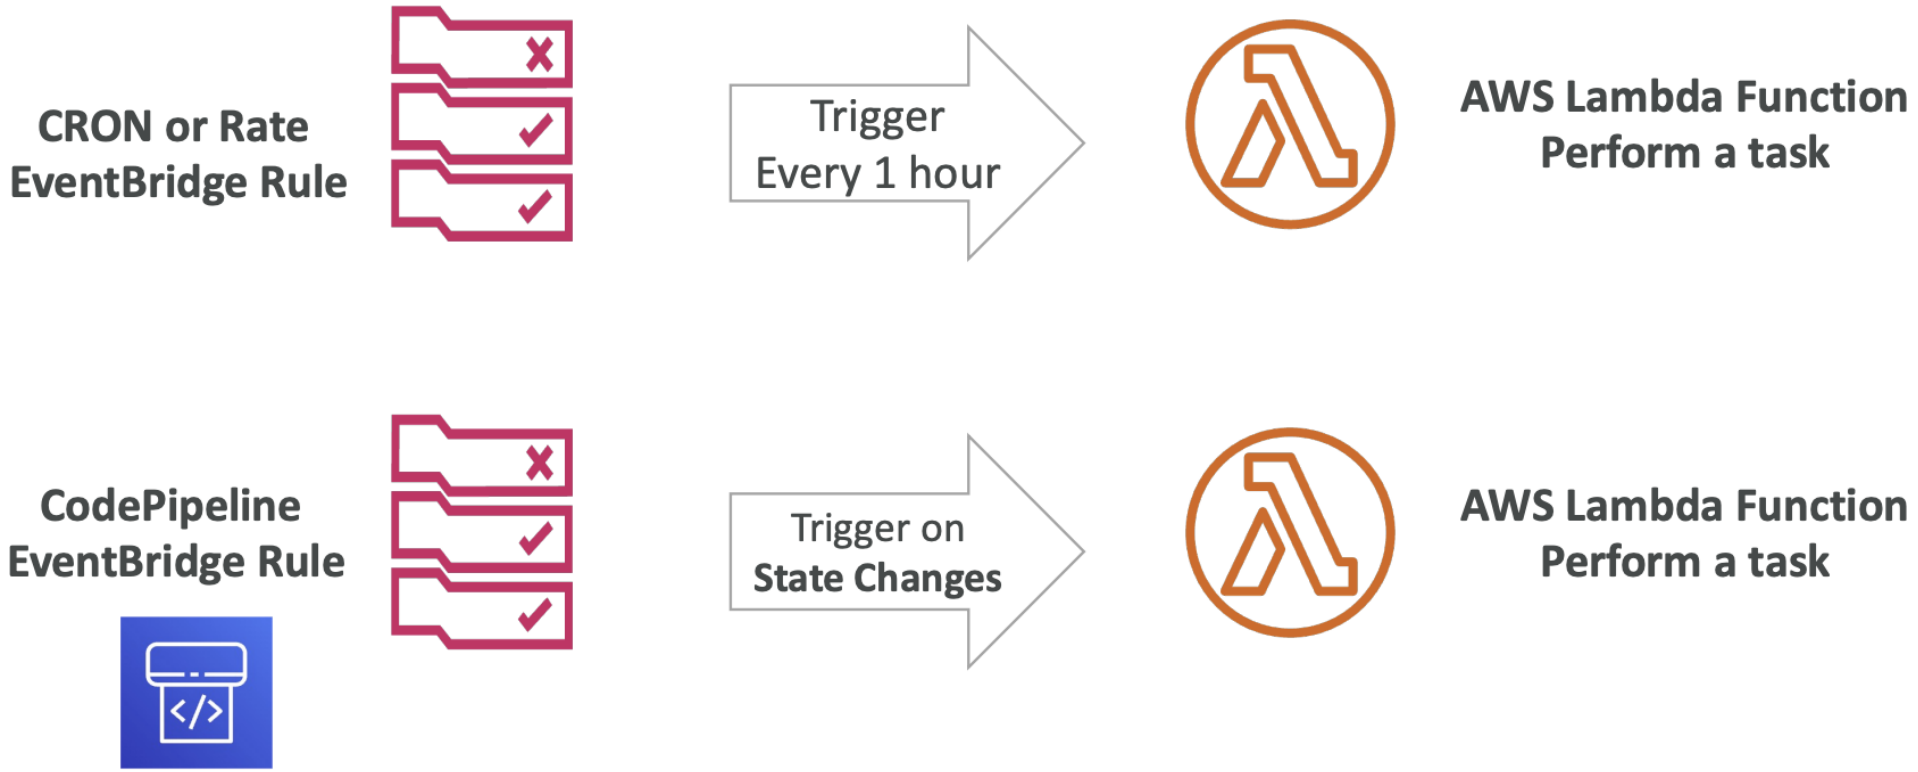

© Copyright by Cloud Mentor Pro | Confidential

S3 Events Notifications

- S3:ObjectCreated, S3:ObjectRemoved, S3:ObjectRestore, S3:Replication…
- Object name filtering possible (*.jpg)
- Use case: generate thumbnails of images uploaded to S3

- S3 event notifications typically deliver events in seconds but can sometimes take a minute or longer
- If two writes are made to a single non- versioned object at the same time, it is possible that only a single event notification will be sent
- If you want to ensure that an event notification is sent for every successful write, you can enable versioning on your bucket.

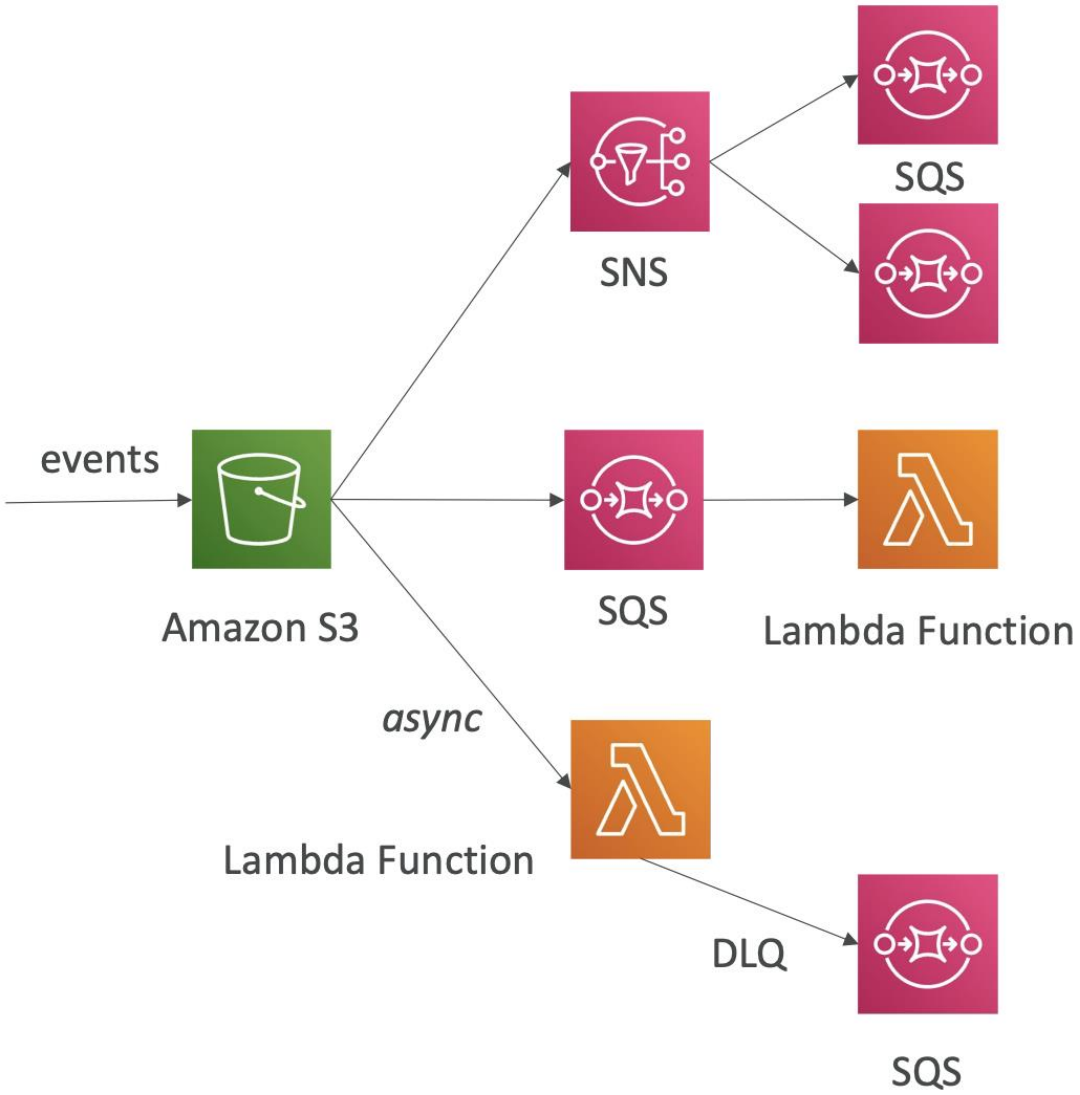

© Copyright by Cloud Mentor Pro | Confidential

Simple S3 Event Pattern – Metadata Sync

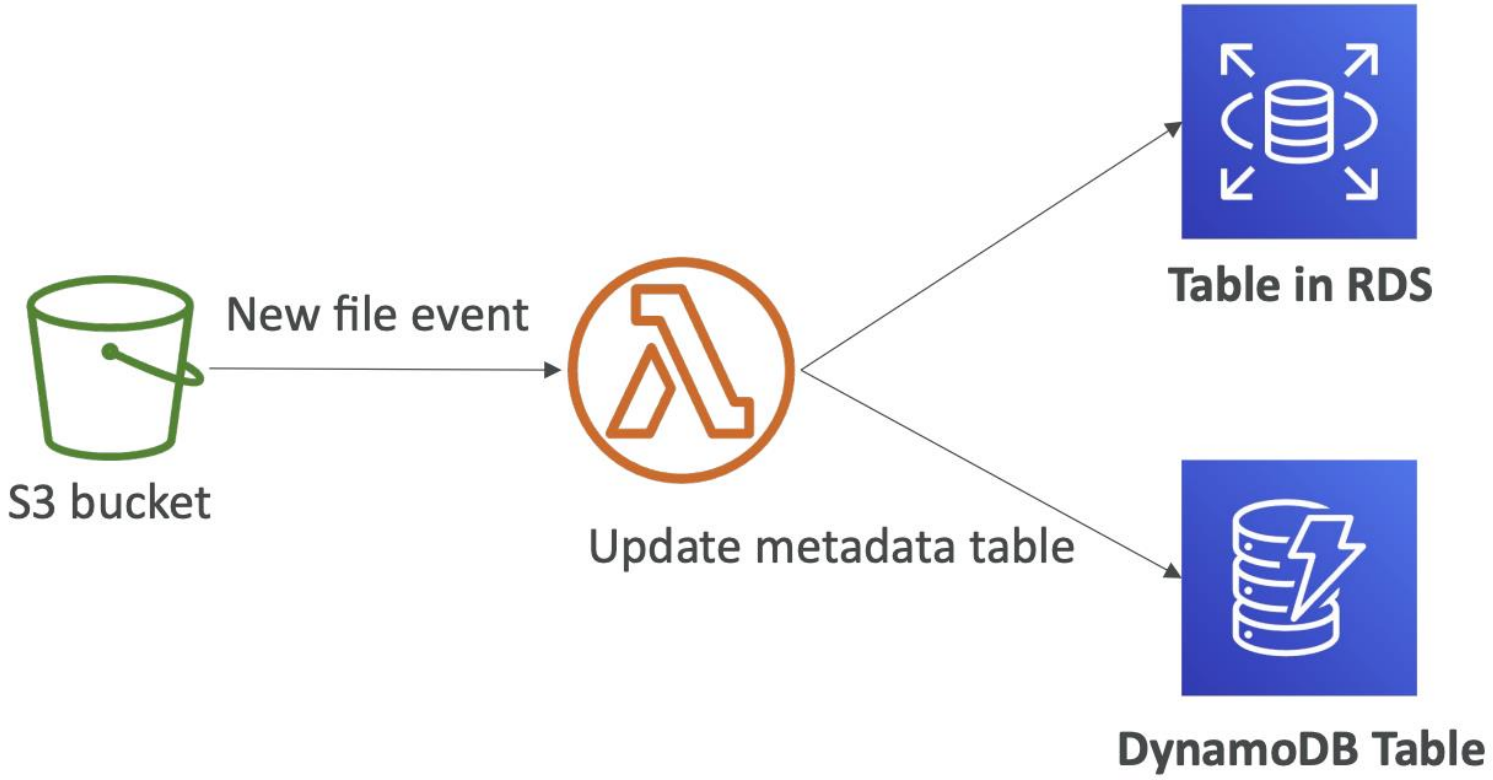

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Event Source Mapping

- Kinesis Data Streams
- SQS & SQS FIFO queue
- DynamoDB Streams

- Common denominator: records need to be polled from the source
- **Your Lambda function is invoked synchronously**

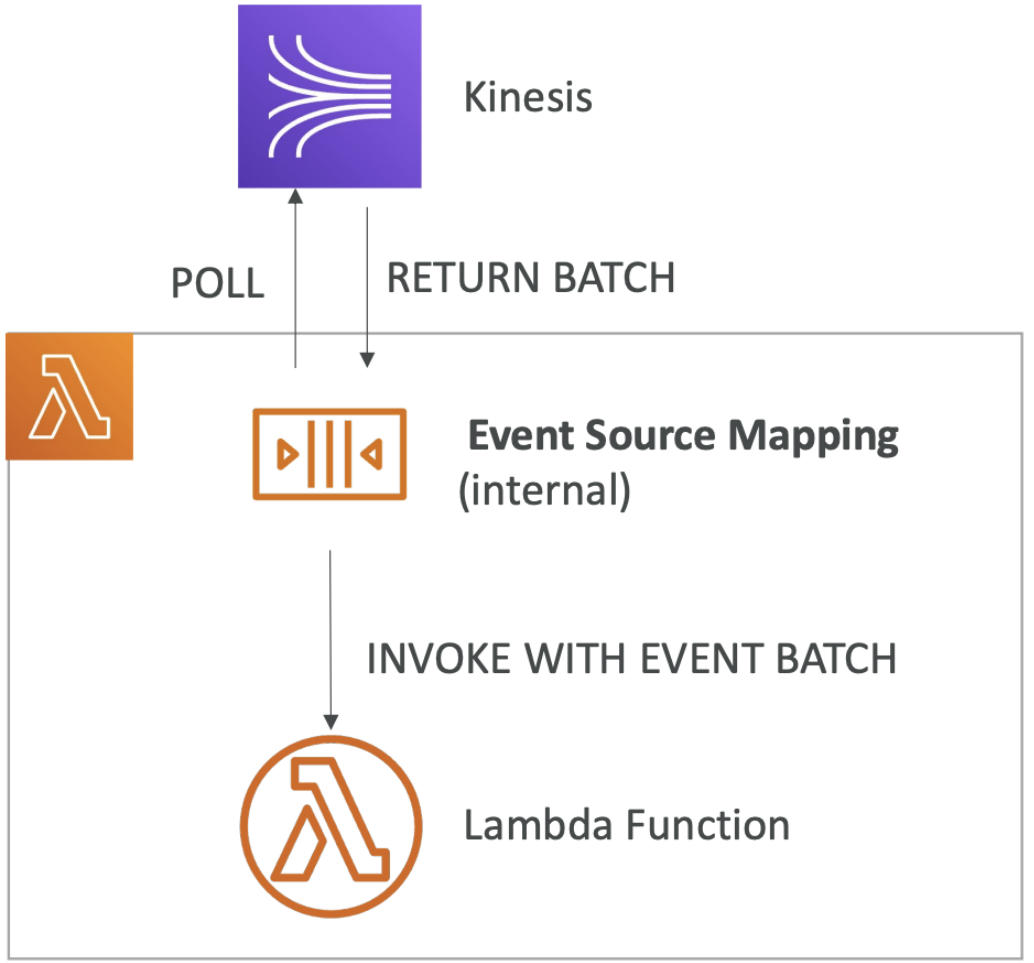

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Event and Context Objects

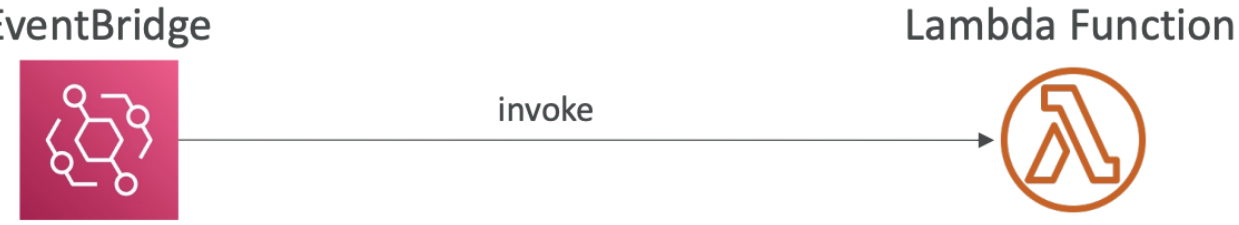


```
{
    "version": "0",
    "id": "fe8d3c65-xmpl-c5c3-2c87-81584709a377",
    "detail-type": "RDS DB Instance Event",
    "source": "aws.rds",
    "account": "123456789012",
    "time": "2020-04-28T07:20:20Z",
    "region": "us-east-2",
    "resources": [
        "arn:aws:rds:us-east-2:123456789012:db:rdz6xmpliljlb1"
    ],
    "detail": {
        "EventCategories": [ "backup" ],
        "SourceType": "DB_INSTANCE",
        "SourceArn": "arn:aws:rds:us-east-2:123456789012:db:rdz6xmpliljlb1",
        "Date": "2020-04-28T07:20:20.112Z",
        "Message": "Finished DB Instance backup",
        "SourceIdentifier": "rdz6xmpliljlb1"
    }
}
```

Event Object

```
{
    "aws_request_id": "a3de505e-f16b-42f4-b3e6-bcd2e4a73903",
    "function_name": "MyFunction-AULC3LB8Q02F",
    "invoked_function_arn": "arn:aws:lambda:us-west-2:123456789012:function:
MyFunction-AULC3LB8Q02F",
    "function_version": "$LATEST",
    "memory_limit_in_mb": "128",
    "log_group_name": "/aws/lambda/MyFunction-AULC3LB8Q02F",
    "log_stream_name": "2015/10/26/[$LATEST]c71058d852474b980f221f73402ad",
    "client_context": "None",
    ...
}
```

Context Object

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Event and Context Objects

- Event Object
  - JSON-formatted document contains data for the function to process
  - Contains information from the invoking service (e.g., EventBridge, custom, …)
  - Lambda runtime converts the event to an object (e.g., ***dict*** type in Python)
  - Example: input arguments, invoking service arguments, …

- Context Object
  - Provides methods and properties that provide information about the invocation, function, and runtime environment
  - Passed to your function by Lambda at runtime
  - Example: aws_request_id, function_name, memory_limit_in_mb, …

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Event and Context Objects

Access Event & Context Objects using Python

```
def lambda_handler(event, context):
    print("Event Source:", event.source)
    print("Event Region:", event.region)


    print("Lambda Request ID:", context.aws_request_id)
    print("Lambda function ARN:", context.function_name)
    print("Lambda function ARN:", context.invoked_function_arn)
    print("Lambda function memory limits in MB:", context.memory_limit_in_mb)
    print("CloudWatch log stream name:", context.log_stream_name)
    print("CloudWatch log group name:",  context.log_group_name)
```

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Destinations

- Nov 2019: Can configure to send result to a **destination**
- **Asynchronous invocations** - can define destinations for successful and failed event:
  - Amazon SQS
  - Amazon SNS
  - AWS Lambda
  - Amazon EventBridge bus
- *Note: AWS recommends you use destinations instead of DLQ now (but both can be used at the same time)*
- **Event Source mapping**: for discarded event batches
  - Amazon SQS
  - Amazon SNS
- *Note: you can send events to a DLQ directly from SQS*

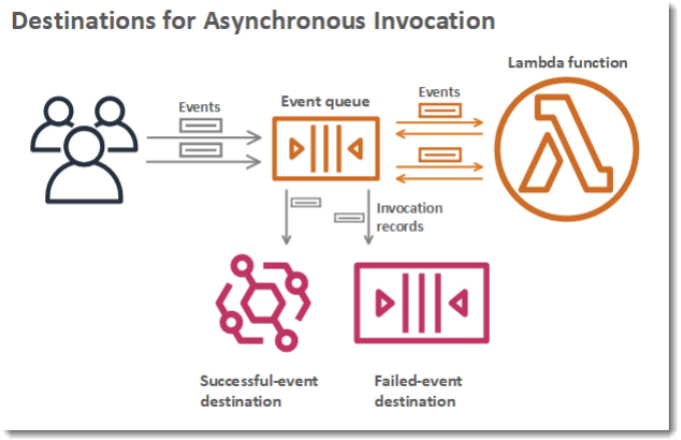

https://docs.aws.amazon.com/lambda/latest/dg/invocation-async.html

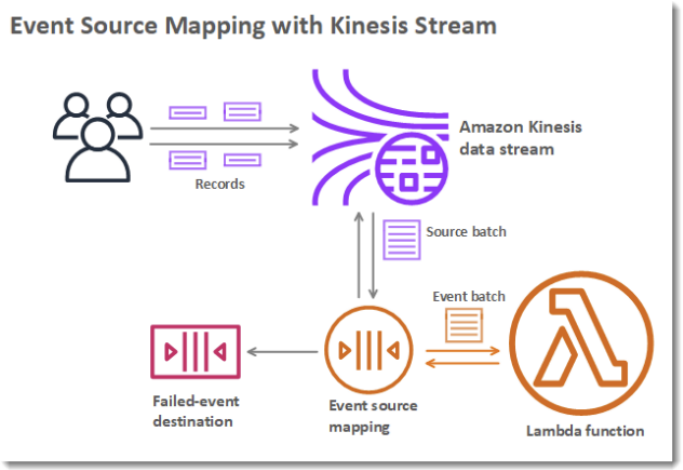

https://docs.aws.amazon.com/lambda/latest/dg/invocation-eventsourcemapping.html

© Copyright by Cloud Mentor Pro | Confidential

Lambda Execution Role (IAM Role)

- Grants the Lambda function permissions to AWS services / resources
- Sample managed policies for Lambda:
  - AWSLambdaBasicExecutionRole – Upload logs to CloudWatch.
  - AWSLambdaKinesisExecutionRole – Read from Kinesis
  - AWSLambdaDynamoDBExecutionRole – Read from DynamoDB Streams
  - AWSLambdaSQSQueueExecutionRole – Read from SQS
  - AWSLambdaVPCAccessExecutionRole – Deploy Lambda function in VPC
  - AWSXRayDaemonWriteAccess – Upload trace data to X-Ray.

- **When you use an event source mapping to invoke your function, Lambda uses the execution role to read event data.**
- Best practice: create one Lambda Execution Role per function

© Copyright by Cloud Mentor Pro | Confidential

Lambda Resource Based Policies

- Use resource-based policies to give other accounts and AWS services permission to use your Lambda resources
- Similar to S3 bucket policies for S3 bucket
- An IAM principal can access Lambda:
  - if the IAM policy attached to the principal authorizes it (e.g. user access)
  - OR if the resource-based policy authorizes (e.g. service access)

- When an AWS service like Amazon S3 calls your Lambda function, the resource-based policy gives it access.

© Copyright by Cloud Mentor Pro | Confidential

Lambda Environment Variables

- Environment variable = key / value pair in “String” form
- Adjust the function behavior without updating code
- The environment variables are available to your code
- Lambda Service adds its own system environment variables as well

- Helpful to store secrets (encrypted by KMS)
- Secrets can be encrypted by the Lambda service key, or your own CMK

© Copyright by Cloud Mentor Pro | Confidential

Lambda Logging & Monitoring

- CloudWatch Logs:
  - AWS Lambda execution logs are stored in AWS CloudWatch Logs
  - **Make sure your AWS Lambda function has an execution role with an IAM policy that authorizes writes to CloudWatch Logs**
- CloudWatch Metrics:
  - AWS Lambda metrics are displayed in AWS CloudWatch Metrics
  - Invocations, Durations, Concurrent Executions
  - Error count, Success Rates, Throttles
  - Async Delivery Failures
  - Iterator Age (Kinesis & DynamoDB Streams)

© Copyright by Cloud Mentor Pro | Confidential

Lambda Tracing with X-Ray

- Enable in Lambda configuration (**Active Tracing**)
- Runs the X-Ray daemon for you
- Use AWS X-Ray SDK in Code
- Ensure Lambda Function has a correct IAM Execution Role
  - The managed policy is called AWSXRayDaemonWriteAccess
- Environment variables to communicate with X-Ray
  - **_X_AMZN_TRACE_ID**: contains the tracing header
  - **AWS_XRAY_CONTEXT_MISSING**: by default, LOG_ERROR
  - **AWS_XRAY_DAEMON_ADDRESS**: the X-Ray Daemon IP_ADDRESS:PORT

© Copyright by Cloud Mentor Pro | Confidential

Customization At The Edge

- Many modern applications execute some form of the logic at the edge
- **Edge Function**:
  - A code that you write and attach to CloudFront distributions
  - Runs close to your users to minimize latency
- CloudFront provides two types: **CloudFront Functions** & **Lambda@Edge**
- You don’t have to manage any servers, deployed globally

- Use case: customize the CDN content
- Pay only for what you use
- Fully serverless

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Functions & Lambda@Edge Use Cases

- Website Security and Privacy
- Dynamic Web Application at the Edge
- Search Engine Optimization (SEO)
- Intelligently Route Across Origins and Data Centers
- Bot Mitigation at the Edge
- Real-time Image Transformation
- A/B Testing
- User Authentication and Authorization
- User Prioritization
- User Tracking and Analytics

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Functions

- Lightweight functions written in JavaScript
- For high-scale, latency-sensitive CDN customizations
- Sub-ms startup times, **millions of requests/second**
- Used to change Viewer requests and responses:
  - **Viewer Request**: after CloudFront receives a request from a viewer
  - **Viewer Response**: before CloudFront forwards the response to the viewer
- Native feature of CloudFront (manage code entirely within CloudFront)

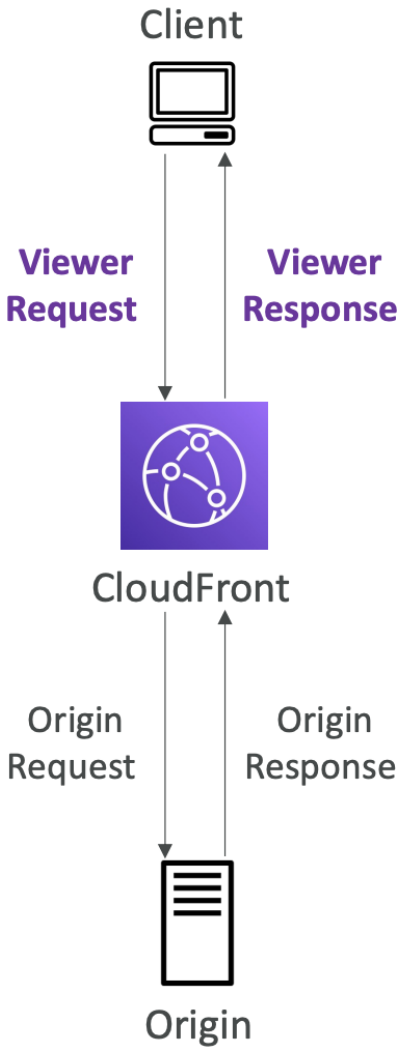

© Copyright by Cloud Mentor Pro | Confidential

Lambda@Edge

- Lambda functions written in NodeJS or Python
- Scales to **1000s of requests/second**
- Used to change CloudFront requests and responses:
  - **Viewer Request** – after CloudFront receives a request from a viewer
  - **Origin Request** – before CloudFront forwards the request to the origin
  - **Origin Response** – after CloudFront receives the response from the origin
  - **Viewer Response** – before CloudFront forwards the response to the viewer
- Author your functions in one AWS Region (us-east-1), then CloudFront replicates to its locations

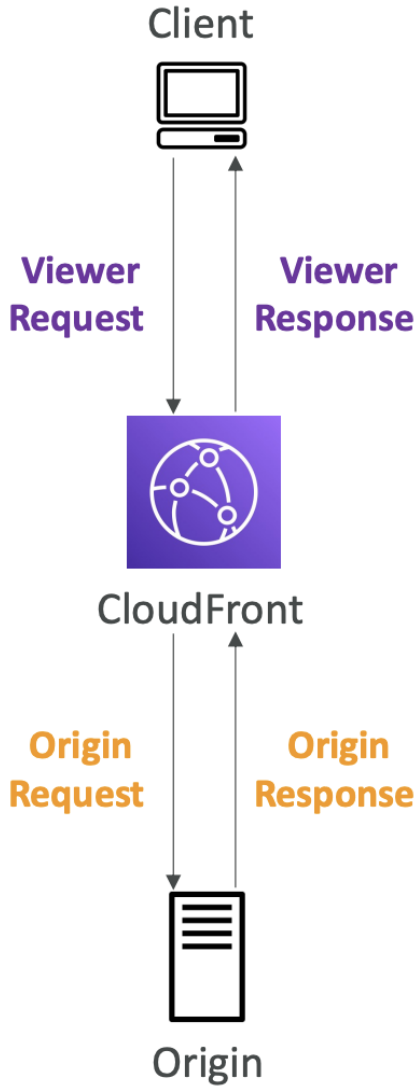

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Functions vs. Lambda@Edge

	CloudFront Functions	Lambda@Edge
Runtime Support	JavaScript	Node.js, Python
# of Requests	**Millions** of requests per second	**Thousands** of requests per second
CloudFront Triggers	- Viewer Request/Response	- Viewer Request/Response - Origin Request/Response
Max. Execution Time	< 1 ms	5 – 10 seconds
Max. Memory	2 MB	128 MB up to 10 GB
Total Package Size	10 KB	1 MB – 50 MB
Network Access, File System Access	No	Yes
Access to the Request Body	No	Yes
Pricing	Free tier available, 1/6th price of @Edge	No free tier, charged per request & duration

© Copyright by Cloud Mentor Pro | Confidential

CloudFront Functions vs. Lambda@Edge - Use Cases

CloudFront Functions

- Cache key normalization
  - Transform request attributes (headers, cookies, query strings, URL) to create an optimal Cache Key
- Header manipulation
  - Insert/modify/delete HTTP headers in the request or response
- URL rewrites or redirects
- Request authentication & authorization
  - Create and validate user-generated tokens (e.g., JWT) to allow/deny requests

Lambda@Edge

- Longer execution time (several ms)
- Adjustable CPU or memory
- Your code depends on a 3rd libraries (e.g., AWS SDK to access other AWS services)
- Network access to use external services for processing
- File system access or access to the body of HTTP requests

© Copyright by Cloud Mentor Pro | Confidential

Lambda by default

- By default, your Lambda function is launched outside your own VPC (in an AWS -owned VPC)
- Therefore it cannot access resources in your VPC (RDS, ElastiCache, internal ELB…)

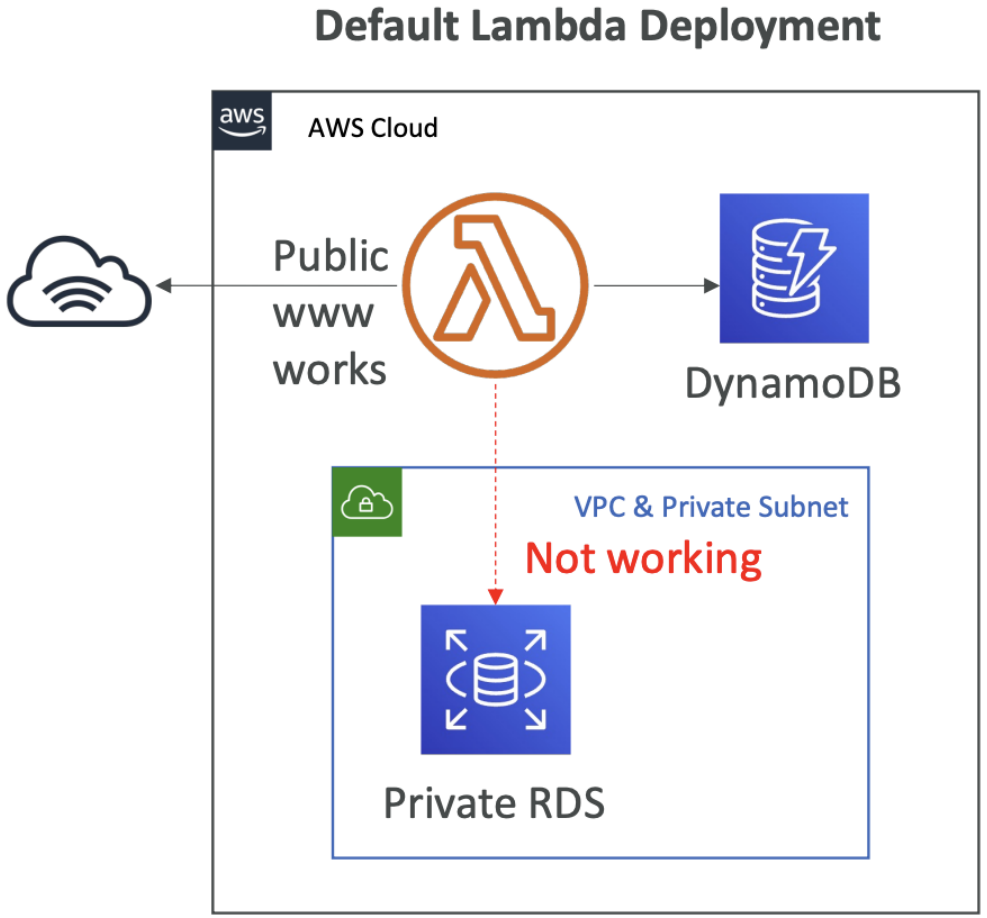

© Copyright by Cloud Mentor Pro | Confidential

Lambda in VPC

- You must define the VPC ID, the Subnets and the Security Groups
- Lambda will create an ENI (Elastic Network Interface) in your subnets
- **AWSLambdaVPCAccessExecutionRole**

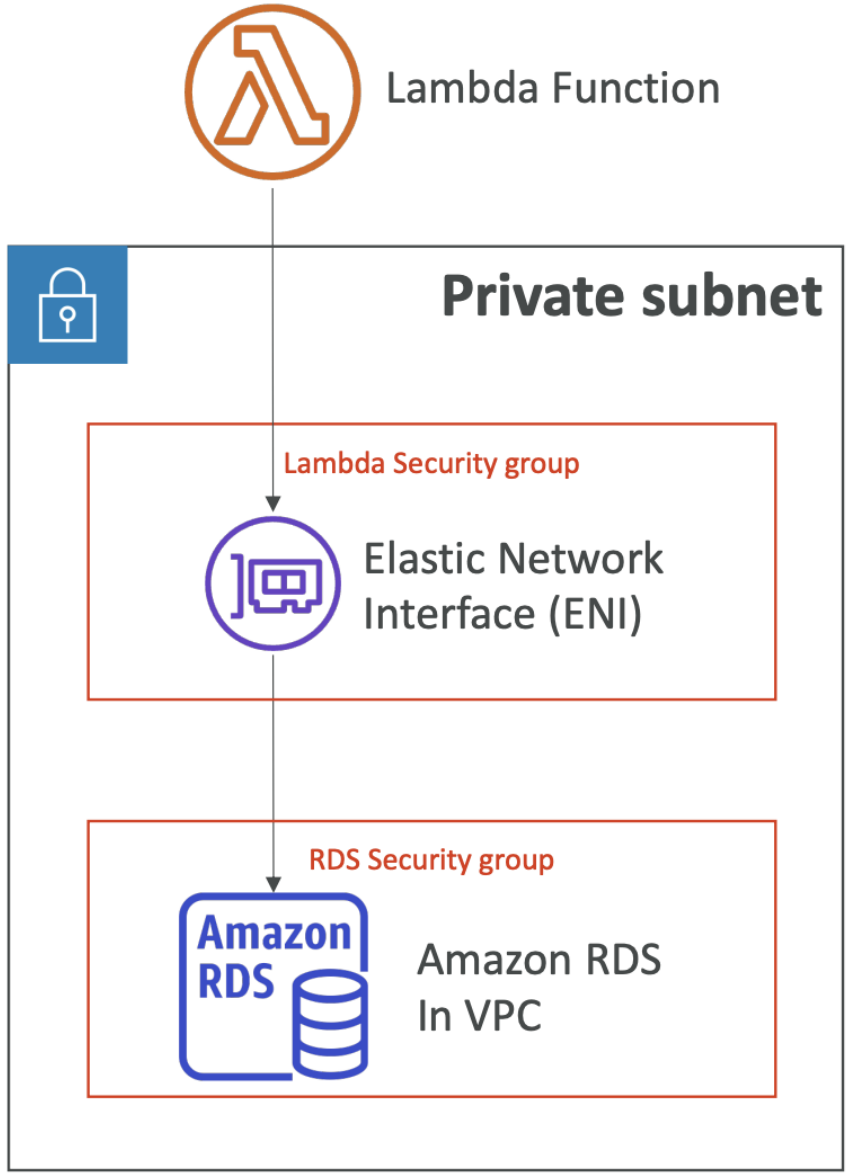

© Copyright by Cloud Mentor Pro | Confidential

Lambda in VPC – Internet Access

- A Lambda function in your VPC does not have internet access
- **Deploying a Lambda function in a public subnet does not give it internet access or a public IP**
- Deploying a Lambda function in a private subnet gives it internet access if you have a **NAT Gateway / Instance**
- You can use **VPC endpoints** to privately access AWS services without a NAT

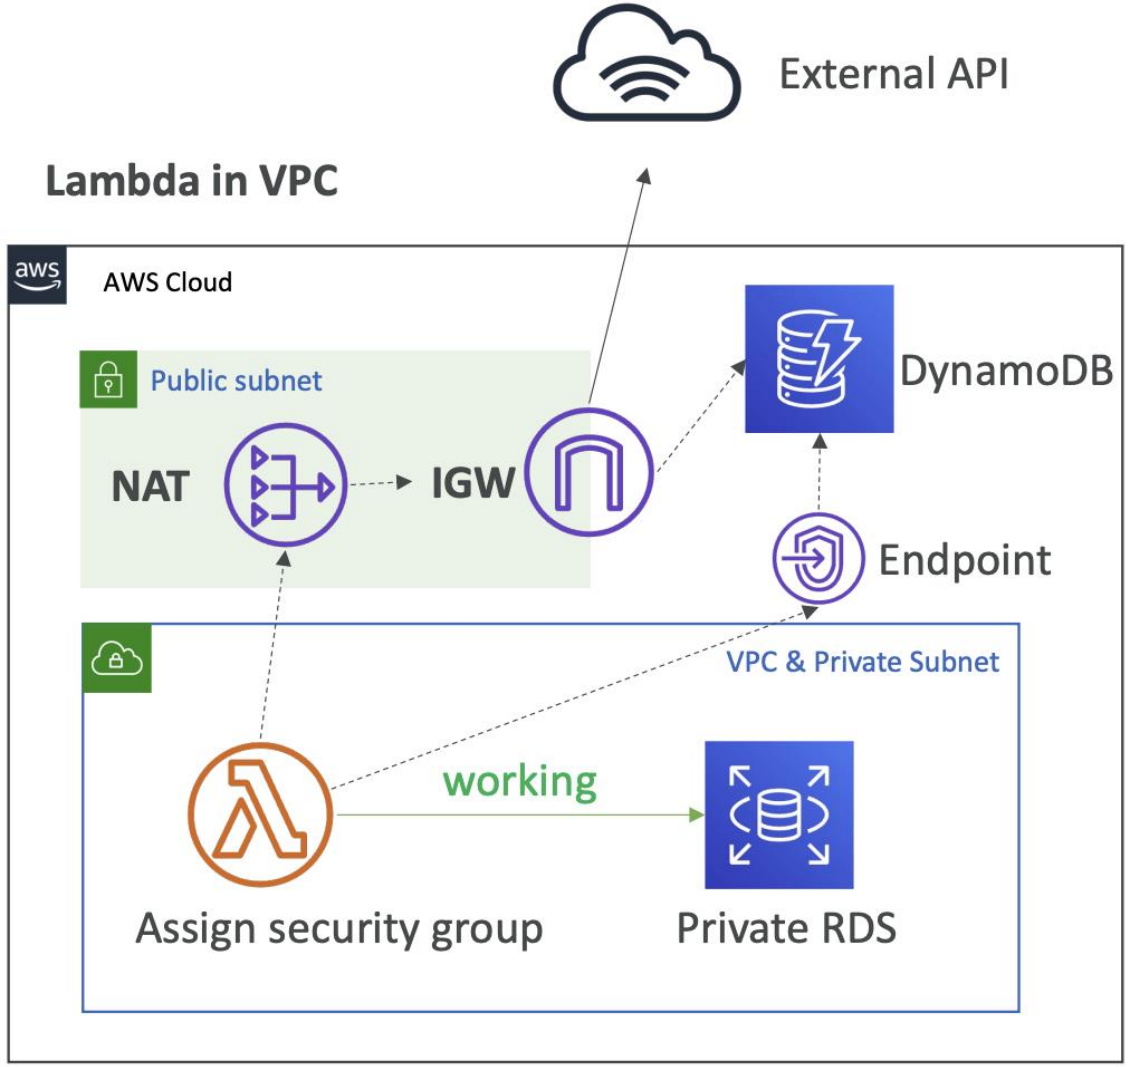

Note: Lambda - CloudWatch Logs works even without endpoint or NAT Gateway

© Copyright by Cloud Mentor Pro | Confidential

Lambda Function Configuration

- **RAM**:
  - From 128MB to 10GB in 1MB increments
  - The more RAM you add, the more vCPU credits you get
  - At 1,792 MB, a function has the equivalent of one full vCPU
  - After 1,792 MB, you get more than one CPU, and need to use multi-threading in your code to benefit from it (up to 6 vCPU)
- **If your application is CPU-bound (computation heavy), increase RAM**

- **Timeout**: default 3 seconds, maximum is 900 seconds (15 minutes)

© Copyright by Cloud Mentor Pro | Confidential

Lambda Execution Context

- The execution context is a temporary runtime environment that initializes any external dependencies of your lambda code
- Great for database connections, HTTP clients, SDK clients…
- The execution context is maintained for some time in anticipation of another Lambda function invocation
- The next function invocation can “re-use” the context to execution time and save time in initializing connections objects
- The execution context includes the /tmp directory

© Copyright by Cloud Mentor Pro | Confidential

Initialize outside the handler

BAD!

```
import os

def get_user_handler(event, context):

    DB_URL = os.getenv("DB_URL")
    db_client = db.connect(DB_URL)
    user = db_client.get(user_id = event["user_id"])

    return user
```

The DB connection is established
At every function invocation

GOOD!

```
import os

DB_URL = os.getenv("DB_URL")
db_client = db.connect(DB_URL)

def get_user_handler(event, context):

    user = db_client.get(user_id = event["user_id"])

    return user
```

The DB connection is established once
And re-used across invocations

© Copyright by Cloud Mentor Pro | Confidential

Lambda Functions /tmp space

- If your Lambda function needs to download a big file to work…
- If your Lambda function needs disk space to perform operations…
- You can use the /tmp directory
- Max size is 10GB
- The directory content remains when the execution context is frozen, providing transient cache that can be used for multiple invocations (helpful to checkpoint your work)
- For permanent persistence of object (non temporary), use S3
- To encrypt content on /tmp, you must generate KMS Data Keys

© Copyright by Cloud Mentor Pro | Confidential

Lambda Layers

- Custom Runtimes
  - Ex: C++ https://github.com/awslabs/aws-lambda-cpp
  - Ex: Rust https://github.com/awslabs/aws-lambda-rust-runtime
- Externalize Dependencies to re-use them:

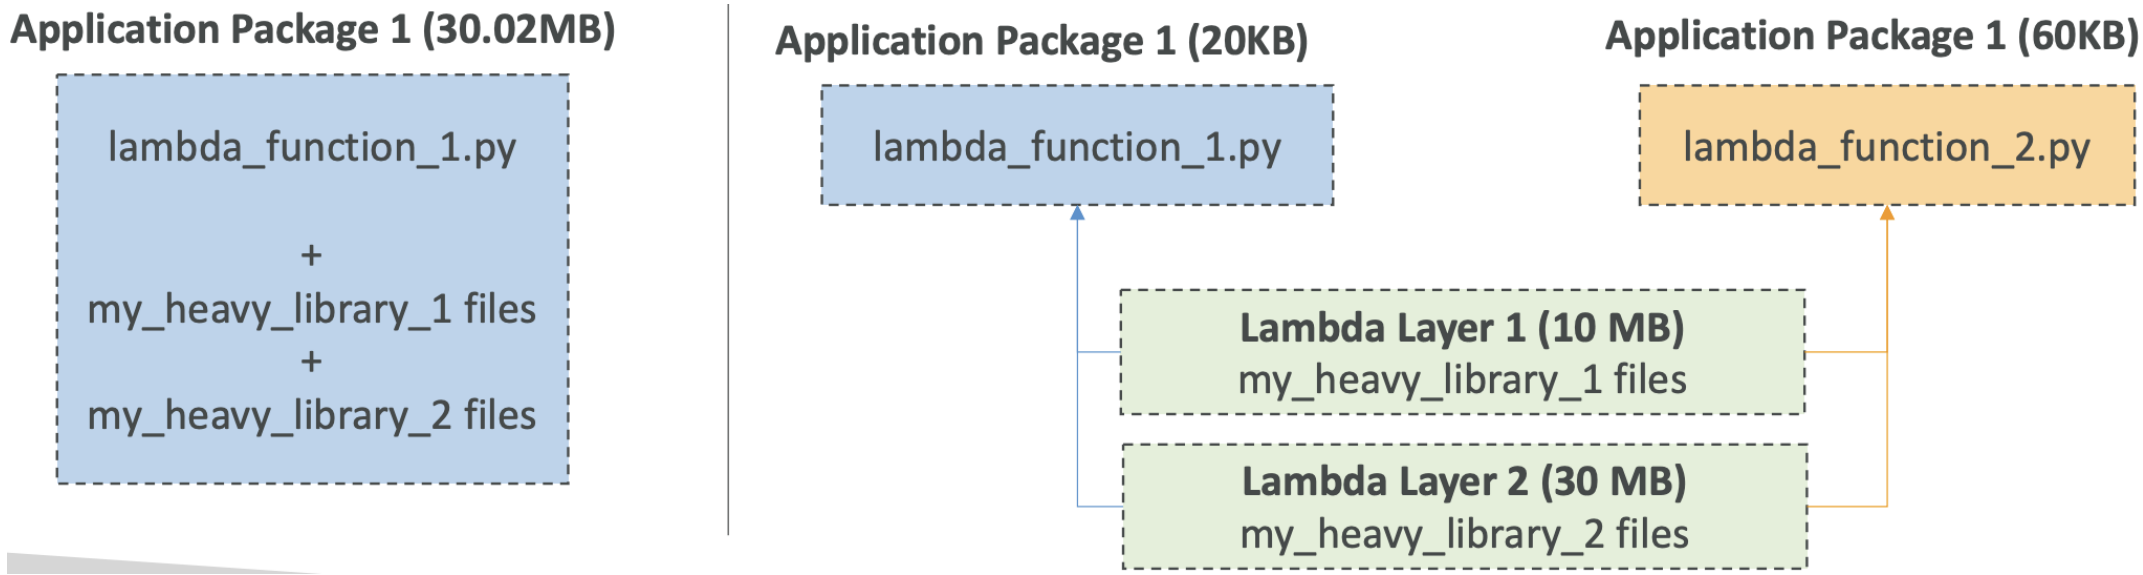

© Copyright by Cloud Mentor Pro | Confidential

Lambda – File Systems Mounting

- Lambda functions can access EFS file systems if they are running in a VPC
- Configure Lambda to mount EFS file systems to local directory during initialization
- Must leverage EFS Access Points
- Limitations: watch out for the EFS connection limits (one function instance = one connection) and connection burst limits

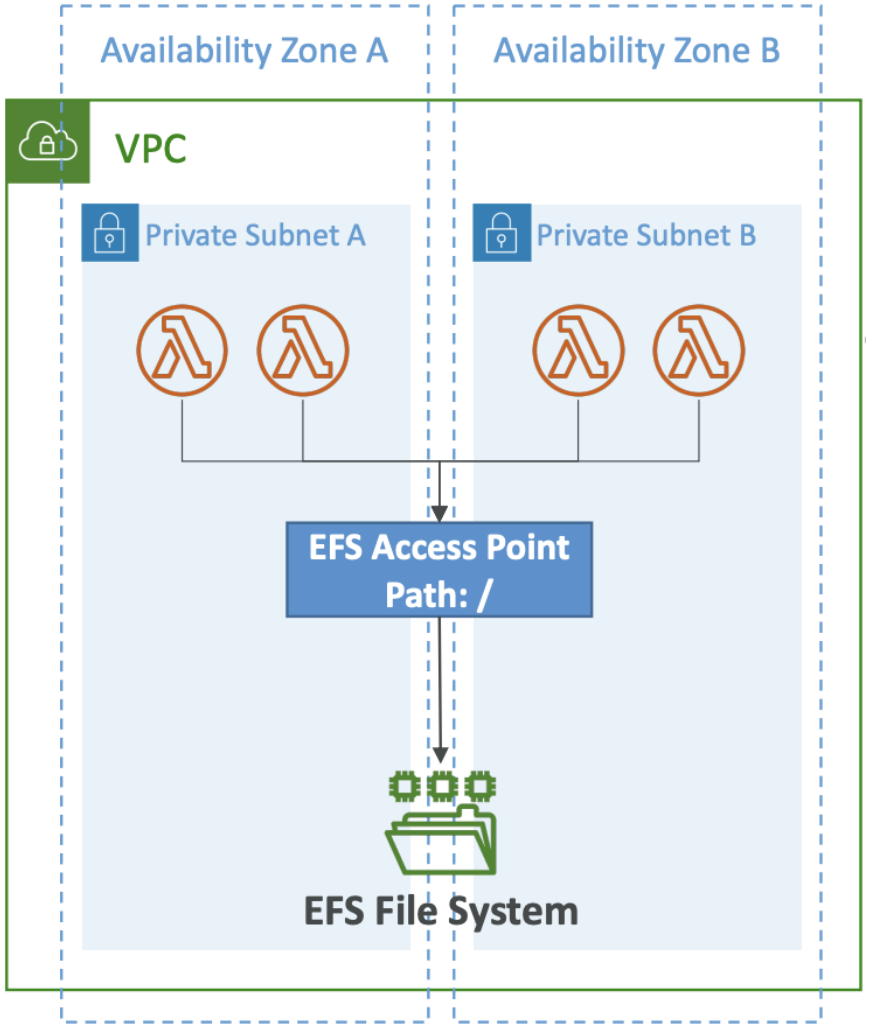

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Storage Options

	Ephemeral Storage /tmp	Lambda Layers	Amazon S3	Amazon EFS
Max. Size	10,240 MB	5 layers per function up to 250MB total	Elastic	Elastic
Persistence	Ephemeral	Durable	Durable	Durable
Content	Dynamic	Static	Dynamic	Dynamic
Storage Type	File System	Archive	Object	File System
Operations supported	any File System operation	Immutable	Atomic with Versioning	any File System operation
Pricing	Included in Lambda	Included in Lambda	Storage + Requests + Data Transfer	Storage + Data Transfer + Throughput
Sharing/Permissions	Function Only	IAM	IAM	IAM + NFS
Relative Data Access Speed from Lambda	Fastest	Fastest	Fast	Very Fast
Shared Across All Invocations	No	Yes	Yes	Yes

© Copyright by Cloud Mentor Pro | Confidential

Lambda Concurrency and Throttling

- Concurrency limit: up to 1000 concurrent executions

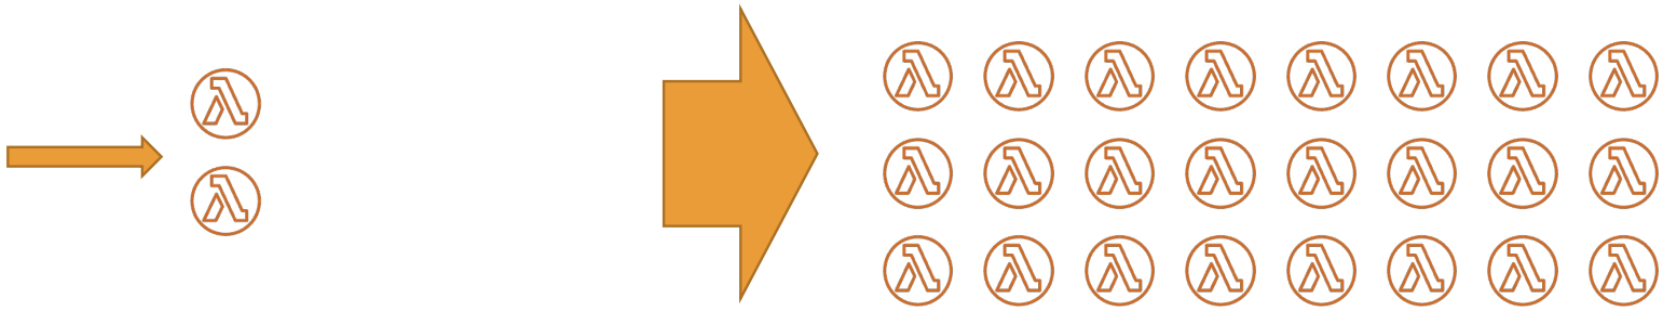

- Can set a "**reserved concurrency**" at the function level (=limit)
- Each invocation over the concurrency limit will trigger a "Throttle"
- Throttle behavior:
  - If synchronous invocation => return ThrottleError - 429
  - If asynchronous invocation => retry automatically and then go to DLQ
- If you need a higher limit, open a support ticket

© Copyright by Cloud Mentor Pro | Confidential

Lambda Concurrency Issue

- If you don’t reserve (=limit) concurrency, the following can happen:

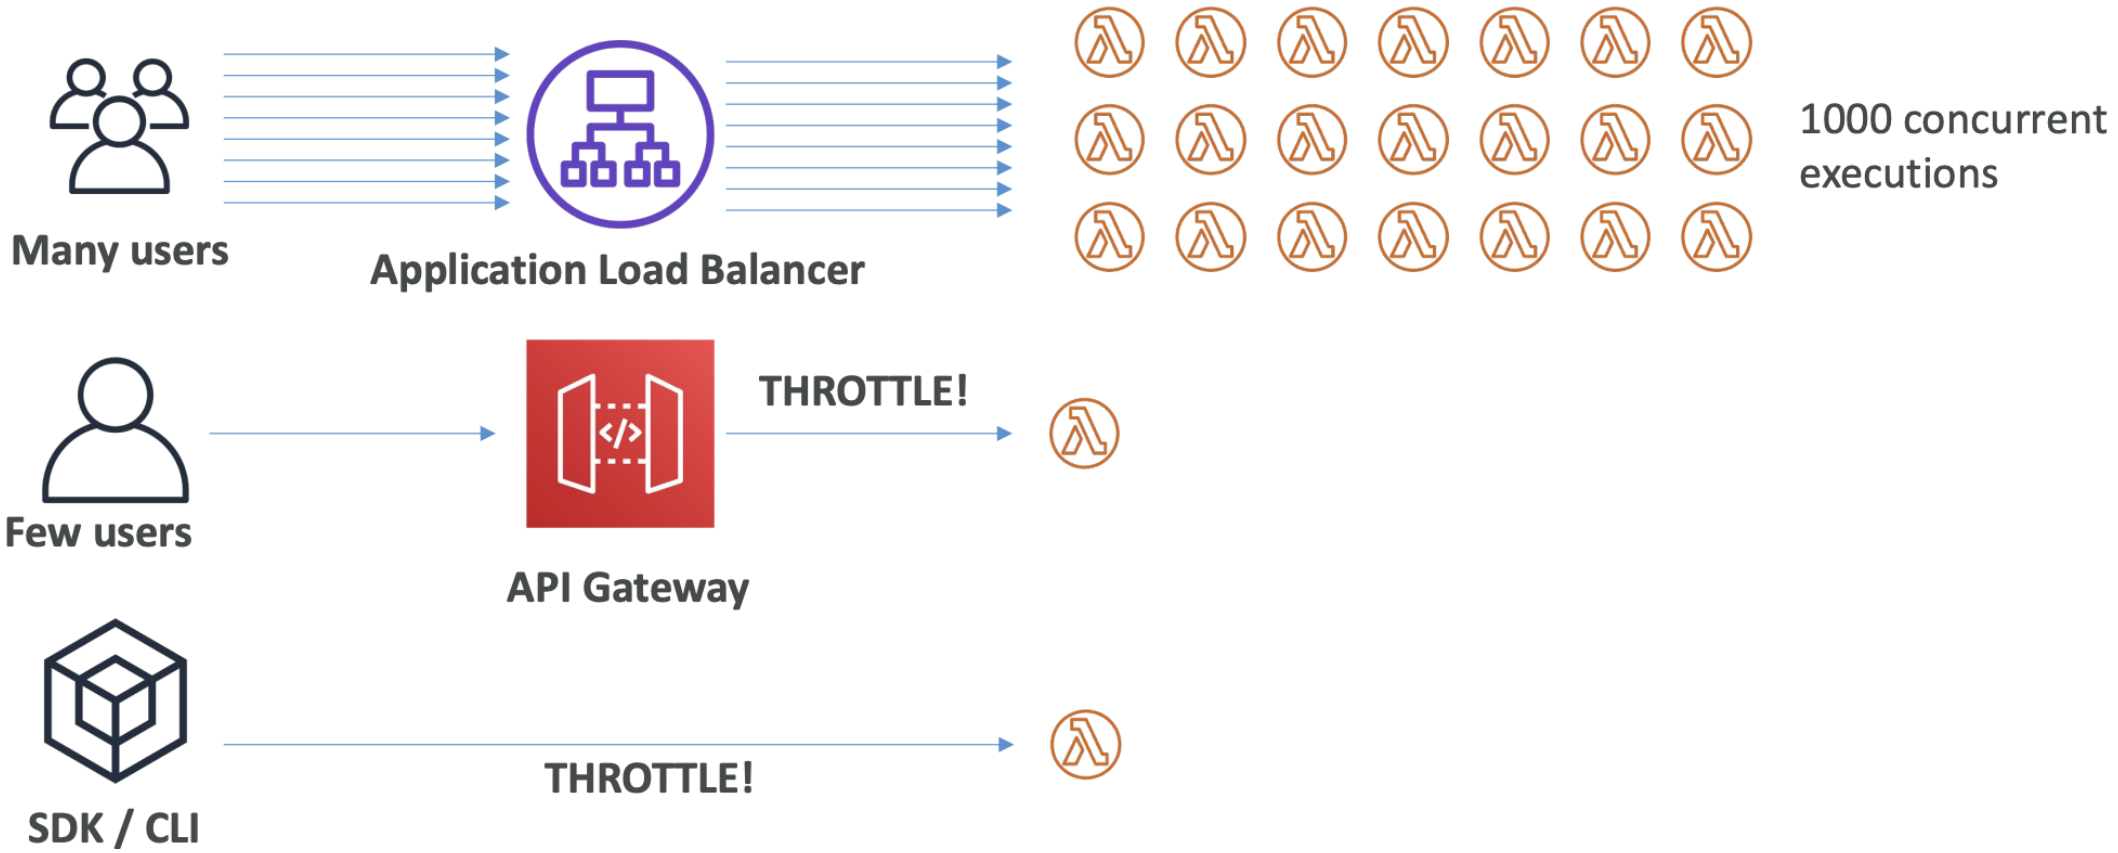

© Copyright by Cloud Mentor Pro | Confidential

Concurrency and Asynchronous Invocations

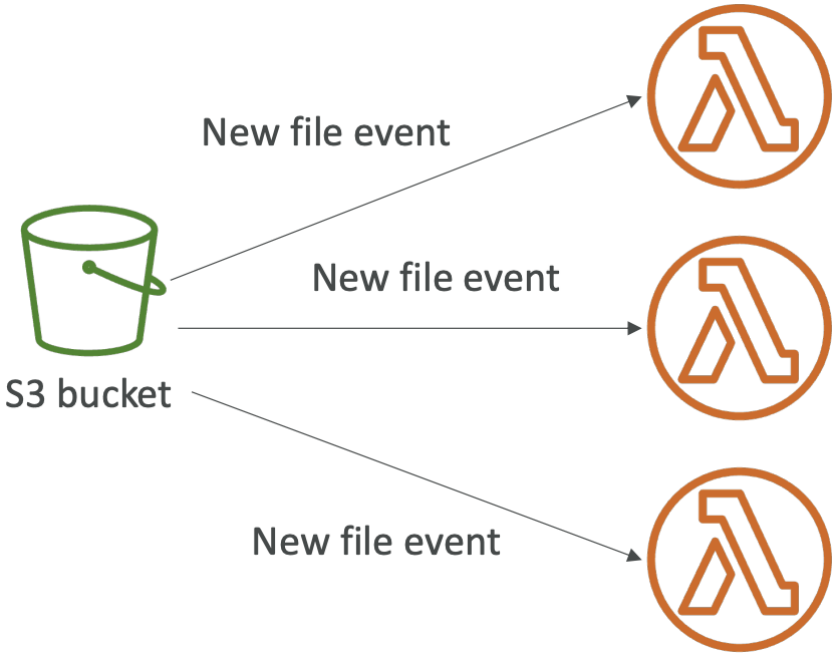

- If the function doesn't have enough concurrency available to process all events, additional requests are throttled.
- For throttling errors (429) and system errors (500-series), Lambda returns the event to the queue and attempts to run the function again for up to 6 hours.
- The retry interval increases exponentially from 1 second after the first attempt to a maximum of 5 minutes.

© Copyright by Cloud Mentor Pro | Confidential

Cold Starts & Provisioned Concurrency

- Cold Start:
  - New instance => code is loaded and code outside the handler run (init)
  - If the init is large (code, dependencies, SDK…) this process can take some time.
  - First request served by new instances has higher latency than the rest
- Provisioned Concurrency:
  - Concurrency is allocated before the function is invoked (in advance)
  - So the cold start never happens and all invocations have low latency
  - Application Auto Scaling can manage concurrency (schedule or target utilization)
- Note:
  - Note: cold starts in VPC have been dramatically reduced in Oct & Nov 2019
  - https://aws.amazon.com/blogs/compute/announcing-improved-vpc-networking-for-aws-lambda-functions/

© Copyright by Cloud Mentor Pro | Confidential

Reserved and Provisioned Concurrency

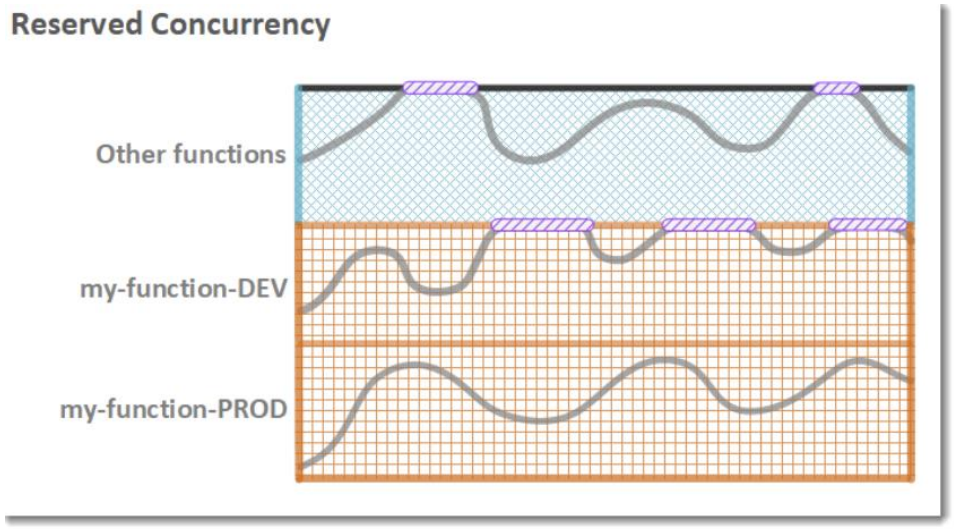

Legend

- Function concurrency
- Reserved concurrency
- Unreserved concurrency
- Throttling

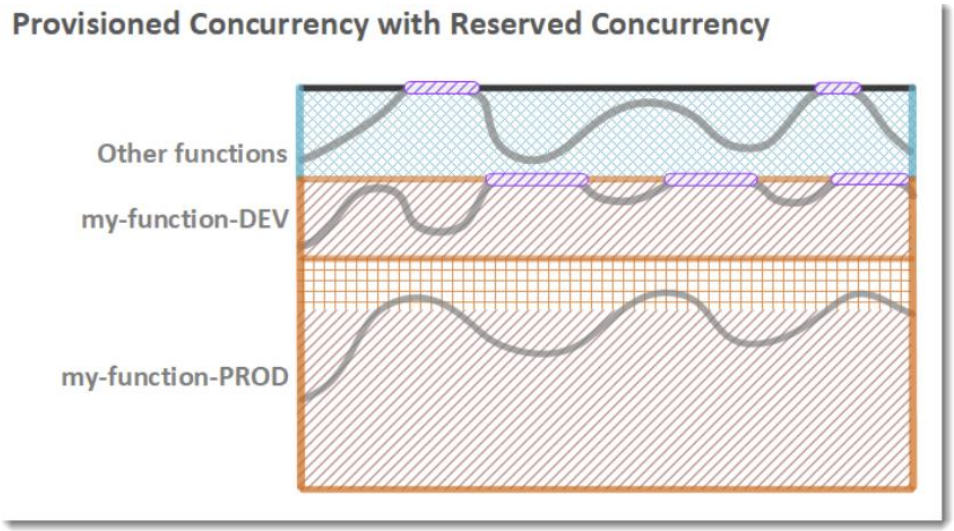

Legend

- Function concurrency
- Reserved concurrency
- Provisioned concurrency
- Unreserved concurrency
- Throttling

https://docs.aws.amazon.com/lambda/latest/dg/configuration-concurrency.html

© Copyright by Cloud Mentor Pro | Confidential

Lambda Function Dependencies

- If your Lambda function depends on external libraries:
  for example AWS X-Ray SDK, Database Clients, etc…
- You need to install the packages alongside your code and zip it together
  - For Node.js, use npm & “node_modules” directory
  - For Python, use pip --target options
  - For Java, include the relevant .jar files
- Upload the zip straight to Lambda if less than 50MB, else to S3 first
- Native libraries work: they need to be compiled on Amazon Linux
- AWS SDK comes by default with every Lambda function

© Copyright by Cloud Mentor Pro | Confidential

Lambda and CloudFormation – inline

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Lambda function inline
Resources:
  primer:
    Type: AWS::Lambda::Function
    Properties:
      Runtime: python3.x
      Role: arn:aws:iam::123456789012:role/lambda-role
      Handler: index.handler
      Code:
        ZipFile: |
          import os

          DB_URL = os.getenv("DB_URL")
          db_client = db.connect(DB_URL)
          def handler(event, context):
              user = db_client.get(user_id = event["user_id"])
              return user
```

- Inline functions are very simple
- Use the **Code.ZipFile** property
- You cannot include function dependencies with inline functions

© Copyright by Cloud Mentor Pro | Confidential

Lambda and CloudFormation – through S3

```
AWSTemplateFormatVersion: '2010-09-09'
Description: Lambda from S3
Resources:
  Function:
    Type: AWS::Lambda::Function
    Properties:
      Handler: index.handler
      Role: arn:aws:iam::123456789012:role/lambda-role
      Code:
        S3Bucket: my-bucket
        S3Key: function.zip
        S3ObjectVersion: String
      Runtime: nodejs12.x
```

- You must store the Lambda zip in S3
- You must refer the S3 zip location in the CloudFormation code
  - S3Bucket
  - S3Key: full path to zip
  - S3ObjectVersion: if versioned bucket
- If you update the code in S3, but don't update S3Bucket, S3Key or S3ObjectVersion, CloudFormation won't update your function

© Copyright by Cloud Mentor Pro | Confidential

Lambda and CloudFormation – through S3 Multiple accounts

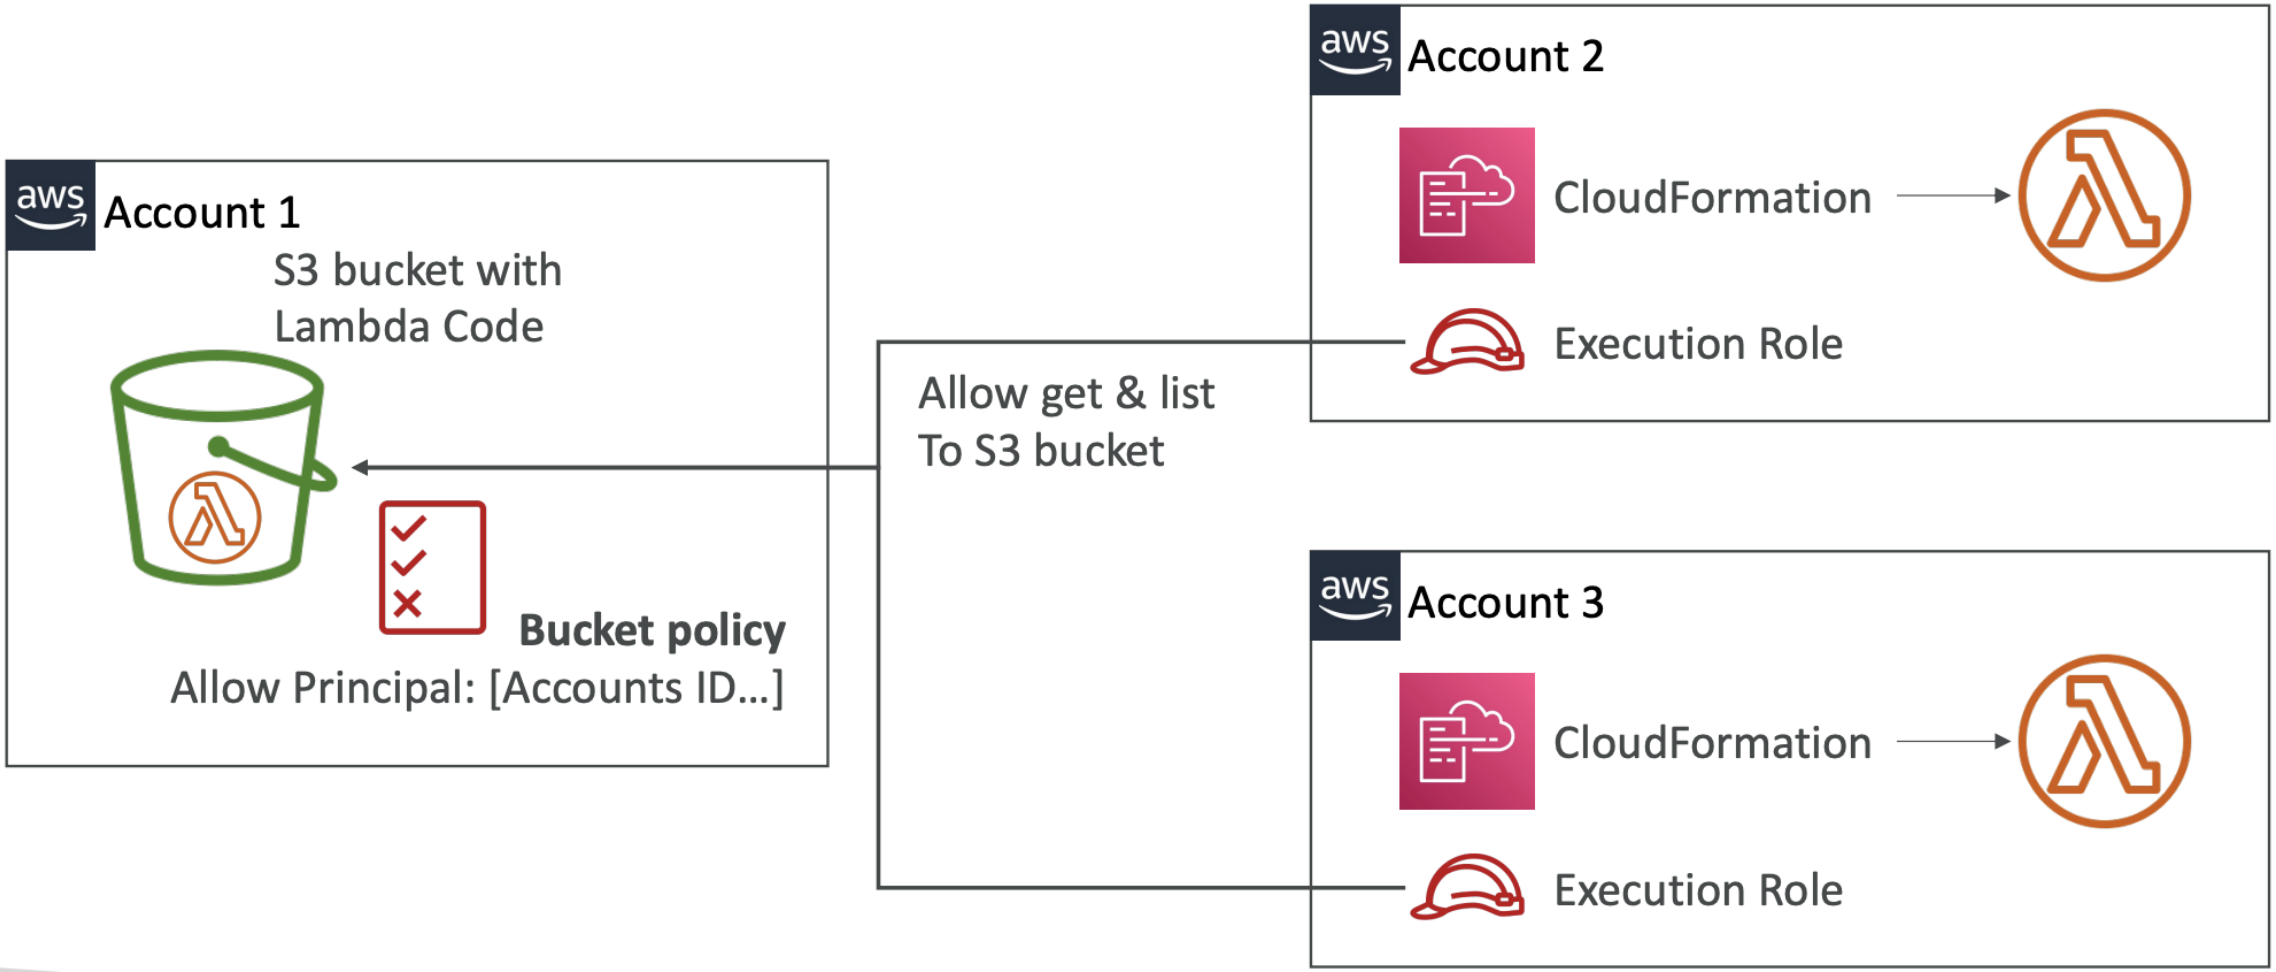

© Copyright by Cloud Mentor Pro | Confidential

Lambda Container Images

- Deploy Lambda function as container images of up to 10GB from ECR
- Pack complex dependencies, large dependencies in a container
- Base images are available for Python, Node.js, Java, .NET, Go, Ruby
- Can create your own image as long as it implements the Lambda Runtime API
- Test the containers locally using the Lambda Runtime Interface Emulator
- Unified workflow to build apps

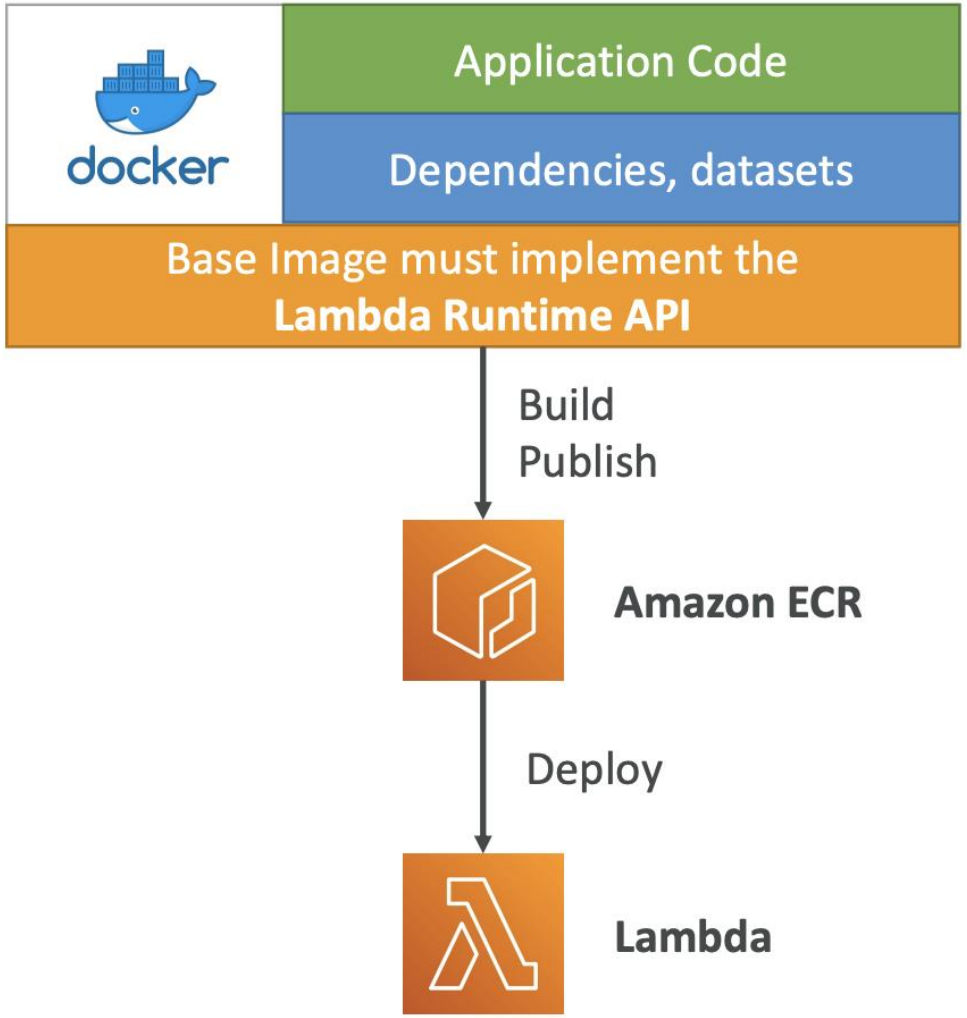

© Copyright by Cloud Mentor Pro | Confidential

Lambda Container Images

- Example: build from the base images provided by AWS

```
# Use an image that implements the Lambda Runtime API
FROM amazon/aws-lambda-nodejs:12

# Copy your application code and files
COPY app.js package*.json ./

# Install the dependencies in the container
RUN npm install

# Function to run when the Lambda function is invoked
CMD [ "app.lambdaHandler" ]
```

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda Versions

- When you work on a Lambda function, we work on **$LATEST**
- When we're ready to publish a Lambda function, we create a version
- Versions are immutable
- Versions have increasing version numbers
- Versions get their own ARN (Amazon Resource Name)
- Version = code + configuration (nothing can be changed - immutable)
- Each version of the lambda function can be accessed

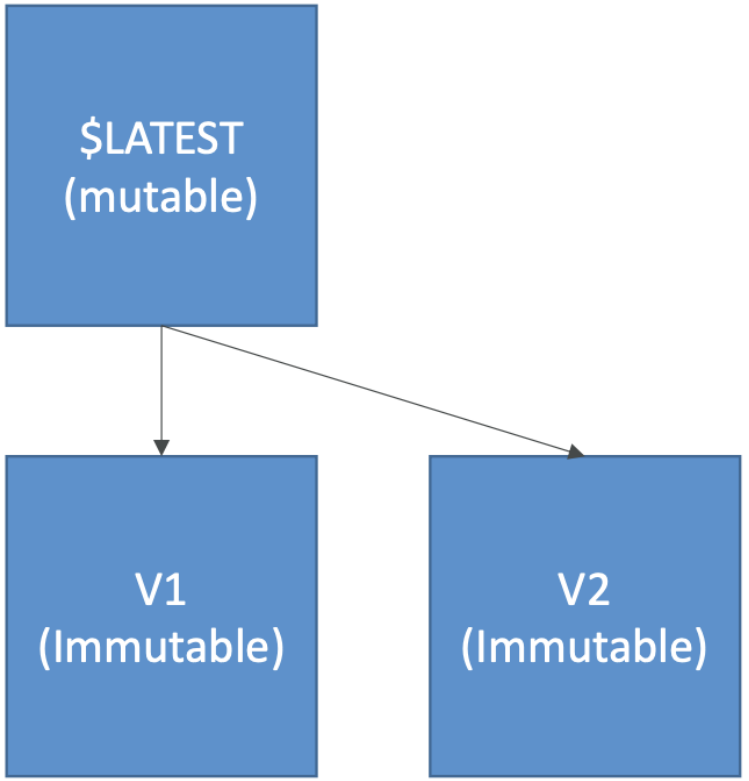

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda Aliases

- Aliases are ”pointers” to Lambda function versions
- We can define a “dev”, ”test”, “prod” aliases and have them point at different lambda versions
- Aliases are mutable
- Aliases enable Canary deployment by assigning weights to lambda functions
- Aliases enable stable configuration of our event triggers / destinations
- Aliases have their own ARNs
- **Aliases cannot reference aliases**

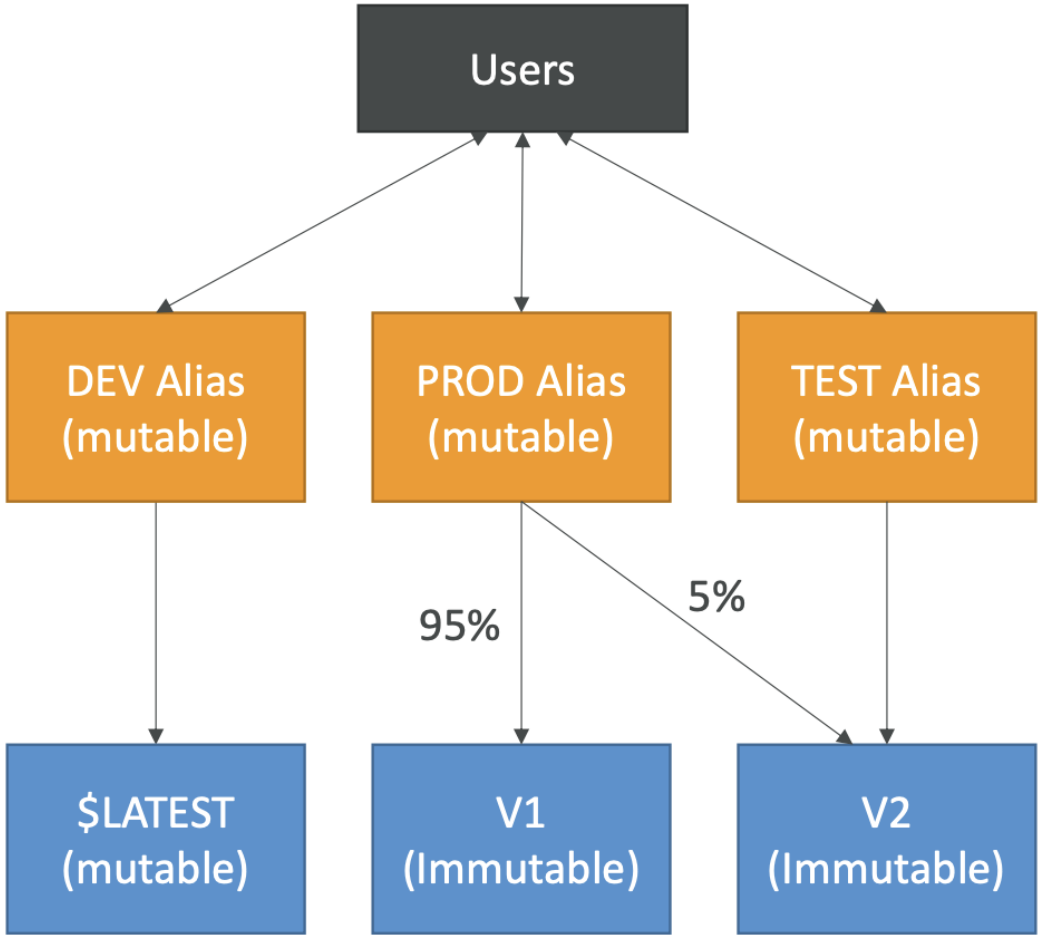

© Copyright by Cloud Mentor Pro | Confidential

Lambda & CodeDeploy

- **CodeDeploy** can help you automate traffic shift for Lambda aliases
- Feature is integrated within the SAM framework
- **Linear**: grow traffic every N minutes until 100%
  - Linear10PercentEvery3Minutes
  - Linear10PercentEvery10Minutes
- **Canary**: try X percent then 100%
  - Canary10Percent5Minutes
  - Canary10Percent30Minutes
- **AllAtOnce**: immediate
- Can create Pre & Post Traffic hooks to check the health of the Lambda function

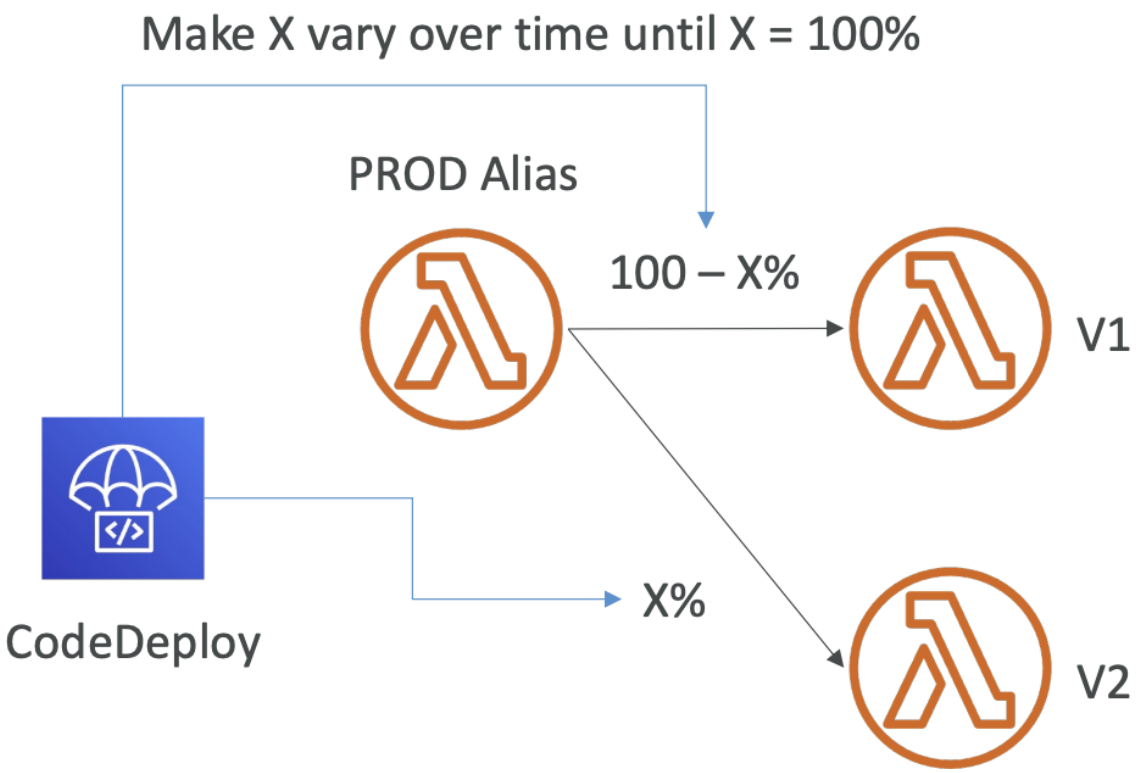

© Copyright by Cloud Mentor Pro | Confidential

Lambda & CodeDeploy – AppSpec.yml

```yaml
version: 0.0

Resources:
  - myLambdaFunction:
      Type: AWS::Lambda::Function
      Properties:
        Name: myLambdaFunction
        Alias: myLambdaFunctionAlias
        CurrentVersion: 1
        TargetVersion: 2
```

- **Name (required)** – the name of the Lambda function to deploy
- **Alias (required)** – the name of the alias to the Lambda function
- **CurrentVersion (required)** – the version of the Lambda function traffic currently points to
- **TargetVersion (required)** – the version of the Lambda function traffic is shifted to

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Function URL

- Dedicated HTTP(S) endpoint for your Lambda function
- A unique URL endpoint is generated for you (never changes)
  - https://<url-id>.lambda-url.<region>.on.aws (dual-stack IPv4 & IPv6)
- Invoke via a web browser, curl, Postman, or any HTTP client
- Access your function URL through the public Internet only
  - Doesn't support PrivateLink (Lambda functions do support)
- Supports Resource-based Policies & CORS configurations
- Can be applied to any function alias or to $LATEST (can't be applied to other function versions)
- Create and configure using AWS Console or AWS API
- Throttle your function by using Reserved Concurrency

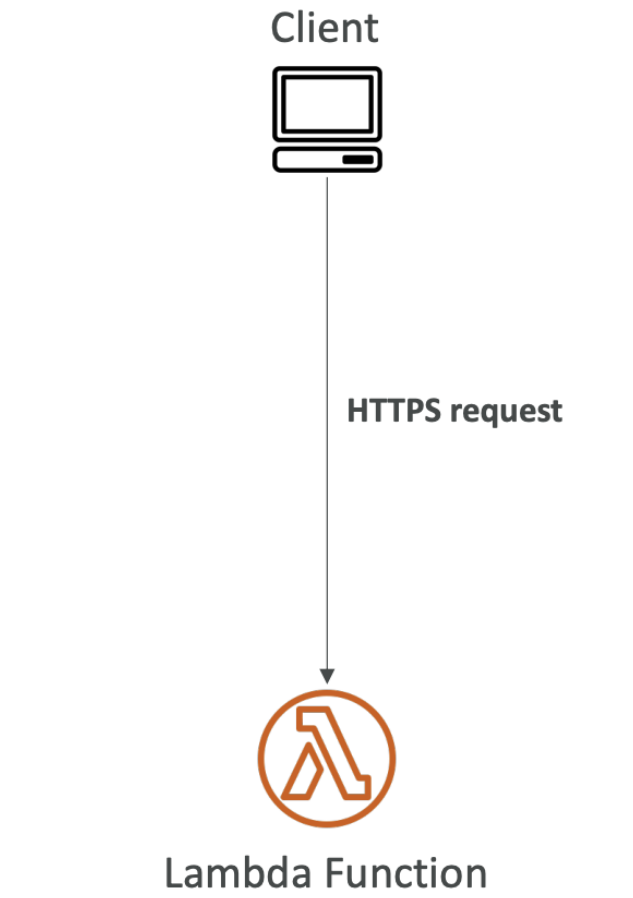

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Function URL Security

- Resource-based Policy
  - Authorize other accounts / specific CIDR / IAM principals

- Cross-Origin Resource Sharing (CORS)
  - If you call your Lambda function URL from a different domain

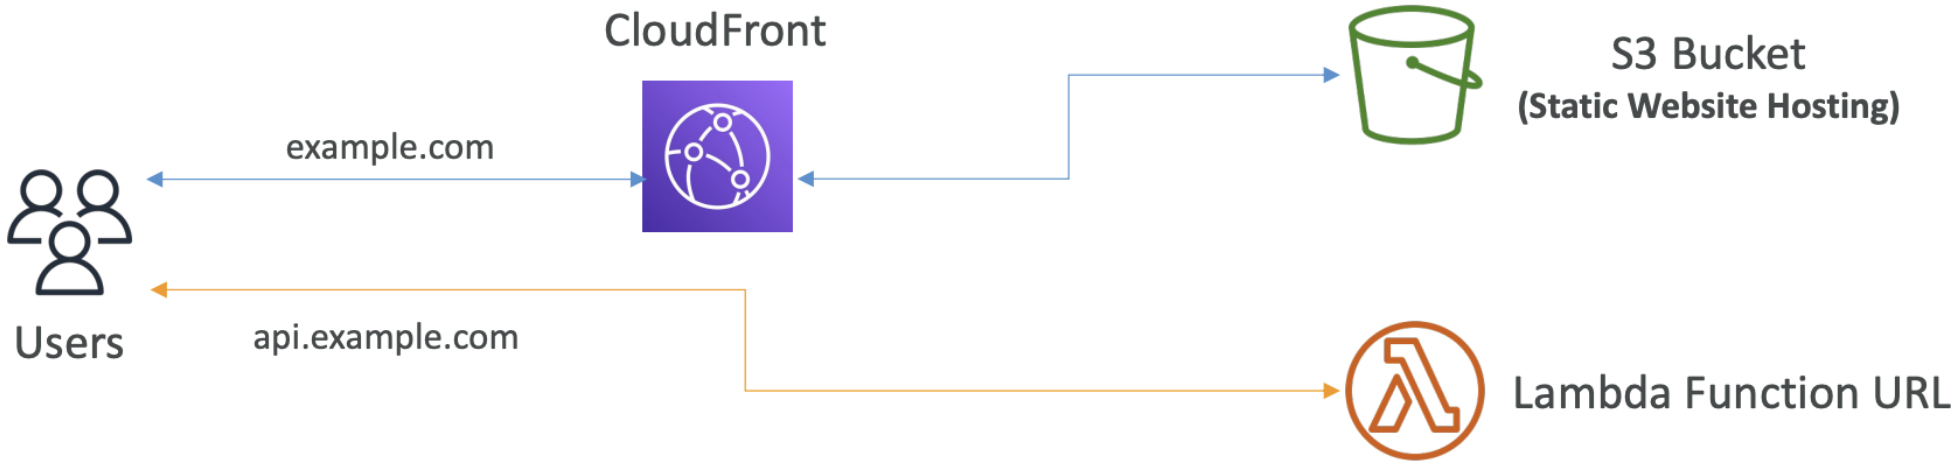

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Function URL Security

- **AuthType NONE** – allow public and unauthenticated access
  - Resource-based Policy is always in effect (must grant public access)

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": "*",
            "Action": "lambda:InvokeFunctionUrl",
            "Resource": "arn:aws:lambda:us-east-1:123456789012:
function:my-function",
            "Condition": {
                "StringEquals": {
                    "lambda:FunctionUrlAuthType": "NONE"
                }
            }
        }
    ]
}
```

Resource-based Policy

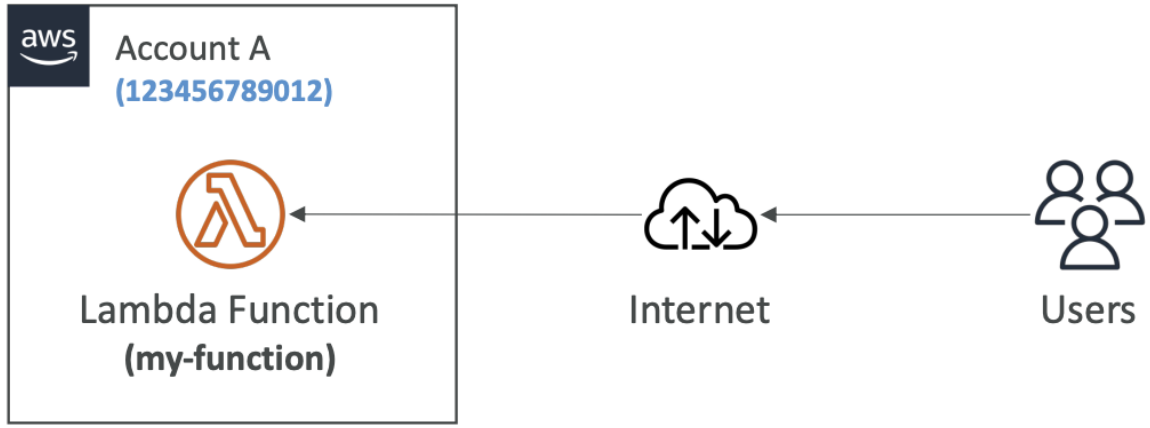

© Copyright by Cloud Mentor Pro | Confidential

Lambda – Function URL Security

- **AuthType AWS_IAM – IAM** is used to authenticate and authorize requests
  - Both Principal's Identity-based Policy & Resource-based Policy are evaluated
  - Principal must have **lambda:InvokeFunctionUrl** permissions
  - **Same account** – Identity-based Policy OR Resource-based Policy as ALLOW
  - **Cross account** – Identity-based Policy AND Resource Based Policy as ALLOW

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "AWS": "arn:aws:iam::444455556666:role/my-role"
            },
            "Action": "lambda:InvokeFunctionUrl",
            "Resource": "arn:aws:lambda:us-east-1:123456789012:
function:my-function",
            "Condition": {
                "StringEquals": {
                    "lambda:FunctionUrlAuthType": "AWS_IAM"
                }
            }
        }
    ]
}
```

Resource-based Policy

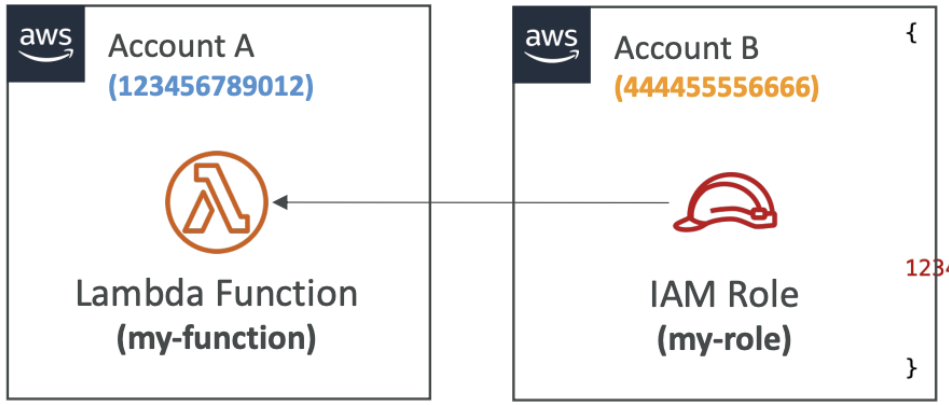


```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": "lambda:InvokeFunctionUrl",
            "Resource": "arn:aws:lambda:us-east-1:
123456789012:function:my-function"
        }
    ]
}
```

Identity-based Policy

© Copyright by Cloud Mentor Pro | Confidential

Lambda and CodeGuru Profiling

- Gain insights into runtime performance of your Lambda functions using CodeGuru Profiler
- CodeGuru creates a Profiler Group for your Lambda function
- Supported for Java and Python runtimes
- Activate from AWS Lambda Console
- When activated, Lambda adds:
  - CodeGuru Profiler layer to your function
  - Environment variables to your function
  - **AmazonCodeGuruProfilerAgentAccess** policy to your function

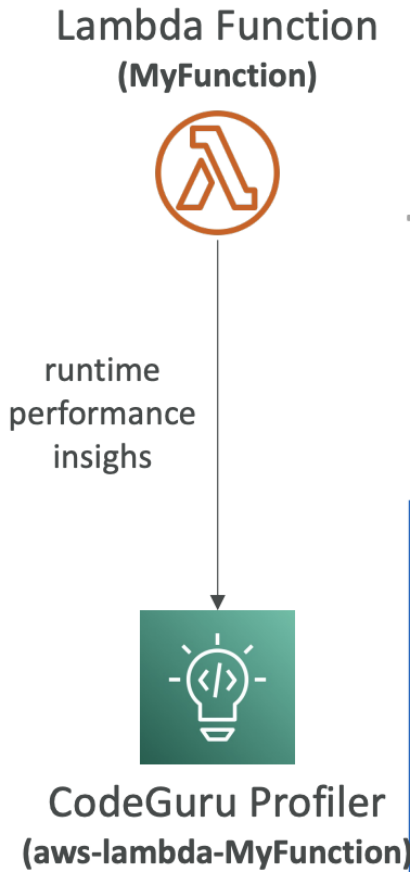

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda Limits to Know - per region

- **Execution:**
  - Memory allocation: 128 MB – 10GB (1 MB increments)
  - Maximum execution time: 900 seconds (15 minutes)
  - Environment variables (4 KB)
  - Disk capacity in the “function container” (in /tmp): 512 MB to 10GB
  - Concurrency executions: 1000 (can be increased)
- **Deployment:**
  - Lambda function deployment size (compressed .zip): 50 MB
  - Size of uncompressed deployment (code + dependencies): 250 MB
  - Can use the /tmp directory to load other files at startup
  - Size of environment variables: 4 KB

© Copyright by Cloud Mentor Pro | Confidential

AWS Lambda Best Practices

- Perform heavy-duty work outside of your function handler
  - Connect to databases outside of your function handler
  - Initialize the AWS SDK outside of your function handler
  - Pull in dependencies or datasets outside of your function handler
- Use environment variables for:
  - Database Connection Strings, S3 bucket, etc… don't put these values in your code
  - Passwords, sensitive values… they can be encrypted using KMS
- Minimize your deployment package size to its runtime necessities.
  - Break down the function if need be
  - Remember the AWS Lambda limits
  - Use Layers where necessary
- Avoid using recursive code, never have a Lambda function call itself

© Copyright by Cloud Mentor Pro | Confidential

Exercise

Amazon S3 trigger Lambda to create thumbnail images

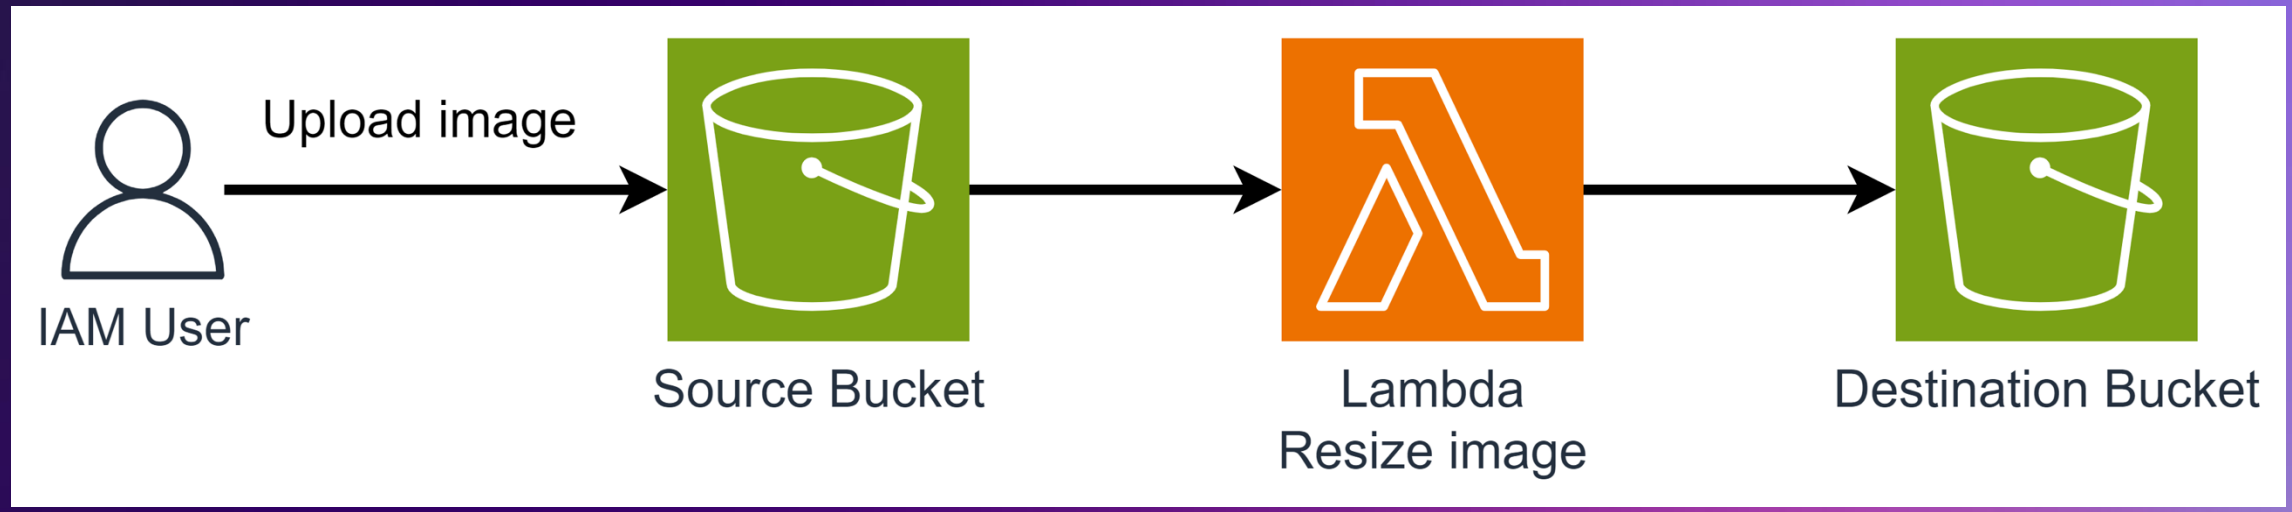

© Copyright by Cloud Mentor Pro | Confidential

Section 12

- *Amazon DynamoDB*

© Copyright by Cloud Mentor Pro | Confidential

Amazon DynamoDB

NoSQL Serverless Database

© Copyright by Cloud Mentor Pro | Confidential

Traditional Architecture

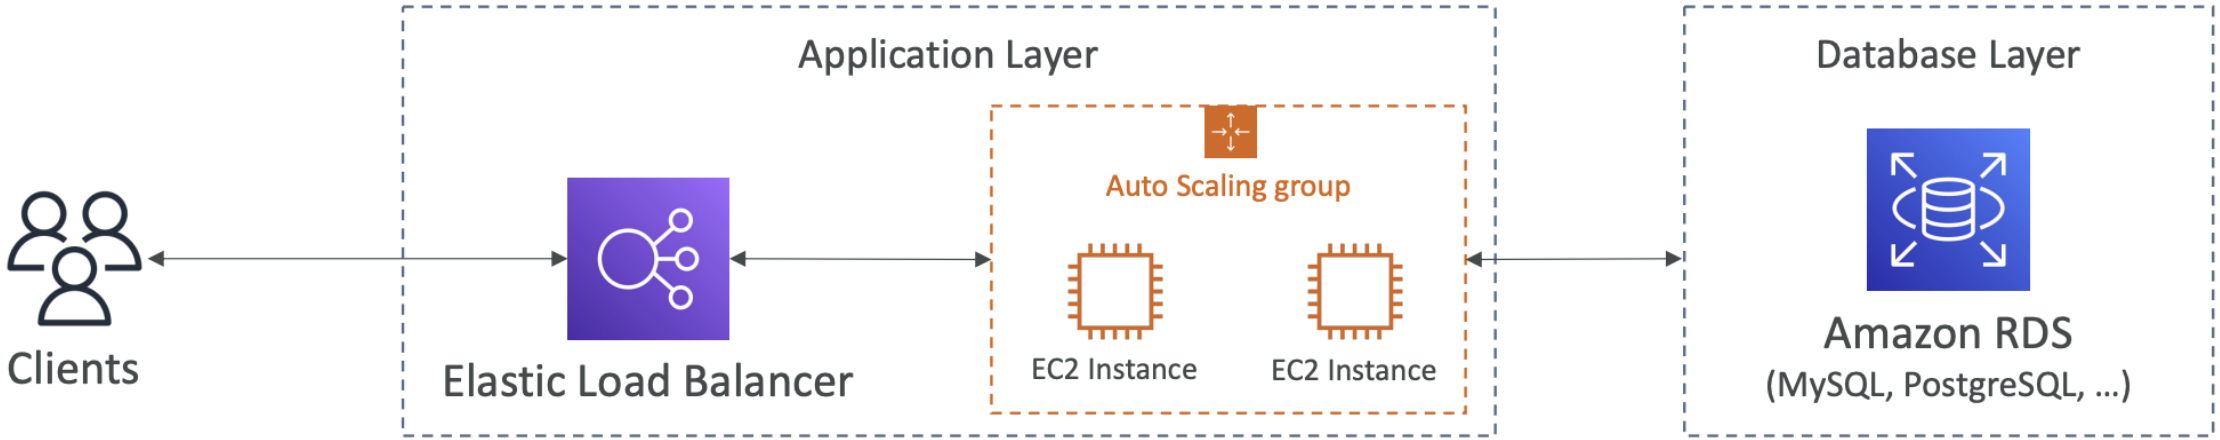

- Traditional applications leverage RDBMS databases
- These databases have the SQL query language
- Strong requirements about how the data should be modeled
- Ability to do query joins, aggregations, complex computations
- Vertical scaling (getting a more powerful CPU / RAM / IO)
- Horizontal scaling (increasing reading capability by adding EC2 / RDS Read Replicas)

© Copyright by Cloud Mentor Pro | Confidential

NoSQL databases

- NoSQL databases are non-relational databases and are **distributed**
- NoSQL databases include MongoDB, DynamoDB, …
- NoSQL databases do not support query joins (or just limited support)
- All the data that is needed for a query is present in one row
- NoSQL databases don't perform aggregations such as "SUM", "AVG", …
- **NoSQL databases scale horizontally**

- There's no "right or wrong" for NoSQL vs SQL, they just require to model the data differently and think about user queries differently

© Copyright by Cloud Mentor Pro | Confidential

Amazon DynamoDB

- Fully managed, highly available with replication across multiple AZs
- NoSQL database - not a relational database
- Scales to massive workloads, distributed database
- Millions of requests per seconds, trillions of row, 100s of TB of storage
- Fast and consistent in performance (low latency on retrieval)
- Integrated with IAM for security, authorization and administration
- Enables event driven programming with DynamoDB Streams
- Low cost and auto-scaling capabilities
- Standard & Infrequent Access (IA) Table Class

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB - Basics

- DynamoDB is made of **Tables**
- Each table has a **Primary Key** (must be decided at creation time)
- Each table can have an infinite number of items (= rows)
- Each item has **attributes** (can be added over time – can be null)
- Maximum size of an item is **400KB**
- Data types supported are:
  - **Scalar Types** – String, Number, Binary, Boolean, Null
  - **Document Types** – List, Map
  - **Set Types** – String Set, Number Set, Binary Set

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Primary Keys

- Option 1: Partition Key (HASH)
  - Partition key must be unique for each item
  - Partition key must be "diverse" so that the data is distributed
  - Example: "User_ID" for a users table

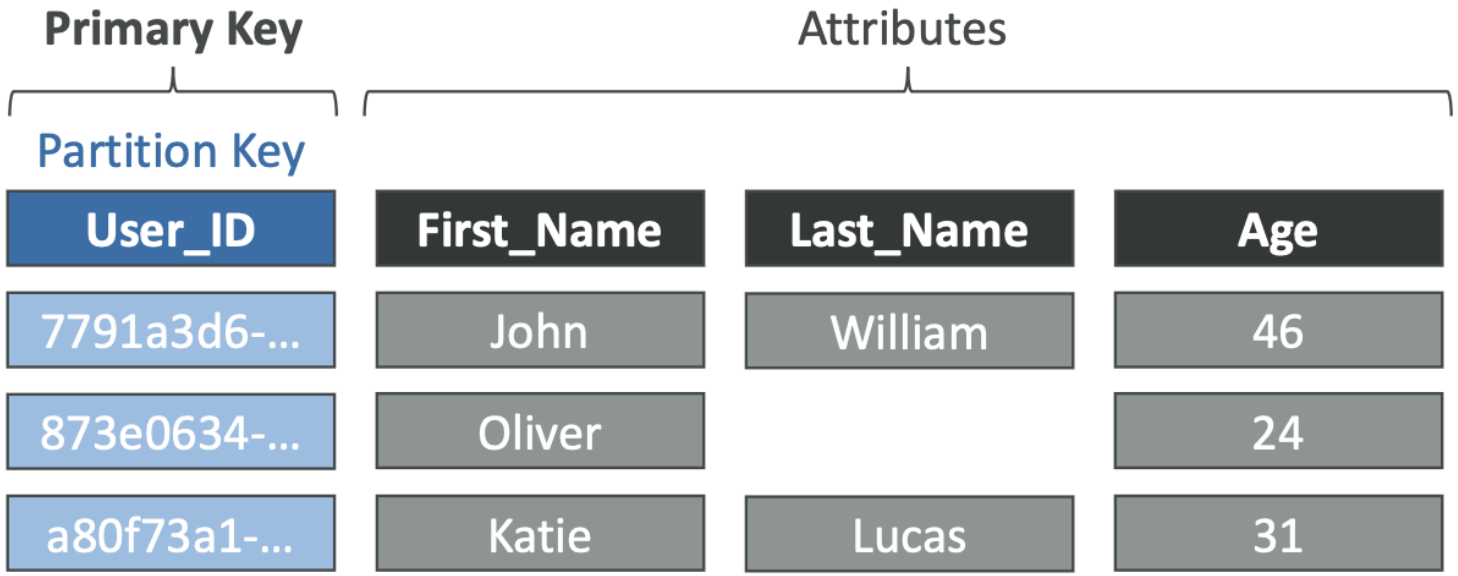

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Primary Keys

- Option 2: Partition Key + Sort Key (HASH + RANGE)
  - The combination must be unique for each item
  - Data is grouped by partition key
  - Example: users-games table, “User_ID” for Partition Key and “Game_ID” for Sort Key

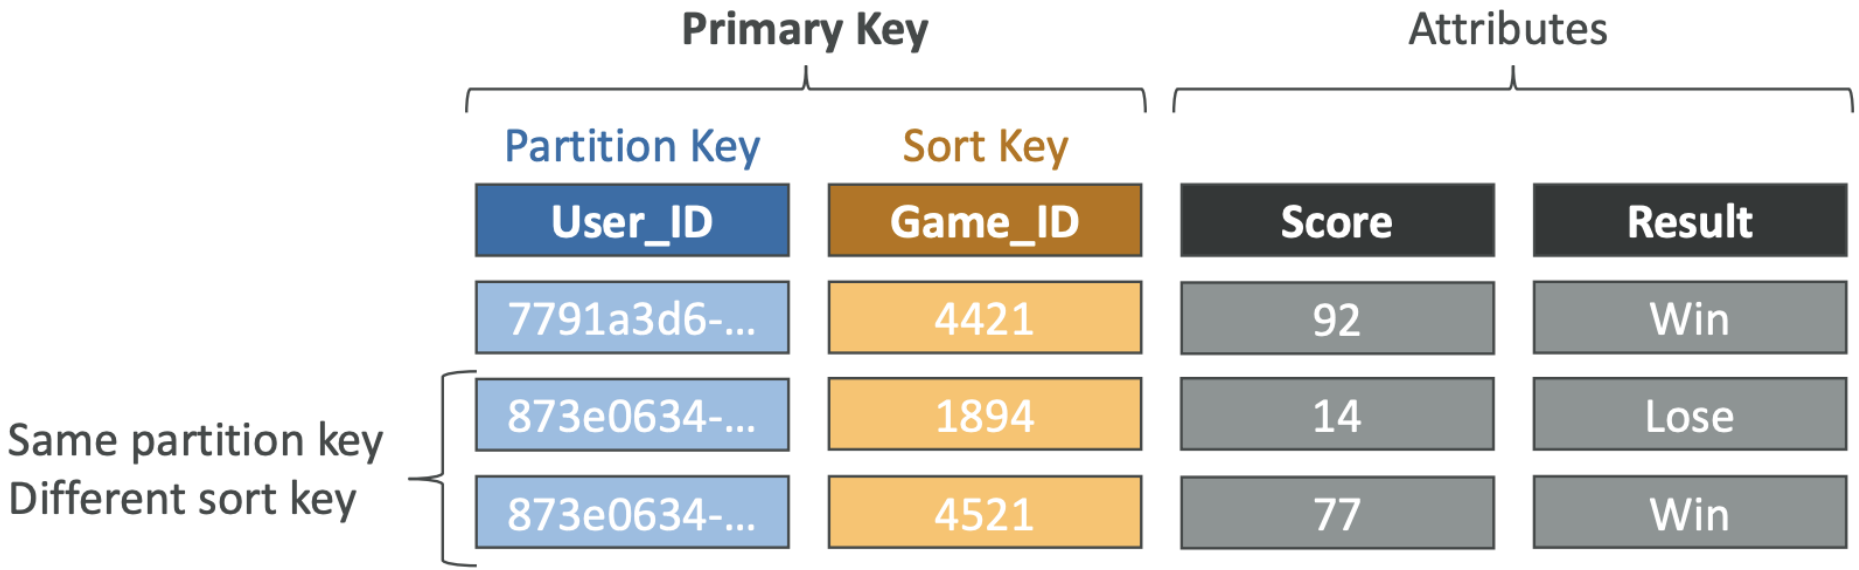

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Partition Keys (Exercise)

- We're building a movie database
- What is the best Partition Key to maximize data distribution?
  - movie_id
  - producer_name
  - leader_actor_name
  - movie_language

- "movie_id" has the highest cardinality so it's a good candidate
- "movie_language" doesn't take many values and may be skewed towards English so it's not a great choice for the Partition Key

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Read/Write Capacity Modes

- Control how you manage your table's capacity (read/write throughput)

- **Provisioned Mode (default)**
  - You specify the number of reads/writes per second
  - You need to plan capacity beforehand
  - Pay for provisioned read & write capacity units

- **On-Demand Mode**
  - Read/writes automatically scale up/down with your workloads
  - No capacity planning needed
  - Pay for what you use, more expensive ($$$)

- You can switch between different modes once every 24 hours

© Copyright by Cloud Mentor Pro | Confidential

R/W Capacity Modes – Provisioned

- Table must have provisioned read and write capacity units
- **Read Capacity Units (RCU)** – throughput for reads
- **Write Capacity Units (WCU)** – throughput for writes
- Option to setup auto-scaling of throughput to meet demand
- Throughput can be exceeded temporarily using "**Burst Capacity**"
- If Burst Capacity has been consumed, you'll get a "**ProvisionedThroughputExceededException**"
- It's then advised to do an **exponential backoff** retry

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Write Capacity Units (WCU)

- One Write Capacity Unit (WCU) represents one write per second for an item up to **1 KB** in size
- If the items are larger than 1 KB, more WCUs are consumed

- **Example 1**: we write 10 items per second, with item size 2 KB
  - We need $10 * (\frac{2\ \text{KB}}{1\ KB})$ = 20 WCUs
- **Example 2**: we write 6 items per second, with item size 4.5 KB
  - We need $6 * (\frac{5\ \text{KB}}{1\ KB})$ = 30 WCUs (4.5 gets rounded to the upper KB)
- **Example 3**: we write 120 items per minute, with item size 2 KB
  - We need $(\frac{120}{60}) * (\frac{2\ \text{KB}}{1\ KB})$ = 4 WCUs

© Copyright by Cloud Mentor Pro | Confidential

Strongly Consistent Read vs. Eventually Consistent Read

- Eventually Consistent Read (default)
  - If we read just after a write, it's possible we'll get some stale data because of replication

- Strongly Consistent Read
  - If we read just after a write, we will get the correct data
  - Set "**ConsistentRead**" parameter to **True** in API calls (GetItem, BatchGetItem, Query, Scan)
  - Consumes twice the RCU

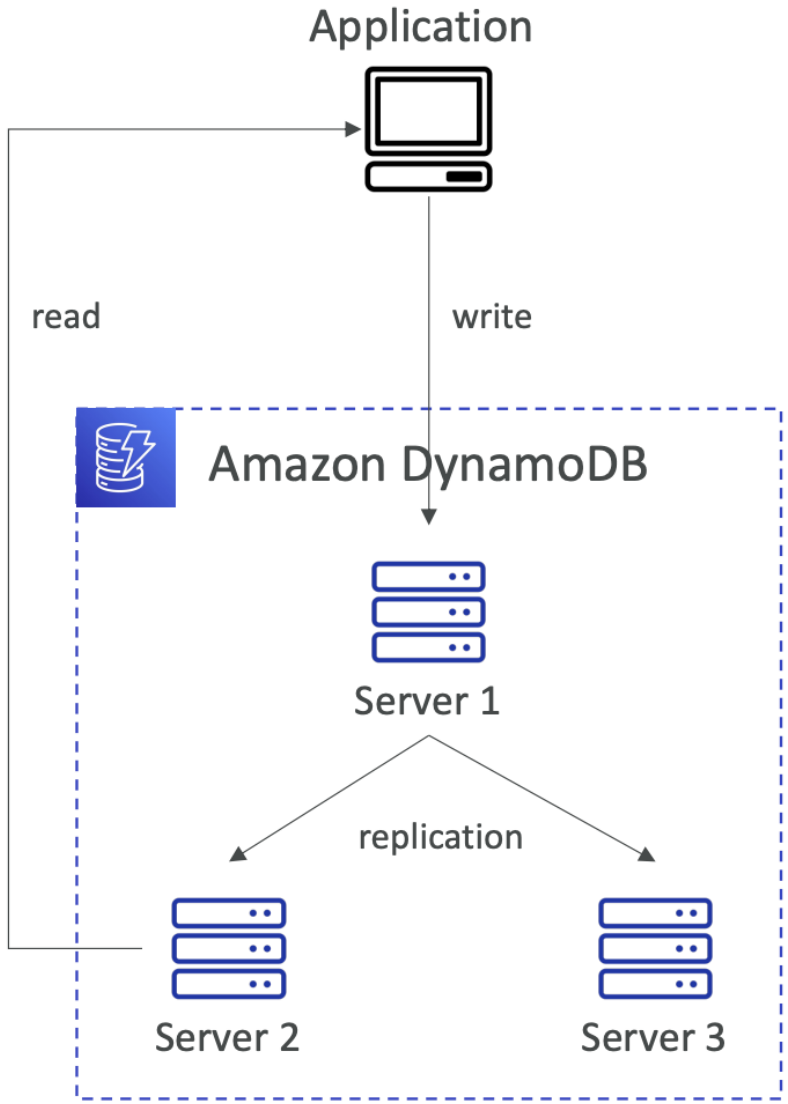

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Read Capacity Units (RCU)

- One Read Capacity Unit (RCU) represents one **Strongly Consistent Read** per second, or **two Eventually Consistent Reads** per second, for an item up to **4 KB** in size
- If the items are larger than 4 KB, more RCUs are consumed

- **Example 1**: 10 Strongly Consistent Reads per second, with item size 4 KB
  - We need $10 * (\frac{4\ \text{KB}}{4\ KB})$ = 10 RCUs
- **Example 2**: 16 Eventually Consistent Reads per second, with item size 12 KB
  - We need $(\frac{16}{2}) * (\frac{12\ \text{KB}}{4\ KB})$ = 24 RCUs
- **Example 3**: 10 Strongly Consistent Reads per second, with item size 6 KB
  - We need $10 * (\frac{8\ \text{KB}}{4\ KB})$ = 20 RCUs (we must round up 6 KB to 8 KB)

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Partitions Internal

- Data is stored in partitions
- Partition Keys go through a hashing algorithm to know to which partition they go to

- To compute the number of partitions:
  - *# of partitions*$_{by\ capacity}$ = $(\frac{RCUs_{Total}}{3000}) + (\frac{WCUs_{Total}}{1000})$
  - *# of partitions*$_{by\ size}$ = $(\frac{Total\ Size}{10\ GB})$
  - *# of partitions = ceil(max(# of partitions*$_{by\ capacity}$*, of partitions*$_{by\ size}$ *))*

- WCUs and RCUs are spread evenly across partitions

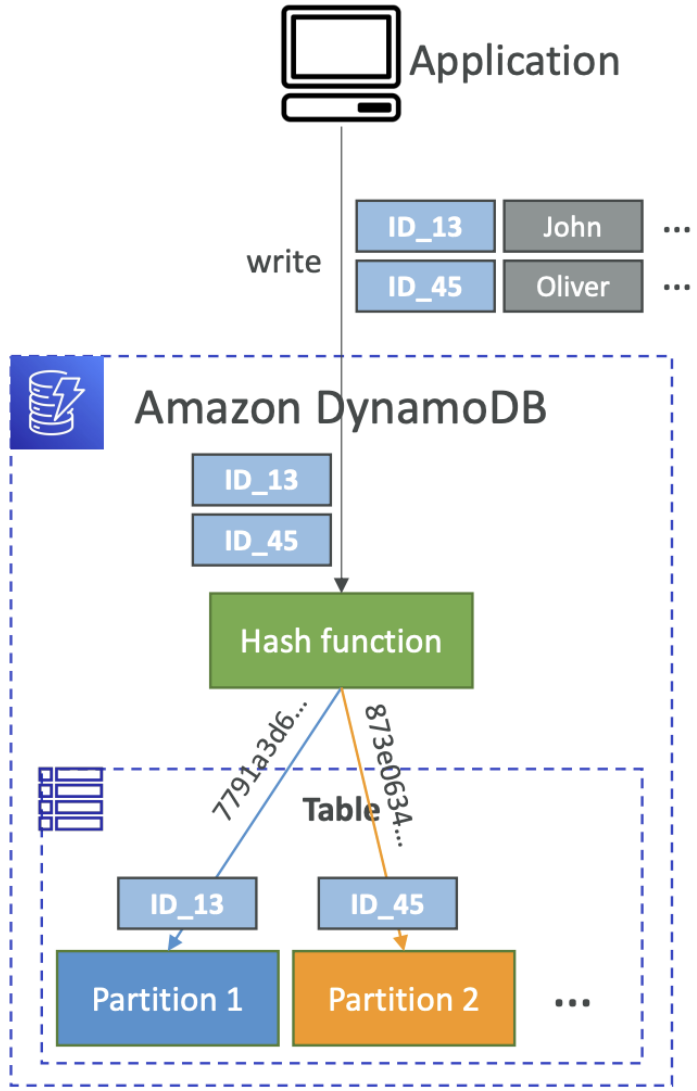

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB –Throttling

- If we exceed provisioned RCUs or WCUs, we get "**ProvisionedThroughputExceededException**"
- Reasons:
  - **Hot Keys** – one partition key is being read too many times (e.g., popular item)
  - **Hot Partitions**
  - **Very large items**, remember RCU and WCU depends on size of items
- Solutions:
  - **Exponential backoff** when exception is encountered (already in SDK)
  - **Distribute partition keys** as much as possible
  - If RCU issue, we can **use DynamoDB Accelerator (DAX)**

© Copyright by Cloud Mentor Pro | Confidential

R/W Capacity Modes – On-Demand

- Read/writes automatically scale up/down with your workloads
- No capacity planning needed (WCU / RCU)
- Unlimited WCU & RCU, no throttle, more expensive
- You're charged for reads/writes that you use in terms of **RRU** and **WRU**
- **Read Request Units (RRU)** – throughput for reads (same as RCU)
- **Write Request Units (WRU)** – throughput for writes (same as WCU)
- 2.5x more expensive than provisioned capacity (use with care)
- Use cases: unknown workloads, unpredictable application traffic, ...

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Writing Data

- PutItem
  - Creates a new item or fully replace an old item (same Primary Key)
  - Consumes WCUs
- UpdateItem
  - Edits an existing item's attributes or adds a new item if it doesn't exist
  - Can be used to implement **Atomic Counters** – a numeric attribute that's unconditionally incremented
- Conditional Writes
  - Accept a write/update/delete only if conditions are met, otherwise returns an error
  - Helps with concurrent access to items
  - No performance impact

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Reading Data

- GetItem
  - Read based on Primary key
  - Primary Key can be **HASH or HASH+RANGE**
  - Eventually Consistent Read (default)
  - Option to use Strongly Consistent Reads (more RCU - might take longer)
  - **ProjectionExpression** can be specified to retrieve only certain attributes

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Reading Data (Query)

- **Query** returns items based on:
  - **KeyConditionExpression**
    - Partition Key value (**must be = operator**) – required
    - Sort Key value (=, <=, >, >=, Between, Begins with) – optional
  - **FilterExpression**
    - Additional filtering after the Query operation (before data returned to you)
    - Use only with non-key attributes (does not allow HASH or RANGE attributes)
- Returns:
  - The number of items specified in **Limit**
  - Or up to 1 MB of data
- Ability to do pagination on the results
- Can query table, a Local Secondary Index, or a Global Secondary Index

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Reading Data (Scan)

- **Scan** the entire table and then filter out data (inefficient)
- Returns up to 1 MB of data – use pagination to keep on reading
- Consumes a lot of RCU
- Limit impact using **Limit** or reduce the size of the result and pause
- For faster performance, use **Parallel Scan**
  - Multiple workers scan multiple data segments at the same time
  - Increases the throughput and RCU consumed
  - Limit the impact of parallel scans just like you would for Scans
- Can use **ProjectionExpression & FilterExpression** (no changes to RCU)

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Deleting Data

- **DeleteItem**
  - Delete an individual item
  - Ability to perform a conditional delete

- **DeleteTable**
  - Delete a whole table and all its items
  - Much quicker deletion than calling **DeleteItem** on all items

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Batch Operations

- Allows you to save in latency by reducing the number of API calls
- Operations are done in parallel for better efficiency
- Part of a batch can fail; in which case we need to try again for the failed items
  -
- **BatchWriteItem**
  - Up to 25 PutItem and/or **DeleteItem** in one call
  - Up to 16 MB of data written, up to 400 KB of data per item
  - Can't update items (use **UpdateItem**)
  - **UnprocessedItems** for failed write operations (exponential backoff or add WCU)

- **BatchGetItem**
  - Return items from one or more tables
  - Up to 100 items, up to 16 MB of data
  - Items are retrieved in parallel to minimize latency
  - **UnprocessedKeys** for failed read operations (exponential backoff or add RCU)

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – PartiQL

- SQL -compatible query language for DynamoDB
- Allows you to select, insert, update, and delete data in DynamoDB using SQL
- Run queries across multiple DynamoDB tables
- Run PartiQL queries from:
  - AWS Management Console
  - NoSQL Workbench for DynamoDB
  - DynamoDB APIs
  - AWS CLI
  - AWS SDK

```
SELECT OrderID, Total
FROM Orders
WHERE OrderID IN [1, 2, 3]
ORDER BY OrderID DESC
```

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Conditional Writes

- For **PutItem, UpdateItem, DeleteItem**, and **BatchWriteItem**
- You can specify a Condition expression to determine which items should be modified:
  - **attribute_exists**
  - **attribute_not_exists**
  - **attribute_type**
  - **contains** (for string)
  - **begins_with** (for string)
  - ProductCategory **IN** (:cat1, :cat2) **and** Price **between** :low **and** :high
  - **size** (string length)

- Note: Filter Expression filters the results of read queries, while Condition Expressions are for write operations

© Copyright by Cloud Mentor Pro | Confidential

Conditional Writes – Example on Update Item

```
aws dynamodb update-item \
    --table-name ProductCatalog \
    --key '{ "Id": { "N": "456" } }' \
    --update-expression "SET Price = Price - :discount" \
    --condition-expression "Price > :limit" \
    --expression-attribute-values file://values.json
```

```
{
    ":discount": {
        "N": "150"
    },
    ":limit": {
        "N": "500"
    }
}
```

values.json

```
{
    "Id": {
        "N": "456"
    },
    "Price": {
        "N": "650"
    },
    "ProductCategory": {
        "S": "Sporting Goods"
    }
}
```



```
{
    "Id": {
        "N": "456"
    },
    "Price": {
        "N": "500"
    },
    "ProductCategory": {
        "S": "Sporting Goods"
    }
}
```

© Copyright by Cloud Mentor Pro | Confidential

Conditional Writes – Example on Delete Item

- attribute_not_exists
  - Only succeeds if the attribute doesn't exist yet (no value)

```
aws dynamodb delete-item \
    --table-name ProductCatalog \
    --key '{ "Id": { "N": "456" } }' \
    --condition-expression "attribute_not_exists(Price)"
```

- attribute_exists
  - Opposite of attribute_not_exists

```
aws dynamodb delete-item \
    --table-name ProductCatalog \
    --key '{ "Id": { "N": "456" } }' \
    --condition-expression "attribute_exists(ProductReviews.OneStar)"
```

© Copyright by Cloud Mentor Pro | Confidential

Conditional Writes – Do Not Overwrite Elements

- attribute_not_exists(partition_key)
  - Make sure the item isn't overwritten

- attribute_not_exists(partition_key) and attribute_not_exists(sort_key)
  - Make sure the partition / sort key combination is not overwritten

© Copyright by Cloud Mentor Pro | Confidential

Conditional Writes – Example Complex Condition

```
aws dynamodb delete-item \
    --table-name ProductCatalog \
    --key '{ "Id": { "N": "456" } }' \
    --condition-expression "(ProductCategory IN (:cat1, :cat2)) and (Price between :lo and :hi)" \
    --expression-attribute-values file://values.json
```

```
{
    ":cat1": {
        "S": "Sporting Goods"
    },
    ":cat2": {
        "S": "Gardening Supplies"
    },
    ":lo": {
        "N": "500"
    },
    ":hi": {
        "N": "600"
    }
}
```

values.json

```
{
    "Id": {
        "N": "456"
    },
    "Price": {
        "N": "650"
    },
    "ProductCategory": {
        "S": "Sporting Goods"
    }
}
```

© Copyright by Cloud Mentor Pro | Confidential

Conditional Writes – Example of String Comparisons

- **begins_with** – check if prefix matches
- **contains** – check if string is contained in another string

```
aws dynamodb delete-item \
    --table-name ProductCatalog \
    --key '{ "Id": { "N": "456" } }' \
    --condition-expression "begins_with(Pictures.FrontView, :v_sub)" \
    --expression-attribute-values file://values.json
```

```
{
    ":v_sub": {
        "S": "http://"
    }
}
```

values.json

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Local Secondary Index (LSI)

- **Alternative Sort Key** for your table (same **Partition Key** as that of base table)
- The Sort Key consists of one scalar attribute (String, Number, or Binary)
- Up to 5 Local Secondary Indexes per table
- **Must be defined at table creation time**
- **Attribute Projections** – can contain some or all the attributes of the base table (KEYS_ONLY, INCLUDE, ALL)

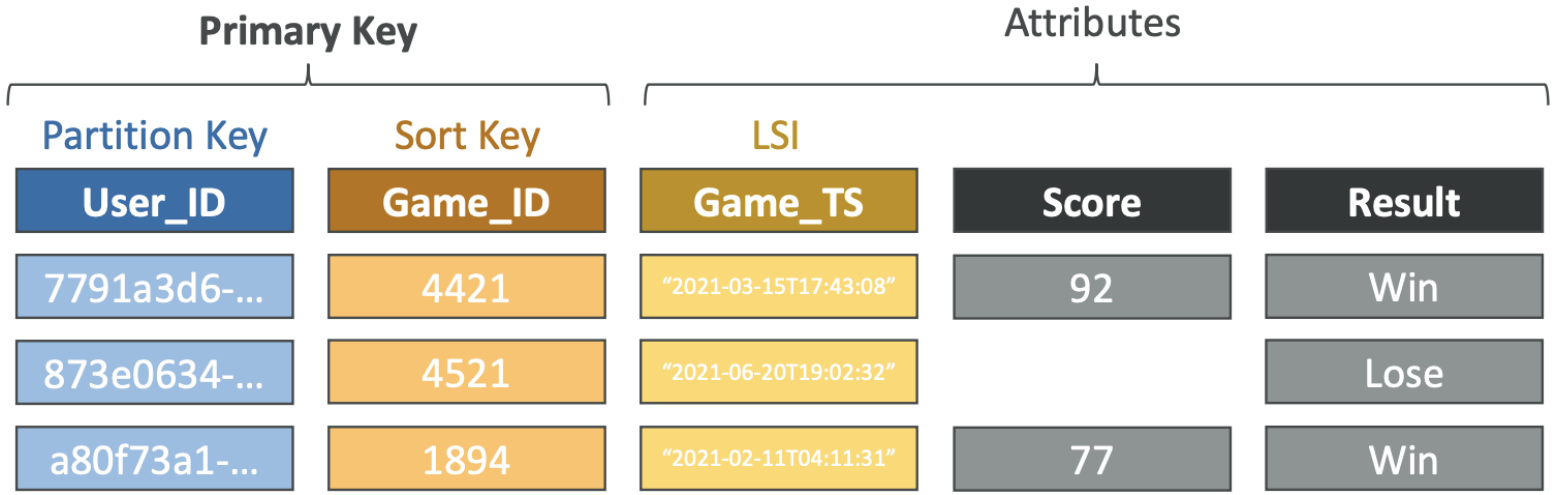

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Global Secondary Index (GSI)

- **Alternative Primary Key (HASH or HASH+RANGE)** from the base table
- Speed up queries on non-key attributes
- The Index Key consists of scalar attributes (String, Number, or Binary)
- **Attribute Projections** – some or all the attributes of the base table (**KEYS_ONLY, INCLUDE, ALL**)
- Must provision RCUs & WCUs for the index
- **Can be added/modified after table creation**

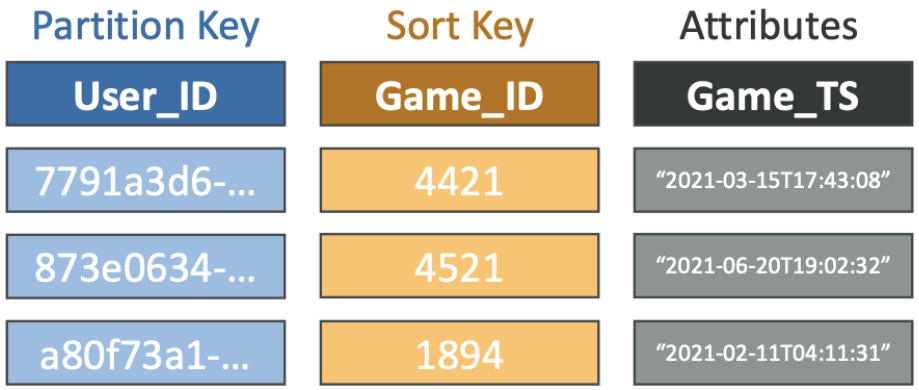

TABLE (query by "User_ID")

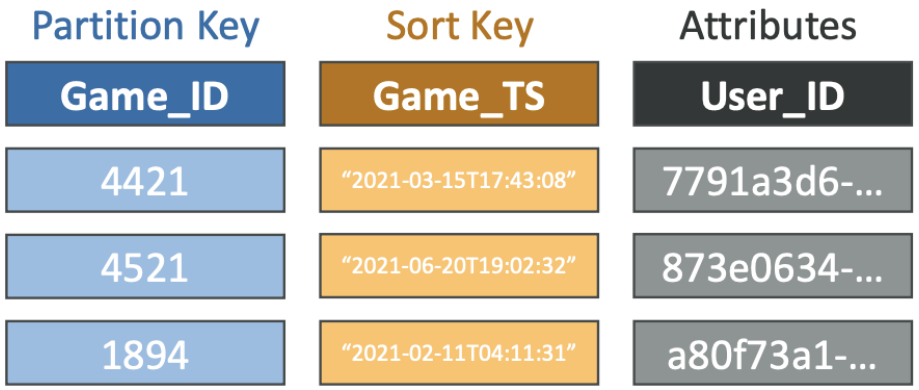

INDEX GSI (query by "Game_ID")

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Indexes and Throttling

- Global Secondary Index (GSI):
  - If the writes are throttled on the GSI, then the main table will be throttled!
  - Even if the WCU on the main tables are fine
  - Choose your GSI partition key carefully!
  - Assign your WCU capacity carefully!
- Local Secondary Index (LSI):
  - Uses the WCUs and RCUs of the main table
  - No special throttling considerations

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB - PartiQL

- Use a SQL -like syntax to manipulate DynamoDB tables

- Supports some (but not all) statements:
  - INSERT
  - UPDATE
  - SELECT
  - DELETE
- It supports Batch operations

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Optimistic Locking

- DynamoDB has a feature called "**Conditional Writes**"
- A strategy to ensure an item hasn't changed before you update/delete it
- Each item has an attribute that acts as a *version number*

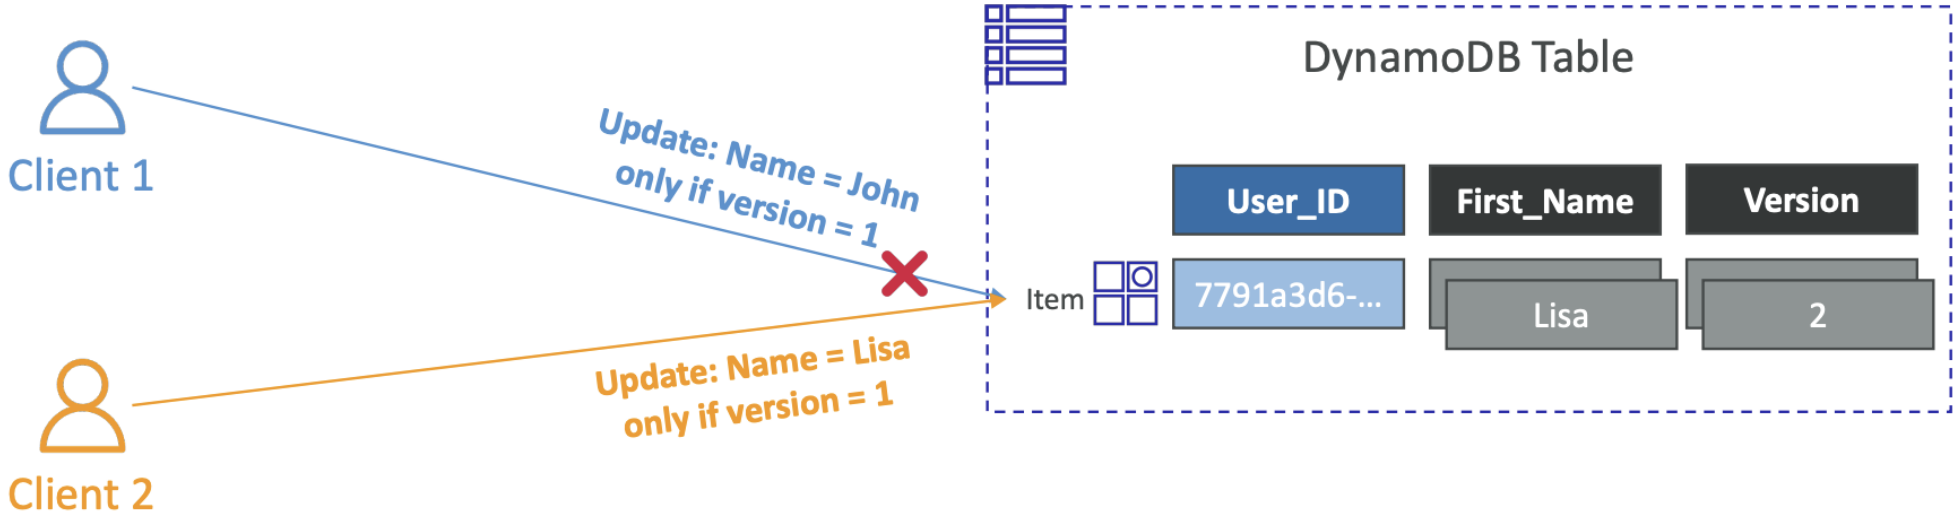

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Accelerator (DAX)

- Fully-managed, highly available, seamless in-memory cache for DynamoDB
- Microseconds latency for cached reads & queries
- Doesn’t require application logic modification (compatible with existing DynamoDB APIs)
- Solves the “Hot Key” problem (too many reads)
- 5 minutes TTL for cache (default)
- Up to 10 nodes in the cluster
- Multi-AZ (3 nodes minimum recommended for production)
- Secure (Encryption at rest with KMS, VPC, IAM, CloudTrail, …)

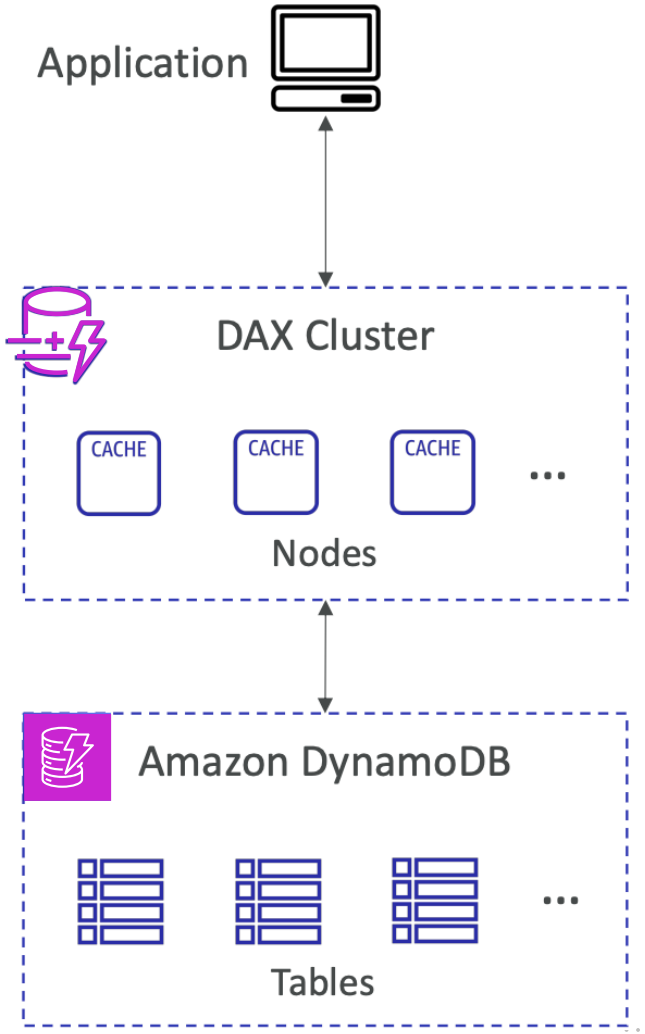

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Accelerator (DAX) vs. ElastiCache

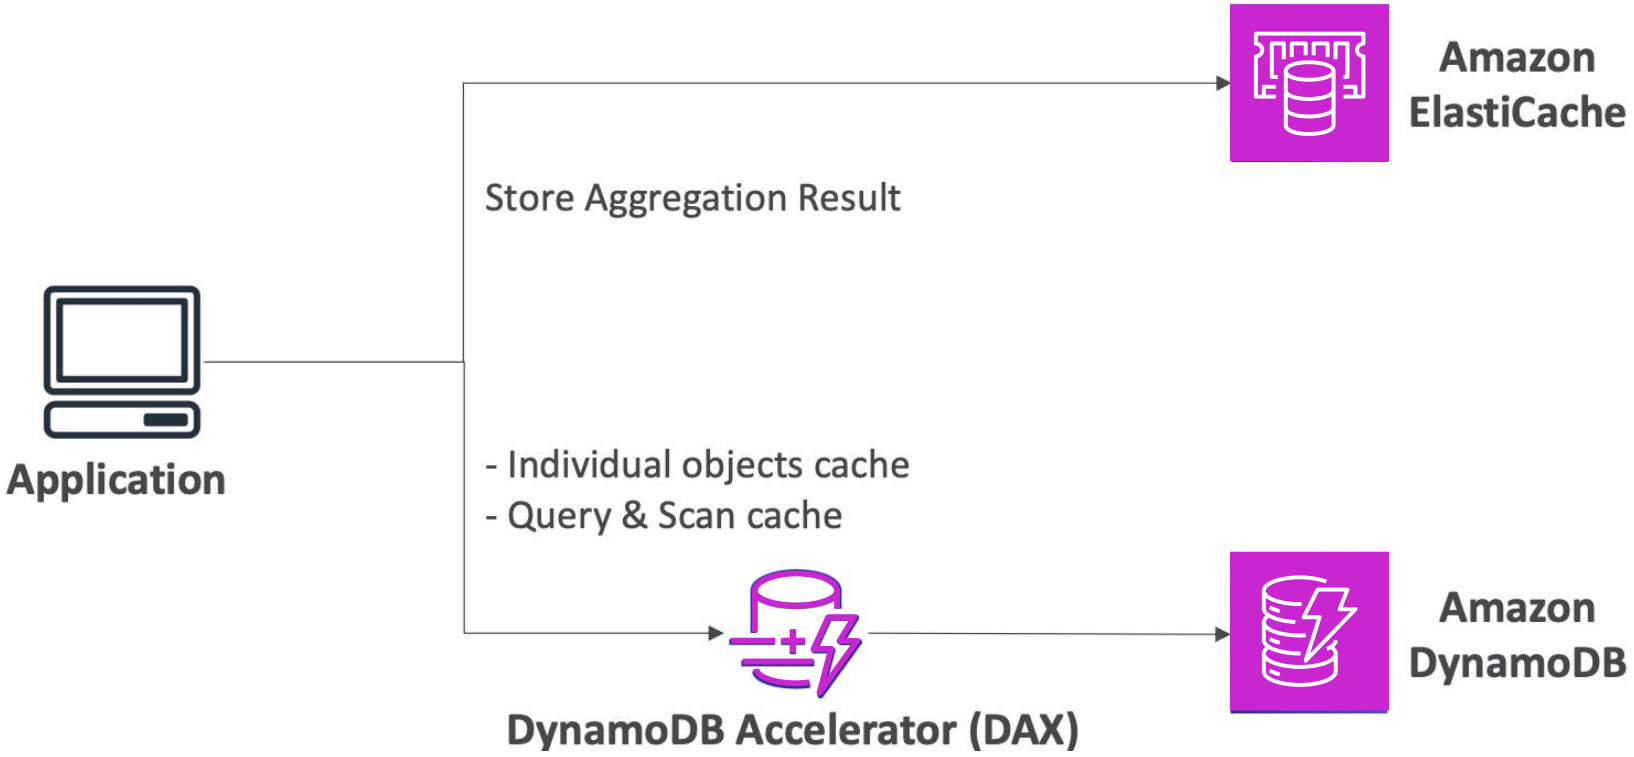

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Streams

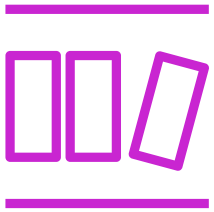

- Ordered stream of item-level modifications (create/update/delete) in a table
- Stream records can be:
  - Sent to Kinesis Data Streams
  - Read by AWS Lambda
  - Read by Kinesis Client Library applications
- Data Retention for up to 24 hours
- Use cases:
  - react to changes in real-time (welcome email to users)
  - Analytics
  - Insert into derivative tables
  - Insert into OpenSearch Service
  - Implement cross-region replication

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Streams

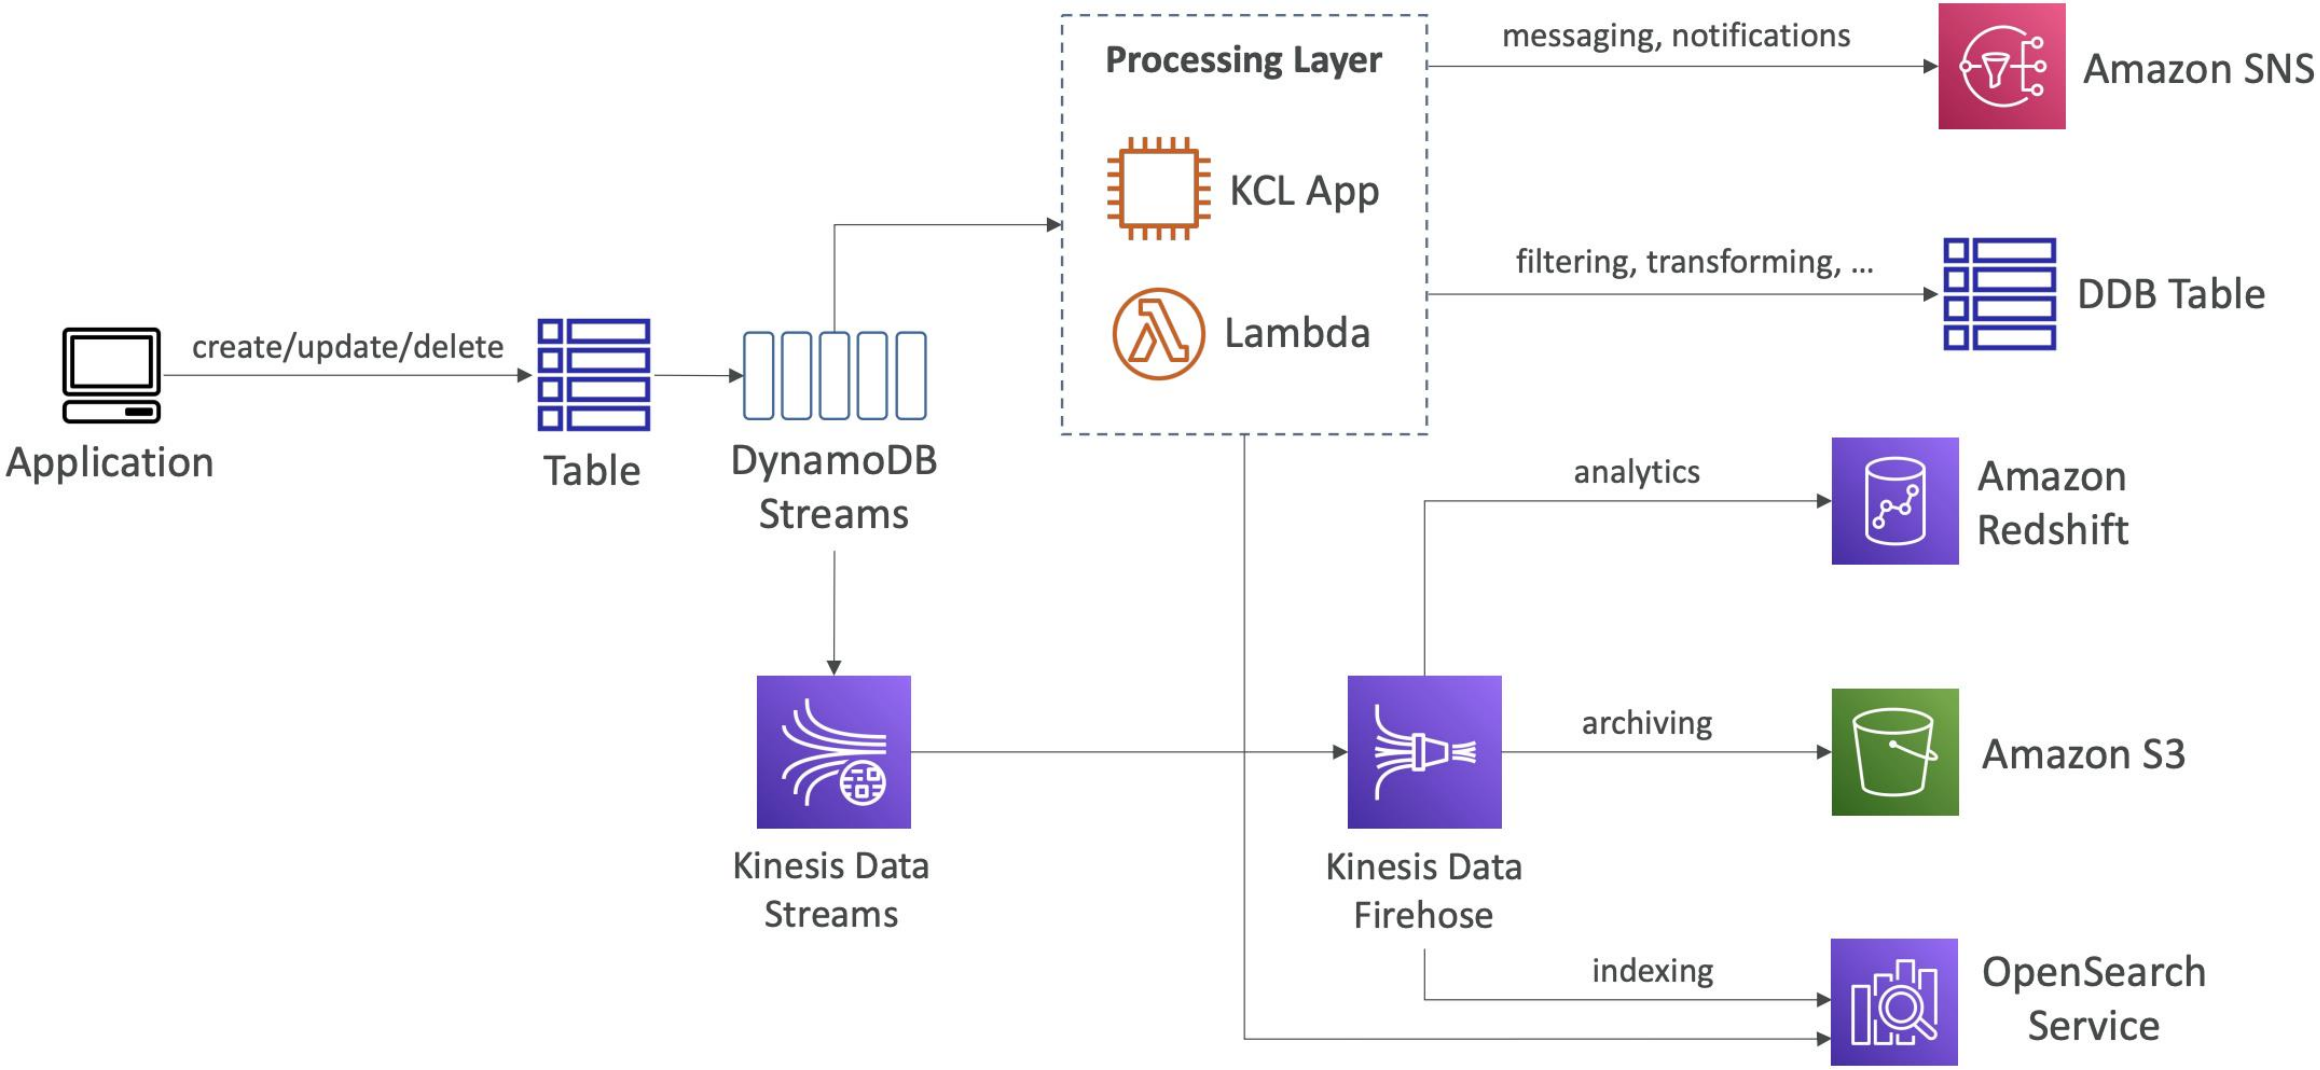

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Streams

- Ability to choose the information that will be written to the stream:
  - **KEYS_ONLY** – only the key attributes of the modified item
  - **NEW_IMAGE** – the entire item, as it appears after it was modified
  - **OLD_IMAGE** – the entire item, as it appeared before it was modified
  - **NEW_AND_OLD_IMAGES** – both the new and the old images of the item
- DynamoDB Streams are made of shards, just like Kinesis Data Streams
- You don't provision shards, this is automated by AWS

- **Records are not retroactively populated in a stream after enabling it**

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Streams & AWS Lambda

- You need to define an **Event Source Mapping** to read from a DynamoDB Streams
- You need to ensure the Lambda function has the **appropriate permissions**
- Your Lambda function is invoked synchronously

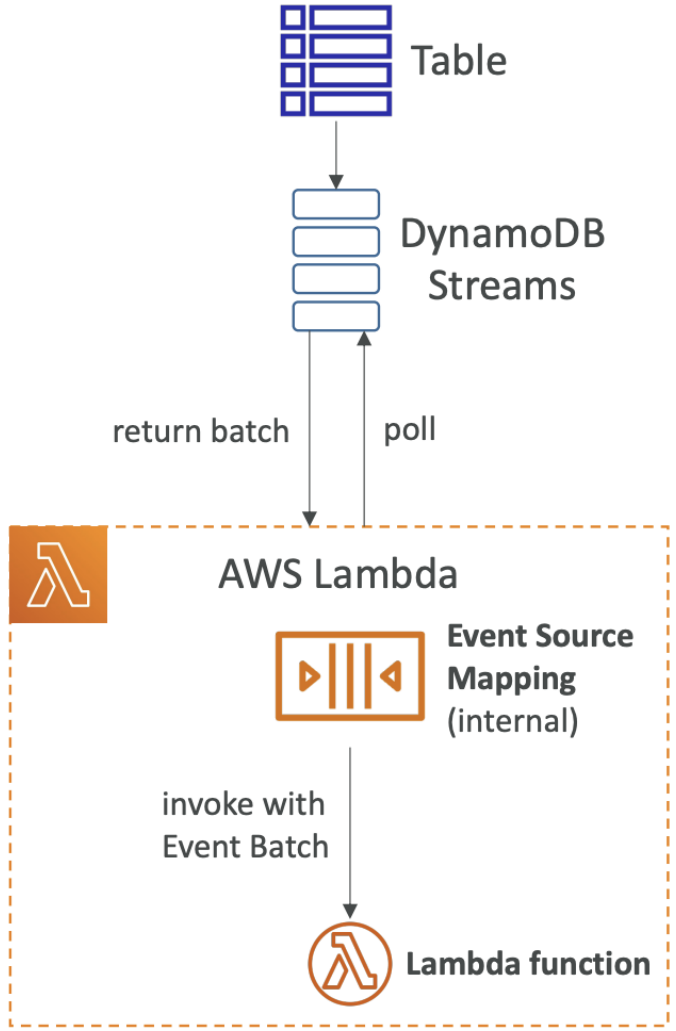

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB –Time To Live (TTL)

- Automatically delete items after an expiry timestamp
- Doesn't consume any WCUs (i.e., no extra cost)
- The TTL attribute must be a "**Number**" data type with "**Unix Epoch timestamp**" value
- Expired items deleted within 48 hours of expiration
- Expired items, that haven't been deleted, appears in reads/queries/scans (if you don't want them, filter them out)
- Expired items are deleted from both LSIs and GSIs
- A delete operation for each expired item enters the DynamoDB Streams (can help recover expired items)
- Use cases: reduce stored data by keeping only current items, adhere to regulatory obligations, …

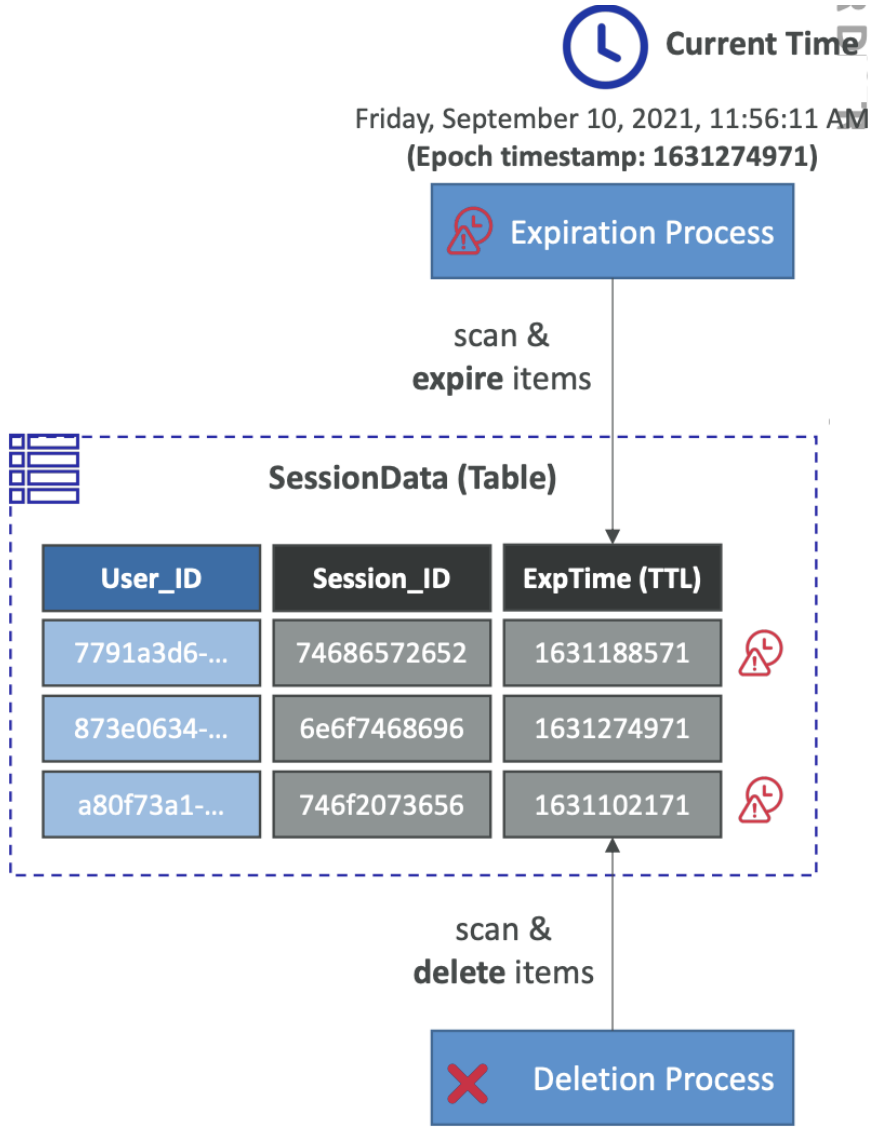

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB CLI – Good to Know

- **--projection-expression**: one or more attributes to retrieve
- **--filter-expression**: filter items before returned to you

- General AWS CLI Pagination options (e.g., DynamoDB, S3, …)
  - **--page-size**: specify that AWS CLI retrieves the full list of items but with a larger number of API calls instead of one API call (default: 1000 items)
  - **--max-items**: max. number of items to show in the CLI (returns **NextToken**)
  - **--starting-token**: specify the last **NextToken** to retrieve the next set of items

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Transactions

- Coordinated, all-or-nothing operations (add/update/delete) to multiple items across one or more tables
- Provides Atomicity, Consistency, Isolation, and Durability (ACID)
- **Read Modes** – Eventual Consistency, Strong Consistency, Transactional
- **Write Modes** – Standard, Transactional
- **Consumes 2x WCUs & RCUs**
  - DynamoDB performs 2 operations for every item (prepare & commit)
- Two operations:
  - **TransactGetItems** – one or more **GetItem** operations
  - **TransactWriteItems** – one or more **PutItem, UpdateItem**, and **DeleteItem** operations
- Use cases: financial transactions, managing orders, multiplayer games, …

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Transactions

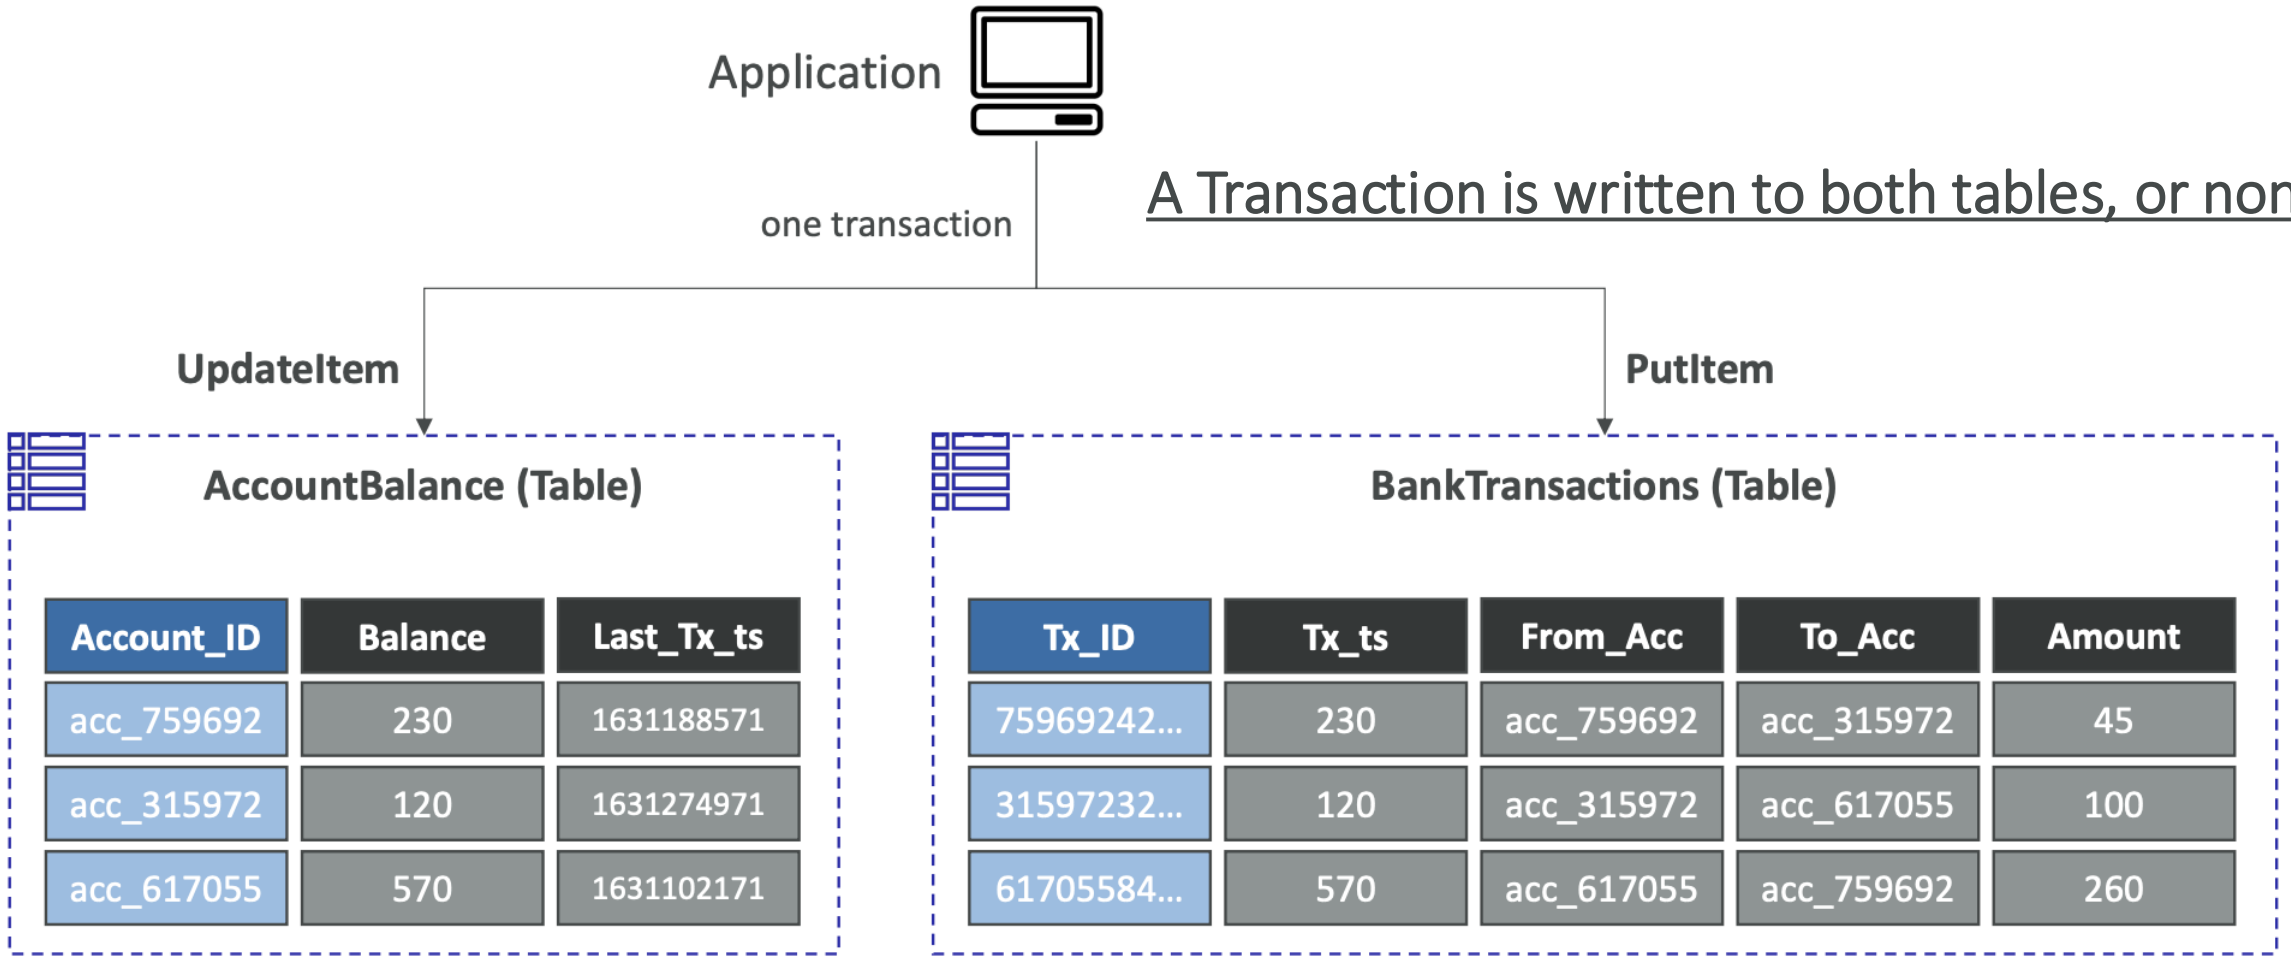

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Transactions – Capacity Computations

- Important for the exam!

- **Example1**: 3 Transactional writes per second, with item size 5 KB
  - We need $3 * (\frac{5\ \text{KB}}{1\ KB})$ * 2 (*transactional cost*) = 30 WCUs

- **Example 2**: 5 Transaction reads per second , with item size 5 KB
  - We need $5 * (\frac{8\ \text{KB}}{4\ KB})$ * 2 (*transactional cost*) = 20 RCUs
  - (5 gets rounded to the upper 4 KB)

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB as Session State Cache

- It’s common to use DynamoDB to store session states

- **vs. ElastiCache**
  - ElastiCache is in-memory, but DynamoDB is serverless
  - Both are key/value stores
- **vs. EFS**
  - EFS must be attached to EC2 instances as a network drive
- **vs. EBS & Instance Store**
  - EBS & Instance Store can only be used for local caching, not shared caching
- **vs. S3**
  - S3 is higher latency, and not meant for small objects

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Write Types

Concurrent Writes

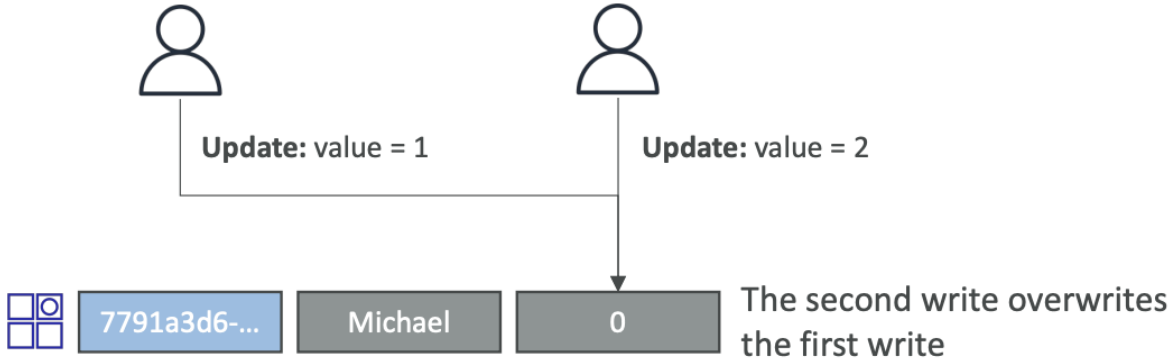

The second write overwrites the first write

Atomic Writes

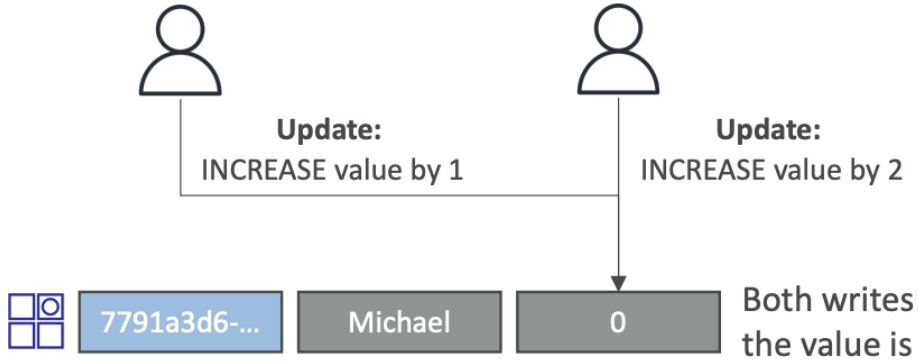

Both writes succeed, the value is increased by 3

Conditional Writes

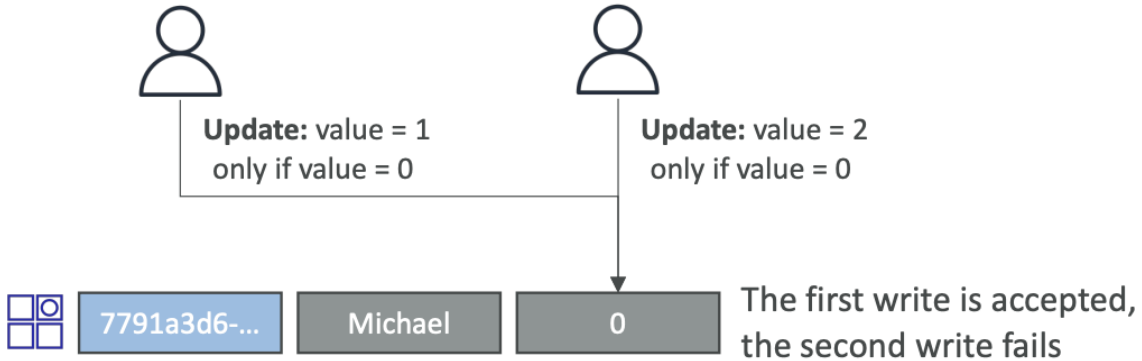

The first write is accepted, the second write fails

Batch Writes

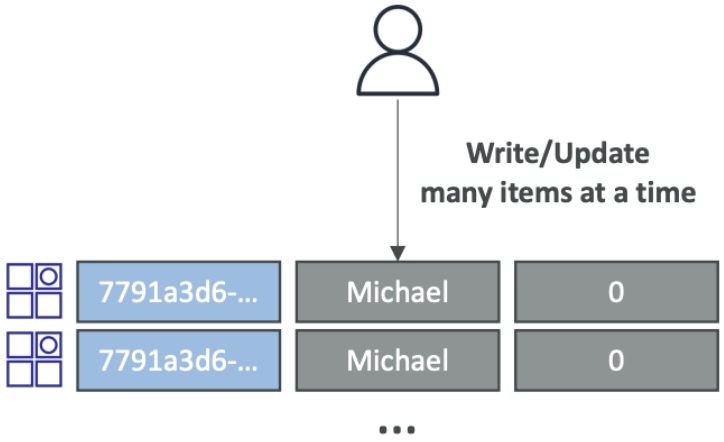

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Large Objects Pattern

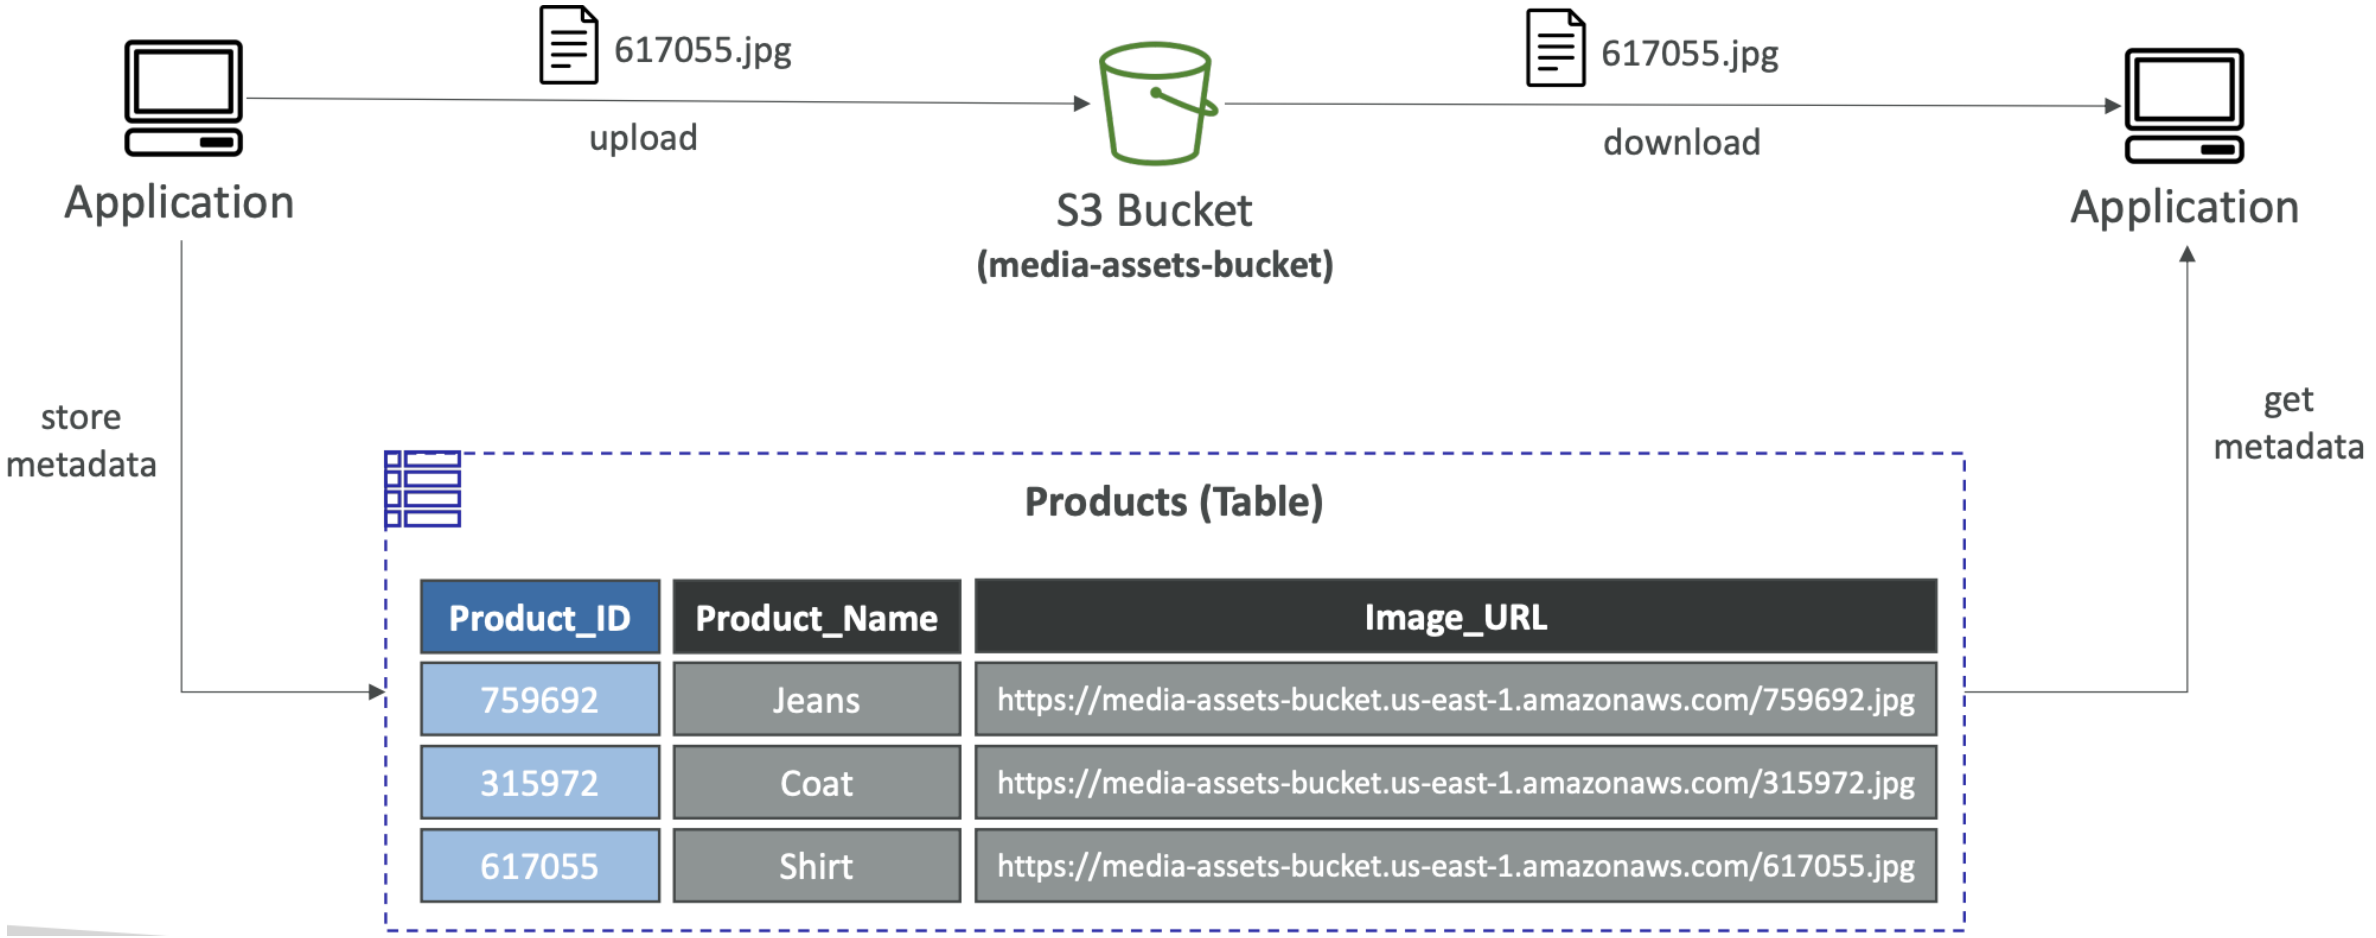

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB Operations

- Table Cleanup
  - Option 1: Scan + DeleteItem
    - Very slow, consumes RCU & WCU, expensive
  - Option 2: Drop Table + Recreate table
    - Fast, efficient, cheap
- Copying a DynamoDB Table
  - Option 1: Using AWS Data Pipeline
  - Option 2: Backup and restore into a new table
    - Takes some time
  - Option 3: Scan + PutItem or BatchWriteItem
    - Write your own code

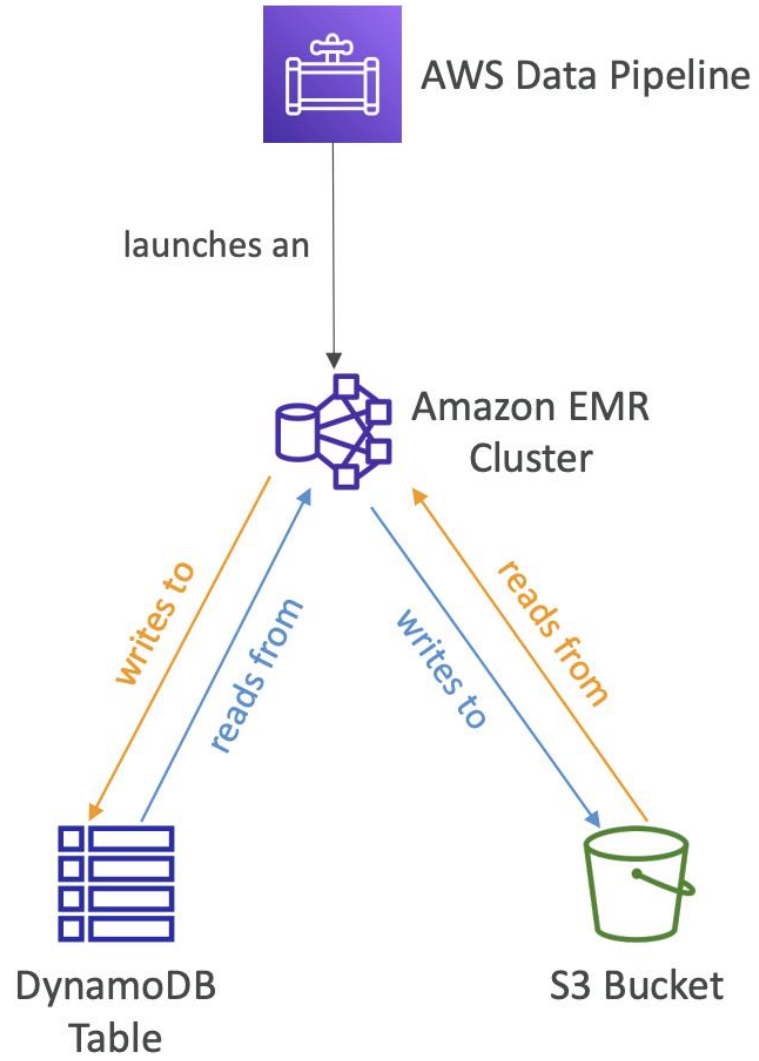

© Copyright by Cloud Mentor Pro | Confidential

DynamoDB – Security & Other Features

- Security
  - VPC Endpoints available to access DynamoDB without using the Internet
  - Access fully controlled by IAM
  - Encryption at rest using AWS KMS and in-transit using SSL/TLS
- Backup and Restore feature available
  - Point-in-time Recovery (PITR) like RDS
  - No performance impact
- Global Tables
  - Multi-region, multi-active, fully replicated, high performance
- DynamoDB Local
  - Develop and test apps locally without accessing the DynamoDB web service (without Internet)
- AWS Database Migration Service (AWS DMS) can be used to migrate to DynamoDB (from MongoDB, Oracle, MySQL, S3, …)

© Copyright by Cloud Mentor Pro | Confidential

Exercise

Use an S3 bucket event to trigger SQS Queue to insert image info into DynamoDB Table

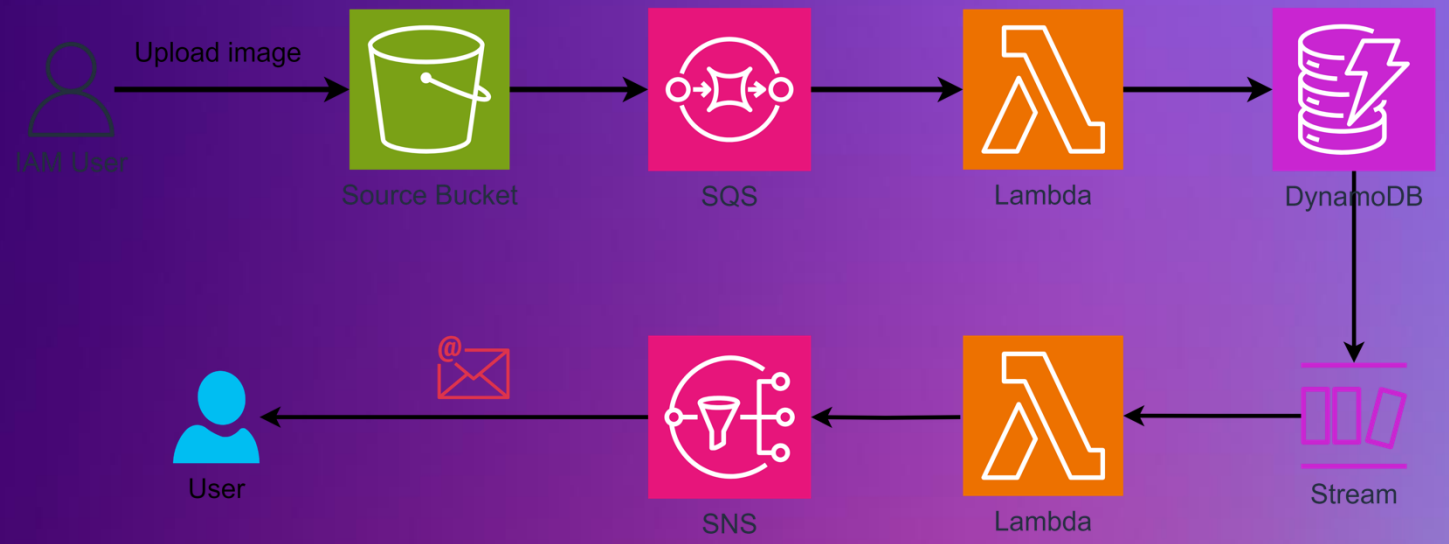

© Copyright by Cloud Mentor Pro | Confidential

Section 13

- *Amazon API Gateway*

© Copyright by Cloud Mentor Pro | Confidential

Amazon API Gateway

Build, Deploy and Manage APIs

© Copyright by Cloud Mentor Pro | Confidential

Example: Building a Serverless API

© Copyright by Cloud Mentor Pro | Confidential

AWS API Gateway

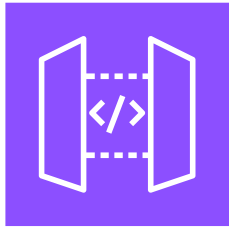

- AWS Lambda + API Gateway: No infrastructure to manage
- Support for the WebSocket Protocol
- Handle API versioning (v1, v2…)
- Handle different environments (dev, test, prod…)
- Handle security (Authentication and Authorization)
- Create API keys, handle request throttling
- Swagger / Open API import to quickly define APIs
- Transform and validate requests and responses
- Generate SDK and API specifications
- Cache API responses

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – HTTP API vs REST API

- HTTP APIs
  - low-latency, cost-effective AWS Lambda proxy, HTTP proxy APIs and private integration (no data mapping)
  - support OIDC and OAuth 2.0 authorization, and built-in support for CORS
  - No usage plans and API keys
- REST APIs
  - All features (except Native OpenID Connect / OAuth 2.0)

Authorizers	HTTP API	REST API
AWS Lambda	✓	✓
IAM	✓	✓
Resource Policies		✓
Amazon Cognito	✓ *	✓
Native OpenID Connect / OAuth 2.0 / JWT	✓	

Full list here: https://docs.aws.amazon.com/apigateway/latest/developerguide/http-api-vs-rest.html

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Integrations High Level

- Lambda Function
  - Invoke Lambda function
  - Easy way to expose REST API backed by AWS Lambda
- HTTP
  - Expose HTTP endpoints in the backend
  - Example: internal HTTP API on premise, Application Load Balancer…
  - Why? Add rate limiting, caching, user authentications, API keys, etc…
- AWS Service
  - Expose any AWS API through the API Gateway
  - Example: start an AWS Step Function workflow, post a message to SQS
  - Why? Add authentication, deploy publicly, rate control…

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Architecture

- Create a single interface for all the microservices in your company
- Use API endpoints with various resources
- Apply a simple domain name and SSL certificates
- Can apply forwarding and transformation rules at the API Gateway level

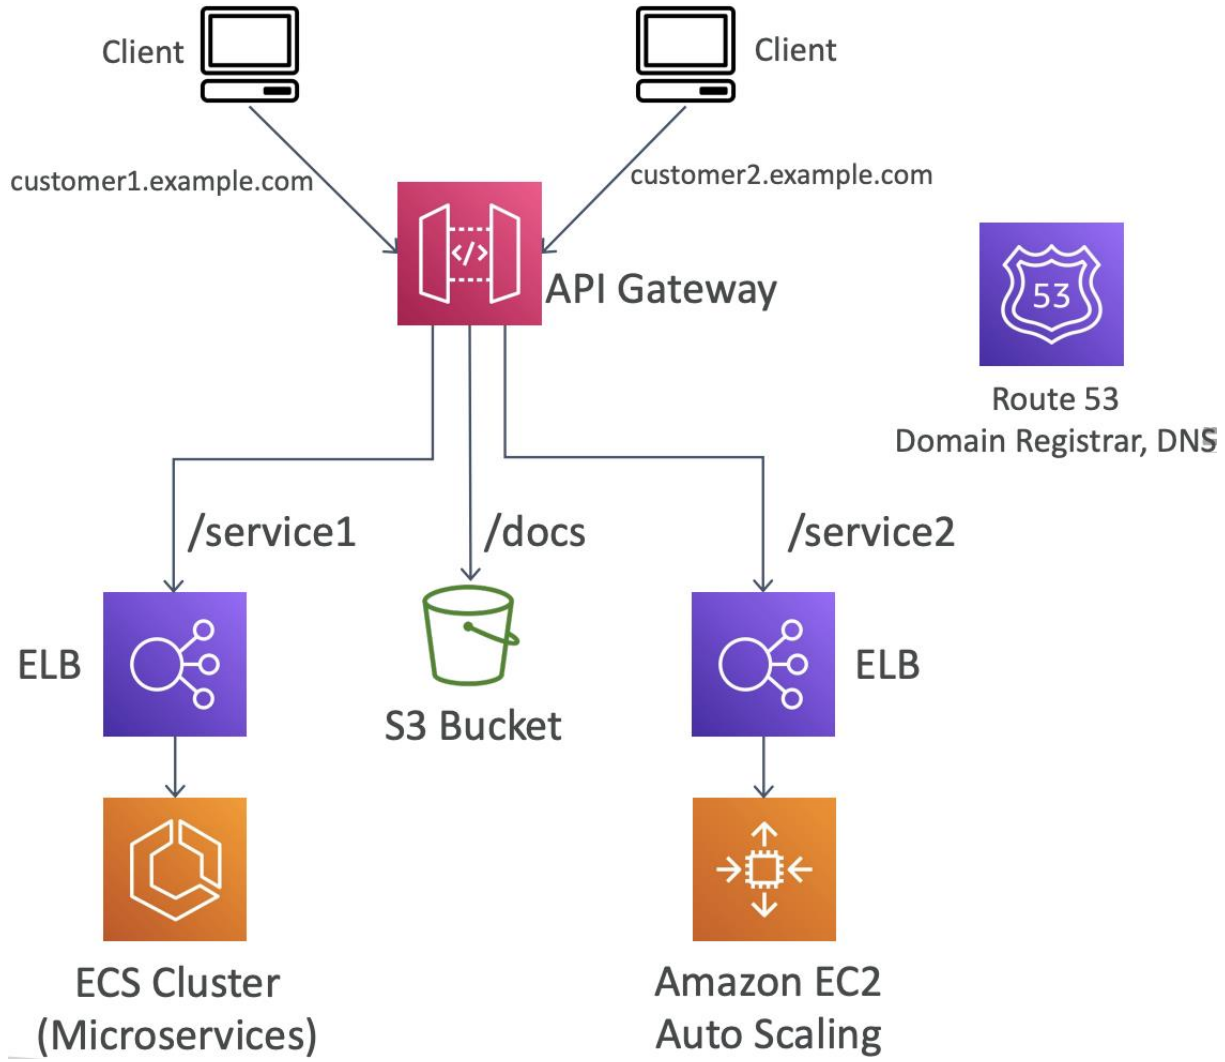

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – AWS Service Integration Kinesis Data Streams example

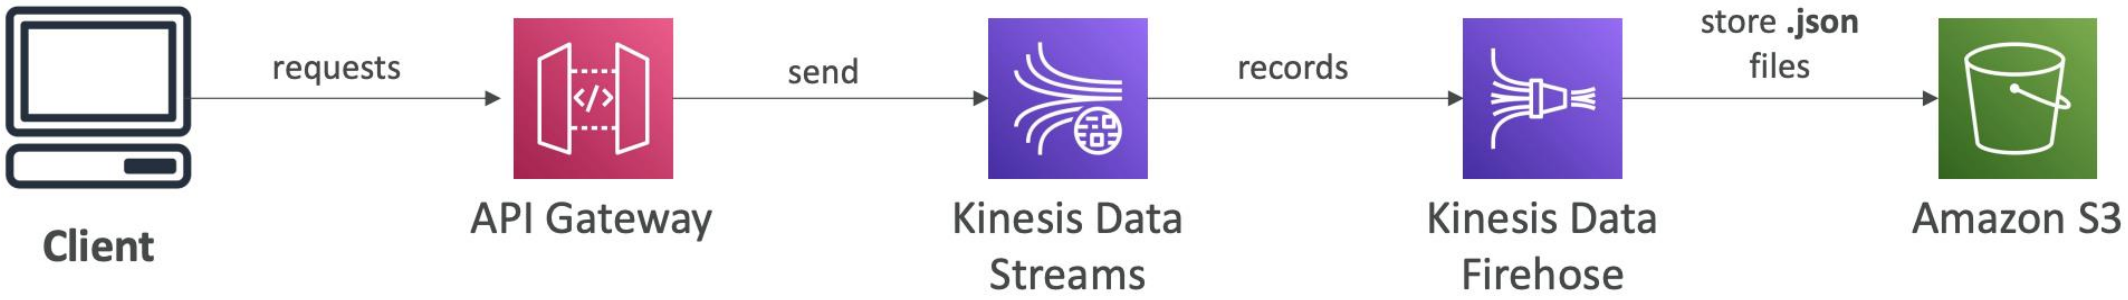

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Endpoint Types

- **Edge-Optimized (default):** For global clients
  - Requests are routed through the CloudFront Edge locations (improves latency)
  - The API Gateway still lives in only one region
- **Regional**:
  - For clients within the same region
  - Could manually combine with CloudFront (more control over the caching strategies and the distribution)
- **Private**:
  - Can only be accessed from your VPC using an interface VPC endpoint (ENI)
  - Use a resource policy to define access

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Security

- **User Authentication through**
  - IAM Roles (useful for internal applications)
  - Cognito (identity for external users – example mobile users)
  - Custom Authorizer (your own logic)
- **Custom Domain Name HTTPS** security through integration with AWS Certificate Manager (ACM)
  - If using Edge-Optimized endpoint, then the certificate must be in **us-east-1**
  - If using Regional endpoint, the certificate must be in the API Gateway region
  - Must setup CNAME or A-alias record in Route 53

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Deployment Stages

- Making changes in the API Gateway does not mean they're effective
- You need to make a "deployment" for them to be in effect
- It's a common source of confusion
- Changes are deployed to "Stages" (as many as you want)
- Use the naming you like for stages (dev, test, prod)
- Each stage has its own configuration parameters
- Stages can be rolled back as a history of deployments is kept

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Stages v1 and v2 API breaking change

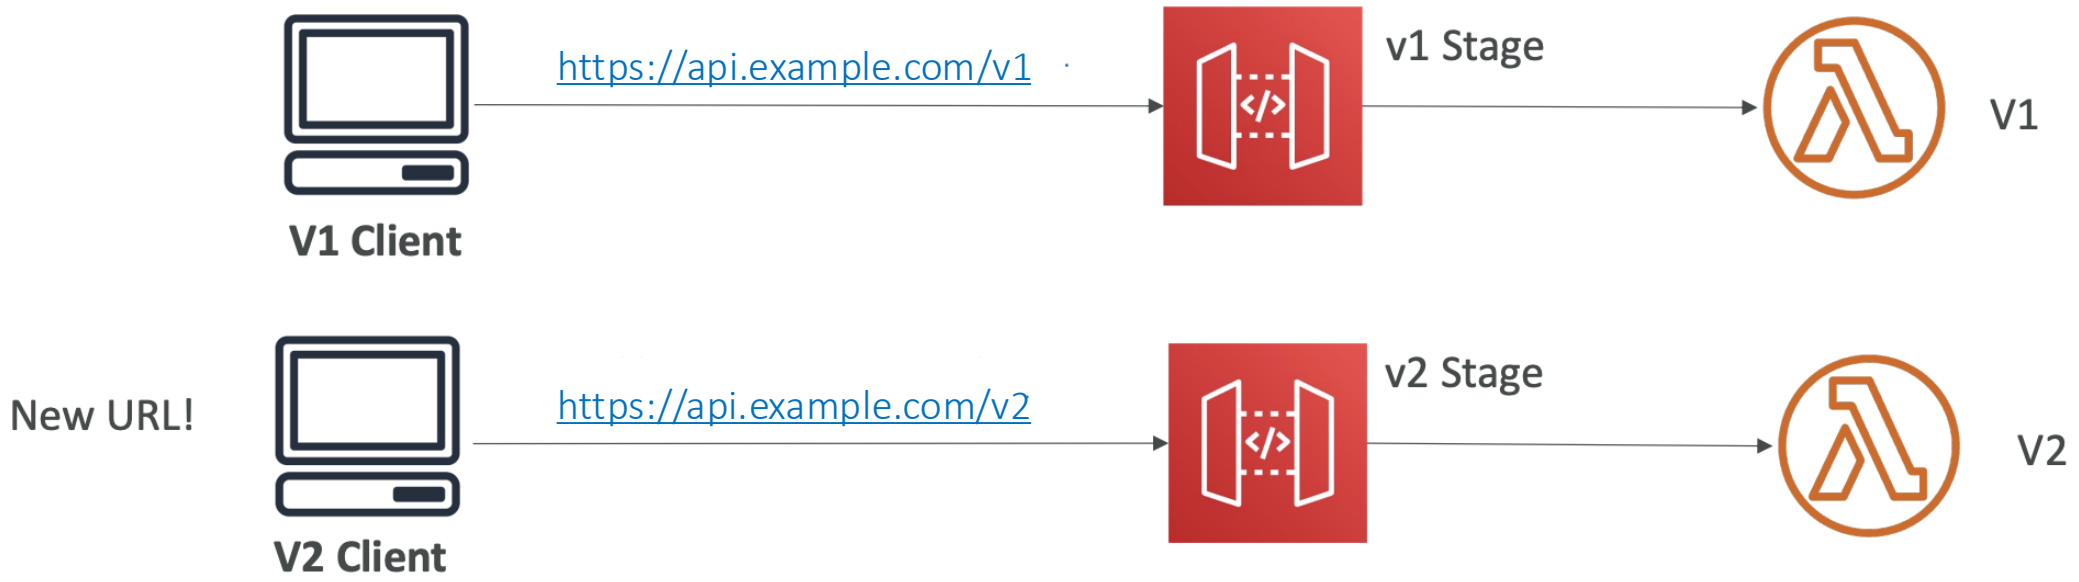

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Stage Variables

- Stage variables are like environment variables for API Gateway
- Use them to change often changing configuration values
- They can be used in:
  - Lambda function ARN
  - HTTP Endpoint
  - Parameter mapping templates
- Use cases:
  - Configure HTTP endpoints your stages talk to (dev, test, prod…)
  - Pass configuration parameters to AWS Lambda through mapping templates
- Stage variables are passed to the ”context” object in AWS Lambda
- Format: ${stageVariables.variableName}

© Copyright by Cloud Mentor Pro | Confidential

API Gateway Stage Variables & Lambda Aliases

- We create a **stage variable** to indicate the corresponding Lambda alias
- Our API gateway will automatically invoke the right Lambda function!

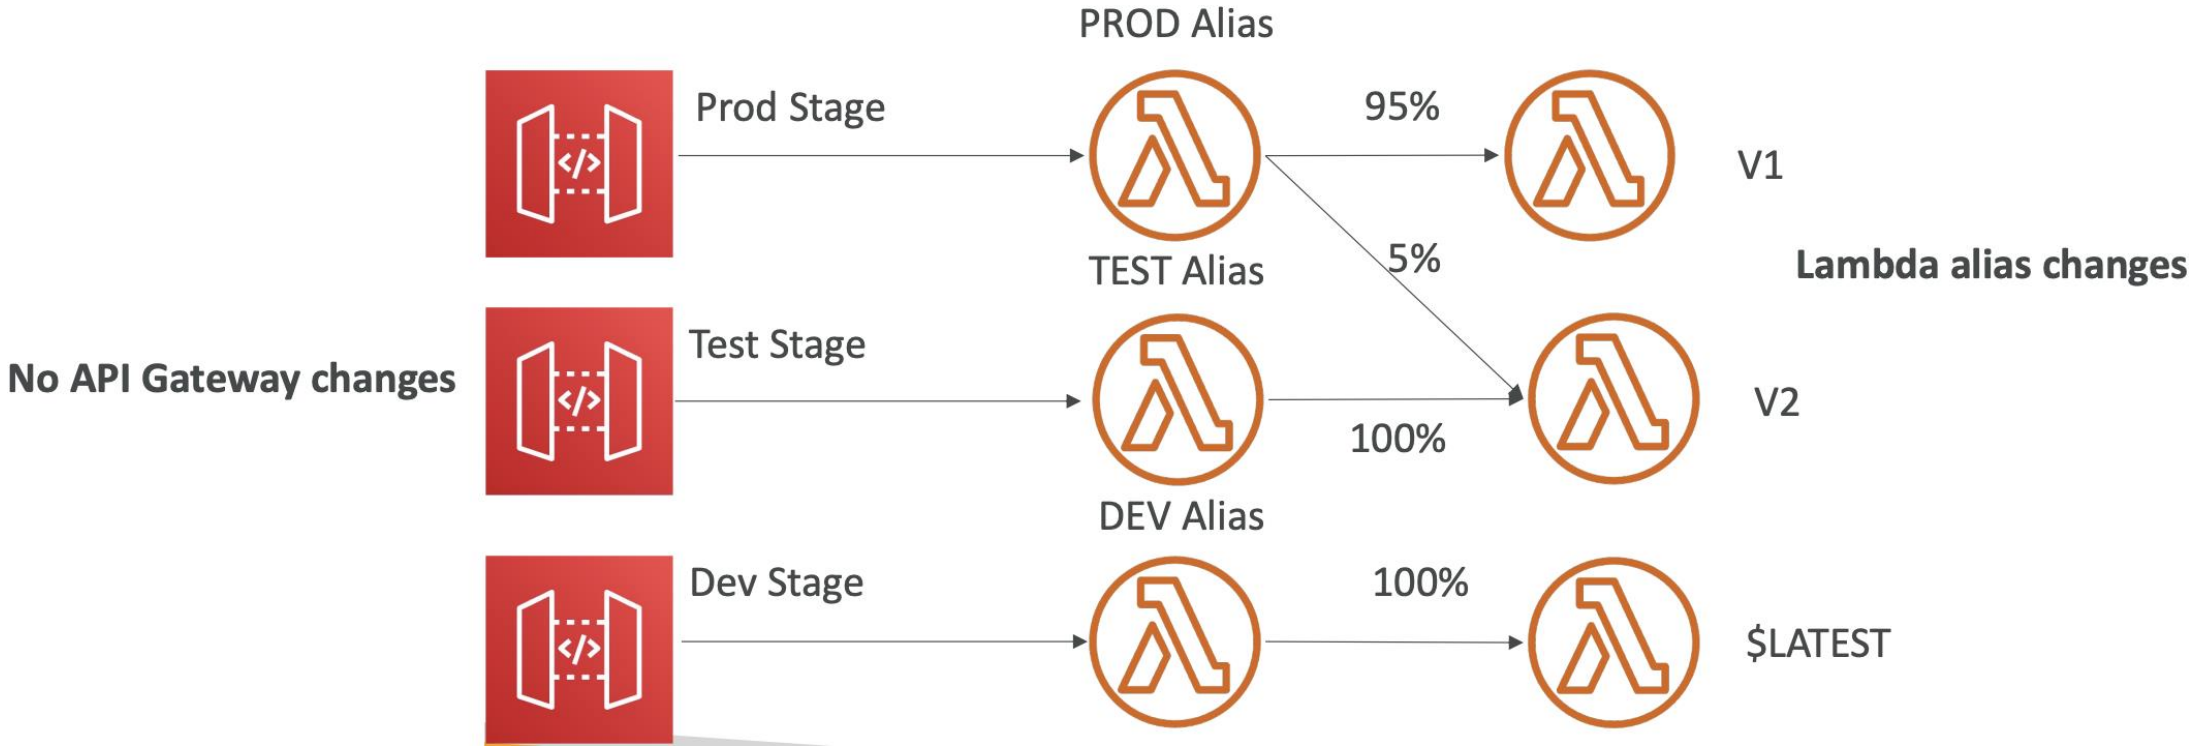

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Canary Deployment

- Possibility to enable canary deployments for any stage (usually prod)
- Choose the % of traffic the canary channel receives

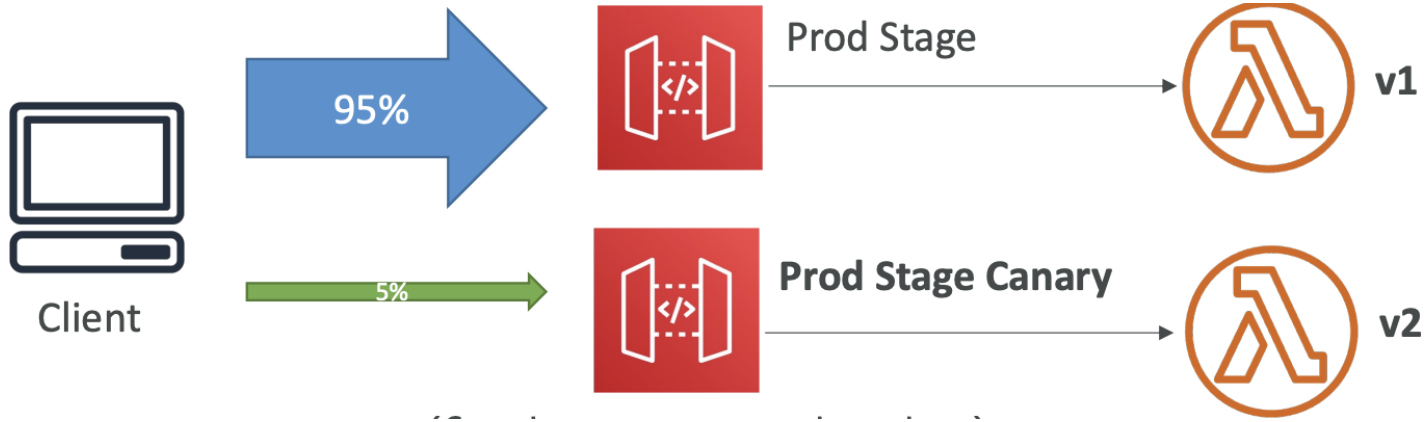

- Metrics & Logs are separate (for better monitoring)
- Possibility to override stage variables for canary
- This is blue / green deployment with AWS Lambda & API Gateway

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Integration Types

- Integration Type **MOCK**
  - API Gateway returns a response without sending the request to the backend

- Integration Type **HTTP / AWS (Lambda & AWS Services)**
  - you must configure both the integration request and integration response
  - Setup data mapping using **mapping templates** for the request & response

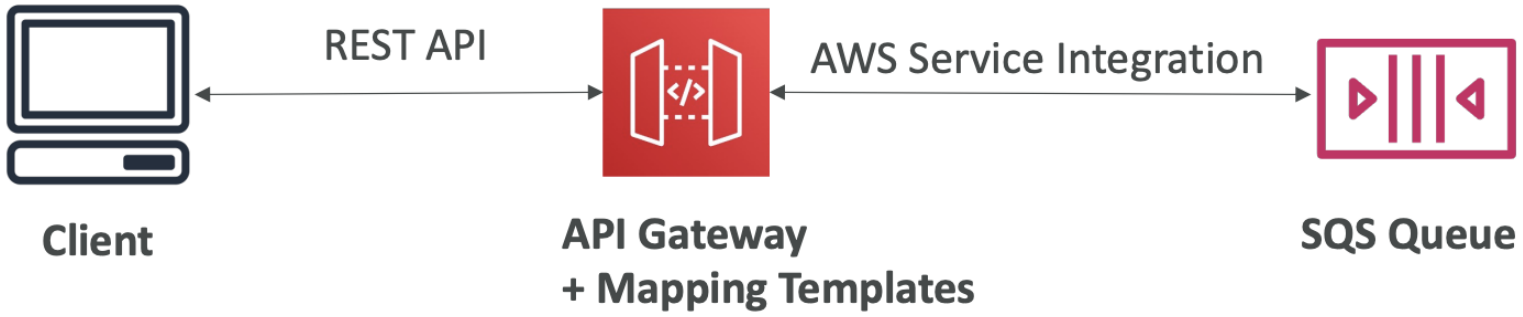

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Integration Types

- Integration Type **AWS_PROXY (Lambda Proxy)**:
  - incoming request from the client is the input to Lambda
  - The function is responsible for the logic of request / response
  - **No mapping template, headers, query string parameters… are passed as arguments**

```
{
    "resource": "Resource path",
    "path": "Path parameter",
    "httpMethod": "Incoming request's method name",
    "headers": "String containing incoming request headers",
    "multiValueHeaders": "List of strings containing incomin
    "queryStringParameters": "query string parameters ",
    "multiValueQueryStringParameters": "List of query string
    "pathParameters":  "path parameters",
    "stageVariables": "Applicable stage variables",
    "requestContext": "Request context, including authorizer
    "body": "A JSON string of the request payload.",
    "isBase64Encoded": "A boolean flag"
}
```

Lambda function invocation payload

```
{
    "isBase64Encoded": "true|false",
    "statusCode": "httpStatusCode",
    "headers": { "headerName": "headerValue", ... },
    "multiValueHeaders": { "headerName": ["headerValue", "heade
    "body": "..."
}
```

Lambda function expected response

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Integration Types

- Integration Type **HTTP_PROXY**
  - No mapping template
  - The HTTP request is passed to the backend
  - The HTTP response from the backend is forwarded by API Gateway
  - Possibility to add HTTP Headers if need be (ex: API key)

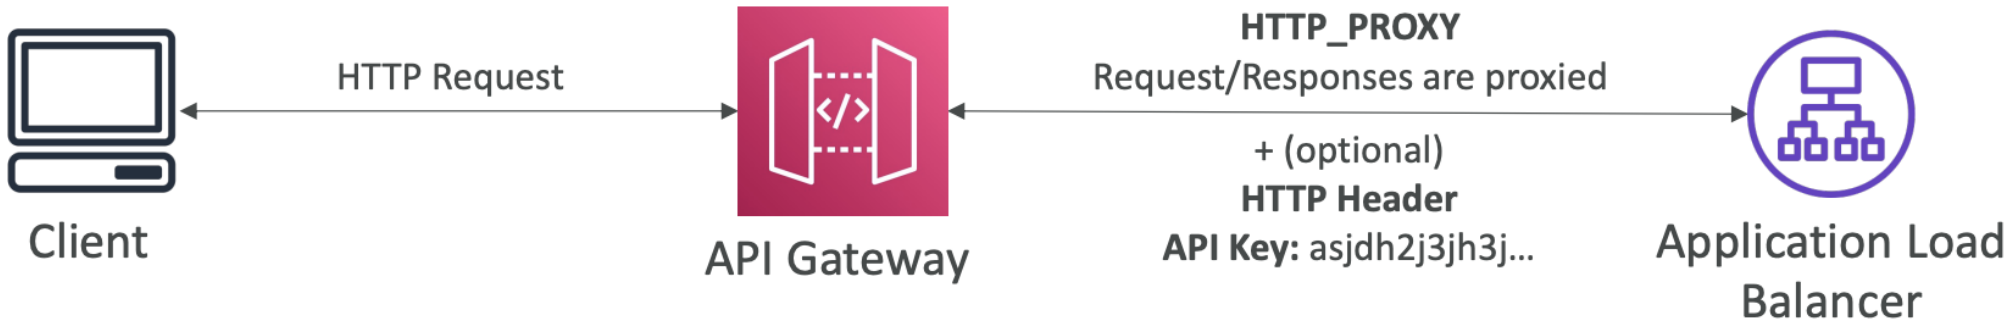

© Copyright by Cloud Mentor Pro | Confidential

Mapping Templates (AWS & HTTP Integration)

- Mapping templates can be used to modify request / responses
- Rename / Modify **query string parameters**
- Modify **body content**
- **Add headers**
- Uses Velocity Template Language (VTL): for loop, if etc…
- Filter output results (remove unnecessary data)
- Content-Type can be set to **application/json** or **application/xml**

© Copyright by Cloud Mentor Pro | Confidential

Mapping Example: JSON to XML with SOAP

- SOAP API are XML based, whereas REST API are JSON based

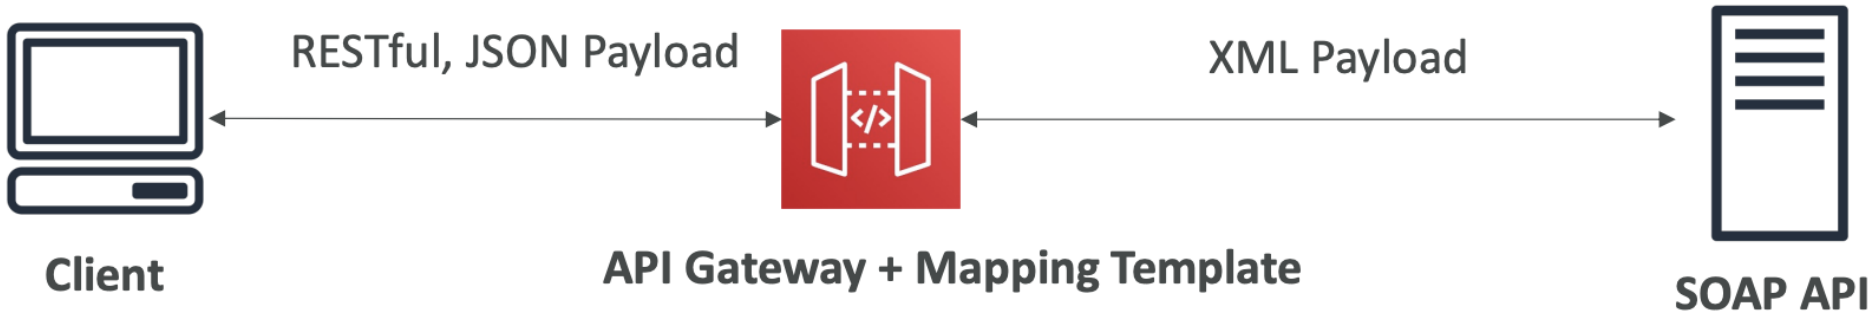

- In this case, API Gateway should:
  - Extract data from the request: either path, payload or header
  - Build SOAP message based on request data (mapping template)
  - Call SOAP service and receive XML response
  - Transform XML response to desired format (like JSON), and respond to the user

© Copyright by Cloud Mentor Pro | Confidential

Mapping Example: Query String parameters

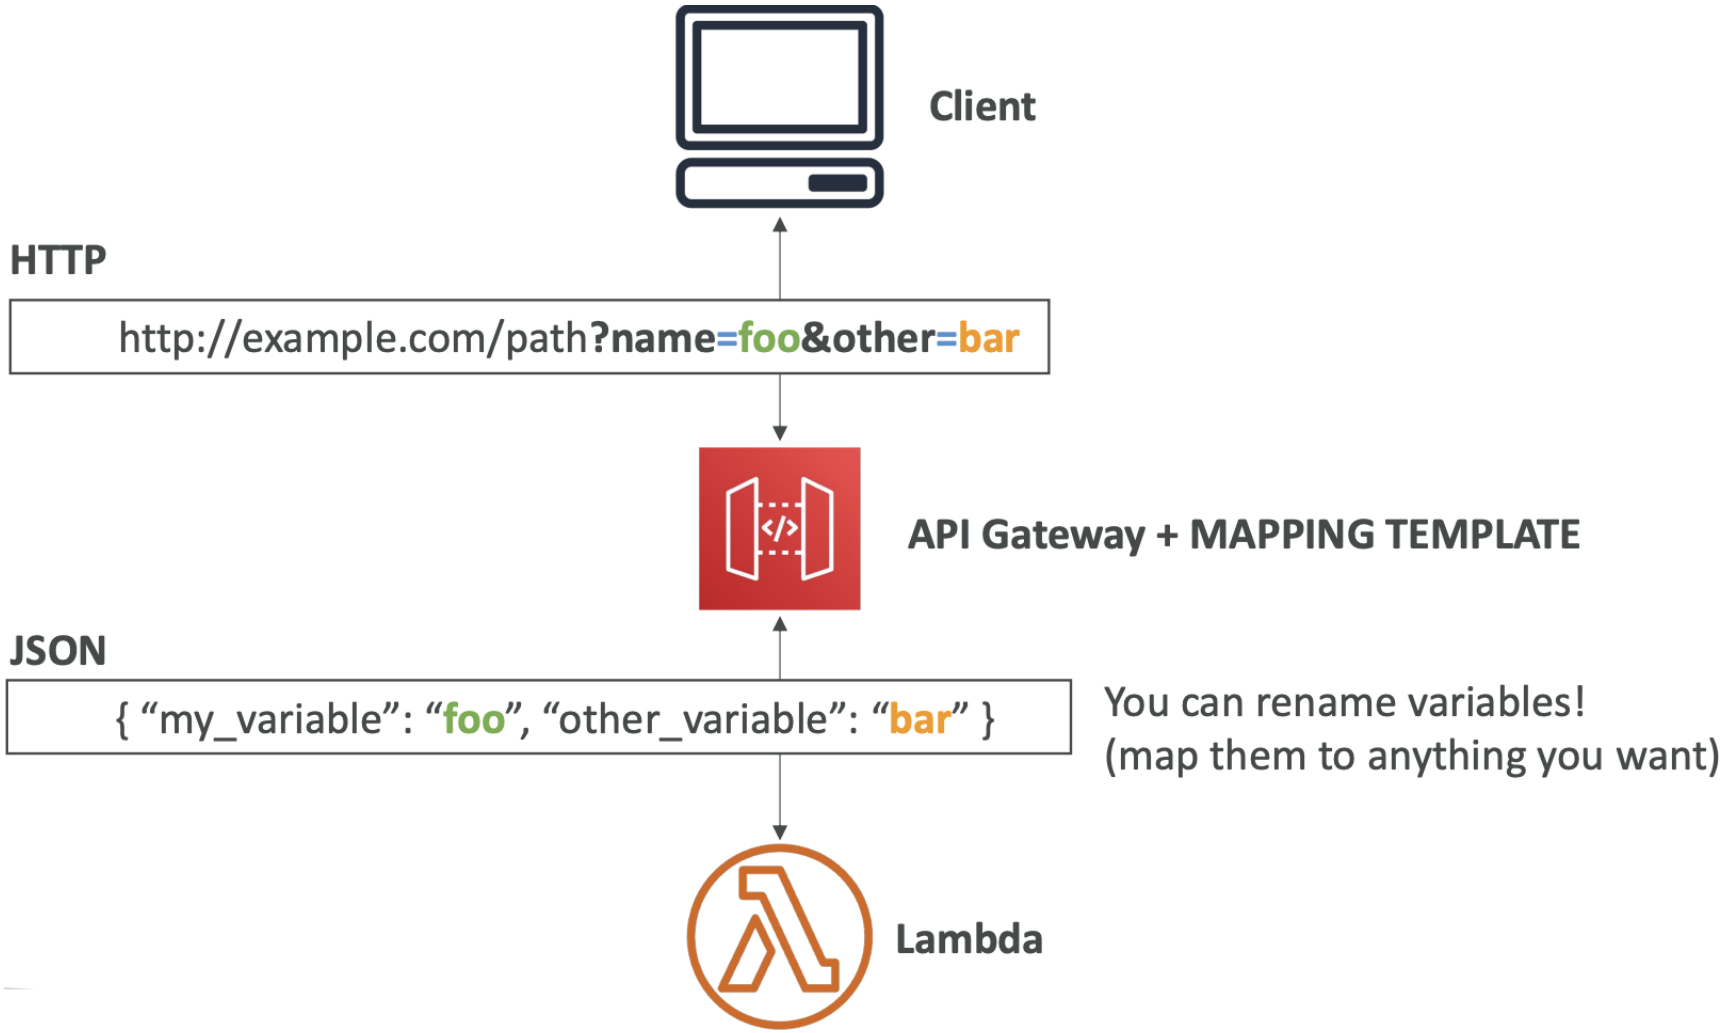

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Open API spec

- Common way of defining REST APIs, using API definition as code
- Import existing OpenAPI 3.0 spec to API Gateway
  - Method
  - Method Request
  - Integration Request
  - Method Response
  - + AWS extensions for API gateway and setup every single option
- Can export current API as OpenAPI spec
- OpenAPI specs can be written in YAML or JSON
- Using OpenAPI we can generate SDK for our applications

© Copyright by Cloud Mentor Pro | Confidential

REST API – Request Validation

- You can configure API Gateway to perform basic validation of an API request before proceeding with the integration request
- When the validation fails, API Gateway immediately fails the request
  - Returns a 400-error response to the caller
- This reduces unnecessary calls to the backend
- Checks:
  - The required request parameters in the URI, query string, and headers of an incoming request are included and non-blank
  - The applicable request payload adheres to the configured JSON Schema request model of the method

© Copyright by Cloud Mentor Pro | Confidential

REST API – RequestValidation – OpenAPI

- Setup request validation by importing OpenAPI definitions file

```
{
    "openapi": "3.0.0",
    "info": {
        "title": "ReqValidation Sample",
        "version": "1.0.0"
    },
    "servers": [ ... ],
    "x-amazon-apigateway-request-validators": {
        "all": {
            "validateRequestBody": true,
            "validateRequestParameters": true
        },
        "params-only": {
            "validateRequestBody": false,
            "validateRequestParameters": true
        }
    }
}
```

Defines the Validators

```
{
    "openapi": "3.0.0",
    "info": {
        "title": "ReqValidation Sample",
        "version": "1.0.0"
    },
    "servers": [ ... ],
    "x-amazon-apigateway-request-validator": "params-only",
    ...
}
```

Enable **params-only** Validator on all API methods

```
{
    "openapi": "3.0.0",
    "info": {
        "title": "ReqValidation Sample",
        "version": "1.0.0"
    },
    "servers": [ ... ],
    "paths": {
        "/validation": {
            "post": {
                "x-amazon-apigateway-request-validator": "all"
            }
        }
    },
    ...
}
```

Enable **all** Validator on the **POST /validation** method (overrides the **params-only** validator inherited from the API)

© Copyright by Cloud Mentor Pro | Confidential

Caching API responses

- Caching reduces the number of calls made to the backend
- Default **TTL** (time to live) is 300 seconds (min: 0s, max: 3600s)
- **Caches are defined per stage**
- Possible to override cache settings **per method**
- Cache encryption option
- Cache capacity between 0.5GB to 237GB
- Cache is expensive, makes sense in production, may not make sense in dev / test

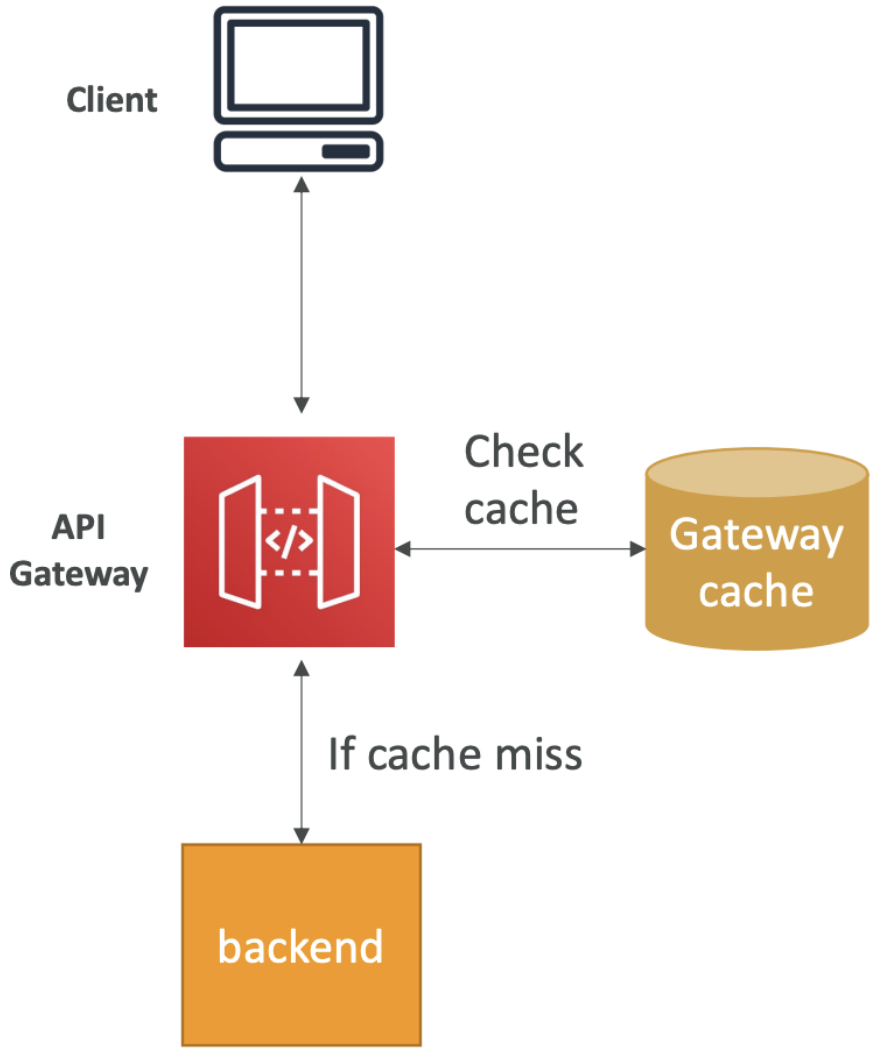

© Copyright by Cloud Mentor Pro | Confidential

API Gateway Cache Invalidation

- Able to flush the entire cache (invalidate it) immediately
- Clients can invalidate the cache with **header: Cache-Control: max-age=0** (with proper IAM authorization)
- If you don't impose an InvalidateCache policy (or choose the Require authorization check box in the console), any client can invalidate the API cache

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "execute-api:InvalidateCache"
            ],
            "Resource": [
                "arn:...:api-id/stage-name/GET/resource-path-specifier"
            ]
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Usage Plans & API Keys

- If you want to make an API available as an offering ($) to your customers
- **Usage Plan:**
  - who can access one or more deployed API stages and methods
  - how much and how fast they can access them
  - uses API keys to identify API clients and meter access
  - configure throttling limits and quota limits that are enforced on individual client
- **API Keys:**
  - alphanumeric string values to distribute to your customers
  - Ex: WBjHxNtoAb4WPKBC7cGm64CBibIb24b4jt8jJHo9
  - Can use with usage plans to control access
  - Throttling limits are applied to the API keys
  - Quotas limits is the overall number of maximum requests

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Correct Order for API keys

- To configure a usage plan

1. Create one or more APIs, configure the methods to require an API key, and deploy the APIs to stages.
2. Generate or import API keys to distribute to application developers (your customers) who will be using your API.
3. Create the usage plan with the desired throttle and quota limits.
4. Associate API stages and API keys with the usage plan.

- Callers of the API must supply an assigned API key in the x-api-key header in requests to the API.

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Logging & Tracing

- CloudWatch Logs
  - Log contains information about request/response body
  - Enable CloudWatch logging at the Stage level (with Log Level - ERROR, DEBUG, INFO)
  - Can override settings on a per API basis

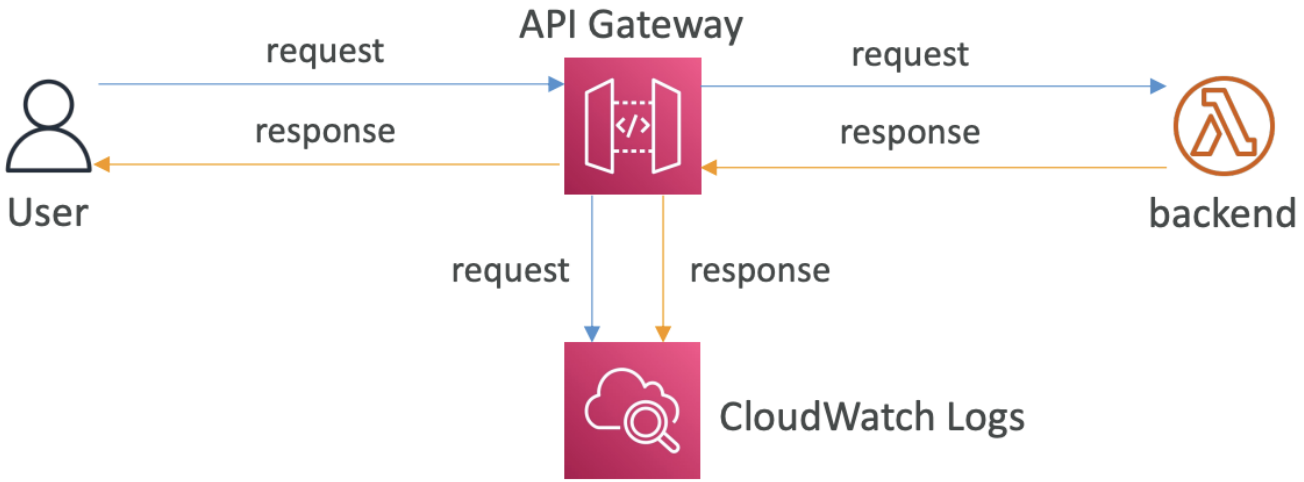

- X-Ray
  - Enable tracing to get extra information about requests in API Gateway
  - X-Ray API Gateway + AWS Lambda gives you the full picture

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – CloudWatch Metrics

- Metrics are by stage, Possibility to enable detailed metrics
- **CacheHitCount** & **CacheMissCount**: efficiency of the cache
- **Count**: The total number API requests in a given period.
- **IntegrationLatency**: The time between when API Gateway relays a request to the backend and when it receives a response from the backend.
- **Latency**: The time between when API Gateway receives a request from a client and when it returns a response to the client. The latency includes the integration latency and other API Gateway overhead.
- **4XXError** (client-side) & **5XXError** (server-side)

© Copyright by Cloud Mentor Pro | Confidential

API Gateway Throttling

- **Account Limit**
  - API Gateway throttles requests at **10,000** rps across all API
  - Soft limit that can be increased upon request
- In case of throttling => **429 Too Many Requests** (retriable error)
- Can set **Stage limit & Method limits** to improve performance
- Or you can define **Usage Plans** to throttle per customer

- Just like Lambda Concurrency, one API that is overloaded, if not limited, can cause the other APIs to be throttled

© Copyright by Cloud Mentor Pro | Confidential

API Gateway - Errors

- 4xx means Client errors
  - 400: Bad Request
  - 403: Access Denied, WAF filtered
  - 429: Quota exceeded, Throttle

- 5xx means Server errors
  - 502: Bad Gateway Exception, usually for an incompatible output returned from a Lambda proxy integration backend and occasionally for out-of-order invocations due to heavy loads.
  - 503: Service Unavailable Exception
  - 504: Integration Failure – ex Endpoint Request Timed-out Exception
    API Gateway requests time out after 29 second maximum

© Copyright by Cloud Mentor Pro | Confidential

AWS API Gateway - CORS

- CORS must be enabled when you receive API calls from another domain.
- The OPTIONS pre-flight request must contain the following headers:
  - Access-Control-Allow-Methods
  - Access-Control-Allow-Headers
  - Access-Control-Allow-Origin
- CORS can be enabled through the console

© Copyright by Cloud Mentor Pro | Confidential

CORS – Enabled on the API Gateway

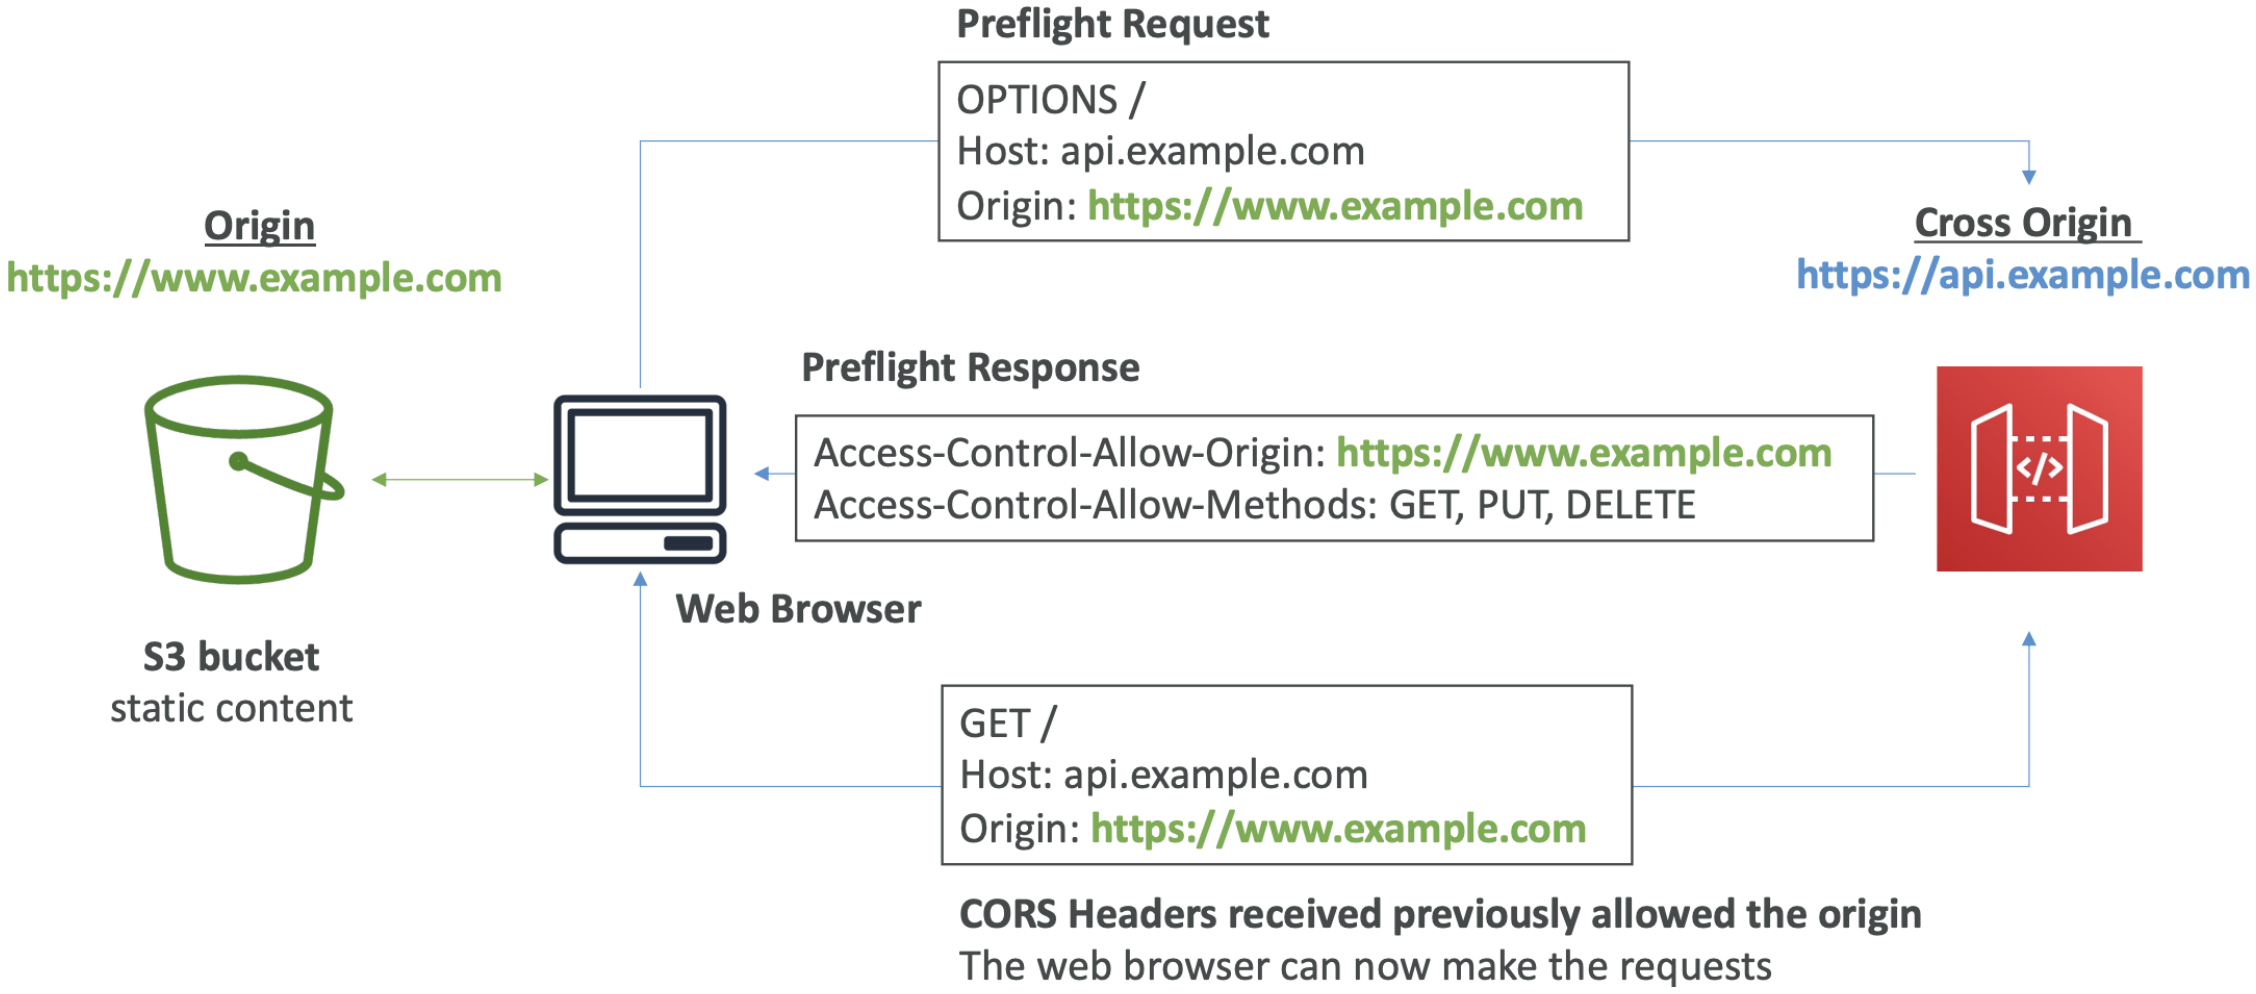

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Security
IAM Permissions

- Create an IAM policy authorization and attach to User / Role
- **Authentication = IAM | Authorization = IAM Policy**
- Good to provide access within AWS (EC2, Lambda, IAM users...)
- Leverages "Sig v4" capability where IAM credential are in headers

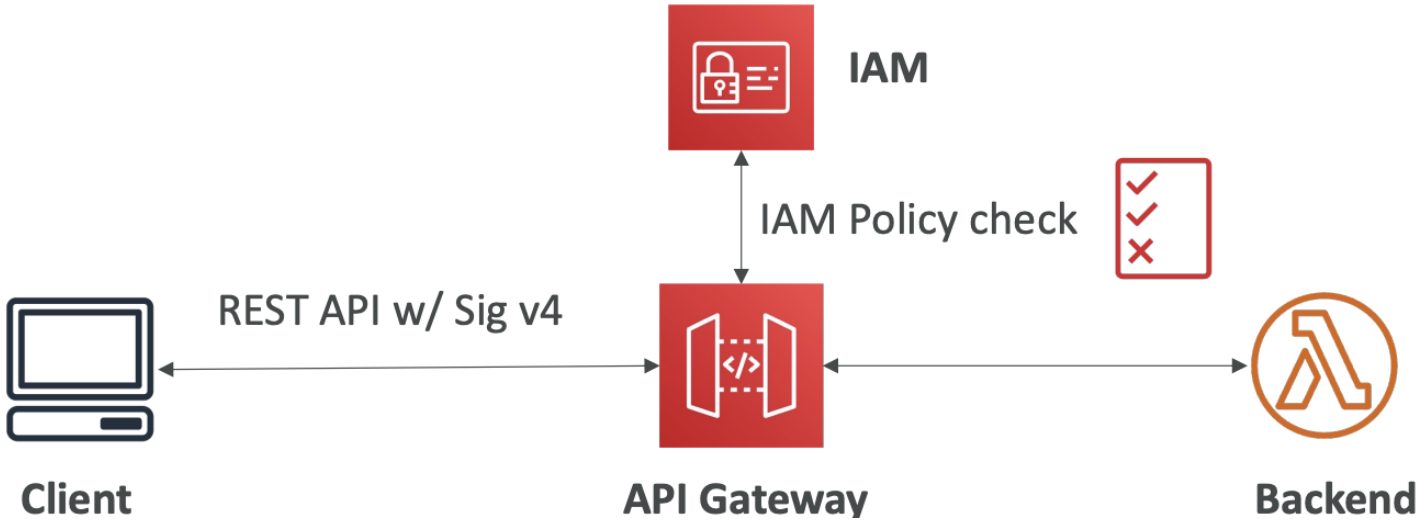

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Resource Policies

- **Resource policies** (similar to Lambda Resource Policy)

- **Allow for Cross Account Access (combined with IAM Security)**
- Allow for a specific source IP address
- Allow for a VPC Endpoint

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "AWS": [
                    "arn:aws:iam::account-id-2:user/Alice",
                    "account-id-2"
                ]
            },
            "Action": "execute-api:Invoke",
            "Resource": [
                "arn:aws:execute-api:region:account-id-1:api-id/stage/GET/pets"
            ]
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Security Cognito User Pools

- Cognito fully manages user lifecycle, token expires automatically
- API gateway verifies identity automatically from AWS Cognito
- No custom implementation required
- Authentication = Cognito User Pools | Authorization = API Gateway Methods

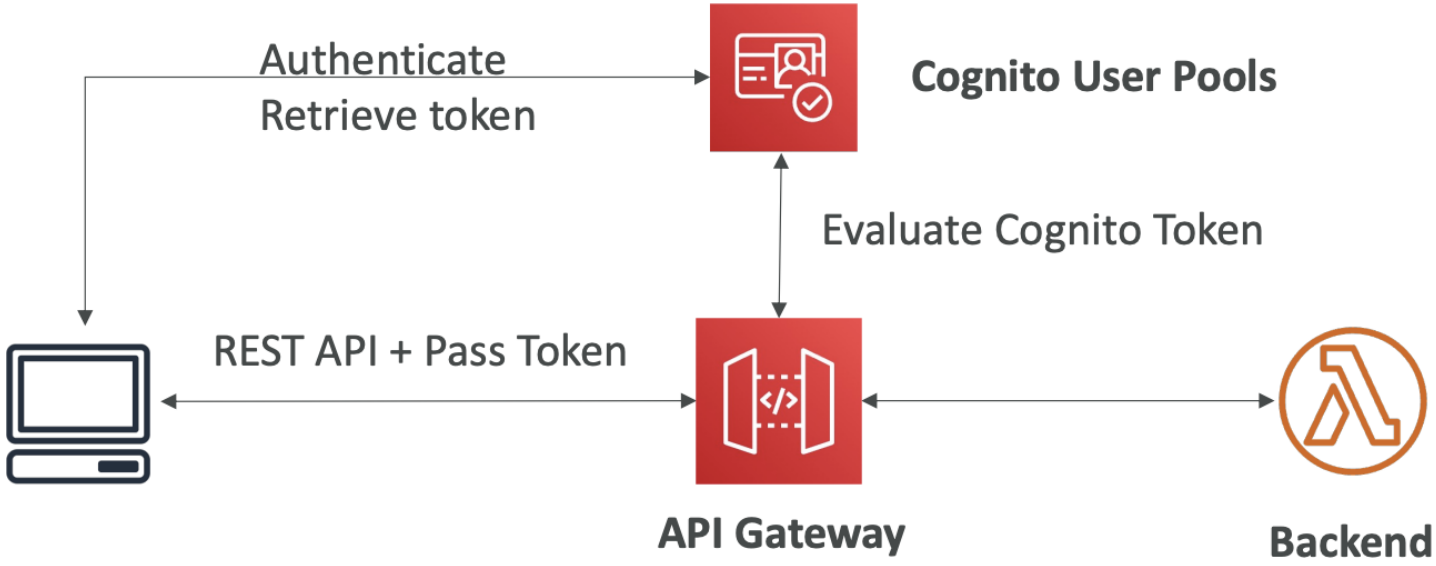

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Security Lambda Authorizer (formerly Custom Authorizers)

- Token-based authorizer (bearer token) – ex JWT (JSON Web Token) or Oauth
- A **request parameter-based** Lambda authorizer (**headers, query string**, stage var)
- Lambda must return an IAM policy for the user, result policy is cached
- **Authentication = External | Authorization = Lambda function**

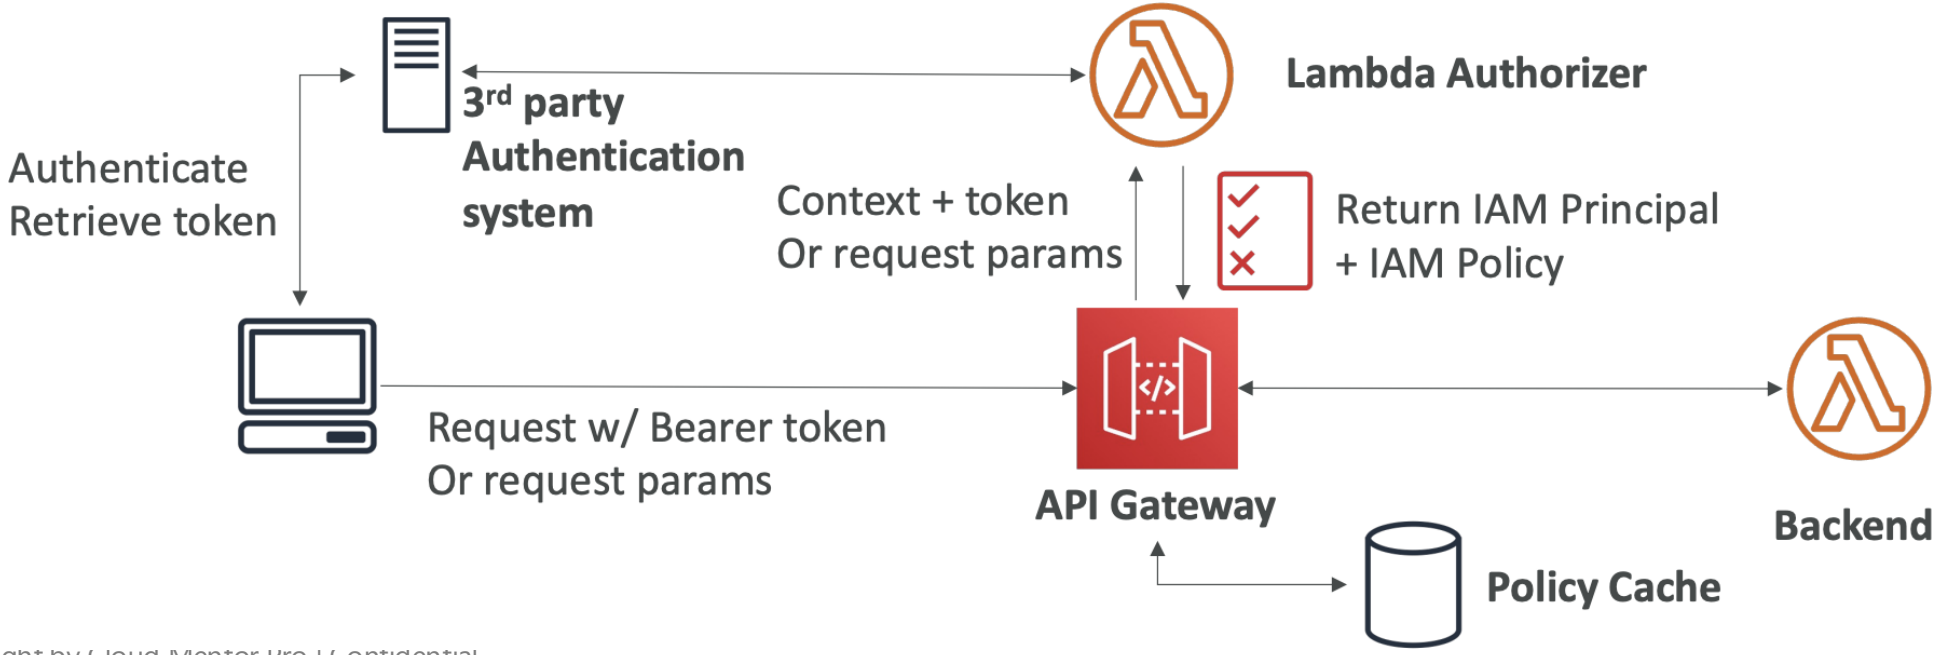

© Copyright by Cloud Mentor Pro | Confidential

API Gateway – Security – Summary

- IAM:
  - Great for users / roles already within your AWS account, **+ resource policy for cross account**
  - Handle authentication + authorization
  - Leverages Signature v4
- Custom Authorizer:
  - Great for 3rd party tokens
  - Very flexible in terms of what IAM policy is returned
  - Handle Authentication verification + Authorization in the Lambda function
  - Pay per Lambda invocation, results are cached
- Cognito User Pool:
  - You manage your own user pool (can be backed by Facebook, Google login etc…)
  - No need to write any custom code
  - Must implement authorization in the backend

© Copyright by Cloud Mentor Pro | Confidential

Exercise

Build API Gateway with Lambda Integration connect DynamoDB

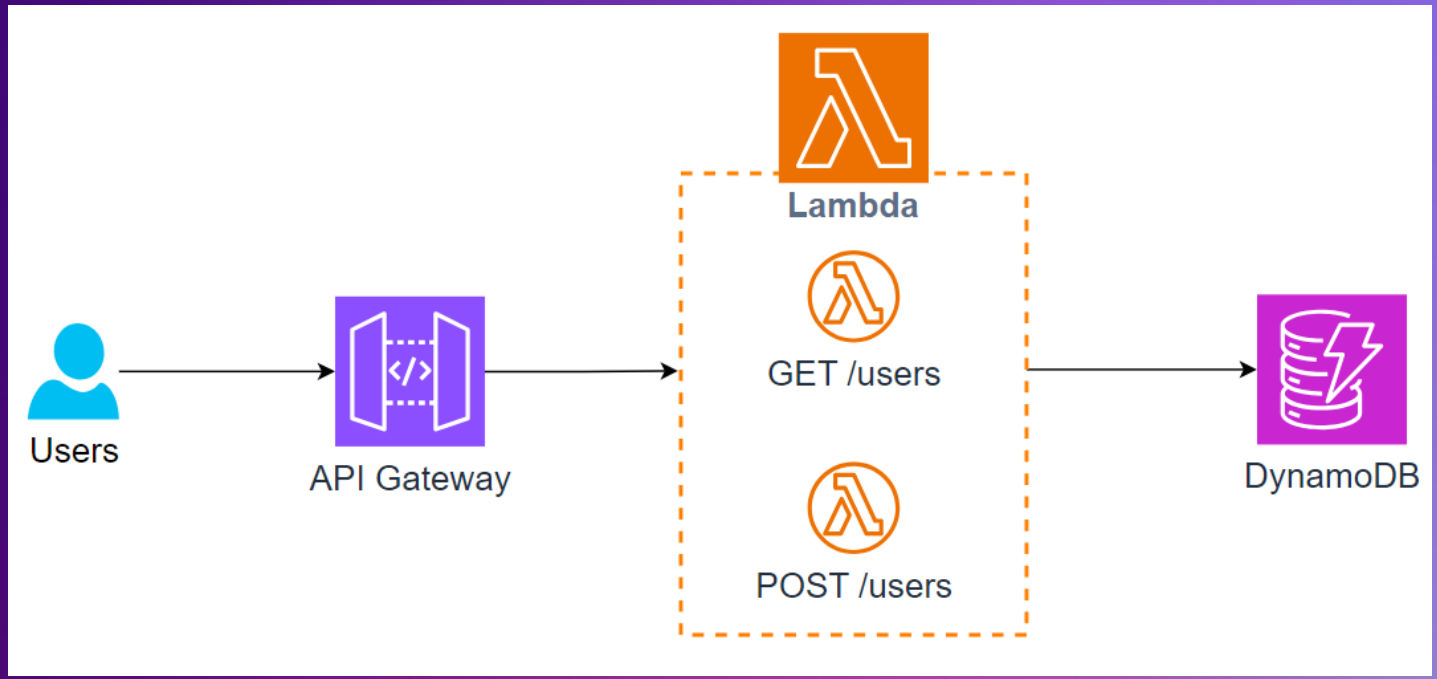

© Copyright by Cloud Mentor Pro | Confidential

Section 14

- *AWS Serverless: SAM - Serverless Application Model*
- *Cloud Development Kit (CDK)*
- *Cognito: Cognito User Pools, Cognito Identity Pools & Cognito Sync*
- *Other Serverless: Step Functions & AppSync*

© Copyright by Cloud Mentor Pro | Confidential

AWS Serverless Application Model (SAM)

Taking your Serverless Development to the next level

© Copyright by Cloud Mentor Pro | Confidential

AWS SAM

- SAM = Serverless Application Model
- Framework for developing and deploying serverless applications
- All the configuration is YAML code
- Generate complex CloudFormation from simple SAM YAML file
- Supports anything from CloudFormation: Outputs, Mappings, Parameters, Resources…
- SAM can use CodeDeploy to deploy Lambda functions
- SAM can help you to run Lambda, API Gateway, DynamoDB locally

© Copyright by Cloud Mentor Pro | Confidential

AWS SAM – Recipe

- **Transform Header indicates it's SAM template:**
  - Transform: 'AWS::Serverless-2016-10-31'
- **Write Code**
  - AWS : : Serverless : : Function
  - AWS : : Serverless : : Api
  - AWS : : Serverless : : SimpleTable
- **Package & Deploy**: sam deploy **(optionally preceded by** " sam package"**)**
- Quickly sync local changes to AWS Lambda (SAM Accelerate): sam sync --watch

© Copyright by Cloud Mentor Pro | Confidential

Deep Dive into SAM Deployment

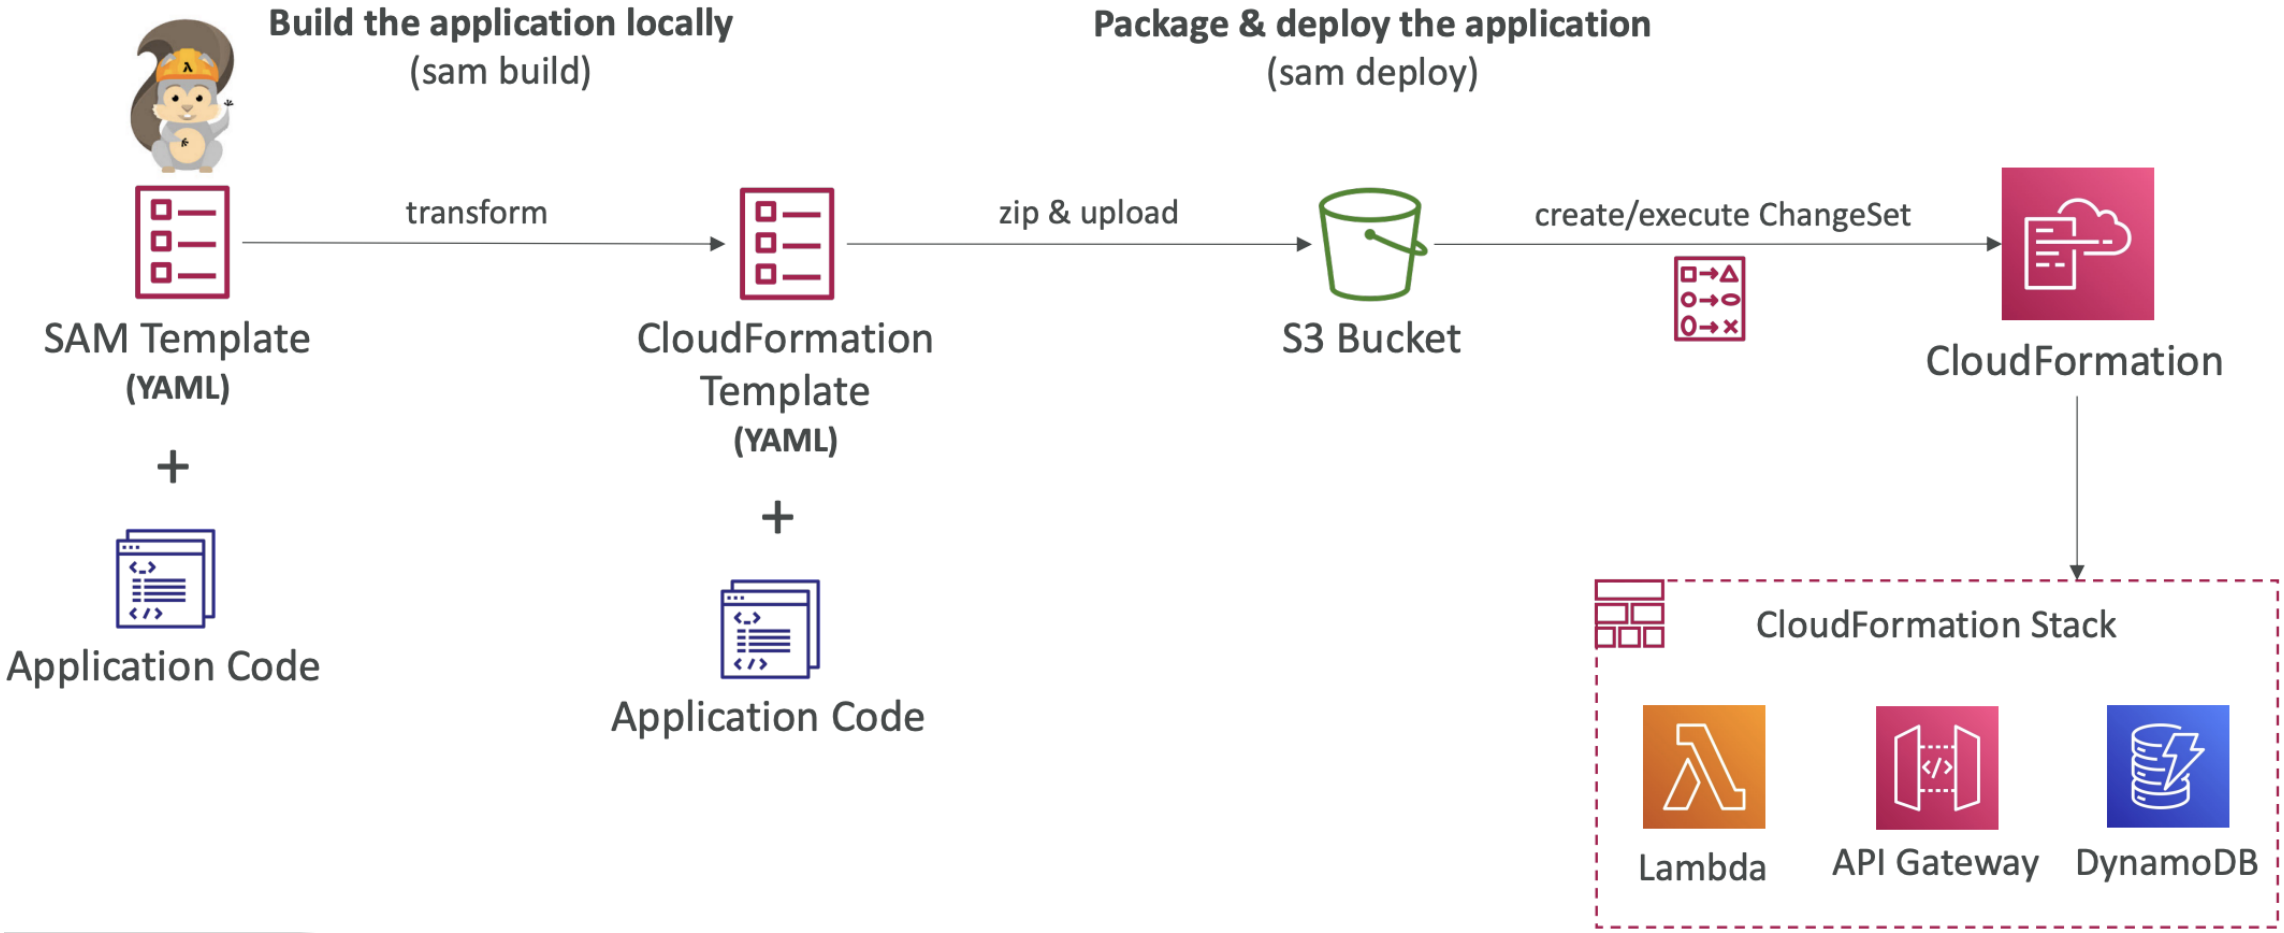

© Copyright by Cloud Mentor Pro | Confidential

SAM Accelerate (sam sync)

- SAM Accelerate is a set of features to reduce latency while deploying resources to AWS
- *sam sync*
  - Synchronizes your project declared in SAM templates to AWS
  - Synchronizes code changes to AWS without updating infrastructure (uses service APIs & **bypass CloudFormation**)

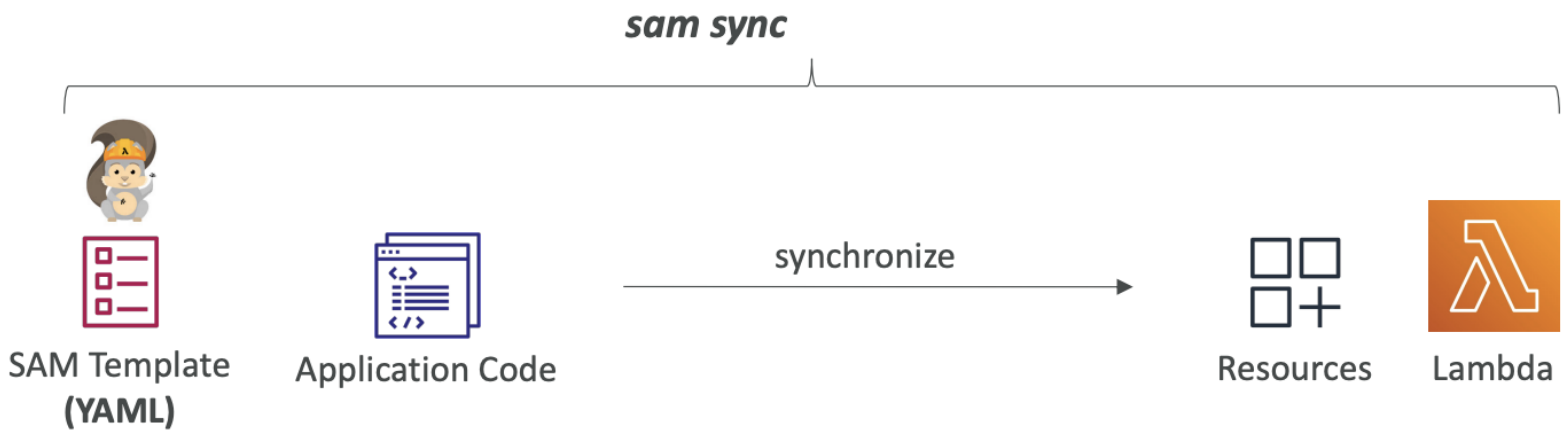

© Copyright by Cloud Mentor Pro | Confidential

SAM Accelerate (sam sync) – Examples

- sam sync (no options)
  - Synchronize code and infrastructure
- sam sync --code
  - Synchronize code changes without updating infrastructure (bypass CloudFormation, update in seconds)
- sam sync --code --resource AWS::Serverless::Function
  - Synchronize only all Lambda functions and their dependencies
- sam sync --code --resource-id HelloWorldLambdaFunction
  - Synchronize only a specific resource by its ID
- sam sync --watch
  - Monitor for file changes and automatically synchronize when changes are detected
  - If changes include configuration, it uses **sam sync**
  - If changes are code only, it uses **sam sync --code**

© Copyright by Cloud Mentor Pro | Confidential

SAM – CLI Debugging

- Locally build, test, and debug your serverless applications that are defined using AWS SAM templates
- Provides a lambda-like execution environment locally
- SAM CLI + AWS Toolkits => step-through and debug your code
- Supported IDEs: AWS Cloud9, Visual Studio Code, JetBrains, PyCharm, IntelliJ, …
- **AWS Toolkits**: IDE plugins which allows you to build, test, debug, deploy, and invoke Lambda functions built using AWS SAM

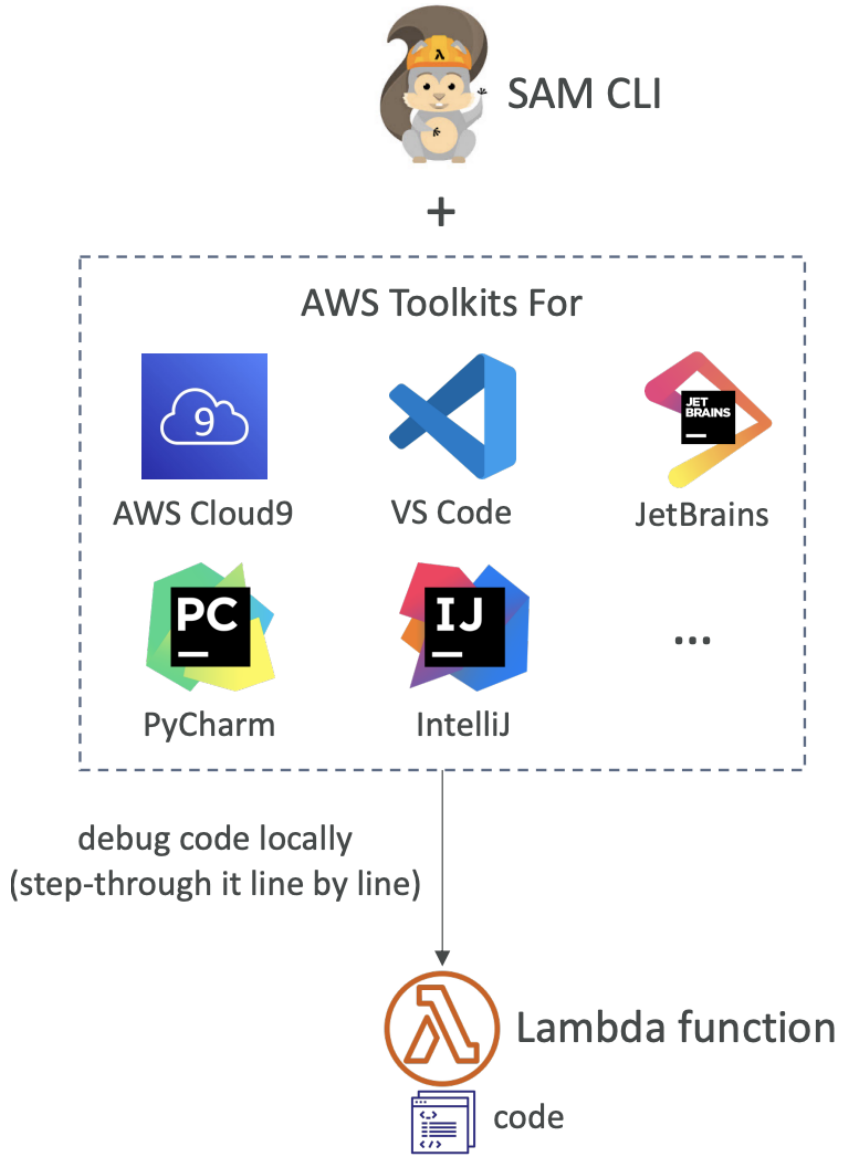

© Copyright by Cloud Mentor Pro | Confidential

SAM Policy Templates

- List of templates to apply permissions to your Lambda Functions
- Full list available here: https://docs.aws.amazon.com/serverless-application-model/latest/developerguide/serverless-policy-templates.html#serverless-policy-template-table
- Important examples:
  - **S3ReadPolicy**: Gives read only permissions to objects in S3
  - **SQSPollerPolicy**: Allows to poll an SQS queue
  - **DynamoDBCrudPolicy**: CRUD = create read update delete

```
MyFunction:
  Type: 'AWS::Serverless::Function'
  Properties:
    CodeUri: ${codeuri}
    Handler: hello.handler
    Runtime: python2.7
    Policies:
      - SQSPollerPolicy:
          QueueName:
            !GetAtt MyQueue.QueueName
```

© Copyright by Cloud Mentor Pro | Confidential

SAM and CodeDeploy

- SAM framework natively uses CodeDeploy to update Lambda functions

- Traffic Shifting feature
- Pre and Post traffic hooks features to validate deployment (before the traffic shift starts and after it ends)
- Easy & automated rollback using CloudWatch Alarms

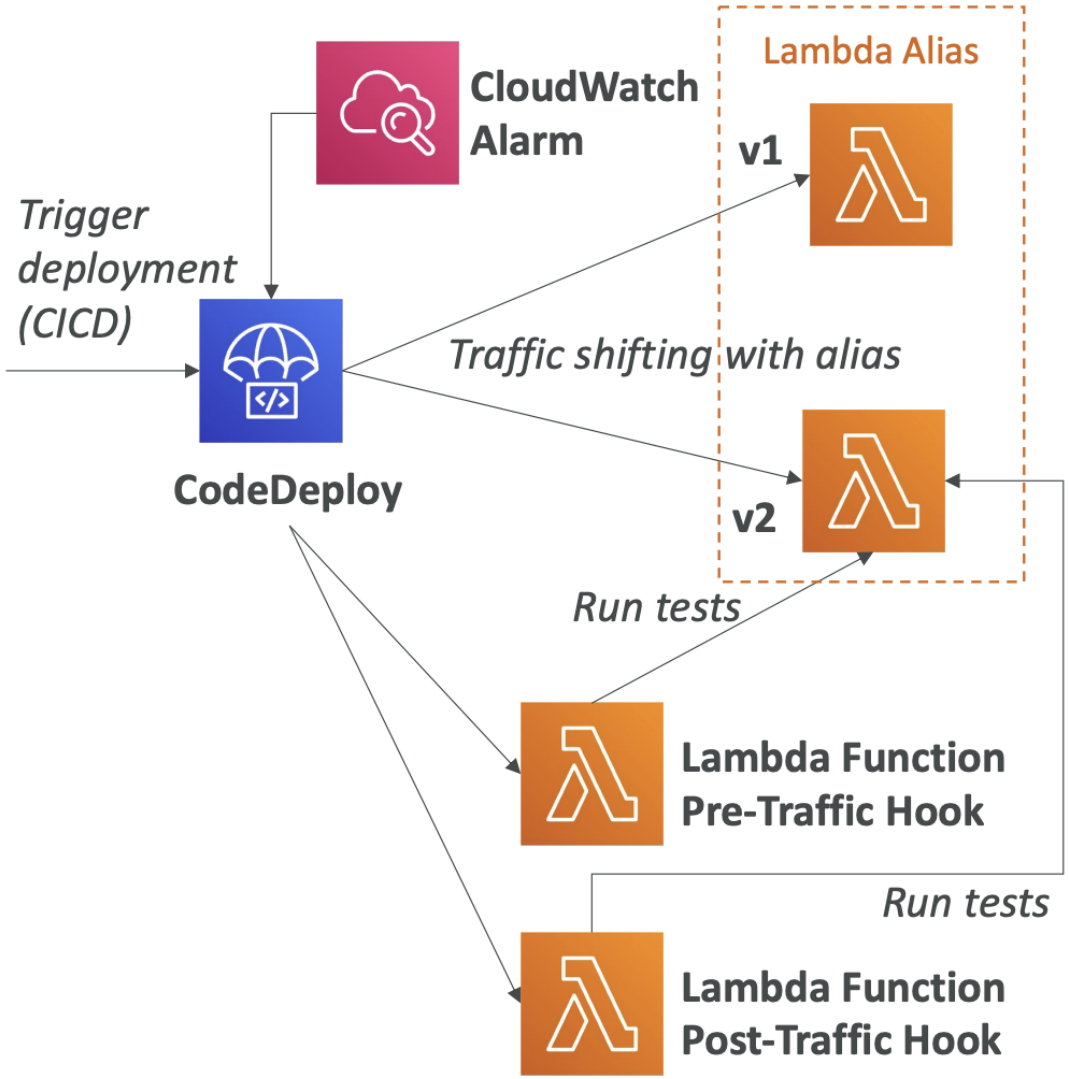

© Copyright by Cloud Mentor Pro | Confidential

SAM and CodeDeploy

- AutoPublishAlias
  - Detects when new code is being deployed
  - Creates and publishes an updated version of that function with the latest code
  - Points the alias to the updated version of the Lambda function
- DeploymentPreference
  - Canary, Linear, AllAtOnce
- Alarms
  - Alarms that can trigger a rollback
- Hooks
  - Pre and post traffic shifting Lambda functions to test your deployment

```yaml
Resources:
MyLambdaFunction:
  Type: AWS::Serverless::Function
  Properties:
    Handler: index.handler
    Runtime: nodejs12.x
    CodeUri: s3://bucket/code.zip

    AutoPublishAlias: live

    DeploymentPreference:
      Type: Canary10Percent10Minutes
      Alarms:
        # A list of alarms that you want to monitor
        - !Ref AliasErrorMetricGreaterThanZeroAlarm
        - !Ref LatestVersionErrorMetricGreaterThanZeroAlarm
      Hooks:
        # Validation Lambda functions that are run before & after traffic shifting
        PreTraffic: !Ref PreTrafficLambdaFunction
        PostTraffic: !Ref PostTrafficLambdaFunction
```

© Copyright by Cloud Mentor Pro | Confidential

SAM – Local Capabilities

- Locally start AWS Lambda
  - *sam local start-lambda*
  - Starts a local endpoint that emulates AWS Lambda
  - Can run automated tests against this local endpoint
- Locally Invoke Lambda Function
  - *sam local invoke*
  - Invoke Lambda function with payload once and quit after invocation completes
  - Helpful for generating test cases
  - If the function make API calls to AWS, make sure you are using the correct --***profile*** option

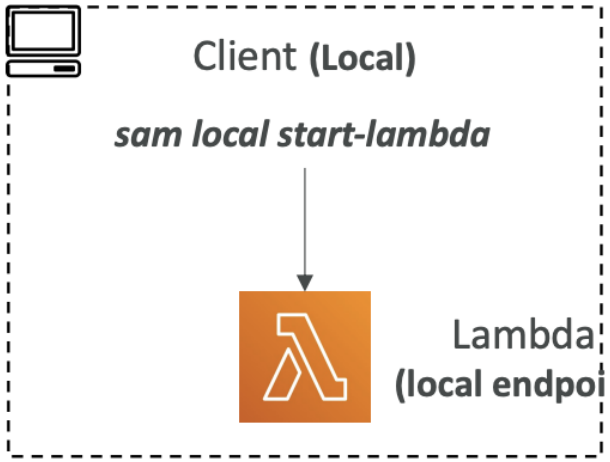

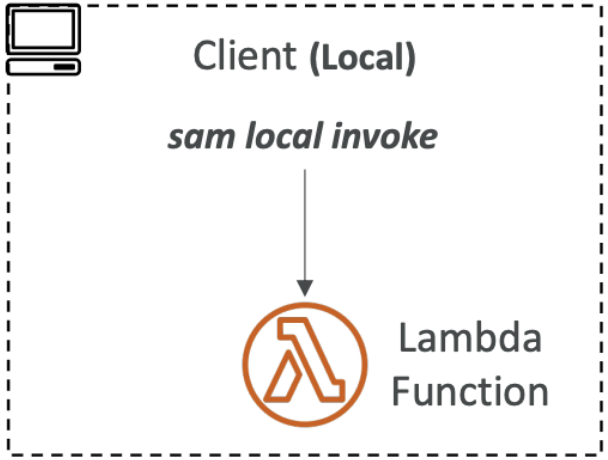

© Copyright by Cloud Mentor Pro | Confidential

SAM – Local Capabilities

- Locally Start an API Gateway Endpoint
  - *sam local start-api*
  - Starts a local HTTP server that hosts all your functions
  - Changes to functions are automatically reloaded

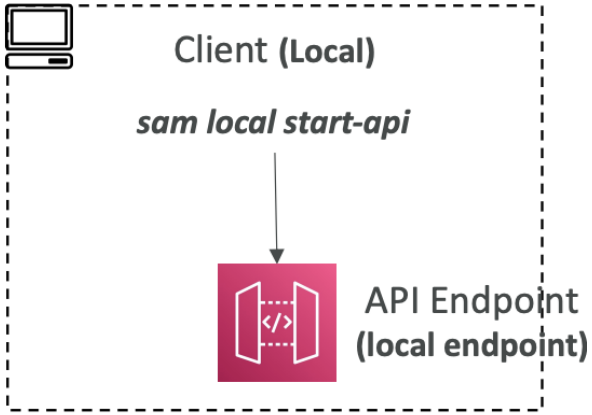

- Generate AWS Events for Lambda Functions
  - *sam local generate-event*
  - Generate sample payloads for event sources
  - S3, API Gateway, SNS, Kinesis, DynamoDB…

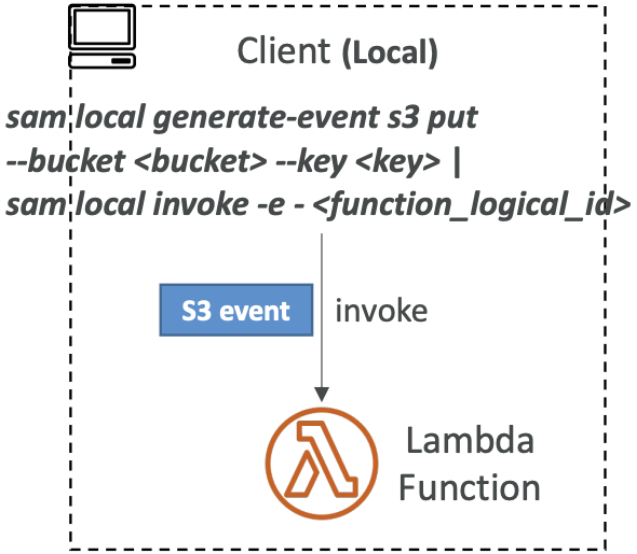

© Copyright by Cloud Mentor Pro | Confidential

SAM – Multiple Environments

samconfig.toml

```
version = 0.1

[dev.deploy.parameters]
stack_name = "my-dev-stack"
s3_bucket = "XXXXX-dev"
s3_prefix = "XXXXX/dev"
region = "us-east-1"
capabilities = "CAPABILITY_IAM"
parameter_overrides = "Environment=Development"

[prod.deploy.parameters]
stack_name = "my-prod-stack"
s3_bucket = "XXXXX-prod"
s3_prefix = "XXXXX/prod"
region = "us-east-1"
capabilities = "CAPABILITY_IAM"
parameter_overrides = "Environment=Production"

[dev.sync.parameters]
watch = true

[prod.sync.parameters]
watch = false
```

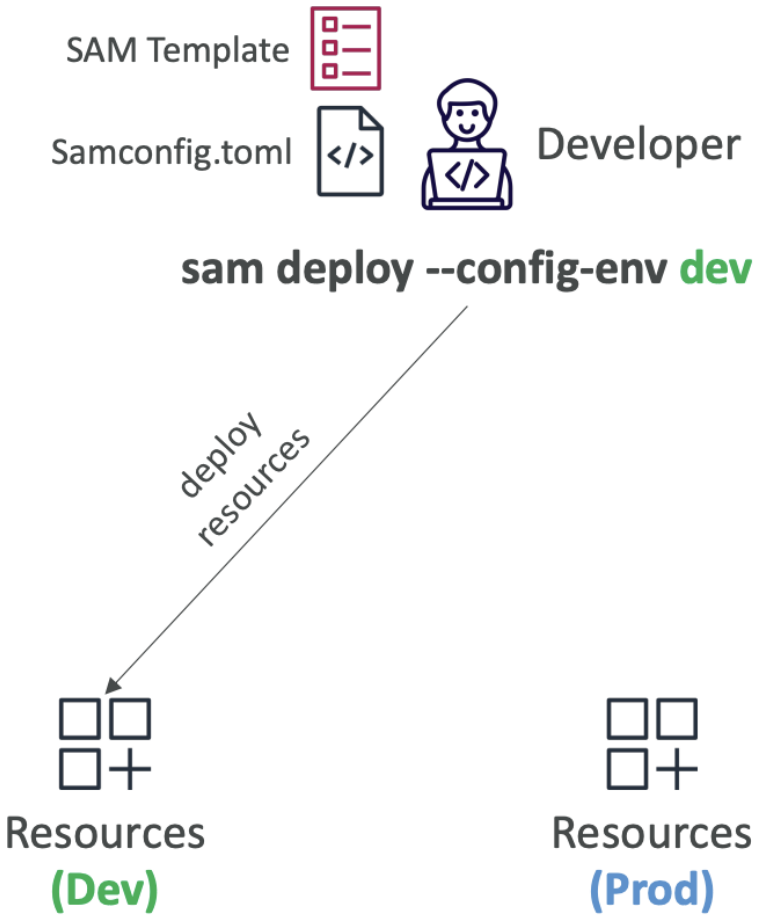

© Copyright by Cloud Mentor Pro | Confidential

AWS Cloud Development Kit

© Copyright by Cloud Mentor Pro | Confidential

AWS Cloud Development Kit (CDK)

- Define your cloud infrastructure using a familiar language:
  - JavaScript/TypeScript, Python, Java, and .NET
- Contains high level components called **constructs**
- The code is "compiled" into a CloudFormation template (JSON/YAML)
- **You can therefore deploy infrastructure and application runtime code together**
  - Great for Lambda functions
  - Great for Docker containers in ECS / EKS

```
export class MyEcsConstructStack extends core.Stack {
  constructor(scope: core.App, id: string, props?: core.StackProps)
    super(scope, id, props);

    const vpc = new ec2.Vpc(this, "MyVpc", {
      maxAzs: 3 // Default is all AZs in region
    });

    const cluster = new ecs.Cluster(this, "MyCluster", {
      vpc: vpc
    });

    // Create a load-balanced Fargate service and make it public
    new ecs_patterns.ApplicationLoadBalancedFargateService(this, "My
      cluster: cluster, // Required
      cpu: 512, // Default is 256
      desiredCount: 6, // Default is 1
      taskImageOptions: { image: ecs.ContainerImage.fromRegistry("am
      memoryLimitMiB: 2048, // Default is 512
      publicLoadBalancer: true // Default is false
    });
  }
}
```

© Copyright by Cloud Mentor Pro | Confidential

CDK in a diagram

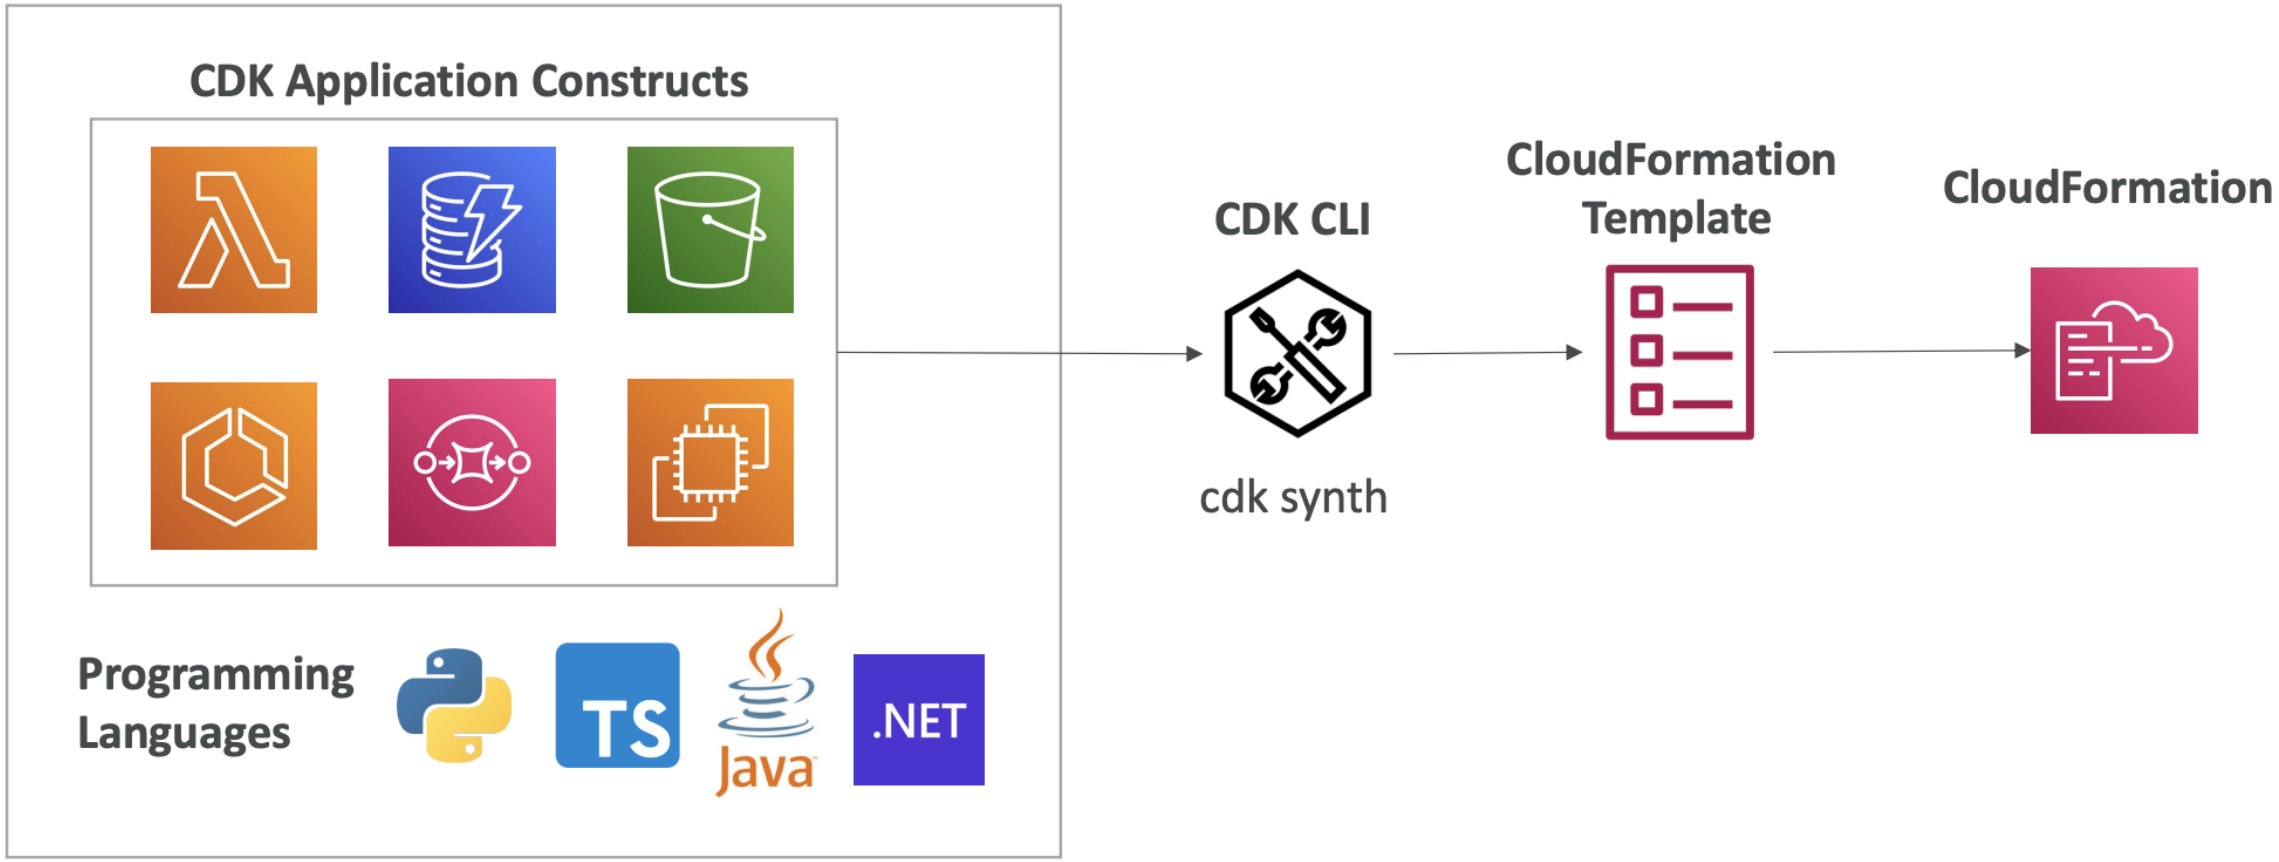

© Copyright by Cloud Mentor Pro | Confidential

CDK vs SAM

- SAM:
  - Serverless focused
  - Write your template declaratively in JSON or YAML
  - Great for quickly getting started with Lambda
  - Leverages CloudFormation
- CDK:
  - All AWS services
  - Write infra in a programming language JavaScript/TypeScript, Python, Java, and .NET
  - Leverages CloudFormation

© Copyright by Cloud Mentor Pro | Confidential

CDK + SAM

- You can use SAM CLI to locally test your CDK apps
- **You must first run *cdk synth***

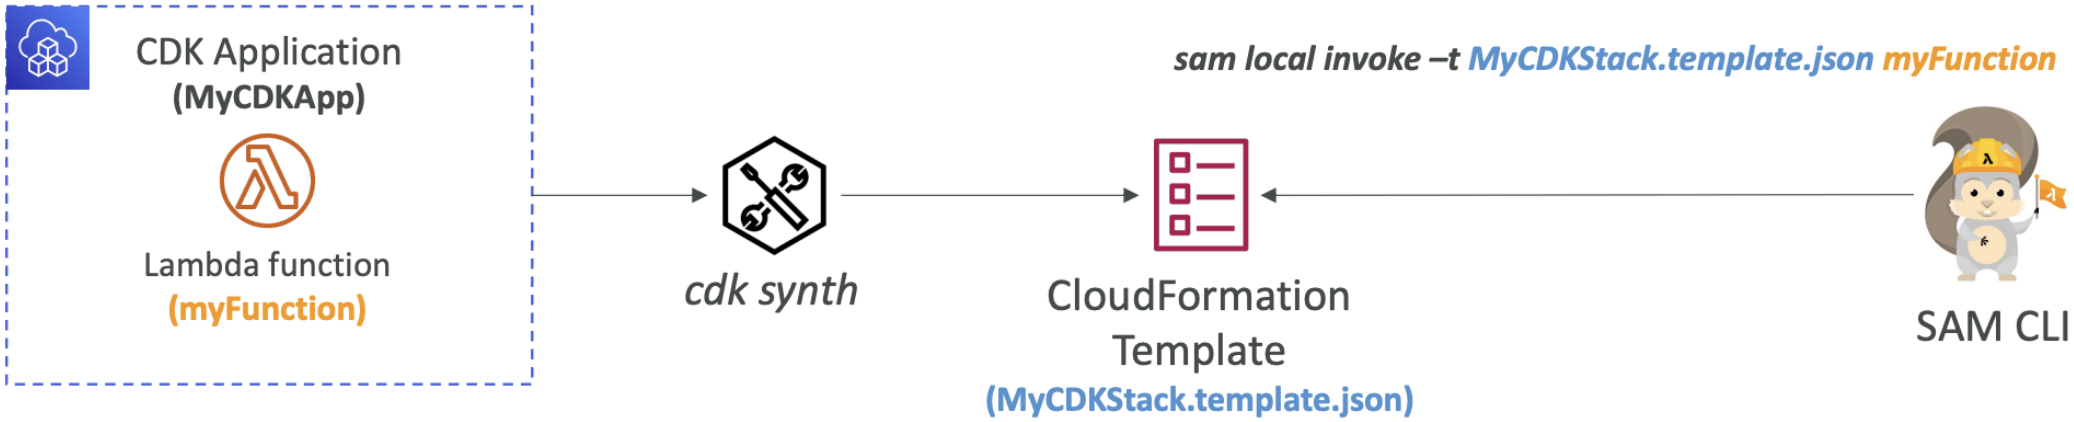

© Copyright by Cloud Mentor Pro | Confidential

CDK – Important Commands to know

Command	Description
npm install -g aws-cdk-lib	Install the CDK CLI and libraries
cdk init app	Create a new CDK project from a specified template
cdk synth	Synthesizes and prints the CloudFormation template
cdk bootstrap	Deploys the CDK Toolkit staging Stack
cdk deploy	Deploy the Stack(s)
cdk diff	View differences of local CDK and deployed Stack
cdk destroy	Destroy the Stack(s)

© Copyright by Cloud Mentor Pro | Confidential

CDK – Bootstrapping

- The process of provisioning resources for CDK before you can deploy CDK apps into an AWS environment
- **AWS Environment = account & region**
- CloudFormation Stack called **CDKToolkit** is created and contains:
  - **S3 Bucket** – to store files
  - **IAM Roles** – to grant permissions to perform deployments
- You must run the following command for each new environment:
  - *cdk bootstrap aws://<aws_account>/<aws_region>*
- Otherwise, you will get an error "**Policy contains a statement with one or more invalid principal**"

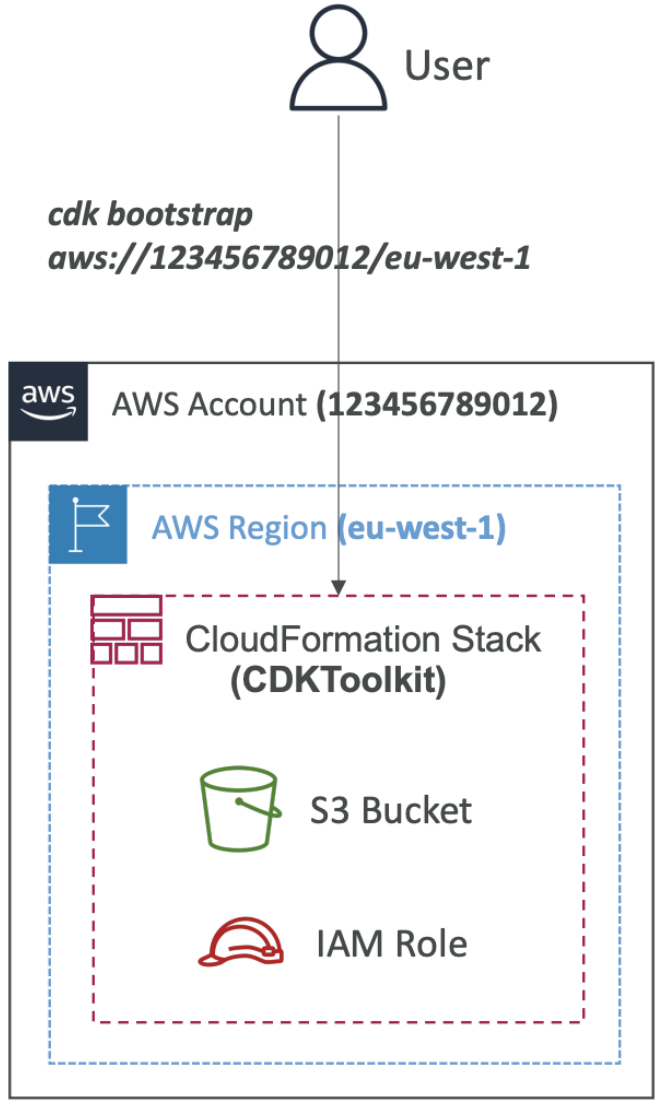

© Copyright by Cloud Mentor Pro | Confidential

CDK –Testing

- To test CDK apps, use **CDK Assertions Module** combined with popular test frameworks such as Jest (JavaScript) or Pytest (Python)
- Verify we have specific resources, rules, conditions, parameters…
- Two types of tests:
  - **Fine-grained Assertions (common)** – test specific aspects of the CloudFormation template (e.g., check if a resource has this property with this value)
  - **Snapshot Tests** – test the synthesized CloudFormation template against a previously stored baseline template
- To import a template
  - Template.fromStack(MyStack) : stack built in CDK
  - Template.fromString(mystring) : stack build outside CDK

```
describe("StateMachineStack", () => {
    test("synthesizes the way we expect", () => {
        ...

        // Prepare the stack for assertions
        const template = Template.fromStack(MyStack);

        // Assert it creates Lambda with correct properties...
        template.hasResourceProperties("AWS::Lambda::Function", {
            Handler: "handler",
            Runtime: "nodejs14.x",
        });

        // Assert it creates the SNS subscription...
        template.resourceCountIs("AWS::SNS::Subscription", 1);

        // Assert the synthesized CloudFormation template
        // against a previously stored baseline template
        expect(template.toJSON()).toMatchSnapshot();
    }
});
```

Fine-grained Assertions

Snapshot Test

© Copyright by Cloud Mentor Pro | Confidential

Amazon Cognito

© Copyright by Cloud Mentor Pro | Confidential

Amazon Cognito

- Give users an identity to interact with our web or mobile application
- Cognito User Pools:
  - Sign in functionality for app users
  - Integrate with API Gateway & Application Load Balancer
- Cognito Identity Pools (Federated Identity):
  - Provide AWS credentials to users so they can access AWS resources directly
  - Integrate with Cognito User Pools as an identity provider
- **Cognito vs IAM**: “hundreds of users”, ”mobile users”, “authenticate with SAML”

© Copyright by Cloud Mentor Pro | Confidential

Cognito User Pools (CUP) – User Features

- **Create a serverless database of user for your web & mobile apps**
- Simple login: Username (or email) / password combination
- Password reset
- Email & Phone Number Verification
- Multi-factor authentication (MFA)
- Federated Identities: users from Facebook, Google, SAML…
- Feature: block users if their credentials are compromised elsewhere
- Login sends back a JSON Web Token (JWT)

© Copyright by Cloud Mentor Pro | Confidential

Cognito User Pools (CUP) – Diagram

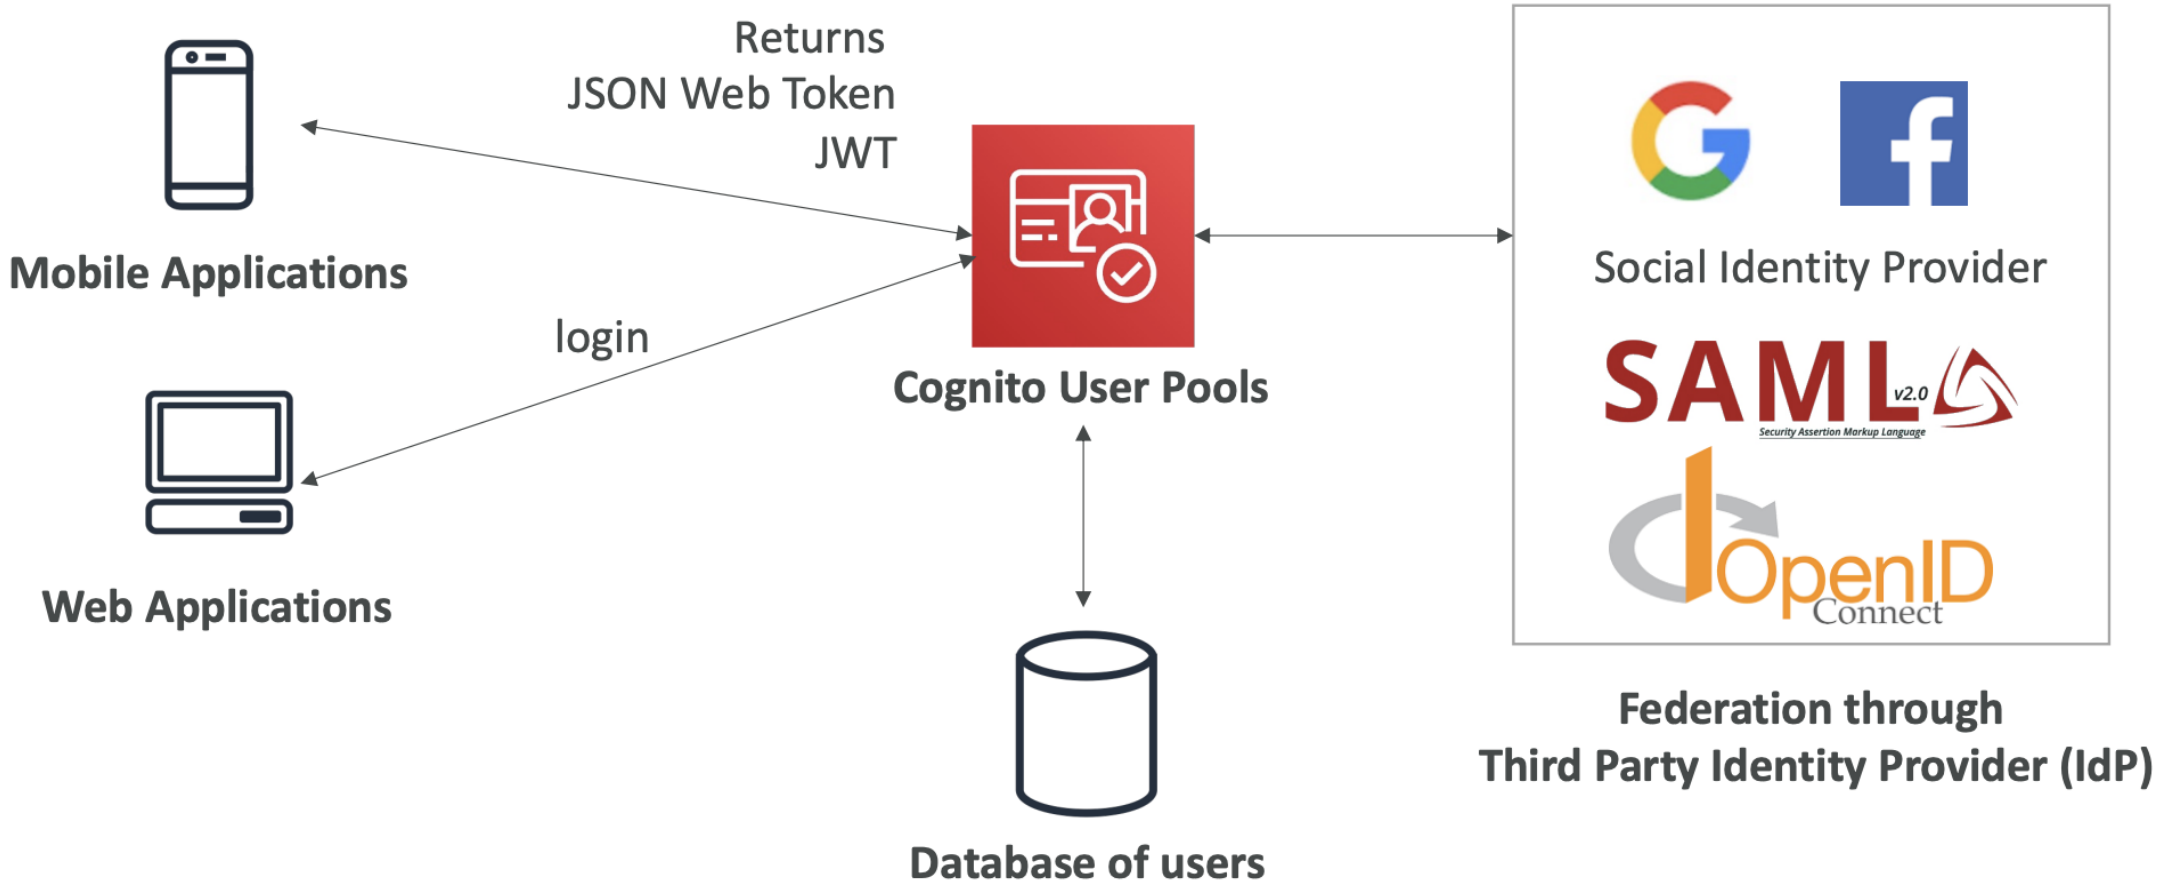

© Copyright by Cloud Mentor Pro | Confidential

Cognito User Pools (CUP) - Integrations

- CUP integrates with **API Gateway** and **Application Load Balancer**

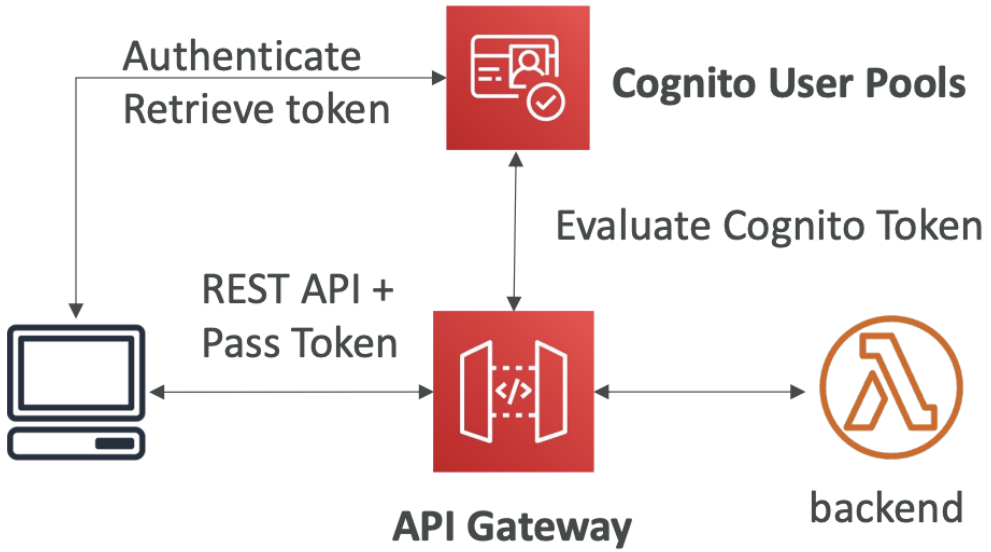

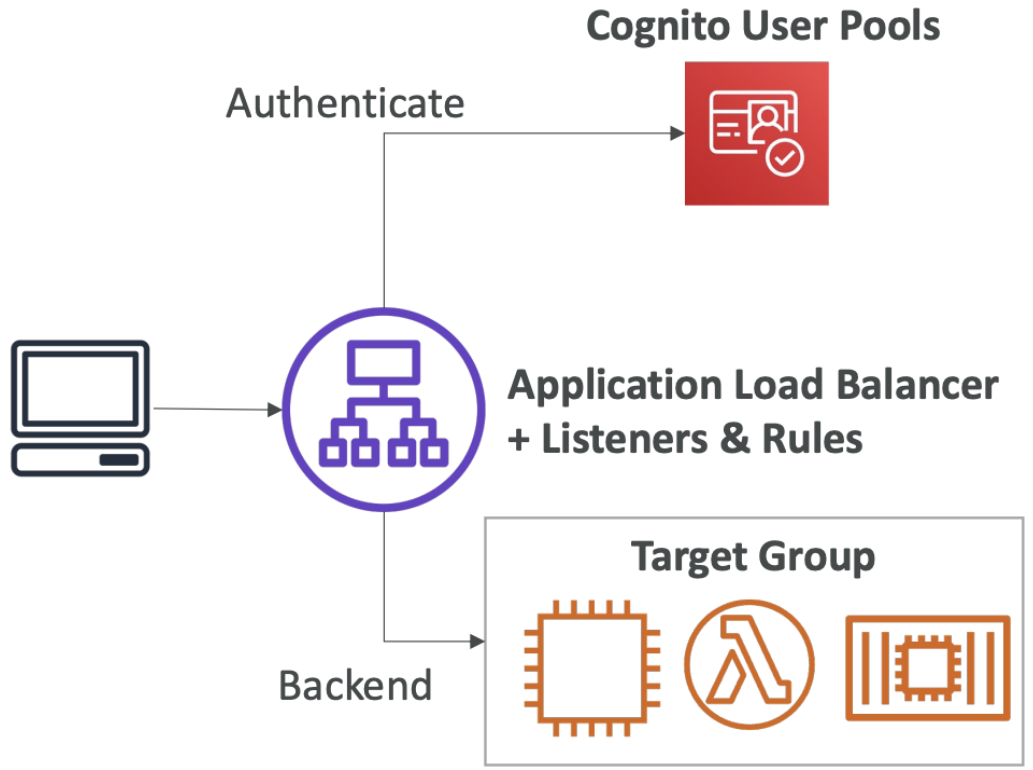

© Copyright by Cloud Mentor Pro | Confidential

Cognito User Pools – Lambda Triggers

- CUP can invoke a Lambda function synchronously on these triggers:

User Pool Flow	Operation	Description
Authentication Events	Pre Authentication Lambda Trigger	Custom validation to accept or deny the sign-in request
	Post Authentication Lambda Trigger	Event logging for custom analytics
	Pre Token Generation Lambda Trigger	Augment or suppress token claims
Sign-Up	Pre Sign-up Lambda Trigger	Custom validation to accept or deny the sign-up request
	Post Confirmation Lambda Trigger	Custom welcome messages or event logging for custom analytics
	Migrate User Lambda Trigger	Migrate a user from an existing user directory to user pools
Messages	Custom Message Lambda Trigger	Advanced customization and localization of messages
Token Creation	Pre Token Generation Lambda Trigger	Add or remove attributes in Id tokens

https://docs.aws.amazon.com/cognito/latest/developerguide/cognito-user-identity-pools-working-with-aws-lambda-triggers.html

© Copyright by Cloud Mentor Pro | Confidential

Cognito User Pools – Hosted Authentication UI

- Cognito has a **hosted authentication UI** that you can add to your app to handle sign-up and sign-in workflows
- Using the hosted UI, you have a foundation for integration with social logins, OIDC or SAML
- Can customize with a **custom logo** and **custom CSS**

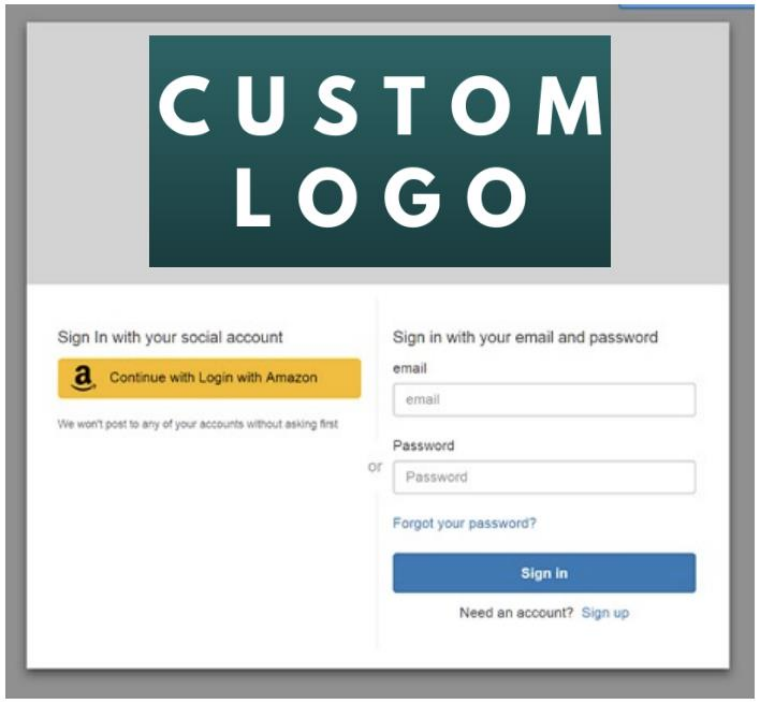

https://aws.amazon.com/blogs/aws/launch-amazon-cognito-user-pools-general-availability-app-integration-and-federation/

© Copyright by Cloud Mentor Pro | Confidential

CUP – Hosted UI Custom Domain

- For custom domains, you must create an ACM certificate in us-east-1
- The custom domain must be defined in the “App Integration” section

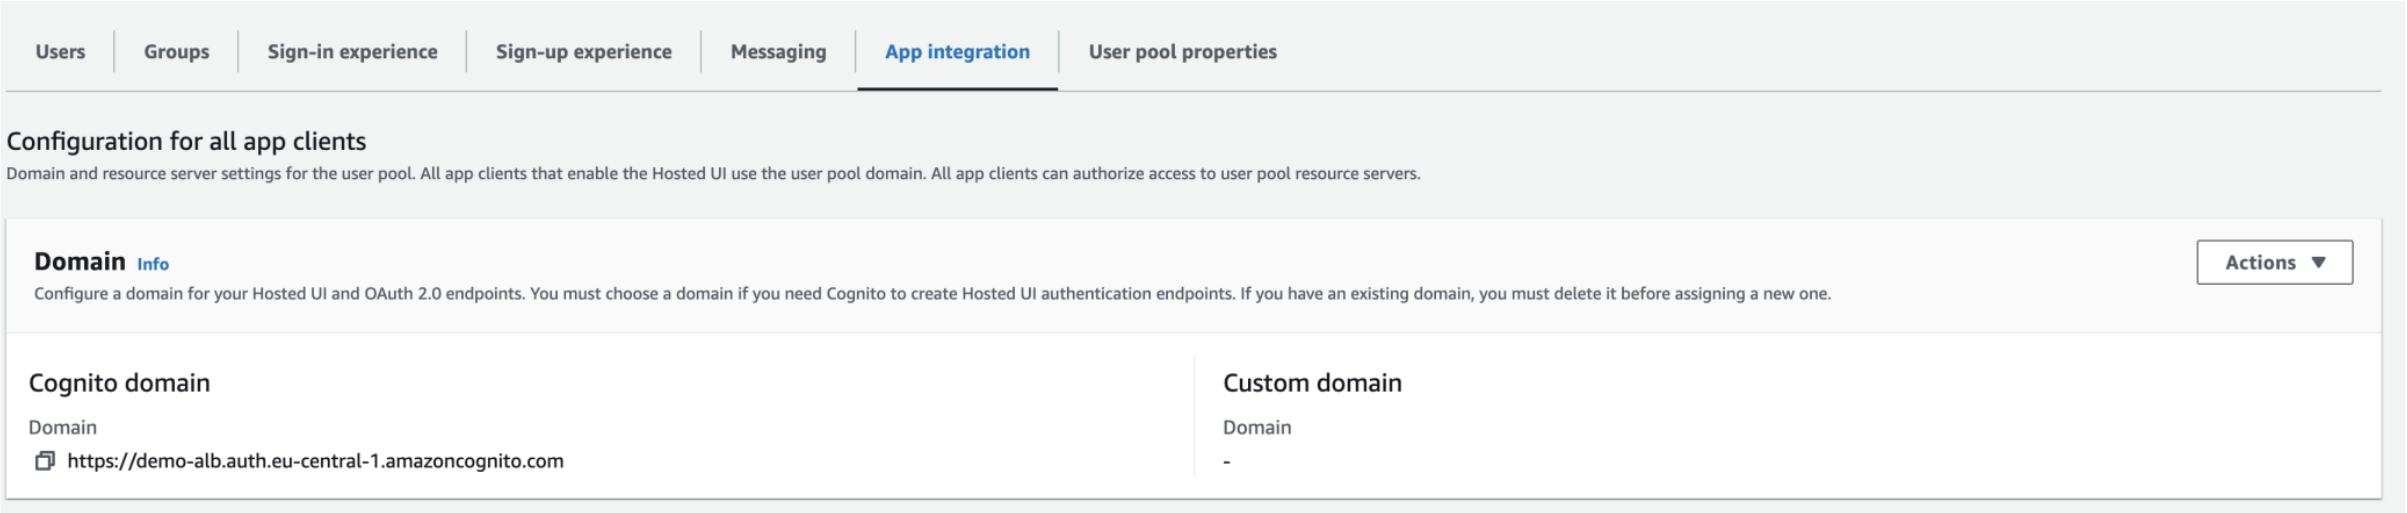

© Copyright by Cloud Mentor Pro | Confidential

Application Load Balancer – Cognito Auth.

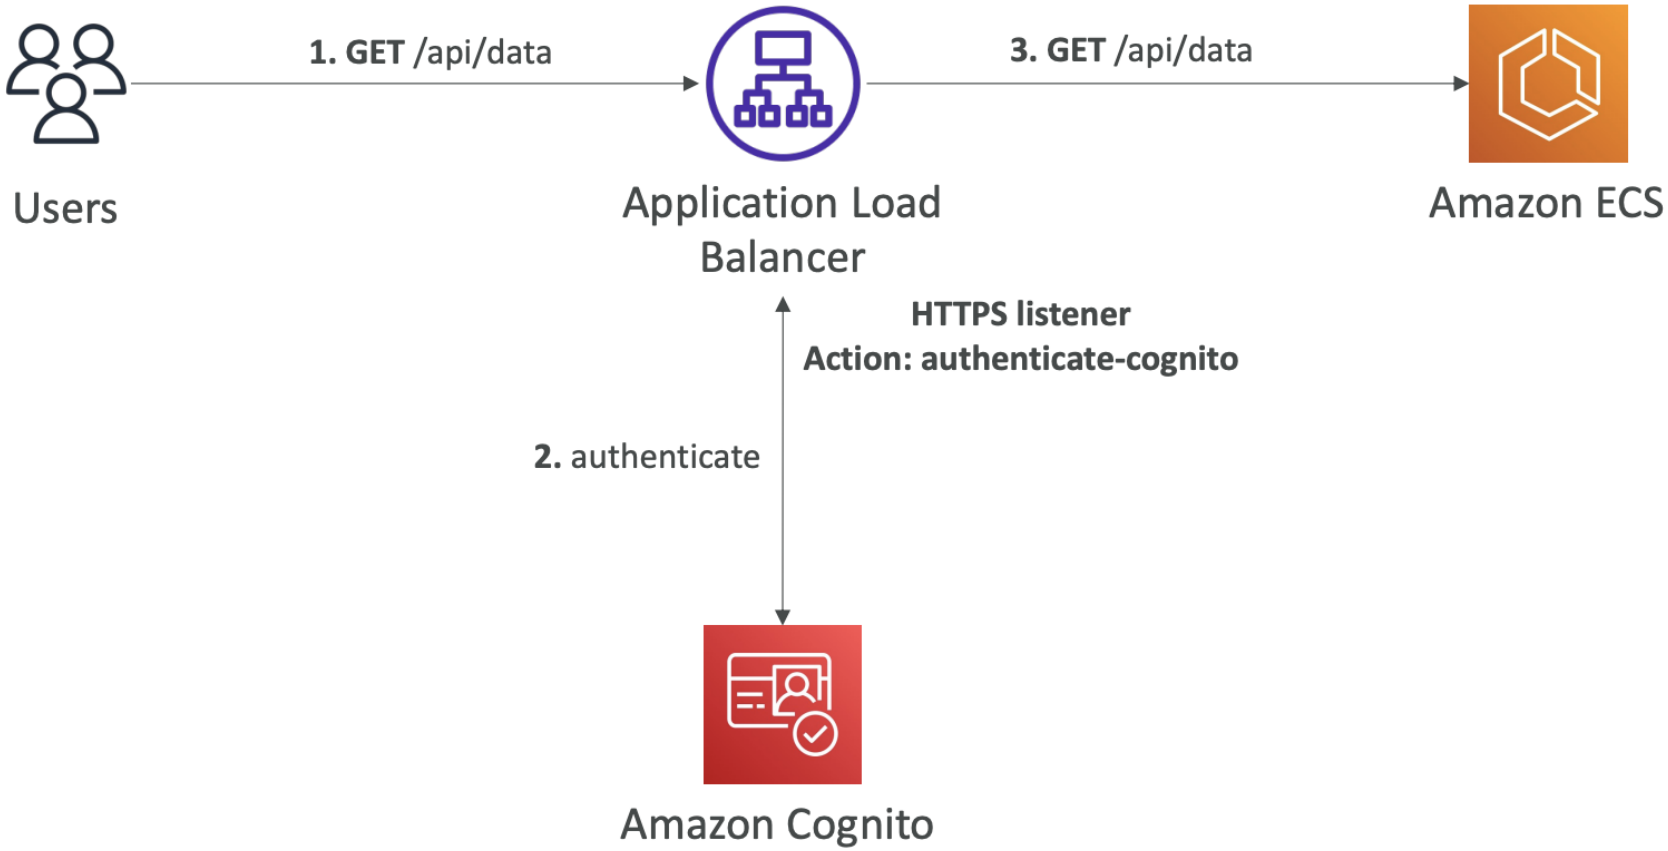

© Copyright by Cloud Mentor Pro | Confidential

ALB – Auth through Cognito User Pools

- Create Cognito User Pool, Client and Domain
- Make sure an ID token is returned
- Add the social or Corporate IdP if needed
- Several URL redirections are necessary
- Allow your Cognito User Pool Domain on your IdP app's callback URL. For example:
  - https://domainprefix.auth.region.amazoncognito.com/saml2/ idpresponse
  - https://user-pool-domain/oauth2/idpresponse

© Copyright by Cloud Mentor Pro | Confidential

Cognito Identity Pools (Federated Identities)

- Get identities for “users” so they obtain temporary AWS credentials
- Your identity pool (e.g identity source) can include:
  - Public Providers (Login with Amazon, Facebook, Google, Apple)
  - Users in an Amazon Cognito user pool
  - OpenID Connect Providers & SAML Identity Providers
  - Developer Authenticated Identities (custom login server)
  - Cognito Identity Pools allow for unauthenticated (guest) access
- Users can then access AWS services directly or through API Gateway
  - The IAM policies applied to the credentials are defined in Cognito
  - They can be customized based on the user_id for fine grained control

© Copyright by Cloud Mentor Pro | Confidential

Cognito Identity Pools – Diagram

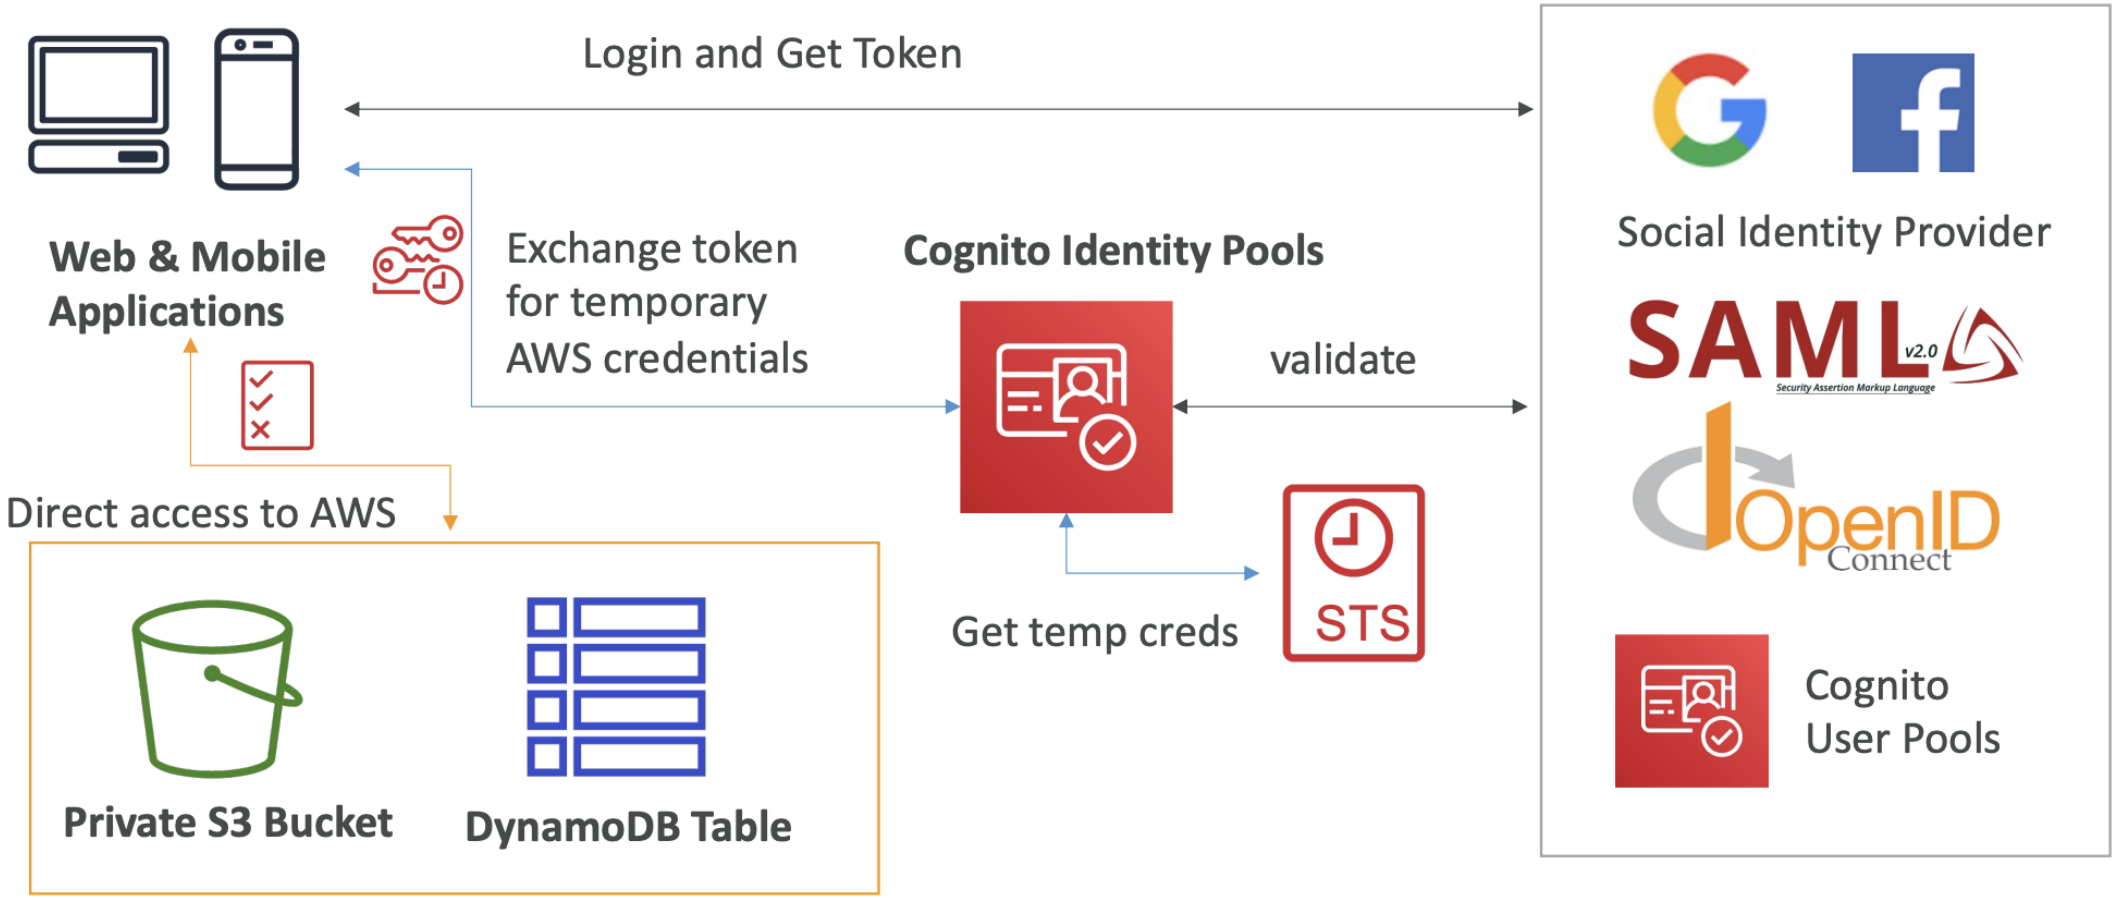

© Copyright by Cloud Mentor Pro | Confidential

Cognito Identity Pools – Diagram with CUP

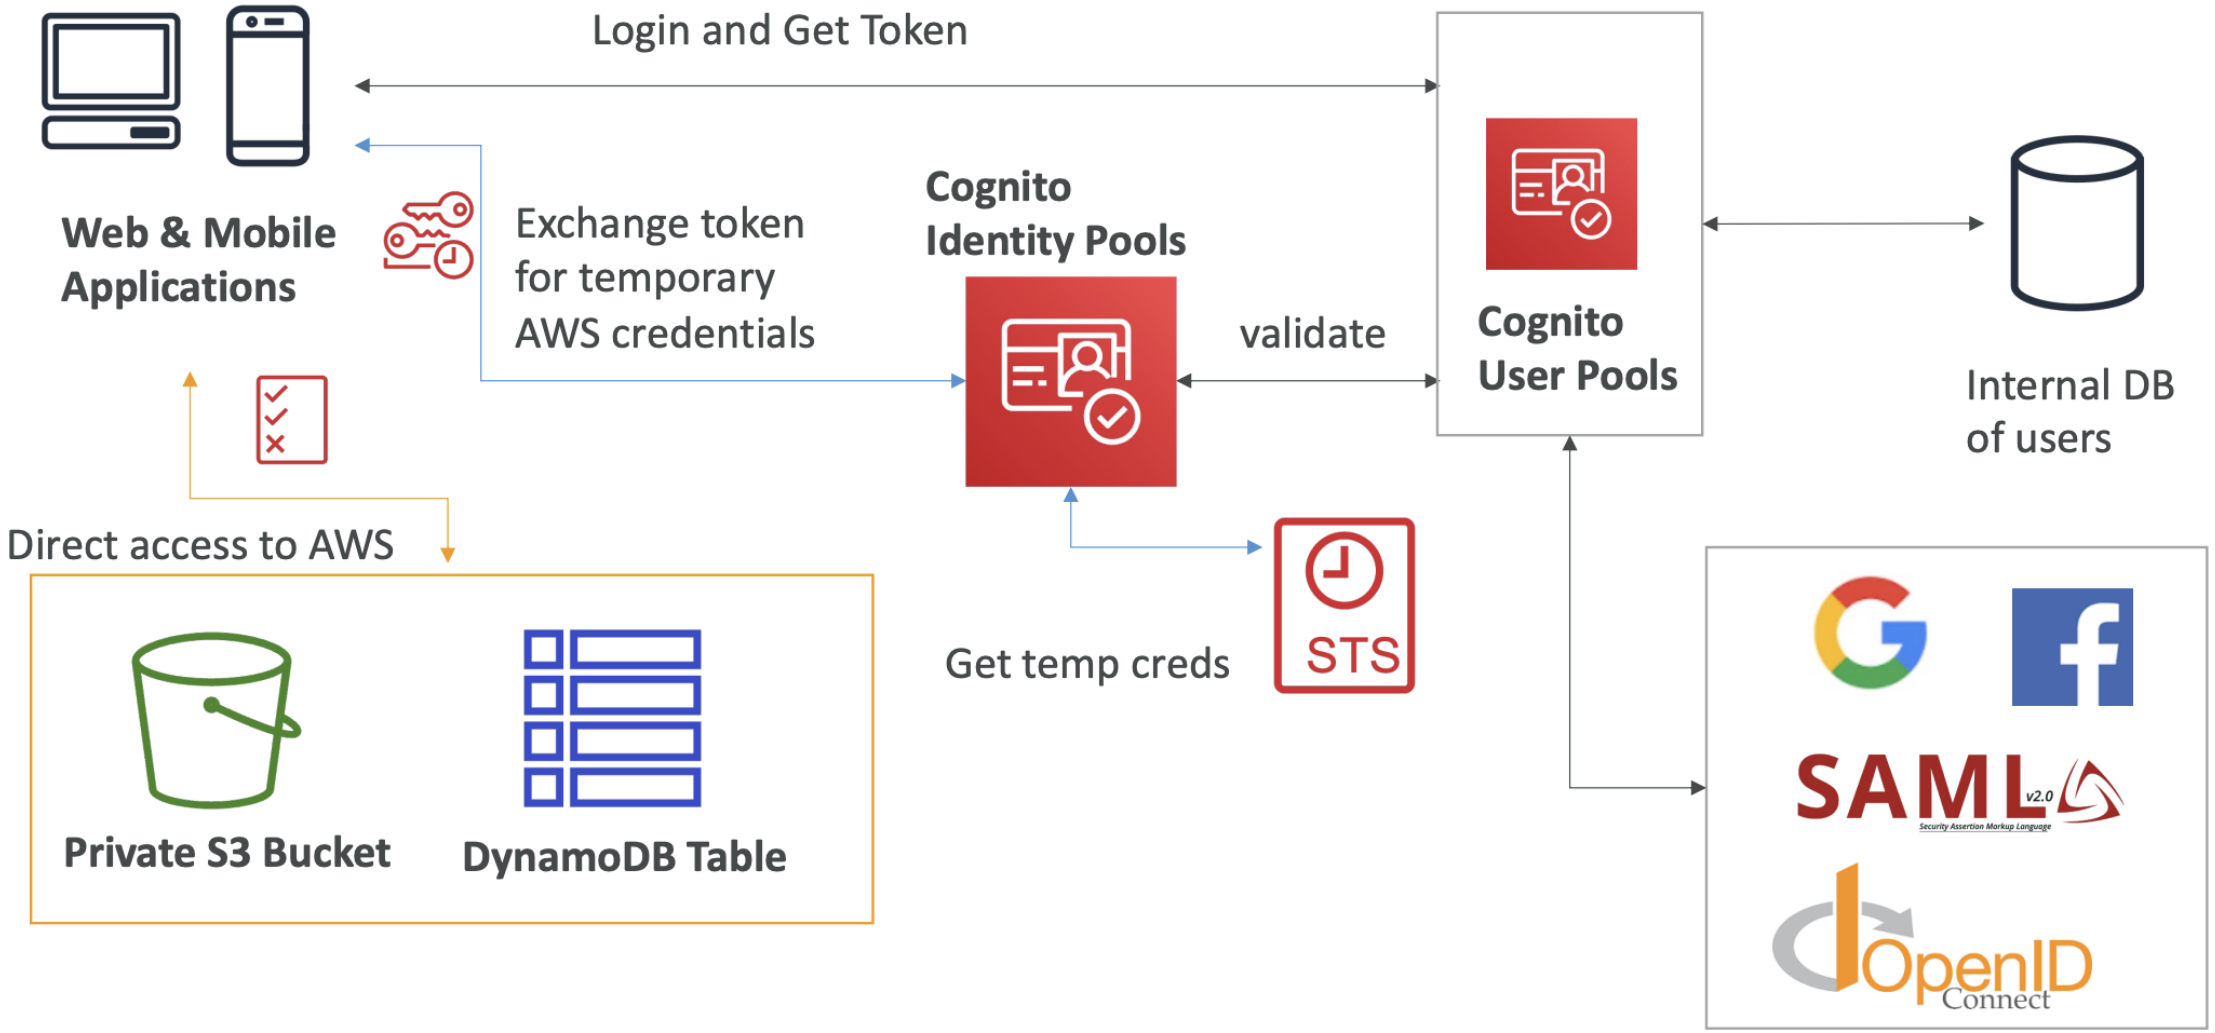

© Copyright by Cloud Mentor Pro | Confidential

Cognito User Pools vs Identity Pools

- Cognito User Pools (for authentication = identity verification)
  - Database of users for your web and mobile application
  - Allows to federate logins through Public Social, OIDC, SAML…
  - Can customize the hosted UI for authentication (including the logo)
  - Has triggers with AWS Lambda during the authentication flow
  - Adapt the sign-in experience to different risk levels (MFA, adaptive authentication, etc…)
- Cognito Identity Pools (for authorization = access control)
  - Obtain AWS credentials for your users
  - Users can login through Public Social, OIDC, SAML & Cognito User Pools
  - Users can be unauthenticated (guests)
  - Users are mapped to IAM roles & policies, can leverage policy variables
- CUP + CIP = authentication + authorization

© Copyright by Cloud Mentor Pro | Confidential

Other Serverless

© Copyright by Cloud Mentor Pro | Confidential

AWS Step Functions

- Model your workflows as **state machines (one per workflow)**
  - Order fulfillment, Data processing
  - Web applications, Any workflow
- Written in JSON
- Visualization of the workflow and the execution of the workflow, as well as history
- Start workflow with SDK call, API Gateway, Event Bridge (CloudWatch Event)

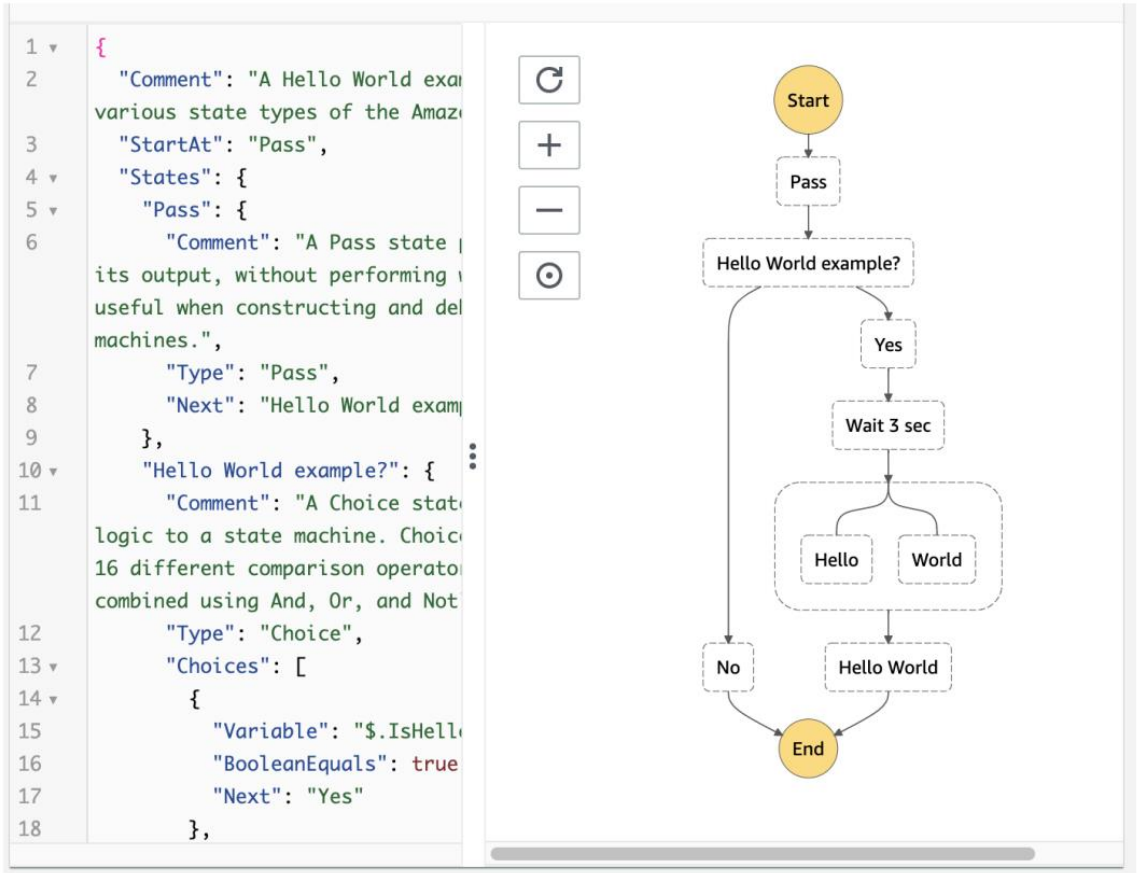

© Copyright by Cloud Mentor Pro | Confidential

Step Function –Task States

- Do some work in your state machine
- Invoke one AWS service
  - Can invoke a **Lambda function**
  - Run an AWS Batch job
  - Run an ECS task and wait for it to complete
  - Insert an item from DynamoDB
  - Publish message to SNS, SQS
  - Launch another Step Function workflow…
- Run an one Activity
  - EC2, Amazon ECS, on-premises
  - Activities poll the Step functions for work
  - Activities send results back to Step Functions

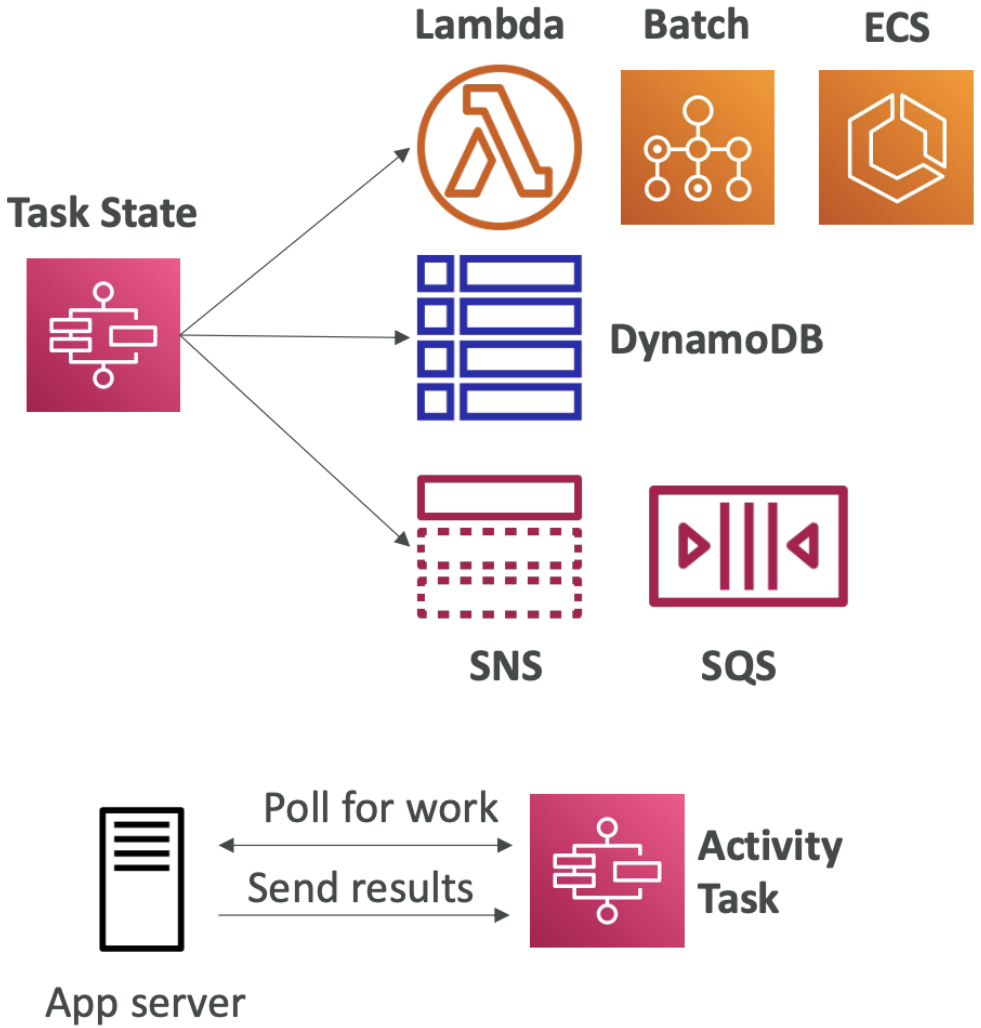

© Copyright by Cloud Mentor Pro | Confidential

Example – Invoke Lambda Function

```
"Invoke Lambda function": {
    "Type": "Task",
    "Resource": "arn:aws:states:::lambda:invoke",
    "Parameters": {
        "FunctionName":
        "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
        "Payload": {
            "Input.$": "$"
        }
    },
    "Next": "NEXT_STATE",
    "TimeoutSeconds": 300
}
```

© Copyright by Cloud Mentor Pro | Confidential

Step Function - States

- **Choice State** -Test for a condition to send to a branch (or default branch)
- **Fail or Succeed State** - Stop execution with failure or success
- **Pass State** - Simply pass its input to its output or inject some fixed data, without performing work.
- **Wait State** - Provide a delay for a certain amount of time or until a specified time/date.
- **Map State** - Dynamically iterate steps.'
- ***Parallel State*** *- Begin parallel branches of execution*.

© Copyright by Cloud Mentor Pro | Confidential

Visual workflow in Step Functions

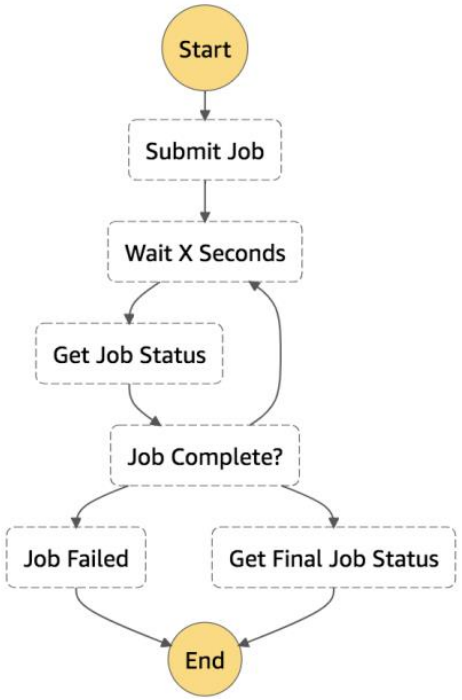

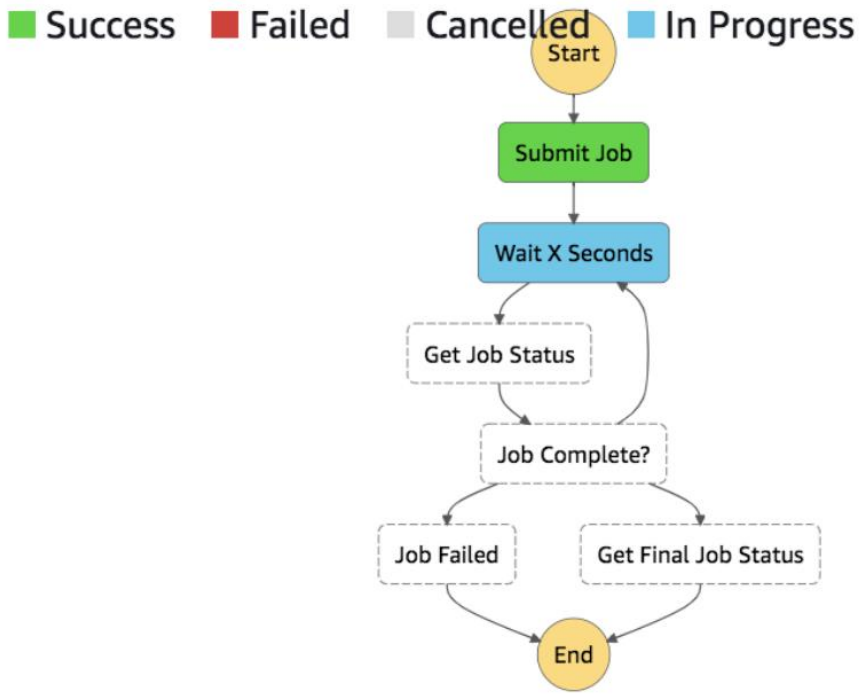

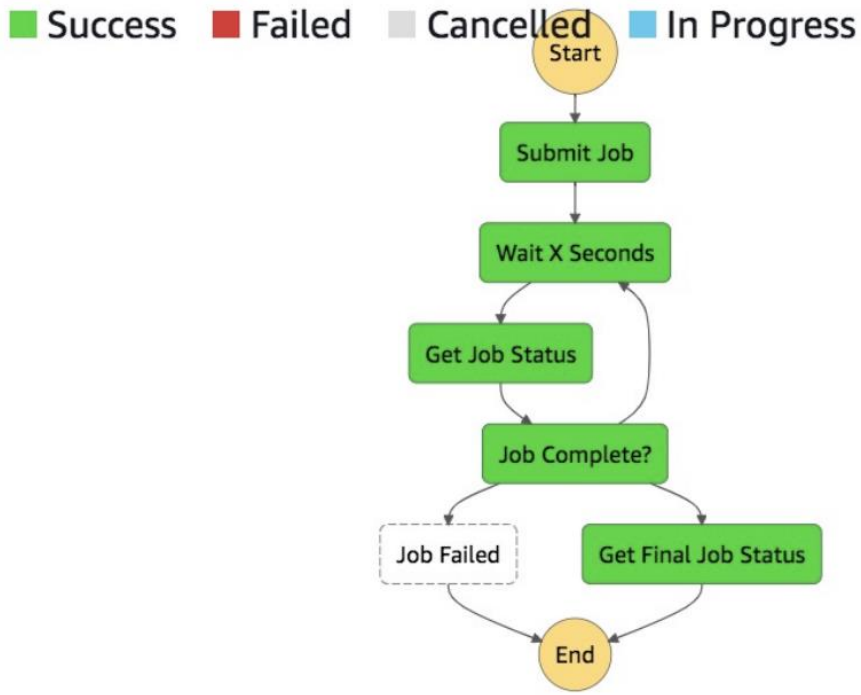

© Copyright by Cloud Mentor Pro | Confidential

Error Handling in Step Functions

- Any state can encounter runtime errors for various reasons:
  - State machine definition issues (for example, no matching rule in a Choice state)
  - Task failures (for example, an exception in a Lambda function)
  - Transient issues (for example, network partition events)
- Use Retry (to retry failed state) and Catch (transition to failure path) in the State Machine to handle the errors instead of inside the Application Code
- Predefined error codes:
  - States.ALL : matches any error name
  - States.Timeout: Task ran longer than TimeoutSeconds or no heartbeat received
  - States.TaskFailed: execution failure
  - States.Permissions: insufficient privileges to execute code
- The state may report is own errors

© Copyright by Cloud Mentor Pro | Confidential

Step Functions – Retry (Task or Parallel State)

```
"HelloWorld": {
  "Type": "Task",
  "Resource": "arn:aws:lambda:REGION:ACCOUNT_ID:function:FUNCTION_NAME",
  "Retry": [
    {
      "ErrorEquals": ["CustomError"],
      "IntervalSeconds": 1,
      "MaxAttempts": 2,
      "BackoffRate": 2.0
    },
    {
      "ErrorEquals": ["States.TaskFailed"],
      "IntervalSeconds": 30,
      "MaxAttempts": 2,
      "BackoffRate": 2.0
    },
    {
      "ErrorEquals": ["States.ALL"],
      "IntervalSeconds": 5,
      "MaxAttempts": 5,
      "BackoffRate": 2.0
    }
  ],
  "End": true
}
```

- Evaluated from top to bottom
- **ErrorEquals**: match a specific kind of error
- **IntervalSeconds**: initial delay before retrying
- **BackoffRate**: multiple the delay after each retry
- **MaxAttempts**: default to 3, set to 0 for never retried
- When max attempts are reached, the **Catch** kicks in

© Copyright by Cloud Mentor Pro | Confidential

Step Functions – Catch (Task or Parallel State)

```
"HelloWorld": {
  "Type": "Task",
  "Resource": "arn:aws:lambda:....",
  "Catch": [
    {
      "ErrorEquals": ["CustomError"],
      "Next": "CustomErrorFallback"
    },
    {
      "ErrorEquals": ["States.TaskFailed"],
      "Next": "ReservedTypeFallback"
    },
    {
       "ErrorEquals": ["States.ALL"],
       "Next": "NextTask",
       "ResultPath": "$.error"
    }
  ],
  "End": true
},
"CustomErrorFallback": {
  "Type": "Pass",
  "Result": "This is a fallback from a custom lambda function exception"
  "End": true
},
```

- Evaluated from top to bottom
- **ErrorEquals**: match a specific kind of error
- **Next**: State to send to
- **ResultPath** - A path that determines what input is sent to the state specified in the Next field.

© Copyright by Cloud Mentor Pro | Confidential

Step Function – ResultPath

```
"HelloWorld": {
  "Type": "Task",
  "Resource": "arn:aws:lambda:....",
  "Catch": [{
      "ErrorEquals": ["States.ALL"],
      "Next": "NextTask",
      "ResultPath": "$.error"
    }],
  "End": true
},
"NextTask": {
  "Type": "Pass",
  "Result": "This is a fallback from a reserved error code",
  "End": true
}
```

- Include the error in the input

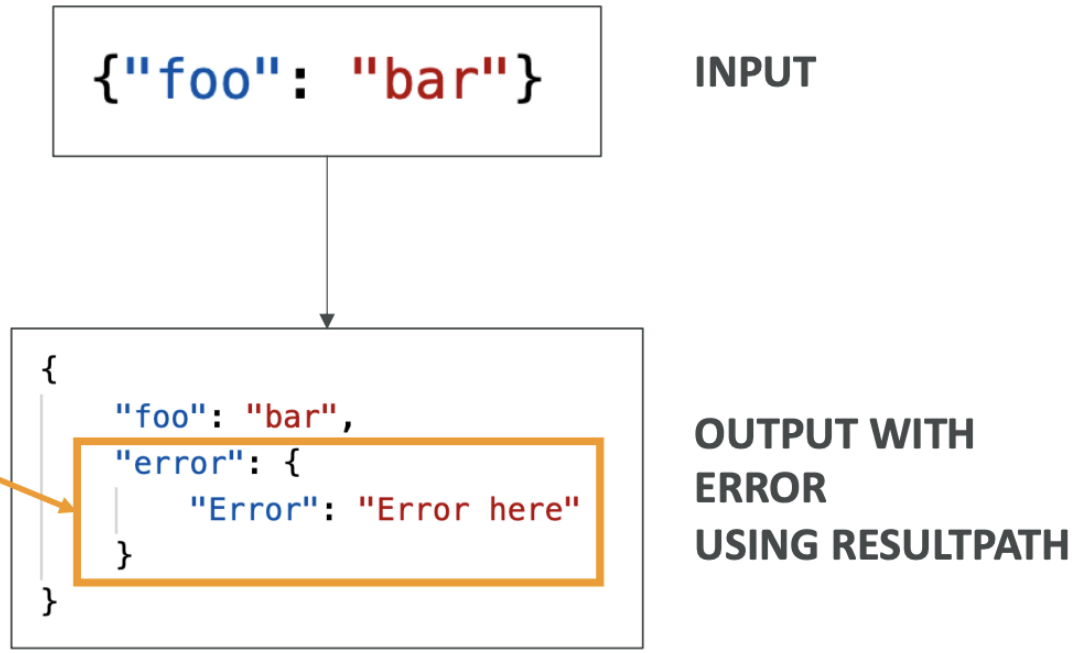

© Copyright by Cloud Mentor Pro | Confidential

Step Functions – Wait for Task Token

- Allows you to pause Step Functions during a Task until a Task Token is returned
- Task might wait for other AWS services, human approval, 3rd party integration, call legacy systems…
- Append .**waitForTaskToken** to the **Resource** field to tell Step Functions to wait for the Task Token to be returned

```
"Resource": "arn:aws:states:::sqs:sendMessage.waitForTaskToken"
```

- Task will pause until it receives that Task Token back with a **SendTaskSuccess** or **SendTaskFailure** API call

© Copyright by Cloud Mentor Pro | Confidential

Step Functions – Wait for Task Token

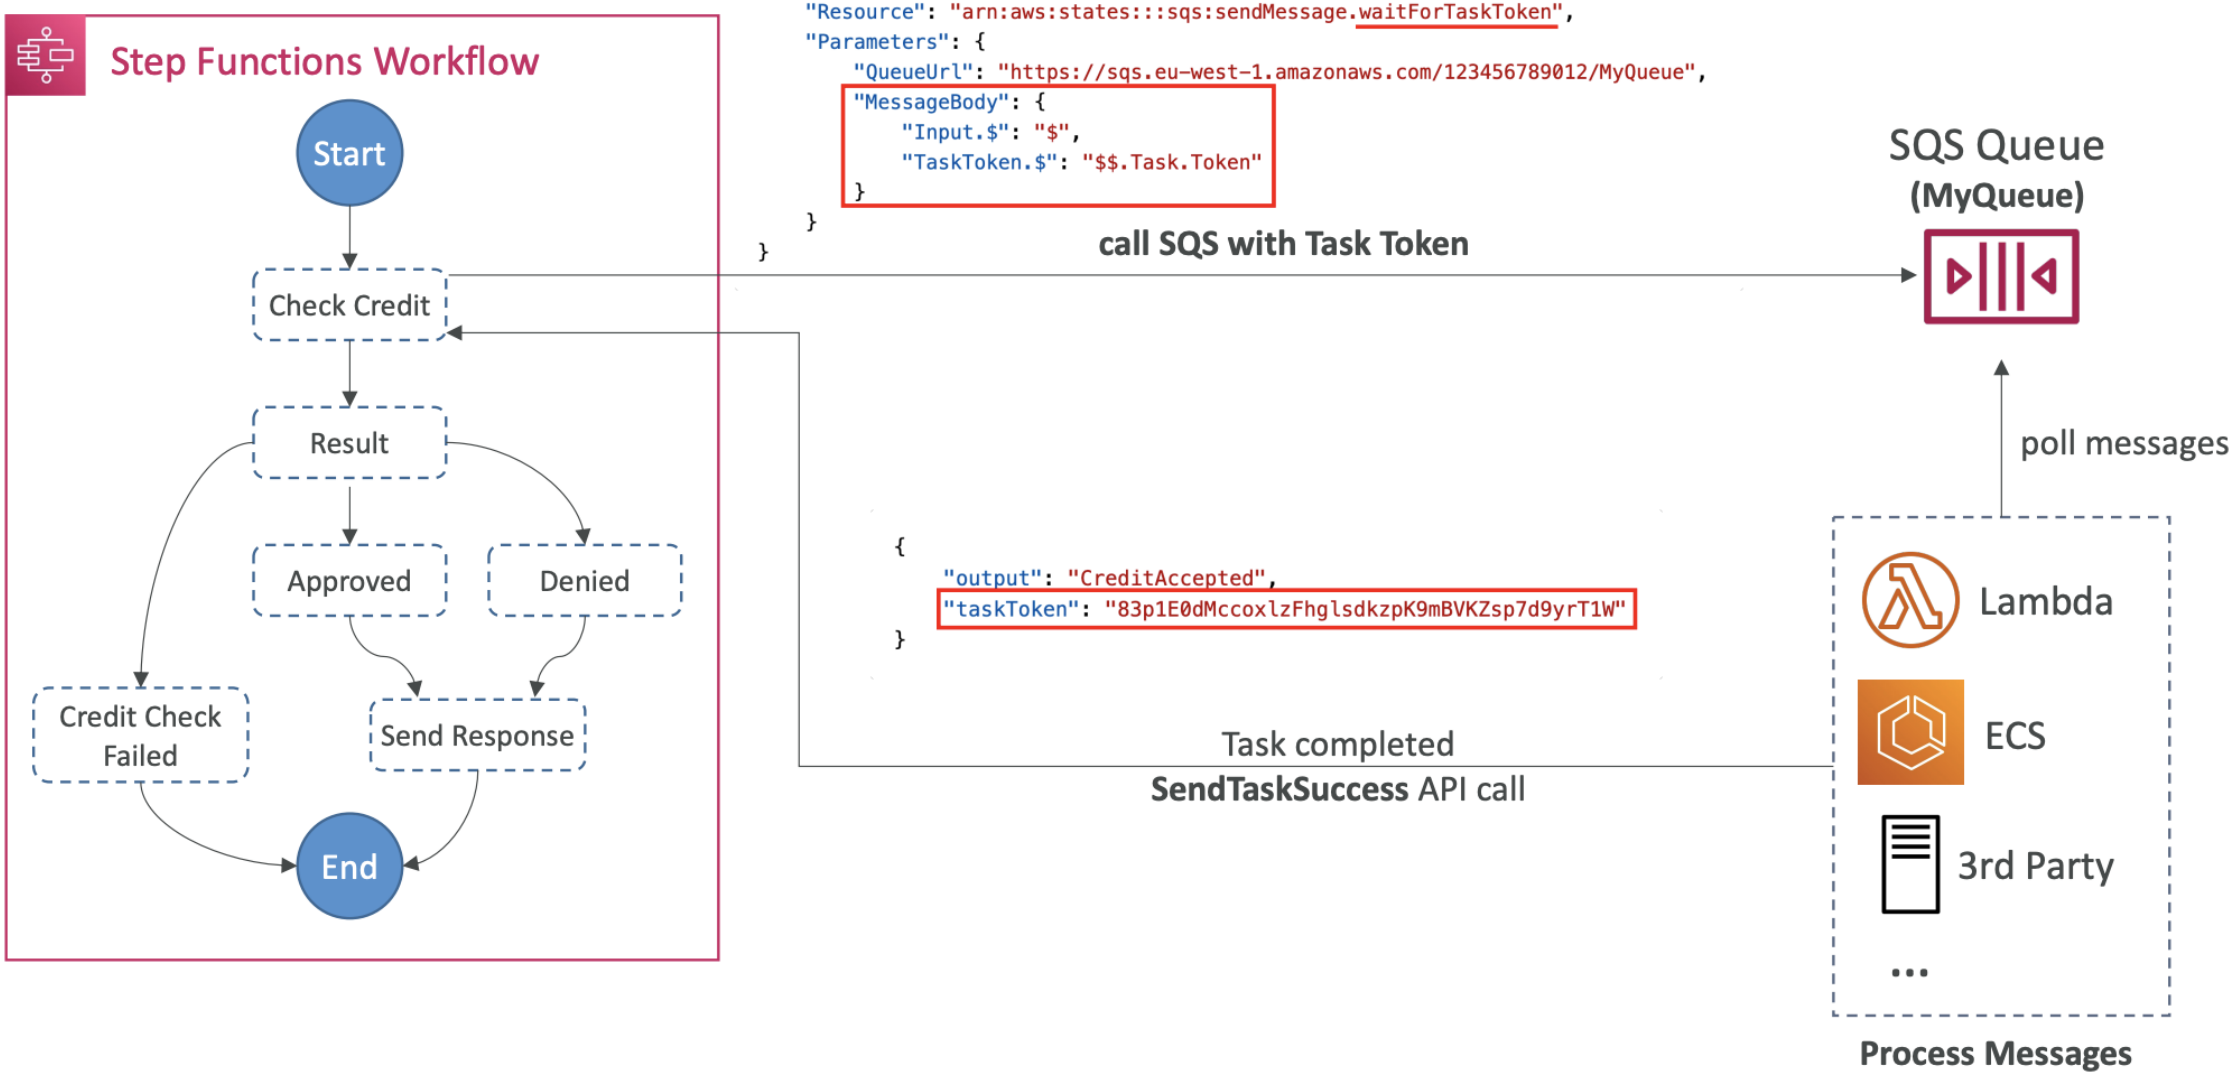

© Copyright by Cloud Mentor Pro | Confidential

Step Functions – Activity Tasks

- Enables you to have the Task work performed by an **Activity Worker**
- **Activity Worker** apps can be running on EC2, Lambda, mobile device…
- Activity Worker poll for a Task using **GetActivityTask** API
- After Activity Worker completes its work, it sends a response of its success/failure using **SendTaskSuccess** or **SendTaskFailure**
- To keep the Task active:
  - Configure how long a task can wait by setting **TimeoutSeconds**
  - Periodically send a heartbeat from your Activity Worker using **SendTaskHeartBeat** within the time you set in **HeartBeatSeconds**
- By configuring a long **TimeoutSeconds** and actively sending a heartbeat, Activity Task can wait up to 1 year

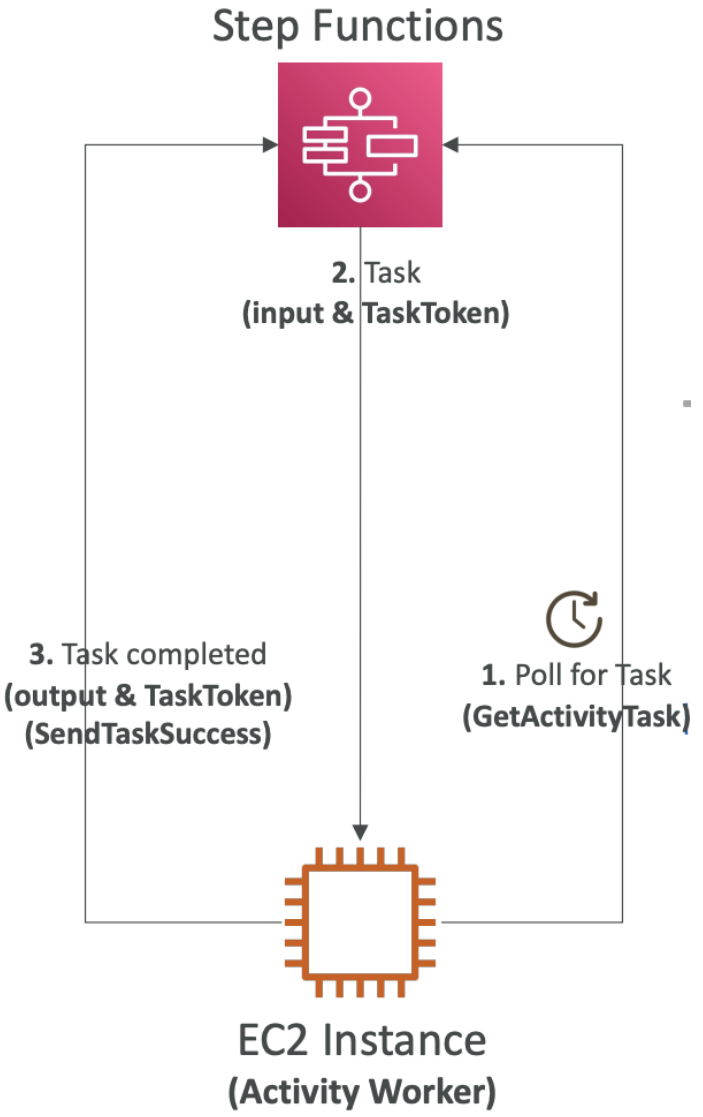

© Copyright by Cloud Mentor Pro | Confidential

Step Functions – Standard vs. Express

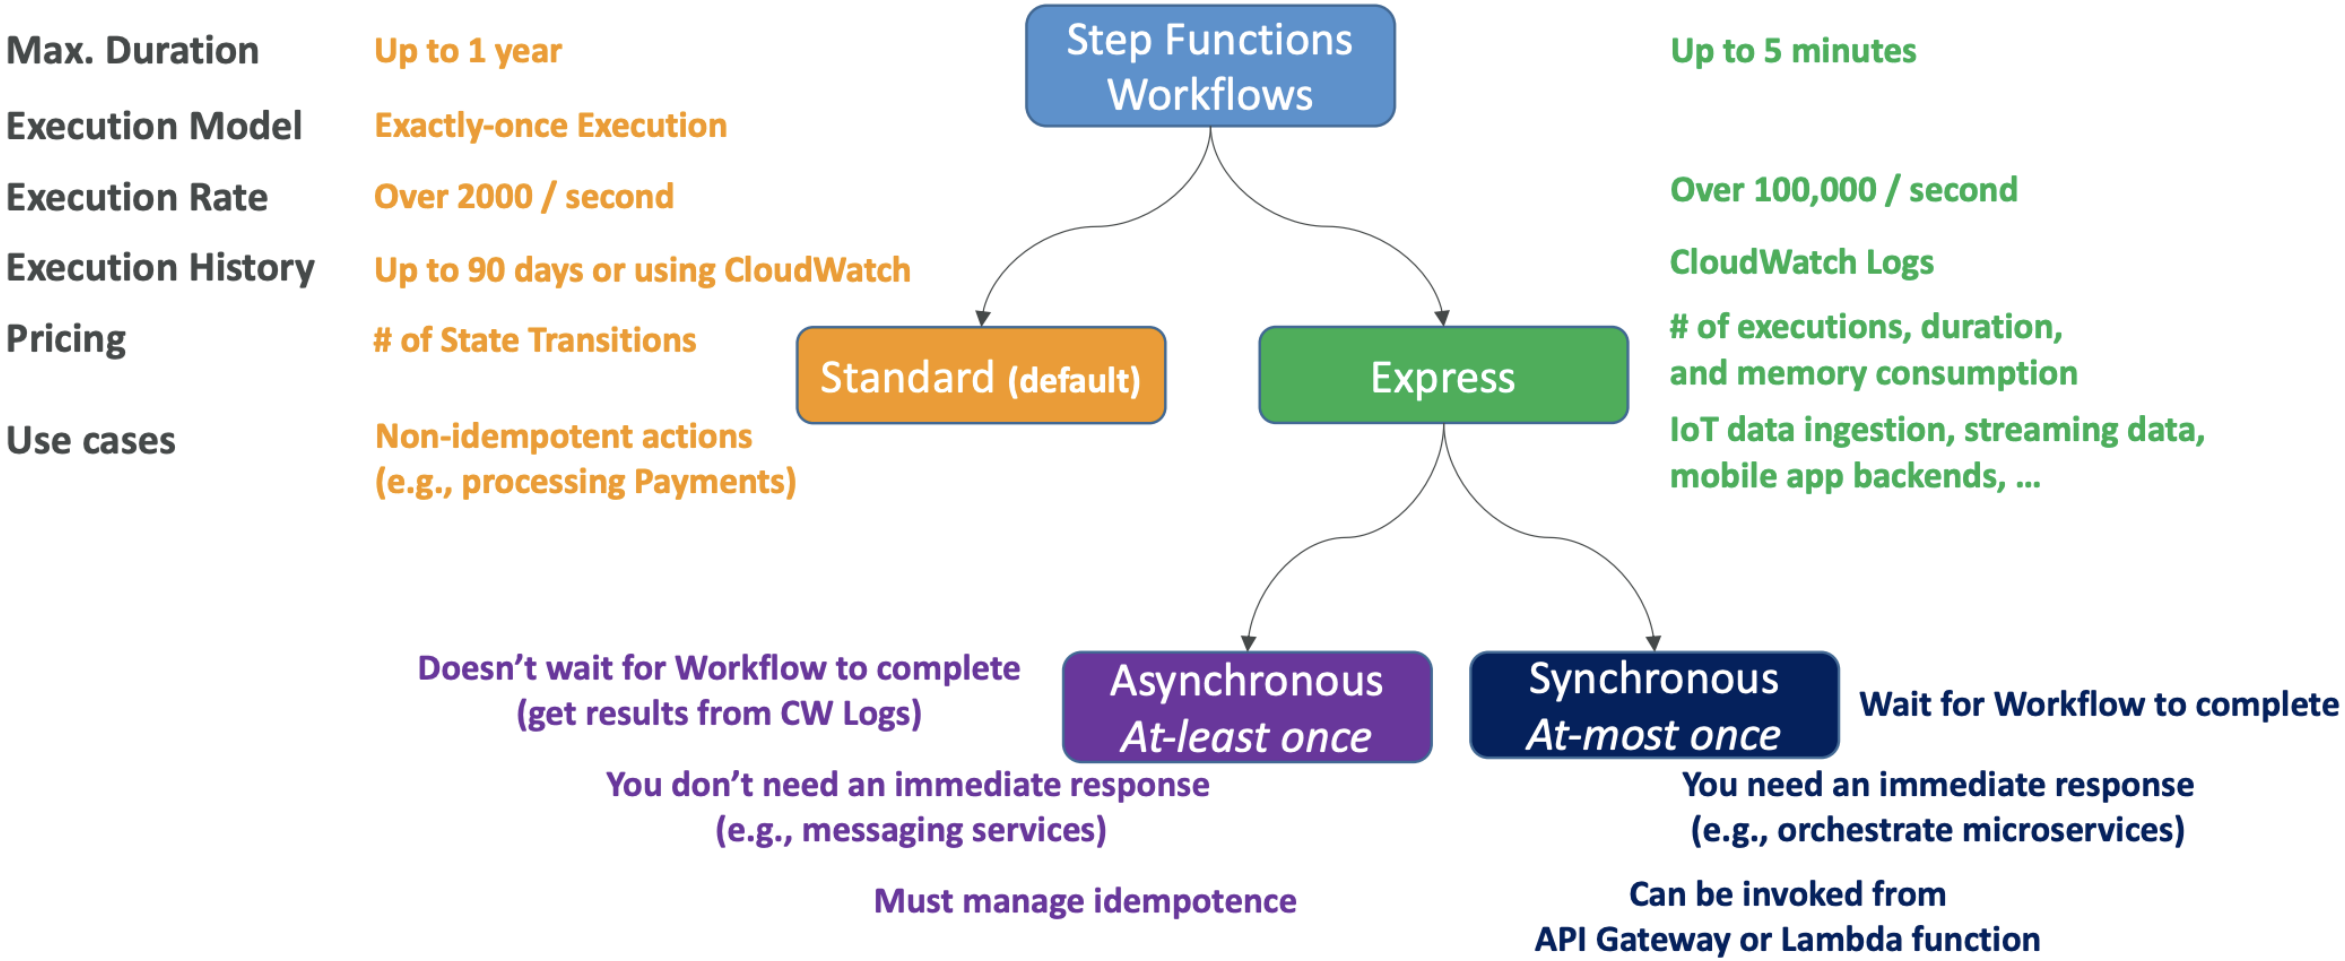

© Copyright by Cloud Mentor Pro | Confidential

AWS AppSync - Overview

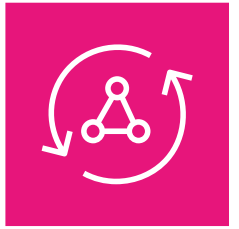

- **AppSync** is a managed service that uses **GraphQL**
- **GraphQL** makes it easy for applications to get exactly the data they need.
- This includes combining data from **one or more sources**
  - NoSQL data stores, Relational databases, HTTP APIs…
  - Integrates with DynamoDB, Aurora, OpenSearch & others
  - Custom sources with AWS Lambda
- Retrieve data in **real-time with WebSocket or MQTT on WebSocket**
- For mobile apps: local data access & data synchronization
- It all starts with uploading one **GraphQL schema**

© Copyright by Cloud Mentor Pro | Confidential

GraphQL Example

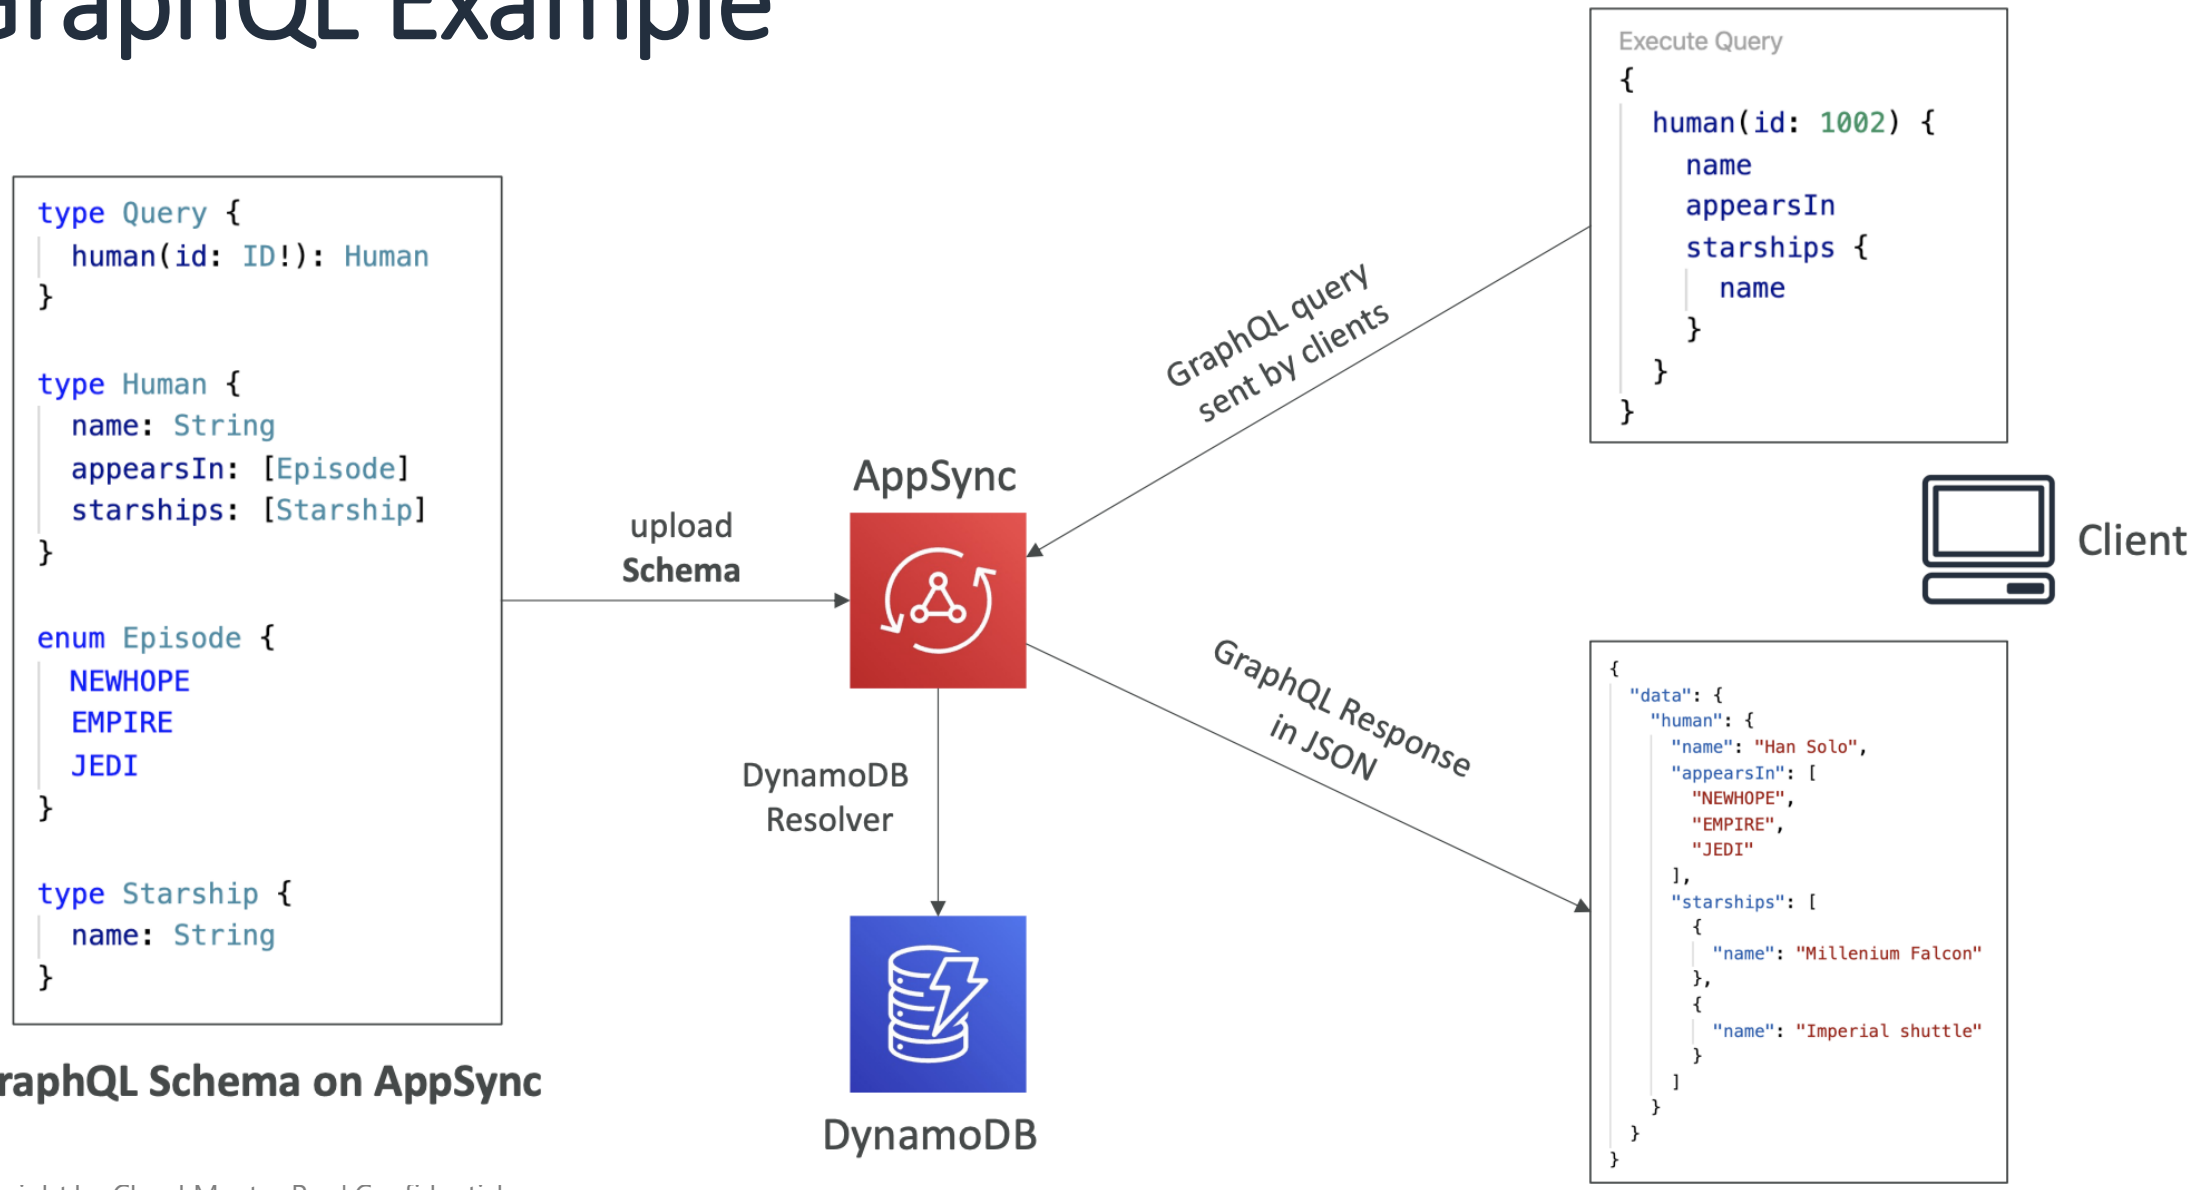

© Copyright by Cloud Mentor Pro | Confidential

AppSync Diagram

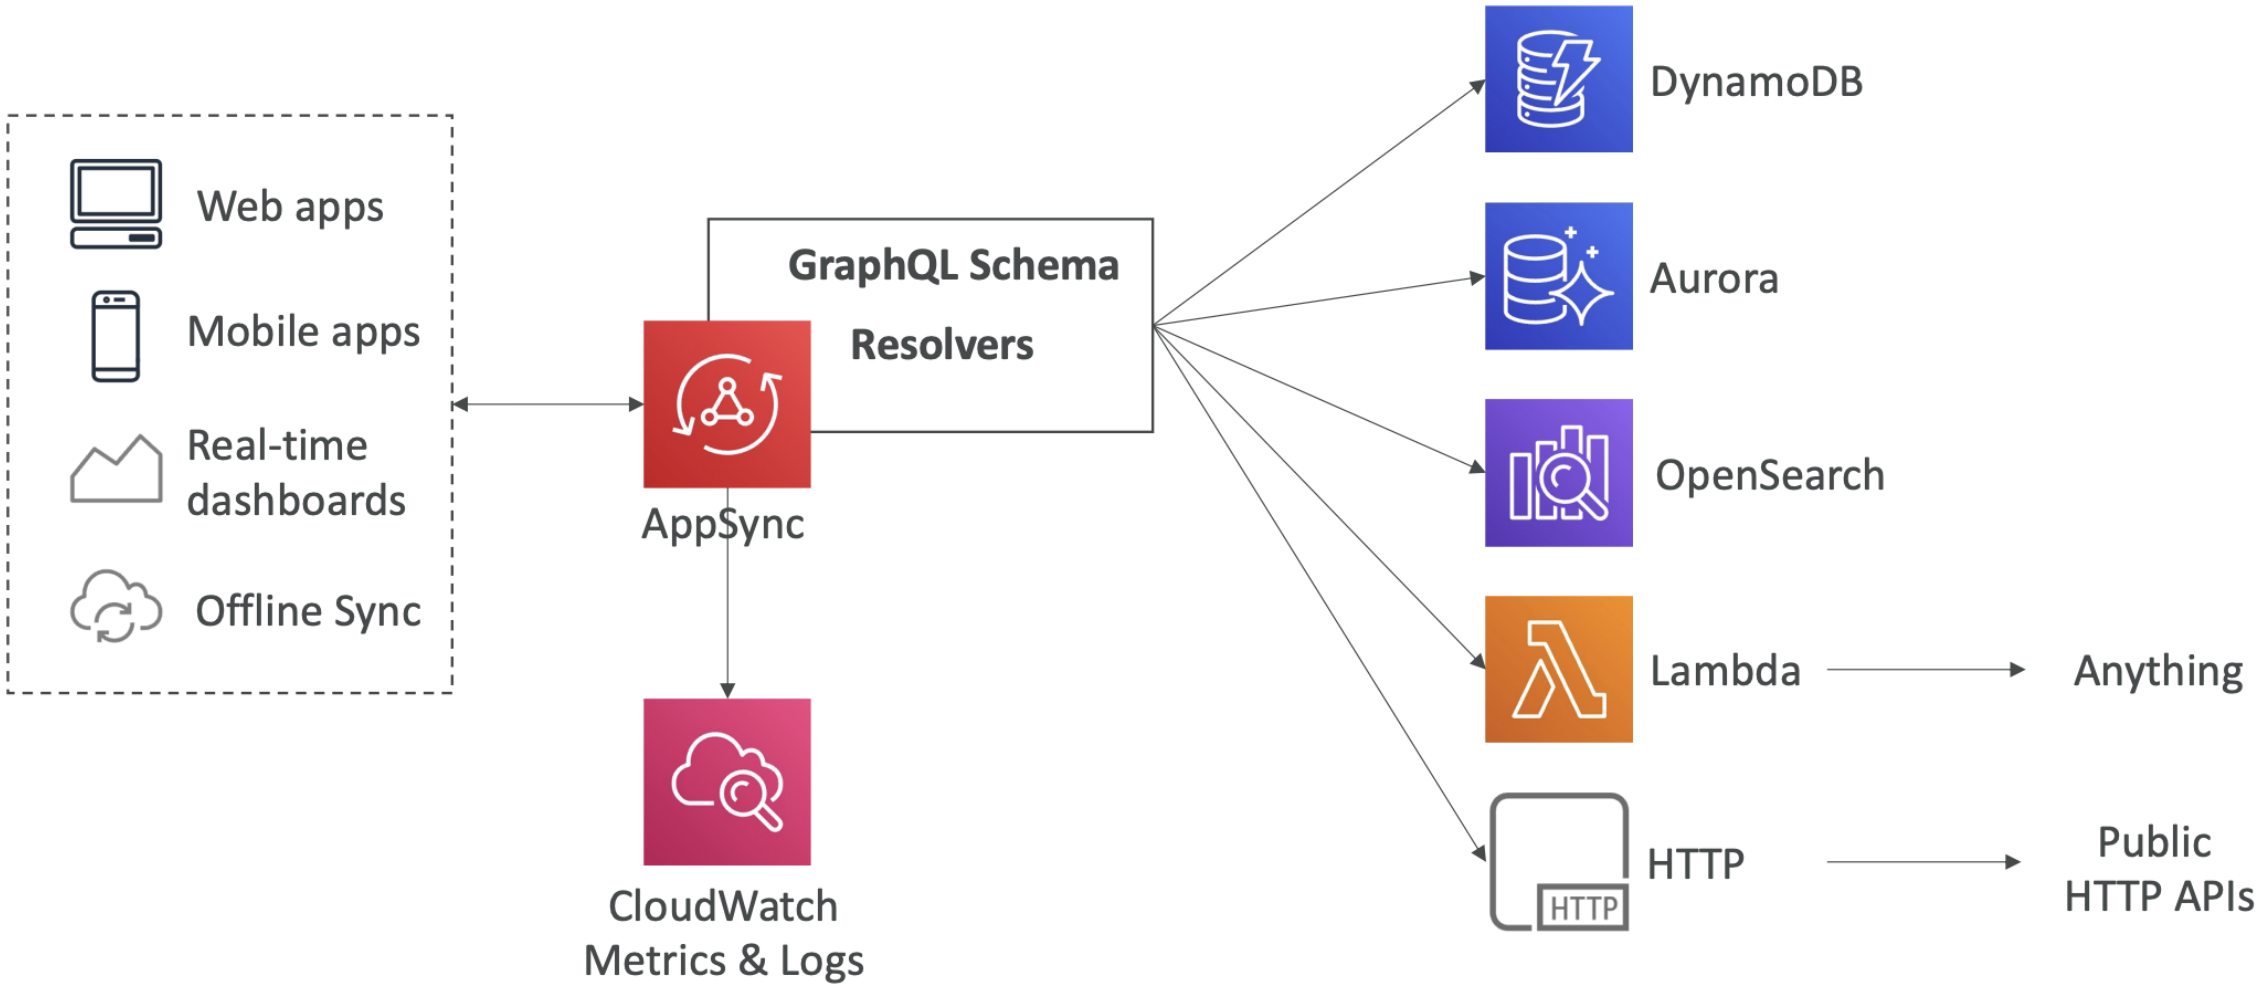

© Copyright by Cloud Mentor Pro | Confidential

AppSync – Security

- There are four ways you can authorize applications to interact with your AWS AppSync GraphQL API:
- **API_KEY**
- **AWS_IA**M: IAM users / roles / cross-account access
- **OPENID_CONNECT**: OpenID Connect provider / JSON Web Token
- **AMAZON_COGNITO_USER_POOLS**

- For custom domain & HTTPS, use CloudFront in front of AppSync

© Copyright by Cloud Mentor Pro | Confidential

AWS Amplify
Create mobile and web applications

Amplify Studio

Visually build a full-stack app,
both front-end UI and a backend.

Amplify CLI

Configure an Amplify backend
With a guided CLI workflow

Amplify Libraries

Connect your app to existing AWS
Services (Cognito, S3 and more)

Amplify Hosting

Host secure, reliable, fast web apps or websites
via the AWS content delivery network.

© Copyright by Cloud Mentor Pro | Confidential

AWS Amplify

- Set of tools to get started with creating **mobile and web applications**
- "Elastic Beanstalk for mobile and web applications"
- Must-have features such as **data storage**, **authentication, storage, and machine-learning**, all powered by AWS services
- **Front-end libraries** with ready-to-use components for React.js, Vue, Javascript, iOS, Android, Flutter, etc…
- Incorporates AWS best practices to for reliability, security, scalability
- Build and deploy with the **Amplify CLI** or **Amplify Studio**

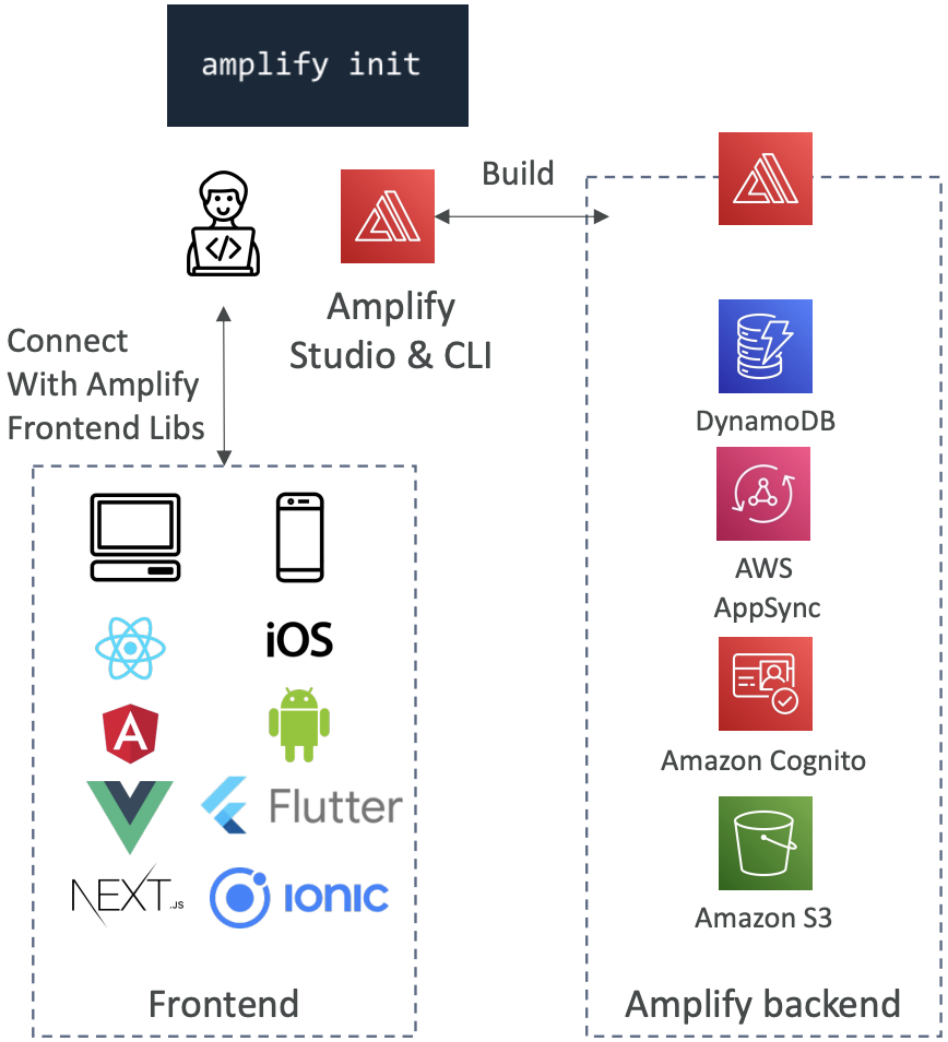

© Copyright by Cloud Mentor Pro | Confidential

AWS Amplify – Important Features

```
amplify add auth
```

AUTHENTICATION

- Leverages Amazon Cognito
- User registration, authentication, account recovery & other operations
- Support MFA, Social Sign-in, etc…
- Pre-built UI components
- Fine-grained authorization

```
amplify add api
```

DATASTORE

- Leverages Amazon AppSync and Amazon DynamoDB
- Work with local data and have **automatic synchronization to the cloud** without complex code
- Powered by GraphQL
- Offline and real-time capabilities
- Visual data modeling w/ Amplify Studio

© Copyright by Cloud Mentor Pro | Confidential

AWS Amplify Hosting

- Build and Host Modern Web Apps
- CICD (build, test, deploy)
- Pull Request Previews
- Custom Domains
- Monitoring
- Redirect and Custom Headers
- Password protection

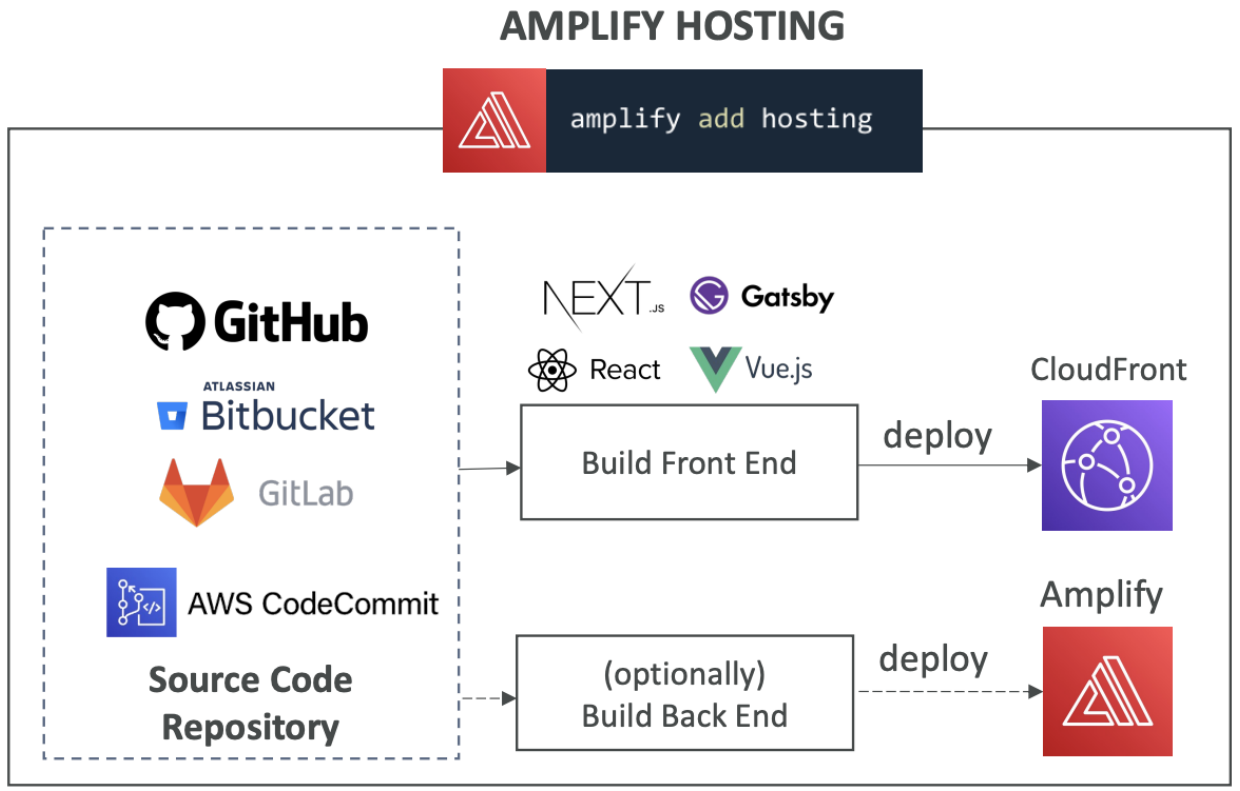

© Copyright by Cloud Mentor Pro | Confidential

AWS Amplify – End-to-End (E2E) Testing

- Run end-to-end (E2E) tests in the **test phase** in Amplify
- Catch regressions before pushing code to production
- Use the test step to run any test commands at build time (**amplify.yml**)
- **Integrated with Cypress testing framework**
  - Allows you to generate UI report for your tests

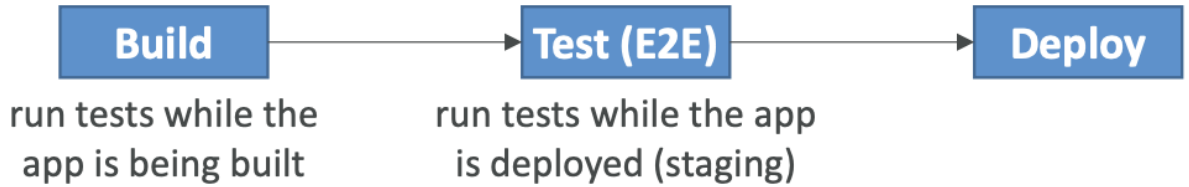


```yaml
test:
  phases:
    preTest:
      commands:
        - npm ci
        - npm install -g pm2
        - npm install -g wait-on
        - npm install mocha mochawesome ...
        - pm2 start npm -- start
        - wait-on http://localhost:3000
    test:
      commands:
        - 'npx cypress run --reporter ...'
    postTest:
      commands:
        - npx mochawesome-merge cypress/...
        - pm2 kill

  artifacts:
    baseDirectory: cypress
    configFilePath: "**/mochawesome.json"
    files:
      - "**/*.png"
      - "**/*.mp4"
```

amplify.yml

© Copyright by Cloud Mentor Pro | Confidential

Exercise

Create, build and Deploy a Sample Hello World App using AWS SAM

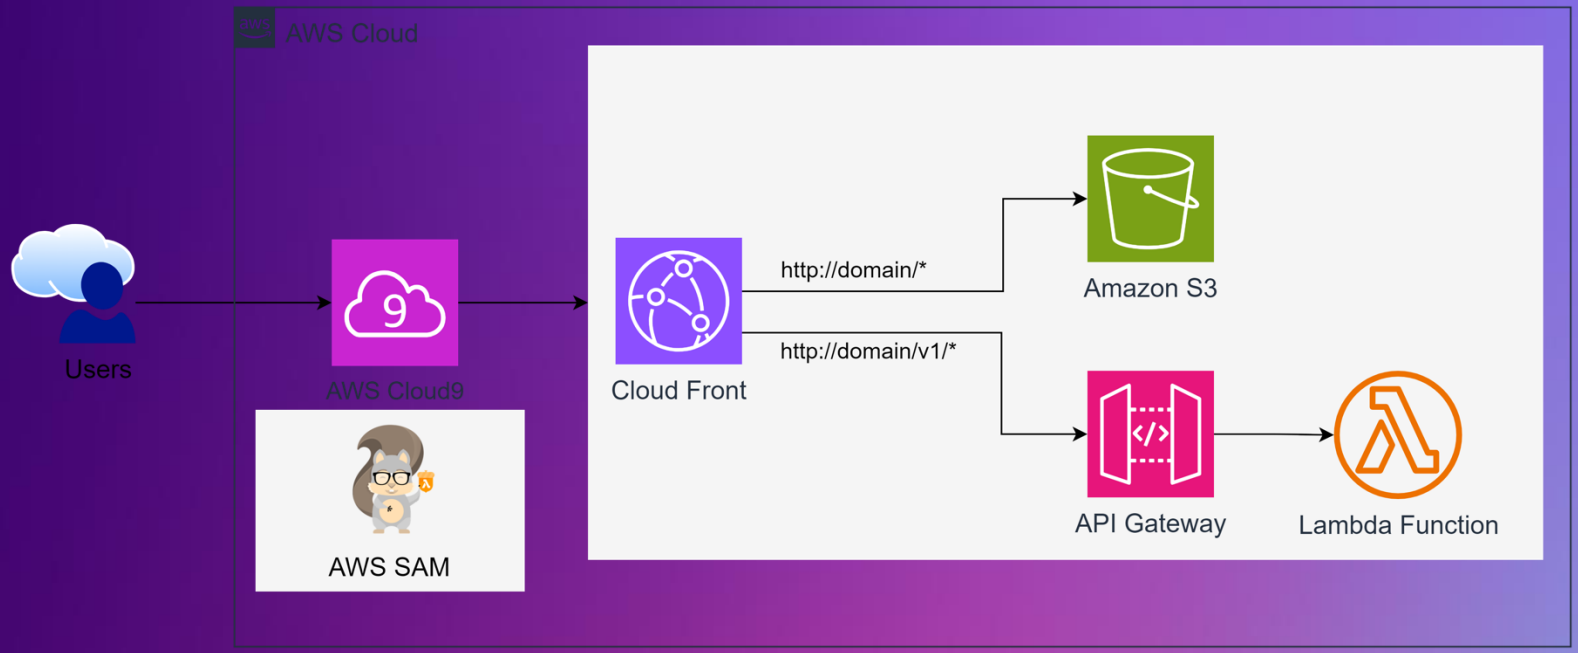

© Copyright by Cloud Mentor Pro | Confidential

Section 15

- *AWS CICD: CodeCommit, CodePipeline, CodeBuild, CodeDeploy*

© Copyright by Cloud Mentor Pro | Confidential

AWS CICD

CodeCommit, CodePipeline, CodeBuild, CodeDeploy, …

© Copyright by Cloud Mentor Pro | Confidential

CICD – Introduction

- We have learned how to:
  - Create AWS resources, manually (fundamentals)
  - Interact with AWS programmatically (AWS CLI)
  - Deploy code to AWS using Elastic Beanstalk
- All these manual steps make it very likely for us to do mistakes!

- We would like our code “in a repository” and have it deployed onto AWS
  - Automatically
  - The right way
  - Making sure it’s tested before being deployed
  - With possibility to go into different stages (dev, test, staging, prod)
  - With manual approval where needed
- To be a proper AWS developer… we need to learn AWS CICD

© Copyright by Cloud Mentor Pro | Confidential

CICD – Introduction

- This section is all about automating the deployment we've done so far while adding increased safety

- We'll learn about:
  - **AWS CodeCommit** – storing our code
  - **AWS CodePipeline** – automating our pipeline from code to Elastic Beanstalk
  - **AWS CodeBuild** – building and testing our code
  - **AWS CodeDeploy** – deploying the code to EC2 instances (not Elastic Beanstalk)
  - **AWS CodeArtifact** – store, publish, and share software packages
  - **AWS CodeGuru** – automated code reviews using Machine Learning

© Copyright by Cloud Mentor Pro | Confidential

Continuous Integration (CI)

- Developers push the code to a code repository often (e.g., GitHub, CodeCommit, Bitbucket…)
- A testing / build server checks the code as soon as it's pushed (CodeBuild, Jenkins CI, …)
- The developer gets feedback about the tests and checks that have passed / failed

- Find bugs early, then fix bugs
- Deliver faster as the code is tested
- Deploy often
- Happier developers, as they're unblocked

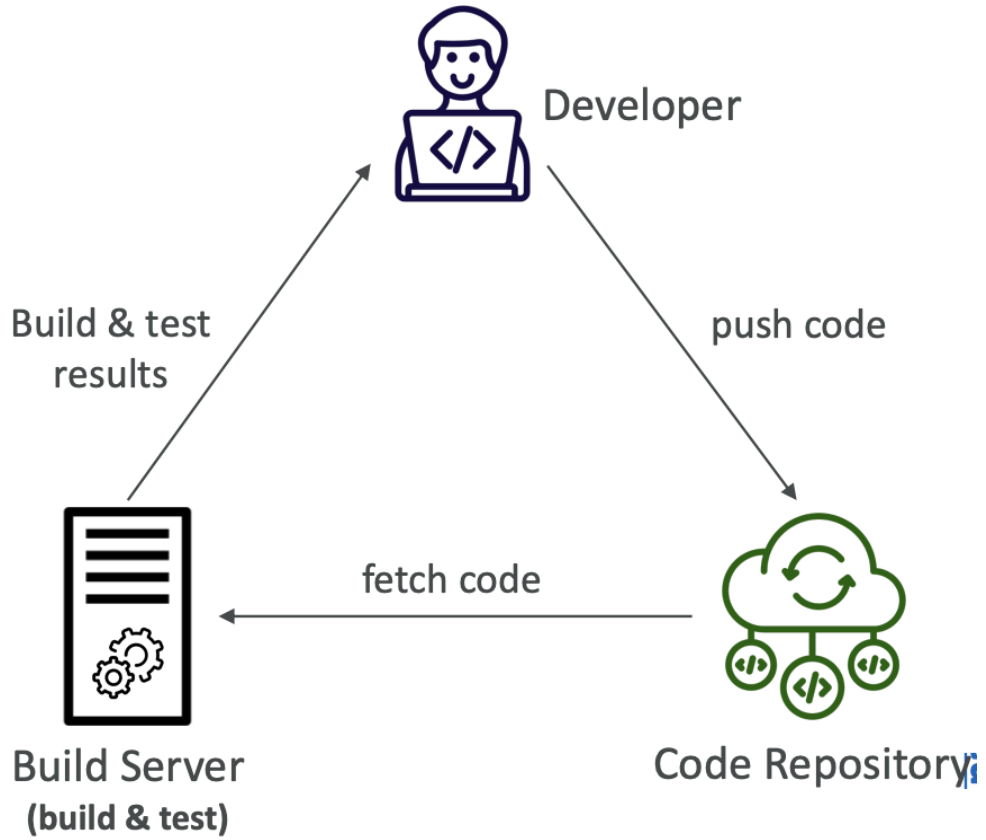

© Copyright by Cloud Mentor Pro | Confidential

Continuous Delivery (CD)

- Ensures that the software can be released reliably whenever needed
- Ensures deployments happen often and are quick
- Shift away from “one release every 3 months” to ”5 releases a day”
- That usually means automated deployment (e.g., CodeDeploy, Jenkins CD, Spinnaker, …)

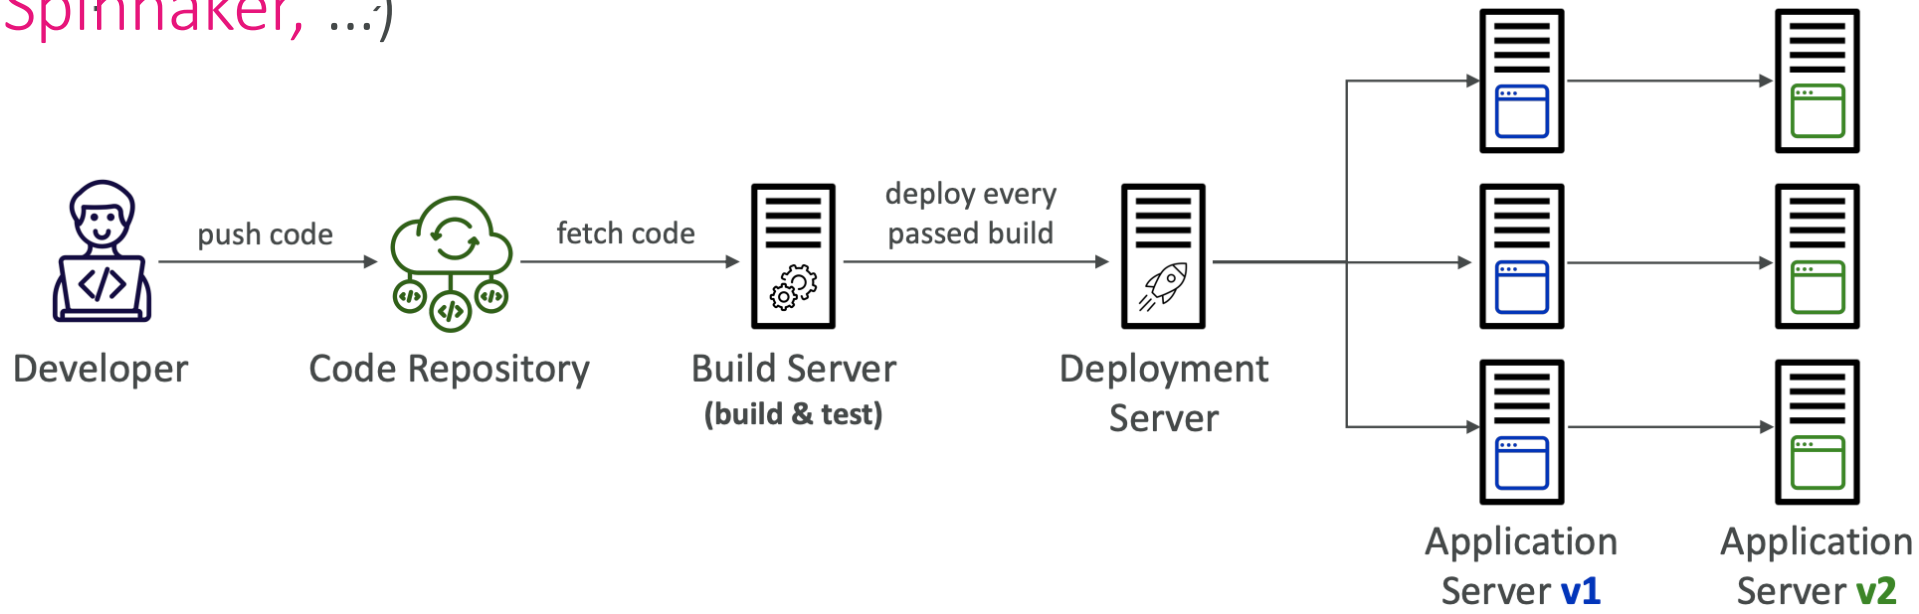

© Copyright by Cloud Mentor Pro | Confidential

Technology Stack for CICD

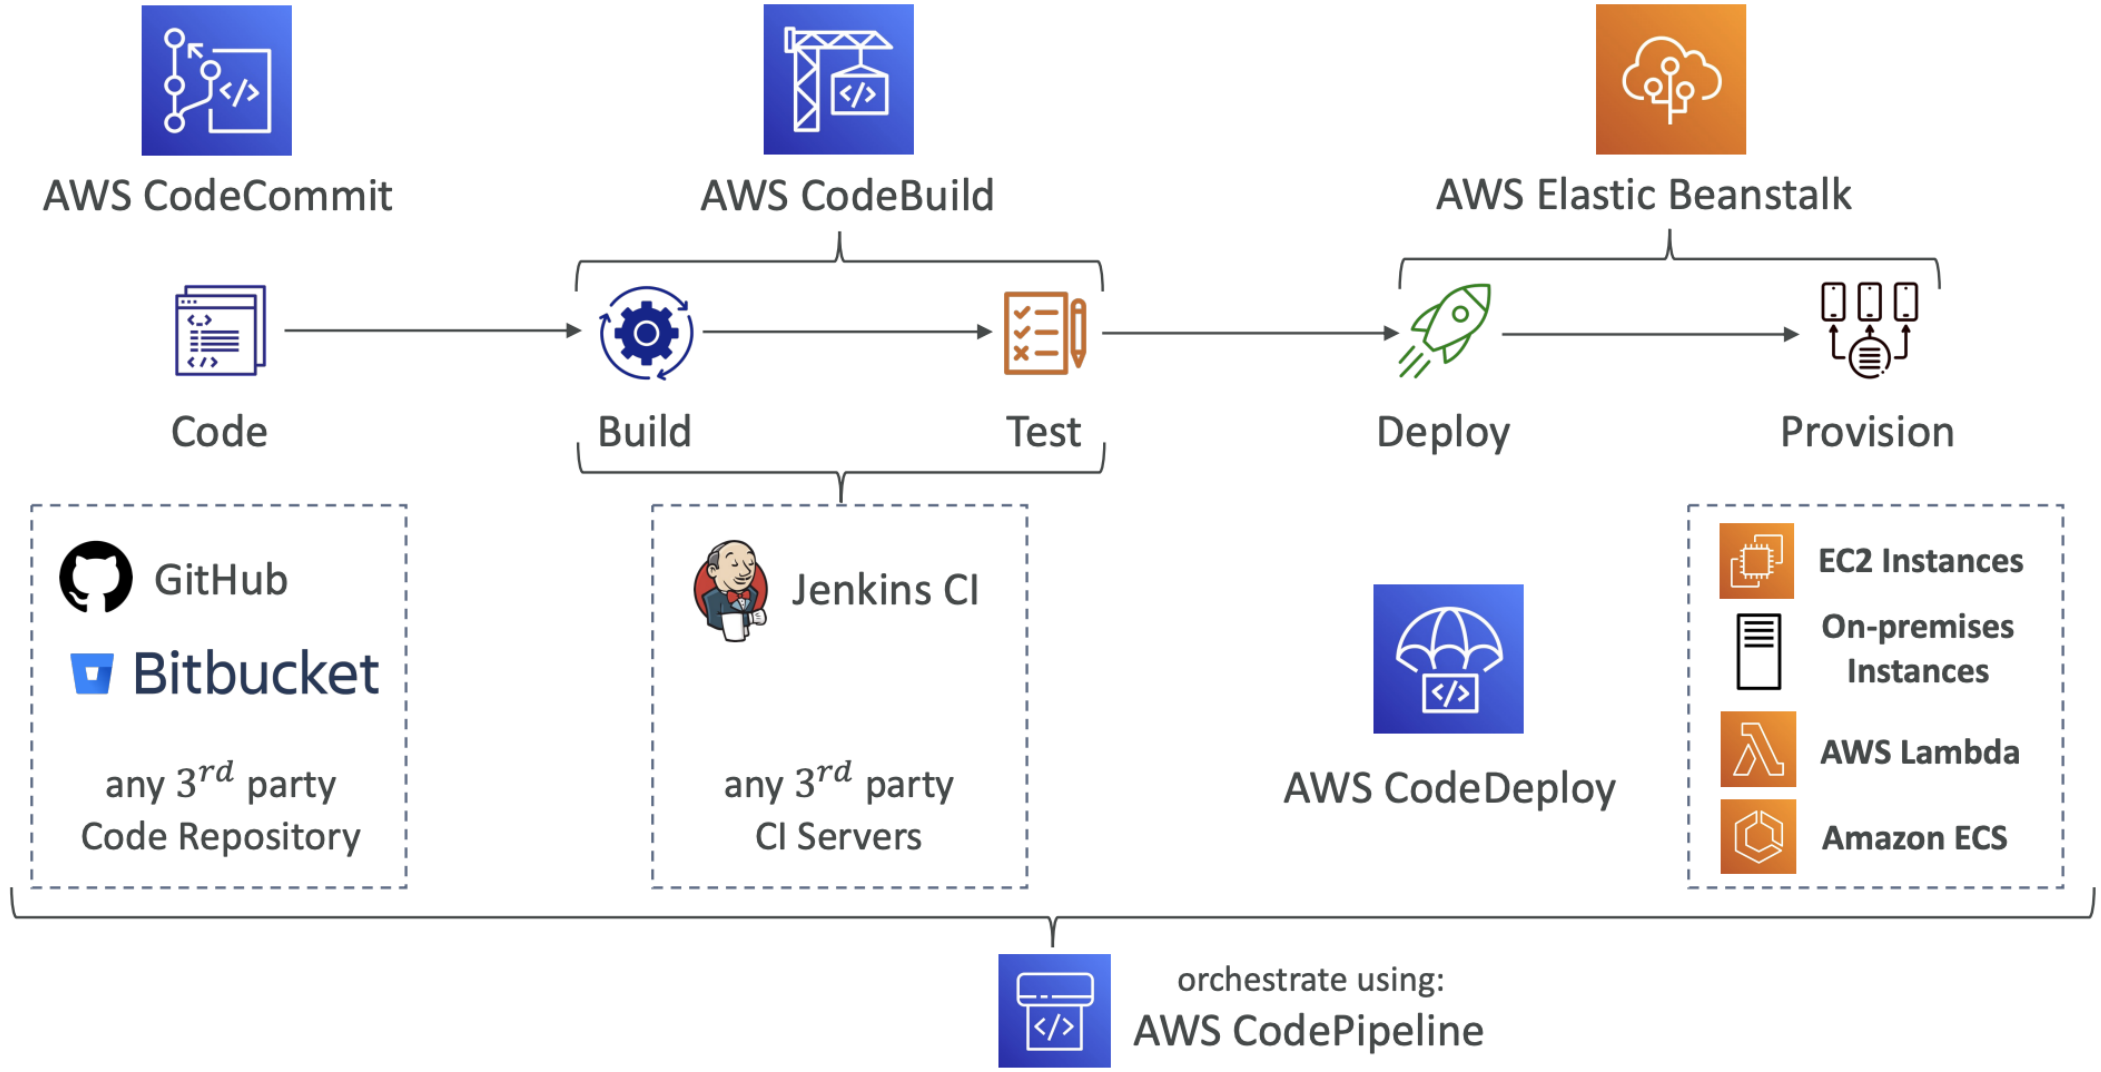

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeCommit

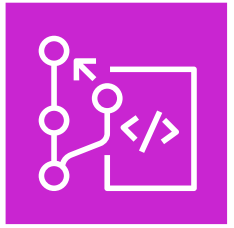

- **Version control** is the ability to understand the various changes that happened to the code over time (and possibly roll back)
- All these are enabled by using a version control system such as **Git**
- A Git repository can be synchronized on your computer, but it usually is uploaded on a central online repository
- Benefits are:
  - Collaborate with other developers
  - Make sure the code is backed-up somewhere
  - Make sure it's fully viewable and auditable

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeCommit

- Git repositories can be expensive
- The industry includes GitHub, GitLab, Bitbucket, …
- And **AWS CodeCommit**:
  - Private Git repositories
  - No size limit on repositories (scale seamlessly)
  - Fully managed, highly available
  - Code only in AWS Cloud account => increased security and compliance
  - Security (encrypted, access control, …)
  - Integrated with Jenkins, AWS CodeBuild, and other CI tools

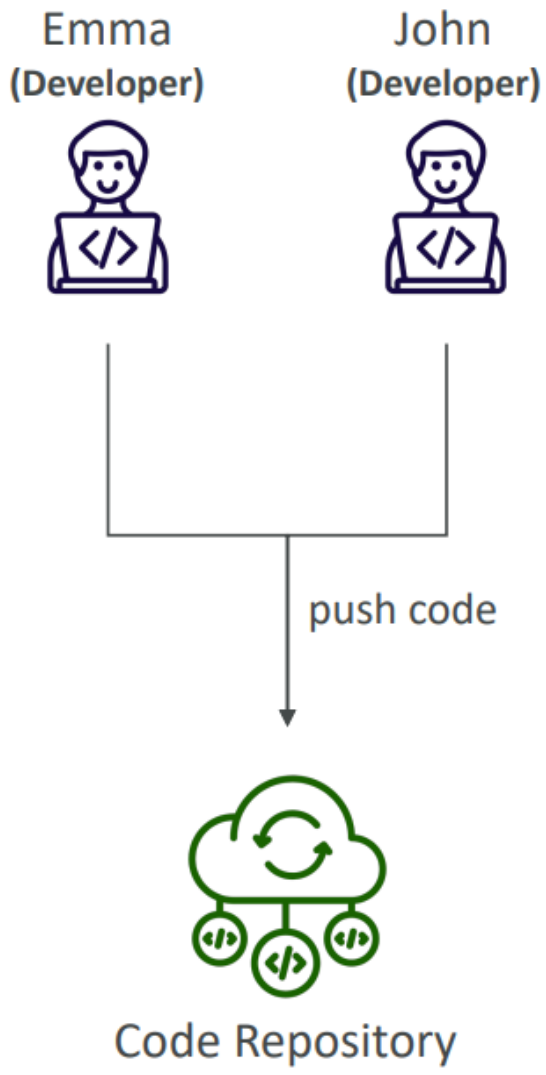

© Copyright by Cloud Mentor Pro | Confidential

CodeCommit – Security

- Interactions are done using Git (standard)
- **Authentication**
  - **SSH Keys** – AWS Users can configure SSH keys in their IAM Console
  - **HTTPS** – with AWS CLI Credential helper or Git Credentials for IAM user
- **Authorization**
  - IAM policies to manage users/roles permissions to repositories
- **Encryption**
  - Repositories are automatically encrypted at rest using AWS KMS
  - Encrypted in transit (can only use HTTPS or SSH – both secure)
- **Cross-account Access**
  - Do NOT share your SSH keys or your AWS credentials
  - Use an IAM Role in your AWS account and use AWS STS (**AssumeRole** API)

© Copyright by Cloud Mentor Pro | Confidential

CodeCommit vs. GitHub

	CodeCommit	GitHub
Support Code Review (Pull Requests)	✔	✔
Integration with AWS CodeBuild	✔	✔
Authentication (SSH & HTTPS)	✔	✔
Security	IAM Users & Roles	GitHub Users
Hosting	Managed & hosted by AWS	- Hosted by GitHub - GitHub Enterprise: self hosted on your servers
UI	Minimal	Fully Featured

© Copyright by Cloud Mentor Pro | Confidential

AWS CodePipeline

- Visual Workflow to orchestrate your CICD
- **Source** – CodeCommit, ECR, S3, Bitbucket, GitHub
- **Build** – CodeBuild, Jenkins, CloudBees, TeamCity
- **Test** – CodeBuild, AWS Device Farm, 3rd party tools, …
- **Deploy** – CodeDeploy, Elastic Beanstalk, CloudFormation, ECS, S3, …
- **Invoke** – Lambda, Step Functions
- Consists of stages:
  - Each stage can have sequential actions and/or parallel actions
  - Example: Build ➡ Test ➡ Deploy ➡ Load Testing ➡ …
  - Manual approval can be defined at any stage

© Copyright by Cloud Mentor Pro | Confidential

CodePipeline – Artifacts

- Each pipeline stage can create **artifacts**
- Artifacts stored in an S3 bucket and passed on to the next stage

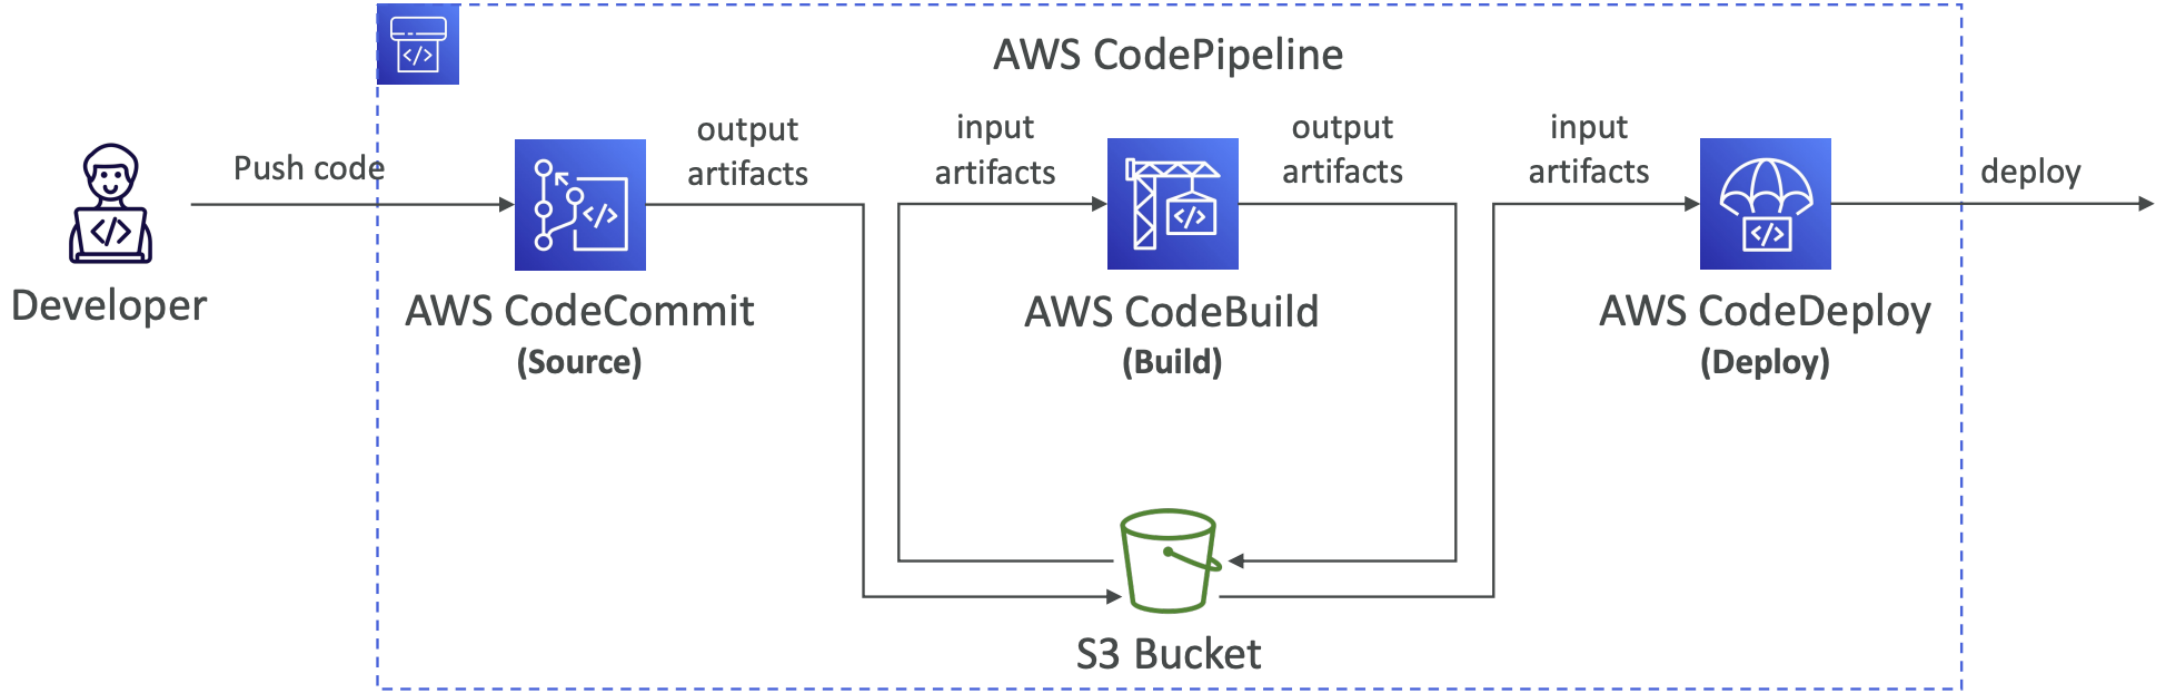

© Copyright by Cloud Mentor Pro | Confidential

CodePipeline –Troubleshooting

- For CodePipeline Pipeline/Action/Stage Execution State Changes
- Use **CloudWatch Events (Amazon EventBridge)**. Example:
  - You can create events for failed pipelines
  - You can create events for cancelled stages

- If CodePipeline fails a stage, your pipeline stops, and you can get information in the console
- If pipeline can't perform an action, make sure the "IAM Service Role" attached does have enough IAM permissions (IAM Policy)
- AWS CloudTrail can be used to audit AWS API calls

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeBuild

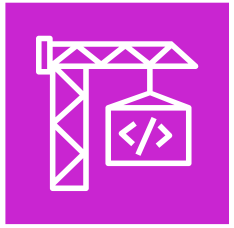

- A fully managed continuous integration (CI) service
- Continuous scaling (no servers to manage or provision – no build queue)
- Compile source code, run tests, produce software packages, …
- Alternative to other build tools (e.g., Jenkins)
- Charged per minute for compute resources (time it takes to complete the builds)
- Leverages Docker under the hood for reproducible builds
- Use prepackaged Docker images or create your own custom Docker image
- Security:
  - Integration with KMS for encryption of build artifacts
  - IAM for CodeBuild permissions, and VPC for network security
  - AWS CloudTrail for API calls logging

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeBuild

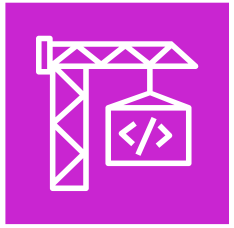

- **Source** – CodeCommit, S3, Bitbucket, GitHub
- **Build instructions**: Code file **buildspec.yml** or insert manually in Console
- **Output logs** can be stored in Amazon S3 & CloudWatch Logs
- Use CloudWatch Metrics to monitor build statistics
- Use EventBridge to detect failed builds and trigger notifications
- Use CloudWatch Alarms to notify if you need “thresholds” for failures

- **Build Projects can be defined within CodePipeline or CodeBuild**

© Copyright by Cloud Mentor Pro | Confidential

CodeBuild – Supported Environments

- Java
- Ruby
- Python
- Go
- Node.js
- Android
- .NET Core
- PHP
- Docker – extend any environment you like

© Copyright by Cloud Mentor Pro | Confidential

CodeBuild – How it Works

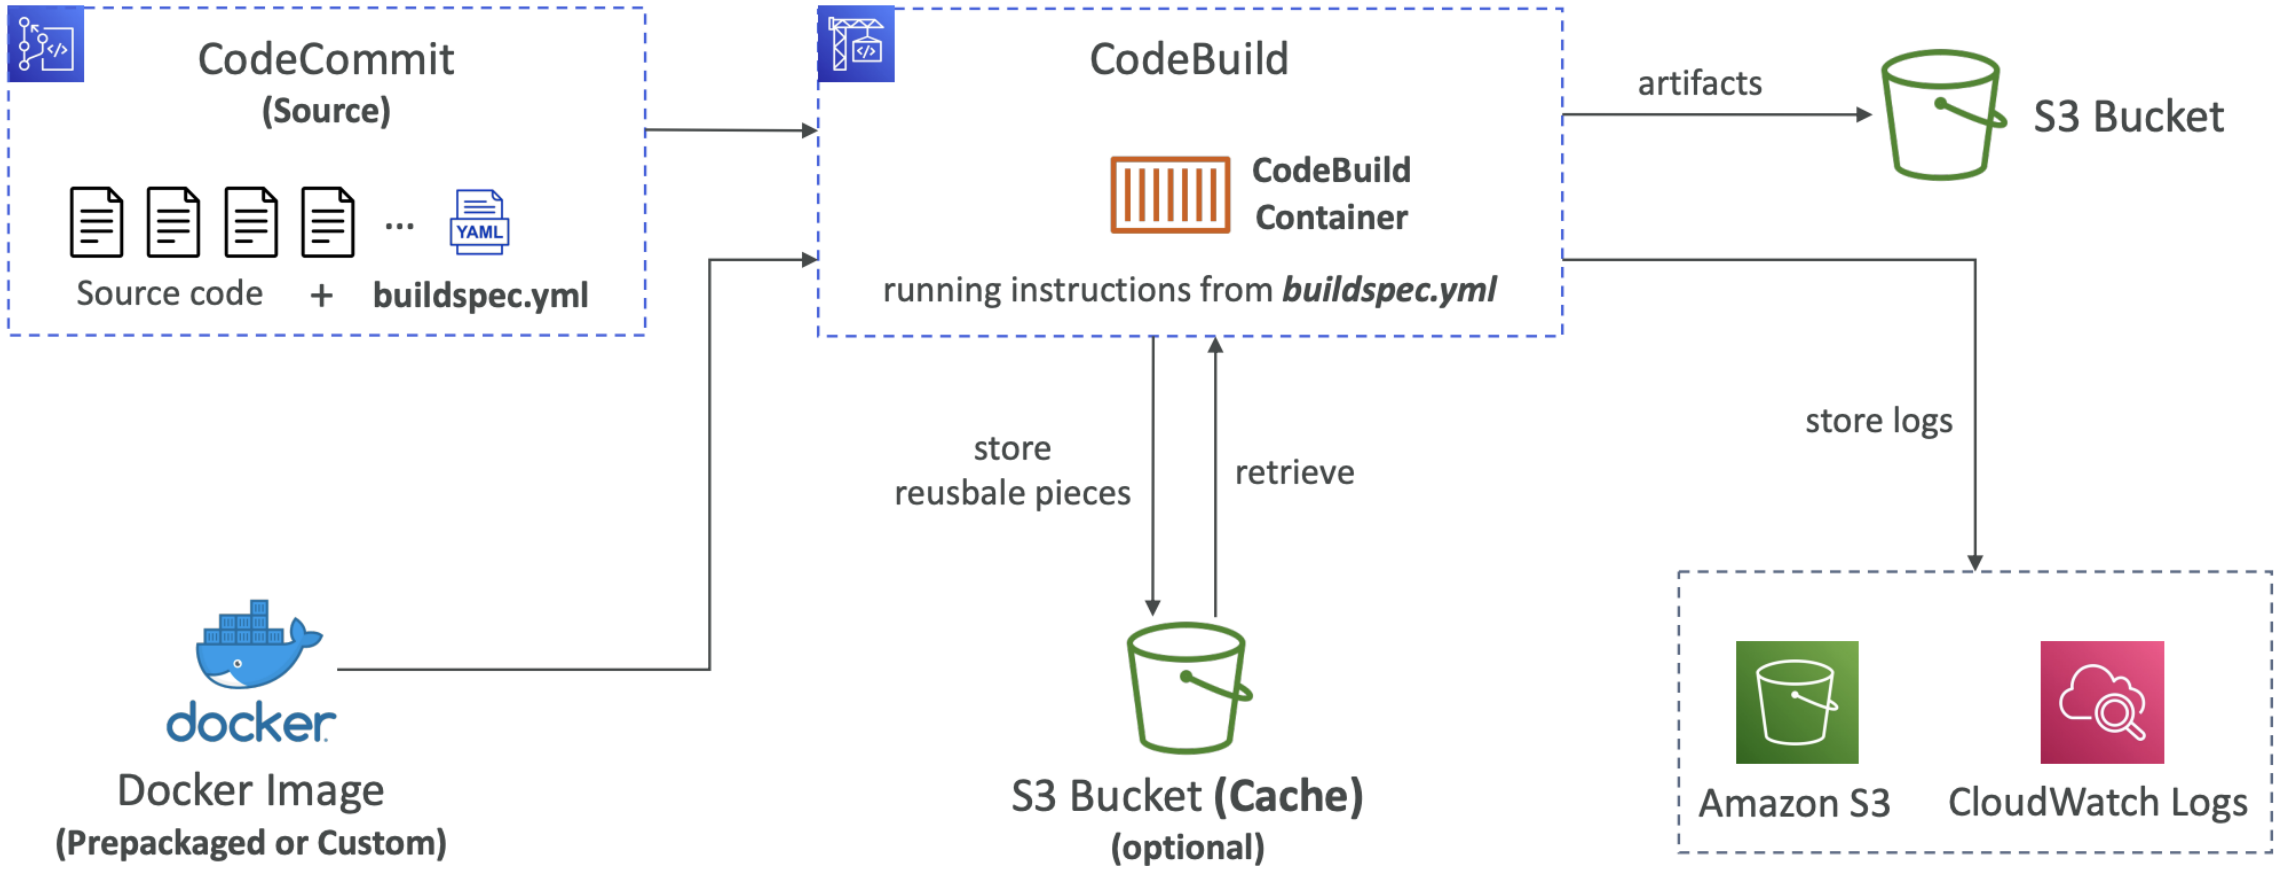

© Copyright by Cloud Mentor Pro | Confidential

CodeBuild – buildspec.yml

- **buildspec.yml** file must be at the **root** of your code
- **env** – define environment variables
  - **variables** – plaintext variables
  - **parameter-store** – variables stored in SSM Parameter Store
  - **secrets-manager** – variables stored in AWS Secrets Manager
- **phases** – specify commands to run:
  - **install** – install dependencies you may need for your build
  - **pre_build** – final commands to execute before build
  - **Build – actual build commands**
  - **post_build** – finishing touches (e.g., zip output)
- **artifacts** – what to upload to S3 (encrypted with KMS)
- **cache** – files to cache (usually dependencies) to S3 for future build speedup

```yaml
version: 0.2

env:
  variables:
    JAVA_HOME: "/usr/lib/jvm/java-8-openjdk-amd64"
  parameter-store:
    LOGIN_PASSWORD: /CodeBuild/dockerLoginPassword

phases:
  install:
    commands:
      - echo "Entered the install phase..."
      - apt-get update -y
      - apt-get install -y maven
  pre_build:
    commands:
      - echo "Entered the pre_build phase..."
      - docker login -u User -p $LOGIN_PASSWORD
  build:
    commands:
      - echo "Entered the build phase..."
      - echo "Build started on `date`"
      - mvn install
  post_build:
    commands:
      - echo "Entered the post_build phase..."
      - echo "Build completed on `date`"

artifacts:
  files:
    - target/messageUtil-1.0.jar

cache:
  paths:
    - "/root/.m2/**/*"
```

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeDeploy

- Deployment service that automates application deployment
- Deploy new applications versions to EC2 Instances, On -premises servers, Lambda functions, ECS Services
- Automated Rollback capability in case of failed deployments, or trigger CloudWatch Alarm
- Gradual deployment control
- A file named appspec.yml defines how the deployment happens

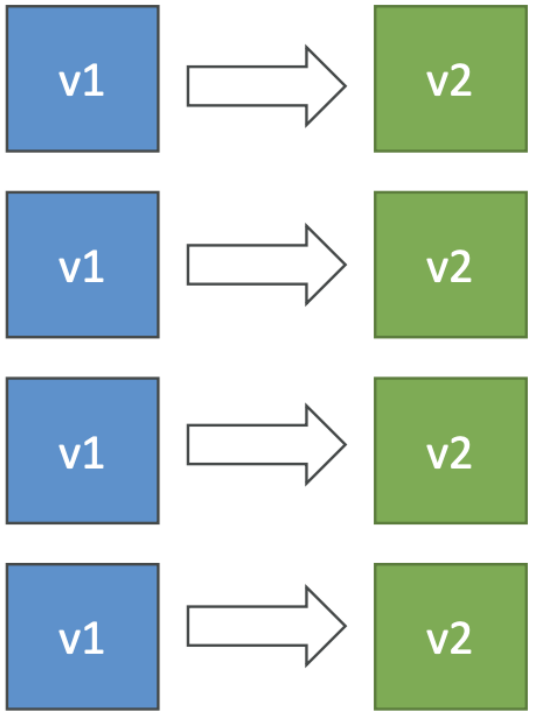

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – EC2/On-premises Platform

- Can deploy to EC2 Instances & on-premises servers
- Perform in-place deployments or blue/green deployments
- Must run the **CodeDeploy Agent** on the target instances
- Define deployment speed
  - **AllAtOnce**: Deploys to all instances at the same time, most downtime but fastest
  - **HalfAtATime**: Deploys to half of the instances at a time, reduces capacity by 50%
  - **OneAtATime**: Deploys to one instance at a time, slowest but least impact capacity
  - **Custom**: define a custom percentage of instances to deploy at a time

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – In -Place Deployment

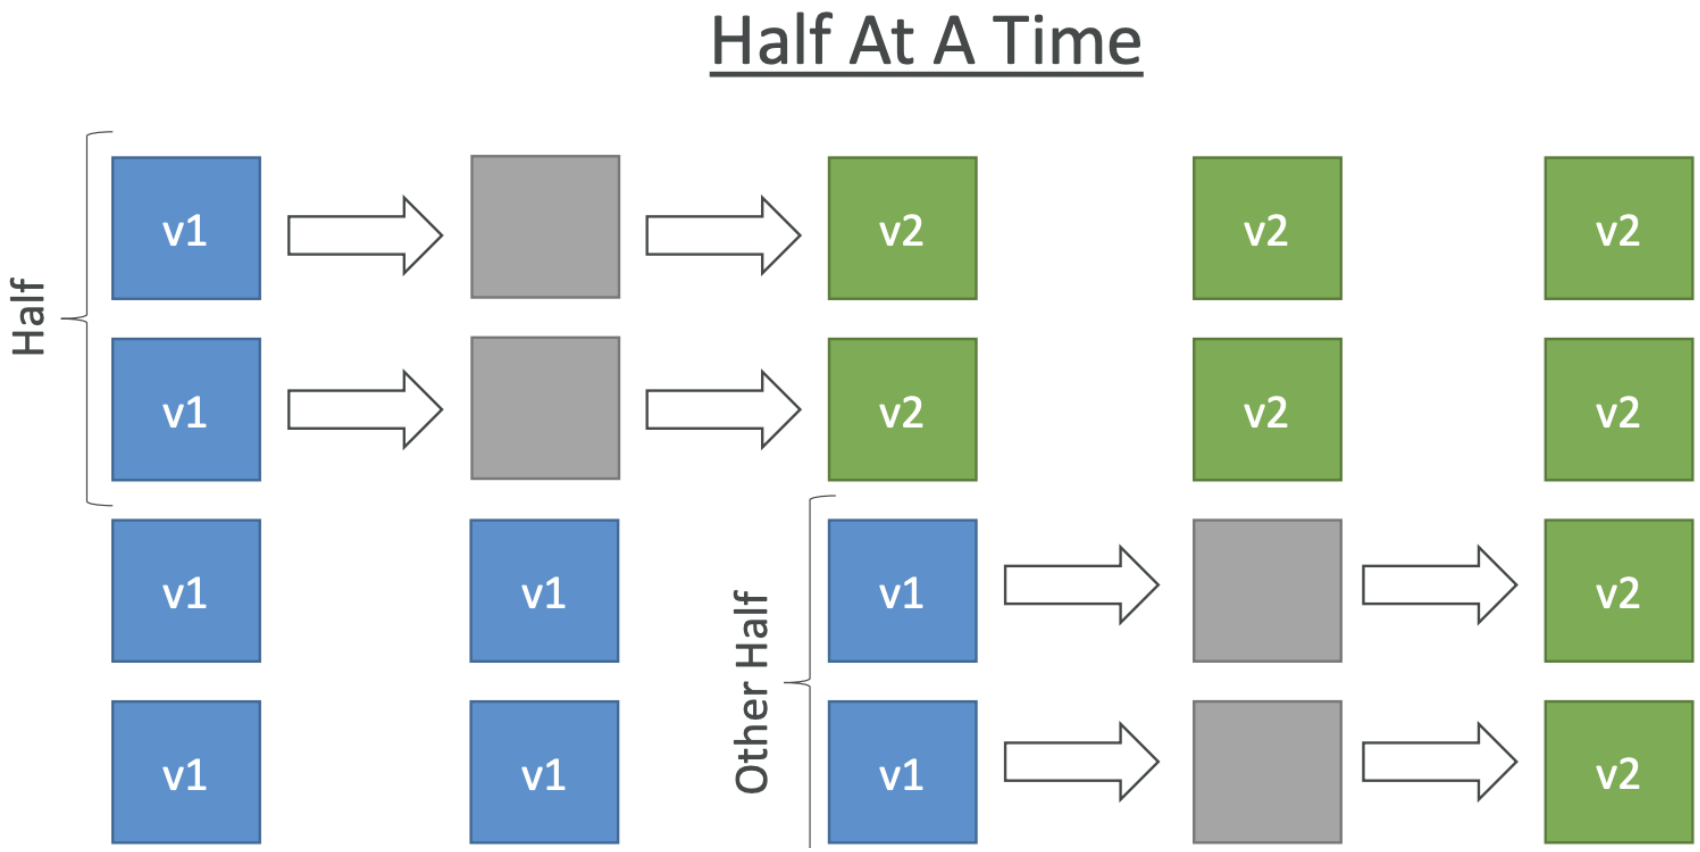

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – Blue-Green Deployment

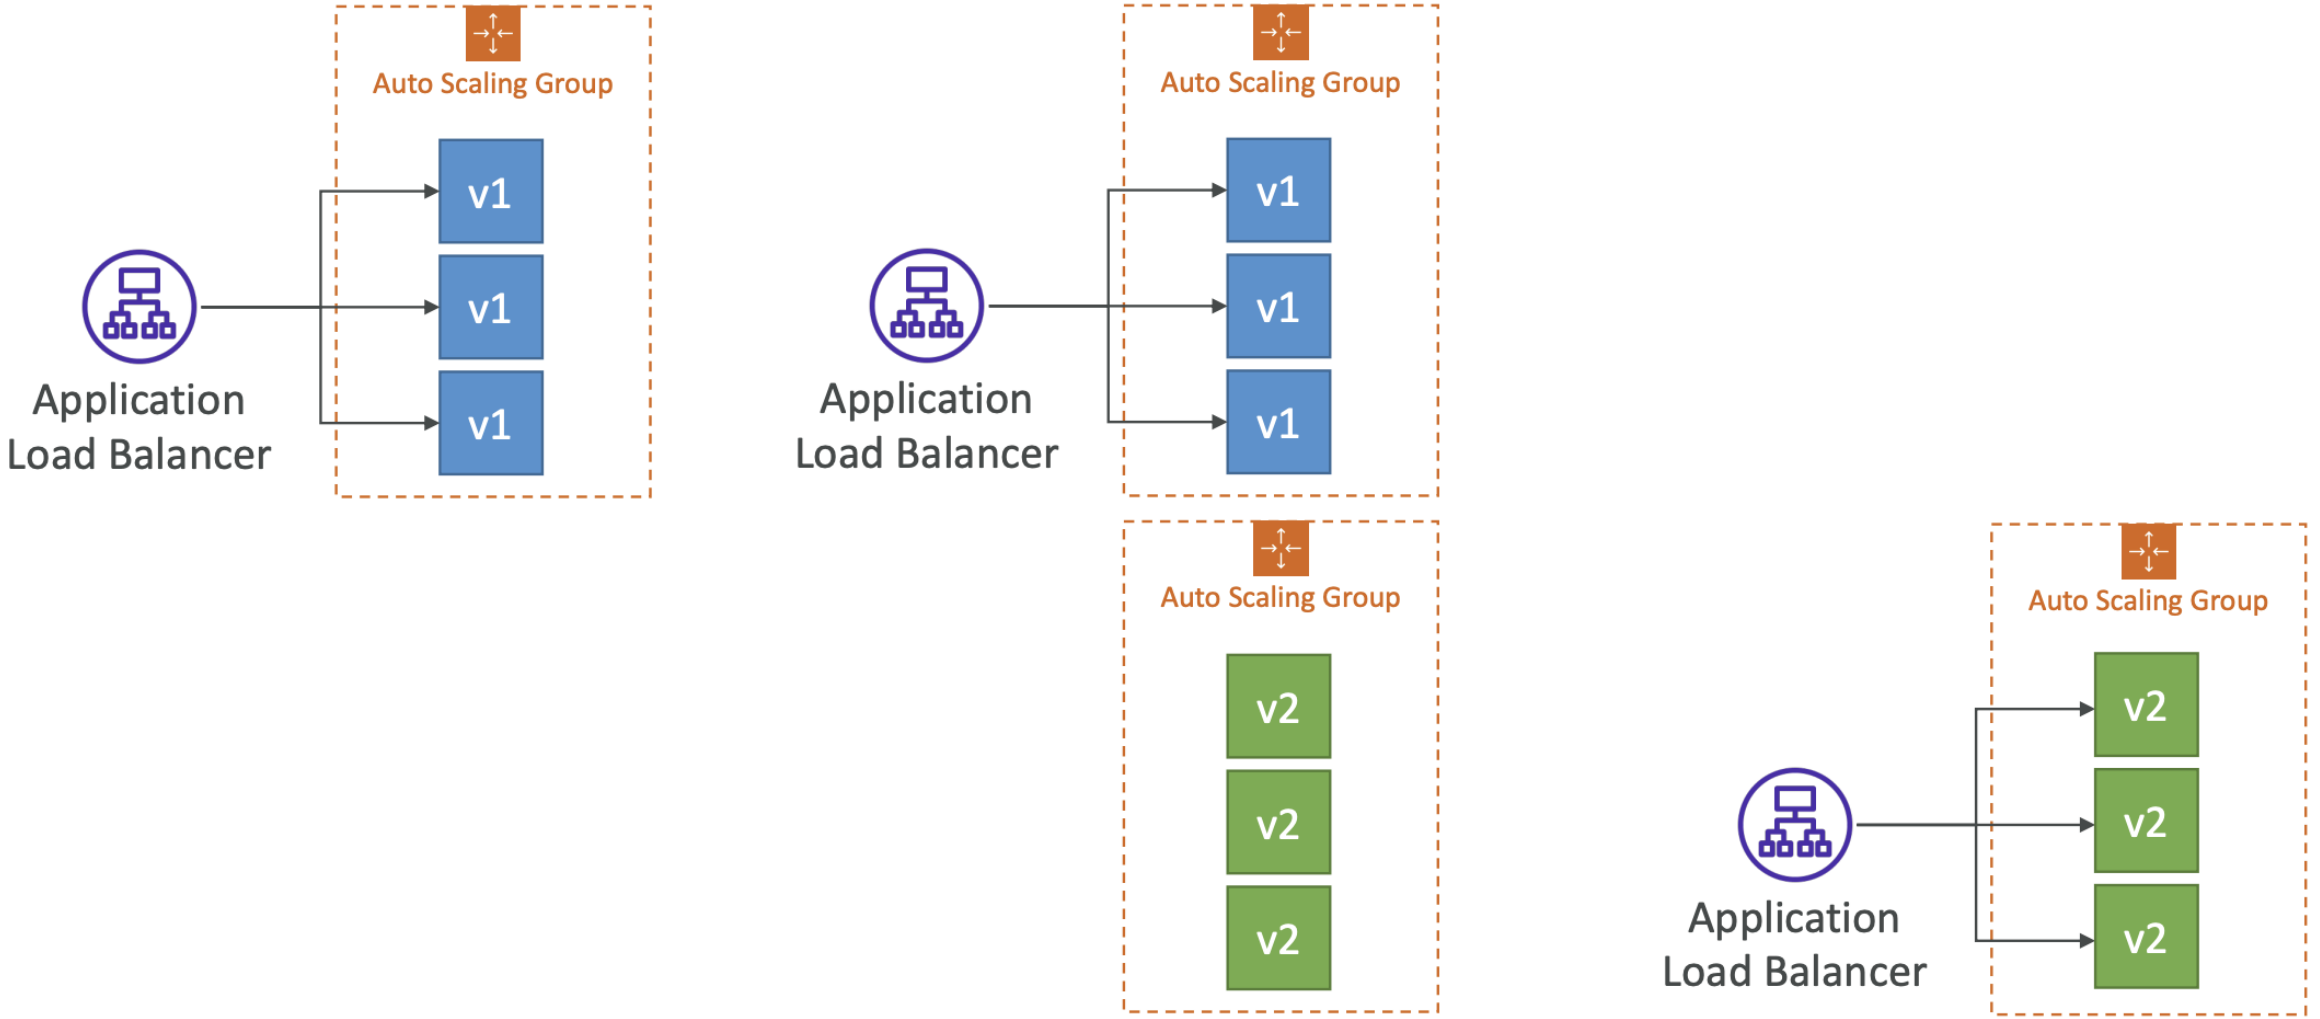

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy Agent

- The CodeDeploy Agent must be running on the EC2 instances as a pre-requisites
- It can be installed and updated automatically if you're using Systems Manager
- The EC2 Instances must have sufficient permissions to access Amazon S3 to get deployment bundles

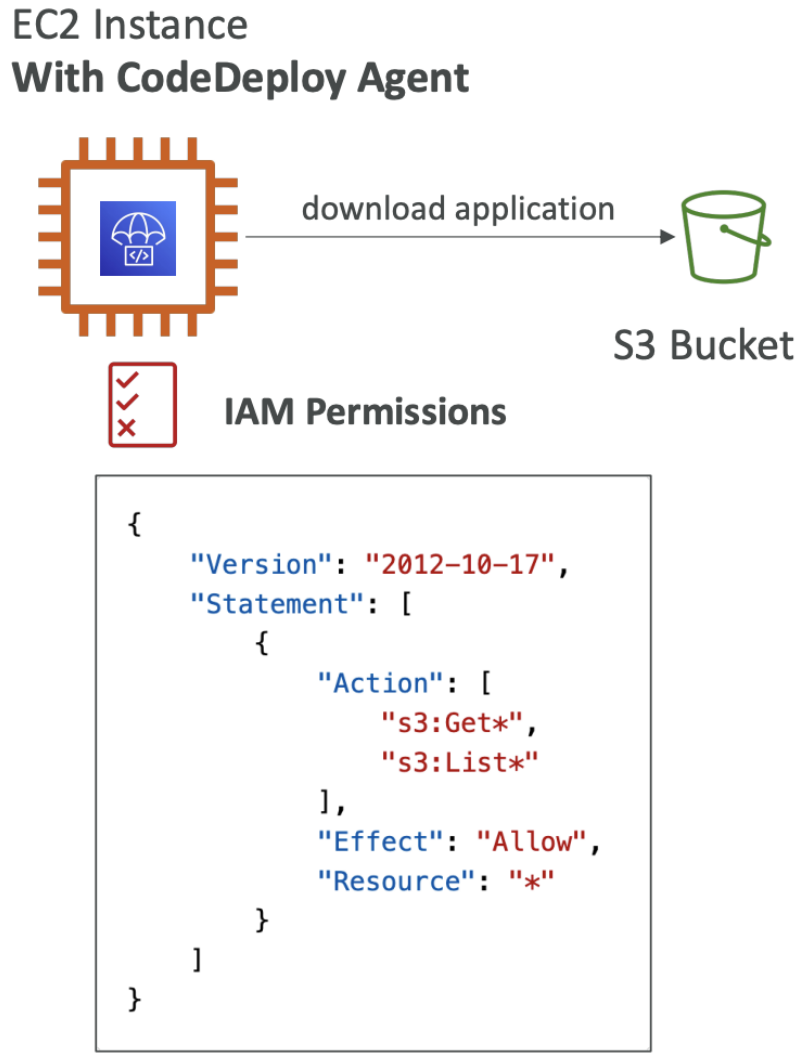

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – Lambda Platform

- **CodeDeploy** can help you automate traffic shift for Lambda aliases
- Feature is integrated within the SAM framework
- **Linear**: grow traffic every N minutes until 100%
  - LambdaLinear10PercentEvery3Minutes
  - LambdaLinear10PercentEvery10Minutes
- **Canary**: try X percent then 100%
  - LambdaCanary10Percent5Minutes
  - LambdaCanary10Percent30Minutes
- **AllAtOnce**: immediate

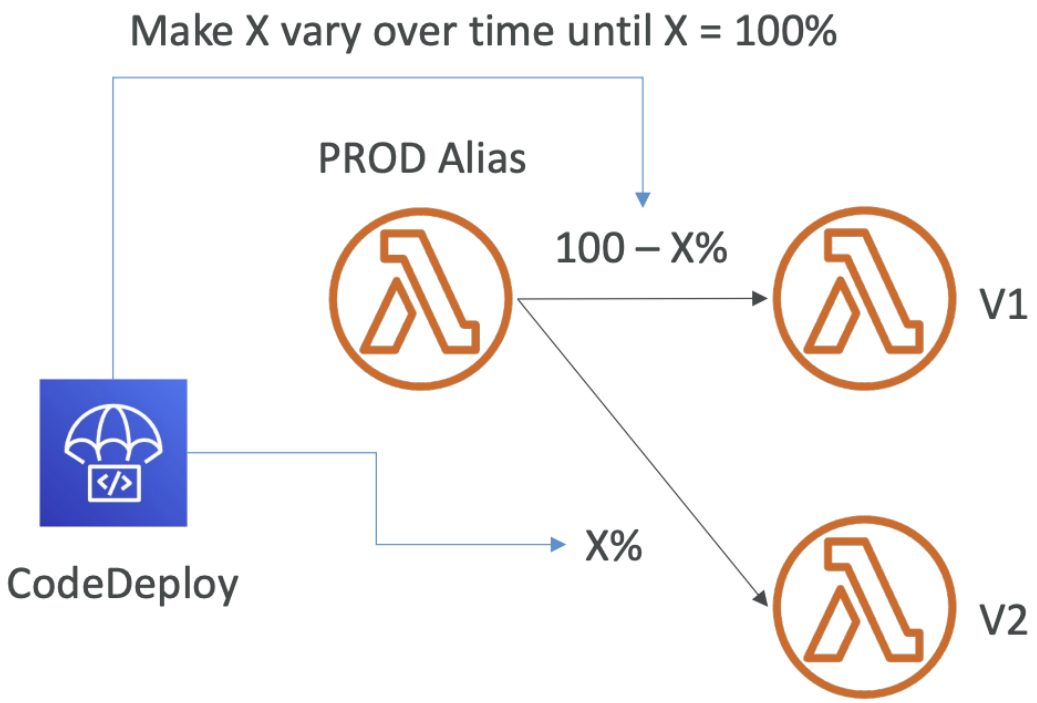

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – ECS Platform

- **CodeDeploy** can help you automate the deployment of a new ECS Task Definition
- Only Blue/Green Deployments
- **Linear**: grow traffic every N minutes until 100%
  - ECSLinear10PercentEvery3Minutes
  - ECSLinear10PercentEvery10Minutes
- **Canary**: try X percent then 100%
  - ECSCanary10Percent5Minutes
  - ECSCanary10Percent30Minutes
- **AllAtOnce**: immediate

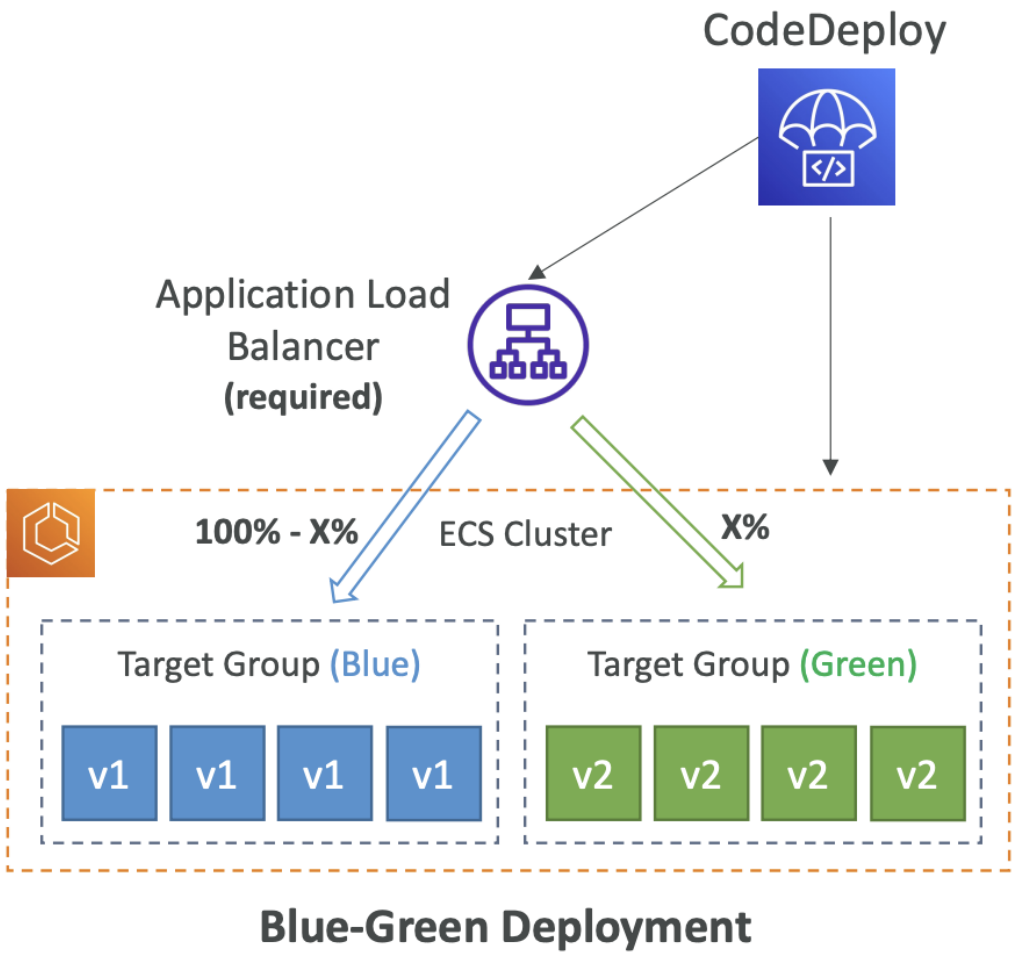

Blue-Green Deployment

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – Deployment to EC2

- Define how to deploy the application using appspec.yml + Deployment Strategy
- Will do In -place update to your fleet of EC2 instances
- Can use hooks to verify the deployment after each deployment phase

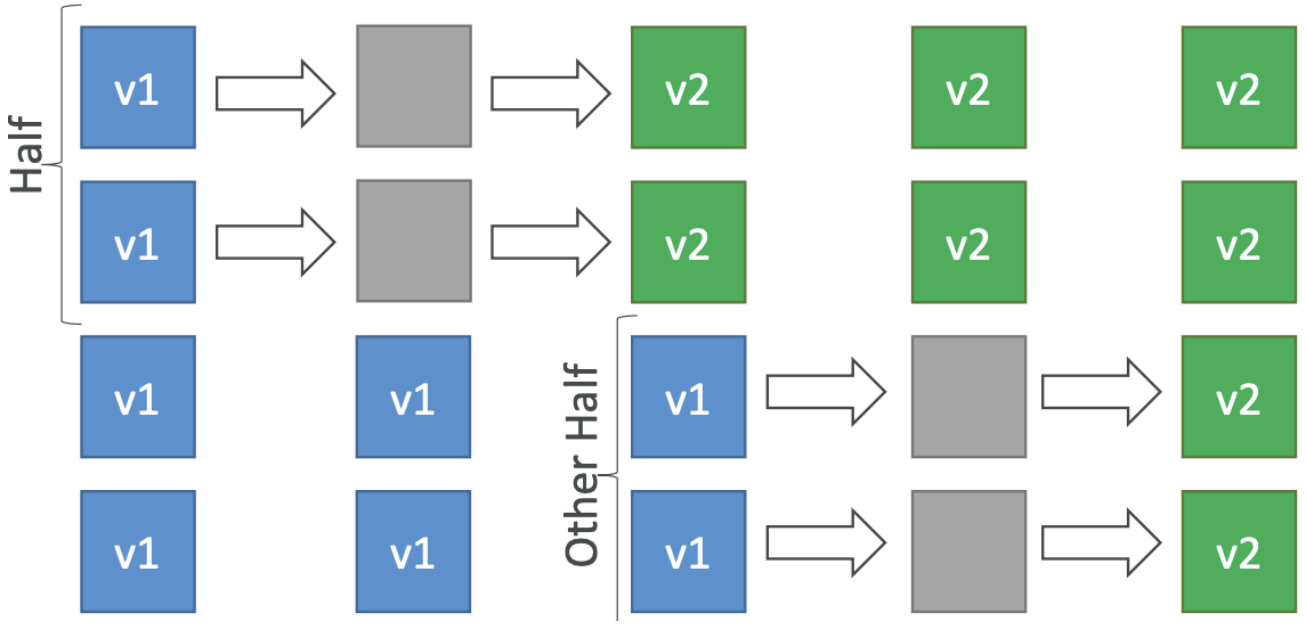

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – Deploy to an ASG

- In-place Deployment
  - Updates existing EC2 instances
  - Newly created EC2 instances by an ASG will also get automated deployments
- Blue/Green Deployment
  - A new Auto-Scaling Group is created (settings are copied)
  - Choose how long to keep the old EC2 instances (old ASG)
  - Must be using an ELB

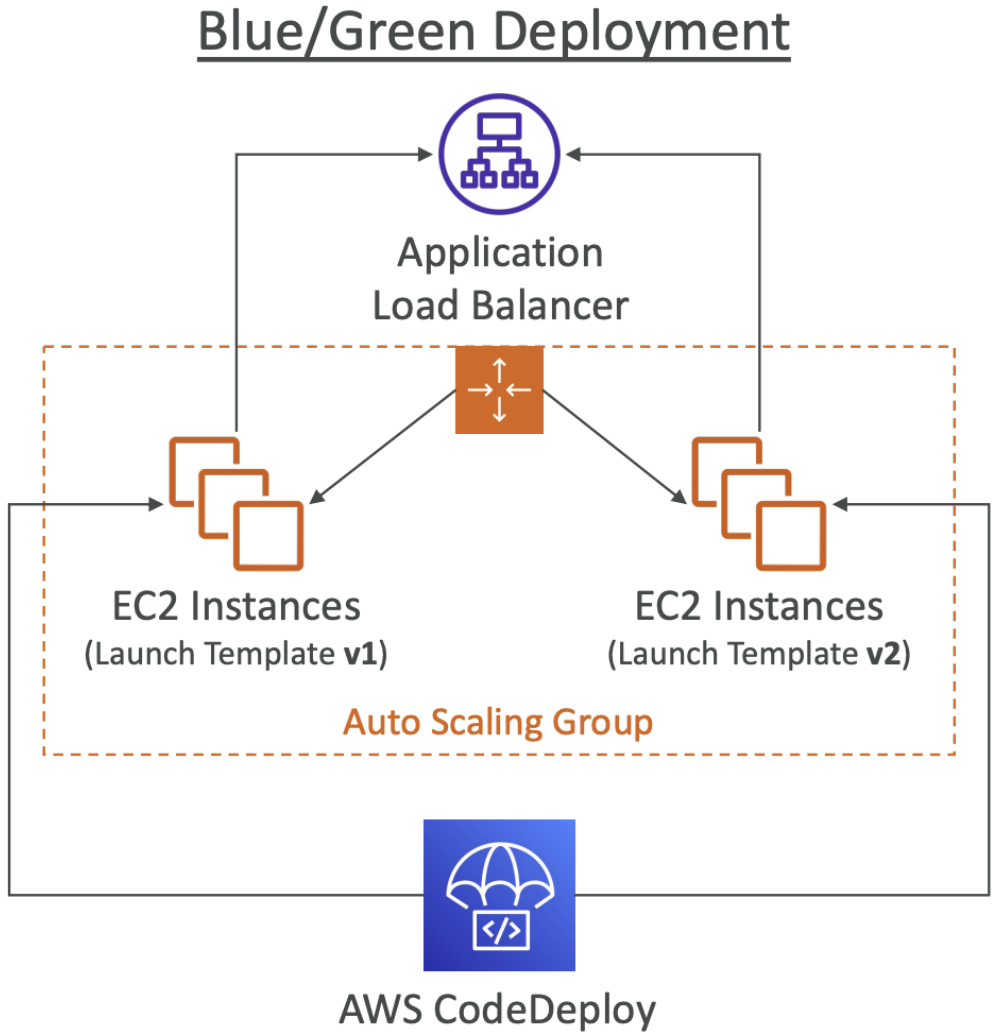

© Copyright by Cloud Mentor Pro | Confidential

CodeDeploy – Redeploy & Rollbacks

- Rollback = redeploy a previously deployed revision of your application
- Deployments can be rolled back:
  - Automatically – rollback when a deployment fails or rollback when a CloudWatch Alarm thresholds are met
  - **Manually**
- Disable Rollbacks — do not perform rollbacks for this deployment

- If a roll back happens, CodeDeploy redeploys the last known good revision as a new deployment (not a restored version)

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeArtifact

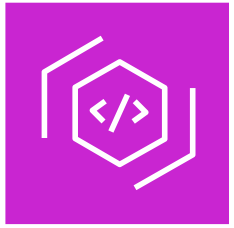

- Software packages depend on each other to be built (also called code dependencies), and new ones are created
- Storing and retrieving these dependencies is called **artifact management**
- Traditionally you need to setup your own artifact management system
- **CodeArtifact** is a secure, scalable, and cost-effective **artifact management** for software development
- Works with common dependency management tools such as Maven, Gradle, npm, yarn, twine, pip, and NuGet
- **Developers and CodeBuild can then retrieve dependencies straight from CodeArtifact**

© Copyright by Cloud Mentor Pro | Confidential

AWS CodeArtifact

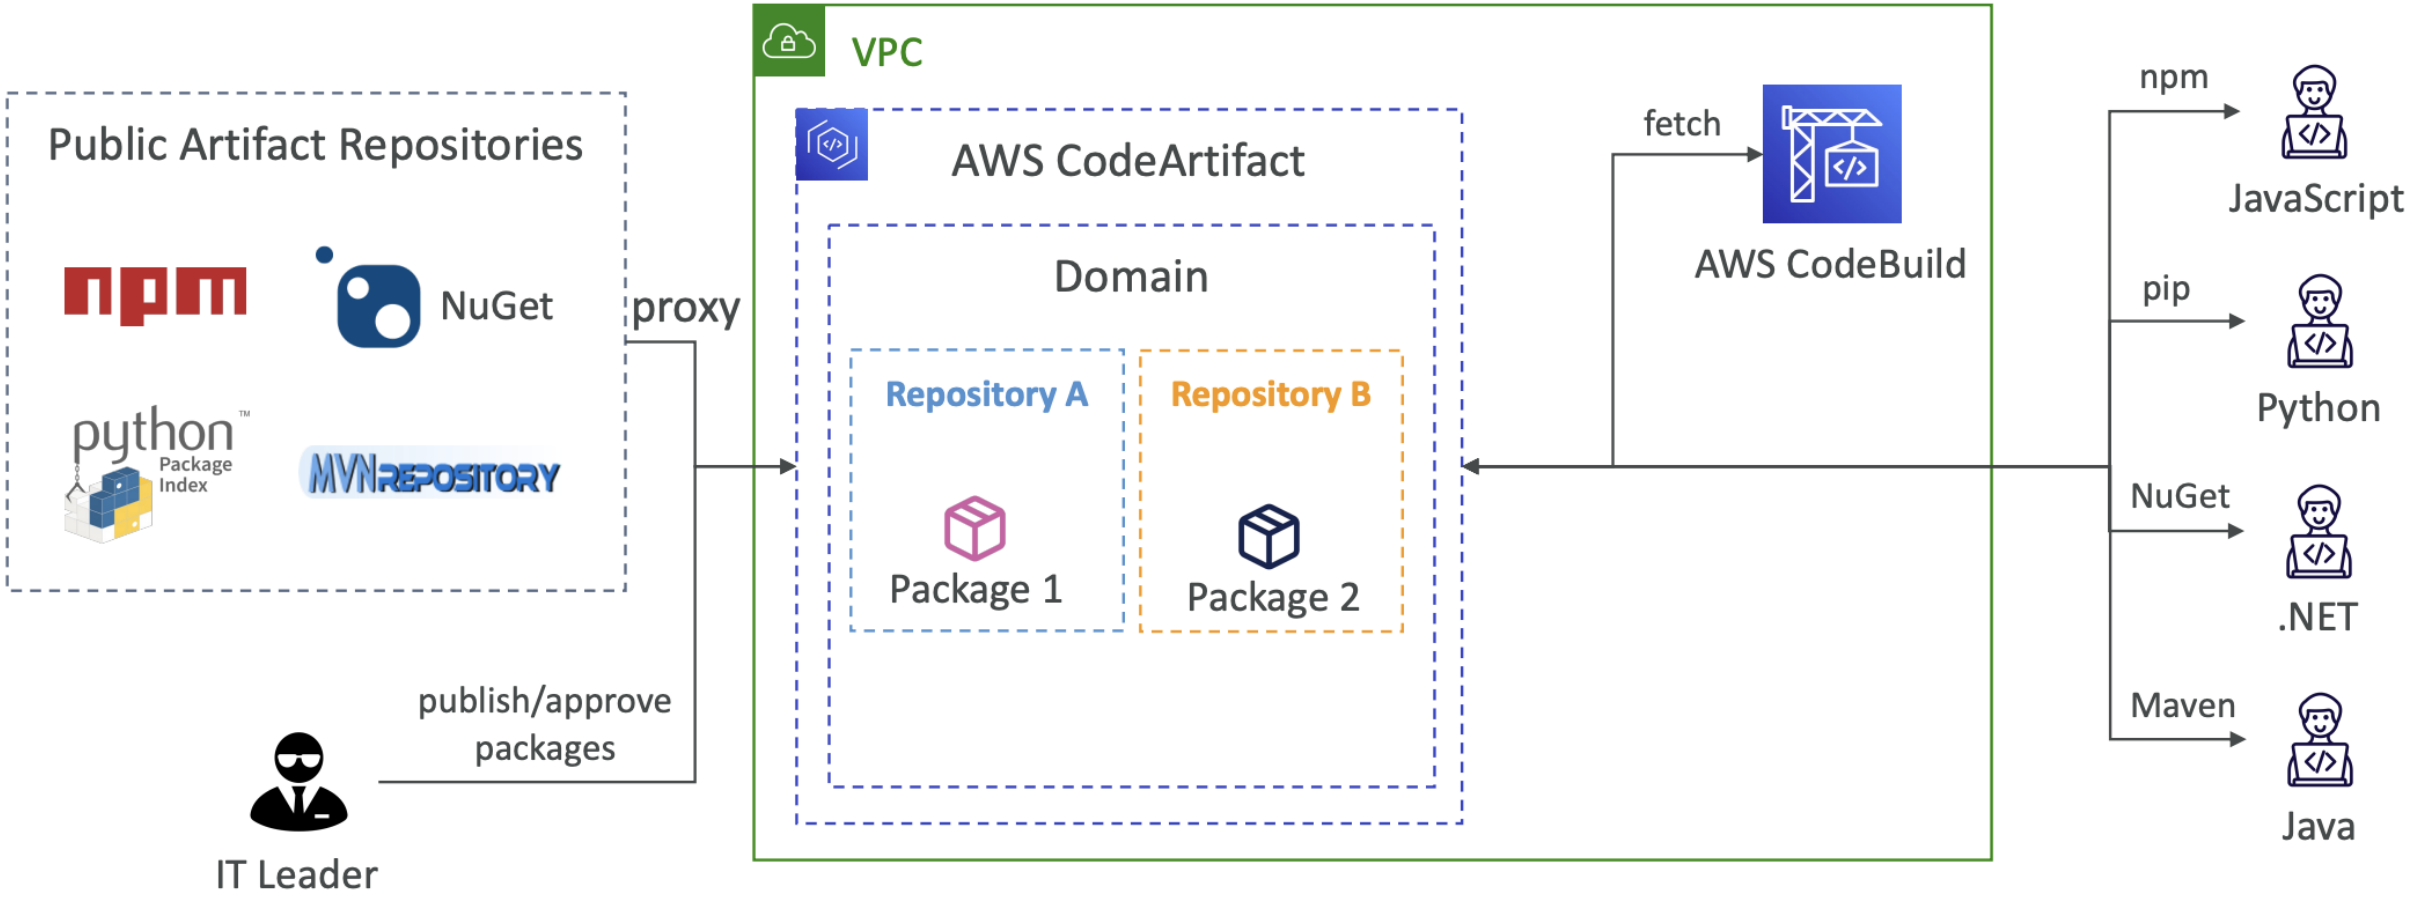

© Copyright by Cloud Mentor Pro | Confidential

CodeArtifact – EventBridge Integration

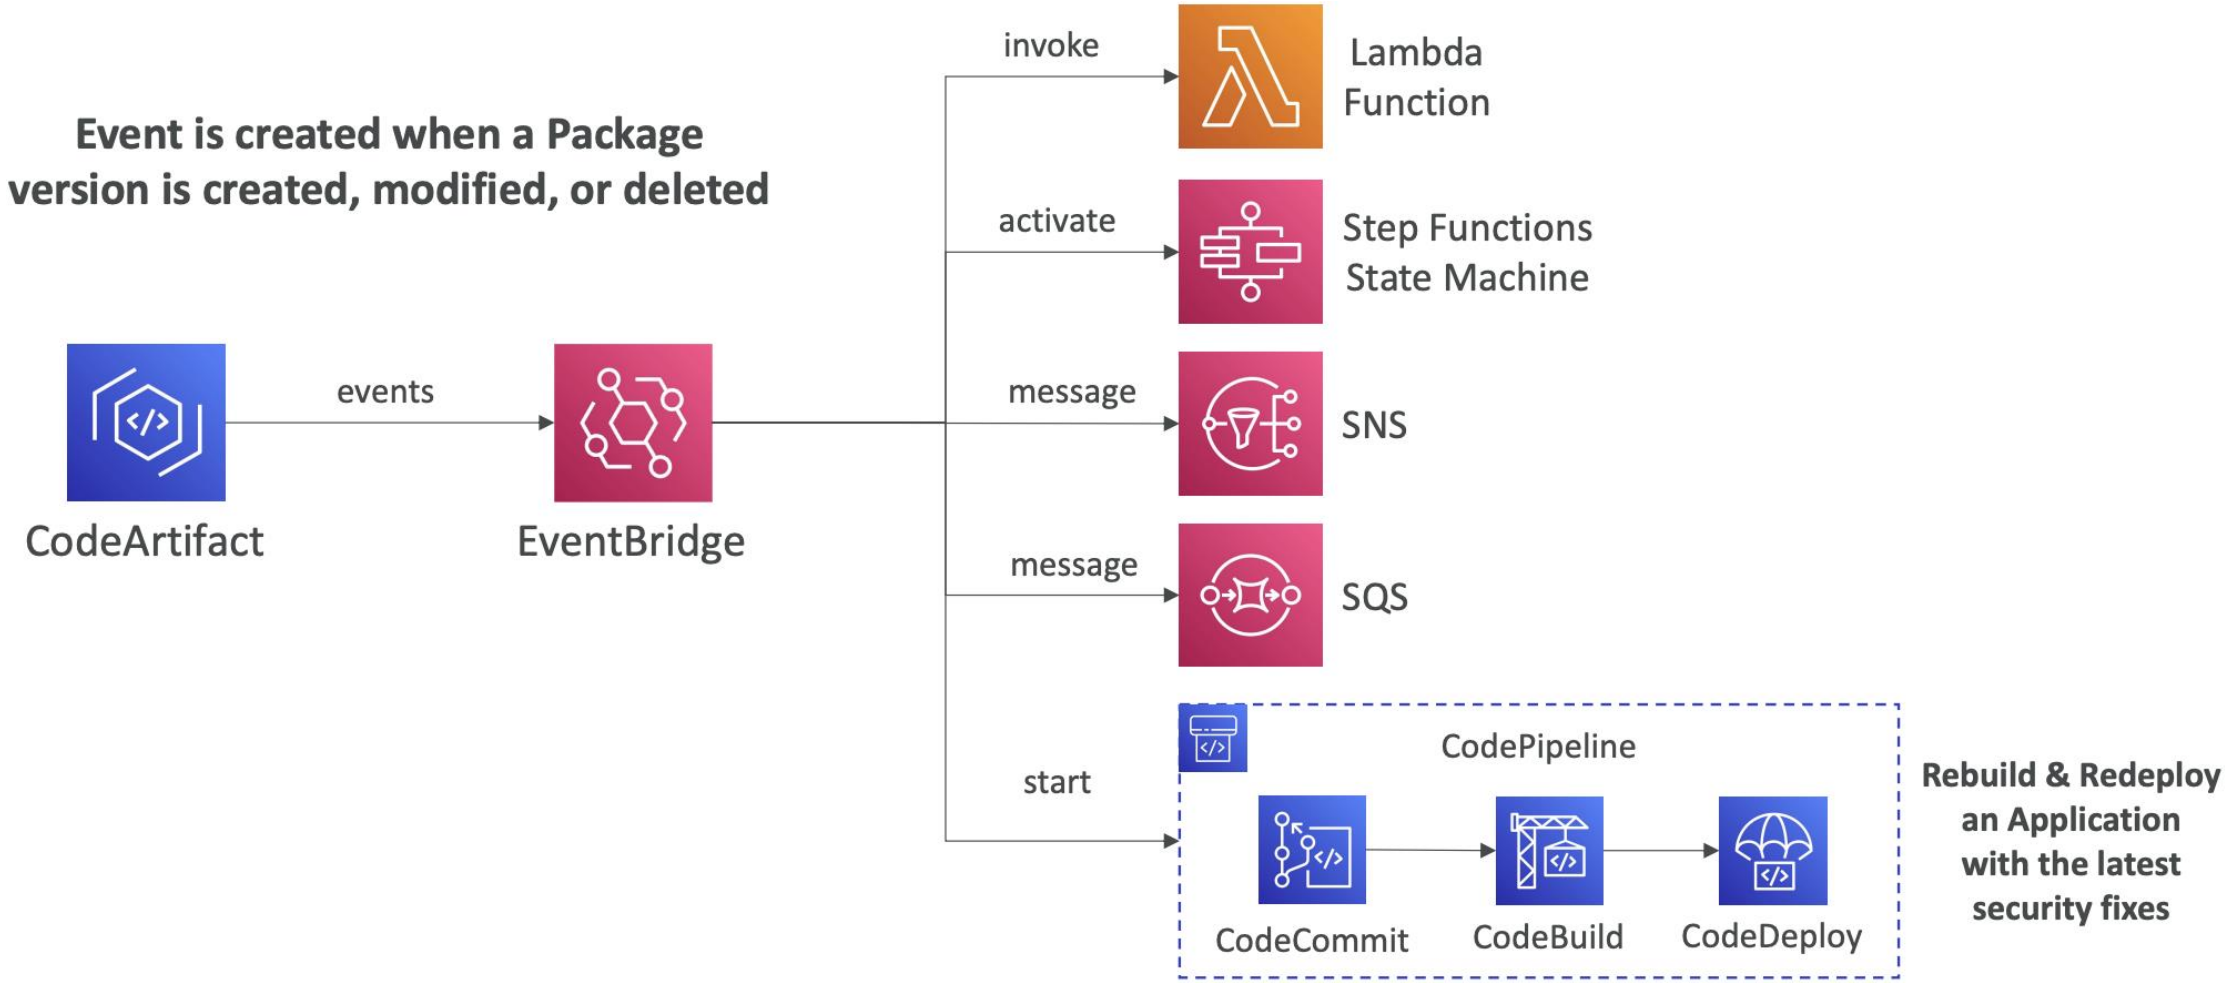

© Copyright by Cloud Mentor Pro | Confidential

CodeArtifact – Resource Policy

- Can be used to authorize another account to access CodeArtifact
- A given principal can either read all the packages in a repository or none of them

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "codeartifact:DescribePackageVersion",
                "codeartifact:DescribeRepository",
                "codeartifact:GetPackageVersionReadme",
                "codeartifact:GetRepositoryEndpoint",
                "codeartifact:ListPackages",
                "codeartifact:ListPackageVersions",
                "codeartifact:ListPackageVersionAssets",
                "codeartifact:ListPackageVersionDependencies",
                "codeartifact:ReadFromRepository"
            ],
            "Principal": {
                "AWS": [
                    "arn:aws:iam::123456789012:root",
                    "arn:aws:iam::222333344555:user/bob"
                ]
            },
            "Resource": "*"
        }
    ]
}
```

Repository Resource Policy

© Copyright by Cloud Mentor Pro | Confidential

Amazon CodeGuru

- An ML-powered service for automated code reviews and application performance recommendations
- Provides two functionalities
  - **CodeGuru Reviewer**: automated code reviews for static code analysis (development)
  - **CodeGuru Profiler**: visibility/recommendations about application performance during runtime (production)

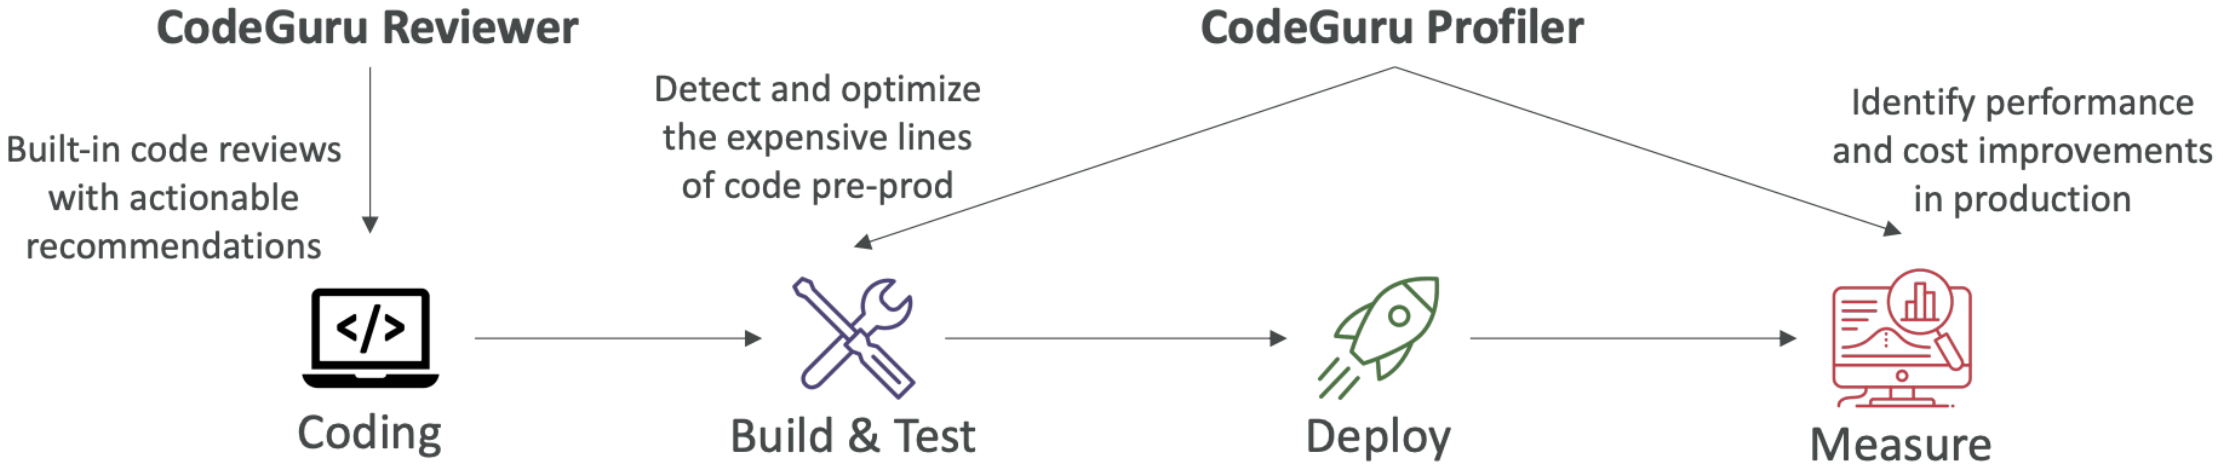

© Copyright by Cloud Mentor Pro | Confidential

Amazon CodeGuru Reviewer

- Identify critical issues, security vulnerabilities, and hard-to-find bugs
- Example: common coding best practices, resource leaks, security detection, input validation
- Uses Machine Learning and automated reasoning
- Hard-learned lessons across millions of code reviews on 1000s of open-source and Amazon repositories
- Supports Java and Python
- Integrates with GitHub, Bitbucket, and AWS CodeCommit

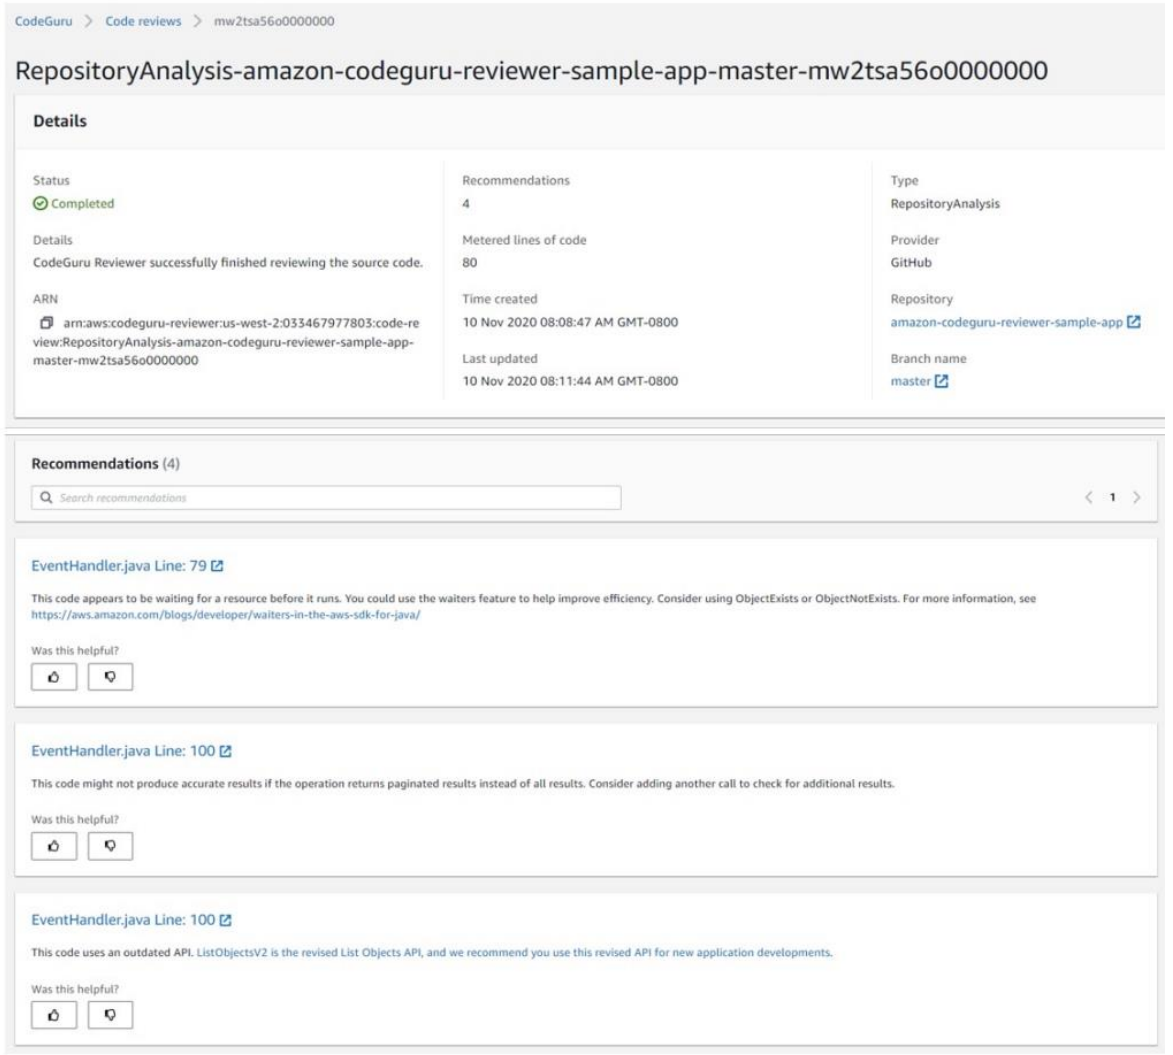

https://aws.amazon.com/codeguru/features/

© Copyright by Cloud Mentor Pro | Confidential

Amazon CodeGuru Profiler

- Helps understand the runtime behavior of your application
- Example: identify if your application is consuming excessive CPU capacity on a logging routine
- Features:
  - Identify and remove code inefficiencies
  - Improve application performance (e.g., reduce CPU utilization)
  - Decrease compute costs
  - Provides heap summary (identify which objects using up memory)
  - Anomaly Detection
- Support applications running on AWS or on-premise
- Minimal overhead on application

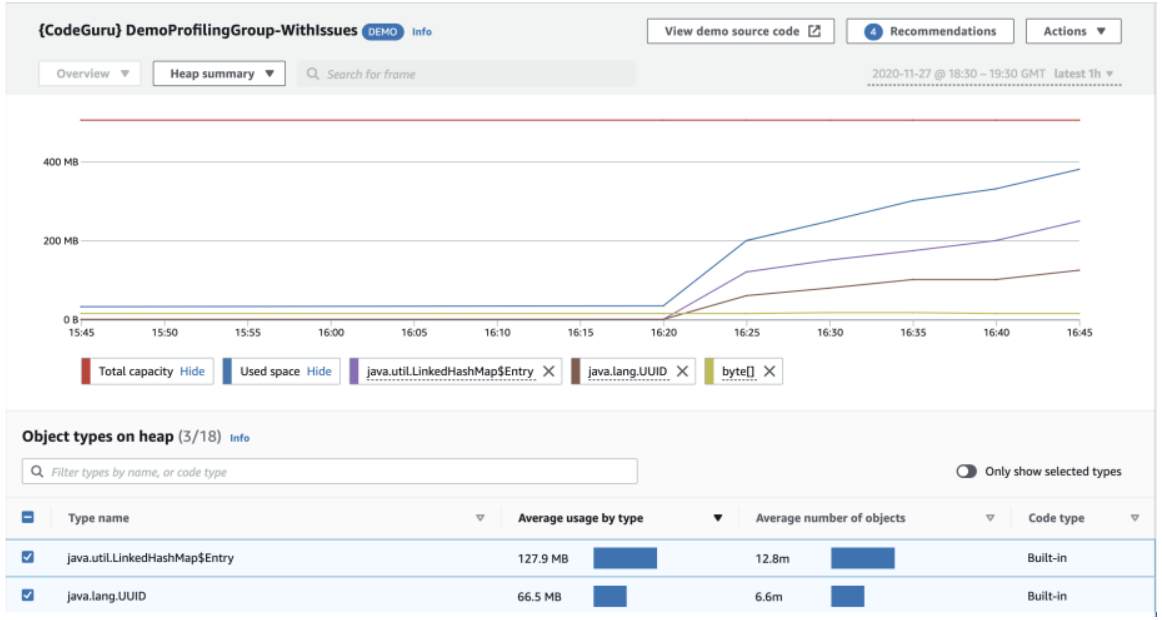

https://aws.amazon.com/codeguru/features/

© Copyright by Cloud Mentor Pro | Confidential

Amazon CodeGuru – Agent Configuration

- **MaxStackDepth** – the maximum depth of the stacks in the code that is represented in the profile
  - Example: if CodeGuru Profiler finds a method A, which calls method B, which calls method C, which calls method D, then the depth is 4
  - If the MaxStackDepth is set to 2, then the profiler evaluates A and B
- **MemoryUsageLimitPercent** – the memory percentage used by the profiler
- **MinimumTimeForReportingInMilliseconds** – the minimum time between sending reports (milliseconds)
- **ReportingIntervalInMilliseconds** – the reporting interval used to report profiles (milliseconds)
- **SamplingIntervalInMilliseconds** – the sampling interval that is used to profile samples (milliseconds)
  - Reduce to have a higher sampling rate

© Copyright by Cloud Mentor Pro | Confidential

Exercise

A Step by Step Guide to Create a CI/CD Pipeline with AWS Services

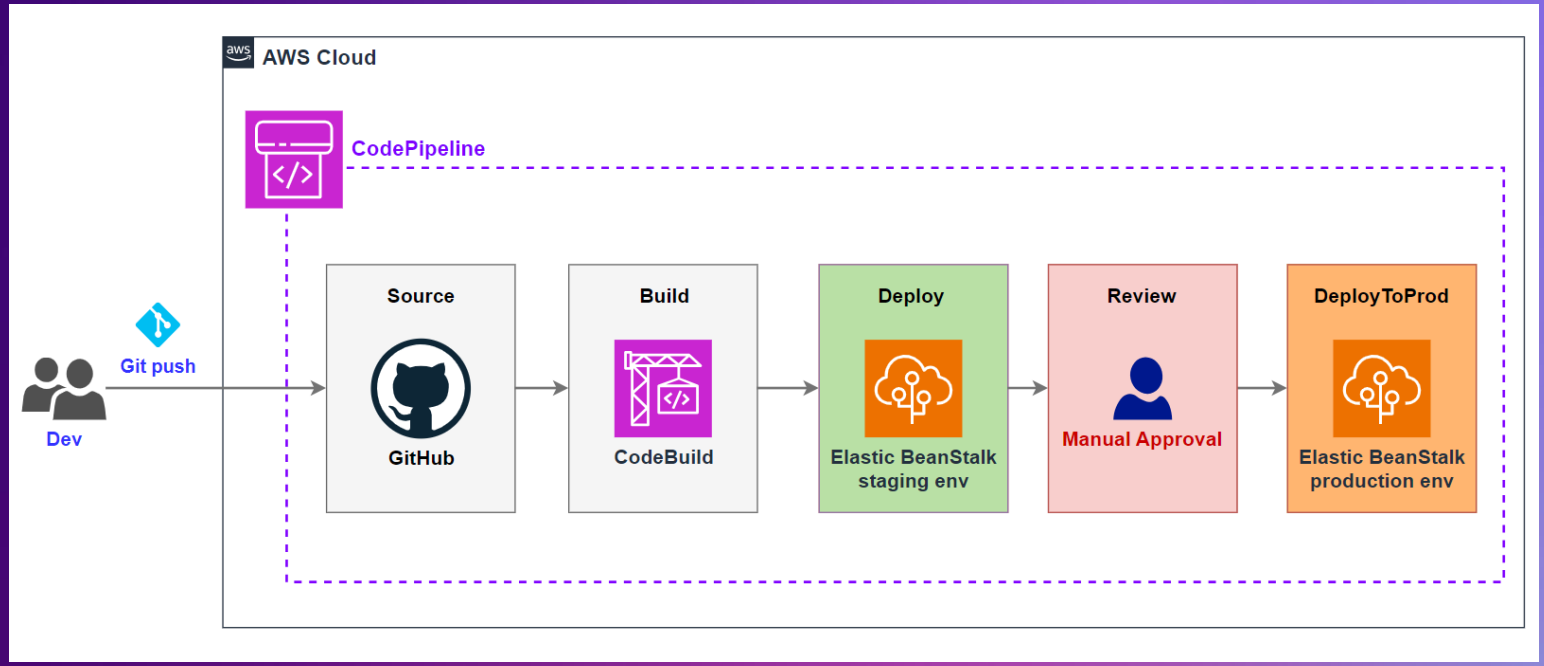

© Copyright by Cloud Mentor Pro | Confidential

Section 16

- *Advanced Identity*
- *AWS Security & Encryption: KMS, Encryption SDK, SSM Parameter Store, IAM & STS*
- *AWS Other Services*

© Copyright by Cloud Mentor Pro | Confidential

Advanced Identity in AWS

© Copyright by Cloud Mentor Pro | Confidential

AWS STS – Security Token Service

- Allows to grant limited and temporary access to AWS resources (up to 1 hour).
- **AssumeRole**: Assume roles within your account or cross account
- **AssumeRoleWithSAML**: return credentials for users logged with SAML
- **AssumeRoleWithWebIdentity**
  - return creds for users logged with an IdP (Facebook Login, Google Login, OIDC compatible…)
  - AWS recommends against using this, and using Cognito Identity Pools instead
- **GetSessionToken**: for MFA, from a user or AWS account root user
- **GetFederationToken**: obtain temporary creds for a federated user
- **GetCallerIdentity**: return details about the IAM user or role used in the API call
- **DecodeAuthorizationMessage**: decode error message when an AWS API is denied

© Copyright by Cloud Mentor Pro | Confidential

Using STS to Assume a Role

- Define an IAM Role within your account or cross-account
- Define which principals can access this IAM Role
- Use AWS STS (Security Token Service) to retrieve credentials and impersonate the IAM Role you have access to (**AssumeRole API**)
- Temporary credentials can be valid between 15 minutes to 1 hour

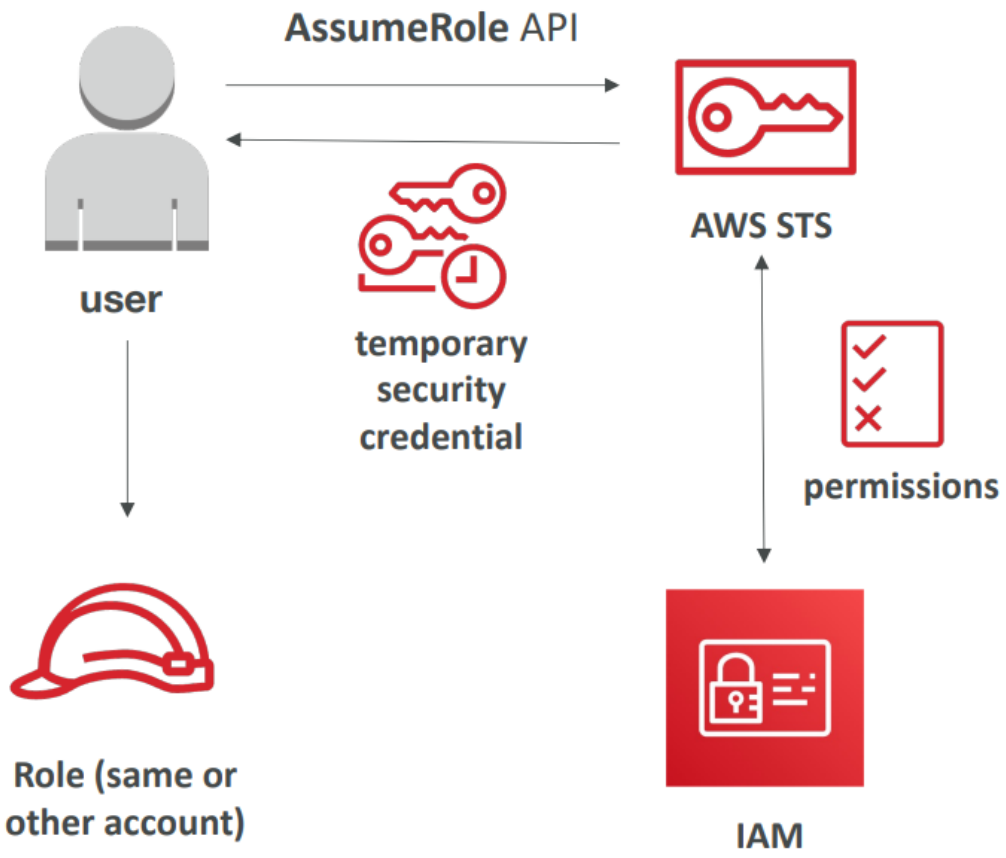

© Copyright by Cloud Mentor Pro | Confidential

Cross account access with STS

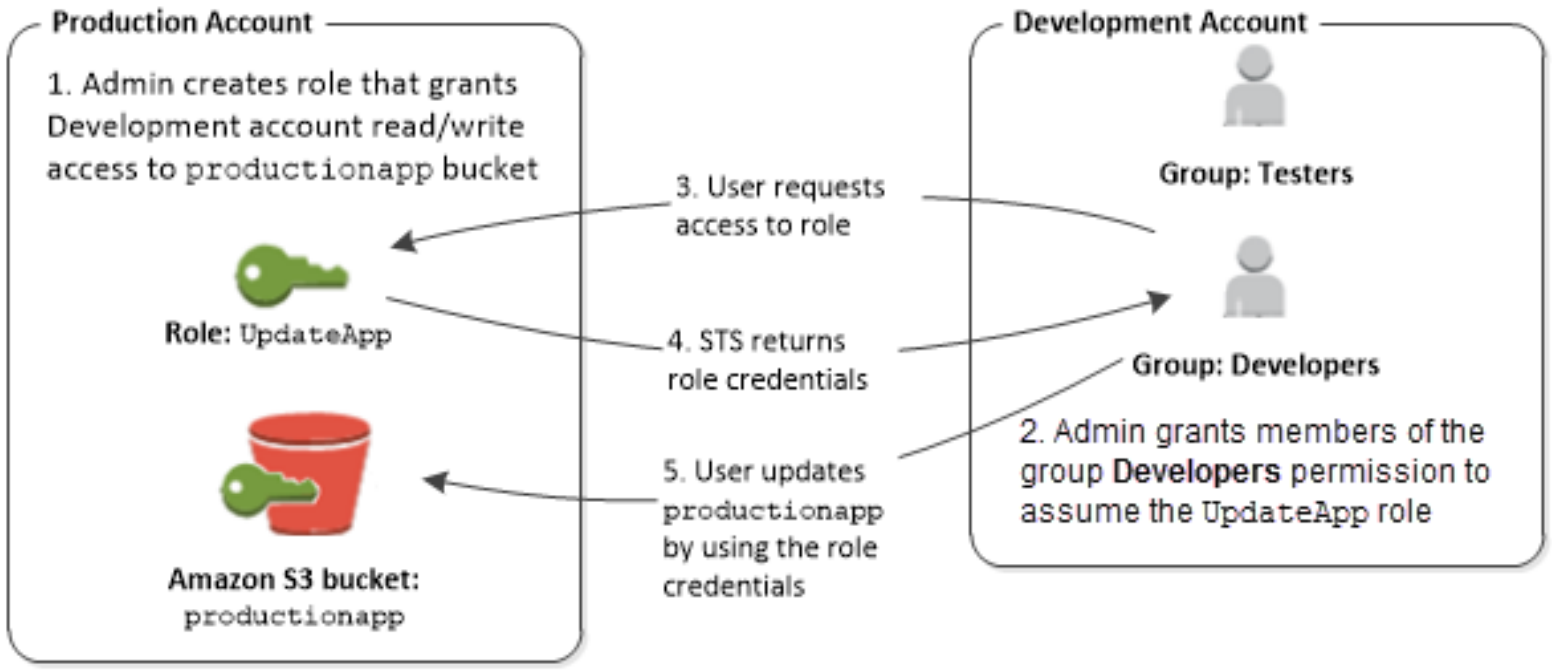

https://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_common-scenarios_aws-accounts.html

© Copyright by Cloud Mentor Pro | Confidential

STS with MFA

- Use **GetSessionToken** from STS
- Appropriate IAM policy using IAM Conditions
- **aws:MultiFactorAuthPresent:true**
- Reminder, GetSessionToken returns:
  - Access ID
  - Secret Key
  - Session Token
  - Expiration date

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "ec2:StopInstances",
                "ec2:TerminateInstances"
            ],
            "Resource": [
                "*"
            ],
            "Condition": {
                "Bool": {
                    "aws:MultiFactorAuthPresent": "true"
                }
            }
        }
    ]
}
```

© Copyright by Cloud Mentor Pro | Confidential

IAM Best Practices – General

- Never use Root Credentials, enable MFA for Root Account
- Grant Least Privilege
  - Each Group / User / Role should only have the minimum level of permission it needs
  - Never grant a policy with "*" access to a service
  - Monitor API calls made by a user in CloudTrail (especially Denied ones)
- Never ever ever store IAM key credentials on any machine but a personal computer or on-premise server
- On premise server best practice is to call STS to obtain temporary security credentials

© Copyright by Cloud Mentor Pro | Confidential

IAM Best Practices – IAM Roles

- EC2 machines should have their own roles
- Lambda functions should have their own roles
- ECS Tasks should have their own roles (ECS_ENABLE_TASK_IAM_ROLE=true)
- CodeBuild should have its own service role
- Create a least-privileged role for any service that requires it
- Create a role per application / lambda function (do not reuse roles)

© Copyright by Cloud Mentor Pro | Confidential

IAM Best Practices – Cross Account Access

- Define an IAM Role for another account to access
- Define which accounts can access this IAM Role
- Use AWS STS (Security Token Service) to retrieve credentials and impersonate the IAM Role you have access to (AssumeRole API)
- Temporary credentials can be valid between 15 minutes to 1 hour

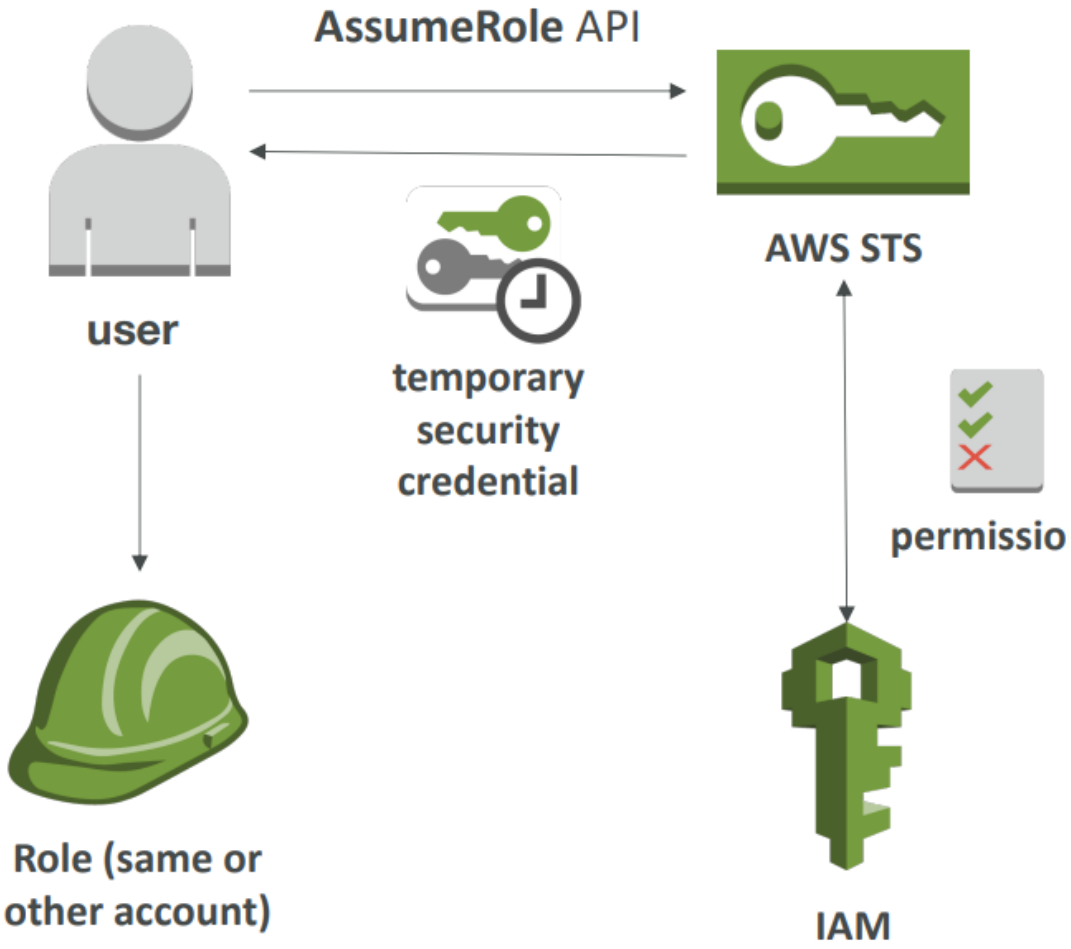

© Copyright by Cloud Mentor Pro | Confidential

Advanced IAM - Authorization Model Evaluation of Policies, simplified

1. If there's an explicit DENY, end decision and DENY
2. If there's an ALLOW, end decision with ALLOW
3. Else DENY

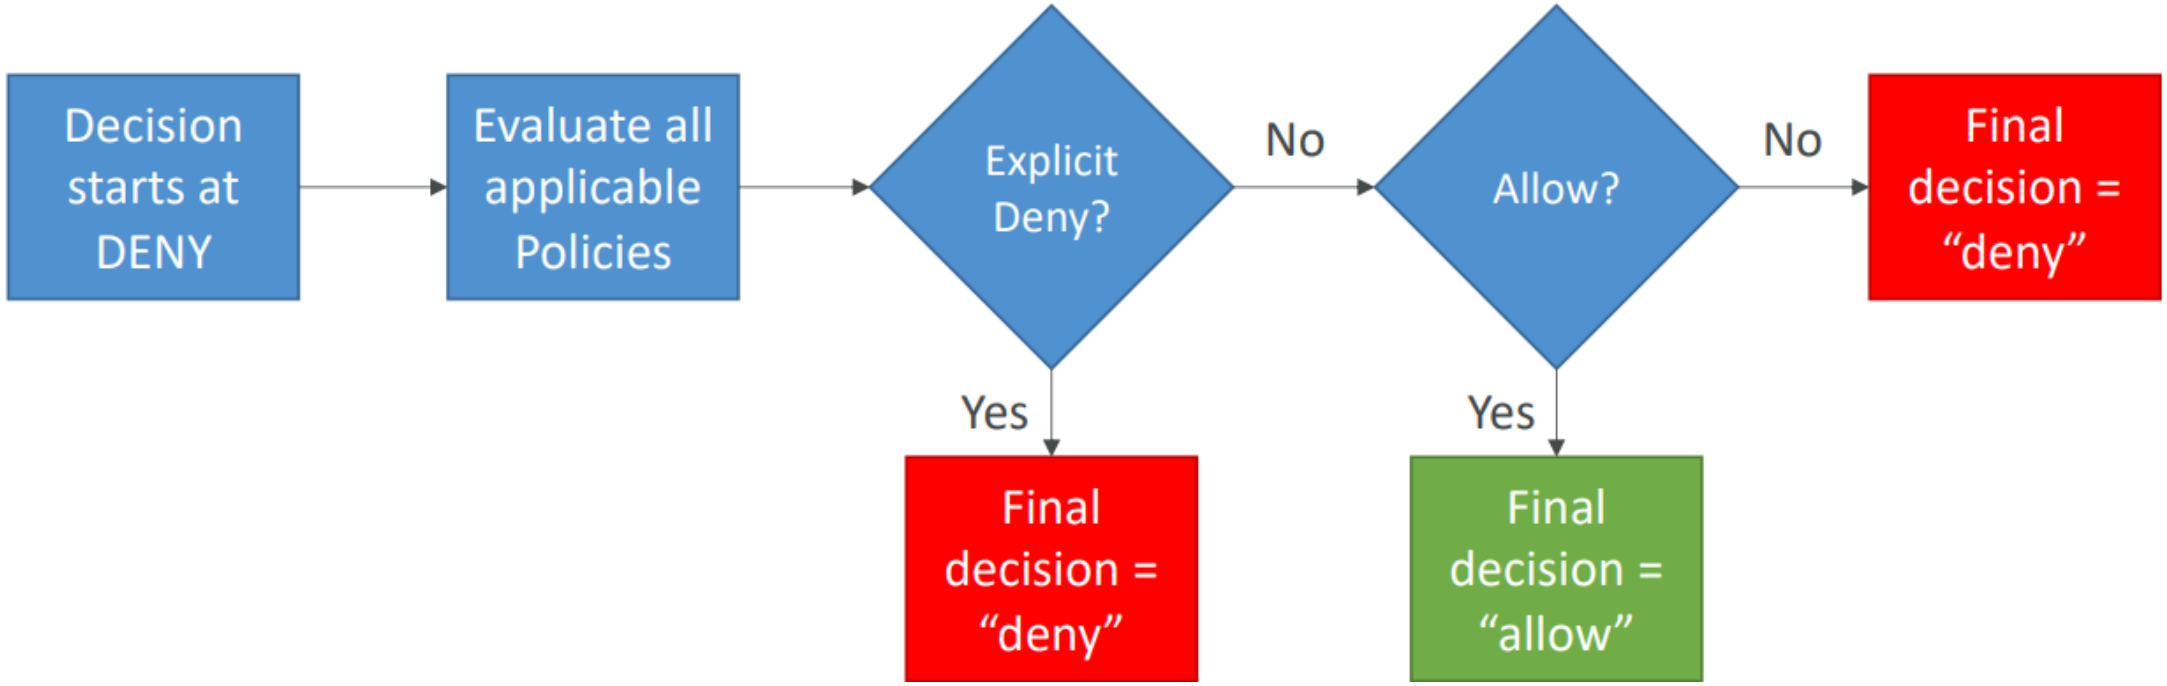

© Copyright by Cloud Mentor Pro | Confidential

IAM Policies & S3 Bucket Policies

- IAM Policies are attached to users, roles, groups
- S3 Bucket Policies are attached to buckets
- When evaluating if an IAM Principal can perform an operation X on a bucket, the **union** of its assigned IAM Policies and S3 Bucket Policies will be evaluated.

© Copyright by Cloud Mentor Pro | Confidential

Example 1

- IAM Role attached to EC2 instance, authorizes RW to “my_bucket”
- No S3 Bucket Policy attached

=> EC2 instance can read and write to “my_bucket”

© Copyright by Cloud Mentor Pro | Confidential

Example 2

- IAM Role attached to EC2 instance, authorizes RW to “my_bucket”
- S3 Bucket Policy attached, explicit deny to the IAM Role

=> EC2 instance **cannot** read and write to “my_bucket”

© Copyright by Cloud Mentor Pro | Confidential

Example 3

- IAM Role attached to EC2 instance, no S3 bucket permissions
- S3 Bucket Policy attached, explicit RW allow to the IAM Role

=> EC2 instance **can** read and write to “my_bucket”

© Copyright by Cloud Mentor Pro | Confidential

Example 4

- IAM Role attached to EC2 instance, explicit deny S3 bucket permissions
- S3 Bucket Policy attached, explicit RW allow to the IAM Role

=> EC2 instance **cannot** read and write to “my_bucket”

© Copyright by Cloud Mentor Pro | Confidential

Dynamic Policies with IAM

- How do you assign each user a /home/<user> folder in an S3 bucket?
- Option 1:
  - Create an IAM policy allowing georges to have access to /home/georges
  - Create an IAM policy allowing sarah to have access to /home/sarah
  - Create an IAM policy allowing matt to have access to /home/matt
  - … One policy per user!
  - This doesn't scale
- Option 2:
  - Create one dynamic policy with IAM
  - Leverage the special policy variable ${aws:username}

© Copyright by Cloud Mentor Pro | Confidential

Dynamic Policy example

```
{
    "Sid": "AllowAllS3ActionsInUserFolder",
    "Action":["s3:*"],
    "Effect":"Allow",
    "Resource": ["arn:aws:s3:::my-company/home/${aws:username}/*"]
}
```

© Copyright by Cloud Mentor Pro | Confidential

Inline vs Managed Policies

- AWS Managed Policy
  - Maintained by AWS
  - Good for power users and administrators
  - Updated in case of new services / new APIs
- Customer Managed Policy
  - Best Practice, re-usable, can be applied to many principals
  - Version Controlled + rollback, central change management
- Inline
  - Strict one-to-one relationship between policy and principal
  - Policy is deleted if you delete the IAM principal

© Copyright by Cloud Mentor Pro | Confidential

What is Microsoft Active Directory (AD)?

- Found on any Windows Server with AD Domain Services
- Database of objects: User Accounts, Computers, Printers, File Shares, Security Groups
- Centralized security management, create account, assign permissions
- Objects are organized in **trees**
- A group of trees is a **forest**

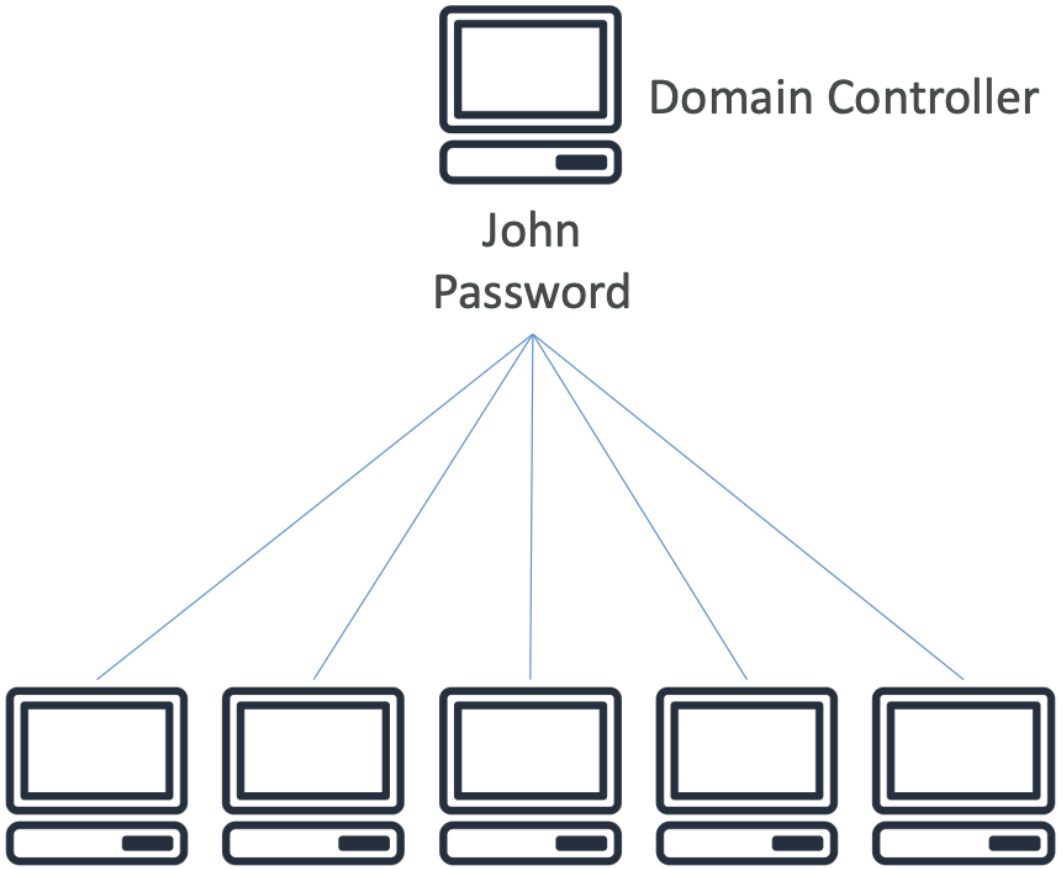

© Copyright by Cloud Mentor Pro | Confidential

AWS Directory Services

- AWS Managed Microsoft AD
  - Create your own AD in AWS, manage users locally, supports MFA
  - Establish “trust” connections with your on- premises AD
- AD Connector
  - Directory Gateway (proxy) to redirect to on- premises AD, supports MFA
  - Users are managed on the on-premises AD
- Simple AD
  - AD-compatible managed directory on AWS
  - Cannot be joined with on-premises AD

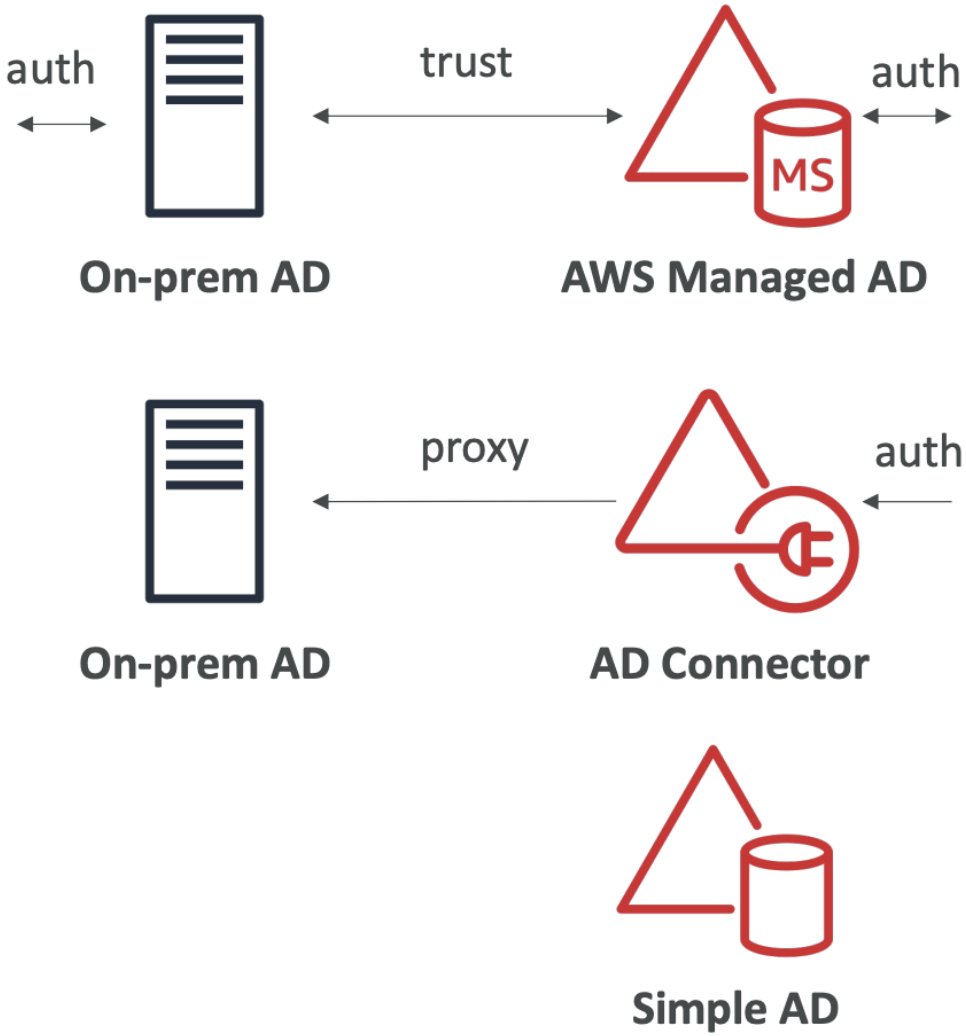

© Copyright by Cloud Mentor Pro | Confidential

AWS Security & Encryption

KMS, Encryption SDK, SSM Parameter Store

© Copyright by Cloud Mentor Pro | Confidential

Why encryption?
Encryption in flight (TLS / SSL)

- Data is encrypted before sending and decrypted after receiving
- TLS certificates help with encryption (HTTPS)
- Encryption in flight ensures no MITM (man in the middle attack) can happen

© Copyright by Cloud Mentor Pro | Confidential

Why encryption?
Server-side encryption at rest

- Data is encrypted after being received by the server
- Data is decrypted before being sent
- It is stored in an encrypted form thanks to a key (usually a data key)
- The encryption / decryption keys must be managed somewhere, and the server must have access to it

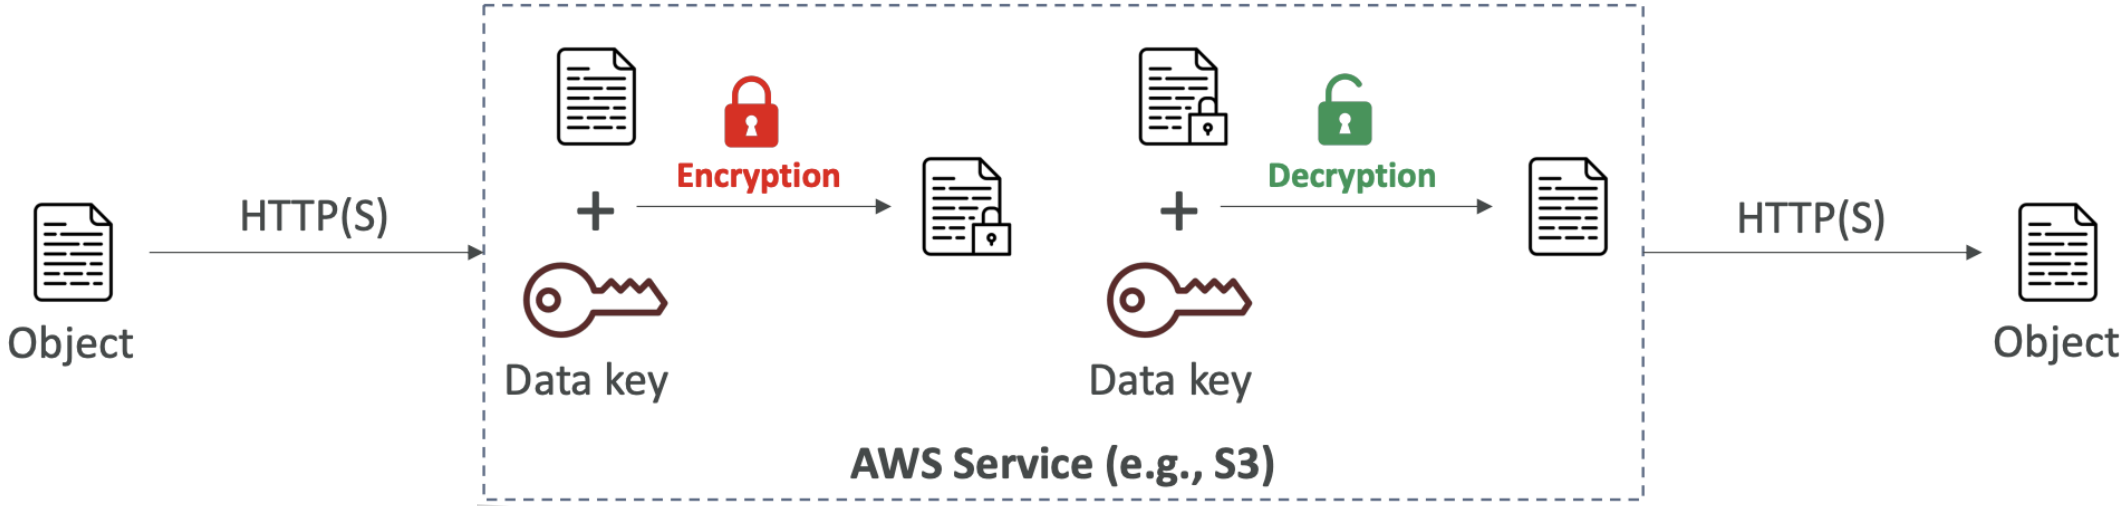

© Copyright by Cloud Mentor Pro | Confidential

Why encryption?
Client-side encryption

- Data is encrypted by the client and never decrypted by the server
- Data will be decrypted by a receiving client
- The server should not be able to decrypt the data
- Could leverage Envelope Encryption

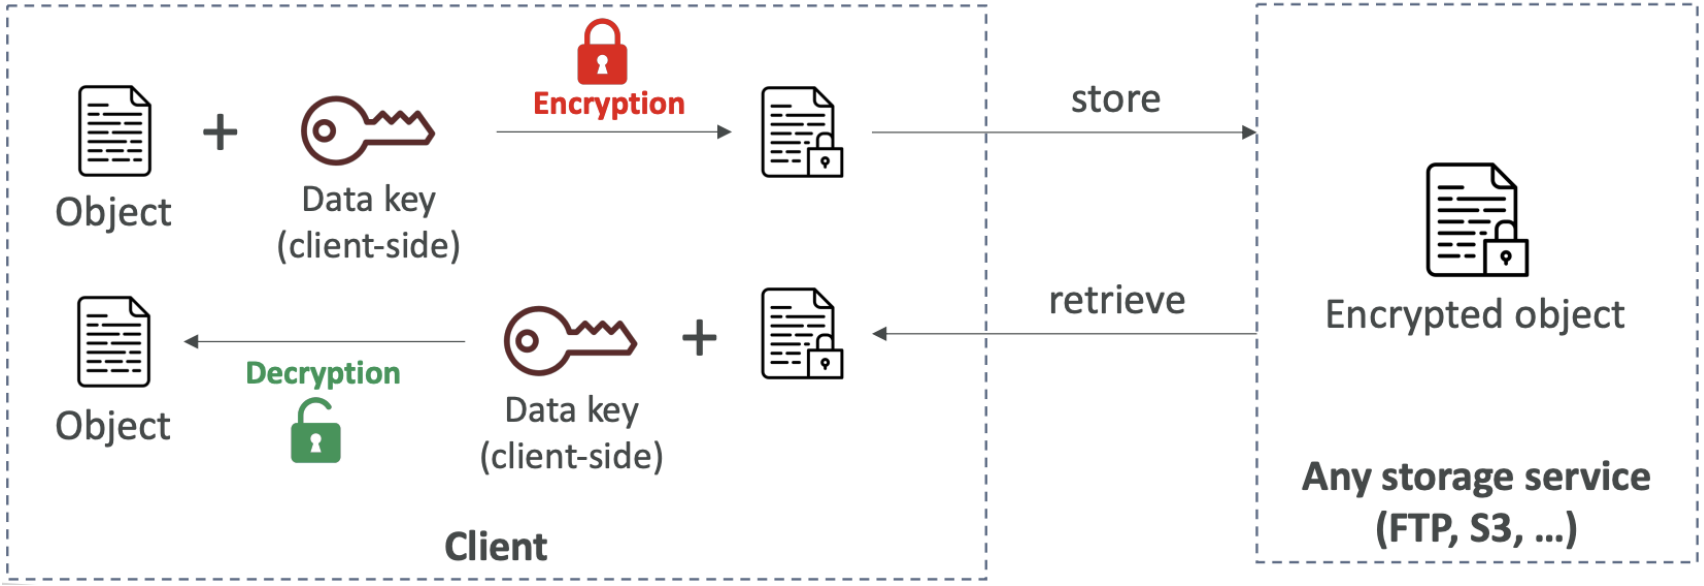

© Copyright by Cloud Mentor Pro | Confidential

AWS KMS (Key Management Service)

- Anytime you hear “encryption” for an AWS service, it’s most likely KMS
- AWS manages encryption keys for us
- Fully integrated with IAM for authorization
- Easy way to control access to your data
- Able to audit KMS Key usage using CloudTrail
- Seamlessly integrated into most AWS services (EBS, S3, RDS, SSM…)
- **Never ever store your secrets in plaintext, especially in your code!**
  - KMS Key Encryption also available through API calls (SDK, CLI)
  - Encrypted secrets can be stored in the code / environment variables

© Copyright by Cloud Mentor Pro | Confidential

KMS Keys Types

- **KMS Keys is the new name of KMS Customer *Master* Key**
- **Symmetric (AES-256 keys)**
  - Single encryption key that is used to Encrypt and Decrypt
  - AWS services that are integrated with KMS use Symmetric CMKs
  - You never get access to the KMS Key unencrypted (must call KMS API to use)
- **Asymmetric (RSA & ECC key pairs)**
  - Public (Encrypt) and Private Key (Decrypt) pair
  - Used for Encrypt/Decrypt, or Sign/Verify operations
  - The public key is downloadable, but you can't access the Private Key unencrypted
  - Use case: encryption outside of AWS by users who can't call the KMS API

© Copyright by Cloud Mentor Pro | Confidential

AWS KMS (Key Management Service)

- Types of KMS Keys:
  - AWS Owned Keys (free): SSE-S3, SSE-SQS, SSE-DDB (default key)
  - AWS Managed Key: free (aws/service-name, example: aws/rds or aws/ebs)
  - Customer managed keys created in KMS: $1 / month
  - Customer managed keys imported (must be symmetric key): $1 / month
  - + pay for API call to KMS ($0.03 / 10000 calls)

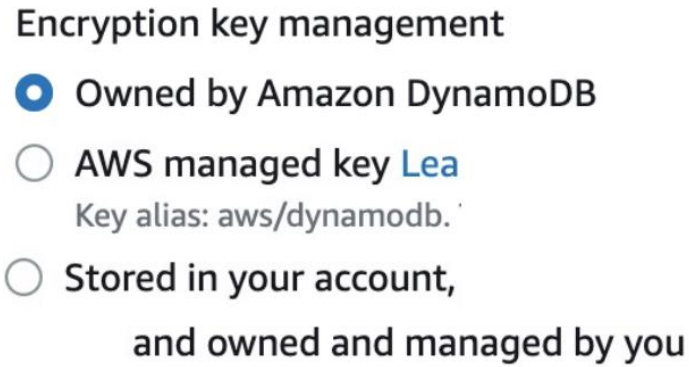

- Automatic Key rotation:
  - AWS-managed KMS Key: automatic every 1 year
  - Customer-managed KMS Key: (must be enabled) automatic every 1 year
  - Imported KMS Key: only manual rotation possible using alias

© Copyright by Cloud Mentor Pro | Confidential

Copying Snapshots across regions

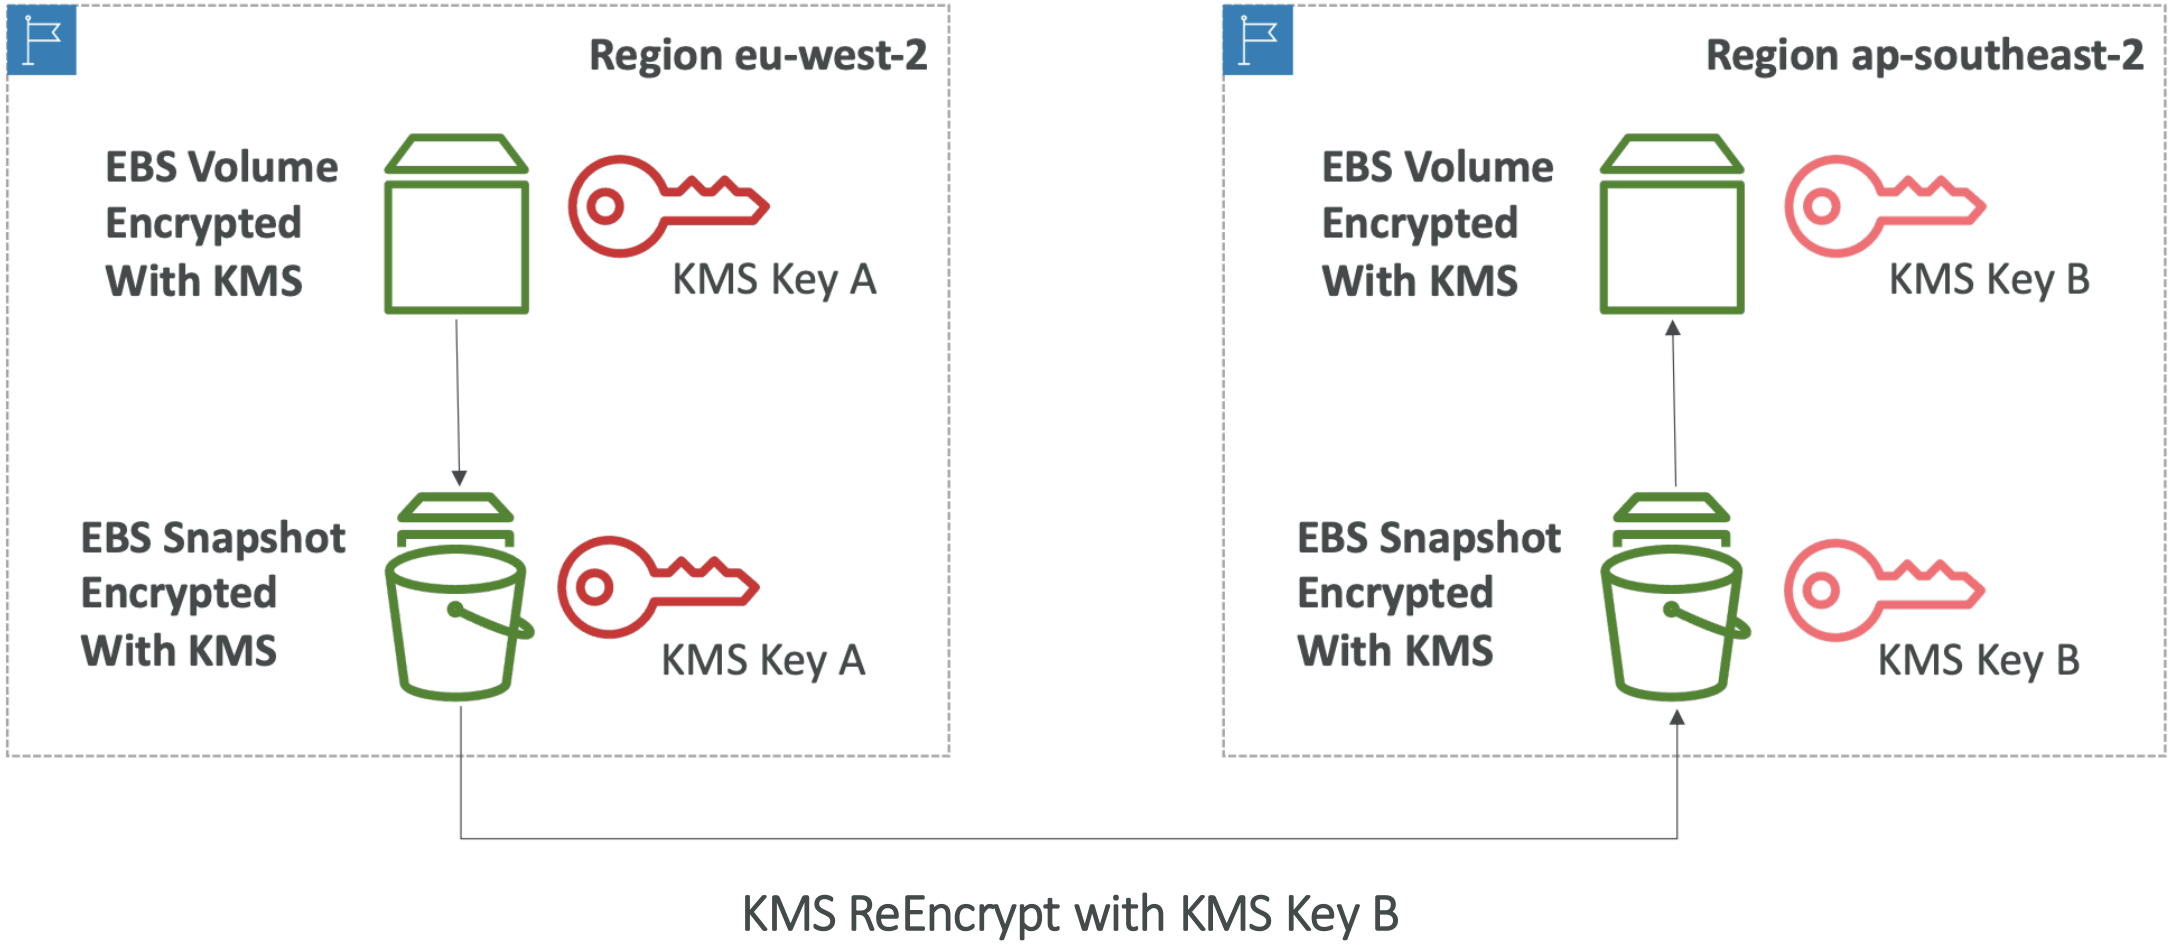

© Copyright by Cloud Mentor Pro | Confidential

KMS Key Policies

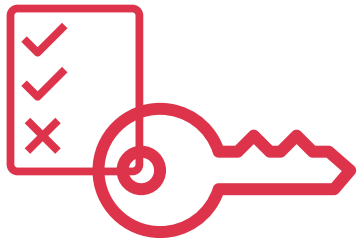

- Control access to KMS keys, “similar” to S3 bucket policies
- Difference: you cannot control access without them

- **Default KMS Key Policy:**
  - Created if you don’t provide a specific KMS Key Policy
  - Complete access to the key to the root user = entire AWS account
- **Custom KMS Key Policy:**
  - Define users, roles that can access the KMS key
  - Define who can administer the key
  - Useful for cross-account access of your KMS key

© Copyright by Cloud Mentor Pro | Confidential

Copying Snapshots across accounts

1. Create a Snapshot, encrypted with your own KMS Key (Customer Managed Key)
2. **Attach a KMS Key Policy to authorize cross-account access**
3. Share the encrypted snapshot
4. (in target) Create a copy of the Snapshot, encrypt it with a CMK in your account
5. Create a volume from the snapshot

```
{
    "Sid": "Allow use of the key with destination account",
    "Effect": "Allow",
    "Principal": {
      "AWS": "arn:aws:iam::TARGET-ACCOUNT-ID:role/ROLENAME"
    },
    "Action": [
      "kms:Decrypt",
      "kms:CreateGrant"
    ],
    "Resource": "*",
    "Condition": {
      "StringEquals": {
        "kms:ViaService": "ec2.REGION.amazonaws.com",
        "kms:CallerAccount": "TARGET-ACCOUNT-ID"
      }
    }
}
```

KMS Key Policy

© Copyright by Cloud Mentor Pro | Confidential

How does KMS work?
API – Encrypt and Decrypt

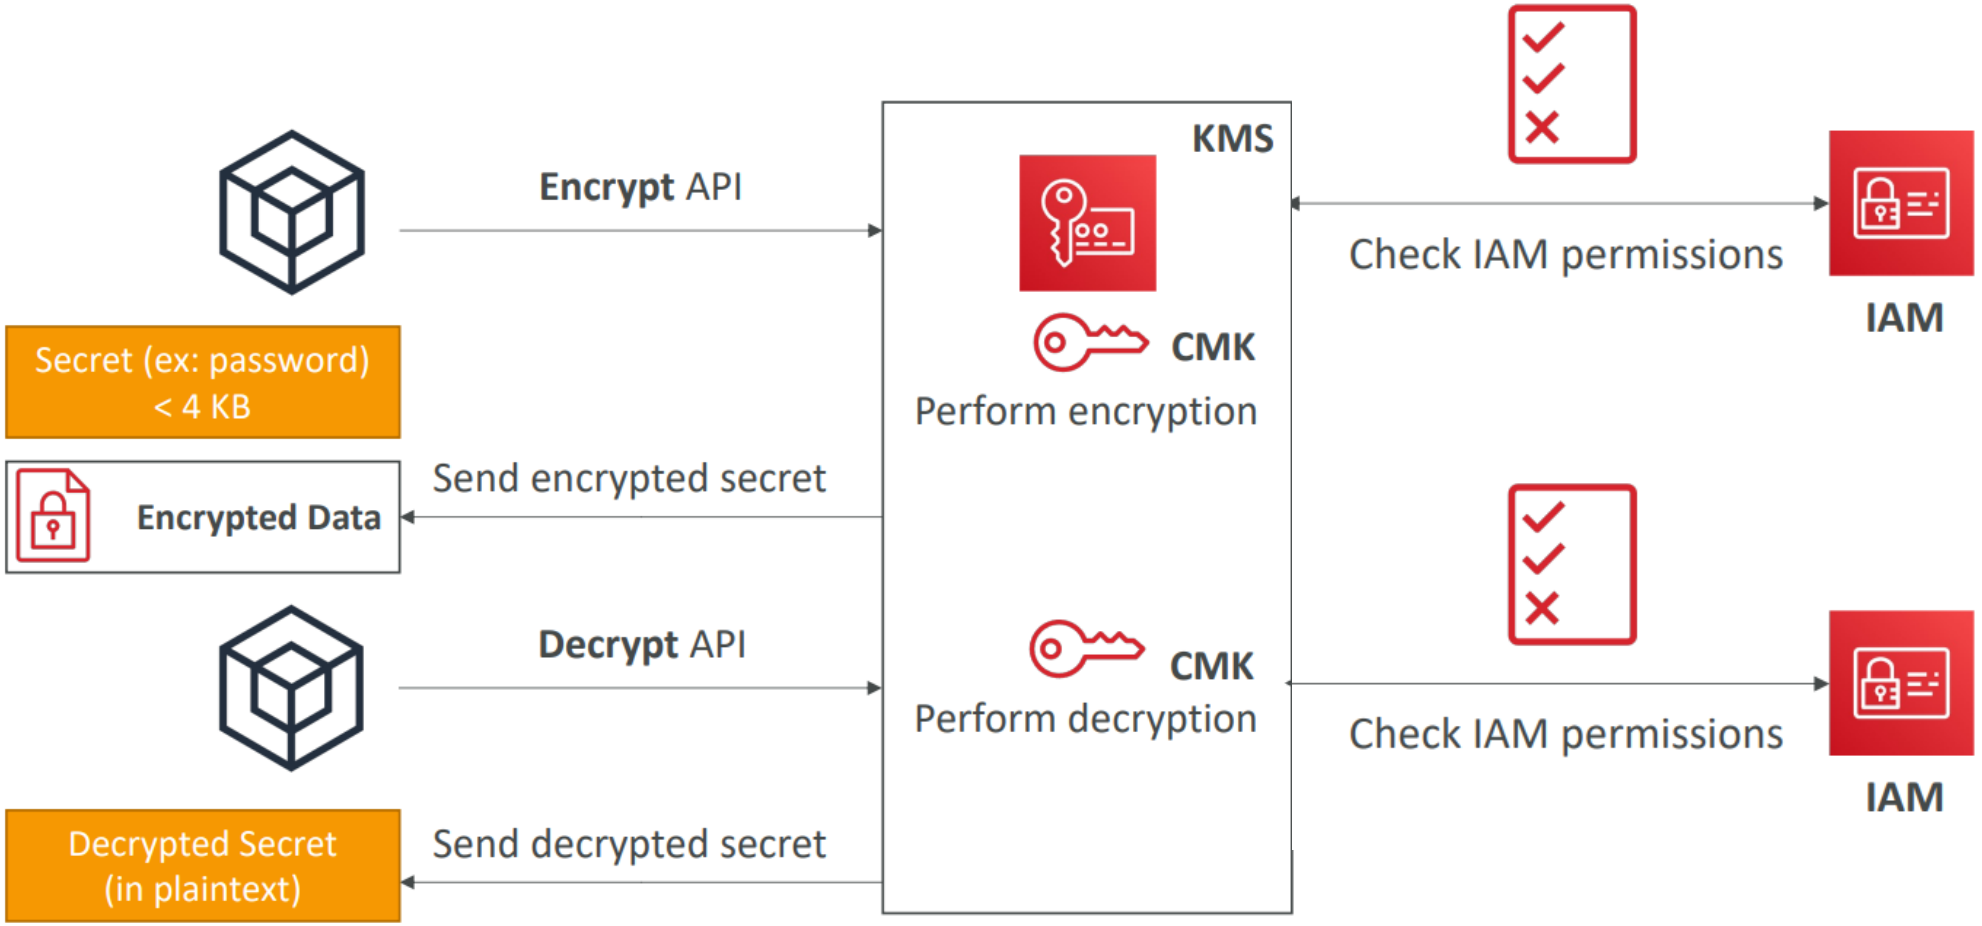

© Copyright by Cloud Mentor Pro | Confidential

Envelope Encryption

- KMS Encrypt API call has a limit of 4 KB
- If you want to encrypt >4 KB, we need to use Envelope Encryption
- The main API that will help us is the **GenerateDataKey** API

- **For the exam: anything over 4 KB of data that needs to be encrypted must use the Envelope Encryption == GenerateDataKey API**

© Copyright by Cloud Mentor Pro | Confidential

Deep dive into Envelope Encryption GenerateDataKey API

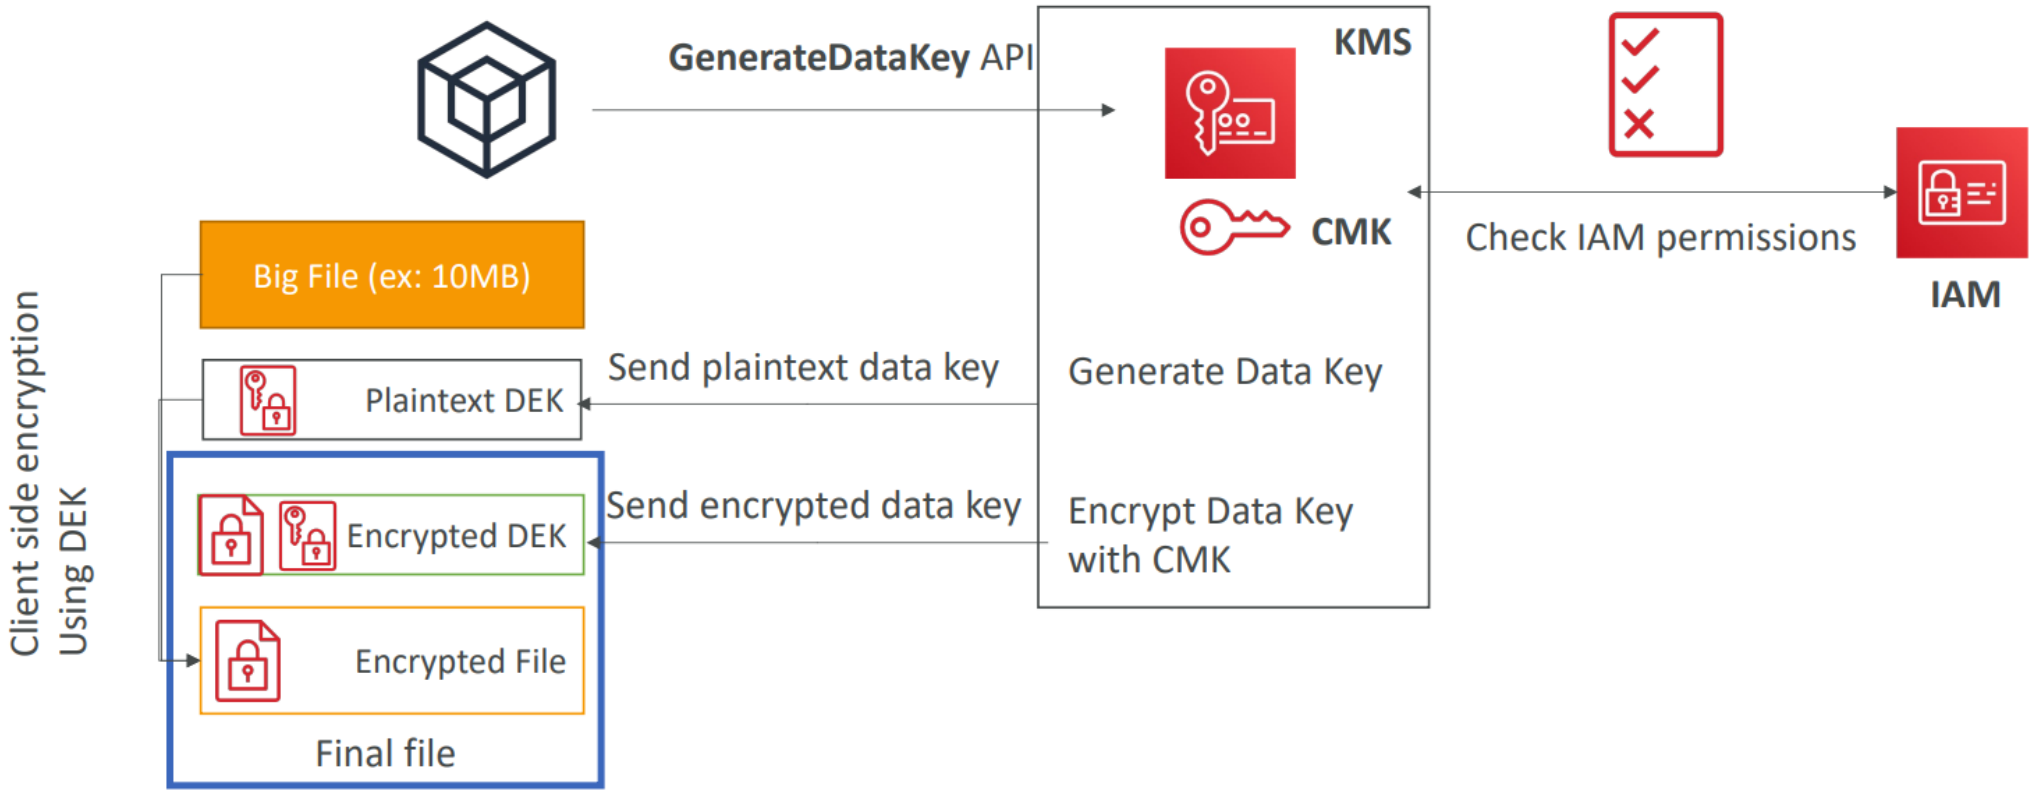

© Copyright by Cloud Mentor Pro | Confidential

CloudHSM

- KMS => AWS manages the software for encryption
- CloudHSM => AWS provisions encryption **hardware**
- Dedicated Hardware (HSM = Hardware Security Module)
- You manage your own encryption keys entirely (not AWS)
- HSM device is tamper resistant, FIPS 140-2 Level 3 compliance
- Supports both symmetric and asymmetric encryption (SSL/TLS keys)
- No free tier available
- Must use the CloudHSM Client Software
- Redshift supports CloudHSM for database encryption and key management
- Good option to use with SSE-C encryption

© Copyright by Cloud Mentor Pro | Confidential

SSM Parameter Store

- Secure storage for configuration and secrets
- Optional Seamless Encryption using KMS
- Serverless, scalable, durable, easy SDK
- Version tracking of configurations / secrets
- Security through IAM
- Notifications with Amazon EventBridge
- Integration with CloudFormation

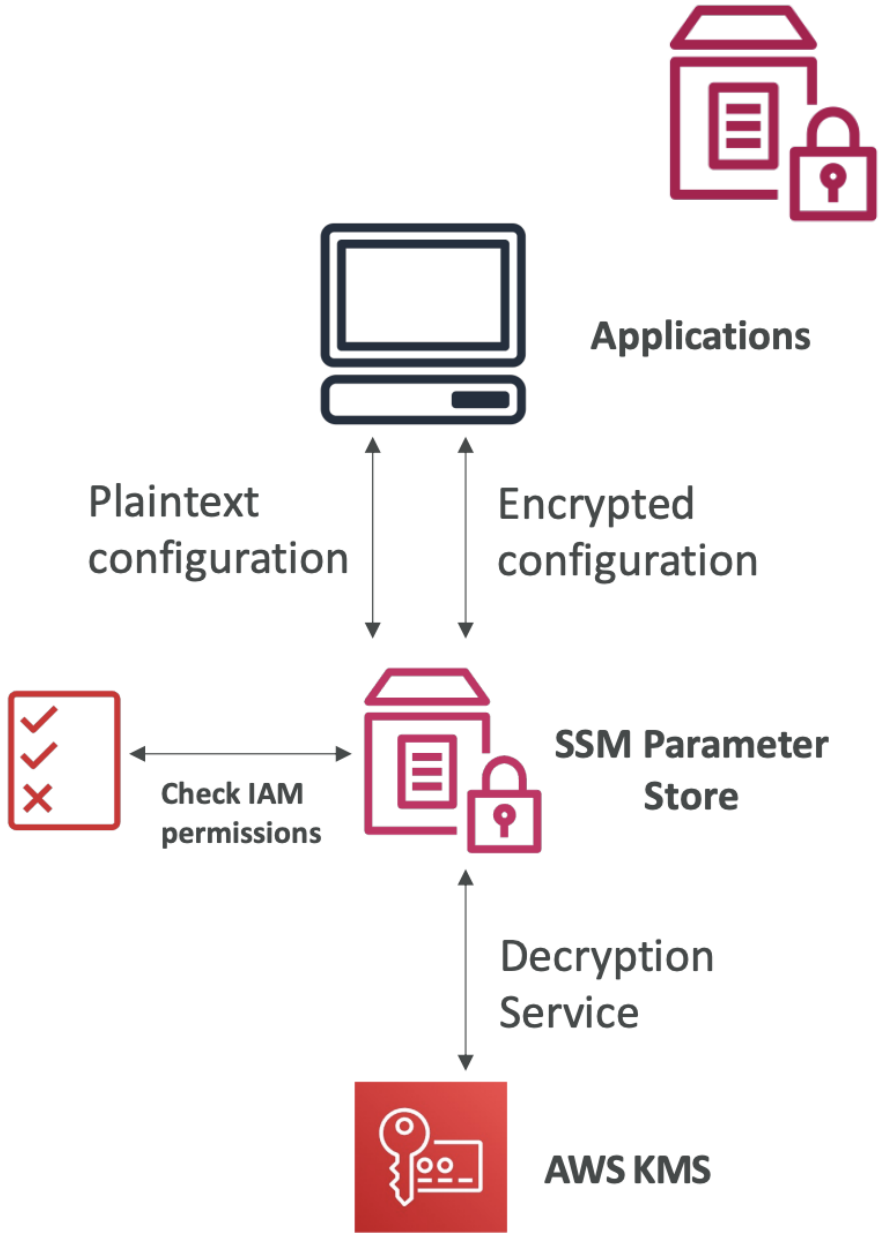

© Copyright by Cloud Mentor Pro | Confidential

SSM Parameter Store Hierarchy

- /my-department/
  - my-app/
    - dev/
      - db-url
      - db-password
    - prod/
      - db-url
      - db-password
  - other-app/
- /other-department/
- /aws/reference/secretsmanager/secret_ID_in_Secrets_Manager
- /aws/service/ami-amazon-linux-latest/amzn2-ami-hvm-x86_64-gp2 (public)

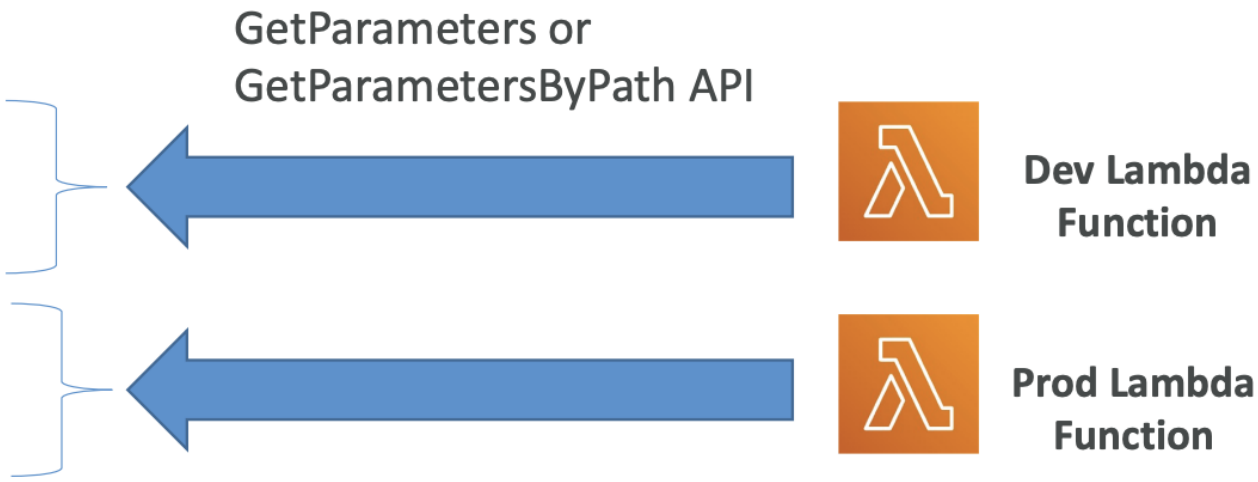

© Copyright by Cloud Mentor Pro | Confidential

Standard and advanced parameter tiers

	Standard	Advanced
Total number of parameters allowed (per AWS account and Region)	10000	100000
Maximum size of a parameter value	4 KB	8 KB
Parameter policies available	No	Yes
Cost	No additional charge	Charges apply
Storage Pricing	Free	$0.05 per advanced parameter per month

© Copyright by Cloud Mentor Pro | Confidential

AWS Secrets Manager

- Newer service, meant for storing secrets
- Capability to force **rotation of secrets** every X days
- Automate generation of secrets on rotation (uses Lambda)
- Integration with **Amazon RDS** (MySQL, PostgreSQL, Aurora)
- Secrets are encrypted using KMS

- Mostly meant for RDS integration

© Copyright by Cloud Mentor Pro | Confidential

AWS Secrets Manager – Multi-Region Secrets

- Replicate Secrets across multiple AWS Regions
- Secrets Manager keeps read replicas in sync with the primary Secret
- Ability to promote a read replica Secret to a standalone Secret
- Use cases: multi-region apps, disaster recovery strategies, multi-region DB…

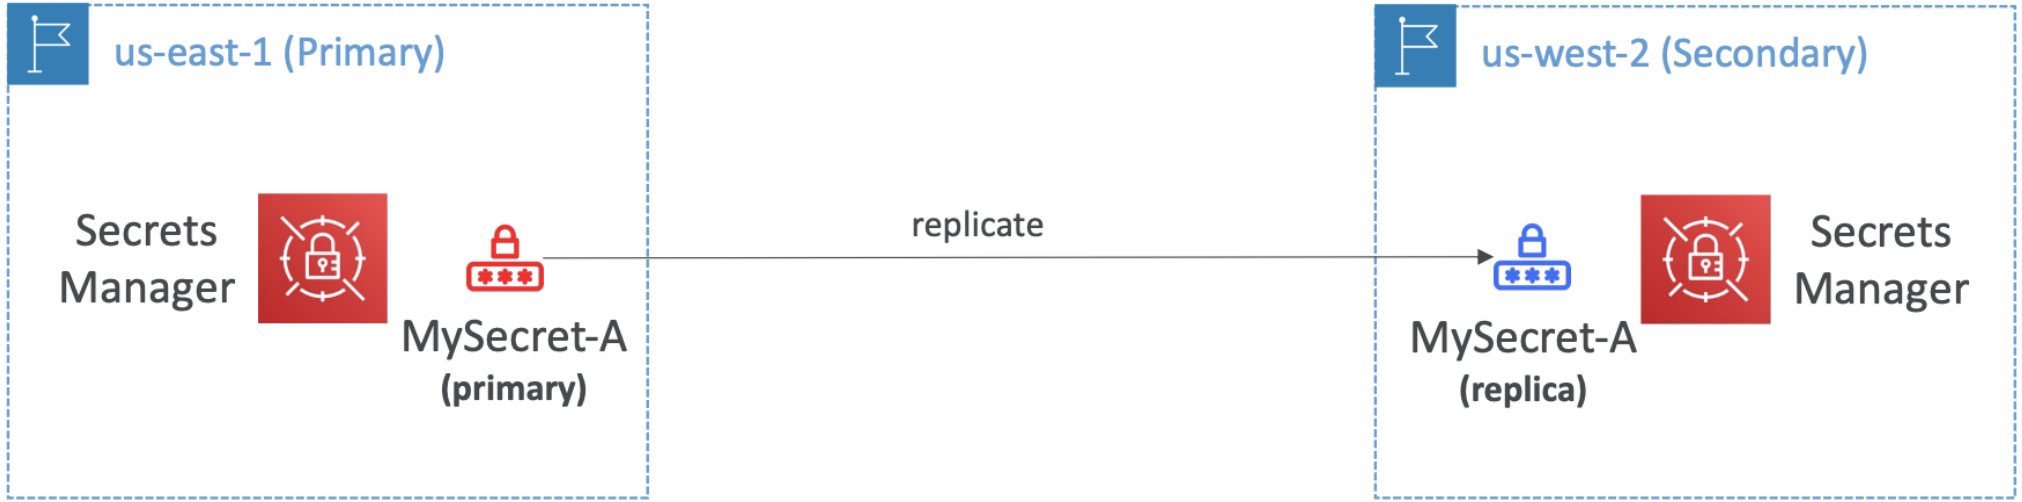

© Copyright by Cloud Mentor Pro | Confidential

SSM Parameter Store vs Secrets Manager

- Secrets Manager ($$$):
  - Automatic rotation of secrets with AWS Lambda
  - Lambda function is provided for RDS, Redshift, DocumentDB
  - KMS encryption is mandatory
  - Can integration with CloudFormation
- SSM Parameter Store ($):
  - Simple API
  - No secret rotation (can enable rotation using Lambda triggered by EventBridge)
  - KMS encryption is optional
  - Can integration with CloudFormation
  - Can pull a Secrets Manager secret using the SSM Parameter Store API

© Copyright by Cloud Mentor Pro | Confidential

SSM Parameter Store vs. Secrets Manager Rotation

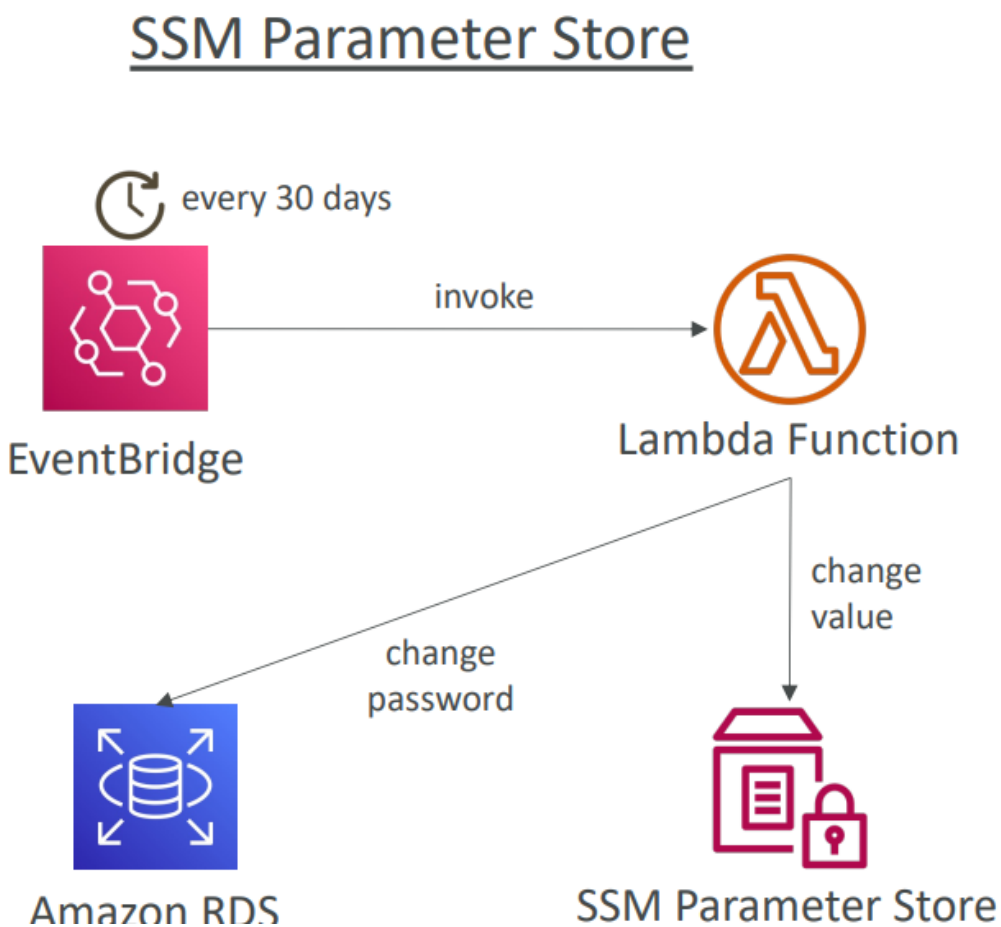

© Copyright by Cloud Mentor Pro | Confidential

Other Services

Quick overview of other services that might have questions on at the exam

© Copyright by Cloud Mentor Pro | Confidential

Amazon Simple Email Service (Amazon SES)

- Fully managed service to send emails securely, globally and at scale
- Allows inbound/outbound emails
- Reputation dashboard, performance insights, anti-spam feedback
- Provides statistics such as email deliveries, bounces, feedback loop results, email open
- Supports DomainKeys Identified Mail (DKIM) and Sender Policy Framework (SPF)
- Flexible IP deployment: shared, dedicated, and customer-owned IPs
- Send emails using your application using AWS Console, APIs, or SMTP
- Use cases: transactional, marketing and bulk email communications

© Copyright by Cloud Mentor Pro | Confidential

Amazon OpenSearch Service

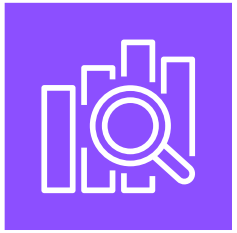

- Amazon OpenSearch is successor to Amazon ElasticSearch
- In DynamoDB, queries only exist by primary key or indexes…
- With OpenSearch, you can search any field, even partially matches
- It’s common to use OpenSearch as a complement to another database
- Two modes: managed cluster or serverless cluster
- Does not natively support SQL (can be enabled via a plugin)
- Ingestion from Kinesis Data Firehose, AWS IoT, and CloudWatch Logs
- Security through Cognito & IAM, KMS encryption, TLS
- Comes with OpenSearch Dashboards (visualization)

© Copyright by Cloud Mentor Pro | Confidential

Amazon Athena

- Serverless query service to analyze data stored in Amazon S3
- Uses standard SQL language to query the files (built on Presto)
- Supports CSV, JSON, ORC, Avro, and Parquet
- Pricing: $5.00 per TB of data scanned
- Commonly used with Amazon Quicksight for reporting/dashboards

- **Use cases:** Business intelligence / analytics / reporting, analyze & query VPC Flow Logs, ELB Logs, **CloudTrail trails**, etc...
- **Exam Tip:** analyze data in S3 using serverless SQL, use Athena

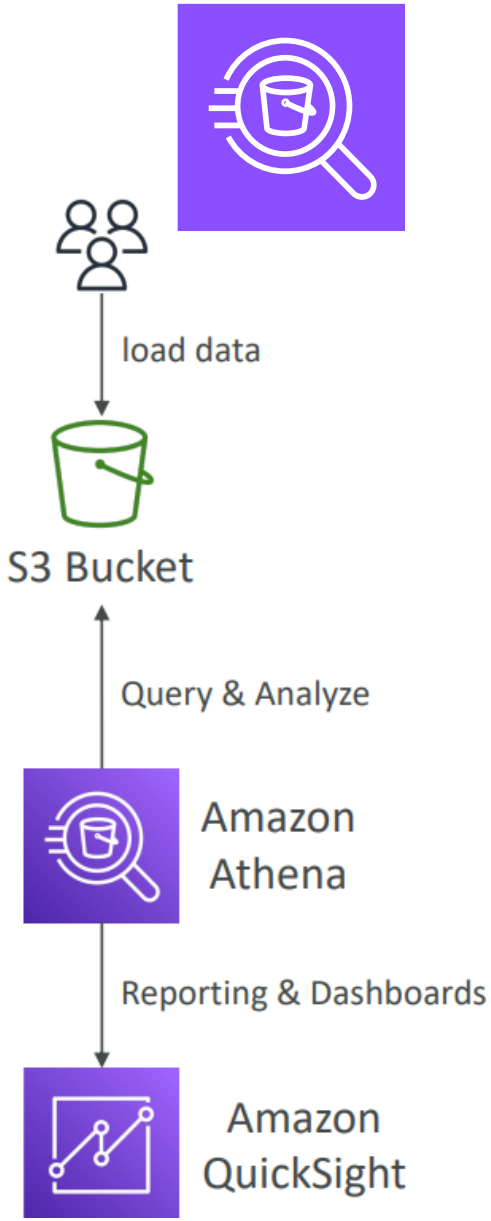

© Copyright by Cloud Mentor Pro | Confidential

Amazon Athena – Federated Query

- Allows you to run SQL queries across data stored in relational, non-relational, object, and custom data sources (AWS or on-premises)
- Uses Data Source Connectors that run on AWS Lambda to run Federated Queries (e.g., CloudWatch Logs, DynamoDB, RDS, …)
- Store the results back in Amazon S3

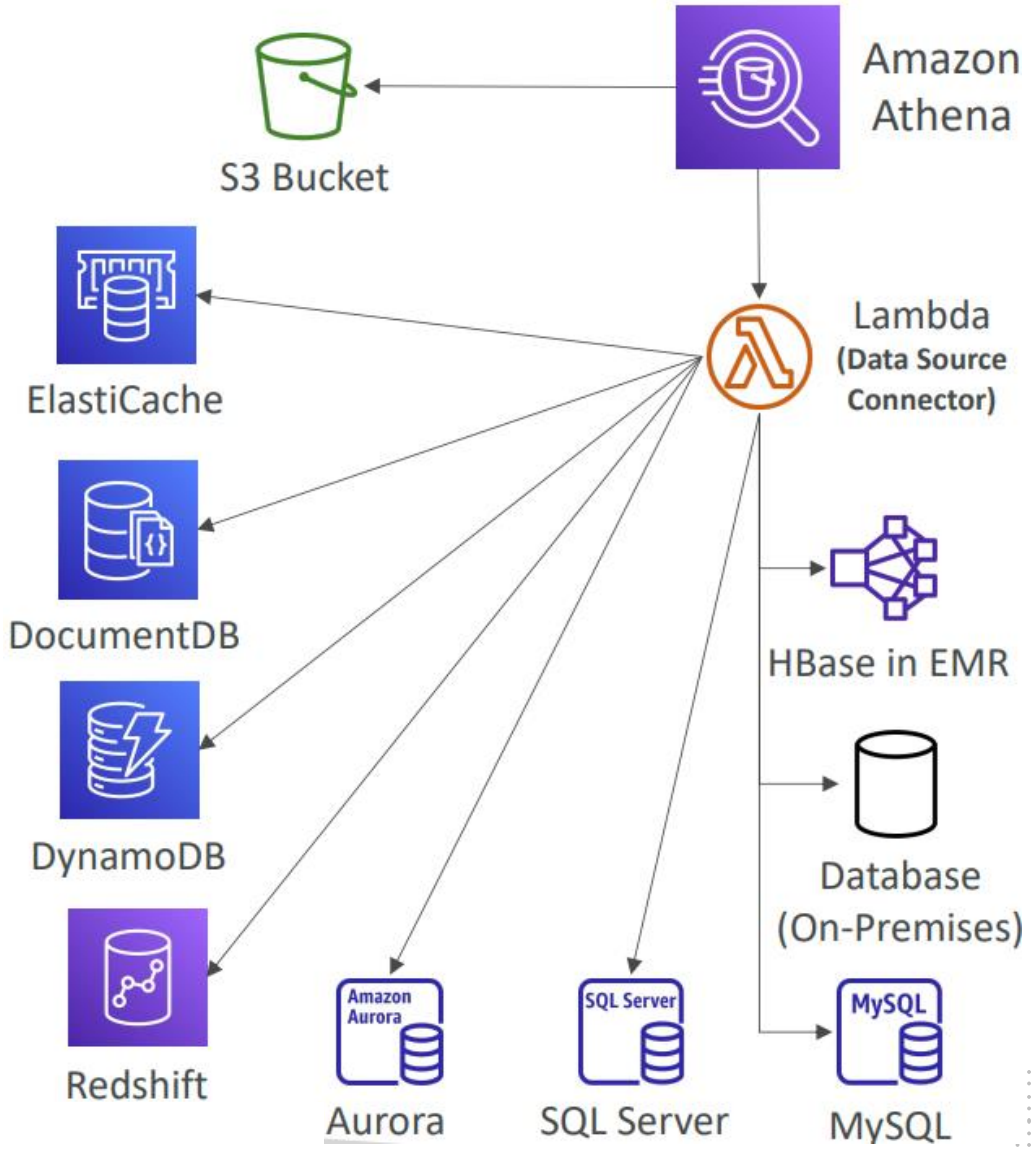

© Copyright by Cloud Mentor Pro | Confidential

AWS Certificate Manager (ACM)

- Let's you easily provision, manage, and deploy SSL/TLS Certificates
- Used to provide in-flight encryption for websites (HTTPS)
- Supports both public and private TLS certificates
- Free of charge for public TLS certificates
- Automatic TLS certificate renewal
- Integrations with (load TLS certificates on)
  - Elastic Load Balancers
  - CloudFront Distributions
  - APIs on API Gateway

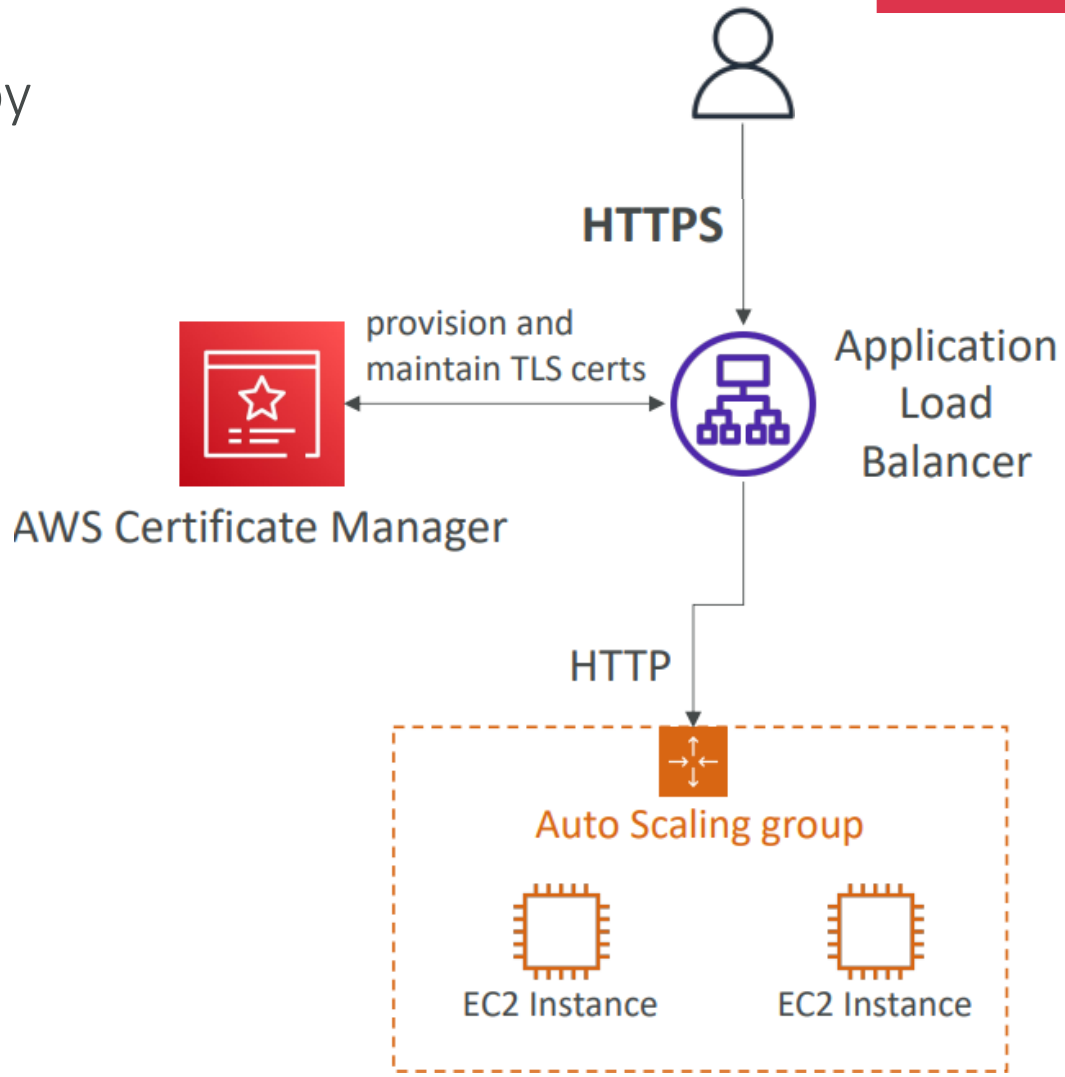

© Copyright by Cloud Mentor Pro | Confidential

AWS Macie

- Amazon Macie is a fully managed data security and data privacy service that uses machine learning and pattern matching to discover and protect your sensitive data in AWS.
- Macie helps identify and alert you to sensitive data, such as personally identifiable information (PII)

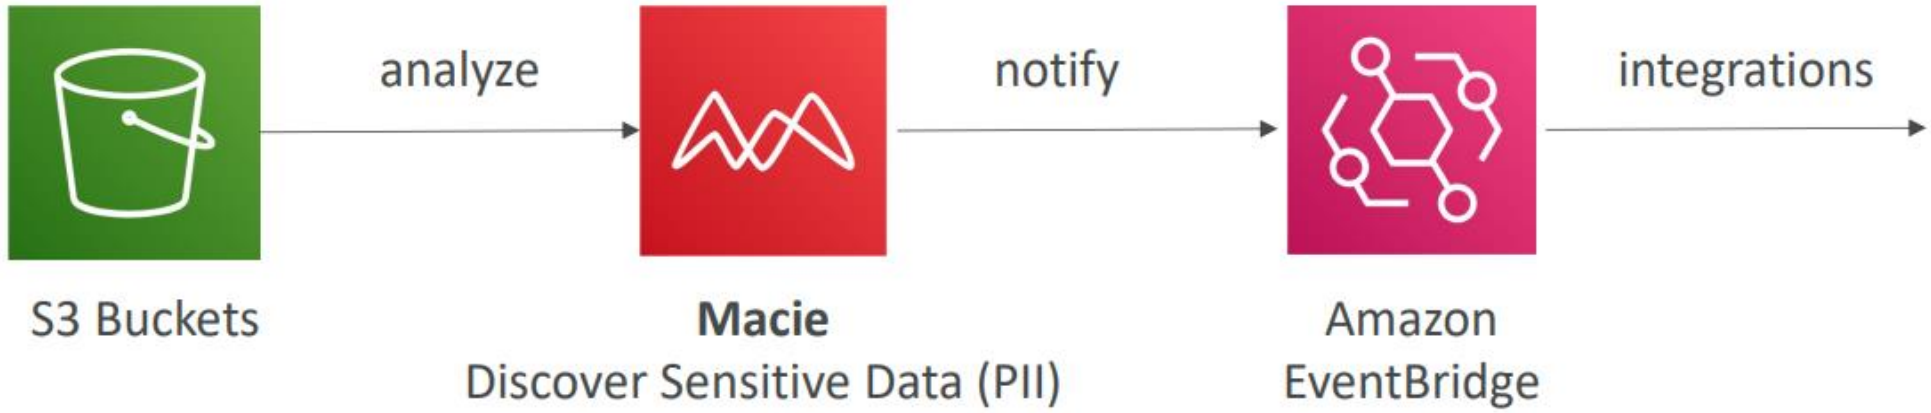

© Copyright by Cloud Mentor Pro | Confidential